
HPX Documentation

master

The STE||AR Group

January 29, 2020

USER DOCUMENTATION

1	What is <i>HPX</i>?	3
2	What's so special about <i>HPX</i>?	5
2.1	Why <i>HPX</i> ?	5
2.2	Quick start	11
2.3	Terminology	17
2.4	Examples	18
2.5	Manual	40
2.6	Additional material	246
2.7	Overview	247
2.8	All modules	247
2.9	API reference	260
2.10	Contributing to <i>HPX</i>	925
2.11	Releases	932
2.12	About <i>HPX</i>	1120
3	Index	1129
Index		1131

If you're new to *HPX* you can get started with the [Quick start](#) guide. Don't forget to read the [Terminology](#) section to learn about the most important concepts in *HPX*. The [Examples](#) give you a feel for how it is to write real *HPX* applications and the [Manual](#) contains detailed information about everything from building *HPX* to debugging it. There are links to blog posts and videos about *HPX* in [Additional material](#).

If you can't find what you're looking for in the documentation, please:

- open an issue on [GitHub](#)¹;
- contact us on [IRC](#), the *HPX* channel on the [C++ Slack](#)², or on our [mailing list](#)³; or
- read or ask questions tagged with *HPX* on [StackOverflow](#)⁴.

¹ <https://github.com/STELLAR-GROUP/hpx/issues>

² <https://cpplang.slack.com>

³ hpx-users@stellar.cct.lsu.edu

⁴ <https://stackoverflow.com/questions/tagged/hpx>

**CHAPTER
ONE**

WHAT IS HPX?

HPX is a C++ Standard Library for Concurrency and Parallelism. It implements all of the corresponding facilities as defined by the C++ Standard. Additionally, in *HPX* we implement functionalities proposed as part of the ongoing C++ standardization process. We also extend the C++ Standard APIs to the distributed case. *HPX* is developed by the STE||AR group (see [People](#)).

The goal of *HPX* is to create a high quality, freely available, open source implementation of a new programming model for conventional systems, such as classic Linux based Beowulf clusters or multi-socket highly parallel SMP nodes. At the same time, we want to have a very modular and well designed runtime system architecture which would allow us to port our implementation onto new computer system architectures. We want to use real-world applications to drive the development of the runtime system, coining out required functionalities and converging onto a stable API which will provide a smooth migration path for developers.

The API exposed by *HPX* is not only modeled after the interfaces defined by the C++11/14/17/20 ISO standard. It also adheres to the programming guidelines used by the Boost collection of C++ libraries. We aim to improve the scalability of today's applications and to expose new levels of parallelism which are necessary to take advantage of the exascale systems of the future.

WHAT'S SO SPECIAL ABOUT HPX?

- HPX exposes a uniform, standards-oriented API for ease of programming parallel and distributed applications.
- It enables programmers to write fully asynchronous code using hundreds of millions of threads.
- HPX provides unified syntax and semantics for local and remote operations.
- HPX makes concurrency manageable with dataflow and future based synchronization.
- It implements a rich set of runtime services supporting a broad range of use cases.
- HPX exposes a uniform, flexible, and extendable performance counter framework which can enable runtime adaptivity
- It is designed to solve problems conventionally considered to be scaling-impaired.
- HPX has been designed and developed for systems of any scale, from hand-held devices to very large scale systems.
- It is the first fully functional implementation of the ParalleX execution model.
- HPX is published under a liberal open-source license and has an open, active, and thriving developer community.

2.1 Why HPX?

Current advances in high performance computing (HPC) continue to suffer from the issues plaguing parallel computation. These issues include, but are not limited to, ease of programming, inability to handle dynamically changing workloads, scalability, and efficient utilization of system resources. Emerging technological trends such as multi-core processors further highlight limitations of existing parallel computation models. To mitigate the aforementioned problems, it is necessary to rethink the approach to parallelization models. ParalleX contains mechanisms such as multi-threading, *parcels*, *global name space* support, percolation and *local control objects (LCO)*. By design, ParalleX overcomes limitations of current models of parallelism by alleviating contention, latency, overhead and starvation. With ParalleX, it is further possible to increase performance by at least an order of magnitude on challenging parallel algorithms, e.g., dynamic directed graph algorithms and adaptive mesh refinement methods for astrophysics. An additional benefit of ParalleX is fine-grained control of power usage, enabling reductions in power consumption.

2.1.1 ParalleX—a new execution model for future architectures

ParalleX is a new parallel execution model that offers an alternative to the conventional computation models, such as message passing. ParalleX distinguishes itself by:

- Split-phase transaction model
- Message-driven
- Distributed shared memory (not cache coherent)
- Multi-threaded
- Futures synchronization
- *Local Control Objects (LCOs)*
- Synchronization for anonymous producer-consumer scenarios
- Percolation (pre-staging of task data)

The ParalleX model is intrinsically latency hiding, delivering an abundance of variable-grained parallelism within a hierarchical namespace environment. The goal of this innovative strategy is to enable future systems delivering very high efficiency, increased scalability and ease of programming. ParalleX can contribute to significant improvements in the design of all levels of computing systems and their usage from application algorithms and their programming languages to system architecture and hardware design together with their supporting compilers and operating system software.

2.1.2 What is HPX?

High Performance ParalleX (*HPX*) is the first runtime system implementation of the ParalleX execution model. The *HPX* runtime software package is a modular, feature-complete, and performance-oriented representation of the ParalleX execution model targeted at conventional parallel computing architectures, such as SMP nodes and commodity clusters. It is academically developed and freely available under an open source license. We provide *HPX* to the community for experimentation and application to achieve high efficiency and scalability for dynamic adaptive and irregular computational problems. *HPX* is a C++ library that supports a set of critical mechanisms for dynamic adaptive resource management and lightweight task scheduling within the context of a global address space. It is solidly based on many years of experience in writing highly parallel applications for HPC systems.

The two-decade success of the communicating sequential processes (CSP) execution model and its message passing interface (MPI) programming model have been seriously eroded by challenges of power, processor core complexity, multi-core sockets, and heterogeneous structures of GPUs. Both efficiency and scalability for some current (strong scaled) applications and future Exascale applications demand new techniques to expose new sources of algorithm parallelism and exploit unused resources through adaptive use of runtime information.

The ParalleX execution model replaces CSP to provide a new computing paradigm embodying the governing principles for organizing and conducting highly efficient scalable computations greatly exceeding the capabilities of today's problems. *HPX* is the first practical, reliable, and performance-oriented runtime system incorporating the principal concepts of the ParalleX model publicly provided in open source release form.

HPX is designed by the STE||AR⁵ Group (Systems Technology, Emergent Parallelism, and Algorithm Research) at Louisiana State University (LSU)⁶'s Center for Computation and Technology (CCT)⁷ to enable developers to exploit the full processing power of many-core systems with an unprecedented degree of parallelism. STE||AR⁸ is a research group focusing on system software solutions and scientific application development for hybrid and many-core hardware architectures.

⁵ <https://stellar-group.org>

⁶ <https://www.lsu.edu>

⁷ <https://www.cct.lsu.edu>

⁸ <https://stellar-group.org>

For more information about the STE||AR⁹ Group, see [People](#).

2.1.3 What makes our systems slow?

Estimates say that we currently run our computers at well below 100% efficiency. The theoretical peak performance (usually measured in **FLOPS**¹⁰—floating point operations per second) is much higher than any practical peak performance reached by any application. This is particularly true for highly parallel hardware. The more hardware parallelism we provide to an application, the better the application must scale in order to efficiently use all the resources of the machine. Roughly speaking, we distinguish two forms of scalability: strong scaling (see [Amdahl's Law](#)¹¹) and weak scaling (see [Gustafson's Law](#)¹²). Strong scaling is defined as how the solution time varies with the number of processors for a fixed **total** problem size. It gives an estimate of how much faster we can solve a particular problem by throwing more resources at it. Weak scaling is defined as how the solution time varies with the number of processors for a fixed problem size **per processor**. In other words, it defines how much more data can we process by using more hardware resources.

In order to utilize as much hardware parallelism as possible an application must exhibit excellent strong and weak scaling characteristics, which requires a high percentage of work executed in parallel, i.e., using multiple threads of execution. Optimally, if you execute an application on a hardware resource with N processors it either runs N times faster or it can handle N times more data. Both cases imply 100% of the work is executed on all available processors in parallel. However, this is just a theoretical limit. Unfortunately, there are more things that limit scalability, mostly inherent to the hardware architectures and the programming models we use. We break these limitations into four fundamental factors that make our systems *SLOW*:

- Starvation occurs when there is insufficient concurrent work available to maintain high utilization of all resources.
- Latencies are imposed by the time-distance delay intrinsic to accessing remote resources and services.
- Overhead is work required for the management of parallel actions and resources on the critical execution path, which is not necessary in a sequential variant.
- Waiting for contention resolution is the delay due to the lack of availability of oversubscribed shared resources.

Each of those four factors manifests itself in multiple and different ways; each of the hardware architectures and programming models expose specific forms. However, the interesting part is that all of them are limiting the scalability of applications no matter what part of the hardware jungle we look at. Hand-helds, PCs, supercomputers, or the cloud, all suffer from the reign of the 4 horsemen: **Starvation**, **Latency**, **Overhead**, and **Contention**. This realization is very important as it allows us to derive the criteria for solutions to the scalability problem from first principles, and it allows us to focus our analysis on very concrete patterns and measurable metrics. Moreover, any derived results will be applicable to a wide variety of targets.

2.1.4 Technology demands new response

Today's computer systems are designed based on the initial ideas of [John von Neumann](#)¹³, as published back in 1945, and later extended by the [Harvard architecture](#)¹⁴. These ideas form the foundation, the execution model, of computer systems we use currently. However, a new response is required in the light of the demands created by today's technology.

So, what are the overarching objectives for designing systems allowing for applications to scale as they should? In our opinion, the main objectives are:

⁹ <https://stellar-group.org>

¹⁰ <http://en.wikipedia.org/wiki/FLOPS>

¹¹ http://en.wikipedia.org/wiki/Amdahl%27s_law

¹² http://en.wikipedia.org/wiki/Gustafson%27s_law

¹³ <http://qss.stanford.edu/~godfrey/vonNeumann/vnedvac.pdf>

¹⁴ http://en.wikipedia.org/wiki/Harvard_architecture

- Performance: as previously mentioned, scalability and efficiency are the main criteria people are interested in.
- Fault tolerance: the low expected mean time between failures ([MTBF¹⁵](#)) of future systems requires embracing faults, not trying to avoid them.
- Power: minimizing energy consumption is a must as it is one of the major cost factors today, and will continue to rise in the future.
- Generality: any system should be usable for a broad set of use cases.
- Programmability: for programmer this is a very important objective, ensuring long term platform stability and portability.

What needs to be done to meet those objectives, to make applications scale better on tomorrow's architectures? Well, the answer is almost obvious: we need to devise a new execution model—a set of governing principles for the holistic design of future systems—targeted at minimizing the effect of the outlined **SLOW** factors. Everything we create for future systems, every design decision we make, every criteria we apply, have to be validated against this single, uniform metric. This includes changes in the hardware architecture we prevalently use today, and it certainly involves new ways of writing software, starting from the operating system, runtime system, compilers, and at the application level. However, the key point is that all those layers have to be co-designed; they are interdependent and cannot be seen as separate facets. The systems we have today have been evolving for over 50 years now. All layers function in a certain way, relying on the other layers to do so. But we do not have the time to wait another 50 years for a new coherent system to evolve. The new paradigms are needed now—therefore, co-design is the key.

2.1.5 Governing principles applied while developing HPX

As it turns out, we do not have to start from scratch. Not everything has to be invented and designed anew. Many of the ideas needed to combat the 4 horsemen already exist, many for more than 30 years. All it takes is to gather them into a coherent approach. We'll highlight some of the derived principles we think to be crucial for defeating **SLOW**. Some of those are focused on high-performance computing, others are more general.

2.1.6 Focus on latency hiding instead of latency avoidance

It is impossible to design a system exposing zero latencies. In an effort to come as close as possible to this goal many optimizations are mainly targeted towards minimizing latencies. Examples for this can be seen everywhere, such as low latency network technologies like [InfiniBand¹⁶](#), caching memory hierarchies in all modern processors, the constant optimization of existing [MPI¹⁷](#) implementations to reduce related latencies, or the data transfer latencies intrinsic to the way we use [GPGPUs¹⁸](#) today. It is important to note that existing latencies are often tightly related to some resource having to wait for the operation to be completed. At the same time it would be perfectly fine to do some other, unrelated work in the meantime, allowing the system to hide the latencies by filling the idle-time with useful work. Modern systems already employ similar techniques (pipelined instruction execution in the processor cores, asynchronous input/output operations, and many more). What we propose is to go beyond anything we know today and to make latency hiding an intrinsic concept of the operation of the whole system stack.

¹⁵ http://en.wikipedia.org/wiki/Mean_time_between_failures

¹⁶ <http://en.wikipedia.org/wiki/InfiniBand>

¹⁷ https://en.wikipedia.org/wiki/Message_Passing_Interface

¹⁸ <http://en.wikipedia.org/wiki/GPGPU>

2.1.7 Embrace fine-grained parallelism instead of heavyweight threads

If we plan to hide latencies even for very short operations, such as fetching the contents of a memory cell from main memory (if it is not already cached), we need to have very lightweight threads with extremely short context switching times, optimally executable within one cycle. Granted, for mainstream architectures, this is not possible today (even if we already have special machines supporting this mode of operation, such as the [Cray XMT](#)¹⁹). For conventional systems, however, the smaller the overhead of a context switch and the finer the granularity of the threading system, the better will be the overall system utilization and its efficiency. For today's architectures we already see a flurry of libraries providing exactly this type of functionality: non-pre-emptive, task-queue based parallelization solutions, such as [Intel Threading Building Blocks \(TBB\)](#)²⁰, [Microsoft Parallel Patterns Library \(PPL\)](#)²¹, [Cilk++](#)²², and many others. The possibility to suspend a current task if some preconditions for its execution are not met (such as waiting for I/O or the result of a different task), seamlessly switching to any other task which can continue, and to reschedule the initial task after the required result has been calculated, which makes the implementation of latency hiding almost trivial.

2.1.8 Rediscover constraint-based synchronization to replace global barriers

The code we write today is riddled with implicit (and explicit) global barriers. By “global barriers,” we mean the synchronization of the control flow between several (very often all) threads (when using [OpenMP](#)²³) or processes ([MPI](#)²⁴). For instance, an implicit global barrier is inserted after each loop parallelized using [OpenMP](#)²⁵ as the system synchronizes the threads used to execute the different iterations in parallel. In [MPI](#)²⁶ each of the communication steps imposes an explicit barrier onto the execution flow as (often all) nodes have to be synchronized. Each of those barriers is like the eye of a needle the overall execution is forced to be squeezed through. Even minimal fluctuations in the execution times of the parallel threads (jobs) causes them to wait. Additionally, it is often only one of the executing threads that performs the actual reduce operation, which further impedes parallelism. A closer analysis of a couple of key algorithms used in science applications reveals that these global barriers are not always necessary. In many cases it is sufficient to synchronize a small subset of the threads. Any operation should proceed whenever the preconditions for its execution are met, and only those. Usually there is no need to wait for iterations of a loop to finish before you can continue calculating other things; all you need is to complete the iterations that produce the required results for the next operation. Good bye global barriers, hello constraint based synchronization! People have been trying to build this type of computing (and even computers) since the 1970s. The theory behind what they did is based on ideas around static and dynamic dataflow. There are certain attempts today to get back to those ideas and to incorporate them with modern architectures. For instance, a lot of work is being done in the area of constructing dataflow-oriented execution trees. Our results show that employing dataflow techniques in combination with the other ideas, as outlined herein, considerably improves scalability for many problems.

2.1.9 Adaptive locality control instead of static data distribution

While this principle seems to be a given for single desktop or laptop computers (the operating system is your friend), it is everything but ubiquitous on modern supercomputers, which are usually built from a large number of separate nodes (i.e., Beowulf clusters), tightly interconnected by a high-bandwidth, low-latency network. Today's prevalent programming model for those is MPI, which does not directly help with proper data distribution, leaving it to the programmer to decompose the data to all of the nodes the application is running on. There are a couple of specialized languages and programming environments based on [PGAS](#)²⁷ (Partitioned Global Address Space) designed to over-

¹⁹ http://en.wikipedia.org/wiki/Cray_XMT

²⁰ <https://www.threadingbuildingblocks.org/>

²¹ <https://msdn.microsoft.com/en-us/library/dd492418.aspx>

²² <https://software.intel.com/en-us/articles/intel-cilk-plus/>

²³ <https://openmp.org/wp/>

²⁴ https://en.wikipedia.org/wiki/Message_Passing_Interface

²⁵ <https://openmp.org/wp/>

²⁶ https://en.wikipedia.org/wiki/Message_Passing_Interface

²⁷ <https://www.pgas.org/>

come this limitation, such as Chapel²⁸, X10²⁹, UPC³⁰, or Fortress³¹. However, all systems based on PGAS rely on static data distribution. This works fine as long as this static data distribution does not result in heterogeneous workload distributions or other resource utilization imbalances. In a distributed system these imbalances can be mitigated by migrating part of the application data to different localities (nodes). The only framework supporting (limited) migration today is Charm++³². The first attempts towards solving related problem go back decades as well, a good example is the Linda coordination language³³. Nevertheless, none of the other mentioned systems support data migration today, which forces the users to either rely on static data distribution and live with the related performance hits or to implement everything themselves, which is very tedious and difficult. We believe that the only viable way to flexibly support dynamic and adaptive *locality* control is to provide a global, uniform address space to the applications, even on distributed systems.

2.1.10 Prefer moving work to the data over moving data to the work

For the best performance it seems obvious to minimize the amount of bytes transferred from one part of the system to another. This is true on all levels. At the lowest level we try to take advantage of processor memory caches, thus, minimizing memory latencies. Similarly, we try to amortize the data transfer time to and from GPGPUs³⁴ as much as possible. At high levels we try to minimize data transfer between different nodes of a cluster or between different virtual machines on the cloud. Our experience (well, it's almost common wisdom) shows that the amount of bytes necessary to encode a certain operation is very often much smaller than the amount of bytes encoding the data the operation is performed upon. Nevertheless, we still often transfer the data to a particular place where we execute the operation just to bring the data back to where it came from afterwards. As an example let's look at the way we usually write our applications for clusters using MPI. This programming model is all about data transfer between nodes. MPI is the prevalent programming model for clusters, and it is fairly straightforward to understand and to use. Therefore, we often write applications in a way that accommodates this model, centered around data transfer. These applications usually work well for smaller problem sizes and for regular data structures. The larger the amount of data we have to churn and the more irregular the problem domain becomes, the worse the overall machine utilization and the (strong) scaling characteristics become. While it is not impossible to implement more dynamic, data driven, and asynchronous applications using MPI, it is somewhat difficult to do so. At the same time, if we look at applications that prefer to execute the code close to the *locality* where the data was placed, i.e., utilizing active messages (for instance based on Charm++³⁵), we see better asynchrony, simpler application codes, and improved scaling.

2.1.11 Favor message driven computation over message passing

Today's prevalently used programming model on parallel (multi-node) systems is MPI. It is based on message passing, as the name implies, which means that the receiver has to be aware of a message about to come in. Both codes, the sender and the receiver, have to synchronize in order to perform the communication step. Even the newer, asynchronous interfaces require explicitly coding the algorithms around the required communication scheme. As a result, everything but the most trivial MPI applications spends a considerable amount of time waiting for incoming messages, thus, causing starvation and latencies to impede full resource utilization. The more complex and more dynamic the data structures and algorithms become, the larger the adverse effects. The community discovered message-driven and data-driven methods of implementing algorithms a long time ago, and systems such as Charm++³⁶ have already integrated active messages demonstrating the validity of the concept. Message-driven computation allows for sending messages without requiring the receiver to actively wait for them. Any incoming message is handled asynchronously and triggers

²⁸ <https://chapel.cray.com/>

²⁹ <https://x10-lang.org/>

³⁰ <https://upc.lbl.gov/>

³¹ <https://labs.oracle.com/projects/plrg/Publications/index.html>

³² <https://charm.cs.uiuc.edu/>

³³ [http://en.wikipedia.org/wiki/Linda_\(coordination_language\)](http://en.wikipedia.org/wiki/Linda_(coordination_language))

³⁴ <http://en.wikipedia.org/wiki/GPGPU>

³⁵ <https://charm.cs.uiuc.edu/>

³⁶ <https://charm.cs.uiuc.edu/>

the encoded action by passing along arguments and—possibly—continuations. *HPX* combines this scheme with work-queue based scheduling as described above, which allows the system to almost completely overlap any communication with useful work, thereby minimizing latencies.

2.2 Quick start

This section is intended to get you to the point of running a basic *HPX* program as quickly as possible. To that end we skip many details but instead give you hints and links to more details along the way.

We assume that you are on a Unix system with access to reasonably recent packages. You should have `cmake` and `make` available for the build system (`pkg-config` is also supported, see [Using HPX with `pkg-config`](#)).

2.2.1 Getting HPX

Download a tarball of the latest release from [HPX Downloads](#)³⁷ and unpack it or clone the repository directly using `git`:

```
git clone https://github.com/STELLAR-GROUP/hpx.git
```

It is also recommended that you check out the latest stable tag:

```
git checkout 1.4.0
```

2.2.2 HPX dependencies

The minimum dependencies needed to use *HPX* are [Boost](#)³⁸ and [Portable Hardware Locality \(HWLOC\)](#)³⁹. If these are not available through your system package manager, see [Installing Boost](#) and [Installing Hwloc](#) for instructions on how to build them yourself. In addition to [Boost](#)⁴⁰ and [Portable Hardware Locality \(HWLOC\)](#)⁴¹, it is recommended that you don't use the system allocator, but instead use either `tcmalloc` from [google-perftools](#)⁴² (default) or `jemalloc`⁴³ for better performance. If you would like to try *HPX* without a custom allocator at this point you can configure *HPX* to use the system allocator in the next step.

A full list of required and optional dependencies, including recommended versions is available at [Prerequisites](#).

2.2.3 Building HPX

Once you have the source code and the dependencies, set up a separate build directory and configure the project. Assuming all your dependencies are in paths known to CMake, the following gets you started:

```
# In the HPX source directory
mkdir build && cd build
cmake -DCMAKE_INSTALL_PREFIX=/install/path ..
make install
```

³⁷ <https://stellar-group.org/downloads/>

³⁸ <https://www.boost.org/>

³⁹ <https://www.open-mpi.org/projects/hwloc/>

⁴⁰ <https://www.boost.org/>

⁴¹ <https://www.open-mpi.org/projects/hwloc/>

⁴² <https://code.google.com/p/gperftools>

⁴³ <http://jemalloc.net>

This will build the core *HPX* libraries and examples, and install them to your chosen location. If you want to install *HPX* to system folders simply leave out the `CMAKE_INSTALL_PREFIX` option. This may take a while. To speed up the process launch more jobs by passing the `-jN` option to `make`.

Tip: Do not set only `-j` (i.e. `-j` without an explicit number of jobs) unless you have a lot of memory available on your machine.

Tip: If you want to change `CMake`⁴⁴ variables for your build it is usually a good idea to start with a clean build directory to avoid configuration problems. It is especially important that you use a clean build directory when changing between `Release` and `Debug` modes.

If your dependencies are in custom locations you may need to tell `CMake`⁴⁵ where to find them by passing one or more of the following options to `CMake`⁴⁶:

```
-DBOOST_ROOT=/path/to/boost  
-DHWLOC_ROOT=/path/to/hwloc  
-DTCMALLOC_ROOT=/path/to/tcmalloc  
-DJEMALLOC_ROOT=/path/to/jemalloc
```

If you want to try *HPX* without using a custom allocator pass `-DHFX_WITH_MALLOC=system` to `CMake`⁴⁷.

Important: If you are building *HPX* for a system with more than 64 processing units you must change the `CMake` variables `HPX_WITH_MORE_THAN_64_THREADS` (to `On`) and `HPX_WITH_MAX_CPU_COUNT` (to a value at least as big as the number of (virtual) cores on your system).

To build the tests run `make tests`. To run the tests run either `make test` or use `ctest` for more control over which tests to run. You can run single tests for example with `ctest --output-on-failure -R tests.unit.parallel.algorithms.for_loop` or a whole group of tests with `ctest --output-on-failure -R tests.unit`.

If you did not run `make install` earlier do so now or build the `hello_world_1` example by running:

```
make hello_world_1
```

HPX executables end up in the `bin` directory in your build directory. You can now run `hello_world_1` and should see the following output:

```
./bin/hello_world_1  
Hello World!
```

You've just run an example which prints `Hello World!` from the *HPX* runtime. The source for the example is in `examples/quickstart/hello_world_1.cpp`. The `hello_world_distributed` example (also available in the `examples/quickstart` directory) is a distributed hello world program which is described in *Remote execution with actions: Hello world*. It provides a gentle introduction to the distributed aspects of *HPX*.

Tip: Most build targets in *HPX* have two names: a simple name and a hierarchical name corresponding to what type of example or test the target is. If you are developing *HPX* it is often helpful to run `make help` to get a list of available targets. For example, `make help | grep hello_world` outputs the following:

⁴⁴ <https://www.cmake.org>

⁴⁵ <https://www.cmake.org>

⁴⁶ <https://www.cmake.org>

⁴⁷ <https://www.cmake.org>

```
... examples.quickstart.hello_world_2
... hello_world_2
... examples.quickstart.hello_world_1
... hello_world_1
... examples.quickstart.hello_world_distributed
... hello_world_distributed
```

It is also possible to build e.g. all quickstart examples using `make examples.quickstart`.

2.2.4 Hello, World!

The following `CMakeLists.txt` is a minimal example of what you need in order to build an executable using `CMake`⁴⁸ and `HPX`:

```
cmake_minimum_required(VERSION 3.3.2)
project(my_hpx_project CXX)
find_package(HPX REQUIRED)
add_hpx_executable(my_hpx_program
    SOURCES main.cpp
    COMPONENT_DEPENDENCIES iostreams)
```

Note: You will most likely have more than one `main.cpp` file in your project. See the section on [Using HPX with CMake-based projects](#) for more details on how to use `add_hpx_executable`.

Note: `COMPONENT_DEPENDENCIES iostreams` is optional for a minimal project but lets us use the `HPX` equivalent of `std::cout`, i.e. the `HPX` [The HPX I/O-streams component](#) functionality in our application.

Create a new project directory and a `CMakeLists.txt` with the contents above. Also create a `main.cpp` with the contents below.

```
// Including 'hpx/hpx_main.hpp' instead of the usual 'hpx/hpx_init.hpp' enables
// to use the plain C-main below as the direct main HPX entry point.
#include <hpx/hpx_main.hpp>
#include <hpx/include/iostreams.hpp>

int main()
{
    // Say hello to the world!
    hpx::cout << "Hello World!\n" << hpx::flush;
    return 0;
}
```

Then, in your project directory run the following:

```
mkdir build && cd build
cmake -DCMAKE_PREFIX_PATH=/path/to/hpx/installation ..
make all
./my_hpx_program
```

⁴⁸ <https://www.cmake.org>

The program looks almost like a regular C++ hello world with the exception of the two includes and `hpx::cout`. When you include `hpx_main.hpp` some things will be done behind the scenes to make sure that `main` actually gets launched on the *HPX* runtime. So while it looks almost the same you can now use futures, `async`, parallel algorithms and more which make use of the *HPX* runtime with lightweight threads. `hpx::cout` is a replacement for `std::cout` to make sure printing never blocks a lightweight thread. You can read more about `hpx::cout` in [The HPX I/O-streams component](#). If you rebuild and run your program now you should see the familiar Hello World!:

```
./my_hpx_program
Hello World!
```

Note: You do not have to let *HPX* take over your main function like in the example. You can instead keep your normal `main` function, and define a separate `hpx_main` function which acts as the entry point to the *HPX* runtime. In that case you start the *HPX* runtime explicitly by calling `hpx::init`:

```
// Copyright (c) 2007-2012 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

///////////////////////////////
// The purpose of this example is to initialize the HPX runtime explicitly and
// execute a HPX-thread printing "Hello World!" once. That's all.

//[hello_world_2_getting_started
#include <hpx/hpx_init.hpp>
#include <hpx/include/iostreams.hpp>

int hpx_main(int, char**)
{
    // Say hello to the world!
    hpx::cout << "Hello World!\n" << hpx::flush;
    return hpx::finalize();
}

int main(int argc, char* argv[])
{
    return hpx::init(argc, argv);
}
//]
```

You can also use `hpx::start` and `hpx::stop` for a non-blocking alternative, or use `hpx::resume` and `hpx::suspend` if you need to combine *HPX* with other runtimes.

See [Starting the HPX runtime](#) for more details on how to initialize and run the *HPX* runtime.

Caution: When including `hpx_main.hpp` the user-defined `main` gets renamed and the real `main` function is defined by *HPX*. This means that the user-defined `main` must include a `return` statement, unlike the real `main`. If you do not include the `return` statement you may end up with confusing compile time errors mentioning `user_main` or even runtime errors.

2.2.5 Writing task-based applications

So far we haven't done anything that can't be done using the C++ standard library. In this section we will give a short overview of what you can do with *HPX* on a single node. The essence is to avoid global synchronization and break up your application into small, composable tasks whose dependencies control the flow of your application. Remember, however, that *HPX* allows you to write distributed applications similarly to how you would write applications for a single node (see [Why HPX?](#) and [Writing distributed HPX applications](#)).

If you are already familiar with `async` and `futures` from the C++ standard library, the same functionality is available in *HPX*.

The following terminology is essential when talking about task-based C++ programs:

- **lightweight thread:** Essential for good performance with task-based programs. Lightweight refers to smaller stacks and faster context switching compared to OS-threads. Smaller overheads allow the program to be broken up into smaller tasks, which in turns helps the runtime fully utilize all processing units.
- **async:** The most basic way of launching tasks asynchronously. Returns a `future<T>`.
- **`future<T>`:** Represents a value of type T that will be ready in the future. The value can be retrieved with `get` (blocking) and one can check if the value is ready with `is_ready` (non-blocking).
- **`shared_future<T>`:** Same as `future<T>` but can be copied (similar to `std::unique_ptr` vs `std::shared_ptr`).
- **continuation:** A function that is to be run after a previous task has run (represented by a `future`). `then` is a method of `future<T>` that takes a function to run next. Used to build up dataflow DAGs (directed acyclic graphs). `shared_futures` help you split up nodes in the DAG and functions like `when_all` help you join nodes in the DAG.

The following example is a collection of the most commonly used functionality in *HPX*:

```
#include <hpx/hpx_main.hpp>
#include <hpx/include/iostreams.hpp>
#include <hpx/include/lcos.hpp>
#include <hpx/include/parallel_generate.hpp>
#include <hpx/include/parallel_sort.hpp>

#include <random>
#include <vector>

void final_task(hpx::future<hpx::util::tuple<hpx::future<double>, hpx::future<void>>>)
{
    hpx::cout << "in final_task" << hpx::endl;
}

// Avoid ABI incompatibilities between C++11/C++17 as std::rand has exception
// specification in libstdc++.
int rand_wrapper()
{
    return std::rand();
}

int main(int, char**)
{
    // A function can be launched asynchronously. The program will not block
    // here until the result is available.
    hpx::future<int> f = hpx::async([]() { return 42; });
    hpx::cout << "Just launched a task!" << hpx::endl;
}
```

(continues on next page)

(continued from previous page)

```

// Use get to retrieve the value from the future. This will block this task
// until the future is ready, but the HPX runtime will schedule other tasks
// if there are tasks available.
hpx::cout << "f contains " << f.get() << hpx::endl;

// Let's launch another task.
hpx::future<double> g = hpx::async([]() { return 3.14; });

// Tasks can be chained using the then method. The continuation takes the
// future as an argument.
hpx::future<double> result = g.then([](hpx::future<double>&& gg)
{
    // This function will be called once g is ready. gg is g moved
    // into the continuation.
    return gg.get() * 42.0 * 42.0;
});

// You can check if a future is ready with the is_ready method.
hpx::cout << "Result is ready? " << result.is_ready() << hpx::endl;

// You can launch other work in the meantime. Let's sort a vector.
std::vector<int> v(1000000);

// We fill the vector synchronously and sequentially.
hpx::parallel::generate(hpx::parallel::execution::seq,
                      std::begin(v), std::end(v), &rand_wrapper);

// We can launch the sort in parallel and asynchronously.
hpx::future<void> done_sorting =
    hpx::parallel::sort(
        hpx::parallel::execution::par( // In parallel.
            hpx::parallel::execution::task), // Asynchronously.
        std::begin(v),
        std::end(v));

// We launch the final task when the vector has been sorted and result is
// ready using when_all.
auto all = hpx::when_all(result, done_sorting).then(&final_task);

// We can wait for all to be ready.
all.wait();

// all must be ready at this point because we waited for it to be ready.
hpx::cout <<
    (all.is_ready() ? "all is ready!" : "all is not ready...") << hpx::endl;

return hpx::finalize();
}

```

Try copying the contents to your main.cpp file and look at the output. It can be a good idea to go through the program step by step with a debugger. You can also try changing the types or adding new arguments to functions to make sure you can get the types to match. The type of the then method can be especially tricky to get right (the continuation needs to take the future as an argument).

Note: HPX programs accept command line arguments. The most important one is `--hpx:threads=N` to set the number of OS-threads used by HPX. HPX uses one thread per core by default. Play around with the example above

and see what difference the number of threads makes on the `sort` function. See [Launching and configuring HPX applications](#) for more details on how and what options you can pass to *HPX*.

Tip: The example above used the construction `hpx::when_all(...).then(...)`. For convenience and performance it is a good idea to replace uses of `hpx::when_all(...).then(...)` with dataflow. See [Dataflow: Interest calculator](#) for more details on dataflow.

Tip: If possible, prefer to use the provided parallel algorithms instead of writing your own implementation. This can save you time and the resulting program is often faster.

2.2.6 Next steps

If you haven't done so already, reading the [Terminology](#) section will help you get familiar with the terms used in *HPX*.

The [Examples](#) section contains small, self-contained walkthroughs of example *HPX* programs. The [Local to remote: 1D stencil](#) example is a thorough, realistic example starting from a single node implementation and going stepwise to a distributed implementation.

The [Manual](#) contains detailed information on writing, building and running *HPX* applications.

2.3 Terminology

This section gives definitions for some of the terms used throughout the *HPX* documentation and source code.

Locality A locality in *HPX* describes a synchronous domain of execution, or the domain of bounded upper response time. This normally is just a single node in a cluster or a NUMA domain in a SMP machine.

Active Global Address Space

AGAS *HPX* incorporates a global address space. Any executing thread can access any object within the domain of the parallel application with the caveat that it must have appropriate access privileges. The model does not assume that global addresses are cache coherent; all loads and stores will deal directly with the site of the target object. All global addresses within a Synchronous Domain are assumed to be cache coherent for those processor cores that incorporate transparent caches. The Active Global Address Space used by *HPX* differs from research PGAS⁴⁹ models. Partitioned Global Address Space is passive in their means of address translation. Copy semantics, distributed compound operations, and affinity relationships are some of the global functionality supported by AGAS.

Process The concept of the “process” in *HPX* is extended beyond that of either sequential execution or communicating sequential processes. While the notion of process suggests action (as do “function” or “subroutine”) it has a further responsibility of context, that is, the logical container of program state. It is this aspect of operation that process is employed in *HPX*. Furthermore, referring to “parallel processes” in *HPX* designates the presence of parallelism within the context of a given process, as well as the coarse grained parallelism achieved through concurrency of multiple processes of an executing user job. *HPX* processes provide a hierarchical name space within the framework of the active global address space and support multiple means of internal state access from external sources.

Parcel The Parcel is a component in *HPX* that communicates data, invokes an action at a distance, and distributes flow-control through the migration of continuations. Parcels bridge the gap of asynchrony between synchronous

⁴⁹ <https://www.pgas.org/>

domains while maintaining symmetry of semantics between local and global execution. Parcels enable message-driven computation and may be seen as a form of “active messages”. Other important forms of message-driven computation predating active messages include [dataflow tokens](#)⁵⁰, the [J-machine’s](#)⁵¹ support for remote method instantiation, and at the coarse grained variations of Unix remote procedure calls, among others. This enables work to be moved to the data as well as performing the more common action of bringing data to the work. A parcel can cause actions to occur remotely and asynchronously, among which are the creation of threads at different system nodes or synchronous domains.

Local Control Object

Lightweight Control Object

LCO A local control object (sometimes called a lightweight control object) is a general term for the synchronization mechanisms used in *HPX*. Any object implementing a certain concept can be seen as an LCO. This concepts encapsulates the ability to be triggered by one or more events which when taking the object into a predefined state will cause a thread to be executed. This could either create a new thread or resume an existing thread.

The LCO is a family of synchronization functions potentially representing many classes of synchronization constructs, each with many possible variations and multiple instances. The LCO is sufficiently general that it can subsume the functionality of conventional synchronization primitives such as spinlocks, mutexes, semaphores, and global barriers. However due to the rich concept an LCO can represent powerful synchronization and control functionality not widely employed, such as dataflow and futures (among others), which open up enormous opportunities for rich diversity of distributed control and operation.

See [Using LCOs](#) for more details on how to use LCOs in *HPX*.

Action An action is a function that can be invoked remotely. In *HPX* a plain function can be made into an action using a macro. See [Applying actions](#) for details on how to use actions in *HPX*.

Component A component is a C++ object which can be accessed remotely. A component can also contain member functions which can be invoked remotely. These are referred to as component actions. See [Writing components](#) for details on how to use components in *HPX*.

2.4 Examples

The following sections analyze some examples to help you get familiar with the *HPX* style of programming. We start off with simple examples that utilize basic *HPX* elements and then begin to expose the reader to the more complex and powerful *HPX* concepts.

2.4.1 Asynchronous execution with `hpx::async`: Fibonacci

The Fibonacci sequence is a sequence of numbers starting with 0 and 1 where every subsequent number is the sum of the previous two numbers. In this example, we will use *HPX* to calculate the value of the n-th element of the Fibonacci sequence. In order to compute this problem in parallel, we will use a facility known as a future.

As shown in the [Fig. 2.1](#) below, a future encapsulates a delayed computation. It acts as a proxy for a result initially not known, most of the time because the computation of the result has not completed yet. The future synchronizes the access of this value by optionally suspending any *HPX*-threads requesting the result until the value is available. When a future is created, it spawns a new *HPX*-thread (either remotely with a [parcel](#) or locally by placing it into the thread queue) which, when run, will execute the function associated with the future. The arguments of the function are bound when the future is created.

Once the function has finished executing, a write operation is performed on the future. The write operation marks the future as completed, and optionally stores data returned by the function. When the result of the delayed computation

⁵⁰ http://en.wikipedia.org/wiki/Dataflow_architecture

⁵¹ <http://en.wikipedia.org/wiki/J%20Machine>

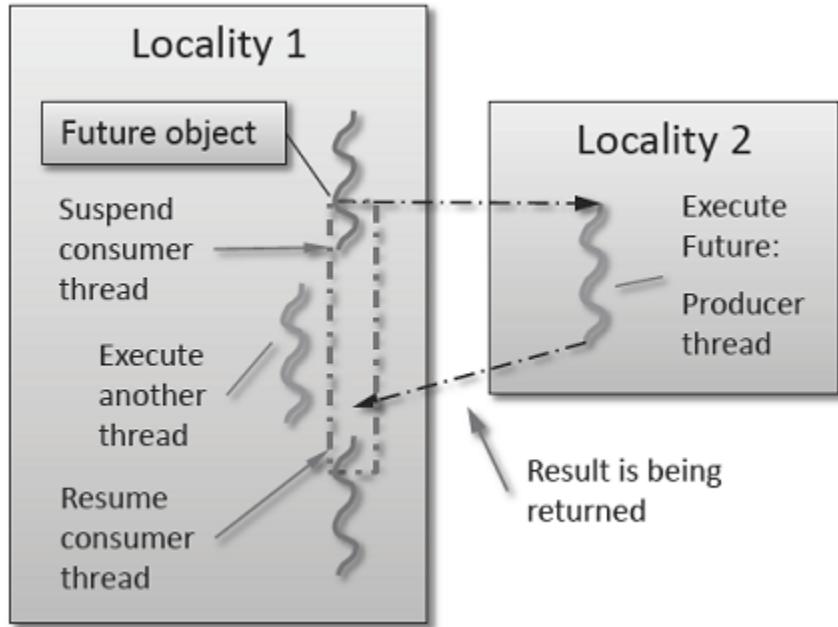


Fig. 2.1: Schematic of a future execution.

is needed, a read operation is performed on the future. If the future's function hasn't completed when a read operation is performed on it, the reader *HPX*-thread is suspended until the future is ready. The future facility allows *HPX* to schedule work early in a program so that when the function value is needed it will already be calculated and available. We use this property in our Fibonacci example below to enable its parallel execution.

Setup

The source code for this example can be found here: `fibonacci_local.cpp`.

To compile this program, go to your *HPX* build directory (see [HPX build system](#) for information on configuring and building *HPX*) and enter:

```
make examples.quickstart.fibonacci_local
```

To run the program type:

```
./bin/fibonacci_local
```

This should print (time should be approximate):

```
fibonacci(10) == 55
elapsed time: 0.002430 [s]
```

This run used the default settings, which calculate the tenth element of the Fibonacci sequence. To declare which Fibonacci value you want to calculate, use the `--n-value` option. Additionally you can use the `--hpx:threads` option to declare how many OS-threads you wish to use when running the program. For instance, running:

```
./bin/fibonacci --n-value 20 --hpx:threads 4
```

Will yield:

```
fibonacci(20) == 6765
elapsed time: 0.062854 [s]
```

Walkthrough

Now that you have compiled and run the code, let's look at how the code works. Since this code is written in C++, we will begin with the `main()` function. Here you can see that in *HPX*, `main()` is only used to initialize the runtime system. It is important to note that application-specific command line options are defined here. *HPX* uses [Boost.Program Options](#)⁵² for command line processing. You can see that our programs `--n-value` option is set by calling the `add_options()` method on an instance of `hpx::program_options::options_description`. The default value of the variable is set to 10. This is why when we ran the program for the first time without using the `--n-value` option the program returned the 10th value of the Fibonacci sequence. The constructor argument of the description is the text that appears when a user uses the `--hpx:help` option to see what command line options are available. `HPX_APPLICATION_STRING` is a macro that expands to a string constant containing the name of the *HPX* application currently being compiled.

In *HPX* `main()` is used to initialize the runtime system and pass the command line arguments to the program. If you wish to add command line options to your program you would add them here using the instance of the Boost class `options_description`, and invoking the public member function `.add_options()` (see [Boost Documentation](#)⁵³ for more details). `hpx::init` calls `hpx_main()` after setting up *HPX*, which is where the logic of our program is encoded.

```
int main(int argc, char* argv[])
{
    // Configure application-specific options
    hpx::program_options::options_description
        desc_commandline("Usage: " HPX_APPLICATION_STRING " [options]");

    desc_commandline.add_options()
        ("n-value",
         hpx::program_options::value<std::uint64_t>()>default_value(10),
         "n value for the Fibonacci function")
        ;

    // Initialize and run HPX
    return hpx::init(desc_commandline, argc, argv);
}
```

The `hpx::init` function in `main()` starts the runtime system, and invokes `hpx_main()` as the first *HPX*-thread. Below we can see that the basic program is simple. The command line option `--n-value` is read in, a timer (`hpx::util::high_resolution_timer`) is set up to record the time it takes to do the computation, the `fibonacci` function is invoked synchronously, and the answer is printed out.

```
int hpx_main(hpx::program_options::variables_map& vm)
{
    // extract command line argument, i.e. fib(N)
    std::uint64_t n = vm["n-value"].as<std::uint64_t>();

    {
        // Keep track of the time required to execute.
        hpx::util::high_resolution_timer t;
```

(continues on next page)

⁵² https://www.boost.org/doc/html/program_options.html

⁵³ <https://www.boost.org/doc/>

(continued from previous page)

```

    std::uint64_t r = fibonacci(n);

    char const* fmt = "fibonacci({1}) == {2}\nelapsed time: {3} [s]\n";
    hpx::util::format_to(std::cout, fmt, n, r, t.elapsed());
}

return hpx::finalize(); // Handles HPX shutdown
}

```

The `fibonacci` function itself is synchronous as the work done inside is asynchronous. To understand what is happening we have to look inside the `fibonacci` function:

```

std::uint64_t fibonacci(std::uint64_t n)
{
    if (n < 2)
        return n;

    // Invoking the Fibonacci algorithm twice is inefficient.
    // However, we intentionally demonstrate it this way to create some
    // heavy workload.

    hpx::future<std::uint64_t> n1 = hpx::async(fibonacci, n - 1);
    hpx::future<std::uint64_t> n2 = hpx::async(fibonacci, n - 2);

    return n1.get() + n2.get(); // wait for the Futures to return their values
}

```

This block of code looks similar to regular C++ code. First, `if (n < 2)`, meaning `n` is 0 or 1, then we return 0 or 1 (recall the first element of the Fibonacci sequence is 0 and the second is 1). If `n` is larger than 1 we spawn two new tasks whose results are contained in `n1` and `n2`. This is done using `hpx::async` which takes as arguments a function (function pointer, object or lambda) and the arguments to the function. Instead of returning a `std::uint64_t` like `fibonacci` does, `hpx::async` returns a future of a `std::uint64_t`, i.e. `hpx::future<std::uint64_t>`. Each of these futures represents an asynchronous, recursive call to `fibonacci`. After we've created the futures, we wait for both of them to finish computing, we add them together, and return that value as our result. We get the values from the futures using the `get` method. The recursive call tree will continue until `n` is equal to 0 or 1, at which point the value can be returned because it is implicitly known. When this termination condition is reached, the futures can then be added up, producing the `n`-th value of the Fibonacci sequence.

Note that calling `get` potentially blocks the calling *HPX*-thread, and lets other *HPX*-threads run in the meantime. There are, however, more efficient ways of doing this. `examples/quickstart/fibonacci_futures.cpp` contains many more variations of locally computing the Fibonacci numbers, where each method makes different tradeoffs in where asynchrony and parallelism is applied. To get started, however, the method above is sufficient and optimizations can be applied once you are more familiar with *HPX*. The example *Dataflow: Interest calculator* presents dataflow, which is a way to more efficiently chain together multiple tasks.

2.4.2 Asynchronous execution with `hpx::async` and `actions`: Fibonacci

This example extends the [previous example](#) by introducing `actions`: functions that can be run remotely. In this example, however, we will still only run the action locally. The mechanism to execute `actions` stays the same: `hpx::async`. Later examples will demonstrate running actions on remote `localities` (e.g. [Remote execution with actions: Hello world](#)).

Setup

The source code for this example can be found here: `fibonacci.cpp`.

To compile this program, go to your *HPX* build directory (see [HPX build system](#) for information on configuring and building *HPX*) and enter:

```
make examples.quickstart.fibonacci
```

To run the program type:

```
./bin/fibonacci
```

This should print (time should be approximate):

```
fibonacci(10) == 55
elapsed time: 0.00186288 [s]
```

This run used the default settings, which calculate the tenth element of the Fibonacci sequence. To declare which Fibonacci value you want to calculate, use the `--n-value` option. Additionally you can use the `--hpx:threads` option to declare how many OS-threads you wish to use when running the program. For instance, running:

```
./bin/fibonacci --n-value 20 --hpx:threads 4
```

Will yield:

```
fibonacci(20) == 6765
elapsed time: 0.233827 [s]
```

Walkthrough

The code needed to initialize the *HPX* runtime is the same as in the [previous example](#):

```
//[fib_main
int main(int argc, char* argv[])
{
    // Configure application-specific options
    hpx::program_options::options_description
        desc_commandline("Usage: " HPX_APPLICATION_STRING " [options]");
    desc_commandline.add_options()
        ("n-value",
         hpx::program_options::value<std::uint64_t>() -> default_value(10),
         "n value for the Fibonacci function");
    ;

    // Initialize and run HPX
    return hpx::init(desc_commandline, argc, argv);
```

The `hpx::init` function in `main()` starts the runtime system, and invokes `hpx_main()` as the first *HPX*-thread. The command line option `--n-value` is read in, a timer (`hpx::util::high_resolution_timer`) is set up to record the time it takes to do the computation, the `fibonacci action` is invoked synchronously, and the answer is printed out.

```
//[fib_hpx_main
int hpx_main(hpx::program_options::variables_map& vm)
{
    // extract command line argument, i.e. fib(N)
    std::uint64_t n = vm["n-value"].as<std::uint64_t>();

    {
        // Keep track of the time required to execute.
        hpx::util::high_resolution_timer t;

        // Wait for fib() to return the value
        fibonacci_action fib;
        std::uint64_t r = fib(hpx::find_here(), n);

        char const* fmt = "fibonacci({1}) == {2}\nelapsed time: {3} [s]\n";
        hpx::util::format_to(std::cout, fmt, n, r, t.elapsed());
    }

    return hpx::finalize(); // Handles HPX shutdown
}
```

Upon a closer look we see that we've created a `std::uint64_t` to store the result of invoking our `fibonacci_action` `fib`. This `action` will launch synchronously (as the work done inside of the `action` will be asynchronous itself) and return the result of the Fibonacci sequence. But wait, what is an `action`? And what is this `fibonacci_action`? For starters, an `action` is a wrapper for a function. By wrapping functions, *HPX* can send packets of work to different processing units. These vehicles allow users to calculate work now, later, or on certain nodes. The first argument to our `action` is the location where the `action` should be run. In this case, we just want to run the `action` on the machine that we are currently on, so we use `hpx::find_here`. To further understand this we turn to the code to find where `fibonacci_action` was defined:

```
//[fib_action
// forward declaration of the Fibonacci function
std::uint64_t fibonacci(std::uint64_t n);

// This is to generate the required boilerplate we need for the remote
// invocation to work.
```

A plain `action` is the most basic form of `action`. Plain `actions` wrap simple global functions which are not associated with any particular object (we will discuss other types of `actions` in *Components and actions: Accumulator*). In this block of code the function `fibonacci()` is declared. After the declaration, the function is wrapped in an `action` in the declaration `HPX_PLAIN_ACTION`. This function takes two arguments: the name of the function that is to be wrapped and the name of the `action` that you are creating.

This picture should now start making sense. The function `fibonacci()` is wrapped in an `action` `fibonacci_action`, which was run synchronously but created asynchronous work, then returns a `std::uint64_t` representing the result of the function `fibonacci()`. Now, let's look at the function `fibonacci()`:

```
//[fib_func
std::uint64_t fibonacci(std::uint64_t n)
{
    if (n < 2)
        return n;
```

(continues on next page)

(continued from previous page)

```
// We restrict ourselves to execute the Fibonacci function locally.
hpx::naming::id_type const locality_id = hpx::find_here();

// Invoking the Fibonacci algorithm twice is inefficient.
// However, we intentionally demonstrate it this way to create some
// heavy workload.

fibonacci_action fib;
hpx::future<std::uint64_t> n1 =
    hpx::async(fib, locality_id, n - 1);
hpx::future<std::uint64_t> n2 =
    hpx::async(fib, locality_id, n - 2);

return n1.get() + n2.get(); // wait for the Futures to return their values
```

This block of code is much more straightforward and should look familiar from the [previous example](#). First, if $(n < 2)$, meaning n is 0 or 1, then we return 0 or 1 (recall the first element of the Fibonacci sequence is 0 and the second is 1). If n is larger than 1 we spawn two tasks using `hpx::async`. Each of these futures represents an asynchronous, recursive call to `fibonacci`. As previously we wait for both futures to finish computing, get the results, add them together, and return that value as our result. The recursive call tree will continue until n is equal to 0 or 1, at which point the value can be returned because it is implicitly known. When this termination condition is reached, the futures can then be added up, producing the n -th value of the Fibonacci sequence.

2.4.3 Remote execution with actions: Hello world

This program will print out a hello world message on every OS-thread on every *locality*. The output will look something like this:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 1 on locality 1
hello world from OS-thread 0 on locality 0
hello world from OS-thread 0 on locality 1
```

Setup

The source code for this example can be found here: `hello_world_distributed.cpp`.

To compile this program, go to your *HPX* build directory (see [HPX build system](#) for information on configuring and building *HPX*) and enter:

```
make examples.quickstart.hello_world_distributed
```

To run the program type:

```
./bin/hello_world_distributed
```

This should print:

```
hello world from OS-thread 0 on locality 0
```

To use more OS-threads use the command line option `--hpx:threads` and type the number of threads that you wish to use. For example, typing:

```
./bin/hello_world_distributed --hpx:threads 2
```

will yield:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
```

Notice how the ordering of the two print statements will change with subsequent runs. To run this program on multiple localities please see the section [How to use HPX applications with PBS](#).

Walkthrough

Now that you have compiled and run the code, let's look at how the code works, beginning with `main()`:

```
// Here is the main entry point. By using the include 'hpx/hpx_main.hpp' HPX
// will invoke the plain old C-main() as its first HPX thread.
int main()
{
    // Get a list of all available localities.
    std::vector<hpx::naming::id_type> localities =
        hpx::find_all_localities();

    // Reserve storage space for futures, one for each locality.
    std::vector<hpx::lcos::future<void>> futures;
    futures.reserve(localities.size());

    for (hpx::naming::id_type const& node : localities)
    {
        // Asynchronously start a new task. The task is encapsulated in a
        // future, which we can query to determine if the task has
        // completed.
        typedef hello_world_foreman_action action_type;
        futures.push_back(hpx::async<action_type>(node));
    }

    // The non-callback version of hpx::lcos::wait_all takes a single parameter,
    // a vector of futures to wait on. hpx::wait_all only returns when
    // all of the futures have finished.
    hpx::wait_all(futures);
    return 0;
}
```

In this excerpt of the code we again see the use of futures. This time the futures are stored in a vector so that they can easily be accessed. `hpx::wait_all` is a family of functions that wait on for an `std::vector<>` of futures to become ready. In this piece of code, we are using the synchronous version of `hpx::wait_all`, which takes one argument (the `std::vector<>` of futures to wait on). This function will not return until all the futures in the vector have been executed.

In [Asynchronous execution with hpx::async and actions: Fibonacci](#) we used `hpx::find_here` to specify the target of our actions. Here, we instead use `hpx::find_all_localities`, which returns an `std::vector<>` containing the identifiers of all the machines in the system, including the one that we are on.

As in [Asynchronous execution with hpx::async and actions: Fibonacci](#) our futures are set using `hpx::async<>`. The `hello_world_foreman_action` is declared here:

```
// Define the boilerplate code necessary for the function 'hello_world_foreman'
// to be invoked as an HPX action.
HPX_PLAIN_ACTION(hello_world_foreman, hello_world_foreman_action);
```

Another way of thinking about this wrapping technique is as follows: functions (the work to be done) are wrapped in actions, and actions can be executed locally or remotely (e.g. on another machine participating in the computation).

Now it is time to look at the `hello_world_foreman()` function which was wrapped in the action above:

```
void hello_world_foreman()
{
    // Get the number of worker OS-threads in use by this locality.
    std::size_t const os_threads = hpx::get_os_thread_count();

    // Find the global name of the current locality.
    hpx::naming::id_type const here = hpx::find_here();

    // Populate a set with the OS-thread numbers of all OS-threads on this
    // locality. When the hello world message has been printed on a particular
    // OS-thread, we will remove it from the set.
    std::set<std::size_t> attendance;
    for (std::size_t os_thread = 0; os_thread < os_threads; ++os_thread)
        attendance.insert(os_thread);

    // As long as there are still elements in the set, we must keep scheduling
    // HPX-threads. Because HPX features work-stealing task schedulers, we have
    // no way of enforcing which worker OS-thread will actually execute
    // each HPX-thread.
    while (!attendance.empty())
    {
        // Each iteration, we create a task for each element in the set of
        // OS-threads that have not said "Hello world". Each of these tasks
        // is encapsulated in a future.
        std::vector<hpx::lcos::future<std::size_t>> futures;
        futures.reserve(attendance.size());

        for (std::size_t worker : attendance)
        {
            // Asynchronously start a new task. The task is encapsulated in a
            // future, which we can query to determine if the task has
            // completed. We give the task a hint to run on a particular worker
            // thread, but no guarantees are given by the scheduler that the
            // task will actually run on that worker thread.
            hpx::parallel::execution::default_executor exec(
                hpx::threads::thread_priority_default,
                hpx::threads::thread_stacksize_default,
                hpx::threads::thread_schedule_hint(
                    hpx::threads::thread_schedule_hint_mode_thread, worker));
            futures.push_back(hpx::async(exec, hello_world_worker, worker));
        }

        // Wait for all of the futures to finish. The callback version of the
        // hpx::lcos::wait_each function takes two arguments: a vector of futures,
        // and a binary callback. The callback takes two arguments; the first
        // is the index of the future in the vector, and the second is the
        // return value of the future. hpx::lcos::wait_each doesn't return until
        // all the futures in the vector have returned.
        hpx::lcos::local::spinlock mtx;
```

(continues on next page)

(continued from previous page)

```

    hpx::lcos::wait_each(
        hpx::util::unwrapping([&](std::size_t t) {
            if (std::size_t(-1) != t)
            {
                std::lock_guard<hpx::lcos::local::spinlock> lk(mtx);
                attendance.erase(t);
            }
        }),
        futures);
    }
}

```

Now, before we discuss `hello_world_foreman()`, let's talk about the `hpx::wait_each` function. The version of `hpx::lcos::wait_each` invokes a callback function provided by the user, supplying the callback function with the result of the future.

In `hello_world_foreman()`, an `std::set` called `attendance` keeps track of which OS-threads have printed out the hello world message. When the OS-thread prints out the statement, the future is marked as ready, and `hpx::lcos::wait_each` in `hello_world_foreman()`. If it is not executing on the correct OS-thread, it returns a value of -1, which causes `hello_world_foreman()` to leave the OS-thread id in `attendance`.

```

std::size_t hello_world_worker(std::size_t desired)
{
    // Returns the OS-thread number of the worker that is running this
    // HPX-thread.
    std::size_t current = hpx::get_worker_thread_num();
    if (current == desired)
    {
        // The HPX-thread has been run on the desired OS-thread.
        char const* msg = "hello world from OS-thread {1} on locality {2}\n";

        hpx::util::format_to(hpx::cout, msg, desired, hpx::get_locality_id())
            << hpx::flush;

        return desired;
    }

    // This HPX-thread has been run by the wrong OS-thread, make the foreman
    // try again by rescheduling it.
    return std::size_t(-1);
}

```

Because *HPX* features work stealing task schedulers, there is no way to guarantee that an action will be scheduled on a particular OS-thread. This is why we must use a guess-and-check approach.

2.4.4 Components and actions: Accumulator

The accumulator example demonstrates the use of components. Components are C++ classes that expose methods as a type of *HPX* action. These actions are called component actions.

Components are globally named, meaning that a component action can be called remotely (e.g. from another machine). There are two accumulator examples in *HPX*:

In the *Asynchronous execution with hpx::async and actions: Fibonacci* and the *Remote execution with actions: Hello world*, we introduced plain actions, which wrapped global functions. The target of a plain action is an identifier which

refers to a particular machine involved in the computation. For plain actions, the target is the machine where the action will be executed.

Component actions, however, do not target machines. Instead, they target component instances. The instance may live on the machine that we've invoked the component action from, or it may live on another machine.

The component in this example exposes three different functions:

- `reset()` - Resets the accumulator value to 0.
- `add(arg)` - Adds `arg` to the accumulators value.
- `query()` - Queries the value of the accumulator.

This example creates an instance of the accumulator, and then allows the user to enter commands at a prompt, which subsequently invoke actions on the accumulator instance.

Setup

The source code for this example can be found here: `accumulator_client.cpp`.

To compile this program, go to your *HPX* build directory (see [HPX build system](#) for information on configuring and building *HPX*) and enter:

```
make examples.accumulators.accumulator
```

To run the program type:

```
./bin/accumulator_client
```

Once the program starts running, it will print the following prompt and then wait for input. An example session is given below:

```
commands: reset, add [amount], query, help, quit
> add 5
> add 10
> query
15
> add 2
> query
17
> reset
> add 1
> query
1
> quit
```

Walkthrough

Now, let's take a look at the source code of the accumulator example. This example consists of two parts: an *HPX* component library (a library that exposes an *HPX* component) and a client application which uses the library. This walkthrough will cover the *HPX* component library. The code for the client application can be found here: `accumulator_client.cpp`.

An *HPX* component is represented by two C++ classes:

- **A server class** - The implementation of the components functionality.
- **A client class** - A high-level interface that acts as a proxy for an instance of the component.

Typically, these two classes all have the same name, but the server class usually lives in different sub-namespaces (server). For example, the full names of the two classes in accumulator are:

- examples::server::accumulator (server class)
- examples::accumulator (client class)

The server class

The following code is from: `accumulator.hpp`.

All *HPX* component server classes must inherit publicly from the *HPX* component base class: `hpx::components::component_base`

The accumulator component inherits from `hpx::components::locking_hook`. This allows the runtime system to ensure that all action invocations are serialized. That means that the system ensures that no two actions are invoked at the same time on a given component instance. This makes the component thread safe and no additional locking has to be implemented by the user. Moreover, accumulator component is a component, because it also inherits from `hpx::components::component_base` (the template argument passed to `locking_hook` is used as its base class). The following snippet shows the corresponding code:

```
//[accumulator_server_inherit
class accumulator
: public hpx::components::locking_hook<
```

Our accumulator class will need a data member to store its value in, so let's declare a data member:

```
//[accumulator_server_data_member
```

The constructor for this class simply initializes `value_` to 0:

```
//[accumulator_server_ctor
```

Next, let's look at the three methods of this component that we will be exposing as component actions:

```
//[accumulator_methods
/// Reset the components value to 0.
void reset()
{
    // set value_ to 0.
    value_ = 0;
}

/// Add the given number to the accumulator.
void add(argument_type arg)
{
    // add value_ to arg, and store the result in value_.
    value_ += arg;
}

/// Return the current value to the caller.
argument_type query() const
{
    // Get the value of value_.
    return value_;
```

Here are the action types. These types wrap the methods we're exposing. The wrapping technique is very similar to the one used in the [Asynchronous execution with hpx::async and actions: Fibonacci](#) and the [Remote execution with actions: Hello world](#):

```
//[accumulator_action_types
HPX_DEFINE_COMPONENT_ACTION(accumulator, reset);
HPX_DEFINE_COMPONENT_ACTION(accumulator, add);
```

The last piece of code in the server class header is the declaration of the action type registration code:

```
//[accumulator_registration_declarations
HPX_REGISTER_ACTION_DECLARATION(
    examples::server::accumulator::reset_action,
    accumulator_reset_action);

HPX_REGISTER_ACTION_DECLARATION(
    examples::server::accumulator::add_action,
    accumulator_add_action);

HPX_REGISTER_ACTION_DECLARATION(
    examples::server::accumulator::query_action,
```

Note: The code above must be placed in the global namespace.

The rest of the registration code is in `accumulator.cpp`

```
//[accumulator_registration_definitions
///////////////////////////////
// Add factory registration functionality.
HPX_REGISTER_COMPONENT_MODULE();

///////////////////////////////
typedef hpx::components::component<
    examples::server::accumulator
> accumulator_type;

HPX_REGISTER_COMPONENT(accumulator_type, accumulator);

///////////////////////////////
// Serialization support for accumulator actions.
HPX_REGISTER_ACTION(
    accumulator_type::wrapped_type::reset_action,
    accumulator_reset_action);
HPX_REGISTER_ACTION(
    accumulator_type::wrapped_type::add_action,
    accumulator_add_action);
HPX_REGISTER_ACTION(
    accumulator_type::wrapped_type::query_action,
```

Note: The code above must be placed in the global namespace.

The client class

The following code is from `accumulator.hpp`.

The client class is the primary interface to a component instance. Client classes are used to create components:

```
// Create a component on this locality.
examples::accumulator c = hpx::new<examples::accumulator>(hpx::find_here());
```

and to invoke component actions:

```
c.add(hpx::launch::apply, 4);
```

Clients, like servers, need to inherit from a base class, this time, `hpx::components::client_base`:

```
//[accumulator_client_inherit
class accumulator
: public hpx::components::client_base<
    accumulator, server::accumulator
```

For readability, we typedef the base class like so:

```
//[accumulator_base_type
typedef hpx::components::client_base<
    accumulator, server::accumulator
```

Here are examples of how to expose actions through a client class:

There are a few different ways of invoking actions:

- **Non-blocking:** For actions which don't have return types, or when we do not care about the result of an action, we can invoke the action using fire-and-forget semantics. This means that once we have asked HPX to compute the action, we forget about it completely and continue with our computation. We use `hpx::apply` to invoke an action in a non-blocking fashion.

```
//[accumulator_client_reset_non_blocking
void reset(hpx::launch::apply_policy)
{
    HPX_ASSERT(this->get_id());

    typedef server::accumulator::reset_action action_type;
    hpx::apply<action_type>(this->get_id());
```

- **Asynchronous:** Futures, as demonstrated in [Asynchronous execution with `hpx::async`: Fibonacci](#), [Asynchronous execution with `hpx::async` and actions: Fibonacci](#), and the [Remote execution with actions: Hello world](#), enable asynchronous action invocation. Here's an example from the accumulator client class:

```
//[accumulator_client_query_async
hpx::future<argument_type> query(hpx::launch::async_policy)
{
    HPX_ASSERT(this->get_id());

    typedef server::accumulator::query_action action_type;
    return hpx::async<action_type>(hpx::launch::async, this->get_id());
```

- **Synchronous:** To invoke an action in a fully synchronous manner, we can simply call `hpx::async().get()` (e.g., create a future and immediately wait on it to be ready). Here's an example from the accumulator client class:

```
//[accumulator_client_add_sync
void add(argument_type arg)
{
    HPX_ASSERT(this->get_id());

    typedef server::accumulator::add_action action_type;
    action_type() (this->get_id(), arg);
```

Note that `this->get_id()` references a data member of the `hpx::components::client_base` base class which identifies the server accumulator instance.

`hpx::naming::id_type` is a type which represents a global identifier in *HPX*. This type specifies the target of an action. This is the type that is returned by `hpx::find_here` in which case it represents the *locality* the code is running on.

2.4.5 Dataflow: Interest calculator

HPX provides its users with several different tools to simply express parallel concepts. One of these tools is a *local control object* (*LCO*) called dataflow. An *LCO* is a type of component that can spawn a new thread when triggered. They are also distinguished from other components by a standard interface which allow users to understand and use them easily. A Dataflow, being an *LCO*, is triggered when the values it depends on become available. For instance, if you have a calculation *X* that depends on the results of three other calculations, you could set up a dataflow that would begin the calculation *X* as soon as the other three calculations have returned their values. Dataflows are set up to depend on other dataflows. It is this property that makes dataflow a powerful parallelization tool. If you understand the dependencies of your calculation, you can devise a simple algorithm which sets up a dependency tree to be executed. In this example, we calculate compound interest. To calculate compound interest, one must calculate the interest made in each compound period, and then add that interest back to the principal before calculating the interest made in the next period. A practical person would of course use the formula for compound interest:

$$F = P(1 + i)^n$$

where *F* is the future value, *P* is the principal value, *i* is the interest rate, and *n* is the number of compound periods.

Nevertheless, we have chosen for the sake of example to manually calculate the future value by iterating:

$$I = Pi$$

and

$$P = P + I$$

Setup

The source code for this example can be found here: `interest_calculator.cpp`.

To compile this program, go to your *HPX* build directory (see *HPX build system* for information on configuring and building *HPX*) and enter:

```
make examples.quickstart.interest_calculator
```

To run the program type:

```
./bin/interest_calculator --principal 100 --rate 5 --cp 6 --time 36
```

This should print:

```
Final amount: 134.01
Amount made: 34.0096
```

Walkthrough

Let us begin with main, here we can see that we again are using Boost.Program Options to set our command line variables (see *Asynchronous execution with hpx::async and actions: Fibonacci* for more details). These options set the principal, rate, compound period, and time. It is important to note that the units of time for cp and time must be the same.

```
//[interest_main
int main(int argc, char ** argv)
{
    options_description cmdline("Usage: " HPX_APPLICATION_STRING " [options]");

    cmdline.add_options()
        ("principal", value<double>()>default_value(1000), "The principal [$]")
        ("rate", value<double>()>default_value(7), "The interest rate [%]")
        ("cp", value<int>()>default_value(12), "The compound period [months]")
        ("time", value<int>()>default_value(12*30),
            "The time money is invested [months]")
;

    return hpx::init(cmdline, argc, argv);
}
```

Next we look at hpx_main.

```
//[interest_hpx_main
int hpx_main(variables_map & vm)
{
    {
        using hpx::shared_future;
        using hpx::make_ready_future;
        using hpx::dataflow;
        using hpx::util::unwrapping;
        hpx::naming::id_type here = hpx::find_here();

        double init_principal=vm["principal"].as<double>(); //Initial principal
        double init_rate=vm["rate"].as<double>(); //Interest rate
        int cp=vm["cp"].as<int>(); //Length of a compound period
        int t=vm["time"].as<int>(); //Length of time money is invested

        init_rate/=100; //Rate is a % and must be converted
        t/=cp; //Determine how many times to iterate interest calculation:
                //How many full compound periods can fit in the time invested

        // In non-dataflow terms the implemented algorithm would look like:
        //
        // int t = 5;      // number of time periods to use
        // double principal = init_principal;
        // double rate = init_rate;
        //
        // for (int i = 0; i < t; ++i)
        // {
        //     double interest = calc(principal, rate);
        //     principal = add(principal, interest);
    }
}
```

(continues on next page)

(continued from previous page)

```

// }

// Please note the similarity with the code below!

shared_future<double> principal = make_ready_future(init_principal);
shared_future<double> rate = make_ready_future(init_rate);

for (int i = 0; i < t; ++i)
{
    shared_future<double> interest = dataflow(unwrapping(calc), principal,_
→rate);
    principal = dataflow(unwrapping(add), principal, interest);
}

// wait for the dataflow execution graph to be finished calculating our
// overall interest
double result = principal.get();

std::cout << "Final amount: " << result << std::endl;
std::cout << "Amount made: " << result-init_principal << std::endl;
}

return hpx::finalize();

```

Here we find our command line variables read in, the rate is converted from a percent to a decimal, the number of calculation iterations is determined, and then our shared_futures are set up. Notice that we first place our principal and rate into shares futures by passing the variables `init_principal` and `init_rate` using `hpx::make_ready_future`.

In this way `hpx::shared_future<double> principal` and `rate` will be initialized to `init_principal` and `init_rate` when `hpx::make_ready_future<double>` returns a future containing those initial values. These shared futures then enter the for loop and are passed to `interest`. Next `principal` and `interest` are passed to the reassignment of `principal` using a `hpx::dataflow`. A dataflow will first wait for its arguments to be ready before launching any callbacks, so `add` in this case will not begin until both `principal` and `interest` are ready. This loop continues for each compound period that must be calculated. To see how `interest` and `principal` are calculated in the loop let us look at `calc_action` and `add_action`:

```

///[interest_calc_add_action
// Calculate interest for one period
double calc(double principal, double rate)
{
    return principal * rate;
}

///////////////////////////////
// Add the amount made to the principal
double add(double principal, double interest)
{
    return principal + interest;
}

```

After the shared future dependencies have been defined in `hpx_main`, we see the following statement:

```
double result = principal.get();
```

This statement calls `hpx::future::get` on the shared future `principal` which had its value calculated by our for loop. The program will wait here until the entire dataflow tree has been calculated and the value assigned to `result`. The program then prints out the final value of the investment and the amount of interest made by subtracting the final

value of the investment from the initial value of the investment.

2.4.6 Local to remote: 1D stencil

When developers write code they typically begin with a simple serial code and build upon it until all of the required functionality is present. The following set of examples were developed to demonstrate this iterative process of evolving a simple serial program to an efficient, fully-distributed *HPX* application. For this demonstration, we implemented a 1D heat distribution problem. This calculation simulates the diffusion of heat across a ring from an initialized state to some user-defined point in the future. It does this by breaking each portion of the ring into discrete segments and using the current segment's temperature and the temperature of the surrounding segments to calculate the temperature of the current segment in the next timestep as shown by Fig. 2.2 below.

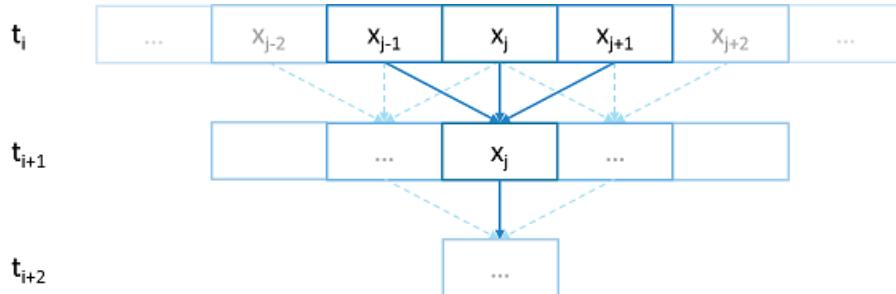


Fig. 2.2: Heat diffusion example program flow.

We parallelize this code over the following eight examples:

- Example 1
- Example 2
- Example 3
- Example 4
- Example 5
- Example 6
- Example 7
- Example 8

The first example is straight serial code. In this code we instantiate a vector \mathbf{U} that contains two vectors of doubles as seen in the structure `stepper`.

```
//[stepper_1
struct stepper
{
    // Our partition type
    typedef double partition;

    // Our data for one time step
    typedef std::vector<partition> space;

    // Our operator
    static double heat(double left, double middle, double right)
    {
        return middle + (k*dt/(dx*dx)) * (left - 2*middle + right);
    }
}
```

(continues on next page)

(continued from previous page)

```

}

// do all the work on 'nx' data points for 'nt' time steps
space do_work(std::size_t nx, std::size_t nt)
{
    // U[t][i] is the state of position i at time t.
    std::vector<space> U(2);
    for (space& s : U)
        s.resize(nx);

    // Initial conditions: f(0, i) = i
    for (std::size_t i = 0; i != nx; ++i)
        U[0][i] = double(i);

    // Actual time step loop
    for (std::size_t t = 0; t != nt; ++t)
    {
        space const& current = U[t % 2];
        space& next = U[(t + 1) % 2];

        next[0] = heat(current[nx-1], current[0], current[1]);

        for (std::size_t i = 1; i != nx-1; ++i)
            next[i] = heat(current[i-1], current[i], current[i+1]);

        next[nx-1] = heat(current[nx-2], current[nx-1], current[0]);
    }

    // Return the solution at time-step 'nt'.
    return U[nt % 2];
}

```

Each element in the vector of doubles represents a single grid point. To calculate the change in heat distribution, the temperature of each grid point, along with its neighbors, is passed to the function `heat`. In order to improve readability, references named `current` and `next` are created which, depending on the time step, point to the first and second vector of doubles. The first vector of doubles is initialized with a simple heat ramp. After calling the `heat` function with the data in the `current` vector, the results are placed into the `next` vector.

In example 2 we employ a technique called futurization. Futurization is a method by which we can easily transform a code that is serially executed into a code that creates asynchronous threads. In the simplest case this involves replacing a variable with a future to a variable, a function with a future to a function, and adding a `.get()` at the point where a value is actually needed. The code below shows how this technique was applied to the `stepper`.

```

///[stepper_2
struct stepper
{
    // Our partition type
    typedef hpx::shared_future<double> partition;

    // Our data for one time step
    typedef std::vector<partition> space;

    // Our operator
    static double heat(double left, double middle, double right)
    {
        return middle + (k*dt/(dx*dx)) * (left - 2*middle + right);
    }
}

```

(continues on next page)

(continued from previous page)

```

}

// do all the work on 'nx' data points for 'nt' time steps
hpx::future<space> do_work(std::size_t nx, std::size_t nt)
{
    using hpx::dataflow;
    using hpx::util::unwrapping;

    // U[t][i] is the state of position i at time t.
    std::vector<space> U(2);
    for (space& s : U)
        s.resize(nx);

    // Initial conditions: f(0, i) = i
    for (std::size_t i = 0; i != nx; ++i)
        U[0][i] = hpx::make_ready_future(double(i));

    auto Op = unwrapping(&stepper::heat);

    // Actual time step loop
    for (std::size_t t = 0; t != nt; ++t)
    {
        space const& current = U[t % 2];
        space& next = U[(t + 1) % 2];

        // WHEN U[t][i-1], U[t][i], and U[t][i+1] have been computed, THEN we
        // can compute U[t+1][i]
        for (std::size_t i = 0; i != nx; ++i)
        {
            next[i] = dataflow(
                hpx::launch::async, Op,
                current[idx(i, -1, nx)], current[i], current[idx(i, +1, nx)])
        }
    }

    // Now the asynchronous computation is running; the above for-loop does not
    // wait on anything. There is no implicit waiting at the end of each timestep;
    // the computation of each U[t][i] will begin as soon as its dependencies
    // are ready and hardware is available.

    // Return the solution at time-step 'nt'.
    return hpx::when_all(U[nt % 2]);
}

```

In example 2, we redefine our partition type as a `shared_future` and, in `main`, create the object `result`, which is a future to a vector of partitions. We use `result` to represent the last vector in a string of vectors created for each timestep. In order to move to the next timestep, the values of a partition and its neighbors must be passed to `heat` once the futures that contain them are ready. In *HPX*, we have an LCO (Local Control Object) named Dataflow that assists the programmer in expressing this dependency. Dataflow allows us to pass the results of a set of futures to a specified function when the futures are ready. Dataflow takes three types of arguments, one which instructs the dataflow on how to perform the function call (async or sync), the function to call (in this case `Op`), and futures to the arguments that will be passed to the function. When called, dataflow immediately returns a future to the result of the specified function. This allows users to string dataflows together and construct an execution tree.

After the values of the futures in dataflow are ready, the values must be pulled out of the future container to be passed to the function `heat`. In order to do this, we use the *HPX* facility `unwrapping`, which underneath calls `.get()` on

each of the futures so that the function `heat` will be passed doubles and not futures to doubles.

By setting up the algorithm this way, the program will be able to execute as quickly as the dependencies of each future are met. Unfortunately, this example runs terribly slow. This increase in execution time is caused by the overheads needed to create a future for each data point. Because the work done within each call to `heat` is very small, the overhead of creating and scheduling each of the three futures is greater than that of the actual useful work! In order to amortize the overheads of our synchronization techniques, we need to be able to control the amount of work that will be done with each future. We call this amount of work per overhead grain size.

In example 3, we return to our serial code to figure out how to control the grain size of our program. The strategy that we employ is to create “partitions” of data points. The user can define how many partitions are created and how many data points are contained in each partition. This is accomplished by creating the `struct partition`, which contains a member object `data_`, a vector of doubles that holds the data points assigned to a particular instance of `partition`.

In example 4, we take advantage of the partition setup by redefining `space` to be a vector of `shared_futures` with each future representing a partition. In this manner, each future represents several data points. Because the user can define how many data points are in each partition, and, therefore, how many data points are represented by one future, a user can control the grainsize of the simulation. The rest of the code is then futurized in the same manner as example 2. It should be noted how strikingly similar example 4 is to example 2.

Example 4 finally shows good results. This code scales equivalently to the OpenMP version. While these results are promising, there are more opportunities to improve the application’s scalability. Currently, this code only runs on one *locality*, but to get the full benefit of *HPX*, we need to be able to distribute the work to other machines in a cluster. We begin to add this functionality in example 5.

In order to run on a distributed system, a large amount of boilerplate code must be added. Fortunately, *HPX* provides us with the concept of a `component`, which saves us from having to write quite as much code. A component is an object that can be remotely accessed using its global address. Components are made of two parts: a server and a client class. While the client class is not required, abstracting the server behind a client allows us to ensure type safety instead of having to pass around pointers to global objects. Example 5 renames example 4’s `struct partition` to `partition_data` and adds serialization support. Next, we add the server side representation of the data in the structure `partition_server`. `Partition_server` inherits from `hpx::components::component_base`, which contains a server-side component boilerplate. The boilerplate code allows a component’s public members to be accessible anywhere on the machine via its Global Identifier (GID). To encapsulate the component, we create a client side helper class. This object allows us to create new instances of our component and access its members without having to know its GID. In addition, we are using the client class to assist us with managing our asynchrony. For example, our client class `partition`’s member function `get_data()` returns a future to `partition_data` `get_data()`. This struct inherits its boilerplate code from `hpx::components::client_base`.

In the structure `stepper`, we have also had to make some changes to accommodate a distributed environment. In order to get the data from a particular neighboring partition, which could be remote, we must retrieve the data from all of the neighboring partitions. These retrievals are asynchronous and the function `heat_part_data`, which, amongst other things, calls `heat`, should not be called unless the data from the neighboring partitions have arrived. Therefore, it should come as no surprise that we synchronize this operation with another instance of `dataflow` (found in `heat_part`). This `dataflow` receives futures to the data in the current and surrounding partitions by calling `get_data()` on each respective partition. When these futures are ready, `dataflow` passes them to the unwrapping function, which extracts the `shared_array` of doubles and passes them to the lambda. The lambda calls `heat_part_data` on the *locality*, which the middle partition is on.

Although this example could run distributed, it only runs on one *locality*, as it always uses `hpx::find_here()` as the target for the functions to run on.

In example 6, we begin to distribute the partition data on different nodes. This is accomplished in `stepper::do_work()` by passing the GID of the *locality* where we wish to create the partition to the `partition` constructor.

```
// Initial conditions: f(0, i) = i
//[do_work_6]
```

We distribute the partitions evenly based on the number of localities used, which is described in the function `locidx`. Because some of the data needed to update the partition in `heat_part` could now be on a new *locality*, we must devise a way of moving data to the *locality* of the middle partition. We accomplished this by adding a switch in the function `get_data()` that returns the end element of the buffer `data_` if it is from the left partition or the first element of the buffer if the data is from the right partition. In this way only the necessary elements, not the whole buffer, are exchanged between nodes. The reader should be reminded that this exchange of end elements occurs in the function `get_data()` and, therefore, is executed asynchronously.

Now that we have the code running in distributed, it is time to make some optimizations. The function `heat_part` spends most of its time on two tasks: retrieving remote data and working on the data in the middle partition. Because we know that the data for the middle partition is local, we can overlap the work on the middle partition with that of the possibly remote call of `get_data()`. This algorithmic change, which was implemented in example 7, can be seen below:

```
//[stepper_7
// The partitioned operator, it invokes the heat operator above on all elements
// of a partition.
static partition heat_part(partition const& left,
                           partition const& middle, partition const& right)
{
    using hpx::dataflow;
    using hpx::util::unwrapping;

    hpx::shared_future<partition_data> middle_data =
        middle.get_data(partition_server::middle_partition);

    hpx::future<partition_data> next_middle = middle_data.then(
        unwrapping(
            [middle] (partition_data const& m) -> partition_data
            {
                HPX_UNUSED(middle);

                // All local operations are performed once the middle data of
                // the previous time step becomes available.
                std::size_t size = m.size();
                partition_data next(size);
                for (std::size_t i = 1; i != size-1; ++i)
                    next[i] = heat(m[i-1], m[i], m[i+1]);
                return next;
            }
        )
    );

    return dataflow(
        hpx::launch::async,
        unwrapping(
            [left, middle, right] (partition_data next, partition_data const& l,
                                   partition_data const& m, partition_data const& r) -> partition
            {
                HPX_UNUSED(left);
                HPX_UNUSED(right);

                // Calculate the missing boundary elements once the
                // corresponding data has become available.
            }
        )
    );
}
```

(continues on next page)

(continued from previous page)

```

    std::size_t size = m.size();
    next[0] = heat(l[size-1], m[0], m[1]);
    next[size-1] = heat(m[size-2], m[size-1], r[0]);

    // The new partition_data will be allocated on the same locality
    // as 'middle'.
    return partition(middle.get_id(), next);
}
),
std::move(next_middle),
left.get_data(partition_server::left_partition),
middle_data,
right.get_data(partition_server::right_partition)

```

Example 8 completes the futurization process and utilizes the full potential of *HPX* by distributing the program flow to multiple localities, usually defined as nodes in a cluster. It accomplishes this task by running an instance of *HPX* main on each *locality*. In order to coordinate the execution of the program, the `struct stepper` is wrapped into a component. In this way, each *locality* contains an instance of stepper that executes its own instance of the function `do_work()`. This scheme does create an interesting synchronization problem that must be solved. When the program flow was being coordinated on the head node, the GID of each component was known. However, when we distribute the program flow, each partition has no notion of the GID of its neighbor if the next partition is on another *locality*. In order to make the GIDs of neighboring partitions visible to each other, we created two buffers to store the GIDs of the remote neighboring partitions on the left and right respectively. These buffers are filled by sending the GID of newly created edge partitions to the right and left buffers of the neighboring localities.

In order to finish the simulation, the solution vectors named `result` are then gathered together on *locality* 0 and added into a vector of spaces `overall_result` using the *HPX* functions `gather_id` and `gather_here`.

Example 8 completes this example series, which takes the serial code of example 1 and incrementally morphs it into a fully distributed parallel code. This evolution was guided by the simple principles of futurization, the knowledge of grainsize, and utilization of components. Applying these techniques easily facilitates the scalable parallelization of most applications.

2.5 Manual

The manual is your comprehensive guide to *HPX*. It contains detailed information on how to build and use *HPX* in different scenarios.

2.5.1 Getting *HPX*

There are *HPX* packages available for a few Linux distributions. The easiest way to get started with *HPX* is to use those packages. We keep an up-to-date list with instructions on the [HPX Downloads⁵⁴](#) page. If you use one of the available packages you can skip the next section, [HPX build system](#), but we still recommend that you look through it as it contains useful information on how you can customize *HPX* at compile-time.

If there isn't a package available for your platform you should either clone our repository:

or download a package with the source files from [HPX Downloads⁵⁵](#).

⁵⁴ <https://stellar-group.org/downloads/>

⁵⁵ <https://stellar-group.org/downloads/>

2.5.2 HPX build system

The build system for *HPX* is based on [CMake](#)⁵⁶. CMake is a cross-platform build-generator tool. CMake does not build the project, it generates the files needed by your build tool (GNU make, Visual Studio, etc.) for building *HPX*.

This section gives an introduction on how to use our build system to build *HPX* and how to use *HPX* in your own projects.

CMake basics

CMake is a cross-platform build-generator tool. CMake does not build the project, it generates the files needed by your build tool (gnu make, visual studio, etc.) for building *HPX*.

In general, the *HPX* CMake scripts try to adhere to the general CMake policies on how to write CMake-based projects.

Basic CMake usage

This section explains basic aspects of CMake, specifically options needed for day-to-day usage.

CMake comes with extensive documentation in the form of html files and on the CMake executable itself. Execute `cmake --help` for further help options.

CMake needs to know which build tool it will generate files for (GNU make, Visual Studio, Xcode, etc.). If not specified on the command line, it will try to guess the build tool based on you environment. Once it has identified the build tool, CMake uses the corresponding generator to create files for your build tool. You can explicitly specify the generator with the command line option `-G "Name of the generator"`. To see the available generators on your platform, execute:

```
cmake --help
```

This will list the generator names at the end of the help text. Generator names are case-sensitive. Example:

```
cmake -G "Visual Studio 9 2008" path/to/hpx
```

For a given development platform there can be more than one adequate generator. If you use Visual Studio "NMake Makefiles" is a generator you can use for building with NMake. By default, CMake chooses the more specific generator supported by your development environment. If you want an alternative generator, you must tell this to CMake with the `-G` option.

Quick start

Here, you will use the command-line, non-interactive CMake interface.

1. Download and install CMake here: [CMake Downloads](#)⁵⁷. Version 3.3.2 is the minimally required version for *HPX*.
2. Open a shell. Your development tools must be reachable from this shell through the `PATH` environment variable.
3. Create a directory for containing the build. Building *HPX* on the source directory is not supported. `cd` to this directory:

```
mkdir mybuilddir
cd mybuilddir
```

⁵⁶ <https://www.cmake.org>

⁵⁷ <https://www.cmake.org/cmake/resources/software.html>

4. Execute this command on the shell replacing `path/to/hpx/` with the path to the root of your *HPX* source tree:

```
cmake path/to/hpx
```

CMake will detect your development environment, perform a series of tests and will generate the files required for building *HPX*. CMake will use default values for all build parameters. See the [CMake variables used to configure HPX](#) section for fine-tuning your build.

This can fail if CMake can't detect your toolset, or if it thinks that the environment is not sane enough. In this case make sure that the toolset that you intend to use is the only one reachable from the shell and that the shell itself is the correct one for your development environment. CMake will refuse to build MinGW makefiles if you have a POSIX shell reachable through the `PATH` environment variable, for instance. You can force CMake to use various compilers and tools. Please visit [CMake Useful Variables](#)⁵⁸ for a detailed overview of specific CMake variables.

Options and variables

Variables customize how the build will be generated. Options are boolean variables, with possible values ON/OFF. Options and variables are defined on the CMake command line like this:

```
cmake -DVARIABLE=value path/to/hpx
```

You can set a variable after the initial CMake invocation for changing its value. You can also undefine a variable:

```
cmake -UVARIABLE path/to/hpx
```

Variables are stored on the CMake cache. This is a file named `CMakeCache.txt` on the root of the build directory. Do not hand-edit it.

Variables are listed here appending its type after a colon. You should write the variable and the type on the CMake command line:

```
cmake -DVARIABLE:TYPE=value path/to/llvm/source
```

CMake supports the following variable types: `BOOL` (options), `STRING` (arbitrary string), `PATH` (directory name), `FILEPATH` (file name).

Prerequisites

Supported platforms

At this time, *HPX* supports the following platforms. Other platforms may work, but we do not test *HPX* with other platforms, so please be warned.

Table 2.1: Supported Platforms for *HPX*

Name	Recommended Version	Minimum Version	Architectures
Linux	3.2	2.6	x86-32, x86-64, k1om
BlueGeneQ	V1R2M0	V1R2M0	PowerPC A2
Windows	7, Server 2008 R2	Any Windows system	x86-32, x86-64
Mac OSX		Any OSX system	x86-64

⁵⁸ <https://gitlab.kitware.com/cmake/community/wikis/doc/cmake/Useful-Variables#Compilers-and-Tools>

Software and libraries

In the simplest case, *HPX* depends on [Boost⁵⁹](#) and [Portable Hardware Locality \(HWLOC\)⁶⁰](#). So, before you read further, please make sure you have a recent version of [Boost⁶¹](#) installed on your target machine. *HPX* currently requires at least Boost V1.61.0 to work properly. It may build and run with older versions, but we do not test *HPX* with those versions, so please be warned.

The installation of Boost is described in detail in Boost's [Getting Started⁶²](#) document. However, if you've never used the Boost libraries (or even if you have), here's a quick primer: [Installing Boost](#).

It is often possible to download the Boost libraries using the package manager of your distribution. Please refer to the corresponding documentation for your system for more information.

In addition, we require a recent version of hwloc in order to support thread pinning and NUMA awareness. See [Installing Hwloc](#) for instructions on building Portable Hardware Locality (HWLOC).

HPX is written in 99.99% Standard C++ (the remaining 0.01% is platform specific assembly code). As such, *HPX* is compilable with almost any standards compliant C++ compiler. A compiler supporting the C++11 Standard is highly recommended. The code base takes advantage of C++11 language features when available (move semantics, rvalue references, magic statics, etc.). This may speed up the execution of your code significantly. We currently support the following C++ compilers: GCC, MSVC, ICPC and clang. For the status of your favorite compiler with *HPX* visit [HPX Buildbot Website⁶³](#).

Table 2.2: Software prerequisites for *HPX* on Linux systems.

Name	Recommended version	Minimum version	Notes
Compilers			
GNU Compiler Collection (g++) ⁶⁴	4.9 or newer	4.9	
Intel Composer XE Suites ⁶⁵	2014 or newer	2014	
clang: a C language family frontend for LLVM ⁶⁶	3.8 or newer	3.8	
Build System			
CMake ⁶⁷	3.9.0	3.3.2	Cuda support 3.9
Required Libraries			
Boost C++ Libraries ⁶⁸	1.67.0 or newer	1.61.0	
Portable Hardware Locality (HWLOC) ⁶⁹	1.11	1.2 (Xeon Phi: 1.6)	

Note: When compiling with the Intel Compiler on Linux systems, we only support C++ Standard Libraries provided by gcc 4.8 and upwards. If the g++ in your path is older than 4.8, please specify the path of a newer g++ by setting CMAKE_CXX_FLAGS=' -cxx-name=/path/to/g++' via [CMake⁷⁰](#).

⁵⁹ <https://www.boost.org/>

⁶⁰ <https://www.open-mpi.org/projects/hwloc/>

⁶¹ <https://www.boost.org/>

⁶² https://www.boost.org/doc/libs/1_71_0/more/getting_started/index.html

⁶³ <http://rostam.cct.lsu.edu/>

⁶⁴ <https://gcc.gnu.org>

⁶⁵ <https://software.intel.com/en-us/intel-composer-xe/>

⁶⁶ <https://clang.llvm.org/>

⁶⁷ <https://www.cmake.org>

⁶⁸ <https://www.boost.org/>

⁶⁹ <https://www.open-mpi.org/projects/hwloc/>

⁷⁰ <https://www.cmake.org>

Note: When building Boost using gcc, please note that it is always a good idea to specify a `cxxflags=-std=c++11` command line argument to `b2` (`bjam`). Note, however, that this is absolutely necessary when using gcc V5.2 and above.

Table 2.3: Software prerequisites for *HPX* on Windows systems

Name	Recommended version	Minimum version	Notes
Compilers			
Visual C++ ⁷¹ (x64)	2015	2015	
Build System			
CMake ⁷²	3.9.0	3.3.2	
Required Libraries			
Boost ⁷³	1.67.0 or newer	1.61.0	
Portable Hardware Locality (HWLOC) ⁷⁴	1.11	1.5	

Note: You need to build the following Boost libraries for *HPX*: `Boost.Filesystem`, `Boost.ProgramOptions`, `Boost.Regex`, and `Boost.System`. The following are not needed by default, but are required in certain configurations: `Boost.Chrono`, `Boost.DateTime`, `Boost.Log`, `Boost.LogSetup`, and `Boost.Thread`.

Depending on the options you chose while building and installing *HPX*, you will find that *HPX* may depend on several other libraries such as those listed below.

Note: In order to use a high speed parcelport, we currently recommend configuring *HPX* to use MPI so that MPI can be used for communication between different localities. Please set the CMake variable `MPI_CXX_COMPILER` to your MPI C++ compiler wrapper if not detected automatically.

Table 2.4: Highly recommended optional software prerequisites for *HPX* on Linux systems

Name	Recommended version	Minimum version	Notes
google-perfetto ⁷⁵	1.7.1	1.7.1	Used as a replacement for the system allocator, and for allocation diagnostics.
libunwind ⁷⁶	0.99	0.97	Dependency of google-perfetto on x86-64, used for stack unwinding.
Open MPI ⁷⁷	1.10.1	1.8.0	Can be used as a highspeed communication library backend for the parcelport.

Note: When using OpenMPI please note that Ubuntu (notably 18.04 LTS) and older Debian ship an OpenMPI 2.x built with `--enable-heterogeneous` which may cause communication failures at runtime and should not be used.

⁷¹ <https://msdn.microsoft.com/en-us/visualc/default.aspx>

⁷² <https://www.cmake.org/>

⁷³ <https://www.boost.org/>

⁷⁴ <https://www.open-mpi.org/projects/hwloc/>

⁷⁵ <https://code.google.com/p/gperftools>

⁷⁶ <https://www.nongnu.org/libunwind>

⁷⁷ <https://www.open-mpi.org>

Table 2.5: Optional software prerequisites for *HPX* on Linux systems

Name	Recommended version	Minimum version	Notes
Performance Application Programming Interface (PAPI)			Used for accessing hardware performance data.
jemalloc ⁷⁸	2.1.2	2.1.0	Used as a replacement for the system allocator.
mi-malloc ⁷⁹	latest	1.0.0	Used as a replacement for the system allocator.
Hierarchical Data Format V5 (HDF5) ⁸⁰	1.8.7	1.6.7	Used for data I/O in some example applications. See important note below.

Table 2.6: Optional software prerequisites for *HPX* on Windows systems

Name	Recommended version	Minimum version	Notes
Hierarchical Data Format V5 (HDF5) ⁸¹	1.8.7	1.6.7	Used for data I/O in some example applications. See important note below.

Important: The C++ HDF5 libraries must be compiled with enabled thread safety support. This has to be explicitly specified while configuring the HDF5 libraries as it is not the default. Additionally, you must set the following environment variables before configuring the HDF5 libraries (this part only needs to be done on Linux):

```
export CFLAGS=' -DHDatexit=""'
export CPPFLAGS=' -DHDatexit=""'
```

Documentation

To build the *HPX* documentation, you need recent versions of the following packages:

- python (2 or 3)
- sphinx (Python package)
- sphinx_rtd_theme (Python package)
- breathe (Python package)
- doxygen

If the Python⁸² dependencies are not available through your system package manager, you can install them using the Python package manager pip:

```
pip install --user sphinx sphinx_rtd_theme breathe
```

You may need to set the following CMake variables to make sure CMake can find the required dependencies.

⁷⁸ <http://jemalloc.net>

⁷⁹ <http://microsoft.github.io/mimalloc/>

⁸⁰ <https://www.hdfgroup.org/HDF5>

⁸¹ <https://www.hdfgroup.org/HDF5>

⁸² <https://www.python.org>

DOXYGEN_ROOT:PATH

Specifies where to look for the installation of the Doxygen⁸³ tool.

SPHINX_ROOT:PATH

Specifies where to look for the installation of the Sphinx⁸⁴ tool.

BREATHE_APIDOC_ROOT:PATH

Specifies where to look for the installation of the Breathe⁸⁵ tool.

Installing Boost

Important: When building Boost using gcc, please note that it is always a good idea to specify a `cxxflags=-std=c++11` command line argument to `b2` (`bjam`). Doing so is absolutely necessary when using gcc V5.2 and above.

Important: On Windows, depending on the installed versions of Visual Studio, you might also want to pass the correct toolset to the `b2` command depending on which version of the IDE you want to use. In addition, passing `address-model=64` is highly recommended. It might also be necessary to add command line argument `--build-type=complete` to the `b2` command on the Windows platform.

The easiest way to create a working Boost installation is to compile Boost from sources yourself. This is particularly important as many high performance resources, even if they have Boost installed, usually only provide you with an older version of Boost. We suggest you download the most recent release of the Boost libraries from here: [Boost Downloads](#)⁸⁶. Unpack the downloaded archive into a directory of your choosing. We will refer to this directory a `$BOOST`.

Building and installing the Boost binaries is simple. Regardless of what platform you are on, the basic instructions are as follows (with possible additional platform-dependent command line arguments):

```
cd $BOOST
bootstrap --prefix=<where to install boost>
./b2 -j<N>
./b2 install
```

where: `<where to install boost>` is the directory the built binaries will be installed to, and `<N>` is the number of cores to use to build the Boost binaries.

After the above sequence of commands has been executed (this may take a while!), you will need to specify the directory where Boost was installed as `BOOST_ROOT` (`<where to install boost>`) while executing CMake for HPX as explained in detail in the sections [How to install HPX on Unix variants](#) and [How to install HPX on Windows](#).

⁸³ <https://www.doxygen.org>

⁸⁴ <http://www.sphinx-doc.org>

⁸⁵ <https://breathe.readthedocs.io/en/latest>

⁸⁶ <https://www.boost.org/users/download/>

Installing Hwloc

Note: These instructions are for everything except Windows. On Windows there is no need to build hwloc. Instead, download the latest release, extract the files, and set `HWLOC_ROOT` during CMake configuration to the directory in which you extracted the files.

We suggest you download the most recent release of hwloc from here: [Hwloc Downloads](#)⁸⁷. Unpack the downloaded archive into a directory of your choosing. We will refer to this directory as `$HWLOC`.

To build hwloc run:

```
cd $HWLOC
./configure --prefix=<where to install hwloc>
make -j<N> install
```

where: `<where to install hwloc>` is the directory the built binaries will be installed to, and `<N>` is the number of cores to use to build hwloc.

After the above sequence of commands has been executed, you will need to specify the directory where hwloc was installed as `HWLOC_ROOT` (`<where to install hwloc>`) while executing CMake for *HPX* as explained in detail in the sections [How to install HPX on Unix variants](#) and [How to install HPX on Windows](#).

Please see [Hwloc Documentation](#)⁸⁸ for more information about hwloc.

Building HPX

Basic information

Once CMake has been run, the build process can be started. The *HPX* build process is highly configurable through CMake, and various CMake variables influence the build process. The build process consists of the following parts:

- The *HPX* core libraries (target core): This forms the basic set of *HPX* libraries. The generated targets are:
 - `hpx`: The core *HPX* library (always enabled).
 - `hpx_init`: The *HPX* initialization library that applications need to link against to define the *HPX* entry points (disabled for static builds).
 - `hpx_wrap`: The *HPX* static library used to determine the runtime behavior of *HPX* code and respective entry points for `hpx_main.h`
 - `iostreams_component`: The component used for (distributed) IO (always enabled).
 - `component_storage_component`: The component needed for migration to persistent storage.
 - `unordered_component`: The component needed for a distributed (partitioned) hash table.
 - `partitioned_vector_component`: The component needed for a distributed (partitioned) vector.
 - `memory_component`: A dynamically loaded plugin that exposes memory based performance counters (only available on Linux).
 - `io_counter_component`: A dynamically loaded plugin that exposes I/O performance counters (only available on Linux).

⁸⁷ <https://www.open-mpi.org/software/hwloc/v1.11>

⁸⁸ <https://www.open-mpi.org/projects/hwloc/doc/>

- `papi_component`: A dynamically loaded plugin that exposes PAPI performance counters (enabled with `HPX_WITH_PAPI:BOOL`, default is Off).
- *HPX Examples* (target `examples`): This target is enabled by default and builds all *HPX* examples (disable by setting `HPX_WITH_EXAMPLES:BOOL=Off`). *HPX* examples are part of the `all` target and are included in the installation if enabled.
- *HPX Tests* (target `tests`): This target builds the *HPX* test suite and is enabled by default (disable by setting `HPX_WITH_TESTS:BOOL=Off`). They are not built by the `all` target and have to be built separately.
- *HPX Documentation* (target `docs`): This target builds the documentation, and is not enabled by default (enable by setting `HPX_WITH_DOCUMENTATION:BOOL=On`. For more information see *Documentation*.

For a complete list of available CMake variables that influence the build of *HPX*, see [CMake variables used to configure HPX](#).

The variables can be used to refine the recipes that can be found at [Platform specific build recipes](#) which show some basic steps on how to build *HPX* for a specific platform.

In order to use *HPX*, only the core libraries are required (the ones marked as optional above are truly optional). When building against *HPX*, the CMake variable `HPX_LIBRARIES` will contain `hpx` and `hpx_init` (for `pkgconfig`, those are added to the `Libs` sections). In order to use the optional libraries, you need to specify them as link dependencies in your build (See [Creating HPX projects](#)).

As *HPX* is a modern C++ library, we require a certain minimum set of features from the C++11 standard. In addition, we make use of certain C++14 features if the used compiler supports them. This means that the *HPX* build system will try to determine the highest support C++ standard flavor and check for availability of those features. That is, the default will be the highest C++ standard version available. If you want to force *HPX* to use a specific C++ standard version, you can use the following CMake variables:

- `HPX_WITH_CXX14`: Enables C++14 support (this is the minimum requirement)
- `HPX_WITH_CXX17`: Enables C++17 support
- `HPX_WITH_CXX2A`: Enables (experimental) C++20 support

Build types

CMake can be configured to generate project files suitable for builds that have enabled debugging support or for an optimized build (without debugging support). The CMake variable used to set the build type is `CMAKE_BUILD_TYPE` (for more information see the [CMake Documentation](#)⁸⁹). Available build types are:

- **Debug**: Full debug symbols are available as well as additional assertions to help debugging. To enable the debug build type for the *HPX* API, the C++ Macro `HPX_DEBUG` is defined.
- **RelWithDebInfo**: Release build with debugging symbols. This is most useful for profiling applications.
- **Release**: Release build. This disables assertions and enables default compiler optimizations.
- **RelMinSize**: Release build with optimizations for small binary sizes.

Important: We currently don't guarantee ABI compatibility between Debug and Release builds. Please make sure that applications built against *HPX* use the same build type as you used to build *HPX*. For CMake builds, this means that the `CMAKE_BUILD_TYPE` variables have to match and for projects not using [CMake](#)⁹⁰, the `HPX_DEBUG` macro has to be set in debug mode.

⁸⁹ https://cmake.org/cmake/help/latest/variable/CMAKE_BUILD_TYPE.html

⁹⁰ <https://www.cmake.org>

Platform specific notes

Some platforms require users to have special link and/or compiler flags specified to build *HPX*. This is handled via CMake's support for different toolchains (see [cmake-toolchains\(7\)](#)⁹¹ for more information). This is also used for cross compilation.

HPX ships with a set of toolchains that can be used for compilation of *HPX* itself and applications depending on *HPX*. Please see [CMake toolchains shipped with HPX](#) for more information.

In order to enable full static linking with the libraries, the CMake variable `HPX_WITH_STATIC_LINKING:BOOL` has to be set to `On`.

Debugging applications using core files

For *HPX* to generate useful core files, *HPX* has to be compiled without signal and exception handlers `HPX_WITH_DISABLED_SIGNAL_EXCEPTION_HANDLERS:BOOL`. If this option is not specified, the signal handlers change the application state. For example, after a segmentation fault the stack trace will show the signal handler. Similarly, unhandled exceptions are also caught by these handlers and the stack trace will not point to the location where the unhandled exception was thrown.

In general, core files are a helpful tool to inspect the state of the application at the moment of the crash (post-mortem debugging), without the need of attaching a debugger beforehand. This approach to debugging is especially useful if the error cannot be reliably reproduced, as only a single crashed application run is required to gain potentially helpful information like a stacktrace.

To debug with core files, the operating system first has to be told to actually write them. On most Unix systems this can be done by calling:

```
ulimit -c unlimited
```

in the shell. Now the debugger can be started up with:

```
gdb <application> <core file name>
```

The debugger should now display the last state of the application. The default file name for core files is `core`.

Platform specific build recipes

Note: The following build recipes are mostly user-contributed and may be outdated. We always welcome updated and new build recipes.

⁹¹ <https://cmake.org/cmake/help/latest/manual/cmake-toolchains.7.html>

How to install HPX on Unix variants

- Create a build directory. *HPX* requires an out-of-tree build. This means you will be unable to run CMake in the *HPX* source tree.

```
cd hpx
mkdir my_hpx_build
cd my_hpx_build
```

- Invoke CMake from your build directory, pointing the CMake driver to the root of your *HPX* source tree.

```
cmake -DBOOST_ROOT=/root/of/boost/installation \
-DHWLOC_ROOT=/root/of/hwloc/installation \
[other CMake variable definitions] \
/path/to/source/tree
```

For instance:

```
cmake -DBOOST_ROOT=~/packages/boost -DHWLOC_ROOT=~/packages/hwloc -DCMAKE_INSTALL_ \
PREFIX=~/packages/hpx ~/downloads/hpx_0.9.10
```

- Invoke GNU make. If you are on a machine with multiple cores, add the `-jN` flag to your make invocation, where N is the number of parallel processes *HPX* gets compiled with.

```
gmake -j4
```

Caution: Compiling and linking *HPX* needs a considerable amount of memory. It is advisable that at least 2 GB of memory per parallel process is available.

Note: Many Linux distributions use `make` as an alias for `gmake`.

- To complete the build and install *HPX*:

```
gmake install
```

Important: These commands will build and install the essential core components of *HPX* only. In order to build and run the tests, please invoke:

```
gmake tests && gmake test
```

and in order to build (and install) all examples invoke:

```
cmake -DHPX_WITH_EXAMPLES=On .
gmake examples
gmake install
```

For more detailed information about using CMake, please refer to its documentation and also the section [Building HPX](#). Please pay special attention to the section about `HPX_WITH_MALLOC:STRING` as this is crucial for getting decent performance.

How to install HPX on OS X (Mac)

This section describes how to build *HPX* for OS X (Mac).

Build (and install) a recent version of Boost, using Clang and libc++

To build Boost with Clang and make it link to libc++ as standard library, you'll need to set up either of the following in your `~/user-config.jam` file:

```
# user-config.jam (put this file into your home directory)
# ...

using clang
:
: "/usr/bin/clang++"
: <cxxflags>"-std=c++11 -fcolor-diagnostics"
<linkflags>"-stdlib=libc++ -L/path/to/libcxx/lib"
;
```

(Again, remember to replace `/path/to` with whatever you used earlier.)

Then, you can use one of the following for your build command:

```
b2 --build-dir=/tmp/build-boost --layout=versioned toolset=clang install -j4
```

or:

```
b2 --build-dir=/tmp/build-boost --layout=versioned toolset=clang install -j4
```

We verified this using Boost V1.53. If you use a different version, just remember to replace `/usr/local/include/boost-1_53` with whatever prefix you used in your installation.

Build HPX, finally

```
cd /path/to
git clone https://github.com/STELLAR-GROUP/hpx.git
mkdir build-hpx && cd build-hpx
```

To build with Clang 3.2, execute:

```
cmake ../hpx \
-DCMAKE_CXX_COMPILER=clang++ \
-DBOOST_INCLUDE_DIR=/usr/local/include/boost-1_53 \
-DBOOST_LIBRARY_DIR=/usr/local/lib \
-DBOOST_SUFFIX=-clang-darwin32-mt-1_53 \
make
```

To build with Clang 3.3 (trunk), execute:

```
cmake ../hpx \
-DCMAKE_CXX_COMPILER=clang++ \
-DBOOST_INCLUDE_DIR=/usr/local/include/boost-1_53 \
-DBOOST_LIBRARY_DIR=/usr/local/lib \
-DBOOST_SUFFIX=-clang-darwin33-mt-1_53 \
make
```

For more detailed information about using CMake, please refer its documentation and to the section *Building HPX*.

Alternative installation method of *HPX* on OS X (Mac)

Alternatively, you can install a recent version of gcc as well as all required libraries via MacPorts:

1. Install MacPorts
2. Install CMake, gcc 4.8, and hwloc:

```
sudo port install gcc48
sudo port install hwloc
```

You may also want:

```
sudo port install cmake
sudo port install git-core
```

3. Make this version of gcc your default compiler:

```
sudo port install gcc_select
sudo port select gcc mp-gcc48
```

4. Build Boost manually (the Boost package of MacPorts is built with Clang, and unfortunately doesn't work with a GCC-build version of *HPX*):

```
wget https://dl.bintray.com/boostorg/release/1.69.0/source/boost_1_69_0.tar.bz2
tar xjf boost_1_69_0.tar.bz2
pushd boost_1_69_0
export BOOST_ROOT=$HOME/boost_1_69_0
./bootstrap.sh --prefix=$BOOST_DIR
./b2 -j8
./b2 -j8 install
export DYLD_LIBRARY_PATH=$DYLD_LIBRARY_PATH:$BOOST_ROOT/lib
popd
```

5. Build *HPX*:

```
git clone https://github.com/STELLAR-GROUP/hpx.git
mkdir hpx-build
pushd hpx-build
export HPX_ROOT=$HOME/hpx
cmake -DCMAKE_C_COMPILER=gcc \
      -DCMAKE_CXX_COMPILER=g++ \
      -DCMAKE_FORTRAN_COMPILER=gfortran \
      -DCMAKE_C_FLAGS="-Wno-unused-local-typedefs" \
      -DCMAKE_CXX_FLAGS="-Wno-unused-local-typedefs" \
      -DBOOST_ROOT=$BOOST_ROOT \
      -DHWLOC_ROOT=/opt/local \
      -DCMAKE_INSTALL_PREFIX=$HOME/hpx \
      $(pwd)/../hpx
make -j8
make -j8 install
export DYLD_LIBRARY_PATH=$DYLD_LIBRARY_PATH:$HDX_ROOT/lib/hpx
popd
```

6. Note that you need to set BOOST_ROOT, HPX_ROOT and DYLD_LIBRARY_PATH (for both BOOST_ROOT and HPX_ROOT) every time you configure, build, or run an *HPX* application.

7. If you want to use *HPX* with MPI, you need to enable the MPI parcelport, and also specify the location of the MPI wrapper scripts. This can be done using the following command:

```
cmake -DHPX_WITH_PARCELPORT_MPI=ON \
-DCMAKE_C_COMPILER=gcc \
-DCMAKE_CXX_COMPILER=g++ \
-DCMAKE_FORTRAN_COMPILER=gfortran \
-DMPI_C_COMPILER=openmpicc \
-DMPI_CXX_COMPILER=openmpic++ \
-DMPI_FORTRAN_COMPILER=openmpif90 \
-DCMAKE_C_FLAGS="-Wno-unused-local-typedefs" \
-DCMAKE_CXX_FLAGS="-Wno-unused-local-typedefs" \
-DBOOST_ROOT=$BOOST_DIR \
-DHWLOC_ROOT=/opt/local \
-DCMAKE_INSTALL_PREFIX=$HOME/hpx
$(pwd) / ../hpx
```

How to install *HPX* on Windows

Installation of required prerequisites

- Download the Boost c++ libraries from Boost Downloads⁹²
- Install the Boost library as explained in the section *Installing Boost*
- Install the hwloc library as explained in the section *Installing Hwloc*
- Download the latest version of CMake binaries, which are located under the platform section of the downloads page at [CMake Downloads](#)⁹³.
- Download the latest version of *HPX* from the STELLAR website: [HPX Downloads](#)⁹⁴.

Installation of the *HPX* library

- Create a build folder. *HPX* requires an out-of-tree-build. This means that you will be unable to run CMake in the *HPX* source folder.
- Open up the CMake GUI. In the input box labelled “Where is the source code:”, enter the full path to the source folder. The source directory is the one where the sources were checked out. CMakeLists.txt files in the source directory as well as the subdirectories describe the build to CMake. In addition to this, there are CMake scripts (usually ending in .cmake) stored in a special CMake directory. CMake does not alter any file in the source directory and doesn’t add new ones either. In the input box labelled “Where to build the binaries:”, enter the full path to the build folder you created before. The build directory is one where all compiler outputs are stored, which includes object files and final executables.
- Add CMake variable definitions (if any) by clicking the “Add Entry” button. There are two required variables you need to define: BOOST_ROOT and HWLOC_ROOT These (PATH) variables need to be set to point to the root folder of your Boost and hwloc installations. It is recommended to set the variable CMAKE_INSTALL_PREFIX as well. This determines where the *HPX* libraries will be built and installed. If this (PATH) variable is set, it has to refer to the directory where the built *HPX* files should be installed to.
- Press the “Configure” button. A window will pop up asking you which compilers to use. Select the Visual Studio 10 (64Bit) compiler (it usually is the default if available). The Visual Studio 2012 (64Bit) and Visual

⁹² <https://www.boost.org/users/download/>

⁹³ <https://www.cmake.org/cmake/resources/software.html>

⁹⁴ <https://stellar-group.org/downloads/>

Studio 2013 (64Bit) compilers are supported as well. Note that while it is possible to build *HPX* for x86, we don't recommend doing so as 32 bit runs are severely restricted by a 32 bit Windows system limitation affecting the number of *HPX* threads you can create.

- Press “Configure” again. Repeat this step until the “Generate” button becomes clickable (and until no variable definitions are marked in red anymore).
- Press “Generate”.
- Open up the build folder, and double-click hpx.sln.
- Build the INSTALL target.

For more detailed information about using CMake⁹⁵ please refer its documentation and also the section *Building HPX*.

How to build *HPX* under Windows 10 x64 with Visual Studio 2015

- Download the CMake V3.4.3 installer (or latest version) from [here](#)⁹⁶
- Download the hwloc V1.11.0 (or the latest version) from [here](#)⁹⁷ and unpack it.
- Download the latest Boost libraries from [here](#)⁹⁸ and unpack them.
- Build the Boost DLLs and LIBs by using these commands from Command Line (or PowerShell). Open CMD/PowerShell inside the Boost dir and type in:

```
bootstrap.bat
```

This batch file will set up everything needed to create a successful build. Now execute:

```
b2.exe link=shared variant=release,debug architecture=x86 address-model=64  
→threading=multi --build-type=complete install
```

This command will start a (very long) build of all available Boost libraries. Please, be patient.

- Open CMake-GUI.exe and set up your source directory (input field ‘Where is the source code’) to the *base directory* of the source code you downloaded from *HPX*'s GitHub pages. Here's an example of CMake path settings, which point to the Documents/GitHub/hpx folder:

Inside ‘Where is the source-code’ enter the base directory of your *HPX* source directory (do not enter the “src” sub-directory!). Inside ‘Where to build the binaries’ you should put in the path where all the building processes will happen. This is important because the building machinery will do an “out-of-tree” build. CMake will not touch or change the original source files in any way. Instead, it will generate Visual Studio Solution Files, which will build *HPX* packages out of the *HPX* source tree.

- Set three new environment variables (in CMake, not in Windows environment): BOOST_ROOT, HWLOC_ROOT, CMAKE_INSTALL_PREFIX. The meaning of these variables is as follows:
 - BOOST_ROOT the *HPX* root directory of the unpacked Boost headers/cpp files.
 - HWLOC_ROOT the *HPX* root directory of the unpacked Portable Hardware Locality files.
 - CMAKE_INSTALL_PREFIX the *HPX* root directory where the future builds of *HPX* should be installed.

⁹⁵ <https://www.cmake.org>

⁹⁶ <https://blog.kitware.com/cmake-3-4-3-available-for-download/>

⁹⁷ <http://www.open-mpi.org/software/hwloc/v1.11/downloads/hwloc-win64-build-1.11.0.zip>

⁹⁸ <https://www.boost.org/users/download/>

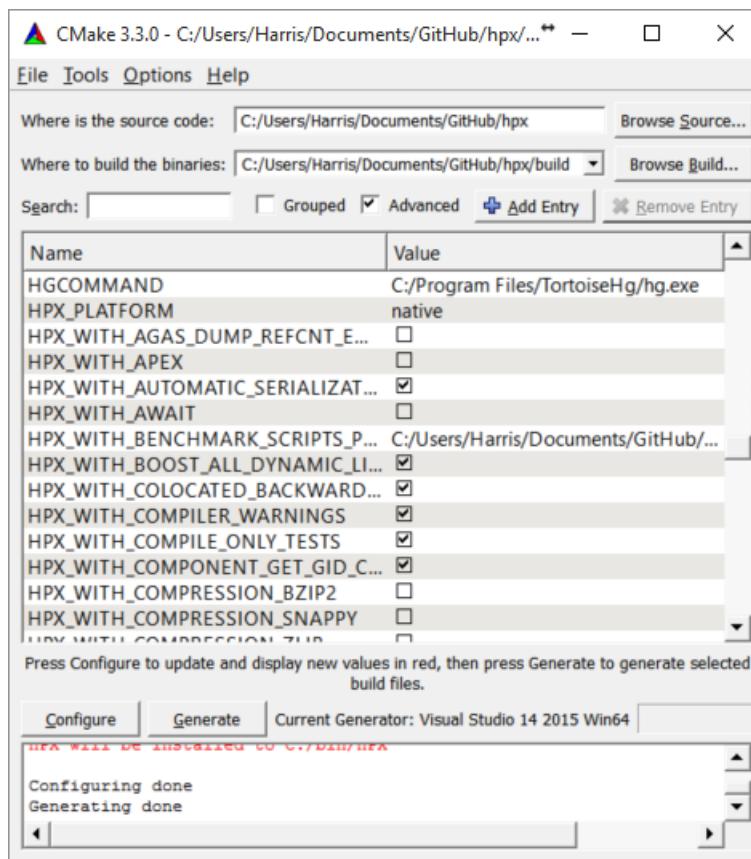


Fig. 2.3: Example CMake path settings.

Note: HPX is a very large software collection, so it is not recommended to use the default C:\Program Files\hpx. Many users may prefer to use simpler paths *without* whitespace, like C:\bin\hpx or D:\bin\hpx etc.

To insert new env-vars click on “Add Entry” and then insert the name inside “Name”, select PATH as Type and put the path-name in the “Path” text field. Repeat this for the first three variables.

This is how variable insertion will look:

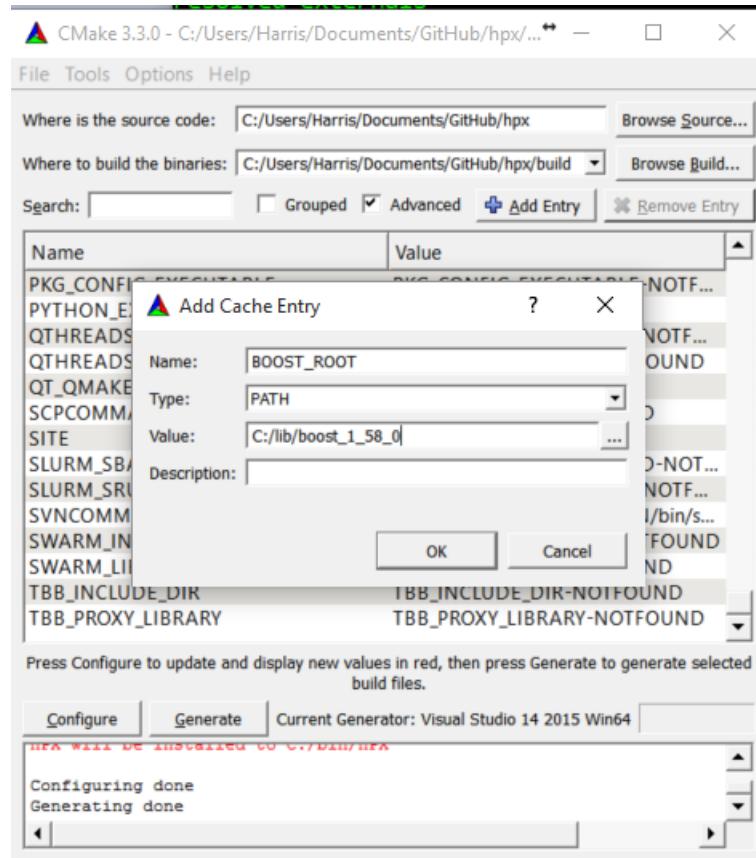


Fig. 2.4: Example CMake adding entry.

Alternatively, users could provide BOOST_LIBRARYDIR instead of BOOST_ROOT; the difference is that BOOST_LIBRARYDIR should point to the subdirectory inside Boost root where all the compiled DLLs/LIBs are. For example,

BOOST_LIBRARYDIR may point to the `bin.v2` subdirectory under the Boost rootdir. It is important to keep the meanings of these two variables separated from each other:

BOOST_DIR points to the ROOT folder of the Boost library. **BOOST_LIBRARYDIR** points to the subdir inside the Boost root folder where the compiled binaries are.

- Click the ‘Configure’ button of CMake-GUI. You will be immediately presented with a small window where you can select the C++ compiler to be used within Visual Studio. This has been tested using the latest v14 (a.k.a C++ 2015) but older versions should be sufficient too. Make sure to select the 64Bit compiler.
- After the generate process has finished successfully, click the ‘Generate’ button. Now, CMake will put new VS Solution files into the BUILD folder you selected at the beginning.

- Open Visual Studio and load the `HPX.sln` from your build folder.
- Go to `CMakePredefinedTargets` and build the `INSTALL` project:

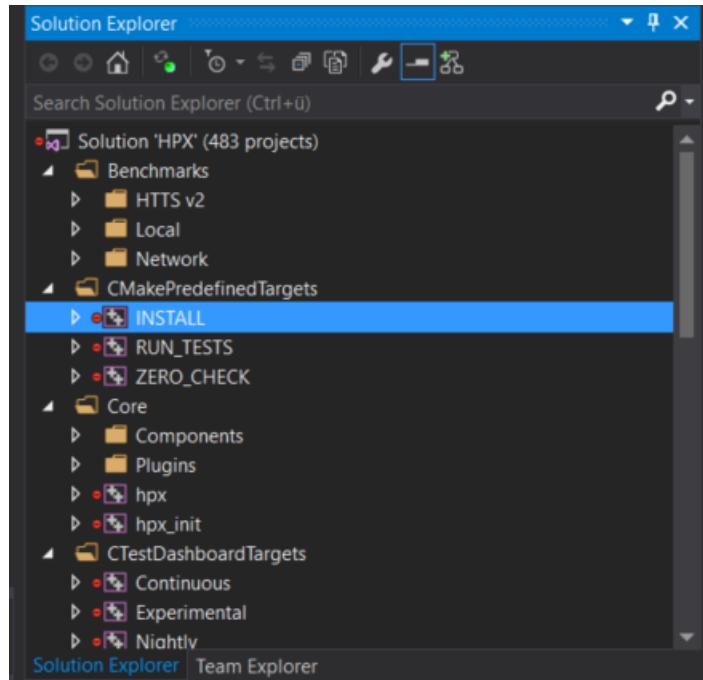


Fig. 2.5: Visual Studio INSTALL target.

It will take some time to compile everything, and in the end you should see an output similar to this one:

How to Install HPX on BlueGene/Q

So far we only support BGClang for compiling *HPX* on the BlueGene/Q.

- Check if BGClang is available on your installation. If not, obtain and install a copy from the [BGClang trac](#)⁹⁹.
- Build (and install) a recent version of [Hwloc Downloads](#)¹⁰⁰. With the following commands:

```
./configure \
--host=powerpc64-bgq-linux \
--prefix=$HOME/install/hwloc \
--disable-shared \
--enable-static \
CPPFLAGS='-I/bgsys/drivers/ppcfloor -I/bgsys/drivers/ppcfloor/spi/include/' \
kernel/cnk'
make
make install
```

- Build (and install) a recent version of Boost, using BGClang. To build Boost with BGClang, you'll need to set up the following in your Boost `~/user-config.jam` file:

⁹⁹ <https://trac.alcf.anl.gov/projects/llvm-bgq>

¹⁰⁰ <https://www.open-mpi.org/software/hwloc/v1.11>

```

Output
Show output from: Build
116> -- Installing: C:/bin/HPX/bin/1d_stencil_2.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_3.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_4.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_4_parallel.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_5.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_6.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_7.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_8.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_1_omp.exe
116> -- Installing: C:/bin/HPX/bin/1d_stencil_3_omp.exe
116> -- Installing: C:/bin/HPX/bin/simple_central_tuplespace_client.exe
116> -- Installing: C:/bin/HPX/lib/hpx_simple_central_tuplespaced.lib
116> -- Installing: C:/bin/HPX/lib/hpx_simple_central_tuplespaced.dll
116> -- Installing: C:/bin/HPX/bintranspose_serial.exe
116> -- Installing: C:/bin/HPX/bintranspose_serial_block.exe
116> -- Installing: C:/bin/HPX/bintranspose_smp.exe
116> -- Installing: C:/bin/HPX/bintranspose_smp_block.exe
116> -- Installing: C:/bin/HPX/bintranspose_block.exe
116> -- Installing: C:/bin/HPX/bintranspose_serial_vector.exe
116> -- Installing: C:/bin/HPX/binhpx_runtime.exe
===== Build: 116 succeeded, 0 failed, 0 up-to-date, 0 skipped =====

```

Error List | **Output** | Find Symbol Results | Package Manager Console | Azure App Service Activity

Fig. 2.6: Visual Studio build output.

```
# user-config.jam (put this file into your home directory)
using clang
:
: bgclang++11
:
;
```

You can then use this as your build command:

```
./bootstrap.sh
./b2 --build-dir=/tmp/build-boost --layout=versioned toolset=clang -j12
```

- Clone the master *HPX* git repository (or a stable tag):

```
git clone git://github.com/STELLAR-GROUP/hpx.git
```

- Generate the *HPX* buildfiles using CMake:

```
cmake -DHDX_PLATFORM=BlueGeneQ \
-DCLMAKE_TOOLCHAIN_FILE=/path/to/hpx/cmake/toolchains/BGQ.cmake \
-DCLMAKE_CXX_COMPILER=bgclang++11 \
-DMPICH_CXX_COMPILER=mpiclang++11 \
-DHWLOC_ROOT=/path/to/hwloc/installation \
-DBOOST_ROOT=/path/to/boost \
-DHDX_WITH_MALLOC=system \
/path/to/hpx
```

- To complete the build and install *HPX*:

```
make -j24
make install
```

This will build and install the essential core components of *HPX* only. Use:

```
make -j24 examples
make -j24 install
```

to build and install the examples.

How to Install *HPX* on the Xeon Phi

Installation of the Boost Libraries

- Download Boost Downloads¹⁰¹ for Linux and unpack the retrieved tarball.
- Adapt your `~/user-config.jam` to contain the following lines:

```
## Toolset to be used for compiling for the host
using intel
    : host
    :
    : <cxxflags>"-std=c++0x"
;

## Toolset to be used for compiling for the Xeon Phi
using intel
    : mic
    :
    : <cxxflags>"-std=c++0x -mmic"
    <linkflags>"-std=c++0x -mmic"
;
```

- Change to the directory you unpacked Boost in (which will be referred to as `$BOOST_ROOT` from now on) and execute the following commands:

```
./bootstrap.sh
./b2 toolset=intel-mic -j<N>
```

You should now have all the required Boost libraries.

Installation of the Hwloc library

- Download Hwloc Downloads¹⁰², unpack the retrieved tarball and change to the newly created directory.
- Run the configure-make-install procedure as follows:

```
CC=icc CFLAGS=-mmic CXX=icpc CXXFLAGS=-mmic LDFLAGS=-mmic ./configure --host=x86_64-k1om-linux --prefix=$HWLOC_ROOT
make
make install
```

¹⁰¹ <https://www.boost.org/users/download/>

¹⁰² <https://www.open-mpi.org/software/hwloc/v1.11>

Important: The minimally required version of the Portable Hardware Locality (HWLOC) library on the Intel Xeon Phi is V1.6.

You now have a working hwloc installation in \$HWLOC_ROOT.

Building HPX

After all the prerequisites have been successfully installed, we can now start building and installing *HPX*. The build procedure is almost the same as the one for [How to install HPX on Unix variants](#) with the sole difference that you have to enable the Xeon Phi in the CMake Build system. This is achieved by invoking CMake in the following way:

```
cmake
  -DCMAKE_TOOLCHAIN_FILE=/path/to/hpx/cmake/toolchains/XeonPhi.cmake \
  -DBOOST_ROOT=$BOOST_ROOT
  -DHWLOC_ROOT=$HWLOC_ROOT
  /path/to/hpx
```

For more detailed information about using CMake, please refer to its documentation and to the section [Building HPX](#). Please pay special attention to the section about [HPX_WITH_MALLOC: STRING](#) as this is crucial for getting decent performance on the Xeon Phi.

How to install HPX on Fedora distributions

Important: There are official *HPX* packages for Fedora. Unless you want to customize your build you may want to start off with the official packages. Instructions can be found on the [HPX Downloads](#)¹⁰³ page.

Note: This section of the manual is based off of our collaborator Patrick Diehl's blog post [Installing hpxl on Fedora 22](#)¹⁰⁴.

- Install all packages for minimal installation:

```
sudo dnf install gcc-c++ cmake boost-build boost boost-devel hwloc-devel \
  hwloc gcc-gfortran papi-devel gperf-tools-devel docbook-dtds \
  docbook-style-xsl libsodium-devel doxygen boost-doc hdf5-devel \
  fop boost-devel boost-openmpi-devel boost-mpich-devel
```

- Get the development branch of *HPX*:

```
git clone https://github.com/STELLAR-GROUP/hpx.git
```

- Configure it with CMake:

```
cd hpx
mkdir build
cd build
cmake -DCMAKE_INSTALL_PREFIX=/opt/hpx ..
```

(continues on next page)

¹⁰³ <https://stellar-group.org/downloads/>

¹⁰⁴ <http://diehlpk.github.io/2015/08/04/hpx-fedora.html>

(continued from previous page)

```
make -j
make install
```

Note: To build HPX without examples use:

```
cmake -DCMAKE_INSTALL_PREFIX=/opt/hpx -DHPX_WITH_EXAMPLES=Off ..
```

- Add the library path of HPX to ldconfig:

```
sudo echo /opt/hpx/lib > /etc/ld.so.conf.d/hpx.conf
sudo ldconfig
```

How to install HPX on Arch distributions

Important: There are HPX packages for Arch in the AUR. Unless you want to customize your build, you may want to start off with those. Instructions can be found on the [HPX Downloads](#)¹⁰⁵ page.

- Install all packages for a minimal installation:

```
sudo pacman -S gcc clang cmake boost hwloc gperftools
```

- For building the documentation, you will need to further install the following:

```
sudo pacman -S doxygen python-pip
pip install --user sphinx sphinx_rtd_theme breathe
```

The rest of the installation steps are the same as those for the Fedora or Unix variants.

How to install HPX on Debian-based distributions

- Install all packages for a minimal installation:

```
sudo apt install cmake libboost-all-dev hwloc libgoogle-perf-tools-dev
```

- To build the documentation you will need to further install the following:

```
sudo apt install doxygen python-pip
pip install --user sphinx sphinx_rtd_theme breathe
```

or the following if you prefer to get Python packages from the Debian repositories:

```
sudo apt install doxygen python-sphinx python-sphinx-rtd-theme python-breathe
```

The rest of the installation steps are same as those for the Fedora or Unix variants.

¹⁰⁵ <https://stellar-group.org/downloads/>

CMake toolchains shipped with HPX

In order to compile *HPX* for various platforms, we provide a variety of toolchain files that take care of setting up various CMake variables like compilers, etc. They are located in the `cmake/toolchains` directory:

- *ARM-gcc*
- *BGION-gcc*
- *BGQ*
- *Cray*
- *CrayKNL*
- *CrayKNLStatic*
- *CrayStatic*
- *XeonPhi*

To use them, pass the `-DCMAKE_TOOLCHAIN_FILE=<toolchain>` argument to the CMake invocation.

ARM-gcc

```
# Copyright (c) 2015 Thomas Heller
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
set(CMAKE_SYSTEM_NAME Linux)
set(CMAKE_CROSSCOMPILING ON)
# Set the gcc Compiler
set(CMAKE_CXX_COMPILER arm-linux-gnueabihf-g++-4.8)
set(CMAKE_C_COMPILER arm-linux-gnueabihf-gcc-4.8)
set(HPX_WITH_GENERIC_CONTEXT_COROUTINES ON CACHE BOOL "enable generic coroutines")
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM NEVER)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
```

BGION-gcc

```
# Copyright (c) 2014 John Biddiscombe
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
# This is the default toolchain file to be used with CNK on a BlueGene/Q. It sets
# the appropriate compile flags and compiler such that HPX will compile.
# Note that you still need to provide Boost, hwloc and other utility libraries
# like a custom allocator yourself.
#
# Usage : cmake -DCMAKE_TOOLCHAIN_FILE=~/src/hpx/cmake/toolchains/BGION-gcc.cmake ~/src/hpx
#
set(CMAKE_SYSTEM_NAME Linux)
```

(continues on next page)

(continued from previous page)

```

# Set the gcc Compiler
set(CMAKE_CXX_COMPILER g++)
set(CMAKE_C_COMPILER gcc)
#set(CMAKE_Fortran_COMPILER)
# Add flags we need for BGAS compilation
set(CMAKE_CXX_FLAGS_INIT
    "-D__powerpc__ -D__bgion__ -I/gpfs/bbp.cscs.ch/home/biddisco/src/bgas/rdmahelper "
    CACHE STRING "Initial compiler flags used to compile for BGAS"
)
# the V1R2M2 includes are necessary for some hardware specific features
#-DHPX_SMALL_STACK_SIZE=0x200000 -DHPX_MEDIUM_STACK_SIZE=0x200000 -DHPX_LARGE_STACK_
#-SIZE=0x200000 -DHPX_HUGE_STACK_SIZE=0x200000
set(CMAKE_EXE_LINKER_FLAGS_INIT "-L/gpfs/bbp.cscs.ch/apps/bgas/tools/gcc/gcc-4.8.2/
    ↪install/lib64 -latomic -lrt" CACHE STRING "BGAS flags")
set(CMAKE_C_FLAGS_INIT "-D__powerpc__ -I/gpfs/bbp.cscs.ch/home/biddisco/src/bgas/
    ↪rdmahelper" CACHE STRING "BGAS flags")
# We do not perform cross compilation here ...
set(CMAKE_CROSSCOMPILING OFF)
# Set our platform name
set(HPX_PLATFORM "native")
# Disable generic coroutines (and use posix version)
set(HPX_WITH_GENERIC_CONTEXT_COROUTINES OFF CACHE BOOL "disable generic coroutines")
# BGAS nodes support ibverbs, but it is deprecated
set(HPX_WITH_PARCELPORT_VERBS OFF CACHE BOOL "")
# Always disable the tcp parcelport as it is non-functional on the BGQ.
set(HPX_WITH_PARCELPORT_TCP ON CACHE BOOL "")
# Always enable the tcp parcelport as it is currently the only way to communicate on
# the BGQ.
set(HPX_WITH_PARCELPORT_MPI ON CACHE BOOL "")
# We have a bunch of cores on the A2 processor ...
set(HPX_WITH_MAX_CPU_COUNT "64" CACHE STRING "")
# We have no custom malloc yet
if(NOT DEFINED HPX_WITH_MALLOC)
    set(HPX_WITH_MALLOC "system" CACHE STRING "")
endif()
set(HPX_HIDDEN_VISIBILITY OFF CACHE BOOL "")
#
# Convenience setup for jb @ bbpbg2.cscs.ch
#
set(BOOST_ROOT "/gpfs/bbp.cscs.ch/home/biddisco/apps/gcc-4.8.2/boost_1_56_0")
set(HWLOC_ROOT "/gpfs/bbp.cscs.ch/home/biddisco/apps/gcc-4.8.2/hwloc-1.8.1")
set(CMAKE_BUILD_TYPE "Debug" CACHE STRING "Default build")
#
# Testing flags
#
set(BUILD_TESTING           ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS           ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS_BENCHMARKS ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS_REGRESSIONS ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS_UNIT      ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS_EXAMPLES  ON   CACHE BOOL "Testing enabled by default")
set(HPX_WITH_TESTS_EXTERNAL_BUILD OFF CACHE BOOL "Turn off build of cmake build tests
    ↪")
set(DART_TESTING_TIMEOUT     45   CACHE STRING "Life is too short")
# HPX_WITH_STATIC_LINKING

```

BGQ

```
# Copyright (c) 2014 Thomas Heller
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
#
# This is the default toolchain file to be used with CNK on a BlueGene/Q. It sets
# the appropriate compile flags and compiler such that HPX will compile.
# Note that you still need to provide Boost, hwloc and other utility libraries
# like a custom allocator yourself.
#
set(CMAKE_SYSTEM_NAME Linux)
# Set the Intel Compiler
set(CMAKE_CXX_COMPILER bgclang++11)
set(CMAKE_C_COMPILER bgclang)
#set(CMAKE_Fortran_COMPILER)
set(MPI_CXX_COMPILER mpiclang++11)
set(MPI_C_COMPILER mpiclang)
#set(MPI_Fortran_COMPILER)
set(CMAKE_C_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_C_COMPILE_OBJECT "<CMAKE_C_COMPILER> -fPIC <DEFINES> <FLAGS> -o <OBJECT> -c
-><SOURCE>" CACHE STRING "")
set(CMAKE_C_LINK_EXECUTABLE "<CMAKE_C_COMPILER> -fPIC -dynamic <FLAGS> <CMAKE_C_LINK_
->FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_C_CREATE_SHARED_LIBRARY "<CMAKE_C_COMPILER> -fPIC -shared <CMAKE_SHARED_
->LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_
->CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET> <OBJECTS> <LINK_
->LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_CXX_COMPILE_OBJECT "<CMAKE_CXX_COMPILER> -fPIC <DEFINES> <FLAGS> -o <OBJECT>
->-c <SOURCE>" CACHE STRING "")
set(CMAKE_CXX_LINK_EXECUTABLE "<CMAKE_CXX_COMPILER> -fPIC -dynamic <FLAGS> <CMAKE_CXX_
->LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_CREATE_SHARED_LIBRARY "<CMAKE_CXX_COMPILER> -fPIC -shared <CMAKE_SHARED_
->LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_
->CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET> <OBJECTS> <LINK_
->LIBRARIES>" CACHE STRING "")
set(CMAKE_Fortran_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_Fortran_COMPILE_OBJECT "<CMAKE_Fortran_COMPILER> -fPIC <DEFINES> <FLAGS> -o
-><OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_Fortran_LINK_EXECUTABLE "<CMAKE_Fortran_COMPILER> -fPIC -dynamic <FLAGS>
-><CMAKE_Fortran_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>")
set(CMAKE_Fortran_CREATE_SHARED_LIBRARY "<CMAKE_Fortran_COMPILER> -fPIC -shared
-><CMAKE_SHARED_LIBRARY_Fortran_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_
->SHARED_LIBRARY_CREATE_Fortran_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET>
-><OBJECTS> <LINK_LIBRARIES>" CACHE STRING ")
# Disable searches in the default system paths. We are cross compiling after all
# and cmake might pick up wrong libraries that way
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM BOTH)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
# We do a cross compilation here ...
set(CMAKE_CROSSCOMPILING ON)
# Set our platform name
```

(continues on next page)

(continued from previous page)

```

set(HPX_PLATFORM "BlueGeneQ")
# Always disable the ibverbs parcelport as it is non-functional on the BGQ.
set(HPX_WITH_PARCELPORT_VERBS OFF)
# Always disable the tcp parcelport as it is non-functional on the BGQ.
set(HPX_WITH_PARCELPORT_TCP OFF)
# Always enable the mpi parcelport as it is currently the only way to communicate on_
# the BGQ.
set(HPX_WITH_PARCELPORT_MPI ON)
# We have a bunch of cores on the BGQ ...
set(HPX_WITH_MAX_CPU_COUNT "64")
# We default to tbbmalloc as our allocator on the MIC
if(NOT DEFINED HPX_WITH_MALLOC)
    set(HPX_WITH_MALLOC "system" CACHE STRING "")
endif()

```

Cray

```

# Copyright (c) 2014 Thomas Heller
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
#
# This is the default toolchain file to be used with Intel Xeon PHIs. It sets
# the appropriate compile flags and compiler such that HPX will compile.
# Note that you still need to provide Boost, hwloc and other utility libraries
# like a custom allocator yourself.
#
#set(CMAKE_SYSTEM_NAME Cray-CNK-Intel)
if(HPX_WITH_STATIC_LINKING)
    set_property(GLOBAL PROPERTY TARGET_SUPPORTS_SHARED_LIBS FALSE)
else()
endif()
# Set the Cray Compiler Wrapper
set(CMAKE_CXX_COMPILER CC)
set(CMAKE_C_COMPILER cc)
set(CMAKE_Fortran_COMPILER ftn)
if (CMAKE_VERSION VERSION_GREATER 3.3.9)
    set(__includes "<INCLUDES>")
endif()
set(CMAKE_C_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_C_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_C_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_C_COMPILE_OBJECT "<CMAKE_C_COMPILER> -shared -fPIC <DEFINES> ${__includes}"
    "<FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_C_LINK_EXECUTABLE "<CMAKE_C_COMPILER> -fPIC -dynamic <FLAGS> <CMAKE_C_LINK_
    <FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_C_CREATE_SHARED_LIBRARY "<CMAKE_C_COMPILER> -fPIC -shared <CMAKE_SHARED_
    <LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_
    <CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET> <OBJECTS> <LINK_
    <LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CXX_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS "-fPIC -shared" CACHE STRING "")

```

(continues on next page)

(continued from previous page)

```

set(CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_CXX_COMPILE_OBJECT "<CMAKE_CXX_COMPILER> -shared -fPIC <DEFINES> ${__  
↳includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_CXX_LINK_EXECUTABLE "<CMAKE_CXX_COMPILER> -fPIC -dynamic <FLAGS> <CMAKE_CXX_  
↳LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_CREATE_SHARED_LIBRARY "<CMAKE_CXX_COMPILER> -fPIC -shared <CMAKE_SHARED_  
↳LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_  
↳CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET_SONAME> -o <TARGET> <OBJECTS> <LINK_  
↳LIBRARIES>" CACHE STRING "")
set(CMAKE_Fortran_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_Fortran_FLAGS "-fPIC" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_Fortran_FLAGS "-shared" CACHE STRING "")
set(CMAKE_Fortran_COMPILE_OBJECT "<CMAKE_Fortran_COMPILER> -shared -fPIC <DEFINES> ${__  
↳includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_Fortran_LINK_EXECUTABLE "<CMAKE_Fortran_COMPILER> -fPIC -dynamic <FLAGS>  
↳<CMAKE_Fortran_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>")
set(CMAKE_Fortran_CREATE_SHARED_LIBRARY "<CMAKE_Fortran_COMPILER> -fPIC -shared  
↳<CMAKE_SHARED_LIBRARY_Fortran_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_  
↳SHARED_LIBRARY_CREATE_Fortran_FLAGS> <SONAME_FLAG><TARGET_SONAME> -o <TARGET>  
↳<OBJECTS> <LINK_LIBRARIES>" CACHE STRING "")
# Disable searches in the default system paths. We are cross compiling after all
# and cmake might pick up wrong libraries that way
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM BOTH)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
set(HPX_WITH_PARCELPORT_TCP ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_MPI ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_MPI_MULTITHREADED OFF CACHE BOOL "")
set(HPX_WITH_PARCELPORT_LIBFABRIC ON CACHE BOOL "")
set(HPX_PARCELPORT_LIBFABRIC_PROVIDER "gni" CACHE STRING
    "See libfabric docs for details, gni,verbs,psm2 etc etc")
set(HPX_PARCELPORT_LIBFABRIC_THROTTLE_SENDS "256" CACHE STRING
    "Max number of messages in flight at once")
set(HPX_PARCELPORT_LIBFABRIC_WITH_DEV_MODE OFF CACHE BOOL
    "Custom libfabric logging flag")
set(HPX_PARCELPORT_LIBFABRIC_WITH_LOGGING OFF CACHE BOOL
    "Libfabric parcelport logging on/off flag")
set(HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD "4096" CACHE STRING
    "The threshhold in bytes to when perform zero copy optimizations (default: 128)")
# We do a cross compilation here ...
set(CMAKE_CROSSCOMPILING ON CACHE BOOL "")

```

CrayKNL

```

# Copyright (c) 2014 Thomas Heller
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
#
# This is the default toolchain file to be used with Intel Xeon PHIs. It sets
# the appropriate compile flags and compiler such that HPX will compile.
# Note that you still need to provide Boost, hwloc and other utility libraries

```

(continues on next page)

(continued from previous page)

```

# like a custom allocator yourself.
#
if(HPX_WITH_STATIC_LINKING)
    set_property(GLOBAL PROPERTY TARGET_SUPPORTS_SHARED_LIBS FALSE)
else()
endif()
# Set the Cray Compiler Wrapper
set(CMAKE_CXX_COMPILER CC)
set(CMAKE_C_COMPILER cc)
set(CMAKE_Fortran_COMPILER ftn)
if (CMAKE_VERSION VERSION_GREATER 3.3.9)
    set(__includes "<INCLUDES>")
endif()
set(CMAKE_C_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CFLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_CFLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_C_COMPILE_OBJECT "<CMAKE_C_COMPILER> -shared -fPIC <DEFINES> ${__includes}
→<FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_C_LINK_EXECUTABLE "<CMAKE_C_COMPILER> -fPIC <FLAGS> <CMAKE_C_LINK_FLAGS>
→<LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_C_CREATE_SHARED_LIBRARY "<CMAKE_C_COMPILER> -fPIC -shared <CMAKE_SHARED_
→LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_
→CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET> <OBJECTS> <LINK_
→LIBRARIES>" CACHE STRING "")
#
set(CMAKE_CXX_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CXX_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_CXX_FLAGS "-fPIC -shared" CACHE STRING "")
set(CMAKE_CXX_COMPILE_OBJECT "<CMAKE_CXX_COMPILER> -shared -fPIC <DEFINES> ${__
→includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING ")
set(CMAKE_CXX_LINK_EXECUTABLE "<CMAKE_CXX_COMPILER> -fPIC -dynamic <FLAGS> <CMAKE_CXX_
→LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_CREATE_SHARED_LIBRARY "<CMAKE_CXX_COMPILER> -fPIC -shared <CMAKE_SHARED_
→LIBRARY_CXX_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_SHARED_LIBRARY_
→CREATE_CXX_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET> <OBJECTS> <LINK_
→LIBRARIES>" CACHE STRING ")
#
set(CMAKE_Fortran_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_Fortran_FLAGS "-fPIC" CACHE STRING "")
set(CMAKE_SHARED_LIBRARY_CREATE_Fortran_FLAGS "-shared" CACHE STRING "")
set(CMAKE_Fortran_COMPILE_OBJECT "<CMAKE_Fortran_COMPILER> -shared -fPIC <DEFINES> ${__
→includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING ")
set(CMAKE_Fortran_LINK_EXECUTABLE "<CMAKE_Fortran_COMPILER> -fPIC <FLAGS> <CMAKE_
→Fortran_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>")
set(CMAKE_Fortran_CREATE_SHARED_LIBRARY "<CMAKE_Fortran_COMPILER> -fPIC -shared
→<CMAKE_SHARED_LIBRARY_Fortran_FLAGS> <LANGUAGE_COMPILE_FLAGS> <LINK_FLAGS> <CMAKE_
→SHARED_LIBRARY_CREATE_Fortran_FLAGS> <SONAME_FLAG><TARGET SONAME> -o <TARGET>
→<OBJECTS> <LINK_LIBRARIES>" CACHE STRING ")
#
# Disable searches in the default system paths. We are cross compiling after all
# and cmake might pick up wrong libraries that way
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM BOTH)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
set(HPX_WITH_PARCELPORT_TCP ON CACHE BOOL "")

```

(continues on next page)

(continued from previous page)

```

set(HPX_WITH_PARCELPORT_MPI ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_MPI_MULTITHREADED OFF CACHE BOOL "")
set(HPX_WITH_PARCELPORT_LIBFABRIC ON CACHE BOOL "")
set(HPX_PARCELPORT_LIBFABRIC_PROVIDER "gni" CACHE STRING
    "See libfabric docs for details, gni,verbs,psm2 etc etc")
set(HPX_PARCELPORT_LIBFABRIC_THROTTLE_SENDS "256" CACHE STRING
    "Max number of messages in flight at once")
set(HPX_PARCELPORT_LIBFABRIC_WITH_DEV_MODE OFF CACHE BOOL
    "Custom libfabric logging flag")
set(HPX_PARCELPORT_LIBFABRIC_WITH_LOGGING OFF CACHE BOOL
    "Libfabric parcelport logging on/off flag")
set(HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD "4096" CACHE STRING
    "The threshhold in bytes to when perform zero copy optimizations (default: 128)")
# Set the TBBMALLOC_PLATFORM correctly so that find_package(TBBMalloc) sets the
# right hints
set(TBBMALLOC_PLATFORM "mic-knl" CACHE STRING "")
# We have a bunch of cores on the MIC ... increase the default
set(HPX_WITH_MAX_CPU_COUNT "512" CACHE STRING "")
# We do a cross compilation here ...
set(CMAKE_CROSSCOMPILING ON CACHE BOOL "")
# RDTSCP is available on Xeon/Phis
set(HPX_WITH_RDTSCP ON CACHE BOOL "")

```

CrayKNLStatic

```

# Copyright (c) 2014-2017 Thomas Heller
# Copyright (c) 2017 Bryce Adelstein Lelbach
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
set(HPX_WITH_STATIC_LINKING ON CACHE BOOL "")
set(HPX_WITH_STATIC_EXE_LINKING ON CACHE BOOL "")
set_property(GLOBAL PROPERTY TARGET_SUPPORTS_SHARED_LIBS FALSE)
# Set the Cray Compiler Wrapper
set(CMAKE_CXX_COMPILER cc)
set(CMAKE_C_COMPILER cc)
set(CMAKE_Fortran_COMPILER ftn)
if (CMAKE_VERSION VERSION_GREATER 3.3.9)
    set(__includes "<INCLUDES>")
endif()
set(CMAKE_C_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_C_COMPILE_OBJECT "<CMAKE_C_COMPILER> -static -fPIC <DEFINES> ${__includes}
    ↪<FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_C_LINK_EXECUTABLE "<CMAKE_C_COMPILER> -fPIC <FLAGS> <CMAKE_C_LINK_FLAGS>
    ↪<LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_CXX_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_CXX_COMPILE_OBJECT "<CMAKE_CXX_COMPILER> -static -fPIC <DEFINES> ${__
    ↪includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_CXX_LINK_EXECUTABLE "<CMAKE_CXX_COMPILER> -fPIC <FLAGS> <CMAKE_CXX_LINK_
    ↪FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_Fortran_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_Fortran_COMPILE_OBJECT "<CMAKE_Fortran_COMPILER> -static -fPIC <DEFINES> ${__
    ↪includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")

```

(continues on next page)

(continued from previous page)

CrayStatic

```

# Copyright (c) 2014-2017 Thomas Heller
# Copyright (c) 2017 Bryce Adelstein Lelbach
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
set(HPX_WITH_STATIC_LINKING ON CACHE BOOL "")
set(HPX_WITH_STATIC_EXE_LINKING ON CACHE BOOL "")
set_property(GLOBAL PROPERTY TARGET_SUPPORTS_SHARED_LIBS FALSE)
# Set the Cray Compiler Wrapper
set(CMAKE_CXX_COMPILER CC)
set(CMAKE_C_COMPILER cc)
set(CMAKE_Fortran_COMPILER ftn)
if (CMAKE_VERSION VERSION_GREATER 3.3.9)
    set(__includes "<INCLUDES>")
endif()
set(CMAKE_C_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_C_COMPILE_OBJECT "<CMAKE_C_COMPILER> -static -fPIC <DEFINES> ${__includes}"
    →<FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_C_LINK_EXECUTABLE "<CMAKE_C_COMPILER> -fPIC <FLAGS> <CMAKE_C_LINK_FLAGS>
    →<LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")

```

(continues on next page)

(continued from previous page)

```

set(CMAKE_CXX_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_CXX_COMPILE_OBJECT "<CMAKE_CXX_COMPILER> -static -fPIC <DEFINES> ${__  
↳includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_CXX_LINK_EXECUTABLE "<CMAKE_CXX_COMPILER> -fPIC <FLAGS> <CMAKE_CXX_LINK_  
↳FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>" CACHE STRING "")
set(CMAKE_Fortran_FLAGS_INIT "" CACHE STRING "")
set(CMAKE_Fortran_COMPILE_OBJECT "<CMAKE_Fortran_COMPILER> -static -fPIC <DEFINES> ${__  
↳includes} <FLAGS> -o <OBJECT> -c <SOURCE>" CACHE STRING "")
set(CMAKE_Fortran_LINK_EXECUTABLE "<CMAKE_Fortran_COMPILER> -fPIC <FLAGS> <CMAKE_  
↳Fortran_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET> <LINK_LIBRARIES>")
# Disable searches in the default system paths. We are cross compiling after all
# and cmake might pick up wrong libraries that way
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM BOTH)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
# We do a cross compilation here ...
set(CMAKE_CROSSCOMPILING ON CACHE BOOL "")
# RDTSCP is available on Xeon/Phis
set(HPX_WITH_RDTSCP ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_TCP ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_MPI ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_MPI_MULTITHREADED ON CACHE BOOL "")
set(HPX_WITH_PARCELPORT_LIBFABRIC ON CACHE BOOL "")
set(HPX_PARCELPORT_LIBFABRIC_PROVIDER "gni" CACHE STRING
    "See libfabric docs for details, gni,verbs,psm2 etc etc")
set(HPX_PARCELPORT_LIBFABRIC_THROTTLE_SENDS "256" CACHE STRING
    "Max number of messages in flight at once")
set(HPX_PARCELPORT_LIBFABRIC_WITH_DEV_MODE OFF CACHE BOOL
    "Custom libfabric logging flag")
set(HPX_PARCELPORT_LIBFABRIC_WITH_LOGGING OFF CACHE BOOL
    "Libfabric parcelport logging on/off flag")
set(HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD "4096" CACHE STRING
    "The threshold in bytes to when perform zero copy optimizations (default: 128)")

```

XeonPhi

```

# Copyright (c) 2014 Thomas Heller
#
# SPDX-License-Identifier: BSL-1.0
# Distributed under the Boost Software License, Version 1.0. (See accompanying
# file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)
#
# This is the default toolchain file to be used with Intel Xeon PHIs. It sets
# the appropriate compile flags and compiler such that HPX will compile.
# Note that you still need to provide Boost, hwloc and other utility libraries
# like a custom allocator yourself.
#
set(CMAKE_SYSTEM_NAME Linux)
# Set the Intel Compiler
set(CMAKE_CXX_COMPILER icpc)
set(CMAKE_C_COMPILER icc)
set(CMAKE_Fortran_COMPILER ifort)
# Add the -mmic compile flag such that everything will be compiled for the correct

```

(continues on next page)

(continued from previous page)

```

# platform
set(CMAKE_CXX_FLAGS_INIT "-mmic" CACHE STRING "Initial compiler flags used to compile_
↪for the Xeon Phi")
set(CMAKE_C_FLAGS_INIT "-mmic" CACHE STRING "Initial compiler flags used to compile_
↪for the Xeon Phi")
set(CMAKE_Fortran_FLAGS_INIT "-mmic" CACHE STRING "Initial compiler flags used to_
↪compile for the Xeon Phi")
# Disable searches in the default system paths. We are cross compiling after all
# and cmake might pick up wrong libraries that way
set(CMAKE_FIND_ROOT_PATH_MODE_PROGRAM BOTH)
set(CMAKE_FIND_ROOT_PATH_MODE_LIBRARY ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_INCLUDE ONLY)
set(CMAKE_FIND_ROOT_PATH_MODE_PACKAGE ONLY)
# We do a cross compilation here ...
set(CMAKE_CROSSCOMPILING ON)
# Set our platform name
set(HPX_PLATFORM "XeonPhi")
# Always disable the ibverbs parcelport as it is non-functional on the BGQ.
set(HPX_WITH_PARCELPORT_VERBS OFF CACHE BOOL "Enable the ibverbs based parcelport._
↪This is currently an experimental feature")
set(HPX_WITH_PARCELPORT_MPI ON CACHE BOOL "Enable the MPI based parcelport.")
# We have a bunch of cores on the MIC ... increase the default
set(HPX_WITH_MAX_CPU_COUNT "256" CACHE STRING "")
# We default to tbbmalloc as our allocator on the MIC
if(NOT DEFINED HPX_WITH_MALLOC)
    set(HPX_WITH_MALLOC "tbbmalloc" CACHE STRING "")
endif()
# Set the TBBMALLOC_PLATFORM correctly so that find_package(TBBMalloc) sets the
# right hints
set(TBBMALLOC_PLATFORM "mic" CACHE STRING "")
set(HPX_HIDDEN_VISIBILITY OFF CACHE BOOL "Use -fvisibility=hidden for builds on_
↪platforms which support it")
# RDTSC is available on Xeon/Phis
set(HPX_WITH_RDTSC ON CACHE BOOL "")

```

CMake variables used to configure HPX

In order to configure *HPX*, you can set a variety of options to allow CMake to generate your specific makefiles/project files.

Variables that influence how *HPX* is built

The options are split into these categories:

- *Generic options*
- *Build Targets options*
- *Thread Manager options*
- *AGAS options*
- *Parcelport options*
- *Profiling options*
- *Debugging options*

- *Modules options*

Generic options

- `HPX_WITH_ACTION_BASE_COMPATIBILITY:BOOL`
- `HPX_WITH_AUTOMATIC_SERIALIZATION_REGISTRATION:BOOL`
- `HPX_WITH_BENCHMARK_SCRIPTS_PATH:PATH`
- `HPX_WITH_BUILD_BINARY_PACKAGE:BOOL`
- `HPX_WITH_COMPILER_WARNINGS:BOOL`
- `HPX_WITH_COMPILER_WARNINGS_AS_ERRORS:BOOL`
- `HPX_WITH_COMPRESSION_BZIP2:BOOL`
- `HPX_WITH_COMPRESSION_SNAPPY:BOOL`
- `HPX_WITH_COMPRESSION_ZLIB:BOOL`
- `HPX_WITH_CUDA:BOOL`
- `HPX_WITH_CUDA_CLANG:BOOL`
- `HPX_WITH_DATAPAR_VC:BOOL`
- `HPX_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_WITH_DISABLED_SIGNAL_EXCEPTION_HANDLERS:BOOL`
- `HPX_WITH_DYNAMIC_HPX_MAIN:BOOL`
- `HPX_WITH_FAULT_TOLERANCE:BOOL`
- `HPX_WITH_FORTRAN:BOOL`
- `HPX_WITH_FULL_RPATH:BOOL`
- `HPX_WITH_GCC_VERSION_CHECK:BOOL`
- `HPX_WITH_GENERIC_CONTEXT_COROUTINES:BOOL`
- `HPX_WITH_HIDDEN_VISIBILITY:BOOL`
- `HPX_WITH_LOGGING:BOOL`
- `HPX_WITH_MALLOC:STRING`
- `HPX_WITH_NATIVE_TLS:BOOL`
- `HPX_WITH_NICE_THREADLEVEL:BOOL`
- `HPX_WITH_PARCEL_COALESCING:BOOL`
- `HPX_WITH_REGISTER_THREAD_COMPATIBILITY:BOOL`
- `HPX_WITH_RUN_MAIN_EVERYWHERE:BOOL`
- `HPX_WITH_STACKOVERFLOW_DETECTION:BOOL`
- `HPX_WITH_STATIC_LINKING:BOOL`
- `HPX_WITH_VIM_YCM:BOOL`
- `HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD:STRING`

HPX_WITH_ACTION_BASE_COMPATIBILITY:BOOL

Enable deprecated action bases (default: ON)

HPX_WITH_AUTOMATIC_SERIALIZATION_REGISTRATION:BOOL

Use automatic serialization registration for actions and functions. This affects compatibility between HPX applications compiled with different compilers (default ON)

HPX_WITH_BENCHMARK_SCRIPTS_PATH:PATH

Directory to place batch scripts in

HPX_WITH_BUILD_BINARY_PACKAGE:BOOL

Build HPX on the build infrastructure on any LINUX distribution (default: OFF).

HPX_WITH_COMPILER_WARNINGS:BOOL

Enable compiler warnings (default: ON)

HPX_WITH_COMPILER_WARNINGS_AS_ERRORS:BOOL

Turn compiler warnings into errors (default: OFF)

HPX_WITH_COMPRESSION_BZIP2:BOOL

Enable bzip2 compression for parcel data (default: OFF).

HPX_WITH_COMPRESSION_SNAPPY:BOOL

Enable snappy compression for parcel data (default: OFF).

HPX_WITH_COMPRESSION_ZLIB:BOOL

Enable zlib compression for parcel data (default: OFF).

HPX_WITH_CUDA:BOOL

Enable CUDA support (default: OFF)

HPX_WITH_CUDA_CLANG:BOOL

Use clang to compile CUDA code (default: OFF)

HPX_WITH_DATAPAR_VC:BOOL

Enable data parallel algorithm support using the external Vc library (default: OFF)

HPX_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: ON)

HPX_WITH_DISABLED_SIGNAL_EXCEPTION_HANDLERS:BOOL

Disables the mechanism that produces debug output for caught signals and unhandled exceptions (default: OFF)

HPX_WITH_DYNAMIC_HPX_MAIN:BOOL

Enable dynamic overload of system `main()` (Linux only, default: ON)

HPX_WITH_FAULT_TOLERANCE:BOOL

Build HPX to tolerate failures of nodes, i.e. ignore errors in active communication channels (default: OFF)

HPX_WITH_FORTRAN:BOOL

Enable or disable the compilation of Fortran examples using HPX

HPX_WITH_FULL_RPATH:BOOL

Build and link HPX libraries and executables with full RPATHs (default: ON)

HPX_WITH_GCC_VERSION_CHECK:BOOL

Don't ignore version reported by `gcc` (default: ON)

HPX_WITH_GENERIC_CONTEXT_COROUTINES:BOOL

Use Boost.Context as the underlying coroutines context switch implementation.

HPX_WITH_HIDDEN_VISIBILITY:BOOL

Use `-fvisibility=hidden` for builds on platforms which support it (default OFF)

HPX_WITH_LOGGING:BOOL

Build HPX with logging enabled (default: ON).

HPX_WITH_MALLOC:STRING

Define which allocator should be linked in. Options are: system, tcmalloc, jemalloc, mimalloc, tbbmalloc, and custom (default is: tcmalloc)

HPX_WITH_NATIVE_TLS:BOOL

Use native TLS support if available (default:)

HPX_WITH_NICE_THREADLEVEL:BOOL

Set HPX worker threads to have high NICE level (may impact performance) (default: OFF)

HPX_WITH_PARCEL_COALESCING:BOOL

Enable the parcel coalescing plugin (default: ON).

HPX_WITH_REGISTER_THREAD_COMPATIBILITY:BOOL

Enable deprecated register_thread/work functions in the hpx::applier namespace (default: ON)

HPX_WITH_RUN_MAIN_EVERYWHERE:BOOL

Run hpx_main by default on all localities (default: OFF).

HPX_WITH_STACKOVERFLOW_DETECTION:BOOL

Enable stackoverflow detection for HPX threads/coroutines. (default: OFF, debug: ON)

HPX_WITH_STATIC_LINKING:BOOL

Compile HPX statically linked libraries (Default: OFF)

HPX_WITH_VIM_YCM:BOOL

Generate HPX completion file for VIM YouCompleteMe plugin

HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD:STRING

The threshold in bytes to when perform zero copy optimizations (default: 128)

Build Targets options

- *HPX_WITH_COMPILE_ONLY_TESTS:BOOL*
- *HPX_WITH_DEFAULT_TARGETS:BOOL*
- *HPX_WITH_DOCUMENTATION:BOOL*
- *HPX_WITH_DOCUMENTATION_OUTPUT_FORMATS:STRING*
- *HPX_WITH_EXAMPLES:BOOL*
- *HPX_WITH_EXAMPLES_HDF5:BOOL*
- *HPX_WITH_EXAMPLES_OPENMP:BOOL*
- *HPX_WITH_EXAMPLES_QT4:BOOL*
- *HPX_WITH_EXAMPLES_QTHREADS:BOOL*
- *HPX_WITH_EXAMPLES_TBB:BOOL*
- *HPX_WITH_EXECUTABLE_PREFIX:STRING*
- *HPX_WITH_FAIL_COMPILE_TESTS:BOOL*
- *HPX_WITH_IO_COUNTERS:BOOL*
- *HPX_WITH_PSEUDO_DEPENDENCIES:BOOL*
- *HPX_WITH_TESTS:BOOL*

- `HPX_WITH_TESTS_BENCHMARKS:BOOL`
- `HPX_WITH_TESTS_EXAMPLES:BOOL`
- `HPX_WITH_TESTS_EXTERNAL_BUILD:BOOL`
- `HPX_WITH_TESTS_HEADERS:BOOL`
- `HPX_WITH_TESTS_REGRESSIONS:BOOL`
- `HPX_WITH_TESTS_UNIT:BOOL`
- `HPX_WITH_TOOLS:BOOL`

`HPX_WITH_COMPILE_ONLY_TESTS:BOOL`

Create build system support for compile time only HPX tests (default ON)

`HPX_WITH_DEFAULT_TARGETS:BOOL`

Associate the core HPX library with the default build target (default: ON).

`HPX_WITH_DOCUMENTATION:BOOL`

Build the HPX documentation (default OFF).

`HPX_WITH_DOCUMENTATION_OUTPUT_FORMATS:STRING`

List of documentation output formats to generate. Valid options are html;singlehtml;latexpdf;man. Multiple values can be separated with semicolons. (default html).

`HPX_WITH_EXAMPLES:BOOL`

Build the HPX examples (default ON)

`HPX_WITH_EXAMPLES_HDF5:BOOL`

Enable examples requiring HDF5 support (default: OFF).

`HPX_WITH_EXAMPLES_OPENMP:BOOL`

Enable examples requiring OpenMP support (default: OFF).

`HPX_WITH_EXAMPLES_QT4:BOOL`

Enable examples requiring Qt4 support (default: OFF).

`HPX_WITH_EXAMPLES_QTHREADS:BOOL`

Enable examples requiring QThreads support (default: OFF).

`HPX_WITH_EXAMPLES_TBB:BOOL`

Enable examples requiring TBB support (default: OFF).

`HPX_WITH_EXECUTABLE_PREFIX:STRING`

Executable prefix (default none), ‘`hpx_`’ useful for system install.

`HPX_WITH_FAIL_COMPILE_TESTS:BOOL`

Create build system support for fail compile HPX tests (default ON)

`HPX_WITH_IO_COUNTERS:BOOL`

Enable IO counters (default: ON)

`HPX_WITH_PSEUDO_DEPENDENCIES:BOOL`

Force creating pseudo targets and pseudo dependencies (default ON).

`HPX_WITH_TESTS:BOOL`

Build the HPX tests (default ON)

`HPX_WITH_TESTS_BENCHMARKS:BOOL`

Build HPX benchmark tests (default: ON)

`HPX_WITH_TESTS_EXAMPLES:BOOL`

Add HPX examples as tests (default: ON)

HPX_WITH_TESTS_EXTERNAL_BUILD:BOOL
Build external cmake build tests (default: ON)

HPX_WITH_TESTS_HEADERS:BOOL
Build HPX header tests (default: OFF)

HPX_WITH_TESTS_REGRESSIONS:BOOL
Build HPX regression tests (default: ON)

HPX_WITH_TESTS_UNIT:BOOL
Build HPX unit tests (default: ON)

HPX_WITH_TOOLS:BOOL
Build HPX tools (default: OFF)

Thread Manager options

- *HPX_SCHEDULER_MAX_TERMINATED_THREADS:STRING*
- *HPX_WITH_COROUTINE_COUNTERS:BOOL*
- *HPX_WITH_IO_POOL:BOOL*
- *HPX_WITH_MAX_CPU_COUNT:STRING*
- *HPX_WITH_MAX_NUMA_DOMAIN_COUNT:STRING*
- *HPX_WITH_MORE_THAN_64_THREADS:BOOL*
- *HPX_WITH_SCHEDULER_LOCAL_STORAGE:BOOL*
- *HPX_WITH_SPINLOCK_DEADLOCK_DETECTION:BOOL*
- *HPX_WITH_SPINLOCK_POOL_NUM:STRING*
- *HPX_WITH_STACKTRACES:BOOL*
- *HPX_WITH_SWAP_CONTEXT_EMULATION:BOOL*
- *HPX_WITH_THREAD_BACKTRACE_DEPTH:STRING*
- *HPX_WITH_THREAD_BACKTRACE_ON_SUSPENSION:BOOL*
- *HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATES:BOOL*
- *HPX_WITH_THREAD_CUMULATIVE_COUNTS:BOOL*
- *HPX_WITH_THREAD_IDLE_RATES:BOOL*
- *HPX_WITH_THREAD_LOCAL_STORAGE:BOOL*
- *HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF:BOOL*
- *HPX_WITH_THREAD_QUEUE_WAITTIME:BOOL*
- *HPX_WITH_THREAD_SCHEDULERS:STRING*
- *HPX_WITH_THREAD_STACK_MMAP:BOOL*
- *HPX_WITH_THREAD_STEALING_COUNTS:BOOL*
- *HPX_WITH_THREAD_TARGET_ADDRESS:BOOL*
- *HPX_WITH_TIMER_POOL:BOOL*

HPX_SCHEDULER_MAX_TERMINATED_THREADS:STRING

[Deprecated] Maximum number of terminated threads collected before those are cleaned up (default: 100)

HPX_WITH_COROUTINE_COUNTERS:BOOL

Enable keeping track of coroutine creation and rebind counts (default: OFF)

HPX_WITH_IO_POOL:BOOL

Disable internal IO thread pool, do not change if not absolutely necessary (default: ON)

HPX_WITH_MAX_CPU_COUNT:STRING

HPX applications will not use more than this number of OS-Threads (empty string means dynamic) (default: 64)

HPX_WITH_MAX_NUMA_DOMAIN_COUNT:STRING

HPX applications will not run on machines with more NUMA domains (default: 8)

HPX_WITH_MORE_THAN_64_THREADS:BOOL

HPX applications will be able to run on more than 64 cores (This variable is deprecated. The value is derived from HPX_WITH_MAX_CPU_COUNT instead.)

HPX_WITH_SCHEDULER_LOCAL_STORAGE:BOOL

Enable scheduler local storage for all HPX schedulers (default: OFF)

HPX_WITH_SPINLOCK_DEADLOCK_DETECTION:BOOL

Enable spinlock deadlock detection (default: OFF)

HPX_WITH_SPINLOCK_POOL_NUM:STRING

Number of elements a spinlock pool manages (default: 128)

HPX_WITH_STACKTRACES:BOOL

Attach backtraces to HPX exceptions (default: ON)

HPX_WITH_SWAP_CONTEXT_EMULATION:BOOL

Emulate SwapContext API for coroutines (default: OFF)

HPX_WITH_THREAD_BACKTRACE_DEPTH:STRING

Thread stack back trace depth being captured (default: 5)

HPX_WITH_THREAD_BACKTRACE_ON_SUSPENSION:BOOL

Enable thread stack back trace being captured on suspension (default: OFF)

HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATES:BOOL

Enable measuring thread creation and cleanup times (default: OFF)

HPX_WITH_THREAD_CUMULATIVE_COUNTS:BOOL

Enable keeping track of cumulative thread counts in the schedulers (default: ON)

HPX_WITH_THREAD_IDLE_RATES:BOOL

Enable measuring the percentage of overhead times spent in the scheduler (default: OFF)

HPX_WITH_THREAD_LOCAL_STORAGE:BOOL

Enable thread local storage for all HPX threads (default: OFF)

HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF:BOOL

HPX scheduler threads do exponential backoff on idle queues (default: ON)

HPX_WITH_THREAD_QUEUE_WAITTIME:BOOL

Enable collecting queue wait times for threads (default: OFF)

HPX_WITH_THREAD_SCHEDULERS:STRING

Which thread schedulers are built. Options are: all, abp-priority, local, static-priority, static, shared-priority. For multiple enabled schedulers, separate with a semicolon (default: all)

HPX_WITH_THREAD_STACK_MMAP:BOOL

Use mmap for stack allocation on appropriate platforms

HPX_WITH_THREAD_STEALING_COUNTS:BOOL

Enable keeping track of counts of thread stealing incidents in the schedulers (default: OFF)

HPX WITH THREAD TARGET ADDRESS:BOOL

Enable storing target address in thread for NUMA awareness (default: OFF)

HPX WITH TIMER POOL:BOOL

Disable internal timer thread pool, do not change if not absolutely necessary (default: ON)

AGAS options

- HPX WITH AGAS DUMP REFCNT ENTRIES:BOOL

HPX WITH AGAS DUMP REFCNT ENTRIES:BOOL

Enable dumps of the AGAS refcnt tables to logs (default: OFF)

Parcelport options

- *HPX_WITH_NETWORKING:BOOL*
 - *HPX_WITH_PARCELPORT_ACTION_COUNTERS:BOOL*
 - *HPX_WITH_PARCELPORT_LIBFABRIC:BOOL*
 - *HPX_WITH_PARCELPORT_MPI:BOOL*
 - *HPX_WITH_PARCELPORT_MPI_ENV:STRING*
 - *HPX_WITH_PARCELPORT_MPI_MULTITHREADED:BOOL*
 - *HPX_WITH_PARCELPORT_TCP:BOOL*
 - *HPX_WITH_PARCELPORT_VERBS:BOOL*
 - *HPX_WITH_PARCEL_PROFILING:BOOL*

HPX WITH NETWORKING: BOOL

Enable support for networking and multi-node runs (default: ON)

HPX WITH PARCELPORT ACTION COUNTERS:BOOL

Enable performance counters reporting parcelport statistics on a per-action basis.

HPX_WITH_PARCELPORT_LIBFABRIC:BOOL

Enable the libfabric based parcelport. This is currently an experimental feature.

HPX_WITH_PARCELPORT_MPI:BOOL

Enable the MPI based parcelport.

HPX_WITH_PARCELPORT_MPI_ENV:STRING

List of environment variables checked to detect MPI (default: MV2_COMM_WORLD_RANK;PMI_RANK;OMPI_COMM_WORLD_RANK)

HPX_WITH_PARCELPORT_MPI_MULTITHREADED:BOOL

Turn on MPI multithreading support (default: ON).

HPX_WITH_PARCELPORT_TCP : BOOL

Enable the TCP based parcelport.

HPX_WITH_PARCELPORT_VERBS:BOOL

Enable the ibverbs based parcelport. This is currently an experimental feature.

HPX_WITH_PARCEL_PROFILING:BOOL

Enable profiling data for parcels

Profiling options

- *HPX_WITH_APEX:BOOL*
- *HPX_WITH_GOOGLE_PERFTOOLS:BOOL*
- *HPX_WITH_ITTNOTIFY:BOOL*
- *HPX_WITH_PAPI:BOOL*

HPX_WITH_APEX:BOOL

Enable APEX instrumentation support.

HPX_WITH_GOOGLE_PERFTOOLS:BOOL

Enable Google Perftools instrumentation support.

HPX_WITH_ITTNOTIFY:BOOL

Enable Amplifier (ITT) instrumentation support.

HPX_WITH_PAPI:BOOL

Enable the PAPI based performance counter.

Debugging options

- *HPX_WITH_ATTACH_DEBUGGER_ON_TEST_FAILURE:BOOL*
- *HPX_WITH_SANITIZERS:BOOL*
- *HPX_WITH_TESTS_DEBUG_LOG:BOOL*
- *HPX_WITH_TESTS_DEBUG_LOG_DESTINATION:STRING*
- *HPX_WITH_THREAD_DEBUG_INFO:BOOL*
- *HPX_WITH_THREAD_DESCRIPTION_FULL:BOOL*
- *HPX_WITH_THREAD_GUARD_PAGE:BOOL*
- *HPX_WITH_VALGRIND:BOOL*
- *HPX_WITH_VERIFY_LOCKS:BOOL*
- *HPX_WITH_VERIFY_LOCKS_BACKTRACE:BOOL*
- *HPX_WITH_VERIFY_LOCKS_GLOBALLY:BOOL*

HPX_WITH_ATTACH_DEBUGGER_ON_TEST_FAILURE:BOOL

Break the debugger if a test has failed (default: OFF)

HPX_WITH_SANITIZERS:BOOL

Configure with sanitizer instrumentation support.

HPX_WITH_TESTS_DEBUG_LOG:BOOL

Turn on debug logs (-hpx:debug-hpx-log) for tests (default: OFF)

HPX_WITH_TESTS_DEBUG_LOG_DESTINATION:STRING

Destination for test debug logs (default: cout)

HPX_WITH_THREAD_DEBUG_INFO:BOOL

Enable thread debugging information (default: OFF, implicitly enabled in debug builds)

HPX_WITH_THREAD_DESCRIPTION_FULL:BOOL

Use function address for thread description (default: OFF)

HPX_WITH_THREAD_GUARD_PAGE:BOOL

Enable thread guard page (default: ON)

HPX_WITH_VALGRIND:BOOL

Enable Valgrind instrumentation support.

HPX_WITH_VERIFY_LOCKS:BOOL

Enable lock verification code (default: OFF, implicitly enabled in debug builds)

HPX_WITH_VERIFY_LOCKS_BACKTRACE:BOOL

Enable thread stack back trace being captured on lock registration (to be used in combination with HPX_WITH_VERIFY_LOCKS=ON, default: OFF)

HPX_WITH_VERIFY_LOCKS_GLOBALLY:BOOL

Enable global lock verification code (default: OFF, implicitly enabled in debug builds)

Modules options

- *HPX_AFFINITY_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_AFFINITY_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_AFFINITY_WITH_TESTS:BOOL*
- *HPX_ALGORITHMS_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_ALGORITHMS_WITH_TESTS:BOOL*
- *HPX_ALLOCATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_ALLOCATOR_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_ALLOCATOR_SUPPORT_WITH_TESTS:BOOL*
- *HPX_ASSERTION_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_ASSERTION_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_ASSERTION_WITH_TESTS:BOOL*
- *HPX_BASIC_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_BASIC_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_BASIC_EXECUTION_WITH_TESTS:BOOL*
- *HPX_BATCH_ENVIRONMENTS_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_BATCH_ENVIRONMENTS_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_BATCH_ENVIRONMENTS_WITH_TESTS:BOOL*
- *HPX_CACHE_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_CACHE_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_CACHE_WITH_TESTS:BOOL*
- *HPX_CHECKPOINT_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_CHECKPOINT_WITH_DEPRECATED_WARNINGS:BOOL*
- *HPX_CHECKPOINT_WITH_TESTS:BOOL*
- *HPX_COLLECTIVES_WITH_COMPATIBILITY_HEADERS:BOOL*
- *HPX_COLLECTIVES_WITH_DEPRECATED_WARNINGS:BOOL*

- `HPX_COLLECTIVES_WITH_TESTS:BOOL`
- `HPX_COMPUTE_CUDA_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_COMPUTE_CUDA_WITH_TESTS:BOOL`
- `HPX_COMPUTE_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_COMPUTE_WITH_TESTS:BOOL`
- `HPX_CONCEPTS_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_CONCEPTS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_CONCEPTS_WITH_TESTS:BOOL`
- `HPX_CONCURRENCY_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_CONCURRENCY_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_CONCURRENCY_WITH_TESTS:BOOL`
- `HPX_CONFIG_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_CONFIG_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_CONFIG_WITH_TESTS:BOOL`
- `HPX_COROUTINES_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_COROUTINES_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_COROUTINES_WITH_TESTS:BOOL`
- `HPX_DATASTRUCTURES_WITH_ADAPTER_STD_TUPLE:BOOL`
- `HPX_DATASTRUCTURES_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_DATASTRUCTURES_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_DATASTRUCTURES_WITH_TESTS:BOOL`
- `HPX_DEBUGGING_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_DEBUGGING_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_DEBUGGING_WITH_TESTS:BOOL`
- `HPX_ERRORS_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_ERRORS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_ERRORS_WITH_TESTS:BOOL`
- `HPX_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_EXECUTION_WITH_TESTS:BOOL`
- `HPX_FILESYSTEM_WITH_BOOST_FILESYSTEM_COMPATIBILITY:BOOL`
- `HPX_FILESYSTEM_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_FILESYSTEM_WITH_TESTS:BOOL`
- `HPX_FORMAT_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_FORMAT_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_FORMAT_WITH_TESTS:BOOL`

- `HPX_FUNCTIONAL_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_FUNCTIONAL_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_FUNCTIONAL_WITH_TESTS:BOOL`
- `HPX_HARDWARE_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_HARDWARE_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_HARDWARE_WITH_TESTS:BOOL`
- `HPX_HASHING_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_HASHING_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_HASHING_WITH_TESTS:BOOL`
- `HPX_ITERATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_ITERATOR_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_ITERATOR_SUPPORT_WITH_TESTS:BOOL`
- `HPX_LOCAL_LCOS_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_LOCAL_LCOS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_LOCAL_LCOS_WITH_TESTS:BOOL`
- `HPX_LOGGING_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_LOGGING_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_LOGGING_WITH_TESTS:BOOL`
- `HPX_MEMORY_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_MEMORY_WITH_TESTS:BOOL`
- `HPX_NAMING_BASE_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_NAMING_BASE_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_NAMING_BASE_WITH_TESTS:BOOL`
- `HPX_PLUGIN_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_PLUGIN_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_PLUGIN_WITH_TESTS:BOOL`
- `HPX_PREPROCESSOR_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_PREPROCESSOR_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_PREPROCESSOR_WITH_TESTS:BOOL`
- `HPX_PROGRAM_OPTIONS_WITH_BOOST_PROGRAM_OPTIONS_COMPATIBILITY:BOOL`
- `HPX_PROGRAM_OPTIONS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_PROGRAM_OPTIONS_WITH_TESTS:BOOL`
- `HPX_RESILIENCY_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_RESILIENCY_WITH_TESTS:BOOL`
- `HPX_RESOURCE_PARTITIONER_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_RESOURCE_PARTITIONER_WITH_DEPRECATED_WARNINGS:BOOL`

- `HPX_RESOURCE_PARTITIONER_WITH_TESTS:BOOL`
- `HPX_SEGMENTED_ALGORITHMS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_SEGMENTED_ALGORITHMS_WITH_TESTS:BOOL`
- `HPX_SERIALIZATION_WITH_BOOST_TYPES:BOOL`
- `HPX_SERIALIZATION_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_SERIALIZATION_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_SERIALIZATION_WITH_TESTS:BOOL`
- `HPX_STATIC_REINIT_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_STATIC_REINIT_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_STATIC_REINIT_WITH_TESTS:BOOL`
- `HPX_STATISTICS_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_STATISTICS_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_STATISTICS_WITH_TESTS:BOOL`
- `HPX_SYNCHRONIZATION_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_SYNCHRONIZATION_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_SYNCHRONIZATION_WITH_TESTS:BOOL`
- `HPX_TESTING_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_TESTING_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_TESTING_WITH_TESTS:BOOL`
- `HPX_THREADMANAGER_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_THREADMANAGER_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_THREADMANAGER_WITH_TESTS:BOOL`
- `HPX_THREAD_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_THREAD_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_THREAD_SUPPORT_WITH_TESTS:BOOL`
- `HPX_TIMING_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_TIMING_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_TIMING_WITH_TESTS:BOOL`
- `HPX_TOPOLOGY_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_TOPOLOGY_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_TOPOLOGY_WITH_TESTS:BOOL`
- `HPX_TYPE_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_TYPE_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL`
- `HPX_TYPE_SUPPORT_WITH_TESTS:BOOL`
- `HPX_UTIL_WITH_COMPATIBILITY_HEADERS:BOOL`
- `HPX_UTIL_WITH_DEPRECATED_WARNINGS:BOOL`

- `HPX_UTIL_WITH_TESTS:BOOL`

`HPX_AFFINITY_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_AFFINITY_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_AFFINITY_WITH_TESTS:BOOL`

Build HPX affinity module tests. (default: ON)

`HPX_ALGORITHMS_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_ALGORITHMS_WITH_TESTS:BOOL`

Build HPX algorithms module tests. (default: ON)

`HPX_ALLOCATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_ALLOCATOR_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_ALLOCATOR_SUPPORT_WITH_TESTS:BOOL`

Build HPX allocator_support module tests. (default: ON)

`HPX_ASSERTION_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_ASSERTION_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_ASSERTION_WITH_TESTS:BOOL`

Build HPX assertion module tests. (default: ON)

`HPX_BASIC_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_BASIC_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_BASIC_EXECUTION_WITH_TESTS:BOOL`

Build HPX basic_execution module tests. (default: ON)

`HPX_BATCH_ENVIRONMENTS_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_BATCH_ENVIRONMENTS_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_BATCH_ENVIRONMENTS_WITH_TESTS:BOOL`

Build HPX batch_environments module tests. (default: ON)

`HPX_CACHE_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

`HPX_CACHE_WITH_DEPRECATED_WARNINGS:BOOL`

Enable warnings for deprecated facilities. (default: On)

`HPX_CACHE_WITH_TESTS:BOOL`

Build HPX cache module tests. (default: ON)

`HPX_CHECKPOINT_WITH_COMPATIBILITY_HEADERS:BOOL`

Enable compatibility headers for old headers. (default: ON)

HPX_CHECKPOINT_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_CHECKPOINT_WITH_TESTS : BOOL
Build HPX checkpoint module tests. (default: ON)

HPX_COLLECTIVES_WITH_COMPATIBILITY_HEADERS : BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_COLLECTIVES_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_COLLECTIVES_WITH_TESTS : BOOL
Build HPX collectives module tests. (default: ON)

HPX_COMPUTE_CUDA_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_COMPUTE_CUDA_WITH_TESTS : BOOL
Build HPX compute_cuda module tests. (default: ON)

HPX_COMPUTE_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_COMPUTE_WITH_TESTS : BOOL
Build HPX compute module tests. (default: ON)

HPX_CONCEPTS_WITH_COMPATIBILITY_HEADERS : BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_CONCEPTS_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_CONCEPTS_WITH_TESTS : BOOL
Build HPX concepts module tests. (default: ON)

HPX_CONCURRENCY_WITH_COMPATIBILITY_HEADERS : BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_CONCURRENCY_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_CONCURRENCY_WITH_TESTS : BOOL
Build HPX concurrency module tests. (default: ON)

HPX_CONFIG_WITH_COMPATIBILITY_HEADERS : BOOL
Enable compatibility headers for old headers. (default: OFF)

HPX_CONFIG_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_CONFIG_WITH_TESTS : BOOL
Build HPX config module tests. (default: ON)

HPX_COROUTINES_WITH_COMPATIBILITY_HEADERS : BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_COROUTINES_WITH_DEPRECATED_WARNINGS : BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_COROUTINES_WITH_TESTS : BOOL
Build HPX coroutines module tests. (default: ON)

HPX_DATASTRUCTURES_WITH_ADAPT_STD_TUPLE:BOOL
Enable compatibility of hpx::util::tuple with std::tuple. (default: ON)

HPX_DATASTRUCTURES_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_DATASTRUCTURES_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_DATASTRUCTURES_WITH_TESTS:BOOL
Build HPX datastructures module tests. (default: ON)

HPX_DEBUGGING_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_DEBUGGING_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_DEBUGGING_WITH_TESTS:BOOL
Build HPX debugging module tests. (default: ON)

HPX_ERRORS_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_ERRORS_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_ERRORS_WITH_TESTS:BOOL
Build HPX errors module tests. (default: ON)

HPX_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_EXECUTION_WITH_TESTS:BOOL
Build HPX execution module tests. (default: ON)

HPX_FILESYSTEM_WITH_BOOST_FILESYSTEM_COMPATIBILITY:BOOL
Enable Boost.FileSystem compatibility. (default: ON)

HPX_FILESYSTEM_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_FILESYSTEM_WITH_TESTS:BOOL
Build HPX filesystem module tests. (default: ON)

HPX_FORMAT_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_FORMAT_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_FORMAT_WITH_TESTS:BOOL
Build HPX format module tests. (default: ON)

HPX_FUNCTIONAL_WITH_COMPATIBILITY_HEADERS:BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_FUNCTIONAL_WITH_DEPRECATED_WARNINGS:BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_FUNCTIONAL_WITH_TESTS : BOOL

Build HPX functional module tests. (default: ON)

HPX_HARDWARE_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_HARDWARE_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_HARDWARE_WITH_TESTS : BOOL

Build HPX hardware module tests. (default: ON)

HPX_HASHING_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_HASHING_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_HASHING_WITH_TESTS : BOOL

Build HPX hashing module tests. (default: ON)

HPX_ITERATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_ITERATOR_SUPPORT_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_ITERATOR_SUPPORT_WITH_TESTS : BOOL

Build HPX iterator_support module tests. (default: ON)

HPX_LOCAL_LCOS_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_LOCAL_LCOS_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_LOCAL_LCOS_WITH_TESTS : BOOL

Build HPX local_lclos module tests. (default: ON)

HPX_LOGGING_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_LOGGING_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_LOGGING_WITH_TESTS : BOOL

Build HPX logging module tests. (default: ON)

HPX_MEMORY_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_MEMORY_WITH_TESTS : BOOL

Build HPX memory module tests. (default: ON)

HPX_NAMING_BASE_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: OFF)

HPX_NAMING_BASE_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_NAMING_BASE_WITH_TESTS : BOOL

Build HPX naming_base module tests. (default: ON)

HPX_PLUGIN_WITH_COMPATIBILITY_HEADERS :BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_PLUGIN_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_PLUGIN_WITH_TESTS :BOOL
Build HPX plugin module tests. (default: ON)

HPX_PREPROCESSOR_WITH_COMPATIBILITY_HEADERS :BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_PREPROCESSOR_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_PREPROCESSOR_WITH_TESTS :BOOL
Build HPX preprocessor module tests. (default: ON)

HPX_PROGRAM_OPTIONS_WITH_BOOST_PROGRAM_OPTIONS_COMPATIBILITY :BOOL
Enable Boost.ProgramOptions compatibility. (default: ON)

HPX_PROGRAM_OPTIONS_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_PROGRAM_OPTIONS_WITH_TESTS :BOOL
Build HPX program_options module tests. (default: ON)

HPX_RESILIENCY_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_RESILIENCY_WITH_TESTS :BOOL
Build HPX resiliency module tests. (default: ON)

HPX_RESOURCE_PARTITIONER_WITH_COMPATIBILITY_HEADERS :BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_RESOURCE_PARTITIONER_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_RESOURCE_PARTITIONER_WITH_TESTS :BOOL
Build HPX resource_partitioner module tests. (default: ON)

HPX_SEGMENTED_ALGORITHMS_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_SEGMENTED_ALGORITHMS_WITH_TESTS :BOOL
Build HPX segmented_algorithms module tests. (default: ON)

HPX_SERIALIZATION_WITH_BOOST_TYPES :BOOL
Enable serialization of certain Boost types. (default: ON)

HPX_SERIALIZATION_WITH_COMPATIBILITY_HEADERS :BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_SERIALIZATION_WITH_DEPRECATED_WARNINGS :BOOL
Enable warnings for deprecated facilities. (default: On)

HPX_SERIALIZATION_WITH_TESTS :BOOL
Build HPX serialization module tests. (default: ON)

HPX_STATIC_REINIT_WITH_COMPATIBILITY_HEADERS :BOOL
Enable compatibility headers for old headers. (default: ON)

HPX_STATIC_REINIT_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_STATIC_REINIT_WITH_TESTS:BOOL

Build HPX static_reinit module tests. (default: ON)

HPX_STATISTICS_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_STATISTICS_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_STATISTICS_WITH_TESTS:BOOL

Build HPX statistics module tests. (default: ON)

HPX_SYNCHRONIZATION_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_SYNCHRONIZATION_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_SYNCHRONIZATION_WITH_TESTS:BOOL

Build HPX synchronization module tests. (default: ON)

HPX_TESTING_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_TESTING_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_TESTING_WITH_TESTS:BOOL

Build HPX testing module tests. (default: ON)

HPX_THREADMANAGER_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_THREADMANAGER_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_THREADMANAGER_WITH_TESTS:BOOL

Build HPX threadmanager module tests. (default: ON)

HPX_THREAD_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_THREAD_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_THREAD_SUPPORT_WITH_TESTS:BOOL

Build HPX thread_support module tests. (default: ON)

HPX_TIMING_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_TIMING_WITH_DEPRECATED_WARNINGS:BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_TIMING_WITH_TESTS:BOOL

Build HPX timing module tests. (default: ON)

HPX_TOPOLOGY_WITH_COMPATIBILITY_HEADERS:BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_TOPOLOGY_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_TOPOLOGY_WITH_TESTS : BOOL

Build HPX topology module tests. (default: ON)

HPX_TYPE_SUPPORT_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_TYPE_SUPPORT_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_TYPE_SUPPORT_WITH_TESTS : BOOL

Build HPX type_support module tests. (default: ON)

HPX_UTIL_WITH_COMPATIBILITY_HEADERS : BOOL

Enable compatibility headers for old headers. (default: ON)

HPX_UTIL_WITH_DEPRECATED_WARNINGS : BOOL

Enable warnings for deprecated facilities. (default: On)

HPX_UTIL_WITH_TESTS : BOOL

Build HPX util module tests. (default: ON)

Additional tools and libraries used by *HPX*

Here is a list of additional libraries and tools that are either optionally supported by the build system or are optionally required for certain examples or tests. These libraries and tools can be detected by the *HPX* build system.

Each of the tools or libraries listed here will be automatically detected if they are installed in some standard location. If a tool or library is installed in a different location, you can specify its base directory by appending _ROOT to the variable name as listed below. For instance, to configure a custom directory for BOOST, specify BOOST_ROOT=/custom/boost/root.

BOOST_ROOT : PATH

Specifies where to look for the Boost installation to be used for compiling *HPX*. Set this if CMake is not able to locate a suitable version of Boost. The directory specified here can be either the root of an installed Boost distribution or the directory where you unpacked and built Boost without installing it (with staged libraries).

HWLOC_ROOT : PATH

Specifies where to look for the hwloc library. Set this if CMake is not able to locate a suitable version of hwloc. Hwloc provides platform-independent support for extracting information about the used hardware architecture (number of cores, number of NUMA domains, hyperthreading, etc.). *HPX* utilizes this information if available.

PAPI_ROOT : PATH

Specifies where to look for the PAPI library. The PAPI library is needed to compile a special component exposing PAPI hardware events and counters as *HPX* performance counters. This is not available on the Windows platform.

AMPLIFIER_ROOT : PATH

Specifies where to look for one of the tools of the Intel Parallel Studio product, either Intel Amplifier or Intel Inspector. This should be set if the CMake variable HPX_USE_ITT_NOTIFY is set to ON. Enabling ITT support in *HPX* will integrate any application with the mentioned Intel tools, which customizes the generated information for your application and improves the generated diagnostics.

In addition, some of the examples may need the following variables:

HDF5_ROOT : PATH

Specifies where to look for the Hierarchical Data Format V5 (HDF5) include files and libraries.

2.5.3 Creating HPX projects

Using HPX with pkg-config

How to build HPX applications with pkg-config

After you are done installing *HPX*, you should be able to build the following program. It prints Hello World! on the *locality* you run it on.

```
// Copyright (c) 2007-2012 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

///////////////////////////////
// The purpose of this example is to execute a HPX-thread printing
// "Hello World!" once. That's all.

//[hello_world_1_getting_started
// Including 'hpx/hpx_main.hpp' instead of the usual 'hpx/hpx_init.hpp' enables
// to use the plain C-main below as the direct main HPX entry point.
#include <hpx/hpx_main.hpp>
#include <hpx/include/iostreams.hpp>

int main()
{
    // Say hello to the world!
    hpx::cout << "Hello World!\n" << hpx::flush;
    return 0;
}
//]
```

Copy the text of this program into a file called hello_world.cpp.

Now, in the directory where you put hello_world.cpp, issue the following commands (where \$HPX_LOCATION is the build directory or CMAKE_INSTALL_PREFIX you used while building *HPX*):

```
export PKG_CONFIG_PATH=$PKG_CONFIG_PATH:$HPX_LOCATION/lib/pkgconfig
c++ -o hello_world hello_world.cpp \
`pkg-config --cflags --libs hpx_application` \
-lhpx_iostreams -DHPX_APPLICATION_NAME=hello_world
```

Important: When using pkg-config with *HPX*, the pkg-config flags must go after the `-o` flag.

Note: *HPX* libraries have different names in debug and release mode. If you want to link against a debug *HPX* library, you need to use the `_debug` suffix for the pkg-config name. That means instead of `hpx_application` or `hpx_component`, you will have to use `hpx_application_debug` or `hpx_component_debug`. Moreover, all referenced *HPX* components need to have an appended `d` suffix. For example, instead of `-lhpx_iostreams` you will need to specify `-lhpx_iostreams_d`.

Important: If the *HPX* libraries are in a path that is not found by the dynamic linker, you will need to add the path

\$HPX_LOCATION/lib to your linker search path (for example LD_LIBRARY_PATH on Linux).

To test the program, type:

```
./hello_world
```

which should print Hello World! and exit.

How to build HPX components with pkg-config

Let's try a more complex example involving an *HPX* component. An *HPX* component is a class that exposes *HPX* actions. *HPX* components are compiled into dynamically loaded modules called component libraries. Here's the source code:

hello_world_component.cpp

```
//[hello_world_cpp_getting_started
#include "hello_world_component.hpp"
#include <hpx/include/iostreams.hpp>

#include <iostream>

namespace examples { namespace server
{
    void hello_world::invoke()
    {
        hpx::cout << "Hello HPX World!" << std::endl;
    }
} }

HPX_REGISTER_COMPONENT_MODULE();

typedef hpx::components::component<
    examples::server::hello_world
> hello_world_type;

HPX_REGISTER_COMPONENT(hello_world_type, hello_world);

HPX_REGISTER_ACTION(
```

hello_world_component.hpp

```
//[hello_world_hpp_getting_started
#ifndef HELLO_WORLD_COMPONENT_HPP
#define HELLO_WORLD_COMPONENT_HPP

#include <hpx/hpx.hpp>
#include <hpx/include/actions.hpp>
#include <hpx/include/lcos.hpp>
#include <hpx/include/components.hpp>
#include <hpx/include/serialization.hpp>

#include <utility>

namespace examples { namespace server
{
```

(continues on next page)

(continued from previous page)

```

struct HPX_COMPONENT_EXPORT hello_world
    : hpx::components::component_base<hello_world>
{
    void invoke();
    HPX_DEFINE_COMPONENT_ACTION(hello_world, invoke);
};
}

HPX_REGISTER_ACTION_DECLARATION(
    examples::server::hello_world::invoke_action, hello_world_invoke_action);

namespace examples
{
    struct hello_world
        : hpx::components::client_base<hello_world, server::hello_world>
    {
        typedef hpx::components::client_base<hello_world, server::hello_world>
            base_type;

        hello_world(hpx::future<hpx::naming::id_type> && f)
            : base_type(std::move(f))
        {}

        hello_world(hpx::naming::id_type && f)
            : base_type(std::move(f))
        {}

        void invoke()
        {
            hpx::async<server::hello_world::invoke_action>(*this->get_id()).get();
        }
    };
}

```

hello_world_client.cpp

```

// Copyright (c) 2012 Bryce Lelbach
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

//[hello_world_client_getting_started
#include "hello_world_component.hpp"
#include <hpx/hpx_init.hpp>

int hpx_main(hpx::program_options::variables_map&)
{
{
    // Create a single instance of the component on this locality.
    examples::hello_world client =
        hpx::new_<examples::hello_world>(hpx::find_here());

    // Invoke the component's action, which will print "Hello World!".
    client.invoke();
}

```

(continues on next page)

(continued from previous page)

```
    return hpx::finalize(); // Initiate shutdown of the runtime system.  
}  
  
int main(int argc, char* argv[])  
{  
    return hpx::init(argc, argv); // Initialize and run HPX.  
}  
//]
```

Copy the three source files above into three files (called `hello_world_component.cpp`, `hello_world_component.hpp` and `hello_world_client.cpp`, respectively).

Now, in the directory where you put the files, run the following command to build the component library. (where `$HPX_LOCATION` is the build directory or `CMAKE_INSTALL_PREFIX` you used while building *HPX*):

```
export PKG_CONFIG_PATH=$PKG_CONFIG_PATH:$HPX_LOCATION/lib/pkgconfig  
c++ -o libhpx_hello_world.so hello_world_component.cpp \  
`pkg-config --cflags --libs hpx_component` \  
-lhpx_iostreams -DHPX_COMPONENT_NAME=hpx_hello_world
```

Now pick a directory in which to install your *HPX* component libraries. For this example, we'll choose a directory named `my_hpx_libs`:

```
mkdir ~./my_hpx_libs  
mv libhpx_hello_world.so ~./my_hpx_libs
```

Note: *HPX* libraries have different names in debug and release mode. If you want to link against a debug *HPX* library, you need to use the `_debug` suffix for the pkg-config name. That means instead of `hpx_application` or `hpx_component` you will have to use `hpx_application_debug` or `hpx_component_debug`. Moreover, all referenced *HPX* components need to have a appended `d` suffix, e.g. instead of `-lhpx_iostreams` you will need to specify `-lhpx_iostreamsd`.

Important: If the *HPX* libraries are in a path that is not found by the dynamic linker. You need to add the path `$HPX_LOCATION/lib` to your linker search path (for example `LD_LIBRARY_PATH` on Linux).

Now, to build the application that uses this component (`hello_world_client.cpp`), we do:

```
export PKG_CONFIG_PATH=$PKG_CONFIG_PATH:$HPX_LOCATION/lib/pkgconfig  
c++ -o hello_world_client hello_world_client.cpp \  
` `pkg-config --cflags --libs hpx_application` \  
-L${HOME}/my_hpx_libs -lhpx_hello_world -lhpx_iostreams
```

Important: When using `pkg-config` with *HPX*, the `pkg-config` flags must go after the `-o` flag.

Finally, you'll need to set your `LD_LIBRARY_PATH` before you can run the program. To run the program, type:

```
export LD_LIBRARY_PATH="$LD_LIBRARY_PATH:${HOME}/my_hpx_libs"  
./hello_world_client
```

which should print `Hello HPX World!` and exit.

Using HPX with CMake-based projects

In addition to the pkg-config support discussed on the previous pages, *HPX* comes with full CMake support. In order to integrate *HPX* into existing or new CMakeLists.txt, you can leverage the `find_package`¹⁰⁶ command integrated into CMake. Following, is a Hello World component example using CMake.

Let's revisit what we have. We have three files that compose our example application:

- `hello_world_component.hpp`
- `hello_world_component.cpp`
- `hello_world_client.hpp`

The basic structure to include *HPX* into your CMakeLists.txt is shown here:

```
# Require a recent version of cmake
cmake_minimum_required(VERSION 3.3.2 FATAL_ERROR)

# This project is C++ based.
project(your_app CXX)

# Instruct cmake to find the HPX settings
find_package(HPX)
```

In order to have CMake find *HPX*, it needs to be told where to look for the `HPXConfig.cmake` file that is generated when *HPX* is built or installed. It is used by `find_package(HPX)` to set up all the necessary macros needed to use *HPX* in your project. The ways to achieve this are:

- Set the `HPX_DIR` CMake variable to point to the directory containing the `HPXConfig.cmake` script on the command line when you invoke CMake:

```
cmake -DHPX_DIR=$HPX_LOCATION/lib/cmake/HPX ...
```

where `$HPX_LOCATION` is the build directory or `CMAKE_INSTALL_PREFIX` you used when building/configuring *HPX*.

- Set the `CMAKE_PREFIX_PATH` variable to the root directory of your *HPX* build or install location on the command line when you invoke CMake:

```
cmake -DCMAKE_PREFIX_PATH=$HPX_LOCATION ...
```

The difference between `CMAKE_PREFIX_PATH` and `HPX_DIR` is that CMake will add common postfixes, such as `lib/cmake/<project>`, to the `MAKE_PREFIX_PATH` and search in these locations too. Note that if your project uses *HPX* as well as other CMake-managed projects, the paths to the locations of these multiple projects may be concatenated in the `CMAKE_PREFIX_PATH`.

- The variables above may be set in the CMake GUI or curses `ccmake` interface instead of the command line.

Additionally, if you wish to require *HPX* for your project, replace the `find_package(HPX)` line with `find_package(HPX REQUIRED)`.

You can check if *HPX* was successfully found with the `HPX_FOUND` CMake variable.

The simplest way to add the *HPX* component is to use the `add_hpx_component` macro and add it to the `CMakeLists.txt` file:

¹⁰⁶ https://www.cmake.org/cmake/help/latest/command/find_package.html

```
# build your application using HPX
add_hpx_component(hello_world
    SOURCES hello_world_component.cpp
    HEADERS hello_world_component.hpp
    COMPONENT_DEPENDENCIES iostreams)
```

Note: `add_hpx_component` adds a `_component` suffix to the target name. In the example above, a `hello_world_component` target will be created.

The available options to `add_hpx_component` are:

- `SOURCES`: The source files for that component
- `HEADERS`: The header files for that component
- `DEPENDENCIES`: Other libraries or targets this component depends on
- `COMPONENT_DEPENDENCIES`: The components this component depends on
- `PLUGIN`: Treats this component as a plugin-able library
- `COMPILE_FLAGS`: Additional compiler flags
- `LINK_FLAGS`: Additional linker flags
- `FOLDER`: Adds the headers and source files to this Source Group folder
- `EXCLUDE_FROM_ALL`: Do not build this component as part of the `all` target

After adding the component, the way you add the executable is as follows:

```
# build your application using HPX
add_hpx_executable(hello_world
    ESSENTIAL
    SOURCES hello_world_client.cpp
    COMPONENT_DEPENDENCIES hello_world)
```

Note: `add_hpx_executable` automatically adds a `_component` suffix to dependencies specified in `COMPONENT_DEPENDENCIES`, meaning you can directly use the name given when adding a component using `add_hpx_component`.

When you configure your application, all you need to do is set the `HPX_DIR` variable to point to the installation of `HPX`.

Note: All library targets built with `HPX` are exported and readily available to be used as arguments to `target_link_libraries`¹⁰⁷ in your targets. The `HPX` include directories are available with the `HPX_INCLUDE_DIRS` CMake variable.

¹⁰⁷ https://www.cmake.org/cmake/help/latest/command/target_link_libraries.html

CMake macros to integrate HPX into existing applications

In addition to the `add_hpx_component` and `add_hpx_executable`, you can use the `hpx_setup_target` macro to have an already existing target to be used with the *HPX* libraries:

```
hpx_setup_target(target)
```

Optional parameters are:

- `EXPORT`: Adds it to the CMake export list `HPXTargets`
- `INSTALL`: Generates an install rule for the target
- `PLUGIN`: Treats this component as a plugin-able library
- `TYPE`: The type can be: `EXECUTABLE`, `LIBRARY` or `COMPONENT`
- `DEPENDENCIES`: Other libraries or targets this component depends on
- `COMPONENT_DEPENDENCIES`: The components this component depends on
- `COMPILE_FLAGS`: Additional compiler flags
- `LINK_FLAGS`: Additional linker flags

If you do not use CMake, you can still build against *HPX*, but you should refer to the section on *How to build HPX components with pkg-config*.

Note: Since *HPX* relies on dynamic libraries, the dynamic linker needs to know where to look for them. If *HPX* isn't installed into a path that is configured as a linker search path, external projects need to either set `RPATH` or adapt `LD_LIBRARY_PATH` to point to where the *HPX* libraries reside. In order to set `RPATHS`, you can include `HPX_SetFullRPATH` in your project after all libraries you want to link against have been added. Please also consult the CMake documentation [here](#)¹⁰⁸.

Using HPX with Makefile

A basic project building with *HPX* is through creating makefiles. The process of creating one can get complex depending upon the use of `cmake` parameter `HPX_WITH_HPX_MAIN` (which defaults to `ON`).

How to build HPX applications with makefile

If *HPX* is installed correctly, you should be able to build and run a simple Hello World program. It prints `Hello World!` on the `locality` you run it on.

```
// Copyright (c) 2007-2012 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

///////////////////////////////
// The purpose of this example is to execute a HPX-thread printing
// "Hello World!" once. That's all.
```

(continues on next page)

¹⁰⁸ <https://gitlab.kitware.com/cmake/community/wikis/doc/cmake/RPATH-handling>

(continued from previous page)

```
//[hello_world_1_getting_started
// Including 'hpx/hpx_main.hpp' instead of the usual 'hpx/hpx_init.hpp' enables
// to use the plain C-main below as the direct main HPX entry point.
#include <hpx/hpx_main.hpp>
#include <hpx/include/iostreams.hpp>

int main()
{
    // Say hello to the world!
    hpx::cout << "Hello World!\n" << hpx::flush;
    return 0;
}
//]
```

Copy the content of this program into a file called hello_world.cpp.

Now, in the directory where you put hello_world.cpp, create a Makefile. Add the following code:

```
CXX=(CXX) # Add your favourite compiler here or let makefile choose default.

CXXFLAGS=-O3 -std=c++17

BOOST_ROOT=/path/to/boost
HWLOC_ROOT=/path/to/hwloc
TCMALLOC_ROOT=/path/to/tcmalloc
HPX_ROOT=/path/to/hpx

INCLUDE_DIRECTIVES=$(HPX_ROOT)/include $(BOOST_ROOT)/include $(HWLOC_ROOT)/include

LIBRARY_DIRECTIVES=-L$(HPX_ROOT)/lib $(HPX_ROOT)/lib/libhpx_init.a $(HPX_ROOT)/lib/
↳ libhpx.so $(BOOST_ROOT)/lib/libboost_atomic-mt.so $(BOOST_ROOT)/lib/libboost_
↳ filesystem-mt.so $(BOOST_ROOT)/lib/libboost_program_options-mt.so $(BOOST_ROOT)/lib/
↳ libboost_regex-mt.so $(BOOST_ROOT)/lib/libboost_system-mt.so -lpthread $(TCMALLOC_
↳ ROOT)/libtcmalloc_minimal.so $(HWLOC_ROOT)/libhwloc.so -ldl -lrt

LINK_FLAGS=$(HPX_ROOT)/lib/libhpx_wrap.a -Wl,-wrap=main # should be left empty for
↳ HPX_WITH_HPX_MAIN=OFF

hello_world: hello_world.o
    $(CXX) $(CXXFLAGS) -o hello_world hello_world.o $(LIBRARY_DIRECTIVES) $(LINK_FLAGS)

hello_world.o:
    $(CXX) $(CXXFLAGS) -c -o hello_world.o hello_world.cpp $(INCLUDE_DIRECTIVES)
```

Important: LINK_FLAGS should be left empty if HPX_WITH_HPX_MAIN is set to OFF. Boost in the above example is build with --layout=tagged. Actual Boost flags may vary on your build of Boost.

To build the program, type:

```
make
```

A successful build should result in hello_world binary. To test, type:

```
./hello_world
```

How to build HPX components with makefile

Let's try a more complex example involving an *HPX* component. An *HPX* component is a class that exposes *HPX* actions. *HPX* components are compiled into dynamically-loaded modules called component libraries. Here's the source code:

hello_world_component.cpp

```
//[hello_world_cpp_getting_started
#include "hello_world_component.hpp"
#include <hpx/include/iostreams.hpp>

#include <iostream>

namespace examples { namespace server
{
    void hello_world::invoke()
    {
        hpx::cout << "Hello HPX World!" << std::endl;
    }
} }

HPX_REGISTER_COMPONENT_MODULE();

typedef hpx::components::component<
    examples::server::hello_world
> hello_world_type;

HPX_REGISTER_COMPONENT(hello_world_type, hello_world);

HPX_REGISTER_ACTION(
```

hello_world_component.hpp

```
//[hello_world_hpp_getting_started
#ifndef HELLO_WORLD_COMPONENT_HPP
#define HELLO_WORLD_COMPONENT_HPP

#include <hpx/hpx.hpp>
#include <hpx/include/actions.hpp>
#include <hpx/include/lcos.hpp>
#include <hpx/include/components.hpp>
#include <hpx/include/serialization.hpp>

#include <utility>

namespace examples { namespace server
{
    struct HPX_COMPONENT_EXPORT hello_world
        : hpx::components::component_base<hello_world>
    {
        void invoke();
        HPX_DEFINE_COMPONENT_ACTION(hello_world, invoke);
    };
} }

HPX_REGISTER_ACTION_DECLARATION(
    examples::server::hello_world::invoke_action, hello_world_invoke_action);
```

(continues on next page)

(continued from previous page)

```

namespace examples
{
    struct hello_world
        : hpx::components::client_base<hello_world, server::hello_world>
    {
        typedef hpx::components::client_base<hello_world, server::hello_world>
            base_type;

        hello_world(hpx::future<hpx::naming::id_type> && f)
            : base_type(std::move(f))
        {}

        hello_world(hpx::naming::id_type && f)
            : base_type(std::move(f))
        {}

        void invoke()
        {
            hpx::async<server::hello_world::invoke_action>(*this->get_id()).get();
        }
    };
}

```

hello_world_client.cpp

```

// Copyright (c) 2012 Bryce Lelbach
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

//[hello_world_client_getting_started
#include "hello_world_component.hpp"
#include <hpx/hpx_init.hpp>

int hpx_main(hpx::program_options::variables_map&)
{
{
    // Create a single instance of the component on this locality.
    examples::hello_world client =
        hpx::new<examples::hello_world>(hpx::find_here());

    // Invoke the component's action, which will print "Hello World!".
    client.invoke();
}

    return hpx::finalize(); // Initiate shutdown of the runtime system.
}

int main(int argc, char* argv[])
{
    return hpx::init(argc, argv); // Initialize and run HPX.
}
//]

```

Now, in the directory, create a Makefile. Add the following code:

```

CXX=(CXX) # Add your favourite compiler here or let makefile choose default.

CXXFLAGS=-O3 -std=c++17

BOOST_ROOT=/path/to/boost
HWLOC_ROOT=/path/to/hwloc
TCMALLOC_ROOT=/path/to/tcmalloc
HPX_ROOT=/path/to/hpx

INCLUDE_DIRECTIVES=$(HPX_ROOT)/include $(BOOST_ROOT)/include $(HWLOC_ROOT)/include

LIBRARY_DIRECTIVES=-L$(HPX_ROOT)/lib $(HPX_ROOT)/lib/libhpx_init.a $(HPX_ROOT)/lib/
→ libhpx.so $(BOOST_ROOT)/lib/libboost_atomic-mt.so $(BOOST_ROOT)/lib/libboost_
→ filesystem-mt.so $(BOOST_ROOT)/lib/libboost_program_options-mt.so $(BOOST_ROOT)/lib/
→ libboost_regex-mt.so $(BOOST_ROOT)/lib/libboost_system-mt.so -lpthread $(TCMALLOC_
→ ROOT)/libtcmalloc_minimal.so $(HWLOC_ROOT)/libhwloc.so -ldl -lrt

LINK_FLAGS=$(HPX_ROOT)/lib/libhpx_wrap.a -Wl,-wrap=main # should be left empty for_
→ HPX_WITH_HPX_MAIN=OFF

hello_world_client: libhpx_hello_world hello_world_client.o
$(CXX) $(CXXFLAGS) -o hello_world_client $(LIBRARY_DIRECTIVES) libhpx_hello_world
→ $(LINK_FLAGS)

hello_world_client.o: hello_world_client.cpp
$(CXX) $(CXXFLAGS) -o hello_world_client.o hello_world_client.cpp $(INCLUDE_
→ DIRECTIVES)

libhpx_hello_world: hello_world_component.o
$(CXX) $(CXXFLAGS) -o libhpx_hello_world hello_world_component.o $(LIBRARY_
→ DIRECTIVES)

hello_world_component.o: hello_world_component.cpp
$(CXX) $(CXXFLAGS) -c -o hello_world_component.o hello_world_component.cpp
→ $(INCLUDE_DIRECTIVES)

```

To build the program, type:

```
make
```

A successful build should result in `hello_world` binary. To test, type:

```
./hello_world
```

Note: Due to high variations in CMake flags and library dependencies, it is recommended to build *HPX* applications and components with `pkg-config` or `CMakeLists.txt`. Writing `Makefile` may result in broken builds if due care is not taken. `pkg-config` files and CMake systems are configured with CMake build of *HPX*. Hence, they are stable when used together and provide better support overall.

2.5.4 Starting the HPX runtime

In order to write an application which uses services from the *HPX* runtime system you need to initialize the *HPX* library by inserting certain calls into the code of your application. Depending on your use case, this can be done in 3 different ways:

- *Minimally invasive*: Re-use the `main()` function as the main *HPX* entry point.
- *Balanced use case*: Supply your own main *HPX* entry point while blocking the main thread.
- *Most flexibility*: Supply your own main *HPX* entry point while avoiding to block the main thread.
- *Suspend and resume*: As above but suspend and resume the *HPX* runtime to allow for other runtimes to be used.

Re-use the `main()` function as the main *HPX* entry point

This method is the least intrusive to your code. It however provides you with the smallest flexibility in terms of initializing the *HPX* runtime system. The following code snippet shows what a minimal *HPX* application using this technique looks like:

```
#include <hpx/hpx_main.hpp>

int main(int argc, char* argv[])
{
    return 0;
}
```

The only change to your code you have to make is to include the file `hpx/hpx_main.hpp`. In this case the function `main()` will be invoked as the first *HPX* thread of the application. The runtime system will be initialized behind the scenes before the function `main()` is executed and will automatically stop after `main()` has returned. All *HPX* API functions can be used from within this function now.

Note: The function `main()` does not need to expect receiving `argc` `argv` as shown above, but could expose the signature `int main()`. This is consistent with the usually allowed prototypes for the function `main()` in C++ applications.

All command line arguments specific to *HPX* will still be processed by the *HPX* runtime system as usual. However, those command line options will be removed from the list of values passed to `argc/argv` of the function `main()`. The list of values passed to `main()` will hold only the commandline options which are not recognized by the *HPX* runtime system (see the section [HPX Command Line Options](#) for more details on what options are recognized by *HPX*).

Note: In this mode all one-letter-shortcuts are disabled which are normally available on the *HPX* command line (such as `-t` or `-l` see [HPX Command Line Options](#)). This is done to minimize any possible interaction between the command line options recognized by the *HPX* runtime system and any command line options defined by the application.

The value returned from the function `main()` as shown above will be returned to the operating system as usual.

Important: To achieve this seamless integration, the header file `hpx/hpx_main.hpp` defines a macro:

```
#define main hpx_startup::user_main
```

which could result in unexpected behavior.

Important: To achieve this seamless integration, we use different implementations for different Operating Systems. In case of Linux or Mac OSX, the code present in `hpx_wrap.cpp` is put into action. We hook into the system function in case of Linux and provide alternate entry point in case of Mac OSX. For other Operating Systems we rely on a macro:

```
#define main hpx_startup::user_main
```

provided in the header file `hpx/hpx_main.hpp`. This implementation can result in unexpected behavior.

Caution: We make use of an *override* variable `include_libhpx_wrap` in the header file `hpx/hpx_main.hpp` to swiftly choose the function call stack at runtime. Therefore, the header file should *only* be included in the main executable. Including it in the components will result in multiple definition of the variable.

Supply your own main HPX entry point while blocking the main thread

With this method you need to provide an explicit main thread function named `hpx_main` at global scope. This function will be invoked as the main entry point of your *HPX* application on the console `locality` only (this function will be invoked as the first *HPX* thread of your application). All *HPX* API functions can be used from within this function.

The thread executing the function `hpx::init` will block waiting for the runtime system to exit. The value returned from `hpx_main` will be returned from `hpx::init` after the runtime system has stopped.

The function `hpx::finalize` has to be called on one of the *HPX* localities in order to signal that all work has been scheduled and the runtime system should be stopped after the scheduled work has been executed.

This method of invoking *HPX* has the advantage of you being able to decide which version of `hpx::init` to call. This allows to pass additional configuration parameters while initializing the *HPX* runtime system.

```
#include <hpx/hpx_init.hpp>

int hpx_main(int argc, char* argv[])
{
    // Any HPX application logic goes here...
    return hpx::finalize();
}

int main(int argc, char* argv[])
{
    // Initialize HPX, run hpx_main as the first HPX thread, and
    // wait for hpx::finalize being called.
    return hpx::init(argc, argv);
}
```

Note: The function `hpx_main` does not need to expect receiving `argc/argv` as shown above, but could expose one of the following signatures:

```
int hpx_main();
int hpx_main(int argc, char* argv[]);
int hpx_main(hpx::program_options::variables_map& vm);
```

This is consistent with (and extends) the usually allowed prototypes for the function `main()` in C++ applications.

The header file to include for this method of using *HPX* is `hpx/hpx_init.hpp`.

There are many additional overloads of `hpx::init` available, such as for instance to provide your own entry point function instead of `hpx_main`. Please refer to the function documentation for more details (see: `hpx/hpx_init.hpp`).

Supply your own main *HPX* entry point while avoiding to block the main thread

With this method you need to provide an explicit main thread function named `hpx_main` at global scope. This function will be invoked as the main entry point of your *HPX* application on the console `locality` only (this function will be invoked as the first *HPX* thread of your application). All *HPX* API functions can be used from within this function.

The thread executing the function `hpx::start` will *not* block waiting for the runtime system to exit, but will return immediately.

Important: You cannot use any of the *HPX* API functions other than `hpx::stop` from inside your `main()` function.

The function `hpx::finalize` has to be called on one of the *HPX* localities in order to signal that all work has been scheduled and the runtime system should be stopped after the scheduled work has been executed.

This method of invoking *HPX* is useful for applications where the main thread is used for special operations, such as GUIs. The function `hpx::stop` can be used to wait for the *HPX* runtime system to exit and should be at least used as the last function called in `main()`. The value returned from `hpx_main` will be returned from `hpx::stop` after the runtime system has stopped.

```
#include <hpx/hpx_start.hpp>

int hpx_main(int argc, char* argv[])
{
    // Any HPX application logic goes here...
    return hpx::finalize();
}

int main(int argc, char* argv[])
{
    // Initialize HPX, run hpx_main.
    hpx::start(argc, argv);

    // ...Execute other code here...

    // Wait for hpx::finalize being called.
    return hpx::stop();
}
```

Note: The function `hpx_main` does not need to expect receiving `argc/argv` as shown above, but could expose one of the following signatures:

```
int hpx_main();
int hpx_main(int argc, char* argv[]);
int hpx_main(hpx::program_options::variables_map& vm);
```

This is consistent with (and extends) the usually allowed prototypes for the function `main()` in C++ applications.

The header file to include for this method of using *HPX* is `hpx/hpx_start.hpp`.

There are many additional overloads of `hpx::start` available, such as for instance to provide your own entry point function instead of `hpx_main`. Please refer to the function documentation for more details (see: `hpx/hpx_start.hpp`).

Suspending and resuming the *HPX* runtime

In some applications it is required to combine *HPX* with other runtimes. To support this use case *HPX* provides two functions: `hpx::suspend` and `hpx::resume`. `hpx::suspend` is a blocking call which will wait for all scheduled tasks to finish executing and then put the thread pool OS threads to sleep. `hpx::resume` simply wakes up the sleeping threads so that they are ready to accept new work. `hpx::suspend` and `hpx::resume` can be found in the header `hpx/hpx_suspend.hpp`.

```
#include <hpx/hpx_start.hpp>
#include <hpx/hpx_suspend.hpp>

int main(int argc, char* argv[])
{
    // Initialize HPX, don't run hpx_main
    hpx::start(nullptr, argc, argv);

    // Schedule a function on the HPX runtime
    hpx::apply(&my_function, ...);

    // Wait for all tasks to finish, and suspend the HPX runtime
    hpx::suspend();

    // Execute non-HPX code here

    // Resume the HPX runtime
    hpx::resume();

    // Schedule more work on the HPX runtime

    // hpx::finalize has to be called from the HPX runtime before hpx::stop
    hpx::apply([]() { hpx::finalize(); });
    return hpx::stop();
}
```

Note: `hpx::suspend` does not wait for `hpx::finalize` to be called. Only call `hpx::finalize` when you wish to fully stop the *HPX* runtime.

HPX also supports suspending individual thread pools and threads. For details on how to do that see the documentation for `hpx::threads::thread_pool_base`.

Automatically suspending worker threads

The previous method guarantees that the worker threads are suspended when you ask for it and that they stay suspended. An alternative way to achieve the same effect is to tweak how quickly *HPX* suspends its worker threads when they run out of work. The following configuration values make sure that *HPX* idles very quickly:

```
hpx.max_idle_backoff_time = 1000
hpx.max_idle_loop_count = 0
```

They can be set on the command line using `--hpx:ini=hpx.max_idle_backoff_time=1000` and `--hpx:ini=hpx.max_idle_loop_count=0`. See [Launching and configuring HPX applications](#) for more details on how to set configuration parameters.

After setting idling parameters the previous example could now be written like this instead:

```
#include <hpx/hpx_start.hpp>

int main(int argc, char* argv[])
{
    // Initialize HPX, don't run hpx_main
    hpx::start(nullptr, argc, argv);

    // Schedule some functions on the HPX runtime
    // NOTE: run_as_hpx_thread blocks until completion.
    hpx::run_as_hpx_thread(&my_function, ...);
    hpx::run_as_hpx_thread(&my_other_function, ...);

    // hpx::finalize has to be called from the HPX runtime before hpx::stop
    hpx::apply([]() { hpx::finalize(); });
    return hpx::stop();
}
```

In this example each call to `hpx::run_as_hpx_thread` acts as a “parallel region”.

Working of `hpx_main.hpp`

In order to initialize *HPX* from `main()`, we make use of linker tricks.

It is implemented differently for different Operating Systems. Method of implementation is as follows:

- *Linux*: Using linker `--wrap` option.
- *Mac OSX*: Using the linker `-e` option.
- *Windows*: Using `#define main hpx_startup::user_main`

Linux implementation

We make use of the Linux linker `ld`'s `--wrap` option to wrap the `main()` function. This way any call to `main()` are redirected to our own implementation of `main`. It is here that we check for the existence of `hpx_main.hpp` by making use of a shadow variable `include_libhpx_wrap`. The value of this variable determines the function stack at runtime.

The implementation can be found in `libhpx_wrap.a`.

Important: It is necessary that `hpx_main.hpp` be not included more than once. Multiple inclusions can result in multiple definition of `include_libhpx_wrap`.

Mac OSX implementation

Here we make use of yet another linker option `-e` to change the entry point to our custom entry function `initialize_main`. We initialize the *HPX* runtime system from this function and call `main` from the initialized system. We determine the function stack at runtime by making use of the shadow variable `include_libhpx_wrap`.

The implementation can be found in `libhpx_wrap.a`.

Important: It is necessary that `hpx_main.hpp` be not included more than once. Multiple inclusions can result in multiple definition of `include_libhpx_wrap`.

Windows implementation

We make use of a macro `#define main hpx_startup::user_main` to take care of the initializations.

This implementation could result in unexpected behaviors.

2.5.5 Launching and configuring *HPX* applications

Configuring *HPX* applications

All *HPX* applications can be configured using special command line options and/or using special configuration files. This section describes the available options, the configuration file format, and the algorithm used to locate possible predefined configuration files. Additionally this section describes the defaults assumed if no external configuration information is supplied.

During startup any *HPX* application applies a predefined search pattern to locate one or more configuration files. All found files will be read and merged in the sequence they are found into one single internal database holding all configuration properties. This database is used during the execution of the application to configure different aspects of the runtime system.

In addition to the ini files, any application can supply its own configuration files, which will be merged with the configuration database as well. Moreover, the user can specify additional configuration parameters on the command line when executing an application. The *HPX* runtime system will merge all command line configuration options (see the description of the `--hpx:ini`, `--hpx:config`, and `--hpx:app-config` command line options).

The HPX INI File Format

All HPX applications can be configured using a special file format which is similar to the well-known Windows INI file format¹⁰⁹. This is a structured text format allowing to group key/value pairs (properties) into sections. The basic element contained in an ini file is the property. Every property has a name and a value, delimited by an equals sign '='. The name appears to the left of the equals sign:

```
name=value
```

The value may contain equal signs as only the first '=' character is interpreted as the delimiter between name and value. Whitespace before the name, after the value and immediately before and after the delimiting equal sign is ignored. Whitespace inside the value is retained.

Properties may be grouped into arbitrarily named sections. The section name appears on a line by itself, in square brackets [and]. All properties after the section declaration are associated with that section. There is no explicit “end of section” delimiter; sections end at the next section declaration, or the end of the file:

```
[section]
```

In HPX sections can be nested. A nested section has a name composed of all section names it is embedded in. The section names are concatenated using a dot '.':

```
[outer_section.inner_section]
```

Here `inner_section` is logically nested within `outer_section`.

It is possible to use the full section name concatenated with the property name to refer to a particular property. For example in:

```
[a.b.c]
d = e
```

the property value of `d` can be referred to as `a.b.c.d=e`.

In HPX ini files can contain comments. Hash signs '#' at the beginning of a line indicate a comment. All characters starting with the '#' until the end of line are ignored.

If a property with the same name is reused inside a section, the second occurrence of this property name will override the first occurrence (discard the first value). Duplicate sections simply merge their properties together, as if they occurred contiguously.

In HPX ini files, a property value `${FOO:default}` will use the environmental variable `FOO` to extract the actual value if it is set and `default` otherwise. No default has to be specified. Therefore `${FOO}` refers to the environmental variable `FOO`. If `FOO` is not set or empty the overall expression will evaluate to an empty string. A property value `${[section.key]:default}` refers to the value held by the property `section.key` if it exists and `default` otherwise. No default has to be specified. Therefore `${[section.key]}` refers to the property `section.key`. If the property `section.key` is not set or empty, the overall expression will evaluate to an empty string.

Note: Any property `${[section.key]:default}` is evaluated whenever it is queried and not when the configuration data is initialized. This allows for lazy evaluation and relaxes initialization order of different sections. The only exception are recursive property values, e.g. values referring to the very key they are associated with. Those property values are evaluated at initialization time to avoid infinite recursion.

¹⁰⁹ https://en.wikipedia.org/wiki/INI_file

Built-in Default Configuration Settings

During startup any *HPX* application applies a predefined search pattern to locate one or more configuration files. All found files will be read and merged in the sequence they are found into one single internal data structure holding all configuration properties.

As a first step the internal configuration database is filled with a set of default configuration properties. Those settings are described on a section by section basis below.

Note: You can print the default configuration settings used for an executable by specifying the command line option `--hpx:dump-config`.

The system configuration section

```
[system]
pid = <process-id>
prefix = <current prefix path of core HPX library>
executable = <current prefix path of executable>
```

Property	Description
system.pid	This is initialized to store the current OS-process id of the application instance.
system.prefix	This is initialized to the base directory <i>HPX</i> has been loaded from.
system.executable_prefix	This is initialized to the base directory the current executable has been loaded from.

The hpx configuration section

```
[hpx]
location = ${HPX_LOCATION:${system.prefix}}
component_path = ${hpx.location}/lib/hpx:${system.executable_prefix}/lib/hpx:${system.
↳ executable_prefix}/../lib/hpx
master_ini_path = ${hpx.location}/share/hpx-<version>:${system.executable_prefix}/
↳ share/hpx-<version>:${system.executable_prefix}/../share/hpx-<version>
ini_path = ${hpx.master_ini_path}/ini
os_threads = 1
localities = 1
program_name =
cmd_line =
lock_detection = ${HPX_LOCK_DETECTION:0}
throw_on_held_lock = ${HPX_THROW_ON_HELD_LOCK:1}
minimal_deadlock_detection = <debug>
spinlock_deadlock_detection = <debug>
spinlock_deadlock_detection_limit = ${HPX_SPINLOCK_DEADLOCK_DETECTION_LIMIT:1000000}
max_background_threads = ${HPX_MAX_BACKGROUND_THREADS:${hpx.os_threads}}
max_idle_loop_count = ${HPX_MAX_IDLE_LOOP_COUNT:<hpx_idle_loop_count_max>}
max_busy_loop_count = ${HPX_MAX_BUSY_LOOP_COUNT:<hpx_busy_loop_count_max>}
max_idle_backoff_time = ${HPX_MAX_IDLE_BACKOFF_TIME:<hpx_idle_backoff_time_max>}

[hpx.stacks]
small_size = ${HPX_SMALL_STACK_SIZE:<hpx_small_stack_size>}
```

(continues on next page)

(continued from previous page)

```
medium_size = ${HPX_MEDIUM_STACK_SIZE:<hpx_medium_stack_size>}
large_size = ${HPX_LARGE_STACK_SIZE:<hpx_large_stack_size>}
huge_size = ${HPX_HUGE_STACK_SIZE:<hpx_huge_stack_size>}
use_guard_pages = ${HPX_THREAD_GUARD_PAGE:1}
```


Property	Description
hpx.location	This is initialized to the id of the <i>locality</i> this application instance is running on.
hpx.component	Duplicates are discarded. This property can refer to a list of directories separated by ':' (Linux, Android, and MacOS) or using ';' (Windows).
hpx.master_ini	This is initialized to the list of default paths of the main hpx.ini configuration files. This property can refer to a list of directories separated by ':' (Linux, Android, and MacOS) or using ';' (Windows).
hpx.ini_path	This is initialized to the default path where <i>HPX</i> will look for more ini configuration files. This property can refer to a list of directories separated by ':' (Linux, Android, and MacOS) or using ';' (Windows).
hpx.os_threads	This setting reflects the number of OS-threads used for running <i>HPX</i> -threads. Defaults to number of detected cores (not hyperthreads/PUs).
hpx.localities	This setting reflects the number of localities the application is running on. Defaults to 1.
hpx.program_name	This setting reflects the program name of the application instance. Initialized from the command line argv[0].
hpx.cmd_line	This setting reflects the actual command line used to launch this application instance.
hpx.lock_detection	This setting verifies that no locks are being held while a <i>HPX</i> thread is suspended. This setting is applicable only if HPX_WITH_VERIFY_LOCKS is set during configuration in CMake.
hpx.throw_on_error	This setting causes an exception if during lock detection at least one lock is being held while a <i>HPX</i> thread is suspended. This setting is applicable only if HPX_WITH_VERIFY_LOCKS is set during configuration in CMake. This setting has no effect if hpx.lock_detection=0.
hpx.minimal_deadlock_detection	This setting enables support for minimal deadlock detection for <i>HPX</i> -threads. By default this is set to 1 (for Debug builds) or to 0 (for Release, RelWithDebInfo, RelMinSize builds), this setting is effective only if HPX_WITH_THREAD_DEADLOCK_DETECTION is set during configuration in CMake.
hpx.spinlock_detection_limit	This setting verifies that spinlocks don't spin longer than specified using the hpx.spinlock_detection_limit. This setting is applicable only if HPX_WITH_SPINLOCK_DEADLOCK_DETECTION is set during configuration in CMake. By default this is set to 1 (for Debug builds) or to 0 (for Release, RelWithDebInfo, RelMinSize builds).
hpx.spinlock_limit	This setting specifies the upper limit of allowed number of spins that spinlocks are allowed to perform. This setting is applicable only if HPX_WITH_SPINLOCK_DEADLOCK_DETECTION is set during configuration in CMake. By default this is set to 1000000.
hpx.max_background_work	This setting defines the number of threads in the scheduler which are used to execute background work. By default this is the same as the number of cores used for the scheduler.
hpx.max_idle_loops	By default this is defined by the preprocessor constant HPX_IDLE_LOOP_COUNT_MAX. This is an internal setting which you should change only if you know exactly what you are doing.
hpx.max_busy_loops	This setting defines the maximum value of the busy-loop counter in the scheduler. By default this is defined by the preprocessor constant HPX_BUSY_LOOP_COUNT_MAX. This is an internal setting which you should change only if you know exactly what you are doing.
hpx.max_idle_time	This setting defines the maximum time (in milliseconds) for the scheduler to sleep after being idle for hpx.max_idle_loop_count iterations. This setting is applicable only if HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF is set during configuration in CMake. By default this is defined by the preprocessor constant HPX_IDLE_BACKOFF_TIME_MAX. This is an internal setting which you should change only if you know exactly what you are doing.
hpx.stack_size.small_size	This is initialized to the small stack size to be used by <i>HPX</i> -threads. Set by default to the value of the compile time preprocessor constant HPX_SMALL_STACK_SIZE (defaults to 0x8000). This value is used for all <i>HPX</i> threads by default, except for the thread running hpx_main (which runs on a large stack).
hpx.stack_size.medium_size	This is initialized to the medium stack size to be used by <i>HPX</i> -threads. Set by default to the value of the compile time preprocessor constant HPX_MEDIUM_STACK_SIZE (defaults to 0x20000).
hpx.stack_size.large_size	This is initialized to the large stack size to be used by <i>HPX</i> -threads. Set by default to the value of the compile time preprocessor constant HPX_LARGE_STACK_SIZE (defaults to 0x200000). This setting is used by default for the thread running hpx_main only.
hpx.stack_size.huge_size	This is initialized to the huge stack size to be used by <i>HPX</i> -threads. Set by default to the value of the compile time preprocessor constant HPX_HUGE_STACK_SIZE (defaults to 0x4000000).

The `hpx.threadpools` configuration section

```
[hpx.threadpools]
io_pool_size = ${HPX_NUM_IO_POOL_SIZE:2}
parcel_pool_size = ${HPX_NUM_PARCEL_POOL_SIZE:2}
timer_pool_size = ${HPX_NUM_TIMER_POOL_SIZE:2}
```

Property	Description
<code>hpx.threadpools.io_pool_size</code>	The value of this property defines the number of OS-threads created for the internal I/O thread pool.
<code>hpx.threadpools.parcel_pool_size</code>	The value of this property defines the number of OS-threads created for the internal parcel thread pool.
<code>hpx.threadpools.timer_pool_size</code>	The value of this property defines the number of OS-threads created for the internal timer thread pool.

The `hpx.thread_queue` configuration section

Important: These setting control internal values used by the thread scheduling queues in the *HPX* scheduler. You should not modify these settings except if you know exactly what you are doing]

```
[hpx.thread_queue]
min_tasks_to_steal_pending = ${HPX_THREAD_QUEUE_MIN_TASKS_TO_STEAL_PENDING:0}
min_tasks_to_steal_staged = ${HPX_THREAD_QUEUE_MIN_TASKS_TO_STEAL_STAGED:10}
min_add_new_count = ${HPX_THREAD_QUEUE_MIN_ADD_NEW_COUNT:10}
max_add_new_count = ${HPX_THREAD_QUEUE_MAX_ADD_NEW_COUNT:10}
max_delete_count = ${HPX_THREAD_QUEUE_MAX_DELETE_COUNT:1000}
```

Property	Description
<code>hpx.thread_queue.min_tasks_to_stea</code>	The value of this property defines the number of pending <i>HPX</i> threads which have to be available before neighboring cores are allowed to steal work. The default is to allow <code>stealing always</code> .
<code>hpx.thread_queue.min_tasks_to_stea</code>	The value of this property defines the number of staged <i>HPX</i> tasks have which to be available before neighboring cores are allowed to steal work. The default is to allow <code>stealing only</code> if there are more tan 10 tasks available.
<code>hpx.thread_queue.min_add_new_count</code>	The value of this property defines the minimal number tasks to be converted into <i>HPX</i> threads whenever the thread queues for a core have run empty.
<code>hpx.thread_queue.max_add_new_count</code>	The value of this property defines the maximal number tasks to be converted into <i>HPX</i> threads whenever the thread queues for a core have run empty.
<code>hpx.thread_queue.max_delete_count</code>	The value of this property defines the number of terminated <i>HPX</i> threads to discard during each invocation of the corresponding function.

The `hpx.components` configuration section

```
[hpx.components]
load_external = ${HPX_LOAD_EXTERNAL_COMPONENTS:1}
```

Property	Description
<code>hpx.components.load_external</code>	This entry defines whether external components will be loaded on this <i>locality</i> . This entry normally is set to 1 and usually there is no need to directly change this value. It is automatically set to 0 for a dedicated <i>AGAS</i> server <i>locality</i> .

Additionally, the section `hpx.components` will be populated with the information gathered from all found components. The information loaded for each of the components will contain at least the following properties:

```
[hpx.components.<component_instance_name>]
name = <component_name>
path = <full_path_of_the_component_module>
enabled = ${hpx.components.load_external}]
```

Property	Description
<code>hpx.components.<component_instance_name>.name</code>	This is the name of a component, usually the same as the second argument to the macro used while registering the component with <i>HPX_REGISTER_COMPONENT</i> . Set by the component factory.
<code>hpx.components.<component_instance_name>.path</code>	This is either the full path file name of the component module or the directory the component module is located in. In this case, the component module name will be derived from the property <code>hpx.components.<component_instance_name>.name</code> . Set by the component factory.
<code>hpx.components.<component_instance_name>.enabled</code>	This setting explicitly enables or disables the component. This is an optional property, <i>HPX</i> assumed that the component is enabled if it is not defined.

The value for `<component_instance_name>` is usually the same as for the corresponding `name` property. However generally it can be defined to any arbitrary instance name. It is used to distinguish between different ini sections, one for each component.

The `hpx.parcel` configuration section

```
[hpx.parcel]
address = ${HPX_PARCEL_SERVER_ADDRESS:<hpx_initial_ip_address>}
port = ${HPX_PARCEL_SERVER_PORT:<hpx_initial_ip_port>}
bootstrap = ${HPX_PARCEL_BOOTSTRAP:<hpx_parcel_bootstrap>}
max_connections = ${HPX_PARCEL_MAX_CONNECTIONS:<hpx_parcel_max_connections>}
max_connections_per_locality = ${HPX_PARCEL_MAX_CONNECTIONS_PER_LOCALITY:<hpx_parcel_max_connections_per_locality>}
max_message_size = ${HPX_PARCEL_MAX_MESSAGE_SIZE:<hpx_parcel_max_message_size>}
max_outbound_message_size = ${HPX_PARCEL_MAX_OUTBOUND_MESSAGE_SIZE:<hpx_parcel_max_outbound_message_size>}
array_optimization = ${HPX_PARCEL_ARRAY_OPTIMIZATION:1}
zero_copy_optimization = ${HPX_PARCEL_ZERO_COPY_OPTIMIZATION:[hpx.parcel.array_optimization]}
```

(continues on next page)

(continued from previous page)

```
async_serialization = ${HPX_PARCEL_ASYNC_SERIALIZATION:1}
message_handlers = ${HPX_PARCEL_MESSAGE_HANDLERS:0}
```

Property	Description
hpx.parcel.address	This property defines the default IP address to be used for the <i>parcel</i> layer to listen to. This IP address will be used as long as no other values are specified (for instance using the <code>--hpx:hpx</code> command line option). The expected format is any valid IP address or domain name format which can be resolved into an IP address. The default depends on the compile time preprocessor constant <code>HPX_INITIAL_IP_ADDRESS</code> ("127.0.0.1").
hpx.parcel.port	This property defines the default IP port to be used for the <i>parcel</i> layer to listen to. This IP port will be used as long as no other values are specified (for instance using the <code>--hpx:hpx</code> command line option). The default depends on the compile time preprocessor constant <code>HPX_INITIAL_IP_PORT</code> (7910).
hpx.parcel.bootstrap	This property defines which parcelport type should be used during application bootstrap. The default depends on the compile time preprocessor constant <code>HPX_PARCEL_BOOTSTRAP</code> ("tcp").
hpx.parcel.max_connections	This property defines how many network connections between different localities are overall kept alive by each of <i>locality</i> . The default depends on the compile time preprocessor constant <code>HPX_PARCEL_MAX_CONNECTIONS</code> (512).
hpx.parcel.max_connections_per_locality	This property defines the maximum number of network connections that one <i>locality</i> will open to another <i>locality</i> . The default depends on the compile time preprocessor constant <code>HPX_PARCEL_MAX_CONNECTIONS_PER_LOCALITY</code> (4).
hpx.parcel.max_message_size	This property defines the maximum allowed message size which will be transferrable through the <i>parcel</i> layer. The default depends on the compile time preprocessor constant <code>HPX_PARCEL_MAX_MESSAGE_SIZE</code> (100000000 bytes).
hpx.parcel.max_outbound_message_size	This property defines the maximum allowed outbound coalesced message size which will be transferrable through the <i>parcel</i> layer. The default depends on the compile time preprocessor constant <code>HPX_PARCEL_MAX_OUTBOUND_MESSAGE_SIZE</code> (1000000 bytes).
hpx.parcel.array_optimization	This property defines whether this <i>locality</i> is allowed to utilize array optimizations during serialization of <i>parcel</i> data. The default is 1.
hpx.parcel.zero_copy_optimization	This property defines whether this <i>locality</i> is allowed to utilize zero copy optimizations during serialization of <i>parcel</i> data. The default is the same value as set for <code>hpx.parcel.array_optimization</code> .
hpx.parcel.async_serialization	This property defines whether this <i>locality</i> is allowed to spawn a new thread for serialization (this is both for encoding and decoding parcels). The default is 1.
hpx.parcel.message_handlers	This property defines whether message handlers are loaded. The default is 0.

The following settings relate to the TCP/IP parcelport.

```
[hpx.parcel.tcp]
enable = ${HPX_HAVE_PARCELPORT_TCP:$[hpx.parcel.enabled]}
array_optimization = ${HPX_PARCEL_TCP_ARRAY_OPTIMIZATION:$[hpx.parcel.array_
optimization]}
zero_copy_optimization = ${HPX_PARCEL_TCP_ZERO_COPY_OPTIMIZATION:$[hpx.parcel.zero_
copy_optimization]}
async_serialization = ${HPX_PARCEL_TCP_ASYNC_SERIALIZATION:$[hpx.parcel.async_
serialization]}
parcel_pool_size = ${HPX_PARCEL_TCP_PARCEL_POOL_SIZE:$[hpx.threadpools.parcel_pool_
size]}
```

(continues on next page)

(continued from previous page)

```
max_connections = ${HPX_PARCEL_TCP_MAX_CONNECTIONS:[hpx.parcel.max_connections]}
max_connections_per_locality = ${HPX_PARCEL_TCP_MAX_CONNECTIONS_PER_LOCALITY:[hpx.
    ↪parcel.max_connections_per_locality]}
max_message_size = ${HPX_PARCEL_TCP_MAX_MESSAGE_SIZE:[hpx.parcel.max_message_size]}
max_outbound_message_size = ${HPX_PARCEL_TCP_MAX_OUTBOUND_MESSAGE_SIZE:[hpx.parcel.
    ↪max_outbound_message_size]}
```

Property	Description
hpx.parcel. tcp.enable	Enable the use of the default TCP parcelport. Note that the initial bootstrap of the overall HPX application will be performed using the default TCP connections. This parcelport is enabled by default. This will be disabled only if MPI is enabled (see below).
hpx. parcel.tcp. array_optimization	This property defines whether this <i>locality</i> is allowed to utilize array optimizations in the TCP/IP parcelport during serialization of parcel data. The default is the same value as set for hpx.parcel.array_optimization.
hpx. parcel.tcp. zero_copy_optimization	This property defines whether this <i>locality</i> is allowed to utilize zero copy optimizations in the TCP/IP parcelport during serialization of parcel data. The default is the same value as set for hpx.parcel.zero_copy_optimization.
hpx. parcel.tcp. async_serialization	This property defines whether this <i>locality</i> is allowed to spawn a new thread for serialization in the TCP/IP parcelport (this is both for encoding and decoding parcels). The default is the same value as set for hpx.parcel.async_serialization.
hpx. parcel.tcp. parcel_pool_size	The value of this property defines the number of OS-threads created for the internal parcel thread pool of the TCP <i>parcel</i> port. The default is taken from hpx.threadpools.parcel_pool_size.
hpx. parcel.tcp. max_connections	This property defines how many network connections between different localities are overall kept alive by each of <i>locality</i> . The default is taken from hpx.parcel.max_connections.
hpx. parcel.tcp. max_connections_per_locality	This property defines the maximum number of network connections that one <i>locality</i> will open to another <i>locality</i> . The default is taken from hpx.parcel.max_connections_per_locality.
hpx. parcel.tcp. max_message_size	This property defines the maximum allowed message size which will be transferrable through the <i>parcel</i> layer. The default is taken from hpx.parcel.max_message_size.
hpx. parcel.tcp. max_outbound_message_size	This property defines the maximum allowed outbound coalesced message size which will be transferrable through the <i>parcel</i> layer. The default is taken from hpx.parcel.max_outbound_message_size.

The following settings relate to the MPI parcelport. These settings take effect only if the compile time constant `HPX_HAVE_PARCELPORT_MPI` is set (the equivalent cmake variable is `HPX_WITH_PARCELPORT_MPI` and has to be set to ON).

```
[hpx.parcel.mpi]
enable = ${HPX_HAVE_PARCELPORT_MPI:[hpx.parcel.enabled]}
env = ${HPX_HAVE_PARCELPORT_MPI_ENV:MV2_COMM_WORLD_RANK,PMI_RANK,OMPI_COMM_WORLD_SIZE,
    ↪ALPS_APP_PE}
multithreaded = ${HPX_HAVE_PARCELPORT_MPI_MULTITHREADED:0}
rank = <MPI_rank>
processor_name = <MPI_processor_name>
array_optimization = ${HPX_HAVE_PARCEL_MPI_ARRAY_OPTIMIZATION:[hpx.parcel.array_
    ↪optimization]}
zero_copy_optimization = ${HPX_HAVE_PARCEL_MPI_ZERO_COPY_OPTIMIZATION:[hpx.parcel.
    ↪zero_copy_optimization]}
use_io_pool = ${HPX_HAVE_PARCEL_MPI_USE_IO_POOL:$1}
```

(continues on next page)

(continued from previous page)

```
async_serialization = ${HPX_HAVE_PARCEL_MPI_ASYNC_SERIALIZATION:[hpx.parcel.async_
˓→serialization]}
parcel_pool_size = ${HPX_HAVE_PARCEL_MPI_PARCEL_POOL_SIZE:[hpx.threadpools.parcel_
˓→pool_size]}
max_connections = ${HPX_HAVE_PARCEL_MPI_MAX_CONNECTIONS:[hpx.parcel.max_
˓→connections]}
max_connections_per_locality = ${HPX_HAVE_PARCEL_MPI_MAX_CONNECTIONS_PER_LOCALITY:
˓→${[hpx.parcel.max_connections_per_locality]}}
max_message_size = ${HPX_HAVE_PARCEL_MPI_MAX_MESSAGE_SIZE:[hpx.parcel.max_message_
˓→size]}
max_outbound_message_size = ${HPX_HAVE_PARCEL_MPI_MAX_OUTBOUND_MESSAGE_SIZE:[hpx.
˓→parcel.max_outbound_message_size]}
```

Property	Description
hpx.parcel.mpi.enable	Enable the use of the MPI parcelport. HPX tries to detect if the application was started within a parallel MPI environment. If the detection was successful, the MPI parcelport is enabled by default. To explicitly disable the MPI parcelport, set to 0. Note that the initial bootstrap of the overall <i>HPX</i> application will be performed using MPI as well.
hpx.parcel.mpi.env	This property influences which environment variables (comma separated) will be analyzed to find out whether the application was invoked by MPI.
hpx.parcel.mpi.multithreaded	This property is used to determine what threading mode to use when initializing MPI. If this setting is 0 <i>HPX</i> will initialize MPI with <code>MPI_THREAD_SINGLE</code> if the value is not equal to 0 <i>HPX</i> will initialize MPI with <code>MPI_THREAD_MULTI</code> .
hpx.parcel.mpi.rank	This property will be initialized to the MPI rank of the <i>locality</i> .
hpx.parcel.mpi.processor_name	This property will be initialized to the MPI processor name of the <i>locality</i> .
hpx.parcel.mpi.array_optimization	This property defines whether this <i>locality</i> is allowed to utilize array optimizations in the MPI parcelport during serialization of <i>parcel</i> data. The default is the same value as set for <code>hpx.parcel.array_optimization</code> .
hpx.parcel.mpi.zero_copy_optimization	This property defines whether this <i>locality</i> is allowed to utilize zero copy optimizations in the MPI parcelport during serialization of parcel data. The default is the same value as set for <code>hpx.parcel.zero_copy_optimization</code> .
hpx.parcel.mpi.use_io_pool	This property can be set to run the progress thread inside of HPX threads instead of a separate thread pool. The default is 1.
hpx.parcel.mpi.async_serialization	This property defines whether this <i>locality</i> is allowed to spawn a new thread for serialization in the MPI parcelport (this is both for encoding and decoding parcels). The default is the same value as set for <code>hpx.parcel.async_serialization</code> .
hpx.parcel.mpi.parcel_pool_size	The value of this property defines the number of OS-threads created for the internal parcel thread pool of the MPI <i>parcel</i> port. The default is taken from <code>hpx.threadpools.parcel_pool_size</code> .
hpx.parcel.mpi.max_connections	This property defines how many network connections between different localities are overall kept alive by each of <i>locality</i> . The default is taken from <code>hpx.parcel.max_connections</code> .
hpx.parcel.mpi.max_connections_per_locality	This property defines the maximum number of network connections that one <i>locality</i> will open to another <i>locality</i> . The default is taken from <code>hpx.parcel.max_connections_per_locality</code> .
hpx.parcel.mpi.max_message_size	This property defines the maximum allowed message size which will be transferrable through the <i>parcel</i> layer. The default is taken from <code>hpx.parcel.max_message_size</code> .
hpx.parcel.mpi.max_outbound_message_size	This property defines the maximum allowed outbound coalesced message size which will be transferrable through the <i>parcel</i> layer. The default is taken from <code>hpx.parcel.max_outbound_message_size</code> .

The `hpx.agas` configuration section

```
[hpx.agas]
address = ${HPX_AGAS_SERVER_ADDRESS:<hpx_initial_ip_address>}
port = ${HPX_AGAS_SERVER_PORT:<hpx_initial_ip_port>}
service_mode = hosted
dedicated_server = 0
max_pending_refcnt_requests = ${HPX_AGAS_MAX_PENDING_REFCNT_REQUESTS:<hpx_initial_
    ↴agас_max_pending_refcnt_requests>}
use_caching = ${HPX_AGAS_USE_CACHING:1}
use_range_caching = ${HPX_AGAS_USE_RANGE_CACHING:1}
local_cache_size = ${HPX_AGAS_LOCAL_CACHE_SIZE:<hpx_agas_local_cache_size>}
```

Property	Description
<code>hpx.agas.address</code>	This property defines the default IP address to be used for the <code>AGAS</code> root server. This IP address will be used as long as no other values are specified (for instance using the <code>--hpx:agas</code> command line option). The expected format is any valid IP address or domain name format which can be resolved into an IP address. The default depends on the compile time preprocessor constant <code>HPX_INITIAL_IP_ADDRESS</code> ("127.0.0.1").
<code>hpx.agas.port</code>	This property defines the default IP port to be used for the <code>AGAS</code> root server. This IP port will be used as long as no other values are specified (for instance using the <code>--hpx:agas</code> command line option). The default depends on the compile time preprocessor constant <code>HPX_INITIAL_IP_PORT</code> (7009).
<code>hpx.agas.service_mode</code>	This property specifies what type of <code>AGAS</code> service is running on this <code>locality</code> . Currently, two modes exist. The <code>locality</code> that acts as the <code>AGAS</code> server runs in <code>bootstrap</code> mode. All other localities are in <code>hosted</code> mode.
<code>hpx.agas.dedicated</code>	This property specifies whether the <code>AGAS</code> server is exclusively running <code>AGAS</code> services and not hosting any application components. It is a boolean value. Set to 1 if <code>separate-run-agas-server-only</code> is present.
<code>hpx.agas.max_pending_refcnt_requests</code>	This property defines the number of reference counting requests (increments or decrements) to buffer. The default depends on the compile time preprocessor constant <code>HPX_DEFAULT_MAX_AGAS_MAX_PENDING_REFCNT_REQUESTS</code> (4096).
<code>hpx.agas.use_caching</code>	This property specifies whether a software address translation cache is used. It is a boolean value. Defaults to 1.
<code>hpx.agas.use_range_caching</code>	This property specifies whether range-based caching is used by the software address translation cache. This property is ignored if <code>hpx.agas.use_caching</code> is false. It is a boolean value. Defaults to 1.
<code>hpx.agas.local_cache_size</code>	This property defines the size of the software address translation cache for <code>AGAS</code> services. This property is ignored if <code>hpx.agas.use_caching</code> is false. Note that if <code>hpx.agas.use_range_caching</code> is true, this size will refer to the maximum number of ranges stored in the cache, not the number of entries spanned by the cache. The default depends on the compile time preprocessor constant <code>HPX_AGAS_LOCAL_CACHE_SIZE</code> (4096).

The `hpx.commandline` configuration section

The following table lists the definition of all pre-defined command line option shortcuts. For more information about commandline options see the section [HPX Command Line Options](#).

```
[hpx.commandline]
aliasing = ${HPX_COMMANDLINE_ALIASING:1}
allow_unknown = ${HPX_COMMANDLINE_ALLOW_UNKNOWN:0}

[hpx.commandline.aliases]
-a = --hpx:agas
-c = --hpx:console
-h = --hpx:help
-I = --hpx:ini
-l = --hpx:localities
-p = --hpx:app-config
-q = --hpx:queuing
-r = --hpx:run-agas-server
-t = --hpx:threads
-v = --hpx:version
-w = --hpx:worker
-x = --hpx:hpx
-0 = --hpx:node=0
-1 = --hpx:node=1
-2 = --hpx:node=2
-3 = --hpx:node=3
-4 = --hpx:node=4
-5 = --hpx:node=5
-6 = --hpx:node=6
-7 = --hpx:node=7
-8 = --hpx:node=8
-9 = --hpx:node=9
```


Property	Description
hpx.commandline. aliasing	Enable command line aliases as defined in the section <code>hpx.commandline.aliases</code> (see below). Defaults to 1.
hpx.commandline. allow_unknown	Allow for unknown command line options to be passed through to <code>hpx_main()</code> . Defaults to 0.
hpx.commandline. aliases.-a	On the commandline, <code>-a</code> expands to: <code>--hpx:agas</code> .
hpx.commandline. aliases.-c	On the commandline, <code>-c</code> expands to: <code>--hpx:console</code> .
hpx.commandline. aliases.-h	On the commandline, <code>-h</code> expands to: <code>--hpx:help</code> .
hpx.commandline. aliases.--help	On the commandline, <code>--help</code> expands to: <code>--hpx:help</code> .
hpx.commandline. aliases.-I	On the commandline, <code>-I</code> expands to: <code>--hpx:ini</code> .
hpx.commandline. aliases.-l	On the commandline, <code>-l</code> expands to: <code>--hpx:localities</code> .
hpx.commandline. aliases.-p	On the commandline, <code>-p</code> expands to: <code>--hpx:app-config</code> .
hpx.commandline. aliases.-q	On the commandline, <code>-q</code> expands to: <code>--hpx:queuing</code> .
hpx.commandline. aliases.-r	On the commandline, <code>-r</code> expands to: <code>--hpx:run-agas-server</code> .
hpx.commandline. aliases.-t	On the commandline, <code>-t</code> expands to: <code>--hpx:threads</code> .
hpx.commandline. aliases.-v	On the commandline, <code>-v</code> expands to: <code>--hpx:version</code> .
hpx.commandline. aliases.--version	On the commandline, <code>--version</code> expands to: <code>--hpx:version</code> .
hpx.commandline. aliases.-w	On the commandline, <code>-w</code> expands to: <code>--hpx:worker</code> .
hpx.commandline. aliases.-x	On the commandline, <code>-x</code> expands to: <code>--hpx:hpx</code> .
hpx.commandline. aliases.-0	On the commandline, <code>-0</code> expands to: <code>--hpx:node=0</code> .
hpx.commandline. aliases.-1	On the commandline, <code>-1</code> expands to: <code>--hpx:node=1</code> .
hpx.commandline. aliases.-2	On the commandline, <code>-2</code> expands to: <code>--hpx:node=2</code> .
hpx.commandline. aliases.-3	On the commandline, <code>-3</code> expands to: <code>--hpx:node=3</code> .
hpx.commandline. aliases.-4	On the commandline, <code>-4</code> expands to: <code>--hpx:node=4</code> .
hpx.commandline. aliases.-5	On the commandline, <code>-5</code> expands to: <code>--hpx:node=5</code> .
hpx.commandline. aliases.-6	On the commandline, <code>-6</code> expands to: <code>--hpx:node=6</code> .
hpx.commandline. aliases.-7	On the commandline, <code>-7</code> expands to: <code>--hpx:node=7</code> .
hpx.commandline. aliases.-8	On the commandline, <code>-8</code> expands to: <code>--hpx:node=8</code> .
hpx.commandline. aliases.-9	On the commandline, <code>-9</code> expands to: <code>--hpx:node=9</code> .

Loading INI files

During startup and after the internal database has been initialized as described in the section [Built-in Default Configuration Settings](#), HPX will try to locate and load additional ini files to be used as a source for configuration properties. This allows for a wide spectrum of additional customization possibilities by the user and system administrators. The sequence of locations where HPX will try loading the ini files is well defined and documented in this section. All ini files found are merged into the internal configuration database. The merge operation itself conforms to the rules as described in the section [The HPX INI File Format](#).

1. Load all component shared libraries found in the directories specified by the property `hpx.component_path` and retrieve their default configuration information (see section [Loading components](#) for more details). This property can refer to a list of directories separated by ':' (Linux, Android, and MacOS) or using ';' (Windows).
2. Load all files named `hpx.ini` in the directories referenced by the property `hpx.master_ini_path`. This property can refer to a list of directories separated by ':' (Linux, Android, and MacOS) or using ';' (Windows).
3. Load a file named `.hpx.ini` in the current working directory, e.g. the directory the application was invoked from.
4. Load a file referenced by the environment variable `HPX_INI`. This variable is expected to provide the full path name of the ini configuration file (if any).
5. Load a file named `/etc/hpx.ini`. This lookup is done on non-Windows systems only.
6. Load a file named `.hpx.ini` in the home directory of the current user, e.g. the directory referenced by the environment variable `HOME`.
7. Load a file named `.hpx.ini` in the directory referenced by the environment variable `PWD`.
8. Load the file specified on the command line using the option `--hpx:config`.
9. Load all properties specified on the command line using the option `--hpx:ini`. The properties will be added to the database in the same sequence as they are specified on the command line. The format for those options is for instance `--hpx:ini=hpx.default_stack_size=0x4000`. In addition to the explicit command line options, this will set the following properties as implied from other settings:
 - `hpx.parcel.address` and `hpx.parcel.port` as set by `--hpx:hpx`
 - `hpx.agas.address`, `hpx.agas.port` and `hpx.agas.service_mode` as set by `--hpx:agas`
 - `hpx.program_name` and `hpx.cmd_line` will be derived from the actual command line
 - **`hpx.os_threads` and `hpx.localities` as set by `--hpx:threads` and `--hpx:localities`**
 - `hpx.runtime_mode` will be derived from any explicit `--hpx:console`, `--hpx:worker`, or `--hpx:connect`, or it will be derived from other settings, such as `--hpx:node=0` which implies `--hpx:console`
10. Load files based on the pattern `*.ini` in all directories listed by the property `hpx.ini_path`. All files found during this search will be merged. The property `hpx.ini_path` can hold a list of directories separated by ':' (on Linux or Mac) or ';' (on Windows).
11. Load the file specified on the command line using the option `--hpx:app-config`. Note that this file will be merged as the content for a top level section [application].

Note: Any changes made to the configuration database caused by one of the steps will influence the loading process for all subsequent steps. For instance, if one of the ini files loaded changes the property `hpx.ini_path` this will

influence the directories searched in step 9 as described above.

Important: The *HPX* core library will verify that all configuration settings specified on the command line (using the `--hpx:ini` option) will be checked for validity. That means that the library will accept only *known* configuration settings. This is to protect the user from unintentional typos while specifying those settings. This behavior can be overwritten by appending a '!' to the configuration key, thus forcing the setting to be entered into the configuration database, for instance: `--hpx:ini=hpx.foo! = 1`

If any of the environment variables or files listed above is not found the corresponding loading step will be silently skipped.

Loading components

HPX relies on loading application specific components during the runtime of an application. Moreover, *HPX* comes with a set of preinstalled components supporting basic functionalities useful for almost every application. Any component in *HPX* is loaded from a shared library, where any of the shared libraries can contain more than one component type. During startup, *HPX* tries to locate all available components (e.g. their corresponding shared libraries) and creates an internal component registry for later use. This section describes the algorithm used by *HPX* to locate all relevant shared libraries on a system. As described, this algorithm is customizable by the configuration properties loaded from the ini files (see section [Loading INI files](#)).

Loading components is a two stage process. First *HPX* tries to locate all component shared libraries, loads those, and generates default configuration section in the internal configuration database for each component found. For each found component the following information is generated:

```
[hpx.components.<component_instance_name>]
name = <name_of_shared_library>
path = ${component_path}
enabled = ${hpx.components.load_external}
default = 1
```

The values in this section correspond to the expected configuration information for a component as described in the section [Built-in Default Configuration Settings](#).

In order to locate component shared libraries, *HPX* will try loading all shared libraries (files with the platform specific extension of a shared library, Linux: *.so, Windows: *.dll, MacOS: *.dylib found in the directory referenced by the ini property `hpx.component_path`).

This first step corresponds to step 1) during the process of filling the internal configuration database with default information as described in section [Loading INI files](#).

After all of the configuration information has been loaded, *HPX* performs the second step in terms of loading components. During this step, *HPX* scans all existing configuration sections `[hpx.component.<some_component_instance_name>]` and instantiates a special factory object for each of the successfully located and loaded components. During the application's life time, these factory objects will be responsible to create new and discard old instances of the component they are associated with. This step is performed after step 11) of the process of filling the internal configuration database with default information as described in section [Loading INI files](#).

Application specific component example

In this section we assume to have a simple application component which exposes one member function as a component action. The header file `app_server.hpp` declares the C++ type to be exposed as a component. This type has a member function `print_greeting()` which is exposed as an action `print_greeting_action`. We assume the source files for this example are located in a directory referenced by `$APP_ROOT`:

```
// file: $APP_ROOT/app_server.hpp
#include <hpx/hpx.hpp>
#include <hpx/include/iostreams.hpp>

namespace app
{
    // Define a simple component exposing one action 'print_greeting'
    class HPX_COMPONENT_EXPORT server
        : public hpx::components::component_base<server>
    {
        void print_greeting ()
        {
            hpx::cout << "Hey, how are you?\n" << hpx::flush;
        }

        // Component actions need to be declared, this also defines the
        // type 'print_greeting_action' representing the action.
        HPX_DEFINE_COMPONENT_ACTION(server, print_greeting, print_greeting_action);
    };
}

// Declare boilerplate code required for each of the component actions.
HPX_REGISTER_ACTION_DECLARATION(app::server::print_greeting_action);
```

The corresponding source file contains mainly macro invocations which define boilerplate code needed for *HPX* to function properly:

```
// file: $APP_ROOT/app_server.cpp
#include "app_server.hpp"

// Define boilerplate required once per component module.
HPX_REGISTER_COMPONENT_MODULE();

// Define factory object associated with our component of type 'app::server'.
HPX_REGISTER_COMPONENT(app::server, app_server);

// Define boilerplate code required for each of the component actions. Use the
// same argument as used for HPX_REGISTER_ACTION_DECLARATION above.
HPX_REGISTER_ACTION(app::server::print_greeting_action);
```

The following gives an example of how the component can be used. We create one instance of the `app::server` component on the current *locality* and invoke the exposed action `print_greeting_action` using the global id of the newly created instance. Note, that no special code is required to delete the component instance after it is not needed anymore. It will be deleted automatically when its last reference goes out of scope, here at the closing brace of the block surrounding the code:

```
// file: $APP_ROOT/use_app_server_example.cpp
#include <hpx/hpx_init.hpp>
#include "app_server.hpp"
```

(continues on next page)

(continued from previous page)

```

int hpx_main()
{
{
    // Create an instance of the app_server component on the current locality.
    hpx::naming:id_type app_server_instance =
        hpx::create_component<app::server>(hpx::find_here());

    // Create an instance of the action 'print_greeting_action'.
    app::server::print_greeting_action print_greeting;

    // Invoke the action 'print_greeting' on the newly created component.
    print_greeting(app_server_instance);
}
return hpx::finalize();
}

int main(int argc, char* argv[])
{
    return hpx::init(argc, argv);
}

```

In order to make sure that the application will be able to use the component `app::server`, special configuration information must be passed to *HPX*. The simplest way to allow *HPX* to ‘find’ the component is to provide special ini configuration files, which add the necessary information to the internal configuration database. The component should have a special ini file containing the information specific to the component `app_server`.

```

# file: $APP_ROOT/app_server.ini
[hpx.components.app_server]
name = app_server
path = $APP_LOCATION/

```

Here `$APP_LOCATION` is the directory where the (binary) component shared library is located. *HPX* will attempt to load the shared library from there. The section name `hpx.components.app_server` reflects the instance name of the component (`app_server` is an arbitrary, but unique name). The property value for `hpx.components.app_server.name` should be the same as used for the second argument to the macro `HPX_REGISTER_COMPONENT` above.

Additionally a file `.hpx.ini` which could be located in the current working directory (see step 3 as described in the section *Loading INI files*) can be used to add to the ini search path for components:

```

# file: $PWD/.hpx.ini
[hpx]
ini_path = ${hpx.ini_path}:$APP_ROOT/

```

This assumes that the above ini file specific to the component is located in the directory `$APP_ROOT`.

Note: It is possible to reference the defined property from inside its value. *HPX* will gracefully use the previous value of `hpx.ini_path` for the reference on the right hand side and assign the overall (now expanded) value to the property.

Logging

HPX uses a sophisticated logging framework allowing to follow in detail what operations have been performed inside the *HPX* library in what sequence. This information proves to be very useful for diagnosing problems or just for improving the understanding what is happening in *HPX* as a consequence of invoking *HPX* API functionality.

Default logging

Enabling default logging is a simple process. The detailed description in the remainder of this section explains different ways to customize the defaults. Default logging can be enabled by using one of the following:

- a command line switch `--hpx:debug-hpx-log`, which will enable logging to the console terminal
- the command line switch `--hpx:debug-hpx-log=<filename>`, which enables logging to a given file `<filename>`, or
- setting an environment variable `HPX_LOGLEVEL=<loglevel>` while running the *HPX* application. In this case `<loglevel>` should be a number between (or equal to) 1 and 5 where 1 means minimal logging and 5 causes to log all available messages. When setting the environment variable the logs will be written to a file named `hpx.<PID>.lo` in the current working directory, where `<PID>` is the process id of the console instance of the application.

Customizing logging

Generally, logging can be customized either using environment variable settings or using by an ini configuration file. Logging is generated in several categories, each of which can be customized independently. All customizable configuration parameters have reasonable defaults, allowing to use logging without any additional configuration effort. The following table lists the available categories.

Table 2.7: Logging categories

Category	Category shortcut	Information to be generated	Environment variable
General	None	Logging information generated by different subsystems of <i>HPX</i> , such as thread-manager, parcel layer, LCOs, etc.	<code>HPX_LOGLEVEL</code>
<i>AGAS</i>	<i>AGAS</i>	Logging output generated by the <i>AGAS</i> subsystem	<code>HPX_AGAS_LOGLEVEL</code>
Application	APP	Logging generated by applications.	<code>HPX_APP_LOGLEVEL</code>

By default, all logging output is redirected to the console instance of an application, where it is collected and written to a file, one file for each logging category.

Each logging category can be customized at two levels, the parameters for each are stored in the ini configuration sections `hpx.logging.CATEGORY` and `hpx.logging.console.CATEGORY` (where `CATEGORY` is the category shortcut as listed in the table above). The former influences logging at the source *locality* and the latter modifies the logging behaviour for each of the categories at the console instance of an application.

Levels

All *HPX* logging output has seven different logging levels. These levels can be set explicitly or through environmental variables in the main *HPX* ini file as shown below. The logging levels and their associated integral values are shown in the table below, ordered from most verbose to least verbose. By default, all *HPX* logs are set to 0, e.g. all logging output is disabled by default.

Table 2.8: Logging levels

Logging level	Integral value
<debug>	5
<info>	4
<warning>	3
<error>	2
<fatal>	1
No logging	0

Tip: The easiest way to enable logging output is to set the environment variable corresponding to the logging category to an integral value as described in the table above. For instance, setting `HPX_LOGLEVEL=5` will enable full logging output for the general category. Please note that the syntax and means of setting environment variables varies between operating systems.

Configuration

Logs will be saved to destinations as configured by the user. By default, logging output is saved on the console instance of an application to `hpx.<CATEGORY>.<PID>.lo` (where `CATEGORY` and `PID` are placeholders for the category shortcut and the OS process id). The output for the general logging category is saved to `hpx.<PID>.log`. The default settings for the general logging category are shown here (the syntax is described in the section [The HPX INI File Format](#)):

```
[hpx.logging]
level = ${HPX_LOGLEVEL:0}
destination = ${HPX_LOGDESTINATION:console}
format = ${HPX_LOGFORMAT:(T%locality%/%hpxthread%.%hpxphase%/%hpxcomponent%) P
    ↪%parentloc%/%hpxparent%.%hpxparentphase% %time%($hh:$mm.$ss.$mil) [%idx%] |\\n}
```

The logging level is taken from the environment variable `HPX_LOGLEVEL` and defaults to zero, e.g. no logging. The default logging destination is read from the environment variable `HPX_LOGDESTINATION`. On any of the localities it defaults to `console` which redirects all generated logging output to the console instance of an application. The following table lists the possible destinations for any logging output. It is possible to specify more than one destination separated by whitespace.

Table 2.9: Logging destinations

Logging destination	Description
file(<filename>)	Direct all output to a file with the given <filename>.
cout	Direct all output to the local standard output of the application instance on this <i>locality</i> .
cerr	Direct all output to the local standard error output of the application instance on this <i>locality</i> .
console	Direct all output to the console instance of the application. The console instance has its logging destinations configured separately.
android_log	Direct all output to the (Android) system log (available on Android systems only).

The logging format is read from the environment variable `HPX_LOGFORMAT` and it defaults to a complex format description. This format consists of several placeholder fields (for instance `%locality%` which will be replaced by concrete values when the logging output is generated. All other information is transferred verbatim to the output. The table below describes the available field placeholders. The separator character `|` separates the logging message prefix formatted as shown and the actual log message which will replace the separator.

Table 2.10: Available field placeholders

Name	Description
<code>locality</code>	The id of the <code>locality</code> on which the logging message was generated.
<code>hpxthread</code>	The id of the <i>HPX</i> -thread generating this logging output.
<code>hpxphase</code>	The phase ¹¹¹ of the <i>HPX</i> -thread generating this logging output.
<code>hpxcomponent</code>	The local virtual address of the component which the current <i>HPX</i> -thread is accessing.
<code>parentloc</code>	The id of the <code>locality</code> where the <i>HPX</i> thread was running which initiated the current <i>HPX</i> -thread. The current <i>HPX</i> -thread is generating this logging output.
<code>hpxparent</code>	The id of the <i>HPX</i> -thread which initiated the current <i>HPX</i> -thread. The current <i>HPX</i> -thread is generating this logging output.
<code>hpxparentphase</code>	The phase of the <i>HPX</i> -thread when it initiated the current <i>HPX</i> -thread. The current <i>HPX</i> -thread is generating this logging output.
<code>time</code>	The time stamp for this logging output line as generated by the source <code>locality</code> .
<code>idx</code>	The sequence number of the logging output line as generated on the source <code>locality</code> .
<code>osthread</code>	The sequence number of the OS-thread which executes the current <i>HPX</i> -thread.

Note: Not all of the field placeholder may be expanded for all generated logging output. If no value is available for a particular field it is replaced with a sequence of '-' characters.]

Here is an example line from a logging output generated by one of the *HPX* examples (please note that this is generated on a single line, without line break):

```
(T00000000/0000000002d46f90.01/00000000009ebc10) P-----/0000000002d46f80.02 17:49.
˓→37.320 [00000000000004d]
    <info> [RT] successfully created component {000000100ff0001, 0000000000030002}_
˓→of type: component_barrier[7(3)]
```

The default settings for the general logging category on the console is shown here:

```
[hpx.logging.console]
level = ${HPX_LOGLEVEL:${hpx.logging.level}}
destination = ${HPX_CONSOLE_LOGDESTINATION:file(hpx.${[system.pid].log})}
format = ${HPX_CONSOLE_LOGFORMAT:|}
```

These settings define how the logging is customized once the logging output is received by the console instance of an application. The logging level is read from the environment variable `HPX_LOGLEVEL` (as set for the console instance of the application). The level defaults to the same values as the corresponding settings in the general logging configuration shown before. The destination on the console instance is set to be a file which name is generated based from its OS process id. Setting the environment variable `HPX_CONSOLE_LOGDESTINATION` allows customization of the naming scheme for the output file. The logging format is set to leave the original logging output unchanged, as received from one of the localities the application runs on.

¹¹¹ The phase of a *HPX*-thread counts how often this thread has been activated.

HPX Command Line Options

The predefined command line options for any application using `hpx::init` are described in the following subsections.

HPX options (allowed on command line only)

```
--hpx:help
    print out program usage (default: this message), possible values: full (additionally prints options from components)
--hpx:version
    print out HPX version and copyright information
--hpx:info
    print out HPX configuration information
--hpx:options-file arg
    specify a file containing command line options (alternatively: @filepath)
```

HPX options (additionally allowed in an options file)

```
--hpx:worker
    run this instance in worker mode
--hpx:console
    run this instance in console mode
--hpx:connect
    run this instance in worker mode, but connecting late
--hpx:run-agas-server
    run AGAS server as part of this runtime instance
--hpx:run-hpx-main
    run the hpx_main function, regardless of locality mode
--hpx:hpx arg
    the IP address the HPX parcelport is listening on, expected format: address:port (default: 127.0.0.1:7910)
--hpx:agas arg
    the IP address the AGAS root server is running on, expected format: address:port (default: 127.0.0.1:7910)
--hpx:run-agas-server-only
    run only the AGAS server
--hpx:nodename arg
    the file name of a node file to use (list of nodes, one node name per line and core)
--hpx:nodes arg
    the (space separated) list of the nodes to use (usually this is extracted from a node file)
--hpx:endnodes
    this can be used to end the list of nodes specified using the option --hpx:nodes
--hpx:ifsuffix arg
    suffix to append to host names in order to resolve them to the proper network interconnect
```

--**hpx:ifprefix** arg
prefix to prepend to host names in order to resolve them to the proper network interconnect

--**hpx:iftransform** arg
sed-style search and replace (s/search/replace/) used to transform host names to the proper network interconnect

--**hpx:localities** arg
the number of localities to wait for at application startup (default: 1)

--**hpx:node** arg
number of the node this *locality* is run on (must be unique)

--**hpx:ignore-batch-env**
ignore batch environment variables

--**hpx:expect-connecting-localities**
this *locality* expects other localities to dynamically connect (this is implied if the number of initial localities is larger than 1)

--**hpx:pu-offset**
the first processing unit this instance of *HPX* should be run on (default: 0)

--**hpx:pu-step**
the step between used processing unit numbers for this instance of *HPX* (default: 1)

--**hpx:threads** arg
the number of operating system threads to spawn for this *HPX locality*. Possible values are: numeric values 1, 2, 3 and so on, all (which spawns one thread per processing unit, includes hyperthreads), or cores (which spawns one thread per core) (default: cores).

--**hpx:cores** arg
the number of cores to utilize for this *HPX locality* (default: all, i.e. the number of cores is based on the number of threads --*hpx:threads* assuming --*hpx:bind*=compact)

--**hpx:affinity** arg
the affinity domain the OS threads will be confined to, possible values: pu, core, numa, machine (default: pu)

--**hpx:bind** arg
the detailed affinity description for the OS threads, see *More details about HPX command line options* for a detailed description of possible values. Do not use with --*hpx:pu-step*, --*hpx:pu-offset* or --*hpx:affinity* options. Implies --*hpx:numa-sensitive* (--*hpx:bind*=none) disables defining thread affinities).

--**hpx:use-process-mask**
use the process mask to restrict available hardware resources (implies --*hpx:ignore-batch-env*)

--**hpx:print-bind**
print to the console the bit masks calculated from the arguments specified to all --*hpx:bind* options.

--**hpx:queuing** arg
the queue scheduling policy to use, options are local, local-priority-fifo, local-priority-lifo, static, static-priority, abp-priority-fifo and abp-priority-lifo (default: local-priority-fifo)

--**hpx:high-priority-threads** arg
the number of operating system threads maintaining a high priority queue (default: number of OS threads), valid for --*hpx:queuing*=abp-priority, --*hpx:queuing*=static-priority and --*hpx:queuing*=local-priority only

```
--hpx: numa-sensitive
    makes the scheduler NUMA sensitive
```

HPX configuration options

```
--hpx:app-config arg
    load the specified application configuration (ini) file
--hpx:config arg
    load the specified hpx configuration (ini) file
--hpx:ini arg
    add a configuration definition to the default runtime configuration
--hpx:exit
    exit after configuring the runtime
```

HPX debugging options

```
--hpx:list-symbolic-names
    list all registered symbolic names after startup
--hpx:list-component-types
    list all dynamic component types after startup
--hpx:dump-config-initial
    print the initial runtime configuration
--hpx:dump-config
    print the final runtime configuration
--hpx:debug-hpx-log [arg]
    enable all messages on the HPX log channel and send all HPX logs to the target destination (default: cout)
--hpx:debug-agas-log [arg]
    enable all messages on the AGAS log channel and send all AGAS logs to the target destination (default: cout)
--hpx:debug-parcel-log [arg]
    enable all messages on the parcel transport log channel and send all parcel transport logs to the target destination
        (default: cout)
--hpx:debug-timing-log [arg]
    enable all messages on the timing log channel and send all timing logs to the target destination (default: cout)
--hpx:debug-app-log [arg]
    enable all messages on the application log channel and send all application logs to the target destination (default:
        cout)
--hpx:debug-clp
    debug command line processing
--hpx:attach-debugger arg
    wait for a debugger to be attached, possible arg values: startup or exception (default: startup)
```

HPX options related to performance counters

--hpx:print-counter

print the specified performance counter either repeatedly and/or at the times specified by --hpx:print-counter-at (see also option --hpx:print-counter-interval)

--hpx:print-counter-reset

print the specified performance counter either repeatedly and/or at the times specified by --hpx:print-counter-at reset the counter after the value is queried. (see also option --hpx:print-counter-interval)

--hpx:print-counter-interval

print the performance counter(s) specified with --hpx:print-counter repeatedly after the time interval (specified in milliseconds), (default: 0, which means print once at shutdown)

--hpx:print-counter-destination

print the performance counter(s) specified with --hpx:print-counter to the given file (default: console)

--hpx:list-counters

list the names of all registered performance counters, possible values: minimal (prints counter name skeletons), full (prints all available counter names)

--hpx:list-counter-infos

list the description of all registered performance counters, possible values: minimal (prints info for counter name skeletons), full (prints all available counter infos)

--hpx:print-counter-format

print the performance counter(s) specified with --hpx:print-counter possible formats in csv format with header or without any header (see option --hpx:no-csv-header, possible values: csv (prints counter values in CSV format with full names as header), csv-short (prints counter values in CSV format with shortnames provided with --hpx:print-counter as --hpx:print-counter shortname, full-countername

--hpx:no-csv-header

print the performance counter(s) specified with --hpx:print-counter and csv or csv-short format specified with --hpx:print-counter-format without header

--hpx:print-counter-at arg

print the performance counter(s) specified with --hpx:print-counter (or --hpx:print-counter-reset at the given point in time, possible argument values: startup, shutdown (default), noshutdown

--hpx:reset-counters

reset all performance counter(s) specified with --hpx:print-counter after they have been evaluated.

--hpx:print-counters-locally

Each *locality* prints only its own local counters. If this is used with --hpx:print-counter-destination=<file>, the code will append a ".<locality_id>" to the file name in order to avoid clashes between localities.

Command line argument shortcuts

Additionally, the following shortcuts are available from every *HPX* application.

Table 2.11: Predefined command line option shortcuts

Shortcut option	Equivalent long option
-a	--hpx:agas
-c	--hpx:console
-h	--hpx:help
-I	--hpx:ini
-l	--hpx:localities
-p	--hpx:app-config
-q	--hpx:queuing
-r	--hpx:run-agas-server
-t	--hpx:threads
-v	--hpx:version
-w	--hpx:worker
-x	--hpx:hpx
-0	--hpx:node=0
-1	--hpx:node=1
-2	--hpx:node=2
-3	--hpx:node=3
-4	--hpx:node=4
-5	--hpx:node=5
-6	--hpx:node=6
-7	--hpx:node=7
-8	--hpx:node=8
-9	--hpx:node=9

It is possible to define your own shortcut options. In fact, all of the shortcuts listed above are pre-defined using the technique described here. Also, it is possible to redefine any of the pre-defined shortcuts to expand differently as well.

Shortcut options are obtained from the internal configuration database. They are stored as key-value properties in a special properties section named `hpx.commandline`. You can define your own shortcuts by adding the corresponding definitions to one of the `ini` configuration files as described in the section *Configuring HPX applications*. For instance, in order to define a command line shortcut `--p` which should expand to `--hpx:print-counter`, the following configuration information needs to be added to one of the `ini` configuration files:

```
[hpx.commandline.aliases]
--pc = --hpx:print-counter
```

Note: Any arguments for shortcut options passed on the command line are retained and passed as arguments to the corresponding expanded option. For instance, given the definition above, the command line option:

```
--pc=/threads{locality#0/total}/count/cumulative
```

would be expanded to:

```
--hpx:print-counter=/threads{locality#0/total}/count/cumulative
```

Important: Any shortcut option should either start with a single '-' or with two '--' characters. Shortcuts

starting with a single '-' are interpreted as short options (i.e. everything after the first character following the '-' is treated as the argument). Shortcuts starting with '--' are interpreted as long options. No other shortcut formats are supported.

Specifying options for single localities only

For runs involving more than one *locality* it is sometimes desirable to supply specific command line options to single localities only. When the *HPX* application is launched using a scheduler (like PBS, for more details see section [How to use HPX applications with PBS](#)), specifying dedicated command line options for single localities may be desirable. For this reason all of the command line options which have the general format `--hpx:<some_key>` can be used in a more general form: `--hpx:<N>:<some_key>`, where <N> is the number of the *locality* this command line options will be applied to, all other localities will simply ignore the option. For instance, the following PBS script passes the option `--hpx:pu-offset=4` to the *locality* '1' only.

```
#!/bin/bash
#
#PBS -l nodes=2:ppn=4

APP_PATH=~/packages/hpx/bin/hello_world_distributed
APP_OPTIONS=

pbsdsh -u $APP_PATH $APP_OPTIONS --hpx:1:pu-offset=4 --hpx:nodes=`cat $PBS_NODEFILE`
```

Caution: If the first application specific argument (inside \$APP_OPTIONS is a non-option (i.e. does not start with a - or a --, then it must be placed before the option `--hpx:nodes`, which, in this case, should be the last option on the command line).

Alternatively, use the option `--hpx:endnodes` to explicitly mark the end of the list of node names:

```
pbsdsh -u $APP_PATH --hpx:1:pu-offset=4 --hpx:nodes=`cat $PBS_NODEFILE` --
→hpx:endnodes $APP_OPTIONS
```

More details about *HPX* command line options

This section documents the following list of the command line options in more detail:

- *The command line option --hpx:bind*

The command line option `--hpx:bind`

This command line option allows one to specify the required affinity of the *HPX* worker threads to the underlying processing units. As a result the worker threads will run only on the processing units identified by the corresponding bind specification. The affinity settings are to be specified using `--hpx:bind=<BINDINGS>`, where <BINDINGS> have to be formatted as described below.

In addition to the syntax described below one can use `--hpx:bind=none` to disable all binding of any threads to a particular core. This is mostly supported for debugging purposes.

The specified affinities refer to specific regions within a machine hardware topology. In order to understand the hardware topology of a particular machine it may be useful to run the lstopo tool which is part of Portable Hardware

Locality (HWLOC) to see the reported topology tree. Seeing and understanding a topology tree will definitely help in understanding the concepts that are discussed below.

Affinities can be specified using HWLOC (Portable Hardware Locality (HWLOC)) tuples. Tuples of HWLOC *objects* and associated *indexes* can be specified in the form `object:index`, `object:index-index` or `object:index,...,index`. HWLOC objects represent types of mapped items in a topology tree. Possible values for objects are `socket`, `numanode`, `core` and `pu` (processing unit). Indexes are non-negative integers that specify a unique physical object in a topology tree using its logical sequence number.

Chaining multiple tuples together in the more general form `object1:index1[.object2:index2[...]]` is permissible. While the first tuple's object may appear anywhere in the topology, the Nth tuple's object must have a shallower topology depth than the (N+1)th tuple's object. Put simply: as you move right in a tuple chain, objects must go deeper in the topology tree. Indexes specified in chained tuples are relative to the scope of the parent object. For example, `socket:0.core:1` refers to the second core in the first socket (all indices are zero based).

Multiple affinities can be specified using several `--hpx:bind` command line options or by appending several affinities separated by a ';' . By default, if multiple affinities are specified, they are added.

"`all`" is a special affinity consisting in the entire current topology.

Note: All 'names' in an affinity specification, such as `thread`, `socket`, `numanode`, `pu` or `all` can be abbreviated. Thus the affinity specification `threads:0-3=socket:0.core:1.pu:1` is fully equivalent to its shortened form `t:0-3=s:0.c:1.p:1`.

Here is a full grammar describing the possible format of mappings:

```
mappings      ::=  distribution | mapping (";" mapping)*
distribution   ::=  "compact" | "scatter" | "balanced" | "numa-balanced"
mapping        ::=  thread_spec "=" pu_specs
thread_spec    ::=  "thread:" range_specs
pu_specs       ::=  pu_spec (".") pu_spec)*
pu_spec        ::=  type ":" range_specs | "~" pu_spec
range_specs    ::=  range_spec ("," range_spec)*
range_spec     ::=  int | int "-" int | "all"
type           ::=  "socket" | "numanode" | "core" | "pu"
```

The following example assumes a system with at least 4 cores, where each core has more than 1 processing unit (hardware threads). Running `hello_world_distributed` with 4 OS-threads (on 4 processing units), where each of those threads is bound to the first processing unit of each of the cores, can be achieved by invoking:

```
hello_world_distributed -t4 --hpx:bind=thread:0-3=core:0-3.pu:0
```

Here `thread:0-3` specifies the OS threads for which to define affinity bindings, and `core:0-3.pu:0` defines that for each of the cores (`core:0-3`) only their first processing unit `pu:0` should be used.

Note: The command line option `--hpx:print-bind` can be used to print the bitmasks generated from the affinity mappings as specified with `--hpx:bind`. For instance, on a system with hyperthreading enabled (i.e. 2 processing units per core), the command line:

```
hello_world_distributed -t4 --hpx:bind=thread:0-3=core:0-3.pu:0 --hpx:print-bind
```

will cause this output to be printed:

```

0: PU L#0(P#0), Core L#0, Socket L#0, Node L#0(P#0)
1: PU L#2(P#2), Core L#1, Socket L#0, Node L#0(P#0)
2: PU L#4(P#4), Core L#2, Socket L#0, Node L#0(P#0)
3: PU L#6(P#6), Core L#3, Socket L#0, Node L#0(P#0)

```

where each bit in the bitmasks corresponds to a processing unit the listed worker thread will be bound to run on.

The difference between the four possible predefined distribution schemes (compact, scatter, balanced and numa-balanced) is best explained with an example. Imagine that we have a system with 4 cores and 4 hardware threads per core on 2 sockets. If we place 8 threads the assignments produced by the compact, scatter, balanced and numa-balanced types are shown in the figure below. Notice that compact does not fully utilize all the cores in the system. For this reason it is recommended that applications are run using the scatter or balanced/numa-balanced options in most cases.

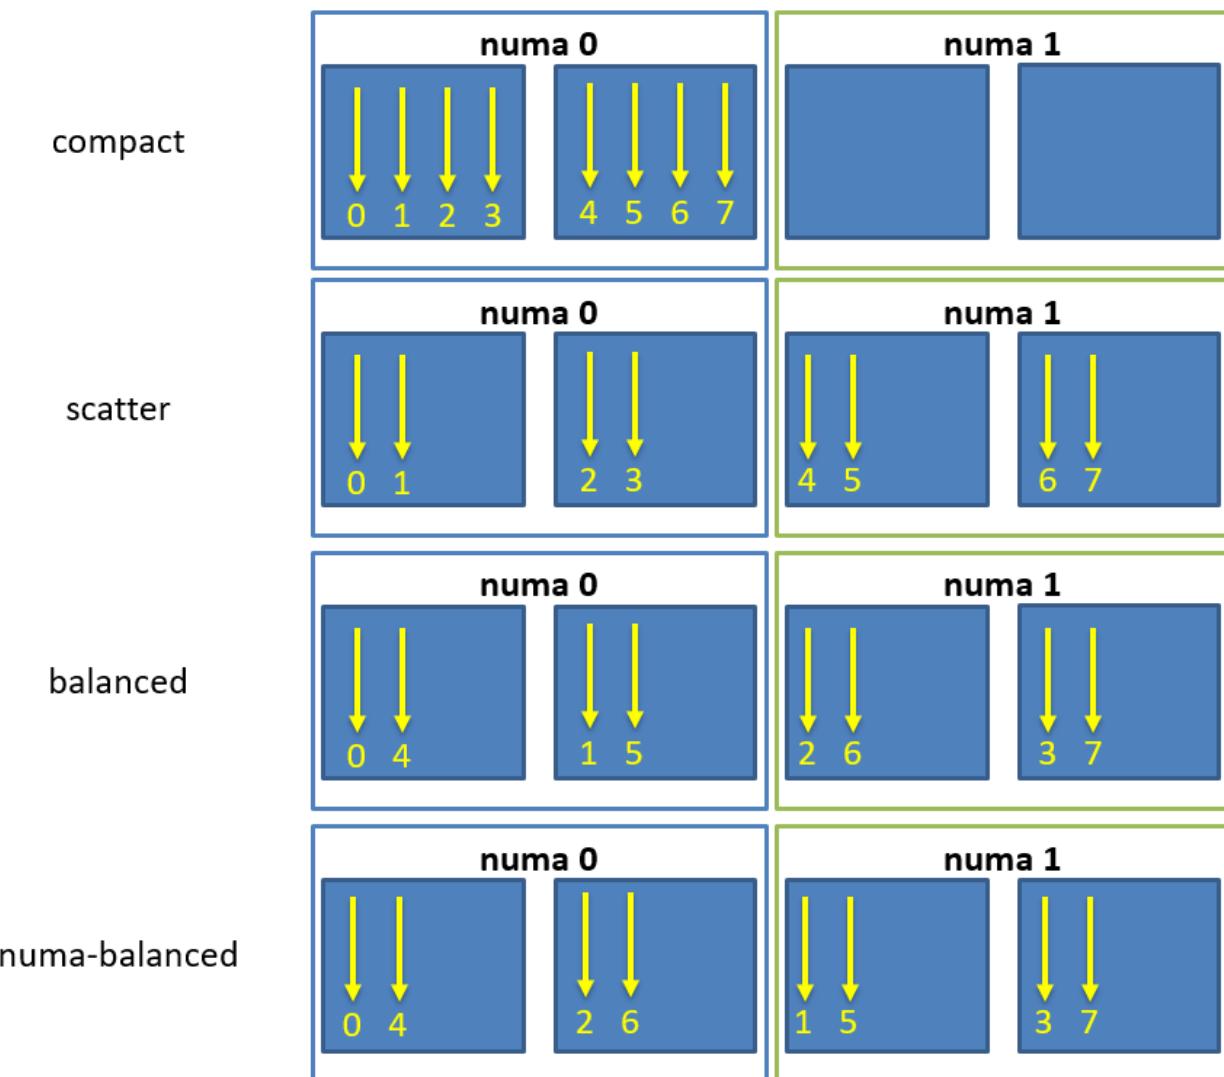


Fig. 2.7: Schematic of thread affinity type distributions.

In addition to the predefined distributions it is possible to restrict the resources used by *HPX* to the process CPU

mask. The CPU mask is typically set by e.g. MPI¹¹⁰ and batch environments. Using the command line option `--hpx:use-process-mask` makes HPX act as if only the processing units in the CPU mask are available for use by HPX. The number of threads is automatically determined from the CPU mask. The number of threads can still be changed manually using this option, but only to a number less than or equal to the number of processing units in the CPU mask. The option `--hpx:print-bind` is useful in conjunction with `--hpx:use-process-mask` to make sure threads are placed as expected.

2.5.6 Writing single-node HPX applications

HPX is a C++ Standard Library for Concurrency and Parallelism. This means that it implements all of the corresponding facilities as defined by the C++ Standard. Additionally, HPX implements functionalities proposed as part of the ongoing C++ standardization process. This section focuses on the features available in HPX for parallel and concurrent computation on a single node, although many of the features presented here are also implemented to work in the distributed case.

Using LCOs

Lightweight Control Objects (LCOs) provide synchronization for HPX applications. Most of them are familiar from other frameworks, but a few of them work in slightly different ways adapted to HPX. The following synchronization objects are available in HPX:

1. future
2. queue
3. object_semaphore
4. barrier

Channels

Channels combine communication (the exchange of a value) with synchronization (guaranteeing that two calculations (tasks) are in a known state). A channel can transport any number of values of a given type from a sender to a receiver:

```
hpx::lcos::local::channel<int> c;
c.set(42);
cout << c.get();           // will print '42'
```

Channels can be handed to another thread (or in case of channel components, to other localities), thus establishing a communication channel between two independent places in the program:

```
void do_something()
{
    hpx::lcos::local::receive_channel<int> c,
    hpx::lcos::local::send_channel<> done)
{
    cout << c.get();           // prints 42
    done.set();                // signal back
}

{
    hpx::lcos::local::channel<int> c;
    hpx::lcos::local::channel<> done;
```

(continues on next page)

¹¹⁰ https://en.wikipedia.org/wiki/Message_Passing_Interface

(continued from previous page)

```

    hpx::apply(&do_something, c, done);

    c.set(42);           // send some value
    done.get();          // wait for thread to be done
}

```

A channel component is created on one *locality* and can be sent to another *locality* using an action. This example also demonstrates how a channel can be used as a range of values:

```

// channel components need to be registered for each used type (not needed
// for hpx::lcos:::local:::channel)
HPX_REGISTER_CHANNEL(double);

void some_action(hpx::lcos:::channel<double> c)
{
    for (double d : c)
        hpx::cout << d << std::endl;
}

HPX_REGISTER_ACTION(some_action);

{
    // create the channel on this locality
    hpx::lcos:::channel<double> c(hpx::find_here());

    // pass the channel to a (possibly remote invoked) action
    hpx::apply(some_action(), hpx::find_here(), c);

    // send some values to the receiver
    std::vector<double> v = { 1.2, 3.4, 5.0 };
    for (double d : v)
        c.set(d);

    // explicitly close the communication channel (implicit at destruction)
    c.close();
}

```

Composable guards

Composable guards operate in a manner similar to locks, but are applied only to asynchronous functions. The guard (or guards) is automatically locked at the beginning of a specified task and automatically unlocked at the end. Because guards are never added to an existing task's execution context, the calling of guards is freely composable and can never deadlock.

To call an application with a single guard, simply declare the guard and call `run_guarded()` with a function (task):

```

hpx::lcos:::local:::guard gu;
run_guarded(gu,task);

```

If a single method needs to run with multiple guards, use a guard set:

```

boost::shared<hpx::lcos:::local:::guard> gu1(new hpx::lcos:::local:::guard());
boost::shared<hpx::lcos:::local:::guard> gu2(new hpx::lcos:::local:::guard());
gs.add(*gu1);
gs.add(*gu2);
run_guarded(gs,task);

```

Guards use two atomic operations (which are not called repeatedly) to manage what they do, so overhead should be extremely low. The following guards are available in *HPX*:

1. conditional_trigger
2. counting_semaphore
3. dataflow
4. event
5. mutex
6. once
7. recursive_mutex
8. spinlock
9. spinlock_no_backoff
10. trigger

Extended facilities for futures

Concurrency is about both decomposing and composing the program from the parts that work well individually and together. It is in the composition of connected and multicore components where today's C++ libraries are still lacking.

The functionality of `std::future` offers a partial solution. It allows for the separation of the initiation of an operation and the act of waiting for its result; however, the act of waiting is synchronous. In communication-intensive code this act of waiting can be unpredictable, inefficient and simply frustrating. The example below illustrates a possible synchronous wait using futures:

```
#include <future>
using namespace std;
int main()
{
    future<int> f = async([]() { return 123; });
    int result = f.get(); // might block
}
```

For this reason, *HPX* implements a set of extensions to `std::future` (as proposed by [__cpp11_n4107__](#)). This proposal introduces the following key asynchronous operations to `hpx::future`, `hpx::shared_future` and `hpx::async`, which enhance and enrich these facilities.

Table 2.13: Facilities extending std::future

Facility	Description
hpx::future	In asynchronous programming, it is very common for one asynchronous operation, on completion, to invoke a second operation and pass data to it. The current C++ standard does not allow one to register a continuation to a future. With <code>then</code> , instead of waiting for the result, a continuation is “attached” to the asynchronous operation, which is invoked when the result is ready. Continuations registered using <code>then</code> function will help to avoid blocking waits or wasting threads on polling, greatly improving the responsiveness and scalability of an application.
unwrapping constructor for hpx::future	In some scenarios, you might want to create a future that returns another future, resulting in nested futures. Although it is possible to write code to unwrap the outer future and retrieve the nested future and its result, such code is not easy to write because users must handle exceptions and it may cause a blocking call. Unwrapping can allow users to mitigate this problem by doing an asynchronous call to unwrap the outermost future.
hpx::future::test	There are often situations where a <code>get()</code> call on a future may not be a blocking call, or is only a blocking call under certain circumstances. This function gives the ability to test for early completion and allows us to avoid associating a continuation, which needs to be scheduled with some non-trivial overhead and near-certain loss of cache efficiency.
hpx::make_ready_future	Some functions may know the value at the point of construction. In these cases the value is immediately available, but needs to be returned as a future. By using <code>hpx::make_ready_future</code> a future can be created that holds a pre-computed result in its shared state. In the current standard it is non-trivial to create a future directly from a value. First a promise must be created, then the promise is set, and lastly the future is retrieved from the promise. This can now be done with one operation.

The standard also omits the ability to compose multiple futures. This is a common pattern that is ubiquitous in other asynchronous frameworks and is absolutely necessary in order to make C++ a powerful asynchronous programming language. Not including these functions is synonymous to Boolean algebra without AND/OR.

In addition to the extensions proposed by N4313¹¹², HPX adds functions allowing users to compose several futures in a more flexible way.

¹¹² <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

Table 2.14: Facilities for composing `hpx::futures`

Facility	Description	Comment
<code>hpx::when_any</code> , <code>hpx::when_any_n</code>	Asynchronously wait for at least one of multiple future or shared_future objects to finish.	N4313 ¹¹³ , ..._n versions are HPX only
<code>hpx::wait_any</code> , <code>hpx::wait_any_n</code>	Synchronously wait for at least one of multiple future or shared_future objects to finish.	HPX only
<code>hpx::when_all</code> , <code>hpx::when_all_n</code>	Asynchronously wait for all future and shared_future objects to finish.	N4313 ¹¹⁴ , ..._n versions are HPX only
<code>hpx::wait_all</code> , <code>hpx::wait_all_n</code>	Synchronously wait for all future and shared_future objects to finish.	HPX only
<code>hpx::when_some</code> , <code>hpx::when_some_n</code>	Asynchronously wait for multiple future and shared_future objects to finish.	HPX only
<code>hpx::wait_some</code> , <code>hpx::wait_some_n</code>	Synchronously wait for multiple future and shared_future objects to finish.	HPX only
<code>hpx::when_each</code>	Asynchronously wait for multiple future and shared_future objects to finish and call a function for each of the future objects as soon as it becomes ready.	HPX only
<code>hpx::wait_each</code> , <code>hpx::wait_each_n</code>	Synchronously wait for multiple future and shared_future objects to finish and call a function for each of the future objects as soon as it becomes ready.	HPX only

High level parallel facilities

In preparation for the upcoming C++ Standards, there are currently several proposals targeting different facilities supporting parallel programming. *HPX* implements (and extends) some of those proposals. This is well aligned with our strategy to align the APIs exposed from *HPX* with current and future C++ Standards.

At this point, *HPX* implements several of the C++ Standardization working papers, most notably N4409¹¹⁵ (Working Draft, Technical Specification for C++ Extensions for Parallelism), N4411¹¹⁶ (Task Blocks), and N4406¹¹⁷ (Parallel Algorithms Need Executors).

Using parallel algorithms

A parallel algorithm is a function template described by this document which is declared in the (inline) namespace `hpx::parallel::v1`.

Note: For compilers that do not support inline namespaces, all of the namespace `v1` is imported into the namespace `hpx::parallel`. The effect is similar to what inline namespaces would do, namely all names defined in `hpx::parallel::v1` are accessible from the namespace `hpx::parallel` as well.

All parallel algorithms are very similar in semantics to their sequential counterparts (as defined in the namespace `std`) with an additional formal template parameter named `ExecutionPolicy`. The execution policy is generally

¹¹³ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

¹¹⁴ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

¹¹⁵ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4409.pdf>

¹¹⁶ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4411.pdf>

¹¹⁷ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4406.pdf>

passed as the first argument to any of the parallel algorithms and describes the manner in which the execution of these algorithms may be parallelized and the manner in which they apply user-provided function objects.

The applications of function objects in parallel algorithms invoked with an execution policy object of type `hpx::parallel::execution::sequenced_policy` or `hpx::parallel::execution::sequenced_task_policy` execute in sequential order. For `hpx::parallel::execution::sequenced_policy` the execution happens in the calling thread.

The applications of function objects in parallel algorithms invoked with an execution policy object of type `hpx::parallel::execution::parallel_policy` or `hpx::parallel::execution::parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and are indeterminately sequenced within each thread.

Important: It is the caller's responsibility to ensure correctness, such as making sure that the invocation does not introduce data races or deadlocks.

The applications of function objects in parallel algorithms invoked with an execution policy of type `hpx::parallel::execution::parallel_unsequenced_policy` is, in HPX, equivalent to the use of the execution policy `hpx::parallel::execution::parallel_policy`.

Algorithms invoked with an execution policy object of type `hpx::parallel::v1::execution_policy` execute internally as if invoked with the contained execution policy object. No exception is thrown when an `hpx::parallel::v1::execution_policy` contains an execution policy of type `hpx::parallel::execution::sequenced_task_policy` or `hpx::parallel::execution::parallel_task_policy` (which normally turn the algorithm into its asynchronous version). In this case the execution is semantically equivalent to the case of passing a `hpx::parallel::execution::sequenced_policy` or `hpx::parallel::execution::parallel_policy` contained in the `hpx::parallel::v1::execution_policy` object respectively.

Parallel exceptions

During the execution of a standard parallel algorithm, if temporary memory resources are required by any of the algorithms and no memory is available, the algorithm throws a `std::bad_alloc` exception.

During the execution of any of the parallel algorithms, if the application of a function object terminates with an uncaught exception, the behavior of the program is determined by the type of execution policy used to invoke the algorithm:

- If the execution policy object is of type `hpx::parallel::execution::parallel_unsequenced_policy`, `hpx::terminate` shall be called.
- If the execution policy object is of type `hpx::parallel::execution::sequenced_policy`, `hpx::parallel::execution::sequenced_task_policy`, `hpx::parallel::execution::parallel_policy` or `hpx::parallel::execution::parallel_task_policy`, the execution of the algorithm terminates with an `hpx::exception_list` exception. All uncaught exceptions thrown during the application of user-provided function objects shall be contained in the `hpx::exception_list`.

For example, the number of invocations of the user-provided function object in `for_each` is unspecified. When `hpx::parallel::v1::for_each` is executed sequentially, only one exception will be contained in the `hpx::exception_list` object.

These guarantees imply that, unless the algorithm has failed to allocate memory and terminated with `std::bad_alloc`, all exceptions thrown during the execution of the algorithm are communicated to the caller. It is unspecified whether an algorithm implementation will "forge ahead" after encountering and capturing a user exception.

The algorithm may terminate with the `std::bad_alloc` exception even if one or more user-provided function objects have terminated with an exception. For example, this can happen when an algorithm fails to allocate memory while creating or adding elements to the `hpx::exception_list` object.

Parallel algorithms

HPX provides implementations of the following parallel algorithms:

Table 2.15: Non-modifying parallel algorithms (in header: <hpx/include/parallel_algorithm.hpp>)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::v1::adjacent_difference</code>	Computes the differences between adjacent elements in a range.	<hpx/include/parallel_adjacent_find.hpp>	adjacent_difference ¹¹⁸
<code>hpx::parallel::v1::all_of</code>	Checks if a predicate is true for all of the elements in a range.	<hpx/include/parallel_all_any_none.hpp>	all_of ¹¹⁹
<code>hpx::parallel::v1::any_of</code>	Checks if a predicate is true for any of the elements in a range.	<hpx/include/parallel_all_any_none.hpp>	any_of ¹²⁰
<code>hpx::parallel::v1::count</code>	Returns the number of elements equal to a given value.	<hpx/include/parallel_count.hpp>	count ¹²¹
<code>hpx::parallel::v1::count_if</code>	Returns the number of elements satisfying a specific criteria.	<hpx/include/parallel_count.hpp>	count_if ¹²²
<code>hpx::parallel::v1::equal</code>	Determines if two sets of elements are the same.	<hpx/include/parallel_equal.hpp>	equal ¹²³
<code>hpx::parallel::v1::exclusive_scan</code>	Does an exclusive parallel scan over a range of elements.	<hpx/include/parallel_scan.hpp>	exclusive_scan ¹²⁴
<code>hpx::parallel::v1::find</code>	Finds the first element equal to a given value.	<hpx/include/parallel_find.hpp>	find ¹²⁵
<code>hpx::parallel::v1::find_end</code>	Finds the last sequence of elements in a certain range.	<hpx/include/parallel_find.hpp>	find_end ¹²⁶
<code>hpx::parallel::v1::find_first_of</code>	Searches for any one of a set of elements.	<hpx/include/parallel_find.hpp>	find_first_of ¹²⁷
<code>hpx::parallel::v1::find_if</code>	Finds the first element satisfying a specific criteria.	<hpx/include/parallel_find.hpp>	find_if ¹²⁸
<code>hpx::parallel::v1::find_if_not</code>	Finds the first element not satisfying a specific criteria.	<hpx/include/parallel_find.hpp>	find_if_not ¹²⁹
<code>hpx::parallel::v1::for_each</code>	Applies a function to a range of elements.	<hpx/include/parallel_for_each.hpp>	for_each ¹³⁰
<code>hpx::parallel::v1::for_each_n</code>	Applies a function to a number of elements.	<hpx/include/parallel_for_each.hpp>	for_each_n ¹³¹
<code>hpx::parallel::v1::inclusive_scan</code>	Does an inclusive parallel scan over a range of elements.	<hpx/include/parallel_scan.hpp>	inclusive_scan ¹³²
<code>hpx::parallel::v1::lexicographical_compare</code>	Checks if a range of values is lexicographically less than another range of values.	<hpx/include/parallel_lexicographical.hpp>	lexicographical_compare ¹³³
<code>hpx::parallel::v1::mismatch</code>	Finds the first position where two ranges differ.	<hpx/include/parallel_mismatch.hpp>	mismatch ¹³⁴
<code>hpx::parallel::v1::none_of</code>	Checks if a predicate is true for none of the elements in a range.	<hpx/include/parallel_all_any_none.hpp>	none_of ¹³⁵
<code>hpx::parallel::v1::search</code>	Searches for a range of elements.	<hpx/include/parallel_search.hpp>	search ¹³⁶
<code>hpx::parallel::v1::search_n</code>	Searches for a number consecutive copies of an element in a range.	<hpx/include/parallel_search.hpp>	search_n ¹³⁷

118 http://en.cppreference.com/w/cpp/algorithm/adjacent_find
119 http://en.cppreference.com/w/cpp/algorithm/all_any_none_of
120 http://en.cppreference.com/w/cpp/algorithm/all_any_none_of
121 <http://en.cppreference.com/w/cpp/algorithm/count>
122 http://en.cppreference.com/w/cpp/algorithm/count_if
123 <http://en.cppreference.com/w/cpp/algorithm/equal>
124 http://en.cppreference.com/w/cpp/algorithm/exclusive_scan
125 <http://en.cppreference.com/w/cpp/algorithm/find>
126 http://en.cppreference.com/w/cpp/algorithm/find_end
127 http://en.cppreference.com/w/cpp/algorithm/find_first_of
128 <http://en.cppreference.com/w/cpp/algorithm/find>
129 http://en.cppreference.com/w/cpp/algorithm/find_if_not
130 http://en.cppreference.com/w/cpp/algorithm/for_each
131 http://en.cppreference.com/w/cpp/algorithm/for_each_n
132 http://en.cppreference.com/w/cpp/algorithm/inclusive_scan
133 http://en.cppreference.com/w/cpp/algorithm/lexicographical_compare
134 <http://en.cppreference.com/w/cpp/algorithm/mismatch>
135 http://en.cppreference.com/w/cpp/algorithm/all_any_none_of
136 <http://en.cppreference.com/w/cpp/algorithm/search>
137 http://en.cppreference.com/w/cpp/algorithm/search_n

Table 2.16: Modifying parallel algorithms (In Header:
`<hpx/include/parallel_algorithm.hpp>`)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::v1_copy</code>	Copies a range of elements to a new location.	<code><hpx/include/parallel_copy.hpp></code>	<code>exclusive_scan</code> ¹³⁸
<code>hpx::parallel::v1_copy_n</code>	Copies a number of elements to a new location.	<code><hpx/include/parallel_copy.hpp></code>	<code>copy_n</code> ¹³⁹
<code>hpx::parallel::v1_copy_if</code>	Copies the elements from a range to a new location for which the given predicate is true	<code><hpx/include/parallel_copy.hpp></code>	<code>copy</code> ¹⁴⁰
<code>hpx::parallel::v1_move</code>	Moves a range of elements to a new location.	<code><hpx/include/parallel_fill.hpp></code>	<code>move</code> ¹⁴¹
<code>hpx::parallel::v1_fill</code>	Assigns a range of elements a certain value.	<code><hpx/include/parallel_fill.hpp></code>	<code>fill</code> ¹⁴²
<code>hpx::parallel::v1_fill_n</code>	Assigns a value to a number of elements.	<code><hpx/include/parallel_fill.hpp></code>	<code>fill_n</code> ¹⁴³
<code>hpx::parallel::v1_generate</code>	Saves the result of a function in a range.	<code><hpx/include/parallel_generate.hpp></code>	<code>generate</code> ¹⁴⁴
<code>hpx::parallel::v1_generate_n</code>	Saves the result of N applications of a function.	<code><hpx/include/parallel_generate.hpp></code>	<code>generate_n</code> ¹⁴⁵
<code>hpx::parallel::v1_remove</code>	Removes the elements from a range that are equal to the given value.	<code><hpx/include/parallel_remove.hpp></code>	<code>remove</code> ¹⁴⁶
<code>hpx::parallel::v1_remove_if</code>	Removes the elements from a range that are equal to the given predicate is false	<code><hpx/include/parallel_remove.hpp></code>	<code>remove</code> ¹⁴⁷
<code>hpx::parallel::v1_remove_copy</code>	Copies the elements from a range to a new location that are not equal to the given value.	<code><hpx/include/parallel_remove_copy.hpp></code>	<code>remove_copy</code> ¹⁴⁸
<code>hpx::parallel::v1_remove_copy_if</code>	Copies the elements from a range to a new location for which the given predicate is false	<code><hpx/include/parallel_remove_copy.hpp></code>	<code>remove_copy</code> ¹⁴⁹
<code>hpx::parallel::v1_replace</code>	Replaces all values satisfying specific criteria with another value.	<code><hpx/include/parallel_replace.hpp></code>	<code>replace</code> ¹⁵⁰
<code>hpx::parallel::v1_replace_all</code>	Replaces all values satisfying specific criteria with another value.	<code><hpx/include/parallel_replace.hpp></code>	<code>replace</code> ¹⁵¹
<code>hpx::parallel::v1_replace_if</code>	Copies a range, replacing elements satisfying specific criteria with another value.	<code><hpx/include/parallel_replace.hpp></code>	<code>replace_copy</code> ¹⁵²
<code>hpx::parallel::v1_replace_if_all</code>	Copies a range, replacing elements satisfying specific criteria with another value.	<code><hpx/include/parallel_replace.hpp></code>	<code>replace_copy</code> ¹⁵³
<code>hpx::parallel::v1_reverse</code>	Reverses the order elements in a range.	<code><hpx/include/parallel_reverse.hpp></code>	<code>reverse</code> ¹⁵⁴
2.5. Manual		<code>parallel_reverse.hpp</code>	147
<code>hpx::parallel::v1_reversed</code>	Creates a copy of a range that is reversed.	<code><hpx/include/parallel_reverse.hpp></code>	<code>reverse_copy</code> ¹⁵⁵

Table 2.17: Set operations on sorted sequences (In Header:
`<hpx/include/parallel_algorithm.hpp>`)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::v1::merge</code> ¹³⁸	Merges two sorted ranges.	<code><hpx/include/parallel_merge.hpp></code>	<code>merge</code> ¹⁶²
<code>hpx::parallel::v1::inplace_merge</code> ¹³⁹	Merges <i>any</i> two ordered ranges in-place.	<code><hpx/include/parallel_merge.hpp></code>	<code>inplace_merge</code> ¹⁶³
<code>hpx::parallel::v1::includes</code> ¹⁴⁰	Returns true if one set is a subset of another.	<code><hpx/include/parallel_set_operations.hpp></code>	<code>includes</code> ¹⁶⁴
<code>hpx::parallel::v1::set_difference</code> ¹⁴¹	Computes the difference between two sets.	<code><hpx/include/parallel_set_operations.hpp></code>	<code>set_difference</code> ¹⁶⁵
<code>hpx::parallel::v1::set_intersection</code> ¹⁴²	Computes the intersection of two sets.	<code><hpx/include/parallel_set_operations.hpp></code>	<code>set_intersection</code> ¹⁶⁶
<code>hpx::parallel::v1::set_symmetric_difference</code> ¹⁴³	Computes the symmetric difference between two sets.	<code><hpx/include/parallel_set_operations.hpp></code>	<code>set_symmetric_difference</code> ¹⁶⁷
<code>hpx::parallel::v1::set_union</code> ¹⁴⁴	Computes the union of two sets.	<code><hpx/include/parallel_set_operations.hpp></code>	<code>set_union</code> ¹⁶⁸

¹³⁸ http://en.cppreference.com/w/cpp/algorithm/exclusive_scan

¹³⁹ http://en.cppreference.com/w/cpp/algorithm/copy_n

¹⁴⁰ <http://en.cppreference.com/w/cpp/algorithm/copy>

141 <http://en.cppreference.com/w/cpp/algorithm/move>142 <http://en.cppreference.com/w/cpp/algorithm/fill>143 http://en.cppreference.com/w/cpp/algorithm/fill_n144 <http://en.cppreference.com/w/cpp/algorithm/generate>145 http://en.cppreference.com/w/cpp/algorithm/generate_n146 <http://en.cppreference.com/w/cpp/algorithm/remove>147 <http://en.cppreference.com/w/cpp/algorithm/remove>148 http://en.cppreference.com/w/cpp/algorithm/remove_copy149 http://en.cppreference.com/w/cpp/algorithm/remove_copy150 <http://en.cppreference.com/w/cpp/algorithm/replace>151 <http://en.cppreference.com/w/cpp/algorithm/replace>152 http://en.cppreference.com/w/cpp/algorithm/replace_copy153 http://en.cppreference.com/w/cpp/algorithm/replace_copy154 <http://en.cppreference.com/w/cpp/algorithm/reverse>155 http://en.cppreference.com/w/cpp/algorithm/reverse_copy156 <http://en.cppreference.com/w/cpp/algorithm/rotate>157 http://en.cppreference.com/w/cpp/algorithm/rotate_copy158 http://en.cppreference.com/w/cpp/algorithm/swap_ranges159 <http://en.cppreference.com/w/cpp/algorithm/transform>160 <http://en.cppreference.com/w/cpp/algorithm/unique>161 http://en.cppreference.com/w/cpp/algorithm/unique_copy162 <http://en.cppreference.com/w/cpp/algorithm/merge>163 http://en.cppreference.com/w/cpp/algorithm/inplace_merge164 <http://en.cppreference.com/w/cpp/algorithm/includes>165 http://en.cppreference.com/w/cpp/algorithm/set_difference166 http://en.cppreference.com/w/cpp/algorithm/set_intersection167 http://en.cppreference.com/w/cpp/algorithm/set_symmetric_difference168 http://en.cppreference.com/w/cpp/algorithm/set_union

Table 2.18: Heap operations (In Header:
<hpx/include/parallel_algorithm.hpp>)

Name	Description	In header	Algorithm page at cppreference.com
<i>hpx::parallel::v1::is_max_heap</i>	Returns true if the range is max heap.	<i><hpx/include/parallel_algorithm.hpp></i>	<i>is_heap</i> ¹⁶⁹
<i>hpx::parallel::v1::is_heap_until</i>	Returns the first element that breaks a max heap.	<i><hpx/include/parallel_algorithm.hpp></i>	<i>is_heap_until</i> ¹⁷⁰

Table 2.19: Minimum/maximum operations (In Header:
<hpx/include/parallel_algorithm.hpp>)

Name	Description	In header	Algorithm page at cppreference.com
<i>hpx::parallel::v1::max_element</i>	Returns the largest element in a range.	<i><hpx/include/parallel_minmax.hpp></i>	<i>max_element</i> ¹⁷¹
<i>hpx::parallel::v1::min_element</i>	Returns the smallest element in a range.	<i><hpx/include/parallel_minmax.hpp></i>	<i>min_element</i> ¹⁷²
<i>hpx::parallel::v1::minmax_element</i>	Returns the smallest and the largest element in a range.	<i><hpx/include/parallel_minmax.hpp></i>	<i>minmax_element</i> ¹⁷³

Table 2.20: Partitioning Operations (In Header:
<hpx/include/parallel_algorithm.hpp>)

Name	Description	In header	Algorithm page at cppreference.com
<i>hpx::parallel::v1::is_partitioned</i>	Returns <code>partie</code> if each true element for a predicate precedes the false elements in a range.	<i><hpx/include/parallel_is_partitioned.hpp></i>	<i>is_partitioned</i> ¹⁷⁴
<i>hpx::parallel::v1::partition</i>	Divides elements into two groups without preserving their relative order.	<i><hpx/include/parallel_partition.hpp></i>	<i>partition</i> ¹⁷⁵
<i>hpx::parallel::v1::partition_copy</i>	Copies a range dividing the elements into two groups.	<i><hpx/include/parallel_partition.hpp></i>	<i>partition_copy</i> ¹⁷⁶
<i>hpx::parallel::v1::stable_partition</i>	Divides elements into two groups while preserving their relative order.	<i><hpx/include/parallel_partition.hpp></i>	<i>stable_partition</i> ¹⁷⁷

¹⁶⁹ http://en.cppreference.com/w/cpp/algorithms/is_heap

¹⁷⁰ http://en.cppreference.com/w/cpp/algorithms/is_heap_until

¹⁷¹ http://en.cppreference.com/w/cpp/algorithms/max_element

¹⁷² http://en.cppreference.com/w/cpp/algorithms/min_element

¹⁷³ http://en.cppreference.com/w/cpp/algorithms/minmax_element

¹⁷⁴ http://en.cppreference.com/w/cpp/algorithms/is_partitioned

¹⁷⁵ <http://en.cppreference.com/w/cpp/algorithms/partition>

¹⁷⁶ http://en.cppreference.com/w/cpp/algorithms/partition_copy

¹⁷⁷ http://en.cppreference.com/w/cpp/algorithms/stable_partition

Table 2.21: Sorting Operations (In Header:
`<hpx/include/parallel_algorithm.hpp>`)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::v1::is_sorted</code>	Returns <code>true</code> if each element in a range is sorted.	<code><hpx/include/parallel_is_sorted.hpp></code>	<code>is_sorted</code> ¹⁷⁸
<code>hpx::parallel::v1::is_sorted_until</code>	Returns <code>true</code> if the first unsorted element.	<code><hpx/include/parallel_is_sorted.hpp></code>	<code>is_sorted_until</code> ¹⁷⁹
<code>hpx::parallel::v1::sort</code>	Sorts the elements in a range.	<code><hpx/include/parallel_sort.hpp></code>	<code>sort</code> ¹⁸⁰
<code>hpx::parallel::v1::sort_one_range</code>	Sorts <u>one</u> range of data using keys supplied in another range.	<code><hpx/include/parallel_sort.hpp></code>	

Table 2.22: Numeric Parallel Algorithms (In Header:
`<hpx/include/parallel_numeric.hpp>`)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::adjacent_difference</code>	Calculates the <u>difference</u> between each element in an input range and the preceding element.	<code><hpx/include/parallel_adjacent_difference.hpp></code>	<code>adjacent_difference</code> ¹⁸¹
<code>hpx::parallel::reduce</code>	Sums up a range of elements.	<code><hpx/include/parallel_reduce.hpp></code>	<code>reduce</code> ¹⁸²
<code>hpx::parallel::reduce</code>	Performs an <u>inclusive scan</u> on consecutive elements with matching keys, with a reduction to output only the final sum for each key. The key sequence {1, 1, 1, 2, 3, 3, 3, 3, 1} and value sequence {2, 3, 4, 5, 6, 7, 8, 9, 10} would be reduced to <code>keys={1, 2, 3, 1}</code> , <code>values={9, 5, 30, 10}</code> .	<code><hpx/include/parallel_reduce.hpp></code>	
<code>hpx::parallel::transform_reduce</code>	Sums up a range of elements after applying a function. Also, accumulates the inner products of two input ranges.	<code><hpx/include/parallel_transform_reduce.hpp></code>	<code>transform_reduce</code> ¹⁸³
<code>hpx::parallel::transform_inclusive_scan</code>	Does an <u>inclusive parallel scan</u> over a range of elements after applying a function.	<code><hpx/include/parallel_scan.hpp></code>	<code>transform_inclusive_scan</code> ¹⁸⁴
<code>hpx::parallel::transform_exclusive_scan</code>	Does an <u>exclusive parallel scan</u> over a range of elements after applying a function.	<code><hpx/include/parallel_scan.hpp></code>	<code>transform_exclusive_scan</code> ¹⁸⁵

¹⁷⁸ http://en.cppreference.com/w/cpp/algorithm/is_sorted

¹⁷⁹ http://en.cppreference.com/w/cpp/algorithm/is_sorted_until

¹⁸⁰ <http://en.cppreference.com/w/cpp/algorithm/sort>

Table 2.23: Dynamic Memory Management (In Header:
`<hpx/include/parallel_memory.hpp>`)

Name	Description	In header	Algorithm page at cppreference.com
<code>hpx::parallel::v1::destroy</code>	Destroys a range of objects.	<code><hpx/include/parallel_destroy.hpp></code>	<code>destroy</code> ¹⁸⁶
<code>hpx::parallel::v1::destroy_n</code>	Destroys a range of objects.	<code><hpx/include/parallel_destroy.hpp></code>	<code>destroy_n</code> ¹⁸⁷
<code>hpx::parallel::v1::uninitialized_copy</code>	Copies a range of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_copy.hpp></code>	<code>uninitialized_copy</code> ¹⁸⁸
<code>hpx::parallel::v1::uninitialized_copy_n</code>	Copies a number of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_copy.hpp></code>	<code>uninitialized_copy_n</code> ¹⁸⁹
<code>hpx::parallel::v1::uninitialized_deallocate</code>	Copies a range of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_deallocate.hpp></code>	<code>uninitialized_deallocate</code> ¹⁹⁰
<code>hpx::parallel::v1::uninitialized_deallocate_n</code>	Copies a number of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_deallocate.hpp></code>	<code>uninitialized_deallocate_n</code> ¹⁹¹
<code>hpx::parallel::v1::uninitialized_fill</code>	Copies an object to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_fill.hpp></code>	<code>uninitialized_fill</code> ¹⁹²
<code>hpx::parallel::v1::uninitialized_fill_n</code>	Copies an object to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_fill.hpp></code>	<code>uninitialized_fill_n</code> ¹⁹³
<code>hpx::parallel::v1::uninitialized_move</code>	Moves a range of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_move.hpp></code>	<code>uninitialized_move</code> ¹⁹⁴
<code>hpx::parallel::v1::uninitialized_move_n</code>	Moves a number of objects to an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_move.hpp></code>	<code>uninitialized_move_n</code> ¹⁹⁵
<code>hpx::parallel::v1::uninitialized_value_construct</code>	Constructs objects in an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_value_construct.hpp></code>	<code>uninitialized_value_construct</code> ¹⁹⁶
<code>hpx::parallel::v1::uninitialized_value_construct_n</code>	Constructs objects in an uninitialized area of memory.	<code><hpx/include/parallel_uninitialized_value_construct.hpp></code>	<code>uninitialized_value_construct_n</code> ¹⁹⁷

¹⁸¹ http://en.cppreference.com/w/cpp/algorithm/adjacent_difference

¹⁸² <http://en.cppreference.com/w/cpp/algorithm/reduce>

¹⁸³ http://en.cppreference.com/w/cpp/algorithm/transform_reduce

¹⁸⁴ http://en.cppreference.com/w/cpp/algorithm/transform_inclusive_scan

¹⁸⁵ http://en.cppreference.com/w/cpp/algorithm/transform_exclusive_scan

Table 2.24: Index-based for-loops (In Header:
<hpx/include/parallel_algorithm.hpp>)

Name	Description	In header
<code>hpx::parallel::v2::for</code>	Implements loop functionality over a range specified by integral or iterator bounds.	<code><hpx/include/parallel_for_loop.hpp></code>
<code>hpx::parallel::v2::for</code>	Implements loop functionality over a range specified by integral or iterator bounds.	<code><hpx/include/parallel_for_loop.hpp></code>
<code>hpx::parallel::v2::for</code>	Implements loop functionality over a range specified by integral or iterator bounds.	<code><hpx/include/parallel_for_loop.hpp></code>
<code>hpx::parallel::v2::for</code>	Implements loop functionality over a range specified by integral or iterator bounds.	<code><hpx/include/parallel_for_loop.hpp></code>

Executor parameters and executor parameter traits

HPX introduces the notion of execution parameters and execution parameter traits. At this point, the only parameter that can be customized is the size of the chunks of work executed on a single HPX thread (such as the number of loop iterations combined to run as a single task).

An executor parameter object is responsible for exposing the calculation of the size of the chunks scheduled. It abstracts the (potentially platform-specific) algorithms of determining those chunk sizes.

The way executor parameters are implemented is aligned with the way executors are implemented. All functionalities of concrete executor parameter types are exposed and accessible through a corresponding `hpx::parallel::executor_parameter_traits` type.

With `executor_parameter_traits`, clients access all types of executor parameters uniformly:

```
std::size_t chunk_size =
    executor_parameter_traits<my_parameter_t>::get_chunk_size(my_parameter,
        my_executor, [](){ return 0; }, num_tasks);
```

This call synchronously retrieves the size of a single chunk of loop iterations (or similar) to combine for execution on a single HPX thread if the overall number of tasks to schedule is given by `num_tasks`. The lambda function exposes a means of test-probing the execution of a single iteration for performance measurement purposes. The execution parameter type might dynamically determine the execution time of one or more tasks in order to calculate the chunk size; see `hpx::parallel::execution::auto_chunk_size` for an example of this executor parameter type.

Other functions in the interface exist to discover whether an executor parameter type should be invoked once (i.e., it returns a static chunk size; see `hpx::parallel::execution::static_chunk_size`) or whether

¹⁸⁶ <http://en.cppreference.com/w/cpp/memory/destroy>

¹⁸⁷ http://en.cppreference.com/w/cpp/memory/destroy_n

¹⁸⁸ http://en.cppreference.com/w/cpp/memory/uninitialized_copy

¹⁸⁹ http://en.cppreference.com/w/cpp/memory/uninitialized_copy_n

¹⁹⁰ http://en.cppreference.com/w/cpp/memory/uninitialized_default_construct

¹⁹¹ http://en.cppreference.com/w/cpp/memory/uninitialized_default_construct_n

¹⁹² http://en.cppreference.com/w/cpp/memory/uninitialized_fill

¹⁹³ http://en.cppreference.com/w/cpp/memory/uninitialized_fill_n

¹⁹⁴ http://en.cppreference.com/w/cpp/memory/uninitialized_move

¹⁹⁵ http://en.cppreference.com/w/cpp/memory/uninitialized_move_n

¹⁹⁶ http://en.cppreference.com/w/cpp/memory/uninitialized_value_construct

¹⁹⁷ http://en.cppreference.com/w/cpp/memory/uninitialized_value_construct_n

it should be invoked for each scheduled chunk of work (i.e., it returns a variable chunk size; for an example, see `hpx::parallel::execution::guided_chunk_size`).

Although this interface appears to require executor parameter type authors to implement all different basic operations, none are required. In practice, all operations have sensible defaults. However, some executor parameter types will naturally specialize all operations for maximum efficiency.

HPX implements the following executor parameter types:

- `hpx::parallel::execution::auto_chunk_size`: Loop iterations are divided into pieces and then assigned to threads. The number of loop iterations combined is determined based on measurements of how long the execution of 1% of the overall number of iterations takes. This executor parameter type makes sure that as many loop iterations are combined as necessary to run for the amount of time specified.
- `hpx::parallel::execution::static_chunk_size`: Loop iterations are divided into pieces of a given size and then assigned to threads. If the size is not specified, the iterations are, if possible, evenly divided contiguously among the threads. This executor parameters type is equivalent to OpenMP's STATIC scheduling directive.
- `hpx::parallel::execution::dynamic_chunk_size`: Loop iterations are divided into pieces of a given size and then dynamically scheduled among the cores; when a core finishes one chunk, it is dynamically assigned another. If the size is not specified, the default chunk size is 1. This executor parameter type is equivalent to OpenMP's DYNAMIC scheduling directive.
- `hpx::parallel::execution::guided_chunk_size`: Iterations are dynamically assigned to cores in blocks as cores request them until no blocks remain to be assigned. This is similar to `dynamic_chunk_size` except that the block size decreases each time a number of loop iterations is given to a thread. The size of the initial block is proportional to `number_of_iterations / number_of_cores`. Subsequent blocks are proportional to `number_of_iterations_remaining / number_of_cores`. The optional chunk size parameter defines the minimum block size. The default minimal chunk size is 1. This executor parameter type is equivalent to OpenMP's GUIDED scheduling directive.

Using task blocks

The `define_task_block`, `run` and the `wait` functions implemented based on N4411¹⁹⁸ are based on the `task_block` concept that is a part of the common subset of the Microsoft Parallel Patterns Library (PPL)¹⁹⁹ and the Intel Threading Building Blocks (TBB)²⁰⁰ libraries.

These implementations adopt a simpler syntax than exposed by those libraries—one that is influenced by language-based concepts, such as `spawn` and `sync` from Cilk++²⁰¹ and `async` and `finish` from X10²⁰². They improve on existing practice in the following ways:

- The exception handling model is simplified and more consistent with normal C++ exceptions.
- Most violations of strict fork-join parallelism can be enforced at compile time (with compiler assistance, in some cases).
- The syntax allows scheduling approaches other than child stealing.

Consider an example of a parallel traversal of a tree, where a user-provided function `compute` is applied to each node of the tree, returning the sum of the results:

¹⁹⁸ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4411.pdf>

¹⁹⁹ <https://msdn.microsoft.com/en-us/library/dd492418.aspx>

²⁰⁰ <https://www.threadingbuildingblocks.org/>

²⁰¹ <https://software.intel.com/en-us/articles/intel-cilk-plus/>

²⁰² <https://x10-lang.org/>

```
template <typename Func>
int traverse(node& n, Func && compute)
{
    int left = 0, right = 0;
    define_task_block(
        [&](task_block<>& tr) {
            if (n.left)
                tr.run([&] { left = traverse(*n.left, compute); });
            if (n.right)
                tr.run([&] { right = traverse(*n.right, compute); });
        });

    return compute(n) + left + right;
}
```

The example above demonstrates the use of two of the functions, `hpx::parallel::define_task_block` and the `hpx::parallel::task_block::run` member function of a `hpx::parallel::task_block`.

The `task_block` function delineates a region in a program code potentially containing invocations of threads spawned by the `run` member function of the `task_block` class. The `run` function spawns an *HPX* thread, a unit of work that is allowed to execute in parallel with respect to the caller. Any parallel tasks spawned by `run` within the task block are joined back to a single thread of execution at the end of the `define_task_block`. `run` takes a user-provided function object `f` and starts it asynchronously—i.e., it may return before the execution of `f` completes. The *HPX* scheduler may choose to run `f` immediately or delay running `f` until compute resources become available.

A `task_block` can be constructed only by `define_task_block` because it has no public constructors. Thus, `run` can be invoked directly or indirectly only from a user-provided function passed to `define_task_block`:

```
void g();

void f(task_block<>& tr)
{
    tr.run(g);           // OK, invoked from within task_block in h
}

void h()
{
    define_task_block(f);
}

int main()
{
    task_block<> tr;   // Error: no public constructor
    tr.run(g);          // No way to call run outside of a define_task_block
    return 0;
}
```

Extensions for task blocks

Using execution policies with task blocks

HPX implements some extensions for `task_block` beyond the actual standards proposal N4411²⁰³. The main addition is that a `task_block` can be invoked with an execution policy as its first argument, very similar to the parallel algorithms.

An execution policy is an object that expresses the requirements on the ordering of functions invoked as a consequence of the invocation of a task block. Enabling passing an execution policy to `define_task_block` gives the user control over the amount of parallelism employed by the created `task_block`. In the following example the use of an explicit `par` execution policy makes the user's intent explicit:

```
template <typename Func>
int traverse(node *n, Func&& compute)
{
    int left = 0, right = 0;

    define_task_block(
        execution::par,                                // execution::parallel_policy
        [&](task_block<>& tb) {
            if (n->left)
                tb.run([&] { left = traverse(n->left, compute); });
            if (n->right)
                tb.run([&] { right = traverse(n->right, compute); });
        });

    return compute(n) + left + right;
}
```

This also causes the `hpx::parallel::v2::task_block` object to be a template in our implementation. The template argument is the type of the execution policy used to create the task block. The template argument defaults to `hpx::parallel::execution::parallel_policy`.

HPX still supports calling `hpx::parallel::v2::define_task_block` without an explicit execution policy. In this case the task block will run using the `hpx::parallel::execution::parallel_policy`.

HPX also adds the ability to access the execution policy that was used to create a given `task_block`.

Using executors to run tasks

Often, users want to be able to not only define an execution policy to use by default for all spawned tasks inside the task block, but also to customize the execution context for one of the tasks executed by `task_block::run`. Adding an optionally passed executor instance to that function enables this use case:

```
template <typename Func>
int traverse(node *n, Func&& compute)
{
    int left = 0, right = 0;

    define_task_block(
        execution::par,                                // execution::parallel_policy
        [&] (auto& tb) {
            if (n->left)
```

(continues on next page)

²⁰³ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4411.pdf>

(continued from previous page)

```

{
    // use explicitly specified executor to run this task
    tb.run(my_executor(), [&] { left = traverse(n->left, compute); });
}
if (n->right)
{
    // use the executor associated with the par execution policy
    tb.run([&] { right = traverse(n->right, compute); });
}
};

return compute(n) + left + right;
}

```

HPX still supports calling `hpx::parallel::v2::task_block::run` without an explicit executor object. In this case the task will be run using the executor associated with the execution policy that was used to call `hpx::parallel::v2::define_task_block`.

2.5.7 Writing distributed HPX applications

This section focuses on the features of *HPX* needed to write distributed applications, namely the *Active Global Address Space (AGAS)*, remotely executable functions (i.e. *actions*), and distributed objects (i.e. *components*).

Global names

HPX implements an *Active Global Address Space (AGAS)* which is exposing a single uniform address space spanning all localities an application runs on. *AGAS* is a fundamental component of the ParalleX execution model. Conceptually, there is no rigid demarcation of local or global memory in *AGAS*; all available memory is a part of the same address space. *AGAS* enables named objects to be moved (migrated) across localities without having to change the object's name, i.e., no references to migrated objects have to be ever updated. This feature has significance for dynamic load balancing and in applications where the workflow is highly dynamic, allowing work to be migrated from heavily loaded nodes to less loaded nodes. In addition, immutability of names ensures that *AGAS* does not have to keep extra indirections ("bread crumbs") when objects move, hence minimizing complexity of code management for system developers as well as minimizing overheads in maintaining and managing aliases.

The *AGAS* implementation in *HPX* does not automatically expose every local address to the global address space. It is the responsibility of the programmer to explicitly define which of the objects have to be globally visible and which of the objects are purely local.

In *HPX* global addresses (global names) are represented using the `hpx::id_type` data type. This data type is conceptually very similar to `void*` pointers as it does not expose any type information of the object it is referring to.

The only predefined global addresses are assigned to all localities. The following *HPX* API functions allow one to retrieve the global addresses of localities:

- `hpx::find_here`: retrieve the global address of the *locality* this function is called on.
- `hpx::find_all_localities`: retrieve the global addresses of all localities available to this application (including the *locality* the function is being called on).
- `hpx::find_remote_localities`: retrieve the global addresses of all remote localities available to this application (not including the *locality* the function is being called on)
- `hpx::get_num_localities`: retrieve the number of localities available to this application.
- `hpx::find_locality`: retrieve the global address of any *locality* supporting the given component type.

- `hpx::get_colocation_id`: retrieve the global address of the *locality* currently hosting the object with the given global address.

Additionally, the global addresses of localities can be used to create new instances of components using the following HPX API function:

- `hpx::components::new_`: Create a new instance of the given Component type on the specified *locality*.

Note: HPX does not expose any functionality to delete component instances. All global addresses (as represented using `hpx::id_type`) are automatically garbage collected. When the last (global) reference to a particular component instance goes out of scope the corresponding component instance is automatically deleted.

Applying actions

Action type definition

Actions are special types we use to describe possibly remote operations. For every global function and every member function which has to be invoked distantly, a special type must be defined. For any global function the special macro `HPX_PLAIN_ACTION` can be used to define the action type. Here is an example demonstrating this:

```
namespace app
{
    void some_global_function(double d)
    {
        cout << d;
    }

    // This will define the action type 'some_global_action' which represents
    // the function 'app::some_global_function'.
    HPX_PLAIN_ACTION(app::some_global_function, some_global_action);
}
```

Important: The macro `HPX_PLAIN_ACTION` has to be placed in global namespace, even if the wrapped function is located in some other namespace. The newly defined action type is placed in the global namespace as well.

If the action type should be defined somewhere not in global namespace, the action type definition has to be split into two macro invocations (`HPX_DEFINE_PLAIN_ACTION` and `HPX_REGISTER_ACTION`) as shown in the next example:

```
namespace app
{
    void some_global_function(double d)
    {
        cout << d;
    }

    // On conforming compilers the following macro expands to:
    //
    //     typedef hpx::actions::make_action<
    //         decltype(&some_global_function), &some_global_function
    //     >::type some_global_action;
    //
    // This will define the action type 'some_global_action' which represents
}
```

(continues on next page)

(continued from previous page)

```
// the function 'some_global_function'.
HPX_DEFINE_PLAIN_ACTION(some_global_function, some_global_action);
}

// The following macro expands to a series of definitions of global objects
// which are needed for proper serialization and initialization support
// enabling the remote invocation of the function ``some_global_function``
HPX_REGISTER_ACTION(app::some_global_action, app_some_global_action);
```

The shown code defines an action type `some_global_action` inside the namespace `app`.

Important: If the action type definition is split between two macros as shown above, the name of the action type to create has to be the same for both macro invocations (here `some_global_action`).

Important: The second argument passed to `HPX_REGISTER_ACTION` (`app_some_global_action`) has to comprise a globally unique C++ identifier representing the action. This is used for serialization purposes.

For member functions of objects which have been registered with [AGAS](#) (e.g. ‘components’) a different registration macro `HPX_DEFINE_COMPONENT_ACTION` has to be utilized. Any component needs to be declared in a header file and have some special support macros defined in a source file. Here is an example demonstrating this. The first snippet has to go into the header file:

```
namespace app
{
    struct some_component
        : hpx::components::component_base<some_component>
    {
        int some_member_function(std::string s)
        {
            return boost::lexical_cast<int>(s);
        }

        // This will define the action type 'some_member_action' which
        // represents the member function 'some_member_function' of the
        // object type 'some_component'.
        HPX_DEFINE_COMPONENT_ACTION(some_component, some_member_function,
                                    some_member_action);
    };
}

// Note: The second argument to the macro below has to be systemwide-unique
//       C++ identifiers
HPX_REGISTER_ACTION_DECLARATION(app::some_component::some_member_action, some_
                                ↪component_some_action);
```

The next snippet belongs into a source file (e.g. the main application source file) in the simplest case:

```
typedef hpx::components::component<app::some_component> component_type;
typedef app::some_component some_component;

HPX_REGISTER_COMPONENT(component_type, some_component);

// The parameters for this macro have to be the same as used in the corresponding
```

(continues on next page)

(continued from previous page)

```
// HPX_REGISTER_ACTION_DECLARATION() macro invocation above
typedef some_component::some_member_action some_component_some_action;
HPX_REGISTER_ACTION(some_component_some_action);
```

Granted, these macro invocations are a bit more complex than for simple global functions, however we believe they are still manageable.

The most important macro invocation is the `HPX_DEFINE_COMPONENT_ACTION` in the header file as this defines the action type we need to invoke the member function. For a complete example of a simple component action see [hpx_link examples/quickstart/component_in_executable.cpp..component_in_executable.cpp]

Action invocation

The process of invoking a global function (or a member function of an object) with the help of the associated action is called ‘applying the action’. Actions can have arguments, which will be supplied while the action is applied. At the minimum, one parameter is required to apply any action - the id of the *locality* the associated function should be invoked on (for global functions), or the id of the component instance (for member functions). Generally, *HPX* provides several ways to apply an action, all of which are described in the following sections.

Generally, *HPX* actions are very similar to ‘normal’ C++ functions except that actions can be invoked remotely. Fig. 2.8 below shows an overview of the main API exposed by *HPX*. This shows the function invocation syntax as defined by the C++ language (dark gray), the additional invocation syntax as provided through C++ Standard Library features (medium gray), and the extensions added by *HPX* (light gray) where:

- f function to invoke,
- p...: (optional) arguments,
- R: return type of f,
- action: action type defined by, `HPX_DEFINE_PLAIN_ACTION` or `HPX_DEFINE_COMPONENT_ACTION` encapsulating f,
- a: an instance of the type `action,
- id: the global address the action is applied to.

R f(p...)	Synchronous Execution (returns R)	Asynchronous Execution (returns <code>future<R></code>)	Fire & Forget Execution (returns <code>void</code>)
Functions (direct invocation)	<code>f(p...)</code> C++	<code>async(f, p...)</code>	<code>apply(f, p...)</code>
Functions (lazy invocation)	<code>bind(f, p...)(...)</code>	<code>async(bind(f, p...), ...)</code> C++ Standard Library	<code>apply(bind(f, p...), ...)</code>
Actions (direct invocation)	<code>HPX_ACTION(f, action)</code> <code>a(id, p...)</code>	<code>HPX_ACTION(f, action)</code> <code>async(a, id, p...)</code>	<code>HPX_ACTION(f, action)</code> <code>apply(a, id, p...)</code>
Actions (lazy invocation)	<code>HPX_ACTION(f, action)</code> <code>bind(a, id, p...)(...)</code>	<code>HPX_ACTION(f, action)</code> <code>async(bind(a, id, p...), ...)</code>	<code>HPX_ACTION(f, action)</code> <code>apply(bind(a, id, p...), ...)</code> HPX

Fig. 2.8: Overview of the main API exposed by *HPX*.

This figure shows that *HPX* allows the user to apply actions with a syntax similar to the C++ standard. In fact, all action types have an overloaded function operator allowing to synchronously apply the action. Further, *HPX* implements `hpx::async` which semantically works similar to the way `std::async` works for plain C++ function.

Note: The similarity of applying an action to conventional function invocations extends even further. *HPX* implements `hpx::bind` and `hpx::function` two facilities which are semantically equivalent to the `std::bind` and `std::function` types as defined by the C++11 Standard. While `hpx::async` extends beyond the conventional semantics by supporting actions and conventional C++ functions, the *HPX* facilities `hpx::bind` and `hpx::function` extend beyond the conventional standard facilities too. The *HPX* facilities not only support conventional functions, but can be used for actions as well.

Additionally, *HPX* exposes `hpx::apply` and `hpx::async_continue` both of which refine and extend the standard C++ facilities.

The different ways to invoke a function in *HPX* will be explained in more detail in the following sections.

Applying an action asynchronously without any synchronization

This method ('fire and forget') will make sure the function associated with the action is scheduled to run on the target *locality*. Applying the action does not wait for the function to start running, instead it is a fully asynchronous operation. The following example shows how to apply the action as defined *in the previous section* on the local *locality* (the *locality* this code runs on):

```
some_global_action act;      // define an instance of some_global_action
hpx::apply(act, hpx::find_here(), 2.0);
```

(the function `hpx::find_here()` returns the id of the local *locality*, i.e. the *locality* this code executes on).

Any component member function can be invoked using the same syntactic construct. Given that `id` is the global address for a component instance created earlier, this invocation looks like:

```
some_component_action act;      // define an instance of some_component_action
hpx::apply(act, id, "42");
```

In this case any value returned from this action (e.g. in this case the integer 42 is ignored. Please look at [Action type definition](#) for the code defining the component action `some_component_action` used.

Applying an action asynchronously with synchronization

This method will make sure the action is scheduled to run on the target *locality*. Applying the action itself does not wait for the function to start running or to complete, instead this is a fully asynchronous operation similar to using `hpx::apply` as described above. The difference is that this method will return an instance of a `hpx::future<>` encapsulating the result of the (possibly remote) execution. The future can be used to synchronize with the asynchronous operation. The following example shows how to apply the action from above on the local *locality*:

```
some_global_action act;      // define an instance of some_global_action
hpx::future<void> f = hpx::async(act, hpx::find_here(), 2.0);
//
// ... other code can be executed here
//
f.get();      // this will possibly wait for the asynchronous operation to 'return'
```

(as before, the function `hpx::find_here()` returns the id of the local *locality* (the *locality* this code is executed on).

Note: The use of a `hpx::future<void>` allows the current thread to synchronize with any remote operation not returning any value.

Note: Any `std::future<>` returned from `std::async()` is required to block in its destructor if the value has not been set for this future yet. This is not true for `hpx::future<>` which will never block in its destructor, even if the value has not been returned to the future yet. We believe that consistency in the behavior of futures is more important than standards conformance in this case.

Any component member function can be invoked using the same syntactic construct. Given that `id` is the global address for a component instance created earlier, this invocation looks like:

```
some_component_action act;      // define an instance of some_component_action
hpx::future<int> f = hpx::async(act, id, "42");
//
// ... other code can be executed here
//
cout << f.get();    // this will possibly wait for the asynchronous operation to
← 'return' 42
```

Note: The invocation of `f.get()` will return the result immediately (without suspending the calling thread) if the result from the asynchronous operation has already been returned. Otherwise, the invocation of `f.get()` will suspend the execution of the calling thread until the asynchronous operation returns its result.

Applying an action synchronously

This method will schedule the function wrapped in the specified action on the target *locality*. While the invocation appears to be synchronous (as we will see), the calling thread will be suspended while waiting for the function to return. Invoking a plain action (e.g. a global function) synchronously is straightforward:

```
some_global_action act;      // define an instance of some_global_action
act(hpx::find_here(), 2.0);
```

While this call looks just like a normal synchronous function invocation, the function wrapped by the action will be scheduled to run on a new thread and the calling thread will be suspended. After the new thread has executed the wrapped global function, the waiting thread will resume and return from the synchronous call.

Equivalently, any action wrapping a component member function can be invoked synchronously as follows:

```
some_component_action act;      // define an instance of some_component_action
int result = act(id, "42");
```

The action invocation will either schedule a new thread locally to execute the wrapped member function (as before, `id` is the global address of the component instance the member function should be invoked on), or it will send a parcel to the remote *locality* of the component causing a new thread to be scheduled there. The calling thread will be suspended until the function returns its result. This result will be returned from the synchronous action invocation.

It is very important to understand that this ‘synchronous’ invocation syntax in fact conceals an asynchronous function call. This is beneficial as the calling thread is suspended while waiting for the outcome of a potentially remote

operation. The *HPX* thread scheduler will schedule other work in the meantime, allowing the application to make further progress while the remote result is computed. This helps overlapping computation with communication and hiding communication latencies.

Note: The syntax of applying an action is always the same, regardless whether the target *locality* is remote to the invocation *locality* or not. This is a very important feature of *HPX* as it frees the user from the task of keeping track what actions have to be applied locally and which actions are remote. If the target for applying an action is local, a new thread is automatically created and scheduled. Once this thread is scheduled and run, it will execute the function encapsulated by that action. If the target is remote, *HPX* will send a parcel to the remote *locality* which encapsulates the action and its parameters. Once the parcel is received on the remote *locality* *HPX* will create and schedule a new thread there. Once this thread runs on the remote *locality*, it will execute the function encapsulated by the action.

Applying an action with a continuation but without any synchronization

This method is very similar to the method described in section [Applying an action asynchronously without any synchronization](#). The difference is that it allows the user to chain a sequence of asynchronous operations, while handing the (intermediate) results from one step to the next step in the chain. Where `hpx::apply` invokes a single function using ‘fire and forget’ semantics, `hpx::apply_continue` triggers a chain of functions without the need for the execution flow ‘to come back’ to the invocation site. Each of the asynchronous functions can be executed on a different *locality*.

Applying an action with a continuation and with synchronization

This method is very similar to the method described in section [Applying an action asynchronously with synchronization](#). In addition to what `hpx::async` can do, the functions `hpx::async_continue` takes an additional function argument. This function will be called as the continuation of the executed action. It is expected to perform additional operations and to make sure that a result is returned to the original invocation site. This method chains operations asynchronously by providing a continuation operation which is automatically executed once the first action has finished executing.

As an example we chain two actions, where the result of the first action is forwarded to the second action and the result of the second action is sent back to the original invocation site:

```
// first action
std::int32_t action1(std::int32_t i)
{
    return i+1;
}
HPX_PLAIN_ACTION(action1);      // defines action1_type

// second action
std::int32_t action2(std::int32_t i)
{
    return i*2;
}
HPX_PLAIN_ACTION(action2);      // defines action2_type

// this code invokes 'action1' above and passes along a continuation
// function which will forward the result returned from 'action1' to
// 'action2'.
action1_type act1;            // define an instance of 'action1_type'
action2_type act2;            // define an instance of 'action2_type'
```

(continues on next page)

(continued from previous page)

```
hpx::future<int> f =
    hpx::async_continue(act1, hpx::make_continuation(act2),
        hpx::find_here(), 42);
hpx::cout << f.get() << "\n"; // will print: 86 ((42 + 1) * 2)
```

By default, the continuation is executed on the same *locality* as `hpx::async_continue` is invoked from. If you want to specify the *locality* where the continuation should be executed, the code above has to be written as:

```
// this code invokes 'action1' above and passes along a continuation
// function which will forward the result returned from 'action1' to
// 'action2'.
action1_type act1; // define an instance of 'action1_type'
action2_type act2; // define an instance of 'action2_type'
hpx::future<int> f =
    hpx::async_continue(act1, hpx::make_continuation(act2, hpx::find_here())),
    hpx::find_here(), 42);
hpx::cout << f.get() << "\n"; // will print: 86 ((42 + 1) * 2)
```

Similarly, it is possible to chain more than 2 operations:

```
action1_type act1; // define an instance of 'action1_type'
action2_type act2; // define an instance of 'action2_type'
hpx::future<int> f =
    hpx::async_continue(act1,
        hpx::make_continuation(act2, hpx::make_continuation(act1)),
        hpx::find_here(), 42);
hpx::cout << f.get() << "\n"; // will print: 87 ((42 + 1) * 2 + 1)
```

The function `hpx::make_continuation` creates a special function object which exposes the following prototype:

```
struct continuation
{
    template <typename Result>
    void operator()(hpx::id_type id, Result&& result) const
    {
        ...
    }
};
```

where the parameters passed to the overloaded function `operator()` are:

- the `id` is the global id where the final result of the asynchronous chain of operations should be sent to (in most cases this is the id of the `hpx::future` returned from the initial call to `hpx::async_continue`. Any custom continuation function should make sure this `id` is forwarded to the last operation in the chain).
- the `result` is the result value of the current operation in the asynchronous execution chain. This value needs to be forwarded to the next operation.

Note: All of those operations are implemented by the predefined continuation function object which is returned from `hpx::make_continuation`. Any (custom) function object used as a continuation should conform to the same interface.

Action error handling

Like in any other asynchronous invocation scheme it is important to be able to handle error conditions occurring while the asynchronous (and possibly remote) operation is executed. In *HPX* all error handling is based on standard C++ exception handling. Any exception thrown during the execution of an asynchronous operation will be transferred back to the original invocation *locality*, where it is rethrown during synchronization with the calling thread.

Important: Exceptions thrown during asynchronous execution can be transferred back to the invoking thread only for the synchronous and the asynchronous case with synchronization. Like with any other unhandled exception, any exception thrown during the execution of an asynchronous action *without* synchronization will result in calling `hpx::terminate` causing the running application to exit immediately.

Note: Even if error handling internally relies on exceptions, most of the API functions exposed by *HPX* can be used without throwing an exception. Please see [Working with exceptions](#) for more information.

As an example, we will assume that the following remote function will be executed:

```
namespace app
{
    void some_function_with_error(int arg)
    {
        if (arg < 0) {
            HPX_THROW_EXCEPTION(bad_parameter, "some_function_with_error",
                "some really bad error happened");
        }
        // do something else...
    }
}

// This will define the action type 'some_error_action' which represents
// the function 'app::some_function_with_error'.
HPX_PLAIN_ACTION(app::some_function_with_error, some_error_action);
```

The use of `HPX_THROW_EXCEPTION` to report the error encapsulates the creation of a `hpx::exception` which is initialized with the error code `hpx::bad_parameter`. Additionally it carries the passed strings, the information about the file name, line number, and call stack of the point the exception was thrown from.

We invoke this action using the synchronous syntax as described before:

```
// note: wrapped function will throw hpx::exception
some_error_action act;           // define an instance of some_error_action
try {
    act(hpx::find_here(), -3);   // exception will be rethrown from here
}
catch (hpx::exception const& e) {
    // prints: 'some really bad error happened: HPX(bad parameter)'
    cout << e.what();
}
```

If this action is invoked asynchronously with synchronization, the exception is propagated to the waiting thread as well and is re-thrown from the future's function `get()`:

```
// note: wrapped function will throw hpx::exception
some_error_action act;           // define an instance of some_error_action
```

(continues on next page)

(continued from previous page)

```

hpx::future<void> f = hpx::async(act, hpx::find_here(), -3);
try {
    f.get();                                // exception will be rethrown from here
}
catch (hpx::exception const& e) {
    // prints: 'some really bad error happened: HPX(bad parameter)'
    cout << e.what();
}

```

For more information about error handling please refer to the section [Working with exceptions](#). There we also explain how to handle error conditions without having to rely on exception.

Writing components

A component in *HPX* is a C++ class which can be created remotely and for which its member functions can be invoked remotely as well. The following sections highlight how components can be defined, created, and used.

Defining components

In order for a C++ class type to be managed remotely in *HPX*, the type must be derived from the `hpx::components::component_base` template type. We call such C++ class types ‘components’.

Note that the component type itself is passed as a template argument to the base class:

```

// header file some_component.hpp

#include <hpx/include/components.hpp>

namespace app
{
    // Define a new component type 'some_component'
    struct some_component
        : hpx::components::component_base<some_component>
    {
        // This member function is has to be invoked remotely
        int some_member_function(std::string const& s)
        {
            return boost::lexical_cast<int>(s);
        }

        // This will define the action type 'some_member_action' which
        // represents the member function 'some_member_function' of the
        // object type 'some_component'.
        HPX_DEFINE_COMPONENT_ACTION(some_component, some_member_function, some_member_
        ↪action);
    };
}

// This will generate the necessary boiler-plate code for the action allowing
// it to be invoked remotely. This declaration macro has to be placed in the
// header file defining the component itself.
// Note: The second argument to the macro below has to be systemwide-unique
//       C++ identifiers

```

(continues on next page)

(continued from previous page)

```
//  
HPX_REGISTER_ACTION_DECLARATION(app::some_component::some_member_action, some_  
    ↪component_some_action);
```

There is more boiler plate code which has to be placed into a source file in order for the component to be usable. Every component type is required to have macros placed into its source file, one for each component type and one macro for each of the actions defined by the component type.

For instance:

```
// source file some_component.cpp  
  
#include "some_component.hpp"  
  
// The following code generates all necessary boiler plate to enable the  
// remote creation of 'app::some_component' instances with 'hpx::new_<>()'  
//  
using some_component = app::some_component;  
using some_component_type = hpx::components::component<some_component>;  
  
// Please note that the second argument to this macro must be a  
// (system-wide) unique C++-style identifier (without any namespaces)  
//  
HPX_REGISTER_COMPONENT(some_component_type, some_component);  
  
// The parameters for this macro have to be the same as used in the corresponding  
// HPX_REGISTER_ACTION_DECLARATION() macro invocation in the corresponding  
// header file.  
//  
// Please note that the second argument to this macro must be a  
// (system-wide) unique C++-style identifier (without any namespaces)  
//  
HPX_REGISTER_ACTION(app::some_component::some_member_action, some_component_some_  
    ↪action);
```

Defining client side representation classes

Often it is very convenient to define a separate type for a component which can be used on the client side (from where the component is instantiated and used). This step might seem as unnecessary duplicating code, however it significantly increases the type safety of the code.

A possible implementation of such a client side representation for the component described in the previous section could look like:

```
#include <hpx/include/components.hpp>  
  
namespace app  
{  
    // Define a client side representation type for the component type  
    // 'some_component' defined in the previous section.  
    //  
    struct some_component_client  
        : hpx::components::client_base<some_component_client, some_component>  
    {  
        using base_type = hpx::components::client_base<
```

(continues on next page)

(continued from previous page)

```

        some_component_client, some_component>;

some_component_client(hpx::future<hpx::id_type> && id)
    : base_type(std::move(id))
{ }

hpx::future<int> some_member_function(std::string const& s)
{
    some_component::some_member_action act;
    return hpx::async(act, get_id(), s);
}
};

}

```

A client side object stores the global id of the component instance it represents. This global id is accessible by calling the function `client_base<>::get_id()`. The special constructor which is provided in the example allows to create this client side object directly using the API function `hpx::new_`.

Creating component instances

Instances of defined component types can be created in two different ways. If the component to create has a defined client side representation type, then this can be used, otherwise use the server type.

The following examples assume that `some_component_type` is the type of the server side implementation of the component to create. All additional arguments (see , ... notation below) are passed through to the corresponding constructor calls of those objects:

```

// create one instance on the given locality
hpx::id_type here = hpx::find_here();
hpx::future<hpx::id_type> f =
    hpx::new_<some_component_type>(here, ...);

// create one instance using the given distribution
// policy (here: hpx::colocating_distribution_policy)
hpx::id_type here = hpx::find_here();
hpx::future<hpx::id_type> f =
    hpx::new_<some_component_type>(hpx::colocated(here), ...);

// create multiple instances on the given locality
hpx::id_type here = find_here();
hpx::future<std::vector<hpx::id_type>> f =
    hpx::new_<some_component_type[]>(here, num, ...);

// create multiple instances using the given distribution
// policy (here: hpx::binpacking_distribution_policy)
hpx::future<std::vector<hpx::id_type>> f = hpx::new_<some_component_type[]>(
    hpx::binpacking(hpx::find_all_localities()), num, ...);

```

The examples below demonstrate the use of the same API functions for creating client side representation objects (instead of just plain ids). These examples assume that `client_type` is the type of the client side representation of the component type to create. As above, all additional arguments (see , ... notation below) are passed through to the corresponding constructor calls of the server side implementation objects corresponding to the `client_type`:

```

// create one instance on the given locality
hpx::id_type here = hpx::find_here();

```

(continues on next page)

(continued from previous page)

```

client_type c = hpx::new<client_type>(here, ...);

// create one instance using the given distribution
// policy (here: hpx:: colocating_distribution_policy)
hpx::id_type here = hpx::find_here();
client_type c = hpx::new<client_type>(hpx::colocated(here), ...);

// create multiple instances on the given locality
hpx::id_type here = hpx::find_here();
hpx::future<std::vector<client_type>> f =
    hpx::new<client_type[]>(here, num, ...);

// create multiple instances using the given distribution
// policy (here: hpx:: binpacking_distribution_policy)
hpx::future<std::vector<client_type>> f = hpx::new<client_type[]>(
    hpx::binpacking(hpx::find_all_localities()), num, ...);

```

Using component instances

Segmented containers

In parallel programming, there is now a plethora of solutions aimed at implementing “partially contiguous” or segmented data structures, whether on shared memory systems or distributed memory systems. *HPX* implements such structures by drawing inspiration from Standard C++ containers.

Using segmented containers

A segmented container is a template class that is described in the namespace `hpx`. All segmented containers are very similar semantically to their sequential counterpart (defined in namespace `std` but with an additional template parameter named `DistPolicy`). The distribution policy is an optional parameter that is passed last to the segmented container constructor (after the container size when no default value is given, after the default value if not). The distribution policy describes the manner in which a container is segmented and the placement of each segment among the available runtime localities.

However, only a part of the `std` container member functions were reimplemented:

- (constructor), (destructor), operator=
- operator[]
- begin, cbegin, end, cend
- size

An example of how to use the `partitioned_vector` container would be:

```

#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(double);

// By default, the number of segments is equal to the current number of
// localities

```

(continues on next page)

(continued from previous page)

```
//  
hpx::partitioned_vector<double> va(50);  
hpx::partitioned_vector<double> vb(50, 0.0);
```

An example of how to use the `partitioned_vector` container with distribution policies would be:

```
#include <hpx/include/partitioned_vector.hpp>  
#include <hpx/runtime/find_localities.hpp>  
  
// The following code generates all necessary boiler plate to enable the  
// remote creation of 'partitioned_vector' segments  
//  
HPX_REGISTER_PARTITIONED_VECTOR(double);  
  
std::size_t num_segments = 10;  
std::vector<hpx::id_type> locs = hpx::find_all_localities();  
  
auto layout =  
    hpx::container_layout( num_segments, locs );  
  
// The number of segments is 10 and those segments are spread across the  
// localities collected in the variable locs in a Round-Robin manner  
//  
hpx::partitioned_vector<double> va(50, layout);  
hpx::partitioned_vector<double> vb(50, 0.0, layout);
```

By definition, a segmented container must be accessible from any thread although its construction is synchronous only for the thread who has called its constructor. To overcome this problem, it is possible to assign a symbolic name to the segmented container:

```
#include <hpx/include/partitioned_vector.hpp>  
  
// The following code generates all necessary boiler plate to enable the  
// remote creation of 'partitioned_vector' segments  
//  
HPX_REGISTER_PARTITIONED_VECTOR(double);  
  
hpx::future<void> fserver = hpx::async(  
[] () {  
    hpx::partitioned_vector<double> v(50);  
  
    // Register the 'partitioned_vector' with the name "some_name"  
    //  
    v.register_as("some_name");  
  
    /* Do some code */  
});  
  
hpx::future<void> fclient =  
hpx::async(  
[] () {  
    // Naked 'partitioned_vector'  
    //  
    hpx::partitioned_vector<double> v;  
  
    // Now the variable v points to the same 'partitioned_vector' that has
```

(continues on next page)

(continued from previous page)

```
// been registered with the name "some_name"
//
v.connect_to("some_name");

/* Do some code */
});
```

Segmented containers

HPX provides the following segmented containers:

Table 2.25: Sequence containers

Name	Description	In header	Class page at cppreference.com
hpx::partitioned	Dynamic segmented contiguous array.	<hpx/include/partitioned_vector.hpp>	vector²⁰⁴

Table 2.26: Unordered associative containers

Name	Description	In header	Class page at cppreference.com
hpx::unordered_map	Segmented collection of key-value pairs, hashed by keys, keys are unique.	<hpx/include/unordered_map.hpp>	unordered_map²⁰⁵

Segmented iterators and segmented iterator traits

The basic iterator used in the STL library is only suitable for one-dimensional structures. The iterators we use in HPX must adapt to the segmented format of our containers. Our iterators are then able to know when incrementing themselves if the next element of type T is in the same data segment or in another segment. In this second case, the iterator will automatically point to the beginning of the next segment.

Note: Note that the dereference operation operator `*` does not directly return a reference of type `T&` but an intermediate object wrapping this reference. When this object is used as an l-value, a remote write operation is performed; When this object is used as an r-value, implicit conversion to `T` type will take care of performing remote read operation.

It is sometimes useful not only to iterate element by element, but also segment by segment, or simply get a local iterator in order to avoid additional construction costs at each referencing operations. To mitigate this need, the `hpx::traits::segmented_iterator_traits` are used.

With `segmented_iterator_traits` users can uniformly get the iterators which specifically iterates over segments (by providing a segmented iterator as a parameter), or get the local begin/end iterators of the nearest local segment (by providing a per-segment iterator as a parameter):

²⁰⁴ <http://en.cppreference.com/w/cpp/container/vector>

²⁰⁵ http://en.cppreference.com/w/cpp/container/unordered_map

```
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(double);

using iterator = hpx::partitioned_vector<T>::iterator;
using traits = hpx::traits::segmented_iterator_traits<iterator>;

hpx::partitioned_vector<T> v;
std::size_t count = 0;

auto seg_begin = traits::segment(v.begin());
auto seg_end = traits::segment(v.end());

// Iterate over segments
for (auto seg_it = seg_begin; seg_it != seg_end; ++seg_it)
{
    auto loc_begin = traits::begin(seg_it)
    auto loc_end = traits::end(seg_it);

    // Iterate over elements inside segments
    for (auto lit = loc_begin; lit != loc_end; ++lit, ++count)
    {
        *lit = count;
    }
}
```

Which is equivalent to:

```
hpx::partitioned_vector<T> v;
std::size_t count = 0;

auto begin = v.begin();
auto end = v.end();

for (auto it = begin; it != end; ++it, ++count)
{
    *it = count;
}
```

Using views

The use of multidimensional arrays is quite common in the numerical field whether to perform dense matrix operations or to process images. It exist many libraries which implement such object classes overloading their basic operators (e.g. ``+``, ``-``, ``*``, ``()`` , etc.). However, such operation becomes more delicate when the underlying data layout is segmented or when it is mandatory to use optimized linear algebra subroutines (i.e. BLAS subroutines).

Our solution is thus to relax the level of abstraction by allowing the user to work not directly on n-dimensionnal data, but on “n-dimensionnal collections of 1-D arrays”. The use of well-accepted techniques on contiguous data is thus preserved at the segment level, and the composability of the segments is made possible thanks to multidimensional array-inspired access mode.

Preface: Why SPMD?

Although *HPX* refutes by design this programming model, the *locality* plays a dominant role when it comes to implement vectorized code. To maximize local computations and avoid unneeded data transfers, a parallel section (or Single Programming Multiple Data section) is required. Because the use of global variables is prohibited, this parallel section is created via the RAII idiom.

To define a parallel section, simply write an action taking a `spmd_block` variable as a first parameter:

```
#include <hpx/collectives/spmd_block.hpp>

void bulk_function(hpx::lcos::spmd_block block /* , arg0, arg1, ... */)
{
    // Parallel section

    /* Do some code */
}
HPX_PLAIN_ACTION(bulk_function, bulk_action);
```

Note: In the following paragraphs, we will use the term “image” several times. An image is defined as a lightweight process whose entry point is a function provided by the user. It’s an “image of the function”.

The `spmd_block` class contains the following methods:

- [def Team information] `get_num_images`, `this_image`, `images_per_locality`
- [def Control statements] `sync_all`, `sync_images`

Here is a sample code summarizing the features offered by the `spmd_block` class:

```
#include <hpx/collectives/spmd_block.hpp>

void bulk_function(hpx::lcos::spmd_block block /* , arg0, arg1, ... */)
{
    std::size_t num_images = block.get_num_images();
    std::size_t this_image = block.this_image();
    std::size_t images_per_locality = block.images_per_locality();

    /* Do some code */

    // Synchronize all images in the team
    block.sync_all();

    /* Do some code */

    // Synchronize image 0 and image 1
    block.sync_images(0,1);

    /* Do some code */

    std::vector<std::size_t> vec_images = {2,3,4};

    // Synchronize images 2, 3 and 4
    block.sync_images(vec_images);

    // Alternative call to synchronize images 2, 3 and 4
    block.sync_images(vec_images.begin(), vec_images.end());
```

(continues on next page)

(continued from previous page)

```

/* Do some code */

// Non-blocking version of sync_all()
hpx::future<void> event =
    block.sync_all(hpx::launch::async);

// Callback waiting for 'event' to be ready before being scheduled
hpx::future<void> cb =
    event.then(
        [] (hpx::future<void>)
    {

        /* Do some code */

    });

// Finally wait for the execution tree to be finished
cb.get();
}
HPX_PLAIN_ACTION(bulk_test_function, bulk_test_action);

```

Then, in order to invoke the parallel section, call the function `define_spmd_block` specifying an arbitrary symbolic name and indicating the number of images per `locality` to create:

```

void bulk_function(hpx::lcos::spmd_block block, /* , arg0, arg1, ... */)
{
}

HPX_PLAIN_ACTION(bulk_test_function, bulk_test_action);

int main()
{
    /* std::size_t arg0, arg1, ... */

    bulk_action act;
    std::size_t images_per_locality = 4;

    // Instantiate the parallel section
    hpx::lcos::define_spmd_block(
        "some_name", images_per_locality, std::move(act) /*, arg0, arg1, ... */);

    return 0;
}

```

Note: In principle, the user should never call the `spmd_block` constructor. The `define_spmd_block` function is responsible of instantiating `spmd_block` objects and broadcasting them to each created image.

SPMD multidimensional views

Some classes are defined as “container views” when the purpose is to observe and/or modify the values of a container using another perspective than the one that characterizes the container. For example, the values of an `std::vector` object can be accessed via the expression `[i]`. Container views can be used, for example, when it is desired for those values to be “viewed” as a 2D matrix that would have been flattened in a `std::vector`. The values would be possibly accessible via the expression `vv(i, j)` which would call internally the expression `v[k]`.

By default, the `partitioned_vector` class integrates 1-D views of its segments:

```
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(double);

using iterator = hpx::partitioned_vector<double>::iterator;
using traits   = hpx::traits::segmented_iterator_traits<iterator>;

hpx::partitioned_vector<double> v;

// Create a 1-D view of the vector of segments
auto vv = traits::segment(v.begin());

// Access segment i
std::vector<double> v = vv[i];
```

Our views are called “multidimensional” in the sense that they generalize to N dimensions the purpose of `segmented_iterator_traits::segment()` in the 1-D case. Note that in a parallel section, the 2-D expression `a(i, j) = b(i, j)` is quite confusing because without convention, each of the images invoked will race to execute the statement. For this reason, our views are not only multidimensional but also “spmd-aware”.

Note: SPMD-awareness: The convention is simple. If an assignment statement contains a view subscript as an l-value, it is only and only the image holding the r-value who is evaluating the statement. (In MPI sense, it is called a Put operation).

Subscript-based operations

Here are some examples of using subscripts in the 2-D view case:

```
#include <hpx/components/containers/partitioned_vector/partitioned_vector_view.hpp>
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(double);

using Vec = hpx::partitioned_vector<double>;
using View_2D = hpx::partitioned_vector_view<double, 2>;

/* Do some code */
```

(continues on next page)

(continued from previous page)

```

Vec v;

// Parallel section (suppose 'block' an spmd_block instance)
{
    std::size_t height, width;

    // Instantiate the view
    View_2D vv(block, v.begin(), v.end(), {height, width});

    // The l-value is a view subscript, the image that owns vv(1,0)
    // evaluates the assignment.
    vv(0,1) = vv(1,0);

    // The l-value is a view subscript, the image that owns the r-value
    // (result of expression 'std::vector<double>(4,1.0)' evaluates the
    // assignment : oops! race between all participating images.
    vv(2,3) = std::vector<double>(4,1.0);
}

```

Iterator-based operations

Here are some examples of using iterators in the 3-D view case:

```

#include <hpx/components/containers/partitioned_vector/partitioned_vector_view.hpp>
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
// 
HPC_REGISTER_PARTITIONED_VECTOR(int);

using Vec = hpx::partitioned_vector<int>;
using View_3D = hpx::partitioned_vector_view<int,3>;

/* Do some code */

Vec v1, v2;

// Parallel section (suppose 'block' an spmd_block instance)
{
    std::size_t size_x, size_y, size_z;

    // Instantiate the views
    View_3D vv1(block, v1.begin(), v1.end(), {size_x, size_y, size_z});
    View_3D vv2(block, v2.begin(), v2.end(), {size_x, size_y, size_z});

    // Save previous segments covered by vv1 into segments covered by vv2
    auto vv2_it = vv2.begin();
    auto vv1_it = vv1.cbegin();

    for(; vv2_it != vv2.end(); vv2_it++, vv1_it++)
    {
        // It's a Put operation
        *vv2_it = *vv1_it;
    }
}

```

(continues on next page)

(continued from previous page)

```
// Ensure that all images have performed their Put operations
block.sync_all();

// Ensure that only one image is putting updated data into the different
// segments covered by vvl
if(block.this_image() == 0)
{
    int idx = 0;

    // Update all the segments covered by vvl
    for(auto i = vvl.begin(); i != vvl.end(); i++)
    {
        // It's a Put operation
        *i = std::vector<float>(elt_size, idx++);
    }
}
}
```

Here is an example that shows how to iterate only over segments owned by the current image:

```
#include <hpx/components/containers/partitioned_vector/partitioned_vector_view.hpp>
#include <hpx/components/containers/partitioned_vector/partitioned_vector_local_view.
↪hpp>
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(float);

using Vec = hpx::partitioned_vector<float>;
using View_1D = hpx::partitioned_vector_view<float,1>;

/* Do some code */

Vec v;

// Parallel section (suppose 'block' an spmd_block instance)
{
    std::size_t num_segments;

    // Instantiate the view
    View_1D vv(block, v.begin(), v.end(), {num_segments});

    // Instantiate the local view from the view
    auto local_vv = hpx::local_view(vv);

    for (auto i = local_vv.begin(); i != local_vv.end(); i++)
    {
        std::vector<float> & segment = *i;

        /* Do some code */
    }
}
```

Instantiating sub-views

It is possible to construct views from other views: we call it sub-views. The constraint nevertheless for the subviews is to retain the dimension and the value type of the input view. Here is an example showing how to create a sub-view:

```
#include <hpx/components/containers/partitioned_vector/partitioned_vector_view.hpp>
#include <hpx/include/partitioned_vector.hpp>

// The following code generates all necessary boiler plate to enable the
// remote creation of 'partitioned_vector' segments
//
HPX_REGISTER_PARTITIONED_VECTOR(float);

using Vec = hpx::partitioned_vector<float>;
using View_2D = hpx::partitioned_vector_view<float,2>;

/* Do some code */

Vec v;

// Parallel section (suppose 'block' an spmd_block instance)
{
    std::size_t N = 20;
    std::size_t tilesize = 5;

    // Instantiate the view
    View_2D vv(block, v.begin(), v.end(), {N,N});

    // Instantiate the subview
    View_2D svv(
        block,&vv(tilesize,0),&vv(2*tilesize-1,tilesize-1),{tilesize,tilesize},{N,N});

    if(block.this_image() == 0)
    {
        // Equivalent to 'vv(tilesize,0) = 2.0f'
        svv(0,0) = 2.0f;

        // Equivalent to 'vv(2*tilesize-1,tilesize-1) = 3.0f'
        svv(tilesize-1,tilesize-1) = 3.0f;
    }
}
```

Note: The last parameter of the subview constructor is the size of the original view. If one would like to create a subview of the subview and so on, this parameter should stay unchanged. `{N, N}` for the above example).

C++ co-arrays

Fortran has extended its scalar element indexing approach to reference each segment of a distributed array. In this extension, a segment is attributed a ?co-index? and lives in a specific *locality*. A co-index provides the application with enough information to retrieve the corresponding data reference. In C++, containers present themselves as a ?smarter? alternative of Fortran arrays but there are still no corresponding standardized features similar to the Fortran co-indexing approach. We present here an implementation of such features in *HPX*.

Preface: co-array, a segmented container tied to a SPMD multidimensional views

As mentioned before, a co-array is a distributed array whose segments are accessible through an array-inspired access mode. We have previously seen that it is possible to reproduce such access mode using the concept of views. Nevertheless, the user must pre-create a segmented container to instantiate this view. We illustrate below how a single constructor call can perform those two operations:

```
#include <hpx/components/containers/coarray/coarray.hpp>
#include <hpx/collectives/spmd_block.hpp>

// The following code generates all necessary boiler plate to enable the
// co-creation of 'coarray'
//
HPX_REGISTER_COARRAY(double);

// Parallel section (suppose 'block' an spmd_block instance)
{
    using hpx::container::placeholders::_;

    std::size_t height=32, width=4, segment_size=10;

    hpx::coarray<double,3> a(block, "a", {height,width,_}, segment_size);

    /* Do some code */
}
```

Unlike segmented containers, a co-array object can only be instantiated within a parallel section. Here is the description of the parameters to provide to the coarray constructor:

Table 2.27: Parameters of coarray constructor

Parameter	Description
block	Reference to a <code>spmd_block</code> object
"a"	Symbolic name of type <code>std::string</code>
{height,width, _}	Dimensions of the coarray object
segment_size	Size of a co-indexed element (i.e. size of the object referenced by the expression <code>a(i,j,k)</code>)

Note that the “last dimension size” cannot be set by the user. It only accepts the `constexpr` variable `hpx::container::placeholders::_`. This size, which is considered private, is equal to the number of current images (value returned by `block.get_num_images()`).

Note: An important constraint to remember about coarray objects is that all segments sharing the same “last dimension index” are located in the same image.

Using co-arrays

The member functions owned by the `coarray` objects are exactly the same as those of spmd multidimensional views. These are:

- * Subscript-based operations
- * Iterator-based operations

However, one additional functionality is provided. Knowing that the element $a(i, j, k)$ is in the memory of the k th image, the use of local subscripts is possible.

Note: For spmd multidimensional views, subscripts are only global as it still involves potential remote data transfers.

Here is an example of using local subscripts:

```
#include <hpx/components/containers/coarray/coarray.hpp>
#include <hpx/collectives/spmd_block.hpp>

// The following code generates all necessary boiler plate to enable the
// co-creation of 'coarray'
//
HPX_REGISTER_COARRAY(double);

// Parallel section (suppose 'block' an spmd_block instance)
{
    using hpx::container::placeholders::_;

    std::size_t height=32, width=4, segment_size=10;

    hpx::coarray<double,3> a(block, "a", {height,width,_}, segment_size);

    double idx = block.this_image()*height*width;

    for (std::size_t j = 0; j<width; j++)
        for (std::size_t i = 0; i<height; i++)
    {
        // Local write operation performed via the use of local subscript
        a(i,j,_) = std::vector<double>(elt_size, idx);
        idx++;
    }

    block.sync_all();
}
```

Note: When the “last dimension index” of a subscript is equal to `hpx::container::placeholders::_`, local subscript (and not global subscript) is used. It is equivalent to a global subscript used with a “last dimension index” equal to the value returned by `block.this_image()`.

2.5.8 Running on batch systems

This section walks you through launching *HPX* applications on various batch systems.

How to use *HPX* applications with PBS

Most *HPX* applications are executed on parallel computers. These platforms typically provide integrated job management services that facilitate the allocation of computing resources for each parallel program. *HPX* includes support for one of the most common job management systems, the Portable Batch System (PBS).

All PBS jobs require a script to specify the resource requirements and other parameters associated with a parallel job. The PBS script is basically a shell script with PBS directives placed within commented sections at the beginning of the file. The remaining (not commented-out) portions of the file executes just like any other regular shell script. While the description of all available PBS options is outside the scope of this tutorial (the interested reader may refer to in-depth documentation²⁰⁶ for more information), below is a minimal example to illustrate the approach. The following test application will use the multithreaded `hello_world_distributed` program, explained in the section *Remote execution with actions: Hello world*.

```
#!/bin/bash
#
#PBS -l nodes=2:ppn=4

APP_PATH=~/packages/hpx/bin/hello_world_distributed
APP_OPTIONS=

pbsdsh -u $APP_PATH $APP_OPTIONS --hpx:nodes=`cat $PBS_NODEFILE`
```

Caution: If the first application specific argument (inside `$APP_OPTIONS`) is a non-option (i.e., does not start with a `-` or a `--`), then the argument has to be placed before the option `--hpx:nodes`, which, in this case, should be the last option on the command line.

Alternatively, use the option `--hpx:endnodes` to explicitly mark the end of the list of node names:

```
pbsdsh -u $APP_PATH --hpx:nodes`cat $PBS_NODEFILE` --hpx:endnodes $APP_OPTIONS
```

The `#PBS -l nodes=2:ppn=4` directive will cause two compute nodes to be allocated for the application, as specified in the option `nodes`. Each of the nodes will dedicate four cores to the program, as per the option `ppn`, short for “processors per node” (PBS does not distinguish between processors and cores). Note that requesting more cores per node than physically available is pointless and may prevent PBS from accepting the script.

On newer PBS versions the PBS command syntax might be different. For instance, the PBS script above would look like:

```
#!/bin/bash
#
#PBS -l select=2:ncpus=4

APP_PATH=~/packages/hpx/bin/hello_world_distributed
APP_OPTIONS=

pbsdsh -u $APP_PATH $APP_OPTIONS --hpx:nodes=`cat $PBS_NODEFILE`
```

²⁰⁶ <http://www.clusterresources.com/torquedocs21/>

`APP_PATH` and `APP_OPTIONS` are shell variables that respectively specify the correct path to the executable (`hello_world_distributed` in this case) and the command line options. Since the `hello_world_distributed` application doesn't need any command line options, `APP_OPTIONS` has been left empty. Unlike in other execution environments, there is no need to use the `--hpx:threads` option to indicate the required number of OS threads per node; the *HPX* library will derive this parameter automatically from PBS.

Finally, `pbsdsh` is a PBS command that starts tasks to the resources allocated to the current job. It is recommended to leave this line as shown and modify only the PBS options and shell variables as needed for a specific application.

Important: A script invoked by `pbsdsh` starts in a very basic environment: the user's `$HOME` directory is defined and is the current directory, the `LANG` variable is set to `C` and the `PATH` is set to the basic `/usr/local/bin:/usr/bin:/bin` as defined in a system-wide file `pbs_environment`. Nothing that would normally be set up by a system shell profile or user shell profile is defined, unlike the environment for the main job script.

Another choice is for the `pbsdsh` command in your main job script to invoke your program via a shell, like `sh` or `bash`, so that it gives an initialized environment for each instance. Users can create a small script `runme.sh`, which is used to invoke the program:

```
#!/bin/bash
# Small script which invokes the program based on what was passed on its
# command line.
#
# This script is executed by the bash shell which will initialize all
# environment variables as usual.
$@
```

Now, the script is invoked using the `pbsdsh` tool:

```
#!/bin/bash
#
#PBS -l nodes=2:ppn=4

APP_PATH=~/packages/hpx/bin/hello_world_distributed
APP_OPTIONS=

pbsdsh -u runme.sh $APP_PATH $APP_OPTIONS --hpx:nodes=`cat $PBS_NODEFILE`
```

All that remains now is submitting the job to the queuing system. Assuming that the contents of the PBS script were saved in the file `pbs_hello_world.sh` in the current directory, this is accomplished by typing:

```
qsub ./pbs_hello_world_pbs.sh
```

If the job is accepted, `qsub` will print out the assigned job ID, which may look like:

```
$ 42.supercomputer.some.university.edu
```

To check the status of your job, issue the following command:

```
qstat 42.supercomputer.some.university.edu
```

and look for a single-letter job status symbol. The common cases include:

- *Q* - signifies that the job is queued and awaiting its turn to be executed.
- *R* - indicates that the job is currently running.
- *C* - means that the job has completed.

The example qstat output below shows a job waiting for execution resources to become available:

Job id	Name	User	Time	Use	S	Queue
42.supercomputer	...ello_world.sh	joe_user		0	Q	batch

After the job completes, PBS will place two files, pbs_hello_world.sh.o42 and pbs_hello_world.sh.e42, in the directory where the job was submitted. The first contains the standard output and the second contains the standard error from all the nodes on which the application executed. In our example, the error output file should be empty and the standard output file should contain something similar to:

```
hello world from OS-thread 3 on locality 0
hello world from OS-thread 2 on locality 0
hello world from OS-thread 1 on locality 1
hello world from OS-thread 0 on locality 0
hello world from OS-thread 3 on locality 1
hello world from OS-thread 2 on locality 1
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 1
```

Congratulations! You have just run your first distributed *HPX* application!

How to use *HPX* applications with SLURM

Just like PBS (described in section [How to use *HPX* applications with PBS](#)), SLURM is a job management system which is widely used on large supercomputing systems. Any *HPX* application can easily be run using SLURM. This section describes how this can be done.

The easiest way to run an *HPX* application using SLURM is to utilize the command line tool srun, which interacts with the SLURM batch scheduling system:

```
srun -p <partition> -N <number-of-nodes> hpx-application <application-arguments>
```

Here, <partition> is one of the node partitions existing on the target machine (consult the machine's documentation to get a list of existing partitions) and <number-of-nodes> is the number of compute nodes that should be used. By default, the *HPX* application is started with one *locality* per node and uses all available cores on a node. You can change the number of localities started per node (for example, to account for NUMA effects) by specifying the -n option of srun. The number of cores per *locality* can be set by -c. The <application-arguments> are any application specific arguments that need to be passed on to the application.

Note: There is no need to use any of the *HPX* command line options related to the number of localities, number of threads, or related to networking ports. All of this information is automatically extracted from the SLURM environment by the *HPX* startup code.

Important: The srun documentation explicitly states: "If -c is specified without -n, as many tasks will be allocated per node as possible while satisfying the -c restriction. For instance on a cluster with 8 CPUs per node, a job request for 4 nodes and 3 CPUs per task may be allocated 3 or 6 CPUs per node (1 or 2 tasks per node) depending upon resource consumption by other jobs." For this reason, it's recommended to always specify -n <number-of-instances>, even if <number-of-instances> is equal to one (1).

Interactive shells

To get an interactive development shell on one of the nodes, users can issue the following command:

```
srun -p <node-type> -N <number-of-nodes> --pty /bin/bash -l
```

After the shell has been opened, users can run their *HPX* application. By default, it uses all available cores. Note that if you requested one node, you don't need to do `srun` again. However, if you requested more than one node, and want to run your distributed application, you can use `srun` again to start up the distributed *HPX* application. It will use the resources that have been requested for the interactive shell.

Scheduling batch jobs

The above mentioned method of running *HPX* applications is fine for development purposes. The disadvantage that comes with `srun` is that it only returns once the application is finished. This might not be appropriate for longer-running applications (for example, benchmarks or larger scale simulations). In order to cope with that limitation, users can use the `sbatch` command.

The `sbatch` command expects a script that it can run once the requested resources are available. In order to request resources, users need to add `#SBATCH` comments in their script or provide the necessary parameters to `sbatch` directly. The parameters are the same as with `run`. The commands you need to execute are the same you would need to start your application as if you were in an interactive shell.

2.5.9 Debugging *HPX* applications

Using a debugger with *HPX* applications

Using a debugger such as `gdb` with *HPX* applications is no problem. However, there are some things to keep in mind to make the experience somewhat more productive.

Call stacks in *HPX* can often be quite unwieldy as the library is heavily templated and the call stacks can be very deep. For this reason it is sometimes a good idea compile *HPX* in `RelWithDebInfo` mode, which applies some optimizations but keeps debugging symbols. This can often compress call stacks significantly. On the other hand, stepping through the code can also be more difficult because of statements being reordered and variables being optimized away. Also, note that because *HPX* implements user-space threads and context switching, call stacks may not always be complete in a debugger.

HPX launches not only worker threads but also a few helper threads. The first thread is the main thread, which typically does no work in an *HPX* application, except at startup and shutdown. If using the default settings, *HPX* will spawn six additional threads (used for service thread pools). The first worker thread is usually the eighth thread, and most user codes will be run on these worker threads. The last thread is a helper thread used for *HPX* shutdown.

Finally, since *HPX* is a multi-threaded runtime, the following `gdb` options can be helpful:

```
set pagination off
set non-stop on
```

Non-stop mode allows users to have a single thread stop on a breakpoint without stopping all other threads as well.

Using sanitizers with *HPX* applications

Warning: Not all parts of *HPX* are sanitizer clean. This means that users may end up with false positives from *HPX* itself when using sanitizers for their applications.

To use sanitizers with *HPX*, turn on `HPX_WITH_SANITIZERS` and turn off `HPX_WITH_STACK_OVERFLOW_DETECTION` during CMake configuration. It's recommended to also build Boost with the same sanitizers that will be used for *HPX*. The appropriate sanitizers can then be enabled using CMake by appending `-fsanitize=address -fno-omit-frame-pointer` to `CMAKE_CXX_FLAGS` and `-fsanitize=address` to `CMAKE_EXE_LINKER_FLAGS`. Replace `address` with the sanitizer that you want to use.

2.5.10 Optimizing *HPX* applications

Performance counters

Performance Counters in *HPX* are used to provide information as to how well the runtime system or an application is performing. The counter data can help determine system bottlenecks and fine-tune system and application performance. The *HPX* runtime system, its networking, and other layers provide counter data that an application can consume to provide users with information of how well the application is performing.

Applications can also use counter data to determine how much system resources to consume. For example, an application that transfers data over the network could consume counter data from a network switch to determine how much data to transfer without competing for network bandwidth with other network traffic. The application could use the counter data to adjust its transfer rate as the bandwidth usage from other network traffic increases or decreases.

Performance Counters are *HPX* parallel processes which expose a predefined interface. *HPX* exposes special API functions that allow one to create, manage, read the counter data, and release instances of Performance Counters. Performance Counter instances are accessed by name, and these names have a predefined structure which is described in the section *Performance counter names*. The advantage of this is that any Performance Counter can be accessed remotely (from a different *locality*) or locally (from the same *locality*). Moreover, since all counters expose their data using the same API, any code consuming counter data can be utilized to access arbitrary system information with minimal effort.

Counter data may be accessed in real time. More information about how to consume counter data can be found in the section *Consuming performance counter data*.

All *HPX* applications provide command line options related to performance counters, such as the ability to list available counter types, or periodically query specific counters to be printed to the screen or save them in a file. For more information, please refer to the section *HPX Command Line Options*.

Performance counter names

All Performance Counter instances have a name uniquely identifying this instance. This name can be used to access the counter, retrieve all related meta data, and to query the counter data (as described in the section *Consuming performance counter data*). Counter names are strings with a predefined structure. The general form of a countername is:

```
/objectname{full_instancename}/countername@parameters
```

where `full_instancename` could be either another (full) counter name or a string formatted as:

```
parentinstancename#parentindex/instancename#instanceindex
```

Each separate part of a countername (e.g. `objectname`, `countername` `parentinstancename`, `instancename`, and `parameters`) should start with a letter ('a'...'z', 'A'...'Z') or an underscore character ('_'), optionally followed by letters, digits ('0'...'9'), hyphen ('-'), or underscore characters. Whitespace is not allowed inside a counter name. The characters '/', '{', '}', '#', and '@' have a special meaning and are used to delimit the different parts of the counter name.

The parts `parentinstanceindex` and `instanceindex` are integers. If an index is not specified *HPX* will assume a default of -1.

Two simple examples

An instance for a well formed (and meaningful) simple counter name would be:

```
/threads{locality#0/total}/count/cumulative
```

This counter returns the current cumulative number of executed (retired) *HPX*-threads for the *locality* 0. The counter type of this counter is `/threads/count/cumulative` and the full instance name is `locality#0/total`. This counter type does not require an `instanceindex` or `parameters` to be specified.

In this case, the `parentindex` (the '0') designates the *locality* for which the counter instance is created. The counter will return the number of *HPX*-threads retired on that particular *locality*.

Another example for a well formed (aggregate) counter name is:

```
/statistics{/threads{locality#0/total}/count/cumulative}/average@500
```

This counter takes the simple counter from the first example, samples its values every 500 milliseconds, and returns the average of the value samples whenever it is queried. The counter type of this counter is `/statistics/average` and the instance name is the full name of the counter for which the values have to be averaged. In this case, the `parameters` (the '500') specify the sampling interval for the averaging to take place (in milliseconds).

Performance counter types

Every Performance Counter belongs to a specific Performance Counter type which classifies the counters into groups of common semantics. The type of a counter is identified by the `objectname` and the `countername` parts of the name.

```
/objectname/countername
```

When an application starts *HPX* will register all available counter types on each of the localities. These counter types are held in a special Performance Counter registration database which can be later used to retrieve the meta data related to a counter type and to create counter instances based on a given counter instance name.

Performance counter instances

The `full_instancename` distinguishes different counter instances of the same counter type. The formatting of the `full_instancename` depends on the counter type. There are two types of counters: simple counters which usually generate the counter values based on direct measurements, and aggregate counters which take another counter and transform its values before generating their own counter values. An example for a simple counter is given [above](#): counting retired HPX-threads. An aggregate counter is shown as an example [above](#) as well: calculating the average of the underlying counter values sampled at constant time intervals.

While simple counters use instance names formatted as `parentinstancename#parentindex/instancename#instanceindex`, most aggregate counters have the full counter name of the embedded counter as its instance name.

Not all simple counter types require specifying all 4 elements of a full counter instance name, some of the parts (`parentinstancename`, `parentindex`, `instancename`, and `instanceindex`) are optional for specific counters. Please refer to the documentation of a particular counter for more information about the formatting requirements for the name of this counter (see [Existing HPX performance counters](#)).

The parameters are used to pass additional information to a counter at creation time. They are optional and they fully depend on the concrete counter. Even if a specific counter type allows additional parameters to be given, those usually are not required as sensible defaults will be chosen. Please refer to the documentation of a particular counter for more information about what parameters are supported, how to specify them, and what default values are assumed (see also [Existing HPX performance counters](#)).

Every `locality` of an application exposes its own set of Performance Counter types and Performance Counter instances. The set of exposed counters is determined dynamically at application start based on the execution environment of the application. For instance, this set is influenced by the current hardware environment for the `locality` (such as whether the `locality` has access to accelerators), and the software environment of the application (such as the number of OS-threads used to execute HPX-threads).

Using wildcards in performance counter names

It is possible to use wildcard characters when specifying performance counter names. Performance counter names can contain 2 types of wildcard characters:

- Wildcard characters in the performance counter type
- Wildcard characters in the performance counter instance name

Wildcard character have a meaning which is very close to usual file name wildcard matching rules implemented by common shells (like bash).

Table 2.28: Wildcard characters in the performance counter type

Wild-card	Description
*	This wildcard character matches any number (zero or more) of arbitrary characters.
?	This wildcard character matches any single arbitrary character.
[...]	This wildcard character matches any single character from the list of specified within the square brackets.

Table 2.29: Wildcard characters in the performance counter instance name

Wildcard	Description
*	This wildcard character matches any <i>locality</i> or any thread, depending on whether it is used for <code>locality#*</code> or <code>worker-thread#*</code> . No other wildcards are allowed in counter instance names.

Consuming performance counter data

You can consume performance data using either the command line interface or via the *HPX* application or the *HPX* API. The command line interface is easier to use, but it is less flexible and does not allow one to adjust the behaviour of your application at runtime. The command line interface provides a convenience abstraction but simplified abstraction for querying and logging performance counter data for a set of performance counters.

Consuming performance counter data from the command line

HPX provides a set of predefined command line options for every application which uses `hpx::init` for its initialization. While there are much more command line options available (see [HPX Command Line Options](#)), the set of options related to Performance Counters allow one to list existing counters, query existing counters once at application termination or repeatedly after a constant time interval.

The following table summarizes the available command line options:

Table 2.30: *HPX* Command Line Options Related to Performance Counters

Command line option	Description
<code>--hpx:print</code>	print the specified performance counter either repeatedly and/or at the times specified by <code>--hpx:print-counter-at</code> (see also option <code>--hpx:print-counter-interval</code>).
<code>--hpx:print</code>	print the specified performance counter either repeatedly and/or at the times specified by <code>--hpx:print-counter-at</code> reset the counter after the value is queried. (see also option <code>--hpx:print-counter-interval</code>).
<code>--hpx:print</code>	print the performance counter(s) specified with <code>--hpx:print-counter</code> repeatedly after the time interval (specified in milliseconds) (default:0 which means print once at shutdown).
<code>--hpx:print</code>	print the performance counter(s) specified with <code>--hpx:print-counter</code> to the given file (default: console)).
<code>--hpx:list</code>	list the names of all registered performance counters.
<code>--hpx:list</code>	list the description of all registered performance counters.
<code>--hpx:print</code>	print the performance counter(s) specified with <code>--hpx:print-counter</code> possible formats in csv format with header or without any header (see option <code>--hpx:no-csv-header</code>), possible values: csv (prints counter values in CSV format with full names as header) csv-short (prints counter values in CSV format with shortnames provided with <code>--hpx:print-counter</code> as <code>--hpx:print-counter shortname,full-countername</code>)
<code>--hpx:no</code>	print the performance counter(s) specified with <code>--hpx:print-counter</code> and csv or csv-short format specified with <code>--hpx:print-counter-format</code> without header.
<code>--hpx:print</code> arg	print the performance counter(s) specified with <code>--hpx:print-counter</code> (or <code>--hpx:print-counter-reset</code>) at the given point in time, possible argument values: startup, shutdown (default), nosutdown.
<code>--hpx:reset</code>	reset the performance counter(s) specified with <code>--hpx:print-counter</code> after they have been evaluated)

While the options `--hpx:list-counters` and `--hpx:list-counter-infos` give a short listing of all available counters, the full documentation for those can be found in the section [Existing HPX performance counters](#).

A simple example

All of the commandline options mentioned above can be for instance tested using the `hello_world_distributed` example.

Listing all available counters `hello_world_distributed --hpx:list-counters` yields:

```
List of available counter instances (replace * below with the appropriate sequence number)
-----
/agas/count/allocate /agas/count/bind /agas/count/bind_gid
/agas/count/bind_name ... /threads{locality#*/allocator#*/}count/objects
/threads{locality#*/total}/count/stack-recycles
/threads{locality#*/total}/idle-rate
/threads{locality#*/worker-thread#*/}idle-rate
```

Providing more information about all available counters `hello_world_distributed --hpx:list-counter-infos` yields:

```
Information about available counter instances (replace * below with the appropriate sequence number)
-----
fullname: /agas/count/allocate helptext: returns the number of invocations of the AGAS service 'allocate' type: counter_raw version: 1.0.0
-----

fullname: /agas/count/bind helptext: returns the number of invocations of the AGAS service 'bind' type: counter_raw version: 1.0.0
-----

fullname: /agas/count/bind_gid helptext: returns the number of invocations of the AGAS service 'bind_gid' type: counter_raw version: 1.0.0
-----

...
```

This command will not only list the counter names but also a short description of the data exposed by this counter.

Note: The list of available counters may differ depending on the concrete execution environment (hardware or software) of your application.

Requesting the counter data for one or more performance counters can be achieved by invoking `hello_world_distributed` with a list of counter names:

```
hello_world_distributed \
  --hpx:print-counter=/threads{locality#0/total}/count/cumulative \
  --hpx:print-counter=/agas{locality#0/total}/count/bind
```

which yields for instance:

```
hello world from OS-thread 0 on locality 0
/threads{locality#0/total}/count/cumulative,1,0.212527,[s],33
/agas{locality#0/total}/count/bind,1,0.212790,[s],11
```

The first line is the normal output generated by `hello_world_distributed` and has no relation to the counter data listed. The last two lines contain the counter data as gathered at application shutdown. These lines have 6 fields, the counter name, the sequence number of the counter invocation, the time stamp at which this information has been sampled, the unit of measure for the time stamp, the actual counter value, and an optional unit of measure for the counter value.

The actual counter value can be represented by a single number (for counters returning singular values) or a list of numbers separated by ' : ' (for counters returning an array of values, like for instance a histogram).

Note: The name of the performance counter will be enclosed in double quotes " " if it contains one or more commas ', '.

Requesting to query the counter data once after a constant time interval with this command line:

```
hello_world_distributed \
  --hpx:print-counter=/threads{locality#0/total}/count/cumulative \
  --hpx:print-counter=/agas{locality#0/total}/count/bind \
  --hpx:print-counter-interval=20
```

yields for instance (leaving off the actual console output of the `hello_world_distributed` example for brevity):

```
threads{locality#0/total}/count/cumulative,1,0.002409,[s],22
agas{locality#0/total}/count/bind,1,0.002542,[s],9
threads{locality#0/total}/count/cumulative,2,0.023002,[s],41
agas{locality#0/total}/count/bind,2,0.023557,[s],10
threads{locality#0/total}/count/cumulative,3,0.037514,[s],46
agas{locality#0/total}/count/bind,3,0.038679,[s],10
```

The command `--hpx:print-counter-destination=<file>` will redirect all counter data gathered to the specified file name, which avoids cluttering the console output of your application.

The command line option `--hpx:print-counter` supports using a limited set of wildcards for a (very limited) set of use cases. In particular, all occurrences of #* as in `locality#*` and in `worker-thread#*` will be automatically expanded to the proper set of performance counter names representing the actual environment for the executed program. For instance, if your program is utilizing 4 worker threads for the execution of HPX threads (see command line option `--hpx:threads`) the following command line

```
hello_world_distributed \
  --hpx:threads=4 \
  --hpx:print-counter=/threads{locality#0/worker-thread#*}/count/cumulative
```

will print the value of the performance counters monitoring each of the worker threads:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
hello world from OS-thread 3 on locality 0
hello world from OS-thread 2 on locality 0
/threads{locality#0/worker-thread#0}/count/cumulative,1,0.0025214,[s],27
/threads{locality#0/worker-thread#1}/count/cumulative,1,0.0025453,[s],33
/threads{locality#0/worker-thread#2}/count/cumulative,1,0.0025683,[s],29
/threads{locality#0/worker-thread#3}/count/cumulative,1,0.0025904,[s],33
```

The command `--hpx:print-counter-format` takes values `csv` and `csv-short` to generate CSV formatted counter values with header.

With format as csv:

```
hello_world_distributed \
--hpx:threads=2 \
--hpx:print-counter-format csv \
--hpx:print-counter /threads{locality#/*/total}/count/cumulative \
--hpx:print-counter /threads{locality#/*/total}/count/cumulative-phases
```

will print the values of performance counters in CSV format with full countername as header:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
/threads{locality#/*/total}/count/cumulative,/threads{locality#/*/total}/count/
cumulative-phases
39,93
```

With format csv-short:

```
hello_world_distributed \
--hpx:threads 2 \
--hpx:print-counter-format csv-short \
--hpx:print-counter cumulative,/threads{locality#/*/total}/count/cumulative \
--hpx:print-counter phases,/threads{locality#/*/total}/count/cumulative-phases
```

will print the values of performance counters in CSV format with short countername as header:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
cumulative,phases
39,93
```

With format csv and csv-short when used with `--hpx:print-counter-interval`:

```
hello_world_distributed \
--hpx:threads 2 \
--hpx:print-counter-format csv-short \
--hpx:print-counter cumulative,/threads{locality#/*/total}/count/cumulative \
--hpx:print-counter phases,/threads{locality#/*/total}/count/cumulative-phases \
--hpx:print-counter-interval 5
```

will print the header only once repeating the performance counter value(s) repeatedly:

```
cum,phases
25,42
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
44,95
```

The command `--hpx:no-csv-header` to be used with `--hpx:print-counter-format` to print performance counter values in CSV format without any header:

```
hello_world_distributed \
--hpx:threads 2 \
--hpx:print-counter-format csv-short \
--hpx:print-counter cumulative,/threads{locality#/*/total}/count/cumulative \

```

(continues on next page)

(continued from previous page)

```
--hpx:print-counter phases,/threads{locality#/*/total}/count/cumulative-phases \
--hpx:no-csv-header
```

will print:

```
hello world from OS-thread 1 on locality 0
hello world from OS-thread 0 on locality 0
37,91
```

Consuming performance counter data using the HPX API

HPX provides an API allowing to discover performance counters and to retrieve the current value of any existing performance counter from any application.

Discover existing performance counters

Retrieve the current value of any performance counter

Performance counters are specialized HPX components. In order to retrieve a counter value, the performance counter needs to be instantiated. HPX exposes a client component object for this purpose:

```
hpx::performance_counters::performance_counter counter(std::string const& name);
```

Instantiating an instance of this type will create the performance counter identified by the given name. Only the first invocation for any given counter name will create a new instance of that counter, all following invocations for a given counter name will reference the initially created instance. This ensures, that at any point in time there is always not more than one active instance of any of the existing performance counters.

In order to access the counter value (or invoking any of the other functionality related to a performance counter, like start, stop or reset) member functions of the created client component instance should be called:

```
// print the current number of threads created on locality 0
hpx::performance_counters::performance_counter count(
    "/threads{locality#0/total}/count/cumulative");
hpx::cout << count.get_value<int>().get() << hpx::endl;
```

For more information about the client component type see [classref hpx::performance_counters::performance_counter].

Note: In the above example `count.get_value()` returns a future. In order to print the result we must append `.get()` to retrieve the value. You could write the above example like this for more clarity:

```
// print the current number of threads created on locality 0
hpx::performance_counters::performance_counter count(
    "/threads{locality#0/total}/count/cumulative");
hpx::future<int> result = count.get_value<int>();
hpx::cout << result.get() << hpx::endl;
```

Providing performance counter data

HPX offers several ways by which you may provide your own data as a performance counter. This has the benefit of exposing additional, possibly application specific information using the existing Performance Counter framework, unifying the process of gathering data about your application.

An application that wants to provide counter data can implement a Performance Counter to provide the data. When a consumer queries performance data, the HPX runtime system calls the provider to collect the data. The runtime system uses an internal registry to determine which provider to call.

Generally, there are two ways of exposing your own Performance Counter data: a simple, function based way and a more complex, but more powerful way of implementing a full Performance Counter. Both alternatives are described in the following sections.

Exposing performance counter data using a simple function

The simplest way to expose arbitrary numeric data is to write a function which will then be called whenever a consumer queries this counter. Currently, this type of Performance Counter can only be used to expose integer values. The expected signature of this function is:

```
std::int64_t some_performance_data(bool reset);
```

The argument `bool reset` (which is supplied by the runtime system when the function is invoked) specifies whether the counter value should be reset after evaluating the current value (if applicable).

For instance, here is such a function returning how often it was invoked:

```
// The atomic variable 'counter' ensures the thread safety of the counter.
boost::atomic<std::int64_t> counter(0);

std::int64_t some_performance_data(bool reset)
{
    std::int64_t result = ++counter;
    if (reset)
        counter = 0;
    return result;
}
```

This example function exposes a linearly increasing value as our performance data. The value is incremented on each invocation, e.g. each time a consumer requests the counter data of this Performance Counter.

The next step in exposing this counter to the runtime system is to register the function as a new raw counter type using the HPX API function `hpx::performance_counters::install_counter_type`. A counter type represents certain common characteristics of counters, like their counter type name, and any associated description information. The following snippet shows an example of how to register the function `some_performance_data` which is shown above for a counter type named "/test/data". This registration has to be executed before any consumer instantiates and queries an instance of this counter type:

```
#include <hpx/include/performance_counters.hpp>

void register_counter_type()
{
    // Call the HPX API function to register the counter type.
    hpx::performance_counters::install_counter_type(
        "/test/data",                                         // counter type name
        &some_performance_data,                            // function providing counter_
        &data);
```

(continues on next page)

(continued from previous page)

```

    "returns a linearly increasing counter value" // description text (optional)
    ""
);
}

```

Now it is possible to instantiate a new counter instance based on the naming scheme "/test{locality#*/total}/data" where * is a zero based integer index identifying the *locality* for which the counter instance should be accessed. The function `hpx::performance_counters::install_counter_type` enables to instantiate exactly one counter instance for each *locality*. Repeated requests to instantiate such a counter will return the same instance, e.g. the instance created for the first request.

If this counter needs to be accessed using the standard *HPX* command line options, the registration has to be performed during application startup, before `hpx_main` is executed. The best way to achieve this is to register an *HPX* startup function using the API function `hpx::register_startup_function` before calling `hpx::init` to initialize the runtime system:

```

int main(int argc, char* argv[])
{
    // By registering the counter type we make it available to any consumer
    // who creates and queries an instance of the type "/test/data".
    //
    // This registration should be performed during startup. The
    // function 'register_counter_type' should be executed as an HPX thread right
    // before hpx_main is executed.
    hpx::register_startup_function(&register_counter_type);

    // Initialize and run HPX.
    return hpx::init(argc, argv);
}

```

Please see the code in [hpx_link examples/performance_counters/simplest_performance_counter.cpp..simplest_performance_counter.cpp] for a full example demonstrating this functionality.

Implementing a full performance counter

Sometimes, the simple way of exposing a single value as a Performance Counter is not sufficient. For that reason, *HPX* provides a means of implementing full Performance Counters which support:

- Retrieving the descriptive information about the Performance Counter
- Retrieving the current counter value
- Resetting the Performance Counter (value)
- Starting the Performance Counter
- Stopping the Performance Counter
- Setting the (initial) value of the Performance Counter

Every full Performance Counter will implement a predefined interface:

```

// Copyright (c) 2007-2018 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

```

(continues on next page)

(continued from previous page)

```
#if !defined(HPX_PERFORMANCE_COUNTERS_PERFORMANCE_COUNTER_JAN_18_2013_0939AM)
#define HPX_PERFORMANCE_COUNTERS_PERFORMANCE_COUNTER_JAN_18_2013_0939AM

#include <hpx/config.hpp>
#include <hpx/lcos/future.hpp>
#include <hpx/runtime/components/client_base.hpp>
#include <hpx/runtime/launch_policy.hpp>
#include <hpx/functional/bind_front.hpp>

#include <hpx/performance_counters/counters_fwd.hpp>
#include <hpx/performance_counters/stubs/performance_counter.hpp>

#include <string>
#include <utility>
#include <vector>

///////////////////////////////
namespace hpx { namespace performance_counters
{
    ///////////////////////////
    struct HPX_EXPORT performance_counter
        : components::client_base<performance_counter, stubs::performance_counter>
    {
        typedef components::client_base<
            performance_counter, stubs::performance_counter
        > base_type;

        performance_counter() {}

        performance_counter(std::string const& name);

        performance_counter(std::string const& name, hpx::id_type const& locality);

        performance_counter(future<id_type> && id)
            : base_type(std::move(id))
        {}

        performance_counter(hpx::future<performance_counter> && c)
            : base_type(std::move(c))
        {}

        ///////////////////////////
        future<counter_info> get_info() const;
        counter_info get_info(launch::sync_policy,
            error_code& ec = throws) const;

        future<counter_value> get_counter_value(bool reset = false);
        counter_value get_counter_value(launch::sync_policy,
            bool reset = false, error_code& ec = throws);

        future<counter_value> get_counter_value() const;
        counter_value get_counter_value(launch::sync_policy,
            error_code& ec = throws) const;

        future<counter_values_array> get_counter_values_array(bool reset = false);
        counter_values_array get_counter_values_array(launch::sync_policy,
            error_code& ec = throws);
    };
}}
```

(continues on next page)

(continued from previous page)

```

bool reset = false, error_code& ec = throws);

future<counter_values_array> get_counter_values_array() const;
counter_values_array get_counter_values_array(launch::sync_policy,
                                             error_code& ec = throws) const;

///////////////////////////////
future<bool> start();
bool start(launch::sync_policy, error_code& ec = throws);

future<bool> stop();
bool stop(launch::sync_policy, error_code& ec = throws);

future<void> reset();
void reset(launch::sync_policy, error_code& ec = throws);

future<void> reinit(bool reset = true);
void reinit(
    launch::sync_policy, bool reset = true, error_code& ec = throws);

/////////////////////////////
future<std::string> get_name() const;
std::string get_name(launch::sync_policy, error_code& ec = throws) const;

private:
    template <typename T>
    static T extract_value(future<counter_value> && value)
    {
        return value.get().get_value<T>();
    }

public:
    template <typename T>
    future<T> get_value(bool reset = false)
    {
        return get_counter_value(reset).then(
            hpx::launch::sync,
            util::bind_front(
                &performance_counter::extract_value<T>));
    }
    template <typename T>
    T get_value(launch::sync_policy, bool reset = false,
                error_code& ec = throws)
    {
        return get_counter_value(launch::sync, reset).get_value<T>(ec);
    }

    template <typename T>
    future<T> get_value() const
    {
        return get_counter_value().then(
            hpx::launch::sync,
            util::bind_front(
                &performance_counter::extract_value<T>));
    }
    template <typename T>
    T get_value(launch::sync_policy, error_code& ec = throws) const

```

(continues on next page)

(continued from previous page)

```

    {
        return get_counter_value(launch::sync).get_value<T>(ec);
    }
};

/// Return all counters matching the given name (with optional wildcards).
HPX_API_EXPORT std::vector<performance_counter> discover_counters(
    std::string const& name, error_code& ec = throws);
}

#endif

```

In order to implement a full Performance Counter you have to create an *HPX* component exposing this interface. To simplify this task, *HPX* provides a ready made base class which handles all the boiler plate of creating a component for you. The remainder of this section will explain the process of creating a full Performance Counter based on the Sine example which you can find in the directory examples/performance_counters/sine/.

The base class is defined in the header file [hpx_link hpx/performance_counters/base_performance_counter.hpp..hpx/performance_counters/base_performance_counter.hpp] as:

```

// Copyright (c) 2007-2018 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

#if !defined(HPX_PERFORMANCE_COUNTERS_BASE_PERFORMANCE_COUNTER_JAN_18_2013_1036AM)
#define HPX_PERFORMANCE_COUNTERS_BASE_PERFORMANCE_COUNTER_JAN_18_2013_1036AM

#include <hpx/config.hpp>
#include <hpx/performance_counters/counters.hpp>
#include <hpx/performance_counters/server/base_performance_counter.hpp>
#include <hpx/runtime/actions/component_action.hpp>
#include <hpx/runtime/components/component_type.hpp>
#include <hpx/runtime/components/server/component_base.hpp>

///////////////////////////////
//[performance_counter_base_class
namespace hpx { namespace performance_counters
{
    template <typename Derived>
    class base_performance_counter;
}}}
///

/////////////////////////////
namespace hpx { namespace performance_counters
{
    template <typename Derived>
    class base_performance_counter
        : public hpx::performance_counters::server::base_performance_counter,
        public hpx::components::component_base<Derived>
    {
        private:
            typedef hpx::components::component_base<Derived> base_type;
    public:

```

(continues on next page)

(continued from previous page)

```

typedef Derived type_holder;
typedef hpx::performance_counters::server::base_performance_counter
    base_type_holder;

base_performance_counter()
{ }

base_performance_counter(hpx::performance_counters::counter_info const& info)
    : base_type_holder(info)
{ }

// Disambiguate finalize() which is implemented in both base classes
void finalize()
{
    base_type_holder::finalize();
    base_type::finalize();
}
};

#endif

```

The single template parameter is expected to receive the type of the derived class implementing the Performance Counter. In the Sine example this looks like:

```

// Copyright (c) 2007-2012 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

#if !defined(PERFORMANCE_COUNTERS_SINE_SEP_20_2011_0112PM)
#define PERFORMANCE_COUNTERS_SINE_SEP_20_2011_0112PM

#include <hpx/hpx.hpp>
#include <hpx/util/interval_timer.hpp>
#include <hpx/synchronization/spinlock.hpp>
#include <hpx/performance_counters/base_performance_counter.hpp>

#include <cstdint>

namespace performance_counters { namespace sine { namespace server
{
    /////////////////////////////////
    //[sine_counter_definition
    class sine_counter
        : public hpx::performance_counters::base_performance_counter<sine_counter>
    [/]
    {
        public:
            sine_counter() : current_value_(0), evaluated_at_(0) {}
            sine_counter(hpx::performance_counters::counter_info const& info);

            // This function will be called in order to query the current value of
            // this performance counter
            hpx::performance_counters::counter_value get_counter_value(bool reset);
    };
}}
```

(continues on next page)

(continued from previous page)

```

/// The functions below will be called to start and stop collecting
/// counter values from this counter.
bool start();
bool stop();

/// finalize() will be called just before the instance gets destructed
void finalize();

protected:
    bool evaluate();

private:
    typedef hpx::lcos::local::spinlock mutex_type;

    mutable mutex_type mtx_;
    double current_value_;
    std::uint64_t evaluated_at_;

    hpx::util::interval_timer timer_;
};

}
}

#endif

```

i.e. the type `sine_counter` is derived from the base class passing the type as a template argument (please see [hpx_link examples/performance_counters/sine/server/sine.hpp..sine.hpp] for the full source code of the counter definition). For more information about this technique (called Curiously Recurring Template Pattern - CRTP), please see for instance the corresponding Wikipedia article²⁰⁷. This base class itself is derived from the `performance_counter` interface described above.

Additionally, a full Performance Counter implementation not only exposes the actual value but also provides information about

- The point in time a particular value was retrieved
- A (sequential) invocation count
- The actual counter value
- An optional scaling coefficient
- Information about the counter status

Existing HPX performance counters

The `HPX` runtime system exposes a wide variety of predefined Performance Counters. These counters expose critical information about different modules of the runtime system. They can help determine system bottlenecks and fine-tune system and application performance.

²⁰⁷ http://en.wikipedia.org/wiki/Curiously_recurring_template_pattern

Table 2.31: AGAS performance counters

Counter type	Counter instance formatting	Description	Parameters
/agas/count/ <agas_service> where: <agas_service> is one of the following: <i>primary namespace services:</i> route, bind_gid, resolve_gid, unbind_gid, increment_credit, decrement_credit, allocate, begin_migration, end_migration <i>component namespace services:</i> bind_prefix, bind_name, resolve_id, unbind_name, iterate_types, get_component_type <i>locality namespace services:</i> free, localities, num_localities, num_threads, resolve_locality, resolved_localities <i>symbol namespace services:</i> bind, resolve, unbind, iterate_names, on_symbol_namespace_event	<agas_instance>/ total where: <agas_instance> is the name of the AGAS service to query. Currently, this value will be locality#0 where 0 is the root <i>locality</i> (the id of the locality hosting the AGAS service). The value for * can be any <i>locality</i> id for the follow- ing <agas_service>: route, bind_gid, resolve_gid, unbind_gid, increment_credit, decrement_credit, bin, resolve, unbind, and iterate_names only the primary and num_localities_type symbol AGAS service components live on all lo- calities, whereas all other AGAS services are avail- able on locality#0 only).	None	Returns the total number of invocations of the spec- ified AGAS service since its creation.
/agas/ <agas_service_category>/ count where: <agas_service_category> is one of the following: primary , locality, component or symbol	<agas_instance>/ total where: <agas_instance> is service to query. Currently, this value will be locality#0 where 0 is the root <i>locality</i> (the id of the <i>locality</i> hosting the AGAS service). Except for <agas_service_category>, primary or symbol for which the value for * can be any <i>locality</i> id (only the primary and symbol	None	Returns the overall total number of invocations of all AGAS services pro- vided by the given AGAS service category since its creation.
2.5. Manual	AGAS service compo- nents live on all localities, whereas all other AGAS services are available on		199

Table 2.32: Parcel layer performance counters

Counter type	Counter instance formatting	Description	Parameters
/data/count/ <connection_type> <operation> where: <operation> is one of the following: sent, received <connection_type> is one of the following: tcp, mpi	locality#*total where: * is the <i>locality</i> id of the <i>locality</i> the overall number of transmitted bytes should be queried for. The <i>locality</i> id is a (zero based) number iden- tifying the <i>locality</i> .	Returns the overall number of raw (un- compressed) bytes sent or received (see <operation>, e.g. en or eceived) for the specified <connection_type>. The performance counters for the connection type mpi are available only if the compile time constant HPX_HAVE_PARCELPORT_MPI was defined while compiling the HPX core li- brary (which is not defined by default, the corresponding cmake configuration constant is HPX_WITH_PARCELPORT_MPI). Please see <i>CMake variables used to configure HPX</i> for more details.	None
/data/time/ <connection_type> <operation> where: <operation> is one of the following: sent, received <connection_type> is one of the following: tcp, mpi	locality#*total where: * is the <i>locality</i> id of the <i>locality</i> the total transmission time should be queried for. The <i>locality</i> id is a (zero based) number iden- tifying the <i>locality</i> .	Returns the total time (in nanoseconds) between the start of each asynchronous transmission op- eration and the end of the corresponding op- eration for the specified <connection_type> the given <i>locality</i> (see <operation>, e.g. en or eceived). The performance counters for the connection type mpi are available only if the compile time constant HPX_HAVE_PARCELPORT_MPI was defined while compiling the HPX core li- brary (which is not defined by default, the corresponding cmake configuration constant is HPX_WITH_PARCELPORT_MPI). Please see <i>CMake variables used to configure HPX</i> for more details.	None
/serialize/ count/ <connection_type> <operation> where: <operation> is one of the following: sent, received <connection_type> is one of the following: tcp, mpi	locality#*total where: * is the <i>locality</i> id of the <i>locality</i> the overall number of transmitted bytes should be queried for. The <i>locality</i> id is a (zero based) number iden- tifying the <i>locality</i> .	Returns the overall number of bytes trans- ferred (see <operation>, e.g. sent or received possibly compressed) for the speci- fied <connection_type> by the given <i>local- ity</i> . The performance counters for the connection type mpi are available only if the compile time constant HPX_HAVE_PARCELPORT_MPI was defined while compiling the HPX core li- brary (which is not defined by default, the corresponding cmake configuration constant is HPX_WITH_PARCELPORT_MPI). Please see <i>CMake variables used to configure HPX</i> for more details.	If the config- time option -DHPX_WITH_PARCELPORT_ACTION was specified, this counter allows one to specify an optional action name as its parameter. In this case the counter will report the number of bytes transmitted for the given action only.
/serialize/ time/ <connection_type> <operation> where:	locality#*total where: * is the <i>locality</i> id of the	Returns the overall time spent performing outgoing data serialization for the specified <connection_type> on the given <i>locality</i> (see <operation>, e.g. sent or received). The performance counters for the connection	If the config- time option -DHPX_WITH_PARCELPORT_ACTION was specified, this counter allows one

Table 2.33: Thread manager performance counters

Counter type	Counter instance formatting	Description	Parameters
/threads/count/cumulative	<p>locality#*/total or locality#*/worker-thread#* or locality#*/pool#*/worker-thread#* where: locality#* is defining the <i>locality</i> for which the overall number of retired HPX-threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>. pool#* is defining the pool for which the current value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the overall number of retired HPX-threads should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the 'default' pool.</p>	<p>Returns the overall number of executed (retired) HPX-threads on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the accumulated number of retired HPX-threads for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the overall number of retired HPX-threads for all worker threads separately which the current value of the available only if the configuration time constant <code>HPX_WITH_THREAD_CUMULATIVE_COUNTS</code></p>	None

Continued on next page

²⁰⁸ A message can potentially consist of more than one *parcel*.

Table 2.33 – continued from previous page

/threads/time/ average	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average time spent executing one HPX-thread should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the average time spent executing one HPX-thread should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the average time spent executing one HPX-thread on the given <i>locality</i> since application start. If the instance name is total the counter returns the average time spent executing one HPX-thread for all worker threads (cores) on that <i>locality</i>. If the instance name is worker-thread#* the counter will return the average time spent executing one HPX-thread for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_WITH_THREAD_CUMULATIVE_COUNTS</code> (default: ON) and <code>HPX_WITH_THREAD_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].</p>	None
---------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ average-overhead	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average overhead spent executing one <i>HPX</i>-thread should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the average overhead spent executing one <i>HPX</i>-thread should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the average time spent on overhead while executing one <i>HPX</i>-thread on the given <i>locality</i> since application start. If the instance name is total the counter returns the average time spent on overhead while executing one <i>HPX</i>-thread for all worker threads (cores) on that <i>locality</i>. If the instance name is worker-thread#* the counter will return the average time spent on overhead executing one <i>HPX</i>-thread for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_WITH_THREAD_CUMULATIVE_COUNTS</code> (default: ON) and <code>HPX_WITH_THREAD_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].</p>	None
------------------------------------	---	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ cumulative-phases	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the overall number of executed HPX-thread phases (invocations) should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the overall number of executed HPX-thread phases (invocations) should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the overall number of executed HPX-thread phases (invocations) on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the accumulated number of executed HPX-thread phases (invocations) for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the overall number of executed HPX-thread phases for all worker threads separately. This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_CUMULATIVE_COUNTS</code> is set to ON (default: ON). The unit of measure for this counter is nanosecond [ns].</p>	None
--------------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ average-phase	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average time spent executing one HPX-thread phase (invocation) should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the average time executing one HPX-thread phase (invocation) should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the average time spent executing one HPX-thread phase (invocation) on the given <i>locality</i> since application start. If the instance name is total the counter returns the average time spent executing one HPX-thread phase (invocation) for all worker threads (cores) on that <i>locality</i>. If the instance name is worker-thread#* the counter will return the average time spent executing one HPX-thread phase for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_WITH_THREAD_CUMULATIVE_COUNTS</code> (default: ON) and <code>HPX_WITH_THREAD_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].</p>	None
---------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ average-phase-overhead	locality#*/total locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#* where: locality#* is defining the <i>locality</i> for which the average time overhead executing one HPX-thread phase (invocation) should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i> . pool#* is defining the pool for which the current value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the average overhead executing one HPX-thread phase (invocation) should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> . If no pool-name is specified the counter refers to the ‘default’ pool.	Returns the average time spent on overhead executing one HPX-thread phase (invocation) on the given <i>locality</i> since application start. If the instance name is total the counter returns the average time spent on overhead while executing one HPX-thread phase (invocation) for all worker threads (cores) on that <i>locality</i> . If the instance name is worker-thread#* the counter will return the average time spent on overhead executing one HPX-thread phase for all worker threads separately. This counter is available only if the configuration time constants HPX_WITH_THREAD_CUMULATIVE_COUNTS (default: ON) and HPX_WITH_THREAD_IDLE_RATES are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].	None
--	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ overall	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the overall time spent running the scheduler should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the overall time spent running the scheduler should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the overall time spent running the scheduler on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the overall time spent running the scheduler for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the overall time spent running the scheduler for all worker threads separately. This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_IDLE_RATES</code> is set to <code>ON</code> (default: <code>OFF</code>). The unit of measure for this counter is nanosecond [ns].</p>	None
---------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ cumulative	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the overall time spent executing all HPX-threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the overall time spent executing all HPX-threads should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the overall time spent executing all HPX-threads on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the overall time spent executing all HPX-threads for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the overall time spent executing all HPX-threads for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_THREAD_MAINTAIN_CUMULATIVE_COUNTS</code> (default: ON) and <code>HPX_THREAD_MAINTAIN_IDLE_RATES</code> are set to ON (default: OFF).</p>	None
------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/ cumulative-overhead	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the overall overhead time incurred by executing all <i>HPX</i>-threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the the overall overhead time incurred by executing all <i>HPX</i>-threads should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the overall overhead time incurred executing all <i>HPX</i>-threads on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the overall overhead time incurred executing all <i>HPX</i>-threads for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the overall overhead time incurred executing all <i>HPX</i>-threads for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_THREAD_MAINTAIN_CUMULATIVE_COUNTS</code> (default: ON) and <code>HPX_THREAD_MAINTAIN_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].</p>	None
---------------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

<p>threads/count/ instantaneous/ <thread-state> where: <thread-state> is one of the following: all, active, pending, suspended, terminated, staged</p>	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#* where: locality#* is defining the <i>locality</i> for which the current number of threads with the given state should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>. pool#* is defining the pool for which the current value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the current number of threads with the given state should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool. The staged thread state refers to registered tasks before they are converted to thread objects.</p>	<p>Returns the current number of HPX-threads having the given thread state on the given <i>locality</i>. If the instance name is total the counter returns the current number of HPX-threads of the given state for all worker threads (cores) on that <i>locality</i>. If the instance name is worker-thread#* the counter will return the current number of HPX-threads in the given state for all worker threads separately.</p>	<p>None</p>
--	---	---	-------------

Continued on next page

Table 2.33 – continued from previous page

<p>threads/ wait-time/ <thread-state> where: <thread-state> is one of the following: pending staged</p>	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#* where: locality#* is defining the <i>locality</i> for which the average wait time of <i>HPX</i>-threads (pending) or thread descriptions (staged) with the given state should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>. pool#* is defining the pool for which the cur- rent value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the average wait time for the given state should be queried for. The worker thread number (given by the * is a (zero based) num- ber identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘de- fault’ pool. The staged thread state refers to the wait time of registered tasks be- fore they are converted into thread objects, while the pending thread state refers to the wait time of threads in any of the scheduling queues.</p>	<p>Returns the average wait time of <i>HPX</i>-threads (if the thread state is pending or of task descriptions (if the thread state is staged on the given <i>locality</i> since application start. If the instance name is <i>total</i> the counter returns the wait time of <i>HPX</i>-threads of the given state for all worker threads (cores) on that <i>locality</i>. If the instance name is <i>worker-thread#*</i> the counter will return the wait time of <i>HPX</i>-threads in the given state for all worker threads separately. These counters are available only if the compile time constant <code>HPX_WITH_THREAD_QUEUE_WAITTIME</code> was defined while com- piling the <i>HPX</i> core library (default: OFF). The unit of measure for this counter is nanosecond [ns].</p>	<p>None</p>
---	--	--	-------------

Continued on next page

Table 2.33 – continued from previous page

/threads/ idle-rate	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average idle rate of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i></p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the averaged idle rate should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the average idle rate for the given worker thread(s) on the given <i>locality</i> . The idle rate is defined as the ratio of the time spent on scheduling and management tasks and the overall time spent executing work since the application started. This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_IDLE_RATES</code> is set to ON (default: OFF).	None
------------------------	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ creation-idle-rate	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average creation idle rate of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the averaged idle rate should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the average idle rate for the given worker thread(s) on the given <i>locality</i> which is caused by creating new threads. The creation idle rate is defined as the ratio of the time spent on creating new threads and the overall time spent executing work since the application started. This counter is available only if the configuration time constants <code>HPX_WITH_THREAD_IDLE_RATES</code> (default: OFF) and <code>HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATES</code> are set to ON.</p>	None
---------------------------------	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ cleanup-idle-rate	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the average cleanup idle rate of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the averaged cleanup idle rate should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	<p>Returns the average idle rate for the given worker thread(s) on the given <i>locality</i> which is caused by cleaning up terminated threads. The cleanup idle rate is defined as the ratio of the time spent on cleaning up terminated thread objects and the overall time spent executing work since the application started. This counter is available only if the configuration time constants <code>HPX_WITH_THREAD_IDLE_RATES</code> (default: OFF) and <code>HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATES</code> are set to ON.</p>	None
--------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threadqueue/ length	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the current length of all thread queues in the scheduler for all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the current length of all thread queues in the scheduler should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the overall length of all queues for the given worker thread(s) on the given <i>locality</i> .	None
/threads/count/ stack-unbinds	<p>locality#*/total where:</p> <p>* is the <i>locality</i> id of the <i>locality</i> the unbind (madvise) operations should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i>.</p>	Returns the total number of HPX-thread unbind (madvise) operations performed for the referenced <i>locality</i> . Note that this counter is not available on Windows based platforms.	None

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ stack-recycles	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the recycling operations should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the total number of HPX-thread recycling operations performed.	None
/threads/count/ stolen-from-pending	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the number of ‘stole’ threads should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the total number of HPX-threads ‘stolen’ from the pending thread queue by a neighboring thread worker thread (these threads are executed by a different worker thread than they were initially scheduled on). This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None

Continued on next page

Table 2.33 – continued from previous page

/threads/count/pending-misses	<p>locality#*/total or locality#*/worker-thread#* or locality#*/pool#*/worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the number of pending queue misses of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i></p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the number of pending queue misses should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the total number of times that the referenced worker-thread on the referenced <i>locality</i> failed to find pending HPX-threads in its associated queue. This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None
-------------------------------	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/pending-accesses	<p>locality#*/total or locality#*/worker-thread#* or locality#*/pool#*/worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the number of pending queue accesses of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i></p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the number of pending queue accesses should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the total number of times that the referenced worker-thread on the referenced <i>locality</i> looked for pending HPX-threads in its associated queue. This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None
---------------------------------	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ stolen-from-staged	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the number of HPX-threads stolen from the staged queue of all (or one) worker threads should be queried for. The <i>locality</i> id (given by *) is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the number of HPX-threads stolen from the staged queue should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the total number of HPX-threads ‘stolen’ from the staged thread queue by a neighboring worker thread (these threads are executed by a different worker thread than they were initially scheduled on). This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None
---------------------------------------	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ stolen-to-pending	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the number of HPX-threads stolen to the pending queue of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the number of HPX-threads stolen to the pending queue should be queried for. The worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the total number of HPX-threads ‘stolen’ to the pending thread queue of the worker thread (these threads are executed by a different worker thread than they were initially scheduled on). This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None
--------------------------------------	--	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ stolen-to-staged	<p>locality#*/total or locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the <i>locality</i> for which the number of HPX-threads stolen to the staged queue of all (or one) worker threads should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i>.</p> <p>pool#* is defining the pool for which the current value of the idle-loop counter should be queried for.</p> <p>worker-thread#* is defining the worker thread for which the number of HPX-threads stolen to the staged queue should be queried for. The worker thread number (given by the * is a (zero based) worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>. If no pool-name is specified the counter refers to the ‘default’ pool.</p>	Returns the total number of HPX-threads ‘stolen’ to the staged thread queue of a neighboring worker thread (these threads are executed by a different worker thread than they were initially scheduled on). This counter is available only if the configuration time constant <code>HPX_WITH_THREAD_STEALING_COUNTS</code> is set to ON (default: ON).	None
-------------------------------------	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/count/ objects	locality#*/total or locality#*/ allocator#* where: locality#* is defining the <i>locality</i> for which the current (cumulative) num- ber of all created <i>HPX</i> - thread objects should be queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i> . allocator#* is defin- ing the number of the allo- cator instance using which the threads have been cre- ated. <i>HPX</i> uses a vary- ing number of allocators to create (and recycle) <i>HPX</i> -thread objects, most likely these counters are of use for debugging pur- poses only. The allocator id (given by * is a (zero based) number identifying the allocator to query.	Returns the total num- ber of <i>HPX</i> -thread ob- jects created. Note that thread objects are reused to improve system perfor- mance, thus this number does not reflect the num- ber of actually executed (retired) <i>HPX</i> -threads.	None
/scheduler/ utilization/ instantaneous	locality#*/total where: locality#* is defining the <i>locality</i> for which the current (instantaneous) scheduler utilization queried for. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i> .	Returns the total (instantaneous) scheduler utilization. This is the current percentage of scheduler threads executing <i>HPX</i> threads.	Percent

Continued on next page

Table 2.33 – continued from previous page

/threads/ idle-loop-count/ instantaneous	locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#* where: locality#* is defining the <i>locality</i> for which the current current accumulated value of all idle-loop counters of all worker threads should be queried. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i> . pool#* is defining the pool for which the current value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the current value of the idle-loop counter should be queried for. The worker thread number (given by the * is a (zero based) worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> . If no pool-name is specified the counter refers to the ‘default’ pool.	Returns the current (instantaneous) idle-loop count for the given HPX-worker thread or the accumulated value for all worker threads.	None
--	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ busy-loop-count/ instantaneous	locality#*/ worker-thread#* or locality#*/ pool#*/ worker-thread#* where: locality#* is defining the <i>locality</i> for which the current current accumulated value of all busy-loop counters of all worker threads should be queried. The <i>locality</i> id (given by * is a (zero based) number identifying the <i>locality</i> . pool#* is defining the pool for which the current value of the idle-loop counter should be queried for. worker-thread#* is defining the worker thread for which the current value of the busy-loop counter should be queried for. The worker thread number (given by the * is a (zero based) worker thread number (given by the * is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> . If no pool-name is specified the counter refers to the ‘default’ pool.	Returns the current (instantaneous) busy-loop count for the given HPX-worker thread or the accumulated value for all worker threads.	None
--	---	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/background-work-duration	locality#*/total locality#*/ worker-thread#* where: locality#* is defining the locality for which the overall time spent performing background work should be queried for. The locality id (given by *) is a (zero based) number identifying the locality. worker-thread#* is defining the worker thread for which the overall time spent performing background work should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> .	Returns the overall time spent performing background work on the given locality since application start. If the instance name is total the counter returns the overall time spent performing background work for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return the overall time spent performing background work for all worker threads separately. This counter is available only if the configuration time constants HPX_WITH_BACKGROUND_THREAD_COUNTERS (default: OFF) and HPX_WITH_THREAD_IDLE_RATES are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns].	None
--	--	--	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ background-overhead	locality#*/total or locality#*/ worker-thread#* where: locality#* is defining the locality for which the background overhead should be queried for. The locality id (given by *) is a (zero based) number identifying the locality. worker-thread#* is defining the worker thread for which the background overhead should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option --hpx:threads .	Returns the background overhead on the given locality since application start. If the instance name is total the counter returns the background overhead for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return background overhead for all worker threads separately. This counter is available only if the configuration time constants HPX_WITH_BACKGROUND_THREAD_COUNTERS (default: OFF) and HPX_WITH_THREAD_IDLE_RATES are set to ON (default: OFF). The unit of measure displayed for this counter is 0.1%.	None
----------------------------------	---	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/background-send-duration	locality#*/total locality#*/ worker-thread#* where: locality#* is defining the locality for which the overall time spent performing background work related to sending parcels should be queried for. The locality id (given by *) is a (zero based) number identifying the locality. worker-thread#* is defining the worker thread for which the overall time spent performing background work related to sending parcels should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> .	Returns the overall time spent performing background work related to sending parcels on the given locality since application start. If the instance name is total the counter returns the overall time spent performing background work for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return the overall time spent performing background work for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_WITH_BACKGROUND_THREAD_COUNTERS</code> (default: OFF) and <code>HPX_WITH_THREAD_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns]. This counter will currently return meaningful values for the MPI parcel-port only.	None
--	--	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ background-send-overhead	locality#*/total locality#*/ worker-thread#* where: locality#* is defining the locality for which the background overhead related to sending parcels should be queried for. The locality id (given by *) is a (zero based) number identifying the locality. worker-thread#* is defining the worker thread for which the background overhead related to sending parcels should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code> .	Returns the background overhead related to sending parcels on the given locality since application start. If the instance name is total the counter returns the background overhead for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return background overhead for all worker threads separately. This counter is available only if the configuration time constants HPX_WITH_BACKGROUND_THREAD_COUNTERS (default: OFF) and HPX_WITH_THREAD_IDLE_RATES are set to ON (default: OFF). The unit of measure displayed for this counter is 0.1%. This counter will currently return meaningful values for the MPI parcel-port only.	None
---------------------------------------	--	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/time/background-receive	<p>locality#*/total duration</p> <p>locality#*/ worker-thread#*</p> <p>where:</p> <p>locality#* is defining the locality for which the overall time spent performing background work related to receiving parcels should be queried for. The locality id (given by *) is a (zero based) number identifying the locality.</p> <p>worker-thread#* is defining the worker thread for which the overall time spent performing background work related to receiving parcels should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <code>--hpx:threads</code>.</p>	<p>Returns the overall time spent performing background work related to receiving parcels on the given locality since application start. If the instance name is total the counter returns the overall time spent performing background work for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return the overall time spent performing background work for all worker threads separately. This counter is available only if the configuration time constants <code>HPX_WITH_BACKGROUND_THREAD_COUNTERS</code> (default: OFF) and <code>HPX_WITH_THREAD_IDLE_RATES</code> are set to ON (default: OFF). The unit of measure for this counter is nanosecond [ns]. This counter will currently return meaningful values for the MPI parcel-port only.</p>	None
----------------------------------	--	---	------

Continued on next page

Table 2.33 – continued from previous page

/threads/ background-receive	locality#*/total overhead locality#*/ worker-thread#* where: locality#* is defining the locality for which the background overhead re- lated to receiving should be queried for. The lo- cality id (given by *) is a (zero based) number iden- tifying the locality. worker-thread#* is defining the worker thread for which the background overhead related to receiving parcels should be queried for. The worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the ap- plication using the option <i>--hpx:threads</i> .	Returns the background overhead related to re- ceiving parcels on the given locality since ap- plication start. If the instance name is total the counter returns the background overhead for all worker threads (cores) on that locality. If the instance name is worker-thread#* the counter will return background overhead for all worker threads separately. This counter is available only if the configura- tion time constants HPX_WITH_BACKGROUND_THREAD_COUNTERS (default: OFF) and HPX_WITH_THREAD_IDLE_RATES are set to ON (default: OFF). The unit of mea- sure displayed for this counter is 0.1%. This counter will cur- rently return meaningful values for the MPI parcel- port only.	None
---------------------------------	--	---	------

Table 2.34: General performance counters exposing characteristics of localities

Counter type	Counter instance formatting	Description	Parameters
/runtime/count/component	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the number of components should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the overall number of currently active components of the specified type on the given <i>locality</i> .	The type of the component. This is the string which has been used while registering the component with HPX, e.g. which has been passed as the second parameter to the macro HPX_REGISTER_COMPONENT .
/runtime/count/action-invocation	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the number of action invocations should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the overall (local) invocation count of the specified action type on the given <i>locality</i> .	The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to the macro HPX_REGISTER_ACTION or HPX_REGISTER_ACTION_ID .
/runtime/count/remote-action-invocation	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the number of action invocations should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the overall (remote) invocation count of the specified action type on the given <i>locality</i> .	The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to the macro HPX_REGISTER_ACTION or HPX_REGISTER_ACTION_ID .
/runtime/uptime	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the system uptime should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the overall time since application start on the given <i>locality</i> in nanoseconds.	None
/runtime/memory/virtual	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the allocated virtual memory should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the amount of virtual memory currently allocated by the referenced <i>locality</i> (in bytes).	None
/runtime/memory/resident	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the allocated resident memory should be queried. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the amount of resident memory currently allocated by the referenced <i>locality</i> (in bytes).	None
2.5. Manual	locality#*/total where: * is the <i>locality</i> id of the <i>locality</i> the number identifying the <i>locality</i> .		231
/runtime/memory/total	locality#*/total where:	Returns the total available memory for use by the referenced <i>locality</i> (in bytes)	None

Table 2.35: Performance counters exposing PAPI hardware counters

Counter type	Counter instance formatting	Description	Pa-ram-e-ters
/papi/<papi_event> where: <papi_event> is the name of the PAPI event to expose as a performance counter (such as PAPI_SR_INS). Note that the list of available PAPI events changes depending on the used architecture. For a full list of available PAPI events and their (short) description use the --hpx:list-counters and --hpx:papi-event-info= command line options.	locality#*/total or locality#*/worker-thread#* where: locality#* is defining the <i>locality</i> for which the current accumulated value of all busy-loop counters of all worker threads should be queried. The <i>locality</i> id (given by *) is a (zero based) number identifying the <i>locality</i> . worker-thread#* is defining the worker thread for which the current value of the busy-loop counter should be queried for. The worker thread number (given by the *) is a (zero based) worker thread number (given by the *) is a (zero based) number identifying the worker thread. The number of available worker threads is usually specified on the command line for the application using the option <i>--hpx:threads</i> .	This counter returns the current count of occurrences of the specified PAPI event. This counter is available only if the configuration time constant HPX_WITH_PAPI is set to ON (default: OFF).	None

Table 2.36: Performance counters for general statistics

Counter type	Counter instance formatting	Description	Parameters
/	Any full performance counter average. The referenced performance counter is queried at fixed time intervals as specified by the first parameter.	Returns the current average (mean) value calculated based on the values queried from the underlying counter (the one specified as the instance name).	Any parameter will be interpreted as a list of up to two comma separated (integer) values, where the first is the time interval (in milliseconds) at which the underlying counter should be queried. If no value is specified, the counter will assume 1000 [ms] as the default. The second value can be either 0 or 1 and specifies whether the underlying counter should be reset during evaluation 1 or not 0. The default value is 0.
/	Any full performance counter rolling average. The referenced performance counter is queried at fixed time intervals as specified by the first parameter.	Returns the current rolling average (mean) value calculated based on the values queried from the underlying counter (the one specified as the instance name).	Any parameter will be interpreted as a list of up to three comma separated (integer) values, where the first is the time interval (in milliseconds) at which the underlying counter should be queried. If no value is specified, the counter will assume 1000 [ms] as the default. The second value will be interpreted as the size of the rolling window (the number of latest values to use to calculate the rolling average). The default value for this is 10. The third value can be either 0 or 1 and specifies whether the underlying counter should be reset during evaluation 1 or not 0. The default value is 0.
/	Any full performance counter standard deviation. The referenced performance counter is queried at fixed time intervals as specified by the first parameter.	Returns the current standard deviation (stddev) value calculated based on the values queried from the underlying counter (the one specified as the instance name).	Any parameter will be interpreted as a list of up to two comma separated (integer) values, where the first is the time interval (in milliseconds) at which the underlying counter should be queried. If no value is specified, the counter will assume 1000 [ms] as the default. The second value can be either 0 or 1 and specifies whether the underlying counter should be reset during evaluation 1 or not 0. The default value is 0.
/	Any full performance counter rolling standard deviation. The referenced performance counter is queried at fixed time intervals as specified by the first parameter.	Returns the current rolling variance (stddev) value calculated based on the values queried from the underlying counter (the one specified as the instance name).	Any parameter will be interpreted as a list of up to three comma separated (integer) values, where the first is the time interval (in milliseconds) at which the underlying counter should be queried. If no value is specified, the counter will assume 1000 [ms] as the default. The second value will be interpreted as the size of the rolling window (the number of latest values to use to calculate the rolling average). The default value for this is 10. The third value can be either 0 or 1 and specifies whether the underlying counter should be reset during evaluation 1 or not 0. The default value is 0.
/	Any full performance counter median. The referenced performance counter is queried at fixed time intervals as specified by the first parameter.	Returns the current (statistically estimated) median value calculated based on the values queried from the underlying counter (the one specified as the instance name).	Any parameter will be interpreted as a list of up to two comma separated (integer) values, where the first is the time interval (in milliseconds) at which the underlying counter should be queried. If no value is specified, the counter will assume 1000 [ms] as the default. The second value can be either 0 or 1 and specifies whether the underlying counter should be reset during evaluation 1 or not 0. The default value is 0.
2.5. Manual	intervals as specified by the first parameter.		

Table 2.37: Performance counters for elementary arithmetic operations

Counter type	Counter instance formatting	Description	Parameters
/arithmetics/add	None	Returns the sum calculated based on the values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/subtract	None	Returns the difference calculated based on the values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/multiply	None	Returns the product calculated based on the values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/divide	None	Returns the result of division of the values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/mean	None	Returns the average value of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/variance	None	Returns the standard deviation of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/median	None	Returns the median value of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/min	None	Returns the minimum value of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/max	None	Returns the maximum value of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.
/arithmetics/count	None	Returns the count value of all values queried from the underlying counters (the ones specified as the parameters).	The parameter will be interpreted as a comma separated list of full performance counter names which are queried whenever this counter is accessed. Any wildcards in the counter names will be expanded.

Note: The /arithmetics counters can consume an arbitrary number of other counters. For this reason those have to be specified as parameters (a comma separated list of counters appended after a '@'. For instance:

```
./bin/hello_world_distributed -t2 \
```

(continues on next page)

(continued from previous page)

```
--hpx:print-counter=/threads{locality#0/worker-thread#*}/count/cumulative \
--hpx:print-counter=/arithmetics/add@/threads{locality#0/worker-thread#*}/count/
↪cumulative
hello world from OS-thread 0 on locality 0
hello world from OS-thread 1 on locality 0
/threads{locality#0/worker-thread#0}/count/cumulative,1,0.515640,[s],25
/threads{locality#0/worker-thread#1}/count/cumulative,1,0.515520,[s],36
/arithmetics/add@/threads{locality#0/worker-thread#*}/count/cumulative,1,0.516445,[s],
↪64
```

Since all wildcards in the parameters are expanded, this example is fully equivalent to specifying both counters separately to `/arithmetics/add`:

```
./bin/hello_world_distributed -t2 \
    --hpx:print-counter=/threads{locality#0/worker-thread#*}/count/cumulative \
    --hpx:print-counter=/arithmetics/add@\
        /threads{locality#0/worker-thread#0}/count/cumulative, \
        /threads{locality#0/worker-thread#1}/count/cumulative
```

Table 2.38: Performance counters tracking parcel coalescing

Counter type	Counter instance formatting	Description	Parameters
/	locality#*/coalescesofad/ countwhere: parcelsis the <i>locality</i> id of the <i>locality</i> the number of parcels for the given action should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the number of parcels handled by the message handler associated with the action which is given by the counter parameter.	The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to the macro <code>HPX_REGISTER_ACTION</code> or <code>HPX_REGISTER_ACTION_ID</code> .
/	locality#*/coalescesofad/ countwhere: messagess is the <i>locality</i> id of the <i>locality</i> the number of messages for the given action should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the number of messages generated by the message handler associated with the action which is given by the counter parameter.	The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to the macro <code>HPX_REGISTER_ACTION</code> or <code>HPX_REGISTER_ACTION_ID</code> .
/	locality#*/coalescesofad/ countwhere: averageis the <i>locality</i> id of the <i>locality</i> the number of messages for the given action should be queried for. The <i>locality</i> id is a (zero based) number identifying the <i>locality</i> .	Returns the average number of parcels sent in a message generated by the message handler associated with the action which is given by the counter parameter.	The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to the macro <code>HPX_REGISTER_ACTION</code> or <code>HPX_REGISTER_ACTION_ID</code>
236	locality#*/coalescesofad/ time/where: averageis the <i>locality</i> arrival	Returns the average time between arriving parcels for the action which is given by the counter parameter.	Chapter 2. What's so special about HPX? The action type. This is the string which has been used while registering the action with HPX, e.g. which has been passed as the second parameter to

Note: The performance counters related to *parcel* coalescing are available only if the configuration time constant `HPX_WITH_PARCEL_COALESCING` is set to ON (default: ON). However, even in this case it will be available only for those actions, which are enabled for *parcel* coalescing (see the macros `HPX_ACTIONUSES_MESSAGE_COALESCING` and `HPX_ACTIONUSES_MESSAGE_COALESCING_NOTHROW`).

APEX integration

HPX provides integration with *APEX*²⁰⁹, which is a framework for application profiling using task timers and various performance counters. It can be added as a `git` submodule by turning on the option `HPX_WITH_APEX:BOOL` during `CMake`²¹⁰ configuration. *TAU*²¹¹ is an optional dependency when using *APEX*²¹².

To build *HPX* with *APEX*²¹³ add `HPX_WITH_APEX=ON`, and, optionally, `TAU_ROOT=$PATH_TO_TAU` to your `CMake`²¹⁴ configuration. In addition, you can override the tag used for *APEX*²¹⁵ with the `HPX_WITH_APEX_TAG` option. Please see the *APEX HPX documentation*²¹⁶ for detailed instructions on using *APEX*²¹⁷ with *HPX*.

2.5.11 *HPX* runtime and resources

HPX thread scheduling policies

The *HPX* runtime has five thread scheduling policies: local-priority, static-priority, local, static and abp-priority. These policies can be specified from the command line using the command line option `--hpx:queuing`. In order to use a particular scheduling policy, the runtime system must be built with the appropriate scheduler flag turned on (e.g. `cmake -DHPX_THREAD_SCHEDULERS=local`, see *CMake variables used to configure HPX* for more information).

Priority local scheduling policy (default policy)

- default or invoke using: `--hpx:queuinglocal-priority-fifo`

The priority local scheduling policy maintains one queue per operating system (OS) thread. The OS thread pulls its work from this queue. By default the number of high priority queues is equal to the number of OS threads; the number of high priority queues can be specified on the command line using `--hpx:high-priority-threads`. High priority threads are executed by any of the OS threads before any other work is executed. When a queue is empty work will be taken from high priority queues first. There is one low priority queue from which threads will be scheduled only when there is no other work.

For this scheduling policy there is an option to turn on NUMA sensitivity using the command line option `--hpx:numa-sensitive`. When NUMA sensitivity is turned on work stealing is done from queues associated with the same NUMA domain first, only after that work is stolen from other NUMA domains.

This scheduler is enabled at build time by default and will be available always.

²⁰⁹ <https://khuck.github.io/xpress-apex/>

²¹⁰ <https://www.cmake.org>

²¹¹ <https://www.cs.uoregon.edu/research/tau/home.php>

²¹² <https://khuck.github.io/xpress-apex/>

²¹³ <https://khuck.github.io/xpress-apex/>

²¹⁴ <https://www.cmake.org>

²¹⁵ <https://khuck.github.io/xpress-apex/>

²¹⁶ <https://khuck.github.io/xpress-apex/usage/#hpx-louisiana-state-university>

²¹⁷ <https://khuck.github.io/xpress-apex/>

This scheduler can be used with two underlying queuing policies (FIFO: first-in-first-out, and LIFO: last-in-first-out). The default is FIFO. In order to use the LIFO policy use the command line option `--hpx:queuing=local-priority-lifo`.

Static priority scheduling policy

- invoke using: `--hpx:queuing=static-priority` (or `-qs`)
- flag to turn on for build: `HPX_THREAD_SCHEDULERS=all` or `HPX_THREAD_SCHEDULERS=static-priority`

The static scheduling policy maintains one queue per OS thread from which each OS thread pulls its tasks (user threads). Threads are distributed in a round robin fashion. There is no thread stealing in this policy.

Local scheduling policy

- invoke using: `--hpx:queuing=local` (or `-ql`)
- flag to turn on for build: `HPX_THREAD_SCHEDULERS=all` or `HPX_THREAD_SCHEDULERS=local`

The local scheduling policy maintains one queue per OS thread from which each OS thread pulls its tasks (user threads).

Static scheduling policy

- invoke using: `--hpx:queuing=static`
- flag to turn on for build: `HPX_THREAD_SCHEDULERS=all` or `HPX_THREAD_SCHEDULERS=static`

The static scheduling policy maintains one queue per OS thread from which each OS thread pulls its tasks (user threads). Threads are distributed in a round robin fashion. There is no thread stealing in this policy.

Priority ABP scheduling policy

- invoke using: `--hpx:queuing=abp-priority-fifo`
- flag to turn on for build: `HPX_THREAD_SCHEDULERS=all` or `HPX_THREAD_SCHEDULERS=abp-priority`

Priority ABP policy maintains a double ended lock free queue for each OS thread. By default the number of high priority queues is equal to the number of OS threads; the number of high priority queues can be specified on the command line using `--hpx:high-priority-threads`. High priority threads are executed by the first OS threads before any other work is executed. When a queue is empty work will be taken from high priority queues first. There is one low priority queue from which threads will be scheduled only when there is no other work. For this scheduling policy there is an option to turn on NUMA sensitivity using the command line option `--hpx:numa-sensitive`. When NUMA sensitivity is turned on work stealing is done from queues associated with the same NUMA domain first, only after that work is stolen from other NUMA domains.

This scheduler can be used with two underlying queuing policies (FIFO: first-in-first-out, and LIFO: last-in-first-out). In order to use the LIFO policy use the command line option `--hpx:queuing=abp-priority-lifo`.

The HPX resource partitioner

The *HPX* resource partitioner lets you take the execution resources available on a system—processing units, cores, and numa domains—and assign them to thread pools. By default *HPX* creates a single thread pool name `default`. While this is good for most use cases, the resource partitioner lets you create multiple thread pools with custom resources and options.

Creating custom thread pools is useful for cases where you have tasks which absolutely need to run without interference from other tasks. An example of this is when using [MPI²¹⁸](#) for distribution instead of the built-in mechanisms in *HPX* (useful in legacy applications). In this case one can create a thread pool containing a single thread for [MPI²¹⁹](#) communication. [MPI²²⁰](#) tasks will then always run on the same thread, instead of potentially being stuck in a queue behind other threads.

Note that *HPX* thread pools are completely independent from each other in the sense that task stealing will never happen between different thread pools. However, tasks running on a particular thread pool can schedule tasks on another thread pool.

Note: It is simpler in some situations to schedule important tasks with high priority instead of using a separate thread pool.

Using the resource partitioner

In order to create custom thread pools the resource partitioner needs to be set up before *HPX* is initialized by creating an instance of `hpx::resource::partitioner`:

Note that we have to pass `argc` and `argv` to the resource partitioner to be able to parse thread binding options passed on the command line. You should pass the same arguments to the `hpx::resource::partitioner` constructor as you would to `hpx::init` or `hpx::start`. Running the above code will have the same effect as not initializing it at all, i.e. a default thread pool will be created with the type and number of threads specified on the command line.

The resource partitioner class is the interface to add thread pools to the *HPX* runtime and to assign resources to the thread pools.

To add a thread pool use the `hpx::resource::partitioner::create_thread_pool` method. If you simply want to use the default scheduler and scheduler options it is enough to call `rp.create_thread_pool("my-thread-pool")`.

Then, to add resources to the thread pool you can use the `hpx::resource::partitioner::add_resource` method. The resource partitioner exposes the hardware topology retrieved using [Portable Hardware Locality \(HWLOC\)](#)²²¹ and lets you iterate through the topology to add the wanted processing units to the thread pool. Below is an example of adding all processing units from the first NUMA domain to a custom thread pool, unless there is only one NUMA domain in which case we leave the first processing unit for the default thread pool:

Note: Whatever processing units not assigned to a thread pool by the time `hpx::init` is called will be added to the default thread pool. It is also possible to explicitly add processing units to the default thread pool, and to create the default thread pool manually (in order to e.g. set the scheduler type).

²¹⁸ https://en.wikipedia.org/wiki/Message_Passing_Interface

²¹⁹ https://en.wikipedia.org/wiki/Message_Passing_Interface

²²⁰ https://en.wikipedia.org/wiki/Message_Passing_Interface

²²¹ <https://www.open-mpi.org/projects/hwloc/>

Tip: The command line option `--hpx:print-bind` is useful for checking that the thread pools have been set up the way you expect.

Advanced usage

It is possible to customize the built in schedulers by passing scheduler options to `hpx::resource::partitioner::create_thread_pool`. It is also possible to create and use custom schedulers.

Note: It is not recommended to create your own scheduler. The *HPX* developers use this to experiment with new scheduler designs before making them available to users via the standard mechanisms of choosing a scheduler (command line options). If you would like to experiment with a custom scheduler the resource partitioner example `shared_priority_queue_scheduler.cpp` contains a fully implemented scheduler with logging etc. to make exploration easier.

To choose a scheduler and custom mode for a thread pool, pass additional options when creating the thread pool like this:

```
rp.create_thread_pool("my-thread-pool",
    hpx::resource::policies::local_priority_lifo,
    hpx::policies::scheduler_mode(
        hpx::policies::scheduler_mode::default_ |
        hpx::policies::scheduler_mode::enable_elasticity));
```

The available schedulers are documented here: `hpx::resource::scheduling_policy`, and the available scheduler modes here: `hpx::threads::policies::scheduler_mode`. Also see the examples folder for examples of advanced resource partitioner usage: `simple_resource_partitioner.cpp` and `oversubscribing_resource_partitioner.cpp`.

2.5.12 Miscellaneous

Error handling

Like in any other asynchronous invocation scheme, it is important to be able to handle error conditions occurring while the asynchronous (and possibly remote) operation is executed. In *HPX* all error handling is based on standard C++ exception handling. Any exception thrown during the execution of an asynchronous operation will be transferred back to the original invocation `locality`, where it will be rethrown during synchronization with the calling thread.

The source code for this example can be found here: `error_handling.cpp`.

Working with exceptions

For the following description assume that the function `raise_exception()` is executed by invoking the plain action `raise_exception_type`.

```
// [error_handling_raise_exception
void raise_exception()
{
    HPX_THROW_EXCEPTION(hpx::no_success, "raise_exception", "simulated error");
```

The exception is thrown using the macro `HPX_THROW_EXCEPTION`. The type of the thrown exception is `hpx::exception`. This associates additional diagnostic information with the exception, such as file name and line number, `locality` id and thread id, and stack backtrace from the point where the exception was thrown.

Any exception thrown during the execution of an action is transferred back to the (asynchronous) invocation site. It will be rethrown in this context when the calling thread tries to wait for the result of the action by invoking either `future<T>::get()` or the synchronous action invocation wrapper as shown here:

```
// Error reporting using exceptions
// [exception_diagnostic_information
hpx::cout << "Error reporting using exceptions\n";
try {
    // invoke raise_exception() which throws an exception
    raise_exception_action do_it;
    do_it(hpx::find_here());
}
catch (hpx::exception const& e) {
    // Print just the essential error information.
    hpx::cout << "caught exception: " << e.what() << "\n\n";

    // Print all of the available diagnostic information as stored with
    // the exception.
    hpx::cout << "diagnostic information:"
        << hpx::diagnostic_information(e) << "\n";
```

Note: The exception is transferred back to the invocation site even if it is executed on a different `locality`.

Additionally, this example demonstrates how an exception thrown by an (possibly remote) action can be handled. It shows the use of `hpx::diagnostic_information`, which retrieves all available diagnostic information from the exception as a formatted string. This includes, for instance, the name of the source file and line number, the sequence number of the OS thread and the `HPX` thread id, the `locality` id and the stack backtrace of the point where the original exception was thrown.

Under certain circumstances it is desirable to output only some of the diagnostics, or to output those using different formatting. For this case, `HPX` exposes a set of lower-level functions as demonstrated in the following code snippet:

```
// Detailed error reporting using exceptions
// [exception_diagnostic_elements
hpx::cout << "Detailed error reporting using exceptions\n";
try {
    // Invoke raise_exception() which throws an exception.
    raise_exception_action do_it;
    do_it(hpx::find_here());
}
catch (hpx::exception const& e) {
```

(continues on next page)

(continued from previous page)

```
// Print the elements of the diagnostic information separately.
hpx::cout << "{what}: "           << hpx::get_error_what(e) << "\n";
hpx::cout << "{locality-id}: "    << hpx::get_error_locality_id(e) << "\n";
hpx::cout << "{hostname}: "       << hpx::get_error_host_name(e) << "\n";
hpx::cout << "{pid}: "           << hpx::get_error_process_id(e) << "\n";
hpx::cout << "{function}: "      << hpx::get_error_function_name(e) << "\n";
hpx::cout << "{file}: "          << hpx::get_error_file_name(e) << "\n";
hpx::cout << "{line}: "          << hpx::get_error_line_number(e) << "\n";
hpx::cout << "{os-thread}: "     << hpx::get_error_os_thread(e) << "\n";
hpx::cout << "{thread-id}: "     << std::hex << hpx::get_error_thread_id(e)
    << "\n";
hpx::cout << "{thread-description}: "
    << hpx::get_error_thread_description(e) << "\n";
hpx::cout << "{state}: "         << std::hex << hpx::get_error_state(e)
    << "\n";
hpx::cout << "{stack-trace}: "   << hpx::get_error_backtrace(e) << "\n";
hpx::cout << "{env}: "          << hpx::get_error_env(e) << "\n";
```

Working with error codes

Most of the API functions exposed by *HPX* can be invoked in two different modes. By default those will throw an exception on error as described above. However, sometimes it is desirable not to throw an exception in case of an error condition. In this case an object instance of the `hpx::error_code` type can be passed as the last argument to the API function. In case of an error, the error condition will be returned in that `hpx::error_code` instance. The following example demonstrates extracting the full diagnostic information without exception handling:

```
{
    // [error_handling_diagnostic_information
    hpx::cout << "Error reporting using error code\n";

    // Create a new error_code instance.
    hpx::error_code ec;

    // If an instance of an error_code is passed as the last argument while
    // invoking the action, the function will not throw in case of an error
    // but store the error information in this error_code instance instead.
    raise_exception_action do_it;
    do_it(hpx::find_here(), ec);

    if (ec) {
        // Print just the essential error information.
        hpx::cout << "returned error: " << ec.get_message() << "\n";

        // Print all of the available diagnostic information as stored with
        // the exception.
        hpx::cout << "diagnostic information:"
            << hpx::diagnostic_information(ec) << "\n";
    }
}
```

Note: The error information is transferred back to the invocation site even if it is executed on a different *locality*.

This example show how an error can be handled without having to resolve to exceptions and that the returned `hpx::error_code` instance can be used in a very similar way as the `hpx::exception` type above. Simply

pass it to the `hpx::diagnostic_information`, which retrieves all available diagnostic information from the error code instance as a formatted string.

As for handling exceptions, when working with error codes, under certain circumstances it is desirable to output only some of the diagnostics, or to output those using different formatting. For this case, *HPX* exposes a set of lower-level functions usable with error codes as demonstrated in the following code snippet:

```
{
    // [error_handling_diagnostic_elements
    hpx::cout << "Detailed error reporting using error code\n";

    // Create a new error_code instance.
    hpx::error_code ec;

    // If an instance of an error_code is passed as the last argument while
    // invoking the action, the function will not throw in case of an error
    // but store the error information in this error_code instance instead.
    raise_exception_action do_it;
    do_it(hpx::find_here(), ec);

    if (ec) {
        // Print the elements of the diagnostic information separately.
        hpx::cout << "{what}: " << hpx::get_error_what(ec) << "\n";
        hpx::cout << "{locality-id}: " << hpx::get_error_locality_id(ec) << "\n";
        hpx::cout << "{hostname}: " << hpx::get_error_host_name(ec) << "\n";
        hpx::cout << "{pid}: " << hpx::get_error_process_id(ec) << "\n";
        hpx::cout << "{function}: " << hpx::get_error_function_name(ec)
            << "\n";
        hpx::cout << "{file}: " << hpx::get_error_file_name(ec) << "\n";
        hpx::cout << "{line}: " << hpx::get_error_line_number(ec) << "\n";
        hpx::cout << "{os-thread}: " << hpx::get_error_os_thread(ec) << "\n";
        hpx::cout << "{thread-id}: " << std::hex
            << hpx::get_error_thread_id(ec) << "\n";
        hpx::cout << "{thread-description}: "
            << hpx::get_error_thread_description(ec) << "\n\n";
        hpx::cout << "{state}: " << std::hex << hpx::get_error_state(ec)
            << "\n";
        hpx::cout << "{stack-trace}: " << hpx::get_error_backtrace(ec) << "\n";
        hpx::cout << "{env}: " << hpx::get_error_env(ec) << "\n";
    }
}
```

For more information please refer to the documentation of `hpx::get_error_what`, `hpx::get_error_locality_id`, `hpx::get_error_host_name`, `hpx::get_error_process_id`, `hpx::get_error_function_name`, `hpx::get_error_file_name`, `hpx::get_error_line_number`, `hpx::get_error_os_thread`, `hpx::get_error_thread_id`, `hpx::get_error_thread_description`, `hpx::get_error_backtrace`, `hpx::get_error_env`, and `hpx::get_error_state`.

Lightweight error codes

Sometimes it is not desirable to collect all the ambient information about the error at the point where it happened as this might impose too much overhead for simple scenarios. In this case, *HPX* provides a lightweight error code facility that will hold the error code only. The following snippet demonstrates its use:

```
{  
    // [lightweight_error_handling_diagnostic_information  
    hpx::cout << "Error reporting using an lightweight error code\n";  
  
    // Create a new error_code instance.  
    hpx::error_code ec(hpx::lightweight);  
  
    // If an instance of an error_code is passed as the last argument while  
    // invoking the action, the function will not throw in case of an error  
    // but store the error information in this error_code instance instead.  
    raise_exception_action do_it;  
    do_it(hpx::find_here(), ec);  
  
    if (ec) {  
        // Print just the essential error information.  
        hpx::cout << "returned error: " << ec.get_message() << "\n";  
  
        // Print all of the available diagnostic information as stored with  
        // the exception.  
        hpx::cout << "error code:" << ec.value() << "\n";  
    }  
}
```

All functions that retrieve other diagnostic elements from the `hpx::error_code` will fail if called with a lightweight error_code instance.

Utilities in *HPX*

In order to ease the burden of programming, *HPX* provides several utilities to users. The following section documents those facilities.

Checkpoint

See *checkpoint*.

The *HPX* I/O-streams component

The *HPX* I/O-streams subsystem extends the standard C++ output streams `std::cout` and `std::cerr` to work in the distributed setting of an *HPX* application. All of the output streamed to `hpx::cout` will be dispatched to `std::cout` on the console *locality*. Likewise, all output generated from `hpx::cerr` will be dispatched to `std::cerr` on the console *locality*.

Note: All existing standard manipulators can be used in conjunction with `hpx::cout` and `hpx::cerr`. Historically, *HPX* also defines `hpx::endl` and `hpx::flush` but those are just aliases for the corresponding standard manipulators.

In order to use either `hpx::cout` or `hpx::cerr`, application codes need to `#include <hpx/include/iostreams.hpp>`. For an example, please see the following ‘Hello world’ program:

```
// Copyright (c) 2007-2012 Hartmut Kaiser
//
// SPDX-License-Identifier: BSL-1.0
// Distributed under the Boost Software License, Version 1.0. (See accompanying
// file LICENSE_1_0.txt or copy at http://www.boost.org/LICENSE_1_0.txt)

///////////////////////////////
// The purpose of this example is to execute a HPX-thread printing
// "Hello World!" once. That's all.

//[hello_world_1_getting_started
// Including 'hpx/hpx_main.hpp' instead of the usual 'hpx/hpx_init.hpp' enables
// to use the plain C-main below as the direct main HPX entry point.
#include <hpx/hpx_main.hpp>
#include <hpx/include/iostreams.hpp>

int main()
{
    // Say hello to the world!
    hpx::cout << "Hello World!\n" << hpx::flush;
    return 0;
}
//]
```

Additionally, those applications need to link with the `iostreams` component. When using CMake this can be achieved by using the `COMPONENT_DEPENDENCIES` parameter; for instance:

```
include(HPX_AddExecutable)

add_hpx_executable(
    hello_world
    SOURCES hello_world.cpp
    COMPONENT_DEPENDENCIES iostreams
)
```

Note: The `hpx::cout` and `hpx::cerr` streams buffer all output locally until a `std::endl` or `std::flush` is encountered. That means that no output will appear on the console as long as either of these is explicitly used.

2.5.13 Troubleshooting

This section contains commonly encountered problems when compiling or using HPX.

Undefined reference to boost::program_options

Boost.ProgramOptions is not ABI compatible between all C++ versions and compilers. Because of this you may see linker errors similar to this:

```
....: undefined reference to `boost::program_options::operator<<(std::ostream&,   
boost::program_options::options_description const&)'
```

if you are not linking to a compatible version of Boost.ProgramOptions. We recommend that you use `hpx::program_options`, which is part of *HPX*, as a replacement for `boost::program_options` (see [program_options](#)). Until you have migrated to use `hpx::program_options` we recommend that you always build Boost²²² libraries and *HPX* with the same compiler and C++ standard.

Undefined reference to hpx::cout

You may see an linker error message that looks a bit like this:

```
hello_world.cpp:(.text+0x5aa): undefined reference to `hpx::cout'  
hello_world.cpp:(.text+0x5c3): undefined reference to `hpx::iostreams::flush'
```

This usually happens if you are trying to use *HPX* iostreams functionality such as `hpx::cout` but are not linking against it. The iostreams functionality is not part of the core *HPX* library, and must be linked to explicitly. Typically this can be solved by adding `COMPONENT_DEPENDENCIES iostreams` to a call to `add_hpx_library/` `add_hpx_executable/hpx_setup_target` if using CMake. See [Creating HPX projects](#) for more details.

2.6 Additional material

- 2-day workshop held at CSCS in 2016
 - Recorded lectures²²³
 - Slides²²⁴
- Tutorials repository²²⁵
- STE||AR Group blog posts²²⁶

²²² <https://www.boost.org/>

²²³ <https://www.youtube.com/playlist?list=PL1tk5lGm7zvSXfS-sqOOmIJ0lFNjKze18>

²²⁴ <https://github.com/STELLAR-GROUP/tutorials/tree/master/csccs2016>

²²⁵ <https://github.com/STELLAR-GROUP/tutorials>

²²⁶ <http://stellar-group.org/blog/>

2.7 Overview

HPX is organized into different sub-libraries. Those libraries can be seen as independent modules, with clear dependencies and no cycles. As an end-user, the use of these modules is completely transparent. If you use e.g. `add_hpx_executable` to create a target in your project you will automatically get all modules as dependencies. See [All modules](#) for a list of the available modules.

2.8 All modules

2.8.1 affinity

2.8.2 algorithms

The algorithms module exposes the full set of algorithms defined by the C++ standard. There is also partial support for C++ ranges.

See the [API reference](#) of the module for more details.

2.8.3 allocator_support

This module provides utilities for allocators. It contains `hpx::util::internal_allocator` which directly forwards allocation calls to `jemalloc`. This utility is mainly useful on Windows.

See the [API reference](#) of the module for more details.

2.8.4 assertion

The assertion library implements the macros `HPX_ASSERT` and `HPX_ASSERT_MSG`. Those two macros can be used to implement assertions which are turned off during a release build.

By default, the location and function where the assert has been called from are displayed when the assertion fires. This behavior can be modified by using `hpx::assertion::set_assertion_handler`. When HPX initializes, it uses this function to specify a more elaborate assertion handler. If your application needs to customize this, it needs to do so before calling `hpx::hpx_init`, `hpx::hpx_main` or using the C-main wrappers.

See the [API reference](#) of the module for more details.

2.8.5 basic_execution

The basic execution module is the main entry point to implement parallel and concurrent operations. It is modeled after P0443²²⁷ with some additions and implementations for the described concepts. Most notably, it provides an abstraction for execution resources, execution contexts and execution agents in such a way, that it provides customization points that those aforementioned concepts can be replaced and combined with ease.

For that purpose, three virtual base classes are provided to be able to provide implementations with different properties:

- **resource_base**: This is the abstraction for execution resources, that is for example CPU cores or an accelerator.
- **context_base**: An execution context uses execution resources and is able to spawn new execution agents, as new threads of executions on the available resources.

²²⁷ <http://wg21.link/p0443>

- **agent_base**: The execution agent represents the thread of execution, and can be used to yield, suspend, resume or abort a thread of execution.

2.8.6 batch_environments

2.8.7 cache

This module provides two cache data structures:

- `hpx::util::cache::local_cache`
- `hpx::util::cache::lru_cache`

See the [API reference](#) of the module for more details.

2.8.8 checkpoint

A common need of users is to periodically backup an application. This practice provides resiliency and potential restart points in code. *HPX* utilizes the concept of a `checkpoint` to support this use case.

Found in `hpx/util/checkpoint.hpp`, checkpoints are defined as objects that hold a serialized version of an object or set of objects at a particular moment in time. This representation can be stored in memory for later use or it can be written to disk for storage and/or recovery at a later point. In order to create and fill this object with data, users must use a function called `save_checkpoint`. In code the function looks like this:

```
hpx::future<hpx::util::checkpoint> hpx::util::save_checkpoint(a, b, c, ...);
```

`save_checkpoint` takes arbitrary data containers, such as `int`, `double`, `float`, `vector`, and `future`, and serializes them into a newly created `checkpoint` object. This function returns a `future` to a `checkpoint` containing the data. Here's an example of a simple use case:

```
using hpx::util::checkpoint;
using hpx::util::save_checkpoint;

std::vector<int> vec{1, 2, 3, 4, 5};
hpx::future<checkpoint> save_checkpoint(vec);
```

Once the future is ready, the `checkpoint` object will contain the `vector` `vec` and its five elements.

It is also possible to modify the launch policy used by `save_checkpoint`. This is accomplished by passing a launch policy as the first argument. It is important to note that passing `hpx::launch::sync` will cause `save_checkpoint` to return a `checkpoint` instead of a `future` to a `checkpoint`. All other policies passed to `save_checkpoint` will return a `future` to a `checkpoint`.

Sometimes `checkpoint`s must be declared before they are used. `save_checkpoint` allows users to move pre-created `checkpoint`s into the function as long as they are the first container passing into the function (In the case where a launch policy is used, the `checkpoint` will immediately follow the launch policy). An example of these features can be found below:

```
char character = 'd';
int integer = 10;
float flt = 10.01f;
bool boolean = true;
std::string str = "I am a string of characters";
std::vector<char> vec(str.begin(), str.end());
checkpoint archive;
```

(continues on next page)

(continued from previous page)

```
// Test 1
// test basic functionality
hpx::shared_future<checkpoint> f_archive = save_checkpoint(
    std::move(archive), character, integer, flt, boolean, str, vec);
```

Once users can create checkpoints they must now be able to restore the objects they contain into memory. This is accomplished by the function `restore_checkpoint`. This function takes a `checkpoint` and fills its data into the containers it is provided. It is important to remember that the containers must be ordered in the same way they were placed into the `checkpoint`. For clarity see the example below:

```
char character2;
int integer2;
float flt2;
bool boolean2;
std::string str2;
std::vector<char> vec2;

restore_checkpoint(
    f_archive.get(), character2, integer2, flt2, boolean2, str2, vec2);
```

The core utility of `checkpoint` is in its ability to make certain data persistent. Often, this means that the data needs to be stored in an object, such as a file, for later use. *HPX* has two solutions for these issues: stream operator overloads and access iterators.

HPX contains two stream overloads, `operator<<` and `operator>>`, to stream data out of and into `checkpoint`. Here is an example of the overloads in use below:

```
double a9 = 1.0, b9 = 1.1, c9 = 1.2;
std::ofstream test_file_9("test_file_9.txt");
hpx::future<checkpoint> f_9 = save_checkpoint(a9, b9, c9);
test_file_9 << f_9.get();
test_file_9.close();

double a9_1, b9_1, c9_1;
std::ifstream test_file_9_1("test_file_9.txt");
checkpoint archive9;
test_file_9_1 >> archive9;
restore_checkpoint(archive9, a9_1, b9_1, c9_1);
```

This is the primary way to move data into and out of a `checkpoint`. It is important to note, however, that users should be cautious when using a stream operator to load data and another function to remove it (or vice versa). Both `operator<<` and `operator>>` rely on a `.write()` and a `.read()` function respectively. In order to know how much data to read from the `std::istream`, the `operator<<` will write the size of the `checkpoint` before writing the `checkpoint` data. Correspondingly, the `operator>>` will read the size of the stored data before reading the data into a new instance of `checkpoint`. As long as the user employs the `operator<<` and `operator>>` to stream the data, this detail can be ignored.

Important: Be careful when mixing `operator<<` and `operator>>` with other facilities to read and write to a `checkpoint`. `operator<<` writes an extra variable, and `operator>>` reads this variable back separately. Used together the user will not encounter any issues and can safely ignore this detail.

Users may also move the data into and out of a `checkpoint` using the exposed `.begin()` and `.end()` iterators. An example of this use case is illustrated below.

```
std::ofstream test_file_7("checkpoint_test_file.txt");
std::vector<float> vec7{1.02f, 1.03f, 1.04f, 1.05f};
hpx::future<checkpoint> fut_7 = save_checkpoint(vec7);
checkpoint archive7 = fut_7.get();
std::copy(archive7.begin(),           // Write data to ostream
          archive7.end(),          // ie. the file
          std::ostream_iterator<char>(test_file_7));
test_file_7.close();

std::vector<float> vec7_1;
std::vector<char> char_vec;
std::ifstream test_file_7_1("checkpoint_test_file.txt");
if (test_file_7_1)
{
    test_file_7_1.seekg(0, test_file_7_1.end);
    auto length = test_file_7_1.tellg();
    test_file_7_1.seekg(0, test_file_7_1.beg);
    char_vec.resize(length);
    test_file_7_1.read(char_vec.data(), length);
}
checkpoint archive7_1(std::move(char_vec));      // Write data to checkpoint
restore_checkpoint(archive7_1, vec7_1);
```

Checkpointing components

`save_checkpoint` and `restore_checkpoint` are also able to store components inside checkpoints. This can be done in one of two ways. First a client of the component can be passed to `save_checkpoint`. When the user wishes to resurrect the component she can pass a client instance to `restore_checkpoint`.

This technique is demonstrated below:

```
// Try to checkpoint and restore a component with a client
std::vector<int> vec3{10, 10, 10, 10, 10};

// Create a component instance through client constructor
data_client D(hpx::find_here(), std::move(vec3));
hpx::future<checkpoint> f3 = save_checkpoint(D);

// Create a new client
data_client E;

// Restore server inside client instance
restore_checkpoint(f3.get(), E);
```

The second way a user can save a component is by passing a `shared_ptr` to the component to `save_checkpoint`. This component can be resurrected by creating a new instance of the component type and passing a `shared_ptr` to the new instance to `restore_checkpoint`.

This technique is demonstrated below:

```
// test checkpoint a component using a shared_ptr
std::vector<int> vec{1, 2, 3, 4, 5};
data_client A(hpx::find_here(), std::move(vec));

// Checkpoint Server
hpx::id_type old_id = A.get_id();
```

(continues on next page)

(continued from previous page)

```

hpx::future<std::shared_ptr<data_server>> f_a_ptr =
    hpx::get_ptr<data_server>(A.get_id());
std::shared_ptr<data_server> a_ptr = f_a_ptr.get();
hpx::future<checkpoint> f = save_checkpoint(a_ptr);

// Restore Server
// Create a new server instance
std::shared_ptr<data_server> b_server;
restore_checkpoint(f.get(), b_server);

```

2.8.9 collectives

The collectives module exposes a set of distributed collective operations. Those can be used to exchange data between participating sites in a coordinated way. At this point the module exposes the following collective primitives:

- `hpx::collectives::all_reduce`: performs a reduction on data from each participating site to each participating site.
- `hpx::collectives::all_to_all`: each participating site provides its element of the data to collect while all participating sites receive the data from every other site.
- `hpx::lcos::barrier`: distributed barrier.
- `hpx::lcos::broadcast`: performs a given action on all given global identifiers.
- `hpx::lcos::fold`: performs a fold with a given action on all given global identifiers.
- `hpx::lcos::gather`: gathers values from all participating sites.
- `hpx::lcos::latch`: distributed latch.
- `hpx::lcos::reduce`: performs a reduction on data from each participating site to a root site.
- `hpx::lcos::spmd_block`: performs the same operation on a local image while providing handles to the other images.

See the [API reference](#) of the module for more details.

2.8.10 compute

The compute module provides utilities for handling task and memory affinity on host systems. The `compute_cuda` for extensions to CUDA programmable GPU devices.

See the [API reference](#) of the module for more details.

2.8.11 compute_cuda

This module extends the `compute` module to handle CUDA programmable GPU devices.

See the [API reference](#) of the module for more details.

2.8.12 concepts

This module provides helpers for emulating concepts. It provides the following macros:

- `HPX_CONCEPT_REQUIRES`
- `HPX_HAS_MEMBER_XXX_TRAIT_DEF`
- `HPX_HAS_XXX_TRAIT_DEF`

See the [API reference](#) of the module for more details.

2.8.13 concurrency

This module provides concurrency primitives useful for multi-threaded programming such as:

- `hpx::util::barrier`
- `hpx::util::cache_line_data` and `hpx::util::cache_aligned_data`: wrappers for aligning and padding data to cache lines.
- various lockfree queue data structures

See the [API reference](#) of the module for more details.

2.8.14 config

The config module contains various configuration options, typically hidden behind macros that choose the correct implementation based on the compiler and other available options.

See the [API reference](#) of the module for more details.

2.8.15 coroutines

The coroutines module provides coroutine (user-space thread) implementations for different platforms.

See the [API reference](#) of the module for more details.

2.8.16 datastructures

The datastructures module provides basic data structures (typically provided for compatibility with older C++ standards):

- `hpx::util::basic_any`
- `hpx::util::optional`
- `hpx::util::tuple`

See the [API reference](#) of the module for more details.

2.8.17 debugging

This module provides helpers for demangling symbol names.

See the [API reference](#) of the module for more details.

2.8.18 errors

This module provides support for exceptions and error codes:

- `hpx::exception`
- `hpx::error_code`
- `:c:enum:`hpx::error``

See the [API reference](#) of the module for more details.

2.8.19 execution

This library implements executors and execution policies for use with parallel algorithms and other facilities related to managing the execution of tasks.

See the [API reference](#) of the module for more details.

2.8.20 filesystem

This module provides a compatibility layer for the C++17 filesystem library. If the filesystem library is available this module will simply forward its contents into the `hpx::filesystem` namespace. If the library is not available it will fall back to Boost.Filesystem instead.

See the [API reference](#) of the module for more details.

2.8.21 format

The format module exposes the `format` and `format_to` functions for formatting strings.

See the [API reference](#) of the module for more details.

2.8.22 functional

This module provides function wrappers and helpers for managing functions and their arguments.

- `hpx::util::function`
- `hpx::util::function_ref`
- `hpx::util::unique_function`
- `hpx::util::bind`
- `hpx::util::bind_back`
- `hpx::util::bind_front`
- `hpx::util::deferred_call`
- `hpx::util::invoke`

- `hpx::util::invoke_fused`
- `hpx::util::mem_fn`
- `hpx::util::one_shot`
- `hpx::util::protect`
- `hpx::util::result_of`

See the [API reference](#) of the module for more details.

2.8.23 hardware

The hardware module abstracts away hardware specific details of timestamps and CPU features.

See the [API reference](#) of the module for more details.

2.8.24 hashing

The hashing module provides two hashing implementations:

- `hpx::util::fibhash`
- `hpx::util::jenkins_hash`

See the [API reference](#) of the module for more details.

2.8.25 iterator_support

This module provides helpers for iterators. It provides `hpx::util::iterator_facade` and `hpx::util::iterator_adaptor` for creating new iterators, and the trait `hpx::util::is_iterator` along with more specific iterator traits.

See the [API reference](#) of the module for more details.

2.8.26 local_lcos

This module provides local LCOs which return futures:

- `hpx::lcos::local::and_gate`
- `hpx::lcos::local::channel`
- `hpx::lcos::local::one_element_channel`
- `hpx::lcos::local::receive_channel`
- `hpx::lcos::local::send_channel`
- `hpx::lcos::local::guard`
- `hpx::lcos::local::guard_set`
- `hpx::lcos::local::run_guarded`
- `hpx::lcos::local::conditional_trigger`
- `hpx::lcos::local::packaged_task`
- `hpx::lcos::local::promise`

- `hpx::lcos::local::receive_buffer`
- `hpx::lcos::local::trigger`

See the [API reference](#) of this module for more details.

2.8.27 logging

This module provides useful macros for logging information.

See the [API reference](#) of the module for more details.

2.8.28 memory

Part of this module is a forked version of boost::intrusive_ptr from Boost.SmartPtr.

See the [API reference](#) of the module for more details.

2.8.29 naming_base

This module provides a forward declaration of `address_type`, `component_type` and `invalid_locality_id`.

See the [API reference](#) of this module for more details.

2.8.30 plugin

This module provides base utilities for creating plugins.

See the [API reference](#) of the module for more details.

2.8.31 preprocessor

This library contains useful preprocessor macros:

- `HPX_PP_CAT`
- `HPX_PP_EXPAND`
- `HPX_PP_NARGS`
- `HPX_PP_STRINGIZE`
- `HPX_PP_STRIP_PARENS`

See the [API reference](#) of the module for more details.

2.8.32 program_options

The module `program_options` is a direct fork of the Boost.ProgramOptions library (Boost V1.70.0). For more information about this library please see [here²²⁸](#). In order to be included as an HPX module, the Boost.ProgramOptions library has been moved to the namespace `hpx::program_options`. We have also replaced all Boost facilities the library depends on with either the equivalent facilities from the standard library or from HPX. As a result, the HPX `program_options` module is fully interface compatible with Boost.ProgramOptions (sans the `hpx` namespace and the `#include <hpx/program_options.hpp>` changes that need to be applied to all code relying on this library).

All credit goes to Vladimir Prus, the author of the excellent Boost.ProgramOptions library. All bugs have been introduced by us.

See the [API reference](#) of the module for more details.

2.8.33 resiliency

In *HPX*, a program failure is a manifestation of a failing task. This module exposes several APIs that allow to manage failing tasks in a convenient way by either replaying a failed task or by replicating a specific task to begin with.

Task replay is analogous to the Checkpoint/Restart mechanism found in conventional execution models. The key difference being localized fault detection. When the runtime detects an error it replays the failing task as opposed to completely rolling back of the entire program to the previous checkpoint.

Task replication is designed to provide reliability enhancements by replicating a set of tasks and evaluating their results to determine a consensus among them. This technique is most effective in situations where there are few tasks in the critical path of the DAG which leaves the system underutilized or where hardware or software failures may result in an incorrect result instead of an error. However, the drawback of this method is the additional computational cost incurred by repeating a task multiple times.

The following API functions are exposed:

- `hpx::resiliency::async_replay`: This version of task replay will catch user defined exceptions and automatically reschedule the task N times before throwing an `hpx::resiliency::abort_replay_exception` if no task is able to complete execution without an exception.
- `hpx::resiliency::async_replay_validate`: This version of replay adds an argument to `async_replay` which receives a user provided validation function to test the result of the task against. If the task's output is validated, the result is returned. If the output fails the check or an exception is thrown, the task is replayed until no errors are encountered or the number of specified retries have been exceeded.
- `hpx::resiliency::async_replicate`: This is the most basic implementation of the task replication. The API returns the first result that runs without detecting any errors.
- `hpx::resiliency::async_replicate_validate`: This API additionally takes a validation function which evaluates the return values produced by the threads. The first task to compute a valid result is returned.
- `hpx::resiliency::async_replicate_vote`: This API adds a vote function to the basic replicate function. Many hardware or software failures are silent errors which do not interrupt program flow. In order to detect errors of this kind, it is necessary to run the task several times and compare the values returned by every version of the task. In order to determine which return value is "correct", the API allows the user to provide a custom consensus function to properly form a consensus. This voting function then returns the "correct" answer.
- `hpx::resiliency::async_replicate_vote_validate`: This combines the features of the previously discussed replicate set. Replicate vote validate allows a user to provide a validation function to filter

²²⁸ https://www.boost.org/doc/libs/1_70_0/doc/html/program_options.html

results. Additionally, as described in replicate vote, the user can provide a “voting function” which returns the consensus formed by the voting logic.

- `hpx::resiliency::dataflow_replay`: This version of dataflow replay will catch user defined exceptions and automatically reschedules the task N times before throwing an `hpx::resiliency::abort_replay_exception` if no task is able to complete execution without an exception. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.
- `hpx::resiliency::dataflow_replay_validate`: This version of replay adds an argument to dataflow replay which receives a user provided validation function to test the result of the task against. If the task’s output is validated, the result is returned. If the output fails the check or an exception is thrown, the task is replayed until no errors are encountered or the number of specified retries have been exceeded. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.
- `hpx::resiliency::dataflow_replicate`: This is the most basic implementation of the task replication. The API returns the first result that runs without detecting any errors. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.
- `hpx::resiliency::dataflow_replicate_validate`: This API additionally takes a validation function which evaluates the return values produced by the threads. The first task to compute a valid result is returned. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.
- `hpx::resiliency::dataflow_replicate_vote`: This API adds a vote function to the basic replicate function. Many hardware or software failures are silent errors which do not interrupt program flow. In order to detect errors of this kind, it is necessary to run the task several times and compare the values returned by every version of the task. In order to determine which return value is “correct”, the API allows the user to provide a custom consensus function to properly form a consensus. This voting function then returns the “correct” answer. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.
- `hpx::resiliency::dataflow_replicate_vote_validate`: This combines the features of the previously discussed replicate set. Replicate vote validate allows a user to provide a validation function to filter results. Additionally, as described in replicate vote, the user can provide a “voting function” which returns the consensus formed by the voting logic. Any arguments for the executed task that are futures will cause the task invocation to be delayed until all of those futures have become ready.

See the [API reference](#) of the module for more details.

2.8.34 resource_partitioner

2.8.35 segmented_algorithms

Segmented algorithms extend the usual parallel algorithms (*algorithms*) by providing overloads that work with distributed containers, such as partitioned vectors.

See the [API reference](#) of the module for more details.

2.8.36 serialization

This module provides serialization primitives and support for many C++ standard library and Boost types.

See the [API reference](#) of the module for more details.

2.8.37 static_reinit

2.8.38 statistics

This module provide some statistics utilities like rolling min/max and histogram.

See the [API reference](#) of the module for more details.

2.8.39 synchronization

This module provides synchronization primitives which should be used rather than the C++ standard ones:

- `hpx::lcos::local::barrier`
- `hpx::lcos::local::condition_variable`
- `hpx::lcos::local::counting_semaphore`
- `hpx::lcos::local::event`
- `hpx::lcos::local::latch`
- `hpx::lcos::local::mutex`
- `hpx::lcos::local::no_mutex`
- `hpx::lcos::local::once_flag`
- `hpx::lcos::local::recursive_mutex`
- `hpx::lcos::local::shared_mutex`
- `hpx::lcos::local::sliding_semaphore`
- `hpx::lcos::local::spinlock` (`std::mutex` compatible spinlock)
- `hpx::lcos::local::spinlock_no_backoff` (`boost::mutex` compatible spinlock)
- `hpx::lcos::local::spinlock_pool`

See the [API reference](#) of this module for more details.

2.8.40 testing

The testing module contains useful macros for testing. The results of tests can be printed with `hpx::util::report_errors`. The following macros are provided:

- `HPX_TEST`
- `HPX_TEST_MSG`
- `HPX_TEST_EQ`
- `HPX_TEST_NEQ`
- `HPX_TEST_LT`

- `HPX_TEST_LTE`
- `HPX_TEST_RANGE`
- `HPX_TEST_EQ_MSG`
- `HPX_TEST_NEQ_MSG`
- `HPX_SANITY`
- `HPX_SANITY_MSG`
- `HPX_SANITY_EQ`
- `HPX_SANITY_NEQ`
- `HPX_SANITY_LT`
- `HPX_SANITY_LTE`
- `HPX_SANITY_RANGE`
- `HPX_SANITY_EQ_MSG`

See the [API reference](#) of the module for more details.

2.8.41 `thread_support`

This module provides miscellaneous utilities for threading and concurrency.

See the [API reference](#) of the module for more details.

2.8.42 `thread_manager`

2.8.43 `timing`

This module provides the timing utilities (clocks and timers).

See the [API reference](#) of the module for more details.

2.8.44 `topology`

This module provides the class `hpx::threads::topology` which represents the hardware resources available on a node. The class is a light wrapper around the Portable Hardware Locality (HWLOC)²²⁹ library. The `hpx::threads::cpu_mask` is a small companion class that represents a set of resources on a node.

See the [API reference](#) of the module for more details.

²²⁹ <https://www.open-mpi.org/projects/hwloc/>

2.8.45 type_support

This module provides helper facilities related to types.

See the [API reference](#) of the module for more details.

2.8.46 util

The util module provides miscellaneous standalone utilities.

See the [API reference](#) of the module for more details.

2.9 API reference

2.9.1 Main HPX library reference

```
template<typename Action>
struct async_result
    #include <colocating_distribution_policy.hpp>
```

Note This function is part of the invocation policy implemented by this class

Public Types

```
template<>
using type = hpx::future<typename traits::promise_local_result<typename hpx::traits::extract_action<Action>::remote_
```

```
template<typename Action>
struct async_result
    #include <default_distribution_policy.hpp>
```

Note This function is part of the invocation policy implemented by this class

Public Types

```
template<>
using type = hpx::future<typename traits::promise_local_result<typename hpx::traits::extract_action<Action>::remote_
```

```
struct binpacking_distribution_policy
    #include <binpacking_distribution_policy.hpp> This class specifies the parameters for a binpacking distribution
    policy to use for creating a given number of items on a given set of localities. The binpacking policy will
    distribute the new objects in a way such that each of the localities will equalize the number of overall objects of
    this type based on a given criteria (by default this criteria is the overall number of objects of this type).
```

Public Functions

`binpacking_distribution_policy()`

Default-construct a new instance of a `binpacking_distribution_policy`. This policy will represent one locality (the local locality).

`binpacking_distribution_policy operator() (std::vector<id_type> const &locs,
char const *perf_counter_name = default_binpacking_counter_name) const`

Create a new `default_distribution` policy representing the given set of localities.

Parameters

- `locs`: [in] The list of localities the new instance should represent
- `perf_counter_name`: [in] The name of the performance counter which should be used as the distribution criteria (by default the overall number of existing instances of the given component type will be used).

`binpacking_distribution_policy operator() (std::vector<id_type> &&locs,
const *perf_counter_name = default_binpacking_counter_name) const`

Create a new `default_distribution` policy representing the given set of localities.

Parameters

- `locs`: [in] The list of localities the new instance should represent
- `perf_counter_name`: [in] The name of the performance counter which should be used as the distribution criteria (by default the overall number of existing instances of the given component type will be used).

`binpacking_distribution_policy operator() (id_type const &loc, char const *perf_counter_name
= default_binpacking_counter_name) const`

Create a new `default_distribution` policy representing the given locality

Parameters

- `loc`: [in] The locality the new instance should represent
- `perf_counter_name`: [in] The name of the performance counter which should be used as the distribution criteria (by default the overall number of existing instances of the given component type will be used).

`template<typename Component, typename ...Ts>`

`hpx::future<hpx::id_type> create(Ts&&... vs) const`

Create one object on one of the localities associated by this policy instance

Return A future holding the global address which represents the newly created object

Parameters

- `vs`: [in] The arguments which will be forwarded to the constructor of the new object.

`template<typename Component, typename ...Ts>`

`hpx::future<std::vector<bulk_locality_result>> bulk_create(std::size_t count, Ts&&... vs) const`

Create multiple objects on the localities associated by this policy instance

Return A future holding the list of global addresses which represent the newly created objects

Parameters

- count: [in] The number of objects to create
- vs: [in] The arguments which will be forwarded to the constructors of the new objects.

`std::string const &get_counter_name() const`

Returns the name of the performance counter associated with this policy instance.

`std::size_t get_num_localities() const`

Returns the number of associated localities for this distribution policy

Note This function is part of the creation policy implemented by this class

`struct colocating_distribution_policy`

`#include <colocating_distribution_policy.hpp>` This class specifies the parameters for a distribution policy to use for creating a given number of items on the locality where a given object is currently placed.

Public Functions

`colocating_distribution_policy()`

Default-construct a new instance of a `colocating_distribution_policy`. This policy will represent the local locality.

`colocating_distribution_policy operator()(id_type const &id) const`

Create a new `colocating_distribution_policy` representing the locality where the given object os current located

Parameters

- id: [in] The global address of the object with which the new instances should be colocated on

`template<typename Client, typename Stub>`

`colocating_distribution_policy operator()(client_base<Client, Stub> const &client) const`

Create a new `colocating_distribution_policy` representing the locality where the given object os current located

Parameters

- client: [in] The client side representation of the object with which the new instances should be colocated on

`template<typename Component, typename ...Ts>`

`hpx::future<hpx::id_type> create(Ts&&... vs) const`

Create one object on the locality of the object this distribution policy instance is associated with

Note This function is part of the placement policy implemented by this class

Return A future holding the global address which represents the newly created object

Parameters

- vs: [in] The arguments which will be forwarded to the constructor of the new object.

`template<typename Component, typename ...Ts>`

`hpx::future<std::vector<bulk_locality_result>> bulk_create (std::size_t count, Ts&&... vs) const`
Create multiple objects colocated with the object represented by this policy instance

Note This function is part of the placement policy implemented by this class

Return A future holding the list of global addresses which represent the newly created objects

Parameters

- `count`: [in] The number of objects to create
- `vs`: [in] The arguments which will be forwarded to the constructors of the new objects.

`template<typename Action, typename ...Ts>`
`async_result<Action>::type async (launch policy, Ts&&... vs) const`

`template<typename Action, typename Callback, typename ...Ts>`
`async_result<Action>::type async_cb (launch policy, Callback &&cb, Ts&&... vs) const`

Note This function is part of the invocation policy implemented by this class

`template<typename Action, typename Continuation, typename ...Ts>`
`bool apply (Continuation &&c, threads::thread_priority priority, Ts&&... vs) const`

Note This function is part of the invocation policy implemented by this class

`template<typename Action, typename ...Ts>`
`bool apply (threads::thread_priority priority, Ts&&... vs) const`

`template<typename Action, typename Continuation, typename Callback, typename ...Ts>`
`bool apply_cb (Continuation &&c, threads::thread_priority priority, Callback &&cb, Ts&&... vs)`
`const`

Note This function is part of the invocation policy implemented by this class

`template<typename Action, typename Callback, typename ...Ts>`
`bool apply_cb (threads::thread_priority priority, Callback &&cb, Ts&&... vs) const`

`std::size_t get_num_localities () const`
Returns the number of associated localities for this distribution policy

Note This function is part of the creation policy implemented by this class

`hpx::id_type get_next_target () const`
Returns the locality which is anticipated to be used for the next async operation

struct default_distribution_policy

`#include <default_distribution_policy.hpp>` This class specifies the parameters for a simple distribution policy to use for creating (and evenly distributing) a given number of items on a given set of localities.

Public Functions

`default_distribution_policy()`

Default-construct a new instance of a `default_distribution_policy`. This policy will represent one locality (the local locality).

`default_distribution_policy operator() (std::vector<id_type> const &locs) const`

Create a new `default_distribution` policy representing the given set of localities.

Parameters

- `locs`: [in] The list of localities the new instance should represent

`default_distribution_policy operator() (std::vector<id_type> &&locs) const`

Create a new `default_distribution` policy representing the given set of localities.

Parameters

- `locs`: [in] The list of localities the new instance should represent

`default_distribution_policy operator() (id_type const &loc) const`

Create a new `default_distribution` policy representing the given locality

Parameters

- `loc`: [in] The locality the new instance should represent

`template<typename Component, typename ...Ts>`

`hpx::future<hpx::id_type> create (Ts&&... vs) const`

Create one object on one of the localities associated by this policy instance

Note This function is part of the placement policy implemented by this class

Return A future holding the global address which represents the newly created object

Parameters

- `vs`: [in] The arguments which will be forwarded to the constructor of the new object.

`template<typename Component, typename ...Ts>`

`hpx::future<std::vector<bulk_locality_result>> bulk_create (std::size_t count, Ts&&... vs) const`

Create multiple objects on the localities associated by this policy instance

Note This function is part of the placement policy implemented by this class

Return A future holding the list of global addresses which represent the newly created objects

Parameters

- `count`: [in] The number of objects to create
- `vs`: [in] The arguments which will be forwarded to the constructors of the new objects.

`template<typename Action, typename ...Ts>`

`async_result<Action>::type async (launch policy, Ts&&... vs) const`

`template<typename Action, typename Callback, typename ...Ts>`

`async_result<Action>::type async_cb (launch policy, Callback &&cb, Ts&&... vs) const`

Note This function is part of the invocation policy implemented by this class

```
template<typename Action, typename Continuation, typename ...Ts>
bool apply(Continuation &&c, threads::thread_priority priority, Ts&&... vs) const
```

Note This function is part of the invocation policy implemented by this class

```
template<typename Action, typename ...Ts>
bool apply(threads::thread_priority priority, Ts&&... vs) const
```

```
template<typename Action, typename Continuation, typename Callback, typename ...Ts>
bool apply_cb(Continuation &&c, threads::thread_priority priority, Callback &&cb, Ts&&... vs)
const
```

Note This function is part of the invocation policy implemented by this class

```
template<typename Action, typename Callback, typename ...Ts>
bool apply_cb(threads::thread_priority priority, Callback &&cb, Ts&&... vs) const
```

```
std::size_t get_num_localities() const
```

Returns the number of associated localities for this distribution policy

Note This function is part of the creation policy implemented by this class

```
hpx::id_type get_next_target() const
```

Returns the locality which is anticipated to be used for the next async operation

```
template<>
struct hash<hpx::threads::thread_id>
#include <thread_data_fwd.hpp>
```

Public Functions

```
std::size_t operator() (hpx::threads::thread_id const &v) const
```

```
struct launch : public detail::policy_holder<>
#include <launch_policy.hpp> Launch policies for hpx::async etc.
```

Public Functions

```
constexpr launch()
```

Default constructor. This creates a launch policy representing all possible launch modes

Public Static Attributes

```
const detail::fork_policy fork
```

Predefined launch policy representing asynchronous execution. The new thread is executed in a preferred way

```
const detail::sync_policy sync
```

Predefined launch policy representing synchronous execution.

```
const detail::deferred_policy deferred
```

Predefined launch policy representing deferred execution.

```
const detail::apply_policy apply
Predefined launch policy representing fire and forget execution.
```

```
const detail::select_policy_generator select
Predefined launch policy representing delayed policy selection.
```

```
template<typename ExPolicy = parallel::execution::parallel_policy>
```

```
class task_block
```

```
#include <task_block.hpp> The class task_block defines an interface for forking and joining parallel tasks. The define_task_block and define_task_block_restore_thread function templates create an object of type task_block and pass a reference to that object to a user-provided callable object.
```

An object of class *task_block* cannot be constructed, destroyed, copied, or moved except by the implementation of the task region library. Taking the address of a *task_block* object via operator& or addressof is ill formed. The result of obtaining its address by any other means is unspecified.

A *task_block* is active if it was created by the nearest enclosing task block, where “task block” refers to an invocation of *define_task_block* or *define_task_block_restore_thread* and “nearest

enclosing” means the most recent invocation that has not yet completed. Code designated for execution in another thread by means other than the facilities in this section (e.g., using *thread* or *async*) are not enclosed in the task region and a *task_block* passed to (or captured by) such code is not active within that code. Performing any operation on a *task_block* that is not active results in undefined behavior.

The *task_block* that is active before a specific call to the *run* member function is not active within the asynchronous function that invoked *run*. (The invoked function should not, therefore, capture the *task_block* from the surrounding block.)

Example:

```
define_task_block([&] (auto& tr) {
    tr.run([&] {
        tr.run([] { f(); }); // Error: tr is not active
        define_task_block([&] (auto& tr) {
            tr.run(f); // OK: inner tr is active
            /// ...
        });
        /// ...
    });
    /// ...
});
```

Template Parameters

- *ExPolicy*: The execution policy an instance of a *task_block* was created with. This defaults to *parallel_policy*.

Public Types

```
typedef ExPolicy execution_policy
```

Refers to the type of the execution policy used to create the *task_block*.

Public Functions

`execution_policy const &get_execution_policy() const`
 Return the execution policy instance used to create this `task_block`

`template<typename F, typename ...Ts>`
`void run (F &&f, Ts&&... ts)`

Causes the expression `f()` to be invoked asynchronously. The invocation of `f` is permitted to run on an unspecified thread in an unordered fashion relative to the sequence of operations following the call to `run(f)` (the continuation), or indeterminately sequenced within the same thread as the continuation.

The call to `run` synchronizes with the invocation of `f`. The completion of `f()` synchronizes with the next invocation of `wait` on the same `task_block` or completion of the nearest enclosing task block (i.e., the `define_task_block` or `define_task_block_restore_thread` that created this task block).

Requires: `F` shall be `MoveConstructible`. The expression, `(void)f()`, shall be well-formed.

Precondition: this shall be the active `task_block`.

Postconditions: A call to `run` may return on a different thread than that on which it was called.

Note The call to `run` is sequenced before the continuation as if `run` returns on the same thread. The invocation of the user-supplied callable object `f` may be immediate or may be delayed until compute resources are available. `run` might or might not return before invocation of `f` completes.

Exceptions

- This: function may throw `task_canceled_exception`, as described in Exception Handling.

`template<typename Executor, typename F, typename ...Ts>`
`void run (Executor &exec, F &&f, Ts&&... ts)`

Causes the expression `f()` to be invoked asynchronously using the given executor. The invocation of `f` is permitted to run on an unspecified thread associated with the given executor and in an unordered fashion relative to the sequence of operations following the call to `run(exec, f)` (the continuation), or indeterminately sequenced within the same thread as the continuation.

The call to `run` synchronizes with the invocation of `f`. The completion of `f()` synchronizes with the next invocation of `wait` on the same `task_block` or completion of the nearest enclosing task block (i.e., the `define_task_block` or `define_task_block_restore_thread` that created this task block).

Requires: `Executor` shall be a type modeling the `Executor` concept. `F` shall be `MoveConstructible`. The expression, `(void)f()`, shall be well-formed.

Precondition: this shall be the active `task_block`.

Postconditions: A call to `run` may return on a different thread than that on which it was called.

Note The call to `run` is sequenced before the continuation as if `run` returns on the same thread. The invocation of the user-supplied callable object `f` may be immediate or may be delayed until compute resources are available. `run` might or might not return before invocation of `f` completes.

Exceptions

- This: function may throw `task_canceled_exception`, as described in Exception Handling.

`void wait ()`

Blocks until the tasks spawned using this `task_block` have finished.

Precondition: this shall be the active `task_block`.

Postcondition: All tasks spawned by the nearest enclosing task region have finished. A call to wait may return on a different thread than that on which it was called.

Example:

```
define_task_block([&] (auto& tr) {
    tr.run([&]{ process(a, w, x); });
    if (y < x) tr.wait();
    process(a, y, z);
});
```

Note The call to *wait* is sequenced before the continuation as if *wait* returns on the same thread.

Exceptions

- This function may throw `task_canceled_exception`, as described in Exception Handling.

ExPolicy &policy()

Returns a reference to the execution policy used to construct this object.

Precondition: this shall be the active *task_block*.

ExPolicy const &policy() const

Returns a reference to the execution policy used to construct this object.

Precondition: this shall be the active *task_block*.

Private Members

```
mutex_type mtx_
std::vector<hpx::future<void>> tasks_
parallel::exception_list errors_
threads::thread_id_type id_
ExPolicy policy_

class task_canceled_exception : public exception
#include <task_block.hpp> The class task_canceled_exception defines the type of objects thrown by task_block::run or task_block::wait if they detect that an exception is pending within the current parallel region.
```

Public Functions

```
task_canceled_exception()

class thread_pool_base : public manage_executor
#include <thread_pool_base.hpp> The base class used to manage a pool of OS threads.
```

Public Functions

```
virtual void suspend_processing_unit_direct (std::size_t virt_core, error_code &ec = throws) = 0
```

Suspends the given processing unit. Blocks until the processing unit has been suspended.

Parameters

- `virt_core`: [in] The processing unit on the pool to be suspended. The processing units are indexed starting from 0.

```
virtual void resume_processing_unit_direct (std::size_t virt_core, error_code &ec = throws) = 0
```

Resumes the given processing unit. Blocks until the processing unit has been resumed.

Parameters

- `virt_core`: [in] The processing unit on the pool to be resumed. The processing units are indexed starting from 0.

```
virtual void resume_direct (error_code &ec = throws) = 0
```

Resumes the thread pool. Blocks until all OS threads on the thread pool have been resumed.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
virtual void suspend_direct (error_code &ec = throws) = 0
```

Suspends the thread pool. Blocks until all OS threads on the thread pool have been suspended.

Note A thread pool cannot be suspended from an HPX thread running on the pool itself.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

Exceptions

- `hpx::exception`: if called from an HPX thread which is running on the pool itself.

```
struct thread_pool_init_parameters  
#include <thread_pool_base.hpp>
```

Public Functions

```
thread_pool_init_parameters(std::string const &name, std::size_t index, policies::scheduler_mode mode, std::size_t num_threads, std::size_t thread_offset, hpx::threads::policies::callback_notifier &notifier, hpx::threads::policies::detail::affinity_data const &affinity_data, hpx::threads::detail::network_background_callback_type const &network_background_callback = hpx::threads::detail::network_background_callback_type(),  
std::size_t max_background_threads = std::size_t(-1), std::size_t max_idle_loop_count = HPX_IDLE_LOOP_COUNT_MAX,  
std::size_t max_busy_loop_count = HPX_BUSY_LOOP_COUNT_MAX)
```

Public Members

```
std::string const &name_  
std::size_t index_  
policies::scheduler_mode mode_  
std::size_t num_threads_  
std::size_t thread_offset_  
hpx::threads::policies::callback_notifier &notifier_  
hpx::threads::policies::detail::affinity_data const &affinity_data_  
hpx::threads::detail::network_background_callback_type const &network_background_callback_  
std::size_t max_background_threads_  
std::size_t max_idle_loop_count_  
std::size_t max_busy_loop_count_  
struct unwrap  
#include <unwrap.hpp> A helper function object for functionally invoking hpx::util::unwrap. For more information please refer to its documentation.  
struct unwrap_all  
#include <unwrap.hpp> A helper function object for functionally invoking hpx::util::unwrap_all. For more information please refer to its documentation.  
template<std::size_t Depth> struct unwrap_n  
#include <unwrap.hpp> A helper function object for functionally invoking hpx::util::unwrap_n. For more information please refer to its documentation.  
template<typename Sequence> struct when_any_result  
#include <when_any.hpp> Result type for when_any, contains a sequence of futures and an index pointing to a ready future.
```

Public Members

`std::size_t index`

The index of a future which has become ready.

Sequence **futures**

The sequence of futures as passed to `hpx::when_any`.

`template<typename Sequence>`
`struct when_some_result`

`#include <when_some.hpp>` Result type for `when_some`, contains a sequence of futures and indices pointing to ready futures.

Public Members

`std::vector<std::size_t> indices`

List of indices of futures which became ready.

Sequence **futures**

The sequence of futures as passed to `hpx::when_some`.

`namespace applier`

The namespace `applier` contains all definitions needed for the class `hpx::applier::applier` and its related functionality. This namespace is part of the HPX core module.

`namespace hpx`

Typedefs

`typedef util::unique_function<void ()> shutdown_function_type`

The type of a function which is registered to be executed as a shutdown or pre-shutdown function.

`typedef util::unique_function<void ()> startup_function_type`

The type of a function which is registered to be executed as a startup or pre-startup function.

Enums

`enum runtime_mode`

A HPX runtime can be executed in two different modes: console mode and worker mode.

Values:

`runtime_mode_invalid = -1`

`runtime_mode_console = 0`

The runtime is the console locality.

`runtime_mode_worker = 1`

The runtime is a worker locality.

`runtime_mode_connect = 2`

The runtime is a worker locality connecting late

`runtime_mode_default = 3`

The runtime mode will be determined based on the command line arguments

`runtime_mode_last`

Functions

```
int init (util::function_nonser<int>) hpx::program_options::variables_map &vmm  
> const &f, hpx::program_options::options_description const &desc_cmdline, int argc, char  
**argv, std::vector<std::string> const &cfg, startup_function_type startup = startup_function_type(),  
shutdown_function_type shutdown = shutdown_function_type(), hpx::runtime_mode mode =  
hpx::runtime_mode_defaultMain entry point for launching the HPX runtime system.
```

This is the main entry point for any HPX application. This function (or one of its overloads below) should be called from the users `main()` function. It will set up the HPX runtime environment and schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied `f`.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (int (*f)) hpx::program_options::variables_map &vmm  
>, hpx::program_options::options_description const &desc_cmdline, int argc, char **argv,  
startup_function_type startup = startup_function_type(), shutdown_function_type shutdown = shutdown_function_type(), hpx::runtime_mode mode = hpx::runtime_mode_defaultMain entry point for launching the HPX runtime system.
```

This is the main entry point for any HPX application. This function (or one of its overloads below) should be called from the users `main()` function. It will set up the HPX runtime environment and schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied `f`.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init(hpx::program_options::options_description const &desc_cmdline, int argc, char **argv,
         startup_function_type startup = startup_function_type(), shutdown_function_type shutdown =
         shutdown_function_type(), hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).

- argv: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- startup: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- shutdown: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- mode: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init(hpx::program_options::options_description const &desc_cmdline, int argc, char **argv,  
         std::vector<std::string> const &cfg, startup_function_type startup = startup_function_type(),  
         shutdown_function_type shutdown = shutdown_function_type(), hpx::runtime_mode mode =  
         hpx::runtime_mode_default)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- desc_cmdline: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- argc: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- argv: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- cfg: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- startup: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- shutdown: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- mode: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (int argc, char **argv, std::vector<std::string> const &cfg, hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (hpx::program_options::options_description const &desc_cmdline, int argc, char **argv, hpx::runtime_mode mode)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note If the parameter `mode` is `runtime_mode_default`, the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).

- mode: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (hpx::program_options::options_description const &desc_cmdline, int argc, char **argv,  
          std::vector<std::string> const &cfg, hpx::runtime_mode mode)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note If the parameter `mode` is `runtime_mode_default`, the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (std::string const &app_name, int argc = 0, char **argv = nullptr, hpx::runtime_mode mode =  
          hpx::runtime_mode_default)
```

Main entry point for launching the HPX runtime system.

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (int argc = 0, char **argv = nullptr, hpx::runtime_mode mode = hpx::runtime_mode_default)  
Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. If no command line arguments are passed, console mode is assumed.

Note If no command line arguments are passed the HPX runtime system will not support any of the default command line options as described in the section ‘HPX Command Line Options’.

Parameters

- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (std::vector<std::string> const &cfg, hpx::runtime_mode mode =  
hpx::runtime_mode_default)  
Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings).

Return The function returns the value, which has been returned from `hpx_main` (or 0 when executed in worker mode).

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. If no command line arguments are passed, console mode is assumed.

Note If no command line arguments are passed the HPX runtime system will not support any of the default command line options as described in the section ‘HPX Command Line Options’.

Parameters

- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (int (*f) hpx::program_options::variables_map &vm  
, std::string const &app_name, int argc, char **argv, hpx::runtime_mode mode =  
hpx::runtime_mode_default)Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (int (*f) hpx::program_options::variables_map &vm  
, int argc, char **argv, hpx::runtime_mode mode = hpx::runtime_mode_default)Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (util::function_nonscoped<int> int, char **  
> const &f, std::string const &app_name, int argc, char **argv, hpx::runtime_mode mode = hpx::runtime_mode_default Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init (util::function_nonscoped<int> int, char **  
> const &f, int argc, char **argv, hpx::runtime_mode mode = hpx::runtime_mode_default Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init(util::function_nonser<int> int, char **  
> const &f, int argc, char **argv, std::vector<std::string> const &cfg, hpx::runtime_mode mode =  
hpx::runtime_mode_default Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int init(util::function_nonser<int> int, char **  
> const &f, std::vector<std::string> const &cfg, hpx::runtime_mode mode =  
hpx::runtime_mode_default Main entry point for launching the HPX runtime system.
```

This is a simplified main entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will schedule the function given by `f` as a HPX thread. This overload will not call `hpx_main`.

Return The function returns the value, which has been returned from the user supplied function `f`.

Note The created runtime system instance will be executed in console or worker mode depending on the configuration passed in `cfg`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (util::function_nonser<int> hpx::program_options::variables_map &vm
> const &f, hpx::program_options::options_description const &desc_cmdline, int argc, char **argv, std::vector<std::string> const &cfg, startup_function_type startup = startup_function_type(), shutdown_function_type shutdown = shutdown_function_type(), hpx::runtime_mode mode = hpx::runtime_mode_default)Main non-blocking entry point for launching the HPX runtime system.
```

This is the main, non-blocking entry point for any HPX application. This function (or one of its overloads below) should be called from the users `main()` function. It will set up the HPX runtime environment and schedule the function given by `f` as a HPX thread. It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system’s execution. This overload will not call `hpx_main`.

Return The function returns `true` if command line processing succeeded and the runtime system was started successfully. It will return `false` otherwise.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).

- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (int (*f)) hpx::program_options::variables_map &vmm  
    , hpx::program_options::options_description const &desc cmdline, int argc, char **argv,  
    startup_function_type startup = startup_function_type(), shutdown_function_type shutdown = shutdown_function_type(), hpx::runtime_mode mode = hpx::runtime_mode_default  
Main non-blocking entry point for launching the HPX runtime system.
```

This is the main, non-blocking entry point for any HPX application. This function (or one of its overloads below) should be called from the users `main()` function. It will set up the HPX runtime environment and schedule the function given by `f` as a HPX thread. It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system’s execution. This overload will not call `hpx_main`.

Return The function returns `true` if command line processing succeeded and the runtime system was started successfully. It will return `false` otherwise.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `desc cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode

(`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (hpx::program_options::options_description const &desc cmdline, int argc, char **argv,  
           startup_function_type startup = startup_function_type(), shutdown_function_type shutdown  
           = shutdown_function_type(), hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx:::wait` and `hpx:::stop` to synchronize with the runtime system's execution.

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (hpx::program_options::options_description const &desc cmdline, int argc, char  
           **argv, std::vector<std::string> const &cfg, startup_function_type startup =  
           startup_function_type(), shutdown_function_type shutdown = shutdown_function_type(),  
           hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx:::wait` and `hpx:::stop` to synchronize with the runtime system's execution.

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note If the parameter `mode` is not given (defaulted), the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `startup`: [in] A function to be executed inside a HPX thread before `f` is called. If this parameter is not given no function will be executed.
- `shutdown`: [in] A function to be executed inside an HPX thread while `hpx::finalize` is executed. If this parameter is not given no function will be executed.
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start(int argc, char **argv, std::vector<std::string> const &cfg, hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system’s execution.

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note If the parameter `mode` is `runtime_mode_default`, the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).

- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (hpx::program_options::options_description const &desc cmdline, int argc, char **argv,  
           hpx::runtime_mode mode)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx:::wait` and `hpx:::stop` to synchronize with the runtime system’s execution.

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note If the parameter `mode` is `runtime_mode_default`, the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (hpx::program_options::options_description const &desc cmdline, int argc, char **argv,  
           std::vector<std::string> const &cfg, hpx::runtime_mode mode)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx:::wait` and `hpx:::stop` to synchronize with the runtime system’s execution.

In console mode it will execute the user supplied function `hpx_main`, in worker mode it will execute an empty `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note If the parameter `mode` is `runtime_mode_default`, the created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. Otherwise it will be executed as specified by the parameter `mode`.

Parameters

- `desc_cmdline`: [in] This parameter may hold the description of additional command line arguments understood by the application. These options will be prepended to the default command line options understood by `hpx::init` (see description below).
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (std::string const &app_name, int argc = 0, char **argv = nullptr, hpx::runtime_mode mode  
           = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system’s execution.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (int argc = 0, char **argv = nullptr, hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. If no command line arguments are passed, console mode is assumed.

Note If no command line arguments are passed the HPX runtime system will not support any of the default command line options as described in the section 'HPX Command Line Options'.

Parameters

- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (std::vector<std::string> const &cfg, hpx::runtime_mode mode = hpx::runtime_mode_default)
```

Main non-blocking entry point for launching the HPX runtime system.

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`. If no command line arguments are passed, console mode is assumed.

Note If no command line arguments are passed the HPX runtime system will not support any of the default command line options as described in the section 'HPX Command Line Options'.

Parameters

- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance 'hpx.component.enabled=1')
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode

(`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (int (*f)) hpx::program_options::variables_map &vm  
, std::string const &app_name, int argc, char **argv, hpx::runtime_mode mode =  
hpx::runtime_mode_defaultMain non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by `f` as a HPX thread. It will not call `hpx_main`.

Return The function returns `true` if command line processing succeeded and the runtime system was started successfully. It will return `false` otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application.
- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (util::function_nonser<int> f, int argc, char **argv,  
> const &f, std::string const &app_name, int argc, char **argv, hpx::runtime_mode mode =  
hpx::runtime_mode_defaultMain non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by `f` as a HPX thread. It will not call `hpx_main`.

Return The function returns `true` if command line processing succeeded and the runtime system was started successfully. It will return `false` otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.

- `app_name`: [in] The name of the application.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (int (*f) hpx::program_options::variables_map &vm, int argc, char **argv, hpx::runtime_mode mode = hpx::runtime_mode_default)  
Main non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by *f* as a HPX thread. It will not call `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If *f* is `nullptr` the HPX runtime environment will be started without invoking *f*.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (util::function_nonser<int> int, char **  
> const &f, int argc, char **argv, hpx::runtime_mode mode = hpx::runtime_mode_default)  
Main non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by *f* as a HPX thread. It will not call `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (util::function_nonser<int> int, char **  
> const &f, int argc, char **argv, std::vector<std::string> const &cfg, hpx::runtime_mode mode =  
hpx::runtime_mode_default Main non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by `f` as a HPX thread. It will not call `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the command line arguments passed in `argc/argv`.

Parameters

- `f`: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If `f` is `nullptr` the HPX runtime environment will be started without invoking `f`.
- `argc`: [in] The number of command line arguments passed in `argv`. This is usually the unchanged value as passed by the operating system (to `main()`).
- `argv`: [in] The command line arguments for this application, usually that is the value as passed by the operating system (to `main()`).
- `cfg`: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- `mode`: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
bool start (util::function_nonser<int> int, char **
    > const &f, std::vector<std::string> const &cfg, hpx::runtime_mode mode =
hpx::runtime_mode_default) Main non-blocking entry point for launching the HPX runtime system.
```

This is a simplified main, non-blocking entry point, which can be used to set up the runtime for an HPX application (the runtime system will be set up in console mode or worker mode depending on the command line settings). It will return immediately after that. Use `hpx::wait` and `hpx::stop` to synchronize with the runtime system's execution. This overload will schedule the function given by *f* as a HPX thread. It will not call `hpx_main`.

Return The function returns *true* if command line processing succeeded and the runtime system was started successfully. It will return *false* otherwise.

Note The created runtime system instance will be executed in console or worker mode depending on the configuration passed in *cfg*.

Parameters

- *f*: [in] The function to be scheduled as an HPX thread. Usually this function represents the main entry point of any HPX application. If *f* is `nullptr` the HPX runtime environment will be started without invoking *f*.
- *cfg*: A list of configuration settings which will be added to the system configuration before the runtime instance is run. Each of the entries in this list must have the format of a fully defined key/value pair from an ini-file (for instance ‘`hpx.component.enabled=1`’)
- *mode*: [in] The mode the created runtime environment should be initialized in. There has to be exactly one locality in each HPX application which is executed in console mode (`hpx::runtime_mode_console`), all other localities have to be run in worker mode (`hpx::runtime_mode_worker`). Normally this is set up automatically, but sometimes it is necessary to explicitly specify the mode.

```
int finalize (double shutdown_timeout, double localwait = -1.0, error_code &ec = throws)
Main function to gracefully terminate the HPX runtime system.
```

The function `hpx::finalize` is the main way to (gracefully) exit any HPX application. It should be called from one locality only (usually the console) and it will notify all connected localities to finish execution. Only after all other localities have exited this function will return, allowing to exit the console locality as well.

During the execution of this function the runtime system will invoke all registered shutdown functions (see `hpx::init`) on all localities.

The default value (-1.0) will try to find a globally set timeout value (can be set as the configuration parameter `hpx.shutdown_timeout`), and if that is not set or -1.0 as well, it will disable any timeout, each connected locality will wait for all existing HPX-threads to terminate.

Parameters

- `shutdown_timeout`: This parameter allows to specify a timeout (in microseconds), specifying how long any of the connected localities should wait for pending tasks to be executed. After this timeout, all suspended HPX-threads will be aborted. Note, that this function will not abort any running HPX-threads. In any case the shutdown will not proceed as long as there is at least one pending/running HPX-thread.

The default value (-1.0) will try to find a globally set wait time value (can be set as the configuration parameter “`hpx.finalize_wait_time`”), and if this is not set or -1.0 as well, it will disable any addition local wait time before proceeding.

Parameters

- `localwait`: This parameter allows to specify a local wait time (in microseconds) before the connected localities will be notified and the overall shutdown process starts.

This function will block and wait for all connected localities to exit before returning to the caller. It should be the last HPX-function called by any application.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

Using this function is an alternative to `hpx::disconnect`, these functions do not need to be called both.

`int finalize(error_code &ec = throws)`

Main function to gracefully terminate the HPX runtime system.

The function `hpx::finalize` is the main way to (gracefully) exit any HPX application. It should be called from one locality only (usually the console) and it will notify all connected localities to finish execution. Only after all other localities have exited this function will return, allowing to exit the console locality as well.

During the execution of this function the runtime system will invoke all registered shutdown functions (see `hpx::init`) on all localities.

This function will block and wait for all connected localities to exit before returning to the caller. It should be the last HPX-function called by any application.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

Using this function is an alternative to `hpx::disconnect`, these functions do not need to be called both.

`HPX_NORETURN void hpx::terminate()`

Terminate any application non-gracefully.

The function `hpx::terminate` is the non-graceful way to exit any application immediately. It can be called from any locality and will terminate all localities currently used by the application.

Note This function will cause HPX to call `std::terminate()` on all localities associated with this application. If the function is called not from an HPX thread it will fail and return an error using the argument `ec`.

`int disconnect(double shutdown_timeout, double localwait = -1.0, error_code &ec = throws)`

Disconnect this locality from the application.

The function `hpx::disconnect` can be used to disconnect a locality from a running HPX application.

During the execution of this function the runtime system will invoke all registered shutdown functions (see `hpx::init`) on this locality. The default value (-1.0) will try to find a globally set timeout value (can be set as the configuration parameter "hpx.shutdown_timeout"), and if that is not set or -1.0 as well, it will disable any timeout, each connected locality will wait for all existing HPX-threads to terminate.

Parameters

- `shutdown_timeout`: This parameter allows to specify a timeout (in microseconds), specifying how long this locality should wait for pending tasks to be executed. After this timeout, all suspended HPX-threads will be aborted. Note, that this function will not abort any running HPX-threads. In any case the shutdown will not proceed as long as there is at least one pending/running HPX-thread.

The default value (`-1.0`) will try to find a globally set wait time value (can be set as the configuration parameter `hpx.finalize_wait_time`), and if this is not set or `-1.0` as well, it will disable any addition local wait time before proceeding.

Parameters

- `localwait`: This parameter allows to specify a local wait time (in microseconds) before the connected localities will be notified and the overall shutdown process starts.

This function will block and wait for this locality to finish executing before returning to the caller. It should be the last HPX-function called by any locality being disconnected.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

int **disconnect** (*error_code &ec = throws*)

Disconnect this locality from the application.

The function `hpx::disconnect` can be used to disconnect a locality from a running HPX application.

During the execution of this function the runtime system will invoke all registered shutdown functions (see `hpx::init`) on this locality.

This function will block and wait for this locality to finish executing before returning to the caller. It should be the last HPX-function called by any locality being disconnected.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

int **stop** (*error_code &ec = throws*)

Stop the runtime system.

This function will block and wait for this locality to finish executing before returning to the caller. It should be the last HPX-function called on every locality. This function should be used only if the runtime system was started using `hpx::start`.

Return The function returns the value, which has been returned from the user supplied main HPX function (usually `hpx_main`).

int **suspend** (*error_code &ec = throws*)

Suspend the runtime system.

The function `hpx::suspend` is used to suspend the HPX runtime system. It can only be used when running HPX on a single locality. It will block waiting for all thread pools to be empty. This function only be called when the runtime is running, or already suspended in which case this function will do nothing.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
int resume (error_code &ec = throws)  
    Resume the HPX runtime system.
```

The function `hpx::resume` is used to resume the HPX runtime system. It can only be used when running HPX on a single locality. It will block waiting for all thread pools to be resumed. This function only be called when the runtime suspended, or already running in which case this function will do nothing.

Return This function will always return zero.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
bool register_thread (runtime *rt, char const *name, error_code &ec = throws)  
    Register the current kernel thread with HPX, this should be done once for each external OS-thread intended  
    to invoke HPX functionality. Calling this function more than once will silently fail.
```

```
void unregister_thread (runtime *rt)  
    Unregister the thread from HPX, this should be done once in the end before the external thread exists.
```

```
naming::gid_type const &get_locality ()  
    The function get_locality returns a reference to the locality prefix.
```

```
std::size_t get_runtime_instance_number ()  
    The function get_runtime_instance_number returns a unique number associated with the runtime instance  
    the current thread is running in.
```

```
bool register_on_exit (util::function_nonser<void>  
    > const&Register a function to be called during system shutdown.
```

```
bool is_starting ()  
    Test whether the runtime system is currently being started.
```

This function returns whether the runtime system is currently being started or not, e.g. whether the current state of the runtime system is `hpx::state_startup`

Note This function needs to be executed on a HPX-thread. It will return false otherwise.

```
bool tolerate_node_faults()
```

Test if HPX runs in fault-tolerant mode.

This function returns whether the runtime system is running in fault-tolerant mode

```
bool is_running()
```

Test whether the runtime system is currently running.

This function returns whether the runtime system is currently running or not, e.g. whether the current state of the runtime system is *hpx::state_running*

Note This function needs to be executed on a HPX-thread. It will return false otherwise.

```
bool is_stopped()
```

Test whether the runtime system is currently stopped.

This function returns whether the runtime system is currently stopped or not, e.g. whether the current state of the runtime system is *hpx::state_stopped*

Note This function needs to be executed on a HPX-thread. It will return false otherwise.

```
bool is_stopped_or_shutting_down()
```

Test whether the runtime system is currently being shut down.

This function returns whether the runtime system is currently being shut down or not, e.g. whether the current state of the runtime system is *hpx::state_stopped* or *hpx::state_shutdown*

Note This function needs to be executed on a HPX-thread. It will return false otherwise.

```
std::size_t get_num_worker_threads()
```

Return the number of worker OS- threads used to execute HPX threads.

This function returns the number of OS-threads used to execute HPX threads. If the function is called while no HPX runtime system is active, it will return zero.

```
std::uint64_t get_system_uptime()
```

Return the system uptime measure on the thread executing this call.

This function returns the system uptime measured in nanoseconds for the thread executing this call. If the function is called while no HPX runtime system is active, it will return zero.

```
void start_active_counters(error_code &ec = throws)
```

Start all active performance counters, optionally naming the section of code.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note The active counters are those which have been specified on the command line while executing the application (see command line option `-hpx:print-counter`)

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void reset_active_counters(error_code &ec = throws)
```

Resets all active performance counters.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note The active counters are those which have been specified on the command line while executing the application (see command line option `-hpx:print-counter`)

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void reinit_active_counters (bool reset = true, error_code &ec = throws)
```

Re-initialize all active performance counters.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note The active counters are those which have been specified on the command line while executing the application (see command line option `-hpx:print-counter`)

Parameters

- *reset*: [in] Reset the current values before re-initializing counters (default: true)
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void stop_active_counters (error_code &ec = throws)
```

Stop all active performance counters.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note The active counters are those which have been specified on the command line while executing the application (see command line option `-hpx:print-counter`)

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void evaluate_active_counters (bool reset = false, char const *description = nullptr, error_code &ec = throws)
```

Evaluate and output all active performance counters, optionally naming the point in code marked by this function.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note The output generated by this function is redirected to the destination specified by the corresponding command line options (see `-hpx:print-counter-destination`).

Note The active counters are those which have been specified on the command line while executing the application (see command line option `-hpx:print-counter`)

Parameters

- *reset*: [in] this is an optional flag allowing to reset the counter value after it has been evaluated.
- *description*: [in] this is an optional value naming the point in the code marked by the call to this function.

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
serialization::binary_filter *create_binary_filter(char const *binary_filter_type, bool compress,
                                                serialization::binary_filter *next_filter
                                                = nullptr, error_code &ec = throws)
```

Create an instance of a binary filter plugin.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `binary_filter_type`: [in] The type of the binary filter to create
- `compress`: [in] The created filter should support compression
- `next_filter`: [in] Use this as the filter to dispatch the invocation into.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
std::vector<Client> find_all_from_basename (std::string base_name, std::size_t num_ids)
```

Return all registered ids from all localities from the given base name.

This function locates all ids which were registered with the given base name. It returns a list of futures representing those ids.

Return all registered clients from all localities from the given base name.

Return A list of futures representing the ids which were registered using the given base name.

Note The futures will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Parameters

- `base_name`: [in] The base name for which to retrieve the registered ids.
- `num_ids`: [in] The number of registered ids to expect.

This function locates all ids which were registered with the given base name. It returns a list of futures representing those ids.

Return A list of futures representing the ids which were registered using the given base name.

Note The futures embedded in the returned client objects will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Template Parameters

- `Client`: The client type to return

Parameters

- `base_name`: [in] The base name for which to retrieve the registered ids.
- `num_ids`: [in] The number of registered ids to expect.

```
std::vector<Client> find_from_basename (std::string base_name, std::vector<std::size_t> const
&ids)
```

Return registered ids from the given base name and sequence numbers.

This function locates the ids which were registered with the given base name and the given sequence numbers. It returns a list of futures representing those ids.

Return registered clients from the given base name and sequence numbers.

Return A list of futures representing the ids which were registered using the given base name and sequence numbers.

Note The futures will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Parameters

- *base_name*: [in] The base name for which to retrieve the registered ids.
- *ids*: [in] The sequence numbers of the registered ids.

This function locates the ids which were registered with the given base name and the given sequence numbers. It returns a list of futures representing those ids.

Return A list of futures representing the ids which were registered using the given base name and sequence numbers.

Note The futures embedded in the returned client objects will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Template Parameters

- Client: The client type to return

Parameters

- *base_name*: [in] The base name for which to retrieve the registered ids.
- *ids*: [in] The sequence numbers of the registered ids.

```
Client find_from_basename (std::string base_name, std::size_t sequence_nr =  
~static_cast<std::size_t>(0))
```

Return registered id from the given base name and sequence number.

This function locates the id which was registered with the given base name and the given sequence number. It returns a future representing those id.

This function locates the id which was registered with the given base name and the given sequence number. It returns a future representing those id.

Return A representing the id which was registered using the given base name and sequence numbers.

Note The future will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Parameters

- *base_name*: [in] The base name for which to retrieve the registered ids.
- *sequence_nr*: [in] The sequence number of the registered id.

Return A representing the id which was registered using the given base name and sequence numbers.

Note The future embedded in the returned client object will become ready even if the event (for instance, binding the name to an id) has already happened in the past. This is important in order to reliably retrieve ids from a name, even if the name was already registered.

Template Parameters

- Client: The client type to return

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- sequence_nr: [in] The sequence number of the registered id.

```
hpx::future<bool> register_with_basename (std::string base_name, hpx::id_type id, std::size_t  
sequence_nr = ~static_cast<std::size_t>(0))
```

Register the given id using the given base name.

The function registers the given ids using the provided base name.

Return A future representing the result of the registration operation itself.

Note The operation will fail if the given sequence number is not unique.

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- id: [in] The id to register using the given base name.
- sequence_nr: [in, optional] The sequential number to use for the registration of the id. This number has to be unique system wide for each registration using the same base name. The default is the current locality identifier. Also, the sequence numbers have to be consecutive starting from zero.

```
hpx::future<bool> register_with_basename (std::string base_name, hpx::future<hpx::id_type>  
f, std::size_t sequence_nr = ~static_cast<std::size_t>(0))
```

Register the id wrapped in the given future using the given base name.

The function registers the object the given future refers to using the provided base name.

Return A future representing the result of the registration operation itself.

Note The operation will fail if the given sequence number is not unique.

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- f: [in] The future which should be registered using the given base name.
- sequence_nr: [in, optional] The sequential number to use for the registration of the id. This number has to be unique system wide for each registration using the same base name. The default is the current locality identifier. Also, the sequence numbers have to be consecutive starting from zero.

```
template<typename Client, typename Stub>  
hpx::future<bool> register_with_basename (std::string base_name, components::client_base<Client, Stub> &client, std::size_t  
sequence_nr = ~static_cast<std::size_t>(0))
```

Register the id wrapped in the given client using the given base name.

The function registers the object the given client refers to using the provided base name.

Return A future representing the result of the registration operation itself.

Note The operation will fail if the given sequence number is not unique.

Template Parameters

- Client: The client type to register

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- client: [in] The client which should be registered using the given base name.
- sequence_nr: [in, optional] The sequential number to use for the registration of the id. This number has to be unique system wide for each registration using the same base name. The default is the current locality identifier. Also, the sequence numbers have to be consecutive starting from zero.

Client **unregister_with_basename** (*std::string* *base_name*, *std::size_t* *sequence_nr* = ~*static_cast<std::size_t>(0)*)

Unregister the given id using the given base name.

The function unregisters the given ids using the provided base name.

Unregister the given base name.

Return A future representing the result of the un-registration operation itself.

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- sequence_nr: [in, optional] The sequential number to use for the un-registration. This number has to be the same as has been used with *register_with_basename* before.

The function unregisters the given ids using the provided base name.

Return A future representing the result of the un-registration operation itself.

Template Parameters

- Client: The client type to return

Parameters

- base_name: [in] The base name for which to retrieve the registered ids.
- sequence_nr: [in, optional] The sequential number to use for the un-registration. This number has to be the same as has been used with *register_with_basename* before.

naming::id_type find_here (error_code &ec = throws)

Return the global id representing this locality.

The function *find_here()* can be used to retrieve the global id usable to refer to the current locality.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global id representing the locality this function has been called on.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note This function will return meaningful results only if called from an HPX-thread. It will return `hpx::naming::invalid_id` otherwise.

See [hpx::find_all_localities\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`naming::id_type find_root_locality(error_code &ec = throws)`

Return the global id representing the root locality.

The function `find_root_locality()` can be used to retrieve the global id usable to refer to the root locality. The root locality is the locality where the main AGAS service is hosted.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global id representing the root locality for this application.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note This function will return meaningful results only if called from an HPX-thread. It will return `hpx::naming::invalid_id` otherwise.

See [hpx::find_all_localities\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::vector<naming::id_type> find_all_localities(error_code &ec = throws)`

Return the list of global ids representing all localities available to this application.

The function `find_all_localities()` can be used to retrieve the global ids of all localities currently available to this application.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global ids representing the localities currently available to this application.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note This function will return meaningful results only if called from an HPX-thread. It will return an empty vector otherwise.

See [hpx::find_here\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::vector<naming::id_type> find_all_localities(components::component_type type, error_code &ec = throws)`

Return the list of global ids representing all localities available to this application which support the given component type.

The function `find_all_localities()` can be used to retrieve the global ids of all localities currently available to this application which support the creation of instances of the given component type.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global ids representing the localities currently available to this application which support the creation of instances of the given component type. If no localities supporting the given component type are currently available, this function will return an empty vector.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note This function will return meaningful results only if called from an HPX-thread. It will return an empty vector otherwise.

See [hpx::find_here\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- `type`: [in] The type of the components for which the function should return the available localities.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::vector<naming::id_type> find_remote_localities(error_code &ec = throws)`

Return the list of locality ids of remote localities supporting the given component type. By default this function will return the list of all remote localities (all but the current locality).

The function `find_remote_localities()` can be used to retrieve the global ids of all remote localities currently available to this application (i.e. all localities except the current one).

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global ids representing the remote localities currently available to this application.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note This function will return meaningful results only if called from an HPX-thread. It will return an empty vector otherwise.

See [hpx::find_here\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::vector<naming::id_type> find_remote_localities(components::component_type type, error_code &ec = throws)`

Return the list of locality ids of remote localities supporting the given component type. By default this function will return the list of all remote localities (all but the current locality).

The function `find_remote_localities()` can be used to retrieve the global ids of all remote localities currently available to this application (i.e. all localities except the current one) which support the creation of instances of the given component type.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global ids representing the remote localities currently available to this application.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note This function will return meaningful results only if called from an HPX-thread. It will return an empty vector otherwise.

See [hpx::find_here\(\)](#), [hpx::find_locality\(\)](#)

Parameters

- *type*: [in] The type of the components for which the function should return the available remote localities.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

*naming::id_type find_locality(*components*::component_type *type*, *error_code* &*ec* = *throws*)*

Return the global id representing an arbitrary locality which supports the given component type.

The function [find_locality\(\)](#) can be used to retrieve the global id of an arbitrary locality currently available to this application which supports the creation of instances of the given component type.

Note Generally, the id of a locality can be used for instance to create new instances of components and to invoke plain actions (global functions).

Return The global id representing an arbitrary locality currently available to this application which supports the creation of instances of the given component type. If no locality supporting the given component type is currently available, this function will return *hpx::naming::invalid_id*.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Note This function will return meaningful results only if called from an HPX-thread. It will return *hpx::naming::invalid_id* otherwise.

See [hpx::find_here\(\)](#), [hpx::find_all_localities\(\)](#)

Parameters

- *type*: [in] The type of the components for which the function should return any available locality.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

*naming::id_type get_colocation_id(*launch*::sync_policy, *naming::id_type const &id*, *error_code* &*ec* = *throws*)*

Return the id of the locality where the object referenced by the given id is currently located on.

The function [hpx::get_colocation_id\(\)](#) returns the id of the locality where the given object is currently located.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

See [hpx::get_colocation_id\(\)](#)

Parameters

- `id`: [in] The id of the object to locate.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`lcos::future<naming::id_type> get_colocation_id (naming::id_type const &id)`

Asynchronously return the id of the locality where the object referenced by the given id is currently located on.

See `hpx::get_colocation_id(launch::sync_policy)`

Parameters

- `id`: [in] The id of the object to locate.

`template<typename Component>`

`hpx::future<std::shared_ptr<Component>> get_ptr (naming::id_type const &id)`

Returns a future referring to the pointer to the underlying memory of a component.

The function `hpx::get_ptr` can be used to extract a future referring to the pointer to the underlying memory of a given component.

Return This function returns a future representing the pointer to the underlying memory for the component instance with the given `id`.

Note This function will successfully return the requested result only if the given component is currently located on the calling locality. Otherwise the function will raise an error.

Note The component instance the returned pointer refers to can not be migrated as long as there is at least one copy of the returned `shared_ptr` alive.

Parameters

- `id`: [in] The global id of the component for which the pointer to the underlying memory should be retrieved.

Template Parameters

- `The`: only template parameter has to be the type of the server side component.

`template<typename Derived, typename Stub>`

`hpx::future<std::shared_ptr<typename components::client_base<Derived, Stub>::server_component_type>> get_ptr (compo`

`Stu`

`cons`

`&c)`

Returns a future referring to the pointer to the underlying memory of a component.

The function `hpx::get_ptr` can be used to extract a future referring to the pointer to the underlying memory of a given component.

Return This function returns a future representing the pointer to the underlying memory for the component instance with the given `id`.

Note This function will successfully return the requested result only if the given component is currently located on the calling locality. Otherwise the function will raise an error.

Note The component instance the returned pointer refers to can not be migrated as long as there is at least one copy of the returned `shared_ptr` alive.

Parameters

- `c`: [in] A client side representation of the component for which the pointer to the underlying memory should be retrieved.

```
template<typename Component>
std::shared_ptr<Component> get_ptr(launch::sync_policy p, naming::id_type const &id, error_code &ec = throws)
```

Returns the pointer to the underlying memory of a component.

The function `hpx::get_ptr_sync` can be used to extract the pointer to the underlying memory of a given component.

Return This function returns the pointer to the underlying memory for the component instance with the given `id`.

Note This function will successfully return the requested result only if the given component is currently located on the requesting locality. Otherwise the function will raise and error.

Note The component instance the returned pointer refers to can not be migrated as long as there is at least one copy of the returned `shared_ptr` alive.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `p`: [in] The parameter `p` represents a placeholder type to turn make the call synchronous.
- `id`: [in] The global id of the component for which the pointer to the underlying memory should be retrieved.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

Template Parameters

- `The`: only template parameter has to be the type of the server side component.

```
template<typename Derived, typename Stub>
std::shared_ptr<typename components::client_base<Derived, Stub>::server_component_type> get_ptr(launch::sync_policy
p,
com-
po-
nents::client_base<
Stub>
const
&c,
er-
ror_code
&ec
=
throws)
```

Returns the pointer to the underlying memory of a component.

The function `hpx::get_ptr_sync` can be used to extract the pointer to the underlying memory of a given component.

Return This function returns the pointer to the underlying memory for the component instance with the given `id`.

Note This function will successfully return the requested result only if the given component is currently located on the requesting locality. Otherwise the function will raise an error.

Note The component instance the returned pointer refers to can not be migrated as long as there is at least one copy of the returned shared_ptr alive.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of hpx::exception.

Parameters

- *p*: [in] The parameter *p* represents a placeholder type to turn make the call synchronous.
- *c*: [in] A client side representation of the component for which the pointer to the underlying memory should be retrieved.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`std::uint32_t get_locality_id(error_code &ec = throws)`

Return the number of the locality this function is being called from.

This function returns the id of the current locality.

Note The returned value is zero based and its maximum value is smaller than the overall number of localities the current application is running on (as returned by [get_num_localities\(\)](#)).

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of hpx::exception.

Note This function needs to be executed on a HPX-thread. It will fail otherwise (it will return -1).

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`std::string get_locality_name()`

Return the name of the locality this function is called on.

This function returns the name for the locality on which this function is called.

Return This function returns the name for the locality on which the function is called. The name is retrieved from the underlying networking layer and may be different for different parcel ports.

See `future<std::string> get_locality_name(naming::id_type const& id)`

`future<std::string> get_locality_name(naming::id_type const &id)`

Return the name of the referenced locality.

This function returns a future referring to the name for the locality of the given id.

Return This function returns the name for the locality of the given id. The name is retrieved from the underlying networking layer and may be different for different parcel ports.

See `std::string get_locality_name()`

Parameters

- *id*: [in] The global id of the locality for which the name should be retrieved

`std::uint32_t get_initial_num_localities()`

Return the number of localities which were registered at startup for the running application.

The function `get_initial_num_localities` returns the number of localities which were connected to the console at application startup.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

See [hpx::find_all_localities](#), [hpx::get_num_localities](#)

`lcos::future<std::uint32_t> get_num_localities()`

Asynchronously return the number of localities which are currently registered for the running application.

The function `get_num_localities` asynchronously returns the number of localities currently connected to the console. The returned future represents the actual result.

Note This function will return meaningful results only if called from an HPX-thread. It will return 0 otherwise.

See [hpx::find_all_localities](#), [hpx::get_num_localities](#)

`std::uint32_t get_num_localities(launch::sync_policy, error_code &ec = throws)`

Return the number of localities which are currently registered for the running application.

The function `get_num_localities` returns the number of localities currently connected to the console.

Note This function will return meaningful results only if called from an HPX-thread. It will return 0 otherwise.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

See [hpx::find_all_localities](#), [hpx::get_num_localities](#)

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`lcos::future<std::uint32_t> get_num_localities(components::component_type t)`

Asynchronously return the number of localities which are currently registered for the running application.

The function `get_num_localities` asynchronously returns the number of localities currently connected to the console which support the creation of the given component type. The returned future represents the actual result.

Note This function will return meaningful results only if called from an HPX-thread. It will return 0 otherwise.

See [hpx::find_all_localities](#), [hpx::get_num_localities](#)

Parameters

- `t`: The component type for which the number of connected localities should be retrieved.

`std::uint32_t get_num_localities(launch::sync_policy, components::component_type t, error_code &ec = throws)`

Synchronously return the number of localities which are currently registered for the running application.

The function `get_num_localities` returns the number of localities currently connected to the console which support the creation of the given component type. The returned future represents the actual result.

Note This function will return meaningful results only if called from an HPX-thread. It will return 0 otherwise.

See [hpx::find_all_localities](#), [hpx::get_num_localities](#)

Parameters

- `t`: The component type for which the number of connected localities should be retrieved.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::size_t get_os_thread_count ()`

Return the number of OS-threads running in the runtime instance the current HPX-thread is associated with.

`std::size_t get_os_thread_count (threads::executor const &exec)`

Return the number of worker OS- threads used by the given executor to execute HPX threads.

This function returns the number of cores used to execute HPX threads for the given executor. If the function is called while no HPX runtime system is active, it will return zero. If the executor is not valid, this function will fall back to retrieving the number of OS threads used by HPX.

Parameters

- `exec`: [in] The executor to be used.

`std::string get_thread_name ()`

Return the name of the calling thread.

This function returns the name of the calling thread. This name uniquely identifies the thread in the context of HPX. If the function is called while no HPX runtime system is active, the result will be “<unknown>”.

`std::size_t get_worker_thread_num ()`

Return the number of the current OS-thread running in the runtime instance the current HPX-thread is executed with.

This function returns the zero based index of the OS-thread which executes the current HPX-thread.

Note The returned value is zero based and its maximum value is smaller than the overall number of OS-threads executed (as returned by `get_os_thread_count()`).

Note This function needs to be executed on a HPX-thread. It will fail otherwise (it will return -1).

`std::size_t get_worker_thread_num (error_code &ec)`

Return the number of the current OS-thread running in the runtime instance the current HPX-thread is executed with.

This function returns the zero based index of the OS-thread which executes the current HPX-thread.

Note The returned value is zero based and its maximum value is smaller than the overall number of OS-threads executed (as returned by `get_os_thread_count()`). It will return -1 if the current thread is not a known thread or if the runtime is not in running state.

Note This function needs to be executed on a HPX-thread. It will fail otherwise (it will return -1).

Parameters

- `ec`: [in,out] this represents the error status on exit.

`void report_error (std::size_t num_thread, std::exception_ptr const &e)`

The function `report_error` reports the given exception to the console.

`void report_error (std::exception_ptr const &e)`

The function `report_error` reports the given exception to the console.

`char const *get_runtime_mode_name (runtime_mode state)`

Get the readable string representing the name of the given `runtime_mode` constant.

`runtime_mode get_runtime_mode_from_name (std::string const &mode)`

Returns the internal representation (`runtime_mode` constant) from the readable string representing the name.

This represents the internal representation from the readable string representing the name.

Parameters

- `mode`: this represents the runtime mode

`void register_pre_shutdown_function (shutdown_function_type f)`

Add a function to be executed by a HPX thread during `hpx::finalize()` but guaranteed before any shutdown function is executed (system-wide)

Any of the functions registered with `register_pre_shutdown_function` are guaranteed to be executed by an HPX thread during the execution of `hpx::finalize()` before any of the registered shutdown functions are executed (see: [hpx::register_shutdown_function\(\)](#)).

Note If this function is called while the pre-shutdown functions are being executed, or after that point, it will raise a `invalid_status` exception.

See [hpx::register_shutdown_function\(\)](#)

Parameters

- `f`: [in] The function to be registered to run by an HPX thread as a pre-shutdown function.

`void register_shutdown_function (shutdown_function_type f)`

Add a function to be executed by a HPX thread during `hpx::finalize()` but guaranteed after any pre-shutdown function is executed (system-wide)

Any of the functions registered with `register_shutdown_function` are guaranteed to be executed by an HPX thread during the execution of `hpx::finalize()` after any of the registered pre-shutdown functions are executed (see: [hpx::register_pre_shutdown_function\(\)](#)).

Note If this function is called while the shutdown functions are being executed, or after that point, it will raise a `invalid_status` exception.

See [hpx::register_pre_shutdown_function\(\)](#)

Parameters

- `f`: [in] The function to be registered to run by an HPX thread as a shutdown function.

`void register_pre_startup_function (startup_function_type f)`

Add a function to be executed by a HPX thread before `hpx_main` but guaranteed before any startup function is executed (system-wide).

Any of the functions registered with `register_pre_startup_function` are guaranteed to be executed by an HPX thread before any of the registered startup functions are executed (see `hpx::register_startup_function()`).

This function is one of the few API functions which can be called before the runtime system has been fully initialized. It will automatically stage the provided startup function to the runtime system during its initialization (if necessary).

Note If this function is called while the pre-startup functions are being executed or after that point, it will raise a `invalid_status` exception.

Parameters

- `f`: [in] The function to be registered to run by an HPX thread as a pre-startup function.

See `hpx::register_startup_function()`

```
void register_startup_function(startup_function_type f)
```

Add a function to be executed by a HPX thread before `hpx_main` but guaranteed after any pre-startup function is executed (system-wide).

Any of the functions registered with `register_startup_function` are guaranteed to be executed by an HPX thread after any of the registered pre-startup functions are executed (see: `hpx::register_pre_startup_function()`), but before `hpx_main` is being called.

This function is one of the few API functions which can be called before the runtime system has been fully initialized. It will automatically stage the provided startup function to the runtime system during its initialization (if necessary).

Note If this function is called while the startup functions are being executed or after that point, it will raise a `invalid_status` exception.

Parameters

- `f`: [in] The function to be registered to run by an HPX thread as a startup function.

See `hpx::register_pre_startup_function()`

```
void trigger_lco_event(naming::id_type const &id, naming::address &&addr, bool move_credits = true)
```

Trigger the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should be triggered.
- `addr`: [in] This represents the addr of the LCO which should be triggered.
- `move_credits`: [in] If this is set to `true` then it is ok to send all credits in `id` along with the generated message. The default value is `true`.

```
void trigger_lco_event(naming::id_type const &id, bool move_credits = true)
```

Trigger the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should be triggered.
- `move_credits`: [in] If this is set to `true` then it is ok to send all credits in `id` along with the generated message. The default value is `true`.

```
void trigger_lco_event (naming::id_type const &id, naming::address &&addr, naming::id_type
const &cont, bool move_credits = true)
Trigger the LCO referenced by the given id.
```

Parameters

- *id*: [in] This represents the id of the LCO which should be triggered.
- *addr*: [in] This represents the addr of the LCO which should be triggered.
- *cont*: [in] This represents the LCO to trigger after completion.
- *move_credits*: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void trigger_lco_event (naming::id_type const &id, naming::id_type const &cont, bool
move_credits = true)
Trigger the LCO referenced by the given id.
```

Parameters

- *id*: [in] This represents the id of the LCO which should be triggered.
- *cont*: [in] This represents the LCO to trigger after completion.
- *move_credits*: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
template<typename Result>
void set_lco_value (naming::id_type const &id, naming::address &&addr, Result &&t, bool
move_credits = true)
Set the result value for the LCO referenced by the given id.
```

Parameters

- *id*: [in] This represents the id of the LCO which should receive the given value.
- *addr*: [in] This represents the addr of the LCO which should be triggered.
- *t*: [in] This is the value which should be sent to the LCO.
- *move_credits*: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
template<typename Result>
std::enable_if<!std::is_same<typename util::decay<Result>::type, naming::address>::value>::type set_lco_value (naming
const &id,
Re-
sult
&&t,
bool
move_credits
=
true)
```

Set the result value for the (managed) LCO referenced by the given id.

Parameters

- *id*: [in] This represents the id of the LCO which should receive the given value.

- `t`: [in] This is the value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in `id` along with the generated message. The default value is *true*.

```
template<typename Result>
std::enable_if<!std::is_same<typename util::decay<Result>::type, naming::address>::value>::type set_lco_value_unman
```

Set the result value for the (unmanaged) LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the given value.
- `t`: [in] This is the value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in `id` along with the generated message. The default value is *true*.

```
template<typename Result>
void set_lco_value(naming::id_type const &id, naming::address &&addr, Result &&t, naming::id_type const &cont, bool move_credits = true)
```

Set the result value for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the given value.
- `addr`: [in] This represents the addr of the LCO which should be triggered.
- `t`: [in] This is the value which should be sent to the LCO.
- `cont`: [in] This represents the LCO to trigger after completion.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in `id` along with the generated message. The default value is *true*.

```
template<typename Result>
```

```
std::enable_if<!std::is_same<typename util::decay<Result>::type, naming::address>>::value>::type set_lco_value (naming::const &id, Result &e, bool move_credits = true)
```

Set the result value for the (managed) LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the given value.
- `t`: [in] This is the value which should be sent to the LCO.
- `cont`: [in] This represents the LCO to trigger after completion.
- `move_credits`: [in] If this is set to `true` then it is ok to send all credits in `id` along with the generated message. The default value is `true`.

```
template<typename Result>
std::enable_if<!std::is_same<typename util::decay<Result>::type, naming::address>>::value>::type set_lco_value_unmanaged (naming::const &id, Result &e, bool move_credits = true)
```

Set the result value for the (unmanaged) LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the given value.
- `t`: [in] This is the value which should be sent to the LCO.
- `cont`: [in] This represents the LCO to trigger after completion.
- `move_credits`: [in] If this is set to `true` then it is ok to send all credits in `id` along with the generated message. The default value is `true`.

```
void set_lco_error (naming::id_type const &id, naming::address &&addr, std::exception_ptr const &e, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the error value.
- `addr`: [in] This represents the addr of the LCO which should be triggered.
- `e`: [in] This is the error value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, naming::address &&addr, std::exception_ptr  
          &&e, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the error value.
- `addr`: [in] This represents the addr of the LCO which should be triggered.
- `e`: [in] This is the error value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, std::exception_ptr const &e, bool move_credits  
                      = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the error value.
- `e`: [in] This is the error value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, std::exception_ptr &&e, bool move_credits =  
                           true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the error value.
- `e`: [in] This is the error value which should be sent to the LCO.
- `move_credits`: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, naming::address &&addr, std::exception_ptr  
          const &e, naming::id_type const &cont, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- `id`: [in] This represents the id of the LCO which should receive the error value.
- `addr`: [in] This represents the addr of the LCO which should be triggered.

- e: [in] This is the error value which should be sent to the LCO.
- cont: [in] This represents the LCO to trigger after completion.
- move_credits: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, naming::address &&addr, std::exception_ptr  

&&e, naming::id_type const &cont, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- id: [in] This represents the id of the LCO which should receive the error value.
- addr: [in] This represents the addr of the LCO which should be triggered.
- e: [in] This is the error value which should be sent to the LCO.
- cont: [in] This represents the LCO to trigger after completion.
- move_credits: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, std::exception_ptr const &e, naming::id_type  

const &cont, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- id: [in] This represents the id of the LCO which should receive the error value.
- e: [in] This is the error value which should be sent to the LCO.
- cont: [in] This represents the LCO to trigger after completion.
- move_credits: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
void set_lco_error(naming::id_type const &id, std::exception_ptr &&e, naming::id_type const  

&cont, bool move_credits = true)
```

Set the error state for the LCO referenced by the given id.

Parameters

- id: [in] This represents the id of the LCO which should receive the error value.
- e: [in] This is the error value which should be sent to the LCO.
- cont: [in] This represents the LCO to trigger after completion.
- move_credits: [in] If this is set to *true* then it is ok to send all credits in *id* along with the generated message. The default value is *true*.

```
template<typename Component, typename ... Ts><unspecified> hpx::new_(id_type const &
```

Create one or more new instances of the given Component type on the specified locality.

This function creates one or more new instances of the given Component type on the specified locality and returns a future object for the global address which can be used to reference the new component instance.

Note This function requires to specify an explicit template argument which will define what type of component(s) to create, for instance:

```
hpx::future<hpx::id_type> f =
    hpx::new_<some_component>(hpx::find_here(), ...);
hpx::id_type id = f.get();
```

Return The function returns different types depending on its use:

- If the explicit template argument *Component* represents a component type (`traits::is_component<Component>::value` evaluates to true), the function will return an `hpx::future` object instance which can be used to retrieve the global address of the newly created component.
- If the explicit template argument *Component* represents a client side object (`traits::is_client<Component>::value` evaluates to true), the function will return a new instance of that type which can be used to refer to the newly created component instance.

Parameters

- `locality`: [in] The global address of the locality where the new instance should be created on.
- `vs`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the constructor of the created component instance.

template<typename Component, typename ... Ts><unspecified> hpx::local_new(Ts &&... vs)
Create one new instance of the given Component type on the current locality.

This function creates one new instance of the given Component type on the current locality and returns a future object for the global address which can be used to reference the new component instance.

Note This function requires to specify an explicit template argument which will define what type of component(s) to create, for instance:

```
hpx::future<hpx::id_type> f =
    hpx::local_new<some_component>(...);
hpx::id_type id = f.get();
```

Return The function returns different types depending on its use:

- If the explicit template argument *Component* represents a component type (`traits::is_component<Component>::value` evaluates to true), the function will return an `hpx::future` object instance which can be used to retrieve the global address of the newly created component. If the first argument is `hpx::launch::sync` the function will directly return an `hpx::id_type`.
- If the explicit template argument *Component* represents a client side object (`traits::is_client<Component>::value` evaluates to true), the function will return a new instance of that type which can be used to refer to the newly created component instance.

Note The difference of this function to `hpx::new_` is that it can be used in cases where the supplied arguments are non-copyable and non-movable. All operations are guaranteed to be local only.

Parameters

- `vs`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the constructor of the created component instance.

template<typename Component, typename ... Ts><unspecified> hpx::new_(id_type const & l
Create multiple new instances of the given Component type on the specified locality.

This function creates multiple new instances of the given Component type on the specified locality and returns a future object for the global address which can be used to reference the new component instance.

Note This function requires to specify an explicit template argument which will define what type of component(s) to create, for instance:

```
hpx::future<std::vector<hpx::id_type>> f =
    hpx::new_<some_component[]>(hpx::find_here(), 10, ...);
hpx::id_type id = f.get();
```

Return The function returns different types depending on its use:

- If the explicit template argument *Component* represents an array of a component type (i.e. *Component[]*, where *traits::is_component<Component>::value* evaluates to true), the function will return an *hpx::future* object instance which holds a *std::vector<hpx::id_type>*, where each of the items in this vector is a global address of one of the newly created components.
- If the explicit template argument *Component* represents an array of a client side object type (i.e. *Component[]*, where *traits::is_client<Component>::value* evaluates to true), the function will return an *hpx::future* object instance which holds a *std::vector<hpx::id_type>*, where each of the items in this vector is a client side instance of the given type, each representing one of the newly created components.

Parameters

- *locality*: [in] The global address of the locality where the new instance should be created on.
- *count*: [in] The number of component instances to create
- *vs*: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the constructor of the created component instance.

template<typename Component, typename DistPolicy, typename ... Ts><unspecified> hpx::n
Create one or more new instances of the given Component type based on the given distribution policy.

This function creates one or more new instances of the given Component type on the localities defined by the given distribution policy and returns a future object for global address which can be used to reference the new component instance(s).

Note This function requires to specify an explicit template argument which will define what type of component(s) to create, for instance:

```
hpx::future<hpx::id_type> f =
    hpx::new_<some_component>(hpx::default_layout, ...);
hpx::id_type id = f.get();
```

Return The function returns different types depending on its use:

- If the explicit template argument *Component* represents a component type (*traits::is_component<Component>::value* evaluates to true), the function will return an *hpx::future* object instance which can be used to retrieve the global address of the newly created component.
- If the explicit template argument *Component* represents a client side object (*traits::is_client<Component>::value* evaluates to true), the function will return a new instance of that type which can be used to refer to the newly created component instance.

Parameters

- *policy*: [in] The distribution policy used to decide where to place the newly created.

- `vs`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the constructor of the created component instance.

```
template<typename Component, typename DistPolicy, typename ... Ts><unspecified> hpx::n
```

Create multiple new instances of the given Component type on the localities as defined by the given distribution policy.

This function creates multiple new instances of the given Component type on the localities defined by the given distribution policy and returns a future object for the global address which can be used to reference the new component instance.

Note This function requires to specify an explicit template argument which will define what type of component(s) to create, for instance:

```
hpx::future<std::vector<hpx::id_type> > f =
    hpx::new_<some_component[]>(hpx::default_layout, 10, ...);
hpx::id_type id = f.get();
```

Return The function returns different types depending on its use:

- If the explicit template argument `Component` represents an array of a component type (i.e. `Component[]`, where `traits::is_component<Component>::value` evaluates to true), the function will return an `hpx::future` object instance which holds a `std::vector<hpx::id_type>`, where each of the items in this vector is a global address of one of the newly created components.
- If the explicit template argument `Component` represents an array of a client side object type (i.e. `Component[]`, where `traits::is_client<Component>::value` evaluates to true), the function will return an `hpx::future` object instance which holds a `std::vector<hpx::id_type>`, where each of the items in this vector is a client side instance of the given type, each representing one of the newly created components.

Parameters

- `policy`: [in] The distribution policy used to decide where to place the newly created.
- `count`: [in] The number of component instances to create
- `vs`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the constructor of the created component instance.

```
template<typename ...Ts>
```

```
tuple<future<Ts...>> split_future(future<tuple<Ts...>> &&f)
```

The function `split_future` is an operator allowing to split a given future of a sequence of values (any tuple, `std::pair`, or `std::array`) into an equivalent container of futures where each future represents one of the values from the original future. In some sense this function provides the inverse operation of `when_all`.

Return Returns an equivalent container (same container type as passed as the argument) of futures, where each future refers to the corresponding value in the input parameter. All of the returned futures become ready once the input future has become ready. If the input future is exceptional, all output futures will be exceptional as well.

Note The following cases are special:

```
tuple<future<void> > split_future(future<tuple<>> && f);
array<future<void>, 1> split_future(future<array<T, 0>> && f);
```

here the returned futures are directly representing the futures which were passed to the function.

Parameters

- `f`: [in] A future holding an arbitrary sequence of values stored in a tuple-like container. This facility supports `hpx::util::tuple<T>`, `std::pair<T1, T2>`, and `std::array<T, N>`

```
template<typename T>
std::vector<future<T>> split_future(future<std::vector<T>> &&f, std::size_t size)
```

The function `split_future` is an operator allowing to split a given future of a sequence of values (any `std::vector`) into a `std::vector` of futures where each future represents one of the values from the original `std::vector`. In some sense this function provides the inverse operation of `when_all`.

Return Returns a `std::vector` of futures, where each future refers to the corresponding value in the input parameter. All of the returned futures become ready once the input future has become ready. If the input future is exceptional, all output futures will be exceptional as well.

Parameters

- `f`: [in] A future holding an arbitrary sequence of values stored in a `std::vector`.
- `size`: [in] The number of elements the vector will hold once the input future has become ready

```
template<typename InputIter>
void wait_all(InputIter first, InputIter last)
```

The function `wait_all` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns after they finished executing.

Note The function `wait_all` returns after all futures have become ready. All input futures are still valid after `wait_all` returns.

Parameters

- `first`: The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `wait_all` should wait.
- `last`: The iterator pointing to the last element of a sequence of `future` or `shared_future` objects for which `wait_all` should wait.

```
template<typename R>
void wait_all(std::vector<future<R>> &&futures)
```

The function `wait_all` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns after they finished executing.

Note The function `wait_all` returns after all futures have become ready. All input futures are still valid after `wait_all` returns.

Parameters

- `futures`: A vector or array holding an arbitrary amount of `future` or `shared_future` objects for which `wait_all` should wait.

```
template<typename R, std::size_t N>
void wait_all(std::array<future<R>, N> &&futures)
```

The function `wait_all` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns after they finished executing.

Note The function `wait_all` returns after all futures have become ready. All input futures are still valid after `wait_all` returns.

Parameters

- **futures**: A vector or array holding an arbitrary amount of *future* or *shared_future* objects for which *wait_all* should wait.

```
template<typename ...T>
void wait_all(T&&... futures)
```

The function *wait_all* is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns after they finished executing.

Note The function *wait_all* returns after all futures have become ready. All input futures are still valid after *wait_all* returns.

Parameters

- **futures**: An arbitrary number of *future* or *shared_future* objects, possibly holding different types for which *wait_all* should wait.

```
template<typename InputIter>
InputIter wait_all_n(InputIter begin, std::size_t count)
```

The function *wait_all_n* is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns after they finished executing.

Return The function *wait_all_n* will return an iterator referring to the first element in the input sequence after the last processed element.

Note The function *wait_all_n* returns after all futures have become ready. All input futures are still valid after *wait_all_n* returns.

Parameters

- **begin**: The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which *wait_all_n* should wait.
- **count**: The number of elements in the sequence starting at *first*.

```
template<typename InputIter, typename Container = std::vector<future<typename std::iterator_traits<InputIter>::value_type>>
future<Container> when_all(InputIter first, InputIter last)
```

The function *when_all* is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after they finished executing.

Return Returns a future holding the same list of futures as has been passed to *when_all*.

- *future<Container<future<R>>>*: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Note Calling this version of *when_all* where *first == last*, returns a future with an empty container that is immediately ready. Each future and *shared_future* is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by *when_all* will not throw an exception, but the futures held in the output collection may.

Parameters

- **first**: [in] The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which *when_all* should wait.
- **last**: [in] The iterator pointing to the last element of a sequence of *future* or *shared_future* objects for which *when_all* should wait.

```
template<typename Range>
future<Range> when_all (Range &&values)
```

The function `when_all` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after they finished executing.

Return Returns a future holding the same list of futures as has been passed to `when_all`.

- `future<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type.

Note Calling this version of `when_all` where the input container is empty, returns a future with an empty container that is immediately ready. Each future and shared_future is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_all` will not throw an exception, but the futures held in the output collection may.

Parameters

- `values`: [in] A range holding an arbitrary amount of `future` or `shared_future` objects for which `when_all` should wait.

```
template<typename ...T>
future<tuple<future<T>...>> when_all (T&&... futures)
```

The function `when_all` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after they finished executing.

Return Returns a future holding the same list of futures as has been passed to `when_all`.

- `future<tuple<future<T0>, future<T1>, future<T2>...>>`: If inputs are fixed in number and are of heterogeneous types. The inputs can be any arbitrary number of future objects.
- `future<tuple<>>` if `when_all` is called with zero arguments. The returned future will be initially ready.

Note Each future and shared_future is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_all` will not throw an exception, but the futures held in the output collection may.

Parameters

- `futures`: [in] An arbitrary number of `future` or `shared_future` objects, possibly holding different types for which `when_all` should wait.

```
template<typename InputIter, typename Container = vector<future<typename std::iterator_traits<InputIter>::value_type>>
future<Container> when_all_n (InputIter begin, std::size_t count)
```

The function `when_all_n` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after they finished executing.

Return Returns a future holding the same list of futures as has been passed to `when_all_n`.

- `future<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output vector will be the same as given by the input iterator.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `begin`: [in] The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which `wait_all_n` should wait.
- `count`: [in] The number of elements in the sequence starting at `first`.

Exceptions

- `This`: function will throw errors which are encountered while setting up the requested operation only. Errors encountered while executing the operations delivering the results to be stored in the futures are reported through the futures themselves.

```
template<typename InputIter>
void wait_any(InputIter first, InputIter last, error_code &ec = throws)
```

The function `wait_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any` returns after at least one future has become ready. All input futures are still valid after `wait_any` returns.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `first`: [in] The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which `wait_any` should wait.
- `last`: [in] The iterator pointing to the last element of a sequence of *future* or *shared_future* objects for which `wait_any` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename R>
void wait_any(std::vector<future<R>> &futures, error_code &ec = throws)
```

The function `wait_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any` returns after at least one future has become ready. All input futures are still valid after `wait_any` returns.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `futures`: [in] A vector holding an arbitrary amount of *future* or *shared_future* objects for which `wait_any` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename R, std::size_t N>void hpx::wait_any(std::array< future< R >, N > & f
```

The function `wait_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any` returns after at least one future has become ready. All input futures are still valid after `wait_any` returns.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `futures`: [in] An array holding an arbitrary amount of `future` or `shared_future` objects for which `wait_any` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename ...T>
```

```
void wait_any(error_code &ec, T&&... futures)
```

The function `wait_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any` returns after at least one future has become ready. All input futures are still valid after `wait_any` returns.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `futures`: [in] An arbitrary number of `future` or `shared_future` objects, possibly holding different types for which `wait_any` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename ...T>
```

```
void wait_any(T&&... futures)
```

The function `wait_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any` returns after at least one future has become ready. All input futures are still valid after `wait_any` returns.

Note None of the futures in the input sequence are invalidated.

Parameters

- `futures`: [in] An arbitrary number of `future` or `shared_future` objects, possibly holding different types for which `wait_any` should wait.

```
template<typename InputIter>
```

```
InputIter wait_any_n(InputIter first, std::size_t count, error_code &ec = throws)
```

The function `wait_any_n` is a non-deterministic choice operator. It OR-composes all future objects given and returns after one future of that list finishes execution.

Note The function `wait_any_n` returns after at least one future has become ready. All input futures are still valid after `wait_any_n` returns.

Return The function `wait_all_n` will return an iterator referring to the first element in the input sequence after the last processed element.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Note None of the futures in the input sequence are invalidated.

Parameters

- `first`: [in] The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `wait_any_n` should wait.
- `count`: [in] The number of elements in the sequence starting at `first`.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename InputIter, typename Container = vector<future<typename std::iterator_traits<InputIter>::value_type>>
future<when_any_result<Container>> when_any (InputIter first, InputIter last)
```

The function `when_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns a new future object representing the same list of futures after one future of that list finishes execution.

Return Returns a `when_any_result` holding the same list of futures as has been passed to `when_any` and an index pointing to a ready future.

- `future<when_any_result<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Parameters

- `first`: [in] The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `when_any` should wait.
- `last`: [in] The iterator pointing to the last element of a sequence of `future` or `shared_future` objects for which `when_any` should wait.

```
template<typename Range>
future<when_any_result<Range>> when_any (Range &values)
```

The function `when_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns a new future object representing the same list of futures after one future of that list finishes execution.

Return Returns a `when_any_result` holding the same list of futures as has been passed to `when_any` and an index pointing to a ready future.

- `future<when_any_result<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Parameters

- `values`: [in] A range holding an arbitrary amount of `futures` or `shared_future` objects for which `when_any` should wait.

```
template<typename ...T>
future<when_any_result<tuple<future<T>...>>> when_any(T&&... futures)
```

The function `when_any` is a non-deterministic choice operator. It OR-composes all future objects given and returns a new future object representing the same list of futures after one future of that list finishes execution.

Return Returns a `when_any_result` holding the same list of futures as has been passed to `when_any` and an index pointing to a ready future..

- `future<when_any_result<tuple<future<T0>, future<T1>...>>>`: If inputs are fixed in number and are of heterogeneous types. The inputs can be any arbitrary number of future objects.
- `future<when_any_result<tuple<>>>` if `when_any` is called with zero arguments. The returned future will be initially ready.

Parameters

- `futures`: [in] An arbitrary number of `future` or `shared_future` objects, possibly holding different types for which `when_any` should wait.

```
template<typename InputIter, typename Container = vector<future<typename std::iterator_traits<InputIter>::value_type>>
future<when_any_result<Container>> when_any_n(InputIter first, std::size_t count)
```

The function `when_any_n` is a non-deterministic choice operator. It OR-composes all future objects given and returns a new future object representing the same list of futures after one future of that list finishes execution.

Return Returns a `when_any_result` holding the same list of futures as has been passed to `when_any` and an index pointing to a ready future.

- `future<when_any_result<Container<future<R>>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Note None of the futures in the input sequence are invalidated.

Parameters

- `first`: [in] The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `when_any_n` should wait.
- `count`: [in] The number of elements in the sequence starting at `first`.

```
template<typename InputIter>
future<vector<future<typename std::iterator_traits<InputIter>::value_type>>> wait_some(std::size_t
n, Iterator
first,
Iterator
last, error_code
&ec =
throws)
```

The function `wait_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The future returned by the function `wait_some` becomes ready when at least `n` argument futures have become ready.

Return Returns a future holding the same list of futures as has been passed to `wait_some`.

- `future<vector<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type.

Note Calling this version of `wait_some` where `first == last`, returns a future with an empty vector that is immediately ready. Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `wait_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `first`: [in] The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `when_all` should wait.
- `last`: [in] The iterator pointing to the last element of a sequence of `future` or `shared_future` objects for which `when_all` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename R>
void wait_some (std::size_t n, std::vector<future<R>> &&futures, error_code &ec = throws)
```

The function `wait_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after *n* of them finished executing.

Note The function `wait_all` returns after *n* futures have become ready. All input futures are still valid after `wait_all` returns.

Note Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `wait_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `futures`: [in] A vector holding an arbitrary amount of `future` or `shared_future` objects for which `wait_some` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename R, std::size_t N>
```

```
void wait_some (std::size_t n, std::array<future<R>, N> &&futures, error_code &ec = throws)
```

The function `wait_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after *n* of them finished executing.

Note The function `wait_all` returns after *n* futures have become ready. All input futures are still valid after `wait_all` returns.

Note Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `wait_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `futures`: [in] An array holding an arbitrary amount of *future* or *shared_future* objects for which `wait_some` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename ...T>
void wait_some (std::size_t n, T&&... futures, error_code &ec = throws)
```

The function `wait_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The function `wait_all` returns after `n` futures have become ready. All input futures are still valid after `wait_all` returns.

Note Calling this version of `wait_some` where `first == last`, returns a future with an empty vector that is immediately ready. Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `wait_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `futures`: [in] An arbitrary number of *future* or *shared_future* objects, possibly holding different types for which `wait_some` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename InputIter>
```

```
InputIter wait_some_n (std::size_t n, Iterator first, std::size_t count, error_code &ec = throws)
```

The function `wait_some_n` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The function `wait_all` returns after `n` futures have become ready. All input futures are still valid after `wait_all` returns.

Return This function returns an Iterator referring to the first element after the last processed input element.

Note Calling this version of `wait_some_n` where `count == 0`, returns a future with the same elements as the arguments that is immediately ready. Possibly none of the futures in that vector are ready. Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `wait_some_n` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `first`: [in] The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which `when_all` should wait.
- `count`: [in] The number of elements in the sequence starting at `first`.

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename InputIter, typename Container = vector<future<typename std::iterator_traits<InputIter>::value_type>
future<when_some_result<Container>> when_some (std::size_t n, Iterator first, Iterator last, error_code &ec = throws)
```

The function `when_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after *n* of them finished executing.

Note The future returned by the function `when_some` becomes ready when at least *n* argument futures have become ready.

Return Returns a `when_some_result` holding the same list of futures as has been passed to `when_some` and indices pointing to ready futures.

- `future<when_some_result<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Note Calling this version of `when_some` where `first == last`, returns a future with an empty container that is immediately ready. Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `first`: [in] The iterator pointing to the first element of a sequence of `future` or `shared_future` objects for which `when_all` should wait.
- `last`: [in] The iterator pointing to the last element of a sequence of `future` or `shared_future` objects for which `when_all` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename Range>
future<when_some_result<Range>> when_some (std::size_t n, Range &&futures, error_code &ec = throws)
```

The function `when_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after *n* of them finished executing.

Note The future returned by the function `when_some` becomes ready when at least *n* argument futures have become ready.

Return Returns a `when_some_result` holding the same list of futures as has been passed to `when_some` and indices pointing to ready futures.

- `future<when_some_result<Container<future<R>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Note Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `futures`: [in] A container holding an arbitrary amount of `future` or `shared_future` objects for which `when_some` should wait.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename ...T>
future<when_some_result<tuple<future<T>...>>> when_some (std::size_t n, error_code &ec, T&&...  
futures)
```

The function `when_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The future returned by the function `when_some` becomes ready when at least `n` argument futures have become ready.

Return Returns a `when_some_result` holding the same list of futures as has been passed to `when_some` and an index pointing to a ready future..

- `future<when_some_result<tuple<future<T0>, future<T1>...>>>`: If inputs are fixed in number and are of heterogeneous types. The inputs can be any arbitrary number of future objects.
- `future<when_some_result<tuple<>>>` if `when_some` is called with zero arguments. The returned future will be initially ready.

Note Each future and `shared_future` is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.
- `futures`: [in] An arbitrary number of `future` or `shared_future` objects, possibly holding different types for which `when_some` should wait.

```
template<typename ...T>
future<when_some_result<tuple<future<T>...>>> when_some (std::size_t n, T&&... futures)
```

The function `when_some` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The future returned by the function `when_some` becomes ready when at least `n` argument futures have become ready.

Return Returns a `when_some_result` holding the same list of futures as has been passed to `when_some` and an index pointing to a ready future..

- `future<when_some_result<tuple<future<T0>, future<T1>...>>>`: If inputs are fixed in number and are of heterogeneous types. The inputs can be any arbitrary number of future objects.

- `future<when_some_result<tuple<>>>` if `when_some` is called with zero arguments. The returned future will be initially ready.

Note Each future and shared_future is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_some` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `futures`: [in] An arbitrary number of *future* or *shared_future* objects, possibly holding different types for which `when_some` should wait.

```
template<typename InputIter, typename Container = vector<future<typename std::iterator_traits<InputIter>::value_type>>
future<when_some_result<Container>> when_some_n (std::size_t n, Iterator first, std::size_t count, error_code &ec = throws)
```

The function `when_some_n` is an operator allowing to join on the result of all given futures. It AND-composes all future objects given and returns a new future object representing the same list of futures after `n` of them finished executing.

Note The future returned by the function `when_some_n` becomes ready when at least `n` argument futures have become ready.

Return Returns a `when_some_result` holding the same list of futures as has been passed to `when_some` and indices pointing to ready futures.

- `future<when_some_result<Container<future<R>>>>`: If the input cardinality is unknown at compile time and the futures are all of the same type. The order of the futures in the output container will be the same as given by the input iterator.

Note Calling this version of `when_some_n` where `count == 0`, returns a future with the same elements as the arguments that is immediately ready. Possibly none of the futures in that container are ready. Each future and shared_future is waited upon and then copied into the collection of the output (returned) future, maintaining the order of the futures in the input collection. The future returned by `when_some_n` will not throw an exception, but the futures held in the output collection may.

Parameters

- `n`: [in] The number of futures out of the arguments which have to become ready in order for the returned future to get ready.
- `first`: [in] The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which `when_all` should wait.
- `count`: [in] The number of elements in the sequence starting at `first`.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
template<typename F, typename Future>
void wait_each (F &&f, std::vector<Future> &&futures)
```

The function `wait_each` is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns after they finished executing. Additionally, the supplied function is called for each of the passed futures as soon as the future has become ready. `wait_each` returns after all futures have been become ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that `std::size_t`

is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *futures*: A vector holding an arbitrary amount of *future* or *shared_future* objects for which *wait_each* should wait.

```
template<typename F, typename Iterator>
void wait_each (F &&f, Iterator begin, Iterator end)
```

The function *wait_each* is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns after they finished executing. Additionally, the supplied function is called for each of the passed futures as soon as the future has become ready. *wait_each* returns after all futures have been become ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that *std::size_t* is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *begin*: The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which *wait_each* should wait.
- *end*: The iterator pointing to the last element of a sequence of *future* or *shared_future* objects for which *wait_each* should wait.

```
template<typename F, typename ...T>
void wait_each (F &&f, T &&... futures)
```

The function *wait_each* is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns after they finished executing. Additionally, the supplied function is called for each of the passed futures as soon as the future has become ready. *wait_each* returns after all futures have been become ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that *std::size_t* is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *futures*: An arbitrary number of *future* or *shared_future* objects, possibly holding different types for which *wait_each* should wait.

```
template<typename F, typename Iterator>
void wait_each_n (F &&f, Iterator begin, std::size_t count)
```

The function *wait_each* is an operator allowing to join on the result of all given futures. It AND-composes

all future objects given and returns after they finished executing. Additionally, the supplied function is called for each of the passed futures as soon as the future has become ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that *std::size_t* is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *begin*: The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which *wait_each_n* should wait.
- *count*: The number of elements in the sequence starting at *first*.

```
template<typename F, typename Future>
future<void> when_each (F &&f, std::vector<Future> &&futures)
```

The function *when_each* is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns a new future object representing the event of all those futures having finished executing. It also calls the supplied callback for each of the futures which becomes ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that *std::size_t* is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Return Returns a future representing the event of all input futures being ready.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *futures*: A vector holding an arbitrary amount of *future* or *shared_future* objects for which *wait_each* should wait.

```
template<typename F, typename Iterator>
future<Iterator> when_each (F &&f, Iterator begin, Iterator end)
```

The function *when_each* is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns a new future object representing the event of all those futures having finished executing. It also calls the supplied callback for each of the futures which becomes ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that *std::size_t* is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Return Returns a future representing the event of all input futures being ready.

Parameters

- *f*: The function which will be called for each of the input futures once the future has become ready.
- *begin*: The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which *wait_each* should wait.

- `end`: The iterator pointing to the last element of a sequence of *future* or *shared_future* objects for which `wait_each` should wait.

```
template<typename F, typename ...Ts>
future<void> when_each (F &&f, Ts&&... futures)
```

The function `when_each` is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns a new future object representing the event of all those futures having finished executing. It also calls the supplied callback for each of the futures which becomes ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that `std::size_t` is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Return Returns a future representing the event of all input futures being ready.

Parameters

- `f`: The function which will be called for each of the input futures once the future has become ready.
- `futures`: An arbitrary number of *future* or *shared_future* objects, possibly holding different types for which `wait_each` should wait.

```
template<typename F, typename Iterator>
future<Iterator> when_each_n (F &&f, Iterator begin, std::size_t count)
```

The function `when_each` is an operator allowing to join on the results of all given futures. It AND-composes all future objects given and returns a new future object representing the event of all those futures having finished executing. It also calls the supplied callback for each of the futures which becomes ready.

Note This function consumes the futures as they are passed on to the supplied function. The callback should take one or two parameters, namely either a *future* to be processed or a type that `std::size_t` is implicitly convertible to as the first parameter and the *future* as the second parameter. The first parameter will correspond to the index of the current *future* in the collection.

Return Returns a future holding the iterator pointing to the first element after the last one.

Parameters

- `f`: The function which will be called for each of the input futures once the future has become ready.
- `begin`: The iterator pointing to the first element of a sequence of *future* or *shared_future* objects for which `wait_each_n` should wait.
- `count`: The number of elements in the sequence starting at `first`.

```
namespace actions
```

```
namespace applier
```

Functions

`applier &get_applier()`

The function `get_applier` returns a reference to the (thread specific) applier instance.

`applier *get_applier_ptr()`

The function `get_applier` returns a pointer to the (thread specific) applier instance. The returned pointer is NULL if the current thread is not known to HPX or if the runtime system is not active.

namespace components

Functions

`template<typename Component>`

```
future<naming::id_type> migrate_from_storage(naming::id_type const &to_resurrect,
                                              naming::id_type const &target = naming::invalid_id)
```

Migrate the component with the given id from the specified target storage (resurrect the object)

The function `migrate_from_storage<Component>` will migrate the component referenced by `to_resurrect` from the storage facility specified where the object is currently stored on. It returns a future referring to the migrated component instance. The component instance is resurrected on the locality specified by `target_locality`.

Return A future representing the global id of the migrated component instance. This should be the same as `to_resurrect`.

Parameters

- `to_resurrect`: [in] The global id of the component to migrate.
- `target`: [in] The optional locality to resurrect the object on. By default the object is resurrected on the locality it was located on last.

Template Parameters

- The: only template argument specifies the component type of the component to migrate from the given storage facility.

`template<typename Component>`

```
future<naming::id_type> migrate_to_storage(naming::id_type const &to_migrate, naming::id_type const &target_storage)
```

Migrate the component with the given id to the specified target storage

The function `migrate_to_storage<Component>` will migrate the component referenced by `to_migrate` to the storage facility specified with `target_storage`. It returns a future referring to the migrated component instance.

Return A future representing the global id of the migrated component instance. This should be the same as `migrate_to`.

Parameters

- `to_migrate`: [in] The global id of the component to migrate.
- `target_storage`: [in] The id of the storage facility to migrate this object to.

Template Parameters

- The: only template argument specifies the component type of the component to migrate to the given storage facility.

```
template<typename Derived, typename Stub>
Derived migrate_to_storage(client_base<Derived, Stub> const &to_migrate,
                           hpx::components::component_storage const &target_storage)
Migrate the given component to the specified target storage
```

The function *migrate_to_storage* will migrate the component referenced by *to_migrate* to the storage facility specified with *target_storage*. It returns a future referring to the migrated component instance.

Return A client side representation of representing of the migrated component instance. This should be the same as *migrate_to*.

Parameters

- *to_migrate*: [in] The client side representation of the component to migrate.
- *target_storage*: [in] The id of the storage facility to migrate this object to.

```
template<typename Component>
future<naming::id_type> copy (naming::id_type const &to_copy)
Copy given component to the specified target locality.
```

The function *copy*<*Component*> will create a copy of the component referenced by *to_copy* on the locality specified with *target_locality*. It returns a future referring to the newly created component instance.

Return A future representing the global id of the newly (copied) component instance.

Note The new component instance is created on the locality of the component instance which is to be copied.

Parameters

- *to_copy*: [in] The global id of the component to copy

Template Parameters

- The: only template argument specifies the component type to create.

```
template<typename Component>
future<naming::id_type> copy (naming::id_type const &to_copy, naming::id_type const &target_locality)
Copy given component to the specified target locality.
```

The function *copy*<*Component*> will create a copy of the component referenced by *to_copy* on the locality specified with *target_locality*. It returns a future referring to the newly created component instance.

Return A future representing the global id of the newly (copied) component instance.

Parameters

- *to_copy*: [in] The global id of the component to copy
- *target_locality*: [in] The locality where the copy should be created.

Template Parameters

- The: only template argument specifies the component type to create.

```
template<typename Derived, typename Stub>
```

```
Derived copy (client_base<Derived, Stub> const &to_copy, naming::id_type const &target_locality  
= naming::invalid_id)
```

Copy given component to the specified target locality.

The function *copy* will create a copy of the component referenced by the client side object *to_copy* on the locality specified with *target_locality*. It returns a new client side object future referring to the newly created component instance.

Return A future representing the global id of the newly (copied) component instance.

Note If the second argument is omitted (or is *invalid_id*) the new component instance is created on the locality of the component instance which is to be copied.

Parameters

- *to_copy*: [in] The client side object representing the component to copy
- *target_locality*: [in, optional] The locality where the copy should be created (default is same locality as source).

Template Parameters

- *Type*: only template argument specifies the component type to create.

```
template<typename Component, typename DistPolicy>  
future<naming::id_type> migrate (naming::id_type const &to_migrate, DistPolicy const &policy)  
Migrate the given component to the specified target locality
```

The function *migrate*<*Component*> will migrate the component referenced by *to_migrate* to the locality specified with *target_locality*. It returns a future referring to the migrated component instance.

Return A future representing the global id of the migrated component instance. This should be the same as *migrate_to*.

Parameters

- *to_migrate*: [in] The client side representation of the component to migrate.
- *policy*: [in] A distribution policy which will be used to determine the locality to migrate this object to.

Template Parameters

- *Component*: Specifies the component type of the component to migrate.
- *DistPolicy*: Specifies the distribution policy to use to determine the destination locality.

```
template<typename Derived, typename Stub, typename DistPolicy>  
Derived migrate (client_base<Derived, Stub> const &to_migrate, DistPolicy const &policy)  
Migrate the given component to the specified target locality
```

The function *migrate*<*Component*> will migrate the component referenced by *to_migrate* to the locality specified with *target_locality*. It returns a future referring to the migrated component instance.

Return A future representing the global id of the migrated component instance. This should be the same as *migrate_to*.

Parameters

- *to_migrate*: [in] The client side representation of the component to migrate.
- *policy*: [in] A distribution policy which will be used to determine the locality to migrate this object to.

Template Parameters

- **Derived:** Specifies the component type of the component to migrate.
- **DistPolicy:** Specifies the distribution policy to use to determine the destination locality.

```
template<typename Component>
future<naming::id_type> migrate(naming::id_type const &to_migrate, naming::id_type const
&target_locality)
Migrate the component with the given id to the specified target locality
```

The function *migrate*<*Component*> will migrate the component referenced by *to_migrate* to the locality specified with *target_locality*. It returns a future referring to the migrated component instance.

Return A future representing the global id of the migrated component instance. This should be the same as *migrate_to*.

Parameters

- **to_migrate:** [in] The global id of the component to migrate.
- **target_locality:** [in] The locality where the component should be migrated to.

Template Parameters

- **Component:** Specifies the component type of the component to migrate.

```
template<typename Derived, typename Stub>
Derived migrate(client_base<Derived, Stub> const &to_migrate, naming::id_type const &tar-
get_locality)
Migrate the given component to the specified target locality
```

The function *migrate*<*Component*> will migrate the component referenced by *to_migrate* to the locality specified with *target_locality*. It returns a future referring to the migrated component instance.

Return A client side representation of representing of the migrated component instance. This should be the same as *migrate_to*.

Parameters

- **to_migrate:** [in] The client side representation of the component to migrate.
- **target_locality:** [in] The id of the locality to migrate this object to.

Template Parameters

- **Derived:** Specifies the component type of the component to migrate.

Variables

```
char const *const default_binpacking_counter_name = "/runtime{locality/total}/count/component@"
binpacking_distribution_policy const binpacked
```

A predefined instance of the binpacking *distribution_policy*. It will represent the local locality and will place all items to create here.

```
colocating_distribution_policy const colocated
```

A predefined instance of the co-locating *distribution_policy*. It will represent the local locality and will place all items to create here.

```
default_distribution_policy const default_layout = {}
```

A predefined instance of the default *distribution_policy*. It will represent the local locality and will place all items to create here.

namespace naming

Functions

```
id_type unmanaged(id_type const &id)
```

The helper function *hpx::unmanaged* can be used to generate a global identifier which does not participate in the automatic garbage collection.

Return This function returns a new global id referencing the same object as the parameter *id*. The only difference is that the returned global identifier does not participate in the automatic garbage collection.

Note This function allows to apply certain optimizations to the process of memory management in HPX. It however requires the user to take full responsibility for keeping the referenced objects alive long enough.

Parameters

- *id*: [in] The id to generated the unmanaged global id from. This parameter can be itself a managed or a unmanaged global id.

namespace parallel

namespace v2

Functions

```
template<typename ExPolicy, typename F>
```

```
util::detail::algorithm_result<ExPolicy>::type define_task_block(ExPolicy &&policy, F &&f)
```

Constructs a *task_block*, *tr*, using the given execution policy *policy*, and invokes the expression *f(tr)* on the user-provided object, *f*.

Postcondition: All tasks spawned from *f* have finished execution. A call to *define_task_block* may return on a different thread than that on which it was called.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the task block may be parallelized.
- *F*: The type of the user defined function to invoke inside the *define_task_block* (deduced). *F* shall be MoveConstructible.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *f*: The user defined function to invoke inside the task block. Given an lvalue *tr* of type *task_block*, the expression, (*void*)*f(tr)*, shall be well-formed.

Note It is expected (but not mandated) that *f* will (directly or indirectly) call *tr.run(callable_object)*.

Exceptions

- An: *exception_list*, as specified in Exception Handling.

```
template<typename F>
void define_task_block(F &&f)
Constructs a task_block, tr, and invokes the expression f(tr) on the user-provided object, f. This version uses parallel_policy for task scheduling.
```

Postcondition: All tasks spawned from `f` have finished execution. A call to `define_task_block` may return on a different thread than that on which it was called.

Template Parameters

- `F`: The type of the user defined function to invoke inside the `define_task_block` (deduced). `F` shall be `MoveConstructible`.

Parameters

- `f`: The user defined function to invoke inside the task block. Given an lvalue `tr` of type `task_block`, the expression, `(void)f(tr)`, shall be well-formed.

Note It is expected (but not mandated) that `f` will (directly or indirectly) call `tr.run(callable_object)`.

Exceptions

- An: `exception_list`, as specified in Exception Handling.

```
template<typename ExPolicy, typename F>
util::detail::algorithm_result<ExPolicy>::type define_task_block_restore_thread(ExPolicy
&&policy,
                                         F
                                         &&f)
```

Constructs a `task_block`, `tr`, and invokes the expression `f(tr)` on the user-provided object, `f`.

Postcondition: All tasks spawned from `f` have finished execution. A call to `define_task_block_restore_thread` always returns on the same thread as that on which it was called.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the task block may be parallelized.
- `F`: The type of the user defined function to invoke inside the `define_task_block` (deduced). `F` shall be `MoveConstructible`.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `f`: The user defined function to invoke inside the `define_task_block`. Given an lvalue `tr` of type `task_block`, the expression, `(void)f(tr)`, shall be well-formed.

Exceptions

- An: `exception_list`, as specified in Exception Handling.

Note It is expected (but not mandated) that `f` will (directly or indirectly) call `tr.run(callable_object)`.

```
template<typename F>
void define_task_block_restore_thread(F &&f)
Constructs a task_block, tr, and invokes the expression f(tr) on the user-provided object, f. This version uses parallel_policy for task scheduling.
```

Postcondition: All tasks spawned from `f` have finished execution. A call to `define_task_block_restore_thread` always returns on the same thread as that on which it was called.

Template Parameters

- *F*: The type of the user defined function to invoke inside the `define_task_block` (deduced). *F* shall be `MoveConstructible`.

Parameters

- *f*: The user defined function to invoke inside the `define_task_block`. Given an lvalue *tr* of type `task_block`, the expression, `(void)f(tr)`, shall be well-formed.

Exceptions

- An: *exception_list*, as specified in Exception Handling.

Note It is expected (but not mandated) that *f* will (directly or indirectly) call *tr.run(callable_object)*.

namespace performance_counters

Functions

```
counter_status install_counter_type(std::string const &name,  
                                  hpx::util::function_nonser<std::int64_t> bool  
> const &counter_value, std::string const &helptext = "", std::string const &uom = "", error_code  
&ec = throws)Install a new generic performance counter type in a way, which will uninstall it automatically  
during shutdown.
```

The function *install_counter_type* will register a new generic counter type based on the provided function. The counter type will be automatically unregistered during system shutdown. Any consumer querying any instance of this counter type will cause the provided function to be called and the returned value to be exposed as the counter value.

The counter type is registered such that there can be one counter instance per locality. The expected naming scheme for the counter instances is: '`/objectname{locality#<*>/total}/countername`' where '`<*>`' is a zero based integer identifying the locality the counter is created on.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Return If successful, this function returns *status_valid_data*, otherwise it will either throw an exception or return an *error_code* from the enum *counter_status* (also, see note related to parameter *ec*).

Note The counter type registry is a locality based service. You will have to register each counter type on every locality where a corresponding performance counter will be created.

Parameters

- *name*: [in] The global virtual name of the counter type. This name is expected to have the format `/objectname/countername`.
- *counter_value*: [in] The function to call whenever the counter value is requested by a consumer.
- *helptext*: [in, optional] A longer descriptive text shown to the user to explain the nature of the counters created from this type.
- *uom*: [in] The unit of measure for the new performance counter type.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
counter_status install_counter_type (std::string const &name,
          hpx::util::function_nonser<std::vector<std::int64_t>> bool
> const &counter_value, std::string const &helptext = "", std::string const &uom = "", error_code
&ec = throws) Install a new generic performance counter type returning an array of values in a way, that will uninstall it automatically during shutdown.
```

The function *install_counter_type* will register a new generic counter type that returns an array of values based on the provided function. The counter type will be automatically unregistered during system shutdown. Any consumer querying any instance of this counter type will cause the provided function to be called and the returned array value to be exposed as the counter value.

The counter type is registered such that there can be one counter instance per locality. The expected naming scheme for the counter instances is: '/objectname{locality#<*>/total}/countername' where '<*>' is a zero based integer identifying the locality the counter is created on.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Return If successful, this function returns *status_valid_data*, otherwise it will either throw an exception or return an *error_code* from the enum *counter_status* (also, see note related to parameter *ec*).

Note The counter type registry is a locality based service. You will have to register each counter type on every locality where a corresponding performance counter will be created.

Parameters

- *name*: [in] The global virtual name of the counter type. This name is expected to have the format /objectname/countername.
- *counter_value*: [in] The function to call whenever the counter value (array of values) is requested by a consumer.
- *helptext*: [in, optional] A longer descriptive text shown to the user to explain the nature of the counters created from this type.
- *uom*: [in] The unit of measure for the new performance counter type.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void install_counter_type (std::string const &name, counter_type type, error_code &ec =
                      throws)
```

Install a new performance counter type in a way, which will uninstall it automatically during shutdown.

The function *install_counter_type* will register a new counter type based on the provided *counter_type_info*. The counter type will be automatically unregistered during system shutdown.

Return If successful, this function returns *status_valid_data*, otherwise it will either throw an exception or return an *error_code* from the enum *counter_status* (also, see note related to parameter *ec*).

Note The counter type registry is a locality based service. You will have to register each counter type on every locality where a corresponding performance counter will be created.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *name*: [in] The global virtual name of the counter type. This name is expected to have the format /objectname/countername.
- *type*: [in] The type of the counters of this *counter_type*.

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
counter_status install_counter_type (std::string const &name, counter_type
                                         type, std::string const &helptext, std::string
                                         const &uom = "", std::uint32_t version =
                                         HPX_PERFORMANCE_COUNTER_V1, error_code
                                         &ec = throws)
```

Install a new performance counter type in a way, which will uninstall it automatically during shutdown.

The function `install_counter_type` will register a new counter type based on the provided `counter_type_info`. The counter type will be automatically unregistered during system shutdown.

Return If successful, this function returns `status_valid_data`, otherwise it will either throw an exception or return an `error_code` from the enum `counter_status` (also, see note related to parameter `ec`).

Note The counter type registry is a locality based service. You will have to register each counter type on every locality where a corresponding performance counter will be created.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `name`: [in] The global virtual name of the counter type. This name is expected to have the format `/objectname/countername`.
- `type`: [in] The type of the counters of this `counter_type`.
- `helptext`: [in] A longer descriptive text shown to the user to explain the nature of the counters created from this type.
- `uom`: [in] The unit of measure for the new performance counter type.
- `version`: [in] The version of the counter type. This is currently expected to be set to `HPX_PERFORMANCE_COUNTER_V1`.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
counter_status install_counter_type (std::string const &name, counter_type type,
                                         std::string const &helptext, create_counter_func
                                         const &create_counter, discover_counters_func
                                         const &discover_counters, std::uint32_t version =
                                         HPX_PERFORMANCE_COUNTER_V1, std::string
                                         const &uom = "", error_code &ec = throws)
```

Install a new generic performance counter type in a way, which will uninstall it automatically during shutdown.

The function `install_counter_type` will register a new generic counter type based on the provided `counter_type_info`. The counter type will be automatically unregistered during system shutdown.

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Return If successful, this function returns `status_valid_data`, otherwise it will either throw an exception or return an `error_code` from the enum `counter_status` (also, see note related to parameter `ec`).

Note The counter type registry is a locality based service. You will have to register each counter type on every locality where a corresponding performance counter will be created.

Parameters

- name: [in] The global virtual name of the counter type. This name is expected to have the format /objectname/countername.
- type: [in] The type of the counters of this counter_type.
- helptext: [in] A longer descriptive text shown to the user to explain the nature of the counters created from this type.
- version: [in] The version of the counter type. This is currently expected to be set to HPX_PERFORMANCE_COUNTER_V1.
- create_counter: [in] The function which will be called to create a new instance of this counter type.
- discover_counters: [in] The function will be called to discover counter instances which can be created.
- uom: [in] The unit of measure of the counter type (default: "")
- ec: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

namespace this_thread

Functions

```
threads::thread_state_ex_enum suspend(threads::thread_state_enum state, threads::thread_id_type
                                      const &id, util::thread_description const &description
                                      = util::thread_description("this_thread::suspend"), error_code
                                      &ec = throws)
```

The function *suspend* will return control to the thread manager (suspends the current thread). It sets the new state of this thread to the thread state passed as the parameter.

Note Must be called from within a HPX-thread.

Exceptions

- If: &ec != &throws, never throws, but will set ec to an appropriate value when an error occurs. Otherwise, this function will throw an *hpx::exception* with an error code of *hpx::yield_aborted* if it is signaled with *wait_aborted*. If called outside of a HPX-thread, this function will throw an *hpx::exception* with an error code of *hpx::null_thread_id*. If this function is called while the thread-manager is not running, it will throw an *hpx::exception* with an error code of *hpx::invalid_status*.

```
threads::thread_state_ex_enum suspend(threads::thread_state_enum state = threads::pending,
                                      util::thread_description const &description =
                                      util::thread_description("this_thread::suspend"), error_code
                                      &ec = throws)
```

The function *suspend* will return control to the thread manager (suspends the current thread). It sets the new state of this thread to the thread state passed as the parameter.

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_state_ex_enum suspend(util::steady_time_point           const      &abs_time,
                                      threads::thread_id_type       const      &id,
                                      util::thread_description     const      &description =
                                      util::thread_description("this_thread::suspend"), error_code
                                      &ec = throws)
```

The function `suspend` will return control to the thread manager (suspends the current thread). It sets the new state of this thread to `suspended` and schedules a wakeup for this threads at the given time.

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_state_ex_enum suspend(util::steady_time_point           const      &abs_time,
                                      util::thread_description const      &description =
                                      util::thread_description("this_thread::suspend"), error_code
                                      &ec = throws)
```

The function `suspend` will return control to the thread manager (suspends the current thread). It sets the new state of this thread to `suspended` and schedules a wakeup for this threads at the given time.

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_state_ex_enum suspend(util::steady_duration            const      &rel_time,
                                      util::thread_description const      &description =
                                      util::thread_description("this_thread::suspend"), error_code
                                      &ec = throws)
```

The function `suspend` will return control to the thread manager (suspends the current thread). It sets the new state of this thread to `suspended` and schedules a wakeup for this threads after the given duration.

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of

`hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_state_ex_enum suspend(util::steady_duration           const      &rel_time,
                                      threads::thread_id_type       const      &id,
                                      util::thread_description     const      &description = util::thread_description("this_thread::suspend"),
                                      &error_code                 &ec = throws)
```

The function `suspend` will return control to the thread manager (suspends the current thread). It sets the new state of this thread to `suspended` and schedules a wakeup for this threads after the given duration.

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_state_ex_enum suspend(std::uint64_t ms, util::thread_description const &description = util::thread_description("this_thread::suspend"),
                                      &error_code &ec = throws)
```

The function `suspend` will return control to the thread manager (suspends the current thread). It sets the new state of this thread to `suspended` and schedules a wakeup for this threads after the given time (specified in milliseconds).

Note Must be called from within a HPX-thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::executors::current_executor get_executor(error_code &ec = throws)
```

Returns a reference to the executor which was used to create the current thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

```
threads::thread_pool_base *get_pool(error_code &ec = throws)
```

Returns a pointer to the pool that was used to run the current thread

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

`namespace threads`

Functions

`thread_self &get_self()`

The function `get_self` returns a reference to the (OS thread specific) self reference to the current HPX thread.

`thread_self *get_self_ptr()`

The function `get_self_ptr` returns a pointer to the (OS thread specific) self reference to the current HPX thread.

`thread_self_impl_type *get_ctx_ptr()`

The function `get_ctx_ptr` returns a pointer to the internal data associated with each coroutine.

`thread_self *get_self_ptr_checked(error_code &ec = throws)`

The function `get_self_ptr_checked` returns a pointer to the (OS thread specific) self reference to the current HPX thread.

`thread_id_type get_self_id()`

The function `get_self_id` returns the HPX thread id of the current thread (or zero if the current thread is not a HPX thread).

`thread_data *get_self_id_data()`

The function `get_self_id_data` returns the data of the HPX thread id associated with the current thread (or `nullptr` if the current thread is not a HPX thread).

`thread_id_type get_parent_id()`

The function `get_parent_id` returns the HPX thread id of the current thread's parent (or zero if the current thread is not a HPX thread).

Note This function will return a meaningful value only if the code was compiled with `HPX_HAVE_THREAD_PARENT_REFERENCE` being defined.

`std::size_t get_parent_phase()`

The function `get_parent_phase` returns the HPX phase of the current thread's parent (or zero if the current thread is not a HPX thread).

Note This function will return a meaningful value only if the code was compiled with `HPX_HAVE_THREAD_PARENT_REFERENCE` being defined.

`std::size_t get_self_stacksize()`

The function `get_self_stacksize` returns the stack size of the current thread (or zero if the current thread is not a HPX thread).

`std::uint32_t get_parent_locality_id()`

The function `get_parent_locality_id` returns the id of the locality of the current thread's parent (or zero if the current thread is not a HPX thread).

Note This function will return a meaningful value only if the code was compiled with `HPX_HAVE_THREAD_PARENT_REFERENCE` being defined.

`std::uint64_t get_self_component_id()`

The function `get_self_component_id` returns the lva of the component the current thread is acting on

Note This function will return a meaningful value only if the code was compiled with `HPX_HAVE_THREAD_TARGET_ADDRESS` being defined.

`std::int64_t get_thread_count (thread_state_enum state = unknown)`

The function `get_thread_count` returns the number of currently known threads.

Note If `state == unknown` this function will not only return the number of currently existing threads, but will add the number of registered task descriptions (which have not been converted into threads yet).

Parameters

- `state`: [in] This specifies the thread-state for which the number of threads should be retrieved.

`std::int64_t get_thread_count (thread_priority priority, thread_state_enum state = unknown)`

The function `get_thread_count` returns the number of currently known threads.

Note If `state == unknown` this function will not only return the number of currently existing threads, but will add the number of registered task descriptions (which have not been converted into threads yet).

Parameters

- `priority`: [in] This specifies the thread-priority for which the number of threads should be retrieved.
- `state`: [in] This specifies the thread-state for which the number of threads should be retrieved.

`bool enumerate_threads (util::function_nonsr<bool> thread_id_type`

`> const &f, thread_state_enum state = unknown)` The function `enumerate_threads` will invoke the given function `f` for each thread with a matching thread state.

Parameters

- `f`: [in] The function which should be called for each matching thread. Returning ‘false’ from this function will stop the enumeration process.
- `state`: [in] This specifies the thread-state for which the threads should be enumerated.

`thread_state set_thread_state (thread_id_type const &id, thread_state_enum state = pending, thread_state_ex_enum stateex = wait_signaled, thread_priority priority = thread_priority_normal, bool retry_on_active = true, hpx::error_code &ec = throws)`

Set the thread state of the `thread` referenced by the `thread_id id`.

Note If the thread referenced by the parameter `id` is in `thread_state::active` state this function schedules a new thread which will set the state of the thread as soon as its not active anymore. The function returns `thread_state::active` in this case.

Return This function returns the previous state of the thread referenced by the *id* parameter. It will return one of the values as defined by the *thread_state* enumeration. If the thread is not known to the thread-manager the return value will be *thread_state::unknown*.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread the state should be modified for.
- *state*: [in] The new state to be set for the thread referenced by the *id* parameter.
- *stateex*: [in] The new extended state to be set for the thread referenced by the *id* parameter.
- *priority*:
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
thread_id_type set_thread_state(thread_id_type const &id, util::steady_time_point const &abs_time, std::atomic<bool> *started, thread_state_enum state = pending, thread_state_ex_enum stateex = wait_timeout, thread_priority priority = thread_priority_normal, bool retry_on_active = true, error_code &ec = throws)
```

Set the thread state of the *thread* referenced by the thread_id *id*.

Set a timer to set the state of the given *thread* to the given new value after it expired (at the given time)

Return

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread the state should be modified for.
- *abs_time*: [in] Absolute point in time for the new thread to be run
- *started*: [in,out] A helper variable allowing to track the state of the timer helper thread
- *state*: [in] The new state to be set for the thread referenced by the *id* parameter.
- *stateex*: [in] The new extended state to be set for the thread referenced by the *id* parameter.
- *priority*:
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
thread_id_type set_thread_state(thread_id_type const &id, util::steady_time_point const &abs_time, thread_state_enum state = pending, thread_state_ex_enum stateex = wait_timeout, thread_priority priority = thread_priority_normal, bool retry_on_active = true, error_code &ec = throws)
```

```
thread_id_type set_thread_state(thread_id_type const &id, util::steady_duration const &rel_time, thread_state_enum state = pending, thread_state_ex_enum stateex = wait_timeout, thread_priority priority = thread_priority_normal, bool retry_on_active = true, error_code &ec = throws)
```

Set the thread state of the *thread* referenced by the thread_id *id*.

Set a timer to set the state of the given *thread* to the given new value after it expired (after the given duration)

Return

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread the state should be modified for.
- *rel_time*: [in] Time duration after which the new thread should be run
- *state*: [in] The new state to be set for the thread referenced by the *id* parameter.
- *stateex*: [in] The new extended state to be set for the thread referenced by the *id* parameter.
- *priority*:
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
util::thread_description get_thread_description(thread_id_type const &id, error_code &ec = throws)
```

The function *get_thread_description* is part of the thread related API allows to query the description of one of the threads known to the thread-manager.

Return This function returns the description of the thread referenced by the *id* parameter. If the thread is not known to the thread-manager the return value will be the string “<unknown>”.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread being queried.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
util::thread_description set_thread_description(thread_id_type const &id,  
util::thread_description const &desc =  
util::thread_description(), error_code &ec = throws)
```

```
util::thread_description get_thread_lco_description(thread_id_type const &id, error_code &ec = throws)
```

```
util::thread_description set_thread_lco_description(thread_id_type const &id,  
util::thread_description const &desc =  
util::thread_description(), error_code &ec = throws)
```

```
thread_state get_thread_state(thread_id_type const &id, error_code &ec = throws)
```

The function *get_thread_backtrace* is part of the thread related API allows to query the currently stored thread back trace (which is captured during thread suspension).

Return This function returns the currently captured stack back trace of the thread referenced by the *id* parameter. If the thread is not known to the thread-manager the return value will be the zero.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*. The function *get_thread_state* is part of the thread related API. It queries the state of one of the threads known to the thread-manager.

Return This function returns the thread state of the thread referenced by the *id* parameter. If the thread is not known to the thread-manager the return value will be *terminated*.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread being queried.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

Parameters

- *id*: [in] The thread id of the thread the state should be modified for.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`std::size_t get_thread_phase(thread_id_type const &id, error_code &ec = throws)`

The function *get_thread_phase* is part of the thread related API. It queries the phase of one of the threads known to the thread-manager.

Return This function returns the thread phase of the thread referenced by the *id* parameter. If the thread is not known to the thread-manager the return value will be ~0.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread the phase should be modified for.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`bool get_thread_interruption_enabled(thread_id_type const &id, error_code &ec = throws)`

Returns whether the given thread can be interrupted at this point.

Return This function returns *true* if the given thread can be interrupted at this point in time. It will return *false* otherwise.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread which should be queried.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`bool set_thread_interruption_enabled(thread_id_type const &id, bool enable, error_code &ec = throws)`

Set whether the given thread can be interrupted at this point.

Return This function returns the previous value of whether the given thread could have been interrupted.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread which should receive the new value.
- *enable*: [in] This value will determine the new interruption enabled status for the given thread.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
bool get_thread_interruption_requested(thread_id_type const &id, error_code &ec = throws)
```

Returns whether the given thread has been flagged for interruption.

Return This function returns *true* if the given thread was flagged for interruption. It will return *false* otherwise.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread which should be queried.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void interrupt_thread(thread_id_type const &id, bool flag, error_code &ec = throws)
```

Flag the given thread for interruption.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread which should be interrupted.
- *flag*: [in] The flag encodes whether the thread should be interrupted (if it is *true*), or ‘uninterrupted’ (if it is *false*).
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void interrupt_thread(thread_id_type const &id, error_code &ec = throws)
```

```
void interruption_point(thread_id_type const &id, error_code &ec = throws)
```

Interrupt the current thread at this point if it was canceled. This will throw a *thread_interrupted* exception, which will cancel the thread.

Note As long as *ec* is not pre-initialized to *hpx::throws* this function doesn't throw but returns the result code using the parameter *ec*. Otherwise it throws an instance of *hpx::exception*.

Parameters

- *id*: [in] The thread id of the thread which should be interrupted.
- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

`threads::thread_priority get_thread_priority (thread_id_type const &id, error_code &ec = throws)`

Return priority of the given thread

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `id`: [in] The thread id of the thread whose priority is queried.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::ptrdiff_t get_stack_size (thread_id_type const &id, error_code &ec = throws)`

Return stack size of the given thread

Note As long as `ec` is not pre-initialized to `hpx::throws` this function doesn't throw but returns the result code using the parameter `ec`. Otherwise it throws an instance of `hpx::exception`.

Parameters

- `id`: [in] The thread id of the thread whose priority is queried.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`threads::executors::current_executor get_executor (thread_id_type const &id, error_code &ec = throws)`

Returns a reference to the executor which was used to create the given thread.

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

`threads::thread_pool_base *get_pool (thread_id_type const &id, error_code &ec = throws)`

Returns a pointer to the pool that was used to run the current thread

Exceptions

- If: `&ec != &throws`, never throws, but will set `ec` to an appropriate value when an error occurs. Otherwise, this function will throw an `hpx::exception` with an error code of `hpx::yield_aborted` if it is signaled with `wait_aborted`. If called outside of a HPX-thread, this function will throw an `hpx::exception` with an error code of `hpx::null_thread_id`. If this function is called while the thread-manager is not running, it will throw an `hpx::exception` with an error code of `hpx::invalid_status`.

namespace policies

Enums

enum scheduler_mode

This enumeration describes the possible modes of a scheduler.

Values:

nothing_special = 0x000

can be used to disable all other options.

As the name suggests, this option

do_background_work = 0x001

The scheduler will periodically call a provided callback function from a special HPX thread to enable performing background-work, for instance driving networking progress or garbage-collect AGAS.

reduce_thread_priority = 0x002

os-thread driving the scheduler will be reduced below normal.

The kernel priority of the

delay_exit = 0x004

The scheduler will wait for some unspecified amount of time before exiting the scheduling loop while being terminated to make sure no other work is being scheduled during processing the shutdown request.

fast_idle_mode = 0x008

Some schedulers have the capability to act as ‘embedded’ schedulers. In this case it needs to periodically invoke a provided callback into the outer scheduler more frequently than normal. This option enables this behavior.

enable_elasticity = 0x010

This option allows for the scheduler to dynamically increase and reduce the number of processing units it runs on. Setting this value not succeed for schedulers that do not support this functionality.

enable_stealing = 0x020

that support work thread/stealing to enable/disable it

This option allows schedulers

enable_stealing numa = 0x040

that support it to disallow stealing between numa domains

This option allows schedulers

assign_work_round_robin = 0x080

that support it to add tasks round robin to queues on each core

This option tells schedulers

assign_work_thread_parent = 0x100

This option tells schedulers that support it to add tasks round to the same core/queue that the parent task is running on

steal_high_priority_first = 0x200

This option tells schedulers that support it to always (try to) steal high priority tasks from other queues before finishing their own lower priority tasks

steal_after_local = 0x400

This option tells schedulers that support it to steal tasks only when their local queues are empty

enable_idle_backoff = 0x800

schedulers to explicitly disable exponential idle-back off

This option allows for certain

```
default_mode = do_background_work | reduce_thread_priority | delay_exit | enable_stealing | enable_stealing_numa | ...  
This option represents the default mode.
```

```
all_flags = do_background_work | reduce_thread_priority | delay_exit | fast_idle_mode | enable_elasticity | enable_stee...  
namespace traits  
namespace util
```

Functions

```
void attach_debugger()
```

Tries to break an attached debugger, if not supported a loop is invoked which gives enough time to attach a debugger manually.

```
void may_attach_debugger(std::string const &category)
```

Attaches a debugger if category is equal to the configuration entry hpx.attach-debugger.

```
template<typename Mapper, typename... T><unspecified> hpx::util::map_pack(Mapper && map...
```

Maps the pack with the given mapper.

This function tries to visit all plain elements which may be wrapped in:

- homogeneous containers (`std::vector`, `std::list`)
- heterogeneous containers (`hpx::tuple`, `std::pair`, `std::array`) and re-assembles the pack with the result of the mapper. Mapping from one type to a different one is supported.

Elements that aren't accepted by the mapper are routed through and preserved through the hierarchy.

```
// Maps all integers to floats  
map_pack([](int value) {  
    return float(value);  
},  
1, hpx::util::make_tuple(2, std::vector<int>{3, 4}), 5);
```

Return The mapped element or in case the pack contains multiple elements, the pack is wrapped into a `hpx::tuple`.

Exceptions

- `std::exception`: like objects which are thrown by an invocation to the mapper.

Parameters

- `mapper`: A callable object, which accept an arbitrary type and maps it to another type or the same one.
- `pack`: An arbitrary variadic pack which may contain any type.

```
template<typename Visitor, typename ...T>  
auto traverse_pack_async(Visitor &&visitor, T &&... pack)
```

Traverses the pack with the given visitor in an asynchronous way.

This function works in the same way as `traverse_pack`, however, we are able to suspend and continue the traversal at later time. Thus we require a visitor callable object which provides three `operator()` overloads as depicted by the code sample below:

```
struct my_async_visitor
{
    template <typename T>
    bool operator()(async_traverse_visit_tag, T&& element)
    {
        return true;
    }

    template <typename T, typename N>
    void operator()(async_traverse_detach_tag, T&& element, N&& next)
    {
    }

    template <typename T>
    void operator()(async_traverse_complete_tag, T&& pack)
    {
    }
};
```

See `traverse_pack` for a detailed description about the traversal behavior and capabilities.

Return A `hpx::intrusive_ptr` that references an instance of the given visitor object.

Parameters

- `visitor`: A visitor object which provides the three `operator()` overloads that were described above. Additionally the visitor must be compatible for referencing it from a `hpx::intrusive_ptr`. The visitor should must have a virtual destructor!
- `pack`: The arbitrary parameter pack which is traversed asynchronously. Nested objects inside containers and tuple like types are traversed recursively.

```
template<typename Allocator, typename Visitor, typename ...T>
auto traverse_pack_async_allocator(Allocator const &alloc, Visitor &&visitor, T&&...
                                pack)
```

Traverses the pack with the given visitor in an asynchronous way.

This function works in the same way as `traverse_pack`, however, we are able to suspend and continue the traversal at later time. Thus we require a visitor callable object which provides three `operator()` overloads as depicted by the code sample below:

```
struct my_async_visitor
{
    template <typename T>
    bool operator()(async_traverse_visit_tag, T&& element)
    {
        return true;
    }

    template <typename T, typename N>
    void operator()(async_traverse_detach_tag, T&& element, N&& next)
    {
    }

    template <typename T>
    void operator()(async_traverse_complete_tag, T&& pack)
    {
    }
};
```

See `traverse_pack` for a detailed description about the traversal behavior and capabilities.

Return A `hpx::intrusive_ptr` that references an instance of the given visitor object.

Parameters

- `visitor`: A visitor object which provides the three `operator()` overloads that were described above. Additionally the visitor must be compatible for referencing it from a `hpx::intrusive_ptr`. The visitor should have a virtual destructor!
- `pack`: The arbitrary parameter pack which is traversed asynchronously. Nested objects inside containers and tuple like types are traversed recursively.
- `alloc`: Allocator instance to use to create the traversal frame.

```
template<typename ...Args>
auto unwrap(Args&&... args)
```

A helper function for retrieving the actual result of any `hpx::lcos::future` like type which is wrapped in an arbitrary way.

Unwraps the given pack of arguments, so that any `hpx::lcos::future` object is replaced by its future result type in the argument pack:

- `hpx::future<int>` -> `int`
- `hpx::future<std::vector<float>>` -> `std::vector<float>`
- `std::vector<future<float>>` -> `std::vector<float>`

The function is capable of unwrapping `hpx::lcos::future` like objects that are wrapped inside any container or tuple like type, see `hpx::util::map_pack()` for a detailed description about which surrounding types are supported. Non `hpx::lcos::future` like types are permitted as arguments and passed through.

```
// Single arguments
int i1 = hpx::util::unwrap(hpx::lcos::make_ready_future(0));

// Multiple arguments
hpx::tuple<int, int> i2 =
    hpx::util::unwrap(hpx::lcos::make_ready_future(1),
                      hpx::lcos::make_ready_future(2));
```

Note This function unwraps the given arguments until the first traversed nested `hpx::lcos::future` which corresponds to an unwrapping depth of one. See `hpx::util::unwrap_n()` for a function which unwraps the given arguments to a particular depth or `hpx::util::unwrap_all()` that unwraps all future like objects recursively which are contained in the arguments.

Return Depending on the count of arguments this function returns a `hpx::util::tuple` containing the unwrapped arguments if multiple arguments are given. In case the function is called with a single argument, the argument is unwrapped and returned.

Parameters

- `args`: the arguments that are unwrapped which may contain any arbitrary future or non future type.

Exceptions

- `std::exception`: like objects in case any of the given wrapped `hpx::lcos::future` objects were resolved through an exception. See `hpx::lcos::future::get()` for details.

```
template<std::size_t Depth, typename ...Args>
```

```
auto unwrap_n(Args&&... args)
```

An alterntive version of [hpx::util::unwrap\(\)](#), which unwraps the given arguments to a certain depth of hpx::lcos::future like objects.

See [unwrap](#) for a detailed description.

Template Parameters

- Depth: The count of hpx::lcos::future like objects which are unwrapped maximally.

```
template<typename ...Args>
```

```
auto unwrap_all(Args&&... args)
```

An alterntive version of [hpx::util::unwrap\(\)](#), which unwraps the given arguments recursively so that all contained hpx::lcos::future like objects are replaced by their actual value.

See [hpx::util::unwrap\(\)](#) for a detailed description.

```
template<typename T>
```

```
auto unwrapping(T &&callable)
```

Returns a callable object which unwraps its arguments upon invocation using the [hpx::util::unwrap\(\)](#) function and then passes the result to the given callable object.

```
auto callable = hpx::util::unwrapping([](int left, int right) {
    return left + right;
});

int i1 = callable(hpx::lcos::make_ready_future(1),
                   hpx::lcos::make_ready_future(2));
```

See [hpx::util::unwrap\(\)](#) for a detailed description.

Parameters

- *callable*: the callable object which which is called with the result of the corresponding [unwrap](#) function.

```
template<std::size_t Depth, typename T>
```

```
auto unwrapping_n(T &&callable)
```

Returns a callable object which unwraps its arguments upon invocation using the [hpx::util::unwrap_n\(\)](#) function and then passes the result to the given callable object.

See [hpx::util::unwrapping\(\)](#) for a detailed description.

```
template<typename T>
```

```
auto unwrapping_all(T &&callable)
```

Returns a callable object which unwraps its arguments upon invocation using the [hpx::util::unwrap_all\(\)](#) function and then passes the result to the given callable object.

See [hpx::util::unwrapping\(\)](#) for a detailed description.

```
namespace functional
```

```
namespace std
```

```
file migrate_from_storage.hpp
```

```
#include <hpx/config.hpp>#include <hpx/lcos/future.hpp>#include <hpx/allocator_support/allocator_deleter.hpp>#include
<hpx/allocator_support/internal_allocator.hpp>#include <hpx/assertion.hpp>#include
<hpx/concepts/concepts.hpp>#include <hpx/errors.hpp>#include <hpx/functional/bind.hpp>#include
<hpx/functional/function.hpp>#include <hpx/functional/invoke.hpp>#include
<hpx/functional/result_of.hpp>#include <hpx/functional/traits/is_callable.hpp>#include
```

```
<hpx/lcos/detail/future_data.hpp>#include
<hpx/functional/unique_function.hpp>#include
<hpx/synchronization/spinlock.hpp>#include
<hpx/runtime/launch_policy.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include
ity>#include <hpx/runtime/naming_fwd.hpp>#include
<hpx/runtime/agas_fwd.hpp>#include
<hpx/runtime/threads/thread_executor.hpp>#include
<hpx/runtime/threads/thread_data_fwd.hpp>#include
<hpx/coroutines/thread_id_type.hpp>#include
def>#include <memory>#include <hpx/runtime/threads/policies/scheduler_mode.hpp>#include
<hpx/thread_support/atomic_count.hpp>#include
<hpx/topology/cpu_mask.hpp>#include
<hpx/util/thread_description.hpp>#include
<hpx/runtime/actions/preassigned_action_id.hpp>#include <hpx/functional/traits/get_function_address.hpp>#include
<hpx/functional/traits/get_function_annotation.hpp>#include <hpx/functional/traits/is_action.hpp>#include
<iostfwd>#include <chrono>#include
<hpx/config/warnings_suffix.hpp>#include
<hpx/basic_execution/register_locks.hpp>#include
<hpx/runtime/threads/register_thread.hpp>#include <hpx/runtime/threads/policies/scheduler_base.hpp>#include
<hpx/concurrency/cache_line_data.hpp>#include
<hpx/runtime/threads/policies/thread_queue_init_parameters.hpp>#include <limits>#include
<hpx/runtime/threads/thread_init_data.hpp>#include <hpx/runtime/threads_fwd.hpp>#include
<hpx/runtime/threads/thread_pool_base.hpp>#include <hpx/concurrency/barrier.hpp>#include
<hpx/runtime/threads/detail/network_background_callback.hpp>#include <hpx/affinity/affinity_data.hpp>#include
<hpx/runtime/threads/policies/callback_notifier.hpp>#include <deque>#include <exception>#include
<hpx/state.hpp>#include <thread>#include <functional>#include <mutex>#include <vector>#include
<atomic>#include <condition_variable>#include <hpx/runtime/threads/thread_data.hpp>#include
<hpx/basic_execution/this_thread.hpp>#include <hpx/concurrency/spinlock_pool.hpp>#include
<hpx/coroutines/coroutine.hpp>#include <hpx/coroutines/detail/combined_tagged_state.hpp>#include
<hpx/logging.hpp>#include <hpx/runtime/get_locality_id.hpp>#include <hpx/util/backtrace.hpp>#include
<forward_list>#include <stack>#include <hpx/runtime/threads/thread_data_stackful.hpp>#include
<hpx/runtime/threads/execution_agent.hpp>#include <hpx/basic_execution/agent_base.hpp>#include
<hpx/basic_execution/context_base.hpp>#include <hpx/basic_execution/resource_base.hpp>#include
<hpx/coroutines/detail/coroutine_impl.hpp>#include <hpx/coroutines/detail/coroutine_stackful_self.hpp>#include
<hpx/runtime/threads/thread_data_stackless.hpp>#include <hpx/coroutines/stackless_coroutine.hpp>#include
<hpx/thread_support/assert_owns_lock.hpp>#include <hpx/traits/future_access.hpp>#include
<hpx/traits/future_traits.hpp>#include <hpx/traits/is_future.hpp>#include <boost/ref.hpp>#include
<hpx/traits/get_remote_result.hpp>#include <hpx/type_support/decay.hpp>#include
<hpx/type_support/unused.hpp>#include <hpx/util/annotated_function.hpp>#include
<boost/container/small_vector.hpp>#include <hpx/lcos/detail/future_traits.hpp>#include
<hpx/type_support/always_void.hpp>#include <iterator>#include <hpx/lcos_fwd.hpp>#include
<hpx/traits/is_component.hpp>#include <hpx/traits/promise_local_result.hpp>#include
<hpx/traits/promise_remote_result.hpp>#include <hpx/serialization/detail/polymorphic_nonintrusive_factory.hpp>#include
<hpx/traits/acquire_shared_state.hpp>#include <hpx/iterator_support/traits/is_range.hpp>#include
<hpx/iterator_support/range.hpp>#include <hpx/traits/is_future_range.hpp>#include
<hpx/util/detail/reserve.hpp>#include <algorithm>#include <hpx/traits/future_then_result.hpp>#include
<hpx/execution/traits/is_executor.hpp>#include <hpx/type_support/identity.hpp>#include
<hpx/type_support/lazy_conditional.hpp>#include <hpx/type_support/pack.hpp>#include
<hpx/traits/is_launch_policy.hpp>#include <hpx/execution/traits/executor_traits.hpp>#include
<hpx/type_support/lazy_enable_if.hpp>#include <hpx/type_support/void_guard.hpp>#include
<hpx/util/serialize_exception.hpp>#include <hpx/lcos/local/packaged_continuation.hpp>#include
<hpx/execution/executors/execution.hpp>#include <hpx/execution/executors/post_policy_dispatch.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/components/component_storage/server/migrate_from_storage.hpp>
```

```

file migrate_to_storage.hpp
#include <hpx/config.hpp>#include <hpx/lcos/future.hpp>#include <hpx/runtime/components/client_base.hpp>#include
<hpx/assertion.hpp>#include <hpx/errors.hpp>#include <hpx/functional/bind_back.hpp>#include
<hpx/memory/intrusive_ptr.hpp>#include <hpx/runtime/agas/interface.hpp>#include
<hpx/runtime/components/component_type.hpp>#include <hpx/functional/unique_function.hpp>#include
<hpx/preprocessor/cat.hpp>#include <hpx/preprocessor/expand.hpp>#include
<hpx/preprocessor/nargs.hpp>#include <hpx/preprocessor/stringize.hpp>#include
<hpx/preprocessor/strip_parens.hpp>#include <hpx/runtime/naming_fwd.hpp>#include
<hpx/thread_support/atomic_count.hpp>#include <hpx/traits/component_type_database.hpp>#include
<cstdint>#include <hpx/type_support/decay.hpp>#include <string>#include
<hpx/runtime/components_fwd.hpp>#include <hpx/traits/managed_component_policies.hpp>#include
<hpx/type_support/always_void.hpp>#include <cstddef>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/name.hpp>#include <hpx/allocator_support/internal_allocator.hpp>#include
<hpx/basic_execution/register_locks.hpp>#include <hpx/concurrency/itt_notify.hpp>#include
<hpx/concurrency/spinlock_pool.hpp>#include <hpx/runtime/naming/id_type.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include <hpx/serialization/traits/is_bitwise_serializable.hpp>#include
<hpx/traits/get_remote_result.hpp>#include <hpx/traits/promise_local_result.hpp>#include
<hpx/synchronization/detail/yield_k.hpp>#include <functional>#include <iostfwd>#include
<mutex>#include <vector>#include <hpx/config/warnings_prefix.hpp>#include
<hpx/runtime/naming/id_type_impl.hpp>#include <hpx/config/warnings_suffix.hpp>#include
<boost/dynamic_bitset.hpp>#include <map>#include <utility>#include <hpx/runtime/components/make_client.hpp>#include
<hpx/traits/is_client.hpp>#include <type_traits>#include <hpx/runtime/components/stubs/stub_base.hpp>#include
<hpx/lcos/async_fwd.hpp>#include <hpx/lcos/detail/async_colocated_fwd.hpp>#include
<hpx/runtime/actions/basic_action_fwd.hpp>#include <hpx/traits/extract_action.hpp>#include
<hpx/traits/is_continuation.hpp>#include <hpx/lcos/detail/async_implementations_fwd.hpp>#include
<hpx/runtime/naming/unmanaged.hpp>#include <hpx/serialization/serialize.hpp>#include
<hpx/traits/acquire_future.hpp>#include <hpx/iterator_support/range.hpp>#include
<hpx/util/detail/reserve.hpp>#include <hpx/concepts/has_member_xxx.hpp>#include
<hpx/traits/is_future.hpp>#include <hpx/traits/is_future_range.hpp>#include
<hpx/iterator_support/traits/is_range.hpp>#include <algorithm>#include <array>#include <iterator>#include
<hpx/traits/action_remote_result.hpp>#include <hpx/traits/future_access.hpp>#include
<hpx/traits/future_traits.hpp>#include <exception>#include <hpx/traits/is_component.hpp>#include
<hpx/components/component_storage/component_storage.hpp>#include <hpx/components/component_storage/server/migrate_t
```

file hpx_finalize.hpp

```
#include <hpx/config.hpp>#include <hpx/errors.hpp>
```

file hpx_init.hpp

```
#include <hpx/config.hpp>#include <hpx/hpx_finalize.hpp>#include <hpx/errors.hpp>#include
<hpx/hpx_suspend.hpp>#include <hpx/program_options.hpp>#include <hpx/runtime/runtime_mode.hpp>#include
<string>#include <hpx/runtime/shutdown_function.hpp>#include <hpx/functional/unique_function.hpp>#include
<hpx/runtime/startup_function.hpp>#include <hpx/functional/function.hpp>#include <cstddef>#include
<vector>
```

file hpx_start.hpp

```
#include <hpx/config.hpp>#include <hpx/hpx_finalize.hpp>#include <hpx/program_options.hpp>#include
<hpx/runtime/runtime_mode.hpp>#include <hpx/runtime/shutdown_function.hpp>#include
<hpx/runtime/startup_function.hpp>#include <hpx/functional/function.hpp>#include <cstddef>#include
<string>#include <vector>
```

file hpx_suspend.hpp

```
#include <hpx/errors.hpp>
```

file split_future.hpp

file wait_all.hpp

```
file wait_any.hpp
file wait_each.hpp
file wait_some.hpp
file when_all.hpp
file when_any.hpp
file when_each.hpp
file when_some.hpp
file task_block.hpp
#include <hpx/config.hpp>#include <hpx/async.hpp>#include <hpx/async_launch_policy_dispatch.hpp>#include
<hpx/lcos/async_fwd.hpp>#include <hpx/lcos_fwd.hpp>#include <hpx/assertion.hpp>#include
<hpx/coroutines/thread_enums.hpp>#include <hpx/functional/deferred_call.hpp>#include
<hpx/functional/invoke.hpp>#include <hpx/functional/traits/is_action.hpp>#include
<hpx/lcos/future.hpp>#include <hpx/lcos/local/futures_factory.hpp>#include
<hpx/allocator_support/allocator_deleter.hpp>#include <hpx/allocator_support/internal_allocator.hpp>#include
<hpx/errors.hpp>#include <hpx/lcos/detail/future_data.hpp>#include <hpx/memory/intrusive_ptr.hpp>#include
<hpx/runtime/get_worker_thread_num.hpp>#include <cstddef>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/threads/thread_data_fwd.hpp>#include <hpx/runtime/threads/thread_helpers.hpp>#include
<hpx/traits/future_access.hpp>#include <hpx/util/thread_description.hpp>#include
<hpx/execution/executors/execution.hpp>#include <cstdint>#include <exception>#include
<memory>#include <type_traits>#include <utility>#include
<hpx/runtime/threads/thread_pool_base.hpp>#include <hpx/lcos/async.hpp>#include
<hpx/functional/bind_back.hpp>#include <hpx/lcos/detail/async_implementations.hpp>#include
<hpx/lcos/detail/async_implementations_fwd.hpp>#include <hpx/lcos/packaged_action.hpp>#include
<hpx/lcos/promise.hpp>#include <hpx/lcos/detail/promise_base.hpp>#include
<hpx/functional/unique_function.hpp>#include <hpx/lcos/detail/promise_lco.hpp>#include
<hpx/lcos/base_lco_with_value.hpp>#include <hpx/lcos/base_lco.hpp>#include
<hpx/runtime/actions/basic_action.hpp>#include <hpx/datastructures/tuple.hpp>#include
<hpx/functional/invoke_fused.hpp>#include <hpx/lcos/sync_fwd.hpp>#include
<hpx/type_support/decay.hpp>#include <hpx/logging.hpp>#include <hpx/preprocessor/cat.hpp>#include
<hpx/preprocessor/expand.hpp>#include <hpx/preprocessor/nargs.hpp>#include
<hpx/preprocessor/stringize.hpp>#include <hpx/runtime/actions_fwd.hpp>#include
<hpx/runtime/actions/continuation_fwd.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/runtime/components/pinned_ptr.hpp>#include <hpx/naming_base.hpp>#include
<hpx/runtime/get_lva.hpp>#include <hpx/runtime/components_fwd.hpp>#include
<hpx/runtime/naming/address.hpp>#include <hpx/runtime/components/component_type.hpp>#include
<hpx/runtime/naming/name.hpp>#include <hpx/runtime/naming_fwd.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include <hpx/serialization/traits/is_bitwise_serializable.hpp>#include
<iostfwd>#include <hpx/config/warnings_prefix.hpp>#include <hpx/config/warnings_suffix.hpp>#include
<hpx/traits/is_component.hpp>#include <hpx/traits/action_decorate_function.hpp>#include
<hpx/concepts/has_xxx.hpp>#include <hpx/type_support/detail/wrap_int.hpp>#include
<hpx/traits/component_pin_support.hpp>#include <hpx/runtime/parcelset_fwd.hpp>#include
<hpx/serialization/base_object.hpp>#include <hpx/serialization/input_archive.hpp>#include
<hpx/serialization/output_archive.hpp>#include <hpx/runtime/threads/thread_init_data.hpp>#include
<hpx/traits/action_remote_result.hpp>#include <hpx/debugging/demangle_helper.hpp>#include
<hpx/runtime/actions/basic_action_fwd.hpp>#include <hpx/runtime/actions/continuation.hpp>#include
<hpx/runtime/actions/action_priority.hpp>#include <hpx/traits/action_priority.hpp>#include
<hpx/traits/extract_action.hpp>#include <hpx/runtime/actions/trigger.hpp>#include
<hpx/type_support/unused.hpp>#include <hpx/runtime/agas/interface.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/serialization/serialize.hpp>#include
<hpx/runtime/trigger_lco.hpp>#include <hpx/runtime/applier/detail/apply_implementations_fwd.hpp>#include
```

```

<hpx/traits/is_continuation.hpp>#include
<hpx/util/serializable_unique_function.hpp>#include <hpx/util/detail/serializable_basic_function.hpp>#include
<hpx/functional/detail/basic_function.hpp>#include <hpx/functional/detail/vtable/function_vtable.hpp>#include
<hpx/util/detail/vtable/serializable_function_vtable.hpp>#include <hpx/serialization/detail/polymorphic_intrusive_factory.hpp>
<hpx/functional/detail/function_registration.hpp>#include <hpx/functional/detail/empty_function.hpp>#include
<hpx/util/detail/vtable/serializable_vtable.hpp>#include <hpx/functional/detail/vtable.vtable.hpp>#include
<new>#include <string>#include <hpx/runtime/actions/detail/action_factory.hpp>#include
<hpx/runtime/actions/detail/invocation_count_registry.hpp>#include <hpx/performance_counters/counters_fwd.hpp>#include
<hpx/functional/function.hpp>#include <vector>#include <hpx/hashing/jenkins_hash.hpp>#include
<hpx/type_support/static.hpp>#include <unordered_map>#include <hpx/runtime/actions/preassigned_action_id.hpp>#include
<hpx/runtime/actions/transfer_action.hpp>#include <hpx/runtime/actions/transfer_continuation_action.hpp>#include
<hpx/runtime/parcelset/detail/per_action_data_counter_registry.hpp>#include
<hpx/serialization/tuple.hpp>#include <hpx/runtime_fwd.hpp>#include <hpx/runtime/basename_registration_fwd.hpp>#include
<hpx/components_fwd.hpp>#include <hpx/runtime/components/make_client.hpp>#include
<hpx/runtime/config_entry.hpp>#include <cstdlib>#include <hpx/runtime/find_localities.hpp>#include
<hpx/runtime/get_colocation_id.hpp>#include <hpx/runtime/get_locality_id.hpp>#include
<hpx/runtime/get_locality_name.hpp>#include <hpx/runtime/get_num_localities.hpp>#include
<hpx/runtime/get_os_thread_count.hpp>#include <hpx/runtime/get_thread_name.hpp>#include
<hpx/runtime/report_error.hpp>#include <hpx/runtime/runtime_fwd.hpp>#include
<hpx/runtime/runtime_mode.hpp>#include <hpx/runtime/set_parcel_write_handler.hpp>#include
<hpx/runtime/shutdown_function.hpp>#include <hpx/runtime/startup_function.hpp>#include
<hpx/util_fwd.hpp>#include <hpx/traits/action_stacksize.hpp>#include <hpx/traits/is_distribution_policy.hpp>#include
<hpx/traits/promise_local_result.hpp>#include <hpx/type_support/pack.hpp>#include
<hpx/util/get_and_reset_value.hpp>#include <boost/utility/string_ref.hpp>#include
<atomic>#include <sstream>#include <hpx/runtime/actions/component_action.hpp>#include
<hpx/traits/is_client.hpp>#include <hpx/traits/is_future.hpp>#include <hpx/runtime/components/server/managed_component_b
<hpx/runtime/components/server/component_heap.hpp>#include <hpx/static_reinit/reinitializable_static.hpp>#include
<hpx/runtime/components/server/create_component_fwd.hpp>#include <hpx/runtime/components/server/wrapper_heap.hpp>#i
<hpx/concurrency/itt_notify.hpp>#include <hpx/synchronization/spinlock.hpp>#include
<hpx/util/generate_unique_ids.hpp>#include <hpx/concurrency/spinlock.hpp>#include <mutex>#include
<hpx/util/wrapper_heap_base.hpp>#include <hpx/runtime/components/server/wrapper_heap_list.hpp>#include
<hpx/util/one_size_heap_list.hpp>#include <list>#include <hpx/thread_support/unlock_guard.hpp>#include
<iostream>#include <hpx/traits/managed_component_policies.hpp>#include <stdext>
cept>#include <hpx/plugins/parcel/coalescing_message_handler_registration.hpp>#include
<hpx/runtime/components/server/component_base.hpp>#include <hpx/util/ini.hpp>#include
<hpx/util/to_string.hpp>#include <map>#include <hpx/type_support/void_guard.hpp>#include
<hpx/thread_support/atomic_count.hpp>#include <hpx/traits/component_type_database.hpp>#include
<hpx/local_lcos/promise.hpp>#include <hpx/runtime/applier/apply.hpp>#include
<hpx/runtime/applier/apply_helper.hpp>#include <hpx/state.hpp>#include <hpx/traits/action_continuation.hpp>#include
<hpx/traits/action_decorate_continuation.hpp>#include <hpx/traits/action_schedule_thread.hpp>#include
<hpx/traits/action_select_direct_execution.hpp>#include <thread>#include <chrono>#include
<hpx/runtime/applier/detail/apply_implementations.hpp>#include <hpx/runtime/parcelset/parcel.hpp>#include
<hpx/traits/action_is_target_valid.hpp>#include <hpx/traits/action_was_object_migrated.hpp>#include
<hpx/traits/component_supports_migration.hpp>#include <hpx/format.hpp>#include
<hpx/runtime/components/client_base.hpp>#include <hpx/runtime/parcelset/detail/parcel_await.hpp>#include
<hpx/runtime/parcelset/put_parcel.hpp>#include <hpx/traits/component_type_is_compatible.hpp>#include
<hpx/traits/is_valid_action.hpp>#include <hpx/runtime/applier/apply_callback.hpp>#include
<boost/asio/error.hpp>#include <hpx/runtime/threads/thread.hpp>#include <hpx/runtime/threads/policies/scheduler_base.hpp>
<hpx/runtime/threads/thread_data.hpp>#include <hpx/timing/steady_clock.hpp>#include
<hpx/traits/is_launch_policy.hpp>#include <hpx/type_support/lazy_enable_if.hpp>#include
<hpx/lcos-sync.hpp>#include <hpx/lcos/detail/sync_implementations.hpp>#include
<hpx/lcos/detail/sync_implementations_fwd.hpp>#include <hpx/lcos/async_continue.hpp>#include
<hpx/lcos/async_continue_fwd.hpp>#include <hpx/functional/result_of.hpp>#include
<hpx/traits/promise_remote_result.hpp>#include <hpx/execution/traits/is_executor.hpp>#include

```

```
<hpx/util/bind_action.hpp>#include <hpx/functional/bind.hpp>#include <hpx/functional/traits/is_bind_expression.hpp>#include
<hpx/functional/traits/is_placeholder.hpp>#include <hpx/execution/executors/parallel_executor.hpp>#include
<functional>#include <hpx/lcos/dataflow.hpp>#include <hpx/coroutines/detail/get_stack_pointer.hpp>#include
<hpx/lcos/detail/future_transforms.hpp>#include <hpx/lcos/detail/future_traits.hpp>#include
<hpx/traits/acquire_future.hpp>#include <hpx/traits/acquire_shared_state.hpp>#include
<hpx/util/detail/reserve.hpp>#include <hpx/algorithm>#include <iterator>#include
<hpx/type_support/always_void.hpp>#include <hpx/util/annotated_function.hpp>#include
<hpx/util/pack_traversal_async.hpp>#include <hpx/util/detail/pack_traversal_async_impl.hpp>#include
<hpx/util/detail/container_category.hpp>#include <hpx/iterator_support/traits/is_range.hpp>#include
<hpx/datastructures/traits/is_tuple_like.hpp>#include <boost/ref.hpp>#include
<hpx/lcos/when_all.hpp>#include <hpx/execution/exception_list.hpp>#include
<hpx/execution/execution_policy.hpp>#include <hpx/parallel/util/detail/algorithm_result.hpp>#include
<boost/utility/addressof.hpp>
```

file manage_counter_type.hpp

```
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <hpx/performance_counters/counters_fwd.hpp>#include
<hpx/functional/function.hpp>#include <cstdint>#include <string>#include <vector>
```

file basic_action.hpp

```
#include <hpx/config.hpp>#include <hpx/datastructures/tuple.hpp>#include <hpx/errors.hpp>#include
<hpx/functional/invoke_fused.hpp>#include <hpx/lcos/sync_fwd.hpp>#include
<hpx/logging.hpp>#include <hpx/preprocessor/cat.hpp>#include <hpx/preprocessor/expand.hpp>#include
<hpx/preprocessor/nargs.hpp>#include <hpx/preprocessor/stringize.hpp>#include
<hpx/runtime/actions_fwd.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/runtime/actions/basic_action_fwd.hpp>#include <hpx/runtime/actions/continuation.hpp>#include
<hpx/runtime/actions/detail/action_factory.hpp>#include <hpx/runtime/actions/detail/invocation_count_registry.hpp>#include
<hpx/runtime/actions/preassigned_action_id.hpp>#include <hpx/runtime/actions/transfer_action.hpp>#include
<hpx/runtime/actions/transfer_continuation_action.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/address.hpp>#include <hpx/runtime/naming/id_type.hpp>#include
<hpx/runtime/parcelset/detail/per_action_data_counter_registry.hpp>#include
<hpx/serialization/tuple.hpp>#include <hpx/runtime/threads/thread_data_fwd.hpp>#include
<hpx/coroutines/thread_enums.hpp>#include <hpx/runtime_fwd.hpp>#include
<hpx/traits/action_decorate_function.hpp>#include <hpx/traits/action_priority.hpp>#include
<hpx/traits/action_remote_result.hpp>#include <hpx/traits/action_stacksize.hpp>#include
<hpx/functional/traits/is_action.hpp>#include <hpx/traits/is_distribution_policy.hpp>#include
<hpx/traits/promise_local_result.hpp>#include <hpx/type_support/pack.hpp>#include
<hpx/util/get_and_reset_value.hpp>#include <boost/utility/string_ref.hpp>#include <atomic>#include
<cstdint>#include <cstdint>#include <exception>#include <iostream>#include <string>#include
<type_traits>#include <utility>
```

Defines

HPX_REGISTER_ACTION_DECLARATION(...)

Declare the necessary component action boilerplate code.

The macro `HPX_REGISTER_ACTION_DECLARATION` can be used to declare all the boilerplate code which is required for proper functioning of component actions in the context of HPX.

The parameter `action` is the type of the action to declare the boilerplate for.

This macro can be invoked with an optional second parameter. This parameter specifies a unique name of the action to be used for serialization purposes. The second parameter has to be specified if the first parameter is not usable as a plain (non-qualified) C++ identifier, i.e. the first parameter contains special characters which cannot be part of a C++ identifier, such as '<', '>', or ':'.

```

namespace app
{
    // Define a simple component exposing one action 'print_greeting'
    class HPX_COMPONENT_EXPORT server
        : public hpx::components::simple_component_base<server>
    {
        void print_greeting ()
        {
            hpx::cout << "Hey, how are you?\n" << hpx::flush;
        }

        // Component actions need to be declared, this also defines the
        // type 'print_greeting_action' representing the action.
        HPX_DEFINE_COMPONENT_ACTION(server,
            print_greeting, print_greeting_action);
    };
}

// Declare boilerplate code required for each of the component actions.
HPX_REGISTER_ACTION_DECLARATION(app::server::print_greeting_action);

```

Example:

Note This macro has to be used once for each of the component actions defined using one of the `HPX_DEFINE_COMPONENT_ACTION` macros. It has to be visible in all translation units using the action, thus it is recommended to place it into the header file defining the component.

`HPX_REGISTER_ACTION_DECLARATION(...)`

`HPX_REGISTER_ACTION_DECLARATION_1(action)`

`HPX_REGISTER_ACTION(...)`

Define the necessary component action boilerplate code.

The macro `HPX_REGISTER_ACTION` can be used to define all the boilerplate code which is required for proper functioning of component actions in the context of HPX.

The parameter `action` is the type of the action to define the boilerplate for.

This macro can be invoked with an optional second parameter. This parameter specifies a unique name of the action to be used for serialization purposes. The second parameter has to be specified if the first parameter is not usable as a plain (non-qualified) C++ identifier, i.e. the first parameter contains special characters which cannot be part of a C++ identifier, such as '<', '>', or ':'.

Note This macro has to be used once for each of the component actions defined using one of the `HPX_DEFINE_COMPONENT_ACTION` or `HPX_DEFINE_PLAIN_ACTION` macros. It has to occur exactly once for each of the actions, thus it is recommended to place it into the source file defining the component.

Note Only one of the forms of this macro `HPX_REGISTER_ACTION` or `HPX_REGISTER_ACTION_ID` should be used for a particular action, never both.

`HPX_REGISTER_ACTION_ID(action, actionname, actionid)`

Define the necessary component action boilerplate code and assign a predefined unique id to the action.

The macro `HPX_REGISTER_ACTION` can be used to define all the boilerplate code which is required for proper functioning of component actions in the context of HPX.

The parameter *action* is the type of the action to define the boilerplate for.

The parameter *actionname* specifies an unique name of the action to be used for serialization purposes. The second parameter has to be usable as a plain (non-qualified) C++ identifier, it should not contain special characters which cannot be part of a C++ identifier, such as '<', '>', or ':'.

The parameter *actionid* specifies an unique integer value which will be used to represent the action during serialization.

Note This macro has to be used once for each of the component actions defined using one of the *HPX_DEFINE_COMPONENT_ACTION* or global actions *HPX_DEFINE_PLAIN_ACTION* macros. It has to occur exactly once for each of the actions, thus it is recommended to place it into the source file defining the component.

Note Only one of the forms of this macro *HPX_REGISTER_ACTION* or *HPX_REGISTER_ACTION_ID* should be used for a particular action, never both.

file component_action.hpp

```
#include <hpx/config.hpp>#include <hpx/preprocessor/cat.hpp>#include <hpx/preprocessor/expand.hpp>#include
<hpx/preprocessor/nargs.hpp>#include <hpx/runtime/actions/basic_action.hpp>#include
<hpx/runtime/components/pinned_ptr.hpp>#include <hpx/runtime/naming/address.hpp>#include
<hpx/traits/is_client.hpp>#include <hpx/traits/is_future.hpp>#include <boost/utility/string_ref.hpp>#include
<cstdlib>#include <sstream>#include <string>#include <type_traits>#include <utility>#include
<hpx/config/warnings_prefix.hpp>#include <hpx/config/warnings_suffix.hpp>
```

Defines

HPX_DEFINE_COMPONENT_ACTION(...)

Registers a member function of a component as an action type with HPX.

The macro *HPX_DEFINE_COMPONENT_ACTION* can be used to register a member function of a component as an action type named *action_type*.

The parameter *component* is the type of the component exposing the member function *func* which should be associated with the newly defined action type. The parameter *action_type* is the name of the action type to register with HPX.

```
namespace app
{
    // Define a simple component exposing one action 'print_greeting'
    class HPX_COMPONENT_EXPORT server
        : public hpx::components::simple_component_base<server>
    {
        void print_greeting() const
        {
            hpx::cout << "Hey, how are you?\n" << hpx::flush;
        }

        // Component actions need to be declared, this also defines the
        // type 'print_greeting_action' representing the action.
        HPX_DEFINE_COMPONENT_ACTION(server, print_greeting,
            print_greeting_action);
    };
}
```

Example:

The first argument must provide the type name of the component the action is defined for.

The second argument must provide the member function name the action should wrap.

The default value for the third argument (the typename of the defined action) is derived from the name of the function (as passed as the second argument) by appending ‘_action’. The third argument can be omitted only if the second argument with an appended suffix ‘_action’ resolves to a valid, unqualified C++ type name.

Note The macro `HPX_DEFINE_COMPONENT_ACTION` can be used with 2 or 3 arguments. The third argument is optional.

```
file plain_action.hpp
#include <hpx/config.hpp>#include <hpx/assertion.hpp>#include <hpx/preprocessor/cat.hpp>#include
<hpx/preprocessor/expand.hpp>#include <hpx/preprocessor/nargs.hpp>#include
<hpx/preprocessor/strip_parens.hpp>#include <hpx/runtime/actions/basic_action.hpp>#include
<hpx/runtime/naming/address.hpp>#include <hpx/traits/component_type_database.hpp>#include
<boost/utility/string_ref.hpp>#include <cstdlib>#include <iostream>#include <stdexcept>#include
<string>#include <utility>#include <hpx/config/warnings_prefix.hpp>#include
<hpx/config/warnings_suffix.hpp>
```

Defines

`HPX_DEFINE_PLAIN_ACTION(...)`

Defines a plain action type.

```
namespace app
{
    void some_global_function(double d)
    {
        cout << d;
    }

    // This will define the action type 'app::some_global_action' which
    // represents the function 'app::some_global_function'.
    HPX_DEFINE_PLAIN_ACTION(some_global_function, some_global_action);
}
```

Example:

Note Usually this macro will not be used in user code unless the intent is to avoid defining the `action_type` in global namespace. Normally, the use of the macro `HPX_PLAIN_ACTION` is recommended.

Note The macro `HPX_DEFINE_PLAIN_ACTION` can be used with 1 or 2 arguments. The second argument is optional. The default value for the second argument (the typename of the defined action) is derived from the name of the function (as passed as the first argument) by appending ‘_action’. The second argument can be omitted only if the first argument with an appended suffix ‘_action’ resolves to a valid, unqualified C++ type name.

`HPX_DECLARE_PLAIN_ACTION(...)`

Declares a plain action type.

`HPX_PLAIN_ACTION(...)`

Defines a plain action type based on the given function `func` and registers it with HPX.

The macro *HPX_PLAIN_ACTION* can be used to define a plain action (e.g. an action encapsulating a global or free function) based on the given function *func*. It defines the action type *name* representing the given function. This macro additionally registers the newly defined action type with HPX.

The parameter *func* is a global or free (non-member) function which should be encapsulated into a plain action. The parameter *name* is the name of the action type defined by this macro.

```
namespace app
{
    void some_global_function(double d)
    {
        cout << d;
    }
}

// This will define the action type 'some_global_action' which represents
// the function 'app::some_global_function'.
HPX_PLAIN_ACTION(app::some_global_function, some_global_action);
```

Example:

Note The macro *HPX_PLAIN_ACTION* has to be used at global namespace even if the wrapped function is located in some other namespace. The newly defined action type is placed into the global namespace as well.

Note The macro *HPX_PLAIN_ACTION_ID* can be used with 1, 2, or 3 arguments. The second and third arguments are optional. The default value for the second argument (the typename of the defined action) is derived from the name of the function (as passed as the first argument) by appending '_action'. The second argument can be omitted only if the first argument with an appended suffix '_action' resolves to a valid, unqualified C++ type name. The default value for the third argument is *hpx::components::factory_check*.

Note Only one of the forms of this macro *HPX_PLAIN_ACTION* or *HPX_PLAIN_ACTION_ID* should be used for a particular action, never both.

HPX_PLAIN_ACTION_ID (*func*, *name*, *id*)

Defines a plain action type based on the given function *func* and registers it with HPX.

The macro *HPX_PLAIN_ACTION_ID* can be used to define a plain action (e.g. an action encapsulating a global or free function) based on the given function *func*. It defines the action type *actionname* representing the given function. The parameter *actionid*

The parameter *actionid* specifies an unique integer value which will be used to represent the action during serialization.

The parameter *func* is a global or free (non-member) function which should be encapsulated into a plain action. The parameter *name* is the name of the action type defined by this macro.

The second parameter has to be usable as a plain (non-qualified) C++ identifier, it should not contain special characters which cannot be part of a C++ identifier, such as '<', '>', or ':'.

```
namespace app
{
    void some_global_function(double d)
    {
        cout << d;
    }
}
```

(continues on next page)

(continued from previous page)

```
// This will define the action type 'some_global_action' which represents
// the function 'app::some_global_function'.
HPX_PLAIN_ACTION_ID(app::some_global_function, some_global_action,
some_unique_id);
```

Example:

Note The macro `HPX_PLAIN_ACTION_ID` has to be used at global namespace even if the wrapped function is located in some other namespace. The newly defined action type is placed into the global namespace as well.

Note Only one of the forms of this macro `HPX_PLAIN_ACTION` or `HPX_PLAIN_ACTION_ID` should be used for a particular action, never both.

```
file applier_fwd.hpp
    #include <hpx/config.hpp>

file basename_registration_fwd.hpp
    #include <hpx/config.hpp>#include <hpx/components_fwd.hpp>#include <hpx/lcos_fwd.hpp>#include
    <hpx/runtime/components/make_client.hpp>#include <hpx/runtime/naming/id_type.hpp>#include <cstd-
    def>#include <string>#include <utility>#include <vector>

file binpacking_distribution_policy.hpp
    #include <hpx/config.hpp>#include <hpx/assertion.hpp>#include <hpx/dataflow.hpp>#include
    <hpx/lcos/future.hpp>#include <hpx/lcos/dataflow.hpp>#include <hpx/performance_counters/performance_counter.hpp>#inclu
    <hpx/runtime/components/client_base.hpp>#include <hpx/runtime/launch_policy.hpp>#include
    <hpx/functional/bind_front.hpp>#include <hpx/performance_counters/counters_fwd.hpp>#include
    <hpx/performance_counters/stubs/performance_counter.hpp>#include <hpx/performance_counters/server/base_performance_c
    <hpx/errors.hpp>#include <hpx/lcos/base_lco_with_value.hpp>#include <hpx/performance_counters/counters.hpp>#include
    <hpx/runtime/naming/name.hpp>#include <hpx/serialization/serialization_fwd.hpp>#include
    <hpx/functional/function.hpp>#include <cstddef>#include <cstdint>#include <string>#include <util-
    ity>#include <vector>#include <hpx/performance_counters/performance_counter_base.hpp>#include
    <hpx/runtime/actions/component_action.hpp>#include <hpx/runtime/components/component_type.hpp>#include
    <hpx/runtime/components/server/component.hpp>#include <hpx/allocator_support/internal_allocator.hpp>#include
    <hpx/traits/component_heap_type.hpp>#include <new>#include <hpx/thread_support/atomic_count.hpp>#include
    <hpx/runtime/components/stubs/stub_base.hpp>#include <hpx/runtime/find_here.hpp>#include
    <hpx/runtime/naming/id_type.hpp>#include <hpx/serialization/string.hpp>#include
    <hpx/serialization/vector.hpp>#include <hpx/traits/is_distribution_policy.hpp>#include
    <hpx/functional/bind_back.hpp>#include <hpx/util/unwrap.hpp>#include <hpx/util/detail/unwrap_impl.hpp>#include
    <hpx/traits/future_traits.hpp>#include <hpx/traits/is_future.hpp>#include <hpx/datastructures/traits/is_tuple_like.hpp>#include
    <hpx/functional/invoke.hpp>#include <hpx/functional/invoke_fused.hpp>#include
    <hpx/util/pack_traversal.hpp>#include <hpx/util/detail/pack_traversal_impl.hpp>#include
    <hpx/datastructures/tuple.hpp>#include <hpx/functional/result_of.hpp>#include
    <hpx/functional/traits/is_callable.hpp>#include <hpx/traits/pack_traversal_rebind_container.hpp>#include
    <array>#include <memory>#include <list>#include <type_traits>#include
    <hpx/type_support/always_void.hpp>#include <hpx/type_support/pack.hpp>#include
    <hpx/util/detail/container_category.hpp>#include <hpx/util/detail/reserve.hpp>#include <iterator>#include
    <algorithm>

file colocating_distribution_policy.hpp
    #include <hpx/config.hpp>#include <hpx/lcos/detail/async_colocated.hpp>#include
    <hpx/assertion.hpp>#include <hpx/datastructures/tuple.hpp>#include <hpx/functional/bind.hpp>#include
    <hpx/functional/unique_function.hpp>#include <hpx/lcos/async_continue_fwd.hpp>#include
    <hpx/lcos/async_fwd.hpp>#include <hpx/lcos/detail/async_colocated_fwd.hpp>#include
```

```
<hpx/lcos/future.hpp>#include
<hpx/runtime/agas/primary_namespace.hpp>#include
<hpx/runtime/parcelset_fwd.hpp>#include
<hpx/runtime/components/client_base.hpp>#include
<cstdint>#include <memory>#include
<hpx/config/warnings_prefix.hpp>#include
<hpx/runtime/agas/server/primary_namespace.hpp>#include <hpx/lcos/base_lco_with_value.hpp>#include
<hpx/synchronization/condition_variable.hpp>#include <hpx/runtime/actions/component_action.hpp>#include
<hpx/runtime/components/server/fixed_component_base.hpp>#include <hpx/runtime/applier/applier.hpp>#include
<hpx/runtime/applier_fwd.hpp>#include <hpx/runtime/components/component_type.hpp>#include
<hpx/runtime/naming/name.hpp>#include <hpx/runtime/parcelset/parcel.hpp>#include
<hpx/runtime/threads/thread_data_fwd.hpp>#include <hpx/threadmanager.hpp>#include
<cstdint>#include <hpx/runtime/applier/bind_naming_wrappers.hpp>#include
<hpx/errors.hpp>#include <hpx/runtime/components/server/create_component_fwd.hpp>#include
<hpx/runtime/components_fwd.hpp>#include <hpx/runtime_fwd.hpp>#include
<hpx/traits/is_component.hpp>#include <sstream>#include <type_traits>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/traits/action_message_handler.hpp>#include
<hpx/traits/action_serialization_filter.hpp>#include <hpx/allocator_support/internal_allocator.hpp>#include
<atomic>#include <list>#include <map>#include <mutex>#include <string>#include
<hpx/traits/extract_action.hpp>#include <hpx/traits/is_continuation.hpp>#include
<hpx/traits/promise_local_result.hpp>#include <hpx/type_support/pack.hpp>#include
<hpx/util/bind_action.hpp>#include <hpx/util/functional/colocated_helpers.hpp>#include
<hpx/format.hpp>#include <hpx/functional/result_of.hpp>#include <hpx/runtime/actions/continuation.hpp>#include
<hpx/serialization/serialize.hpp>#include <hpx/serialization/unique_ptr.hpp>#include
<hpx/type_support/decay.hpp>#include <hpx/type_support/unused.hpp>#include
<hpx/lcos/detail/async_colocated_callback.hpp>#include <hpx/lcos/async_continue_callback.hpp>#include
<hpx/lcos/async_callback_fwd.hpp>#include <hpx/runtime/actions/basic_action_fwd.hpp>#include
<hpx/runtime/launch_policy.hpp>#include <hpx/lcos/async_continue.hpp>#include
<hpx/runtime/applier/apply_callback.hpp>#include <hpx/traits/is_distribution_policy.hpp>#include
<hpx/traits/promise_remote_result.hpp>#include <hpx/lcos/detail/async_colocated_callback_fwd.hpp>#include
<hpx/lcos/detail/async_implementations.hpp>#include <hpx/runtime/applier/detail/apply_colocated_callback_fwd.hpp>#include
<hpx/runtime/applier/detail/apply_colocated_fwd.hpp>#include <hpx/runtime/applier/detail/apply_implementations.hpp>#include
<hpx/runtime/components/stubs/stub_base.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/find_here.hpp>#include <hpx/serialization/serialization_fwd.hpp>#include <algorithm>
```

file **component_factory.hpp**

Defines

HPX_REGISTER_COMPONENT (type, name, mode)

Define a component factory for a component type.

This macro is used to create and register a minimal component factory for a component type which allows it to be remotely created using the `hpx::new_<>` function.

This macro can be invoked with one, two or three arguments

Parameters

- `type`: The `type` parameter is a (fully decorated) type of the component type for which a factory should be defined.
- `name`: The `name` parameter specifies the name to use to register the factory. This should uniquely (system-wide) identify the component type. The `name` parameter must conform to the C++ identifier rules (without any namespace). If this parameter is not given, the first parameter is used.

- mode: The *mode* parameter has to be one of the defined enumeration values of the enumeration `hpx::components::factory_state_enum`. The default for this parameter is `hpx::components::factory_enabled`.

file copy_component.hpp

```
#include <hpx/config.hpp>#include <hpx/lcos/async.hpp>#include <hpx/lcos/detail/async_colocated.hpp>#include
<hpx/lcos/future.hpp>#include <hpx/runtime/actions/plain_action.hpp>#include
<hpx/assertion.hpp>#include <hpx/preprocessor/cat.hpp>#include <hpx/preprocessor/expand.hpp>#include
<hpx/preprocessor/nargs.hpp>#include <hpx/runtime/actions/basic_action.hpp>#include
<hpx/traits/component_type_database.hpp>#include <boost/utility/string_ref.hpp>#include
<cstdlib>#include <sstream>#include <stdexcept>#include <string>#include <utility>#include
<hpx/config/warnings_prefix.hpp>#include <hpx/config/warnings_suffix.hpp>#include
<hpx/runtime/components/server/copy_component.hpp>#include <hpx/runtime/actions/plain_action.hpp>#include
<hpx/runtime/components/stubs/runtime_support.hpp>#include <hpx/async.hpp>#include
<hpx/errors.hpp>#include <hpx/lcos/detail/async_colocated_fwd.hpp>#include
<hpx/runtime/actions/manage_object_action.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/serialization/array.hpp>#include <hpx/serialization/base_object.hpp>#include
<hpx/serialization/serialize_buffer.hpp>#include <hpx/static_reinit/reinitializable_static.hpp>#include <cstdint>#include
<hpx/runtime/components/server/runtime_support.hpp>#include <hpx/synchronization/condition_variable.hpp>#include
<hpx/synchronization/mutex.hpp>#include <hpx/synchronization/spinlock.hpp>#include
<hpx/performance_counters/counters.hpp>#include <hpx/plugin.hpp>#include
<hpx/plugins/plugin_factory_base.hpp>#include <hpx/plugins/plugin_registry_base.hpp>#include
<hpx/util_fwd.hpp>#include <vector>#include <hpx/type_support/pack.hpp>#include
<hpx/util_ini.hpp>#include <hpx/program_options.hpp>#include <hpx/runtime/actions/component_action.hpp>#include
<hpx/runtime/components/server/create_component.hpp>#include <hpx/runtime/components/server/create_component_fwd.hpp>
<hpx/runtime/components/server/component_heap.hpp>#include <hpx/runtime/components/static_factory_data.hpp>#include
<hpx/datastructures/any.hpp>#include <hpx/preprocessor/stringize.hpp>#include <map>#include
<hpx/runtime/find_here.hpp>#include <hpx/runtime/parcelset/locality.hpp>#include
<hpx/runtime/parcelset_fwd.hpp>#include <hpx/serialization/map.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include <hpx/iterator_support/traits/is_iterator.hpp>#include
<memory>#include <type_traits>#include <hpx/traits/action_does_termination_detection.hpp>#include
<hpx/traits/is_component.hpp>#include <atomic>#include <condition_variable>#include <list>#include
<mutex>#include <set>#include <thread>#include <hpx/runtime/naming/name.hpp>#include
<hpx/serialization/vector.hpp>#include <hpx/type_support/decay.hpp>#include
<hpx/runtime/get_ptr.hpp>#include <hpx/runtime/agas/gva.hpp>#include <hpx/runtime/components/client_base.hpp>#include
<hpx/runtime/get_lva.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime_fwd.hpp>#include <hpx/traits/component_pin_support.hpp>#include
<hpx/traits/component_type_is_compatible.hpp>#include <hpx/functional/bind_back.hpp>#include
<hpx/traits/get_remote_result.hpp>
```

file default_distribution_policy.hpp

```
#include <hpx/config.hpp>#include <hpx/assertion.hpp>#include <hpx/lcos/dataflow.hpp>#include
<hpx/lcos/future.hpp>#include <hpx/lcos/packaged_action.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/runtime/applier/apply.hpp>#include <hpx/runtime/components/stubs/stub_base.hpp>#include
<hpx/runtime/find_here.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/runtime/naming/name.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include <hpx/serialization/shared_ptr.hpp>#include
<hpx/serialization/vector.hpp>#include <hpx/traits/extract_action.hpp>#include
<hpx/traits/is_distribution_policy.hpp>#include <hpx/traits/promise_local_result.hpp>#include <algorithm>#include
<cstdint>#include <memory>#include <type_traits>#include <utility>#include <vector>
```

file migrate_component.hpp

```
#include <hpx/config.hpp>#include <hpx/lcos/async.hpp>#include <hpx/lcos/detail/async_colocated.hpp>#include
```

```
<hpx/lcos/future.hpp>#include <hpx/runtime/actions/plain_action.hpp>#include
<hpx/runtime/components/client_base.hpp>#include <hpx/runtime/components/server/migrate_component.hpp>#include
<hpx/runtime/actions/plain_action.hpp>#include <hpx/runtime/agas/interface.hpp>#include
<hpx/runtime/components/stubs/runtime_support.hpp>#include <hpx/runtime/get_ptr.hpp>#include
<hpx/runtime/naming/name.hpp>#include <hpx/traits/component_supports_migration.hpp>#include
<hpx/traits/is_component.hpp>#include <cstdint>#include <memory>#include <utility>#include
<hpx/lcos/detail/asynchronous_implementations_fwd.hpp>#include
<hpx/lcos/packaged_action.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/runtime/applier/detail/apply_implementations_fwd.hpp>#include <hpx/runtime/components/stubs/stub_base.hpp>#include
<hpx/runtime/find_here.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/serialization/serialization_fwd.hpp>#include
<hpx/traits/extract_action.hpp>#include <hpx/traits/is_distribution_policy.hpp>#include
<hpx/traits/promise_local_result.hpp>#include <algorithm>#include <cstddef>#include
<type_traits>#include <vector>

file new.hpp
#include <hpx/config.hpp>#include <hpx/lcos/future.hpp>#include <hpx/runtime/components/client_base.hpp>#include
<hpx/runtime/components/default_distribution_policy.hpp>#include <hpx/assertion.hpp>#include
<hpx/lcos/dataflow.hpp>#include <hpx/lcos/packaged_action.hpp>#include <hpx/runtime/actions/action_support.hpp>#include
<hpx/runtime/applier/apply.hpp>#include <hpx/runtime/components/stubs/stub_base.hpp>#include
<hpx/runtime/find_here.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <hpx/runtime/naming/name.hpp>#include
<hpx/serialization/serialization_fwd.hpp>#include <hpx/serialization/shared_ptr.hpp>#include
<hpx/serialization/vector.hpp>#include <hpx/traits/extract_action.hpp>#include
<hpx/traits/is_distribution_policy.hpp>#include <hpx/traits/promise_local_result.hpp>#include
<algorithm>#include <cstddef>#include <memory>#include <type_traits>#include <utility>#include
<vector>#include <hpx/runtime/components/server/create_component.hpp>#include
<hpx/runtime/launch_policy.hpp>#include <hpx/traits/is_client.hpp>#include
<hpx/traits/is_component.hpp>#include <hpx/type_support/lazy_enable_if.hpp>

file find_here.hpp
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <hpx/runtime/naming/id_type.hpp>

file find_localities.hpp
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <hpx/runtime/components/component_type.hpp>#include
<hpx/runtime/naming/id_type.hpp>#include <vector>

file get_colocation_id.hpp
#include <hpx/errors.hpp>#include <hpx/lcos_fwd.hpp>#include <hpx/runtime/launch_policy.hpp>#include
<hpx/runtime/naming/id_type.hpp>

file get_locality_id.hpp
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <cstdint>

file get_locality_name.hpp
#include <hpx/config.hpp>#include <hpx/lcos_fwd.hpp>#include <hpx/runtime/naming/id_type.hpp>#include
<string>

file get_num_localities.hpp
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <hpx/lcos_fwd.hpp>#include
<hpx/runtime/launch_policy.hpp>#include <hpx/runtime/components/component_type.hpp>#include <cstdint>

file get_os_thread_count.hpp
#include <hpx/config.hpp>#include <hpx/runtime/threads/thread_data_fwd.hpp>#include <cstddef>

file get_ptr.hpp
```

```

#include <hpx/config.hpp>#include <hpx/assertion.hpp>#include <hpx/runtime/agas/gva.hpp>#include
<hpx/runtime/components/component_type.hpp>#include <hpx/runtime/naming/name.hpp>#include
<hpx/errors.hpp>#include <boost/io/ios_state.hpp>#include <cstdint>#include
<hpx/runtime/components/client_base.hpp>#include <hpx/runtime/get_lva.hpp>#include
<hpx/runtime/launch_policy.hpp>#include <hpx/runtime/naming/address.hpp>#include
<hpx/runtime_fwd.hpp>#include <hpx/traits/component_pin_support.hpp>#include
<hpx/traits/component_type_is_compatible.hpp>#include <hpx/functional/bind_back.hpp>#include <mem-
ory>

file get_thread_name.hpp
    #include <hpx/config.hpp>#include <string>

file get_worker_thread_num.hpp
    #include <hpx/config.hpp>#include <hpx/errors.hpp>#include <cstddef>

file launch_policy.hpp
    #include <hpx/config.hpp>#include <hpx/coroutines/thread enums.hpp>#include
    <hpx/serialization/serialization_fwd.hpp>#include <type_traits>#include <utility>

file unmanaged.hpp
    #include <hpx/runtime/naming/name.hpp>

file report_error.hpp
    #include <hpx/config.hpp>#include <cstddef>#include <exception>

file runtime_mode.hpp
    #include <hpx/config.hpp>#include <string>

file set_parcel_write_handler.hpp
    #include <hpx/config.hpp>

file shutdown_function.hpp
    #include <hpx/config.hpp>#include <hpx/functional/unique_function.hpp>

file startup_function.hpp
    #include <hpx/config.hpp>#include <hpx/functional/unique_function.hpp>

file scheduler_mode.hpp
    #include <cstdint>

file thread_data_fwd.hpp
    #include <hpx/config.hpp>#include <hpx/coroutines/coroutine_fwd.hpp>#include
    <hpx/thread enums.hpp>#include <hpx/coroutines/thread_id_type.hpp>#include
    <hpx/errors.hpp>#include <hpx/functional/function.hpp>#include <hpx/functional/unique_function.hpp>#include
    <hpx/util_fwd.hpp>#include <cstddef>#include <cstdint>#include <memory>#include <utility>

file thread_helpers.hpp
    #include <hpx/config.hpp>#include <hpx/basic_execution/register_locks.hpp>#include
    <hpx/coroutines/thread enums.hpp>#include <hpx/errors.hpp>#include <hpx/functional/unique_function.hpp>#include
    <hpx/runtime/thread_pool_helpers.hpp>#include <hpx/runtime/threads/policies/scheduler_mode.hpp>#include
    <hpx/runtime/threads/register_thread.hpp>#include <hpx/runtime/threads/thread_data_fwd.hpp>#include
    <hpx/runtime/threads/thread_pool_base.hpp>#include <hpx/runtime/threads/fwd.hpp>#include
    <hpx/timing/steady_clock.hpp>#include <hpx/util/thread_description.hpp>#include
    <hpx/util_fwd.hpp>#include <atomic>#include <chrono>#include <cstddef>#include <cstdint>#include
    <functional>#include <type_traits>#include <utility>

file thread_pool_base.hpp
    #include <hpx/config.hpp>#include <hpx/concurrency/barrier.hpp>#include
    <hpx/errors.hpp>#include <hpx/runtime/thread_pool_helpers.hpp>#include
    <hpx/runtime/threads/detail/network_background_callback.hpp>#include <hpx/affinity/affinity_data.hpp>#include

```

```
<hpx/runtime/threads/policies/callback_notifier.hpp>#include <hpx/runtime/threads/policies/scheduler_mode.hpp>#include
<hpx/runtime/threads/thread_executor.hpp>#include   <hpx/runtime/threads/thread_init_data.hpp>#include
<hpx/topology/cpu_mask.hpp>#include           <hpx/topology/topology.hpp>#include
<hpx/timing/steady_clock.hpp>#include      <hpx/state.hpp>#include      <hpx/util_fwd.hpp>#include
<thread>#include <cstddef>#include <cstdint>#include <exception>#include <functional>#include
<iostfwd>#include <memory>#include <mutex>#include <string>#include <vector>#include
<hpx/config/warnings_prefix.hpp>#include <hpx/config/warnings_suffix.hpp>

file trigger_lco.hpp
#include <hpx/config.hpp>#include <hpx/assertion.hpp>#include <hpx/lcos_fwd.hpp>#include
<hpx/runtime/actions/action_priority.hpp>#include <hpx/runtime/actions/continuation_fwd.hpp>#include
<hpx/runtime/actions_fwd.hpp>#include <hpx/runtime/applier/detail/apply_implementations_fwd.hpp>#include
<hpx/runtime/naming/address.hpp>#include          <hpx/runtime/naming/id_type.hpp>#include
<hpx/runtime/naming/name.hpp>#include <hpx/type_support/decay.hpp>#include <exception>#include
<type_traits>#include <utility>

file runtime_fwd.hpp
#include <hpx/config.hpp>#include <hpx/errors.hpp>#include <hpx/runtime/basename_registration_fwd.hpp>#include
<hpx/runtime/config_entry.hpp>#include
<hpx/runtime/get_colocation_id.hpp>#include
<hpx/runtime/get_locality_name.hpp>#include
<hpx/runtime/get_os_thread_count.hpp>#include
<hpx/runtime/get_worker_thread_num.hpp>#include
<hpx/runtime/report_error.hpp>#include
<hpx/runtime/runtime_mode.hpp>#include
<hpx/runtime/shutdown_function.hpp>#include
<hpx/functional/function.hpp>#include <hpx/util_fwd.hpp>#include <cstddef>#include <cstdint>#include
<string>

file debugging.hpp
#include <hpx/config.hpp>#include <string>

file pack_traversal.hpp
#include <hpx/util/detail/pack_traversal_impl.hpp>#include <hpx/datastructures/tuple.hpp>#include
<type_traits>#include <utility>

file pack_traversal_async.hpp
#include <hpx/util/detail/pack_traversal_async_impl.hpp>#include <utility>

file unwrap.hpp
#include <hpx/config.hpp>#include <hpx/util/detail/unwrap_impl.hpp>#include <cstddef>#include <utility>

dir /hpx/source/hpx/runtime/actions
dir /hpx/source/components/component_storage
dir /hpx/source/components/component_storage/include/hpx/components/component_storage
dir /hpx/source/components/component_storage/include/hpx/components
dir /hpx/source/components
dir /hpx/source/hpx/runtime/components
dir /hpx/source/components/component_storage/include/hpx
dir /hpx/source/hpx
dir /hpx/source/components/component_storage/include
dir /hpx/source/hpx/lcos
```

```
dir /hpx/source/hpx/runtime/naming
dir /hpx/source/hpx/parallel
dir /hpx/source/hpx/performance_counters
dir /hpx/source/hpx/runtime/threads/policies
dir /hpx/source/hpx/runtime
dir /hpx/source
dir /hpx/source/hpx/runtime/threads
dir /hpx/source/hpx/util
```

2.9.2 Modules reference

compute_cuda

```
#include <hpx/compute/cuda.hpp>

#include <hpx/compute/cuda/transfer.hpp>

#include <hpx/compute/cuda/default_executor.hpp>

#include <hpx/compute/cuda/concurrent_executor.hpp>

#include <hpx/compute/cuda/get_targets.hpp>

#include <hpx/compute/cuda/default_executor_parameters.hpp>

#include <hpx/compute/cuda/value_proxy.hpp>

#include <hpx/compute/cuda/target.hpp>

#include <hpx/compute/cuda/target_distribution_policy.hpp>

#include <hpx/compute/cuda/concurrent_executor_parameters.hpp>

#include <hpx/compute/cuda/target_ptr.hpp>

#include <hpx/compute/cuda/allocator.hpp>

#include <hpx/compute/cuda/traits/access_target.hpp>

#include <hpx/compute/cuda/serialization/value_proxy.hpp>
```

util

```
#include <hpx/util/get_entry_as.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
template<typename DestType, typename Config, typename std::enable_if<!std::is_same<
```

```
#include <hpx/util/await_traits.hpp>
```

```
#include <hpx/util/calculate_fanout.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
std::size_t calculate_fanout (std::size_t size, std::size_t local_fanout)
```

```
#include <hpx/util/get_and_reset_value.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
std::uint64_t get_and_reset_value (std::uint64_t &value, bool reset)
```

```
std::int64_t get_and_reset_value (std::int64_t &value, bool reset)
```

```
template<typename T>
```

```
T get_and_reset_value (std::atomic<T> &value, bool reset)
```

```
std::vector<std::int64_t> get_and_reset_value (std::vector<std::int64_t> &value, bool reset)
```

```
#include <hpx/util/insert_checked.hpp>

namespace hpx

namespace util
```

Functions

```
template<typename Iterator>
bool insert_checked(std::pair<Iterator, bool> const &r)
```

Helper function for writing predicates that test whether an std::map insertion succeeded. This inline template function negates the need to explicitly write the sometimes lengthy std::pair<Iterator, bool> type.

Return This function returns `r.second`.

Parameters

- `r`: [in] The return value of a std::map insert operation.

```
template<typename Iterator>
bool insert_checked(std::pair<Iterator, bool> const &r, Iterator &it)
```

Helper function for writing predicates that test whether an std::map insertion succeeded. This inline template function negates the need to explicitly write the sometimes lengthy std::pair<Iterator, bool> type.

Return This function returns `r.second`.

Parameters

- `r`: [in] The return value of a std::map insert operation.
- `r`: [out] A reference to an Iterator, which is set to `r.first`.

```
#include <hpx/util/manage_config.hpp>

namespace hpx

namespace util
```

```
struct manage_config
```

Public Types

```
typedef std::map<std::string, std::string> map_type
```

Public Functions

```
manage_config (std::vector<std::string> const &cfg)
void add (std::vector<std::string> const &cfg)
template<typename T>
T get_value (std::string const &key, T dflt = T()) const
```

Public Members

```
map_type config_
```

datastructures

```
#include <compatibility/hpx/util/any.hpp>

#include <compatibility/hpx/util/optional.hpp>

#include <compatibility/hpx/util/tuple.hpp>

#include <compatibility/hpx/traits/is_tuple_like.hpp>

#include <compatibility/hpx/traits/supports_streaming_with_any.hpp>

#include <hpx/datastructures/any.hpp>

template<>
class basic_any<void, void, void, std::true_type>
```

Public Functions

```
constexpr basic_any()
basic_any (basic_any const &x)
basic_any (basic_any &&x)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::type>>::value>
basic_any (T &&x, typename std::enable_if<!std::is_copy_constructible<typename std::decay<T>::type>>::value * = nullptr)

~basic_any()

basic_any &operator= (basic_any const &x)
basic_any &operator= (basic_any &&rhs)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::type>>::value>
basic_any &operator= (T &&rhs)

basic_any &swap (basic_any &x)
```

```
std::type_info const &type() const
template<typename T>
T const &cast() const
bool has_value() const
void reset()
bool equal_to(basic_any const &rhs) const
```

Private Functions

```
basic_any &assign(basic_any const &x)
```

Private Members

```
detail::any::fxn_ptr_table<void, void, void, std::true_type> *table
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>
static void new_object(void *&object, std::true_type, Ts&&... ts)
template<typename T, typename ...Ts>
static void new_object(void *&object, std::false_type, Ts&&... ts)
template<typename Char>
class basic_any<void, void, Char, std::true_type>
```

Public Functions

```
constexpr basic_any()
basic_any(basic_any const &x)
basic_any(basic_any &&x)
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>>::type>
basic_any(T &&x, typename std::enable_if<!std::is_copy_constructible<typename std::decay<T>::value>>::type * = nullptr)
~basic_any()
basic_any &operator=(basic_any const &x)
basic_any &operator=(basic_any &&rhs)
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>>::type>
basic_any &operator=(T &&rhs)
basic_any &swap(basic_any &x)
```

```
std::type_info const &type () const  
  
template<typename T>  
T const &cast () const  
  
bool has_value () const  
  
void reset ()  
  
bool equal_to (basic_any const &rhs) const
```

Private Functions

```
basic_any &assign (basic_any const &x)
```

Private Members

```
detail::any::fxn_ptr_table<void, void, Char, std::true_type> *table  
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::true_type, Ts&&... ts)  
  
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::false_type, Ts&&... ts)  
  
template<>  
class basic_any<void, void, void, std::false_type>
```

Public Functions

```
constexpr basic_any ()  
  
basic_any (basic_any &&x)  
  
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>>::value, basic_any (T &&x, typename std::enable_if<std::is_move_constructible<typename std::decay<T>>::value>>::type * = nullptr)  
  
basic_any (basic_any const &x)  
  
basic_any &operator= (basic_any const &x)  
  
~basic_any ()  
  
basic_any &operator= (basic_any &&rhs)  
  
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>>::value, basic_any &operator= (T &&rhs)  
  
basic_any &swap (basic_any &x)
```

```
std::type_info const &type() const
template<typename T>
T const &cast() const
bool has_value() const
void reset()
bool equal_to(basic_any const &rhs) const
```

Private Members

```
detail::any::fxn_ptr_table<void, void, void, std::false_type> *table
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>
static void new_object(void *&object, std::true_type, Ts&&... ts)

template<typename T, typename ...Ts>
static void new_object(void *&object, std::false_type, Ts&&... ts)

template<typename Char>
class basic_any<void, void, Char, std::false_type>
```

Public Functions

```
constexpr basic_any()
basic_any(basic_any &&x)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>::value>>
basic_any(T &&x, typename std::enable_if<!std::is_move_constructible<typename std::decay<T>::value>::value * = nullptr)

basic_any(basic_any const &x)

basic_any &operator=(basic_any const &x)
~basic_any()

basic_any &operator=(basic_any &&rhs)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>::value>>
basic_any &operator=(T &&rhs)

basic_any &swap(basic_any &x)

std::type_info const &type() const
template<typename T>
T const &cast() const
bool has_value() const
```

```
void reset ()  
bool equal_to (basic_any const &rhs) const
```

Private Members

```
detail::any::fxn_ptr_table<void, void, Char, std::false_type> *table  
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::true_type, Ts&&... ts)
```

```
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::false_type, Ts&&... ts)
```

```
namespace hpx
```

```
namespace util
```

Typedefs

```
using any_nonser = basic_any<void, void, void, std::true_type>  
using streamable_any_nonser = basic_any<void, void, char, std::true_type>  
using streamable_wany_nonser = basic_any<void, void, wchar_t, std::true_type>  
using unique_any_nonser = basic_any<void, void, void, std::false_type>  
using streamable_unique_any_nonser = basic_any<void, void, char, std::false_type>  
using streamable_unique_wany_nonser = basic_any<void, void, wchar_t, std::false_type>
```

Functions

```
template<typename IArch, typename OArch, typename Char, typename Copyable, typename Enable = typename  
std::basic_istream<Char> &operator>> (std::basic_istream<Char> &i, basic_any<IArch, OArch,  
Char, Copyable> &obj)
```

```
template<typename IArch, typename OArch, typename Char, typename Copyable, typename Enable = typename  
std::basic_ostream<Char> &operator<< (std::basic_ostream<Char> &o, basic_any<IArch,  
OArch, Char, Copyable> const &obj)
```

```
template<typename IArch, typename OArch, typename Char, typename Copyable>  
void swap (basic_any<IArch, OArch, Char, Copyable> &lhs, basic_any<IArch, OArch, Char, Copy-  
able> &rhs)
```

```
template<typename T, typename IArch, typename OArch, typename Char, typename Copyable>  
T *any_cast (basic_any<IArch, OArch, Char, Copyable> *operand)
```

```
template<typename T, typename IArch, typename OArch, typename Char, typename Copyable>  
T const *any_cast (basic_any<IArch, OArch, Char, Copyable> const *operand)
```

```

template<typename T, typename IArch, typename OArch, typename Char, typename Copyable>
T any_cast(basic_any<IArch, OArch, Char, Copyable> &operand)

template<typename T, typename IArch, typename OArch, typename Char, typename Copyable>
T const &any_cast(basic_any<IArch, OArch, Char, Copyable> const &operand)

template<typename T>
basic_any<void, void, void, std::true_type> make_any_nonsr(T &&t)

template<typename T, typename Char>
basic_any<void, void, Char, std::true_type> make_streamable_any_nonsr(T &&t)

template<typename T>
basic_any<void, void, void, std::false_type> make_unique_any_nonsr(T &&t)

template<typename T, typename Char>
basic_any<void, void, Char, std::false_type> make_streamable_unique_any_nonsr(T
&&t)

struct bad_any_cast : public bad_cast

```

Public Functions

```

bad_any_cast (std::type_info const &src, std::type_info const &dest)
const char *what() const

```

Public Members

```

const char *from
const char *to

```

```

template<typename Char>
class basic_any<void, void, Char, std::false_type>

```

Public Functions

```
constexpr basic_any()
```

```
basic_any(basic_any &&x)
```

```

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::type>::value>::type * = nullptr>
basic_any(T &&x, typename std::enable_if<std::is_move_constructible<typename std::decay<T>::type>::value>::type * = nullptr)

```

```
basic_any(basic_any const &x)
```

```
basic_any &operator=(basic_any const &x)
```

```
~basic_any()
```

```
basic_any &operator=(basic_any &&rhs)
```

```

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::type>::value>::type * = nullptr>
basic_any &operator=(T &&rhs)

```

```
basic_any &swap (basic_any &x)

std::type_info const &type () const

template<typename T>
T const &cast () const

bool has_value () const

void reset ()

bool equal_to (basic_any const &rhs) const
```

Private Members

```
detail::any::fxn_ptr_table<void, void, Char, std::false_type> *table
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>
static void new_object (void *&object, std::true_type, Ts&&... ts)

template<typename T, typename ...Ts>
static void new_object (void *&object, std::false_type, Ts&&... ts)

template<typename Char>
class basic_any<void, void, Char, std::true_type>
```

Public Functions

```
constexpr basic_any ()

basic_any (basic_any const &x)

basic_any (basic_any &&x)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>>::value>>
basic_any (T &&x, typename std::enable_if<std::is_copy_constructible<typename std::decay<T>>::value>>::type * = nullptr)

~basic_any ()

basic_any &operator= (basic_any const &x)

basic_any &operator= (basic_any &&rhs)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>>::value>>
basic_any &operator= (T &&rhs)

basic_any &swap (basic_any &x)

std::type_info const &type () const

template<typename T>
```

```
T const &cast () const
bool has_value () const
void reset ()
bool equal_to (basic_any const &rhs) const
```

Private Functions

```
basic_any &assign (basic_any const &x)
```

Private Members

```
detail::any::fxn_ptr_table<void, void, Char, std::true_type> *table
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>
static void new_object (void *&object, std::true_type, Ts&&... ts)

template<typename T, typename ...Ts>
static void new_object (void *&object, std::false_type, Ts&&... ts)

template<>
class basic_any<void, void, void, std::false_type>
```

Public Functions

```
constexpr basic_any ()

basic_any (basic_any &&x)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>::type * = nullptr)
basic_any (T &&x, typename std::enable_if<std::is_move_constructible<typename std::decay<T>::value>::type * = nullptr)

basic_any (basic_any const &x)

basic_any &operator= (basic_any const &x)

~basic_any ()

basic_any &operator= (basic_any &&rhs)

template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>::type * = nullptr)
basic_any &operator= (T &&rhs)

basic_any &swap (basic_any &x)

std::type_info const &type () const

template<typename T>
```

```
T const &cast () const  
bool has_value () const  
void reset ()  
bool equal_to (basic_any const &rhs) const
```

Private Members

```
detail::any::fxn_ptr_table<void, void, void, std::false_type> *table  
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::true_type, Ts&&... ts)  
  
template<typename T, typename ...Ts>  
static void new_object (void *&object, std::false_type, Ts&&... ts)  
  
template<>  
class basic_any<void, void, void, std::true_type>
```

Public Functions

```
constexpr basic_any ()  
basic_any (basic_any const &x)  
basic_any (basic_any &&x)  
  
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>>  
basic_any (T &&x, typename std::enable_if<std::is_copy_constructible<typename  
std::decay<T>::value>>::type * = nullptr)  
  
~basic_any ()  
  
basic_any &operator= (basic_any const &x)  
basic_any &operator= (basic_any &&rhs)  
  
template<typename T, typename Enable = typename std::enable_if<!std::is_same<basic_any, typename std::decay<T>::value>>  
basic_any &operator= (T &&rhs)  
  
basic_any &swap (basic_any &x)  
  
std::type_info const &type () const  
  
template<typename T>  
T const &cast () const  
  
bool has_value () const  
void reset ()  
bool equal_to (basic_any const &rhs) const
```

Private Functions

```
basic_any &assign (basic_any const &x)
```

Private Members

```
detail::any::fxn_ptr_table<void, void, void, std::true_type> *table
void *object
```

Private Static Functions

```
template<typename T, typename ...Ts>
```

```
static void new_object (void *&object, std::true_type, Ts&&... ts)
```

```
template<typename T, typename ...Ts>
```

```
static void new_object (void *&object, std::false_type, Ts&&... ts)
```

```
#include <hpx/datastructures/optional.hpp>
```

```
template<typename T>
struct hash<hpx::util::optional<T>>
```

Public Functions

```
constexpr std::size_t operator () (hpx::util::optional<T> const &arg) const
namespace hpx
```

```
namespace util
```

Functions

```
template<typename T>
```

```
constexpr bool operator== (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
```

```
constexpr bool operator!= (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
```

```
constexpr bool operator< (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
```

```
constexpr bool operator>= (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
```

```
constexpr bool operator> (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
```

```
constexpr bool operator<= (optional<T> const &lhs, optional<T> const &rhs)
```

```
template<typename T>
constexpr bool operator==(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator==(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator!=(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator!=(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator<(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator<(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator>=(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator>=(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator>(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator>(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator<=(optional<T> const &opt, nullopt_t)

template<typename T>
constexpr bool operator<=(nullopt_t, optional<T> const &opt)

template<typename T>
constexpr bool operator==(optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator==(T const &value, optional<T> const &opt)

template<typename T>
constexpr bool operator!=(optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator!=(T const &value, optional<T> const &opt)

template<typename T>
constexpr bool operator<(optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator<(T const &value, optional<T> const &opt)

template<typename T>
constexpr bool operator>=(optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator>=(T const &value, optional<T> const &opt)
```

```

template<typename T>
constexpr bool operator> (optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator> (T const &value, optional<T> const &opt)

template<typename T>
constexpr bool operator<= (optional<T> const &opt, T const &value)

template<typename T>
constexpr bool operator<= (T const &value, optional<T> const &opt)

template<typename T>
void swap (optional<T> &x, optional<T> &y)

template<typename T>
constexpr optional<typename std::decay<T>::type> make_optional (T &&v)

template<typename T, typename ...Ts>
constexpr optional<T> make_optional (Ts&&... ts)

template<typename T, typename U, typename ...Ts>
constexpr optional<T> make_optional (std::initializer_list<U> il, Ts&&... ts)

```

Variables

```

constexpr nullopt_t nullopt = {nullopt_t::init()}

class bad_optional_access : public logic_error

```

Public Functions

```

bad_optional_access (std::string const &what_arg)

bad_optional_access (char const *what_arg)

struct nullopt_t

```

Public Functions

```

constexpr nullopt_t (nullopt_t::init)

template<typename T>
class optional

```

Public Types

```
template<>
using value_type = T
```

Public Functions

```
constexpr optional()
constexpr optional(nullopt_t)
optional(optional const &other)
optional(optional &&other)
optional(T const &val)
optional(T &&val)

template<typename ...Ts>
optional(in_place_t, Ts&&... ts)

template<typename U, typename ...Ts>
optional(in_place_t, std::initializer_list<U> il, Ts&&... ts)

~optional()

optional &operator=(optional const &other)
optional &operator=(optional &&other)
optional &operator=(T const &other)
optional &operator=(T &&other)
optional &operator=(nullopt_t)

constexpr T const *operator->() const
T *operator->()

constexpr T const &operator*() const
T &operator*()

constexpr operator bool() const
constexpr bool has_value() const

T &value()

T const &value() const

template<typename U>
constexpr T value_or(U &&value) const

template<typename ...Ts>
void emplace(Ts&&... ts)
```

```
void swap (optional &other)
void reset ()
```

Private Members

```
std::aligned_storage<sizeof(T), alignof(T)>::type storage_
bool empty_

namespace _optional_swap
```

Functions

```
template<typename T>
void check_swap ()
```

```
namespace std
```

```
template<typename T>
struct hash<hpx::util::optional<T>>
```

Public Functions

```
constexpr std::size_t operator() (hpx::util::optional<T> const &arg) const
```

```
#include <hpx/datastructures/tuple.hpp>
```

```
template<typename T0, typename T1>
struct tuple_element<0, std::pair<T0, T1>>
```

Public Types

```
template<>
using type = T0
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::pair< T0, T1 >)
```

```
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std::pair< T0, T1 >)
```

```
template<typename T0, typename T1>
struct tuple_element<1, std::pair<T0, T1>>
```

Public Types

```
template<>
using type = T1
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::pair< T0, T1> const& p)
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std::pair< T0, T1> const& p)
template<std::size_t I, typename Type, std::size_t Size>
struct tuple_element<I, boost::array<Type, Size>>
```

Public Types

```
template<>
using type = Type
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(boost::array< Type, Size> &p)
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(boost::array< Type, Size> const& p)
template<std::size_t I, typename Type, std::size_t Size>
struct tuple_element<I, std::array<Type, Size>>
```

Public Types

```
template<>
using type = Type
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::array< Type, Size> &p)
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std::array< Type, Size> const& p)
namespace hpx
{
    namespace util
    {
        template<std::size_t I, typename Type, std::size_t Size>
        struct tuple_element<I, std::array<Type, Size>>
```

Functions

```
template<typename... Ts>constexpr HPX_HOST_DEVICE tuple<typename decay_unwrap<Ts>::value_type...> tuple(Ts... vs)
template<typename... Ts>constexpr HPX_HOST_DEVICE tuple<Ts&&...> hpx::util::forward_as_tuple(Ts... vs)
template<typename... Ts>constexpr HPX_HOST_DEVICE tuple<Ts&...> hpx::util::tie(Ts &... vs)
template<typename... Tuples>constexpr HPX_HOST_DEVICE auto hpx::util::tuple_cat(Tuples... vs) {
    return tuple_cat(hpx::util::decay<std::decay_t<Tuples>>(vs)...);
}
template<typename... Ts, typename... Us>constexpr HPX_HOST_DEVICE std::enable_if<std::is_constructible<Us, Ts>::value> tuple_cat(tuple<Ts...> ts, tuple<Us...> us) {
    return tuple_cat(ts, us);
}
template<typename... Ts, typename... Us>constexpr HPX_HOST_DEVICE std::enable_if<std::is_constructible<Us, Ts>::value> tuple_cat(tuple<Us...> us, tuple<Ts...> ts) {
    return tuple_cat(us, ts);
}
template<typename... Ts, typename... Us>constexpr HPX_HOST_DEVICE std::enable_if<std::is_constructible<Us, Ts>::value> tuple_cat(tuple<Ts...> ts, tuple<Us...> us) {
    return tuple_cat(us, ts);
}
template<typename... Ts, typename... Us>constexpr HPX_HOST_DEVICE std::enable_if<std::is_constructible<Us, Ts>::value> tuple_cat(tuple<Us...> us, tuple<Ts...> ts) {
    return tuple_cat(us, ts);
}
template<typename... Ts>HPX_HOST_DEVICE void hpx::util::swap(tuple<Ts...> & x, tuple<Ts...> & y) {
    std::swap(x, y);
}
```

Variables

```
detail::ignore_type const ignore = {}

template<typename ...Ts>
class tuple
```

Public Functions

```
template<typename Dependent = void, typename Enable = typename std::enable_if<!util::all_of<std::is_constructible<Dependent, tuple::value_type>>::value> tuple() {
    if constexpr (Dependent) {
        static_assert(util::all_of<std::is_constructible<Dependent, tuple::value_type>>::value, "tuple must be constructible from its elements");
    }
}

constexpr HPX_HOST_DEVICE tuple(Ts const&... vs) {
    return tuple{vs...};
}

template<typename U, typename ...Us, typename Enable = typename std::enable_if<!std::is_same<tuple, typename std::decay<U>>::value> tuple(U &&v, Us&&... vs) {
    return tuple{v, vs...};
}

tuple(tuple const&)

tuple(tuple&&)

template<typename UTuple, typename Enable = typename std::enable_if<!std::is_same<tuple, typename std::decay<UTuple>>::value> tuple(UTuple &&other) {
    return tuple{other};
}

HPX_HOST_DEVICE tuple& hpx::util::tuple::operator=(tuple const & other) {
    return *this = other;
}

HPX_HOST_DEVICE tuple& hpx::util::tuple::operator=(tuple && other) {
    return *this = move(other);
}

template<typename UTuple>HPX_HOST_DEVICE std::enable_if< tuple_size<typename std::decay<UTuple>>::value> tuple(UTuple other) {
    return tuple{other};
}

HPX_HOST_DEVICE void hpx::util::tuple::swap(tuple & other) {
    std::swap(*this, other);
}
```

Public Members

```
detail::tuple_impl<typename util::make_index_pack<sizeof...(Ts)>::type, Ts...> _impl  
template<>  
class tuple<>
```

Public Functions

```
constexpr HPX_HOST_DEVICE tuple()  
constexpr HPX_HOST_DEVICE tuple(tuple const&)  
constexpr HPX_HOST_DEVICE tuple(tuple&&)  
HPX_HOST_DEVICE tuple& hpx::util::tuple::operator=(tuple const &)  
HPX_HOST_DEVICE tuple& hpx::util::tuple::operator=(tuple &&)  
HPX_HOST_DEVICE void hpx::util::tuple::swap(tuple &)  
template<typename T0, typename T1>  
struct tuple_element<0, std::pair<T0, T1>>
```

Public Types

```
template<>  
using type = T0
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::pair<  
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std:  
template<typename T0, typename T1>  
struct tuple_element<1, std::pair<T0, T1>>
```

Public Types

```
template<>  
using type = T1
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::pair<  
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std:  
template<std::size_t I, typename Type, std::size_t Size>  
struct tuple_element<I, boost::array<Type, Size>>
```

Public Types

```
template<>
using type = Type
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(boost::array<...> & t, std::size_t i);
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(boost::array<...> const& t, std::size_t i);

template<std::size_t I, typename Type, std::size_t Size>
struct tuple_element<I, std::array<Type, Size>>
```

Public Types

```
template<>
using type = Type
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(std::array<...> & t, std::size_t i);
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(std::array<...> const& t, std::size_t i);

template<std::size_t I, typename ...Ts>
struct tuple_element<I, tuple<Ts...>>
```

Public Types

```
template<>
using type = typename util::at_index::type
```

Public Static Functions

```
static constexpr HPX_HOST_DEVICE type& hpx::util::tuple_element::get(tuple < Ts... > & t, std::size_t i);
static constexpr HPX_HOST_DEVICE type const& hpx::util::tuple_element::get(tuple < Ts... > const& t, std::size_t i);

template<class T>
struct tuple_size
    Subclassed by hpx::util::tuple_size< const T >, hpx::util::tuple_size< const volatile T >, hpx::util::tuple_size< volatile T >

namespace adl_barrier
```

Functions

```
template<std::size_t I, typename Tuple>constexpr HPX_HOST_DEVICE tuple_element<

#include <hpx/datastructures/traits/is_tuple_like.hpp>

namespace hpx

    namespace traits

        template<typename T>
        struct is_tuple_like : public hpx::traits::detail::is_tuple_like_impl<std::remove_cv<T>::type>
            #include <is_tuple_like.hpp> Deduces to a true type if the given parameter T has a specific tuple like
            size.
```



```
#include <hpx/datastructures/traits/supports_streaming_with_any.hpp>

thread_support

#include <compatibility/hpx/util/thread_specific_ptr.hpp>

#include <compatibility/hpx/util/unlock_guard.hpp>

#include <compatibility/hpx/util/assert_owns_lock.hpp>

#include <compatibility/hpx/util/set_thread_name.hpp>

#include <compatibility/hpx/util/atomic_count.hpp>

#include <hpx/thread_support/thread_specific_ptr.hpp>
```

Defines

```
HPX_EXPORT_THREAD_SPECIFIC_PTR
```

```
namespace hpx
```

```
    namespace util
```

```
        template<typename T, typename Tag>
        struct thread_specific_ptr
```

Public Types

```
typedef boost::thread_specific_ptr<T>::element_type element_type
```

Public Functions

```
T *get () const  
T *operator-> () const  
T &operator* () const  
void reset (T *new_value = nullptr)
```

Private Static Attributes

```
boost::thread_specific_ptr<T> ptr_
```

```
#include <hpx/thread_support/unlock_guard.hpp>

namespace hpx

    namespace util

        template<typename Mutex>
        class unlock_guard
```

Public Types

```
template<>
using mutex_type = Mutex
```

Public Functions

```
HPX_NON_COPYABLE (unlock_guard)  
unlock_guard (Mutex &m)  
~unlock_guard ()
```

Private Members

Mutex &m_

```
#include <hpx/thread_support/assert_owns_lock.hpp>
```

Defines

```
HPX_ASSERT_OWNS_LOCK(l)
```

```
#include <hpx/thread_support/set_thread_name.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
void set_thread_name(char const *threadName)
```

```
#include <hpx/thread_support/atomic_count.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
class atomic_count
```

Public Functions

```
HPX_NON_COPYABLE(atomic_count)
```

```
atomic_count(long value)
```

```
atomic_count &operator=(long value)
```

```
long operator++()
```

```
long operator--()
```

```
atomic_count &operator+=(long n)
```

```
atomic_count &operator-=(long n)
```

```
operator long() const
```

Private Members

`std::atomic<long> value_`

preprocessor

```
#include <hpx/preprocessor/strip_parens.hpp>
```

Defines

HPX_PP_STRIP_PARENS (X)

For any symbol X, this macro returns the same symbol from which potential outer parens have been removed. If no outer parens are found, this macro evaluates to X itself without error.

The original implementation of this macro is from Steven Watanbe as shown in <http://boost.2283326.n4.nabble.com/preprocessor-removing-parentheses-td2591973.html#a2591976>

```
HPX_PP_STRIP_PARENS (no_parens)
HPX_PP_STRIP_PARENS ( (with_parens) )
```

Example Usage:

Parameters

- X: Symbol to strip parens from

This produces the following output

```
no_parens
with_parens
```

```
#include <hpx/preprocessor/expand.hpp>
```

Defines

HPX_PP_EXPAND (X)

The HPX_PP_EXPAND macro performs a double macro-expansion on its argument. This macro can be used to produce a delayed preprocessor expansion.

Parameters

- X: Token to be expanded twice

Example:

```
#define MACRO(a, b, c) (a) (b) (c)
#define ARGS() (1, 2, 3)

HPX_PP_EXPAND(MACRO ARGS()) // expands to (1) (2) (3)
```

```
#include <hpx/preprocessor/stringize.hpp>
```

Defines

HPX_PP_STRINGIZE (X)

The `HPX_PP_STRINGIZE` macro stringizes its argument after it has been expanded.

The passed argument X will expand to "X". Note that the stringizing operator (#) prevents arguments from expanding. This macro circumvents this shortcoming.

Parameters

- X: The text to be converted to a string literal

```
#include <hpx/preprocessor/cat.hpp>
```

Defines

HPX_PP_CAT (A, B)

Concatenates the tokens A and B into a single token. Evaluates to AB

Parameters

- A: First token
- B: Second token

```
#include <hpx/preprocessor/nargs.hpp>
```

Defines

HPX_PP_NARGS (...)

Expands to the number of arguments passed in

Example Usage:

```
HPX_PP_NARGS(hpx, pp, nargs)
HPX_PP_NARGS(hpx, pp)
HPX_PP_NARGS(hpx)
```

Parameters

- ...: The variadic number of arguments

Expands to:

```
3
2
1
```

execution

```
#include <compatibility/hpx/traits/is_execution_policy.hpp>

#include <compatibility/hpx/traits/is_timed_executor.hpp>

#include <compatibility/hpx/traits/executor_traits.hpp>

#include <compatibility/hpx/traits/is_executor.hpp>

#include <compatibility/hpx/traits/is_executor_parameters.hpp>

#include <compatibility/hpx/parallel/execution.hpp>

#include <compatibility/hpx/parallel/execution_policy.hpp>

#include <compatibility/hpx/parallel/execution_policy_fwd.hpp>

#include <compatibility/hpx/parallel/exception_list.hpp>

#include <compatibility/hpx/parallel/executors.hpp>

#include <compatibility/hpx/parallel/executor_parameters.hpp>

#include <compatibility/hpx/parallel/traits/vector_pack_load_store.hpp>

#include <compatibility/hpx/parallel/traits/vector_pack_alignment_size.hpp>

#include <compatibility/hpx/parallel/traits/vector_pack_type.hpp>

#include <compatibility/hpx/parallel/traits/vector_pack_count_bits.hpp>

#include <compatibility/hpx/parallel/executors/default_executor.hpp>

#include <compatibility/hpx/parallel/executors/timed_executors.hpp>

#include <compatibility/hpx/parallel/executors/thread_execution_information.hpp>

#include <compatibility/hpx/parallel/executors/thread_timed_execution.hpp>

#include <compatibility/hpx/parallel/executors/sequenced_executor.hpp>

#include <compatibility/hpx/parallel/executors/execution_information.hpp>

#include <compatibility/hpx/parallel/executors/rebind_executor.hpp>
```

```
#include <compatibility/hpx/parallel/executors/pool_executor.hpp>

#include <compatibility/hpx/parallel/executors/thread_pool_os_executors.hpp>

#include <compatibility/hpx/parallel/executors/static_chunk_size.hpp>

#include <compatibility/hpx/parallel/executors/distribution_policy_executor.hpp>

#include <compatibility/hpx/parallel/executors/execution_information_fwd.hpp>

#include <compatibility/hpx/parallel/executors/thread_pool_executor.hpp>

#include <compatibility/hpx/parallel/executors/timed_execution.hpp>

#include <compatibility/hpx/parallel/executors/fused_bulk_execute.hpp>

#include <compatibility/hpx/parallel/executors/execution.hpp>

#include <compatibility/hpx/parallel/executors/timed_execution_fwd.hpp>

#include <compatibility/hpx/parallel/executors/post_policy_dispatch.hpp>

#include <compatibility/hpx/parallel/executors/parallel_executor.hpp>

#include <compatibility/hpx/parallel/executors/service_executors.hpp>

#include <compatibility/hpx/parallel/executors/this_thread_executors.hpp>

#include <compatibility/hpx/parallel/executors/thread_pool_attached_executors.hpp>

#include <compatibility/hpx/parallel/executors/persistent_auto_chunk_size.hpp>

#include <compatibility/hpx/parallel/executors/dynamic_chunk_size.hpp>

#include <compatibility/hpx/parallel/executors/execution_parameters_fwd.hpp>

#include <compatibility/hpx/parallel/executors/execution_fwd.hpp>

#include <compatibility/hpx/parallel/executors/guided_chunk_size.hpp>

#include <compatibility/hpx/parallel/executors/execution_parameters.hpp>

#include <compatibility/hpx/parallel/executors/parallel_executor_aggregated.hpp>
```

```
#include <compatibility/hpx/parallel/executors/thread_execution.hpp>

#include <compatibility/hpx/parallel/executors/thread_pool_executors.hpp>

#include <compatibility/hpx/parallel/executors/auto_chunk_size.hpp>

#include <compatibility/hpx/parallel/datapar/execution_policy.hpp>

#include <compatibility/hpx/parallel/datapar/execution_policy_fwd.hpp>

#include <hpx/execution/execution.hpp>

#include <hpx/execution/execution_policy.hpp>

namespace hpx

namespace parallel

namespace execution
```

Variables

constexpr task_policy_tag **task**

Default sequential execution policy object.

HPX_STATIC_CONSTEXPR sequenced_policy **hpx::parallel::execution::seq**

Default sequential execution policy object.

HPX_STATIC_CONSTEXPR parallel_policy **hpx::parallel::execution::par**

Default parallel execution policy object.

HPX_STATIC_CONSTEXPR parallel_unsequenced_policy **hpx::parallel::execution::par_u**

Default vector execution policy object.

struct parallel_policy

#include <execution_policy.hpp> The class *parallel_policy* is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be parallelized.

Subclassed by *hpx::parallel::execution::parallel_policy_shim< Executor, Parameters >*

Public Types

typedef *parallel_executor executor_type*

The type of the executor associated with this execution policy.

typedef *execution::extract_executor_parameters<executor_type>::type executor_parameters_type*

The type of the associated executor parameters object which is associated with this execution policy

typedef *parallel_execution_tag execution_category*

The category of the execution agents created by this execution policy.

Public Functions

constexpr parallel_task_policy operator() (task_policy_tag) **const**
Create a new *parallel_policy* referencing a chunk size.

Return The new *parallel_policy*

Parameters

- **tag:** [in] Specify that the corresponding asynchronous execution policy should be used

```
template<typename Executor>
rebind_executor<parallel_policy, Executor, executor_parameters_type>::type on(Executor
    &&exec)
const
```

Create a new *parallel_policy* referencing an executor and a chunk size.

Return The new *parallel_policy*

Parameters

- **exec**: [in] The executor to use for the execution of the parallel algorithm the returned execution policy is used with

```
template<typename ...Parameters, typename ParametersType = typename executor_parameters_join<Parameters...>, rebind_executor<parallel_policy, executor_type, ParametersType>::type with (Parameters&...)
```

params)
const

Create a new *parallel_policy* from the given execution parameters

Note Requires: `is_executor_parameters<Parameters>::value` is true

Return The new *parallel_policy*

Template Parameters

- **Parameters:** The type of the executor parameters to associate with this execution policy.

Parameters

- **params**: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

executor_type & **executor()**

Return the associated executor object.

constexpr *executor_type* **const &executor()** **const**

Return the associated executor object.

executor_parameters_type &**parameters** ()

Return the associated executor parameters object.

```
constexpr executor_parameters_type const &parameters () const
```

Return the associated executor parameters object.

Private Functions

```
template<typename Archive>
```

```
void serialize(Archive &ar, const unsigned int version)
```

Private Members

```
executor_type exec_
executor_parameters_type params_
```

Friends

```
friend hpx::parallel::execution::hpx::serialization::access
template<typename Executor_, typename Parameters_>
struct rebind
#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.
The execution category of Executor shall not be weaker than that of this execution policy
```

Public Types

typedef parallel_policy_shim<Executor_, Parameters_> **type**
 The type of the rebound execution policy.

```
template<typename Executor, typename Parameters>
struct parallel_policy_shim : public hpx::parallel::execution::parallel_policy
#include <execution_policy.hpp> The class parallel_policy_shim is an execution policy type used
as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algo-
rithm's execution may be parallelized.
```

Public Types

typedef Executor **executor_type**
 The type of the executor associated with this execution policy.

typedef Parameters **executor_parameters_type**
 The type of the associated executor parameters object which is associated with this execution policy

typedef *hpx::traits::executor_execution_category<executor_type>::type* **execution_category**
 The category of the execution agents created by this execution policy.

Public Functions

```
constexpr parallel_task_policy_shim<Executor, Parameters> operator() (task_policy_tag
tag)
const
```

Create a new *parallel_policy* referencing a chunk size.

Return The new *parallel_policy*
Parameters

- **tag:** [in] Specify that the corresponding asynchronous execution policy should be used

```
template<typename Executor_>
rebind_executor<parallel_policy_shim, Executor_, executor_parameters_type>::type on (Executor_
&&exec)
const
```

Create a new *parallel_policy* from the given executor

Note Requires: `is_executor<Executor>::value` is true

Return The new `parallel_policy`

Template Parameters

- `Executor`: The type of the executor to associate with this execution policy.

Parameters

- `exec`: [in] The executor to use for the execution of the parallel algorithm the returned execution policy is used with.

```
template<typename ...Parameters_, typename ParametersType = typename executor_parameters_join<
rebind_executor<parallel_policy_shim, executor_type, ParametersType>::type with (Parameters_&&...
                                         params)
                                         const
```

Create a new `parallel_policy_shim` from the given execution parameters

Note Requires: `is_executor_parameters<Parameters>::value` is true

Return The new `parallel_policy_shim`

Template Parameters

- `Parameters`: The type of the executor parameters to associate with this execution policy.

Parameters

- `params`: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

`Executor &executor()`

Return the associated executor object.

`constexpr Executor const &executor() const`

Return the associated executor object.

`Parameters ¶meters()`

Return the associated executor parameters object.

`constexpr Parameters const ¶meters() const`

Return the associated executor parameters object.

`template<typename Executor_, typename Parameters_>`

`struct rebind`

```
#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.  
The execution category of Executor shall not be weaker than that of this execution policy
```

Public Types

`template<>`

`typedef parallel_policy_shim<Executor_, Parameters_> type`

The type of the rebound execution policy.

`struct parallel_task_policy`

```
#include <execution_policy.hpp> Extension: The class parallel_task_policy is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be parallelized.
```

The algorithm returns a future representing the result of the corresponding algorithm when invoked with the `parallel_policy`.

Subclassed by `hpx::parallel::execution::parallel_task_policy_shim< Executor, Parameters >`

Public Types

typedef *parallel_executor* **executor_type**

The type of the executor associated with this execution policy.

typedef *execution::extract_executor_parameters<executor_type>*::type **executor_parameters_type**

The type of the associated executor parameters object which is associated with this execution policy

typedef *parallel_execution_tag* **execution_category**

The category of the execution agents created by this execution policy.

Public Functions

constexpr *parallel_task_policy* **operator()** (*task_policy_tag*) **const**

Create a new *parallel_task_policy* from itself

Return The new *parallel_task_policy*

Parameters

- *tag*: [in] Specify that the corresponding asynchronous execution policy should be used

template<typename **Executor**>

rebind_executor<*parallel_task_policy*, *Executor*, *executor_parameters_type*>::type **on** (*Executor*
 $\&\&exec$)

const

Create a new *parallel_task_policy* from given executor

Note Requires: *is_executor<Executor>*::value is true

Return The new *parallel_task_policy*

Template Parameters

- *Executor*: The type of the executor to associate with this execution policy.

Parameters

- *exec*: [in] The executor to use for the execution of the parallel algorithm the returned execution policy is used with.

template<typename ...**Parameters**, typename **ParametersType** = *typename executor_parameters_join<Parameters...>*>

rebind_executor<*parallel_task_policy*, *executor_type*, *ParametersType*>::type **with** (*Parameters*&&...
 $params$)

const

Create a new *parallel_policy_shim* from the given execution parameters

Note Requires: all parameters are *executor_parameters*, different parameter types can't be duplicated

Return The new *parallel_policy_shim*

Template Parameters

- *Parameters*: The type of the executor parameters to associate with this execution policy.

Parameters

- *params*: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

executor_type &**executor**()

Return the associated executor object.

constexpr *executor_type* **const** &**executor**() **const**

Return the associated executor object.

`executor_parameters_type ¶meters()`
Return the associated executor parameters object.

`constexpr executor_parameters_type const ¶meters() const`
Return the associated executor parameters object.

Private Functions

```
template<typename Archive>
void serialize(Archive &ar, const unsigned int version)
```

Private Members

`executor_type exec_`
`executor_parameters_type params_`

Friends

```
friend hpx::parallel::execution::hpx::serialization::access
```

```
template<typename Executor_, typename Parameters_>
struct rebind
#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.
The execution category of Executor shall not be weaker than that of this execution policy
```

Public Types

`typedef parallel_task_policy_shim<Executor_, Parameters_> type`
The type of the rebound execution policy.

```
template<typename Executor, typename Parameters>
struct parallel_task_policy_shim : public hpx::parallel::execution::parallel_task_policy
#include <execution_policy.hpp> Extension: The class parallel_task_policy_shim is an execution policy type used as a unique type to disambiguate parallel algorithm overloading based on combining a underlying parallel_task_policy and an executor and indicate that a parallel algorithm's execution may be parallelized.
```

Public Types

`typedef Executor executor_type`
The type of the executor associated with this execution policy.

`typedef Parameters executor_parameters_type`
The type of the associated executor parameters object which is associated with this execution policy

`typedef hpx::traits::executor_execution_category<executor_type>::type execution_category`
The category of the execution agents created by this execution policy.

Public Functions

constexpr parallel_task_policy_shim **operator()** (task_policy_tag tag) **const**
Create a new *parallel_task_policy_shim* from itself

Return The new *sequenced_task_policy*

Parameters

- **tag:** [in] Specify that the corresponding asynchronous execution policy should be used

```
template<typename Executor_>
rebind_executor<parallel_task_policy_shim, Executor_, executor_parameters_type>::type on (Executor_
&&exec)
const
```

Create a new *parallel_task_policy* from the given executor

Note Requires: `is_executor<Executor>::value` is true

Return The new *para*

- ## Template Parameters

- Execu

- `exec`: [in] The executor to use for the execution of the parallel algorithm the returned `Future` will be submitted to.

```
template<typename ...Parameters_, typename ParametersType = typename executor_parameters_join<  
rebind_executor<parallel_task_policy_shim, executor_type, ParametersType>::type with (Parameters_&&...  
                                params)  
                                const
```

Create a new *parallel_policy_shim* from the given execution parameters

Note Requires: all parameters are executor_parameters, different parameter types can't be duplicated

Return The new *parallel_policy_shim*

Template Parameters

- **Parameters:** The type of the executor parameters to associate with this execution policy.

Parameters

- **params:** [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

Executor &**executor**()

Return the associated executor object.

```
constexpr Executor const &executor() const
```

Return the associated executor object.

Parameters ¶meters ()

Return the associated executor parameters object.

constexpr Parameters **const ¶meters()** **const**

Return the associated executor parameters object.

```
template<typename Executor_, typename Parameters_>
```

struct rebind

#include <execution_policy.hpp> Rebind the type of executor used by this execution policy. The execution category of Executor shall not be weaker than that of this execution policy

Public Types

```
template<>
typedef parallel_task_policy_shim<Executor_, Parameters_> type
    The type of the rebound execution policy.
```

```
struct parallel_unsequenced_policy
```

#include <execution_policy.hpp> The class *parallel_unsequenced_policy* is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be vectorized.

Public Types

```
typedef parallel_executor executor_type
```

The type of the executor associated with this execution policy.

```
typedef execution::extract_executor_parameters<executor_type>::type executor_parameters_type
```

The type of the associated executor parameters object which is associated with this execution policy

```
typedef parallel_execution_tag execution_category
```

The category of the execution agents created by this execution policy.

Public Functions

```
parallel_unsequenced_policy operator() (task_policy_tag) const
```

Create a new *parallel_unsequenced_policy* from itself

Return The new *parallel_unsequenced_policy*

Parameters

- **tag:** [in] Specify that the corresponding asynchronous execution policy should be used

```
executor_type &executor()
```

Return the associated executor object.

```
constexpr executor_type const &executor() const
```

Return the associated executor object.

```
executor_parameters_type &parameters()
```

Return the associated executor parameters object.

```
constexpr executor_parameters_type const &parameters() const
```

Return the associated executor parameters object.

Private Functions

```
template<typename Archive>
```

```
void serialize (Archive &ar, const unsigned int version)
```

Private Members

```
executor_type exec_  

executor_parameters_type params_
```

Friends

```
friend hpx::parallel::execution::hpx::serialization::access  

struct sequenced_policy  

#include <execution_policy.hpp> The class sequenced_policy is an execution policy type used  

as a unique type to disambiguate parallel algorithm overloading and require that a parallel algo-  

rithm's execution may not be parallelized.
```

Subclassed by *hpx::parallel::execution::sequenced_policy_shim<Executor, Parameters >*

Public Types

typedef *sequenced_executor executor_type*
The type of the executor associated with this execution policy.

typedef *execution::extract_executor_parameters<executor_type>::type executor_parameters_type*
The type of the associated executor parameters object which is associated with this execution
policy

typedef *sequenced_execution_tag execution_category*
The category of the execution agents created by this execution policy.

Public Functions

constexpr sequenced_task_policy operator() (*task_policy_tag*) **const**
Create a new *sequenced_task_policy*.

Return The new *sequenced_task_policy*
Parameters

- *tag*: [in] Specify that the corresponding asynchronous execution policy should be used

template<typename Executor>
rebind_executor<sequenced_policy, Executor, executor_parameters_type>::type on (*Executor*
&&exec)
const
Create a new *sequenced_policy* from the given executor

Note Requires: *is_executor<Executor>::value* is true

Return The new *sequenced_policy*

Template Parameters

- *Executor*: The type of the executor to associate with this execution policy.

Parameters

- *exec*: [in] The executor to use for the execution of the parallel algorithm the returned
execution policy is used with.

template<typename ...Parameters, typename ParametersType = typename executor_parameters_join<Parameters...>::type> **join** (*Parameters...*)

```
rebind_executor<sequenced_policy, executor_type, ParametersType>::type with (Parameters&&...  
                                params)  
const
```

Create a new *sequenced_policy* from the given execution parameters

Note Requires: all parameters are executor_parameters, different parameter types can't be duplicated

Return The new *sequenced_policy*

Template Parameters

- **Parameters:** The type of the executor parameters to associate with this execution policy.

Parameters

- **params:** [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

executor_type &executor()

Return the associated executor object. Return the associated executor object.

constexpr *executor_type const &executor() const*

Return the associated executor object.

executor_parameters_type ¶meters()

Return the associated executor parameters object.

constexpr *executor_parameters_type const ¶meters() const*

Return the associated executor parameters object.

Private Functions

```
template<typename Archive>  
void serialize (Archive &ar, const unsigned int version)
```

Private Members

executor_type exec_

executor_parameters_type params_

Friends

```
friend hpx::parallel::execution::hpx::serialization::access
```

```
template<typename Executor_, typename Parameters_>  
struct rebind
```

#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.
The execution category of Executor shall not be weaker than that of this execution policy

Public Types

typedef sequenced_policy_shim<Executor_, Parameters_> **type**
The type of the rebound execution policy.

```
template<typename Executor, typename Parameters>
struct sequenced_policy_shim: public hpx::parallel::execution::sequenced_policy
#include <execution_policy.hpp> The class sequenced_policy is an execution policy type used
as a unique type to disambiguate parallel algorithm overloading and require that a parallel algo-
rithm's execution may not be parallelized.
```

Public Types

typedef Executor **executor_type**
The type of the executor associated with this execution policy.

typedef Parameters **executor_parameters_type**
The type of the associated executor parameters object which is associated with this execution
policy

typedef *hpx::traits::executor_execution_category<executor_type>::type* **execution_category**
The category of the execution agents created by this execution policy.

Public Functions

```
constexpr sequenced_task_policy_shim<Executor, Parameters> operator() (task_policy_tag
tag)
const
```

Create a new *sequenced_task_policy*.

Return The new *sequenced_task_policy_shim*

Parameters

- *tag*: [in] Specify that the corresponding asynchronous execution policy should be used

```
template<typename Executor_>
rebind_executor<sequenced_policy_shim, Executor_, executor_parameters_type>::type on (Executor_
&&exec)
const
```

Create a new *sequenced_policy* from the given executor

Note Requires: *is_executor<Executor>::value* is true

Return The new *sequenced_policy*

Template Parameters

- *Executor*: The type of the executor to associate with this execution policy.

Parameters

- *exec*: [in] The executor to use for the execution of the parallel algorithm the returned
execution policy is used with.

```
template<typename ...Parameters_, typename ParametersType = typename executor_parameters_join<...
rebind_executor<sequenced_policy_shim, executor_type, ParametersType>::type with (Parameters_&&...
params)
const
```

Create a new *sequenced_policy_shim* from the given execution parameters

Note Requires: all parameters are *executor_parameters*, different parameter types can't be
duplicated

Return The new *sequenced_policy_shim*

Template Parameters

- **Parameters**: The type of the executor parameters to associate with this execution policy.

Parameters

- **params**: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

Executor &executor ()

Return the associated executor object.

constexpr Executor const &executor () const

Return the associated executor object.

Parameters ¶meters ()

Return the associated executor parameters object.

constexpr Parameters const ¶meters () const

Return the associated executor parameters object.

template<typename Executor_, typename Parameters_>**struct rebind**

#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.

The execution category of Executor shall not be weaker than that of this execution policy

Public Types**template<>****typedef sequenced_policy_shim<Executor_, Parameters_> type**

The type of the rebound execution policy.

struct sequenced_task_policy

#include <execution_policy.hpp> Extension: The class *sequenced_task_policy* is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may not be parallelized (has to run sequentially).

The algorithm returns a future representing the result of the corresponding algorithm when invoked with the *sequenced_policy*.

Subclassed by *hpx::parallel::execution::sequenced_task_policy_shim<Executor, Parameters>*

Public Types**typedef sequenced_executor executor_type**

The type of the executor associated with this execution policy.

typedef execution::extract_executor_parameters<executor_type>::type executor_parameters_type

The type of the associated executor parameters object which is associated with this execution policy

typedef sequenced_execution_tag execution_category

The category of the execution agents created by this execution policy.

Public Functions

constexpr sequenced_task_policy operator() (task_policy_tag) **const**
 Create a new *sequenced_task_policy* from itself

Return The new *sequenced_task_policy*

Parameters

- tag: [in] Specify that the corresponding asynchronous execution policy should be used

template<typename Executor>
rebind_executor<sequenced_task_policy, Executor, executor_parameters_type>::type **on** (*Executor*
&&exec)
const

Create a new *sequenced_task_policy* from the given executor

Note Requires: *is_executor<Executor>::value* is true

Return The new *sequenced_task_policy*

Template Parameters

- Executor: The type of the executor to associate with this execution policy.

Parameters

- exec: [in] The executor to use for the execution of the parallel algorithm the returned execution policy is used with.

template<typename ...Parameters, typename ParametersType = typename executor_parameters_join<Parameters...>::type>
rebind_executor<sequenced_task_policy, executor_type, ParametersType>::type **with** (*Parameters**&&...*
params)
const

Create a new *sequenced_task_policy* from the given execution parameters

Note Requires: all parameters are *executor_parameters*, different parameter types can't be duplicated

Return The new *sequenced_task_policy*

Template Parameters

- Parameters: The type of the executor parameters to associate with this execution policy.

Parameters

- params: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

executor_type &executor()

Return the associated executor object.

constexpr executor_type const &executor() **const**

Return the associated executor object.

executor_parameters_type ¶meters()

Return the associated executor parameters object.

constexpr executor_parameters_type const ¶meters() **const**

Return the associated executor parameters object.

Private Functions

```
template<typename Archive>
void serialize (Archive &ar, const unsigned int version)
```

Private Members

```
executor_type exec_
executor_parameters_type params_
```

Friends

```
friend hpx::parallel::execution::hpx::serialization::access

template<typename Executor_, typename Parameters_>
struct rebind
    #include <execution_policy.hpp> Rebind the type of executor used by this execution policy.
    The execution category of Executor shall not be weaker than that of this execution policy
```

Public Types

```
typedef sequenced_task_policy_shim<Executor_, Parameters_> type
    The type of the rebound execution policy.
```

```
template<typename Executor, typename Parameters>
struct sequenced_task_policy_shim : public hpx::parallel::execution::sequenced_task_policy
    #include <execution_policy.hpp> Extension: The class sequenced_task_policy_shim is an execution policy type used as a unique type to disambiguate parallel algorithm overloading based on combining a underlying sequenced_task_policy and an executor and indicate that a parallel algorithm's execution may not be parallelized (has to run sequentially).
```

The algorithm returns a future representing the result of the corresponding algorithm when invoked with the *sequenced_policy*.

Public Types

```
typedef Executor executor_type
    The type of the executor associated with this execution policy.
```

```
typedef Parameters executor_parameters_type
    The type of the associated executor parameters object which is associated with this execution policy
```

```
typedef hpx::traits::executor_execution_category<executor_type>::type execution_category
    The category of the execution agents created by this execution policy.
```

Public Functions

```
constexpr sequenced_task_policy_shim const &operator() (task_policy_tag tag)
const
```

Create a new *sequenced_task_policy* from itself

Return The new *sequenced_task_policy*

Parameters

- tag: [in] Specify that the corresponding asynchronous execution policy should be used

```
template<typename Executor_>
rebind_executor<sequenced_task_policy_shim, Executor_, executor_parameters_type>::type on (Executor_
&&exec)
const
```

Create a new *sequenced_task_policy* from the given executor

Note Requires: *is_executor<Executor>::value* is true

Return The new *sequenced_task_policy*

Template Parameters

- Executor: The type of the executor to associate with this execution policy.

Parameters

- exec: [in] The executor to use for the execution of the parallel algorithm the returned execution policy is used with.

```
template<typename ...Parameters_, typename ParametersType = typename executor_parameters_join<
rebind_executor<sequenced_task_policy_shim, executor_type, ParametersType>::type with (Parameters_&&...
params)
const
```

Create a new *sequenced_task_policy_shim* from the given execution parameters

Note Requires: all parameters are *executor_parameters*, different parameter types can't be duplicated

Return The new *sequenced_task_policy_shim*

Template Parameters

- Parameters: The type of the executor parameters to associate with this execution policy.

Parameters

- params: [in] The executor parameters to use for the execution of the parallel algorithm the returned execution policy is used with.

Executor &executor()

Return the associated executor object.

constexpr Executor **const &executor()** **const**

Return the associated executor object.

Parameters **¶meters()**

Return the associated executor parameters object.

constexpr Parameters **const ¶meters()** **const**

Return the associated executor parameters object.

```
template<typename Executor_, typename Parameters_>
```

struct rebind

#include <execution_policy.hpp> Rebind the type of executor used by this execution policy.
The execution category of Executor shall not be weaker than that of this execution policy

Public Types

```
template<typename T>
typedef sequenced_task_policy_shim<Executor_, Parameters_> type
    The type of the rebound execution policy.

#include <hpx/execution/execution_policy_fwd.hpp>

#include <hpx/execution/exception_list.hpp>

#include <hpx/execution/executors.hpp>

#include <hpx/execution/executor_parameters.hpp>

#include <hpx/execution/traits/is_execution_policy.hpp>

namespace hpx

namespace parallel

namespace execution

template<typename T>
struct is_async_execution_policy : public execution::detail::is_async_execution_policy<hpx::util::decay<T>::type>
    #include <is_execution_policy.hpp> Extension: Detect whether given execution policy makes
    algorithms asynchronous
    1. The type is_async_execution_policy can be used to detect asynchronous execution policies for
        the purpose of excluding function signatures from otherwise ambiguous overload resolution
        participation.
    2. If T is the type of a standard or implementation-defined execution policy,
        is_async_execution_policy<T> shall be publicly derived from integral_constant<bool,
        true>, otherwise from integral_constant<bool, false>.
    3. The behavior of a program that adds specializations for is_async_execution_policy is undefined.

template<typename T>
struct is_execution_policy : public execution::detail::is_execution_policy<hpx::util::decay<T>::type>
    #include <is_execution_policy.hpp>
    1. The type is_execution_policy can be used to detect execution policies for the purpose of ex-
        cluding function signatures from otherwise ambiguous overload resolution participation.
    2. If T is the type of a standard or implementation-defined execution policy,
        is_execution_policy<T> shall be publicly derived from integral_constant<bool, true>,
        otherwise from integral_constant<bool, false>.
    3. The behavior of a program that adds specializations for is_execution_policy is undefined.

template<typename T>
struct is_parallel_execution_policy : public execution::detail::is_parallel_execution_policy<hpx::util::decay<T>::type>
    #include <is_execution_policy.hpp> Extension: Detect whether given execution policy enables
    parallelization
```

1. The type `is_parallel_execution_policy` can be used to detect parallel execution policies for the purpose of excluding function signatures from otherwise ambiguous overload resolution participation.
2. If `T` is the type of a standard or implementation-defined execution policy, `is_parallel_execution_policy<T>` shall be publicly derived from `integral_constant<bool, true>`, otherwise from `integral_constant<bool, false>`.
3. The behavior of a program that adds specializations for `is_parallel_execution_policy` is undefined.

```
template<typename T>
struct is_sequenced_execution_policy : public execution::detail::is_sequenced_execution_policy<hp
#include <is_execution_policy.hpp> Extension: Detect whether given execution policy does not
enable parallelization
```

1. The type `is_sequenced_execution_policy` can be used to detect non-parallel execution policies for the purpose of excluding function signatures from otherwise ambiguous overload resolution participation.
2. If `T` is the type of a standard or implementation-defined execution policy, `is_sequenced_execution_policy<T>` shall be publicly derived from `integral_constant<bool, true>`, otherwise from `integral_constant<bool, false>`.
3. The behavior of a program that adds specializations for `is_sequenced_execution_policy` is undefined.

```
#include <hpx/execution/traits/is_timed_executor.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

Typedefs

```
using is_timed_executor_t = typename is_timed_executor<T>::type
```

```
#include <hpx/execution/traits/executor_traits.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

```
template<typename Executor>
struct executor_context
```

Public Types

```
template<>
using type = typename std::decay::type

template<typename Executor>
struct executor_execution_category
```

Public Types

```
template<>
using type = hpx::util::detected_or_t<unsequenced_execution_tag, execution_category, Executor>
```

Private Types

```
template<>
using execution_category = typename T::execution_category

template<typename Executor>
struct executor_index
```

Public Types

```
template<>
using type = hpx::util::detected_or_t<typename executor_shape<Executor>::type, index_type, Executor>
```

Private Types

```
template<>
using index_type = typename T::index_type

template<typename Executor>
struct executor_parameters_type
```

Public Types

```
template<>
using type = hpx::util::detected_or_t<parallel::execution::static_chunk_size, parameters_type, Executor>
```

Private Types

```
template<>
using parameters_type = typename T::parameters_type

template<typename Executor>
struct executor_shape
```

Public Types

```
template<>
using type = hpx::util::detected_or_t<std::size_t, shape_type, Executor>
```

Private Types

```
template<>
using shape_type = typename T::shape_type
```

```
namespace traits
```

Typedefs

```
using executor_context_t = typename executor_context<Executor>::type
using executor_execution_category_t = typename executor_execution_category<Executor>::type
using executor_shape_t = typename executor_shape<Executor>::type
using executor_index_t = typename executor_index<Executor>::type
using executor_parameters_type_t = typename executor_parameters_type<Executor>::type
```

```
#include <hpx/execution/traits/is_executor.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

Typedefs

```
using is_one_way_executor_t = typename is_one_way_executor<T>::type
using is_never_blocking_one_way_executor_t = typename is_never_blocking_one_way_executor<T>::type
using is_bulk_one_way_executor_t = typename is_bulk_one_way_executor<T>::type
using is_two_way_executor_t = typename is_two_way_executor<T>::type
using is_bulk_two_way_executor_t = typename is_bulk_two_way_executor<T>::type
```

```
namespace traits
```

TypeDefs

```
using is_one_way_executor_t = typename is_one_way_executor<T>::type
using is_never_blocking_one_way_executor_t = typename is_never_blocking_one_way_executor<T>::type
using is_bulk_one_way_executor_t = typename is_bulk_one_way_executor<T>::type
using is_two_way_executor_t = typename is_two_way_executor<T>::type
using is_bulk_two_way_executor_t = typename is_bulk_two_way_executor<T>::type
using is_executor_any_t = typename is_executor_any<T>::type

#include <hpx/execution/traits/is_executor_parameters.hpp>

template<typename Executor>
struct extract_executor_parameters<Executor, typename hpx::util::always_void<typename Executor::executor_params>>
```

Public Types

```
template<>
using type = typename Executor::executor_parameters_type

template<typename Parameters>
struct extract_has_variable_chunk_size<Parameters, typename hpx::util::always_void<typename Parameters::has_variable_chunk_size>>
```

Public Types

```
template<>
using type = typename Parameters::has_variable_chunk_size

namespace hpx

    namespace parallel

        namespace execution
```

TypeDefs

```
using is_executor_parameters_t = typename is_executor_parameters<T>::type

template<typename Executor, typename Enable = void>
struct extract_executor_parameters
```

Public Types

```
template<>
using type = sequential_executor_parameters

template<typename Executor>
struct extract_executor_parameters<Executor, typename hpx::util::always_void<typename Executor>
```

Public Types

```
template<>
using type = typename Executor::executor_parameters_type

template<typename Parameters, typename Enable = void>
struct extract_has_variable_chunk_size
```

Public Types

```
template<>
using type = std::false_type

template<typename Parameters>
struct extract_has_variable_chunk_size<Parameters, typename hpx::util::always_void<typename Parameters>
```

Public Types

```
template<>
using type = typename Parameters::has_variable_chunk_size

namespace traits
```

Typedefs

```
using is_executor_parameters_t = typename is_executor_parameters<T>::type

#include <hpx/execution/executors/default_executor.hpp>

namespace hpx

namespace parallel

namespace execution
```

TypeDefs

```
using default_executor = threads::executors::default_executor  
Refers to the currently used base-executor.
```

```
#include <hpx/execution/executors/timed_executors.hpp>  
  
namespace hpx  
  
    namespace parallel  
  
        namespace execution
```

TypeDefs

```
using sequenced_timed_executor = timed_executor<execution::sequenced_executor>  
using parallel_timed_executor = timed_executor<execution::parallel_executor>  
  
template<typename BaseExecutor>  
struct timed_executor
```

Public Types

```
typedef std::decay<BaseExecutor>::type base_executor_type  
typedef hpx::traits::executor_execution_category<base_executor_type>::type execution_category  
typedef hpx::traits::executor_parameters_type<base_executor_type>::type parameters_type
```

Public Functions

```
timed_executor(hpx::util::steady_time_point const &abs_time)  
timed_executor(hpx::util::steady_duration const &rel_time)  
  
template<typename Executor>  
timed_executor(Executor &&exec, hpx::util::steady_time_point const &abs_time)  
  
template<typename Executor>  
timed_executor(Executor &&exec, hpx::util::steady_duration const &rel_time)  
  
template<typename F, typename ...Ts>  
hpx::util::detail::invoke_deferred_result<F, Ts...>::type sync_execute(F &&f, Ts&&... ts)  
  
template<typename F, typename ...Ts>  
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_execute(F &&f,  
                                         Ts&&... ts)  
  
template<typename F, typename ...Ts>  
void post(F &&f, Ts&&... ts)
```

Public Members

```
BaseExecutor exec_
    std::chrono::steady_clock::time_point execute_at_
}

#include <hpx/execution/executors/thread_execution_information.hpp>

namespace hpx

    namespace threads
```

Functions

```
template<typename Executor, typename Parameters>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, std::size_t>::type processing_units_count (Executor &&exec, Parameters &&params) { ... }

template<typename Executor>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, bool>::type has_pending_closures (Executor &&exec) { ... }

template<typename Executor>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, threads::mask_cref_type>::type get_pu_mask (Executor &&exec, thread_topology &&topo, std::size_t &&size) { ... }

template<typename Executor, typename Mode>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value>::type set_scheduler_mode (Executor &&exec, Mode mode) { ... }
```

```
#include <hpx/execution/executors/thread_timed_execution.hpp>

namespace hpx

    namespace threads
```

Functions

```
template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value>::type post_at (Executor
    &&exec,
    hpx::util::steady_time_point
    const
    &abs_time, F
    &&f, Ts&&...
    ts)

template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value>::type post_after (Executor
    &&exec,
    hpx::util::steady_duration
    const
    &rel_time,
    F &&f,
    Ts&&...
    ts)

template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, hpx::future<typename hpx::util::detail::invoke_deferred>>::type invoke_deferred (Executor
    &&exec,
    hpx::util::detail::invoke_deferred_params<F, Ts>
    &params)

template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts>>>::type invoke_deferred_result (Executor
    &&exec,
    hpx::util::detail::invoke_deferred_params<F, Ts>
    &params)

template<typename Executor, typename F, typename ...Ts>
```

```
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, typename hpx::util::detail::invoke_deferred_result<F, T>>::type invoke_deferred_result(F& f, T& t)
```

```
#include <hpx/execution/executors/sequenced_executor.hpp>

namespace hpx

    namespace parallel

        namespace execution

            struct sequenced_executor
                #include <sequenced_executor.hpp> A sequential_executor creates groups of sequential execution agents which execute in the calling thread. The sequential order is given by the lexicographical order of indices in the index space.
```

```
#include <hpx/execution/executors/execution_information.hpp>

namespace hpx
```

```
    namespace threads
```

Functions

```
threads::mask_cref_type get_pu_mask(threads::topology & topo, std::size_t thread_num)
```

```
#include <hpx/execution/executors/rebind_executor.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

```
            template<typename ExecutionPolicy, typename Executor, typename Parameters>
            struct rebind_executor
```

`#include <rebind_executor.hpp>` Rebind the type of executor used by an execution policy. The execution category of Executor shall not be weaker than that of ExecutionPolicy.

Public Types

typedef ExecutionPolicy::template rebind<executor_type, parameters_type>::type **type**
The type of the rebound execution policy.

```
#include <hpx/execution/executors/pool_executor.hpp>

namespace hpx

    namespace parallel

        namespace execution
```

Typedefs

using pool_executor = *threads*::executors::pool_executor

```
#include <hpx/execution/executors/thread_pool_os_executors.hpp>

namespace hpx

    namespace parallel

        namespace execution
```

Typedefs

using local_priority_queue_os_executor = *threads*::executors::local_priority_queue_os_executor
Creates a new local_priority_queue_executor

Parameters

- max_punits: [in] The maximum number of processing units to associate with the newly created executor.
- min_punits: [in] The minimum number of processing units to associate with the newly created executor (default: 1).

```
#include <hpx/execution/executors/static_chunk_size.hpp>

namespace hpx

    namespace parallel

        namespace execution
```

```
struct static_chunk_size
#include <static_chunk_size.hpp> Loop iterations are divided into pieces of size chunk_size and
then assigned to threads. If chunk_size is not specified, the iterations are evenly (if possible)
divided contiguously among the threads.
```

Note This executor parameters type is equivalent to OpenMP's STATIC scheduling directive.

Public Functions

```
constexpr static_chunk_size()
Construct a static_chunk_size executor parameters object
```

Note By default the number of loop iterations is determined from the number of available cores and the overall number of loop iterations to schedule.

```
constexpr static_chunk_size(std::size_t chunk_size)
Construct a static_chunk_size executor parameters object
```

Parameters

- *chunk_size*: [in] The optional chunk size to use as the number of loop iterations to run on a single thread.

```
#include <hpx/execution/executors/distribution_policy_executor.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

Functions

```
template<typename DistPolicy>
distribution_policy_executor<typename hpx::util::decay<DistPolicy>::type> make_distribution_policy_executor
```

Create a new *distribution_policy_executor* from the given distribution_policy.

Parameters

- *policy*: The distribution_policy to create an executor from

```
template<typename DistPolicy>
class distribution_policy_executor
#include <distribution_policy_executor.hpp> A distribution_policy_executor creates groups of
parallel execution agents which execute in threads implicitly created by the executor and placed
on any of the associated localities.
```

Template Parameters

- *DistPolicy*: The distribution policy type for which an executor should be created. The expression *hpx::traits::is_distribution_policy<DistPolicy>::value* must evaluate to true.

Public Functions

```
template<typename DistPolicy_, typename Enable = typename std::enable_if<!std::is_same<distribution_
distribution_policy_executor (DistPolicy_ &&policy)
Create a new distribution_policy executor from the given distribution policy
```

Parameters

- policy: The distribution_policy to create an executor from

Private Members

```
DistPolicy policy_
```

```
#include <hpx/execution/executors/execution_information_fwd.hpp>
```

```
#include <hpx/execution/executors/thread_pool_executor.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

```
struct thread_pool_executor
```

```
#include <thread_pool_executor.hpp> A thread_pool_executor creates groups of parallel execution agents which execute in threads implicitly created by the executor. This executor prefers continuing with the creating thread first before executing newly created threads.
```

This executor conforms to the concepts of a TwoWayExecutor, and a BulkTwoWayExecutor

Public Types

```
typedef parallel_execution_tag execution_category
```

Associate the `parallel_execution_tag` executor tag type as a default with this executor.

```
typedef static_chunk_size executor_parameters_type
```

Associate the `static_chunk_size` executor parameters type as a default with this executor.

Public Functions

```
thread_pool_executor (threads::thread_pool_base *pool)
```

Create a new parallel executor.

Private Members

```
threads::thread_pool_base *pool_
```

```
#include <hpx/execution/executors/timed_execution.hpp>
```

```
#include <hpx/execution/executors/fused_bulk_execute.hpp>
```

```
#include <hpx/execution/executors/execution.hpp>
```

```
#include <hpx/execution/executors/timed_execution_fwd.hpp>
```

```
#include <hpx/execution/executors/post_policy_dispatch.hpp>
```

```
#include <hpx/execution/executors/parallel_executor.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

TypeDefs

```
using parallel_executor = parallel_policy_executor<hpx::launch>
```

```
template<typename Policy>
struct parallel_policy_executor
```

`#include <parallel_executor.hpp>` A `parallel_executor` creates groups of parallel execution agents which execute in threads implicitly created by the executor. This executor prefers continuing with the creating thread first before executing newly created threads.

This executor conforms to the concepts of a `TwoWayExecutor`, and a `BulkTwoWayExecutor`

Public Types

```
typedef parallel_execution_tag execution_category
```

Associate the `parallel_execution_tag` executor tag type as a default with this executor.

```
typedef static_chunk_size executor_parameters_type
```

Associate the `static_chunk_size` executor parameters type as a default with this executor.

Public Functions

```
constexpr parallel_policy_executor(Policy l = detail::get_default_policy<Policy>::call(),
                                   std::size_t spread = 4, std::size_t tasks =
                                   std::size_t(-1))
```

Create a new parallel executor.

```
#include <hpx/execution/executors/service_executors.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

TypeDefs

```
using service_executor = threads::executors::service_executor
```

A *service_executor* exposes one of the predefined HPX thread pools through an executor interface.

Note All tasks executed by one of these executors will run on one of the OS-threads dedicated for the given thread pool. The tasks will not run as HPX-threads.

```
using io_pool_executor = threads::executors::io_pool_executor
```

A *io_pool_executor* exposes the predefined HPX IO thread pool through an executor interface.

Note All tasks executed by one of these executors will run on one of the OS-threads dedicated for the IO thread pool. The tasks will not run as HPX-threads.

```
using parcel_pool_executor = threads::executors::parcel_pool_executor
```

A *io_pool_executor* exposes the predefined HPX parcel thread pool through an executor interface.

Note All tasks executed by one of these executors will run on one of the OS-threads dedicated for the parcel thread pool. The tasks will not run as HPX-threads.

```
using timer_pool_executor = threads::executors::timer_pool_executor
```

A *io_pool_executor* exposes the predefined HPX timer thread pool through an executor interface.

Note All tasks executed by one of these executors will run on one of the OS-threads dedicated for the timer thread pool. The tasks will not run as HPX-threads.

```
using main_pool_executor = threads::executors::main_pool_executor
```

A *io_pool_executor* exposes the predefined HPX main thread pool through an executor interface.

Note All tasks executed by one of these executors will run on one of the OS-threads dedicated for the main thread pool. The tasks will not run as HPX-threads.

```
#include <hpx/execution/executors/this_thread_executors.hpp>

#include <hpx/execution/executors/thread_pool_attached_executors.hpp>

namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

Typedefs

```
using local_priority_queue_attached_executor = threads::executors::local_priority_queue_attached
```

```
#include <hpx/execution/executors/persistent_auto_chunk_size.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

```
struct persistent_auto_chunk_size
```

#include <persistent_auto_chunk_size.hpp> Loop iterations are divided into pieces and then assigned to threads. The number of loop iterations combined is determined based on measurements of how long the execution of 1% of the overall number of iterations takes. This executor parameters type makes sure that as many loop iterations are combined as necessary to run for the amount of time specified.

Public Functions

```
constexpr persistent_auto_chunk_size()
```

Construct an `persistent_auto_chunk_size` executor parameters object

Note Default constructed `persistent_auto_chunk_size` executor parameter types will use 0 microseconds as the execution time for each chunk and 80 microseconds as the minimal time for which any of the scheduled chunks should run.

```
persistent_auto_chunk_size(hpx::util::steady_duration const &time_cs)
```

Construct an `persistent_auto_chunk_size` executor parameters object

Parameters

- `time_cs`: The execution time for each chunk.

```
persistent_auto_chunk_size(hpx::util::steady_duration const &time_cs,
                           hpx::util::steady_duration const &rel_time)
```

Construct an `persistent_auto_chunk_size` executor parameters object

Parameters

- `rel_time`: [in] The time duration to use as the minimum to decide how many loop iterations should be combined.
- `time_cs`: The execution time for each chunk.

```
#include <hpx/execution/executors/dynamic_chunk_size.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

```
            struct dynamic_chunk_size
```

`#include <dynamic_chunk_size.hpp>` Loop iterations are divided into pieces of size `chunk_size` and then dynamically scheduled among the threads; when a thread finishes one chunk, it is dynamically assigned another. If `chunk_size` is not specified, the default chunk size is 1.

Note This executor parameters type is equivalent to OpenMP's DYNAMIC scheduling directive.

Public Functions

```
constexpr dynamic_chunk_size(std::size_t chunk_size = 1)
```

Construct a `dynamic_chunk_size` executor parameters object

Parameters

- `chunk_size`: [in] The optional chunk size to use as the number of loop iterations to schedule together. The default chunk size is 1.

```
#include <hpx/execution/executors/execution_parameters_fwd.hpp>
```

```
#include <hpx/execution/executors/execution_fwd.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

```
            struct parallel_execution_tag
```

`#include <execution_fwd.hpp>` Function invocations executed by a group of parallel execution agents execute in unordered fashion. Any such invocations executing in the same thread are indeterminately sequenced with respect to each other.

Note `parallel_execution_tag` is weaker than `sequenced_execution_tag`.

```
            struct sequenced_execution_tag
```

`#include <execution_fwd.hpp>` Function invocations executed by a group of sequential execution agents execute in sequential order.

```
struct unsequenced_execution_tag
#include <execution_fwd.hpp> Function invocations executed by a group of vector execution agents are permitted to execute in unordered fashion when executed in different threads, and un-sequenced with respect to one another when executed in the same thread.
```

Note `unsequenced_execution_tag` is weaker than `parallel_execution_tag`.

```
#include <hpx/execution/executors/guided_chunk_size.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

```
struct guided_chunk_size
```

```
#include <guided_chunk_size.hpp> Iterations are dynamically assigned to threads in blocks as threads request them until no blocks remain to be assigned. Similar to dynamic_chunk_size except that the block size decreases each time a number of loop iterations is given to a thread. The size of the initial block is proportional to number_of_iterations / number_of_cores. Subsequent blocks are proportional to number_of_iterations_remaining / number_of_cores. The optional chunk size parameter defines the minimum block size. The default chunk size is 1.
```

Note This executor parameters type is equivalent to OpenMP's GUIDED scheduling directive.

Public Functions

```
constexpr guided_chunk_size (std::size_t min_chunk_size = 1)
```

Construct a `guided_chunk_size` executor parameters object

Parameters

- `min_chunk_size`: [in] The optional minimal chunk size to use as the minimal number of loop iterations to schedule together. The default minimal chunk size is 1.

```
#include <hpx/execution/executors/execution_parameters.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace execution
```

Functions

```
template<typename ...Params>
executor_parameters_join<Params...>::type join_executor_parameters (Params&&...
                                                               params)

template<typename Param>
Param &&join_executor_parameters (Param &&param)

template<typename ...Params>
struct executor_parameters_join
```

Public Types

```
template<>
using type = detail::executor_parameters<typename hpx::util::decay<Params>::type...>

template<typename Param>
struct executor_parameters_join<Param>
```

Public Types

```
template<>
using type = Param
```

```
#include <hpx/execution/executors/parallel_executor_aggregated.hpp>

template<>
struct parallel_policy_executor_aggregated<hpx::launch>
```

Public Types

```
template<>
using execution_category = parallel_execution_tag
    Associate the parallel_execution_tag executor tag type as a default with this executor.

template<>
using executor_parameters_type = static_chunk_size
    Associate the static_chunk_size executor parameters type as a default with this executor.
```

Public Functions

```
constexpr parallel_policy_executor_aggregated(hpx::launch           l           =
                                               hpx::launch::async_policy{},
                                               std::size_t spread = 4, std::size_t
                                               tasks = std::size_t(-1))
```

Create a new parallel executor.

```
template<typename F, typename S, typename ...Ts>
std::vector<hpx::future<void>> bulk_async_execute (F &&f, S const &shape, Ts&&... ts)
                                         const
```

```
namespace hpx

namespace parallel

namespace execution
```

Typedefs

```
using parallel_executor_aggregated = parallel_policy_executor_aggregated<hpx::launch::async_policy>
```

```
template<typename Policy = hpx::launch::async_policy>
struct parallel_policy_executor_aggregated
```

#include <parallel_executor_aggregated.hpp> A `parallel_executor_aggregated` creates groups of parallel execution agents that execute in threads implicitly created by the executor. This executor prefers continuing with the creating thread first before executing newly created threads.

This executor conforms to the concepts of a `TwoWayExecutor`, and a `BulkTwoWayExecutor`

Public Types

```
template<>
```

```
using execution_category = parallel_execution_tag
```

Associate the `parallel_execution_tag` executor tag type as a default with this executor.

```
template<>
```

```
using executor_parameters_type = static_chunk_size
```

Associate the `static_chunk_size` executor parameters type as a default with this executor.

Public Functions

```
constexpr parallel_policy_executor_aggregated(std::size_t spread = 4, std::size_t tasks = std::size_t(-1))
```

Create a new parallel executor.

```
template<typename F, typename S, typename ...Ts>
std::vector<hpx::future<void>> bulk_async_execute(F &&f, S const &shape, Ts&&... ts) const
```

```
template<>
```

```
struct parallel_policy_executor_aggregated<hpx::launch>
```

Public Types

```
template<>
```

```
using execution_category = parallel_execution_tag
```

Associate the `parallel_execution_tag` executor tag type as a default with this executor.

```
template<>
```

```
using executor_parameters_type = static_chunk_size
```

Associate the `static_chunk_size` executor parameters type as a default with this executor.

Public Functions

```
constexpr parallel_policy_executor_aggregated(hpx::launch l = hpx::launch::async_policy{}, std::size_t spread = 4, std::size_t tasks = std::size_t(-1))
```

Create a new parallel executor.

```
template<typename F, typename S, typename ...Ts>
std::vector<hpx::future<void>> bulk_async_execute(F &&f, S const &shape,
Ts&&... ts) const
```

```
#include <hpx/execution/executors/thread_execution.hpp>
```

```
namespace hpx
```

```
namespace threads
```

Functions

```
template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, hpx::lcos::future<typename hpx::util::detail::invoke_de-
```

```
template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, typename hpx::util::detail::invoke_deferred_result<F, T>
```

```
template<typename Executor, typename F, typename Future, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, hpx::lcos::future<typename hpx::util::detail::invoke_de-
```

```
template<typename Executor, typename F, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value>::type post(Executor &&exec,
F &&f, Ts&&... ts)
```

```
template<typename Executor, typename F, typename Hint, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value && std::is_same<typename hpx::util::decay<Hint>::type,
```

```
template<typename Executor, typename F, typename Shape, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, std::vector<hpx::lcos::future<typename parallel::execu
```

```
template<typename Executor, typename F, typename Shape, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, typename parallel::execution::detail::bulk_execute_res
```

```
template<typename Executor, typename F, typename Shape, typename Future, typename ...Ts>
std::enable_if<hpx::traits::is_threads_executor<Executor>::value, hpx::future<typename parallel::execution::detail::bulk_
```

```
#include <hpx/execution/executors/thread_pool_executors.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

Typedefs

```
using local_priority_queue_executor = threads::executors::local_priority_queue_executor
```

Creates a new local_priority_queue_executor

Parameters

- max_punits: [in] The maximum number of processing units to associate with the newly created executor.
- min_punits: [in] The minimum number of processing units to associate with the newly created executor (default: 1).

```
#include <hpx/execution/executors/auto_chunk_size.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace execution
```

struct auto_chunk_size

#include <auto_chunk_size.hpp> Loop iterations are divided into pieces and then assigned to threads. The number of loop iterations combined is determined based on measurements of how long the execution of 1% of the overall number of iterations takes. This executor parameters type makes sure that as many loop iterations are combined as necessary to run for the amount of time specified.

Public Functions

```
constexpr auto_chunk_size()
```

Construct an `auto_chunk_size` executor parameters object

Note Default constructed `auto_chunk_size` executor parameter types will use 80 microseconds as the minimal time for which any of the scheduled chunks should run.

```
auto_chunk_size(hpx::util::steady_duration const &rel_time)
```

Construct an `auto_chunk_size` executor parameters object

Parameters

- rel_time: [in] The time duration to use as the minimum to decide how many loop iterations should be combined.

```
#include <hpx/execution/datapar/execution_policy.hpp>

#include <hpx/execution/datapar/execution_policy_fwd.hpp>

#include <hpx/execution/parallel/traits/vector_pack_load_store.hpp>

#include <hpx/execution/parallel/traits/vector_pack_alignment_size.hpp>

#include <hpx/execution/parallel/traits/vector_pack_type.hpp>

#include <hpx/execution/parallel/traits/vector_pack_count_bits.hpp>

namespace hpx

    namespace parallel

        namespace traits
```

Functions

```
HPX_HOST_DEVICE std::size_t hpx::parallel::traits::count_bits(bool value)
```

topology

```
#include <compatibility/hpx/util/cache_aligned_data.hpp>

#include <compatibility/hpx/runtime/threads/cpu_mask.hpp>

#include <compatibility/hpx/runtime/threads/topology.hpp>

#include <hpx/topology/cpu_mask.hpp>

#include <hpx/topology/topology.hpp>

namespace hpx

    namespace threads
```

TypeDefs

```
using hwloc_bitmap_ptr = std::shared_ptr<hpx_hwloc_bitmap_wrapper>
```

Enums

```
enum hpx_hwloc_membind_policy
```

Please see hwloc documentation for the corresponding enums HWLOC_MEMBIND_XXX.

Values:

```
membind_default = HWLOC_MEMBIND_DEFAULT  
membind_firsttouch = HWLOC_MEMBIND_FIRSTTOUCH  
membind_bind = HWLOC_MEMBIND_BIND  
membind_interleave = HWLOC_MEMBIND_INTERLEAVE  
membind_replicate = HWLOC_MEMBIND_REPLICATE  
membind_nexttouch = HWLOC_MEMBIND_NEXTTOUCH  
membind_mixed = HWLOC_MEMBIND_MIXED  
membind_user = HWLOC_MEMBIND_MIXED + 256
```

Functions

```
topology &create_topology()  
std::size_t hardware_concurrency()  
std::size_t get_memory_page_size()  
struct hpx_hwloc_bitmap_wrapper
```

Public Functions

```
HPX_NON_COPYABLE(hpx_hwloc_bitmap_wrapper)  
hpx_hwloc_bitmap_wrapper()  
hpx_hwloc_bitmap_wrapper(void *bmp)  
~hpx_hwloc_bitmap_wrapper()  
void reset(hwloc_bitmap_t bmp)  
operator bool() const  
hwloc_bitmap_t get bmp() const
```

Private Members

hwloc_bitmap_t **bmp_**

Friends

```
std::ostream &operator<< (std::ostream &os, hpx_hwloc_bitmap_wrapper const *bmp)

struct topology
```

Public Functions

topology()

~topology()

`std::size_t get_socket_number (std::size_t num_thread, error_code& = throws) const`
Return the Socket number of the processing unit the given thread is running on.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`std::size_t get numa_node_number (std::size_t num_thread, error_code& = throws) const`
Return the NUMA node number of the processing unit the given thread is running on.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`mask_cref_type get_machine_affinity_mask (error_code &ec = throws) const`
Return a bit mask where each set bit corresponds to a processing unit available to the application.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`mask_type get_service_affinity_mask (mask_cref_type used_processing_units, error_code &ec = throws) const`
Return a bit mask where each set bit corresponds to a processing unit available to the service threads in the application.

Parameters

- `used_processing_units`: [in] This is the mask of processing units which are not available for service threads.
- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

`mask_cref_type get_socket_affinity_mask (std::size_t num_thread, error_code &ec = throws) const`
Return a bit mask where each set bit corresponds to a processing unit available to the given thread inside the socket it is running on.

Parameters

- `ec`: [in,out] this represents the error status on exit, if this is pre-initialized to `hpx::throws` the function will throw on error instead.

```
mask_cref_type get numa_node_affinity_mask(std::size_t num_thread, error_code &ec = throws) const
Return a bit mask where each set bit corresponds to a processing unit available to the given thread inside the NUMA domain it is running on.
```

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
mask_type get numa_node_affinity_mask_from_numa_node(std::size_t num_node) const
Return a bit mask where each set bit corresponds to a processing unit associated with the given NUMA node.
```

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
mask_cref_type get core affinity mask(std::size_t num_thread, error_code &ec = throws) const
Return a bit mask where each set bit corresponds to a processing unit available to the given thread inside the core it is running on.
```

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
mask_cref_type get thread affinity mask(std::size_t num_thread, error_code &ec = throws) const
Return a bit mask where each set bit corresponds to a processing unit available to the given thread.
```

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void set thread affinity mask(mask_cref_type mask, error_code &ec = throws)
const
```

Use the given bit mask to set the affinity of the given thread. Each set bit corresponds to a processing unit the thread will be allowed to run on.

Note Use this function on systems where the affinity must be set from inside the thread itself.

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
mask_type get thread affinity mask from lva(void const *lva, error_code &ec = throws) const
Return a bit mask where each set bit corresponds to a processing unit co-located with the memory the given address is currently allocated on.
```

Parameters

- *ec*: [in,out] this represents the error status on exit, if this is pre-initialized to *hpx::throws* the function will throw on error instead.

```
void print affinity mask(std::ostream &os, std::size_t num_thread, mask_cref_type m, const std::string &pool_name) const
```

Prints the.

Parameters

- *m*: to *os* in a human readable form

```

bool reduce_thread_priority(error_code &ec = throws) const
    Reduce thread priority of the current thread.

Parameters
    • ec: [in,out] this represents the error status on exit, if this is pre-initialized to hpx::throws the
        function will throw on error instead.

std::size_t get_number_of_sockets() const
    Return the number of available NUMA domains.

std::size_t get_number_of numa_nodes() const
    Return the number of available NUMA domains.

std::size_t get_number_of_cores() const
    Return the number of available cores.

std::size_t get_number_of_pus() const
    Return the number of available hardware processing units.

std::size_t get_number_of numa_node_cores(std::size_t numa) const
    Return number of cores in given numa domain.

std::size_t get_number_of numa_node_pus(std::size_t numa) const
    Return number of processing units in a given numa domain.

std::size_t get_number_of_socket_pus(std::size_t socket) const
    Return number of processing units in a given socket.

std::size_t get_number_of_core_pus(std::size_t core) const
    Return number of processing units in given core.

std::size_t get_number_of_socket_cores(std::size_t socket) const
    Return number of cores units in given socket.

std::size_t get_core_number(std::size_t num_thread, error_code& = throws) const

std::size_t get_pu_number(std::size_t num_core, std::size_t num_pu, error_code &ec = throws) const

mask_type get_cpubind_mask(error_code &ec = throws) const

mask_type get_cpubind_mask(std::thread &handle, error_code &ec = throws) const

hwloc_bitmap_ptr cpuset_to_nodeset(mask_cref_type cpuset) const
    convert a cpu mask into a numa node mask in hwloc bitmap form

void write_to_log() const

void *allocate(std::size_t len) const
    This is equivalent to malloc(), except that it tries to allocate page-aligned memory from the OS.

void *allocate_membind(std::size_t len, hwloc_bitmap_ptr bitmap,
                      hpx_hwloc_membind_policy policy, int flags) const
    allocate memory with binding to a numa node set as specified by the policy and flags (see hwloc
    docs)

threads::mask_type get_area_membind_nodeset(const void *addr, std::size_t len)
                                              const

bool set_area_membind_nodeset(const void *addr, std::size_t len, void *nodeset)
                                 const

```

```
int get_numa_domain(const void *addr) const
void deallocate(void *addr, std::size_t len) const
    Free memory that was previously allocated by allocate.

void print_vector(std::ostream &os, std::vector<std::size_t> const &v) const
void print_mask_vector(std::ostream &os, std::vector<mask_type> const &v) const
void print_hwloc(std::ostream&) const

mask_type init_socket_affinity_mask_from_socket(std::size_t num_socket)
const
mask_type init numa_node_affinity_mask_from numa_node(std::size_t num numa_node)
const
mask_type init core affinity mask from core(std::size_t num_core,
mask_cref_type default_mask
= empty mask) const
mask_type init thread affinity mask(std::size_t num_thread) const
mask_type init thread affinity mask(std::size_t num_core, std::size_t num_pu)
const
hwloc_bitmap_t mask_to_bitmap(mask_cref_type mask, hwloc_obj_type_t htype) const
mask_type bitmap_to_mask(hwloc_bitmap_t bitmap, hwloc_obj_type_t htype) const
```

Private Types

```
using mutex_type = hpx::util::spinlock
```

Private Functions

```
std::size_t init_node_number(std::size_t num_thread, hwloc_obj_type_t type)

std::size_t init_socket_number(std::size_t num_thread)

std::size_t init numa_node_number(std::size_t num_thread)

std::size_t init core number(std::size_t num_thread)

void extract_node_mask(hwloc_obj_t parent, mask_type &mask) const

std::size_t extract_node_count(hwloc_obj_t parent, hwloc_obj_type_t type, std::size_t count) const

mask_type init machine affinity mask() const

mask_type init socket affinity mask(std::size_t num_thread) const

mask_type init numa node affinity mask(std::size_t num_thread) const

mask_type init core affinity mask(std::size_t num_thread) const

void init num of pus()
```

Private Members

```
hwloc_topology_t topo
std::size_t num_of_pus_
mutex_type topo_mtx
std::vector<std::size_t> socket_numbers_
std::vector<std::size_t> numa_node_numbers_
std::vector<std::size_t> core_numbers_
mask_type machine_affinity_mask_
std::vector<mask_type> socket_affinity_masks_
std::vector<mask_type> numa_node_affinity_masks_
std::vector<mask_type> core_affinity_masks_
std::vector<mask_type> thread_affinity_masks_
```

Private Static Attributes

```
mask_type empty_mask
std::size_t memory_page_size_
const std::size_t pu_offset = 0
const std::size_t core_offset = 0
```

Friends

```
std::size_t get_memory_page_size()
```

allocator_support

```
#include <compatibility/hpx/util/internal_allocator.hpp>

#include <compatibility/hpx/util/allocator_deleter.hpp>

#include <hpx/allocator_support/internal_allocator.hpp>

namespace hpx

namespace util
```

Typedefs

```
using internal_allocator = std::allocator<T>

#include <hpx/allocator_support/allocator_deleter.hpp>

namespace hpx

namespace util

template<typename Allocator>
struct allocator_deleter
```

Public Functions

```
template<typename SharedState>
void operator()(SharedState *state)
```

Public Members

Allocator **alloc_**

segmented_algorithms

```
#include <hpx/parallel/segmented_algorithm.hpp>

#include <hpx/parallel/segmented_algorithms/adjacent_difference.hpp>

#include <hpx/parallel/segmented_algorithms/transform_inclusive_scan.hpp>

#include <hpx/parallel/segmented_algorithms/reduce.hpp>

#include <hpx/parallel/segmented_algorithms/find.hpp>

#include <hpx/parallel/segmented_algorithms/exclusive_scan.hpp>

#include <hpx/parallel/segmented_algorithms/inclusive_scan.hpp>

#include <hpx/parallel/segmented_algorithms/transform_reduce.hpp>

#include <hpx/parallel/segmented_algorithms/fill.hpp>

#include <hpx/parallel/segmented_algorithms/generate.hpp>

#include <hpx/parallel/segmented_algorithms/transform_exclusive_scan.hpp>
```

```
#include <hpx/parallel/segmented_algorithms/all_any_none.hpp>

#include <hpx/parallel/segmented_algorithms/minmax.hpp>

#include <hpx/parallel/segmented_algorithms/adjacent_find.hpp>

#include <hpx/parallel/segmented_algorithms/transform.hpp>

#include <hpx/parallel/segmented_algorithms/count.hpp>

#include <hpx/parallel/segmented_algorithms/for_each.hpp>

testing

#include <compatibility/hpx/util/lightweight_test.hpp>

#include <hpx/testing.hpp>

Defines

HPX_TEST(expr)

HPX_TEST_MSG(expr, msg)

HPX_TEST_EQ(expr1, expr2)

HPX_TEST_NEQ(expr1, expr2)

HPX_TEST_LT(expr1, expr2)

HPX_TEST_LTE(expr1, expr2)

HPX_TEST_RANGE(expr1, expr2, expr3)

HPX_TEST_EQ_MSG(expr1, expr2, msg)

HPX_TEST_NEQ_MSG(expr1, expr2, msg)

HPX_SANITY(expr)

HPX_SANITY_MSG(expr, msg)

HPX_SANITY_EQ(expr1, expr2)

HPX_SANITY_NEQ(expr1, expr2)

HPX_SANITY_LT(expr1, expr2)

HPX_SANITY_LTE(expr1, expr2)

HPX_SANITY_RANGE(expr1, expr2, expr3)

HPX_SANITY_EQ_MSG(expr1, expr2, msg)

HPX_TEST_THROW(expression, exception)

namespace hpx

namespace util
```

TypeDefs

```
using test_failure_handler_type = std::function<void ()>
```

Enums

```
enum counter_type
```

Values:

```
    counter_sanity
```

```
    counter_test
```

Functions

```
void set_test_failure_handler (test_failure_handler_type f)
```

```
int report_errors (std::ostream &stream = std::cerr)
```

```
void print_cdash_timing (const char *name, double time)
```

```
void print_cdash_timing (const char *name, std::uint64_t time)
```

program_options

```
#include <hpx/program_options.hpp>
```

```
#include <hpx/program_options/environment_iterator.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

```
class environment_iterator : public hpx::program_options::eof_iterator<environment_iterator, std::pair<std::string,
```

Public Functions

```
environment_iterator (char **environment)
```

```
environment_iterator ()
```

```
void get ()
```

Private Members

```
char **m_environment

#include <hpx/program_options/positional_options.hpp>

namespace hpx

namespace program_options

class positional_options_description
    #include <positional_options.hpp> Describes positional options.
```

The class allows to guess option names for positional options, which are specified on the command line and are identified by the position. The class uses the information provided by the user to associate a name with every positional option, or tell that no name is known.

The primary assumption is that only the relative order of the positional options themselves matters, and that any interleaving ordinary options don't affect interpretation of positional options.

The user initializes the class by specifying that first N positional options should be given the name X1, following M options should be given the name X2 and so on.

Public Functions

```
positional_options_description()

positional_options_description &add(const char *name, int max_count)
    Species that up to ‘max_count’ next positional options should be given the ‘name’. The value of
    ‘-1’ means ‘unlimited’. No calls to ‘add’ can be made after call with ‘max_value’ equal to ‘-1’.

unsigned max_total_count() const
    Returns the maximum number of positional options that can be present. Can return (numeric_limits<unsigned>::max)() to indicate unlimited number.

const std::string &name_for_position(unsigned position) const
    Returns the name that should be associated with positional options at ‘position’. Precondition:
    position < max_total_count()
```

Private Members

```
std::vector<std::string> m_names
std::string m_trailing
```

```
#include <hpx/program_options/variables_map.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

Functions

```
void store(const basic_parsed_options<char> &options, variables_map &m, bool utf8 = false)
```

Stores in ‘m’ all options that are defined in ‘options’. If ‘m’ already has a non-defaulted value of an option, that value is not changed, even if ‘options’ specify some value.

```
void store(const basic_parsed_options<wchar_t> &options, variables_map &m)
```

Stores in ‘m’ all options that are defined in ‘options’. If ‘m’ already has a non-defaulted value of an option, that value is not changed, even if ‘options’ specify some value. This is wide character variant.

```
void notify(variables_map &m)
```

Runs all ‘notify’ function for options in ‘m’.

```
class abstract_variables_map
```

#include <variables_map.hpp> Implements string->string mapping with convenient value casting facilities.

Subclassed by *hpx::program_options::variables_map*

Public Functions

```
abstract_variables_map()
```

```
abstract_variables_map(const abstract_variables_map *next)
```

```
virtual ~abstract_variables_map()
```

```
const variable_value &operator[](const std::string &name) const
```

Obtains the value of variable ‘name’, from *this and possibly from the chain of variable maps.

- if there’s no value in *this.
 - if there’s next variable map, returns value from it
 - otherwise, returns empty value
- if there’s defaulted value
 - if there’s next variable map, which has a non-defaulted value, return that
 - otherwise, return value from *this
- if there’s a non-defaulted value, returns it.

```
void next(abstract_variables_map *next)
```

Sets next variable map, which will be used to find variables not found in *this.

Private Functions

virtual const variable_value &get (const std::string &name) const = 0
 Returns value of variable ‘name’ stored in *this, or empty value otherwise.

Private Members

```
const abstract_variables_map *m_next

template<class Char>
class basic_parsed_options
#include <parsers.hpp> Results of parsing an input source. The primary use of this class is passing
information from parsers component to value storage component. This class does not makes much
sense itself.

class variable_value
#include <variables_map.hpp> Class holding value of option. Contains details about how the value
is set and allows to conveniently obtain the value.
```

Public Functions

```
variable_value()

variable_value (const hpx::util::any_nonser &xv, bool xdefaulted)

template<class T>
const T &as () const
    If stored value if of type T, returns that value. Otherwise, throws boost::bad_any_cast exception.

template<class T>
T &as ()
    This is an overloaded member function, provided for convenience. It differs from the above
    function only in what argument(s) it accepts.

bool empty () const
    Returns true if no value is stored.

bool defaulted () const
    Returns true if the value was not explicitly given, but has default value.

const hpx::util::any_nonser &value () const
    Returns the contained value.



```

Private Members

```
hpx::util::any_nonser v
bool m_defaulted
std::shared_ptr<const value_semantic> m_value_semantic
```

Friends

```
friend hpx::program_options::variables_map
```

```
void store(const basic_parsed_options<char> &options, variables_map &m, bool utf8)
```

Stores in ‘m’ all options that are defined in ‘options’. If ‘m’ already has a non-defaulted value of an option, that value is not changed, even if ‘options’ specify some value.

```
class variables_map : public hpx::program_options::abstract_variables_map, public std::map<std::string, var
```

#include <variables_map.hpp> Concrete variables map which store variables in real map.

This class is derived from std::map<std::string, variable_value>, so you can use all map operators to examine its content.

Public Functions

```
variables_map()
```

```
variables_map(const abstract_variables_map *next)
```

```
const variable_value &operator[](const std::string &name) const
```

```
void clear()
```

```
void notify()
```

Private Functions

```
const variable_value &get(const std::string &name) const
```

Implementation of abstract_variables_map::get which does ‘find’ in *this.

Private Members

```
std::set<std::string> m_final
```

Names of option with ‘final’ values – which should not be changed by subsequent assignments.

```
std::map<std::string, std::string> m_required
```

Names of required options, filled by parser which has access to *options_description*. The map values are the “canonical” names for each corresponding option. This is useful in creating diagnostic messages when the option is absent.

Friends

void `store` (`const basic_parsed_options<char> &options, variables_map &xm, bool utf8)`
 Stores in ‘m’ all options that are defined in ‘options’. If ‘m’ already has a non-defaulted value of an option, that value is not changed, even if ‘options’ specify some value.

```
#include <hpx/program_options/errors.hpp>

namespace hpx
```

```
namespace program_options
```

Functions

std::string `strip_prefixes` (`const std::string &text`)
class `ambiguous_option` : public `hpx::program_options::error_with_no_option_name`
#include <errors.hpp> Class thrown when there’s ambiguity among several options.

Public Functions

```
ambiguous_option (const std::vector<std::string> &xalternatives)  

~ambiguous_option ()  

const std::vector<std::string> &alternatives () const
```

Protected Functions

void `substitute_placeholders` (`const std::string &error_template`) **const**
 Makes all substitutions using the template

Private Members

```
std::vector<std::string> m_alternatives  

class error : public logic_error  

#include <errors.hpp> Base class for all errors in the library.  

Subclassed by hpx::program_options::duplicate_option_error, hpx::program_options::error_with_option_name, hpx::program_options::invalid_command_line_style, hpx::program_options::reading_file, hpx::program_options::too_many_positional_options_error
```

Public Functions

```
error (const std::string &xwhat)

class error_with_no_option_name : public hpx::program_options::error_with_option_name
#include <errors.hpp> Base class of un-parsable options, when the desired option cannot be identified.
```

It makes no sense to have an option name, when we can't match an option to the parameter

Having this a part of the `error_with_option_name` hierarchy makes error handling a lot easier, even if the name indicates some sort of conceptual dissonance!

Subclassed by `hpx::program_options::ambiguous_option`, `hpx::program_options::unknown_option`

Public Functions

```
error_with_no_option_name (const std::string &template_, const std::string &original_token = "")
```

```
void set_option_name (const std::string&)
```

Does NOT set option name, because no option name makes sense

```
~error_with_no_option_name ()
```

```
class error_with_option_name : public hpx::program_options::error
#include <errors.hpp> Base class for most exceptions in the library.
```

Substitutes the values for the parameter name placeholders in the template to create the human readable error message

Placeholders are surrounded by % signs: example% Poor man's version of boost::format

If a parameter name is absent, perform default substitutions instead so ugly placeholders are never left in-place.

Options are displayed in “canonical” form This is the most unambiguous form of the *parsed* option name and would correspond to `option_description::format_name()` i.e. what is shown by `print_usage()`

The “canonical” form depends on whether the option is specified in short or long form, using dashes or slashes or without a prefix (from a configuration file)

Subclassed by `hpx::program_options::error_with_no_option_name`,
`hpx::program_options::invalid_syntax`, `hpx::program_options::multiple_occurrences`,
`hpx::program_options::multiple_values`, `hpx::program_options::required_option`,
`hpx::program_options::validation_error`

Public Functions

```
error_with_option_name (const std::string &template_, const std::string &option_name = "", const std::string &original_token = "", int option_style = 0)
```

```
~error_with_option_name ()
```

gcc says that throw specification on dtor is loosened without this line

```

void set_substitute (const std::string &parameter_name, const std::string &value)
    Substitute parameter_name->value to create the error message from the error template

void set_substitute_default (const std::string &parameter_name, const std::string &from, const std::string &to)
    If the parameter is missing, then make the from->to substitution instead

void add_context (const std::string &option_name, const std::string &original_token, int
                  option_style)
    Add context to an exception

void set_prefix (int option_style)

virtual void set_option_name (const std::string &option_name)
    Overridden in error_with_no_option_name

std::string get_option_name () const

void set_original_token (const std::string &original_token)

const char *what () const
    Creates the error_message on the fly Currently a thin wrapper for substitute_placeholders()

```

Public Members

std::string **m_error_template**
 template with placeholders

Protected Types

using string_pair = std::pair<std::string, std::string>

Protected Functions

```

virtual void substitute_placeholders (const std::string &error_template)
                                         const
    Makes all substitutions using the template

void replace_token (const std::string &from, const std::string &to) const

std::string get_canonical_option_name () const
    Construct option name in accordance with the appropriate prefix style: i.e. long dash or short
    slash etc

std::string get_canonical_option_prefix () const

```

Protected Attributes

```
int m_option_style
    can be 0 = no prefix (config file options) allow_long allow_dash_for_short allow_slash_for_short
    allow_long_disguise

std::map<std::string, std::string> m_substitutions
    substitutions from placeholders to values

std::map<std::string, string_pair> m_substitution_defaults

std::string m_message
    Used to hold the error text returned by what()
```

```
class invalid_bool_value : public hpx::program_options::validation_error
#include <errors.hpp> Class thrown if there is an invalid bool value given
```

Public Functions

```
invalid_bool_value(const std::string &value)

class invalid_command_line_style : public hpx::program_options::error
#include <errors.hpp> Class thrown when there are programming error related to style
```

Public Functions

```
invalid_command_line_style(const std::string &msg)

class invalid_command_line_syntax : public hpx::program_options::invalid_syntax
#include <errors.hpp> Class thrown when there are syntax errors in given command line
```

Public Functions

```
invalid_command_line_syntax(kind_t kind, const std::string &option_name = "",
                           const std::string &original_token = "", int option_style = 0)

~invalid_command_line_syntax()

class invalid_config_file_syntax : public hpx::program_options::invalid_syntax
```

Public Functions

```
invalid_config_file_syntax(const std::string &invalid_line, kind_t kind)

~invalid_config_file_syntax()

std::string tokens() const
    Convenience functions for backwards compatibility

class invalid_option_value : public hpx::program_options::validation_error
#include <errors.hpp> Class thrown if there is an invalid option value given
```

Public Functions

```
invalid_option_value(const std::string &value)
invalid_option_value(const std::wstring &value)

class invalid_syntax : public hpx::program_options::error_with_option_name
    #include <errors.hpp> Class thrown when there's syntax error either for command line or config file
    options. See derived children for concrete classes.

Subclassed by hpx::program_options::invalid_command_line_syntax,
hpx::program_options::invalid_config_file_syntax
```

Public Types

```
enum kind_t
Values:
long_not_allowed = 30
long_adjacent_not_allowed
short_adjacent_not_allowed
empty_adjacent_parameter
missing_parameter
extra_parameter
unrecognized_line
```

Public Functions

```
invalid_syntax(kind_t kind, const std::string &option_name = "", const std::string &original_token = "", int option_style = 0)
~invalid_syntax()

kind_t kind() const
virtual std::string tokens() const
    Convenience functions for backwards compatibility
```

Protected Functions

```
std::string get_template(kind_t kind)
    Used to convert kind_t to a related error text
```

Protected Attributes

`kind_t m_kind`

```
class multiple_occurrences : public hpx::program_options::error_with_option_name
#include <errors.hpp> Class thrown when there are several occurrences of an option, but user called a method which cannot return them all.
```

Public Functions

`multiple_occurrences()`

`~multiple_occurrences()`

```
class multiple_values : public hpx::program_options::error_with_option_name
#include <errors.hpp> Class thrown when there are several option values, but user called a method which cannot return them all.
```

Public Functions

`multiple_values()`

`~multiple_values()`

```
class reading_file : public hpx::program_options::error
#include <errors.hpp> Class thrown if config file can not be read
```

Public Functions

`reading_file(const char *filename)`

```
class required_option : public hpx::program_options::error_with_option_name
#include <errors.hpp> Class thrown when a required/mandatory option is missing
```

Public Functions

`required_option(const std::string &option_name)`

`~required_option()`

```
class too_many_positional_options_error : public hpx::program_options::error
#include <errors.hpp> Class thrown when there are too many positional options. This is a programming error.
```

Public Functions

```
too_many_positional_options_error()

class unknown_option : public hpx::program_options::error_with_no_option_name
    #include <errors.hpp> Class thrown when option name is not recognized.
```

Public Functions

```
unknown_option(const std::string &original_token = "")

~unknown_option()

class validation_error : public hpx::program_options::error_with_option_name
    #include <errors.hpp> Class thrown when value of option is incorrect.

Subclassed by hpx::program_options::invalid_bool_value, hpx::program_options::invalid_option_value
```

Public Types

```
enum kind_t
    Values:
        multiple_values_not_allowed = 30
        at_least_one_value_required
        invalid_bool_value
        invalid_option_value
        invalid_option
```

Public Functions

```
validation_error(kind_t kind, const std::string &option_name = "", const std::string
    &original_token = "", int option_style = 0)

~validation_error()

kind_t kind() const
```

Protected Functions

```
std::string get_template(kind_t kind)
    Used to convert kind_t to a related error text
```

Protected Attributes

`kind_t m_kind`

```
#include <hpx/program_options/options_description.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

```
class duplicate_option_error : public hpx::program_options::error
```

`#include <options_description.hpp>` Class thrown when duplicate option description is found.

Public Functions

`duplicate_option_error(const std::string &xwhat)`

```
class option_description
```

`#include <options_description.hpp>` Describes one possible command line/config file option. There are two kinds of properties of an option. First describe it syntactically and are used only to validate input. Second affect interpretation of the option, for example default value for it or function that should be called when the value is finally known. Routines which perform parsing never use second kind of properties – they are side effect free.

See `options_description`

Public Types

```
enum match_result
```

Values:

`no_match`

`full_match`

`approximate_match`

Public Functions

`option_description()`

`option_description(const char *name, const value_semantic *s)`

Initializes the object with the passed data.

Note: it would be nice to make the second parameter auto_ptr, to explicitly pass ownership. Unfortunately, it's often needed to create objects of types derived from '`value_semantic`': `options_description d; d.add_options()("a", parameter<int>("n")->default_value(1));` Here, the static type returned by 'parameter' should be derived from `value_semantic`.

Alas, derived->base conversion for auto_ptr does not really work, see <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2000/n1232.pdf> http://www.open-std.org/jtc1/sc22/wg21/docs/cwg_defects.html#84

So, we have to use plain old pointers. Besides, users are not expected to use the constructor directly.

The ‘name’ parameter is interpreted by the following rules:

- if there’s no “,” character in ‘name’, it specifies long name
- otherwise, the part before “,” specifies long name and the part after – short name.

```
option_description(const char *name, const value_semantic *s, const char *de-  
scription)
```

Initializes the class with the passed data.

```
virtual ~option_description()
```

```
match_result match(const std::string &option, bool approx, bool long_ignore_case, bool  
short_ignore_case) const
```

Given ‘option’, specified in the input source, returns ‘true’ if ‘option’ specifies `*this`.

```
const std::string &key(const std::string &option) const
```

Returns the key that should identify the option, in particular in the `variables_map` class. The ‘option’ parameter is the option spelling from the input source. If option name contains ‘*’, returns ‘option’. If long name was specified, it’s the long name, otherwise it’s a short name with pre-pended ‘-’.

```
std::string canonical_display_name(int canonical_option_style = 0) const
```

Returns the canonical name for the option description to enable the user to recognize a matching option. 1) For short options ('-', '/'), returns the short name prefixed. 2) For long options ('' / '-') returns the first long name prefixed 3) All other cases, returns the first long name (if present) or the short name, un-prefixed.

```
const std::string &long_name() const
```

```
const std::pair<const std::string *, std::size_t> long_names() const
```

```
const std::string &description() const
```

Explanation of this option.

```
std::shared_ptr<const value_semantic> semantic() const
```

Semantic of option’s value.

```
std::string format_name() const
```

Returns the option name, formatted suitably for usage message.

```
std::string format_parameter() const
```

Returns the parameter name and properties, formatted suitably for usage message.

Private Functions

```
option_description &set_names(const char *name)
```

Private Members

`std::string m_short_name`
a one-character “switch” name - with its prefix, so that this is either empty or has length 2 (e.g. “-c”)

`std::vector<std::string> m_long_names`
one or more names by which this option may be specified on a command-line or in a config file, which are not a single-letter switch. The names here are *without* any prefix.

`std::string m_description`

`std::shared_ptr<const value_semantic> m_value_semantic`

`class options_description`
`#include <options_description.hpp>` A set of option descriptions. This provides convenient interface for adding new option (the add_options) method, and facilities to search for options by name.

See here for option adding interface discussion.
See [option_description](#)

Public Functions

`options_description(unsigned line_length = m_default_line_length, unsigned min_description_length = m_default_line_length / 2)`

Creates the instance.

`options_description(const std::string &caption, unsigned line_length = m_default_line_length, unsigned min_description_length = m_default_line_length / 2)`

Creates the instance. The ‘caption’ parameter gives the name of this ‘[options_description](#)’ instance. Primarily useful for output. The ‘description_length’ specifies the number of columns that should be reserved for the description text; if the option text encroaches into this, then the description will start on the next line.

`void add(std::shared_ptr<option_description> desc)`

Adds new variable description. Throws duplicate_variable_error if either short or long name matches that of already present one.

`options_description &add(const options_description &desc)`

Adds a group of option description. This has the same effect as adding all option_descriptions in ‘desc’ individually, except that output operator will show a separate group. Returns *this.

`std::size_t get_option_column_width() const`

Find the maximum width of the option column, including options in groups.

`options_description_easy_init add_options()`

Returns an object of implementation-defined type suitable for adding options to [options_description](#). The returned object will have overloaded operator() with parameter type matching ‘[option_description](#)’ constructors. Calling the operator will create new [option_description](#) instance and add it.

`const option_description &find(const std::string &name, bool approx, bool long_ignore_case = false, bool short_ignore_case = false) const`

```
const option_description *find_nothrow(const std::string &name, bool approx,
                                         bool long_ignore_case = false, bool
                                         short_ignore_case = false) const

const std::vector<std::shared_ptr<option_description>> &options() const

void print(std::ostream &os, std::size_t width = 0) const
    Outputs ‘desc’ to the specified stream, calling ‘f’ to output each option_description element.
```

Public Static Attributes

```
const unsigned m_default_line_length
```

Private Types

```
using name2index_iterator = std::map<std::string, int>::const_iterator
using approximation_range = std::pair<name2index_iterator, name2index_iterator>
```

Private Members

```
std::string m_caption
std::size_t const m_line_length
std::size_t const m_min_description_length
std::vector<std::shared_ptr<option_description>> m_options
std::vector<char> belong_to_group
std::vector<std::shared_ptr<options_description>> groups
```

Friends

```
std::ostream &operator<<(std::ostream &os, const options_description &desc)
    Produces a human readable output of ‘desc’, listing options, their descriptions and allowed parameters. Other options_description instances previously passed to add will be output separately.
```

```
class options_description_easy_init
    #include <options_description.hpp> Class which provides convenient creation syntax to option_description.
```

Public Functions

```
options_description_easy_init(options_description *owner)
options_description_easy_init &operator() (const char *name, const char *description)
options_description_easy_init &operator() (const char *name, const value_semantic *s)
options_description_easy_init &operator() (const char *name, const value_semantic *s,
                                         const char *description)
```

Private Members

options_description ***owner**

```
#include <hpx/program_options/parsers.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

TypeDefs

```
using parsed_options = basic_parsed_options<char>
using wparsed_options = basic_parsed_options<wchar_t>
using ext_parser = std::function<std::pair<std::string, std::string> (const std::string&) >
    Augments basic_parsed_options<wchar_t> with conversion from 'parsed_options'
using command_line_parser = basic_command_line_parser<char>
using wcommand_line_parser = basic_command_line_parser<wchar_t>
```

Enums

```
enum collect_unrecognized_mode
```

Controls if the 'collect_unregistered' function should include positional options, or not.

Values:

```
include_positional
exclude_positional
```

Functions

```
template<class Char>
basic_parsed_options<Char> parse_command_line (int argc, const Char *const argv[],
                                                const options_description&, int style
                                                = 0, std::function<std::pair<std::string,
                                                std::string> () const std::string&
                                                > ext = ext_parser())Creates instance of 'command_line_parser', passes parameters to it, and returns
the result of calling the 'run' method.
```

```
template<class Char>
basic_parsed_options<Char> parse_config_file (std::basic_istream<Char>&,           const
                                                options_description&,           bool      al-
                                                low_unregistered = false)
```

Parse a config file.

Read from given stream.

```
template<class Char = char>
```

```
basic_parsed_options<Char> parse_config_file(const char *filename, const
                                              options_description&, bool allow_unregistered = false)
```

Parse a config file.

Read from file with the given name. The character type is passed to the file stream.

```
template<class Char>
std::vector<std::basic_string<Char>> collect_unrecognized(const
                                                          std::vector<basic_option<Char>>
                                                          &options, enum collect_unrecognized_mode mode)
```

Collects the original tokens for all named options with ‘unregistered’ flag set. If ‘mode’ is ‘include_positional’ also collects all positional options. Returns the vector of original tokens for all collected options.

```
parsed_options parse_environment (const options_description&, const
                                  std::function<std::string> std::string
                                  > &name_mapper) Parse environment.
```

For each environment variable, the ‘name_mapper’ function is called to obtain the option name. If it returns empty string, the variable is ignored.

This is done since naming of environment variables is typically different from the naming of command line options.

```
parsed_options parse_environment (const options_description&, const std::string &prefix)
```

Parse environment.

Takes all environment variables which start with ‘prefix’. The option name is obtained from variable name by removing the prefix and converting the remaining string into lower case.

```
parsed_options parse_environment (const options_description&, const char *prefix)
```

This is an overloaded member function, provided for convenience. It differs from the above function only in what argument(s) it accepts. This function exists to resolve ambiguity between the two above functions when second argument is of ‘char*’ type. There’s implicit conversion to both std::function and string.

```
std::vector<std::string> split_unix (const std::string &cmdline)
```

Splits a given string to a collection of single strings which can be passed to command_line_parser. The second parameter is used to specify a collection of possible separator chars used for splitting. The separator is defaulted to space “ ”. Splitting is done in a unix style way, with respect to quotes “” and escape characters “”.

```
std::vector<std::wstring> split_unix (const std::wstring &cmdline)
```

This is an overloaded member function, provided for convenience. It differs from the above function only in what argument(s) it accepts.

```
template<class Char>
class basic_command_line_parser : private cmdline
#include <parsers.hpp> Command line parser.
```

The class allows one to specify all the information needed for parsing and to parse the command line. It is primarily needed to emulate named function parameters – a regular function with 5 parameters will be hard to use and creating overloads with a smaller number of parameters will be confusing.

For the most common case, the function parse_command_line is a better alternative.

There are two typedefs – command_line_parser and wcommand_line_parser, for charT == char and charT == wchar_t cases.

Public Functions

basic_command_line_parser (**const std::vector<std::basic_string<Char>> &args**)

Creates a command line parser for the specified arguments list. The ‘args’ parameter should not include program name.

basic_command_line_parser (**int argc, const Char *const argv[]**)

Creates a command line parser for the specified arguments list. The parameters should be the same as passed to ‘main’.

basic_command_line_parser &options (**const options_description &desc**)

Sets options descriptions to use.

basic_command_line_parser &positional (**const positional_options_description &desc**)

Sets positional options description to use.

basic_command_line_parser &style (**int**)

Sets the command line style.

basic_command_line_parser &extra_parser (**ext_parser**)

Sets the extra parsers.

basic_parsed_options<Char> run ()

Parses the options and returns the result of parsing. Throws on error.

basic_command_line_parser &allow_unregistered()

Specifies that unregistered options are allowed and should be passed though. For each command like token that looks like an option but does not contain a recognized name, an instance of basic_option<charT> will be added to result, with ‘unrecognized’ field set to ‘true’. It’s possible to collect all unrecognized options with the ‘collect_unrecognized’ function.

basic_command_line_parser &extra_style_parser (**style_parser s**)

Private Members

const options_description *m_desc

template<>

class basic_parsed_options<wchar_t>

#include <parsers.hpp> Specialization of *basic_parsed_options* which:

- provides convenient conversion from *basic_parsed_options<char>*
- stores the passed char-based options for later use.

Public Functions

basic_parsed_options (**const basic_parsed_options<char>** &*po*)
 Constructs wrapped options from options in UTF8 encoding.

Public Members

```
std::vector<basic_option<wchar_t>> options
const options_description *description
basic_parsed_options<char> utf8_encoded_options
    Stores UTF8 encoded options that were passed to constructor, to avoid reverse conversion in some cases.
int m_options_prefix
    Mainly used for the diagnostic messages in exceptions. The canonical option prefix for the parser which generated these results, depending on the settings for basic_command_line_parser::style() or cmdline::style(). In order of precedence of command_line_style enums: allow_long allow_long_disguise allow_dash_for_short allow_slash_for_short
```

```
#include <hpx/program_options/version.hpp>
```

Defines

HPX_PROGRAM_OPTIONS_VERSION

The version of the source interface. The value will be incremented whenever a change is made which might cause compilation errors for existing code.

HPX_PROGRAM_OPTIONS_IMPLICIT_VALUE_NEXT_TOKEN

```
#include <hpx/program_options/config.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

Typedefs

```
using any = hpx::util::any_nonser
using optional = hpx::util::optional<T>
```

```
#include <hpx/program_options/force_linking.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

Typedefs

```
using parse_environment1_type = basic_parsed_options<char> (*) (options_description  
const&, char  
const *)  
using parse_environment2_type = basic_parsed_options<char> (*) (options_description  
const&, std::function<std::string> std::string  
> const&  
using parse_environment3_type = basic_parsed_options<char> (*) (options_description  
const&, std::string  
const&)  
using parse_config_file_char1_type = basic_parsed_options<char> (*) (char const  
*, options_description  
const&, bool)  
using parse_config_file_char2_type = basic_parsed_options<char> (*) (std::basic_istream<char>  
struct  
std::char_traits<char>>&,  
options_description  
const&, bool)  
using parse_config_file_wchar1_type = basic_parsed_options<wchar_t> (*) (std::basic_istream<wchar_t>  
struct  
std::char_traits<wchar_t>>&,  
options_description  
const&, bool)  
using split_unix_type = std::vector<std::string> (*) (std::string const&, std::string  
const&, std::string const&, std::string const&)
```

Functions

```
force_linking_helper &force_linking()
```

```
struct force_linking_helper
```

Public Members

```
parse_environment1_type parse_environment1
parse_environment2_type parse_environment2
parse_environment3_type parse_environment3
parse_config_file_char1_type parse_config_file_char1
parse_config_file_char2_type parse_config_file_char2
parse_config_file_wchar1_type parse_config_file_wchar1
split_unix_type split_unix
```

```
#include <hpx/program_options/cmdline.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

```
namespace command_line_style
```

Enums

```
enum style_t
```

Various possible styles of options.

There are “long” options, which start with “–” and “short”, which start with either “–” or “/”. Both kinds can be allowed or disallowed, see `allow_long` and `allow_short`. The allowed character for short options is also configurable.

Option’s value can be specified in the same token as name (“`–foo=bar`”), or in the next token.

It’s possible to introduce long options by the same character as short options, see `allow_long_disguise`.

Finally, guessing (specifying only prefix of option) and case insensitive processing are supported.

Values:

```
allow_long = 1
```

Allow “`–long_name`” style.

```
allow_short = allow_long << 1
```

Allow “`-<single character`” style.

```
allow_dash_for_short = allow_short << 1
```

Allow “`-`” in short options.

```
allow_slash_for_short = allow_dash_for_short << 1
    Allow “/” in short options.

long_allow_adjacent = allow_slash_for_short << 1
    Allow option parameter in the same token for long option, like in
    --foo=10

long_allow_next = long_allow_adjacent << 1
    Allow option parameter in the next token for long options.

short_allow_adjacent = long_allow_next << 1
    Allow option parameter in the same token for short options.

short_allow_next = short_allow_adjacent << 1
    Allow option parameter in the next token for short options.

allow_sticky = short_allow_next << 1
    Allow to merge several short options together, so that “-s -k” become “-sk”. All of the options
    but last should accept no parameter. For example, if “-s” accept a parameter, then “k” will be
    taken as parameter, not another short option. Dos-style short options cannot be sticky.

allow_guessing = allow_sticky << 1
    Allow abbreviated spellings for long options, if they unambiguously identify long option. No
    long option name should be prefix of other long option name if guessing is in effect.

long_case_insensitive = allow_guessing << 1
    Ignore the difference in case for long options.

short_case_insensitive = long_case_insensitive << 1
    Ignore the difference in case for short options.

case_insensitive = (long_case_insensitive | short_case_insensitive)
    Ignore the difference in case for all options.

allow_long_disguise = short_case_insensitive << 1
    Allow long options with single option starting character, e.g. -f foo=10

unix_style = (allow_short | short_allow_adjacent | short_allow_next | allow_long | long_allow_adjacent | long_allow_next)
    The more-or-less traditional unix style.

default_style = unix_style
    The default style.
```

```
#include <hpx/program_options/eof_iterator.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

```
template<class Derived, class ValueType>
class eof_iterator : public util::iterator_facade<Derived, ValueType const, std::forward_iterator_tag>
{
    #include <eof_iterator.hpp> The ‘eof_iterator’ class is useful for constructing forward iterators in
    cases where iterator extract data from some source and it’s easy to detect ‘eof’ – i.e. the situation
    where there’s no data. One apparent example is reading lines from a file.
```

Implementing such iterators using ‘iterator_facade’ directly would require to create class with three core operation, a couple of constructors. When using ‘*eof_iterator*’, the derived class should define only one method to get new value, plus a couple of constructors.

The basic idea is that iterator has ‘eof’ bit. Two iterators are equal only if both have their ‘eof’ bits set. The ‘get’ method either obtains the new value or sets the ‘eof’ bit.

Specifically, derived class should define:

1. A default constructor, which creates iterator with ‘eof’ bit set. The constructor body should call ‘found_eof’ method defined here.
2. Some other constructor. It should initialize some ‘data pointer’ used in iterator operation and then call ‘get’.
3. The ‘get’ method. It should operate this way:
 - look at some ‘data pointer’ to see if new element is available; if not, it should call ‘found_eof’.
 - extract new element and store it at location returned by the ‘value’ method.
 - advance the data pointer.

Essentially, the ‘get’ method has the functionality of both ‘increment’ and ‘dereference’. It’s very good for the cases where data extraction implicitly moves data pointer, like for stream operation.

Public Functions

`eof_iterator()`

Protected Functions

`ValueType &value()`

Returns the reference which should be used by derived class to store the next value.

`void found_eof()`

Should be called by derived class to indicate that it can’t produce next element.

Private Functions

`void increment()`

`bool equal (const eof_iterator &other) const`

`const ValueType &dereference () const`

Private Members

`bool m_at_eof`

`ValueType m_value`

Friends

```
friend hpx::program_options::hpx::util::iterator_core_access

#include <hpx/program_options/option.hpp>

namespace hpx

namespace program_options
```

Typedefs

```
using option = basic_option<char>
using woption = basic_option<wchar_t>

template<class Char>
class basic_option
    #include <option.hpp> Option found in input source. Contains a key and a value. The key, in turn,
    can be a string (name of an option), or an integer (position in input source) – in case no name is
    specified. The latter is only possible for command line. The template parameter specifies the type of
    char used for storing the option's value.
```

Public Functions

```
basic_option()
basic_option(const std::string &xstring_key, const std::vector<std::string> &xvalue)
```

Public Members

`std::string string_key`

String key of this option. Intentionally independent of the template parameter.

`int position_key`

Position key of this option. All options without an explicit name are sequentially numbered starting from 0. If an option has explicit name, ‘position_key’ is equal to -1. It is possible that both position_key and string_key is specified, in case name is implicitly added.

`std::vector<std::basic_string<Char>> value`

Option's value

`std::vector<std::basic_string<Char>> original_tokens`

The original unchanged tokens this option was created from.

`bool unregistered`

True if option was not recognized. In that case, ‘string_key’ and ‘value’ are results of purely syntactic parsing of source. The original tokens can be recovered from the “original_tokens” member.

`bool case_insensitive`

True if string_key has to be handled case insensitive.

```
#include <hpx/program_options/value_semantic.hpp>
```

```
namespace hpx
```

```
namespace program_options
```

Functions

```
template<class T>
```

```
typed_value<T> *value()
```

Creates a typed_value<T> instance. This function is the primary method to create *value_semantic* instance for a specific type, which can later be passed to ‘*option_description*’ constructor. The second overload is used when it’s additionally desired to store the value of option into program variable.

```
template<class T>
```

```
typed_value<T> *value(T *v)
```

This is an overloaded member function, provided for convenience. It differs from the above function only in what argument(s) it accepts.

```
template<class T>
```

```
typed_value<T, wchar_t> *wvalue()
```

Creates a typed_value<T> instance. This function is the primary method to create *value_semantic* instance for a specific type, which can later be passed to ‘*option_description*’ constructor.

```
template<class T>
```

```
typed_value<T, wchar_t> *wvalue(T *v)
```

This is an overloaded member function, provided for convenience. It differs from the above function only in what argument(s) it accepts.

```
typed_value<bool> *bool_switch()
```

Works the same way as the ‘value<bool>’ function, but the created *value_semantic* won’t accept any explicit value. So, if the option is present on the command line, the value will be ‘true’.

```
typed_value<bool> *bool_switch(bool *v)
```

This is an overloaded member function, provided for convenience. It differs from the above function only in what argument(s) it accepts.

```
template<class T, class Char = char>
```

```
class typed_value : public hpx::program_options::value_semantic_codecvt_helper<Char>, public hpx::program
```

#include <value_semantic.hpp> Class which handles value of a specific type.

Public Functions

```
typed_value(T *store_to)
```

Ctor. The ‘store_to’ parameter tells where to store the value when it’s known. The parameter can be NULL.

```
typed_value *default_value(const T &v)
```

Specifies default value, which will be used if none is explicitly specified. The type ‘T’ should provide operator<< for ostream.

`typed_value *default_value(const T &v, const std::string &textual)`

Specifies default value, which will be used if none is explicitly specified. Unlike the above overload, the type ‘T’ need not provide operator<< for ostream, but textual representation of default value must be provided by the user.

`typed_value *implicit_value(const T &v)`

Specifies an implicit value, which will be used if the option is given, but without an adjacent value. Using this implies that an explicit value is optional,

`typed_value *value_name(const std::string &name)`

Specifies the name used to to the value in help message.

`typed_value *implicit_value(const T &v, const std::string &textual)`

Specifies an implicit value, which will be used if the option is given, but without an adjacent value. Using this implies that an explicit value is optional, but if given, must be strictly adjacent to the option, i.e.: ‘-ovalue’ or ‘option=value’. Giving ‘-o’ or ‘option’ will cause the implicit value to be applied. Unlike the above overload, the type ‘T’ need not provide operator<< for ostream, but textual representation of default value must be provided by the user.

`typed_value *notifier(std::function<void>) const T&`

>/Specifies a function to be called when the final value is determined.

`typed_value *composing()`

Specifies that the value is composing. See the ‘is_composing’ method for explanation.

`typed_value *multitoken()`

Specifies that the value can span multiple tokens.

`typed_value *zero_tokens()`

Specifies that no tokens may be provided as the value of this option, which means that only presence of the option is significant. For such option to be useful, either the ‘validate’ function should be specialized, or the ‘implicit_value’ method should be also used. In most cases, you can use the ‘bool_switch’ function instead of using this method.

`typed_value *required()`

Specifies that the value must occur.

`std::string name() const`

`bool is_composing() const`

`unsigned min_tokens() const`

`unsigned max_tokens() const`

`bool is_required() const`

`void xparse(hpx::util::any_nonser &value_store, const std::vector<std::basic_string<Char>> &new_tokens) const`

Creates an instance of the ‘validator’ class and calls its operator() to perform the actual conversion.

`virtual bool apply_default(hpx::util::any_nonser &value_store) const`

If default value was specified via previous call to ‘default_value’, stores that value into ‘value_store’. Returns true if default value was stored.

`void notify(const hpx::util::any_nonser &value_store) const`

If an address of variable to store value was specified when creating *this, stores the value there. Otherwise, does nothing.

`const std::type_info &value_type() const`

Private Members

```
T *m_store_to
std::string m_value_name
hpx::util::any_nonsr m_default_value
std::string m_default_value_as_text
hpx::util::any_nonsr m_implicit_value
std::string m_implicit_value_as_text
bool m_composing
bool m_implicit
bool m_multitoken
bool m_zero_tokens
bool m_required
std::function<void (const T&)> m_notifier
```

class typed_value_base

#include <value_semantic.hpp> Base class for all option that have a fixed type, and are willing to announce this type to the outside world. Any ‘value_semantics’ for which you want to find out the type can be dynamic_cast-ed to *typed_value_base*. If conversion succeeds, the ‘type’ method can be called.

Subclassed by *hpx::program_options::typed_value< T, Char >*

Public Functions

```
virtual const std::type_info &value_type() const = 0
```

```
virtual ~typed_value_base()
```

class untyped_value : public hpx::program_options::value_semantic_codecvt_helper<char>

#include <value_semantic.hpp> Class which specifies a simple handling of a value: the value will have string type and only one token is allowed.

Public Functions

```
untitled_value(bool zero_tokens = false)
```

```
std::string name() const
```

Returns the name of the option. The name is only meaningful for automatic help message.

```
unsigned min_tokens() const
```

The minimum number of tokens for this option that should be present on the command line.

```
unsigned max_tokens() const
```

The maximum number of tokens for this option that should be present on the command line.

```
bool is_composing() const
```

Returns true if values from different sources should be composed. Otherwise, value from the first source is used and values from other sources are discarded.

```
bool is_required() const
    Returns true if value must be given. Non-optional value

void xparse(hpx::util::any_nonser &value_store, const std::vector<std::string>
    &new_tokens) const
    If ‘value_store’ is already initialized, or new_tokens has more than one elements, throws. Otherwise, assigns the first string from ‘new_tokens’ to ‘value_store’, without any modifications.

bool apply_default(hpx::util::any_nonser&) const
    Does nothing.

void notify(const hpx::util::any_nonser&) const
    Does nothing.
```

Private Members

```
bool m_zero_tokens

class value_semantic
#include <value_semantic.hpp> Class which specifies how the option’s value is to be parsed and converted into C++ types.

Subclassed by hpx::program_options::value_semantic_codecvt_helper<char>, hpx::program_options::value_semantic_codecvt_helper<wchar_t>
```

Public Functions

```
virtual std::string name() const = 0
    Returns the name of the option. The name is only meaningful for automatic help message.

virtual unsigned min_tokens() const = 0
    The minimum number of tokens for this option that should be present on the command line.

virtual unsigned max_tokens() const = 0
    The maximum number of tokens for this option that should be present on the command line.

virtual bool is_composing() const = 0
    Returns true if values from different sources should be composed. Otherwise, value from the first source is used and values from other sources are discarded.

virtual bool is_required() const = 0
    Returns true if value must be given. Non-optional value

virtual void parse(hpx::util::any_nonser &value_store, const std::vector<std::string>
    &new_tokens, bool utf8) const = 0
    Parses a group of tokens that specify a value of option. Stores the result in ‘value_store’, using whatever representation is desired. May be called several times if value of the same option is specified more than once.

virtual bool apply_default(hpx::util::any_nonser &value_store) const = 0
    Called to assign default value to ‘value_store’. Returns true if default value is assigned, and false if no default value exists.

virtual void notify(const hpx::util::any_nonser &value_store) const = 0
    Called when final value of an option is determined.
```

```

virtual ~value_semantic()

template<class Char>
class value_semantic_codecvt_helper
    #include <value_semantic.hpp> Helper class which perform necessary character conversions in the
    ‘parse’ method and forwards the data further.

    Subclassed by hpx::program\_options::typed\_value< T, Char >

template<>
class value_semantic_codecvt_helper<char> : public hpx::program_options::value_semantic
    #include <value_semantic.hpp> Helper conversion class for values that accept ascii strings as input.
    Overrides the ‘parse’ method and defines new ‘xparse’ method taking std::string. Depending
    on whether input to parse is ascii or UTF8, will pass it to xparse unmodified, or with UTF8->ascii
    conversion.

    Subclassed by hpx::program\_options::untyped\_value

```

Protected Functions

```
virtual void xparse(hpx::util::any_nonser &value_store, const std::vector<std::string>
&new_tokens) const = 0
```

Private Functions

```
void parse(hpx::util::any_nonser &value_store, const std::vector<std::string> &new_tokens,
bool utf8) const
```

Parses a group of tokens that specify a value of option. Stores the result in ‘value_store’, using whatever representation is desired. May be called several times if value of the same option is specified more than once.

```

template<>
class value_semantic_codecvt_helper<wchar_t> : public hpx::program_options::value_semantic
    #include <value_semantic.hpp> Helper conversion class for values that accept ascii strings as input.
    Overrides the ‘parse’ method and defines new ‘xparse’ method taking std::wstring. Depending on
    whether input to parse is ascii or UTF8, will recode input to Unicode, or pass it unmodified.

```

Protected Functions

```
virtual void xparse(hpx::util::any_nonser &value_store, const std::vector<std::wstring>
&new_tokens) const = 0
```

Private Functions

```
void parse(hpx::util::any_nonser &value_store, const std::vector<std::string> &new_tokens,
bool utf8) const
```

Parses a group of tokens that specify a value of option. Stores the result in ‘value_store’, using whatever representation is desired. May be called several times if value of the same option is specified more than once.

basic_execution

```
#include <compatibility/hpx/util/register_locks.hpp>

#include <hpx/basic_execution/this_thread.hpp>

namespace hpx

namespace basic_execution

namespace this_thread
```

Functions

```
hpx::basic_execution::agent_ref agent ()

void yield (char const *desc = "hpx::basic_execution::this_thread::yield")

void yield_k (std::size_t k, char const *desc = "hpx::basic_execution::this_thread::yield_k")

void suspend (char const *desc = "hpx::basic_execution::this_thread::suspend")

template<typename Rep, typename Period>
void sleep_for (std::chrono::duration<Rep, Period> const &sleep_duration, char const
               *desc = "hpx::basic_execution::this_thread::sleep_for")

template<class Clock, class Duration>
void sleep_until (std::chrono::time_point<Clock, Duration> const &sleep_time, char const
                  *desc = "hpx::basic_execution::this_thread::sleep_for")

struct reset_agent
```

Public Functions

```
reset_agent (detail::agent_storage *, agent_base &impl)

reset_agent (agent_base &impl)

~reset_agent ()
```

Public Members

```
detail::agent_storage *storage_

agent_base *old_
```

```
#include <hpx/basic_execution/agent_ref.hpp>

namespace hpx
```

```
namespace basic_execution
```

```
class agent_ref
```

Public Functions

```
constexpr agent_ref()
constexpr agent_ref(agent_base *impl)
constexpr agent_ref(agent_ref const&)
constexpr agent_ref &operator=(agent_ref const&)
constexpr agent_ref (agent_ref&&)
constexpr agent_ref &operator=(agent_ref&&)
constexpr operator bool() const
void reset (agent_base *impl = nullptr)
void yield (char const *desc = "hpx::basic_execution::agent_ref::yield")
void yield_k (std::size_t k, char const *desc = "hpx::basic_execution::agent_ref::yield_k")
void suspend (char const *desc = "hpx::basic_execution::agent_ref::suspend")
void resume (char const *desc = "hpx::basic_execution::agent_ref::resume")
void abort (char const *desc = "hpx::basic_execution::agent_ref::abort")
template<typename Rep, typename Period>
void sleep_for (std::chrono::duration<Rep, Period> const &sleep_duration, char const
               *desc = "hpx::basic_execution::agent_ref::sleep_for")
template<typename Clock, typename Duration>
void sleep_until (std::chrono::time_point<Clock, Duration> const &sleep_time, char const
                  *desc = "hpx::basic_execution::agent_ref::sleep_until")
agent_base &ref()
```

Private Functions

```
void sleep_for (hpx::util::steady_duration const &sleep_duration, char const *desc)
void sleep_until (hpx::util::steady_time_point const &sleep_time, char const *desc)
```

Private Members

```
agent_base *impl_
```

Friends

```
constexpr bool operator==(agent_ref const &lhs, agent_ref const &rhs)  
constexpr bool operator!=(agent_ref const &lhs, agent_ref const &rhs)  
std::ostream &operator<< (std::ostream&, agent_ref const&)
```

```
#include <hpx/basic_execution/resource_base.hpp>
```

```
namespace hpx
```

```
namespace basic_execution
```

```
struct resource_base
```

```
#include <resource_base.hpp> TODO: implement, this is currently just a dummy.
```

Public Functions

```
virtual ~resource_base()
```

```
#include <hpx/basic_execution/register_locks.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
bool register_lock(void const *, util::register_lock_data * = nullptr)  
bool unregister_lock(void const *)  
void verify_no_locks()  
void force_error_on_lock()  
void enable_lock_detection()  
void ignore_lock(void const *)  
void reset_ignored(void const *)  
void ignore_all_locks()  
void reset_ignored_all()
```

```
struct ignore_all_while_checking
```

Public Functions

```
ignore_all_while_checking()
```

```
template<typename Lock, typename Enable>
struct ignore_while_checking
```

Public Functions

```
ignore_while_checking(void const *)
```

```
#include <hpx/basic_execution/context_base.hpp>
```

```
namespace hpx
```

```
namespace basic_execution
```

```
struct context_base
```

Public Functions

```
virtual ~context_base()
```

```
virtual resource_base const &resource() const = 0
```

```
#include <hpx/basic_execution/agent_base.hpp>
```

```
namespace hpx
```

```
namespace basic_execution
```

```
struct agent_base
```

Public Functions

```
virtual ~agent_base()
```

```
virtual std::string description() const = 0
```

```
virtual context_base const &context() const = 0
```

```
virtual void yield(char const *desc) = 0
```

```
virtual void yield_k(std::size_t k, char const *desc) = 0
```

```
virtual void suspend(char const *desc) = 0
virtual void resume(char const *desc) = 0
virtual void abort(char const *desc) = 0
virtual void sleep_for(hpx::util::steady_duration const &sleep_duration, char const
                      *desc) = 0
virtual void sleep_until(hpx::util::steady_time_point const &sleep_time, char const
                        *desc) = 0
```

memory

```
#include <hpx/memory/intrusive_ptr.hpp>

template<typename T>
struct hash<hpx::memory::intrusive_ptr<T>>
```

Public Types

```
template<>
using result_type = std::size_t
```

Public Functions

```
result_type operator()(hpx::memory::intrusive_ptr<T> const &p) const
namespace hpx

namespace memory
```

Functions

```
template<typename T, typename U>
bool operator==(intrusive_ptr<T> const &a, intrusive_ptr<U> const &b)

template<typename T, typename U>
bool operator!=(intrusive_ptr<T> const &a, intrusive_ptr<U> const &b)

template<typename T, typename U>
bool operator==(intrusive_ptr<T> const &a, U *b)

template<typename T, typename U>
bool operator!=(intrusive_ptr<T> const &a, U *b)

template<typename T, typename U>
bool operator==(T *a, intrusive_ptr<U> const &b)

template<typename T, typename U>
bool operator!=(T *a, intrusive_ptr<U> const &b)

template<typename T>
```

```

bool operator== (intrusive_ptr<T> const &p, std::nullptr_t)

template<typename T>
bool operator== (std::nullptr_t, intrusive_ptr<T> const &p)

template<typename T>
bool operator!= (intrusive_ptr<T> const &p, std::nullptr_t)

template<typename T>
bool operator!= (std::nullptr_t, intrusive_ptr<T> const &p)

template<typename T>
bool operator< (intrusive_ptr<T> const &a, intrusive_ptr<T> const &b)

template<typename T>
void swap (intrusive_ptr<T> &lhs, intrusive_ptr<T> &rhs)

template<typename T>
T *get_pointer (intrusive_ptr<T> const &p)

template<typename T, typename U>
intrusive_ptr<T> static_pointer_cast (intrusive_ptr<U> const &p)

template<typename T, typename U>
intrusive_ptr<T> const_pointer_cast (intrusive_ptr<U> const &p)

template<typename T, typename U>
intrusive_ptr<T> dynamic_pointer_cast (intrusive_ptr<U> const &p)

template<typename T, typename U>
intrusive_ptr<T> static_pointer_cast (intrusive_ptr<U> &&p)

template<typename T, typename U>
intrusive_ptr<T> const_pointer_cast (intrusive_ptr<U> &&p)

template<typename T, typename U>
intrusive_ptr<T> dynamic_pointer_cast (intrusive_ptr<U> &&p)

template<typename Y>
std::ostream &operator<< (std::ostream &os, intrusive_ptr<Y> const &p)

template<typename T>
class intrusive_ptr

```

Public Types

```

template<>
using element_type = T

```

Public Functions

```
constexpr intrusive_ptr()

intrusive_ptr(T *p, bool add_ref = true)

template<typename U, typename Enable = typename std::enable_if<memory::detail::sp_convertible<U, T>::value>
intrusive_ptr(intrusive_ptr<U> const &rhs)

intrusive_ptr(intrusive_ptr const &rhs)

~intrusive_ptr()

template<typename U>
intrusive_ptr &operator=(intrusive_ptr<U> const &rhs)

intrusive_ptr(intrusive_ptr &&rhs)

intrusive_ptr &operator=(intrusive_ptr &&rhs)

template<typename U, typename Enable = typename std::enable_if<memory::detail::sp_convertible<U, T>::value>
intrusive_ptr(intrusive_ptr<U> &&rhs)

template<typename U>
intrusive_ptr &operator=(intrusive_ptr<U> &&rhs)

intrusive_ptr &operator=(intrusive_ptr const &rhs)

intrusive_ptr &operator=(T *rhs)

void reset()

void reset(T *rhs)

void reset(T *rhs, bool add_ref)

T *get() const

T *detach()

T &operator*() const

T *operator->() const

operator bool() const

void swap(intrusive_ptr &rhs)
```

Private Types

```
template<>
using this_type = intrusive_ptr
```

Private Members

```
T *px
```

Friends

```
friend hpx::memory::intrusive_ptr
namespace std
```

```
template<typename T>
struct hash<hpx::memory::intrusive_ptr<T>>
```

Public Types

```
template<>
using result_type = std::size_t
```

Public Functions

```
result_type operator() (hpx::memory::intrusive_ptr<T> const &p) const
```

```
#include <hpx/memory/serialization/intrusive_ptr.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T>
void load(input_archive &ar, hpx::intrusive_ptr<T> &ptr, unsigned)
```

```
template<typename T>
void save(output_archive &ar, hpx::intrusive_ptr<T> const &ptr, unsigned)
```

```
hpx::serialization::Hpx_serialization_Split_Free_Template((template< typename T >),
```

```
config
```

```
#include <hpx/config.hpp>
```

Defines

HPX_INITIAL_IP_PORT

This is the default ip/port number used by the parcel subsystem.

HPX_CONNECTING_IP_PORT

HPX_INITIAL_IP_ADDRESS

HPX_RUNTIME_INSTANCE_LIMIT

This defines the maximum number of possible runtime instances in one executable

HPX_PARCEL_BOOTSTRAP

This defines the type of the parcelport to be used during application bootstrap. This value can be changed at runtime by the configuration parameter:

`hpx.parcel.bootstrap = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_BOOTSTRAP`).

HPX_PARCEL_MAX_CONNECTIONS

This defines the number of outgoing (parcel-) connections kept alive (to all other localities). This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.max_connections = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_MAX_CONNECTIONS`).

HPX_PARCEL_IPC_DATA_BUFFER_CACHE_SIZE

This defines the number of outgoing ipc (parcel-) connections kept alive (to each of the other localities on the same node). This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.ipc.data_buffer_cache_size = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_IPC_DATA_BUFFER_CACHE_SIZE`).

HPX_PARCEL_MPI_MAX_REQUESTS

This defines the number of MPI requests in flight. This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.mpi.max_requests = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_MPI_MAX_REQUESTS`).

HPX_PARCEL_MAX_CONNECTIONS_PER_LOCALITY

This defines the number of outgoing (parcel-) connections kept alive (to each of the other localities). This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.max_connections_per_locality = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_MAX_CONNECTIONS_PER_LOCALITY`).

HPX_PARCEL_MAX_MESSAGE_SIZE

This defines the maximally allowed message size for messages transferred between localities. This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.max_message_size = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_MAX_MESSAGE_SIZE`).

HPX_PARCEL_MAX_OUTBOUND_MESSAGE_SIZE

This defines the maximally allowed outbound message size for coalescing messages transferred between localities. This value can be changed at runtime by setting the configuration parameter:

`hpx.parcel.max_outbound_message_size = ...`

(or by setting the corresponding environment variable `HPX_PARCEL_MAX_OUTBOUND_MESSAGE_SIZE`).

HPX_PARCEL_SERIALIZATION_OVERHEAD

HPX_AGAS_LOCAL_CACHE_SIZE

This defines the number of AGAS address translations kept in the local cache. This is just the initial size which

may be adjusted depending on the load of the system (not implemented yet), etc. It must be a minimum of 3 for AGAS v3 bootstrapping.

This value can be changes at runtime by setting the configuration parameter:

`hpx.agas.local_cache_size = ...`

(or by setting the corresponding environment variable `HPX_AGAS_LOCAL_CACHE_SIZE`)

HPX_INITIAL_AGAS_MAX_PENDING_REFCNT_REQUESTS

HPX_GLOBALCREDIT_INITIAL

This defines the initial global reference count associated with any created object.

HPX_NUM_IO_POOL_SIZE

This defines the default number of OS-threads created for the different internal thread pools

HPX_NUM_PARCEL_POOL_SIZE

HPX_NUM_TIMER_POOL_SIZE

HPX_SPINLOCK_DEADLOCK_DETECTION_LIMIT

By default, enable minimal thread deadlock detection in debug builds only.

HPX_COROUTINE_NUM_HEAPS

This defines the default number of coroutine heaps.

HPX_HAVE_THREAD_BACKTRACE_DEPTH

By default, enable storing the thread phase in debug builds only.

By default, enable storing the parent thread information in debug builds only. By default, enable storing the thread description in debug builds only. By default, enable storing the target address of the data the thread is accessing in debug builds only. By default we do not maintain stack back-traces on suspension. This is a pure debugging aid to be able to see in the debugger where a suspended thread got stuck. By default we capture only 5 levels of stack back trace on suspension

HPX_MAX_NETWORK_RETRIES

HPX_NETWORK_RETRIES_SLEEP

HPX_INI_PATH_DELIMITER

HPX_PATH_DELIMITERS

HPX_SHARED_LIB_EXTENSION

HPX_EXECUTABLE_EXTENSION

HPX_MAKE_DLL_STRING(n)

HPX_MANGLE_NAME(n)

HPX_MANGLE_STRING(n)

HPX_COMPONENT_NAME

HPX_COMPONENT_STRING

HPX_PLUGIN_COMPONENT_PREFIX

HPX_PLUGIN_NAME

HPX_PLUGIN_STRING

HPX_PLUGIN_PLUGIN_PREFIX

HPX_APPLICATION_STRING

```
HPX_IDLE_LOOP_COUNT_MAX
HPX_BUSY_LOOP_COUNT_MAX
HPX_THREAD_QUEUE_MAX_THREAD_COUNT
HPX_THREAD_QUEUE_MIN_TASKS_TO_STEAL_PENDING
HPX_THREAD_QUEUE_MIN_TASKS_TO_STEAL_STAGED
HPX_THREAD_QUEUE_MIN_ADD_NEW_COUNT
HPX_THREAD_QUEUE_MAX_ADD_NEW_COUNT
HPX_THREAD_QUEUE_MIN_DELETE_COUNT
HPX_THREAD_QUEUE_MAX_DELETE_COUNT
HPX_THREAD_QUEUE_MAX_TERMINATED_THREADS
HPX_IDLE_BACKOFF_TIME_MAX
HPX_WRAPPER_HEAP_STEP
HPX_INITIAL_GID_RANGE
HPX_CONTINUATION_MAX_RECURSION_DEPTH
HPX_AGAS_BOOTSTRAP_PREFIX
HPX_AGAS_NS_MSB
HPX_AGAS_PRIMARY_NS_MSB
HPX_AGAS_PRIMARY_NS_LSB
HPX_AGAS_COMPONENT_NS_MSB
HPX_AGAS_COMPONENT_NS_LSB
HPX_AGAS_SYMBOL_NS_MSB
HPX_AGAS_SYMBOL_NS_LSB
HPX_AGAS_LOCALITY_NS_MSB
HPX_AGAS_LOCALITY_NS_LSB

#include <hpx/config/compiler_specific.hpp>
```

Defines

HPX_GCC_VERSION

Returns the GCC version HPX is compiled with. Only set if compiled with GCC.

HPX_CLANG_VERSION

Returns the Clang version HPX is compiled with. Only set if compiled with Clang.

HPX_INTEL_VERSION

Returns the Intel Compiler version HPX is compiled with. Only set if compiled with the Intel Compiler.

HPX_MSVC

This macro is set if the compilation is with MSVC.

HPX_MINGW

This macro is set if the compilation is with Mingw.

HPX_WINDOWS

This macro is set if the compilation is for Windows.

HPX_NATIVE_MIC

This macro is set if the compilation is for Intel Knights Landing.

```
#include <hpx/config/constexpr.hpp>
```

Defines**HPX_CONSTEXPR**

This macro evaluates to `constexpr` if the compiler supports it.

This macro is deprecated. It is always replaced with the `constexpr` keyword. Prefer using `constexpr` directly instead.

HPX_CONSTEXPR_OR_CONST

This macro evaluates to `constexpr` if the compiler supports it, `const` otherwise.

This macro is deprecated. It is always replaced with the `constexpr` keyword. Prefer using `constexpr` directly instead.

HPX_STATIC_CONSTEXPR

This macro evaluates to `static constexpr` if the compiler supports it, `static const` otherwise.

This macro is deprecated. It is always replaced with the `static constexpr` keyword. Prefer using `static constexpr` directly instead.

```
#include <hpx/config/debug.hpp>
```

Defines**HPX_DEBUG**

Defined if HPX is compiled in debug mode.

HPX_BUILD_TYPE

Evaluates to `debug` if compiled in debug mode, `release` otherwise.

```
#include <hpx/config/asio.hpp>
```

```
#include <hpx/config/autolink.hpp>
```

```
#include <hpx/config/threads_stack.hpp>
```

Defines**HPX_THREADS_STACK_OVERHEAD****HPX_SMALL_STACK_SIZE****HPX_MEDIUM_STACK_SIZE****HPX_LARGE_STACK_SIZE****HPX_HUGE_STACK_SIZE**

```
#include <hpx/config/emulate_deleted.hpp>
```

Defines

HPX_NON_COPYABLE (cls)

Marks a class as non-copyable and non-movable.

```
#include <hpx/config/branch_hints.hpp>
```

Defines

HPX_LIKELY (expr)

Hint at the compiler that `expr` is likely to be true.

HPX_UNLIKELY (expr)

Hint at the compiler that `expr` is likely to be false.

```
#include <hpx/config/lambda_capture.hpp>
```

Defines

HPX_CAPTURE_FORWARD (var)

Evaluates to `var = std::forward<decltype(var)>(var)` if the compiler supports C++14 Lambdas.
Defaults to `var`.

This macro is deprecated. Prefer using `var = std::forward<decltype(var)>(var)` directly instead.

HPX_CAPTURE_MOVE (var)

Evaluates to `var = std::move(var)` if the compiler supports C++14 Lambdas. Defaults to `var`.

This macro is deprecated. Prefer using `var = std::move(var)` directly instead.

```
#include <hpx/config/warnings_prefix.hpp>
```

```
#include <hpx/config/weak_symbol.hpp>
```

Defines

HPX_WEAK_SYMBOL

```
#include <hpx/config/force_linking.hpp>
```

```
namespace hpx
```

```
    namespace config
```

Functions

```
force_linking_helper &force_linking()  
struct force_linking_helper
```

Public Members

```
const char *const hpx_version  
const char *const boost_version
```

```
#include <hpx/config/export_definitions.hpp>
```

Defines

HPX_EXPORT

Marks a class or function to be exported from HPX or imported if it is consumed.

```
#include <hpx/config/attributes.hpp>
```

Defines

HPX_NOINLINE

Function attribute to tell compiler not to inline the function.

HPX_NORETURN

Function attribute to tell compiler that the function does not return.

HPX_DEPRECATED (x)

Marks an entity as deprecated. The argument x specifies a custom message that is included in the compiler warning. For more details see <>__.

HPX_FALLTHROUGH

Indicates that the fall through from the previous case label is intentional and should not be diagnosed by a compiler that warns on fallthrough. For more details see <>__.

```
#include <hpx/config/warnings_suffix.hpp>
```

```
#include <hpx/config/forceinline.hpp>
```

Defines

HPX_FORCEINLINE

Marks a function to be forced inline.

```
#include <hpx/config/compiler_native_tls.hpp>
```

Defines

HPX_NATIVE_TLS

This macro is replaced with the compiler specific keyword attribute to mark a variable as thread local. For more details see `<___.h>`.

This macro is deprecated. It is always replaced with the `thread_local` keyword. Prefer using `thread_local` directly instead.

```
#include <hpx/config/manual_profiling.hpp>
```

Defines

HPX_SUPER_PURE

HPX_PURE

HPX_HOT

HPX_COLD

```
#include <hpx/config/compiler_fence.hpp>
```

Defines

HPX_COMPILER_FENCE

Generates assembly that serves as a fence to the compiler CPU to disable optimization. Usually implemented in the form of a memory barrier.

HPX_SMT_PAUSE

Generates assembly that executes a “pause” instruction. Useful in spinning loops.

timing

```
#include <compatibility/hpx/util/steady_clock.hpp>

#include <compatibility/hpx/util/scoped_timer.hpp>

#include <compatibility/hpx/util/high_resolution_clock.hpp>

#include <compatibility/hpx/util/high_resolution_timer.hpp>

#include <hpx/timing/steady_clock.hpp>

namespace hpx

namespace util
```

```
class steady_duration
```

Public Functions

```
steady_duration(value_type const &rel_time)
template<typename Rep, typename Period>
steady_duration(std::chrono::duration<Rep, Period> const &rel_time)
value_type const &value() const
steady_clock::time_point from_now() const
```

Private Types

```
typedef steady_clock::duration value_type
```

Private Members

```
value_type _rel_time
class steady_time_point
```

Public Functions

```
steady_time_point(value_type const &abs_time)
template<typename Clock, typename Duration>
steady_time_point(std::chrono::time_point<Clock, Duration> const &abs_time)
value_type const &value() const
```

Private Types

```
typedef steady_clock::time_point value_type
```

Private Members

```
value_type _abs_time
```

```
#include <hpx/timing/scoped_timer.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
template<typename T>
struct scoped_timer
```

Public Functions

```
scoped_timer (T &t, bool enabled = true)
scoped_timer (scoped_timer const&)
scoped_timer (scoped_timer &&rhs)
~scoped_timer()

scoped_timer &operator= (scoped_timer const &rhs)
scoped_timer &operator= (scoped_timer &&rhs)
bool enabled() const
```

Private Members

```
std::uint64_t started_at_
T *t_
```

```
#include <hpx/timing/high_resolution_clock.hpp>

namespace hpx

namespace util

struct high_resolution_clock
```

Public Static Functions

```
static std::uint64_t now()
static std::uint64_t() hpx::util::high_resolution_clock::min()
static std::uint64_t() hpx::util::high_resolution_clock::max()
```

```
#include <hpx/timing/high_resolution_timer.hpp>

namespace hpx

namespace util

class high_resolution_timer
```

Public Functions

```
high_resolution_timer()  
high_resolution_timer(double t)  
void restart()  
double elapsed() const  
std::int64_t elapsed_microseconds() const  
std::int64_t elapsed_nanoseconds() const  
double elapsed_max() const  
double elapsed_min() const
```

Public Static Functions

```
static double now()
```

Protected Static Functions

```
static std::uint64_t take_time_stamp()
```

Private Members

```
std::uint64_t start_time_
```

batch_environments

```
#include <compatibility/hpx/util/batch_environment.hpp>  
  
#include <compatibility/hpx/util/batch_environments/pbs_environment.hpp>  
  
#include <compatibility/hpx/util/batch_environments/slurm_environment.hpp>  
  
#include <compatibility/hpx/util/batch_environments/alps_environment.hpp>  
  
#include <hpx/batch_environments/batch_environment.hpp>  
  
namespace hpx  
  
namespace util  
  
struct batch_environment
```

Public Types

```
typedef std::map<boost::asio::ip::tcp::endpoint, std::pair<std::string, std::size_t>> node_map_type
```

Public Functions

```
batch_environment (std::vector<std::string> &nodelist, bool have_mpi = false, bool debug  
= false, bool enable = true)  
std::string init_from_nodelist (std::vector<std::string> const &nodes, std::string  
const &agas_host)  
std::size_t retrieve_number_of_threads () const  
std::size_t retrieve_number_of_localities () const  
std::size_t retrieve_node_number () const  
std::string host_name () const  
std::string host_name (std::string const &def_hpx_name) const  
std::string agas_host_name (std::string const &def_agas) const  
std::size_t agas_node () const  
bool found_batch_environment () const  
std::string get_batch_name () const
```

Public Members

```
std::string agas_node_  
std::size_t agas_node_num_  
std::size_t node_num_  
std::size_t num_threads_  
node_map_type nodes_  
std::size_t num_localities_  
std::string batch_name_  
bool debug_
```

```
#include <hpx/batch_environments/pbs_environment.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace batch_environments
```

```
struct pbs_environment
```

Public Functions

```
pbs_environment (std::vector<std::string> &nodelist, bool have_mpi, bool debug)
    bool valid() const
    std::size_t node_num() const
    std::size_t num_threads() const
    std::size_t num_localities() const
```

Private Functions

```
void read_nodefile (std::vector<std::string> &nodelist, bool have_mpi, bool debug)
void read_nodelist (std::vector<std::string> &nodelist, bool debug)
```

Private Members

```
std::size_t node_num_
std::size_t num_localities_
std::size_t num_threads_
bool valid_
```

```
#include <hpx/batch_environments/slurm_environment.hpp>
namespace hpx
```

```
namespace util
namespace batch_environments
```

```
struct slurm_environment
```

Public Functions

```
slurm_environment (std::vector<std::string> &nodelist, bool debug)
    bool valid() const
    std::size_t node_num() const
    std::size_t num_threads() const
    std::size_t num_localities() const
```

Private Functions

```
void retrieve_number_of_localities(bool debug)
void retrieve_number_of_tasks(bool debug)
void retrieve_nodelist(std::vector<std::string> &nodelist, bool debug)
void retrieve_number_of_threads()
```

Private Members

```
std::size_t node_num_
std::size_t num_threads_
std::size_t num_tasks_
std::size_t num_localities_
bool valid_
```

```
#include <hpx/batch_environments/alps_environment.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace batch_environments
```

```
struct alps_environment
```

Public Functions

```
alps_environment(std::vector<std::string> &nodelist, bool debug)
bool valid() const
std::size_t node_num() const
std::size_t num_threads() const
std::size_t num_localities() const
```

Private Members

```
std::size_t node_num_
std::size_t num_threads_
std::size_t num_localities_
bool valid_
```

cache

```
#include <compatibility/hpx/util/cache/local_cache.hpp>

#include <compatibility/hpx/util/cache/lru_cache.hpp>

#include <compatibility/hpx/util/cache/entries/fifo_entry.hpp>

#include <compatibility/hpx/util/cache/entries/size_entry.hpp>

#include <compatibility/hpx/util/cache/entries/lfu_entry.hpp>

#include <compatibility/hpx/util/cache/entries/lru_entry.hpp>

#include <compatibility/hpx/util/cache/entries/entry.hpp>

#include <compatibility/hpx/util/cache/statistics/local_full_statistics.hpp>

#include <compatibility/hpx/util/cache/statistics/local_statistics.hpp>

#include <compatibility/hpx/util/cache/statistics/no_statistics.hpp>

#include <compatibility/hpx/util/cache/policies/always.hpp>

#include <hpx/cache/local_cache.hpp>

namespace hpx

namespace util

namespace cache
```

template<typename **Key**, typename **Entry**, typename **UpdatePolicy** = *std::less<Entry>*, typename **InsertPolicy**

class local_cache

#include <hpx/cache/local_cache.hpp> The `local_cache` implements the basic functionality needed for a local (non-distributed) cache.

Template Parameters

- **Key:** The type of the keys to use to identify the entries stored in the cache

- **Entry**: The type of the items to be held in the cache, must model the CacheEntry concept
- **UpdatePolicy**: A (optional) type specifying a (binary) function object used to sort the cache entries based on their ‘age’. The ‘oldest’ entries (according to this sorting criteria) will be discarded first if the maximum capacity of the cache is reached. The default is std::less<Entry>. The function object will be invoked using 2 entry instances of the type *Entry*. This type must model the UpdatePolicy model.
- **InsertPolicy**: A (optional) type specifying a (unary) function object used to allow global decisions whether a particular entry should be added to the cache or not. The default is policies::always, imposing no global insert related criteria on the cache. The function object will be invoked using the entry instance to be inserted into the cache. This type must model the InsertPolicy model.
- **CacheStorage**: A (optional) container type used to store the cache items. The container must be an associative and STL compatible container. The default is a std::map<Key, Entry>.
- **Statistics**: A (optional) type allowing to collect some basic statistics about the operation of the cache instance. The type must conform to the CacheStatistics concept. The default value is the type statistics::no_statistics which does not collect any numbers, but provides empty stubs allowing the code to compile.

Public Types

```
typedef Key key_type
typedef Entry entry_type
typedef UpdatePolicy update_policy_type
typedef InsertPolicy insert_policy_type
typedef CacheStorage storage_type
typedef Statistics statistics_type
typedef entry_type::value_type value_type
typedef storage_type::size_type size_type
typedef storage_type::value_type storage_value_type
```

Public Functions

```
local_cache(size_type max_size = 0, update_policy_type const &up = update_policy_type(), insert_policy_type const &ip = insert_policy_type())
Construct an instance of a local_cache.
```

Parameters

- **max_size**: [in] The maximal size this cache is allowed to reach any time. The default is zero (no size limitation). The unit of this value is usually determined by the unit of the values returned by the entry’s *get_size* function.
- **up**: [in] An instance of the *UpdatePolicy* to use for this cache. The default is to use a default constructed instance of the type as defined by the *UpdatePolicy* template parameter.
- **ip**: [in] An instance of the *InsertPolicy* to use for this cache. The default is to use a default constructed instance of the type as defined by the *InsertPolicy* template parameter.

```
local_cache(local_cache &&other)
```

```
size_type size() const
```

Return current size of the cache.

Return The current size of this cache instance.

`size_type capacity() const`

Access the maximum size the cache is allowed to grow to.

Note The unit of this value is usually determined by the unit of the return values of the entry's function `entry::get_size`.

Return The maximum size this cache instance is currently allowed to reach. If this number is zero the cache has no limitation with regard to a maximum size.

`bool reserve(size_type max_size)`

Change the maximum size this cache can grow to.

Return This function returns *true* if successful. It returns *false* if the new `max_size` is smaller than the current limit and the cache could not be shrunk to the new maximum size.

Parameters

- `max_size`: [in] The new maximum size this cache will be allowed to grow to.

`bool holds_key(key_type const &k) const`

Check whether the cache currently holds an entry identified by the given key.

Note This function does not call the entry's function `entry::touch`. It just checks if the cache contains an entry corresponding to the given key.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `k`: [in] The key for the entry which should be looked up in the cache.

`bool get_entry(key_type const &k, key_type &realkey, entry_type &val)`

Get a specific entry identified by the given key.

Note The function will call the entry's `entry::touch` function if the value corresponding to the provided key is found in the cache.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `k`: [in] The key for the entry which should be retrieved from the cache.
- `val`: [out] If the entry indexed by the key is found in the cache this value on successful return will be a copy of the corresponding entry.

`bool get_entry(key_type const &k, entry_type &val)`

Get a specific entry identified by the given key.

Note The function will call the entry's `entry::touch` function if the value corresponding to the provided key is found in the cache.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `k`: [in] The key for the entry which should be retrieved from the cache.
- `val`: [out] If the entry indexed by the key is found in the cache this value on successful return will be a copy of the corresponding entry.

`bool get_entry(key_type const &k, value_type &val)`

Get a specific entry identified by the given key.

Note The function will call the entry's `entry::touch` function if the value corresponding to the provided is found in the cache.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `k`: [in] The key for the entry which should be retrieved from the cache
- `val`: [out] If the entry indexed by the key is found in the cache this value on successful return will be a copy of the corresponding value.

`bool insert (key_type const &k, value_type const &val)`

Insert a new element into this cache.

Note This function invokes both, the insert policy as provided to the constructor and the function `entry::insert` of the newly constructed entry instance. If either of these functions returns false the key/value pair doesn't get inserted into the cache and the `insert` function will return `false`. Other reasons for this function to fail (return `false`) are a) the key/value pair is already held in the cache or b) inserting the new value into the cache maxed out its capacity and it was not possible to free any of the existing entries.

Return This function returns `true` if the entry has been successfully added to the cache, otherwise it returns `false`.

Parameters

- `k`: [in] The key for the entry which should be added to the cache.
- `value`: [in] The value which should be added to the cache.

`bool insert (key_type const &k, entry_type &e)`

Insert a new entry into this cache.

Note This function invokes both, the insert policy as provided to the constructor and the function `entry::insert` of the provided entry instance. If either of these functions returns false the key/value pair doesn't get inserted into the cache and the `insert` function will return `false`. Other reasons for this function to fail (return `false`) are a) the key/value pair is already held in the cache or b) inserting the new value into the cache maxed out its capacity and it was not possible to free any of the existing entries.

Return This function returns `true` if the entry has been successfully added to the cache, otherwise it returns `false`.

Parameters

- `k`: [in] The key for the entry which should be added to the cache.
- `value`: [in] The entry which should be added to the cache.

`bool update (key_type const &k, value_type const &val)`

Update an existing element in this cache.

Note The function will call the entry's `entry::touch` function if the indexed value is found in the cache.

Note The difference to the other overload of the `insert` function is that this overload replaces the cached value only, while the other overload replaces the whole cache entry, updating the cache entry properties.

Return This function returns `true` if the entry has been successfully updated, otherwise it returns `false`. If the entry currently is not held by the cache it is added and the return value reflects the outcome of the corresponding insert operation.

Parameters

- `k`: [in] The key for the value which should be updated in the cache.
- `value`: [in] The value which should be used as a replacement for the existing value in the cache. Any existing cache entry is not changed except for its value.

`template<typename F>`

`bool update_if (key_type const &k, value_type const &val, F f)`

Update an existing element in this cache.

Note The function will call the entry's `entry::touch` function if the indexed value is found in the cache.

Note The difference to the other overload of the `insert` function is that this overload replaces the cached value only, while the other overload replaces the whole cache entry, updating the cache entry properties.

Return This function returns `true` if the entry has been successfully updated, otherwise it returns `false`. If the entry currently is not held by the cache it is added and the return value reflects the outcome of the corresponding insert operation.

Parameters

- `k`: [in] The key for the value which should be updated in the cache.
- `value`: [in] The value which should be used as a replacement for the existing value in the cache. Any existing cache entry is not changed except for its value.
- `f`: [in] A callable taking two arguments, `k` and the key found in the cache (in that order). If `f` returns true, then the update will continue. If `f` returns false, then the update will not succeed.

```
bool update (key_type const &k, entry_type &e)
```

Update an existing entry in this cache.

Note The function will call the entry's `entry::touch` function if the indexed value is found in the cache.

Note The difference to the other overload of the `insert` function is that this overload replaces the whole cache entry, while the other overload replaces the cached value only, leaving the cache entry properties untouched.

Return This function returns `true` if the entry has been successfully updated, otherwise it returns `false`. If the entry currently is not held by the cache it is added and the return value reflects the outcome of the corresponding insert operation.

Parameters

- `k`: [in] The key for the entry which should be updated in the cache.
- `value`: [in] The entry which should be used as a replacement for the existing entry in the cache. Any existing entry is first removed and then this entry is added.

```
template<typename Func>
```

```
size_type erase (Func const &ep = policies::always<storage_value_type>())
```

Remove stored entries from the cache for which the supplied function object returns true.

Return This function returns the overall size of the removed entries (which is the sum of the values returned by the `entry::get_size` functions of the removed entries).

Parameters

- `ep`: [in] This parameter has to be a (unary) function object. It is invoked for each of the entries currently held in the cache. An entry is considered for removal from the cache whenever the value returned from this invocation is `true`. Even then the entry might not be removed from the cache as its `entry::remove` function might return false.

```
size_type erase ()
```

Remove all stored entries from the cache.

Note All entries are considered for removal, but in the end an entry might not be removed from the cache as its `entry::remove` function might return false. This function is very useful for instance in conjunction with an entry's `entry::remove` function enforcing additional criteria like entry expiration, etc.

Return This function returns the overall size of the removed entries (which is the sum of the values returned by the `entry::get_size` functions of the removed entries).

```
void clear ()
```

Clear the cache.

Unconditionally removes all stored entries from the cache.

```
statistics_type const &get_statistics () const  
Allow to access the embedded statistics instance.
```

Return This function returns a reference to the statistics instance embedded inside this cache

```
statistics_type &get_statistics ()
```

Protected Functions

```
bool free_space (long num_free)
```

Private Types

```
typedef storage_type::iterator iterator  
typedef storage_type::const_iterator const_iterator  
typedef std::deque<iterator> heap_type  
typedef heap_type::iterator heap_iterator  
typedef adapt<UpdatePolicy, iterator> adapted_update_policy_type  
typedef statistics_type::update_on_exit update_on_exit
```

Private Members

```
size_type max_size_  
size_type current_size_  
storage_type store_  
heap_type entry_heap_  
adapted_update_policy_type update_policy_  
insert_policy_type insert_policy_  
statistics_type statistics_  
  
template<typename Func, typename Iterator>  
struct adapt
```

Public Functions

```
template<>  
adapt (Func f)
```

```
template<>  
bool operator () (Iterator const &lhs, Iterator const &rhs) const
```

Public Members

template<>
Func **f_**

```
#include <hpx/cache/lru_cache.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace cache
```

```
template<typename Key, typename Entry, typename Statistics = statistics::no_statistics>
class lru_cache
    #include <hpx/cache/lru_cache.hpp> The lru_cache implements the basic functionality needed
    for a local (non-distributed) LRU cache.
```

Template Parameters

- **Key**: The type of the keys to use to identify the entries stored in the cache
- **Entry**: The type of the items to be held in the cache.
- **Statistics**: A (optional) type allowing to collect some basic statistics about the operation of the cache instance. The type must conform to the CacheStatistics concept. The default value is the type `statistics::no_statistics` which does not collect any numbers, but provides empty stubs allowing the code to compile.

Public Types

```
typedef Key key_type
typedef Entry entry_type
typedef Statistics statistics_type
typedef std::pair<key_type, entry_type> entry_pair
typedef std::list<entry_pair> storage_type
typedef std::map<Key, typename storage_type::iterator> map_type
typedef std::size_t size_type
```

Public Functions

```
lru_cache (size_type max_size = 0)
```

Construct an instance of a `lru_cache`.

Parameters

- `max_size`: [in] The maximal size this cache is allowed to reach any time. The default is zero (no size limitation). The unit of this value is usually determined by the unit of the values returned by the entry's `get_size` function.

```
lru_cache (lru_cache &&other)
```

`size_type size() const`

Return current size of the cache.

Return The current size of this cache instance.

`size_type capacity() const`

Access the maximum size the cache is allowed to grow to.

Note The unit of this value is usually determined by the unit of the return values of the entry's function `entry::get_size`.

Return The maximum size this cache instance is currently allowed to reach. If this number is zero the cache has no limitation with regard to a maximum size.

`void reserve(size_type max_size)`

Change the maximum size this cache can grow to.

Parameters

- `max_size`: [in] The new maximum size this cache will be allowed to grow to.

`bool holds_key(key_type const &key)`

Check whether the cache currently holds an entry identified by the given key.

Note This function does not call the entry's function `entry::touch`. It just checks if the cache contains an entry corresponding to the given key.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `k`: [in] The key for the entry which should be looked up in the cache.

`bool get_entry(key_type const &key, key_type &realkey, entry_type &entry)`

Get a specific entry identified by the given key.

Note The function will "touch" the entry and mark it as recently used if the key was found in the cache.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `key`: [in] The key for the entry which should be retrieved from the cache.
- `entry`: [out] If the entry indexed by the key is found in the cache this value on successful return will be a copy of the corresponding entry.

`bool get_entry(key_type const &key, entry_type &entry)`

Get a specific entry identified by the given key.

Note The function will "touch" the entry and mark it as recently used if the key was found in the cache.

Return This function returns *true* if the cache holds the referenced entry, otherwise it returns *false*.

Parameters

- `key`: [in] The key for the entry which should be retrieved from the cache.
- `entry`: [out] If the entry indexed by the key is found in the cache this value on successful return will be a copy of the corresponding entry.

`bool insert(key_type const &key, entry_type const &entry)`

Insert a new entry into this cache.

Note This function assumes that the entry is not in the cache already. Inserting an already existing entry is considered undefined behavior

Parameters

- **key:** [in] The key for the entry which should be added to the cache.
- **entry:** [in] The entry which should be added to the cache.

`void insert_nonexist (key_type const &key, entry_type const &entry)`

`void update (key_type const &key, entry_type const &entry)`

Update an existing element in this cache.

Note The function will “touch” the entry and mark it as recently used if the key was found in the cache.

Note The difference to the other overload of the *insert* function is that this overload replaces the cached value only, while the other overload replaces the whole cache entry, updating the cache entry properties.

Parameters

- **key:** [in] The key for the value which should be updated in the cache.
- **entry:** [in] The entry which should be used as a replacement for the existing value in the cache. Any existing cache entry is not changed except for its value.

`template<typename F>`

`bool update_if (key_type const &key, entry_type const &entry, F &&f)`

Update an existing element in this cache.

Note The function will “touch” the entry and mark it as recently used if the key was found in the cache.

Note The difference to the other overload of the *insert* function is that this overload replaces the cached value only, while the other overload replaces the whole cache entry, updating the cache entry properties.

Return This function returns *true* if the entry has been successfully updated, otherwise it returns *false*. If the entry currently is not held by the cache it is added and the return value reflects the outcome of the corresponding insert operation.

Parameters

- **key:** [in] The key for the value which should be updated in the cache.
- **entry:** [in] The value which should be used as a replacement for the existing value in the cache. Any existing cache entry is not changed except for its value.
- **f:** [in] A callable taking two arguments, *k* and the key found in the cache (in that order). If *f* returns true, then the update will continue. If *f* returns false, then the update will not succeed.

`template<typename Func>`

`size_type erase (Func const &ep)`

Remove stored entries from the cache for which the supplied function object returns true.

Return This function returns the overall size of the removed entries (which is the sum of the values returned by the *entry::get_size* functions of the removed entries).

Parameters

- **ep:** [in] This parameter has to be a (unary) function object. It is invoked for each of the entries currently held in the cache. An entry is considered for removal from the cache whenever the value returned from this invocation is *true*.

`size_type erase ()`

Remove all stored entries from the cache.

Return This function returns the overall size of the removed entries (which is the sum of the values returned by the *entry::get_size* functions of the removed entries).

`size_type clear ()`

Clear the cache.

Unconditionally removes all stored entries from the cache.

`statistics_type const &get_statistics() const`

Allow to access the embedded statistics instance.

Return This function returns a reference to the statistics instance embedded inside this cache

`statistics_type &get_statistics()`

Private Types

`typedef statistics_type::update_on_exit update_on_exit`

Private Functions

`void touch (typename storage_type::iterator it)`

`void evict ()`

Private Members

`size_type max_size_`
`size_type current_size_`
`storage_type storage_`
`map_type map_`
`statistics_type statistics_`

```
#include <hpx/cache/entries/fifo_entry.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace cache
```

```
namespace entries
```

```
template<typename Value>
class fifo_entry : public hpx::util::cache::entries::entry<Value,fifo_entry<Value>>
#include <hpx/cache/entries/fifo_entry.hpp> The fifo_entry type can be used to store arbitrary values in a cache. Using this type as the cache's entry type makes sure that the least recently inserted entries are discarded from the cache first.
```

Note The `fifo_entry` conforms to the CacheEntry concept.

Note This type can be used to model a ‘last in first out’ cache policy if it is used with a `std::greater` as the caches’ UpdatePolicy (instead of the default `std::less`).

Template Parameters

- **Value:** The data type to be stored in a cache. It has to be default constructible, copy constructible and less_than_comparable.

Public Functions

`fifo_entry()`

Any cache entry has to be default constructible.

`fifo_entry(Value const &val)`

Construct a new instance of a cache entry holding the given value.

`bool insert()`

The function `insert` is called by a cache whenever it is about to be inserted into the cache.

Note This function is part of the CacheEntry concept

Return This function should return `true` if the entry should be added to the cache, otherwise it should return `false`.

`std::chrono::steady_clock::time_point const &get_creation_time() const`

Private Types

`typedef entry<Value, fifo_entry<Value>> base_type`

Private Members

`std::chrono::steady_clock::time_point insertion_time_`

Friends

`bool operator<(fifo_entry const &lhs, fifo_entry const &rhs)`

Compare the ‘age’ of two entries. An entry is ‘older’ than another entry if it has been created earlier (FIFO).

```
#include <hpx/cache/entries/size_entry.hpp>

namespace hpx

namespace util

namespace cache

namespace entries

template<typename Value, typename Derived>
```

```
class size_entry : public hpx::util::cache::entries::entry<Value, detail::size_derived<Value, Derived>>{
```

#include <hpx/cache/entries/size_entry.hpp> The `size_entry` type can be used to store values in a cache which have a size associated (such as files, etc.). Using this type as the cache's entry type makes sure that the entries with the biggest size are discarded from the cache first.

Note The `size_entry` conforms to the CacheEntry concept.

Note This type can be used to model a ‘discard smallest first’ cache policy if it is used with a `std::greater` as the caches’ UpdatePolicy (instead of the default `std::less`).

Template Parameters

- `Value`: The data type to be stored in a cache. It has to be default constructible, copy constructible and `less_than_comparable`.
- `Derived`: The (optional) type for which this type is used as a base class.

Public Functions

`size_entry()`

Any cache entry has to be default constructible.

`size_entry(Value const &val, std::size_t size)`

Construct a new instance of a cache entry holding the given value.

`std::size_t get_size() const`

Return the ‘size’ of this entry.

Private Types

`typedef detail::size_derived<Value, Derived>::type derived_type`

`typedef entry<Value, derived_type> base_type`

Private Members

`std::size_t size_`

Friends

`bool operator<(size_entry const &lhs, size_entry const &rhs)`

Compare the ‘age’ of two entries. An entry is ‘older’ than another entry if it has a bigger size.

```
#include <hpx/cache/entries/lfu_entry.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace cache
```

namespace entries

```
template<typename Value>
class lfu_entry : public hpx::util::cache::entries::entry<Value, lfu_entry<Value>>
#include <hpx/cache/entries/lfu_entry.hpp> The lfu_entry type can be used to store arbitrary values in a cache. Using this type as the cache's entry type makes sure that the least frequently used entries are discarded from the cache first.
```

Note The `lfu_entry` conforms to the CacheEntry concept.

Note This type can be used to model a ‘most frequently used’ cache policy if it is used with a `std::greater` as the caches’ UpdatePolicy (instead of the default `std::less`).

Template Parameters

- `Value`: The data type to be stored in a cache. It has to be default constructible, copy constructible and less_than_comparable.

Public Functions**lfu_entry()**

Any cache entry has to be default constructible.

lfu_entry (Value const &val)

Construct a new instance of a cache entry holding the given value.

bool touch()

The function `touch` is called by a cache holding this instance whenever it has been requested (touched).

In the case of the LFU entry we store the reference count tracking the number of times this entry has been requested. This which will be used to compare the age of an entry during the invocation of the `operator<()`.

Return This function should return true if the cache needs to update its internal heap.

Usually this is needed if the entry has been changed by `touch()` in a way influencing the sort order as mandated by the cache’s UpdatePolicy

unsigned long const &get_access_count () const**Private Types****typedef entry<Value, lfu_entry<Value>> base_type****Private Members****unsigned long ref_count_**

Friends

```
bool operator< (lfu_entry const &lhs, lfu_entry const &rhs)
    Compare the ‘age’ of two entries. An entry is ‘older’ than another entry if it has been
    accessed less frequently (LFU).
```

```
#include <hpx/cache/entries/lru_entry.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace cache
```

```
            namespace entries
```

```
template<typename Value>
class lru_entry : public hpx::util::cache::entries::entry<Value, lru_entry<Value>>
    #include <hpx/cache/entries/lru_entry.hpp> The lru_entry type can be used to store arbitrary
    values in a cache. Using this type as the cache’s entry type makes sure that the least recently
    used entries are discarded from the cache first.
```

Note The [lru_entry](#) conforms to the CacheEntry concept.

Note This type can be used to model a ‘most recently used’ cache policy if it is used with a `std::greater` as the caches’ UpdatePolicy (instead of the default `std::less`).

Template Parameters

- **Value:** The data type to be stored in a cache. It has to be default constructible, copy constructible and less_than_comparable.

Public Functions

```
lru_entry()
```

Any cache entry has to be default constructible.

```
lru_entry (Value const &val)
```

Construct a new instance of a cache entry holding the given value.

```
bool touch()
```

The function *touch* is called by a cache holding this instance whenever it has been requested (touched).

In the case of the LRU entry we store the time of the last access which will be used to compare the age of an entry during the invocation of the [operator<\(\)](#).

Return This function should return true if the cache needs to update its internal heap.

Usually this is needed if the entry has been changed by [touch\(\)](#) in a way influencing the sort order as mandated by the cache’s UpdatePolicy

```
std::chrono::steady_clock::time_point const &get_access_time() const
```

Returns the last access time of the entry.

Private Types

`typedef entry<Value, lru_entry<Value>> base_type`

Private Members

`std::chrono::steady_clock::time_point access_time_`

Friends

`bool operator< (lru_entry const &lhs, lru_entry const &rhs)`

Compare the ‘age’ of two entries. An entry is ‘older’ than another entry if it has been accessed less recently (LRU).

```
#include <hpx/cache/entries/entry.hpp>

namespace hpx

namespace util

namespace cache

namespace entries

template<typename Value, typename Derived>
class entry : private boost::less_than_comparable<detail::derived<Value, Derived>::type>
# include <hpx/cache/entries/entry.hpp>
```

Template Parameters

- **Value:** The data type to be stored in a cache. It has to be default constructible, copy constructible and less_than_comparable.
- **Derived:** The (optional) type for which this type is used as a base class.

Public Types

`typedef Value value_type`

Public Functions

`entry()`

Any cache entry has to be default constructible.

`entry (value_type const &val)`

Construct a new instance of a cache entry holding the given value.

`bool touch()`

The function *touch* is called by a cache holding this instance whenever it has been requested (touched).

Note It is possible to change the entry in a way influencing the sort criteria mandated by the UpdatePolicy. In this case the function should return *true* to indicate this to the cache, forcing to reorder the cache entries.

Note This function is part of the CacheEntry concept

Return This function should return true if the cache needs to update it's internal heap.

Usually this is needed if the entry has been changed by `touch()` in a way influencing the sort order as mandated by the cache's UpdatePolicy

`bool insert ()`

The function *insert* is called by a cache whenever it is about to be inserted into the cache.

Note This function is part of the CacheEntry concept

Return This function should return *true* if the entry should be added to the cache, otherwise it should return *false*.

`bool remove ()`

The function *remove* is called by a cache holding this instance whenever it is about to be removed from the cache.

Note This function is part of the CacheEntry concept

Return The return value can be used to avoid removing this instance from the cache. If the value is *true* it is ok to remove the entry, other wise it will stay in the cache.

`std::size_t get_size () const`

Return the 'size' of this entry. By default the size of each entry is just one (1), which is sensible if the cache has a limit (capacity) measured in number of entries.

`value_type &get ()`

Get a reference to the stored data value.

Note This function is part of the CacheEntry concept

`value_type const &get () const`

Private Members

`value_type value_`

Friends

`bool operator< (entry const &lhs, entry const &rhs)`

Forwarding operator< allowing to compare entries in stead of the values.

```
#include <hpx/cache/statistics/local_full_statistics.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace cache
```

```
namespace statistics

class local_full_statistics : public hpx::util::cache::statistics::local_statistics
```

Public Functions

`std::int64_t get_get_entry_count (bool reset)`

The function `get_get_entry_count` returns the number of invocations of the `get_entry()` API function of the cache.

`std::int64_t get_insert_entry_count (bool reset)`

The function `get_insert_entry_count` returns the number of invocations of the `insert_entry()` API function of the cache.

`std::int64_t get_update_entry_count (bool reset)`

The function `get_update_entry_count` returns the number of invocations of the `update_entry()` API function of the cache.

`std::int64_t get_erase_entry_count (bool reset)`

The function `get_erase_entry_count` returns the number of invocations of the `erase()` API function of the cache.

`std::int64_t get_get_entry_time (bool reset)`

The function `get_get_entry_time` returns the overall time spent executing of the `get_entry()` API function of the cache.

`std::int64_t get_insert_entry_time (bool reset)`

The function `get_insert_entry_time` returns the overall time spent executing of the `insert_entry()` API function of the cache.

`std::int64_t get_update_entry_time (bool reset)`

The function `get_update_entry_time` returns the overall time spent executing of the `update_entry()` API function of the cache.

`std::int64_t get_erase_entry_time (bool reset)`

The function `get_erase_entry_time` returns the overall time spent executing of the `erase()` API function of the cache.

Private Functions

`std::int64_t get_and_reset_value (std::int64_t &value, bool reset)`

Private Members

`api_counter_data get_entry_`

`api_counter_data insert_entry_`

`api_counter_data update_entry_`

`api_counter_data erase_entry_`

Friends

```
friend hpx::util::cache::statistics::update_on_exit  
struct api_counter_data
```

Public Functions

```
api_counter_data()
```

Public Members

```
std::int64_t count_  
std::int64_t time_  
  
struct update_on_exit  
#include <local_full_statistics.hpp> Helper class to update timings and counts on function  
exit.
```

Public Functions

```
update_on_exit (local_full_statistics &stat, method m)  
~update_on_exit()
```

Public Members

```
std::int64_t started_at_  
api_counter_data &data_
```

Private Static Functions

```
static api_counter_data &get_api_counter_data (local_full_statistics &stat,  
method m)  
  
static std::uint64_t now()  
  
#include <hpx/cache/statistics/local_statistics.hpp>  
  
namespace hpx  
  
    namespace util  
  
        namespace cache  
  
            namespace statistics
```

```
class local_statistics : public hpx::util::cache::statistics::no_statistics
Subclassed by hpx::util::cache::statistics::local_full_statistics
```

Public Functions

```
local_statistics()

std::size_t get_and_reset (std::size_t &value, bool reset)

std::size_t hits() const

std::size_t misses() const

std::size_t insertions() const

std::size_t evictions() const

std::size_t hits(bool reset)

std::size_t misses(bool reset)

std::size_t insertions(bool reset)

std::size_t evictions(bool reset)

void got_hit()
    The function got_hit will be called by a cache instance whenever a entry got touched.

void got_miss()
    The function got_miss will be called by a cache instance whenever a requested entry has not been found in the cache.

void got_insertion()
    The function got_insertion will be called by a cache instance whenever a new entry has been inserted.

void got_eviction()
    The function got_eviction will be called by a cache instance whenever an entry has been removed from the cache because a new inserted entry let the cache grow beyond its capacity.

void clear()
    Reset all statistics.
```

Private Members

```
std::size_t hits_
std::size_t misses_
std::size_t insertions_
std::size_t evictions_
```

```
#include <hpx/cache/statistics/no_statistics.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace cache
```

```
            namespace statistics
```

Enums

```
enum method
```

Values:

```
method_get_entry = 0  
method_insert_entry = 1  
method_update_entry = 2  
method_erase_entry = 3
```

```
class no_statistics
```

Subclassed by [hpx::util::cache::statistics::local_statistics](#)

Public Functions

```
void got_hit()
```

The function *got_hit* will be called by a cache instance whenever a entry got touched.

```
void got_miss()
```

The function *got_miss* will be called by a cache instance whenever a requested entry has not been found in the cache.

```
void got_insertion()
```

The function *got_insertion* will be called by a cache instance whenever a new entry has been inserted.

```
void got_eviction()
```

The function *got_eviction* will be called by a cache instance whenever an entry has been removed from the cache because a new inserted entry let the cache grow beyond its capacity.

```
void clear()
```

Reset all statistics.

```
std::int64_t get_get_entry_count(bool)
```

The function *get_get_entry_count* returns the number of invocations of the *get_entry()* API function of the cache.

```
std::int64_t get_insert_entry_count(bool)
```

The function *get_insert_entry_count* returns the number of invocations of the *insert_entry()* API function of the cache.

`std::int64_t get_update_entry_count (bool)`

The function `get_update_entry_count` returns the number of invocations of the `update_entry()` API function of the cache.

`std::int64_t get_erase_entry_count (bool)`

The function `get_erase_entry_count` returns the number of invocations of the `erase()` API function of the cache.

`std::int64_t get_get_entry_time (bool)`

The function `get_get_entry_time` returns the overall time spent executing of the `get_entry()` API function of the cache.

`std::int64_t get_insert_entry_time (bool)`

The function `get_insert_entry_time` returns the overall time spent executing of the `insert_entry()` API function of the cache.

`std::int64_t get_update_entry_time (bool)`

The function `get_update_entry_time` returns the overall time spent executing of the `update_entry()` API function of the cache.

`std::int64_t get_erase_entry_time (bool)`

The function `get_erase_entry_time` returns the overall time spent executing of the `erase()` API function of the cache.

`struct update_on_exit`

`#include <no_statistics.hpp>` Helper class to update timings and counts on function exit.

Public Functions

`update_on_exit (no_statistics const&, method)`

```
#include <hpx/cache/policies/always.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace cache
```

```
namespace policies
```

```
template<typename Entry>
```

```
struct always
```

Public Functions

```
bool operator() (Entry const&)
```

hashing

```
#include <compatibility/hpx/util/jenkins_hash.hpp>

#include <compatibility/hpx/util/fibhash.hpp>

#include <hpx/hashing/jenkins_hash.hpp>

namespace hpx

namespace util

class jenkins_hash
    #include <jenkins_hash.hpp> The jenkins_hash class encapsulates a hash calculation function published by Bob Jenkins here: http://burtleburtle.net/bob/hash
```

Public Types

enum seedenum

The seedenum is used as a dummy parameter to distinguish the different constructors

Values:

seed = 1

typedef std::uint32_t size_type

this is the type representing the result of this hash

Public Functions

jenkins_hash()

constructors and destructor

jenkins_hash(size_type size)

jenkins_hash(size_type seedval, seedenum)

~jenkins_hash()

size_type operator()(std::string const &key) const

calculate the hash value for the given key

size_type operator()(char const *key) const

bool reset(size_type size)

re-seed the hash generator

```
void set_seed(size_type seedval)
    initialize the hash generator to a specific seed

void swap(jenkins_hash &rhs)
    support for std::swap
```

Protected Functions

```
size_type hash(const char *k, std::size_t length) const
```

Private Members

```
size_type seed_
```

```
#include <hpx/hashing/fibhash.hpp>

namespace hpx

    namespace util
```

Functions

```
template<std::uint64_t N>
constexpr std::uint64_t fibhash(std::uint64_t i)
```

serialization

```
#include <compatibility/hpx/traits/brace_initializable_traits.hpp>

#include <compatibility/hpx/traits/needs_automatic_registration.hpp>

#include <compatibility/hpx/traits/polymorphic_traits.hpp>

#include <compatibility/hpx/traits/is_bitwise_serializable.hpp>

#include <compatibility/hpx/traits/serialization_access_data.hpp>

#include <compatibility/hpx/runtime/serialization/map.hpp>

#include <compatibility/hpx/runtime/serialization/basic_archive.hpp>

#include <compatibility/hpx/runtime/serialization/serialize.hpp>

#include <compatibility/hpx/runtime/serialization/output_container.hpp>
```

```
#include <compatibility/hpx/runtime/serialization/dynamic_bitset.hpp>

#include <compatibility/hpx/runtime/serialization/unordered_map.hpp>

#include <compatibility/hpx/runtime/serialization/container.hpp>

#include <compatibility/hpx/runtime/serialization/list.hpp>

#include <compatibility/hpx/runtime/serialization/set.hpp>

#include <compatibility/hpx/runtime/serialization/optional.hpp>

#include <compatibility/hpx/runtime/serialization/multi_array.hpp>

#include <compatibility/hpx/runtime/serialization/brace_initializable.hpp>

#include <compatibility/hpx/runtime/serialization/shared_ptr.hpp>

#include <compatibility/hpx/runtime/serialization/serialization_chunk.hpp>

#include <compatibility/hpx/runtime/serialization/output_archive.hpp>

#include <compatibility/hpx/runtime/serialization/serialization_fwd.hpp>

#include <compatibility/hpx/runtime/serialization/binary_filter.hpp>

#include <compatibility/hpx/runtime/serialization/input_container.hpp>

#include <compatibility/hpx/runtime/serialization/base_object.hpp>

#include <compatibility/hpx/runtime/serialization/access.hpp>

#include <compatibility/hpx/runtime/serialization/tuple.hpp>

#include <compatibility/hpx/runtime/serialization/string.hpp>

#include <compatibility/hpx/runtime/serialization/variant.hpp>

#include <compatibility/hpx/runtime/serialization/vector.hpp>

#include <compatibility/hpx/runtime/serialization/unique_ptr.hpp>

#include <compatibility/hpx/runtime/serialization/deque.hpp>

#include <compatibility/hpx/runtime/serialization/array.hpp>
```

```
#include <compatibility/hpx/runtime/serialization/brace_initializable_fwd.hpp>

#include <compatibility/hpx/runtime/serialization/serialize_buffer.hpp>

#include <compatibility/hpx/runtime/serialization/bitset.hpp>

#include <compatibility/hpx/runtime/serialization/valarray.hpp>

#include <compatibility/hpx/runtime/serialization/input_archive.hpp>

#include <compatibility/hpx/runtime/serialization/complex.hpp>

#include <compatibility/hpx/runtime/serialization/intrusive_ptr.hpp>

#include <hpx/serialization/map.hpp>

namespace hpx

    namespace serialization
```

Functions

```
template<typename Key, typename Value>
void serialize(input_archive &ar, std::pair<Key, Value> &t, unsigned)

template<typename Key, typename Value>
void serialize(output_archive &ar, const std::pair<Key, Value> &t, unsigned)

template<typename Key, typename Value, typename Comp, typename Alloc>
void serialize(input_archive &ar, std::map<Key, Value, Comp, Alloc> &t, unsigned)

template<typename Key, typename Value, typename Comp, typename Alloc>
void serialize(output_archive &ar, std::map<Key, Value, Comp, Alloc> const &t, unsigned)

#include <hpx/serialization/basic_archive.hpp>

namespace hpx

    namespace serialization
```

Enums

```
enum archive_flags
Values:
no_archive_flags = 0x00000000
enable_compression = 0x00002000
endian_big = 0x00004000
endian_little = 0x00008000
disable_array_optimization = 0x00010000
disable_data_chunking = 0x00020000
all_archive_flags = 0x0003e000
```

Functions

```
void reverse_bytes (std::size_t size, char *address)

template<typename Archive>
void save_binary (Archive &ar, void const *address, std::size_t count)

template<typename Archive>
void load_binary (Archive &ar, void *address, std::size_t count)

template<typename Archive>
std::size_t current_pos (const Archive &ar)

template<typename Archive>
struct basic_archive
```

Public Functions

```
virtual ~basic_archive ()

template<typename T>
void invoke (T &t)

bool enable_compression () const
bool endian_big () const
bool endian_little () const
bool disable_array_optimization () const
bool disable_data_chunking () const
std::uint32_t flags () const
bool is_preprocessing () const
std::size_t current_pos () const

void save_binary (void const *address, std::size_t count)
```

```
void load_binary(void *address, std::size_t count)
void reset()

template<typename T>
T &get_extra_data()

template<typename T>
T *try_get_extra_data()
```

Public Static Attributes

```
const std::uint64_t npos = std::uint64_t(-1)
```

Protected Functions

```
basic_archive (std::uint32_t flags)
basic_archive (basic_archive const&)
basic_archive &operator= (basic_archive const&)
```

Protected Attributes

```
std::uint32_t flags_
std::size_t size_
detail::extra_archive_data extra_data_
```

```
#include <hpx/serialization/serialize.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T>
output_archive &operator<< (output_archive &ar, T const &t)

template<typename T>
input_archive &operator>> (input_archive &ar, T &t)

template<typename T>
output_archive &operator& (output_archive &ar, T const &t)

template<typename T>
input_archive &operator& (input_archive &ar, T &t)
```

```
#include <hpx/serialization/output_container.hpp>

namespace hpx

namespace serialization

template<typename Container, typename Chunker>
struct filtered_output_container : public hpx::serialization::output_container<Container, Chunker>
```

Public Types

```
template<>
using access_traits = traits::serialization_access_data<Container>

template<>
using base_type = output_container<Container, Chunker>
```

Public Functions

```
filtered_output_container(Container &cont, std::vector<serialization_chunk>
                          *chunks = nullptr)

~filtered_output_container()

void flush()

void set_filter(binary_filter *filter)

void save_binary(void const *address, std::size_t count)

std::size_t save_binary_chunk(void const *address, std::size_t count)
```

Protected Attributes

```
std::size_t start_compressing_at_
binary_filter *filter_

template<typename Container, typename Chunker>
struct output_container : public hpx::serialization::erased_output_container
    Subclassed by hpx::serialization::filtered_output_container<Container, Chunker>
```

Public Types

```
template<>
using access_traits = traits::serialization_access_data<Container>
```

Public Functions

```
output_container (Container &cont, std::vector<serialization_chunk> *chunks = nullptr)
~output_container()

void flush()

std::size_t get_num_chunks () const

void reset()

void set_filter (binary_filter *filter)

void save_binary (void const *address, std::size_t count)

std::size_t save_binary_chunk (void const *address, std::size_t count)

bool is_preprocessing () const
```

Protected Attributes

```
Container &cont_
std::size_t current_
Chunker chunker_
```

```
#include <hpx/serialization/dynamic_bitset.hpp>

#include <hpx/serialization/unordered_map.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Key, typename Value, typename Hash, typename KeyEqual, typename Allocserialize (input_archive &ar, std::unordered_map<Key, Value, Hash, KeyEqual, Alloc> &t,
               unsigned)
```

```
template<typename Key, typename Value, typename Hash, typename KeyEqual, typename Allocserialize (output_archive &ar, const std::unordered_map<Key, Value, Hash, KeyEqual,
                Alloc> &t, unsigned)
```

```
#include <hpx/serialization/container.hpp>

namespace hpx

namespace serialization

struct erased_input_container
    Subclassed by hpx::serialization::input_container<Container>
```

Public Functions

```
virtual ~erased_input_container()

virtual bool is_preprocessing() const

virtual void set_filter(binary_filter *filter) = 0

virtual void load_binary(void *address, std::size_t count) = 0

virtual void load_binary_chunk(void *address, std::size_t count) = 0
```

struct erased_output_container

```
    Subclassed by hpx::serialization::output_container<Container, Chunker>
```

Public Functions

```
virtual ~erased_output_container()

virtual bool is_preprocessing() const

virtual void set_filter(binary_filter *filter) = 0

virtual void save_binary(void const *address, std::size_t count) = 0

virtual std::size_t save_binary_chunk(void const *address, std::size_t count) = 0

virtual void reset() = 0

virtual std::size_t get_num_chunks() const = 0

virtual void flush() = 0
```

```
#include <hpx/serialization/list.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T, typename Allocator>
void serialize(input_archive &ar, std::list<T, Allocator> &ls, unsigned)

template<typename T, typename Allocator>
void serialize(output_archive &ar, const std::list<T, Allocator> &ls, unsigned)
```

```
#include <hpx/serialization/set.hpp>
```

```
namespace hpx
```

```
    namespace serialization
```

Functions

```
template<typename T, typename Compare, typename Allocator>
void serialize(input_archive &ar, std::set<T, Compare, Allocator> &set, unsigned)

template<typename T, typename Compare, typename Allocator>
void serialize(output_archive &ar, std::set<T, Compare, Allocator> const &set, unsigned)
```

```
#include <hpx/serialization/optional.hpp>
```

```
namespace hpx
```

```
    namespace serialization
```

Functions

```
template<typename T>
void save(output_archive &ar, hpx::util::optional<T> const &o, unsigned)

template<typename T>
void load(input_archive &ar, hpx::util::optional<T> &o, unsigned)

hpx::serialization::HPX_SERIALIZATION_SPLIT_FREE_TEMPLATE((template< typename T >,
```

```
#include <hpx/serialization/multi_array.hpp>
```

```
#include <hpx/serialization/brace_initializable.hpp>
```

```
#include <hpx/serialization/shared_ptr.hpp>
```

```
namespace hpx
```

```
    namespace serialization
```

Functions

```
template<typename T>
void load(input_archive &ar, std::shared_ptr<T> &ptr, unsigned)

template<typename T>
void save(output_archive &ar, std::shared_ptr<T> const &ptr, unsigned)
```

```
#include <hpx/serialization/serialization_chunk.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Enums

```
enum chunk_type
```

Values:

```
chunk_type_index = 0
chunk_type_pointer = 1
```

Functions

```
serialization_chunk create_index_chunk (std::size_t index, std::size_t size)
```

```
serialization_chunk create_pointer_chunk (void const *pos, std::size_t size, std::uint64_t rkey = 0)
```

```
union chunk_data
```

Public Members

```
std::size_t index_
```

```
void const *cpos_
```

```
void *pos_
```

```
struct serialization_chunk
```

Public Members

```
chunk_data data_
```

```
std::size_t size_
```

```
std::uint64_t rkey_
```

```
std::uint8_t type_
```

```
#include <hpx/serialization/output_archive.hpp>

namespace hpx

namespace serialization

struct output_archive : public hpx::serialization::basic_archive<output_archive>
```

Public Types

```
using base_type = basic_archive<output_archive>
```

Public Functions

```
template<typename Container>
output_archive(Container &buffer, std::uint32_t flags = 0U,
               std::vector<serialization_chunk> *chunks = nullptr, binary_filter *filter =
               nullptr)

std::size_t bytes_written() const

std::size_t get_num_chunks() const

std::size_t current_pos() const

void reset()

void flush()

bool is_preprocessing() const
```

Protected Functions

```
template<typename T>
void invoke_impl(T const &t)

template<typename T>
std::enable_if<!std::is_integral<T>::value && !std::is_enum<T>::value>::type save(T const
&t)

template<typename T>
std::enable_if<std::is_integral<T>::value || std::is_enum<T>::value>::type save(T t)

void save(float f)

void save(double d)

void save(char c)

void save(bool b)

template<typename T>
void save_bitwise(T const &t, std::false_type)
```

```
template<typename T>
void save_bitwise(T const &t, std::true_type)

template<typename T>
void save_nonintrusively_polymorphic(T const &t, std::false_type)

template<typename T>
void save_nonintrusively_polymorphic(T const &t, std::true_type)

template<typename T>
void save_integral(T val, std::false_type)

template<typename T>
void save_integral(T val, std::true_type)

template<class Promoted>
void save_integral_impl(Promoted l)

void save_binary(void const *address, std::size_t count)
void save_binary_chunk(void const *address, std::size_t count)
```

Protected Attributes

```
std::unique_ptr<erased_output_container> buffer_
```

Private Static Functions

```
static std::uint32_t make_flags(std::uint32_t flags, std::vector<serialization_chunk> *chunks)
```

Friends

```
friend hpx::serialization::basic_archive< output_archive >
friend hpx::serialization::array
```

```
#include <hpx/serialization/serialization_fwd.hpp>
```

Defines

```
HPX_SERIALIZATION_SPLIT_MEMBER()
HPX_SERIALIZATION_SPLIT_FREE(T)
HPX_SERIALIZATION_SPLIT_FREE_TEMPLATE(TEMPLATE, ARGS)
namespace hpx

namespace serialization
```

Functions

```
template<typename T>
output_archive &operator& (output_archive &ar, T const &t)

template<typename T>
input_archive &operator& (input_archive &ar, T &t)

#include <hpx/serialization/binary_filter.hpp>

namespace hpx

namespace serialization

struct binary_filter
```

Public Functions

```
virtual void set_max_length(std::size_t size) = 0

virtual void save(void const *src, std::size_t src_count) = 0

virtual bool flush(void *dst, std::size_t dst_count, std::size_t &written) = 0

virtual std::size_t init_data(char const *buffer, std::size_t size, std::size_t buffer_size)
    = 0

virtual void load(void *dst, std::size_t dst_count) = 0

template<class T>
void serialize(T&, unsigned)

HPX_SERIALIZATION_POLYMORPHIC_ABSTRACT(binary_filter)

virtual ~binary_filter()

#include <hpx/serialization/input_container.hpp>

namespace hpx

namespace serialization

template<typename Container>
struct input_container : public hpx::serialization::erased_input_container
```

Public Functions

```
input_container(Container const &cont, std::size_t inbound_data_size)
input_container(Container const &cont, std::vector<serialization_chunk> const
                *chunks, std::size_t inbound_data_size)
void set_filter(binary_filter *filter)
void load_binary(void *address, std::size_t count)
void load_binary_chunk(void *address, std::size_t count)
```

Public Members

```
Container const &cont_
std::size_t current_
std::unique_ptr<binary_filter> filter_
std::size_t decompressed_size_
std::vector<serialization_chunk> const *chunks_
std::size_t current_chunk_
std::size_t current_chunk_size_
```

Private Types

```
template<>
using access_traits = traits::serialization_access_data<Container>
```

Private Functions

```
std::size_t get_chunk_size(std::size_t chunk) const
std::uint8_t get_chunk_type(std::size_t chunk) const
chunk_data get_chunk_data(std::size_t chunk) const
std::size_t get_num_chunks() const
```

```
#include <hpx/serialization/base_object.hpp>
```

```
template<typename Derived, typename Base>
struct base_object_type<Derived, Base, std::true_type>
```

Public Functions

```
base_object_type (Derived &d)
template<class Archive>
void save (Archive &ar, unsigned) const

template<class Archive>
void load (Archive &ar, unsigned)

HPX_SERIALIZATION_SPLIT_MEMBER()
```

Public Members

```
Derived &d
namespace hpx

namespace serialization
```

Functions

```
template<typename Base, typename Derived>
base_object_type<Derived, Base> base_object (Derived &d)

template<typename D, typename B>
output_archive &operator<< (output_archive &ar, base_object_type<D, B> t)

template<typename D, typename B>
input_archive &operator>> (input_archive &ar, base_object_type<D, B> t)

template<typename D, typename B>
output_archive &operator& (output_archive &ar, base_object_type<D, B> t)

template<typename D, typename B>
input_archive &operator& (input_archive &ar, base_object_type<D, B> t)

template<typename Derived, typename Base, typename Enable = typename hpx::traits::is_intrusive_polymorphic<Derived>
struct base_object_type
```

Public Functions

```
base_object_type (Derived &d)
template<typename Archive>
void serialize (Archive &ar, unsigned)
```

Public Members

Derived &**d**_

```
template<typename Derived, typename Base>
struct base_object_type<Derived, Base, std::true_type>
```

Public Functions

base_object_type (Derived &**d**)

```
template<class Archive>
void save (Archive &ar, unsigned) const
```

```
template<class Archive>
void load (Archive &ar, unsigned)
```

HPX_SERIALIZATION_SPLIT_MEMBER ()

Public Members

Derived &**d**_

```
#include <hpx/serialization/access.hpp>
```

```
template<typename T>
struct serialize_non_intrusive<T, typename std::enable_if<has_serialize_adl<T>::value>::type>
```

Public Static Functions

```
template<typename Archive>
static void call (Archive &ar, T &t, unsigned)
```

```
namespace hpx
```

```
namespace serialization
```

```
class access
```

Public Static Functions

```
template<class Archive, class T>
static void serialize (Archive &ar, T &t, unsigned)
```

```
template<typename Archive, typename T>
static void save_base_object (Archive &ar, T const &t, unsigned)
```

```
template<typename Archive, typename T>
static void load_base_object (Archive &ar, T &t, unsigned)
```

```
template<typename T>
static std::string get_name (T const *t)

template<class T>
class has_serialize
```

Public Static Attributes

```
constexpr bool value = decltype(test<T>(0))::value
```

Private Static Functions

```
template<class T1>
static std::false_type test (...)

template<class T1, class = decltype(std::declval<typename std::remove_const<T1>::type&>().serialize(std::declval<T1>()))
static std::true_type test (int)

template<class T>
class serialize_dispatcher
```

Public Types

```
template<>
using type = typename std::conditional::type

struct empty
```

Public Static Functions

```
template<class Archive>
static void call (Archive&, T&, unsigned)

struct intrusive_polymorphic
```

Public Static Functions

```
template<>
static void call (hpx::serialization::input_archive &ar, T &t, unsigned)

template<>
static void call (hpx::serialization::output_archive &ar, T const &t, unsigned)

struct intrusive_usual
```

Public Static Functions

```
template<class Archive>
static void call (Archive &ar, T &t, unsigned)

struct non_intrusive
```

Public Static Functions

```
template<class Archive>
static void call (Archive &ar, T &t, unsigned)

template<typename T>
class has_serialize_adl
```

Public Static Attributes

```
constexpr bool value = decltype(test<T>(0))::value
```

Private Static Functions

```
template<typename T1>
static std::false_type test (...)

template<typename T1, typename = decltype(serialize(std::declval<hpx::serialization::output_archive&>(), , ))>
static std::true_type test (int)

template<typename T>
struct serialize_non_intrusive<T, typename std::enable_if<has_serialize_adl<T>::value>::type>
```

Public Static Functions

```
template<typename Archive>
static void call (Archive &ar, T &t, unsigned)

#include <hpx/serialization/tuple.hpp>

namespace hpx

namespace serialization
```

Functions

```
template<typename Archive, typename ...Ts>
void serialize (Archive &ar, hpx::util::tuple<Ts...> &t, unsigned int version)
```

```
template<typename Archive>
void serialize (Archive &ar, hpx::util::tuple<> &, unsigned)
```

```
template<typename Archive, typename ...Ts>
void load_construct_data (Archive &ar, hpx::util::tuple<Ts...> *t, unsigned int version)
```

```
template<typename Archive, typename ...Ts>
void save_construct_data (Archive &ar, hpx::util::tuple<Ts...> const *t, unsigned int version)
```

```
#include <hpx/serialization/string.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Char, typename CharTraits, typename Allocator>
void serialize (input_archive &ar, std::basic_string<Char, CharTraits, Allocator> &s, unsigned)
```

```
template<typename Char, typename CharTraits, typename Allocator>
void serialize (output_archive &ar, std::basic_string<Char, CharTraits, Allocator> const &s, unsigned)
```

```
#include <hpx/serialization/variant.hpp>
```

```
#include <hpx/serialization/vector.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Allocator>
void serialize (input_archive &ar, std::vector<bool, Allocator> &v, unsigned)
```

```
template<typename T, typename Allocator>
void serialize (input_archive &ar, std::vector<T, Allocator> &v, unsigned)
```

```
template<typename Allocator>
void serialize (output_archive &ar, std::vector<bool, Allocator> const &v, unsigned)
```

```
template<typename T, typename Allocator>
void serialize(output_archive &ar, std::vector<T, Allocator> const &v, unsigned)
```

```
#include <hpx/serialization/unique_ptr.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T>
void load(input_archive &ar, std::unique_ptr<T> &ptr, unsigned)
```

```
template<typename T>
void save(output_archive &ar, const std::unique_ptr<T> &ptr, unsigned)
```

```
#include <hpx/serialization/deque.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T, typename Allocator>
void serialize(input_archive &ar, std::deque<T, Allocator> &d, unsigned)
```

```
template<typename T, typename Allocator>
void serialize(output_archive &ar, std::deque<T, Allocator> const &d, unsigned)
```

```
#include <hpx/serialization/array.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<class T>
array<T> make_array(T *begin, std::size_t size)
```

```
template<typename Archive, typename T, std::size_t N>
void serialize(Archive &ar, std::array<T, N> &a, const unsigned int)
```

```
template<typename T>
output_archive &operator<<(output_archive &ar, array<T> t)
```

```

template<typename T>
class array

```

Public Types

```

template<>
using value_type = T

```

Public Functions

```

array (value_type *t, std::size_t s)
value_type *address () const
std::size_t count () const

template<class Archive>
void serialize_optimized (Archive &ar, unsigned int, std::false_type)
void serialize_optimized (output_archive &ar, unsigned int, std::true_type)
void serialize_optimized (input_archive &ar, unsigned int, std::true_type)

template<class Archive>
void serialize (Archive &ar, unsigned int v)

```

Private Members

```
value_type *m_t
std::size_t m_element_count

#include <hpx/serialization/brace_initializable_fwd.hpp>

#include <hpx/serialization/serialize_buffer.hpp>

namespace hpx

namespace serialization

template<typename T, typename Allocator = std::allocator<T>>
class serialize_buffer
```

Public Types

```
enum init_mode
Values:
copy = 0
reference = 1
take = 2

template<>
using value_type = T
```

Public Functions

```
serialize_buffer(allocator_type const &alloc = allocator_type())
serialize_buffer(std::size_t size, allocator_type const &alloc = allocator_type())
serialize_buffer(T *data, std::size_t size, init_mode mode = copy, allocator_type const
&alloc = allocator_type())

template<typename Deallocator>
serialize_buffer(T *data, std::size_t size, allocator_type const &alloc, Deallocator
const &dealloc)

template<typename Deleter>
serialize_buffer(T *data, std::size_t size, init_mode mode, Deleter const &deleter,
allocator_type const &alloc = allocator_type())

template<typename Deleter>
serialize_buffer(T const *data, std::size_t size, init_mode mode, Deleter const
&deleter, allocator_type const &alloc = allocator_type())

template<typename Deallocator, typename Deleter>
serialize_buffer(T *data, std::size_t size, allocator_type const &alloc, Deallocator
const &dealloc, Deleter const &deleter)
```

```

serialize_buffer(T const *data, std::size_t size, allocator_type const &alloc = allocator_type())
template<typename Deleter>
serialize_buffer(T const *data, std::size_t size, Deleter const &deleter, allocator_type const &alloc = allocator_type())
serialize_buffer(T const *data, std::size_t size, init_mode mode, allocator_type const &alloc = allocator_type())

T *data()
T const *data() const
T *begin()
T *end()
T &operator[](std::size_t idx)
T operator[](std::size_t idx) const
boost::shared_array<T> data_array() const
std::size_t size() const

```

Private Types

```

template<>
using allocator_type = Allocator

```

Private Functions

```

template<typename Archive>
void save(Archive &ar, unsigned int const version) const

template<typename Archive>
void load(Archive &ar, unsigned int const version)

```

Private Members

```

boost::shared_array<T> data_
std::size_t size_
Allocator alloc_

```

Private Static Functions

```
static void no_deleter(T *)  
  
template<typename Deallocator>  
static void deleter(T *p, Deallocator dealloc, std::size_t size)
```

Friends

```
friend hpx::serialization::hpx::serialization::access  
bool operator==(serialize_buffer const &rhs, serialize_buffer const &lhs)
```

```
#include <hpx/serialization/bitset.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<std::size_t N>  
void serialize(input_archive &ar, std::bitset<N> &d, unsigned)  
  
template<std::size_t N>  
void serialize(output_archive &ar, std::bitset<N> const &bs, unsigned)
```

```
#include <hpx/serialization/valarray.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T>  
void serialize(input_archive &ar, std::valarray<T> &arr, int)  
  
template<typename T>  
void serialize(output_archive &ar, std::valarray<T> const &arr, int)
```

```
#include <hpx/serialization/input_archive.hpp>

namespace hpx

namespace serialization

struct input_archive : public hpx::serialization::basic_archive<input_archive>
```

Public Types

```
using base_type = basic_archive<input_archive>
```

Public Functions

```
template<typename Container>
input_archive(Container &buffer, std::size_t inbound_data_size = 0, const
std::vector<serialization_chunk> *chunks = nullptr)

template<typename T>
void invoke_impl(T &t)

template<typename T>
std::enable_if<!std::is_integral<T>::value && !std::is_enum<T>::value>::type load(T &t)

template<typename T>
std::enable_if<std::is_integral<T>::value || std::is_enum<T>::value>::type load(T &t)

void load(float &f)
void load(double &d)
void load(char &c)
void load(bool &b)

std::size_t bytes_read() const
std::size_t current_pos() const
```

Private Functions

```
template<typename T>
void load_bitwise(T &t, std::false_type)

template<typename T>
void load_bitwise(T &t, std::true_type)

template<class T>
void load_nonintrusively_polymerphic(T &t, std::false_type)

template<class T>
void load_nonintrusively_polymerphic(T &t, std::true_type)
```

```
template<typename T>
void load_integral (T &val, std::false_type)

template<typename T>
void load_integral (T &val, std::true_type)

template<class Promoted>
void load_integral_impl (Promoted &l)

void load_binary (void *address, std::size_t count)
void load_binary_chunk (void *address, std::size_t count)
```

Private Members

```
std::unique_ptr<erased_input_container> buffer_
```

Friends

```
friend hpx::serialization::basic_archive< input_archive >
friend hpx::serialization::array
```

```
#include <hpx/serialization/complex.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T>
void serialize (input_archive &ar, std::complex<T> &c, unsigned)

template<typename T>
void serialize (output_archive &ar, std::complex<T> const &c, unsigned)
```

```
#include <hpx/serialization/intrusive_ptr.hpp>
```

```
#include <hpx/serialization/traits/brace_initializable_traits.hpp>
```

```
#include <hpx/serialization/traits/needs_automatic_registration.hpp>
```

```
#include <hpx/serialization/traits/polymorphic_traits.hpp>
```

Defines

HPX_TRAITS_NONINTRUSIVE_POLYMORPHIC (Class)

HPX_TRAITS_NONINTRUSIVE_POLYMORPHIC_TEMPLATE (TEMPLATE, ARG_LIST)

HPX_TRAITS_SERIALIZED_WITH_ID (Class)

HPX_TRAITS_SERIALIZED_WITH_ID_TEMPLATE (TEMPLATE, ARG_LIST)

```
#include <hpx/serialization/traits/is_bitwise_serializable.hpp>
```

Defines

HPX_IS_BITWISE_SERIALIZABLE (T)

```
#include <hpx/serialization/traits/serialization_access_data.hpp>
```

```
namespace hpx
```

```
namespace traits
```

```
template<typename Container>
struct default_serialization_access_data
    Subclassed by hpx::traits::serialization_access_data<Container>
```

Public Types

```
template<>
using preprocessing_only = std::false_type
```

Public Static Functions

```
static constexpr bool is_preprocessing()
```

```
static constexpr void write(Container &cont, std::size_t count, std::size_t current, void
                           const *address)
```

```
static bool flush(serialization::binary_filter *filter, Container &cont, std::size_t current,
                  std::size_t size, std::size_t &written)
```

```
static constexpr void read(Container const &cont, std::size_t count, std::size_t current,
                           void *address)
```

```
static constexpr std::size_t init_data(Container const &cont, serialization::binary_filter
                                         *filter, std::size_t current,
                                         std::size_t decompressed_size)
```

```
static constexpr void reset(Container &cont)
```

```
template<typename Container>
struct serialization_access_data : public hpx::traits::default_serialization_access_data<Container>
    Subclassed by hpx::traits::serialization_access_data<Container const >
```

Public Static Functions

```
static std::size_t size(Container const &cont)

static void resize(Container &cont, std::size_t count)

static void write(Container &cont, std::size_t count, std::size_t current, void const *address)

static bool flush(serialization::binary_filter *filter, Container &cont, std::size_t current,
                 std::size_t size, std::size_t &written)

static void read(Container const &cont, std::size_t count, std::size_t current, void *address)

static std::size_t init_data(Container const &cont, serialization::binary_filter *filter,
                           std::size_t current, std::size_t decompressed_size)
```

concepts

```
#include <compatibility/hpx/traits/has_xxx.hpp>

#include <compatibility/hpx/traits/concepts.hpp>

#include <compatibility/hpx/traits/has_member_xxx.hpp>

#include <hpx/concepts/has_xxx.hpp>
```

Defines

HPX_HAS_XXX_TRAIT_DEF (Name)

This macro creates a boolean unary metafunction such that for any type X, has_name<X>::value == true if and only if X is a class type and has a nested type member x::name. The generated trait ends up in a namespace where the macro itself has been placed.

```
#include <hpx/concepts/concepts.hpp>
```

Defines

```
HPX_CONCEPTQUIRES_(...)

HPX_CONCEPTQUIRES(...)

HPX_CONCEPTASSERT(...)
```

```
#include <hpx/concepts/has_member_xxx.hpp>
```

Defines

HPX_HAS_MEMBER_XXX_TRAIT_DEF (MEMBER)

This macro creates a boolean unary metafunction which result is true if and only if its parameter type has member function with MEMBER name (no matter static it is or not). The generated trait ends up in a namespace where the macro itself has been placed.

local_lcoss

```
#include <compatibility/hpx/lcos/local/channel.hpp>

#include <compatibility/hpx/lcos/local/packaged_task.hpp>

#include <compatibility/hpx/lcos/local/receive_buffer.hpp>

#include <compatibility/hpx/lcos/local/trigger.hpp>

#include <compatibility/hpx/lcos/local/composable_guard.hpp>

#include <compatibility/hpx/lcos/local/and_gate.hpp>

#include <compatibility/hpx/lcos/local/conditional_trigger.hpp>

#include <compatibility/hpx/lcos/local/promise.hpp>

#include <hpx/local_lcos/channel.hpp>

namespace hpx

namespace lcos

namespace local

template<typename T>
class channel : protected hpx::lcos::local::detail::channel_base<T>
```

Public Types

```
typedef T value_type
```

Public Functions

```
channel()
```

Private Types

```
typedef detail::channel_base<T> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< T >
```

```
friend hpx::lcos::local::receive_channel< T >
```

```
friend hpx::lcos::local::send_channel< T >
```

```
template<>
```

```
class channel<void> : protected hpx::lcos::local::detail::channel_base<void>
```

Public Types

```
typedef void value_type
```

Public Functions

```
channel()
```

Private Types

```
typedef detail::channel_base<void> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< void >
```

```
friend hpx::lcos::local::receive_channel< void >
```

```
friend hpx::lcos::local::send_channel< void >
```

```
template<typename T>
```

```
class channel_async_iterator : public hpx::util::iterator_facade<channel_async_iterator<T>, hpx::futur
```

Public Functions

```
channel_async_iterator()  
channel_async_iterator(detail::channel_base<T> const *c)
```

Private Types

```
typedef hpx::util::iterator_facade<channel_async_iterator<T>, hpx::future<T>, std::input_iterator_tag, hpx::fut
```

Private Functions

```
std::pair<hpx::future<T>, bool> get_checked() const  
bool equal(channel_async_iterator const &rhs) const  
void increment()  
base_type::reference dereference() const
```

Private Members

```
hpx::intrusive_ptr<detail::channel_impl_base<T>> channel_  
std::pair<hpx::future<T>, bool> data_
```

Friends

```
friend hpx::lcos::local::hpx::util::iterator_core_access  
template<typename T>  
class channel_iterator : public hpx::util::iterator_facade<channel_iterator<T>, T const, std::input_iterator
```

Public Functions

```
channel_iterator()  
channel_iterator(detail::channel_base<T> const *c)  
channel_iterator(receive_channel<T> const *c)
```

Private Types

```
typedef hpx::util::iterator_facade<channel_iterator<T>, T const, std::input_iterator_tag> base_type
```

Private Functions

```
std::pair<T, bool> get_checked() const  
bool equal(channel_iterator const &rhs) const  
void increment()  
base_type::reference dereference() const
```

Private Members

```
hpx::intrusive_ptr<detail::channel_impl_base<T>> channel_  
std::pair<T, bool> data_
```

Friends

```
friend hpx::lcos::local::hpx::util::iterator_core_access  
template<>  
class channel_iterator<void> : public hpx::util::iterator_facade<channel_iterator<void>, util::unused_type
```

Public Functions

```
channel_iterator()  
channel_iterator(detail::channel_base<void> const *c)  
channel_iterator(receive_channel<void> const *c)
```

Private Types

```
typedef hpx::util::iterator_facade<channel_iterator<void>, util::unused_type const, std::input_iterator_tag b
```

Private Functions

```
bool get_checked()  
bool equal(channel_iterator const &rhs) const  
void increment()  
base_type::reference dereference() const
```

Private Members

```
hpx::intrusive_ptr<detail::channel_impl_base<util::unused_type>> channel_
bool data_
```

Friends

```
friend hpx::lcos::local::hpx::util::iterator_core_access
template<typename T>
class one_element_channel : protected hpx::lcos::local::detail::channel_base<T>
```

Public Types

```
typedef T value_type
```

Public Functions

```
one_element_channel()
```

Private Types

```
typedef detail::channel_base<T> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< T >
friend hpx::lcos::local::receive_channel< T >
friend hpx::lcos::local::send_channel< T >
```

```
template<>
class one_element_channel<void> : protected hpx::lcos::local::detail::channel_base<void>
```

Public Types

```
typedef void value_type
```

Public Functions

```
one_element_channel()
```

Private Types

```
typedef detail::channel_base<void> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< void >
friend hpx::lcos::local::receive_channel< void >
friend hpx::lcos::local::send_channel< void >

template<typename T>
class receive_channel : protected hpx::lcos::local::detail::channel_base<T>
```

Public Functions

```
receive_channel(channel<T> const &c)
receive_channel(one_element_channel<T> const &c)
```

Private Types

```
typedef detail::channel_base<T> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< T >
friend hpx::lcos::local::send_channel< T >

template<>
class receive_channel<void> : protected hpx::lcos::local::detail::channel_base<void>
```

Public Functions

```
receive_channel(channel<void> const &c)
receive_channel(one_element_channel<void> const &c)
```

Private Types

```
typedef detail::channel_base<void> base_type
```

Friends

```
friend hpx::lcos::local::channel_iterator< void >
friend hpx::lcos::local::send_channel< void >

template<typename T>
class send_channel : private hpx::lcos::local::detail::channel_base<T>
```

Public Functions

```
send_channel (channel<T> const &c)
send_channel (one_element_channel<T> const &c)
```

Private Types

```
typedef detail::channel_base<T> base_type

template<>
class send_channel<void> : private hpx::lcos::local::detail::channel_base<void>
```

Public Functions

```
send_channel (channel<void> const &c)
send_channel (one_element_channel<void> const &c)
```

Private Types

```
typedef detail::channel_base<void> base_type
```

```
#include <hpx/local_lcos/packaged_task.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
template<typename R, typename ...Ts>
class packaged_task<R (Ts...) >
```

Public Functions

```
packaged_task()

template<typename F, typename FD = typename std::decay<F>::type, typename Enable = typename std::enable_if<std::is_invocable<FD, Args...>::value, void> F> packaged_task(F &&f)

template<typename Allocator, typename F, typename FD = typename std::decay<F>::type, typename Enable = typename std::enable_if<std::is_invocable<FD, Args...>::value, void> F> packaged_task(Allocator const &a, F &&f)

packaged_task(packaged_task &&rhs)

packaged_task &operator=(packaged_task &&rhs)

void swap(packaged_task &rhs)

void operator()(Ts... vs)

lcos::future<R> get_future(error_code &ec = throws)

bool valid() const

void reset(error_code &ec = throws)

void set_exception(std::exception_ptr const &e)
```

Private Types

```
typedef util::unique_function_nonscoped<R (Ts...)> function_type
```

Private Functions

```
template<typename ...Vs>
void invoke_impl(std::false_type, Vs&&... vs)

template<typename ...Vs>
void invoke_impl(std::true_type, Vs&&... vs)
```

Private Members

```
function_type function_
local::promise<R> promise_
```

```
#include <hpx/local_lcos/receive_buffer.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
template<typename T, typename Mutex = lcos::local::spinlock>
struct receive_buffer
```

Public Functions

```
receive_buffer()
receive_buffer(receive_buffer &&other)
~receive_buffer()
receive_buffer &operator=(receive_buffer &&other)
hpx::future<T> receive(std::size_t step)
bool try_receive(std::size_t step, hpx::future<T> *f = nullptr)

template<typename Lock = hpx::lcos::local::no_mutex>
void store_received(std::size_t step, T &&val, Lock *lock = nullptr)
bool empty() const
std::size_t cancel_waiting(std::exception_ptr const &e, bool force_delete_entries =
false)
```

Protected Types

```
typedef Mutex mutex_type
typedef hpx::lcos::local::promise<T> buffer_promise_type
typedef std::map<std::size_t, std::shared_ptr<entry_data>> buffer_map_type
typedef buffer_map_type::iterator iterator
```

Protected Functions

```
iterator get_buffer_entry(std::size_t step)
```

Private Members

```
mutex_type mtx_
buffer_map_type buffer_map_
struct entry_data
```

Public Functions

```
template<>
HPX_NON_COPYABLE (entry_data)

template<>
entry_data ()

template<>
hpx::future<T> get_future ()

template<typename Val>
void set_value (Val &&val)

template<>
bool cancel (std::exception_ptr const &e)
```

Public Members

```
template<>
buffer_promise_type promise_

template<>
bool can_be_deleted_

template<>
bool value_set_

struct erase_on_exit
```

Public Functions

```
template<>
erase_on_exit (buffer_map_type &buffer_map, iterator it)

template<>
~erase_on_exit ()
```

Public Members

```
template<>
buffer_map_type &buffer_map_

template<>
iterator it_

template<typename Mutex>
struct receive_buffer<void, Mutex>
```

Public Functions

```
receive_buffer()
receive_buffer(receive_buffer &&other)
~receive_buffer()
receive_buffer &operator=(receive_buffer &&other)
hpx::future<void> receive(std::size_t step)
bool try_receive(std::size_t step, hpx::future<void> *f = nullptr)
template<typename Lock = hpx::lcos::local::no_mutex>
void store_received(std::size_t step, Lock *lock = nullptr)
bool empty() const
std::size_t cancel_waiting(std::exception_ptr const &e, bool force_delete_entries =
false)
```

Protected Types

```
typedef Mutex mutex_type
typedef hpx::lcos::local::promise<void> buffer_promise_type
typedef std::map<std::size_t, std::shared_ptr<entry_data>> buffer_map_type
typedef buffer_map_type::iterator iterator
```

Protected Functions

```
iterator get_buffer_entry(std::size_t step)
```

Private Members

```
mutex_type mtx_
buffer_map_type buffer_map_
template<>
struct entry_data
```

Public Functions

```
template<>
HPX_NON_COPYABLE(entry_data)
template<>
entry_data()
template<>
hpx::future<void> get_future()
```

```
template<>
void set_value()

template<>
bool cancel (std::exception_ptr const &e)
```

Public Members

```
template<>
buffer.promise_type promise_

template<>
bool can_be_deleted_

template<>
bool value_set_

template<>
struct erase_on_exit
```

Public Functions

```
template<>
erase_on_exit (buffer_map_type &buffer_map, iterator it)

template<>
~erase_on_exit ()
```

Public Members

```
template<>
buffer_map_type &buffer_map_

template<>
iterator it_
```

```
#include <hpx/local_lcos/trigger.hpp>

namespace hpx

    namespace lcos

        namespace local

            template<typename Mutex = lcos::local::spinlock>
            struct base_trigger
```

Public Functions

```
base_trigger()
base_trigger(base_trigger &&rhs)
base_trigger &operator=(base_trigger &&rhs)
future<void> get_future(std::size_t *generation_value = nullptr, error_code &ec =
                           hpx::throws)
    get a future allowing to wait for the trigger to fire
bool set(error_code &ec = throws)
    Trigger this object.
void synchronize(std::size_t generation_value, char const *function_name =
                  "base_and_gate<>::synchronize", error_code &ec = throws)
    Wait for the generational counter to reach the requested stage.
std::size_t next_generation()
std::size_t generation() const
```

Protected Types

```
typedef Mutex mutex_type
```

Protected Functions

```
bool trigger_conditions(error_code &ec = throws)
template<typename Lock>
void synchronize(std::size_t generation_value, Lock &l, char const *function_name =
                  "base_and_gate<>::synchronize", error_code &ec = throws)
```

Private Types

```
typedef std::list<conditional_trigger *> condition_list_type
```

Private Functions

```
bool test_condition(std::size_t generation_value)
```

Private Members

```
mutex_type mtx_
lcos::local::promise<void> promise_
std::size_t generation_
condition_list_type conditions_
struct manage_condition
```

Public Functions

```
template<>
manage_condition(base_trigger &gate, conditional_trigger &cond)

template<>
~manage_condition()

template<typename Condition>
future<void> get_future(Condition &&func, error_code &ec = hpx::throws)
```

Public Members

```
template<>
base_trigger &this_

template<>
condition_list_type::iterator it_

struct trigger : public hpx::lcos::local::base_trigger<no_mutex>
```

Public Functions

```
trigger()

trigger(trigger &&rhs)

trigger &operator=(trigger &&rhs)

template<typename Lock>
void synchronize(std::size_t generation_value, Lock &l, char const *function_name =
    "trigger::synchronize", error_code &ec = throws)
```

Private Types

```
typedef base_trigger<no_mutex> base_type

#include <hpx/local_lcos/force_linking.hpp>

namespace hpx

namespace local_lcos
```

Functions

```
force_linking_helper &force_linking()
```

```
struct force_linking_helper
```

Public Members

```
void (*free)(lcos::local::detail::guard_task *task)
```

```
#include <hpx/local_lcos/composable_guard.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

Functions

```
void run_guarded(guard &guard, detail::guard_function task)
```

Conceptually, a guard acts like a mutex on an asynchronous task. The mutex is locked before the task runs, and unlocked afterwards.

```
template<typename F, typename ...Args>
```

```
void run_guarded(guard &guard, F &&f, Args&&... args)
```

```
void run_guarded(guard_set &guards, detail::guard_function task)
```

Conceptually, a `guard_set` acts like a set of mutexes on an asynchronous task. The mutexes are locked before the task runs, and unlocked afterwards.

```
template<typename F, typename ...Args>
```

```
void run_guarded(guard_set &guards, F &&f, Args&&... args)
```

```
class guard : public hpx::lcos::local::detail::debug_object
```

Public Functions

```
guard()
```

```
~guard()
```

Public Members

detail::guard_atomic **task**

```
class guard_set : public hpx::lcos::local::detail::debug_object
```

Public Functions

guard_set ()

~guard_set ()

std::shared_ptr<*guard*> **get** (*std*::size_t *i*)

void add (*std*::shared_ptr<*guard*> **const** &*guard_ptr*)

std::size_t **size** ()

Private Functions

void sort ()

Private Members

std::vector<*std*::shared_ptr<*guard*>> **guards**

bool sorted

Friends

void run_guarded (*guard_set* &*guards*, detail::guard_function *task*)

Conceptually, a *guard_set* acts like a set of mutexes on an asynchronous task. The mutexes are locked before the task runs, and unlocked afterwards.

```
#include <hpx/local_lcos/and_gate.hpp>

namespace hpx

namespace lcos

namespace local

struct and_gate : public hpx::lcos::local::base_and_gate<no_mutex>
```

Public Functions

```

and_gate (std::size_t count = 0)

and_gate (and_gate &&rhs)

and_gate &operator= (and_gate &&rhs)

template<typename Lock>
future<void> get_future (Lock &l, std::size_t count = std::size_t(-1), std::size_t *generation_value = nullptr, error_code &ec = hpx::throws)

template<typename Lock>
shared_future<void> get_shared_future (Lock &l, std::size_t count = std::size_t(-1),
                                         std::size_t *generation_value = nullptr, error_code &ec = hpx::throws)

template<typename Lock>
bool set (std::size_t which, Lock &l, error_code &ec = throws)

template<typename Lock>
void synchronize (std::size_t generation_value, Lock &l, char const *function_name =
                  "and_gate::synchronize", error_code &ec = throws)

```

Private Types

```

typedef base_and_gate<no_mutex> base_type

template<typename Mutex = lcos::local::spinlock>
struct base_and_gate

```

Public Functions

```

base_and_gate (std::size_t count = 0)
    This constructor initializes the base_and_gate object from the the number of participants to
    synchronize the control flow with.

base_and_gate (base_and_gate &&rhs)

base_and_gate &operator= (base_and_gate &&rhs)

future<void> get_future (std::size_t count = std::size_t(-1), std::size_t *generation_value
                        = nullptr, error_code &ec = hpx::throws)

shared_future<void> get_shared_future (std::size_t count = std::size_t(-1), std::size_t
                                         *generation_value = nullptr, error_code &ec
                                         = hpx::throws)

bool set (std::size_t which, error_code &ec = throws)

void synchronize (std::size_t generation_value, char const *function_name =
                  "base_and_gate<>::synchronize", error_code &ec = throws)
    Wait for the generational counter to reach the requested stage.

std::size_t next_generation ()

std::size_t generation () const

```

Protected Types

```
typedef Mutex mutex_type
```

Protected Functions

```
bool trigger_conditions(error_code &ec = throws)
```

```
template<typename OuterLock>
future<void> get_future(OuterLock &outer_lock, std::size_t count = std::size_t(-1),
                       std::size_t *generation_value = nullptr, error_code &ec =
                           hpx::throws)
```

get a future allowing to wait for the gate to fire

```
template<typename OuterLock>
shared_future<void> get_shared_future(OuterLock &outer_lock, std::size_t count =
                                         std::size_t(-1), std::size_t *generation_value =
                                         nullptr, error_code &ec = hpx::throws)
```

get a shared future allowing to wait for the gate to fire

```
template<typename OuterLock>
bool set(std::size_t which, OuterLock &outer_lock, error_code &ec = throws)
    Set the data which has to go into the segment which.
```

```
bool test_condition(std::size_t generation_value)
```

```
template<typename Lock>
void synchronize(std::size_t generation_value, Lock &l, char const *function_name =
                 "base_and_gate<>::synchronize", error_code &ec = throws)
```

```
template<typename OuterLock, typename Lock>
void init_locked(OuterLock &outer_lock, Lock &l, std::size_t count, error_code &ec =
                  throws)
```

Private Types

```
typedef std::list<conditional_trigger *> condition_list_type
```

Private Members

```
mutex_type mtx_
boost::dynamic_bitset received_segments_
lcos::local::promise<void> promise_
std::size_t generation_
condition_list_type conditions_
struct manage_condition
```

Public Functions

```
template<>
manage_condition(base_and_gate &gate, conditional_trigger &cond)

template<>
~manage_condition()

template<typename Condition>
future<void> get_future(Condition &&func, error_code &ec = hpx::throws)
```

Public Members

```
template<>
base_and_gate &this_

template<>
condition_list_type::iterator it_
```

```
#include <hpx/local_lcos/conditional_trigger.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
struct conditional_trigger
```

Public Functions

```
conditional_trigger()
conditional_trigger(conditional_trigger &&rhs)
conditional_trigger &operator=(conditional_trigger &&rhs)

template<typename Condition>
future<void> get_future(Condition &&func, error_code &ec = hpx::throws)
    get a future allowing to wait for the trigger to fire

void reset()

bool set(error_code &ec = throws)
    Trigger this object.
```

Private Members

```
lcos::local::promise<void> promise_
util::function_nonser<bool ()> cond_
```

```
#include <hpx/local_lcos/promise.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

Functions

```
template<typename R>
void swap(promise<R> &x, promise<R> &y)

template<typename R>
class promise : public hpx::lcos::local::detail::promise_base<R>
```

Public Functions

```
promise()

template<typename Allocator>
promise(std::allocator_arg_t, Allocator const &a)

promise(promise &&other)

~promise()

promise &operator=(promise &&other)

void swap(promise &other)

bool valid() const

void set_value(R const &r)

void set_value(R &&r)

template<typename ...Ts>
void set_value(Ts&&... ts)

void set_exception(std::exception_ptr e)
```

Private Types

```
typedef detail::promise_base<R> base_type

template<typename R>
class promise<R&> : public hpx::lcos::local::detail::promise_base<R&>
```

Public Functions

```
promise()

template<typename Allocator>
promise(std::allocator_arg_t, Allocator const &a)

promise(promise &&other)

~promise()

promise &operator=(promise &&other)

void swap(promise &other)

bool valid() const

void set_value(R &r)

void set_exception(std::exception_ptr e)
```

Private Types

```
typedef detail::promise_base<void> base_type

template<>
class promise<void> : public hpx::lcos::local::detail::promise_base<void>
```

Public Functions

```
promise()

template<typename Allocator>
promise(std::allocator_arg_t, Allocator const &a)

promise(promise &&other)

~promise()

promise &operator=(promise &&other)

void swap(promise &other)

bool valid() const

void set_value()

void set_exception(std::exception_ptr e)
```

Private Types

```
typedef detail::promise_base<void> base_type
```

threadmanager

```
#include <compatibility/hpx/runtime/threads/threadmanager.hpp>

#include <hpx/threadmanager.hpp>

namespace hpx

namespace threads

class threadmanager
    #include <threadmanager.hpp> The thread-manager class is the central instance of management for
    all (non-depleted) threads
```

Public Types

```
typedef threads::policies::callback_notifier notification_policy_type
typedef std::unique_ptr<thread_pool_base> pool_type
typedef threads::policies::scheduler_base scheduler_type
typedef std::vector<pool_type> pool_vector
```

Public Functions

```
threadmanager(notification_policy_type &notifier, detail::network_background_callback_type
              network_background_callback = detail::network_background_callback_type())

~threadmanager()

void init()

void create_pools()

void print_pools(std::ostream&)
    FIXME move to private and add hpx:printpools cmd line option.

thread_pool_base &default_pool() const
scheduler_type &default_scheduler() const
thread_pool_base &get_pool(std::string const &pool_name) const
thread_pool_base &get_pool(pool_id_type const &pool_id) const
thread_pool_base &get_pool(std::size_t thread_index) const
```

```
void register_work(thread_init_data &data, thread_state_enum initial_state = pending, error_code &ec = throws)
```

The function *register_work* adds a new work item to the thread manager. It doesn't immediately create a new *thread*, it just adds the task parameters (function, initial state and description) to the internal management data structures. The thread itself will be created when the number of existing threads drops below the number of threads specified by the constructors `max_count` parameter.

Parameters

- `func`: [in] The function or function object to execute as the thread's function. This must have a signature as defined by `thread_function_type`.
- `description`: [in] The value of this parameter allows to specify a description of the thread to create. This information is used for logging purposes mainly, but might be useful for debugging as well. This parameter is optional and defaults to an empty string.
- `initial_state`: [in] The value of this parameter defines the initial state of the newly created *thread*. This must be one of the values as defined by the `thread_state` enumeration (`thread_state::pending`, or `thread_state::suspended`, any other value will throw a `hpx::bad_parameter` exception).

```
void register_thread(thread_init_data &data, thread_id_type &id, thread_state_enum initial_state = pending, bool run_now = true, error_code &ec = throws)
```

The function *register_thread* adds a new work item to the thread manager. It creates a new *thread*, adds it to the internal management data structures, and schedules the new thread, if appropriate.

Parameters

- `func`: [in] The function or function object to execute as the thread's function. This must have a signature as defined by `thread_function_type`.
- `id`: [out] This parameter will hold the id of the created thread. This id is guaranteed to be validly initialized before the thread function is executed.
- `description`: [in] The value of this parameter allows to specify a description of the thread to create. This information is used for logging purposes mainly, but might be useful for debugging as well. This parameter is optional and defaults to an empty string.
- `initial_state`: [in] The value of this parameter defines the initial state of the newly created *thread*. This must be one of the values as defined by the `thread_state` enumeration (`thread_state::pending`, or `thread_state::suspended`, any other value will throw a `hpx::bad_parameter` exception).
- `run_now`: [in] If this parameter is `true` and the initial state is given as `thread_state::pending` the thread will be run immediately, otherwise it will be scheduled to run later (either this function is called for another thread using `true` for the parameter `run_now` or the function `threadmanager::do_some_work` is called). This parameter is optional and defaults to `true`.

```
bool run()
```

Run the thread manager's work queue. This function instantiates the specified number of OS threads in each pool. All OS threads are started to execute the function *tfunc*.

Return The function returns `true` if the thread manager has been started successfully, otherwise it returns `false`.

```
void stop(bool blocking = true)
```

Forcefully stop the thread-manager.

Parameters

- `blocking`:

```
void suspend()
```

```
void resume()
```

```
state status () const
    Return whether the thread manager is still running This returns the “minimal state”, i.e. the state
    of the least advanced thread pool.

std::int64_t get_thread_count (thread_state_enum state = unknown, thread_priority pri-
ority = thread_priority_default, std::size_t num_thread =
std::size_t(-1), bool reset = false)
    return the number of HPX-threads with the given state

Note This function lock the internal OS lock in the thread manager

std::int64_t get_background_thread_count ()

bool enumerate_threads (util::function_nonser<bool> thread_id_type
> const &f, thread_state_enum state = unknown const

void abort_all_suspended_threads ()

bool cleanup_terminated (bool delete_all)

std::size_t get_os_thread_count () const
    Return the number of OS threads running in this thread-manager.

    This function will return correct results only if the thread-manager is running.

std::thread &get_os_thread_handle (std::size_t num_thread) const

void report_error (std::size_t num_thread, std::exception_ptr const &e)
    API functions forwarding to notification policy.

    This notifies the thread manager that the passed exception has been raised. The exception will be
    routed through the notifier and the scheduler (which will result in it being passed to the runtime
    object, which in turn will report it to the console, etc.).

mask_type get_used_processing_units () const
    Returns the mask identifying all processing units used by this thread manager.

hwloc_bitmap_ptr get_pool numa_bitmap (const std::string &pool_name) const

void set_scheduler_mode (threads::policies::scheduler_mode mode)

void add_scheduler_mode (threads::policies::scheduler_mode mode)

void add_remove_scheduler_mode (threads::policies::scheduler_mode
to_add_mode, threads::policies::scheduler_mode
to_remove_mode)

void remove_scheduler_mode (threads::policies::scheduler_mode mode)

void reset_thread_distribution ()

void init_tss (std::size_t global_thread_num)

void deinit_tss ()

std::size_t shrink_pool (std::string const &pool_name)

std::size_t expand_pool (std::string const &pool_name)

std::int64_t get_queue_length (bool reset)
```

```
std::int64_t get_cumulative_duration(bool reset)
std::int64_t get_thread_count_unknown(bool reset)
std::int64_t get_thread_count_active(bool reset)
std::int64_t get_thread_count_pending(bool reset)
std::int64_t get_thread_count_suspended(bool reset)
std::int64_t get_thread_count_terminated(bool reset)
std::int64_t get_thread_count_staged(bool reset)
```

Private Types

```
typedef std::mutex mutex_type
```

Private Members

```
mutex_type mtx_
std::size_t num_threads_
std::vector<pool_id_type> threads_lookup_
pool_vector pools_
notification_policy_type &notifier_
detail::network_background_callback_type network_background_callback_
```

format

```
#include <compatibility/hpx/util/format.hpp>
#include <hpx/format.hpp>
```

Defines

```
DECL_TYPE_SPECIFIER(Type, Spec)
```

```
HPX_FORMAT_EXPORT
```

```
namespace hpx
```

```
namespace util
```

Functions

```
template<typename ...Args>
std::string format (boost::string_ref format_str, Args const&... args)

template<typename ...Args>
std::ostream &format_to (std::ostream &os, boost::string_ref format_str, Args const&... args)
```

```
#include <hpx/util/from_string.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
template<typename T>
T from_string (std::string const &v)

template<typename T, typename U>
T from_string (std::string const &v, U &&default_value)
```

```
#include <hpx/util/to_string.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
template<typename T>
std::string to_string (T const &v)
```

```
#include <hpx/util/bad_lexical_cast.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        class bad_lexical_cast : public bad_cast
```

Public Functions

```
bad_lexical_cast()

const char *what() const

virtual ~bad_lexical_cast()

bad_lexical_cast(std::type_info const &source_type_arg, std::type_info const &target_type_arg)

std::type_info const &source_type() const

std::type_info const &target_type() const
```

Private Members

```
std::type_info const *source

std::type_info const *target
```

collectives

```
#include <compatibility/hpx/lcos/broadcast.hpp>

#include <compatibility/hpx/lcos/reduce.hpp>

#include <compatibility/hpx/lcos/barrier.hpp>

#include <compatibility/hpx/lcos/gather.hpp>

#include <compatibility/hpx/lcos/fold.hpp>

#include <compatibility/hpx/lcos/latch.hpp>

#include <compatibility/hpx/lcos/spmd_block.hpp>

#include <hpx/collectives/broadcast.hpp>

namespace hpx

namespace lcos
```

Functions

```
template<typename Action, typename ArgN, ...>hpx::future<std::vector<decltype(Action...)>>
```

Perform a distributed broadcast operation.

The function `hpx::lcos::broadcast` performs a distributed broadcast operation resulting in action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The given action is invoked asynchronously on all given identifiers, and the arguments `ArgN` are passed along to those invocations.

Return This function returns a future representing the result of the overall reduction operation.

Note If `decltype(Action(...))` is void, then the result of this function is `future<void>`.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
template<typename Action, typename ArgN, ...>void hpx::lcos::broadcast_apply(std::vector<decltype(Action...)>)
```

Perform an asynchronous (fire&forget) distributed broadcast operation.

The function `hpx::lcos::broadcast_apply` performs an asynchronous (fire&forget) distributed broadcast operation resulting in action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The given action is invoked asynchronously on all given identifiers, and the arguments `ArgN` are passed along to those invocations.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
template<typename Action, typename ArgN, ...>hpx::future< std::vector<decltype(Action...)>>
```

Perform a distributed broadcast operation.

The function `hpx::lcos::broadcast_with_index` performs a distributed broadcast operation resulting in action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The given action is invoked asynchronously on all given identifiers, and the arguments `ArgN` are passed along to those invocations.

The function passes the index of the global identifier in the given list of identifiers as the last argument to the action.

Return This function returns a future representing the result of the overall reduction operation.

Note If `decltype(Action(...))` is void, then the result of this function is `future<void>`.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
template<typename Action, typename ArgN, ...>void hpx::lcos::broadcast_apply_with_index
```

Perform an asynchronous (fire&forget) distributed broadcast operation.

The function `hpx::lcos::broadcast_apply_with_index` performs an asynchronous (fire&forget) distributed broadcast operation resulting in action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The given action is invoked asynchronously on all given identifiers, and the arguments `ArgN` are passed along to those invocations.

The function passes the index of the global identifier in the given list of identifiers as the last argument to the action.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
#include <hpx/collectives/reduce.hpp>

namespace hpx
```

```
namespace lcos
```

Functions

```
template<typename Action, typename ReduceOp, typename ArgN, ...>hpx::future<decay_type
```

Perform a distributed reduction operation.

The function `hpx::lcos::reduce` performs a distributed reduction operation over results returned from action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

Return This function returns a future representing the result of the overall reduction operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `reduce_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the reduction operation performed on its arguments.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
template<typename Action, typename ReduceOp, typename ArgN, ...>hpx::future<decay_type
```

Perform a distributed reduction operation.

The function `hpx::lcos::reduce_with_index` performs a distributed reduction operation over results returned from action invocations on a given set of global identifiers. The action can be either plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The function passes the index of the global identifier in the given list of identifiers as the last argument to the action.

Return This function returns a future representing the result of the overall reduction operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `reduce_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the reduction operation performed on its arguments.
- `argN`: [in] Any number of arbitrary arguments (passed by const reference) which will be forwarded to the action invocation.

```
#include <hpx/collectives/barrier.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

class barrier

`#include <barrier.hpp>` The barrier is an implementation performing a barrier over a number of participating threads. The different threads don't have to be on the same locality. This barrier can be invoked in a distributed application.

For a local only barrier

See `hpx::lcos::local::barrier`.

Public Functions

`barrier (std::string const &base_name)`

Creates a barrier, rank is locality id, size is number of localities

A barrier `base_name` is created. It expects that `hpx::get_num_localities()` participate and the local rank is `hpx::get_locality_id()`.

Parameters

- `base_name`: The name of the barrier

`barrier (std::string const &base_name, std::size_t num)`

Creates a barrier with a given size, rank is locality id

A barrier `base_name` is created. It expects that `num` participate and the local rank is `hpx::get_locality_id()`.

Parameters

- `base_name`: The name of the barrier
- `num`: The number of participating threads

`barrier (std::string const &base_name, std::size_t num, std::size_t rank)`

Creates a barrier with a given size and rank

A barrier `base_name` is created. It expects that `num` participate and the local rank is `rank`.

Parameters

- `base_name`: The name of the barrier
- `num`: The number of participating threads
- `rank`: The rank of the calling site for this invocation

barrier (*std::string const &base_name*, *std::vector<std::size_t> const &ranks*, *std::size_t rank*)

Creates a barrier with a vector of ranks

A barrier *base_name* is created. It expects that *ranks.size()* and the local rank is *rank* (must be contained in *ranks*).

Parameters

- *base_name*: The name of the barrier
- *ranks*: Gives a list of participating ranks (this could be derived from a list of locality ids)
- *rank*: The rank of the calling site for this invocation

void wait()

Wait until each participant entered the barrier. Must be called by all participants

Return This function returns once all participants have entered the barrier (have called *wait*).

hpx::future<void> wait (hpx::launch::async_policy)

Wait until each participant entered the barrier. Must be called by all participants

Return a future that becomes ready once all participants have entered the barrier (have called *wait*).

Public Static Functions

static void synchronize()

Perform a global synchronization using the default global barrier. The barrier is created once at startup and can be reused throughout the lifetime of an HPX application.

Note This function currently does not support dynamic connection and disconnection of localities.

```
#include <hpx/collectives/gather.hpp>
```

Defines

HPX_REGISTER_GATHER_DECLARATION (*type, name*)

Declare a gather object named *name* for a given data type *type*.

The macro *HPX_REGISTER_GATHER_DECLARATION* can be used to declare all facilities necessary for a (possibly remote) gather operation.

The parameter *type* specifies for which data type the gather operations should be enabled.

The (optional) parameter *name* should be a unique C-style identifier which will be internally used to identify a particular gather operation. If this defaults to *<type>_gather* if not specified.

Note The macro *HPX_REGISTER_GATHER_DECLARATION* can be used with 1 or 2 arguments. The second argument is optional and defaults to *<type>_gather*.

HPX_REGISTER_GATHER (*type, name*)

Define a gather object named *name* for a given data type *type*.

The macro *HPX_REGISTER_GATHER* can be used to define all facilities necessary for a (possibly remote) gather operation.

The parameter *type* specifies for which data type the gather operations should be enabled.

The (optional) parameter *name* should be a unique C-style identifier which will be internally used to identify a particular gather operation. If this defaults to *<type>_gather* if not specified.

Note The macro *HPX_REGISTER_GATHER* can be used with 1 or 2 arguments. The second argument is optional and defaults to *<type>_gather*.

namespace hpx

namespace lcos

Functions

```
template<typename T>
hpx::future<std::vector<T>> gather_here (char const *basename, hpx::future<T> result,
                                              std::size_t num_sites = std::size_t(-1), std::size_t
                                              generation = std::size_t(-1), std::size_t this_site =
                                              std::size_t(-1))
```

Gather a set of values from different call sites

This function receives a set of values from all call sites operating on the given base name.

Note Each gather operation has to be accompanied with a unique usage of the *HPX_REGISTER_GATHER* macro to define the necessary internal facilities used by *gather_here* and *gather_there*

Return This function returns a future holding a vector with all gathered values. It will become ready once the gather operation has been completed.

Parameters

- *basename*: The base name identifying the gather operation
- *result*: A future referring to the value to transmit to the central gather point from this call site.
- *num_sites*: The number of participating sites (default: all localities).
- *generation*: The generational counter identifying the sequence number of the gather operation performed on the given base name. This is optional and needs to be supplied only if the gather operation on the given base name has to be performed more than once.
- *this_site*: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever *hpx::get_locality_id()* returns.

```
template<typename T>
hpx::future<void> gather_there (char const *basename, hpx::future<T> result, std::size_t
                                 generation = std::size_t(-1), std::size_t root_site = 0, std::size_t
                                 this_site = std::size_t(-1))
```

Gather a given value at the given call site

This function transmits the value given by *result* to a central gather site (where the corresponding *gather_here* is executed)

Note Each gather operation has to be accompanied with a unique usage of the *HPX_REGISTER_GATHER* macro to define the necessary internal facilities used by *gather_here* and *gather_there*

Return This function returns a future which will become ready once the gather operation has been completed.

Parameters

- *basename*: The base name identifying the gather operation
- *result*: A future referring to the value to transmit to the central gather point from this call site.

- **generation:** The generational counter identifying the sequence number of the gather operation performed on the given base name. This is optional and needs to be supplied only if the gather operation on the given base name has to be performed more than once.
- **root_site:** The sequence number of the central gather point (usually the locality id). This value is optional and defaults to 0.
- **this_site:** The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever hpx::get_locality_id() returns.

```
template<typename T>
hpx::future<std::vector<typename std::decay<T>::type>> gather_here (char const *base-
name, T &&result,
std::size_t num_sites =
= std::size_t(-1),
std::size_t generation
= std::size_t(-1),
std::size_t this_site =
std::size_t(-1))
```

Gather a set of values from different call sites

This function receives a set of values from all call sites operating on the given base name.

Note Each gather operation has to be accompanied with a unique usage of the *HPX_REGISTER_GATHER* macro to define the necessary internal facilities used by *gather_here* and *gather_there*

Return This function returns a future holding a vector with all gathered values. It will become ready once the gather operation has been completed.

Parameters

- **basename:** The base name identifying the gather operation
- **result:** The value to transmit to the central gather point from this call site.
- **num_sites:** The number of participating sites (default: all localities).
- **generation:** The generational counter identifying the sequence number of the gather operation performed on the given base name. This is optional and needs to be supplied only if the gather operation on the given base name has to be performed more than once.
- **this_site:** The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever hpx::get_locality_id() returns.

```
template<typename T>
hpx::future<void> gather_there (char const *basename, T &&result, std::size_t generation =
std::size_t(-1), std::size_t root_site = 0, std::size_t this_site =
std::size_t(-1))
```

Gather a given value at the given call site

This function transmits the value given by *result* to a central gather site (where the corresponding *gather_here* is executed)

Note Each gather operation has to be accompanied with a unique usage of the *HPX_REGISTER_GATHER* macro to define the necessary internal facilities used by *gather_here* and *gather_there*

Return This function returns a future which will become ready once the gather operation has been completed.

Parameters

- **basename:** The base name identifying the gather operation
- **result:** The value to transmit to the central gather point from this call site.
- **generation:** The generational counter identifying the sequence number of the gather operation performed on the given base name. This is optional and needs to be supplied only if the gather operation on the given base name has to be performed more than once.

- `root_site`: The sequence number of the central gather point (usually the locality id). This value is optional and defaults to 0.
- `this_site`: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever `hpx::get_locality_id()` returns.

```
#include <hpx/collectives/all_reduce.hpp>
```

Defines

HPX_REGISTER_ALLREDUCE_DECLARATION (`type, name`)

Declare a `all_reduce` object named `name` for a given data type `type`.

The macro `HPX_REGISTER_ALLREDUCE_DECLARATION` can be used to declare all facilities necessary for a (possibly remote) `all_reduce` operation.

The parameter `type` specifies for which data type the `all_reduce` operations should be enabled.

The (optional) parameter `name` should be a unique C-style identifier that will be internally used to identify a particular `all_reduce` operation. If this defaults to `<type>_all_reduce` if not specified.

Note The macro `HPX_REGISTER_ALLREDUCE_DECLARATION` can be used with 1 or 2 arguments. The second argument is optional and defaults to `<type>_all_reduce`.

HPX_REGISTER_ALLREDUCE (`type, name`)

Define a `all_reduce` object named `name` for a given data type `type`.

The macro `HPX_REGISTER_ALLREDUCE` can be used to define all facilities necessary for a (possibly remote) `all_reduce` operation.

The parameter `type` specifies for which data type the `all_reduce` operations should be enabled.

The (optional) parameter `name` should be a unique C-style identifier that will be internally used to identify a particular `all_reduce` operation. If this defaults to `<type>_all_reduce` if not specified.

Note The macro `HPX_REGISTER_ALLREDUCE` can be used with 1 or 2 arguments. The second argument is optional and defaults to `<type>_all_reduce`.

```
namespace hpx
```

```
namespace lcos
```

Functions

```
template<typename T, typename F>
hpx::future<T> all_reduce (char const *basename, hpx::future<T> result, F && op, std::size_t
num_sites = std::size_t(-1), std::size_t generation = std::size_t(-1),
std::size_t this_site = std::size_t(-1), std::size_t root_site = 0)
```

AllReduce a set of values from different call sites

This function receives a set of values that are the result of applying a given operator on values supplied from all call sites operating on the given base name.

Note Each `all_reduce` operation has to be accompanied with a unique usage of the `HPX_REGISTER_ALLREDUCE` macro to define the necessary internal facilities used by `all_reduce`.

Return This function returns a future holding a vector with all values send by all participating sites.

It will become ready once the all_reduce operation has been completed.

Parameters

- `basename`: The base name identifying the all_reduce operation
- `local_result`: A future referring to the value to transmit to all participating sites from this call site.
- `op`: Reduction operation to apply to all values supplied from all participating sites
- `num_sites`: The number of participating sites (default: all localities).
- `generation`: The generational counter identifying the sequence number of the all_reduce operation performed on the given base name. This is optional and needs to be supplied only if the all_reduce operation on the given base name has to be performed more than once.
- `this_site`: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever `hpx::get_locality_id()` returns. `root_site` The site that is responsible for creating the all_reduce support object. This value is optional and defaults to '0' (zero).

```
template<typename T, typename F>
hpx::future<std::decay_t<T>> all_reduce(char const *basename, T &&result, F &&op,
                                             std::size_t num_sites = std::size_t(-1), std::size_t
                                             generation = std::size_t(-1), std::size_t this_site =
                                             std::size_t(-1), std::size_t root_site = 0)
```

AllReduce a set of values from different call sites

This function receives a set of values from all call sites operating on the given base name.

Note Each all_reduce operation has to be accompanied with a unique usage of the `HPX_REGISTER_ALLREDUCE` macro to define the necessary internal facilities used by `all_reduce`.

Return This function returns a future holding a vector with all values send by all participating sites.

It will become ready once the all_reduce operation has been completed.

Parameters

- `basename`: The base name identifying the all_reduce operation
- `local_result`: The value to transmit to all participating sites from this call site.
- `op`: Reduction operation to apply to all values supplied from all participating sites
- `num_sites`: The number of participating sites (default: all localities).
- `generation`: The generational counter identifying the sequence number of the all_reduce operation performed on the given base name. This is optional and needs to be supplied only if the all_reduce operation on the given base name has to be performed more than once.
- `this_site`: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever `hpx::get_locality_id()` returns. `root_site` The site that is responsible for creating the all_reduce support object. This value is optional and defaults to '0' (zero).

```
#include <hpx/collectives/fold.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

Functions

```
template<typename Action, typename FoldOp, typename Init, typename ArgN, ...>hpx::f
```

Perform a distributed fold operation.

The function `hpx::lcos::fold` performs a distributed folding operation over results returned from action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

Note The type of the initial value must be convertible to the result type returned from the invoked action.

Return This function returns a future representing the result of the overall folding operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `fold_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the folding operation performed on its arguments.
- `init`: [in] The initial value to be used for the folding operation
- `argN`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the action invocation.

```
template<typename Action, typename FoldOp, typename Init, typename ArgN, ...>hpx::f
```

Perform a distributed folding operation.

The function `hpx::lcos::fold_with_index` performs a distributed folding operation over results returned from action invocations on a given set of global identifiers. The action can be either plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The function passes the index of the global identifier in the given list of identifiers as the last argument to the action.

Note The type of the initial value must be convertible to the result type returned from the invoked action.

Return This function returns a future representing the result of the overall folding operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `fold_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the folding operation performed on its arguments.
- `init`: [in] The initial value to be used for the folding operation
- `argN`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the action invocation.

```
template<typename Action, typename FoldOp, typename Init, typename ArgN, ...>hpx::f
```

Perform a distributed inverse folding operation.

The function `hpx::lcos::inverse_fold` performs an inverse distributed folding operation over results returned from action invocations on a given set of global identifiers. The action can be either a plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

Note The type of the initial value must be convertible to the result type returned from the invoked action.

Return This function returns a future representing the result of the overall folding operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `fold_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the folding operation performed on its arguments.
- `init`: [in] The initial value to be used for the folding operation
- `argN`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the action invocation.

```
template<typename Action, typename FoldOp, typename Init, typename ArgN, ...> hpx::future<...> inverse_fold_with_index(...)
```

Perform a distributed inverse folding operation.

The function `hpx::lcos::inverse_fold_with_index` performs an inverse distributed folding operation over results returned from action invocations on a given set of global identifiers. The action can be either plain action (in which case the global identifiers have to refer to localities) or a component action (in which case the global identifiers have to refer to instances of a component type which exposes the action).

The function passes the index of the global identifier in the given list of identifiers as the last argument to the action.

Note The type of the initial value must be convertible to the result type returned from the invoked action.

Return This function returns a future representing the result of the overall folding operation.

Parameters

- `ids`: [in] A list of global identifiers identifying the target objects for which the given action will be invoked.
- `fold_op`: [in] A binary function expecting two results as returned from the action invocations. The function (or function object) is expected to return the result of the folding operation performed on its arguments.
- `init`: [in] The initial value to be used for the folding operation
- `argN`: [in] Any number of arbitrary arguments (passed by value, by const reference or by rvalue reference) which will be forwarded to the action invocation.

```
#include <hpx/collectives/latch.hpp>

namespace hpx

namespace lcos

class latch : public components::client_base<latch, lcos::server::latch>
```

Public Functions

`latch()`

`latch(std::ptrdiff_t count)`

Initialize the latch

Requires: `count >= 0`. Synchronization: None Postconditions: `counter_ == count`.

latch (*naming*::id_type **const** &*id*)

Extension: Create a client side representation for the existing *server*::*latch* instance with the given global id *id*.

latch (*hpx*::future<*naming*::id_type> &&*f*)

Extension: Create a client side representation for the existing *server*::*latch* instance with the given global id *id*.

latch (*hpx*::shared_future<*naming*::id_type> **const** &*id*)

Extension: Create a client side representation for the existing *server*::*latch* instance with the given global id *id*.

latch (*hpx*::shared_future<*naming*::id_type> &&*id*)**void count_down_and_wait()**

Decrements counter_ by 1. Blocks at the synchronization point until counter_ reaches 0.

Requires: counter_ > 0.

Synchronization: Synchronizes with all calls that block on this latch and with all is_ready calls on this latch that return true.

Exceptions

- Nothing..:

void count_down (*std*::ptrdiff_t *n*)

Decrements counter_ by n. Does not block.

Requires: counter_ >= n and n >= 0.

Synchronization: Synchronizes with all calls that block on this latch and with all is_ready calls on this latch that return true .

Exceptions

- Nothing..:

bool is_ready () const

Returns: counter_ == 0. Does not block.

Exceptions

- Nothing..:

void wait () const

If counter_ is 0, returns immediately. Otherwise, blocks the calling thread at the synchronization point until counter_ reaches 0.

Exceptions

- Nothing..:

Private Types

typedef *components*::client_base<*latch*, *lcos*::server::*latch*> **base_type**

```
#include <hpx/collectives/spmd_block.hpp>

namespace hpx

namespace lcos
```

Functions

```
template<typename F, typename ...Args>
hpx::future<void> define_spmd_block(std::string &&name, std::size_t images_per_locality, F
&&f, Args&&... args)

struct spmd_block
#include <spmd_block.hpp> The class spmd_block defines an interface for launching multiple images while giving handles to each image to interact with the remaining images. The define_spmd_block function templates create multiple images of a user-defined action and launches them in a possibly separate thread. A temporary spmd block object is created and diffused to each image. The constraint for the action given to the define_spmd_block function is to accept a spmd_block as first parameter.
```

Public Functions

```
spmd_block()

spmd_block(std::string const &name, std::size_t images_per_locality, std::size_t
num_images, std::size_t image_id)

std::size_t get_images_per_locality() const

std::size_t get_num_images() const

std::size_t this_image() const

void sync_all() const

hpx::future<void> sync_all(hpx::launch::async_policy const&) const

void sync_images(std::set<std::size_t> const &images) const

void sync_images(std::vector<std::size_t> const &input_images) const

template<typename Iterator>
std::enable_if<traits::is_input_iterator<Iterator>::value>::type sync_images(Iterator begin,
Iterator end)
const

template<typename ...I>
std::enable_if<util::all_of<typename std::is_integral<I>::value...>::value>::type sync_images(I...
i)

hpx::future<void> sync_images(hpx::launch::async_policy const &policy,
std::set<std::size_t> const &images) const

hpx::future<void> sync_images(hpx::launch::async_policy const &policy,
std::vector<std::size_t> const &input_images) const

template<typename Iterator>
```

```
std::enable_if<traits::is_input_iterator<Iterator>::value, hpx::future<void>>::type sync_images (hpx::launch::async
    const
    &policy,
    Iterator begin,
    Iterator end)
const

template<typename ...I>
std::enable_if<util::all_of<ttypename std::is_integral<I>::type...>::value, hpx::future<void>>::type sync_images (
```

Private Types

```
using barrier_type = hpx::lcos::barrier
using table_type = std::map<std::set<std::size_t>, std::shared_ptr<barrier_type>>
```

Private Functions

```
template<typename Archive>
void serialize (Archive&, unsigned)
```

Private Members

```
std::string name_
std::size_t images_per_locality_
std::size_t num_images_
std::size_t image_id_
hpx::util::jenkins_hash hash_
std::shared_ptr<hpx::lcos::barrier> barrier_
table_type barriers_
```

Friends

```
friend hpx::lcos::hpx::serialization::access

#include <hpx/collectives/all_to_all.hpp>
```

Defines

HPX_REGISTER_ALLTOALL_DECLARATION (*type, name*)

Declare a all_to_all object named *name* for a given data type *type*.

The macro *HPX_REGISTER_ALLTOALL_DECLARATION* can be used to declare all facilities necessary for a (possibly remote) all_to_all operation.

The parameter *type* specifies for which data type the all_to_all operations should be enabled.

The (optional) parameter *name* should be a unique C-style identifier that will be internally used to identify a particular all_to_all operation. If this defaults to *<type>_all_to_all* if not specified.

Note The macro *HPX_REGISTER_ALLTOALL_DECLARATION* can be used with 1 or 2 arguments. The second argument is optional and defaults to *<type>_all_to_all*.

HPX_REGISTER_ALLTOALL (*type, name*)

Define a all_to_all object named *name* for a given data type *type*.

The macro *HPX_REGISTER_ALLTOALL* can be used to define all facilities necessary for a (possibly remote) all_to_all operation.

The parameter *type* specifies for which data type the all_to_all operations should be enabled.

The (optional) parameter *name* should be a unique C-style identifier that will be internally used to identify a particular all_to_all operation. If this defaults to *<type>_all_to_all* if not specified.

Note The macro *HPX_REGISTER_ALLTOALL* can be used with 1 or 2 arguments. The second argument is optional and defaults to *<type>_all_to_all*.

```
namespace hpx
```

```
namespace lcos
```

Functions

```
template<typename T>
hpx::future<std::vector<T>> all_to_all (char const *basename, hpx::future<T> result,
                                             std::size_t num_sites = std::size_t(-1), std::size_t
                                             generation = std::size_t(-1), std::size_t this_site =
                                             std::size_t(-1), std::size_t root_site = 0)
```

AllToAll a set of values from different call sites

This function receives a set of values from all call sites operating on the given base name.

Note Each all_to_all operation has to be accompanied with a unique usage of the *HPX_REGISTER_ALLTOALL* macro to define the necessary internal facilities used by *all_to_all*.

Return This function returns a future holding a vector with all values send by all participating sites.

It will become ready once the all_to_all operation has been completed.

Parameters

- `basename`: The base name identifying the all_to_all operation
- `local_result`: A future referring to the value to transmit to all participating sites from this call site.
- `num_sites`: The number of participating sites (default: all localities).
- `generation`: The generational counter identifying the sequence number of the all_to_all operation performed on the given base name. This is optional and needs to be supplied only if the all_to_all operation on the given base name has to be performed more than once.
- `this_site`: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever `hpx::get_locality_id()` returns. `root_site` The site that is responsible for creating the all_to_all support object. This value is optional and defaults to '0' (zero).

```
template<typename T>
hpx::future<std::vector<typename std::decay<T>::type>> all_to_all(char const *base-
name, T &&result,
std::size_t num_sites
= std::size_t(-1),
std::size_t generation
= std::size_t(-1),
std::size_t this_site
= std::size_t(-1),
std::size_t root_site
= 0)
```

AllToAll a set of values from different call sites

This function receives a set of values from all call sites operating on the given base name.

Note Each all_to_all operation has to be accompanied with a unique usage of the `HPX_REGISTER_ALLTOALL` macro to define the necessary internal facilities used by `all_to_all`.

Return This function returns a future holding a vector with all values send by all participating sites.

It will become ready once the all_to_all operation has been completed.

Parameters

- `basename`: The base name identifying the all_to_all operation
- `local_result`: The value to transmit to all participating sites from this call site.
- `num_sites`: The number of participating sites (default: all localities).
- `generation`: The generational counter identifying the sequence number of the all_to_all operation performed on the given base name. This is optional and needs to be supplied only if the all_to_all operation on the given base name has to be performed more than once.
- `this_site`: The sequence number of this invocation (usually the locality id). This value is optional and defaults to whatever `hpx::get_locality_id()` returns. `root_site` The site that is responsible for creating the all_to_all support object. This value is optional and defaults to '0' (zero).

statistics

```
#include <compatibility/hpx/util/max.hpp>

#include <compatibility/hpx/util/rolling_max.hpp>

#include <compatibility/hpx/util/rolling_min.hpp>

#include <compatibility/hpx/util/histogram.hpp>
```

```
#include <compatibility/hpx/util/min.hpp>

#include <hpx/statistics/max.hpp>

namespace hpx

namespace util
```

Functions

```
template<typename T>HPX_HOST_DEVICE constexpr T const&() hpx::util::max(T const &
```



```
#include <hpx/statistics/rolling_max.hpp>

namespace boost

namespace accumulators

namespace extract
```

Variables

```
extractor<tag::rolling_max> const rolling_max = {}

namespace tag
```

```
struct rolling_max : public depends_on<rolling_window>

struct impl
```

```
template<typename Sample, typename Weight>
struct apply
```

Public Types

```
typedef hpx::util::detail::rolling_max_impl<Sample> type
```

```
#include <hpx/statistics/rolling_min.hpp>
```

```
namespace boost
```

```
    namespace accumulators
```

```
        namespace extract
```

Variables

```
extractor<tag::rolling_min> const rolling_min = {}
```

```
namespace tag
```

```
struct rolling_min : public depends_on<rolling_window>
```

```
    struct impl
```

```
        template<typename Sample, typename Weight>
        struct apply
```

Public Types

```
typedef hpx::util::detail::rolling_min_impl<Sample> type
```

```
#include <hpx/statistics/histogram.hpp>
```

```
namespace boost
```

```
    namespace accumulators
```

```
        namespace extract
```

Variables

```
extractor<tag::histogram> const histogram = {}
```

```
namespace tag
```

```
struct histogram : public depends_on<count>, public histogram_num_bins, public histogram_min_range
```

```
    struct impl
```

```
        template<typename Sample, typename Weight>
        struct apply
```

Public Types

```
typedef hpx::util::detail::histogram_impl<Sample> type
```

```
#include <hpx/statistics/min.hpp>
```

```
namespace hpx
```

```
    namespace util
```

Functions

```
template<typename T>HPX_HOST_DEVICE constexpr T const&() hpx::util::min(T const & a,
```

synchronization

```
#include <compatibility/hpx/lcos/local/no_mutex.hpp>
```

```
#include <compatibility/hpx/lcos/local/spinlock.hpp>
```

```
#include <compatibility/hpx/lcos/local/counting_semaphore.hpp>
```

```
#include <compatibility/hpx/lcos/local/sliding_semaphore.hpp>
```

```
#include <compatibility/hpx/lcos/local/shared_mutex.hpp>
```

```
#include <compatibility/hpx/lcos/local/spinlock_no_backoff.hpp>
```

```
#include <compatibility/hpx/lcos/local/barrier.hpp>
```

```
#include <compatibility/hpx/lcos/local/once.hpp>
```

```
#include <compatibility/hpx/lcos/local/latch.hpp>
```

```
#include <compatibility/hpx/lcos/local/condition_variable.hpp>
```

```
#include <compatibility/hpx/lcos/local/event.hpp>
```

```
#include <compatibility/hpx/lcos/local/recursive_mutex.hpp>
```

```
#include <compatibility/hpx/lcos/local/spinlock_pool.hpp>
```

```
#include <compatibility/hpx/lcos/local/mutex.hpp>
```

```
#include <hpx/synchronization/no_mutex.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

```
            struct no_mutex
```

Public Functions

```
    void lock()
```

```
    bool try_lock()
```

```
    void unlock()
```

```
#include <hpx/synchronization/channel_spse.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

```
            template<typename T>
            class channel_spse
```

Public Functions

```
    channel_spse(std::size_t size)
```

```
    channel_spse(channel_spse &&rhs)
```

```
    channel_spse &operator=(channel_spse &&rhs)
```

```
    ~channel_spse()
```

```
    bool get(T *val = nullptr) const
```

```
    bool set(T &&t)
```

```
    std::size_t close()
```

```
    std::size_t capacity() const
```

Private Functions

```
bool is_full (std::size_t tail) const  
bool is_empty (std::size_t head) const
```

Private Members

```
hpx::util::cache_aligned_data<std::atomic<std::size_t>> head_  
hpx::util::cache_aligned_data<std::atomic<std::size_t>> tail_  
std::size_t size_  
std::unique_ptr<T[]> buffer_  
std::atomic<bool> closed_
```

```
#include <hpx/synchronization/spinlock.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
struct spinlock
```

Public Functions

```
HPX_NON_COPYABLE (spinlock)  
spinlock (char const *const desc = "hpx::lcos::local::spinlock")  
~spinlock ()  
void lock ()  
bool try_lock ()  
void unlock ()
```

Private Functions

```
bool acquire_lock ()  
void relinquish_lock ()
```

Private Members

```
std::uint64_t v_
```

```
#include <hpx/synchronization/counting_semaphore.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

TypeDefs

```
typedef counting_semaphore_var counting_semaphore
```

```
template<typename Mutex = hpx::lcos::local::spinlock, int N = 0>
```

```
class counting_semaphore_var
```

#include <counting_semaphore.hpp> A semaphore is a protected variable (an entity storing a value) or abstract data type (an entity grouping several variables that may or may not be numerical) which constitutes the classic method for restricting access to shared resources, such as shared memory, in a multiprogramming environment. Semaphores exist in many variants, though usually the term refers to a counting semaphore, since a binary semaphore is better known as a mutex. A counting semaphore is a counter for a set of available resources, rather than a locked/unlocked flag of a single resource. It was invented by Edsger Dijkstra. Semaphores are the classic solution to preventing race conditions in the dining philosophers problem, although they do not prevent resource deadlocks.

Counting semaphores can be used for synchronizing multiple threads as well: one thread waiting for several other threads to touch (signal) the semaphore, or several threads waiting for one other thread to touch this semaphore.

Public Functions

```
counting_semaphore_var (std::int64_t value = N)
```

Construct a new counting semaphore.

Parameters

- **value:** [in] The initial value of the internal semaphore lock count. Normally this value should be zero (which is the default), values greater than zero are equivalent to the same number of signals pre-set, and negative values are equivalent to the same number of waits pre-set.

```
void wait (std::int64_t count = 1)
```

Wait for the semaphore to be signaled.

Parameters

- **count:** [in] The value by which the internal lock count will be decremented. At the same time this is the minimum value of the lock count at which the thread is not yielded.

```
bool try_wait (std::int64_t count = 1)
```

Try to wait for the semaphore to be signaled.

Return The function returns true if the calling thread was able to acquire the requested amount of credits. The function returns false if not sufficient credits are available at this point in time.

Parameters

- `count`: [in] The value by which the internal lock count will be decremented. At the same time this is the minimum value of the lock count at which the thread is not yielded.

```
void signal (std::int64_t count = 1)
```

Signal the semaphore.

```
std::int64_t signal_all ()
```

Private Types

```
typedef Mutex mutex_type
```

Private Members

```
mutex_type mtx_
```

```
detail::counting_semaphore sem_
```

```
#include <hpx/synchronization/sliding_semaphore.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

Typedefs

```
typedef sliding_semaphore_var sliding_semaphore
```

```
template<typename Mutex = hpx::lcos::local::spinlock>
```

```
class sliding_semaphore_var
```

#include <sliding_semaphore.hpp> A semaphore is a protected variable (an entity storing a value) or abstract data type (an entity grouping several variables that may or may not be numerical) which constitutes the classic method for restricting access to shared resources, such as shared memory, in a multiprogramming environment. Semaphores exist in many variants, though usually the term refers to a counting semaphore, since a binary semaphore is better known as a mutex. A counting semaphore is a counter for a set of available resources, rather than a locked/unlocked flag of a single resource. It was invented by Edsger Dijkstra. Semaphores are the classic solution to preventing race conditions in the dining philosophers problem, although they do not prevent resource deadlocks.

Sliding semaphores can be used for synchronizing multiple threads as well: one thread waiting for several other threads to touch (signal) the semaphore, or several threads waiting for one other thread to touch this semaphore. The difference to a counting semaphore is that a sliding semaphore will not limit the number of threads which are allowed to proceed, but will make sure that the difference between the (arbitrary) number passed to set and wait does not exceed a given threshold.

Public Functions

sliding_semaphore_var (*std::int64_t max_difference, std::int64_t lower_limit = 0*)
Construct a new sliding semaphore.

Parameters

- *max_difference*: [in] The max difference between the upper limit (as set by [wait\(\)](#)) and the lower limit (as set by [signal\(\)](#)) which is allowed without suspending any thread calling [wait\(\)](#).
- *lower_limit*: [in] The initial lower limit.

void set_max_difference (*std::int64_t max_difference, std::int64_t lower_limit = 0*)
Set/Change the difference that will cause the semaphore to trigger.

Parameters

- *max_difference*: [in] The max difference between the upper limit (as set by [wait\(\)](#)) and the lower limit (as set by [signal\(\)](#)) which is allowed without suspending any thread calling [wait\(\)](#).
- *lower_limit*: [in] The initial lower limit.

void wait (*std::int64_t upper_limit*)
Wait for the semaphore to be signaled.

Parameters

- *upper_limit*: [in] The new upper limit. The calling thread will be suspended if the difference between this value and the largest lower_limit which was set by [signal\(\)](#) is larger than the *max_difference*.

bool try_wait (*std::int64_t upper_limit = 1*)
Try to wait for the semaphore to be signaled.

Return The function returns true if the calling thread would not block if it was calling [wait\(\)](#).

Parameters

- *upper_limit*: [in] The new upper limit. The calling thread will be suspended if the difference between this value and the largest lower_limit which was set by [signal\(\)](#) is larger than the *max_difference*.

void signal (*std::int64_t lower_limit*)
Signal the semaphore.

Parameters

- *lower_limit*: [in] The new lower limit. This will update the current lower limit of this semaphore. It will also re-schedule all suspended threads for which their associated upper limit is not larger than the lower limit plus the *max_difference*.

std::int64_t signal_all ()

Private Types

typedef Mutex mutex_type

Private Members

```
mutex_type mtx_
detail::sliding_semaphore sem_
```

```
#include <hpx/synchronization/shared_mutex.hpp>

namespace hpx

namespace lcos

namespace local
```

TypeDefs

```
typedef detail::shared_mutex shared_mutex

#include <hpx/synchronization/spinlock_no_backoff.hpp>

namespace hpx

namespace lcos

namespace local

struct spinlock_no_backoff
#include <spinlock_no_backoff.hpp> boost::mutex-compatible spinlock class
```

Public Functions

```
HPX_NON_COPYABLE(spinlock_no_backoff)

spinlock_no_backoff()
~spinlock_no_backoff()

void lock()
bool try_lock()
void unlock()
```

Private Members

```
std::uint64_t v_
```

```
#include <hpx/synchronization/barrier.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

class barrier

#include <barrier.hpp> A barrier can be used to synchronize a specific number of threads, blocking all of the entering threads until all of the threads have entered the barrier.

Note A *barrier* is not a LCO in the sense that it has no global id and it can't be triggered using the action (parcel) mechanism. It is just a low level synchronization primitive allowing to synchronize a given number of *threads*.

Public Functions

```
barrier (std::size_t number_of_threads)
```

```
~barrier ()
```

```
void wait ()
```

The function *wait* will block the number of entering *threads* (as given by the constructor parameter *number_of_threads*), releasing all waiting threads as soon as the last *thread* entered this function.

```
void count_up ()
```

The function *count_up* will increase the number of *threads* to be waited in *wait* function.

```
void reset (std::size_t number_of_threads)
```

The function *reset* will reset the number of *threads* as given by the function parameter *number_of_threads*. the newer coming *threads* executing the function *wait* will be waiting until *total_* is equal to *barrier_flag*. The last *thread* exiting the *wait* function will notify the newer *threads* waiting and the newer *threads* will get the reset *number_of_threads_*. The function *reset* can be executed while previous *threads* executing waiting after they have been waken up. Thus *total_* can not be reset to *barrier_flag* which will break the comparison condition under the function *wait*.

Private Types

```
typedef lcos::local::spinlock mutex_type
```

Private Members

```
HPX_STATIC_CONSTEXPR std::size_t hpx::lcos::local::barrier::barrier_flag_
std::size_t number_of_threads_
std::size_t total_
mutex_type mtx_
local::detail::condition_variable cond_
```

```
#include <hpx/synchronization/once.hpp>
```

Defines

```
HPX_ONCE_INIT
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

Functions

```
template<typename F, typename ...Args>
void call_once(once_flag &flag, F &&f, Args&&... args)

struct once_flag
```

Public Functions

```
HPX_NON_COPYABLE(once_flag)
once_flag()
```

Private Members

```
std::atomic<long> status_
lcos::local::event event_
```

Friends

```
template<typename F, typename ...Args>
void call_once (once_flag &flag, F &&f, Args&&... args)
```

```
#include <hpx/synchronization/channel_mpsc.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

TypeDefs

```
using channel_mpsc = base_channel_mpsc<T, hpx::lcos::local::spinlock>

template<typename T, typename Mutex = util::spinlock>
class base_channel_mpsc
```

Public Functions

```
base_channel_mpsc (std::size_t size)

base_channel_mpsc (base_channel_mpsc &&rhs)

base_channel_mpsc &operator= (base_channel_mpsc &&rhs)

~base_channel_mpsc ()

bool get (T *val = nullptr) const

bool set (T &&t)

std::size_t close ()

std::size_t capacity () const
```

Private Types

```
template<>
using mutex_type = Mutex
```

Private Functions

```
bool is_full (std::size_t tail) const  
bool is_empty (std::size_t head) const
```

Private Members

```
hpx::util::cache_aligned_data<std::atomic<std::size_t>> head_  
hpx::util::cache_aligned_data<tail_data> tail_  
std::size_t size_  
std::unique_ptr<T[]> buffer_  
std::atomic<bool> closed_  
struct tail_data
```

Public Members

```
template<>  
mutex_type mtx_  
template<>  
std::atomic<std::size_t> tail_
```

```
#include <hpx/synchronization/latch.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
class latch
```

#include <latch.hpp> Latches are a thread coordination mechanism that allow one or more threads to block until an operation is completed. An individual latch is a singleuse object; once the operation has been completed, the latch cannot be reused.

A latch maintains an internal counter_ that is initialized when the latch is created. Threads may block at a synchronization point waiting for counter_ to be decremented to 0. When counter_ reaches 0, all such blocked threads are released.

Calls to *countdown_and_wait()* , *count_down()* , *wait()* , *is_ready()* , *count_up()* , and *reset()* behave as atomic operations.

Note A *local::latch* is not a LCO in the sense that it has no global id and it can't be triggered using the action (parcel) mechanism. Use *lcos::latch* instead if this is required. It is just a low level synchronization primitive allowing to synchronize a given number of *threads*.

Public Functions

HPX_NON_COPYABLE (*latch*)

latch (*std*::ptrdiff_t *count*)

Initialize the latch

Requires: count >= 0. Synchronization: None Postconditions: counter_ == count.

~latch()

Requires: No threads are blocked at the synchronization point.

Note May be called even if some threads have not yet returned from *wait()* or *count_down_and_wait()*, provided that counter_ is 0.

Note The destructor might not return until all threads have exited *wait()* or *count_down_and_wait()*.

Note It is the caller's responsibility to ensure that no other thread enters *wait()* after one thread has called the destructor. This may require additional coordination.

void count_down_and_wait()

Decrements counter_ by 1. Blocks at the synchronization point until counter_ reaches 0.

Requires: counter_ > 0.

Synchronization: Synchronizes with all calls that block on this latch and with all *is_ready* calls on this latch that return true.

Exceptions

- Nothing.:

void count_down (std::ptrdiff_t *n*)

Decrements counter_ by n. Does not block.

Requires: counter_ >= n and n >= 0.

Synchronization: Synchronizes with all calls that block on this latch and with all *is_ready* calls on this latch that return true .

Exceptions

- Nothing.:

bool is_ready () const

Returns: counter_ == 0. Does not block.

Exceptions

- Nothing.:

void wait () const

If counter_ is 0, returns immediately. Otherwise, blocks the calling thread at the synchronization point until counter_ reaches 0.

Exceptions

- Nothing.:

void abort_all ()

void count_up (std::ptrdiff_t *n*)

Increments counter_ by n. Does not block.

Requires: n >= 0.

Exceptions

- Nothing.:

```
void reset (std::ptrdiff_t n)
    Reset counter_ to n. Does not block.

    Requires: n >= 0.
```

Exceptions

- Nothing.:

Protected Attributes

```
util::cache_line_data<mutex_type> mtx_
util::cache_line_data<local::detail::condition_variable> cond_
std::atomic<std::ptrdiff_t> counter_
bool notified_
```

Private Types

```
typedef lcos::local::spinlock mutex_type
```

```
#include <hpx/synchronization/condition_variable.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

Enums

```
enum cv_status
```

```
    Values:
```

```
        no_timeout
```

```
        timeout
```

```
        error
```

```
    class condition_variable
```

Public Functions

```
void notify_one (error_code &ec = throws)
void notify_all (error_code &ec = throws)
void wait (std::unique_lock<mutex> &lock, error_code &ec = throws)
template<class Predicate>
void wait (std::unique_lock<mutex> &lock, Predicate pred, error_code& = throws)
cv_status wait_until (std::unique_lock<mutex> &lock, util::steady_time_point const
&abs_time, error_code &ec = throws)
template<typename Predicate>
bool wait_until (std::unique_lock<mutex> &lock, util::steady_time_point const
&abs_time, Predicate pred, error_code &ec = throws)
cv_status wait_for (std::unique_lock<mutex> &lock, util::steady_duration const
&rel_time, error_code &ec = throws)
template<typename Predicate>
bool wait_for (std::unique_lock<mutex> &lock, util::steady_duration const &rel_time,
Predicate pred, error_code &ec = throws)
```

Private Types

```
typedef lcos::local::spinlock mutex_type
```

Private Members

```
util::cache_line_data<mutex_type> mtx_
util::cache_line_data<detail::condition_variable> cond_
class condition_variable_any
```

Public Functions

```
void notify_one (error_code &ec = throws)
void notify_all (error_code &ec = throws)
template<class Lock>
void wait (Lock &lock, error_code &ec = throws)
template<class Lock, class Predicate>
void wait (Lock &lock, Predicate pred, error_code &ec = throws)
template<typename Lock>
cv_status wait_until (Lock &lock, util::steady_time_point const &abs_time, error_code
&ec = throws)
template<typename Lock, typename Predicate>
bool wait_until (Lock &lock, util::steady_time_point const &abs_time, Predicate pred,
error_code &ec = throws)
```

```
template<typename Lock>
cv_status wait_for(Lock &lock, util::steady_duration const &rel_time, error_code &ec = throws)

template<typename Lock, typename Predicate>
bool wait_for(Lock &lock, util::steady_duration const &rel_time, Predicate pred, error_code &ec = throws)
```

Private Types

```
typedef lcos::local::spinlock mutex_type
```

Private Members

```
util::cache_line_data<mutex_type> mtx_
util::cache_line_data<detail::condition_variable> cond_
```

```
#include <hpx/synchronization/event.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

class event

#include <event.hpp> Event semaphores can be used for synchronizing multiple threads that need to wait for an event to occur. When the event occurs, all threads waiting for the event are woken up.

Public Functions

```
event()
```

Construct a new event semaphore.

```
bool occurred()
```

Check if the event has occurred.

```
void wait()
```

Wait for the event to occur.

```
void set()
```

Release all threads waiting on this semaphore.

```
void reset()
```

Reset the event.

Private Types

```
typedef lcos::local::spinlock mutex_type
```

Private Functions

```
void wait_locked(std::unique_lock<mutex_type> &l)
```

```
void set_locked(std::unique_lock<mutex_type> l)
```

Private Members

```
mutex_type mtx_
```

This mutex protects the queue.

```
local::detail::condition_variable cond_
```

```
std::atomic<bool> event_
```

```
#include <hpx/synchronization/recursive_mutex.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

TypeDefs

```
using recursive_mutex = detail::recursive_mutex_impl<>
```

```
#include <hpx/synchronization/spinlock_pool.hpp>
```

```
namespace hpx
```

```
namespace lcos
```

```
namespace local
```

```
template<typename Tag, std::size_t N = HPX_HAVE_SPINLOCK_POOL_NUM>
class spinlock_pool
```

Public Static Functions

```
static lcos::local::spinlock &spinlock_for(void const *pv)
```

Private Static Attributes

```
util::cache_aligned_data<lcos::local::spinlock> pool_
class scoped_lock
```

Public Functions

```
template<>
HPX_NON_COPYABLE(scoped_lock)
template<>
scoped_lock(void const *pv)
template<>
~scoped_lock()
template<>
void lock()
template<>
void unlock()
```

Private Members

```
template<>
hpx::lcos::local::spinlock &sp_
```

```
#include <hpx/synchronization/mutex.hpp>

namespace hpx

namespace lcos

namespace local

class mutex
    Subclassed by hpx::lcos::local::timed_mutex
```

Public Functions

```
HPX_NON_COPYABLE (mutex)

mutex (char const *const description = "")

~mutex ()

void lock (char const *description, error_code &ec = throws)

void lock (error_code &ec = throws)

bool try_lock (char const *description, error_code &ec = throws)

bool try_lock (error_code &ec = throws)

void unlock (error_code &ec = throws)
```

Protected Types

```
typedef lcos::local::spinlock mutex_type
```

Protected Attributes

```
mutex_type mtx_
threads::thread_id_type owner_id_
detail::condition_variable cond_

class timed_mutex : private hpx::lcos::local::mutex
```

Public Functions

```
HPX_NON_COPYABLE (timed_mutex)

timed_mutex (char const *const description = "")

~timed_mutex ()

bool try_lock_until (util::steady_time_point const &abs_time, char const *description, error_code &ec = throws)

bool try_lock_until (util::steady_time_point const &abs_time, error_code &ec = throws)

bool try_lock_for (util::steady_duration const &rel_time, char const *description, error_code &ec = throws)

bool try_lock_for (util::steady_duration const &rel_time, error_code &ec = throws)

namespace threads
```

TypeDefs

```
using thread_id_type = thread_id
using thread_self = coroutines::detail::coroutine_self
```

Functions

thread_id_type **get_self_id()**

The function *get_self_id* returns the HPX thread id of the current thread (or zero if the current thread is not a HPX thread).

thread_self ***get_self_ptr()**

The function *get_self_ptr* returns a pointer to the (OS thread specific) self reference to the current HPX thread.

```
#include <hpx/synchronization/channel_mpmc.hpp>
```

```
namespace hpx
```

```
    namespace lcos
```

```
        namespace local
```

TypeDefs

```
using channel_mpmc = bounded_channel<T, hpx::lcos::local::spinlock>
```

```
template<typename T, typename Mutex = util::spinlock>
class bounded_channel
```

Public Functions

```
bounded_channel(std::size_t size)
bounded_channel(bounded_channel &&rhs)
bounded_channel &operator=(bounded_channel &&rhs)
~bounded_channel()
bool get(T *val = nullptr) const
bool set(T &&t)
std::size_t close()
std::size_t capacity() const
```

Protected Functions

```
std::size_t close (std::unique_lock<mutex_type> &l)
```

Private Types

```
template<>
using mutex_type = Mutex
```

Private Functions

```
bool is_full (std::size_t tail) const
```

```
bool is_empty (std::size_t head) const
```

Private Members

```
hpx::util::cache_aligned_data<mutex_type> mtx_
hpx::util::cache_aligned_data<std::size_t> head_
hpx::util::cache_aligned_data<std::size_t> tail_
std::size_t size_
std::unique_ptr<T[]> buffer_
bool closed_
```

algorithms

```
#include <hpx/traits/is_value_proxy.hpp>

#include <hpx/traits/segmented_iterator_traits.hpp>

namespace hpx

namespace traits

template<typename Iterator, typename Enable = void>
struct segmented_iterator_traits
```

Public Types

```
typedef std::false_type is_segmented_iterator

template<typename Iterator, typename Enable = void>
struct segmented_local_iterator_traits
```

Public Types

```
typedef std::false_type is_segmented_local_iterator

typedef Iterator iterator
typedef Iterator local_iterator
typedef Iterator local_raw_iterator
```

Public Static Functions

```
static local_raw_iterator const &local (local_iterator const &it)
static local_iterator const &remote (local_raw_iterator const &it)
static local_raw_iterator local (local_iterator &&it)
static local_iterator remote (local_raw_iterator &&it)
```

```
#include <hpx/parallel/numeric.hpp>

#include <hpx/parallel/algorithm.hpp>

#include <hpx/parallel/datapar.hpp>

#include <hpx/parallel/tagspec.hpp>

#include <hpx/parallel/container_algorithms.hpp>

#include <hpx/parallel/memory.hpp>

#include <hpx/parallel/util/loop.hpp>

namespace hpx

    namespace parallel

        namespace util
```

Functions

```
template<typename ExPolicy, typename VecOnly, typename F, typename... Iters>HPX_
template<typename ExPolicy, typename Iter>HPX_HOST_DEVICE constexpr std::enable_
template<typename ExPolicy, typename Begin, typename End, typename F>HPX_HOST_DE
template<typename ExPolicy, typename Begin, typename End, typename CancelToken,
template<typename ExPolicy, typename VecOnly, typename Begin1, typename End1, ty
template<typename ExPolicy, typename Iter, typename F>HPX_HOST_DEVICE std::enable_
template<typename ExPolicy, typename Iter, typename CancelToken, typename F>HPX_
template<typename Iter, typename F, typename Cleanup>
Iter loop_with_cleanup (Iter it, Iter last, F &&f, Cleanup &&cleanup)

template<typename Iter, typename FwdIter, typename F, typename Cleanup>
FwdIter loop_with_cleanup (Iter it, Iter last, FwdIter dest, F &&f, Cleanup &&cleanup)

template<typename Iter, typename F, typename Cleanup>
Iter loop_with_cleanup_n (Iter it, std::size_t count, F &&f, Cleanup &&cleanup)

template<typename Iter, typename FwdIter, typename F, typename Cleanup>
FwdIter loop_with_cleanup_n (Iter it, std::size_t count, FwdIter dest, F &&f, Cleanup
&&cleanup)

template<typename Iter, typename CancelToken, typename F, typename Cleanup>
Iter loop_with_cleanup_n_with_token (Iter it, std::size_t count, CancelToken &tok, F
&&f, Cleanup &&cleanup)

template<typename Iter, typename FwdIter, typename CancelToken, typename F, typename Cleanup>
FwdIter loop_with_cleanup_n_with_token (Iter it, std::size_t count, FwdIter dest,
CancelToken &tok, F &&f, Cleanup
&&cleanup)

template<typename Iter, typename F>
Iter loop_idx_n (std::size_t base_idx, Iter it, std::size_t count, F &&f)

template<typename Iter, typename CancelToken, typename F>
Iter loop_idx_n (std::size_t base_idx, Iter it, std::size_t count, CancelToken &tok, F &&f)

template<typename Iter, typename T, typename Pred>
T accumulate_n (Iter it, std::size_t count, T init, Pred &&f)

template<typename T, typename Iter, typename Reduce, typename Conv = util::projection_identity>
T accumulate (Iter first, Iter last, Reduce &&r, Conv &&conv = Conv())

template<typename T, typename Iter1, typename Iter2, typename Reduce, typename Conv>
T accumulate (Iter1 first1, Iter1 last1, Iter2 first2, Reduce &&r, Conv &&conv)
```

```
#include <hpx/parallel/util/transfer.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace util
```

Functions

```
template<typename InIter, typename OutIter>
std::pair<InIter, OutIter> copy (InIter first, InIter last, OutIter dest)
```

```
template<typename InIter, typename OutIter> HPX_HOST_DEVICE std::pair<InIter, OutIter> move (InIter first, InIter last, OutIter dest)
```

```
template<typename InIter, typename OutIter>
void copy_synchronize (InIter const &first, OutIter const &dest)
```

```
template<typename InIter, typename OutIter>
std::pair<InIter, OutIter> move_n (InIter first, std::size_t count, OutIter dest)
```

```
#include <hpx/parallel/util/projection_identity.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace util
```

```
            struct projection_identity
```

Public Functions

```
template<typename T> HPX_HOST_DEVICE T&& hpx::parallel::util::projection_identity::operator= (T&& other)
```

```
#include <hpx/parallel/util/cancellation_token.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace util
```

```
template<typename T = detail::no_data, typename Pred = std::less_equal<T>>
class cancellation_token
```

Public Functions

```
cancellation_token(T data)
bool was_cancelled(T data) const
void cancel(T data)
T get_data() const
```

Private Types

```
typedef std::atomic<T> flag_type
```

Private Members

```
std::shared_ptr<flag_type> was_cancelled_
```

```
#include <hpx/parallel/util/transform_loop.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace util
```

Functions

```
template<typename ExPolicy, typename Iter, typename OutIter, typename F>HPX_HOST
template<typename ExPolicy, typename InIter1, typename InIter2, typename OutIter>HPX_HOST
template<typename ExPolicy, typename InIter1, typename InIter2, typename OutIter>HPX_HOST
template<typename ExPolicy, typename Iter, typename OutIter, typename F>HPX_HOST
template<typename ExPolicy, typename InIter1, typename InIter2, typename OutIter>HPX_HOST
```

```
#include <hpx/parallel/util/partitioner.hpp>
```

```
#include <hpx/parallel/util/scan_partitioner.hpp>
```

```
#include <hpx/parallel/util/invoke_projected.hpp>
```

```
namespace hpx
```

```
namespace parallel

namespace util

template<typename Pred, typename Proj>
struct invoke_projected
```

Public Types

```
typedef hpx::util::decay<Pred>::type pred_type
typedef hpx::util::decay<Proj>::type proj_type
```

Public Functions

```
template<typename Pred_, typename Proj_>
invoke_projected(Pred_ &&pred, Proj_ &&proj)

template<typename T>
auto operator() (T &&t)
```

Public Members

```
pred_type pred_
proj_type proj_
```

```
#include <hpx/parallel/util/compare_projected.hpp>
```

```
namespace hpx
```

```
namespace parallel

namespace util

template<typename Compare, typename Proj>
struct compare_projected<Compare, Proj>
```

Public Functions

```
template<typename Compare_, typename Proj_>
compare_projected(Compare_ &&comp, Proj_ &&proj)

template<typename T1, typename T2>
bool operator() (T1 &&t1, T2 &&t2) const
```

Public Members

Compare `comp_`

Proj `proj_`

```
template<typename Compare, typename Proj1, typename Proj2>
struct compare_projected<Compare, Proj1, Proj2>
```

Public Functions

```
template<typename Compare_, typename Proj1_, typename Proj2_>
compare_projected(Compare_ &&comp, Proj1_ &&proj1, Proj2_ &&proj2)
```

```
template<typename T1, typename T2>
bool operator()(T1 &&t1, T2 &&t2) const
```

Public Members

Compare `comp_`

Proj1 `proj1_`

Proj2 `proj2_`

```
#include <hpx/parallel/util/foreach_partitioner.hpp>

#include <hpx/parallel/util/zip_iterator.hpp>

#include <hpx/parallel/util/partitioner_with_cleanup.hpp>

#include <hpx/parallel/util/prefetching.hpp>

namespace hpx

namespace parallel

namespace util
```

Functions

```
template<typename Itr, typename ...Ts>
detail::prefetcher_context<Itr, Ts const...> make_prefetcher_context(Itr base_begin,
                                                                     Itr base_end,
                                                                     std::size_t p_factor,
                                                                     Ts const&...
                                                                     rngs)
```

```
#include <hpx/parallel/traits/projected_range.hpp>

template<typename Proj, typename Rng>
struct projected_range<Proj, Rng, typename std::enable_if<hpx::traits::is_range<Rng>::value>::type>
```

Public Types

```
typedef hpx::util::decay<Proj>::type projector_type
typedef hpx::traits::range_iterator<Rng>::type iterator_type

namespace hpx
```

```
namespace parallel
```

```
namespace traits
```

```
template<typename Proj, typename Rng>
struct projected_range<Proj, Rng, typename std::enable_if<hpx::traits::is_range<Rng>::value>::type>
```

Public Types

```
typedef hpx::util::decay<Proj>::type projector_type
typedef hpx::traits::range_iterator<Rng>::type iterator_type
```

```
#include <hpx/parallel/traits/projected.hpp>
```

```
template<typename Iterator>
struct projected_iterator<Iterator, typename std::enable_if<is_segmented_iterator<Iterator>::value>::type>
```

Public Types

```
typedef segmented_iterator_traits<Iterator>::local_iterator local_iterator
typedef segmented_local_iterator_traits<local_iterator>::local_raw_iterator type
```

```
template<typename Iterator>
struct projected_iterator<Iterator, typename hpx::util::always_void<typename hpx::util::decay<Iterator>::type::proxy>>
```

Public Types

```
typedef hpx::util::decay<Iterator>::type::proxy_type type
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace traits
```

```
template<typename Proj, typename Iter>
struct projected
```

Public Types

```
typedef hpx::util::decay<Proj>::type projector_type
```

```
typedef hpx::traits::projected_iterator<Iter>::type iterator_type
```

```
namespace traits
```

```
template<typename T, typename Enable = void>
struct projected_iterator
```

Public Types

```
typedef hpx::util::decay<T>::type type
```

```
template<typename Iterator>
```

```
struct projected_iterator<Iterator, typename hpx::util::always_void<typename hpx::util::decay<Iterator>::type>>
```

Public Types

```
typedef hpx::util::decay<Iterator>::type::proxy_type type
```

```
template<typename Iterator>
```

```
struct projected_iterator<Iterator, typename std::enable_if<is_segmented_iterator<Iterator>::value>::type>
```

Public Types

```
typedef segmented_iterator_traits<Iterator>::local_iterator local_iterator
```

```
typedef segmented_local_iterator_traits<local_iterator>::local_raw_iterator type
```

```
#include <hpx/parallel/datapar/loop.hpp>
```

```
#include <hpx/parallel/datapar/iterator_helpers.hpp>
```

```
#include <hpx/parallel/datapar/transform_loop.hpp>
```

```
#include <hpx/parallel/datapar/zip_iterator.hpp>
```

```
#include <hpx/parallel/algorithms/remove_copy.hpp>
```

```
namespace hpx
```

```
namespace parallel

namespace v1
```

Functions

template<typename **ExPolicy**, typename **FwdIter1**, typename **FwdIter2**, typename **T**, typename **Proj** = *util::parallel_if*
util::detail::algorithm_result<*ExPolicy*, *hpx::util*::tagged_pair<tag::in(*FwdIter1*) , tag::out
FwdIter2>>::type **remove_copy***ExPolicy* &&*policy*, *FwdIter1* *first*, *FwdIter1* *last*, *FwdIter2*
dest, *T* **const** &*val*, *Proj* &&*proj* = *Proj*()) Copies the elements in the range, defined by [*first*,
last), to another range beginning at *dest*. Copies only the elements for which the comparison
operator returns false when compare to *val*. The order of the elements that are not removed is
preserved.

Effects: Copies all the elements referred to by the iterator it in the range [*first*,*last*) for which the
following corresponding conditions do not hold: *INVOKE*(*proj*, **it*) == *value*

The assignments in the parallel *remove_copy* algorithm invoked with an execution policy object
of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications
of the predicate *f*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in
which the execution of the algorithm may be parallelized and the manner in which it executes
the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet
the requirements of an forward iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This
iterator type must meet the requirements of an forward iterator.
- **T**: The type that the result of dereferencing *FwdIter1* is compared to.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied
to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **val**: Value to be removed.
- **proj**: Specifies the function (or function object) which will be invoked for each of the
elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove_copy* algorithm invoked with an execution policy object
of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion
in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove_copy* algorithm returns a *hpx::future<tagged_pair<tag::in(*FwdIter1*),
tag::out(*FwdIter2*)>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::in(*FwdIter1*), tag::out(*FwdIter2*)>* otherwise.
The *copy* algorithm returns the pair of the input iterator forwarded to the first element after
the last in the input sequence and the output iterator to the element in the destination range, one
past the last element copied.

template<typename **ExPolicy**, typename **FwdIter1**, typename **FwdIter2**, typename **F**, typename **Proj** = *util::parallel_if*
util::detail::algorithm_result<*ExPolicy*, *hpx::util*::tagged_pair<tag::in(*FwdIter1*) , tag::out
FwdIter2>>::type **remove_copy_if***ExPolicy* &&*policy*, *FwdIter1* *first*, *FwdIter1* *last*,

FwdIter2 *dest*, *F* $\&& f, *Proj* $\&& proj = Proj()$ Copies the elements in the range, defined by [first, last), to another range beginning at *dest*. Copies only the elements for which the predicate *f* returns false. The order of the elements that are not removed is preserved.$

Effects: Copies all the elements referred to by the iterator it in the range [first, last) for which the following corresponding conditions do not hold: INVOKE(pred, INVOKE(proj, *it)) != false.

The assignments in the parallel *remove_copy_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate *f*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *copy_if* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns *true* for the elements to be removed. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter1* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove_copy_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove_copy_if* algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)> >` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The *copy* algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/adjacent_difference.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Assigns each value in the range given by result its corresponding element in the range [first, last] and the one preceding it except *result, which is assigned *first

The difference operations in the parallel *adjacent_difference* invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly (last - first) - 1 application of the binary operator and (last - first) assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the input range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the output range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements of the range the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements of the range the algorithm will be applied to.
- *dest*: Refers to the beginning of the sequence of elements the results will be assigned to.

The difference operations in the parallel *adjacent_difference* invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

This overload of *adjacent_find* is available if the user decides to provide their algorithm their own binary predicate *op*.

Return The *adjacent_difference* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *adjacent_find* algorithm returns an iterator to the last element in the output range.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Op>
```

```
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Assigns each value in the range given by result its corresponding element in the range [first, last] and the one preceding it except *result, which is assigned *first

The difference operations in the parallel *adjacent_difference* invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly (last - first) - 1 application of the binary operator and (last - first) assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the input range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the output range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Op*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *adjacent_difference* requires *Op* to meet the requirements of *CopyConstructible*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements of the range the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements of the range the algorithm will be applied to.
- *dest*: Refers to the beginning of the sequence of elements the results will be assigned to.
- *op*: The binary operator which returns the difference of elements. The signature should be equivalent to the following:

```
bool op(const Type1 &a, const Type1 &b);
```

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The types *Type1* must be such that objects of type *FwdIter1* can be dereferenced and then implicitly converted to the dereferenced type of *dest*.

The difference operations in the parallel *adjacent_difference* invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *adjacent_difference* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *adjacent_find* algorithm returns an iterator to the last element in the output range.

```
#include <hpx/parallel/algorithms/is_sorted.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Pred = std::less<>,
         std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, b>
```

Determines if the range [first, last) is sorted. Uses pred to compare elements.

The comparison operations in the parallel *is_sorted* algorithm invoked with an execution policy object of type *sequenced_policy* executes in sequential order in the calling thread.

Note Complexity: at most $(N+S-1)$ comparisons where $N = \text{distance}(\text{first}, \text{last})$. $S = \text{number of partitions}$

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used for the This iterator type must meet the requirements of a forward iterator.
- **Pred**: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *is_sorted* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::less<>`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements of that the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements of that the algorithm will be applied to.
- **pred**: Refers to the binary predicate which returns true if the first argument should be treated as less than the second argument. The signature of the function should be equivalent to

<code>bool pred(const Type &a, const Type &b);</code>

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type *Type* must be such that objects of types *FwdIter* can be dereferenced and then implicitly converted to *Type*.

The comparison operations in the parallel *is_sorted* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_sorted* algorithm returns a *hpx::future<bool>* if the execution policy is of type *task_execution_policy* and returns *bool* otherwise. The *is_sorted* algorithm returns a *bool* if each element in the sequence [first, last) satisfies the predicate passed. If the range [first, last) contains less than two elements, the function always returns true.

```
template<typename ExPolicy, typename FwdIter, typename Pred = detail::less>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Returns the first element in the range [first, last) that is not sorted. Uses a predicate to compare elements or the less than operator.

The comparison operations in the parallel *is_sorted_until* algorithm invoked with an execution policy object of type *sequenced_policy* executes in sequential order in the calling thread.

Note Complexity: at most $(N+S-1)$ comparisons where $N = \text{distance}(\text{first}, \text{last})$. $S = \text{number of partitions}$

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter*: The type of the source iterators used for the This iterator type must meet the requirements of a forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *is_sorted_until* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::less<>*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements of that the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements of that the algorithm will be applied to.
- *pred*: Refers to the binary predicate which returns true if the first argument should be treated as less than the second argument. The signature of the function should be equivalent to

```
bool pred(const Type &a, const Type &b);
```

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The type *Type* must be such that objects of types *FwdIter* can be dereferenced and then implicitly converted to *Type*.

The comparison operations in the parallel *is_sorted_until* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_sorted_until* algorithm returns a *hpx::future<FwdIter>* if the execution policy is

of type *task_execution_policy* and returns *FwdIter* otherwise. The *is_sorted_until* algorithm returns the first unsorted element. If the sequence has less than two elements or the sequence is sorted, last is returned.

```
#include <hpx/parallel/algorithms/uninitialized_value_construct.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter>
util::detail::algorithm_result<ExPolicy>::type uninitialized_value_construct(ExPolicy
    &&policy,
    FwdIter
    first,
    FwdIter
    last)
```

Constructs objects of type typename iterator_traits<ForwardIt>::value_type in the uninitialized storage designated by the range [first, last) by default-initialization. If an exception is thrown during the initialization, the function has no effects.

The assignments in the parallel *uninitialized_value_construct* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.

The assignments in the parallel *uninitialized_value_construct* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_value_construct* algorithm returns a *hpx::future<void>*, if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *void* otherwise.

```
template<typename ExPolicy, typename FwdIter, typename Size>
```

```
util::detail::algorithm_result<ExPolicy, FwdIter>::type uninitialized_value_construct_n(ExPolicy
&&policy,
FwdIter
first,
Size
count)
```

Constructs objects of type typename iterator_traits<ForwardIt>::value_type in the uninitialized storage designated by the range [first, first + count) by default-initialization. If an exception is thrown during the initialization, the function has no effects.

The assignments in the parallel *uninitialized_value_construct_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* assignments, if *count* > 0, no assignments otherwise.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *Size*: The type of the argument specifying the number of elements to apply *f* to.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *count*: Refers to the number of elements starting at *first* the algorithm will be applied to.

The assignments in the parallel *uninitialized_value_construct_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_value_construct_n* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *uninitialized_value_construct_n* algorithm returns the iterator to the element in the source range, one past the last element constructed.

```
#include <hpx/parallel/algorithms/uninitialized_move.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2ExPolicy, FwdIter2>::type uninitialized_move (ExPolicy
&&policy,
FwdIter1
first,
FwdIter1
last,
FwdIter2
dest)
```

Moves the elements in the range, defined by [first, last), to an uninitialized memory area beginning at dest. If an exception is thrown during the initialization, some objects in [first, last) are left in a valid but unspecified state.

The assignments in the parallel *uninitialized_move* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* move operations.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of a forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.

The assignments in the parallel *uninitialized_move* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_move* algorithm returns a *hpx::future<FwdIter2>*, if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *uninitialized_move* algorithm returns the output iterator to the element in the destination range, one past the last element moved.

```
template<typename ExPolicy, typename FwdIter1, typename Size, typename FwdIter2>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out
FwdIter2>>::type uninitialized_move_nExPolicy &&policy, FwdIter1 first, Size count,
FwdIter2 dest Moves the elements in the range [first, first + count), starting from first and proceeding to first + count - 1., to another range beginning at dest. If an exception is thrown during the initialization, some objects in [first, first + count) are left in a valid but unspecified state.
```

The assignments in the parallel *uninitialized_move_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* movements, if *count > 0*, no move operations otherwise.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.

- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `Size`: The type of the argument specifying the number of elements to apply f to.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of a forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `count`: Refers to the number of elements starting at `first` the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.

The assignments in the parallel `uninitialized_move_n` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `uninitialized_move_n` algorithm returns a `hpx::future<std::pair<FwdIter1, FwdIter2>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `std::pair<FwdIter1, FwdIter2>` otherwise. The `uninitialized_move_n` algorithm returns the pair of the input iterator to the element past in the source range and an output iterator to the element in the destination range, one past the last element moved.

```
#include <hpx/parallel/algorithms/set_symmetric_difference.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename Pre
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Constructs a sorted range beginning at `dest` consisting of all elements present in either of the sorted ranges $[first1, last1]$ and $[first2, last2]$, but not in both of them are copied to the range

beginning at *dest*. The resulting range is also sorted. This algorithm expects both input ranges to be sorted with the given binary predicate *f*.

If some element is found *m* times in [first1, last1) and *n* times in [first2, last2), it will be copied to *dest* exactly std::abs(*m-n*) times. If *m>n*, then the last *m-n* of those elements are copied from [first1, last1), otherwise the last *n-m* elements are copied from [first2, last2). The resulting range cannot overlap with either of the input ranges.

Note Complexity: At most $2*(N1 + N2 - 1)$ comparisons, where *N1* is the length of the first sequence and *N2* is the length of the second sequence.

The resulting range cannot overlap with either of the input ranges.

The application of function objects in parallel algorithm invoked with a sequential execution policy object execute in sequential order in the calling thread (*sequenced_policy*) or in a single new thread spawned from the current thread (for *sequenced_task_policy*).

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *FwdIter1*: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- *FwdIter3*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *set_symmetric_difference* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to std::less<>

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last1*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *first2*: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- *last2*: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *op*: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have const &, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *InIter* can be dereferenced and then implicitly converted to *Type1*

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *set_symmetric_difference* algorithm returns a *hpx::future<FwdIter3>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter3* otherwise. The *set_symmetric_difference* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/is_heap.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename RandIter, typename Comp = detail::less, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type is_heap(ExPolicy &&policy, RandIter
first, RandIter last, Comp
&&comp = Comp(), Proj
&&proj = Proj())
```

Returns whether the range is max heap. That is, true if the range is max heap, false otherwise. The function uses the given comparison function object *comp* (defaults to using operator<()).

comp has to induce a strict weak ordering on the values.

Note Complexity: Performs at most N applications of the comparison *comp*, at most $2 * N$ applications of the projection *proj*, where $N = \text{last} - \text{first}$.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **RandIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- **Comp**: The type of the function/function object to use (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the INVOKE operation applied to an object of type *Comp*, when contextually converted to *bool*, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_heap* algorithm returns a `hpx::future<bool>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *is_heap* algorithm returns whether the range is max heap. That is, true if the range is max heap, false otherwise.

```
template<typename ExPolicy, typename RandIter, typename Comp = detail::less, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type is_heap(ExPolicy &&policy, RandIter
first, RandIter last, Comp
&&comp = Comp(), Proj
&&proj = Proj())
```

```
util::detail::algorithm_result<ExPolicy, RandIter>::type is_heap_until (ExPolicy &&policy,
                                                               RandIter first,
                                                               RandIter last,
                                                               Comp &&comp
                                                               = Comp(), Proj
                                                               &&proj = Proj())
```

Returns the upper bound of the largest range beginning at *first* which is a max heap. That is, the last iterator *it* for which range [*first*, *it*) is a max heap. The function uses the given comparison function object *comp* (defaults to using operator<()).

comp has to induce a strict weak ordering on the values.

Note Complexity: Performs at most N applications of the comparison *comp*, at most 2 * N applications of the projection *proj*, where N = *last* - *first*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *RandIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- *Comp*: The type of the function/function object to use (deduced).
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the INVOKE operation applied to an object of type *Comp*, when contextually converted to bool, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_heap_until* algorithm returns a `hpx::future<RandIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *RandIter* otherwise. The *is_heap_until* algorithm returns the upper bound of the largest range beginning at *first* which is a max heap. That is, the last iterator *it* for which range [*first*, *it*) is a max heap.

```
#include <hpx/parallel/algorithms/transform_inclusive_scan.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Op, typename Conv, typename util::detail::algorithm_result<ExPolicy, FwdIter2>::type transform_inclusive_scan (ExPolicy
    &&policy,
    FwdIter1
    first,
    FwdIter1
    last,
    FwdIter2
    dest,
    Op
    &&op,
    Conv
    &&conv,
    T
    init)
```

Assigns through each iterator i in $[result, result + (last - first))$ the value of GENERALIZED_NONCOMMUTATIVE_SUM($op, init, conv(*first), \dots, conv(*(first + (i - result)))$)).

The reduce operations in the parallel *transform_inclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate op .

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Conv**: The type of the unary function object used for the conversion operation.
- **T**: The type of the value to be used as initial (and intermediate) values (deduced).
- **Op**: The type of the binary function object used for the reduction operation.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **conv**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by $[first, last)$. This is a unary predicate. The signature of this predicate should be equivalent to:

```
R fun (const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter1* can be dereferenced and then implicitly converted to *Type*. The type *R* must be such that an object of this type can be implicitly converted to *T*.

- **init**: The initial value for the generalized sum.
- **op**: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1` and `Ret` must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel `transform_inclusive_scan` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Neither `conv` nor `op` shall invalidate iterators or subranges, or modify elements in the ranges `[first,last)` or `[result,result + (last - first))`.

Return The `copy_n` algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `transform_inclusive_scan` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note `GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aN)` is defined as:

- `a1` when $N = 1$
- `op(GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aK), GENERALIZED_NONCOMMUTATIVE_SUM(op, aM, ..., aN))` where $1 < K+1 = M \leq N$.

The difference between `exclusive_scan` and `transform_inclusive_scan` is that `transform_inclusive_scan` includes the i th input element in the i th sum. If `op` is not mathematically associative, the behavior of `transform_inclusive_scan` may be non-deterministic.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Conv, typename OpExPolicy, FwdIter2>::type transform_inclusive_scan(ExPolicy
    &&pol-
    icy,
    FwdIter1
    first,
    FwdIter1
    last,
    FwdIter2
    dest,
    Op
    &&op,
    Conv
    &&conv)
```

Assigns through each iterator i in `[result, result + (last - first))` the value of `GENERALIZED_NONCOMMUTATIVE_SUM(op, conv(*first), ..., conv(*first + (i - result))))`.

The reduce operations in the parallel `transform_inclusive_scan` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate `op`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `Conv`: The type of the unary function object used for the conversion operation.
- `T`: The type of the value to be used as initial (and intermediate) values (deduced).
- `Op`: The type of the binary function object used for the reduction operation.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **conv**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a unary predicate. The signature of this predicate should be equivalent to:

```
R fun(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`. The type `R` must be such that an object of this type can be implicitly converted to `T`.

- **op**: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1` and `Ret` must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel `transform_inclusive_scan` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Neither `conv` nor `op` shall invalidate iterators or subranges, or modify elements in the ranges `[first,last)` or `[result,result + (last - first))`.

Return The `copy_n` algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `transform_inclusive_scan` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note `GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aN)` is defined as:

- `a1` when `N` is 1
- `op(GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aK), GENERALIZED_NONCOMMUTATIVE_SUM(op, aM, ..., aN))` where $1 < K+1 = M \leq N$.

The difference between `exclusive_scan` and `transform_inclusive_scan` is that `transform_inclusive_scan` includes the `i`th input element in the `i`th sum.

```
#include <hpx/parallel/algorithms/equal.hpp>

namespace hpx

    namespace parallel

        namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, b>
```

Returns true if the range [first1, last1) is equal to the range [first2, last2), and false otherwise.

The comparison operations in the parallel *equal* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most $\min(\text{last1} - \text{first1}, \text{last2} - \text{first2})$ applications of the predicate *f*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Pred**: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *equal* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first1**: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- **last1**: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- **first2**: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- **last2**: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- **op**: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively. The comparison operations in the parallel *equal* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The two ranges are considered equal if, for every iterator *i* in the range [first1, last1), **i* equals *(*first2 + (i - first1))*. This overload of equal uses operator== to determine if two elements are equal.

Return The *equal* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *equal* algorithm returns true if the elements in the two ranges are equal, otherwise it returns false. If the length of the range [first1, last1) does not equal the length of the range [first2, last2), it returns false.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy>, b
```

Returns true if the range [first1, last1) is equal to the range starting at first2, and false otherwise.

The comparison operations in the parallel *equal* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most *last1 - first1* applications of the predicate *f*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *equal* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::equal_to<>*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last1*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *first2*: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- *op*: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively. The comparison operations in the parallel *equal* algorithm invoked with an execution policy object

of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The two ranges are considered equal if, for every iterator *i* in the range [first1,last1), **i* equals **(first2 + (i - first1))*. This overload of equal uses operator== to determine if two elements are equal.

Return The *equal* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *equal* algorithm returns true if the elements in the two ranges are equal, otherwise it returns false.

```
#include <hpx/parallel/algorithms/reduce.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIterB, typename FwdIterE, typename T, typename F>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, T>> GENERALIZED_SUM(FwdIterB first, FwdIterE last, T init, ExPolicy exPolicy)
```

Returns GENERALIZED_SUM(*f*, *init*, **first*, ..., **(first + (last - first) - 1)*).

The reduce operations in the parallel *reduce* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate *f*.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIterB:** The type of the source begin iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIterE:** The type of the source end iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- **F:** The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *copy_if* requires *F* to meet the requirements of *CopyConstructible*.
- **T:** The type of the value to be used as initial (and intermediate) values (deduced).

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **first:** Refers to the beginning of the sequence of elements the algorithm will be applied to.

- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`. The types `Type1` `Ret` must be such that an object of type `FwdIterB` can be dereferenced and then implicitly converted to any of those types.

- `init`: The initial value for the generalized sum.

The reduce operations in the parallel `copy_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between `reduce` and `accumulate` is that the behavior of `reduce` may be non-deterministic for non-associative or non-commutative binary predicate.

Return The `reduce` algorithm returns a `hpx::future<T>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `T` otherwise. The `reduce` algorithm returns the result of the generalized sum over the elements given by the input range [first, last).

Note `GENERALIZED_SUM(op, a1, ..., aN)` is defined as follows:

- a_1 when N is 1
- $op(GENERALIZED_SUM(op, b_1, \dots, b_K), GENERALIZED_SUM(op, b_{M+1}, \dots, b_N))$, where:
 - b_1, \dots, b_N may be any permutation of a_1, \dots, a_N and
 - $1 < K+1 = M \leq N$.

```
template<typename ExPolicy, typename FwdIterB, typename FwdIterE, typename T>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, T>> reduce(...);
```

Returns `GENERALIZED_SUM(+, init, *first, ..., *(first + (last - first) - 1))`.

The reduce operations in the parallel `reduce` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the operator `+`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIterB`: The type of the source begin iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIterE`: The type of the source end iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- `T`: The type of the value to be used as initial (and intermediate) values (deduced).

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `init`: The initial value for the generalized sum.

The reduce operations in the parallel *copy_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *reduce* and *accumulate* is that the behavior of *reduce* may be non-deterministic for non-associative or non-commutative binary predicate.

Return The *reduce* algorithm returns a *hpx::future<T>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *T* otherwise. The *reduce* algorithm returns the result of the generalized sum (applying operator+()) over the elements given by the input range [first, last).

Note GENERALIZED_SUM(+, a1, ..., aN) is defined as follows:

- a1 when N is 1
- op(GENERALIZED_SUM(+, b1, ..., bK), GENERALIZED_SUM(+, bM, ..., bN)), where:
 - b1, ..., bN may be any permutation of a1, ..., aN and
 - 1 < K+1 = M <= N.

```
template<typename ExPolicy, typename FwdIterB, typename FwdIterE>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, t
```

Returns GENERALIZED_SUM(+, T(), *first, ..., *(first + (last - first) - 1)).

The reduce operations in the parallel *reduce* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: O(*last - first*) applications of the operator+().

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIterB*: The type of the source begin iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIterE*: The type of the source end iterator used (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.

The reduce operations in the parallel *copy_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *reduce* and *accumulate* is that the behavior of *reduce* may be non-deterministic for non-associative or non-commutative binary predicate.

Return The *reduce* algorithm returns a *hpx::future<T>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *T* otherwise (where *T* is the *value_type* of *FwdIterB*). The *reduce* algorithm returns the result of the generalized sum (applying operator+()) over the elements given by the input range [first, last).

Note The type of the initial value (and the result type) *T* is determined from the *value_type* of the used *FwdIterB*.

Note GENERALIZED_SUM(+, a1, ..., aN) is defined as follows:

- a1 when N is 1
- op(GENERALIZED_SUM(+, b1, ..., bK), GENERALIZED_SUM(+, bM, ..., bN)), where:
 - b1, ..., bN may be any permutation of a1, ..., aN and
 - 1 < K+1 = M <= N.

```
#include <hpx/parallel/algorithms/find.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename T>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Returns the first element in the range [first, last) that is equal to value

The comparison operations in the parallel *find* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most last - first applications of the operator==().

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- **T**: The type of the value to find (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- **val**: the value to compare the elements to

The comparison operations in the parallel *find* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *find* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *find* algo-

rithm returns the first element in the range [first,last) that is equal to *val*. If no such element in the range of [first,last) is equal to *val*, then the algorithm returns *last*.

```
template<typename ExPolicy, typename FwdIter, typename F>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F>
```

Returns the first element in the range [first, last) for which predicate *f* returns true

The comparison operations in the parallel *find_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most last - first applications of the predicate.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of a forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- **f**: The unary predicate which returns true for the required element. The signature of the predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have **const** &, but the function must not modify the objects passed to it. The type *Type* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

The comparison operations in the parallel *find_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *find_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *find_if* algorithm returns the first element in the range [first,last) that satisfies the predicate *f*. If no such element exists that satisfies the predicate *f*, the algorithm returns *last*.

```
template<typename ExPolicy, typename FwdIter, typename F>
```

`std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter1, FwdIter2, Pred>::type>`

Returns the first element in the range [first, last) for which predicate *f* returns false

The comparison operations in the parallel *find_if_not* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most last - first applications of the predicate.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of a forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *f*: The unary predicate which returns false for the required element. The signature of the predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type *Type* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

The comparison operations in the parallel *find_if_not* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *find_if_not* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *find_if_not* algorithm returns the first element in the range [first, last) that does **not** satisfy the predicate *f*. If no such element exists that does not satisfy the predicate *f*, the algorithm returns *last*.

template<typename **ExPolicy**, typename **FwdIter1**, typename **FwdIter2**, typename **Pred** = detail::equal_to, typ

`std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter1, FwdIter2, Pred, Proj>::type>`

Returns the last subsequence of elements [first2, last2) found in the range [first, last) using the given predicate *f* to compare elements.

The comparison operations in the parallel *find_end* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most $S^*(N-S+1)$ comparisons where $S = \text{distance}(\text{first2}, \text{last2})$ and $N = \text{distance}(\text{first1}, \text{last1})$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *replace* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of type dereferenced *FwdIter1* and dereferenced *FwdIter2*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last1*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *first2*: Refers to the beginning of the sequence of elements the algorithm will be searching for.
- *last2*: Refers to the end of the sequence of elements of the algorithm will be searching for.
- *op*: The binary predicate which returns *true* if the elements should be treated as equal. The signature should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and

FwdIter2 can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *FwdIter1* and dereferenced *FwdIter2* as a projection operation before the function *f* is invoked.

The comparison operations in the parallel *find_end* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

This overload of *find_end* is available if the user decides to provide the algorithm their own predicate *f*.

Return The *find_end* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *find_end* algorithm returns an iterator to the beginning of the last subsequence [first2, last2] in range [first, last]. If the length of the subsequence [first2, last2] is greater than the length of the range [first1, last1], *last1* is returned. Additionally if the size of the subsequence is empty or no subsequence is found, *last1* is also returned.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to, typename std::enable_if<execution::is_execution_policy<ExPolicy>::value, util::detail::algorithm_result<ExPolicy, FwdIter2>> result_type(ExPolicy exPolicy, FwdIter1 first, FwdIter1 last, FwdIter2 sFirst, FwdIter2 sLast, Pred pred)
```

Searches the range [first, last) for any elements in the range [s_first, s_last). Uses binary predicate *p* to compare elements

The comparison operations in the parallel *find_first_of* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most ($S \cdot N$) comparisons where $S = \text{distance}(s_first, s_last)$ and $N = \text{distance}(first, last)$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This

iterator type must meet the requirements of an forward iterator.

- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *equal* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to std::equal_to<>
- *Proj1*: The type of an optional projection function. This defaults to util::projection_identity and is applied to the elements of type dereferenced *FwdIter1*.
- *Proj2*: The type of an optional projection function. This defaults to util::projection_identity and is applied to the elements of type dereferenced *FwdIter2*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *s_first*: Refers to the beginning of the sequence of elements the algorithm will be searching for.
- *s_last*: Refers to the end of the sequence of elements of the algorithm will be searching for.
- *op*: The binary predicate which returns *true* if the elements should be treated as equal. The signature should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have const &, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively.

- *proj1*: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *FwdIter1* as a projection operation before the function *op* is invoked.
- *proj2*: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *FwdIter2* as a projection operation before the function *op* is invoked.

The comparison operations in the parallel *find_first_of* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *find_first_of* algorithm returns a *hpx::future<FwdIter1>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter1* otherwise. The *find_first_of* algorithm returns an iterator to the first element in the range [first, last) that is equal to an element from the range [s_first, s_last). If the length of the subsequence [s_first, s_last) is greater than the length of the range [first, last), last is returned. Additionally if the size of the subsequence is empty or no subsequence is found, last is also returned. This overload of *find_end* is available if the user decides to provide the algorithm their own predicate *f*.

```
#include <hpx/parallel/algorithms/reduce_by_key.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename RanIter, typename RanIter2, typename FwdIter1, typename FwdIter2>
util::detail::algorithm_result<ExPolicy, std::pair<FwdIter1, FwdIter2>>::type reduce_by_key (ExPolicy
    &&pol-
    icy,
    Ran-
    Iter
    key_first,
    Ran-
    Iter
    key_last,
    Ran-
    Iter2
    val-
    ues_first,
    FwdIter1
    keys_output,
    FwdIter2
    val-
    ues_output,
    Com-
    pare
    &&comp
    =
    Com-
    pare(),
    Func
    &&func
    =
    Func())
```

Reduce by Key performs an inclusive scan reduction operation on elements supplied in key/value pairs. The algorithm produces a single output value for each set of equal consecutive keys in [key_first, key_last). the value being the GENERALIZED_NONCOMMUTATIVE_SUM(op, init, *first, ..., *(first + (i - result))). for the run of consecutive matching keys. The number of keys supplied must match the number of values.

comp has to induce a strict weak ordering on the values.

Note Complexity: $O(last - first)$ applications of the predicate *op*.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **RanIter:** The type of the key iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- **RanIter2:** The type of the value iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- **FwdIter1:** The type of the iterator representing the destination key range (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2:** The type of the iterator representing the destination value range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Compare:** The type of the optional function/function object to use to compare keys (deduced). Assumed to be `std::equal_to` otherwise.
- **Func:** The type of the function/function object to use (deduced). Unlike its sequential form,

the parallel overload of `copy_if` requires F to meet the requirements of `CopyConstructible`.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `key_first`: Refers to the beginning of the sequence of key elements the algorithm will be applied to.
- `key_last`: Refers to the end of the sequence of key elements the algorithm will be applied to.
- `values_first`: Refers to the beginning of the sequence of value elements the algorithm will be applied to.
- `keys_output`: Refers to the start output location for the keys produced by the algorithm.
- `values_output`: Refers to the start output location for the values produced by the algorithm.
- `comp`: `comp` is a callable object. The return value of the INVOKE operation applied to an object of type `Comp`, when contextually converted to `bool`, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that `comp` will not apply any non-constant function through the dereferenced iterator.
- `func`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by $[first, last]$. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`. The types `Type1` `Ret` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to any of those types.

The application of function objects in parallel algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `reduce_by_key` algorithm returns a `hpx::future<pair<Iter1,Iter2>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `pair<Iter1,Iter2>` otherwise.

```
#include <hpx/parallel/algorithms/move.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out
FwdIter2>>::type moveExPolicy &&policy, FwdIter1 first, FwdIter1 last, FwdIter2 destMoves
the elements in the range [first, last), to another range beginning at dest. After this operation the
elements in the moved-from range will still contain valid values of the appropriate type, but not
necessarily the same values as before the move.
```

The move assignments in the parallel *move* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* move assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the move assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.

The move assignments in the parallel *move* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *move* algorithm returns a *hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>* otherwise. The *move* algorithm returns the pair of the input iterator *last* and the output iterator to the element in the destination range, one past the last element moved.

```
#include <hpx/parallel/algorithms/for_loop_induction.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v2
```

Functions

```
template<typename T>
```

```
detail::induction_stride_helper<T> induction(T &&value, std::size_t stride)
```

The function template returns an induction object of unspecified type having a value type and encapsulating an initial value *value* of that type and, optionally, a stride.

For each element in the input range, a looping algorithm over input sequence *S* computes an induction value from an induction variable and ordinal position *p* within *S* by the formula $i + p * \text{stride}$ if a stride was specified or $i + p$ otherwise. This induction value is passed to the element access function.

If the *value* argument to *induction* is a non-const lvalue, then that lvalue becomes the live-out object for the returned induction object. For each induction object that has a live-out object, the looping algorithm assigns the value of $i + n * \text{stride}$ to the live-out object upon return, where *n* is the number of elements in the input range.

Return This returns an induction object with value type *T*, initial value *value*, and (if specified) stride *stride*. If *T* is an lvalue of non-const type, *value* is used as the live-out object for the induction object; otherwise there is no live-out object.

Template Parameters

- *T*: The value type to be used by the induction object.

Parameters

- *value*: [in] The initial value to use for the induction object
- *stride*: [in] The (optional) stride to use for the induction object (default: 1)

```
#include <hpx/parallel/algorithms/swap_ranges.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
```

```
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Exchanges elements between range [*first1*, *last1*) and another range starting at *first2*.

The swap operations in the parallel *swap_ranges* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between *first1* and *last1*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the swap operations.
- **FwdIter1**: The type of the first range of iterators to swap (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the second range of iterators to swap (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first1**: Refers to the beginning of the first sequence of elements the algorithm will be applied to.
- **last1**: Refers to the end of the first sequence of elements the algorithm will be applied to.
- **first2**: Refers to the beginning of the second sequence of elements the algorithm will be applied to.

The swap operations in the parallel *swap_ranges* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *swap_ranges* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *parallel_task_policy* and returns *FwdIter2* otherwise. The *swap_ranges* algorithm returns iterator to the element past the last element exchanged in the range beginning with *first2*.

```
#include <hpx/parallel/algorithms/transform_reduce_binary.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename T>
util::detail::algorithm_result<ExPolicy, T>::type transform_reduce(ExPolicy      &&policy,
                                                               FwdIter1 first1,
                                                               FwdIter1 last1,
                                                               FwdIter2 first2,   T
                                                               init)
```

Returns the result of accumulating *init* with the inner products of the pairs formed by the elements of two ranges starting at *first1* and *first2*.

The operations in the parallel *transform_reduce* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate *op2*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the first source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

- `FwdIter2`: The type of the second source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `T`: The type of the value to be used as return) values (deduced).

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the first sequence of elements the result will be calculated with.
- `last1`: Refers to the end of the first sequence of elements the algorithm will be applied to.
- `first2`: Refers to the beginning of the second sequence of elements the result will be calculated with.
- `init`: The initial value for the sum.

The operations in the parallel `transform_reduce` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform_reduce` algorithm returns a `hpx::future<T>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `T` otherwise.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename T, typename Reduce, typename Convert>
util::detail::algorithm_result<ExPolicy, T>::type transform_reduce(ExPolicy      &&policy,
                                                               FwdIter1    first1,
                                                               FwdIter1    last1,
                                                               FwdIter2   first2, T init,
                                                               Reduce     &&red_op,
                                                               Convert   &&conv_op)
```

Returns the result of accumulating `init` with the inner products of the pairs formed by the elements of two ranges starting at `first1` and `first2`.

The operations in the parallel `transform_reduce` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate `op2`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the first source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the second source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `T`: The type of the value to be used as return) values (deduced).
- `Reduce`: The type of the binary function object used for the multiplication operation.
- `Convert`: The type of the unary function object used to transform the elements of the input sequence before invoking the reduce function.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the first sequence of elements the result will be calculated with.
- `last1`: Refers to the end of the first sequence of elements the algorithm will be applied to.
- `first2`: Refers to the beginning of the second sequence of elements the result will be calculated with.
- `init`: The initial value for the sum.
- `red_op`: Specifies the function (or function object) which will be invoked for the initial value and each of the return values of `op2`. This is a binary predicate. The signature of this predicate should be equivalent to should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Ret` must be such that it can be implicitly converted to a type of `T`.

- `conv_op`: Specifies the function (or function object) which will be invoked for each of the input values of the sequence. This is a binary predicate. The signature of this predicate should be equivalent to

```
Ret fun(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Ret` must be such that it can be implicitly converted to an object for the second argument type of `op1`.

The operations in the parallel `transform_reduce` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform_reduce` algorithm returns a `hpx::future<T>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `T` otherwise.

```
#include <hpx/parallel/algorithms/uninitialized_copy.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter2>> uninitialized_copy(FwdIter1 first, FwdIter1 last, FwdIter2 dest);
```

Copies the elements in the range, defined by `[first, last)`, to an uninitialized memory area beginning at `dest`. If an exception is thrown during the copy operation, the function has no effects.

The assignments in the parallel `uninitialized_copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` assignments.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of a forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.

The assignments in the parallel `uninitialized_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `uninitialized_copy` algorithm returns a `hpx::future<FwdIter2>`, if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `uninitialized_copy` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename Size, typename FwdIter2>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter2>> uninitialized_copy(FwdIter1 first, FwdIter1 last, FwdIter2 dest, ExPolicy exPolicy);
```

Copies the elements in the range $[first, first + count]$, starting from `first` and proceeding to `first + count - 1`, to another range beginning at `dest`. If an exception is thrown during the copy operation, the function has no effects.

The assignments in the parallel `uninitialized_copy_n` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `count` assignments, if `count > 0`, no assignments otherwise.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an input iterator.
- `Size`: The type of the argument specifying the number of elements to apply f to.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of a forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `count`: Refers to the number of elements starting at `first` the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.

The assignments in the parallel `uninitialized_copy_n` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `uninitialized_copy_n` algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `uninitialized_copy_n` algorithm returns the output iterator to the element in the des-

tination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/lexicographical_compare.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = std::less<>
        std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, b>
```

Checks if the first range [first1, last1) is lexicographically less than the second range [first2, last2).
uses a provided predicate to compare elements.

The comparison operations in the parallel *lexicographical_compare* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most $2 * \min(N1, N2)$ applications of the comparison operation, where $N1 = \text{std}::distance(first1, last)$ and $N2 = \text{std}::distance(first2, last2)$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *lexicographical_compare* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::less<>*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.

- `last1`: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
 - `first2`: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
 - `last2`: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
 - `pred`: Refers to the comparison function that the first and second ranges will be applied to
- The comparison operations in the parallel *lexicographical_compare* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note Lexicographical comparison is an operation with the following properties

- Two ranges are compared element by element
- The first mismatching element defines which range is lexicographically *less* or *greater* than the other
- If one range is a prefix of another, the shorter range is lexicographically *less* than the other
- If two ranges have equivalent elements and are of the same length, then the ranges are lexicographically *equal*
- An empty range is lexicographically *less* than any non-empty range
- Two empty ranges are lexicographically *equal*

Return The *lexicographically_compare* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *lexicographically_compare* algorithm returns true if the first range is lexicographically less, otherwise it returns false. range [first2, last2), it returns false.

```
#include <hpx/parallel/algorithms/replace.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename T1, typename T2, typename Proj = util::projection_util::detail::algorithm_result<ExPolicy, FwdIter>::type replace(ExPolicy &&policy, FwdIter
first, FwdIter last, T1 const
&old_value, T2 const
&new_value, Proj &&proj =
Proj())
```

Replaces all elements satisfying specific criteria with *new_value* in the range [first, last).

Effects: Substitutes elements referred by the iterator *it* in the range [first, last) with *new_value*, when the following corresponding conditions hold: *INVOKE(proj, *it) == old_value*

The assignments in the parallel *replace* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- **T1**: The type of the old value to replace (deduced).
- **T2**: The type of the new values to replace (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **old_value**: Refers to the old value of the elements to replace.
- **new_value**: Refers to the new value to use as the replacement.
- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *f* is invoked.

The assignments in the parallel *replace* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *void* otherwise. It returns *last*.

```
template<typename ExPolicy, typename FwdIter, typename F, typename T, typename Proj = util::projection_identity
util::detail::algorithm_result<ExPolicy, FwdIter>::type replace_if(ExPolicy &&policy,
FwdIter first, FwdIter
last, F &&f, T const
&new_value, Proj
&&proj = Proj())
```

Replaces all elements satisfying specific criteria (for which predicate *f* returns true) with *new_value* in the range [*first*, *last*].

Effects: Substitutes elements referred by the iterator *f* in the range [*first*, *last*] with *new_value*, when the following corresponding conditions hold: `INVOKE(f, INVOKE(proj, *it)) != false`

The assignments in the parallel *replace_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* applications of the predicate.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*. (deduced).
- **T**: The type of the new values to replace (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements

in the sequence specified by [first, last). This is an unary predicate which returns *true* for the elements which need to be replaced. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have **const&**, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *new_value*: Refers to the new value to use as the replacement.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *replace_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *last*.

template<typename **ExPolicy**, typename **FwdIter1**, typename **FwdIter2**, typename **T1**, typename **T2**, typename **util**::detail::algorithm_result<*ExPolicy*, *hpx::util*::tagged_pair<tag::in (*FwdIter1*) , tag::out

FwdIter2>>::type **replace_copy***ExPolicy* &&*policy*, *FwdIter1* *first*, *FwdIter1* *last*, *FwdIter2* *dest*, *T1* **const** &*old_value*, *T2* **const** &*new_value*, *Proj* &&*proj* = *Proj*()Copies the all elements from the range [*first*, *last*) to another range beginning at *dest* replacing all elements satisfying a specific criteria with *new_value*.

Effects: Assigns to every iterator it in the range [*result*, *result* + (*last* - *first*)) either *new_value* or *(*first* + (*i* - *result*)) depending on whether the following corresponding condition holds: INVOKE(*proj*, *(*first* + (*i* - *result*))) == *old_value*

The assignments in the parallel *replace_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last* - *first* applications of the predicate.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- *T1*: The type of the old value to replace (deduced).
- *T2*: The type of the new values to replace (deduced).
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *old_value*: Refers to the old value of the elements to replace.
- *new_value*: Refers to the new value to use as the replacement.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *replace_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion

in unspecified threads, and indeterminately sequenced within each thread.

Return The `replace_copy` algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)> >` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `copy` algorithm returns the pair of the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename F, typename T, typename Proj>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>::type replace_copy_ifExPolicy &&policy, FwdIter1 first, FwdIter1 last,
FwdIter2 dest, F &&f, T const &new_value, Proj &&proj = Proj()Copies the all elements from the range [first, last) to another range beginning at dest replacing all elements satisfying a specific criteria with new_value.
```

Effects: Assigns to every iterator it in the range [result, result + (last - first)) either `new_value` or `*(first + (it - result))` depending on whether the following corresponding condition holds: `INVOKE(f, INVOKE(proj, *(first + (i - result)))) != false`

The assignments in the parallel `replace_copy_if` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` applications of the predicate.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `equal` requires `F` to meet the requirements of `CopyConstructible`. (deduced).
- `T`: The type of the new values to replace (deduced).
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns `true` for the elements which need to replaced. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`.

- `new_value`: Refers to the new value to use as the replacement.
- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `replace_copy_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `replace_copy_if` algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `replace_copy_if` algorithm returns the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/destroy.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

template<typename **ExPolicy**, typename **FwdIter**>
`util`::detail::algorithm_result<**ExPolicy**>::type **destroy** (**ExPolicy** `&&policy`, **FwdIter** `first`,
`FwdIter` `last`)

Destroys objects of type typename iterator_traits<ForwardIt>::value_type in the range [first, last).

The operations in the parallel `destroy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` operations.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.

The operations in the parallel `destroy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `destroy` algorithm returns a `hpx::future<void>`, if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `void` otherwise.

template<typename **ExPolicy**, typename **FwdIter**, typename **Size**>
`util`::detail::algorithm_result<**ExPolicy**, **FwdIter**>::type **destroy_n** (**ExPolicy** `&&policy`,
`FwdIter` `first`, `Size` `count`)

Destroys objects of type typename iterator_traits<ForwardIt>::value_type in the range [first, first + count).

The operations in the parallel `destroy_n` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `count` operations, if `count > 0`, no assignments otherwise.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **Size**: The type of the argument specifying the number of elements to apply this algorithm to.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.

The operations in the parallel *destroy_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *destroy_n* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *destroy_n* algorithm returns the iterator to the element in the source range, one past the last element constructed.

```
#include <hpx/parallel/algorithms/set_intersection.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename Pre  
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Constructs a sorted range beginning at *dest* consisting of all elements present in both sorted ranges

[first1, last1) and [first2, last2). This algorithm expects both input ranges to be sorted with the given binary predicate f .

If some element is found m times in [first1, last1) and n times in [first2, last2), the first $\text{std}::\text{min}(m, n)$ elements will be copied from the first range to the destination range. The order of equivalent elements is preserved. The resulting range cannot overlap with either of the input ranges.

Note Complexity: At most $2*(N1 + N2 - 1)$ comparisons, where $N1$ is the length of the first sequence and $N2$ is the length of the second sequence.

The resulting range cannot overlap with either of the input ranges.

The application of function objects in parallel algorithm invoked with a sequential execution policy object execute in sequential order in the calling thread (*sequenced_policy*) or in a single new thread spawned from the current thread (for *sequenced_task_policy*).

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- `FwdIter1`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `Pred`: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of `set_intersection` requires `Pred` to meet the requirements of *CopyConstructible*. This defaults to `std::less<>`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- `last1`: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- `first2`: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- `last2`: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `op`: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type `Type1` must be such that objects of type `InIter` can be dereferenced and then implicitly converted to `Type1`

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `set_intersection` algorithm returns a `hpx::future<FwdIter3>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `FwdIter3` otherwise. The `set_intersection` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/exclusive_scan.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename T, typename Op>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Assigns through each iterator i in $[result, result + (last - first))$ the value of GENERALIZED_NONCOMMUTATIVE_SUM(binary_op, init, *first, ..., *(first + (i - result) - 1)).

The reduce operations in the parallel *exclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate *op*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- *T*: The type of the value to be used as initial (and intermediate) values (deduced).
- *Op*: The type of the binary function object used for the reduction operation.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *init*: The initial value for the generalized sum.
- *op*: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1` and `Ret` must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel `exclusive_scan` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between `exclusive_scan` and `inclusive_scan` is that `inclusive_scan` includes the *i*th input element in the *i*th sum. If `op` is not mathematically associative, the behavior of `inclusive_scan` may be non-deterministic.

Return The `copy_n` algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `exclusive_scan` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note `GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aN)` is defined as:

- a_1 when $N = 1$
- $op(GENERALIZED_NONCOMMUTATIVE_SUM(op, a_1, \dots, a_K), GENERALIZED_NONCOMMUTATIVE_SUM(op, a_{M+1}, \dots, a_N))$ where $1 < K+1 = M \leq N$.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename T>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter2>> copy_n(FwdIter1 first, FwdIter1 last, FwdIter2 dest, T init, ExPolicy exPolicy);
```

Assigns through each iterator *i* in $[result, result + (last - first))$ the value of `GENERALIZED_NONCOMMUTATIVE_SUM(+, init, *first, \dots, *(first + (i - result) - 1))`

The reduce operations in the parallel `exclusive_scan` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate `std::plus<T>`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `T`: The type of the value to be used as initial (and intermediate) values (deduced).

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `init`: The initial value for the generalized sum.

The reduce operations in the parallel *exclusive_scan* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *exclusive_scan* and *inclusive_scan* is that *inclusive_scan* includes the *i*th input element in the *i*th sum.

Return The *copy_n* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *exclusive_scan* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note *GENERALIZED_NONCOMMUTATIVE_SUM(+, a1, ..., aN)* is defined as:

- *a1* when *N* is 1
- *GENERALIZED_NONCOMMUTATIVE_SUM(+, a1, ..., aK)*
 - *GENERALIZED_NONCOMMUTATIVE_SUM(+, aM, ..., aN)* where $1 < K+1 = M \leq N$.

```
#include <hpx/parallel/algorithms/partition.hpp>
```

```
namespace hpx
```

```
    namespace parallel
```

```
        namespace v1
```

Functions

```
template<typename ExPolicy, typename BidirIter, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, BidirIter>::type stable_partition(ExPolicy
    &&policy,
    BidirIter first,
    BidirIter last,
    F &&f, Proj
    &&proj =
    Proj())
```

Permutates the elements in the range [*first*, *last*) such that there exists an iterator *i* such that for every iterator *j* in the range [*first*, *i*] *INVOKE(f, INVOKE(proj, *j)) != false*, and for every iterator *k* in the range [*i*, *last*), *INVOKE(f, INVOKE(proj, *k)) == false*

The invocations of *f* in the parallel *stable_partition* algorithm invoked with an execution policy object of type *sequenced_policy* executes in sequential order in the calling thread.

Note Complexity: At most (*last - first*) * $\log(\text{last} - \text{first})$ swaps, but only linear number of swaps if there is enough extra memory. Exactly *last - first* applications of the predicate and projection.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of *f*.
- *BidirIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an input iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *transform* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **f**: Unary predicate which returns true if the element should be ordered before other elements. Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool fun(const Type &a);
```

The signature does not need to have `const&`. The type `Type` must be such that an object of type `BidirIter` can be dereferenced and then implicitly converted to `Type`.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `stable_partition` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `stable_partition` algorithm returns an iterator `i` such that for every iterator `j` in the range `[first, i)`, `f(*j) != false` `INVOKE(f, INVOKE(proj, *j)) != false`, and for every iterator `k` in the range `[i, last)`, `f(*k) == false` `INVOKE(f, INVOKE(proj, *k)) == false`. The relative order of the elements in both groups is preserved. If the execution policy is of type `parallel_task_policy` the algorithm returns a future`<>` referring to this iterator.

```
template<typename ExPolicy, typename FwdIter, typename Pred, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type partition(ExPolicy &&policy,
                                                               FwdIter first, FwdIter
                                                               last, Pred &&pred, Proj
                                                               &&proj = Proj())
```

Reorders the elements in the range `[first, last)` in such a way that all elements for which the predicate `pred` returns true precede the elements for which the predicate `pred` returns false. Relative order of the elements is not preserved.

The assignments in the parallel `partition` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: At most $2 * (\text{last} - \text{first})$ swaps. Exactly $\text{last} - \text{first}$ applications of the predicate and projection.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **Pred**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `partition` requires `Pred` to meet the requirements of `CopyConstructible`.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **pred**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate for partitioning the source iterators. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `InIter` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate `is` invoked.

The assignments in the parallel `partition` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `partition` algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `parallel_task_policy` and returns `FwdIter` otherwise. The `partition` algorithm returns the iterator to the first element of the second group.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename Pre util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in(FwdIter1), tag::out1 FwdIter2, tag::out2FwdIter3>>::type partition_copyExPolicy &&policy, FwdIter1 first, FwdIter1 last, FwdIter2 dest_true, FwdIter3 dest_false, Pred &&pred, Proj &&proj = Proj()Copies the elements in the range, defined by [first, last), to two different ranges depending on the value returned by the predicate pred. The elements, that satisfy the predicate pred, are copied to the range beginning at dest_true. The rest of the elements are copied to the range beginning at dest_false. The order of the elements is preserved.
```

The assignments in the parallel `partition_copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than `last - first` assignments, exactly `last - first` applications of the predicate `f`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range for the elements that satisfy the predicate `pred` (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range for the elements that don't satisfy the predicate `pred` (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `partition_copy` requires `Pred` to meet the requirements of `CopyConstructible`.
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest_true`: Refers to the beginning of the destination range for the elements that satisfy the predicate `pred`.
- `dest_false`: Refers to the beginning of the destination range for the elements that don't satisfy the predicate `pred`.
- `pred`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by `[first, last)`. This is an unary predicate for partitioning

the source iterators. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `partition_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `partition_copy` algorithm returns a `hpx::future<tagged_tuple<tag::in(InIter), tag::out1(OutIter1), tag::out2(OutIter2)> >` if the execution policy is of type `parallel_task_policy` and returns `tagged_tuple<tag::in(InIter), tag::out1(OutIter1), tag::out2(OutIter2)>` otherwise. The `partition_copy` algorithm returns the tuple of the source iterator `last`, the destination iterator to the end of the `dest_true` range, and the destination iterator to the end of the `dest_false` range.

```
#include <hpx/parallel/algorithms/for_loop.hpp>

namespace hpx

namespace parallel

namespace v2
```

Functions

```
template<typename ExPolicy, typename I, typename ...Args>
util::detail::algorithm_result<ExPolicy>::type for_loop(ExPolicy &&policy, typename
std::decay<I>::type first, I last,
Args&&... args)
```

The `for_loop` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

Requires: `I` shall be an integral type or meet the requirements of an input iterator type. The `args` parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, `f`. `f` shall meet the requirements of `MoveConstructible`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- `I`: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- `Args`: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.

- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *args*: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies *f* to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the *args* parameter pack. The length of the input sequence is *last - first*.

The first element in the input sequence is specified by *first*. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the *args* parameter pack excluding *f*, an additional argument is passed to each application of *f* as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of *f*, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of *f* in the input sequence.

Complexity: Applies *f* exactly once for each element of the input sequence.

Remarks: If *f* returns a result, the result is ignored.

Return The *for_loop* algorithm returns a `hpx::future<void>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `void` otherwise.

```
template<typename I, typename ...Args>
void for_loop(typename std::decay<I>::type first, I last, Args&&... args)
```

The *for_loop* implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble *for_each* from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

The execution of *for_loop* without specifying an execution policy is equivalent to specifying `parallel::execution::seq` as the execution policy.

Requires: *I* shall be an integral type or meet the requirements of an input iterator type. The *args* parameter pack shall have at least one element, comprising objects returned by invocations of *reduction* and/or *induction* function templates followed by exactly one element invocable element-access function, *f*. *f* shall meet the requirements of `MoveConstructible`.

Template Parameters

- *I*: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- *Args*: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *args*: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies *f* to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the *args* parameter pack. The length of the input sequence is *last - first*.

The first element in the input sequence is specified by *first*. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the *args* parameter pack excluding *f*, an additional argument is passed to each application of *f* as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using *advance* and *distance*.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of *f*, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of *f* in the input sequence.

Complexity: Applies *f* exactly once for each element of the input sequence.

Remarks: If *f* returns a result, the result is ignored.

template<typename ExPolicy, typename I, typename S, typename... Args, &&std::is_

The `for_loop_strided` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

Requires: *I* shall be an integral type or meet the requirements of an input iterator type. The *args* parameter pack shall have at least one element, comprising objects returned by invocations of *reduction* and/or *induction* function templates followed by exactly one element invocable element-access function, *f*. *f* shall meet the requirements of `MoveConstructible`.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *I*: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- *S*: The type of the stride variable. This should be an integral type.
- *Args*: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.

- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `stride`: Refers to the stride of the iteration steps. This shall have non-zero value and shall be negative only if `I` has integral type or meets the requirements of a bidirectional iterator.
- `args`: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies `f` to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the `args` parameter pack. The length of the input sequence is `last - first`.

The first element in the input sequence is specified by `first`. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the `args` parameter pack excluding `f`, an additional argument is passed to each application of `f` as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of `f`, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of `f` in the input sequence.

Complexity: Applies `f` exactly once for each element of the input sequence.

Remarks: If `f` returns a result, the result is ignored.

Return The `for_loop_strided` algorithm returns a `hpx::future<void>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `void` otherwise.

```
template<typename I, typename S, typename... Args, &&std::is_integral< S >::value
```

The `for_loop_strided` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

The execution of `for_loop` without specifying an execution policy is equivalent to specifying `parallel::execution::seq` as the execution policy.

Requires: `I` shall be an integral type or meet the requirements of an input iterator type. The `args` parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, `f`. `f` shall meet the requirements of `MoveConstructible`.

Template Parameters

- `I`: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- `S`: The type of the stride variable. This should be an integral type.
- `Args`: A parameter pack, it's last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *stride*: Refers to the stride of the iteration steps. This shall have non-zero value and shall be negative only if *I* has integral type or meets the requirements of a bidirectional iterator.
- *args*: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [*first*, *last*] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies *f* to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the *args* parameter pack. The length of the input sequence is *last - first*.

The first element in the input sequence is specified by *first*. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the *args* parameter pack excluding *f*, an additional argument is passed to each application of *f* as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of *f*, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of *f* in the input sequence.

Complexity: Applies *f* exactly once for each element of the input sequence.

Remarks: If *f* returns a result, the result is ignored.

```
template<typename ExPolicy, typename I, typename Size, typename... Args, &&std::
```

The `for_loop_n` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

Requires: *I* shall be an integral type or meet the requirements of an input iterator type. The *args* parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, *f*. *f* shall meet the requirements of `MoveConstructible`.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *I*: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- *Size*: The type of a non-negative integral value specifying the number of items to iterate over.
- *Args*: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **size**: Refers to the number of items the algorithm will be applied to.
- **args**: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies f to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the `args` parameter pack. The length of the input sequence is last - first.

The first element in the input sequence is specified by `first`. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the `args` parameter pack excluding f , an additional argument is passed to each application of f as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of f , even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of f in the input sequence.

Complexity: Applies f exactly once for each element of the input sequence.

Remarks: If f returns a result, the result is ignored.

Return The `for_loop_n` algorithm returns a `hpx::future<void>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `void` otherwise.

```
template<typename I, typename Size, typename... Args, &&std::is_integral<Size>
```

The `for_loop` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

The execution of `for_loop` without specifying an execution policy is equivalent to specifying `parallel::execution::seq` as the execution policy.

Requires: I shall be an integral type or meet the requirements of an input iterator type. The `args` parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, f . f shall meet the requirements of `MoveConstructible`.

Template Parameters

- I : The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- $Size$: The type of a non-negative integral value specifying the number of items to iterate over.
- $Args$: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *size*: Refers to the number of items the algorithm will be applied to.
- *args*: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies *f* to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the *args* parameter pack. The length of the input sequence is last - first.

The first element in the input sequence is specified by *first*. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the *args* parameter pack excluding *f*, an additional argument is passed to each application of *f* as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of *f*, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of *f* in the input sequence.

Complexity: Applies *f* exactly once for each element of the input sequence.

Remarks: If *f* returns a result, the result is ignored.

```
template<typename ExPolicy, typename I, typename Size, typename S, typename... Args>
```

The `for_loop_n_strided` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

Requires: *I* shall be an integral type or meet the requirements of an input iterator type. The *args* parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, *f*. *f* shall meet the requirements of `MoveConstructible`.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *I*: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- *Size*: The type of a non-negative integral value specifying the number of items to iterate over.
- *S*: The type of the stride variable. This should be an integral type.
- *Args*: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **size**: Refers to the number of items the algorithm will be applied to.
- **stride**: Refers to the stride of the iteration steps. This shall have non-zero value and shall be negative only if *I* has integral type or meets the requirements of a bidirectional iterator.
- **args**: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies *f* to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the *args* parameter pack. The length of the input sequence is last - first.

The first element in the input sequence is specified by *first*. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the *args* parameter pack excluding *f*, an additional argument is passed to each application of *f* as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of *f*, even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of *f* in the input sequence.

Complexity: Applies *f* exactly once for each element of the input sequence.

Remarks: If *f* returns a result, the result is ignored.

Return The `for_loop_n_strided` algorithm returns a `hpx::future<void>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `void` otherwise.

```
template<typename I, typename Size, typename S, typename... Args, &&std::is_int
```

The `for_loop_n_strided` implements loop functionality over a range specified by integral or iterator bounds. For the iterator case, these algorithms resemble `for_each` from the Parallelism TS, but leave to the programmer when and if to dereference the iterator.

The execution of `for_loop` without specifying an execution policy is equivalent to specifying `parallel::execution::seq` as the execution policy.

Requires: *I* shall be an integral type or meet the requirements of an input iterator type. The *args* parameter pack shall have at least one element, comprising objects returned by invocations of `reduction` and/or `induction` function templates followed by exactly one element invocable element-access function, *f*. *f* shall meet the requirements of `MoveConstructible`.

Template Parameters

- *I*: The type of the iteration variable. This could be an (forward) iterator type or an integral type.
- *Size*: The type of a non-negative integral value specifying the number of items to iterate over.

- S: The type of the stride variable. This should be an integral type.
- Args: A parameter pack, its last element is a function object to be invoked for each iteration, the others have to be either conforming to the induction or reduction concept.

Parameters

- first: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- size: Refers to the number of items the algorithm will be applied to.
- stride: Refers to the stride of the iteration steps. This shall have non-zero value and shall be negative only if it has integral type or meets the requirements of a bidirectional iterator.
- args: The last element of this parameter pack is the function (object) to invoke, while the remaining elements of the parameter pack are instances of either induction or reduction objects. The function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last] should expose a signature equivalent to:

```
<ignored> pred(I const& a, ...);
```

The signature does not need to have `const&`. It will receive the current value of the iteration variable and one argument for each of the induction or reduction objects passed to the algorithms, representing their current values.

Effects: Applies f to each element in the input sequence, with additional arguments corresponding to the reductions and inductions in the `args` parameter pack. The length of the input sequence is `last - first`.

The first element in the input sequence is specified by `first`. Each subsequent element is generated by incrementing the previous element.

Along with an element from the input sequence, for each member of the `args` parameter pack excluding f , an additional argument is passed to each application of f as follows:

Note As described in the C++ standard, arithmetic on non-random-access iterators is performed using `advance` and `distance`.

Note The order of the elements of the input sequence is important for determining ordinal position of an application of f , even though the applications themselves may be unordered.

If the pack member is an object returned by a call to a reduction function listed in section, then the additional argument is a reference to a view of that reduction object. If the pack member is an object returned by a call to induction, then the additional argument is the induction value for that induction object corresponding to the position of the application of f in the input sequence.

Complexity: Applies f exactly once for each element of the input sequence.

Remarks: If f returns a result, the result is ignored.

```
#include <hpx/parallel/algorithms/inclusive_scan.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Op, typename T>
util::detail::algorithm_result<ExPolicy, FwdIter2>::type inclusive_scan (ExPolicy &&policy,
                                                               FwdIter1 first,
                                                               FwdIter1 last,
                                                               FwdIter2 dest,
                                                               Op &&op,
                                                               T init)
```

Assigns through each iterator i in $[result, result + (last - first))$ the value of GENERALIZED_NONCOMMUTATIVE_SUM($op, init, *first, \dots, *(first + (i - result))$).

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate op .

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- **T**: The type of the value to be used as initial (and intermediate) values (deduced).
- **Op**: The type of the binary function object used for the reduction operation.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **init**: The initial value for the generalized sum.
- **op**: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type1 &b);
```

The signature does not need to have **const&**, but the function must not modify the objects passed to it. The types *Type1* and *Ret* must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *exclusive_scan* and *inclusive_scan* is that *inclusive_scan* includes the i th input element in the i th sum. If op is not mathematically associative, the behavior of *inclusive_scan* may be non-deterministic.

Return The *copy_n* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *inclusive_scan* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note GENERALIZED_NONCOMMUTATIVE_SUM(op, a_1, \dots, a_N) is defined as:

- a_1 when N is 1
- $op(GENERALIZED_NONCOMMUTATIVE_SUM(op, a_1, \dots, a_K), GENERAL-$

IZED_NONCOMMUTATIVE_SUM(op, aM, ..., aN)) where $1 < K+1 = M \leq N$.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Op>
util::detail::algorithm_result<ExPolicy, FwdIter2>::type inclusive_scan (ExPolicy &&policy,
                                                               FwdIter1 first,
                                                               FwdIter1 last,
                                                               FwdIter2 dest,
                                                               Op &&op)
```

Assigns through each iterator i in $[result, result + (last - first))$ the value of GENERALIZED_NONCOMMUTATIVE_SUM(op, *first, ..., *(first + (i - result))).

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate *op*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Op**: The type of the binary function object used for the reduction operation.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **op**: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have **const&**, but the function must not modify the objects passed to it. The types *Type1* and *Ret* must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *exclusive_scan* and *inclusive_scan* is that *inclusive_scan* includes the *i*th input element in the *i*th sum.

Return The *copy_n* algorithm returns a *hpx::future<FwdIter2>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *inclusive_scan* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note GENERALIZED_NONCOMMUTATIVE_SUM(+, a1, ..., aN) is defined as:

- a1 when N is 1
- GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aK)
 - GENERALIZED_NONCOMMUTATIVE_SUM(+, aM, ..., aN) where $1 < K+1 = M \leq N$.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
```

`std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F`

Assigns through each iterator i in $[result, result + (last - first))$ the value of $\text{GENERALIZED_NONCOMMUTATIVE_SUM}(+, *first, \dots, *(first + (i - result)))$.

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicate *op*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.

The reduce operations in the parallel *inclusive_scan* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between *exclusive_scan* and *inclusive_scan* is that *inclusive_scan* includes the i th input element in the i th sum.

Return The *copy_n* algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter2* otherwise. The *inclusive_scan* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note *GENERALIZED_NONCOMMUTATIVE_SUM*($+, a_1, \dots, a_N$) is defined as:

- a_1 when N is 1
- *GENERALIZED_NONCOMMUTATIVE_SUM*($+, a_1, \dots, a_K$)
 - *GENERALIZED_NONCOMMUTATIVE_SUM*($+, a_M, \dots, a_N$) where $1 < K+1 = M \leq N$.

```
#include <hpx/parallel/algorithms/rotate.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

template<typename **ExPolicy**, typename **FwdIter**>
`util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::begin(FwdIter), tag::end(FwdIter)>>`::type **rotate**
ExPolicy &&policy, **FwdIter** first, **FwdIter** new_first, **FwdIter** last
Performs a left rotation on a range of elements. Specifically, *rotate* swaps the elements in the range [first, last) in such a way that the element new_first becomes the first element of the new range and new_first - 1 becomes the last element.

The assignments in the parallel *rotate* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between *first* and *last*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **new_first**: Refers to the element that should appear at the beginning of the rotated range.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.

The assignments in the parallel *rotate* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The type of dereferenced *FwdIter* must meet the requirements of *MoveAssignable* and *MoveConstructible*.

Return The *rotate* algorithm returns a `hpx::future<tagged_pair<tag::begin(FwdIter), tag::end(FwdIter)>>` if the execution policy is of type *parallel_task_policy* and returns `tagged_pair<tag::begin(FwdIter), tag::end(FwdIter)>` otherwise. The *rotate* algorithm returns the iterator equal to pair(first + (last - new_first), last).

template<typename **ExPolicy**, typename **FwdIter1**, typename **FwdIter2**>
`util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>`::type **rotate_copy**
ExPolicy &&policy, **FwdIter1** first, **FwdIter1** new_first, **FwdIter1** last, **FwdIter2** dest_first
Copies the elements from the range [first, last), to another range beginning at *dest_first* in such a way, that the element *new_first* becomes the first element of the new range and *new_first* - 1 becomes the last element.

The assignments in the parallel *rotate_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an bidirectional iterator.
- **FwdIter2**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **new_first**: Refers to the element that should appear at the beginning of the rotated range.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest_first**: Refers to the begin of the destination range.

The assignments in the parallel *rotate_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *rotate_copy* algorithm returns a *hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>* if the execution policy is of type *parallel_task_policy* and returns *tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>* otherwise. The *rotate_copy* algorithm returns the output iterator to the element past the last element copied.

```
#include <hpx/parallel/algorithms/transform_reduce.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename T, typename Reduce, typename Convert>
util::detail::algorithm_result<ExPolicy, T>::type transform_reduce(ExPolicy &&policy,
                                                               FwdIter first, FwdIter
                                                               last, T init, Reduce
                                                               &&red_op, Convert
                                                               &&conv_op)
```

Returns GENERALIZED_SUM(red_op, init, conv_op(*first), ..., conv_op(*(first + (last - first) - 1))).

The reduce operations in the parallel *transform_reduce* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: O(*last - first*) applications of the predicates *red_op* and *conv_op*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an input iterator.

- F : The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `copy_if` requires F to meet the requirements of `CopyConstructible`.
- T : The type of the value to be used as initial (and intermediate) values (deduced).
- `Reduce`: The type of the binary function object used for the reduction operation.
- `Convert`: The type of the unary function object used to transform the elements of the input sequence before invoking the reduce function.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `conv_op`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by `[first, last]`. This is a unary predicate. The signature of this predicate should be equivalent to:

```
R fun (const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to `Type`. The type `R` must be such that an object of this type can be implicitly converted to `T`.

- `init`: The initial value for the generalized sum.
- `red_op`: Specifies the function (or function object) which will be invoked for each of the values returned from the invocation of `conv_op`. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1`, `Type2`, and `Ret` must be such that an object of a type as returned from `conv_op` can be implicitly converted to any of those types.

The reduce operations in the parallel `transform_reduce` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

The difference between `transform_reduce` and `accumulate` is that the behavior of `transform_reduce` may be non-deterministic for non-associative or non-commutative binary predicate.

Return The `transform_reduce` algorithm returns a `hpx::future<T>` if the execution policy is of type `parallel_task_policy` and returns `T` otherwise. The `transform_reduce` algorithm returns the result of the generalized sum over the values returned from `conv_op` when applied to the elements given by the input range `[first, last]`.

Note `GENERALIZED_SUM(op, a1, ..., aN)` is defined as follows:

- $a1$ when N is 1
- $op(GENERALIZED_SUM(op, b1, \dots, bK), GENERALIZED_SUM(op, bM, \dots, bN))$, where:
 - $b1, \dots, bN$ may be any permutation of $a1, \dots, aN$ and
 - $1 < K+1 = M \leq N$.

```
#include <hpx/parallel/algorithms/fill.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename T>
util::detail::algorithm_result<ExPolicy>::type fill(ExPolicy &&policy, FwdIter first, FwdIter
last, T value)
```

Assigns the given value to the elements in the range [first, last).

The comparisons in the parallel *fill* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **value**: The value to be assigned.

The comparisons in the parallel *fill* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *fill* algorithm returns a *hpx::future<void>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by *void*).

```
template<typename ExPolicy, typename FwdIter, typename Size, typename T>
util::detail::algorithm_result<ExPolicy, FwdIter>::type fill_n(ExPolicy &&policy, FwdIter
first, Size count, T value)
```

Assigns the given value value to the first count elements in the range beginning at first if count > 0. Does nothing otherwise.

The comparisons in the parallel *fill_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* assignments, for count > 0.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an output iterator.

- **Size**: The type of the argument specifying the number of elements to apply f to.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.
- **value**: The value to be assigned.

The comparisons in the parallel *fill_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *fill_n* algorithm returns a *hpx::future<void>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by *void*).

```
#include <hpx/parallel/algorithms/set_union.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename Pred, std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, FwdIter3>>::type dest, typename Compare> void set_union(FwdIter1 first1, FwdIter1 last1, FwdIter2 first2, FwdIter2 last2, Pred pred, Compare compare, ExPolicy exPolicy);
```

Constructs a sorted range beginning at *dest* consisting of all elements present in one or both sorted ranges $[first1, last1]$ and $[first2, last2]$. This algorithm expects both input ranges to be sorted with the given binary predicate f .

If some element is found m times in $[first1, last1]$ and n times in $[first2, last2]$, then all m elements will be copied from $[first1, last1]$ to *dest*, preserving order, and then exactly $\text{std}::max(n-m, 0)$ elements will be copied from $[first2, last2]$ to *dest*, also preserving order.

Note Complexity: At most $2*(N1 + N2 - 1)$ comparisons, where $N1$ is the length of the first sequence and $N2$ is the length of the second sequence.

The resulting range cannot overlap with either of the input ranges.

The application of function objects in parallel algorithm invoked with a sequential execution policy object execute in sequential order in the calling thread (*sequenced_policy*) or in a single new thread spawned from the current thread (for *sequenced_task_policy*).

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- `FwdIter1`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `Op`: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *set_union* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::less<>`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- `last1`: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- `first2`: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- `last2`: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `op`: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *InIter* can be dereferenced and then implicitly converted to *Type1*.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *set_union* algorithm returns a `hpx::future<FwdIter3>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter3* otherwise. The *set_union* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/generate.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename F>
util::detail::algorithm_result<ExPolicy, FwdIter>::type generate(ExPolicy      &&policy,
                                                               FwdIter first, FwdIter last,
                                                               F &&f)
```

Assign each element in range [first, last) a value generated by the given function object f

The assignments in the parallel *generate* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly *distance(first, last)* invocations of f and assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *f*: generator function that will be called. signature of function should be equivalent to the following:

```
Ret fun();
```

The type *Ret* must be such that an object of type *FwdIter* can be dereferenced and assigned a value of type *Ret*.

The assignments in the parallel *generate* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *last*.

```
template<typename ExPolicy, typename FwdIter, typename Size, typename F>
util::detail::algorithm_result<ExPolicy, FwdIter>::type generate_n(ExPolicy      &&policy,
                                                               FwdIter first, Size count,
                                                               F &&f)
```

Assigns each element in range [first, first+count) a value generated by the given function object g.

The assignments in the parallel *generate_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly *count* invocations of *f* and assignments, for *count* > 0.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements in the sequence the algorithm will be applied to.
- **f**: Refers to the generator function object that will be called. The signature of the function should be equivalent to

```
Ret fun();
```

The type *Ret* must be such that an object of type *OutputIt* can be dereferenced and assigned a value of type *Ret*.

The assignments in the parallel *generate_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *last*.

```
#include <hpx/parallel/algorithms/uninitialized_default_construct.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter>
util::detail::algorithm_result<ExPolicy>::type uninitialized_default_construct(ExPolicy
    &&policy,
    FwdIter
    first,
    FwdIter
    last)
```

Constructs objects of type typename iterator_traits<ForwardIt>::value_type in the uninitialized storage designated by the range [first, last) by default-initialization. If an exception is thrown during the initialization, the function has no effects.

The assignments in the parallel *uninitialized_default_construct* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.

The assignments in the parallel *uninitialized_default_construct* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_default_construct* algorithm returns a *hpx::future<void>*, if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *void* otherwise.

```
template<typename ExPolicy, typename FwdIter, typename Size>
util::detail::algorithm_result<ExPolicy, FwdIter>::type uninitialized_default_construct_n(ExPolicy
&&policy,
FwdIter
first,
Size
count)
```

Constructs objects of type *typename iterator_traits<ForwardIt>::value_type* in the uninitialized storage designated by the range [first, first + count) by default-initialization. If an exception is thrown during the initialization, the function has no effects.

The assignments in the parallel *uninitialized_default_construct_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* assignments, if *count > 0*, no assignments otherwise.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **Size**: The type of the argument specifying the number of elements to apply *f* to.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.

The assignments in the parallel *uninitialized_default_construct_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_default_construct_n* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *uninitialized_default_construct_n* algorithm returns the iterator to the element

in the source range, one past the last element constructed.

```
#include <hpx/parallel/algorithms/search.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename FwdIter2, typename Pred = detail::equal_to, type
util::detail::algorithm_result<ExPolicy, FwdIter>::type search(ExPolicy &&policy, FwdIter
first, FwdIter last, FwdIter2
s_first, FwdIter2 s_last,
Pred &&op = Pred{}, Proj1
&&proj1 = Proj1{}, Proj2
&&proj2 = Proj2{})
```

Searches the range [first, last) for any elements in the range [s_first, s_last). Uses a provided predicate to compare elements.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most (S^*N) comparisons where $S = \text{distance}(s_{\text{first}}, s_{\text{last}})$ and $N = \text{distance}(\text{first}, \text{last})$.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter:** The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an input iterator.
- **FwdIter2:** The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Pred:** The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *adjacent_find* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- **Proj1:** The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of type dereferenced *FwdIter*.
- **Proj2:** The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of type dereferenced *FwdIter2*.

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **first:** Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- **last:** Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- **s_first:** Refers to the beginning of the sequence of elements the algorithm will be searching for.
- **s_last:** Refers to the end of the sequence of elements of the algorithm will be searching for.

- **op:** Refers to the binary predicate which returns true if the elements should be treated as equal. the signature of the function should be equivalent to

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have **const &**, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively

- **proj1:** Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *FwdIter1* as a projection operation before the actual predicate *is* invoked.
- **proj2:** Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *FwdIter2* as a projection operation before the actual predicate *is* invoked.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *search* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *task_execution_policy* and returns *FwdIter* otherwise. The *search* algorithm returns an iterator to the beginning of the first subsequence [*s_first*, *s_last*) in range [first, last). If the length of the subsequence [*s_first*, *s_last*) is greater than the length of the range [first, last), *last* is returned. Additionally if the size of the subsequence is empty *first* is returned. If no subsequence is found, *last* is returned.

```
template<typename ExPolicy, typename FwdIter, typename FwdIter2, typename Pred = detail::equal_to, type
util::detail::algorithm_result<ExPolicy, FwdIter>::type search_n (ExPolicy &&policy,
FwdIter first, std::size_t
count, FwdIter2 s_first,
FwdIter2 s_last, Pred
&&op = Pred(), Proj1
&&proj1 = Proj1(), Proj2
&&proj2 = Proj2())
```

Searches the range [first, last) for any elements in the range [*s_first*, *s_last*). Uses a provided predicate to compare elements.

The comparison operations in the parallel *search_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most (S*N) comparisons where S = distance(*s_first*, *s_last*) and N = count.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter:** The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an input iterator.
- **FwdIter2:** The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Pred:** The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *adjacent_find* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::equal_to*<>

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **first:** Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- **count:** Refers to the range of elements of the first range the algorithm will be applied to.

- `s_first`: Refers to the beginning of the sequence of elements the algorithm will be searching for.
- `s_last`: Refers to the end of the sequence of elements of the algorithm will be searching for.
- `op`: Refers to the binary predicate which returns true if the elements should be treated as equal. the signature of the function should be equivalent to

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types `Type1` and `Type2` must be such that objects of types `FwdIter1` and `FwdIter2` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced `FwdIter1` as a projection operation before the actual predicate is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced `FwdIter2` as a projection operation before the actual predicate is invoked.

The comparison operations in the parallel `search_n` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `search_n` algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `task_execution_policy` and returns `FwdIter` otherwise. The `search_n` algorithm returns an iterator to the beginning of the last subsequence `[s_first, s_last]` in range `[first, first+count]`. If the length of the subsequence `[s_first, s_last]` is greater than the length of the range `[first, first+count]`, `first` is returned. Additionally if the size of the subsequence is empty or no subsequence is found, `first` is also returned.

```
#include <hpx/parallel/algorithms/copy.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out
FwdIter2>>::type copyExPolicy &&policy, FwdIter1 first, FwdIter1 last, FwdIter2 destCopies
the elements in the range, defined by [first, last), to another range beginning at dest.
```

The assignments in the parallel `copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` assignments.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.

- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.

The assignments in the parallel `copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `copy` algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `copy` algorithm returns the pair of the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename Size, typename FwdIter2>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out
FwdIter2>>::type copy_nExPolicy &&policy, FwdIter1 first, Size count, FwdIter2 destCopies
the elements in the range [first, first + count), starting from first and proceeding to first + count - 1., to another range beginning at dest.
```

The assignments in the parallel `copy_n` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `count` assignments, if `count > 0`, no assignments otherwise.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `Size`: The type of the argument specifying the number of elements to apply `f` to.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `count`: Refers to the number of elements starting at `first` the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.

The assignments in the parallel `copy_n` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `copy_n` algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `copy` algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename F, typename Proj = util::proj
```

`util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>`::type **copy_if** *ExPolicy* &&*policy*, *FwdIter1* *first*, *FwdIter1* *last*, *FwdIter2* *dest*, *F* &&*f*, *Proj* &&*proj* = *Proj*()Copies the elements in the range, defined by [first, last), to another range beginning at dest. Copies only the elements for which the predicate *f* returns true. The order of the elements that are not removed is preserved.

The assignments in the parallel *copy_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate *f*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *copy_if* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter1* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *copy_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *copy_if* algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The *copy* algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/transform_exclusive_scan.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename T, typename Op, typename Conv>
util::detail::algorithm_result<ExPolicy, FwdIter2>::type transform_exclusive_scan(ExPolicy
    &&policy,
    FwdIter1
    first,
    FwdIter1
    last,
    FwdIter2
    dest,
    T
    init,
    Op
    &&op,
    Conv
    &&conv)
```

Assigns through each iterator i in $[result, result + (last - first))$ the value of GENERALIZED_NONCOMMUTATIVE_SUM(binary_op, init, conv(*first), ..., conv(*(first + (i - result) - 1))).

The reduce operations in the parallel *transform_exclusive_scan* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: $O(last - first)$ applications of the predicates *op* and *conv*.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter1:** The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIter2:** The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- **Conv:** The type of the unary function object used for the conversion operation.
- **T:** The type of the value to be used as initial (and intermediate) values (deduced).
- **Op:** The type of the binary function object used for the reduction operation.

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **first:** Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last:** Refers to the end of the sequence of elements the algorithm will be applied to.
- **dest:** Refers to the beginning of the destination range.
- **conv:** Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a unary predicate. The signature of

this predicate should be equivalent to:

```
R fun(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`. The type `R` must be such that an object of this type can be implicitly converted to `T`.

- `init`: The initial value for the generalized sum.
- `op`: Specifies the function (or function object) which will be invoked for each of the values of the input sequence. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1` and `Ret` must be such that an object of a type as given by the input sequence can be implicitly converted to any of those types.

The reduce operations in the parallel `transform_exclusive_scan` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Neither `conv` nor `op` shall invalidate iterators or subranges, or modify elements in the ranges `[first,last)` or `[result,result + (last - first))`.

Return The `copy_n` algorithm returns a `hpx::future<FwdIter2>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter2` otherwise. The `transform_exclusive_scan` algorithm returns the output iterator to the element in the destination range, one past the last element copied.

Note `GENERALIZED_NONCOMMUTATIVE_SUM(op, a1, ..., aN)` is defined as:

- a_1 when N is 1
- $op(GENERALIZED_NONCOMMUTATIVE_SUM(op, a_1, \dots, a_K), GENERALIZED_NONCOMMUTATIVE_SUM(op, a_M, \dots, a_N))$ where $1 < K+1 = M \leq N$.

The behavior of `transform_exclusive_scan` may be non-deterministic for a non-associative predicate.

```
#include <hpx/parallel/algorithms/all_any_none.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type none_of (ExPolicy &&policy, FwdIter
                                                               first, FwdIter last, F &&f, Proj
                                                               &&proj = Proj())
```

Checks if unary predicate *f* returns true for no elements in the range [*first*, *last*).

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most *last - first* applications of the predicate *f*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *none_of* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [*first*, *last*). The signature of this predicate should be equivalent to:

```
bool pred (const Type &a);
```

The signature does not need to have *const&*, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *none_of* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *none_of* algorithm returns true if the unary predicate *f* returns true for no elements in the range, false otherwise. It returns true if the range is empty.

```
template<typename ExPolicy, typename FwdIter, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type any_of (ExPolicy &&policy, FwdIter first,
                                                               FwdIter last, F &&f, Proj &&proj
                                                               = Proj())
```

Checks if unary predicate *f* returns true for at least one element in the range [*first*, *last*).

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most *last - first* applications of the predicate *f*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *any_of* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *any_of* algorithm returns a `hpx::future<bool>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `bool` otherwise. The *any_of* algorithm returns true if the unary predicate *f* returns true for at least one element in the range, false otherwise. It returns false if the range is empty.

```
template<typename ExPolicy, typename FwdIter, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type all_of(ExPolicy &&policy, FwdIter first,
                                                       FwdIter last, F &&f, Proj &&proj
                                                       = Proj())
```

Checks if unary predicate *f* returns true for all elements in the range [first, last).

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most *last - first* applications of the predicate *f*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *all_of* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.

- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `all_of` algorithm returns a `hpx::future<bool>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `bool` otherwise. The `all_of` algorithm returns true if the unary predicate *f* returns true for all elements in the range, false otherwise. It returns true if the range is empty.

```
#include <hpx/parallel/algorithms/reverse.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename BidirIter>
util::detail::algorithm_result<ExPolicy, BidirIter>::type reverse(ExPolicy      &&policy,
                                                               BidirIter   first,   BidirIter
                                                               last)
```

Reverses the order of the elements in the range [first, last). Behaves as if applying `std::iter_swap` to every pair of iterators `first+i, (last-i) - 1` for each non-negative $i < (\text{last}-\text{first})/2$.

The assignments in the parallel `reverse` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between `first` and `last`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `BidirIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an bidirectional iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.

The assignments in the parallel *reverse* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *reverse* algorithm returns a *hpx::future<BidirIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *BidirIter* otherwise. It returns *last*.

```
template<typename ExPolicy, typename BidirIter, typename FwdIter>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(BidirIter), tag::out(FwdIter)>>::type reverse_copyExPolicy &&policy, BidirIter first, BidirIter last, FwdIter dest_firstCopies the elements from the range [first, last) to another range beginning at dest_first in such a way that the elements in the new range are in reverse order. Behaves as if by executing the assignment *(dest_first + (last - first) - 1 - i) = *(first + i) once for each non-negative i < (last - first). If the source and destination ranges (that is, [first, last) and [dest_first, dest_first+(last-first)) respectively) overlap, the behavior is undefined.
```

The assignments in the parallel *reverse_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *BidirIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an bidirectional iterator.
- *FwdIter*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *dest_first*: Refers to the begin of the destination range.

The assignments in the parallel *reverse_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *reverse_copy* algorithm returns a *hpx::future<tagged_pair<tag::in(BidirIter), tag::out(FwdIter)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::in(BidirIter), tag::out(FwdIter)>* otherwise. The *copy* algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/uninitialized_fill.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename T>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy>>
```

Copies the given *value* to an uninitialized memory area, defined by the range [first, last). If an exception is thrown during the initialization, the function has no effects.

The initializations in the parallel *uninitialized_fill* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between *first* and *last*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **value**: The value to be assigned.

The initializations in the parallel *uninitialized_fill* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_fill* algorithm returns a *hpx::future<void>*, if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns nothing otherwise.

```
template<typename ExPolicy, typename FwdIter, typename Size, typename T>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy>>
```

Copies the given *value* value to the first count elements in an uninitialized memory area beginning at *first*. If an exception is thrown during the initialization, the function has no effects.

The initializations in the parallel *uninitialized_fill_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* assignments, if *count* > 0, no assignments otherwise.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- **Size**: The type of the argument specifying the number of elements to apply *f* to.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.
- **value**: The value to be assigned.

The initializations in the parallel *uninitialized_fill_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *uninitialized_fill_n* algorithm returns a *hpx::future<void>*, if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns nothing otherwise.

```
#include <hpx/parallel/algorithms/set_difference.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename Pre  
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Constructs a sorted range beginning at *dest* consisting of all elements present in the range [*first1*,

`last1`) and not present in the range $[first2, last2]$. This algorithm expects both input ranges to be sorted with the given binary predicate f .

Equivalent elements are treated individually, that is, if some element is found m times in $[first1, last1]$ and n times in $[first2, last2]$, it will be copied to $dest$ exactly $\text{std}::\text{max}(m-n, 0)$ times. The resulting range cannot overlap with either of the input ranges.

Note Complexity: At most $2*(N1 + N2 - 1)$ comparisons, where $N1$ is the length of the first sequence and $N2$ is the length of the second sequence.

The resulting range cannot overlap with either of the input ranges.

The application of function objects in parallel algorithm invoked with a sequential execution policy object execute in sequential order in the calling thread (*sequenced_policy*) or in a single new thread spawned from the current thread (for *sequenced_task_policy*).

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- `FwdIter1`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators used (deduced) representing the first sequence. This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `Pred`: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *set_difference* requires `Pred` to meet the requirements of *CopyConstructible*. This defaults to `std::less<>`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- `last1`: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- `first2`: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- `last2`: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `op`: The binary predicate which returns true if the elements should be treated as equal. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type `Type1` must be such that objects of type `InIter` can be dereferenced and then implicitly converted to `Type1`

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *set_difference* algorithm returns a `hpx::future<FwdIter3>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `FwdIter3` otherwise. The *set_difference* algorithm returns the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/is_partitioned.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Pred>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, b
```

Determines if the range [first, last) is partitioned.

The predicate operations in the parallel *is_partitioned* algorithm invoked with an execution policy object of type *sequenced_policy* executes in sequential order in the calling thread.

Note Complexity: at most (N) predicate evaluations where $N = \text{distance}(\text{first}, \text{last})$.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used for the This iterator type must meet the requirements of a forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements of that the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements of that the algorithm will be applied to.
- **pred**: Refers to the binary predicate which returns true if the first argument should be treated as less than the second argument. The signature of the function should be equivalent to

```
bool pred(const Type &a, const Type &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type *Type* must be such that objects of types *FwdIter* can be dereferenced and then implicitly converted to *Type*.

The comparison operations in the parallel *is_partitioned* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_partitioned* algorithm returns a `hpx::future<bool>` if the execution policy is of type *task_execution_policy* and returns `bool` otherwise. The *is_partitioned* algorithm returns true if each element in the sequence for which *pred* returns true precedes those for which *pred*

returns false. Otherwise `is_partitioned` returns false. If the range [first, last) contains less than two elements, the function is always true.

```
#include <hpx/parallel/algorithms/mismatch.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, s>
```

Returns true if the range [first1, last1) is mismatch to the range [first2, last2), and false otherwise.

The comparison operations in the parallel *mismatch* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most $\min(\text{last1} - \text{first1}, \text{last2} - \text{first2})$ applications of the predicate *f*. If *FwdIter1* and *FwdIter2* meet the requirements of *RandomAccessIterator* and $(\text{last1} - \text{first1}) \neq (\text{last2} - \text{first2})$ then no applications of the predicate *f* are made.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *mismatch* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.

- `last1`: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- `first2`: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- `last2`: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- `op`: The binary predicate which returns true if the elements should be treated as mismatch. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types `Type1` and `Type2` must be such that objects of types `FwdIter1` and `FwdIter2` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively. The comparison operations in the parallel *mismatch* algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The two ranges are considered mismatch if, for every iterator `i` in the range `[first1, last1)`, `*i mismatchs *(first2 + (i - first1))`. This overload of *mismatch* uses operator`==` to determine if two elements are mismatch.

Return The *mismatch* algorithm returns a `hpx::future<bool>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `bool` otherwise. The *mismatch* algorithm returns true if the elements in the two ranges are mismatch, otherwise it returns false. If the length of the range `[first1, last1)` does not mismatch the length of the range `[first2, last2)`, it returns false.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, s
```

Returns `std::pair` with iterators to the first two non-equivalent elements.

The comparison operations in the parallel *mismatch* algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: At most `last1 - first1` applications of the predicate `f`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *mismatch* requires `Pred` to meet the requirements of `CopyConstructible`.

structible. This defaults to std::equal_to<>

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last1*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *first2*: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- *op*: The binary predicate which returns true if the elements should be treated as mismatch. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have const &, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively. The comparison operations in the parallel *mismatch* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *mismatch* algorithm returns a *hpx::future<std::pair<FwdIter1, FwdIter2> >* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *std::pair<FwdIter1, FwdIter2>* otherwise. The *mismatch* algorithm returns the first mismatching pair of elements from two ranges: one defined by [first1, last1) and another defined by [first2, last2).

```
#include <hpx/parallel/algorithms/sort_by_key.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename KeyIter, typename ValueIter, typename Compare = detail::less<
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in1 (KeyIter), tag::in2
ValueIter>>::type sort_by_keyExPolicy &&policy, KeyIter key_first, KeyIter key_last, ValueIter value_first, Compare &&comp = Compare()>>Sorts one range of data using keys supplied in another range. The key elements in the range [key_first, key_last) are sorted in ascending order with the corresponding elements in the value range moved to follow the sorted order. The algorithm is not stable, the order of equal elements is not guaranteed to be preserved. The function uses the given comparison function object comp (defaults to using operator<()).
```

A sequence is sorted with respect to a comparator *comp* and a projection *proj* if for every iterator *i* pointing to the sequence and every non-negative integer *n* such that *i + n* is a valid iterator pointing to an element of the sequence, and *INVOKE(comp, INVOKE(proj, *(i + n)), INVOKE(proj, *i)) == false*.

Note Complexity: O(Nlog(N)), where N = std::distance(first, last) comparisons. *comp* has to induce a strict weak ordering on the values.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **KeyIter**: The type of the key iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- **ValueIter**: The type of the value iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- **Comp**: The type of the function/function object to use (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **key_first**: Refers to the beginning of the sequence of key elements the algorithm will be applied to.
- **key_last**: Refers to the end of the sequence of key elements the algorithm will be applied to.
- **value_first**: Refers to the beginning of the sequence of value elements the algorithm will be applied to, the range of elements must match [key_first, key_last)
- **comp**: comp is a callable object. The return value of the INVOKE operation applied to an object of type Comp, when contextually converted to bool, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that comp will not apply any non-constant function through the dereferenced iterator.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *sort_by-key* algorithm returns a *hpx::future<tagged_pair<tag::in1(KeyIter), tag::in2(ValueIter)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *otherwise*. The algorithm returns a pair holding an iterator pointing to the first element after the last element in the input key sequence and an iterator pointing to the first element after the last element in the input value sequence.

```
#include <hpx/parallel/algorithms/remove.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Pred, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type remove_if(ExPolicy           &&policy,
                                                               FwdIter      first,   FwdIter
                                                               last,        Pred   &&pred, Proj
                                                               &&proj = Proj())
```

Removes all elements satisfying specific criteria from the range [first, last) and returns a past-the-end iterator for the new end of the range. This version removes all elements for which predicate *pred* returns true.

The assignments in the parallel `remove_if` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate `pred` and the projection `proj`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `remove_if` requires `Pred` to meet the requirements of `CopyConstructible`.
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `pred`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an unary predicate which returns `true` for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `remove_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `remove_if` algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter` otherwise. The `remove_if` algorithm returns the iterator to the new end of the range.

```
template<typename ExPolicy, typename FwdIter, typename T, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type remove(ExPolicy &&policy, FwdIter
first, FwdIter last, T const
&value, Proj &&proj =
Proj())
```

Removes all elements satisfying specific criteria from the range [first, last) and returns a past-the-end iterator for the new end of the range. This version removes all elements that are equal to `value`.

The assignments in the parallel `remove` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the operator`==()` and the projection `proj`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.

- `FwdIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.

- `T`: The type of the value to remove (deduced). This value type must meet the requirements of `CopyConstructible`.

- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `value`: Specifies the value of elements to remove.
- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove* algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter` otherwise. The *remove* algorithm returns the iterator to the new end of the range.

```
#include <hpx/parallel/algorithms/minmax.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Proj = util::projection_identity, typename F = detail::util::detail::algorithm_result<ExPolicy, FwdIter>::type> min_element(ExPolicy &&policy,
    FwdIter first, FwdIter last, F &&f = F(), Proj
    &&proj = Proj())
```

Finds the smallest element in the range [first, last) using the given comparison function *f*.

The comparisons in the parallel *min_element* algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly $\max(N-1, 0)$ comparisons, where $N = \text{std::distance(first, last)}$.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *min_element* requires *F* to meet the requirements of `CopyConstructible`.
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.

- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **f**: The binary predicate which returns true if the left argument is less than the right element. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type `Type1` must be such that objects of type `FwdIter` can be dereferenced and then implicitly converted to `Type1`.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *f* is invoked.

The comparisons in the parallel `min_element` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `min_element` algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter` otherwise. The `min_element` algorithm returns the iterator to the smallest element in the range [first, last). If several elements in the range are equivalent to the smallest element, returns the iterator to the first such element. Returns last if the range is empty.

```
template<typename ExPolicy, typename FwdIter, typename Proj = util::projection_identity, typename F = detailed::projection_fn<ExPolicy, FwdIter>>::type max_element (ExPolicy &&policy,  
                                FwdIter first, FwdIter  
                                last, F &&f = F(), Proj  
                                &&proj = Proj())
```

Finds the greatest element in the range [first, last) using the given comparison function *f*.

The comparisons in the parallel `max_element` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly $\max(N-1, 0)$ comparisons, where $N = \text{std::distance}(\text{first}, \text{last})$.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `max_element` requires *F* to meet the requirements of `CopyConstructible`.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **f**: The binary predicate which returns true if the left argument is less than the right element. This argument is optional and defaults to `std::less`. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type `Type1` must be such that objects of type `FwdIter` can be dereferenced and then implicitly converted to `Type1`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The comparisons in the parallel *max_element* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *max_element* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *max_element* algorithm returns the iterator to the smallest element in the range [first, last). If several elements in the range are equivalent to the smallest element, returns the iterator to the first such element. Returns last if the range is empty.

```
template<typename ExPolicy, typename FwdIter, typename Proj = util::projection_identity, typename F = detail::util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::min(FwdIter), tag::max(FwdIter)>>::type minmax_element(ExPolicy &&policy, FwdIter first, FwdIter last, F &&f = F(), Proj &&proj = Proj())
```

Finds the greatest element in the range [first, last) using the given comparison function *f*.

The comparisons in the parallel *minmax_element* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most $\max(\text{floor}(3/2*(N-1)), 0)$ applications of the predicate, where $N = \text{std::distance(first, last)}$.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of a forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *minmax_element* requires *F* to meet the requirements of *CopyConstructible*.
- `Proj`: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `f`: The binary predicate which returns true if the left argument is less than the right element. This argument is optional and defaults to `std::less`. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type1*.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The comparisons in the parallel *minmax_element* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *minmax_element* algorithm returns a *hpx::future<tagged_pair<tag::min(FwdIter), tag::max(FwdIter)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::min(FwdIter), tag::max(FwdIter)>* otherwise. The *minmax_element* algorithm returns a pair consisting of an iterator to the smallest el-

ement as the first element and an iterator to the greatest element as the second. Returns std::make_pair(first, first) if the range is empty. If several elements are equivalent to the smallest element, the iterator to the first such element is returned. If several elements are equivalent to the largest element, the iterator to the last such element is returned.

```
#include <hpx/parallel/algorithms/unique.hpp>

namespace hpx

    namespace parallel

        namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Pred = detail::equal_to, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type unique(ExPolicy &&policy, FwdIter
                                                               first, FwdIter last, Pred
                                                               &&pred = Pred(), Proj
                                                               &&proj = Proj())
```

Eliminates all but the first element from every consecutive group of equivalent elements from the range [first, last) and returns a past-the-end iterator for the new logical end of the range.

The assignments in the parallel *unique* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first - 1* applications of the predicate *pred* and no more than twice as many applications of the projection *proj*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **Pred**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *unique* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to std::equal_to<>
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last**: Refers to the end of the sequence of elements the algorithm will be applied to.
- **pred**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an binary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have *const&*, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter* can be dereferenced and then implicitly converted to both *Type1* and *Type2*

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *unique* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *unique* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *unique* algorithm returns the iterator to the new end of the range.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::equal_to, typename util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out FwdIter2>>::type unique_copy ExPolicy &&policy, FwdIter1 first, FwdIter1 last, FwdIter2 dest, Pred &&pred = Pred(), Proj &&proj = Proj()Copies the elements from the range [first, last), to another range beginning at dest in such a way that there are no consecutive equal elements. Only the first element of each group of equal elements is copied.
```

The assignments in the parallel *unique_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first - 1* applications of the predicate *pred* and no more than twice as many applications of the projection *proj*

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *unique_copy* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `pred`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a binary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a, const Type &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter1* can be dereferenced and then implicitly converted to *Type*.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *unique_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *unique_copy* algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1),`

`tag::out(FwdIter2) > >` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `unique_copy` algorithm returns the pair of the source iterator to `last`, and the destination iterator to the end of the `dest` range.

```
#include <hpx/parallel/algorithms/adjacent_find.hpp>

namespace hpx

    namespace parallel

        namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Pred = detail::equal_to>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, F
```

Searches the range [first, last) for two consecutive identical elements. This version uses the given binary predicate op

The comparison operations in the parallel `adjacent_find` invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly the smaller of (`result - first`) + 1 and (`last - first`) - 1 application of the predicate where `result` is the value returned

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `FwdIter`: The type of the source iterators used for the range (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of `adjacent_find` requires `Pred` to meet the requirements of `CopyConstructible`. This defaults to `std::equal_to<>`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements of the range the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements of the range the algorithm will be applied to.
- `op`: The binary predicate which returns `true` if the elements should be treated as equal. The signature should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types `Type1` must be such that objects of type `FwdIter` can be dereferenced and then implicitly converted to `Type1`.

The comparison operations in the parallel `adjacent_find` invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

This overload of `adjacent_find` is available if the user decides to provide their algorithm their own binary predicate `op`.

Return The `adjacent_find` algorithm returns a `hpx::future<InIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `InIter` otherwise. The `adjacent_find` algorithm returns an iterator to the first of the identical elements. If no such elements are found, `last` is returned.

```
#include <hpx/parallel/algorithms/transform.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename F, typename Proj = util::proj_t<util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>>::type transformExPolicy &&policy, FwdIter1 first, FwdIter1 last, FwdIter2 dest, F &&f, Proj &&proj = Proj()Applies the given function f to the range [first, last) and stores the result in another range, beginning at dest.
```

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly $last - first$ applications of `f`

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of `f`.
- `FwdIter1`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `transform` requires `F` to meet the requirements of `CopyConstructible`.
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first`: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- `last`: Refers to the end of the sequence of elements the algorithm will be applied to.

- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type &a);
```

The signature does not need to have `const&`. The type `Type` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to `Type`. The type `Ret` must be such that an object of type `FwdIter2` can be dereferenced and assigned a value of type `Ret`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform` algorithm returns a `hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>` if the execution policy is of type `parallel_task_policy` and returns `tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>` otherwise. The `transform` algorithm returns a tuple holding an iterator referring to the first element after the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename F, typename util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1(FwdIter1), tag::in2(FwdIter2), tag::out(FwdIter3)>>>::type transformExPolicy &&policy, FwdIter1 first1, FwdIter1 last1, FwdIter2 first2, FwdIter3 dest, F &&f, Proj1 &&proj1 = Proj1(), Proj2 &&proj2 = Proj2()Applies the given function f to pairs of elements from two ranges: one defined by [first1, last1) and the other beginning at first2, and stores the result in another range, beginning at dest.
```

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly `last - first` applications of `f`

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of `f`.
- `FwdIter1`: The type of the source iterators for the first range used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators for the second range used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `transform` requires `F` to meet the requirements of `CopyConstructible`.
- `Proj1`: The type of an optional projection function to be used for elements of the first sequence. This defaults to `util::projection_identity`
- `Proj2`: The type of an optional projection function to be used for elements of the second sequence. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `first1`: Refers to the beginning of the first sequence of elements the algorithm will be applied to.
- `last1`: Refers to the end of the first sequence of elements the algorithm will be applied to.
- `first2`: Refers to the beginning of the second sequence of elements the algorithm will be applied to.

- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`. The types `Type1` and `Type2` must be such that objects of types `FwdIter1` and `FwdIter2` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively. The type `Ret` must be such that an object of type `FwdIter3` can be dereferenced and assigned a value of type `Ret`.

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of the first sequence as a projection operation before the actual predicate `f` is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of the second sequence as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform` algorithm returns a `hpx::future<tagged_tuple<tag::in1(FwdIter1), tag::in2(FwdIter2), tag::out(FwdIter3)>>` if the execution policy is of type `parallel_task_policy` and returns `tagged_tuple<tag::in1(FwdIter1), tag::in2(FwdIter2), tag::out(FwdIter3)>` otherwise. The `transform` algorithm returns a tuple holding an iterator referring to the first element after the first input sequence, an iterator referring to the first element after the second input sequence, and the output iterator referring to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename FwdIter3, typename F, typename util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1 (FwdIter1), tag::in2
```

```
FwdIter2, tag::outFwdIter3>>::type transformExPolicy &&policy, FwdIter1 first1, FwdIter1 last1, FwdIter2 first2, FwdIter2 last2, FwdIter3 dest, F &&f, Proj1 &&&proj1 = Proj1(), Proj2 &&&proj2 = Proj2()Applies the given function f to pairs of elements from two ranges: one defined by [first1, last1) and the other beginning at first2, and stores the result in another range, beginning at dest.
```

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly $\min(\text{last2}-\text{first2}, \text{last1}-\text{first1})$ applications of `f`

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of `f`.
- `FwdIter1`: The type of the source iterators for the first range used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the source iterators for the second range used (deduced). This iterator type must meet the requirements of an forward iterator.
- `FwdIter3`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `transform` requires `F` to meet the requirements of `CopyConstructible`.
- `Proj1`: The type of an optional projection function to be used for elements of the first sequence. This defaults to `util::projection_identity`
- `Proj2`: The type of an optional projection function to be used for elements of the second sequence. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first1**: Refers to the beginning of the first sequence of elements the algorithm will be applied to.
- **last1**: Refers to the end of the first sequence of elements the algorithm will be applied to.
- **first2**: Refers to the beginning of the second sequence of elements the algorithm will be applied to.
- **last2**: Refers to the end of the second sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively. The type *Ret* must be such that an object of type *FwdIter3* can be dereferenced and assigned a value of type *Ret*.

- **proj1**: Specifies the function (or function object) which will be invoked for each of the elements of the first sequence as a projection operation before the actual predicate *f* is invoked.
- **proj2**: Specifies the function (or function object) which will be invoked for each of the elements of the second sequence as a projection operation before the actual predicate *f* is invoked.

The invocations of *f* in the parallel *transform* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The algorithm will invoke the binary predicate until it reaches the end of the shorter of the two given input sequences

Return The *transform* algorithm returns a `hpx::future<tagged_tuple<tag::in1(FwdIter1), tag::in2(FwdIter2), tag::out(FwdIter3)>>` if the execution policy is of type *parallel_task_policy* and returns `tagged_tuple<tag::in1(FwdIter1), tag::in2(FwdIter2), tag::out(FwdIter3)>` otherwise. The *transform* algorithm returns a tuple holding an iterator referring to the first element after the first input sequence, an iterator referring to the first element after the second input sequence, and the output iterator referring to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/algorithms/count.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIterB, typename FwdIterE, typename T, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename std::iterator_traits<FwdIterB>::difference_type>::type count(ExPolicy const& policy, FwdIterB first, FwdIterE last, T value, Proj proj = Proj{})
```

Returns the number of elements in the range [first, last) satisfying a specific criteria. This version counts the elements that are equal to the given *value*.

The comparisons in the parallel *count* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* comparisons.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the comparisons.
- **FwdIterB:** The type of the source begin iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- **FwdIterE:** The type of the source end iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- **T:** The type of the value to search for (deduced).
- **Proj:** The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **first:** Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **last:** Refers to the end of the sequence of elements the algorithm will be applied to.
- **value:** The value to search for.
- **proj:** Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

Note The comparisons in the parallel *count* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *count* algorithm returns a `hpx::future<difference_type>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by `std::iterator_traits<FwdIterB>::difference_type`). The *count* algorithm returns the number of elements satisfying the given criteria.

```
template<typename ExPolicy, typename FwdIterB, typename FwdIterE, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, F, typename std::iterator_traits<FwdIterB>::difference_type>::type count(ExPolicy const& policy, FwdIterB first, FwdIterE last, F value, Proj proj = Proj{})
```

`util::detail::algorithm_result<ExPolicy, typename std::iterator_traits<FwdIterB>::difference_type>::type count_if`

Returns the number of elements in the range [first, last) satisfying a specific criteria. This version counts elements for which predicate *f* returns true.

Note Complexity: Performs exactly *last - first* applications of the predicate.

Note The assignments in the parallel *count_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note The assignments in the parallel *count_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *count_if* algorithm returns *hpx::future<difference_type>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by *std::iterator_traits<FwdIterB>::difference_type*). The *count* algorithm returns the number of elements satisfying the given criteria.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the comparisons.
- *FwdIterB*: The type of the source begin iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIterE*: The type of the source end iterator used (deduced). This iterator type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *count_if* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have *const&*, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIterB* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

```
#include <hpx/parallel/algorithms/for_each.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter, typename Size, typename F, typename Proj = util::projection_identities::identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type for_each_n(ExPolicy      &&policy,
                                                               FwdIter first, Size count,
                                                               F &&f, Proj &&proj =
                                                               Proj())
```

Applies f to the result of dereferencing every iterator in the range [first, first + count), starting from first and proceeding to first + count - 1.

If f returns a result, the result is ignored.

Note Complexity: Applies f exactly $count$ times.

If the type of $first$ satisfies the requirements of a mutable iterator, f may apply non-constant functions through the dereferenced iterator.

Unlike its sequential form, the parallel overload of *for_each* does not return a copy of its *Function* parameter, since parallelization may not permit efficient state accumulation.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **FwdIter**: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- **Size**: The type of the argument specifying the number of elements to apply f to.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *for_each* requires F to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first**: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
<ignored> pred(const Type &a);
```

The signature does not need to have `const&`. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate f is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *for_each_n* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *first + count* for non-negative values of *count* and *first* for negative values.

```
template<typename ExPolicy, typename FwdIter, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, FwdIter>::type for_each(ExPolicy &&policy,
                                                               FwdIter first, FwdIter last,
                                                               F &&f, Proj &&proj =
                                                               Proj())
```

Applies *f* to the result of dereferencing every iterator in the range [first, last).

If *f* returns a result, the result is ignored.

Note Complexity: Applies *f* exactly *last - first* times.

If the type of *first* satisfies the requirements of a mutable iterator, *f* may apply non-constant functions through the dereferenced iterator.

Unlike its sequential form, the parallel overload of *for_each* does not return a copy of its *Function* parameter, since parallelization may not permit efficient state accumulation.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *FwdIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *for_each* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
<ignored> pred(const Type &a);
```

The signature does not need to have *const&*. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *f* is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *for_each* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *last*.

```
#include <hpx/parallel/algorithms/sort.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename RandomIt, typename Proj = util::projection_identity, typename Compare = util::less<RandomIt>, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, RandomIt>::type sort(ExPolicy &&policy, RandomIt first, RandomIt last, Compare &&comp = Compare{}, Proj &&proj = Proj{})
```

Sorts the elements in the range [first, last) in ascending order. The order of equal elements is not guaranteed to be preserved. The function uses the given comparison function object comp (defaults to using operator<()).

A sequence is sorted with respect to a comparator *comp* and a projection *proj* if for every iterator *i* pointing to the sequence and every non-negative integer *n* such that *i + n* is a valid iterator pointing to an element of the sequence, and $\text{INVOKE}(\text{comp}, \text{INVOKED}(\text{proj}, *(\text{i} + \text{n})), \text{INVOKED}(\text{proj}, *\text{i})) == \text{false}$.

Note Complexity: $O(N\log(N))$, where $N = \text{std}:\text{distance}(\text{first}, \text{last})$ comparisons.

comp has to induce a strict weak ordering on the values.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *Iter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of a random access iterator.
- *Comp*: The type of the function/function object to use (deduced).
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the sequence of elements the algorithm will be applied to.
- *last*: Refers to the end of the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the `INVOKED` operation applied to an object of type *Comp*, when contextually converted to `bool`, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each pair of elements as a projection operation before the actual predicate *comp* is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *sort* algorithm returns a `hpx::future<RandomIt>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *RandomIt* otherwise. The algorithm

returns an iterator pointing to the first element after the last element in the input sequence.

```
#include <hpx/parallel/algorithms/merge.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

template<typename **ExPolicy**, typename **RandIter1**, typename **RandIter2**, typename **RandIter3**, typename **util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1 (RandIter1), tag::in2**

RandIter2, tag::outRandIter3>>::type **merge>&&policy, RandIter1 first1, RandIter1 last1, RandIter2 first2, RandIter2 last2, RandIter3 dest, Comp &&comp = **Comp()**, Proj1 &&proj1 = **Proj1()**, Proj2 &&proj2 = **Proj2()**)Merges two sorted ranges [first1, last1) and [first2, last2) into one sorted range beginning at dest. The order of equivalent elements in the each of original two ranges is preserved. For equivalent elements in the original two ranges, the elements from the first range precede the elements from the second range. The destination range cannot overlap with either of the input ranges.**

The assignments in the parallel *merge* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs $O(\text{std}::\text{distance}(\text{first1}, \text{last1}) + \text{std}::\text{distance}(\text{first2}, \text{last2}))$ applications of the comparison *comp* and the each projection.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **RandIter1**: The type of the source iterators used (deduced) representing the first sorted range. This iterator type must meet the requirements of an random access iterator.
- **RandIter2**: The type of the source iterators used (deduced) representing the second sorted range. This iterator type must meet the requirements of an random access iterator.
- **RandIter3**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an random access iterator.
- **Comp**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *merge* requires *Comp* to meet the requirements of *CopyConstructible*. This defaults to *std::less*<
- **Proj1**: The type of an optional projection function to be used for elements of the first range. This defaults to *util::projection_identity*
- **Proj2**: The type of an optional projection function to be used for elements of the second range. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **first1**: Refers to the beginning of the first range of elements the algorithm will be applied to.
- **last1**: Refers to the end of the first range of elements the algorithm will be applied to.
- **first2**: Refers to the beginning of the second range of elements the algorithm will be applied to.
- **last2**: Refers to the end of the second range of elements the algorithm will be applied to.

- *dest*: Refers to the beginning of the destination range.
- *comp*: *comp* is a callable object which returns true if the first argument is less than the second, and false otherwise. The signature of this comparison should be equivalent to:

```
bool comp(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *RandIter1* and *RandIter2* can be dereferenced and then implicitly converted to both *Type1* and *Type2*.

- *proj1*: Specifies the function (or function object) which will be invoked for each of the elements of the first range as a projection operation before the actual comparison *comp* is invoked.
- *proj2*: Specifies the function (or function object) which will be invoked for each of the elements of the second range as a projection operation before the actual comparison *comp* is invoked.

The assignments in the parallel *merge* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *merge* algorithm returns a `hpx::future<tagged_tuple<tag::in1(RandIter1), tag::in2(RandIter2), tag::out(RandIter3)>>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `tagged_tuple<tag::in1(RandIter1), tag::in2(RandIter2), tag::out(RandIter3)>` otherwise. The *merge* algorithm returns the tuple of the source iterator *last1*, the source iterator *last2*, the destination iterator to the end of the *dest* range.

```
template<typename ExPolicy, typename RandIter, typename Comp = std::less, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, RandIter>::type inplace_merge(ExPolicy &&policy, RandIter first, RandIter middle, RandIter last, Comp &&comp = Comp{}, Proj &&proj = Proj{})
```

Merges two consecutive sorted ranges [*first*, *middle*] and [*middle*, *last*] into one sorted range [*first*, *last*]. The order of equivalent elements in the each of original two ranges is preserved. For equivalent elements in the original two ranges, the elements from the first range precede the elements from the second range.

The assignments in the parallel *inplace_merge* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs $O(\text{std}::\text{distance}(\text{first}, \text{last}))$ applications of the comparison *comp* and the each projection.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *RandIter*: The type of the source iterators used (deduced). This iterator type must meet the requirements of an random access iterator.
- *Comp*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *inplace_merge* requires *Comp* to meet the requirements of *CopyConstructible*. This defaults to `std::less<>`
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first*: Refers to the beginning of the first sorted range the algorithm will be applied to.

- **middle**: Refers to the end of the first sorted range and the beginning of the second sorted range the algorithm will be applied to.
- **last**: Refers to the end of the second sorted range the algorithm will be applied to.
- **comp**: *comp* is a callable object which returns true if the first argument is less than the second, and false otherwise. The signature of this comparison should be equivalent to:

```
bool comp(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *RandIter* can be dereferenced and then implicitly converted to both *Type1* and *Type2*

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *inplace_merge* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *inplace_merge* algorithm returns a `hpx::future<RandIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *RandIter* otherwise. The *inplace_merge* algorithm returns the source iterator *last*

```
#include <hpx/parallel/algorithms/includes.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename FwdIter1, typename FwdIter2, typename Pred = detail::less>
std::enable_if<execution::is_execution_policy<ExPolicy>::value, typename util::detail::algorithm_result<ExPolicy, b>
```

Returns true if every element from the sorted range [first2, last2) is found within the sorted range [first1, last1). Also returns true if [first2, last2) is empty. The version expects both ranges to be sorted with the user supplied binary predicate *f*.

The comparison operations in the parallel *includes* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note At most $2*(N1+N2-1)$ comparisons, where $N1 = \text{std}::\text{distance}(\text{first1}, \text{last1})$ and $N2 = \text{std}::\text{distance}(\text{first2}, \text{last2})$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *FwdIter1*: The type of the source iterators used for the first range (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the source iterators used for the second range (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *includes* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::less<>*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *first1*: Refers to the beginning of the sequence of elements of the first range the algorithm will be applied to.
- *last1*: Refers to the end of the sequence of elements of the first range the algorithm will be applied to.
- *first2*: Refers to the beginning of the sequence of elements of the second range the algorithm will be applied to.
- *last2*: Refers to the end of the sequence of elements of the second range the algorithm will be applied to.
- *op*: The binary predicate which returns true if the elements should be treated as includes. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively. The comparison operations in the parallel *includes* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *includes* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *includes* algorithm returns true every element from the sorted range $[\text{first2}, \text{last2}]$ is found within the sorted range $[\text{first1}, \text{last1}]$. Also returns true if $[\text{first2}, \text{last2}]$ is empty.

```
#include <hpx/parallel/algorithms/for_loop_reduction.hpp>

namespace hpx

    namespace parallel

        namespace v2
```

Functions

```
template<typename T, typename Op>
detail::reduction_helper<T, typename std::decay<Op>::type> reduction(T &var, T
                                                               const &identity, Op &&com-
                                                               biner)
```

The function template returns a reduction object of unspecified type having a value type and encapsulating an identity value for the reduction, a combiner function object, and a live-out object from which the initial value is obtained and into which the final value is stored.

A parallel algorithm uses reduction objects by allocating an unspecified number of instances, called views, of the reduction's value type. Each view is initialized with the reduction object's identity value, except that the live-out object (which was initialized by the caller) comprises one of the views. The algorithm passes a reference to a view to each application of an element-access function, ensuring that no two concurrently-executing invocations share the same view. A view can be shared between two applications that do not execute concurrently, but initialization is performed only once per view.

Modifications to the view by the application of element access functions accumulate as partial results. At some point before the algorithm returns, the partial results are combined, two at a time, using the reduction object's combiner operation until a single value remains, which is then assigned back to the live-out object.

`T` shall meet the requirements of `CopyConstructible` and `MoveAssignable`. The expression `var = combiner(var, var)` shall be well formed.

Template Parameters

- `T`: The value type to be used by the induction object.
- `Op`: The type of the binary function (object) used to perform the reduction operation.

Parameters

- `var`: [in,out] The life-out value to use for the reduction object. This will hold the reduced value after the algorithm is finished executing.
- `identity`: [in] The identity value to use for the reduction operation.
- `combiner`: [in] The binary function (object) used to perform a pairwise reduction on the elements.

Note In order to produce useful results, modifications to the view should be limited to commutative operations closely related to the combiner operation. For example if the combiner is `plus<T>`, incrementing the view would be consistent with the combiner but doubling it or assigning to it would not.

Return This returns a reduction object of unspecified type having a value type of `T`. When the return value is used by an algorithm, the reference to `var` is used as the live-out object, new views are initialized to a copy of `identity`, and views are combined by invoking the copy of `combiner`, passing it the two views to be combined.

```
#include <hpx/parallel/container_algorithms/remove_copy.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename OutIter, typename T, typename Proj = util::projection_util>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(<typename
OutIter>::type remove_copyExPolicy &&policy, Rng &&rng, OutIter dest, T const &val,
Proj &&proj = Proj())Copies the elements in the range, defined by [first, last), to another range
beginning at dest. Copies only the elements for which the comparison operator returns false when
compare to val. The order of the elements that are not removed is preserved.
```

Effects: Copies all the elements referred to by the iterator it in the range [first, last) for which the following corresponding conditions do not hold: INVOKE(proj, *it) == value

The assignments in the parallel *remove_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate *f*.

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng:** The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **OutIter:** The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- **T:** The type that the result of dereferencing InIter is compared to.
- **Proj:** The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **rng:** Refers to the sequence of elements the algorithm will be applied to.
- **dest:** Refers to the beginning of the destination range.
- **val:** Value to be removed.
- **proj:** Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove_copy* algorithm returns a `hpx::future<tagged_pair<tag::in(InIter), tag::out(OutIter)> >` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `tagged_pair<tag::in(InIter), tag::out(OutIter)>` otherwise. The *copy* algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename Rng, typename OutIter, typename F, typename Proj = util::projection_util>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(<typename
OutIter>::type remove_copy_ifExPolicy &&policy, Rng &&rng, OutIter dest, F &&f, Proj
&&proj = Proj())Copies the elements in the range, defined by [first, last), to another range
beginning at dest. Copies only the elements for which the predicate f returns false. The order of the
elements that are not removed is preserved.
```

Effects: Copies all the elements referred to by the iterator it in the range [first, last) for which the

following corresponding conditions do not hold: `INVOKE(pred, INVOKE(proj, *it)) != false`.

The assignments in the parallel `remove_copy_if` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate *f*.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- `OutIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `copy_if` requires *F* to meet the requirements of `CopyConstructible`.
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns *true* for the elements to be removed. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `InIter` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `remove_copy_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `remove_copy_if` algorithm returns a `hpx::future<tagged_pair<tag::in(InIter), tag::out(OutIter)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(InIter), tag::out(OutIter)>` otherwise. The `copy` algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/container_algorithms/is_heap.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Comp = detail::less, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type is_heap(ExPolicy &&policy, Rng &&rng, Comp &&comp =
    Comp(), Proj &&proj = Proj())
```

Returns whether the range is max heap. That is, true if the range is max heap, false otherwise. The function uses the given comparison function object *comp* (defaults to using operator<()).

comp has to induce a strict weak ordering on the values.

Note Complexity: Performs at most N applications of the comparison *comp*, at most 2 * N applications of the projection *proj*, where N = last - first.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an random access iterator.
- **Comp**: The type of the function/function object to use (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the INVOKE operation applied to an object of type *Comp*, when contextually converted to *bool*, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_heap* algorithm returns a `hpx::future<bool>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *is_heap* algorithm returns whether the range is max heap. That is, true if the range is max heap, false otherwise.

```
template<typename ExPolicy, typename Rng, typename Comp = detail::less, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type is_heap_until(
```

Returns the upper bound of the largest range beginning at *first* which is a max heap. That is, the

last iterator *it* for which range [first, *it*) is a max heap. The function uses the given comparison function object *comp* (defaults to using operator<()).

comp has to induce a strict weak ordering on the values.

Note Complexity: Performs at most N applications of the comparison *comp*, at most $2 * N$ applications of the projection *proj*, where $N = \text{last} - \text{first}$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an random access iterator.
- *Comp*: The type of the function/function object to use (deduced).
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the INVOKE operation applied to an object of type *Comp*, when contextually converted to bool, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *is_heap_until* algorithm returns a `hpx::future<RandIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *RandIter* otherwise. The *is_heap_until* algorithm returns the upper bound of the largest range beginning at *first* which is a max heap. That is, the last iterator *it* for which range [first, *it*) is a max heap.

```
#include <hpx/parallel/container_algorithms/find.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Rng2, typename Pred = detail::equal_to, typename Proj = util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type find_end(ExPolicy &&policy, Rng &&rng, Rng2 &&rng2, Pred &&op = Pred(), Proj &&proj = Proj())
```

Returns the last subsequence of elements *rng2* found in the range *rng* using the given predicate *f* to compare elements.

The comparison operations in the parallel *find_end* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most $S^*(N-S+1)$ comparisons where $S = \text{distance}(\text{begin}(\text{rng2}), \text{end}(\text{rng2}))$ and $N = \text{distance}(\text{begin}(\text{rng}), \text{end}(\text{rng}))$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the first source range (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- *Rng2*: The type of the second source range (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *replace* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the first sequence of elements the algorithm will be applied to.
- *rng2*: Refers to the second sequence of elements the algorithm will be applied to.
- *op*: The binary predicate which returns *true* if the elements should be treated as equal. The signature should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *iterator_t<Rng>* and *iterator_t<Rng2>* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced *iterator_t<Rng>* and dereferenced *iterator_t<Rng2>* as a projection operation before the function *op* is invoked.

The comparison operations in the parallel *find_end* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered

fashion in unspecified threads, and indeterminately sequenced within each thread.

This overload of *find_end* is available if the user decides to provide the algorithm their own predicate *op*.

Return The *find_end* algorithm returns a *hpx::future<iterator_t<Rng>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *iterator_t<Rng>* otherwise. The *find_end* algorithm returns an iterator to the beginning of the last subsequence *rng2* in range *rng*. If the length of the subsequence *rng2* is greater than the length of the range *rng*, *end(rng)* is returned. Additionally if the size of the subsequence is empty or no subsequence is found, *end(rng)* is also returned.

```
template<typename ExPolicy, typename Rng1, typename Rng2, typename Pred = detail::equal_to, typename Proj1 = util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng1>::type>::type find_first_of
```

Searches the range *rng1* for any elements in the range *rng2*. Uses binary predicate *p* to compare elements

The comparison operations in the parallel *find_first_of* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most ($S \cdot N$) comparisons where $S = \text{distance}(\text{begin}(\text{rng2}), \text{end}(\text{rng2}))$ and $N = \text{distance}(\text{begin}(\text{rng1}), \text{end}(\text{rng1}))$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng1*: The type of the first source range (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- *Rng2*: The type of the second source range (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *replace* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to *std::equal_to<>*
- *Proj1*: The type of an optional projection function. This defaults to *util::projection_identity* and is applied to the elements in *rng1*.
- *Proj2*: The type of an optional projection function. This defaults to *util::projection_identity* and is applied to the elements in *rng2*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.

- `rng1`: Refers to the first sequence of elements the algorithm will be applied to.
- `rng2`: Refers to the second sequence of elements the algorithm will be applied to.
- `op`: The binary predicate which returns `true` if the elements should be treated as equal. The signature should be equivalent to the following:

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types `Type1` and `Type2` must be such that objects of types `iterator_t<Rng1>` and `iterator_t<Rng2>` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively.

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced `iterator_t<Rng1>` before the function `op` is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of type dereferenced `iterator_t<Rng2>` before the function `op` is invoked.

The comparison operations in the parallel `find_first_of` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

This overload of `find_first_of` is available if the user decides to provide the algorithm their own predicate `op`.

Return The `find_end` algorithm returns a `hpx::future<iterator_t<Rng1>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `iterator_t<Rng1>` otherwise. The `find_first_of` algorithm returns an iterator to the first element in the range `rng1` that is equal to an element from the range `rng2`. If the length of the subsequence `rng2` is greater than the length of the range `rng1`, `end(rng1)` is returned. Additionally if the size of the subsequence is empty or no subsequence is found, `end(rng1)` is also returned.

```
#include <hpx/parallel/container_algorithms/move.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename OutIter>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in (typename
                                                                  hpx::traits::range_traits<Rng>::iterator_type),
                                                                  tag::out
                                                                OutIter>>::type move(ExPolicy &&policy, Rng &&rng, OutIter dest) Moves the elements in the
range rng to another range beginning at dest. After this operation the elements in the moved-from
range will still contain valid values of the appropriate type, but not necessarily the same values as
before the move.
```

The assignments in the parallel `copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `std::distance(begin(rng), end(rng))` assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **OutIter**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **dest**: Refers to the beginning of the destination range.

The assignments in the parallel *copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *move* algorithm returns a *hpx::future<tagged_pair<tag::in(iterator_t<Rng>), tag::out(FwdIter2)>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::in(iterator_t<Rng>), tag::out(FwdIter2)>* otherwise. The *move* algorithm returns the pair of the input iterator *last* and the output iterator to the element in the destination range, one past the last element moved.

```
#include <hpx/parallel/container_algorithms/replace.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename T1, typename T2, typename Proj = util::projection_identity,
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_traits<Rng>::iterator_type>::type replace(ExPolicy, Rng, T1, T2, Proj)
```

Replaces all elements satisfying specific criteria with *new_value* in the range [first, last).

Effects: Substitutes elements referred by the iterator *it* in the range [first, last) with *new_value*, when the following corresponding conditions hold: *INVOKE(proj, *i) == old_value*

Note Complexity: Performs exactly *last - first* assignments.

The assignments in the parallel *replace* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- **T1**: The type of the old value to replace (deduced).
- **T2**: The type of the new values to replace (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **old_value**: Refers to the old value of the elements to replace.
- **new_value**: Refers to the new value to use as the replacement.
- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *f* is invoked.

The assignments in the parallel *replace* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace* algorithm returns a `hpx::future<void>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *void* otherwise.

```
template<typename ExPolicy, typename Rng, typename F, typename T, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_traits<Rng>::iterator_type>::type replace_if
```

Replaces all elements satisfying specific criteria (for which predicate *f* returns true) with *new_value* in the range [first, last].

Effects: Substitutes elements referred by the iterator *it* in the range [first, last] with *new_value*, when the following corresponding conditions hold: `INVOKE(f, INVOKE(proj, *it)) != false`

Note Complexity: Performs exactly *last - first* applications of the predicate.

The assignments in the parallel *replace_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form,

the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*. (deduced).

- *T*: The type of the new values to replace (deduced).
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an unary predicate which returns *true* for the elements which need to be replaced. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *new_value*: Refers to the new value to use as the replacement.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *replace_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace_if* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *void* otherwise. It returns *last*.

```
template<typename ExPolicy, typename Rng, typename OutIter, typename T1, typename T2, typename Proj = Proj()
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(<typename
                                         hpx::traits::range_traits<Rng>::iterator_type),
                                         tag::out>>::type replace_copyExPolicy &&policy, Rng &&rng, OutIter dest, T1 const
&old_value, T2 const &new_value, Proj &&proj = Proj()Copies the all elements from the
range [first, last) to another range beginning at dest replacing all elements satisfying a specific
criteria with new_value.
```

Effects: Assigns to every iterator *it* in the range [result, result + (last - first)) either *new_value* or *(first + (it - result)) depending on whether the following corresponding condition holds: `INVOKE(proj, *(first + (i - result))) == old_value`

The assignments in the parallel *replace_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* applications of the predicate.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *OutIter*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- *T1*: The type of the old value to replace (deduced).
- *T2*: The type of the new values to replace (deduced).
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.

- `dest`: Refers to the beginning of the destination range.
- `old_value`: Refers to the old value of the elements to replace.
- `new_value`: Refers to the new value to use as the replacement.
- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate is invoked.

The assignments in the parallel `replace_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `replace_copy` algorithm returns a `hpx::future<tagged_pair<tag::in(InIter), tag::out(OutIter)> >` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(InIter), tag::out(OutIter)>` otherwise. The `copy` algorithm returns the pair of the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename Rng, typename OutIter, typename F, typename T, typename Proj = util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(typename hpx::traits::range_traits<Rng>::iterator_type), tag::out(OutIter)>>::type replace_copy_if(ExPolicy &&policy, Rng &&rng, OutIter dest, F &&f, T const &new_value, Proj &&proj = Proj())Copies the all elements from the range [first, last) to another range beginning at dest replacing all elements satisfying a specific criteria with new_value.
```

Effects: Assigns to every iterator it in the range `[result, result + (last - first))` either `new_value` or `*(first + (it - result))` depending on whether the following corresponding condition holds: `INVOKE(f, INVOKE(proj, *(first + (i - result)))) != false`

The assignments in the parallel `replace_copy_if` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` applications of the predicate.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- `OutIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `equal` requires `F` to meet the requirements of `CopyConstructible`. (deduced).
- `T`: The type of the new values to replace (deduced).
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by `[first, last)`. This is an unary predicate which returns `true` for the elements which need to replaced. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter` can be dereferenced

and then implicitly converted to *Type*.

- `new_value`: Refers to the new value to use as the replacement.
- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `replace_copy_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `replace_copy_if` algorithm returns a `hpx::future<tagged_pair<tag::in(Iter), tag::out(OutIter)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(Iter), tag::out(OutIter)>` otherwise. The `replace_copy_if` algorithm returns the input iterator *last* and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/container_algorithms/partition.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Pred, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type partition(ExPol
    &&po
    icy,
    Rng
    &&rn
    Pred
    &&pr
    Proj
    &&pr
    =
    Proj()
```

Reorders the elements in the range *rng* in such a way that all elements for which the predicate *pred* returns true precede the elements for which the predicate *pred* returns false. Relative order of the elements is not preserved.

The assignments in the parallel `partition` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs at most $2 * N$ swaps, exactly N applications of the predicate and projection, where $N = \text{std}::\text{distance}(\text{begin}(\text{rng}), \text{end}(\text{rng}))$.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.

- *Pred*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *partition* requires *Pred* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *pred*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by the range *rng*. This is an unary predicate for partitioning the source iterators. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *partition* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *partition* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *parallel_task_policy* and returns *FwdIter* otherwise. The *partition* algorithm returns the iterator to the first element of the second group.

```
template<typename ExPolicy, typename Rng, typename FwdIter2, typename FwdIter3, typename Pred, typename util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in (typename hpx::traits::range_iterator<Rng>::type), tag::out1 FwdIter2, tag::out2 FwdIter3>>::type partition_copyExPolicy &&policy, Rng &&rng, FwdIter2 dest_true, FwdIter3 dest_false, Pred &&pred, Proj &&proj = Proj()Copies the elements in the range rng, to two different ranges depending on the value returned by the predicate pred. The elements, that satisfy the predicate pred, are copied to the range beginning at dest_true. The rest of the elements are copied to the range beginning at dest_false. The order of the elements is preserved.
```

The assignments in the parallel *partition_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than N assignments, exactly N applications of the predicate *pred*, where N = `std::distance(begin(rng), end(rng))`.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- *FwdIter2*: The type of the iterator representing the destination range for the elements that satisfy the predicate *pred* (deduced). This iterator type must meet the requirements of an forward iterator.
- *FwdIter3*: The type of the iterator representing the destination range for the elements that don't satisfy the predicate *pred* (deduced). This iterator type must meet the requirements of an forward iterator.
- *Pred*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *partition_copy* requires *Pred* to meet the requirements of *CopyConstructible*.

- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `dest_true`: Refers to the beginning of the destination range for the elements that satisfy the predicate `pred`.
- `dest_false`: Refers to the beginning of the destination range for the elements that don't satisfy the predicate `pred`.
- `pred`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by the range `rng`. This is an unary predicate for partitioning the source iterators. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate `is` invoked.

The assignments in the parallel `partition_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `partition_copy` algorithm returns a `hpx::future<tagged_tuple<tag::in(InIter), tag::out1(OutIter1), tag::out2(OutIter2)> >` if the execution policy is of type `parallel_task_policy` and returns `tagged_tuple<tag::in(InIter), tag::out1(OutIter1), tag::out2(OutIter2)>` otherwise. The `partition_copy` algorithm returns the tuple of the source iterator `last`, the destination iterator to the end of the `dest_true` range, and the destination iterator to the end of the `dest_false` range.

```
#include <hpx/parallel/container_algorithms/rotate.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::begin (typename
hpx::traits::range_iterator<Rng>::type),
tag::end
typename hpx::traits::range_iterator<Rng>::type>>::type rotate(ExPolicy &&policy, Rng
&&rng, typename hpx::traits::range_iterator<Rng>::type middle)
```

Performs a left rotation on a range of elements. Specifically, `rotate` swaps the elements in the range `[first, last)` in such a way that the element `new_first` becomes the first element of the new range and `new_first - 1` becomes the last element.

The assignments in the parallel `rotate` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between *first* and *last*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **middle**: Refers to the element that should appear at the beginning of the rotated range.

The assignments in the parallel *rotate* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The type of dereferenced *FwdIter* must meet the requirements of *MoveAssignable* and *MoveConstructible*.

Return The *rotate* algorithm returns a *hpx::future<tagged_pair<tag::begin(FwdIter), tag::end(FwdIter)>>* if the execution policy is of type *parallel_task_policy* and returns *tagged_pair<tag::begin(FwdIter), tag::end(FwdIter)>* otherwise. The *rotate* algorithm returns the iterator equal to *pair(first + (last - new_first), last)*.

```
template<typename ExPolicy, typename Rng, typename OutIter>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(typename
                                         hpx::traits::range_iterator<Rng>::type),
                                         tag::out
                                         OutIter>>::type rotate_copy(ExPolicy &&policy, Rng &&rng, typename
                                         hpx::traits::range_iterator<Rng>::type middle, OutIter dest_first)Copies the elements from
                                         the range [first, last), to another range beginning at dest_first in such a way, that the element
                                         new_first becomes the first element of the new range and new_first - 1 becomes the last element.
```

The assignments in the parallel *rotate_copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a forward iterator.
- **OutIter**: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **middle**: Refers to the element that should appear at the beginning of the rotated range.
- **dest_first**: Refers to the begin of the destination range.

The assignments in the parallel *rotate_copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *rotate_copy* algorithm returns a *hpx::future<tagged_pair<tag::in(FwdIter), tag::out(OutIter)>>* if the execution policy is of type *parallel_task_policy* and returns *tagged_pair<tag::in(FwdIter), tag::out(OutIter)>* otherwise. The *rotate_copy* algorithm returns the output iterator to the element past the last element copied.

```
#include <hpx/parallel/container_algorithms/fill.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

`template<typename ExPolicy, typename Rng, typename T>
util::detail::algorithm_result<ExPolicy>::type fill (ExPolicy &&policy, Rng &&rng, T value)`
Assigns the given value to the elements in the range [first, last).

The comparisons in the parallel *fill* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **value**: The value to be assigned.

The comparisons in the parallel *fill* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *fill* algorithm returns a `hpx::future<void>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by *void*).

`template<typename ExPolicy, typename Rng, typename Size, typename T>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_traits<Rng>::iterator_type>::type fill_n (ExPolicy &&policy, Rng &&rng, Size count, T value)`

Assigns the given value *value* to the first *count* elements in the range beginning at *first* if *count* > 0. Does nothing otherwise.

The comparisons in the parallel *fill_n* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *count* assignments, for *count* > 0.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **Size**: The type of the argument specifying the number of elements to apply *f* to.
- **T**: The type of the value to be assigned (deduced).

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **count**: Refers to the number of elements starting at *first* the algorithm will be applied to.
- **value**: The value to be assigned.

The comparisons in the parallel *fill_n* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *fill_n* algorithm returns a *hpx::future<void>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by *void*).

```
#include <hpx/parallel/container_algorithms/generate.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename F>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type generate(ExPolicy
&&policy,
Rng
&&rng
F
&&f)
```

Assign each element in range [first, last) a value generated by the given function object *f*

The assignments in the parallel *generate* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly *distance(first, last)* invocations of *f* and assignments.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *equal* requires *F* to meet the requirements of *CopyConstructible*.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **f**: generator function that will be called. signature of function should be equivalent to the following:

```
Ret fun();
```

The type *Ret* must be such that an object of type *FwdIter* can be dereferenced and assigned a value of type *Ret*.

The assignments in the parallel *generate* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *replace_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. It returns *last*.

```
#include <hpx/parallel/container_algorithms/search.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng1, typename Rng2, typename Pred = detail::equal_to, typename Proj1 = util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng1>::type>::type search(ExPolicy &&policy, Rng1 &&rng1, Rng2 &&rng2, Pred &&op = Pred(), Proj1 &&proj1 = Proj1(), Proj2 &&proj2 = Proj2())
```

Searches the range [first, last) for any elements in the range [s_first, s_last). Uses a provided predicate to compare elements.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most (S^*N) comparisons where $S = \text{distance}(\text{s_first}, \text{s_last})$ and $N = \text{distance}(\text{first}, \text{last})$.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng1**: The type of the examine range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **Rng2**: The type of the search range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **Pred**: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *adjacent_find* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- **Proj1**: The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of *Rng1*.
- **Proj2**: The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of *Rng2*.

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng1**: Refers to the sequence of elements the algorithm will be examining.
- **rng2**: Refers to the sequence of elements the algorithm will be searching for.
- **op**: Refers to the binary predicate which returns true if the elements should be treated as equal. the signature of the function should be equivalent to

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and *FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively

- **proj1**: Specifies the function (or function object) which will be invoked for each of the elements of *rng1* as a projection operation before the actual predicate *is* invoked.
- **proj2**: Specifies the function (or function object) which will be invoked for each of the elements of *rng2* as a projection operation before the actual predicate *is* invoked.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *search* algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type *task_execution_policy* and returns *FwdIter* otherwise. The *search* algorithm returns an iterator to the beginning of the first subsequence $[\text{s_first}, \text{s_last}]$ in range $[\text{first}, \text{last}]$. If the length of the subsequence $[\text{s_first}, \text{s_last}]$ is greater than the length of the range $[\text{first}, \text{last}]$, *last* is returned. Additionally if the size of the subsequence is empty *first* is returned. If no subsequence is found, *last* is returned.

```
template<typename ExPolicy, typename Rng1, typename Rng2, typename Pred = detail::equal_to, typename Proj1 = util::projection_identity, typename Proj2 = util::projection_identity> FwdIter adjacent_find(ExPolicy policy, Rng1 rng1, Rng2 rng2, Pred op, Proj1 proj1 = {}, Proj2 proj2 = {});
```

```
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng1>::type>::type search_n(ExPolicy
&&policy,
Rng1
&&rng1,
std::size_t count,
Rng2
&&rng2,
Pred
&&op
=
Pred()
Proj1
&&proj1
=
Proj1()
Proj2
&&proj2
=
Proj2()
```

Searches the range [first, last) for any elements in the range [s_first, s_last). Uses a provided predicate to compare elements.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: at most (S^*N) comparisons where $S = \text{distance}(\text{s_first}, \text{s_last})$ and $N = \text{distance}(\text{first}, \text{last})$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng1*: The type of the examine range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *Rng2*: The type of the search range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *Pred*: The type of an optional function/function object to use. Unlike its sequential form, the parallel overload of *adjacent_find* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to `std::equal_to<>`
- *Proj1*: The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of *Rng1*.
- *Proj2*: The type of an optional projection function. This defaults to `util::projection_identity` and is applied to the elements of *Rng2*.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng1*: Refers to the sequence of elements the algorithm will be examining.
- *count*: The number of elements to apply the algorithm on.
- *rng2*: Refers to the sequence of elements the algorithm will be searching for.
- *op*: Refers to the binary predicate which returns true if the elements should be treated as equal. the signature of the function should be equivalent to

```
bool pred(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const &`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *FwdIter1* and

- FwdIter2* can be dereferenced and then implicitly converted to *Type1* and *Type2* respectively
- *proj1*: Specifies the function (or function object) which will be invoked for each of the elements of *rng1* as a projection operation before the actual predicate *is* invoked.
 - *proj2*: Specifies the function (or function object) which will be invoked for each of the elements of *rng2* as a projection operation before the actual predicate *is* invoked.

The comparison operations in the parallel *search* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *search* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *task_execution_policy* and returns *FwdIter* otherwise. The *search* algorithm returns an iterator to the beginning of the first subsequence [*s_first*, *s_last*] in range [*first*, *last*). If the length of the subsequence [*s_first*, *s_last*] is greater than the length of the range [*first*, *last*), *last* is returned. Additionally if the size of the subsequence is empty *first* is returned. If no subsequence is found, *last* is returned.

```
#include <hpx/parallel/container_algorithms/copy.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename OutIter>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in (typename
                                                               hpx::traits::range_traits<Rng>::iterator_type),
                                                               tag::out
                                                               OutIter>>::type copy<ExPolicy &&policy, Rng &&rng, OutIter destCopies the elements in the
                                                               range rng to another range beginning at dest.
```

The assignments in the parallel *copy* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly `std::distance(begin(rng), end(rng))` assignments.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *OutIter*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.

The assignments in the parallel *copy* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `copy` algorithm returns a `hpx::future<tagged_pair<tag::in(iterator_t<Rng>, tag::out(FwdIter2)> >` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(iterator_t<Rng>), tag::out(FwdIter2)>` otherwise. The `copy` algorithm returns the pair of the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename Rng, typename OutIter, typename F, typename Proj = util::projection  
        util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in (typename  
                                hpx::traits::range_traits<Rng>::iterator_type ),  
                                tag::out  
                                )>>
```

`OutIter>>::type copy_if(ExPolicy &&policy, Rng &&rng, OutIter dest, F &&f, Proj &&proj = Proj())`Copies the elements in the range *rng* to another range beginning at *dest*. Copies only the elements for which the predicate *f* returns true. The order of the elements that are not removed is preserved.

The assignments in the parallel `copy_if` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than `std::distance(begin(rng), end(rng))` assignments, exactly `std::distance(begin(rng), end(rng))` applications of the predicate `f`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
 - `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
 - `OutIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
 - `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `copy_if` requires `F` to meet the requirements of `CopyConstructible`.
 - `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
 - **rng**: Refers to the sequence of elements the algorithm will be applied to.
 - **dest**: Refers to the beginning of the destination range.
 - **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an unary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `Iter` can be dereferenced and then implicitly converted to `Type`.

- proj: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `copy_if` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `copy_if` algorithm returns a `hpx::future<tagged_pair<tag::in(iterator_t<Rng>), tag::out(FwdIter2)>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(iterator_t<Rng>), tag::out(FwdIter2)>` otherwise. The `copy_if` algorithm returns the pair of the input iterator `last` and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/container_algorithms/all_any_none.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, bool>::type none_of(ExPolicy &&policy, Rng
&&rng, F &&f, Proj &&proj =
Proj())
```

Checks if unary predicate *f* returns true for no elements in the range *rng*.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most std::distance(begin(rng), end(rng)) applications of the predicate *f*

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *none_of* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have *const&*, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *none_of* algorithm returns a *hpx::future<bool>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *bool* otherwise. The *none_of* algorithm returns true if the unary predicate *f* returns true for no elements in the range, false otherwise. It returns true if the range is empty.

```
template<typename ExPolicy, typename Rng, typename F, typename Proj = util::projection_identity>
```

`util::detail::algorithm_result<ExPolicy, bool>::type any_of (ExPolicy &&policy, Rng &&rng, F &&f, Proj &&proj = Proj())`

Checks if unary predicate *f* returns true for at least one element in the range *rng*.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most std::distance(begin(rng), end(rng)) applications of the predicate *f*

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *none_of* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *f*: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *any_of* algorithm returns a `hpx::future<bool>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns `bool` otherwise. The *any_of* algorithm returns true if the unary predicate *f* returns true for at least one element in the range, false otherwise. It returns false if the range is empty.

`template<typename ExPolicy, typename Rng, typename F, typename Proj = util::projection_identity> util::detail::algorithm_result<ExPolicy, bool>::type all_of (ExPolicy &&policy, Rng &&rng, F &&f, Proj &&proj = Proj())`

Checks if unary predicate *f* returns true for all elements in the range *rng*.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most std::distance(begin(rng), end(rng)) applications of the predicate *f*

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *none_of* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `all_of` algorithm returns a `hpx::future<bool>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `bool` otherwise. The `all_of` algorithm returns true if the unary predicate *f* returns true for all elements in the range, false otherwise. It returns true if the range is empty.

```
#include <hpx/parallel/container_algorithms/reverse.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type reverse(ExPolicy &&policy, Rng &&rng)
```

Reverses the order of the elements in the range [first, last). Behaves as if applying `std::iter_swap` to every pair of iterators `first+i`, `(last-i) - 1` for each non-negative $i < (last-first)/2$.

The assignments in the parallel `reverse` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Linear in the distance between `first` and `last`.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a bidirectional iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.

The assignments in the parallel `reverse` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `reverse` algorithm returns a `hpx::future<BidirIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `BidirIter` otherwise. It returns `last`.

```
template<typename ExPolicy, typename Rng, typename OutIter>
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(<typename
    hpx::traits::range_iterator<Rng>::type),
    tag::out
OutIter>>::type reverse_copy(ExPolicy &&policy, Rng &&rng, OutIter dest_first)Copies the
elements from the range [first, last) to another range beginning at dest_first in such a way that
the elements in the new range are in reverse order. Behaves as if by executing the assignment
*(dest_first + (last - first) - 1 - i) = *(first + i) once for each non-negative i < (last - first). If the
source and destination ranges (that is, [first, last) and [dest_first, dest_first+(last-first)) respec-
tively) overlap, the behavior is undefined.
```

The assignments in the parallel `reverse_copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs exactly `last - first` assignments.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of a bidirectional iterator.
- `OutputIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `dest_first`: Refers to the begin of the destination range.

The assignments in the parallel `reverse_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `reverse_copy` algorithm returns a `hpx::future<tagged_pair<tag::in(BidirIter),
 tag::out(OutIter)> >` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_pair<tag::in(BidirIter), tag::out(OutIter)>` otherwise. The `copy` algorithm returns the pair of the input iterator forwarded to the first element after the last in the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/container_algorithms/remove.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename T, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type remove(ExPolicy
    &&policy,
    Rng
    &&rng,
    T
    const
    &&value,
    Proj
    &&proj
    =
    Proj())
```

Removes all elements satisfying specific criteria from the range [first, last) and returns a past-the-end iterator for the new end of the range. This version removes all elements that are equal to *value*.

The assignments in the parallel *remove* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the operator==() and the projection *proj*.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- *T*: The type of the value to remove (deduced). This value type must meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *value*: Specifies the value of elements to remove.
- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *remove* algorithm returns the iterator to the new end of the range.

```
template<typename ExPolicy, typename Rng, typename Pred, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type remove_if(ExPolicy&&  

ExPolicy,  

Rng&&  

Pred&&  

Proj&&  

= Proj()
```

Removes all elements satisfying specific criteria from the range [first, last) and returns a past-the-end iterator for the new end of the range. This version removes all elements for which predicate *pred* returns true.

The assignments in the parallel *remove_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than *last - first* assignments, exactly *last - first* applications of the predicate *pred* and the projection *proj*.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- **Pred**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *remove_if* requires *Pred* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **pred**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). This is an unary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have *const&*, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *remove_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *remove_if* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *remove_if* algorithm returns the iterator to the new end of the range.

```
#include <hpx/parallel/container_algorithms/minmax.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Proj = util::projection_identity, typename F = detail::less<  
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_traits<Rng>::iterator_type>>::type min_element
```

Finds the smallest element in the range [first, last) using the given comparison function *f*.

The comparisons in the parallel *min_element* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly $\max(N-1, 0)$ comparisons, where $N = \text{std}::\text{distance}(\text{first}, \text{last})$.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *min_element* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **f**: The binary predicate which returns true if the the left argument is less than the right element. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have **const &**, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type1*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The comparisons in the parallel *min_element* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *min_element* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *min_element* algorithm returns the iterator to the smallest element in the range [first, last). If several elements in the range are equivalent to the smallest element, returns the iterator to the first such element. Returns last if the range is empty.

```
template<typename ExPolicy, typename Rng, typename Proj = util::projection_identity, typename F = detail::less<util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_traits<Rng>::iterator_type>::type max_element>
```

Finds the greatest element in the range [first, last) using the given comparison function *f*.

The comparisons in the parallel *max_element* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly $\max(N-1, 0)$ comparisons, where $N = \text{std}::\text{distance}(\text{first}, \text{last})$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *max_element* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *f*: The binary predicate which returns true if the This argument is optional and defaults to *std::less*. the left argument is less than the right element. The signature of the predicate function should be equivalent to the following:

<code>bool pred(const Type1 &a, const Type1 &b);</code>

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type1*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The comparisons in the parallel *max_element* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *max_element* algorithm returns a *hpx::future<FwdIter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *FwdIter* otherwise. The *max_element* algorithm returns the iterator to the smallest element in the range [first, last). If several elements in the range are equivalent to the smallest element, returns the iterator to the first such element. Returns last if the range is empty.

```
template<typename ExPolicy, typename Rng, typename Proj = util::projection_identity, typename F = detail::less<util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::min (typename hpx::traits::range_traits<Rng>::iterator_type), tag::max (typename hpx::traits::range_traits<Rng>::iterator_type>>::type minmax_element ExPolicy &&policy, Rng &&rng, F &&f = F(), Proj &&proj = Proj()> Finds the greatest element in the range [first, last) using the given comparison function f.
```

The comparisons in the parallel *minmax_element* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: At most $\max(\text{floor}(3/2*(N-1)), 0)$ applications of the predicate, where $N = \text{std::distance(first, last)}$.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- *F*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *minmax_element* requires *F* to meet the requirements of *CopyConstructible*.
- *Proj*: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *f*: The binary predicate which returns true if the the left argument is less than the right element. This argument is optional and defaults to *std::less*. The signature of the predicate function should be equivalent to the following:

```
bool pred(const Type1 &a, const Type1 &b);
```

The signature does not need to have *const &*, but the function must not modify the objects passed to it. The type *Type1* must be such that objects of type *FwdIter* can be dereferenced and then implicitly converted to *Type1*.

- *proj*: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The comparisons in the parallel *minmax_element* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *minmax_element* algorithm returns a *hpx::future<tagged_pair<tag::min(FwdIter), tag::max(FwdIter)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::min(FwdIter), tag::max(FwdIter)>* otherwise. The *minmax_element* algorithm returns a pair consisting of an iterator to the smallest element as the first element and an iterator to the greatest element as the second. Returns *std::make_pair(first, first)* if the range is empty. If several elements are equivalent to the smallest element, the iterator to the first such element is returned. If several elements are equivalent to the largest element, the iterator to the last such element is returned.

```
#include <hpx/parallel/container_algorithms/unique.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Pred = detail::equal_to, typename Proj = util::projection_identities::identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type unique(ExPolicy
    &&policy,
    Rng
    &&rng,
    Pred
    &&pred
    =
    Pred(),
    Proj
    &&proj
    =
    Proj())
```

Eliminates all but the first element from every consecutive group of equivalent elements from the range *rng* and returns a past-the-end iterator for the new logical end of the range.

The assignments in the parallel *unique* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs not more than N assignments, exactly N - 1 applications of the predicate *pred* and no more than twice as many applications of the projection *proj*, where N = std::distance(begin(rng), end(rng)).

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- **Pred**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *unique* requires *Pred* to meet the requirements of *CopyConstructible*. This defaults to std::equal_to<>
- **Proj**: The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **pred**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an binary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

<code>bool pred(const Type &a, const Type &b);</code>

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `unique` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `unique` algorithm returns a `hpx::future<FwdIter>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `FwdIter` otherwise. The `unique` algorithm returns the iterator to the new end of the range.

```
template<typename ExPolicy, typename Rng, typename FwdIter2, typename Pred = detail::equal_to, typename util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in (typename hpx::traits::range_iterator<Rng>::type), tag::out FwdIter2>>::type unique_copy ExPolicy &&policy, Rng &&rng, FwdIter2 dest, Pred &&pred = Pred(), Proj &&proj = Proj()Copies the elements from the range rng, to another range beginning at dest in such a way that there are no consecutive equal elements. Only the first element of each group of equal elements is copied.
```

The assignments in the parallel `unique_copy` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs not more than N assignments, exactly $N - 1$ applications of the predicate `pred`, where $N = \text{std}:\text{distance}(\text{begin}(\text{rng}), \text{end}(\text{rng}))$.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an forward iterator.
- `FwdIter2`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an forward iterator.
- `Pred`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `unique_copy` requires `Pred` to meet the requirements of `CopyConstructible`. This defaults to `std::equal_to<>`
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `pred`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by the range `rng`. This is an binary predicate which returns `true` for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a, const Type &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type `Type` must be such that an object of type `FwdIter1` can be dereferenced and then implicitly converted to `Type`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel `unique_copy` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion

in unspecified threads, and indeterminately sequenced within each thread.

Return The *unique_copy* algorithm returns a *hpx::future<tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *tagged_pair<tag::in(FwdIter1), tag::out(FwdIter2)>* otherwise. The *unique_copy* algorithm returns the pair of the source iterator to *last*, and the destination iterator to the end of the *dest* range.

```
#include <hpx/parallel/container_algorithms/transform.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename OutIter, typename F, typename Proj = util::projection_util::detail::algorithm_result<ExPolicy, hpx::util::tagged_pair<tag::in(typename hpx::traits::range_iterator<Rng>::type), tag::out OutIter>>::type transformExPolicy &&policy, Rng &&rng, OutIter dest, F &&f, Proj &&proj = Proj()Applies the given function f to the given range rng and stores the result in another range, beginning at dest.
```

The invocations of *f* in the parallel *transform* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Exactly *size(rng)* applications of *f*

Template Parameters

- **ExPolicy:** The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of *f*.
- **Rng:** The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **OutIter:** The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- **F:** The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *transform* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj:** The type of an optional projection function. This defaults to *util::projection_identity*

Parameters

- **policy:** The execution policy to use for the scheduling of the iterations.
- **rng:** Refers to the sequence of elements the algorithm will be applied to.
- **dest:** Refers to the beginning of the destination range.
- **f:** Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an unary predicate. The signature of this predicate should be equivalent to:

Ret fun (const Type &a);

The signature does not need to have `const&`. The type `Type` must be such that an object of type `InIter` can be dereferenced and then implicitly converted to `Type`. The type `Ret` must be such that an object of type `OutIter` can be dereferenced and assigned a value of type `Ret`.

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform` algorithm returns a `hpx::future<tagged_pair<tag::in(InIter), tag::out(OutIter)>` if the execution policy is of type `parallel_task_policy` and returns `tagged_pair<tag::in(InIter), tag::out(OutIter)>` otherwise. The `transform` algorithm returns a tuple holding an iterator referring to the first element after the input sequence and the output iterator to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename Rng, typename InIter2, typename OutIter, typename F, typename P  
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1 (typename
```

```
          hpx::traits::range_iterator<Rng>::type),  
          tag::in2
```

```
InIter2, tag::outOutIter>>::type transformExPolicy &&policy, Rng &&rng, InIter2 first2,  
OutIter dest, F &&f, Proj1 &&proj1 = Proj1(), Proj2 &&proj2 = Proj2())Applies the given func-  
tion f to pairs of elements from two ranges: one defined by rng and the other beginning at first2,  
and stores the result in another range, beginning at dest.
```

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly `size(rng)` applications of `f`

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of `f`.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- `InIter2`: The type of the source iterators for the second range used (deduced). This iterator type must meet the requirements of an input iterator.
- `OutIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `transform` requires `F` to meet the requirements of `CopyConstructible`.
- `Proj1`: The type of an optional projection function to be used for elements of the first sequence. This defaults to `util::projection_identity`
- `Proj2`: The type of an optional projection function to be used for elements of the second sequence. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the sequence of elements the algorithm will be applied to.
- `first2`: Refers to the beginning of the second sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun (const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`. The types `Type1` and `Type2` must be such that

objects of types `InIter1` and `InIter2` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively. The type `Ret` must be such that an object of type `OutIter` can be dereferenced and assigned a value of type `Ret`.

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of the first sequence as a projection operation before the actual predicate `f` is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of the second sequence as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `transform` algorithm returns a `hpx::future<tagged_tuple<tag::in1(InIter1), tag::in2(InIter2), tag::out(OutIter)>` if the execution policy is of type `parallel_task_policy` and returns `tagged_tuple<tag::in1(InIter1), tag::in2(InIter2), tag::out(OutIter)>` otherwise. The `transform` algorithm returns a tuple holding an iterator referring to the first element after the first input sequence, an iterator referring to the first element after the second input sequence, and the output iterator referring to the element in the destination range, one past the last element copied.

```
template<typename ExPolicy, typename Rng1, typename Rng2, typename OutIter, typename F, typename Pro
util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1(typename
typename hpx::traits::range_iterator<Rng1>::type),
tag::in2
typename hpx::traits::range_iterator<Rng2>::type, tag::outOutIter>>::type,
transformExPolicy &&policy, Rng1 &&rng1, Rng2 &&rng2, OutIter dest, F &&f,
Proj1 &&proj1 = Proj1(), Proj2 &&proj2 = Proj2())Applies the given function f to pairs of
elements from two ranges: one defined by [first1, last1) and the other beginning at first2, and
stores the result in another range, beginning at dest.
```

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Exactly `min(last2-first2, last1-first1)` applications of `f`

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the invocations of `f`.
- `Rng1`: The type of the first source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- `Rng2`: The type of the second source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- `OutIter`: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an output iterator.
- `F`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `transform` requires `F` to meet the requirements of `CopyConstructible`.
- `Proj1`: The type of an optional projection function to be used for elements of the first sequence. This defaults to `util::projection_identity`
- `Proj2`: The type of an optional projection function to be used for elements of the second sequence. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng1`: Refers to the first sequence of elements the algorithm will be applied to.
- `rng2`: Refers to the second sequence of elements the algorithm will be applied to.
- `dest`: Refers to the beginning of the destination range.
- `f`: Specifies the function (or function object) which will be invoked for each of the elements

in the sequence specified by [first, last). This is a binary predicate. The signature of this predicate should be equivalent to:

```
Ret fun(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`. The types `Type1` and `Type2` must be such that objects of types `InIter1` and `InIter2` can be dereferenced and then implicitly converted to `Type1` and `Type2` respectively. The type `Ret` must be such that an object of type `OutIter` can be dereferenced and assigned a value of type `Ret`.

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of the first sequence as a projection operation before the actual predicate `f` is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of the second sequence as a projection operation before the actual predicate `f` is invoked.

The invocations of `f` in the parallel `transform` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Note The algorithm will invoke the binary predicate until it reaches the end of the shorter of the two given input sequences

Return The `transform` algorithm returns a `hpx::future<tagged_tuple<tag::in1(InIter1), tag::in2(InIter2), tag::out(OutIter)>>` if the execution policy is of type `parallel_task_policy` and returns `tagged_tuple<tag::in1(InIter1), tag::in2(InIter2), tag::out(OutIter)>` otherwise. The `transform` algorithm returns a tuple holding an iterator referring to the first element `r` the first input sequence, an iterator referring to the first element after the second input sequence, and the output iterator referring to the element in the destination range, one past the last element copied.

```
#include <hpx/parallel/container_algorithms/count.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename T, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename std::iterator_traits<typename hpx::traits::range_traits<Rng>::iterator>
```

Returns the number of elements in the range [first, last) satisfying a specific criteria. This version

counts the elements that are equal to the given *value*.

The comparisons in the parallel *count* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs exactly *last - first* comparisons.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the comparisons.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **T**: The type of the value to search for (deduced).
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **value**: The value to search for.
- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

Note The comparisons in the parallel *count* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *count* algorithm returns a `hpx::future<difference_type>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by `std::iterator_traits<FwdIter>::difference_type`). The *count* algorithm returns the number of elements satisfying the given criteria.

```
template<typename ExPolicy, typename Rng, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename std::iterator_traits<typename hpx::traits::range_traits<Rng>::item>::difference_type> count_if(ExPolicy exec, Rng rng, F f, Proj proj = util::projection_identity);
```

Returns the number of elements in the range [first, last) satisfying a specific criteria. This version counts elements for which predicate *f* returns true.

Note Complexity: Performs exactly *last - first* applications of the predicate.

Note The assignments in the parallel *count_if* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note The assignments in the parallel *count_if* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *count_if* algorithm returns `hpx::future<difference_type>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *difference_type* otherwise (where *difference_type* is defined by `std::iterator_traits<FwdIter>::difference_type`). The *count* algorithm returns the number of elements satisfying the given criteria.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes

the comparisons.

- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *count_if* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last]. This is an unary predicate which returns *true* for the required elements. The signature of this predicate should be equivalent to:

```
bool pred(const Type &a);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The type *Type* must be such that an object of type *FwdIter* can be dereferenced and then implicitly converted to *Type*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

```
#include <hpx/parallel/container_algorithms/for_each.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename F, typename Proj = util::projection_identity>
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type for_each(ExPolicy
    &&policy,
    Rng
    &&rng,
    F
    &&f,
    Proj
    &&proj =
    Proj())
```

Applies *f* to the result of dereferencing every iterator in the given range *rng*.

If *f* returns a result, the result is ignored.

Note Complexity: Applies *f* exactly *size(rng)* times.

If the type of *first* satisfies the requirements of a mutable iterator, *f* may apply non-constant functions through the dereferenced iterator.

Unlike its sequential form, the parallel overload of *for_each* does not return a copy of its *Function* parameter, since parallelization may not permit efficient state accumulation.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Template Parameters

- **ExPolicy**: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- **Rng**: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- **F**: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *for_each* requires *F* to meet the requirements of *CopyConstructible*.
- **Proj**: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- **policy**: The execution policy to use for the scheduling of the iterations.
- **rng**: Refers to the sequence of elements the algorithm will be applied to.
- **f**: Specifies the function (or function object) which will be invoked for each of the elements in the sequence specified by [first, last). The signature of this predicate should be equivalent to:

```
<ignored> pred(const Type &a);
```

The signature does not need to have `const&`. The type *Type* must be such that an object of type *InIter* can be dereferenced and then implicitly converted to *Type*.

- **proj**: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *for_each* algorithm returns a `hpx::future<InIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *InIter* otherwise. It returns *last*.

```
#include <hpx/parallel/container_algorithms/sort.hpp>
```

```
namespace hpx
```

```
namespace parallel
```

```
namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng, typename Proj = util::projection_identity, typename Compare = de
```

```
util::detail::algorithm_result<ExPolicy, typename hpx::traits::range_iterator<Rng>::type>::type sort (ExPolicy
&&policy,
Rng
&&rng,
Com-
pare
&&comp
=
Com-
pare(),
Proj
&&proj
=
Proj())
```

Sorts the elements in the range *rng* in ascending order. The order of equal elements is not guaranteed to be preserved. The function uses the given comparison function object *comp* (defaults to using operator<()).

A sequence is sorted with respect to a comparator *comp* and a projection *proj* if for every iterator *i* pointing to the sequence and every non-negative integer *n* such that *i + n* is a valid iterator pointing to an element of the sequence, and $\text{INVOKED}(\text{comp}, \text{INVOKED}(\text{proj}, *(\text{i} + \text{n})), \text{INVOKED}(\text{proj}, *\text{i})) == \text{false}$.

Note Complexity: $O(N\log(N))$, where $N = \text{std}:\text{distance}(\text{begin}(\text{rng}), \text{end}(\text{rng}))$ comparisons. *comp* has to induce a strict weak ordering on the values.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it applies user-provided function objects.
- *Rng*: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an input iterator.
- *Comp*: The type of the function/function object to use (deduced).
- *Proj*: The type of an optional projection function. This defaults to [util::projection_identity](#)

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng*: Refers to the sequence of elements the algorithm will be applied to.
- *comp*: *comp* is a callable object. The return value of the *INVOKED* operation applied to an object of type *Comp*, when contextually converted to *bool*, yields true if the first argument of the call is less than the second, and false otherwise. It is assumed that *comp* will not apply any non-constant function through the dereferenced iterator.
- *proj*: Specifies the function (or function object) which will be invoked for each pair of elements as a projection operation before the actual predicate *comp* is invoked.

The application of function objects in parallel algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

The application of function objects in parallel algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *sort* algorithm returns a *hpx::future<Iter>* if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *Iter* otherwise. It returns *last*.

```
#include <hpx/parallel/container_algorithms/merge.hpp>

namespace hpx

namespace parallel

namespace v1
```

Functions

```
template<typename ExPolicy, typename Rng1, typename Rng2, typename RandIter3, typename Comp = detail::util::detail::algorithm_result<ExPolicy, hpx::util::tagged_tuple<tag::in1 (typename hpx::traits::range_iterator<Rng1>::type), tag::in2 (typename hpx::traits::range_iterator<Rng2>::type, tag::out<RandIter3>>::type mergeExPolicy &&policy, Rng1 &&rng1, Rng2 &&rng2, RandIter3 dest, Comp &&comp = Comp(), Proj1 &&proj1 = Proj1(), Proj2 &&proj2 = Proj2()>> Merges two sorted ranges [first1, last1) and [first2, last2) into one sorted range beginning at dest. The order of equivalent elements in the each of original two ranges is preserved. For equivalent elements in the original two ranges, the elements from the first range precede the elements from the second range. The destination range cannot overlap with either of the input ranges.
```

The assignments in the parallel *merge* algorithm invoked with an execution policy object of type *sequenced_policy* execute in sequential order in the calling thread.

Note Complexity: Performs $O(\text{std}::\text{distance}(\text{first1}, \text{last1}) + \text{std}::\text{distance}(\text{first2}, \text{last2}))$ applications of the comparison *comp* and the each projection.

Template Parameters

- *ExPolicy*: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- *Rng1*: The type of the first source range used (deduced). The iterators extracted from this range type must meet the requirements of an random access iterator.
- *Rng2*: The type of the second source range used (deduced). The iterators extracted from this range type must meet the requirements of an random access iterator.
- *RandIter3*: The type of the iterator representing the destination range (deduced). This iterator type must meet the requirements of an random access iterator.
- *Comp*: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of *merge* requires *Comp* to meet the requirements of *CopyConstructible*. This defaults to *std::less<>*
- *Proj1*: The type of an optional projection function to be used for elements of the first range. This defaults to *util::projection_identity*
- *Proj2*: The type of an optional projection function to be used for elements of the second range. This defaults to *util::projection_identity*

Parameters

- *policy*: The execution policy to use for the scheduling of the iterations.
- *rng1*: Refers to the first range of elements the algorithm will be applied to.
- *rng2*: Refers to the second range of elements the algorithm will be applied to.
- *dest*: Refers to the beginning of the destination range.
- *comp*: *comp* is a callable object which returns true if the first argument is less than the second, and false otherwise. The signature of this comparison should be equivalent to:

```
bool comp(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types `Type1` and `Type2` must be such that objects of types `RandIter1` and `RandIter2` can be dereferenced and then implicitly converted to both `Type1` and `Type2`.

- `proj1`: Specifies the function (or function object) which will be invoked for each of the elements of the first range as a projection operation before the actual comparison `comp` is invoked.
- `proj2`: Specifies the function (or function object) which will be invoked for each of the elements of the second range as a projection operation before the actual comparison `comp` is invoked.

The assignments in the parallel `merge` algorithm invoked with an execution policy object of type `parallel_policy` or `parallel_task_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The `merge` algorithm returns a `hpx::future<tagged_tuple<tag::in1(RandIter1), tag::in2(RandIter2), tag::out(RandIter3)>>` if the execution policy is of type `sequenced_task_policy` or `parallel_task_policy` and returns `tagged_tuple<tag::in1(RandIter1), tag::in2(RandIter2), tag::out(RandIter3)>>` otherwise. The `merge` algorithm returns the tuple of the source iterator `last1`, the source iterator `last2`, the destination iterator to the end of the `dest` range.

```
template<typename ExPolicy, typename Rng, typename RandIter, typename Comp = std::less, typename Proj = std::projection_identity>
util::detail::algorithm_result<ExPolicy, RandIter>::type inplace_merge(ExPolicy &&policy,
                                                                     Rng &&rng,
                                                                     RandIter middle,
                                                                     Comp &&comp =
                           Comp(), Proj &&proj = Proj())
```

Merges two consecutive sorted ranges [`first`, `middle`) and [`middle`, `last`] into one sorted range [`first`, `last`). The order of equivalent elements in the each of original two ranges is preserved. For equivalent elements in the original two ranges, the elements from the first range precede the elements from the second range.

The assignments in the parallel `inplace_merge` algorithm invoked with an execution policy object of type `sequenced_policy` execute in sequential order in the calling thread.

Note Complexity: Performs $O(\text{std}::\text{distance}(\text{first}, \text{last}))$ applications of the comparison `comp` and the each projection.

Template Parameters

- `ExPolicy`: The type of the execution policy to use (deduced). It describes the manner in which the execution of the algorithm may be parallelized and the manner in which it executes the assignments.
- `Rng`: The type of the source range used (deduced). The iterators extracted from this range type must meet the requirements of an random access iterator.
- `RandIter`: The type of the source iterators used (deduced). This iterator type must meet the requirements of an random access iterator.
- `Comp`: The type of the function/function object to use (deduced). Unlike its sequential form, the parallel overload of `inplace_merge` requires `Comp` to meet the requirements of `CopyConstructible`. This defaults to `std::less<>`
- `Proj`: The type of an optional projection function. This defaults to `util::projection_identity`

Parameters

- `policy`: The execution policy to use for the scheduling of the iterations.
- `rng`: Refers to the range of elements the algorithm will be applied to.
- `middle`: Refers to the end of the first sorted range and the beginning of the second sorted range the algorithm will be applied to.

- `comp`: *comp* is a callable object which returns true if the first argument is less than the second, and false otherwise. The signature of this comparison should be equivalent to:

```
bool comp(const Type1 &a, const Type2 &b);
```

The signature does not need to have `const&`, but the function must not modify the objects passed to it. The types *Type1* and *Type2* must be such that objects of types *RandIter* can be dereferenced and then implicitly converted to both *Type1* and *Type2*

- `proj`: Specifies the function (or function object) which will be invoked for each of the elements as a projection operation before the actual predicate *is* invoked.

The assignments in the parallel *inplace_merge* algorithm invoked with an execution policy object of type *parallel_policy* or *parallel_task_policy* are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread.

Return The *inplace_merge* algorithm returns a `hpx::future<RandIter>` if the execution policy is of type *sequenced_task_policy* or *parallel_task_policy* and returns *RandIter* otherwise. The *inplace_merge* algorithm returns the source iterator *last*

assertion

```
#include <compatibility/hpx/util/assert.hpp>

#include <hpx/assertion.hpp>
```

Defines

HPX_ASSERT (expr)

This macro asserts that *expr* evaluates to true.

If *expr* evaluates to false, The source location and *msg* is being printed along with the expression and additional. Afterwards the program is being aborted. The assertion handler can be customized by calling `hpx::assertion::set_assertion_handler()`.

Parameters

- `expr`: The expression to assert on. This can either be an expression that's convertible to `bool` or a callable which returns `bool`
- `msg`: The optional message that is used to give further information if the assert fails. This should be convertible to a `std::string`

Asserts are enabled if `HPX_DEBUG` is set. This is the default for `CMAKE_BUILD_TYPE=Debug`

HPX_ASSERT_MSG (expr, msg)
HPX_ASSERT

See

namespace hpx

namespace assertion

Typedefs

```
using assertion_handler = void (*) (source_location const &loc, const char *expr,  
                                     std::string const &msg)
```

The signature for an assertion handler.

Functions

```
void set_assertion_handler (assertion_handler handler)
```

Set the assertion handler to be used within a program. If the handler has been set already once, the call to this function will be ignored.

Note This function is not thread safe

```
#include <hpx/assertion/evaluate_assert.hpp>
```

```
#include <hpx/assertion/current_function.hpp>
```

Defines

```
HPX_ASSERT_CURRENT_FUNCTION
```

```
#include <hpx/assertion/force_linking.hpp>
```

```
namespace hpx
```

```
namespace assertion
```

Functions

```
force_linking_helper &force_linking()
```

```
struct force_linking_helper
```

Public Members

```
void (*handle_assert) (source_location const&, const char *, std::string const&)
```

```
#include <hpx/assertion/source_location.hpp>
```

```
namespace hpx
```

```
namespace assertion
```

Functions

```
std::ostream &operator<< (std::ostream &os, source_location const &loc)

struct source_location
#include <source_location.hpp> This contains the location information where HPX_ASSERT has
been called
```

Public Members

```
const char *file_name
unsigned line_number
const char *function_name
```

filesystem

```
#include <hpx/filesystem.hpp>
```

This file provides a compatibility layer using Boost.Filesystem for the C++17 filesystem library. It is *not* intended to be a complete compatibility layer. It only contains functions required by the HPX codebase. It also provides some functions only available in Boost.Filesystem when using C++17 filesystem.

```
namespace hpx
```

```
namespace filesystem
```

Functions

```
path initial_path()
path basename (path const &p)
path canonical (path const &p, path const &base)
path canonical (path const &p, path const &base, error_code &ec)
```

concurrency

```
#include <compatibility/hpx/util/spinlock.hpp>
#include <compatibility/hpx/util/itt_notify.hpp>
#include <compatibility/hpx/util/spinlock_pool.hpp>
#include <compatibility/hpx/util/lockfree/freelist.hpp>
#include <compatibility/hpx/util/lockfree/deque.hpp>
```

```
#include <compatibility/hpx/util/lockfree/concurrentqueue.hpp>

#include <compatibility/hpx/compat/barrier.hpp>

#include <hpx/concurrency/thread_name.hpp>

#include <hpx/concurrency/spinlock.hpp>

namespace hpx

namespace util

struct spinlock
    #include <spinlock.hpp> boost::mutex-compatible spinlock class
```

Public Types

```
typedef boost::detail::spinlock *native_handle_type
```

Public Functions

```
HPX_NON_COPYABLE (spinlock)

spinlock (char const * = nullptr)
~spinlock ()
void lock ()
bool try_lock ()
void unlock ()

native_handle_type native_handle ()
```

Private Members

```
boost::detail::spinlock m = BOOST_DETAIL_SPINLOCK_INIT
```

```
#include <hpx/concurrency/barrier.hpp>

namespace hpx

namespace util

class barrier
```

Public Functions

```
barrier (std::size_t number_of_threads)
~barrier ()
void wait ()
```

Private Types

```
typedef std::mutex mutex_type
```

Private Members

```
HPX_STATIC_CONSTEXPR std::size_t hpx::util::barrier::barrier_flag=
std::size_t const number_of_threads_
std::size_t total_
mutex_type mtx_
std::condition_variable cond_
```

```
#include <hpx/concurrency/itt_notify.hpp>
```

Defines

```
HPX_ITT_SYNC_CREATE (obj, type, name)
HPX_ITT_SYNC_RENAME (obj, name)
HPX_ITT_SYNC_PREPARE (obj)
HPX_ITT_SYNC_CANCEL (obj)
HPX_ITT_SYNC_ACQUIRED (obj)
HPX_ITT_SYNC_RELEASEING (obj)
HPX_ITT_SYNC_RELEASED (obj)
HPX_ITT_SYNC_DESTROY (obj)
HPX_ITT_STACK_CREATE (ctx)
HPX_ITT_STACK_CALLEE_ENTER (ctx)
HPX_ITT_STACK_CALLEE_LEAVE (ctx)
HPX_ITT_STACK_DESTROY (ctx)
HPX_ITT_FRAME_BEGIN (frame, id)
HPX_ITT_FRAME_END (frame, id)
HPX_ITT_MARK_CREATE (mark, name)
HPX_ITT_MARK_OFF (mark)
HPX_ITT_MARK (mark, parameter)
```

```
HPX_ITT_THREAD_SET_NAME(name)
HPX_ITT_THREAD_IGNORE()
HPX_ITT_TASK_BEGIN(domain, name)
HPX_ITT_TASK_BEGIN_ID(domain, id, name)
HPX_ITT_TASK_END(domain)
HPX_ITT_DOMAIN_CREATE(name)
HPX_ITT_STRING_HANDLE_CREATE(name)
HPX_ITT_MAKE_ID(addr, extra)
HPX_ITT_ID_CREATE(domain, id)
HPX_ITT_ID_DESTROY(id)
HPX_ITT_HEAP_FUNCTION_CREATE(name, domain)
HPX_ITT_HEAP_ALLOCATE_BEGIN(f, size, initialized)
HPX_ITT_HEAP_ALLOCATE_END(f, addr, size, initialized)
HPX_ITT_HEAP_FREE_BEGIN(f, addr)
HPX_ITT_HEAP_FREE_END(f, addr)
HPX_ITT_HEAP_REALLOCATE_BEGIN(f, addr, new_size, initialized)
HPX_ITT_HEAP_REALLOCATE_END(f, addr, new_addr, new_size, initialized)
HPX_ITT_HEAP_INTERNAL_ACCESS_BEGIN()
HPX_ITT_HEAP_INTERNAL_ACCESS_END()
HPX_ITT_COUNTER_CREATE(name, domain)
HPX_ITT_COUNTER_CREATE_TYPED(name, domain, type)
HPX_ITT_COUNTER_SET_VALUE(id, value_ptr)
HPX_ITT_COUNTER_DESTROY(id)
HPX_ITT_METADATA_ADD(domain, id, key, data)
```

Typedefs

```
typedef void *__itt_heap_function
```

Functions

```
void itt_sync_create(void *, const char *, const char *)
void itt_sync_rename(void *, const char *)
void itt_sync_prepare(void *)
void itt_sync_acquired(void *)
void itt_sync_cancel(void *)
void itt_sync_releasing(void *)
void itt_sync_released(void *)
```

```
void itt_sync_destroy(void *)
__itt_caller *__itt_stack_create()
void itt_stack_enter(__itt_caller *)
void itt_stack_leave(__itt_caller *)
void itt_stack_destroy(__itt_caller *)
void itt_frame_begin(__itt_domain const *, __itt_id *)
void itt_frame_end(__itt_domain const *, __itt_id *)
int itt_mark_create(char const *)
void itt_mark_off(int)
void itt_mark(int, char const *)
void itt_thread_set_name(char const *)
void itt_thread_ignore()
void itt_task_begin(__itt_domain const *, __itt_string_handle *)
void itt_task_begin(__itt_domain const *, __itt_id *, __itt_string_handle *)
void itt_task_end(__itt_domain const *)
__itt_domain *__itt_domain_create(char const *)
__itt_string_handle *__itt_string_handle_create(char const *)
__itt_id *__itt_make_id(void *, unsigned long)
void itt_id_create(__itt_domain const *, __itt_id *)
void itt_id_destroy(__itt_id *)
__itt_heap_function itt_heap_function_create(const char *, const char *)
void itt_heap_allocate_begin(__itt_heap_function, std::size_t, int)
void itt_heap_allocate_end(__itt_heap_function, void **, std::size_t, int)
void itt_heap_free_begin(__itt_heap_function, void *)
void itt_heap_free_end(__itt_heap_function, void *)
void itt_heap_reallocate_begin(__itt_heap_function, void *, std::size_t, int)
void itt_heap_reallocate_end(__itt_heap_function, void *, void **, std::size_t, int)
void itt_heap_internal_access_begin()
void itt_heap_internal_access_end()
__itt_counter *__itt_counter_create(char const *, char const *)
__itt_counter *__itt_counter_create_typed(char const *, char const *, int)
void itt_counter_destroy(__itt_counter *)
void itt_counter_set_value(__itt_counter *, void *)
int itt_event_create(char const *, int)
int itt_event_start(int)
int itt_event_end(int)
```

```
void itt_metadata_add(__itt_domain *, __itt_id *, __itt_string_handle *, std::uint64_t const&)
void itt_metadata_add(__itt_domain *, __itt_id *, __itt_string_handle *, double const&)
void itt_metadata_add(__itt_domain *, __itt_id *, __itt_string_handle *, char const *)
void itt_metadata_add(__itt_domain *, __itt_id *, __itt_string_handle *, void const *)

namespace hpx

namespace util

namespace itt
```

Functions

```
void event_tick(event const&)
struct caller_context
```

Public Functions

```
caller_context (stack_context&)
~caller_context ()
struct counter
```

Public Functions

```
counter (char const *, char const *)
~counter ()

struct domain
Subclassed by hpx::util::itt::thread_domain
```

Public Functions

```
HPX_NON_COPYABLE (domain)
domain (char const *)
domain ()

struct event
```

Public Functions

```
event (char const *)
struct frame_context
```

Public Functions

```
frame_context (domain const&, id * = nullptr)
~frame_context ()
struct heap_allocate
```

Public Functions

```
template<typename T>
heap_allocate (heap_function&, T **, std::size_t, int)
~heap_allocate ()
struct heap_free
```

Public Functions

```
heap_free (heap_function&, void *)
~heap_free ()
struct heap_function
```

Public Functions

```
heap_function (char const *, char const *)
~heap_function ()
struct heap_internal_access
```

Public Functions

```
heap_internal_access ()
~heap_internal_access ()
struct id
```

Public Functions

```
id(domain const&, void *, unsigned long = 0)  
~id()  
struct mark_context
```

Public Functions

```
mark_context(char const *)  
~mark_context()  
struct mark_event
```

Public Functions

```
mark_event(event const&)  
~mark_event()  
struct stack_context
```

Public Functions

```
stack_context()  
~stack_context()  
struct string_handle
```

Public Functions

```
string_handle(char const * = nullptr)  
struct task
```

Public Functions

```
task(domain const&, string_handle const&, std::uint64_t)  
task(domain const&, string_handle const&)  
~task()  
struct thread_domain : public hpx::util::itt::domain
```

Public Functions

```
HPX_NON_COPYABLE (thread_domain)  
thread_domain()  
struct undo_frame_context
```

Public Functions

```
undo_frame_context (frame_context const&)  
~undo_frame_context ()  
struct undo_mark_context
```

Public Functions

```
undo_mark_context (mark_context const&)  
~undo_mark_context ()
```

```
#include <hpx/concurrency/cache_line_data.hpp>
```

```
template<typename Data>  
struct cache_aligned_data<Data, std::false_type>
```

Public Functions

```
cache_aligned_data()  
cache_aligned_data(Data &&data)  
cache_aligned_data(Data const &data)
```

Public Members

```
Data data_  
namespace hpx  
  
namespace threads
```

Functions

```
constexpr std::size_t get_cache_line_size()  
namespace util
```

TypeDefs

```
using cache_line_data = cache_aligned_data<Data>  
template<typename Data, typename NeedsPadding = typename detail::needs_padding<Data>::type>  
struct cache_aligned_data
```

Public Functions

```
cache_aligned_data()  
cache_aligned_data(Data &&data)  
cache_aligned_data(Data const &data)
```

Public Members

```
Data data_  
template<>  
char cacheline_pad[get_cache_line_padding_size(sizeof(Data))]  
template<typename Data>  
struct cache_aligned_data<Data, std::false_type>
```

Public Functions

```
cache_aligned_data()  
cache_aligned_data(Data &&data)  
cache_aligned_data(Data const &data)
```

Public Members

```
Data data_
```

```
#include <hpx/concurrency/deque.hpp>

namespace boost

namespace lockfree
```

Enums

```
enum deque_status_type
```

Values:

```
stable
```

```
rpush
```

```
lpush
```

```
template<typename T, typename freelist_t = caching_freelist_t, typename Alloc = std::allocator<T>>
struct deque
```

Public Types

```
typedef deque_node<T> node
```

```
typedef node::pointer node_pointer
```

```
typedef node::atomic_pointer atomic_node_pointer
```

```
typedef node::tag_t tag_t
```

```
typedef deque_anchor<T> anchor
```

```
typedef anchor::pair anchor_pair
```

```
typedef anchor::atomic_pair atomic_anchor_pair
```

```
typedef Alloc::template rebind<node>::other node_allocator
```

```
typedef std::conditional<std::is_same<freelist_t, caching_freelist_t>::value, caching_freelist<node, node_allocator>
```

Public Functions

```
HPX_NON_COPYABLE(deque)
```

```
deque(std::size_t initial_nodes = 128)
```

```
~deque()
```

```
bool empty() const
```

```
bool is_lock_free() const
```

```
bool push_left(T const &data)
```

```
bool push_right(T const &data)
```

```
bool pop_left(T &r)
```

```
bool pop_left (T *r)
bool pop_right (T &r)
bool pop_right (T *r)
```

Private Functions

```
node *alloc_node (node *lptr, node *rptr, T const &v, tag_t ltag = 0, tag_t rtag = 0)
void dealloc_node (node *n)
void stabilize_left (anchor_pair &lrs)
void stabilize_right (anchor_pair &lrs)
void stabilize (anchor_pair &lrs)
```

Private Members

```
anchor anchor_
pool pool_
HPX_STATIC_CONSTEXPR std::size_t boost::lockfree::dequeue::padding_size=
template<>
char padding[padding_size]

template<typename T>
struct deque_anchor
```

Public Types

```
typedef dequeue_node<T> node
typedef node::pointer node_pointer
typedef node::atomic_pointer atomic_node_pointer
typedef node::tag_t tag_t
typedef tagged_ptr_pair<node, node> pair
typedef std::atomic<pair> atomic_pair
```

Public Functions

```
deque_anchor ()
deque_anchor (deque_anchor const &p)
deque_anchor (pair const &p)
deque_anchor (node *lptr, node *rptr, tag_t status = stable, tag_t tag = 0)
pair lrs () volatile const
```

```

node *left () volatile const
node *right () volatile const
tag_t status () volatile const
tag_t tag () volatile const

bool cas (deque_anchor &expected, deque_anchor const &desired) volatile
bool cas (pair &expected, deque_anchor const &desired) volatile
bool cas (deque_anchor &expected, pair const &desired) volatile
bool cas (pair &expected, pair const &desired) volatile
bool operator==(volatile deque_anchor const &rhs) const
bool operator!=(volatile deque_anchor const &rhs) const
bool operator==(volatile pair const &rhs) const
bool operator!=(volatile pair const &rhs) const
bool is_lock_free () const

```

Private Members

```

atomic_pair pair_
template<typename T>
struct deque_node

```

Public Types

```

typedef detail::tagged_ptr<deque_node> pointer
typedef std::atomic<pointer> atomic_pointer
typedef pointer::tag_t tag_t

```

Public Functions

```

deque_node ()
deque_node (deque_node const &p)
deque_node (deque_node *lptr, deque_node *rptr, T const &v, tag_t ltag = 0, tag_t rtag = 0)

```

Public Members

```
atomic_pointer left
atomic_pointer right
T data

#include <hpx/concurrency/spinlock_pool.hpp>

namespace hpx

namespace util

template<typename Tag, std::size_t N = HPX_HAVE_SPINLOCK_POOL_NUM>
class spinlock_pool
```

Public Static Functions

```
static boost::detail::spinlock &spinlock_for(void const *pv)
```

Private Static Attributes

```
cache_aligned_data<boost::detail::spinlock> pool_
class scoped_lock
```

Public Functions

```
template<>
HPX_NON_COPYABLE(scoped_lock)

template<>
scoped_lock(void const *pv)

template<>
~scoped_lock()

template<>
void lock()

template<>
void unlock()
```

Private Members

```
template<>
boost::detail::spinlock &sp_
```

```
#include <hpx/concurrency/concurrentqueue.hpp>
```

Defines

```
MOODYCAMEL_THREADLOCAL
MOODYCAMEL_EXCEPTIONS_ENABLED
MOODYCAMEL_TRY
MOODYCAMEL_CATCH(...)
MOODYCAMEL_RETHROW
MOODYCAMEL_THROW(expr)
MOODYCAMEL_NOEXCEPT
MOODYCAMEL_NOEXCEPTCTOR(type, valueType, expr)
MOODYCAMEL_NOEXCEPT_ASSIGN(type, valueType, expr)
MOODYCAMEL_DELETE_FUNCTION
namespace moodycamel
```

Functions

```
template<typename T, typename Traits>
void swap(typename ConcurrentQueue<T, Traits>::ImplicitProducerKVP &a, typename ConcurrentQueue<T, Traits>::ImplicitProducerKVP &b)

template<typename T, typename Traits>
void swap(ConcurrentQueue<T, Traits> &a, ConcurrentQueue<T, Traits> &b)

void swap(ProducerToken &a, ProducerToken &b)
void swap(ConsumerToken &a, ConsumerToken &b)

template<typename T, typename Traits = ConcurrentQueueDefaultTraits>
class ConcurrentQueue
```

Public Types

```
typedef moodycamel::ProducerToken producer_token_t
typedef moodycamel::ConsumerToken consumer_token_t
typedef Traits::index_t index_t
typedef Traits::size_t size_t
```

Public Functions

```
ConcurrentQueue (size_t capacity = 6 * BLOCK_SIZE)

ConcurrentQueue (size_t minCapacity, size_t maxExplicitProducers, size_t maxImplicitProducers)

~ConcurrentQueue ()

ConcurrentQueue (ConcurrentQueue const&)

ConcurrentQueue &operator= (ConcurrentQueue const&)

ConcurrentQueue (ConcurrentQueue &&other)

ConcurrentQueue &operator= (ConcurrentQueue &&other)

void swap (ConcurrentQueue &other)

bool enqueue (T const &item)

bool enqueue (T &&item)

bool enqueue (producer_token_t const &token, T const &item)

bool enqueue (producer_token_t const &token, T &&item)

template<typename It>
bool enqueue_bulk (It itemFirst, size_t count)

bool try_enqueue (T const &item)

bool try_enqueue (T &&item)

bool try_enqueue (producer_token_t const &token, T const &item)

bool try_enqueue (producer_token_t const &token, T &&item)

template<typename It>
bool try_enqueue_bulk (It itemFirst, size_t count)

bool try_enqueue_bulk (producer_token_t const &token, It itemFirst, size_t count)

template<typename U>
bool try_dequeue (U &item)

template<typename U>
bool try_dequeue_non_interleaved (U &item)

template<typename U>
bool try_dequeue (consumer_token_t &token, U &item)

template<typename It>
size_t try_dequeue_bulk (It itemFirst, size_t max)

template<typename It>
```

```

size_t try_dequeue_bulk (consumer_token_t &token, It itemFirst, size_t max)

template<typename U>
bool try_dequeue_from_producer (producer_token_t const &producer, U &item)

template<typename It>
size_t try_dequeue_bulk_from_producer (producer_token_t const &producer, It itemFirst, size_t max)

size_t size_approx () const

```

Public Static Functions

```
static bool is_lock_free ()
```

Public Static Attributes

```

const size_t BLOCK_SIZE = static_cast<size_t>(Traits::BLOCK_SIZE)
const size_t EXPLICIT_BLOCK_EMPTY_COUNTER_THRESHOLD = static_cast<size_t>(Traits::EXPLICIT_BLOCK_E
const size_t EXPLICIT_INITIAL_INDEX_SIZE = static_cast<size_t>(Traits::EXPLICIT_INITIAL_INDEX_SIZE)
const size_t IMPLICIT_INITIAL_INDEX_SIZE = static_cast<size_t>(Traits::IMPLICIT_INITIAL_INDEX_SIZE)
const size_t INITIAL_IMPLICIT_PRODUCER_HASH_SIZE = static_cast<size_t>(Traits::INITIAL_IMPLICIT_PR
const std::uint32_t EXPLICIT_CONSUMER_CONSUMPTION_QUOTA_BEFORE_ROTATE = static_cast<std::uint32_t>
const size_t moodycamel::ConcurrentQueue::MAX_SUBQUEUE_SIZE = (details::const_number_of_subque

```

Private Types

```

enum AllocationMode
    Values:
        CanAlloc
        CannotAlloc

enum InnerQueueContext
    Values:
        implicit_context = 0
        explicit_context = 1

```

Private Functions

```

ConcurrentQueue &swap_internal (ConcurrentQueue &other)

template<AllocationMode canAlloc, typename U>
bool inner_enqueue (producer_token_t const &token, U &&element)

template<AllocationMode canAlloc, typename U>
bool inner_enqueue (U &&element)

template<AllocationMode canAlloc, typename It>

```

```
bool inner_enqueue_bulk (producer_token_t const &token, It itemFirst, size_t count)  
template<AllocationMode canAlloc, typename It>  
bool inner_enqueue_bulk (It itemFirst, size_t count)  
  
bool update_current_producer_after_rotation (consumer_token_t &token)  
  
void populate_initial_block_list (size_t blockCount)  
  
Block *try_get_block_from_initial_pool ()  
  
void add_block_to_free_list (Block *block)  
  
void add_blocks_to_free_list (Block *block)  
  
Block *try_get_block_from_free_list ()  
  
template<AllocationMode canAlloc>  
Block *requisition_block ()  
  
ProducerBase *recycle_or_create_producer (bool isExplicit)  
  
ProducerBase *recycle_or_create_producer (bool isExplicit, bool &recycled)  
  
ProducerBase *add_producer (ProducerBase *producer)  
  
void reown_producers ()  
  
void populate_initial_implicit_producer_hash ()  
  
void swap_implicit_producer_hashes (ConcurrentQueue &other)  
  
ImplicitProducer *get_or_add_implicit_producer ()
```

Private Members

```
std::atomic<ProducerBase *> producerListTail  
std::atomic<std::uint32_t> producerCount  
std::atomic<size_t> initialBlockPoolIndex  
Block *initialBlockPool  
size_t initialBlockPoolSize  
FreeList<Block> freeList  
std::atomic<ImplicitProducerHash *> implicitProducerHash  
std::atomic<size_t> implicitProducerHashCount  
ImplicitProducerHash initialImplicitProducerHash  
std::array<ImplicitProducerKVP, INITIAL_IMPLICIT_PRODUCER_HASH_SIZE> initialImplicitProducerHas  
std::atomic_flag implicitProducerHashResizeInProgress  
std::atomic<std::uint32_t> nextExplicitConsumerId  
std::atomic<std::uint32_t> globalExplicitConsumerOffset
```

Private Static Functions

```
template<typename U>
static U *create_array (size_t count)
```

```
template<typename U>
static void destroy_array (U *p, size_t count)
```

```
template<typename U>
static U *create ()
```

```
template<typename U, typename A1>
static U *create (A1 &&a1)
```

```
template<typename U>
static void destroy (U *p)
```

Friends

```
friend moodycamel::ProducerToken
friend moodycamel::ConsumerToken
friend moodycamel::ExplicitProducer
friend moodycamel::ImplicitProducer
friend moodycamel::ConcurrentQueueTests
```

```
template<typename XT, typename XTraits>
void swap (typename ConcurrentQueue<XT, XTraits>::ImplicitProducerKVP&, typename ConcurrentQueue<XT, XTraits>::ImplicitProducerKVP&)
```

```
struct Block
```

Public Functions

```
template<>
Block ()
```

```
template<InnerQueueContext context>
bool is_empty () const
```

```
template<InnerQueueContext context>
bool set_empty (index_t i)
```

```
template<InnerQueueContext context>
bool set_many_empty (index_t i, size_t count)
```

```
template<InnerQueueContext context>
void set_all_empty ()
```

```
template<InnerQueueContext context>
void reset_empty ()
```

```
template<>
T *operator[] (index_t idx)
```

```
template<>
T const *operator[](index_t idx) const
```

Public Members

```
template<>
char elements[sizeof(T) * BLOCK_SIZE]

template<>
details::max_align_t dummy

template<>
Block *next

template<>
std::atomic<size_t> elementsCompletelyDequeued

std::atomic<bool> moodycamel::ConcurrentQueue< T, Traits >::Block::emptyFlags[BLOCK_SIZE]

template<>
std::atomic<uint32_t> freeListRefs

template<>
std::atomic<Block *> freeListNext

template<>
std::atomic<bool> shouldBeOnFreeList

template<>
bool dynamicallyAllocated
```

Private Members

```
template<>
union moodycamel::ConcurrentQueue::Block::[anonymous] [anonymous]

struct ExplicitProducer : public moodycamel::ConcurrentQueue<T, Traits>::ProducerBase
```

Public Functions

```
template<>
ExplicitProducer(ConcurrentQueue *parent)

template<>
~ExplicitProducer()

template<AllocationMode allocMode, typename U>
bool enqueue(U &&element)

template<typename U>
bool dequeue(U &element)

template<AllocationMode allocMode, typename It>
bool enqueue_bulk(It itemFirst, size_t count)

template<typename It>
size_t dequeue_bulk(It &itemFirst, size_t max)
```

Private Functions

```
template<>
bool new_block_index(size_t numberOfFilledSlotsToExpose)
```

Private Members

```
template<>
std::atomic<BlockIndexHeader *> blockIndex

template<>
size_t pr_blockIndexSlotsUsed

template<>
size_t pr_blockIndexSize

template<>
size_t pr_blockIndexFront

template<>
BlockIndexEntry *pr_blockIndexEntries

template<>
void *pr_blockIndexRaw

struct BlockIndexEntry
```

Public Members

```
template<>
index_t base

template<>
Block *block

struct BlockIndexHeader
```

Public Members

```
template<>
size_t size

template<>
std::atomic<size_t> front

template<>
BlockIndexEntry *entries

template<>
void *prev

template<typename N>
struct FreeList
```

Public Functions

```
template<>
FreeList ()

template<>
FreeList (FreeList &&other)

template<>
void swap (FreeList &other)

template<>
FreeList (FreeList const&)

template<>
FreeList &operator= (FreeList const&)

template<>
void add (N *node)

template<>
N *try_get ()

template<>
N *head_unsafe () const
```

Private Functions

```
template<>
void add_knowing_refcount_is_zero (N *node)
```

Private Members

```
template<>
std::atomic<N *> freeListHead
```

Private Static Attributes

```
template<>
const std::uint32_t REFS_MASK = 0x7FFFFFFF

template<>
const std::uint32_t SHOULD_BE_ON_FREELIST = 0x80000000

template<typename N>
struct FreeListNode
```

Public Functions

```
template<>
FreeListNode ()
```

Public Members

```
template<>
std::atomic<std::uint32_t> freeListRefs

template<>
std::atomic<N *> freeListNext

struct ImplicitProducer : public moodycamel::ConcurrentQueue<T, Traits>::ProducerBase
```

Public Functions

```
template<>
ImplicitProducer (ConcurrentQueue *parent)

template<>
~ImplicitProducer ()

template<AllocationMode allocMode, typename U>
bool enqueue (U &&element)

template<typename U>
bool dequeue (U &element)

template<AllocationMode allocMode, typename It>
bool enqueue_bulk (It itemFirst, size_t count)

template<typename It>
size_t dequeue_bulk (It &itemFirst, size_t max)
```

Private Functions

```
template<AllocationMode allocMode>
bool insert_block_index_entry (BlockIndexEntry *&idxEntry, index_t blockStartIndex)

template<>
void rewind_block_index_tail ()

template<>
BlockIndexEntry *get_block_index_entry_for_index (index_t index) const

template<>
size_t get_block_index_index_for_index (index_t index, BlockIndexHeader *&localBlockIndex) const

template<>
bool new_block_index ()
```

Private Members

```
template<>
size_t nextBlockIndexCapacity

template<>
std::atomic<BlockIndexHeader *> blockIndex
```

Private Static Attributes

```
template<>
const index_t INVALID_BLOCK_BASE = 1

struct BlockIndexEntry
```

Public Members

```
template<>
std::atomic<index_t> key

template<>
std::atomic<Block *> value

struct BlockIndexHeader
```

Public Members

```
template<>
size_t capacity

template<>
std::atomic<size_t> tail

template<>
BlockIndexEntry *entries

template<>
BlockIndexEntry **index

template<>
BlockIndexHeader *prev

struct ImplicitProducerHash
```

Public Members

```
template<>
size_t capacity

template<>
ImplicitProducerKVP *entries

template<>
ImplicitProducerHash *prev

struct ImplicitProducerKVP
```

Public Functions

```
template<>
ImplicitProducerKVP ()

template<>
ImplicitProducerKVP (ImplicitProducerKVP &&other)

template<>
ImplicitProducerKVP &operator= (ImplicitProducerKVP &&other)

template<>
void swap (ImplicitProducerKVP &other)
```

Public Members

```
template<>
std::atomic<details::thread_id_t> key

template<>
ImplicitProducer *value

struct ProducerBase : public moodycamel::details::ConcurrentQueueProducerTypelessBase
```

Public Functions

```
template<>
ProducerBase (ConcurrentQueue *parent_, bool isExplicit_)

template<>
virtual ~ProducerBase ()

template<typename U>
bool dequeue (U &element)

template<typename It>
size_t dequeue_bulk (It &itemFirst, size_t max)

template<>
ProducerBase *next_prod () const

template<>
size_t size_approx () const

template<>
index_t getTail () const
```

Public Members

```
template<>
bool isExplicit

template<>
ConcurrentQueue *parent
```

Protected Attributes

```
template<>
std::atomic<index_t> tailIndex

template<>
std::atomic<index_t> headIndex

template<>
std::atomic<index_t> dequeueOptimisticCount

template<>
std::atomic<index_t> dequeueOvercommit

template<>
Block *tailBlock

struct ConcurrentQueueDefaultTraits
```

Public Types

```
typedef std::size_t size_t
typedef std::size_t index_t
```

Public Static Functions

```
static void *malloc (size_t size)
static void free (void *ptr)
```

Public Static Attributes

```
const size_t BLOCK_SIZE = 32
const size_t EXPLICIT_BLOCK_EMPTY_COUNTER_THRESHOLD = 32
const size_t EXPLICIT_INITIAL_INDEX_SIZE = 32
const size_t IMPLICIT_INITIAL_INDEX_SIZE = 32
const size_t INITIAL_IMPLICIT_PRODUCER_HASH_SIZE = 32
const std::uint32_t EXPLICIT_CONSUMER_CONSUMPTION_QUOTA_BEFORE_ROTATE = 256
const size_t MAX_SUBQUEUE_SIZE = details::const_numeric_max<size_t>::value

struct ConsumerToken
```

Public Functions

```
template<typename T, typename Traits>
ConsumerToken (ConcurrentQueue<T, Traits> &q)
```

```
template<typename T, typename Traits>
ConsumerToken (BlockingConcurrentQueue<T, Traits> &q)
```

```
ConsumerToken (ConsumerToken &&other)
```

```
ConsumerToken &operator= (ConsumerToken &&other)
```

```
void swap (ConsumerToken &other)
```

```
ConsumerToken (ConsumerToken const&)
```

```
ConsumerToken &operator= (ConsumerToken const&)
```

Private Members

```
std::uint32_t initialOffset
```

```
std::uint32_t lastKnownGlobalOffset
```

```
std::uint32_t itemsConsumedFromCurrent
```

```
details::ConcurrentQueueProducerTypelessBase *currentProducer
```

```
details::ConcurrentQueueProducerTypelessBase *desiredProducer
```

Friends

```
friend moodycamel::ConcurrentQueue
```

```
friend moodycamel::ConcurrentQueueTests
```

```
struct ProducerToken
```

Public Functions

```
template<typename T, typename Traits>
ProducerToken (ConcurrentQueue<T, Traits> &queue)
```

```
template<typename T, typename Traits>
ProducerToken (BlockingConcurrentQueue<T, Traits> &queue)
```

```
ProducerToken (ProducerToken &&other)
```

```
ProducerToken &operator= (ProducerToken &&other)
```

```
void swap (ProducerToken &other)
```

```
bool valid () const
```

```
~ProducerToken ()
```

```
ProducerToken (ProducerToken const&)
```

```
ProducerToken &operator= (ProducerToken const&)
```

Protected Attributes

```
details::ConcurrentQueueProducerTypelessBase *producer
```

Friends

```
friend moodycamel::ConcurrentQueue  
friend moodycamel::ConcurrentQueueTests
```

```
namespace details
```

Typedefs

```
typedef std::uintptr_t thread_id_t  
typedef std::max_align_t std_max_align_t
```

Functions

```
static thread_id_t thread_id()  
static bool() moodycamel::details::likely(bool x)  
static bool() moodycamel::details::unlikely(bool x)  
static size_t hash_thread_id(thread_id_t id)  
template<typename T>  
static bool circular_less_than(T a, T b)  
template<typename U>  
static char *align_for(char *ptr)  
template<typename T>  
static T ceil_to_pow_2(T x)  
template<typename T>  
static void swap_relaxed(std::atomic<T> &left, std::atomic<T> &right)  
template<typename T>  
static T const &nomove(T const &x)  
template<typename It>  
static auto deref_noexcept(It &it)
```

Variables

```
const thread_id_t invalid_thread_id = 0
const thread_id_t invalid_thread_id2 = 1

template<bool use32>
struct _hash_32_or_64
```

Public Static Functions

```
static std::uint32_t hash (std::uint32_t h)
template<>
struct _hash_32_or_64<1>
```

Public Static Functions

```
static std::uint64_t hash (std::uint64_t h)
struct ConcurrentQueueProducerTypelessBase
```

Public Functions

```
ConcurrentQueueProducerTypelessBase ()
```

Public Members

```
ConcurrentQueueProducerTypelessBase *next
std::atomic<bool> inactive
ProducerToken *token

template<typename T>
struct const_numeric_max
```

Public Static Attributes

```
const T moodycamel::details::const_numeric_max::value = std::numeric_limits<T>::
union max_align_t
```

Public Members

```
std_max_align_t x
```

```
long long y
```

```
void *z
```

```
template<bool Enable>
struct nomove_if
```

Public Static Functions

```
template<typename T>
static T const &eval(T const &x)
```

```
template<>
struct nomove_if<false>
```

Public Static Functions

```
template<typename U>
static auto eval(U &&x)
```

```
template<>
struct static_is_lock_free<bool>
```

Public Types

```
enum [anonymous]
```

Values:

```
    value = ATOMIC_BOOL_LOCK_FREE
```

```
template<typename U>
struct static_is_lock_free<U *>
```

Public Types

```
enum [anonymous]
```

Values:

```
    value = ATOMIC_POINTER_LOCK_FREE
```

```
template<typename T>
struct static_is_lock_free_num
```

Public Types

```
enum [anonymous]
Values:
    value = 0

template<>
struct static_is_lock_free_num<int>
```

Public Types

```
enum [anonymous]
Values:
    value = ATOMIC_INT_LOCK_FREE

template<>
struct static_is_lock_free_num<long>
```

Public Types

```
enum [anonymous]
Values:
    value = ATOMIC_LONG_LOCK_FREE

template<>
struct static_is_lock_free_num<long long>
```

Public Types

```
enum [anonymous]
Values:
    value = ATOMIC_LLONG_LOCK_FREE

template<>
struct static_is_lock_free_num<short>
```

Public Types

```
enum [anonymous]
Values:
    value = ATOMIC_SHORT_LOCK_FREE

template<>
struct static_is_lock_free_num<signed char>
```

Public Types

```
enum [anonymous]
Values:
    value = ATOMIC_CHAR_LOCK_FREE
```

```
template<typename thread_id_t>
struct thread_id_converter
```

Public Types

```
typedef thread_id_t thread_id_numeric_size_t
typedef thread_id_t thread_id_hash_t
```

Public Static Functions

```
static thread_id_hash_t prehash(thread_id_t const &x)
```

hardware

```
#include <compatibility/hpx/util/hardware/cpuid.hpp>

#include <compatibility/hpx/util/hardware/timestamp.hpp>

#include <compatibility/hpx/util/hardware/bit_manipulation.hpp>

#include <compatibility/hpx/util/hardware/cpuid/msvc.hpp>

#include <compatibility/hpx/util/hardware/cpuid/linux_x86.hpp>

#include <compatibility/hpx/util/hardware/timestamp/msvc.hpp>

#include <compatibility/hpx/util/hardware/timestamp/linux_generic.hpp>

#include <compatibility/hpx/util/hardware/timestamp/bgq.hpp>

#include <compatibility/hpx/util/hardware/timestamp/linux_x86_64.hpp>

#include <compatibility/hpx/util/hardware/timestamp/linux_x86_32.hpp>

#include <hpx/hardware/cpuid.hpp>

#include <hpx/hardware/timestamp.hpp>
```

```
#include <hpx/hardware/bit_manipulation.hpp>

namespace hpx

namespace util

namespace hardware
```

Functions

```
template<typename T, typename U>
bool has_bit_set (T value, U bit)

template<std::size_t N, typename T>
T unbounded_shl (T x)

template<std::size_t N, typename T>
T unbounded_shr (T x)

template<std::size_t Low, std::size_t High, typename Result, typename T>
Result get_bit_range (T x)

template<std::size_t Low, typename Result, typename T>
Result pack_bits (T x)

template<std::size_t N, typename T>
struct unbounded_shifter
```

Public Static Functions

```
static T shl (T x)

static T shr (T x)

template<typename T>
struct unbounded_shifter<0, T>
```

Public Static Functions

```
static T shl (T x)

static T shr (T x)
```

```
#include <hpx/hardware/cpuid/msvc.hpp>

#include <hpx/hardware/cpuid/linux_x86.hpp>

namespace hpx

namespace util

namespace hardware
```

Functions

```
void cpuid(std::uint32_t &cpuinfo)[4], std::uint32_t eax)
void cpuidex(std::uint32_t &cpuinfo)[4], std::uint32_t eax, std::uint32_t ecx)

struct cpuid_register
```

Public Types

```
enum info
  Values:
    eax = 0
    ebx = 1
    ecx = 2
    edx = 3
```

```
#include <hpx/hardware/timestamp/msvc.hpp>

#include <hpx/hardware/timestamp/linux_generic.hpp>

#include <hpx/hardware/timestamp/bqq.hpp>

#include <hpx/hardware/timestamp/linux_x86_64.hpp>

namespace hpx

namespace util

namespace hardware
```

Functions

```
std::uint64_t timestamp()

#include <hpx/hardware/timestamp/linux_x86_32.hpp>

naming_base

#include <hpx/naming_base.hpp>

namespace hpx

namespace naming
```

TypeDefs

```
using component_type = std::int32_t
using address_type = std::uint64_t
```

Variables

```
constexpr std::uint32_t invalid_locality_id = ~static_cast<std::uint32_t>(0)
```

static_reinit

```
#include <compatibility/hpx/util/reinitializable_static.hpp>

#include <compatibility/hpx/util/static_reinit.hpp>

#include <hpx/static_reinit/reinitializable_static.hpp>
```

Defines

```
HPX_EXPORT_REINITIALIZABLE_STATIC

namespace hpx

namespace util
```

Variables

```
template<typename T, typename Tag = T, std::size_t N = 1>
struct HPX_EXPORT_REINITIALIZABLE_STATIC reinitializable_static

template<typename T, typename Tag, std::size_t N>
struct reinitializable_static
```

Public Types

```
typedef T value_type
typedef T &reference
typedef T const &const_reference
```

Public Functions

```
HPX_NON_COPYABLE(reinitializable_static)

reinitializable_static()

template<typename U>
reinitializable_static(U const &val)

operator reference()

operator const_reference() const

reference get(std::size_t item = 0)

const_reference get(std::size_t item = 0) const
```

Private Types

```
typedef std::add_pointer<value_type>::type pointer
typedef std::aligned_storage<sizeof(value_type), std::alignment_of<value_type>::value>::type storage_type
```

Private Static Functions

```
static void default_construct()

template<typename U>
static void value_construct(U const &v)

static void destruct()

static void default_constructor()

template<typename U>
static void value_constructor(U const *pv)

static pointer get_address(std::size_t item)
```

Private Static Attributes

```
reinitializable_static<T, Tag, N>::storage_type data_
std::once_flag constructed_

#include <hpx/static_reinit/static_reinit.hpp>

namespace hpx

namespace util
```

Functions

```
void reinit_register(util::function_nonsc<void>
    > const &constructutil::function_nonsc<void> const &destruct)

void reinit_construct()

void reinit_destruct()
```

resiliency

```
#include <hpx/resiliency/async_replay.hpp>

namespace hpx

namespace resiliency
```

Functions

```
template<typename Pred, typename F, typename ...TsF, Ts...>::type> async_replay_validate(std::size_t
    n,
    Pred
    &&pred,
    F
    &&f,
    Ts&&...
    ts)
```

Asynchronously launch given function *f*. Verify the result of those invocations using the given predicate *pred*. Repeat launching on error exactly *n* times (except if *abort_replay_exception* is thrown).

```
template<typename F, typename ...TsF, Ts...>::type> async_replay(std::size_t
    n,
    F
    &&f,
    Ts&&...
    ts)
```

Asynchronously launch given function f . Repeat launching on error exactly n times (except if `abort_replay_exception` is thrown).

namespace functional

Functional version of `hpx::resiliency::async_replay`.

Functional version of `hpx::resiliency::async_replicate_validate` and `hpx::resiliency::async_replicate`

struct async_replay**Public Functions**

```
template<typename F, typename ...Ts>
auto operator()(std::size_t n, F &&f, Ts&&... ts) const
```



```
struct async_replay_validate
```

Public Functions

```
template<typename Pred, typename F, typename ...Ts>
auto operator()(std::size_t n, Pred &&pred, F &&f, Ts&&... ts) const
```

```
#include <hpx/resiliency/resiliency.hpp>

#include <hpx/resiliency/async_replicate.hpp>
```



```
namespace hpx
```

namespace resiliency**Functions**

```
template<typename Vote, typename Pred, typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_replicate_vote_validate
```

Asynchronously launch given function f exactly n times. Verify the result of those invocations using the given predicate $pred$. Run all the valid results against a user provided voting function. Return the valid output.

```
template<typename Vote, typename F, typename ...Ts>
```

```
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_replicate_vote(std::size_t
n,
Vote
&&vote,
F
&&f,
Ts&&...
ts)
```

Asynchronously launch given function f exactly n times. Verify the result of those invocations using the given predicate $pred$. Run all the valid results against a user provided voting function. Return the valid output.

```
template<typename Pred, typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_replicate_validate(std::size_t
n,
Pred
&&
F
&&
Ts...
ts)
```

Asynchronously launch given function f exactly n times. Verify the result of those invocations using the given predicate $pred$. Return the first valid result.

```
template<typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_replicate(std::size_t
n,
F
&&f,
Ts&&...
ts)
```

Asynchronously launch given function f exactly n times. Verify the result of those invocations by checking for exception. Return the first valid result.

namespace functional

Functional version of `hpx::resiliency::async_replay`.

Functional version of `hpx::resiliency::async_replicate_validate` and `hpx::resiliency::async_replicate`

struct async_replicate

Public Functions

```
template<typename F, typename ...Ts>
auto operator()(std::size_t n, F &&f, Ts&&... ts) const
```

struct async_replicate_validate

Public Functions

```
template<typename Pred, typename F, typename ...Ts>
auto operator() (std::size_t n, Pred &&pred, F &&f, Ts&&... ts) const
struct async_replicate_vote
```

Public Functions

```
template<typename Vote, typename F, typename ...Ts>
auto operator() (std::size_t n, Vote &&vote, F &&f, Ts&&... ts) const
struct async_replicate_vote_validate
```

Public Functions

```
template<typename Vote, typename Pred, typename F, typename ...Ts>
auto operator() (std::size_t n, Vote &&vote, Pred &&pred, F &&f, Ts&&... ts) const
```

```
#include <hpx/resiliency/version.hpp>
```

Defines

```
HPX_RESILIENCY_VERSION_FULL
HPX_RESILIENCY_VERSION_MAJOR
HPX_RESILIENCY_VERSION_MINOR
HPX_RESILIENCY_VERSION_SUBMINOR
HPX_RESILIENCY_VERSION_DATE
namespace hpx

namespace resiliency
```

Functions

```
unsigned int major_version()
unsigned int minor_version()
unsigned int subminor_version()
unsigned long full_version()
std::string full_version_str()
```

```
#include <hpx/resiliency/config.hpp>
#include <hpx/resiliency/force_linking.hpp>

namespace hpx

namespace resiliency
```

Functions

```
force_linking_helper &force_linking()

struct force_linking_helper
```

Public Members

```
unsigned int (*major_version)()
unsigned int (*minor_version)()
unsigned int (*subminor_version)()
unsigned long (*full_version)()
std::string (*full_version_str)()
```

```
#include <hpx/resiliency/dataflow_replicate.hpp>
```

```
namespace hpx
```

```
namespace resiliency
```

Functions

```
template<typename Vote, typename Pred, typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replicate_vote_vali
```

Launch given function f exactly n times. Run all the valid results against a user provided voting function. Return the valid output.

Delay the invocation of f if any of the arguments to f are futures.

```
template<typename Vote, typename F, typename ...Ts>
```

```
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replicate_vote(std::function<void(F, Ts...)> f, std::function<bool(F, Ts...)> pred, std::vector<Ts> ts, std::function<void(Ts, Ts)> merge_fn)
```

Launch given function f exactly n times. Run all the valid results against a user provided voting function. Return the valid output.

Delay the invocation of f if any of the arguments to f are futures.

```
template<typename Pred, typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replicate_validate(std::function<void(F, Ts...)> f, std::function<bool(F, Ts...)> pred, std::vector<Ts> ts, std::function<void(Ts, Ts)> merge_fn)
```

Launch given function f exactly n times. Verify the result of those invocations using the given predicate $pred$. Return the first valid result.

Delay the invocation of f if any of the arguments to f are futures.

```
template<typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replicate(std::size_t n, std::function<void(F, Ts...)> f, std::vector<Ts> ts, std::function<void(Ts, Ts)> merge_fn)
```

Launch given function f exactly n times. Return the first valid result.

Delay the invocation of f if any of the arguments to f are futures.

```
#include <hpx/resiliency/dataflow_replay.hpp>

namespace hpx

    namespace resiliency
```

Functions

```
template<typename Pred, typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replay_validate(std::function<void(F, Ts...)> f, std::function<bool(F, Ts...)> pred, std::vector<Ts> ts, std::function<void(Ts, Ts)> merge_fn)
```

Asynchronously launch given function *f*. Verify the result of those invocations using the given predicate *pred*. Repeat launching on error exactly *n* times.

Delay the invocation of *f* if any of the arguments to *f* are futures.

```
template<typename F, typename ...Ts>
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> dataflow_replay(std::size_t
n,
F
&&f,
Ts&&...
ts)
```

Asynchronously launch given function *f*. Repeat launching on error exactly *n* times.

Delay the invocation of *f* if any of the arguments to *f* are futures.

errors

```
#include <compatibility/hpx/error_code.hpp>

#include <compatibility/hpx/exception.hpp>

#include <compatibility/hpx/exception_fwd.hpp>

#include <compatibility/hpx/exception_list.hpp>

#include <compatibility/hpx/throw_exception.hpp>

#include <compatibility/hpx/error.hpp>

#include <compatibility/hpx/exception_info.hpp>

#include <hpx/errors/error_code.hpp>
```

namespace hpx

Unnamed Group

error_code **make_error_code** (*error e, throwmode mode = plain*)

Returns a new *error_code* constructed from the given parameters.

error_code **make_error_code** (*error e, char const *func, char const *file, long line, throwmode mode = plain*)

error_code **make_error_code** (*error e, char const *msg, throwmode mode = plain*)

Returns *error_code*(*e, msg, mode*).

error_code **make_error_code** (*error e, char const *msg, char const *func, char const *file, long line, throwmode mode = plain*)

error_code **make_error_code** (*error e, std::string const &msg, throwmode mode = plain*)

Returns *error_code*(*e, msg, mode*).

```
error_code make_error_code(error e, std::string const &msg, char const *func, char const
                           *file, long line, throwmode mode = plain)

error_code make_error_code(std::exception_ptr const &e)
```

Functions

`boost::system::error_category const &get_hpx_category()`

Returns generic HPX error category used for new errors.

`boost::system::error_category const &get_hpx_rethrow_category()`

Returns generic HPX error category used for errors re-thrown after the exception has been de-serialized.

`error_code make_success_code(throwmode mode = plain)`

Returns `error_code`(hpx::success, “success”, mode).

`class error_code : public error_code`

#include <error_code.hpp> A `hpx::error_code` represents an arbitrary error condition.

The class `hpx::error_code` describes an object used to hold error code values, such as those originating from the operating system or other low-level application program interfaces.

Note Class `hpx::error_code` is an adjunct to error reporting by exception

Public Functions

`error_code(throwmode mode = plain)`

Construct an object of type `error_code`.

Parameters

- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- nothing:

`error_code(error e, throwmode mode = plain)`

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- nothing:

`error_code(error e, char const *func, char const *file, long line, throwmode mode = plain)`

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- func: The name of the function where the error was raised.
- file: The file name of the code where the error was raised.
- line: The line number of the code line where the error was raised.

- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- nothing:

error_code (`error e, char const *msg, throwmode mode = plain`)

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- msg: The parameter msg holds the error message the new exception should encapsulate.
- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- `std::bad_alloc`: (if allocation of a copy of the passed string fails).

error_code (`error e, char const *msg, char const *func, char const *file, long line, throwmode mode = plain`)

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- msg: The parameter msg holds the error message the new exception should encapsulate.
- func: The name of the function where the error was raised.
- file: The file name of the code where the error was raised.
- line: The line number of the code line where the error was raised.
- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- `std::bad_alloc`: (if allocation of a copy of the passed string fails).

error_code (`error e, std::string const &msg, throwmode mode = plain`)

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- msg: The parameter msg holds the error message the new exception should encapsulate.
- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is `plain`, this is the default) or to the category `hpx_category_rethrow` (if mode is `rethrow`).

Exceptions

- `std::bad_alloc`: (if allocation of a copy of the passed string fails).

error_code (`error e, std::string const &msg, char const *func, char const *file, long line, throwmode mode = plain`)

Construct an object of type `error_code`.

Parameters

- e: The parameter e holds the hpx::error code the new exception should encapsulate.
- msg: The parameter msg holds the error message the new exception should encapsulate.
- func: The name of the function where the error was raised.
- file: The file name of the code where the error was raised.
- line: The line number of the code line where the error was raised.

- mode: The parameter mode specifies whether the constructed `hpx::error_code` belongs to the error category `hpx_category` (if mode is *plain*, this is the default) or to the category `hpx_category_rethrow` (if mode is *rethrow*).

Exceptions

- `std::bad_alloc`: (if allocation of a copy of the passed string fails).

`std::string get_message() const`

Return a reference to the error message stored in the `hpx::error_code`.

Exceptions

- nothing:

`void clear()`

Clear this `error_code` object. The postconditions of invoking this method are.

- `value() == hpx::success` and `category() == hpx::get_hpx_category()`

`error_code(error_code const &rhs)`

Copy constructor for `error_code`

Note This function maintains the error category of the left hand side if the right hand side is a success code.

`error_code &operator=(error_code const &rhs)`

Assignment operator for `error_code`

Note This function maintains the error category of the left hand side if the right hand side is a success code.

Private Functions

`error_code(int err, hpx::exception const &e)`

`error_code(std::exception_ptr const &e)`

Private Members

`std::exception_ptr exception_`

Friends

`friend hpx::exception`

`error_code make_error_code(std::exception_ptr const &e)`

`#include <hpx/errors/exception.hpp>`

`namespace hpx`

TypeDefs

```
using custom_exception_info_handler_type = std::function<hpx::exception_info (std::string
    const&, std::string const&, long, std::string const&) >

using pre_exception_handler_type = std::function<void () >
```

Functions

void **set_custom_exception_info_handler** (*custom_exception_info_handler_type f*)

void **set_pre_exception_handler** (*pre_exception_handler_type f*)

std::string get_error_what (*exception_info const &xi*)

Return the error message of the thrown exception.

The function *hpx::get_error_what* can be used to extract the diagnostic information element representing the error message as stored in the given exception instance.

Return The error message stored in the exception If the exception instance does not hold this information, the function will return an empty string.

See *hpx::diagnostic_information()*, *hpx::get_error_host_name()*, *hpx::get_error_process_id()*, *hpx::get_error_function_name()*, *hpx::get_error_file_name()*, *hpx::get_error_line_number()*, *hpx::get_error_os_thread()*, *hpx::get_error_thread_id()*, *hpx::get_error_thread_description()*, *hpx::get_error()* *hpx::get_error_backtrace()*, *hpx::get_error_env()*, *hpx::get_error_config()*, *hpx::get_error_state()*

Parameters

- *xi*: The parameter *e* will be inspected for the requested diagnostic information elements which have been stored at the point where the exception was thrown. This parameter can be one of the following types: *hpx::exception_info*, *hpx::error_code*, *std::exception*, or *std::exception_ptr*.

Exceptions

- *std::bad_alloc*: (if one of the required allocations fails)

error get_error (*hpx::exception const &e*)

Return the error code value of the exception thrown.

The function *hpx::get_error* can be used to extract the diagnostic information element representing the error value code as stored in the given exception instance.

Return The error value code of the locality where the exception was thrown. If the exception instance does not hold this information, the function will return *hpx::naming::invalid_locality_id*.

See *hpx::diagnostic_information()*, *hpx::get_error_host_name()*, *hpx::get_error_process_id()*, *hpx::get_error_function_name()*, *hpx::get_error_file_name()*, *hpx::get_error_line_number()*, *hpx::get_error_os_thread()*, *hpx::get_error_thread_id()*, *hpx::get_error_thread_description()*, *hpx::get_error_backtrace()*, *hpx::get_error_env()*, *hpx::get_error_what()*, *hpx::get_error_config()*, *hpx::get_error_state()*

Parameters

- *e*: The parameter *e* will be inspected for the requested diagnostic information elements which have been stored at the point where the exception was thrown. This parameter can be one of the following types: `hpx::exception`, `hpx::error_code`, or `std::exception_ptr`.

Exceptions

- nothing:

`error get_error(hpx::error_code const &e)`

`std::string get_error_function_name(hpx::exception_info const &xi)`

Return the function name from which the exception was thrown.

The function `hpx::get_error_function_name` can be used to extract the diagnostic information element representing the name of the function as stored in the given exception instance.

Return The name of the function from which the exception was thrown. If the exception instance does not hold this information, the function will return an empty string.

See `hpx::diagnostic_information()`, `hpx::get_error_host_name()`, `hpx::get_error_process_id()`,
`hpx::get_error_file_name()`, `hpx::get_error_line_number()`, `hpx::get_error_os_thread()`,
`hpx::get_error_thread_id()`, `hpx::get_error_thread_description()`, `hpx::get_error()`,
`hpx::get_error_backtrace()`, `hpx::get_error_env()`, `hpx::get_error_what()`, `hpx::get_error_config()`,
`hpx::get_error_state()`

Parameters

- *xi*: The parameter *e* will be inspected for the requested diagnostic information elements which have been stored at the point where the exception was thrown. This parameter can be one of the following types: `hpx::exception_info`, `hpx::error_code`, `std::exception`, or `std::exception_ptr`.

Exceptions

- `std::bad_alloc`: (if one of the required allocations fails)

`std::string get_error_file_name(hpx::exception_info const &xi)`

Return the (source code) file name of the function from which the exception was thrown.

The function `hpx::get_error_file_name` can be used to extract the diagnostic information element representing the name of the source file as stored in the given exception instance.

Return The name of the source file of the function from which the exception was thrown. If the exception instance does not hold this information, the function will return an empty string.

See `hpx::diagnostic_information()`, `hpx::get_error_host_name()`, `hpx::get_error_process_id()`,
`hpx::get_error_function_name()`, `hpx::get_error_line_number()`, `hpx::get_error_os_thread()`,
`hpx::get_error_thread_id()`, `hpx::get_error_thread_description()`, `hpx::get_error()`,
`hpx::get_error_backtrace()`, `hpx::get_error_env()`, `hpx::get_error_what()`, `hpx::get_error_config()`,
`hpx::get_error_state()`

Parameters

- *xi*: The parameter *e* will be inspected for the requested diagnostic information elements which have been stored at the point where the exception was thrown. This parameter can be one of the following types: `hpx::exception_info`, `hpx::error_code`, `std::exception`, or `std::exception_ptr`.

Exceptions

- `std::bad_alloc`: (if one of the required allocations fails)

```
long get_error_line_number(hpx::exception_info const &xi)
```

Return the line number in the (source code) file of the function from which the exception was thrown.

The function *hpx::get_error_line_number* can be used to extract the diagnostic information element representing the line number as stored in the given exception instance.

Return The line number of the place where the exception was thrown. If the exception instance does not hold this information, the function will return -1.

See *hpx::diagnostic_information()*, *hpx::get_error_host_name()*, *hpx::get_error_process_id()*, *hpx::get_error_function_name()*, *hpx::get_error_file_name()*, *hpx::get_error_os_thread()*, *hpx::get_error_thread_id()*, *hpx::get_error_thread_description()*, *hpx::get_error()*, *hpx::get_error_backtrace()*, *hpx::get_error_env()*, *hpx::get_error_what()*, *hpx::get_error_config()*, *hpx::get_error_state()*

Parameters

- *xi*: The parameter *e* will be inspected for the requested diagnostic information elements which have been stored at the point where the exception was thrown. This parameter can be one of the following types: *hpx::exception_info*, *hpx::error_code*, *std::exception*, or *std::exception_ptr*.

Exceptions

- nothing:

```
class exception : public system_error
```

#include <exception.hpp> A *hpx::exception* is the main exception type used by HPX to report errors.

The *hpx::exception* type is the main exception type used by HPX to report errors. Any exceptions thrown by functions in the HPX library are either of this type or of a type derived from it. This implies that it is always safe to use this type only in catch statements guarding HPX library calls.

Subclassed by *hpx::exception_list*

Public Functions

exception(error e = success)

Construct a *hpx::exception* from a *hpx::error*.

Parameters

- *e*: The parameter *e* holds the *hpx::error* code the new exception should encapsulate.

exception(boost::system::system_error const &e)

Construct a *hpx::exception* from a *boost::system_error*.

exception(boost::system::error_code const &e)

Construct a *hpx::exception* from a *boost::system::error_code* (this is new for Boost V1.69). This constructor is required to compensate for the changes introduced as a resolution to LWG3162 (<https://cplusplus.github.io/LWG/issue3162>).

exception(error e, char const *msg, throwmode mode = plain)

Construct a *hpx::exception* from a *hpx::error* and an error message.

Parameters

- *e*: The parameter *e* holds the *hpx::error* code the new exception should encapsulate.
- *msg*: The parameter *msg* holds the error message the new exception should encapsulate.
- *mode*: The parameter *mode* specifies whether the returned *hpx::error_code* belongs to the error category *hpx_category* (if *mode* is *plain*, this is the default) or to the category *hpx_category_rethrow* (if *mode* is *rethrow*).

```
exception(error e, std::string const &msg, throwmode mode = plain)
```

Construct a `hpx::exception` from a `hpx::error` and an error message.

Parameters

- `e`: The parameter `e` holds the `hpx::error` code the new exception should encapsulate.
- `msg`: The parameter `msg` holds the error message the new exception should encapsulate.
- `mode`: The parameter `mode` specifies whether the returned `hpx::error_code` belongs to the error category `hpx_category` (if `mode` is `plain`, this is the default) or to the category `hpx_category_rethrow` (if `mode` is `rethrow`).

```
~exception()
```

Destruct a `hpx::exception`

Exceptions

- nothing:

```
error get_error() const
```

The function `get_error()` returns the `hpx::error` code stored in the referenced instance of a `hpx::exception`. It returns the `hpx::error` code this exception instance was constructed from.

Exceptions

- nothing:

```
error_code get_error_code(throwmode mode = plain) const
```

The function `get_error_code()` returns a `hpx::error_code` which represents the same error condition as this `hpx::exception` instance.

Parameters

- `mode`: The parameter `mode` specifies whether the returned `hpx::error_code` belongs to the error category `hpx_category` (if `mode` is `plain`, this is the default) or to the category `hpx_category_rethrow` (if `mode` is `rethrow`).

```
struct thread_interrupted : public exception
```

#include <exception.hpp> A `hpx::thread_interrupted` is the exception type used by HPX to interrupt a running HPX thread.

The `hpx::thread_interrupted` type is the exception type used by HPX to interrupt a running thread.

A running thread can be interrupted by invoking the `interrupt()` member function of the corresponding `hpx::thread` object. When the interrupted thread next executes one of the specified interruption points (or if it is currently blocked whilst executing one) with interruption enabled, then a `hpx::thread_interrupted` exception will be thrown in the interrupted thread. If not caught, this will cause the execution of the interrupted thread to terminate. As with any other exception, the stack will be unwound, and destructors for objects of automatic storage duration will be executed.

If a thread wishes to avoid being interrupted, it can create an instance of `hpx::this_thread::disable_interruption`. Objects of this class disable interruption for the thread that created them on construction, and restore the interruption state to whatever it was before on destruction.

```
void f()
{
    // interruption enabled here
    {
        hpx::this_thread::disable_interruption di;
        // interruption disabled
        {
            hpx::this_thread::disable_interruption di2;
            // interruption still disabled
        } // di2 destroyed, interruption state restored
    }
}
```

(continues on next page)

(continued from previous page)

```
// interruption still disabled
} // di destroyed, interruption state restored
// interruption now enabled
}
```

The effects of an instance of `hpx::this_thread::disable_interruption` can be temporarily reversed by constructing an instance of `hpx::this_thread::restore_interruption`, passing in the `hpx::this_thread::disable_interruption` object in question. This will restore the interruption state to what it was when the `hpx::this_thread::disable_interruption` object was constructed, and then disable interruption again when the `hpx::this_thread::restore_interruption` object is destroyed.

```
void g()
{
    // interruption enabled here
    {
        hpx::this_thread::disable_interruption di;
        // interruption disabled
        {
            hpx::this_thread::restore_interruption ri(di);
            // interruption now enabled
        } // ri destroyed, interruption disable again
    } // di destroyed, interruption state restored
    // interruption now enabled
}
```

At any point, the interruption state for the current thread can be queried by calling `hpx::this_thread::interruption_enabled()`.

```
#include <hpx/errors/exception_fwd.hpp>

namespace hpx
```

Enums

enum throwmode

Encode error category for new `error_code`.

Values:

```
plain = 0
rethrow = 1
lightweight = 0x80
```

Variables

`error_code throws`

Predefined `error_code` object used as “throw on error” tag.

The predefined `hpx::error_code` object `hpx::throws` is supplied for use as a “throw on error” tag.

Functions that specify an argument in the form ‘`error_code& ec=throws`’ (with appropriate namespace qualifiers), have the following error handling semantics:

If `&ec != &throws` and an error occurred: `ec.value()` returns the implementation specific error number for the particular error that occurred and `ec.category()` returns the `error_category` for `ec.value()`.

If `&ec != &throws` and an error did not occur, `ec.clear()`.

If an error occurs and `&ec == &throws`, the function throws an exception of type `hpx::exception` or of a type derived from it. The exception’s `get_errorcode()` member function returns a reference to an `hpx::error_code` object with the behavior as specified above.

```
#include <hpx/errors/exception_list.hpp>
```

```
namespace hpx
```

`class exception_list : public hpx::exception`

`#include <exception_list.hpp>` The class `exception_list` is a container of `exception_ptr` objects parallel algorithms may use to communicate uncaught exceptions encountered during parallel execution to the caller of the algorithm

The type `exception_list::const_iterator` fulfills the requirements of a forward iterator.

Public Types

```
typedef exception_list_type::const_iterator iterator  
bidirectional iterator
```

Public Functions

`std::size_t size() const`

The number of `exception_ptr` objects contained within the `exception_list`.

Note Complexity: Constant time.

`exception_list_type::const_iterator begin() const`

An iterator referring to the first `exception_ptr` object contained within the `exception_list`.

`exception_list_type::const_iterator end() const`

An iterator which is the past-the-end value for the `exception_list`.

Private Types

```
typedef boost::detail::spinlock mutex_type
typedef std::list<std::exception_ptr> exception_list_type
```

Private Members

```
exception_list_type exceptions_
mutex_type mtx_
```

```
#include <hpx/errors/throw_exception.hpp>
```

Defines

HPX_THROW_EXCEPTION (errcode, f, msg)

Throw a *hpx::exception* initialized from the given parameters.

The macro *HPX_THROW_EXCEPTION* can be used to throw a *hpx::exception*. The purpose of this macro is to prepend the source file name and line number of the position where the exception is thrown to the error message. Moreover, this associates additional diagnostic information with the exception, such as file name and line number, locality id and thread id, and stack backtrace from the point where the exception was thrown.

The parameter *errcode* holds the *hpx::error* code the new exception should encapsulate. The parameter *f* is expected to hold the name of the function exception is thrown from and the parameter *msg* holds the error message the new exception should encapsulate.

```
void raise_exception()
{
    // Throw a hpx::exception initialized from the given parameters.
    // Additionally associate with this exception some detailed
    // diagnostic information about the throw-site.
    HPX_THROW_EXCEPTION(hpx::no_success, "raise_exception", "simulated error");
}
```

Example:

HPX_THROWS_IF (ec, errcode, f, msg)

Either throw a *hpx::exception* or initialize *hpx::error_code* from the given parameters.

The macro *HPX_THROWS_IF* can be used to either throw a *hpx::exception* or to initialize a *hpx::error_code* from the given parameters. If *&ec == &hpx::throws*, the semantics of this macro are equivalent to *HPX_THROW_EXCEPTION*. If *&ec != &hpx::throws*, the *hpx::error_code* instance *ec* is initialized instead.

The parameter *errcode* holds the *hpx::error* code from which the new exception should be initialized. The parameter *f* is expected to hold the name of the function exception is thrown from and the parameter *msg* holds the error message the new exception should encapsulate.

```
#include <hpx/errors/error.hpp>

namespace hpx
```

Enums

`enum error`

Possible error conditions.

This enumeration lists all possible error conditions which can be reported from any of the API functions.

Values:

`success = 0`

The operation was successful.

`no_success = 1`

The operation did failed, but not in an unexpected manner.

`not_implemented = 2`

The operation is not implemented.

`out_of_memory = 3`

The operation caused an out of memory condition.

`bad_action_code = 4`

`bad_component_type = 5`

The specified component type is not known or otherwise invalid.

`network_error = 6`

A generic network error occurred.

`version_too_new = 7`

The version of the network representation for this object is too new.

`version_too_old = 8`

The version of the network representation for this object is too old.

`version_unknown = 9`

The version of the network representation for this object is unknown.

`unknown_component_address = 10`

`duplicate_component_address = 11`

The given global id has already been registered.

`invalid_status = 12`

The operation was executed in an invalid status.

`bad_parameter = 13`

One of the supplied parameters is invalid.

`internal_server_error = 14`

`service_unavailable = 15`

`bad_request = 16`

`repeated_request = 17`

`lock_error = 18`

```
duplicate_console = 19
    There is more than one console locality.

no_registered_console = 20
    There is no registered console locality available.

startup_timed_out = 21

uninitialized_value = 22

bad_response_type = 23

deadlock = 24

assertion_failure = 25

null_thread_id = 26
    Attempt to invoke a API function from a non-HPX thread.

invalid_data = 27

yield_aborted = 28
    The yield operation was aborted.

dynamic_link_failure = 29

commandline_option_error = 30
    One of the options given on the command line is erroneous.

serialization_error = 31
    There was an error during serialization of this object.

unhandled_exception = 32
    An unhandled exception has been caught.

kernel_error = 33
    The OS kernel reported an error.

broken_task = 34
    The task associated with this future object is not available anymore.

task_moved = 35
    The task associated with this future object has been moved.

task_already_started = 36
    The task associated with this future object has already been started.

future_already_retrieved = 37
    The future object has already been retrieved.

promise_already_satisfied = 38
    The value for this future object has already been set.

future_does_not_support_cancellation = 39
    The future object does not support cancellation.

future_can_not_be_cancelled = 40
    The future can't be canceled at this time.

no_state = 41
    The future object has no valid shared state.

broken_promise = 42
    The promise has been deleted.

thread_resource_error = 43
```

```
future_cancelled = 44
thread_cancelled = 45
thread_not_interruptable = 46
duplicate_component_id = 47
    The component type has already been registered.

unknown_error = 48
    An unknown error occurred.

bad_plugin_type = 49
    The specified plugin type is not known or otherwise invalid.

filesystem_error = 50
    The specified file does not exist or other filesystem related error.

bad_function_call = 51
    equivalent of std::bad_function_call

task_canceled_exception = 52
    parallel::v2::task_canceled_exception

task_block_not_active = 53
    task_region is not active

out_of_range = 54
    Equivalent to std::out_of_range.

length_error = 55
    Equivalent to std::length_error.

migration_needs_retry = 56
    migration failed because of global race, retry
```

```
#include <hpx/errors/exception_info.hpp>
```

Defines

```
HPX_DEFINE_ERROR_INFO(NAME, TYPE)
namespace hpx
```

Functions

```
template<typename E> HPX_NORETURN void hpx::throw_with_info(E && e, exception_info && ei)
template<typename E> HPX_NORETURN void hpx::throw_with_info(E && e, exception_info const & ei)

template<typename E>
exception_info *get_exception_info(E &e)

template<typename E>
exception_info const *get_exception_info(E const &e)

template<typename E, typename F>
auto invoke_with_exception_info(E const &e, F &&f)

template<typename F>
```

```
auto invoke_with_exception_info (std::exception_ptr const &p, F &&f)
template<typename F>
auto invoke_with_exception_info (hpx::error_code const &ec, F &&f)
template<typename Tag, typename Type>
struct error_info
```

Public Types

```
template<>
using tag = Tag
template<>
using type = Type
```

Public Functions

```
error_info (Type const &value)
error_info (Type &&value)
```

Public Members

Type **_value**

```
class exception_info
Subclassed by hpx::detail::exception_with_info_base
```

Public Functions

```
exception_info ()
exception_info (exception_info const &other)
exception_info (exception_info &&other)
exception_info &operator= (exception_info const &other)
exception_info &operator= (exception_info &&other)
virtual ~exception_info ()
template<typename ...ErrorInfo>
exception_info &set (ErrorInfo&&... tagged_values)
template<typename Tag>
Tag::type const *get () const
```

Private Types

```
using node_ptr = std::shared_ptr<detail::exception_info_node_base>
```

Private Members

```
node_ptr _data
```

logging

```
#include <compatibility/hpx/util/logging.hpp>

#include <compatibility/hpx/util/logging/format_fwd.hpp>

#include <compatibility/hpx/util/logging/logging.hpp>

#include <compatibility/hpx/util/logging/format.hpp>

#include <compatibility/hpx/util/logging/writer/named_write.hpp>

#include <compatibility/hpx/util/logging/writer/format_write.hpp>

#include <compatibility/hpx/util/logging/format/named_write.hpp>

#include <compatibility/hpx/util/logging/format/optimize.hpp>

#include <compatibility/hpx/util/logging/format/op_equal.hpp>

#include <compatibility/hpx/util/logging/format/array.hpp>

#include <compatibility/hpx/util/logging/format/named_write_fwd.hpp>

#include <compatibility/hpx/util/logging/format/formatter/high_precision_time.hpp>

#include <compatibility/hpx/util/logging/format/formatter/time.hpp>

#include <compatibility/hpx/util/logging/format/formatter/convert_format.hpp>

#include <compatibility/hpx/util/logging/format/formatter/named_spacer.hpp>

#include <compatibility/hpx/util/logging/format/formatter/spacer.hpp>

#include <compatibility/hpx/util/logging/format/formatter/thread_id.hpp>

#include <compatibility/hpx/util/logging/format/formatter/defaults.hpp>
```

```
#include <compatibility/hpx/util/logging/format/formatter/time_strf.hpp>

#include <compatibility/hpx/util/logging/format/destination/file.hpp>

#include <compatibility/hpx/util/logging/format/destination/
convert_destination.hpp>

#include <compatibility/hpx/util/logging/format/destination/named.hpp>

#include <compatibility/hpx/util/logging/format/destination/defaults.hpp>

#include <hpx/logging.hpp>
```

Defines

```
LAGAS_(lvl)
LPT_(lvl)
LTIM_(lvl)
LPROGRESS_
LHPX_(lvl, cat)
LAPP_(lvl)
LDEB_
LTM_(lvl)
LRT_(lvl)
LOSH_(lvl)
LERR_(lvl)
LLCO_(lvl)
LPCS_(lvl)
LAS_(lvl)
LBT_(lvl)
LFATAL_
LAGAS_CONSOLE_(lvl)
LPT_CONSOLE_(lvl)
LTIM_CONSOLE_(lvl)
LHPX_CONSOLE_(lvl)
LAPP_CONSOLE_(lvl)
LDEB_CONSOLE_
LAGAS_ENABLED(lvl)
LPT_ENABLED(lvl)
LTIM_ENABLED(lvl)
```

```
LHPX_ENABLED (lvl)
LAPP_ENABLED (lvl)
LDEB_ENABLED
```

Functions

```
template<typename T>
bootstrap_logging const &operator<< (bootstrap_logging const &l, T&&)
```

Variables

```
constexpr bootstrap_logging lbt_
struct bootstrap_logging
```

Public Functions

```
constexpr bootstrap_logging()

#include <hpx/logging/format_fwd.hpp>

#include <hpx/logging/logging.hpp>
```

Include this file when you're using the logging lib, but don't necessarily want to use formatters and destinations. If you want to use formatters and destinations, then you can include this one instead:

```
#include <hpx/logging/format_fwd.hpp>
```

```
#include <hpx/logging/format.hpp>
```

Include this file when you're using formatters and destinations, and you want to define the logger classes, in a source file (using HPX_DEFINE_LOG)

```
namespace hpx

namespace util

namespace logging
```

namespace format_and_write

The `format_and_write` classes know how to call the formatter and destination objects.

Usually you'll be happy with the `format_and_write::simple` class - which simply calls `operator()` on the formatters, and `operator()` on the destinations.

Note that usually the formatter and destination class just have an `operator()`, which when called, formats the message or writes it to a destination. In case your formatters/destinations are

more complex than that (for instance, more than a member function needs to be called), you'll have to implement your own `format_and_write` class.

`struct simple`

#include <format.hpp> Formats the message, and writes it to destinations.

- calls `operator()` on the formatters , and `operator()` on the destinations. Ignores `clear_format()` commands.

If you derive from `destination::base`, this type can be `destination::base::raw_param`(see below).

Parameters

- `msg_type`: The message to pass to the formatter. This is the type that is passed to the formatter objects and to the destination objects. Thus, it needs to be convertible to the argument to be sent to the formatter objects and to the argument to be sent to the destination objects. Usually, it's the argument you pass on to your destination classes.

Example:

```
typedef destination::base<const std::string &> dest_base;
// in this case : msg_type = std::string = dest_base::raw_param
struct write_to_cout : dest_base {
    void operator()(param msg) const {
        std::cout << msg ;
    }
};

typedef destination::base<const std::string &> dest_base;
// in this case : msg_type = cache_string = dest_base::raw_param
struct write_to_file : dest_base, ... {
    void operator()(param msg) const {
        context() << msg ;
    }
};
```

Public Functions

```
simple(msg_type &msg)

template<class formatter_ptr>
void format(const formatter_ptr &fmt)

template<class destination_ptr>
void write(const destination_ptr &dest)

void clear_format()
```

Protected Attributes

msg_type &**m_msg**

namespace msg_route

Specifies the route : how formatting and writing to destinations take place.

Classes in this namespace specify when formatters and destinations are to be called.

See [msg_route::simple](#)

template<class **formatter_array**, class **destination_array**>

struct formatter_and_destination_array_holder

#include <format.hpp> Recommended base class for message routers that need access to the underlying formatter and/or destination array.

Protected Functions

```
formatter_and_destination_array_holder(const formatter_array &formats_, const destination_array &destinations_)  
const formatter_array &formats() const  
const destination_array &destinations() const
```

Private Members

const **formatter_array** &**m_formats**

const **destination_array** &**m_destinations**

struct simple

#include <format.hpp> Represents a simple router - first calls all formatters.

- in the order they were added, then all destinations - in the order they were added

Example:

```
typedef logger< format_write > logger_type;  
HPX_DEFINE_LOG_FILTER(g_log_filter, filter::no_ts )  
HPX_DEFINE_LOG(g_l, logger_type)  
define L_ HPX_LOG_USE_LOG_IF_FILTER(g_l(), g_log_filter()>is_  
→enabled() )  
  
// add formatters : [idx] [time] message [enter]  
g_l()>writer().add_formatter( write_idx() );  
g_l()>writer().add_formatter( write_time() );  
g_l()>writer().add_formatter( append_newline() );  
  
// write to cout and file  
g_l()>writer().add_destination( write_to_cout() );  
g_l()>writer().add_destination( write_to_file("out.txt") );  
  
// usage  
int i = 1;  
L_ << "testing " << i << i+1 << i+2;
```

In the above case:

- First, the formatters are called: `write_idx()` is called, then `write_time()`, then `append_newline()`.
- Then, the destinations are called: `write_to_cout()`, and then `write_to_file()`.

Parameters

- `format_base`: The base class for all formatter classes from your application. See manipulator.
- `destination_base`: The base class for all destination classes from your application. See manipulator.

Public Types

```
typedef formatter::base::ptr_type formatter_ptr
typedef destination::base::ptr_type destination_ptr
typedef std::vector<formatter_ptr> f_array
typedef std::vector<destination_ptr> d_array
```

Public Functions

```
template<class formatter_array, class destination_arraysimple(const formatter_array&, const destination_array&)

void append_formatter(formatter_ptr fmt)
void del_formatter(formatter_ptr fmt)
void append_destination(destination_ptr dest)
void del_destination(destination_ptr dest)

template<class format_and_write>
void write(msg_type &msg) const
```

Private Members

```
write_info m_to_write

struct write_info
```

Public Members

```
f_array formats
d_array destinations
```

```
#include <hpx/logging/writer/named_write.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace logging
```

```
            namespace writer
```

```
                struct named_write
```

```
#include <named_write.hpp> Composed of a named formatter and a named destinations.  
Thus, you can specify the formatting and destinations as strings.
```

```
#include <hpx/logging/format/named_write.hpp>
```

Contains a very easy interface for using formatters and destinations:

- at construction, specify 2 params: the formatter string and the destinations string
- Setting the formatters and destinations to write to is extremely simple:

```
// Set the formatters (first param) and destinatins (second step) in  
// one step  
g_l() ->writer().write("%time%($hh:$mm.$ss.$mili) [%idx%] |\n",  
    "cout file(out.txt) debug");  
  
// set the formatter(s)  
g_l() ->writer().format("%time%($hh:$mm.$ss.$mili) [%idx%] |\n");  
  
// set the destination(s)  
g_l() ->writer().destination("cout file(out.txt) debug");
```

Public Functions

```
named_write()
```

```
void format(const std::string &format_str)
```

sets the format string: what should be before, and what after the original message, separated by “|”

Example: “[%idx%] \n” - this writes “[%idx%] ” before the message, and “\n” after the message

If “|” is not present, the whole message is prepended to the message

```
void format(const std::string &format_before_str, const std::string &for-  
mat_after_str)
```

sets the format strings (what should be before, and what after the original message)

```
void destination(const std::string &destination_str)
```

sets the destinations string - where should logged messages be outputted

```
void write(const std::string &format_str, const std::string &destination_str)
```

Specifies the formats and destinations in one step.

```

const std::string &format () const
const std::string &destination () const
void operator () (msg_type &msg) const

template<class destination>
void replace_destination (const std::string &name, destination d)
    Replaces a destination from the named destination.

    You can use this, for instance, when you want to share a destination between multiple named
    writers.

template<class formatter>
void replace_formatter (const std::string &name, formatter d)
    Replaces a formatter from the named formatter.

    You can use this, for instance, when you want to share a formatter between multiple named
    writers.

template<class formatter>
void add_formatter (formatter fmt)

template<class destination>
void add_destination (const std::string &name, destination d)

```

Private Functions

```

template<class manipulator, class parser_type>
void set_and_configure (manipulator &manip, const std::string &name,
                        parser_type parser)

void init ()

```

Private Members

```

formatter::named_spacer_t<formatter::do_convert_format::prepend> m_format_before
formatter::named_spacer_t<formatter::do_convert_format::append> m_format_after
destination::named m_destination
format_write m_writer
std::string m_format_str
std::string m_format_before_str
std::string m_format_after_str
std::string m_destination_str
struct parse_destination

```

Public Functions

```
bool has_manipulator_name() const  
std::string get_manipulator_name() const  
void clear()  
void add(char c)
```

Private Members

```
std::string m_manipulator  
struct parse_formatter
```

Public Functions

```
bool has_manipulator_name() const  
std::string get_manipulator_name() const  
void clear()  
void add(char c)
```

Private Members

```
std::string m_manipulator
```

```
#include <hpx/logging/writer/format_write.hpp>
```

```
#include <hpx/logging/format/named_write.hpp>
```

```
#include <hpx/logging/format/optimize.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace logging
```

```
namespace optimize
```

Gathering the message: contains optimizers for formatting and/or destinations: for example, caching techniques.

Functions

```
template<class stream>
stream &operator<< (stream &out, const cache_string_one_str &val)
```

struct cache_string_one_str
#include <optimize.hpp> Optimizes the formatting for prepending and/or appending strings to the original message.

It keeps all the modified message in one string. Useful if some formatter needs to access the whole string at once.

reserve_prepend() - the size that is reserved for prepending (similar to `string::reserve` function)
reserve_append() - the size that is reserved for appending (similar to `string::reserve` function)

Note : as strings are prepended, `reserve_prepend()` shrinks. Same goes for append.

Public Types

```
typedef cache_string_one_str self_type
```

Public Functions

```
cache_string_one_str (std::size_t reserve_prepend_, std::size_t reserve_append_,  

std::size_t grow_size_ = 10)
```

Parameters

- `reserve_prepend`: - how many chars to have space to prepend by default
- `reserve_append`: - how many chars to have space to append by default
- `grow_size`: - in case we add a string and there's no room for it, with how much should we grow? We'll grow this much in addition to the added string
 - in the needed direction

```
cache_string_one_str (const std::string &msg, std::size_t reserve_prepend_ = 10,  

std::size_t reserve_append_ = 10, std::size_t grow_size_ =  

10)
```

Parameters

- `msg`: - the message that is originally cached
- `reserve_prepend`: - how many chars to have space to prepend by default
- `reserve_append`: - how many chars to have space to append by default
- `grow_size`: - in case we add a string and there's no room for it, with how much should we grow? We'll grow this much in addition to the added string
 - in the needed direction

```
cache_string_one_str (cache_string_one_str &&other)
```

```
cache_string_one_str ()
```

```
void set_string (const std::string &str)
```

```
std::size_t reserve_prepend () const
```

```
std::size_t reserve_append () const
```

```
std::size_t grow_size () const
```

```
void reserve_prepend (std::size_t new_size)
```

```
void reserve_append(std::size_t new_size)
void grow_size(std::size_t new_size)
void prepend_string(const char *str)
void append_string(const char *str)
void prepend_string(const std::string &str)
    pre-pends a string (inserts it at the beginning)
void append_string(const std::string &str)
    appends a string (inserts it at the end)

template<class stream_type>
void to_stream(stream_type &stream) const
    writes the current cached contents to a stream

const std::string &full_string() const
    returns the full string

operator const std::string&() const
```

Private Functions

```
void resize_string(std::size_t reserve_prepend_, std::size_t reserve_append_)
bool is_string_set() const
```

Private Members

```
std::size_t m_reserve_prepend
std::size_t m_reserve_append
std::size_t m_grow_size
std::string m_str
bool m_full_msg_computed
std::string m_full_msg
```

Private Static Functions

```
static std::size_t str_len(const char *str)
static std::size_t str_len(const wchar_t *str)
```

```
#include <hpx/logging/format/op_equal.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace logging
```

```
            namespace op_equal
```

Implements operator== for manipulators.

Functions

```
bool operator==(const same_type_op_equal_top &a, const same_type_op_equal_top
&b)
```

```
template<class type>
```

```
struct same_type_op_equal : public hpx::util::logging::op_equal::same_type_op_equal_base
#include <op_equal.hpp> Implements operator==, which compares two objects. If they have
the same type, it will compare them using the type's member operator==.
```

The only constraint is that operator== must be a *member* function

Public Functions

```
virtual bool equals(const same_type_op_equal_top &other) const
```

```
struct same_type_op_equal_base : public virtual hpx::util::logging::op_equal::same_type_op_eq
#include <op_equal.hpp> Base class when you want to implement operator== that will com-
pare based on type and member operator==.
```

See [same_type_op_equal](#)

Subclassed by [hpx::util::logging::op_equal::same_type_op_equal< type >](#)

```
struct same_type_op_equal_top
```

Subclassed by [hpx::util::logging::op_equal::same_type_op_equal_base](#)

Public Functions

```
virtual bool equals(const same_type_op_equal_top&) const = 0
```

Protected Functions

```
same_type_op_equal_top()

virtual ~same_type_op_equal_top()

same_type_op_equal_top(const same_type_op_equal_top&)

#include <hpx/logging/format/array.hpp>

namespace hpx

namespace util

namespace logging

namespace array

template<class base_type>
class ptr_holder
    #include <array.hpp> Holds an array of manipulators (formatters or destinations). It owns them, holding them internally as smart pointers Each function call is locked.

The base_type must implement operator==

When you call get_ptr() or del(), the type you provide, must implement operator==(const type&, const base_type&)
```

Public Types

```
typedef base_type value_type
typedef std::unique_ptr<value_type> ptr_type
typedef std::vector<ptr_type> array_type
```

Public Functions

```
template<class derived>
base_type *append(derived val)

template<class derived>
base_type *get_ptr(derived val) const

template<class derived>
void del(derived val)

void del(base_type *p)
```

Private Members

array_type **m_array**

```
#include <hpx/logging/format/named_write_fwd.hpp>

#include <hpx/logging/format/formatter/high_precision_time.hpp>

namespace hpx

namespace util

namespace logging

namespace formatter
```

TypeDefs

typedef `high_precision_time_t` **high_precision_time**
high_precision_time_t with default values. See *high_precision_time_t*

Prefixes the message with a high-precision time (. You pass the format string at construction.

```
#include <hpx/logging/format/formatter/high_precision_time.hpp>
```

Internally, it uses `hpx::util::date_time::microsec_time_clock`. So, our precision matches this class.

The format can contain escape sequences: \$dd - day, 2 digits \$MM - month, 2 digits \$yy - year, 2 digits \$yyyy - year, 4 digits \$hh - hour, 2 digits \$mm - minute, 2 digits \$ss - second, 2 digits \$mili - milliseconds \$micro - microseconds (if the high precision clock allows; otherwise, it pads zeros) \$nano - nanoseconds (if the high precision clock allows; otherwise, it pads zeros)

Example:

```
high_precision_time(" $mm:$ss:$micro");
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See `convert_format`. For instance, you might use *a cached_string class* (see *optimize namespace*).

```
template<class convert = do_convert_format::prepend>
struct high_precision_time_t : public is_generic, public non_const_context<hpx::util::logging::d
```

#include <high_precision_time.hpp> Prefixes the message with a high-precision time (. You pass the format string at construction.

```
#include <hpx/logging/format/formatter/high_precision_time.hpp>
```

Internally, it uses `hpx::util::date_time::microsec_time_clock`. So, our precision matches this class.

The format can contain escape sequences: \$dd - day, 2 digits \$MM - month, 2 digits \$yy - year, 2 digits \$yyyy - year, 4 digits \$hh - hour, 2 digits \$mm - minute, 2 digits \$ss - second, 2 digits \$mili - milliseconds \$micro - microseconds (if the high precision clock allows; otherwise, it pads zeros) \$nano - nanoseconds (if the high precision clock allows; otherwise, it pads zeros)

Example:

```
high_precision_time("$mm:$ss:$micro");
```

Parameters

- `convert`: [optional] In case there needs to be a conversion between `std::(w)string` and the string that holds your logged message. See `convert_format`. For instance, you might use *a cached_string class* (see *optimize namespace*).

Public Types

```
typedef convert convert_type
```

```
typedef non_const_context<hpx::util::logging::detail::time_format_holder> non_const_context_base
```

Public Functions

```
high_precision_time_t(const std::string &format)
```

constructs a `high_precision_time` object

```
void write_high_precision_time(msg_type &msg,  
                               std::chrono::time_point<std::chrono::system_clock>  
                               val) const
```

```
void operator()(msg_type &msg) const
```

```
bool operator==(const high_precision_time_t &other) const
```

```
void configure(const std::string &str)
```

configure through script

the string = the time format

```
#include <hpx/logging/format/formatter/time.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace logging
```

```
            namespace formatter
```

TypeDefs

typedef time_t **time**
time_t with default values. See [time_t](#)

Prefixes the message with the time. You pass the format string at construction. It's friendlier than write_time_strf (which uses strftime).

The format can contain escape sequences: \$dd - day, 2 digits \$MM - month, 2 digits \$yy - year, 2 digits \$yyyy - year, 4 digits \$hh - hour, 2 digits \$mm - minute, 2 digits \$ss - second, 2 digits

Example: time("Today is \$dd/\$MM/\$yyyy");

Note: for a high precision clock, try high_precision_time (uses hpx::util::date_time)

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

template<class **convert** = [do_convert_format::prepend](#)>

struct **time_t** : **public** [is_generic](#), **public** [non_const_context<hpx::util::logging::detail::time_format_holder>](#)
#include <time.hpp> Prefixes the message with the time. You pass the format string at construction.

It's friendlier than write_time_strf (which uses strftime).

The format can contain escape sequences: \$dd - day, 2 digits \$MM - month, 2 digits \$yy - year, 2 digits \$yyyy - year, 4 digits \$hh - hour, 2 digits \$mm - minute, 2 digits \$ss - second, 2 digits

Example: time("Today is \$dd/\$MM/\$yyyy");

Note: for a high precision clock, try high_precision_time (uses hpx::util::date_time)

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

Public Types

typedef *convert* **convert_type**

typedef [non_const_context<hpx::util::logging::detail::time_format_holder>](#) **non_const_context_base**

Public Functions

time_t (**const** *std*::string &*format*)
constructs a time object

void **write_time** (*msg_type* &*msg*, **time_t** *val*) **const**

void **operator()** (*msg_type* &*msg*) **const**

bool **operator==** (**const** **time_t** &*other*) **const**

void **configure** (**const** *std*::string &*str*)
configure through script

the string = the time format

```
#include <hpx/logging/format/formatter/convert_format.hpp>

namespace hpx

    namespace util

        namespace logging

            namespace formatter

                struct do_convert_format

                    struct append
```

Public Static Functions

```
template<class string>
static const std::string &get_underlying_string(const string &str)

template<class string>
static void write(const char *src, string &dest)

template<class src_type, class string>
static void write(const src_type &src, string &dest)

template<class src_type, class string>
static void write(src_type &src, string &dest)

struct prepend
```

Public Static Functions

```
template<class string>
static const std::string &get_underlying_string(const string &str)

template<class string>
static void write(const char *src, string &dest)

template<class src_type, class string>
static void write(const src_type &src, string &dest)

template<class src_type, class string>
static void write(src_type &src, string &dest)
```

namespace convert

Allows format conversions.

- In case you're using a formatter that does not match your string type

In case you want to use a formatter developed by someone else (for instance, a formatter provided by this lib), perhaps you're using another type of string to hold the message

- thus, you need to provide a conversion function

Example: FIXME

> convert_format::prepend

explain that you can extend the following - since they're namespaces!!! so that you can "inject" your own write function in the convert_format::prepend/orwhatever namespace, and then it'll be automatically used!

namespace append

Functions

```
void write (const std::string &src, std::string &dest)
```

```
void write (const std::string &src, hpx::util::logging::optimize::cache_string_one_str  
&dest)
```

```
void write (const char *src, std::string &dest)
```

```
void write (const char *src, hpx::util::logging::optimize::cache_string_one_str  
&dest)
```

namespace prepend

Example : write_time

Functions

```
void write (const char *src, std::string &dest)
```

```
void write (const std::string &src, std::string &dest)
```

```
void write (const std::string &src, hpx::util::logging::optimize::cache_string_one_str  
&dest)
```

```
void write (const char *src, hpx::util::logging::optimize::cache_string_one_str  
&dest)
```

```
#include <hpx/logging/format/formatter/named_spacer.hpp>
```

namespace hpx

namespace util

namespace logging

namespace formatter

```
template<class convert>
```

```
struct named_spacer_t : public is_generic, public non_const_context<detail::named_spacer_context>
```

```
#include <named_spacer.hpp> Allows you to contain multiple formatters, and specify a spacer  
between them. You have a spacer string, and within it, you can escape your contained formatters.
```

```
#include <hpx/logging/format/formatter/named_spacer.hpp>
```

This allows you:

- to hold multiple formatters
- each formatter is given a name, when being added
- you have a spacer string, which contains what is to be prepended or appended to the string (by default, prepended)
- a formatter is escaped with '%' chars, like this "%name%"
- if you want to write the '%', just double it, like this: "this %% gets written"

Example:

```
#define L_ HPX_LOG_USE_LOG_IF_FILTER(g_l(), g_log_filter()->is_
→enabled() )

g_l() ->writer().add_formatter( formatter::named_spacer("[%index%]
→%time% (T%thread%) ")
    .add( "index", formatter::idx())
    .add( "thread", formatter::thread_id())
    .add( "time", formatter::time("$mm")) );
```

Assuming you'd use the above in code

```
int i = 1;
L_ << "this is so cool " << i++;
L_ << "this is so cool again " << i++;
```

You could have an output like this:

```
[1] 53 (T3536) this is so cool 1
[2] 54 (T3536) this is so cool again 2
```

Public Types

```
typedef non_const_context<detail::named_spacer_context<convert>> context_base
```

Public Functions

```
named_spacer_t (const std::string &str = std::string())

named_spacer_t &string (const std::string &str)

template<class formatter>
named_spacer_t &add (const std::string &name, formatter fmt)

void del (const std::string &name)

void configure_inner (const std::string &name, const std::string &configure_str)

void operator() (msg_type &msg) const

bool operator== (const named_spacer_t &other) const
```

```
#include <hpx/logging/format/formatter/spacer.hpp>

namespace hpx

namespace util

namespace logging

namespace formatter
```

Functions

```
template<class original_formatter>
detail::find_spacer<original_formatter>::type spacer(const original_formatter &fmt,
                                                       const char *format_str)
```

Prepends some info, and appends some info to an existing formatter.

The syntax is simple: construct a spacer by passing the original formatter, and the text to space (prepend and append). Use:

- % to mean the original formatter text
- anything before "%" is prepended before
- anything after "%" is appended after

Examples:

```
// prefix "[" before index, and append "] " after it.
formatter::spacer( formatter::idx(), "[%] ");

// prefix "{T" before thread_id, and append "} " after it
formatter::spacer( formatter::thread_id(), "{T%} ");
```

When adding a spacer formatter, you'll do something similar to:

```
g_l()->writer().add_formatter( formatter::spacer( formatter::idx(),
    →"[%] " ) );
```

However, to make this even simpler, I allow an ever easier syntax:

```
// equivalent to the above
g_l()->writer().add_formatter( formatter::idx(), "[%] " );
```

```
template<class convert, class original_formatter, bool is_generic_formatter>
struct spacer_t : public original_formatter
#include <spacer.hpp>
```

Prepends some info, and appends some info to an existing formatter.

The syntax is simple: construct a spacer by passing the original formatter, and the text to space (prepend and append). Use:

- % to mean the original formatter text
- anything before "%" is prepended before
- anything after "%" is appended after

Examples:

```
// prefix "[" before index, and append "] " after it.
formatter::spacer( formatter::idx(), "[%] ");

// prefix "{T" before thread_id, and append "} " after it
formatter::spacer( formatter::thread_id(), "{T%} ");
```

When adding a spacer formatter, you'll do something similar to:

```
g_l()->writer().add_formatter( formatter::spacer( formatter::idx(),
    →"[%] " ) );
```

However, to make this even simpler, I allow an ever easier syntax:

```
// equivalent to the above
g_l()->writer().add_formatter( formatter::idx(), "[%] " );
```

Public Types

```
typedef original_formatter::param param
typedef original_formatter spacer_base
```

Public Functions

```
spacer_t (const original_formatter &fmt, const char *format_str)
void operator() (param msg) const
```

Private Functions

```
void parse_format (const std::string &format_str)
```

Private Members

```
std::string m_prefix
std::string m_suffix

template<class convert, class original_formatter>
struct spacer_t<convert, original_formatter, true> : public original_formatter
```

Public Types

```
typedef original_formatter spacer_base
```

Public Functions

```
spacer_t (const original_formatter &fmt, const char *format_str)
void operator() (msg_type &msg) const
```

Private Functions

```
void parse_format (const std::string &format_str)
```

Private Members

```
std::string m_prefix
std::string m_suffix
```

```
#include <hpx/logging/format/formatter/thread_id.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace logging
```

```
namespace formatter
```

Typedefs

```
typedef thread_id_t thread_id
thread_id_t with default values. See thread\_id\_t
```

Writes the thread_id to the log.

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use a [cached_string class](#) (see [optimize namespace](#)).

```
template<class convert = do\_convert\_format::prepend>
struct thread_id_t : public is\_generic
#include <thread\_id.hpp> Writes the thread_id to the log.
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use a [cached_string class](#) (see [optimize namespace](#)).

Public Types

```
typedef convert convert_type
```

Public Functions

```
void operator() (msg_type &msg) const
```

```
bool operator==(const thread_id_t&) const
```

```
#include <hpx/logging/format/formatter/defaults.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace logging
```

```
namespace formatter
```

Typedefs

```
typedef idx_t idx
```

idx_t with default values. See *idx_t*

prefixes each message with an index. Example:

```
L_ << "my message";
L_ << "my 2nd message";
```

This will output something similar to:

```
[1] my message
[2] my 2nd message
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use a *cached_string class* (see *optimize namespace*).

```
typedef append_newline_t append_newline
```

append_newline_t with default values. See *append_newline_t*

Appends a new line.

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use a *cached_string class* (see *optimize namespace*).

```
typedef append_newline_if_needed_t append_newline_if_needed
```

append_newline_if_needed_t with default values. See *append_newline_if_needed_t*

Appends a new line, if not already there.

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

```
template<class convert = do_convert_format::append>
struct append_newline_if_needed_t : public is_generic
#include <defaults.hpp> Appends a new line, if not already there.
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

Public Types

```
typedef convert convert_type
```

Public Functions

```
void operator() (msg_type &str) const
bool operator==(const append_newline_if_needed_t&) const

template<class convert = do_convert_format::append>
struct append_newline_t : public is_generic
#include <defaults.hpp> Appends a new line.
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

Public Types

```
typedef convert convert_type
```

Public Functions

```
void operator() (msg_type &str) const
bool operator==(const append_newline_t&) const

template<class convert = do_convert_format::prepend>
struct idx_t : public is_generic, public formatter::non_const_context<std::uint64_t>
#include <defaults.hpp> prefixes each message with an index.
```

Example:

```
L_ << "my message";
L_ << "my 2nd message";
```

This will output something similar to:

```
[1] my message
[2] my 2nd message
```

Parameters

- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

Public Types

```
typedef formatter::non_const_context<std::uint64_t> non_const_context_base
typedef convert convert_type
```

Public Functions

```
idx_t()

void operator() (msg_type &str) const
bool operator==(const idx_t&) const
```

```
#include <hpx/logging/format/formatter/time_strf.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace logging
```

```
namespace formatter
```

Typedefs

```
typedef time_strf_t time_strf
time_strf_t with default values. See time\_strf\_t
```

Prefixes the message with the time, by using strftime function. You pass the format string at construction.

Parameters

- msg_type: The type that holds your logged message.
- convert: [optional] In case there needs to be a conversion between std::(w)string and the string that holds your logged message. See convert_format. For instance, you might use [a cached_string class](#) (see [optimize namespace](#)).

```
template<class convert = do_convert_format::prepend>
struct time_strf_t : public is_generic
```

`#include <time_strf.hpp>` Prefixes the message with the time, by using strftime function. You pass the format string at construction.

Parameters

- `msg_type`: The type that holds your logged message.
- `convert`: [optional] In case there needs to be a conversion between `std::(w)string` and the string that holds your logged message. See `convert_format`. For instance, you might use *a cached_string class* (see `optimize namespace`).

Public Types

```
typedef convert convert_type
```

Public Functions

```
time_strf_t (const std::string &format, bool localtime)  
constructs a time_strf object
```

Parameters

- `format`: the time format , strftime-like
- `localtime`: if true, use localtime, otherwise global time

```
void operator() (msg_type &msg) const
```

```
bool operator==(const time_strf_t &other) const
```

Private Members

```
std::string m_format
```

```
bool m_localtime
```

```
#include <hpx/logging/format/destination/file.hpp>  
  
namespace hpx  
  
namespace util  
  
namespace logging  
  
namespace destination  
  
struct file : public is_generic, public non_const_context<detail::file_info>  
#include <file.hpp> Writes the string to a file.
```

Public Types

```
typedef non_const_context<detail::file_info> non_const_context_base
typedef boost::detail::spinlock mutex_type
```

Public Functions

```
file (const std::string &file_name, file_settings set = file_settings())
constructs the file destination
```

Parameters

- file_name: name of the file
- set: [optional] file settings - see *file_settings* class, and dealing_with_flags

```
void operator() (const msg_type &msg) const
```

```
bool operator== (const file &other) const
```

```
void configure (const std::string &str)
```

configure through script right now, you can only specify the file name

Public Static Attributes

```
mutex_type mtx_
```

```
struct file_settings
```

#include <file.hpp> settings for when constructing a file class. To see how it's used, see dealing_with_flags.

Public Types

```
typedef hpx::util::logging::detail::flag<file_settings> flag
```

Public Functions

```
file_settings ()
```

Public Members

```
flag::t<bool> flush_each_time
if true (default), flushes after each write
```

```
flag::t<bool> initial_overwrite
```

```
flag::t<bool> do_append
```

```
flag::t<std::ios_base::openmode> extra_flags
just in case you have some extra flags to pass, when opening the file
```

```
#include <hpx/logging/format/destination/convert_destination.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace logging
```

```
            namespace destination
```

```
                namespace convert
```

Allows writing messages to destinations.

It has 2 function overloads:

- `write(message, output)` - writes the given message, to the given output
- `do_convert(message, into<other_type>()`);

`FIXME`

Functions

```
template<class obj>
```

```
void write(const obj &m, std::ostream &out)
```

```
void write(const char *m, std::ostream &out)
```

```
const char *do_convert(const char *c, const into<const char *>&)
```

```
const char *do_convert(const std::string &s, const into<const char *>&)
```

```
const std::string &do_convert(const std::string &s, const into<std::string>&)
```

```
#include <hpx/logging/format/destination/named.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace logging
```

```
            namespace destination
```

```
                struct named : public is_generic, public non_const_context<detail::named_context>
                    #include <named.hpp> Allows you to contain multiple destinations, give each such destination a name. Then, at run-time, you can specify a format string which will specify which destinations to be called, and on what order.
```

This allows you:

- to hold multiple destinations

- each destination is given a name, when being added. The name **must not** contain spaces and must not start with '+'/-` signs
- you have a format string, which contains what destinations to be called, and on which order The format string contains destination names, separated by space.

When a message is written to this destination, I parse the format string. When a name is encountered, if there's a destination corresponding to this name, I will call it.

Example:

```
g_1() -> writer().add_destination(
    destination::named("cout out debug")
        .add("cout", destination::cout())
        .add("debug", destination::dbg_window())
        .add("out", destination::file("out.txt"))
);
```

In the above code, we'll write to 3 destinations, in the following order:

- first, to the console
- second, to the out.txt file
- third, to the debug window

Public Types

typedef non_const_context<detail::named_context> **non_const_context_base**

Public Functions

named (**const std::string &format_string = std::string()**)
constructs the named destination

Parameters

- named_name: name of the named
- set: [optional] named settings - see named_settings class, and dealing_with_flags

void **operator()** (**const msg_type &msg**) **const**

named &string (**const std::string &str**)

template<class **destination**>

named &add (**const std::string &name, destination dest**)

void **del** (**const std::string &name**)

void **configure_inner** (**const std::string &name, const std::string &configure_str**)

bool **operator==** (**const named &other**) **const**

```
#include <hpx/logging/format/destination/defaults.hpp>

namespace hpx

namespace util

namespace logging

namespace destination

struct cerr : public is_generic
#include <defaults.hpp> Writes the string to cerr.
```

Public Functions

```
void operator() (const msg_type &msg) const
bool operator== (const cerr&) const

struct cout : public is_generic
#include <defaults.hpp> Writes the string to console.
```

Public Functions

```
void operator() (const msg_type &msg) const
bool operator== (const cout&) const

struct dbg_window : public is_generic
#include <defaults.hpp> Writes the string to output debug window.

For non-Windows systems, this is the console.
```

Public Functions

```
void operator() (const msg_type &msg) const
bool operator== (const dbg_window&) const

struct stream : public is_generic, public non_const_context<std::ostream *>
#include <defaults.hpp> writes to stream.

Note : The stream must outlive this object! Or, clear() the stream, before the stream is deleted.
```

Public Types

```
typedef std::ostream stream_type  
typedef non_const_context<stream_type *> non_const_context_base
```

Public Functions

```
stream(stream_type *s)  
stream(stream_type &s)  
void operator()(const msg_type &msg) const  
bool operator==(const stream &other) const  
void set_stream(stream_type *p)  
    resets the stream. Further output will be written to this stream  
void clear()  
    clears the stream. Further output will be ignored
```

iterator_support

```
#include <compatibility/hpx/util/iterator_facade.hpp>  
  
#include <compatibility/hpx/util/transform_iterator.hpp>  
  
#include <compatibility/hpx/util/range.hpp>  
  
#include <compatibility/hpx/util/iterator_range.hpp>  
  
#include <compatibility/hpx/util/iterator_adaptor.hpp>  
  
#include <compatibility/hpx/traits/is_iterator.hpp>  
  
#include <compatibility/hpx/traits/is_range.hpp>  
  
#include <hpx/iterator_support/iterator_facade.hpp>
```

Defines

```
HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(prefix, op, result_type)  
namespace hpx  
  
namespace util
```

Functions

```
template<typename Derived, typename T, typename Category, typename Reference, typename Distance>
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline, bool)
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline, !, bool)
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline)
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline, <=, bool)
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline, >=, bool)
hpx::util::HPX_UTIL_ITERATOR_FACADE_INTEROP_HEAD(inline, -, typename std::iterator<Derived, T, Category, Reference, Distance>)
template<typename Derived, typename T, typename Category, typename Reference, typename Distance>
template<typename Derived, typename T, typename Category, typename Reference, typename Distance>
class iterator_core_access
```

Public Static Functions

```
template<typename Iterator1, typename Iterator2>static HPX_HOST_DEVICE bool hpx::util::operator<(Iterator1, Iterator2)
template<typename Iterator>static HPX_HOST_DEVICE void hpx::util::operator=(Iterator, Iterator)
template<typename Iterator>static HPX_HOST_DEVICE void hpx::util::operator++(Iterator)
template<typename Reference, typename Iterator>static HPX_HOST_DEVICE Reference hpx::util::operator*(Iterator, Reference)
template<typename Iterator, typename Distance>static HPX_HOST_DEVICE void hpx::util::operator+=(Iterator, Distance)
template<typename Iterator1, typename Iterator2>static HPX_HOST_DEVICE std::iterator<Derived, T, Category, Reference, Distance> hpx::util::operator+(Iterator1, Iterator2)
template<typename Derived, typename T, typename Category, typename Reference = T&, typename Distance = std::ptrdiff_t>
struct iterator_facade : public hpx::util::detail::iterator_facade_base<Derived, T, Category, Reference, Distance>
    Subclassed by hpx::util::iterator\_adaptor< Derived, Base, Value, Category, Reference, Difference, Pointer >
```

Public Functions

HPX_HOST_DEVICE **iterator_facade()**

Protected Types

typedef iterator_facade<Derived, T, Category, Reference, Distance, Pointer> iterator_adaptor_

Private Types

```
typedef detail::iterator_facade_base<Derived, T, Category, Reference, Distance, Pointer> base_type

#include <hpx/iterator_support/transform_iterator.hpp>

namespace hpx

namespace util
```

Functions

```
template<typename Transformer, typename Iterator>
transform_iterator<Iterator, Transformer> make_transform_iterator(Iterator const &it,
                                                               Transformer const
                                                               &f)

template<typename Transformer, typename Iterator>
transform_iterator<Iterator, Transformer> make_transform_iterator(Iterator const &it)

template<typename Iterator, typename Transformer, typename Reference, typename Value, typename Category>
class transform_iterator : public hpx::util::iterator_adaptor<transform_iterator<Iterator, Transformer, Reference, Value, Category>>
```

Public Functions

```
transform_iterator()

transform_iterator(Iterator const &it)

transform_iterator(Iterator const &it, Transformer const &f)

template<typename OtherIterator, typename OtherTransformer, typename OtherReference, typename OtherValue, typename OtherCategory, typename OtherDifference>
transform_iterator(transform_iterator<OtherIterator, OtherTransformer, OtherReference, OtherValue, OtherCategory, OtherDifference> const &t,
                  std::enable_if<std::is_convertible<OtherIterator, Iterator>::value && std::is_convertible<OtherTransformer, Transformer>::value && std::is_convertible<OtherCategory, Category>::value && std::is_convertible<OtherDifference, Difference>::value * = nullptr)

Transformer const &transformer() const
```

Private Types

```
typedef detail::transform_iterator_base<Iterator, Transformer, Reference, Value, Category, Difference>::type base
```

Private Functions

```
base_type::reference dereference() const
```

Private Members

```
Transformer transformer_
```

Friends

```
friend hpx::util::hpx::util::iterator_core_access
```

```
#include <hpx/iterator_support/range.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace range_adl
```

Functions

```
template<typename C, typename Iterator = typename detail::iterator<C>::type>HPX_
template<typename C, typename Iterator = typename detail::iterator<C const>::type>HPX_
template<typename C, typename Sentinel = typename detail::sentinel<C>::type>HPX_
template<typename C, typename Sentinel = typename detail::sentinel<C const>::type>HPX_
template<typename C, typename Iterator = typename detail::iterator<C const>::type>HPX_
template<typename C, typename Iterator = typename detail::iterator<C const>::type>HPX_
```

```
#include <hpx/iterator_support/iterator_range.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
template<typename Range, typename Iterator = typename traits::range_iterator<Range>::type, typename Sentinel = typename traits::range_end_iterator<Range>::type> make_iterator_range(Range &r)
template<typename Range, typename Iterator = typename traits::range_iterator<Range const>::type, typename Sentinel = typename traits::range_end_iterator<Range const>::type> make_iterator_range(Range const &r)
template<typename Iterator, typename Sentinel = Iterator>
std::enable_if<traits::is_iterator<Iterator>::value, iterator_range<Iterator, Sentinel>>::type make_iterator_range(Iterator &i, Sentinel &s)
template<typename Iterator, typename Sentinel = Iterator>
class iterator_range
```

Public Functions

```
iterator_range()
iterator_range(Iterator iterator, Sentinel sentinel)
Iterator begin() const
Iterator end() const
std::ptrdiff_t size() const
bool empty() const
```

Private Members

```
Iterator _iterator
Sentinel _sentinel
```

```
#include <hpx/iterator_support/zip_iterator.hpp>

template<typename F, typename ...Ts>
struct lift_zipped_iterators<F, util::zip_iterator<Ts...>>
```

Public Types

```
typedef util::zip_iterator<Ts...>::iterator_tuple_type tuple_type
typedef util::tuple<typename element_result_of<typename F::template apply<Ts>, Ts>::type...> result_type
```

Public Static Functions

```
template<std::size_t... Is, typename ...Ts_>
static result_type call (util::index_pack<Is...>, util::tuple<Ts_...> const &t)
```

```
template<typename ...Ts_>
static result_type call (util::zip_iterator<Ts_...> const &iter)
```

```
namespace hpx
```

```
namespace traits
```

```
namespace functional
```

```
template<typename F, typename ...Ts>
struct lift_zipped_iterators<F, util::zip_iterator<Ts...>>
```

Public Types

```
typedef util::zip_iterator<Ts...>::iterator_tuple_type tuple_type
```

```
typedef util::tuple<typename element_result_of<typename F::template apply<Ts>, Ts>::type...> result_type
```

Public Static Functions

```
template<std::size_t... Is, typename ...Ts_>
static result_type call (util::index_pack<Is...>, util::tuple<Ts_...> const &t)
```

```
template<typename ...Ts_>
static result_type call (util::zip_iterator<Ts_...> const &iter)
```

```
namespace util
```

Functions

```
template<typename... Ts> HPX_HOST_DEVICE zip_iterator<typename decay<Ts>::type...> hpx::util::zip_iterator<decay<Ts>::type...>
```

```
template<typename ...Ts>
class zip_iterator : public hpx::util::detail::zip_iterator_base<tuple<Ts...>, zip_iterator<Ts...>>
```

Public Functions

```
HPX_HOST_DEVICE zip_iterator()  
HPX_HOST_DEVICE zip_iterator(Ts const&... vs)  
HPX_HOST_DEVICE zip_iterator(tuple<Ts...> &&t)  
HPX_HOST_DEVICE zip_iterator(zip_iterator const &other)  
HPX_HOST_DEVICE zip_iterator(zip_iterator &&other)  
HPX_HOST_DEVICE zip_iterator& hpx::util::zip_iterator::operator=(zip_iterator &other) const  
HPX_HOST_DEVICE zip_iterator& hpx::util::zip_iterator::operator=(zip_iterator &&other) const  
template<typename... Ts>HPX_HOST_DEVICE std::enable_if< std::is_assignable<type> < type> > zip_iterator::operator=(zip_iterator &other) const  
template<typename... Ts>HPX_HOST_DEVICE std::enable_if< std::is_assignable<type> < type> > zip_iterator::operator=(zip_iterator &&other) const
```

Private Types

```
typedef detail::zip_iterator_base<tuple<Ts...>, zip_iterator<Ts...>> base_type  
template<typename ...Ts>  
class zip_iterator<tuple<Ts...>> : public hpx::util::detail::zip_iterator_base<tuple<Ts...>, zip_iterator<tuple<Ts...>>
```

Public Functions

```
HPX_HOST_DEVICE zip_iterator()  
HPX_HOST_DEVICE zip_iterator(Ts const&... vs)  
HPX_HOST_DEVICE zip_iterator(tuple<Ts...> &&t)  
HPX_HOST_DEVICE zip_iterator(zip_iterator const &other)  
HPX_HOST_DEVICE zip_iterator(zip_iterator &&other)  
HPX_HOST_DEVICE zip_iterator& hpx::util::zip_iterator::operator=(zip_iterator &other) const  
HPX_HOST_DEVICE zip_iterator& hpx::util::zip_iterator::operator=(zip_iterator &&other) const  
template<typename... Ts>HPX_HOST_DEVICE std::enable_if< std::is_assignable<type> < type> > zip_iterator::operator=(zip_iterator &other) const  
template<typename... Ts>HPX_HOST_DEVICE std::enable_if< std::is_assignable<type> < type> > zip_iterator::operator=(zip_iterator &&other) const
```

Private Types

```
typedef detail::zip_iterator_base<tuple<Ts...>, zip_iterator<tuple<Ts...>>> base_type
```

```
#include <hpx/iterator_support/iterator_adaptor.hpp>

namespace hpx

namespace util

template<typename Derived, typename Base, typename Value = void, typename Category = void, typename Reference = void>
class iterator_adaptor : public hpx::util::iterator_facade<Derived, value_type, iterator_category, reference_type>
    Subclassed by hpx::util::transform_iterator< Iterator, Transformer, Reference, Value, Category, Difference >
```

Public Types

typedef Base **base_type**

Public Functions

```
HPX_HOST_DEVICE iterator_adaptor()
HPX_HOST_DEVICE iterator_adaptor(Base const &iter)
HPX_HOST_DEVICE Base const& hpx::util::iterator_adaptor::base() const
```

Protected Types

typedef *hpx::util::detail::iterator_adaptor_base<Derived, Base, Value, Category, Reference, Difference, Pointer>::type*

typedef iterator_adaptor<Derived, Base, Value, Category, Reference, Difference, Pointer> **iterator_adaptor**

Protected Functions

```
HPX_HOST_DEVICE Base const& hpx::util::iterator_adaptor::base_reference() const
HPX_HOST_DEVICE Base& hpx::util::iterator_adaptor::base_reference()
```

Private Functions

```
HPX_HOST_DEVICE base_adaptor_type::reference hpx::util::iterator_adaptor::dereference()
template<typename OtherDerived, typename OtherIterator, typename V, typename C,
template<typename DifferenceType>HPX_HOST_DEVICE void hpx::util::iterator_adaptor::difference()
HPX_HOST_DEVICE void hpx::util::iterator_adaptor::increment()
template<typename Iterator = Base, typename Enable = typename std::enable_if<
template<typename OtherDerived, typename OtherIterator, typename V, typename C,
```

Private Members

Base `iterator_`

Friends

`friend hpx::util::hpx::util::iterator_core_access`

```
#include <hpx/iterator_support/traits/is_iterator.hpp>
```

```
#include <hpx/iterator_support/traits/is_range.hpp>
```

```
namespace hpx
```

```
namespace traits
```

```
template<typename R>
struct range_traits<R, true> : public std::iterator_traits<util::detail::iterator<R>::type>
```

Public Types

`typedef util::detail::iterator<R>::type iterator_type`

`typedef util::detail::sentinel<R>::type sentinel_type`

`compute`

```
#include <hpx/compute/vector.hpp>
```

```
namespace hpx
```

```
namespace compute
```

Functions

```
template<typename T, typename Allocator>
```

```
void swap (vector<T, Allocator> &x, vector<T, Allocator> &y)
```

Effects: `x.swap(y);`

```
template<typename T, typename Allocator = std::allocator<T>>
```

```
class vector
```

Public Types

```
typedef T value_type
    Member types (FIXME: add reference to std.

typedef Allocator allocator_type
typedef alloc_traits::access_target access_target
typedef std::size_t size_type
typedef std::ptrdiff_t difference_type
typedef alloc_traits::reference reference
typedef alloc_traits::const_reference const_reference
typedef alloc_traits::pointer pointer
typedef alloc_traits::const_pointer const_pointer
typedef detail::iterator<T, Allocator> iterator
typedef detail::iterator<T const, Allocator> const_iterator
typedef detail::reverse_iterator<T, Allocator> reverse_iterator
typedef detail::const_reverse_iterator<T, Allocator> const_reverse_iterator
```

Public Functions

```
vector (Allocator const &alloc = Allocator())
vector (size_type count, T const &value, Allocator const &alloc = Allocator())
vector (size_type count, Allocator const &alloc = Allocator())
template<typename InIter, typename Enable = typename std::enable_if<hpx::traits::is_input_iterator<InIter>::value>
vector (InIter first, InIter last, Allocator const &alloc)
vector (vector const &other)
vector (vector const &other, Allocator const &alloc)
vector (vector &&other)
vector (vector &&other, Allocator const &alloc)
vector (std::initializer_list<T> init, Allocator const &alloc)
~vector ()
vector &operator= (vector const &other)
vector &operator= (vector &&other)
allocator_type get_allocator() const
    Returns the allocator associated with the container.
HPX_HOST_DEVICE reference hpx::compute::vector::operator[](size_type pos)
HPX_HOST_DEVICE const_reference hpx::compute::vector::operator[](size_type pos)
```

pointer **data()**

Returns pointer to the underlying array serving as element storage. The pointer is such that range [*data()*; *data()* + *size()*] is always a valid range, even if the container is empty (*data()* is not dereferenceable in that case).

const_pointer **data() const**

Returns pointer to the underlying array serving as element storage. The pointer is such that range [*data()*; *data()* + *size()*] is always a valid range, even if the container is empty (*data()* is not dereferenceable in that case).

T *device_data() const

Returns a raw pointer corresponding to the address of the data allocated on the device.

std::size_t **size() const**

std::size_t **capacity() const**

bool empty() const

Returns: *size()* == 0.

void resize(*size_type* size)

Effects: If *size* <= *size()*, equivalent to calling *pop_back()* *size()* - *size* times. If *size()* < *size*, appends *size* - *size* default-inserted elements to the sequence.

Requires: *T* shall be *MoveInsertable* and *DefaultInsertable* into **this*.

Remarks: If an exception is thrown other than by the move constructor of a non-*CopyInsertable* *T* there are no effects.

void resize(*size_type* size, *T const &val*)

Effects: If *size* <= *size()*, equivalent to calling *pop_back()* *size()* - *size* times. If *size()* < *size*, appends *size* - *size* copies of *val* to the sequence.

Requires: *T* shall be *CopyInsertable* into **this*.

Remarks: If an exception is thrown there are no effects.

iterator **begin()**

iterator **end()**

const_iterator **cbegin() const**

const_iterator **cend() const**

const_iterator **begin() const**

const_iterator **end() const**

void swap(vector &other)

Effects: Exchanges the contents and *capacity()* of **this* with that of *x*.

Complexity: Constant time.

void clear()

Effects: Erases all elements in the range [*begin()*,*end()*). Destroys all elements in *a*. Invalidates all references, pointers, and iterators referring to the elements of *a* and may invalidate the past-the-end iterator.

Post: *a.empty()* returns true.

Complexity: Linear.

Private Types

```
typedef traits::allocator_traits<Allocator> alloc_traits
```

Private Members

```
size_type size_
size_type capacity_
allocator_type alloc_
pointer data_
```

```
#include <hpx/compute/traits.hpp>

#include <hpx/compute/host.hpp>

#include <hpx/compute/host/block_executor.hpp>

template<typename Executor>
struct executor_execution_category<compute::host::block_executor<Executor>>
```

Public Types

```
typedef parallel::execution::parallel_execution_tag type
namespace hpx

namespace compute

namespace host

template<typename Executor = hpx::threads::executors::local_priority_queue_attached_executor>
struct block_executor
#include <block_executor.hpp> The block executor can be used to build NUMA aware programs.
It will distribute work evenly across the passed targets
```

Template Parameters

- Executor: The underlying executor to use

Public Types

```
typedef hpx::parallel::execution::static_chunk_size executor_parameters_type
```

Public Functions

```
block_executor (std::vector<host::target> const &targets)
```

```
block_executor (std::vector<host::target> &&targets)
```

```
block_executor (block_executor const &other)
```

```
block_executor (block_executor &&other)
```

```
block_executor &operator= (block_executor const &other)
```

```
block_executor &operator= (block_executor &&other)
```

```
template<typename F, typename ...Ts>
```

```
void post (F &&f, Ts&&... ts)
```

```
template<typename F, typename ...Ts>
```

```
hpx::future<typename hpx::util::detail::invoke_deferred_result<F, Ts...>::type> async_execute (F  
  &&f,  
  Ts&&...  
  ts)
```

```
template<typename F, typename ...Ts>
```

```
hpx::util::detail::invoke_deferred_result<F, Ts...>::type sync_execute (F &&f, Ts&&...  
  ts)
```

```
template<typename F, typename Shape, typename ...Ts>
```

```
std::vector<hpx::future<typename parallel::execution::detail::bulk_function_result<F, Shape, Ts...>::type>> bulk_execute (F  
  &&f,  
  Shape  
  const  
  &shape,  
  Ts&&...  
  ts)
```

```
template<typename F, typename Shape, typename ...Ts>
```

```
parallel::execution::detail::bulk_execute_result<F, Shape, Ts...>::type bulk_sync_execute (F  
  &&f,  
  Shape  
  const  
  &shape,  
  Ts&&...  
  ts)
```

```
std::vector<host::target> const &targets () const
```

Private Functions

```
void init_executors()
```

Private Members

```
std::vector<host::target> targets_
std::atomic<std::size_t> current_
std::vector<Executor> executors_
```

```
namespace parallel
```

```
namespace execution
```

```
template<typename Executor>
struct executor_execution_category<compute::host::block_executor<Executor>>
```

Public Types

```
typedef parallel::execution::parallel_execution_tag type
```

```
#include <hpx/compute/host/get_targets.hpp>
```

```
namespace hpx
```

```
namespace compute
```

```
namespace host
```

Functions

```
std::vector<target> get_local_targets()
```

```
hpx::future<std::vector<target>> get_targets(hpx::id_type const &locality)
```

```
#include <hpx/compute/host/target.hpp>
```

```
namespace hpx
```

```
namespace compute
```

```
namespace host
```

```
struct target
```

Public Functions

```
target()
target(hpx::threads::mask_type mask)
target(hpx::id_type const &locality)
target(hpx::id_type const &locality, hpx::threads::mask_type mask)
native_handle_type &native_handle()
native_handle_type const &native_handle() const
hpx::id_type const &get_locality() const
std::pair<std::size_t, std::size_t> num_pus() const
void synchronize() const
hpx::future<void> get_future() const
```

Public Static Functions

```
static std::vector<target> get_local_targets()
static hpx::future<std::vector<target>> get_targets(hpx::id_type const &locality)
```

Private Functions

```
void serialize(serialization::input_archive &ar, const unsigned int)
void serialize(serialization::output_archive &ar, const unsigned int)
```

Private Members

```
native_handle_type handle_
hpx::id_type locality_
```

Friends

```
friend hpx::compute::host::hpx::serialization::access
bool operator==(target const &lhs, target const &rhs)
struct native_handle_type
```

Public Functions

```
native_handle_type()
native_handle_type (hpx::threads::mask_type mask)
hpx::threads::mask_type &get_device()
hpx::threads::mask_type const &get_device() const
```

Private Members

```
hpx::threads::mask_type mask_
```

Friends

```
friend hpx::compute::host::target
```

```
#include <hpx/compute/host/target_distribution_policy.hpp>

namespace hpx

namespace compute

namespace host
```

Variables

```
target_distribution_policy const target_layout
```

A predefined instance of the `target_distribution_policy` for localities. It will represent all NUMA domains of the given locality and will place all items to create here.

```
struct target_distribution_policy : public compute::detail::target_distribution_policy<host::target>
#include <target_distribution_policy.hpp> A target_distribution_policy used for CPU bound localities.
```

Public Types

```
typedef compute::detail::target_distribution_policy<host::target> base_type
```

Public Functions

`target_distribution_policy()`

Default-construct a new instance of a `target_distribution_policy`. This policy will represent all devices on the current locality.

`target_distribution_policy operator()(std::vector<target_type> const &targets, std::size_t num_partitions = std::size_t(-1)) const`

Create a new `target_distribution_policy` representing the given set of targets

Parameters

- `targets`: [in] The targets the new instances should represent

`target_distribution_policy operator()(std::vector<target_type> &&targets, std::size_t num_partitions = std::size_t(-1)) const`

Create a new `target_distribution_policy` representing the given set of targets

Parameters

- `targets`: [in] The targets the new instances should represent

`target_distribution_policy operator()(target_type const &target, std::size_t num_partitions = 1) const`

Create a new `target_distribution_policy` representing the given target

Parameters

- `target`: [in] The target the new instances should represent

`target_distribution_policy operator()(target_type &&target, std::size_t num_partitions = 1) const`

Create a new `target_distribution_policy` representing the given target

Parameters

- `target`: [in] The target the new instances should represent

`template<typename Component, typename ...Ts>`

`hpx::future<hpx::id_type> create(Ts&&... ts) const`

Create one object on one of the localities associated by this policy instance

Note This function is part of the placement policy implemented by this class

Return A future holding the global address which represents the newly created object

Parameters

- `ts`: [in] The arguments which will be forwarded to the constructor of the new object.

`template<typename Component, typename ...Ts>`

`hpx::future<std::vector<bulk_locality_result>> bulk_create(std::size_t count, Ts&&... ts) const`

Create multiple objects on the localities associated by this policy instance

Note This function is part of the placement policy implemented by this class

Return A future holding the list of global addresses which represent the newly created objects

Parameters

- `count`: [in] The number of objects to create
- `vs`: [in] The arguments which will be forwarded to the constructors of the new objects.

```
#include <hpx/compute/host/numa_domains.hpp>
```

```
namespace hpx
```

```
    namespace compute
```

```
        namespace host
```

Functions

```
    std::vector<target> numa_domains()
```

```
#include <hpx/compute/host/block_allocator.hpp>
```

```
namespace hpx
```

```
    namespace compute
```

```
        namespace host
```

```
template<typename T, typename Executor = hpx::parallel::execution::local_priority_queue_attached_executor>
struct block_allocator
```

#include <block_allocator.hpp> The `block_allocator` allocates blocks of memory evenly divided onto the passed vector of targets. This is done by using first touch memory placement. (maybe better methods will be used in the future...);

This allocator can be used to write NUMA aware algorithms:

```
typedef      hpx::compute::host::block_allocator<int>      allocator_type;      typedef
```

```
hpx::compute::vector<int, allocator_type> vector_type;
```

```
auto numa_nodes = hpx::compute::host::numa_domains(); std::size_t N = 2048; vector_type v(N,
```

```
allocator_type(numa_nodes));
```

Public Types

```
typedef T value_type
```

```
typedef T *pointer
```

```
typedef const T *const_pointer
```

```
typedef T &reference
```

```
typedef T const &const_reference
```

```
typedef std::size_t size_type
```

```
typedef std::ptrdiff_t difference_type
```

```
typedef Executor executor_type
```

```
typedef std::false_type is_always_equal
```

```
typedef std::true_type propagate_on_container_move_assignment
typedef std::vector<host::target> target_type
```

Public Functions

```
block_allocator()

block_allocator(target_type const &targets)

block_allocator(target_type &&targets)

block_allocator(block_allocator const &alloc)

block_allocator(block_allocator &&alloc)

template<typename U>
block_allocator(block_allocator<U> const &alloc)

template<typename U>
block_allocator(block_allocator<U> &&alloc)

block_allocator &operator=(block_allocator const &rhs)

block_allocator &operator=(block_allocator &&rhs)

pointer address(reference x) const

const_pointer address(const_reference x) const

pointer allocate(size_type n, std::allocator<void>::const_pointer hint = nullptr)

void deallocate(pointer p, size_type n)

size_type max_size() const

template<typename U, typename ...Args>
void bulk_construct(U *p, std::size_t count, Args&&... args)

template<typename U, typename ...Args>
void construct(U *p, Args&&... args)

template<typename U>
void bulk_destroy(U *p, std::size_t count)

template<typename U>
void destroy(U *p)

target_type const &target() const
```

Private Members

```
block_executor<executor_type> executor_
template<typename U>
struct rebind
```

Public Types

```
template<>
typedef block_allocator<U> other
```

```
#include <hpx/compute/host/traits/access_target.hpp>

template<>
struct access_target<host::target>
```

Public Types

```
typedef host::target target_type
```

Public Static Functions

```
template<typename T>
static T const &read(target_type const &tgt, T const *t)
```

```
template<typename T>
static void write(target_type const &tgt, T *dst, T const *src)
```

```
template<>
struct access_target<std::vector<host::target>>
```

Public Types

```
typedef std::vector<host::target> target_type
```

Public Static Functions

```
template<typename T>
static T const &read(target_type const &tgt, T const *t)
```

```
template<typename T>
static void write(target_type const &tgt, T *dst, T const *src)
```

```
namespace hpx
```

```
namespace compute
```

```
namespace traits
```

```
template<>
struct access_target<host::target>
```

Public Types

```
typedef host::target target_type
```

Public Static Functions

```
template<typename T>
static T const &read(target_type const &tgt, T const *t)
```

```
template<typename T>
static void write(target_type const &tgt, T *dst, T const *src)
```

```
template<>
struct access_target<std::vector<host::target>>
```

Public Types

```
typedef std::vector<host::target> target_type
```

Public Static Functions

```
template<typename T>
static T const &read(target_type const &tgt, T const *t)
```

```
template<typename T>
static void write(target_type const &tgt, T *dst, T const *src)
```

```
#include <hpx/compute/traits/access_target.hpp>
```

```
#include <hpx/compute/traits/allocator_traits.hpp>
```

```
namespace hpx
```

```
namespace compute
```

```
namespace traits
```

```
template<typename Allocator>
struct allocator_traits : public std::allocator_traits<Allocator>
```

Public Types

```
typedef detail::get_reference_type<Allocator>::type reference
typedef detail::get_const_reference_type<Allocator>::type const_reference
typedef detail::get_target_traits<Allocator>::type access_target
typedef access_target::target_type target_type
```

Public Static Functions

```
static HPX_HOST_DEVICE auto hpx::compute::traits::allocator_traits::target(Allocator&)
template<typename... Ts> static HPX_HOST_DEVICE void hpx::compute::traits::allocator_traits::bulk_destory(Ts...)
static HPX_HOST_DEVICE void hpx::compute::traits::allocator_traits::bulk_destory()
```

Private Types

```
typedef std::allocator_traits<Allocator> base_type
```

```
#include <hpx/compute/serialization/vector.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename T, typename Allocator>
void serialize(input_archive &ar, compute::vector<T, Allocator> &v, unsigned)

template<typename T, typename Allocator>
void serialize(output_archive &ar, compute::vector<T, Allocator> const &v, unsigned)
```

debugging

```
#include <compatibility/hpx/util/debug/demangle_helper.hpp>
```

```
#include <hpx/debugging/demangle_helper.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
namespace debug
```

TypeDefs

```
using cxxabi_demangle_helper = demangle_helper<T>
using cxx_type_id = type_id<T>
```

Functions

```
template<typename T = void>
std::string print_type(const char *delim = "")

template<>
std::string print_type(const char *)

template<typename T, typename ...Args>
std::enable_if<sizeof...(Args) != 0, std::string>::type print_type(const char *delim = "")

template<typename T>
struct demangle_helper
```

Public Functions

```
char const *type_id() const

template<typename T>
struct type_id
```

Public Static Attributes

```
demangle_helper<T> typeid_ = demangle_helper<T>()
```

```
#include <hpx/debugging/print.hpp>
```

coroutines

```
#include <compatibility/hpx/runtime/threads/thread_enums.hpp>

#include <compatibility/hpx/runtime/threads/thread_id_type.hpp>

#include <compatibility/hpx/runtime/threads/coroutines/coroutine_fwd.hpp>

#include <compatibility/hpx/runtime/threads/coroutines/coroutine.hpp>

#include <hpx/coroutines/coroutine_fwd.hpp>

#include <hpx/coroutines/stackless_coroutine.hpp>

namespace hpx
```

```
namespace threads

namespace coroutines

class stackless_coroutine
```

Public Types

```
using thread_id_type = hpx::threads::thread_id
using result_type = std::pair<thread_state_enum, thread_id_type>
using arg_type = thread_state_ex_enum
using functor_type = util::unique_function_nonsig<result_type (arg_type)>
```

Public Functions

```
stackless_coroutine (functor_type &&f, thread_id_type id, std::ptrdiff_t stack_size =
                     default_stack_size)

~stackless_coroutine()

stackless_coroutine (stackless_coroutine const &src)

stackless_coroutine &operator= (stackless_coroutine const &src)

stackless_coroutine (stackless_coroutine &&src)

stackless_coroutine &operator= (stackless_coroutine &&src)

thread_id_type get_thread_id() const

std::size_t get_thread_data() const

std::size_t set_thread_data (std::size_t data)

void rebind (functor_type &&f, thread_id_type id)

void reset_tss()

void reset()

stackless_coroutine::result_type operator() (arg_type arg = arg_type())

operator bool() const

bool is_ready() const

std::ptrdiff_t get_available_stack_space()

std::size_t &get_continuation_recursion_count()
```

Protected Attributes

```
functor_type f_
context_state state_
thread_id_type id_
std::size_t thread_data_
std::size_t continuation_recursion_count_
```

Private Types

```
enum context_state
Values:
ctx_running
ctx_ready
ctx_exited
```

Private Functions

```
bool running() const
bool exited() const
```

Private Members

```
HPX_STATIC_CONSTEXPR std::ptrdiff_t hpx::threads::coroutines::stackless_corou
```

Friends

```
friend hpx::threads::coroutines::reset_on_exit
struct reset_on_exit
```

Public Functions

```
reset_on_exit(stackless Coroutine &this__)
~reset_on_exit()
```

Public Members

stacklessCoroutine &**this**

```
#include <hpx/coroutines/thread_enums.hpp>
```

```
namespace hpx
```

```
namespace threads
```

Enums

enum `thread_state_enum`

The `thread_state_enum` enumerator encodes the current state of a `thread` instance

Values:

unknown = 0

active = 1

thread is currently active (running, has resources)

pending = 2

thread is pending (ready to run, but no hardware resource available)

suspended = 3

thread has been suspended (waiting for synchronization event, but still known and under control of the thread-manager)

depleted = 4

thread has been depleted (deeply suspended, it is not known to the thread-manager)

terminated = 5

thread has been stopped an may be garbage collected

staged = 6

this is not a real thread state, but allows to reference staged task descriptions, which eventually will be converted into thread objects

pending_do_not_schedule = 7

pending_boost = 8

enum `thread_priority`

This enumeration lists all possible thread-priorities for HPX threads.

Values:

thread_priority_unknown = -1

thread_priority_default = 0

Will assign the priority of the task to the default (normal) priority.

thread_priority_low = 1

Task goes onto a special low priority queue and will not be executed until all high/normal priority tasks are done, even if they are added after the low priority task.

thread_priority_normal = 2

Task will be executed when it is taken from the normal priority queue, this is usually a first-in-first-out ordering of tasks (depending on scheduler choice). This is the default priority.

thread_priority_high_recursive = 3

The task is a high priority task and any child tasks spawned by this task will be made high priority as well - unless they are specifically flagged as non default priority.

thread_priority_boost = 4

Same as *thread_priority_high* except that the thread will fall back to *thread_priority_normal* if resumed after being suspended.

thread_priority_high = 5

Task goes onto a special high priority queue and will be executed before normal/low priority tasks are taken (some schedulers modify the behavior slightly and the documentation for those should be consulted).

thread_priority_bound = 6

Task goes onto a special high priority queue and will never be stolen by another thread after initial assignment. This should be used for thread placement tasks such as OpenMP type for loops.

enum thread_state_ex_enum

The *thread_state_ex_enum* enumerator encodes the reason why a thread is being restarted

Values:

wait_unknown = 0**wait_signaled = 1**

The thread has been signaled.

wait_timeout = 2

The thread has been reactivated after a timeout

wait_terminate = 3

The thread needs to be terminated.

wait_abort = 4

The thread needs to be aborted.

enum thread_stacksize

A *thread_stacksize* references any of the possible stack-sizes for HPX threads.

Values:

thread_stacksize_unknown = -1**thread_stacksize_small = 1**

use small stack size

thread_stacksize_medium = 2

use medium sized stack size

thread_stacksize_large = 3

use large stack size

thread_stacksize_huge = 4

use very large stack size

thread_stacksize_nostack = 5

this thread does not suspend (does not need a stack)

thread_stacksize_current = 6

use size of current thread's stack

thread_stacksize_default = *thread_stacksize_small*

use default stack size

thread_stacksize_minimal = *thread_stacksize_small*
use minimally stack size

thread_stacksize_maximal = *thread_stacksize_huge*
use maximally stack size

enum thread_schedule_hint_mode
The type of hint given when creating new tasks.

Values:

thread_schedule_hint_mode_none = 0

thread_schedule_hint_mode_thread = 1

thread_schedule_hint_mode_numa = 2

Functions

char const *get_thread_state_name (thread_state_enum state)

Returns the name of the given state.

Get the readable string representing the name of the given thread_state constant.

Parameters

- *state*: this represents the thread state.

char const *get_thread_priority_name (thread_priority priority)

Return the thread priority name.

Get the readable string representing the name of the given thread_priority constant.

Parameters

- *this*: represents the thread priority.

char const *get_thread_state_ex_name (thread_state_ex_enum state)

Get the readable string representing the name of the given thread_state_ex_enum constant.

char const *get_thread_state_name (thread_state state)

Get the readable string representing the name of the given thread_state constant.

char const *get_stack_size_name (std::ptrdiff_t size)

Returns the stack size name.

Get the readable string representing the given stack size constant.

Parameters

- *size*: this represents the stack size

struct thread_schedule_hint

Public Functions

```
thread_schedule_hint()  
thread_schedule_hint(std::int16_t thread_hint)  
thread_schedule_hint(thread_schedule_hint_mode mode, std::int16_t hint)
```

Public Members

```
thread_schedule_hint_mode mode  
std::int16_t hint
```

```
#include <hpx/coroutines/thread_id_type.hpp>  
  
namespace hpx
```

```
namespace threads
```

Variables

```
constexpr thread_id invalid_thread_id  
struct thread_id
```

Public Functions

```
constexpr thread_id()  
constexpr thread_id(thread_id_repr thrd)  
thread_id(thread_id const&)  
thread_id &operator=(thread_id const&)  
thread_id(thread_id &&rhs)  
thread_id &operator=(thread_id &&rhs)  
constexpr operator bool() const  
constexpr thread_id_repr get() const  
constexpr void reset()
```

Private Types

```
using thread_id_repr = void *
```

Private Members

```
thread_id_repr thrd_
```

Friends

```
constexpr bool operator== (std::nullptr_t, thread_id const &rhs)
constexpr bool operator!= (std::nullptr_t, thread_id const &rhs)
constexpr bool operator== (thread_id const &lhs, std::nullptr_t)
constexpr bool operator!= (thread_id const &lhs, std::nullptr_t)
constexpr bool operator== (thread_id const &lhs, thread_id const &rhs)
constexpr bool operator!= (thread_id const &lhs, thread_id const &rhs)
constexpr bool operator< (thread_id const &lhs, thread_id const &rhs)
constexpr bool operator> (thread_id const &lhs, thread_id const &rhs)
constexpr bool operator<= (thread_id const &lhs, thread_id const &rhs)
constexpr bool operator>= (thread_id const &lhs, thread_id const &rhs)
template<typename Char, typename Traits>
std::basic_ostream<Char, Traits> &operator<< (std::basic_ostream<Char, Traits> &os,
thread_id const &id)
```

```
#include <hpx/coroutines/coroutine.hpp>
```

```
namespace hpx
```

```
    namespace threads
```

```
        namespace coroutines
```

```
            class coroutine
```

Public Types

```
using impl_type = detail::coroutine_impl
using thread_id_type = impl_type::thread_id_type
using result_type = impl_type::result_type
using arg_type = impl_type::arg_type
using functor_type = util::unique_function_nonser<result_type (arg_type)>
```

Public Functions

```
coroutine (functor_type &&f, thread_id_type id, std::ptrdiff_t stack_size = detail::default_stack_size)
coroutine (coroutine const &src)

coroutine &operator=(coroutine const &src)

coroutine (coroutine &&src)

coroutine &operator=(coroutine &&src)

thread_id_type get_thread_id() const
std::size_t get_thread_data() const
std::size_t set_thread_data(std::size_t data)

void rebind(functor_type &&f, thread_id_type id)

result_type operator()(arg_type arg = arg_type())
bool is_ready() const
std::ptrdiff_t get_available_stack_space()

impl_type *impl()
```

Private Members

```
impl_type impl_
```

functional

```
#include <compatibility/hpx/util/unique_function.hpp>

#include <compatibility/hpx/util/bind_back.hpp>

#include <compatibility/hpx/util/bind_front.hpp>

#include <compatibility/hpx/util/result_of.hpp>
```

```
#include <compatibility/hpx/util/deferred_call.hpp>

#include <compatibility/hpx/util/one_shot.hpp>

#include <compatibility/hpx/util/invoke_fused.hpp>

#include <compatibility/hpx/util/function.hpp>

#include <compatibility/hpx/util/function_ref.hpp>

#include <compatibility/hpx/util/bind.hpp>

#include <compatibility/hpx/util/protect.hpp>

#include <compatibility/hpx/util/invoke.hpp>

#include <compatibility/hpx/util/first_argument.hpp>

#include <compatibility/hpx/util/mem_fn.hpp>

#include <hpx/functional/unique_function.hpp>
```

Defines

```
HPX_UTIL_REGISTER_UNIQUE_FUNCTION_DECLARATION(Sig, F, Name)
```

```
HPX_UTIL_REGISTER_UNIQUE_FUNCTION(Sig, F, Name)
```

```
namespace hpx
```

```
    namespace util
```

TypeDefs

```
using unique_function_nonser = unique_function<Sig, false>

template<typename R, typename ...Ts, bool Serializable>
class unique_function<R(Ts...), Serializable> : public detail::basic_function<R
    Ts..., false, Serializable>
```

Public Types

```
typedef R result_type
```

Public Functions

```
constexpr unique_function(std::nullptr_t = nullptr)
```

```
unique_function(unique_function&&)
```

```
unique_function &operator=(unique_function&&)
```

```
template<typename F, typename FD = typename std::decay<F>::type, typename Enable1 = typename std::enable_if<std::is_invocable<FD>::value, FD> FD>
unique_function(F &&f)
```

```
template<typename F, typename FD = typename std::decay<F>::type, typename Enable1 = typename std::enable_if<std::is_invocable<FD>::value, FD> FD>
unique_function &operator=(F &&f)
```

Private Types

```
template<>
using base_type = detail::basic_function<R (Ts...), false, Serializable>
```

```
#include <hpx/functional/bind_back.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Archive, typename F, typename ...Ts>
```

```
void serialize(Archive &ar, hpx::util::detail::bound_back<F, Ts...> &bound, unsigned int
const version = 0)
```

```
namespace util
```

Functions

```
template<typename F, typename ...Ts>
```

```
constexpr detail::bound_back<typename std::decay<F>::type, typename util::make_index_pack<sizeof...(Ts)>::type>
```

```
template<typename F>
```

```
constexpr std::decay<F>::type bind_back(F &&f)
```

```
#include <hpx/functional/bind_front.hpp>

namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Archive, typename F, typename ...Ts>
void serialize(Archive &ar, hpx::util::detail::bound_front<F, Ts...> &bound, unsigned int
               const version = 0)

namespace util
```

Functions

```
template<typename F, typename ...Ts>
constexpr detail::bound_front<typename std::decay<F>::type, typename util::make_index_pack<sizeof...(Ts)>::type>
```

```
template<typename F>
constexpr std::decay<F>::type bind_front (F &&f)
```

```
#include <hpx/functional/result_of.hpp>
```

```
#include <hpx/functional/deferred_call.hpp>
```

```
namespace hpx
```

```
namespace serialization
```

Functions

```
template<typename Archive, typename F, typename ...Ts>
void serialize(Archive &ar, hpx::util::detail::deferred<F, Ts...> &d, unsigned int const version
               = 0)
```

```
namespace util
```

Functions

```
template<typename F, typename ...Ts>
detail::deferred<typename std::decay<F>::type, typename util::make_index_pack<sizeof...(Ts)>::type, typename util::make_index_pack<sizeof...(Ts)>::type> make_index_pack(F f, Ts... ts)

template<typename F>
std::decay<F>::type deferred_call(F f, Ts... ts)

#include <hpx/functional/one_shot.hpp>

namespace hpx

namespace serialization
```

Functions

```
template<typename Archive, typename F>
void serialize(Archive &ar, hpx::util::detail::one_shot_wrapper<F> &one_shot_wrapper, unsigned int const version = 0)
```

```
namespace util
```

Functions

```
template<typename F>
constexpr detail::one_shot_wrapper<typename std::decay<F>::type> one_shot(F f)
```

```
#include <hpx/functional/invoke_fused.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
template<typename F, typename Tuple> constexpr HPX_HOST_DEVICE detail::invoke_fused(F f, Tuple t)
```

Invokes the given callable object f with the content of the sequenced type t (tuples, pairs)

Return The result of the callable object when it's called with the content of the given sequenced type.

Note This function is similar to `std::apply` (C++17)

Parameters

- f: Must be a callable object. If f is a member function pointer, the first argument in the sequenced type will be treated as the callee (this object).
- t: A type which is content accessible through a call to `hpx::util::get`.

Exceptions

- std::exception: like objects thrown by call to object f with the arguments contained in the sequenceable type t.

```
template<typename R, typename F, typename Tuple>constexpr HPX_HOST_DEV TemplateParameter
```

- R: The result type of the function when it's called with the content of the given sequenced type.

```
#include <hpx/functional/function.hpp>
```

Defines

```
HPX_UTIL_REGISTER_FUNCTION_DECLARATION(Sig, F, Name)
```

```
HPX_UTIL_REGISTER_FUNCTION(Sig, F, Name)
```

```
namespace hpx
```

```
namespace util
```

TypeDefs

```
using function_nonser = function<Sig, false>

template<typename R, typename ...Ts, bool Serializable>
class function<R(Ts...), Serializable> : public detail::basic_function<R
    Ts..., true, Serializable>
```

Public Types

```
typedef R result_type
```

Public Functions

```
constexpr function(std::nullptr_t = nullptr)
```

```
function(function const&)
```

```
function(function&&)
```

```
function &operator=(function const&)
```

```
function &operator=(function&&)
```

```
template<typename F, typename FD = typename std::decay<F>::type, typename Enable1 = typename std::enable_if<FD == function>::type>
function(F &&f)
```

```
template<typename F, typename FD = typename std::decay<F>::type, typename Enable1 = typename std::enable_if<FD == function>::type>
function &operator=(F &&f)
```

Private Types

```
template<>
using base_type = detail::basic_function<R (Ts...) , true, Serializable>

#include <hpx/functional/function_ref.hpp>

namespace hpx

namespace util

template<typename R, typename ...Ts>
class function_ref<R (Ts...) >
```

Public Functions

```
template<typename F, typename FD = typename std::decay<F>::type, typename Enable = typename std::enable_if<FD>>
function_ref(F &&f)

function_ref(function_ref const &other)

template<typename F, typename FD = typename std::decay<F>::type, typename Enable = typename std::enable_if<FD>
function_ref &operator=(F &&f)

function_ref &operator=(function_ref const &other)

template<typename F, typename T = typename std::remove_reference<F>::type, typename Enable = typename std::enable_if<FD>
void assign(F &&f)

template<typename T>
void assign(std::reference_wrapper<T> f_ref)

template<typename T>
void assign(T *f_ptr)

void swap(function_ref &f)

R operator() (Ts... vs) const

std::size_t get_function_address() const

char const *get_function_annotation() const

util::itt::string_handle get_function_annotation_itt() const
```

Protected Attributes

```
template<>
R (*vptr) (void *, Ts&&...)
```

Private Types

```
template<>
using vTable = detail::function_ref_vtable<R (Ts...) >;
```

Private Static Functions

```
template<typename T>  
static VTable const *get_vtable()
```

```
#include <hpx/functional/bind.hpp>
```

```
namespace hpx
```

namespace serialization

Functions

```
template<typename Archive, typename F, typename ...Ts>
void serialize(Archive &ar, hpx::util::detail::bound<F, Ts...> &bound, unsigned int const version = 0)
```

```
template<typename Archive, std::size_t I>
void serialize(Archive &ar, hpx::util::detail::placeholder<I>&, unsigned int const = 0)
```

namespace util

Functions

```
template<typename F, typename ...Ts, typename Enable = typename std::enable_if<!traits::is_action<typename std::decay<F>::type, typename util::make_index_pack<sizeof...(Ts)>::type, typename util::make_index_pack<sizeof...(Ts)>::type>::value> constexpr detail::bound<typename std::decay<F>::type, typename util::make_index_pack<sizeof...(Ts)>::type, typename util::make_index_pack<sizeof...(Ts)>::type> bound(F f, Ts... ts) {
```

namespace placeholders

Variables

```
HPX_STATIC_CONSTEXPR detail::placeholder<1> hpx::util::placeholders::_1 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<2> hpx::util::placeholders::_2 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<3> hpx::util::placeholders::_3 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<4> hpx::util::placeholders::_4 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<5> hpx::util::placeholders::_5 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<6> hpx::util::placeholders::_6 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<7> hpx::util::placeholders::_7 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<8> hpx::util::placeholders::_8 = {}
HPX_STATIC_CONSTEXPR detail::placeholder<9> hpx::util::placeholders::_9 = {}

#include <hpx/functional/protect.hpp>

namespace hpx

    namespace util
```

Functions

```
template<typename T>HPX_HOST_DEVICE std::enable_if< traits::is_bind_expression<type>
template<typename T>HPX_HOST_DEVICE std::enable_if< !traits::is_bind_expression<type>
```

```
#include <hpx/functional/invoke.hpp>
```

Defines

```
HPX_INVOKE (F, ...)
HPX_INVOKE_R (R, F, ...)
namespace hpx

    namespace util
```

Functions

```
template<typename F, typename... Ts>constexpr HPX_HOST_DEVICE util::invoke_result<F>
    Invokes the given callable object f with the content of the argument pack vs
```

Return The result of the callable object when it's called with the given argument types.

Note This function is similar to `std::invoke` (C++17)

Parameters

- `f`: Requires to be a callable object. If `f` is a member function pointer, the first argument in the pack will be treated as the callee (this object).
- `vs`: An arbitrary pack of arguments

Exceptions

- std::exception: like objects thrown by call to object f with the argument types vs.

```
template<typename R, typename F, typename... Ts>constexpr HPX_HOST_DEVICE TemplateParam
```

- R: The result type of the function when it's called with the content of the given argument types vs.

```
namespace functional
```

```
struct invoke
```

Public Functions

```
template<typename F, typename... Ts>constexpr HPX_HOST_DEVICE util::invoke_re
```

```
template<typename R>  
struct invoke_r
```

Public Functions

```
template<typename F, typename... Ts>constexpr HPX_HOST_DEVICE R hpx::util::fu
```

```
#include <hpx/functional/first_argument.hpp>
```

```
#include <hpx/functional/mem_fn.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
template<typename M, typename C>  
constexpr detail::mem_fn<M C::*> mem_fn (M C::*pm)
```

```
template<typename R, typename C, typename ...Ps>  
constexpr detail::mem_fn<R (C::*)(Ps...) > mem_fn  
R (C::*pm)Ps...
```

```
template<typename R, typename C, typename ...Ps>  
constexpr detail::mem_fn<R (C::*)(Ps...) const> mem_fn  
R (C::*pm)Ps... const
```

```
#include <hpx/functional/traits/is_placeholder.hpp>

namespace hpx

namespace traits

struct
    Subclassed by hpx::traits::is_placeholder< T const >

#include <hpx/functional/traits/is_callable.hpp>

#include <hpx/functional/traits/is_bind_expression.hpp>

namespace hpx

namespace traits

template<typename T>
struct is_bind_expression : public std::is_bind_expression<T>
    Subclassed by hpx::traits::is_bind_expression< T const >

#include <hpx/functional/traits/is_action.hpp>

#include <hpx/functional/traits/get_function_address.hpp>

template<typename R, typename Obj, typename ...Ts>
struct get_function_address<R (Obj::*) (Ts...) >
```

Public Static Functions

```
static std::size_t call (R (Obj::*)f) Ts...
template<typename R, typename Obj, typename ...Ts>
struct get_function_address<R (Obj::*) (Ts...) const>
```

Public Static Functions

```
static std::size_t call (R (Obj::*)f) Ts...
    const
namespace hpx

namespace traits

template<typename F, typename Enable = void>
struct get_function_address
```

Public Static Functions

```
static std::size_t call (F const &f)
template<typename R, typename ...Ts>
struct get_function_address<R (*) (Ts...) >
```

Public Static Functions

```
static std::size_t call (R (*f)) Ts...
template<typename R, typename Obj, typename ...Ts>
struct get_function_address<R (Obj::*) (Ts...) const>
```

Public Static Functions

```
static std::size_t call (R (Obj::*f)) Ts...
const
template<typename R, typename Obj, typename ...Ts>
struct get_function_address<R (Obj::*) (Ts...) >
```

Public Static Functions

```
static std::size_t call (R (Obj::*f)) Ts...
#include <hpx/functional/traits/get_function_annotation.hpp>
namespace hpx
namespace traits
template<typename F, typename Enable = void>
struct get_function_annotation
```

Public Static Functions

```
static char const *call (F const&)
```

affinity

```
#include <compatibility/hpx/runtime/threads/policies/affinity_data.hpp>

#include <compatibility/hpx/runtime/threads/policies/parse_affinity_options.hpp>

#include <hpx/affinity/affinity_data.hpp>

#include <hpx/affinity/parse_affinity_options.hpp>

namespace hpx

namespace threads
```

Functions

```
void parse_affinity_options(std::string const &spec, std::vector<mask_type> &affinities, std::size_t used_cores, std::size_t max_cores, std::size_t num_threads, std::vector<std::size_t> &num_pus, bool use_process_mask, error_code &ec = throws)

void parse_affinity_options(std::string const &spec, std::vector<mask_type> &affinities, error_code &ec = throws)
```

type_support

```
#include <compatibility/hpx/util/lazy_enable_if.hpp>

#include <compatibility/hpx/util/unwrap_ref.hpp>

#include <compatibility/hpx/util/identity.hpp>

#include <compatibility/hpx/util/lazy_conditional.hpp>

#include <compatibility/hpx/util/unused.hpp>

#include <compatibility/hpx/util/detected.hpp>

#include <compatibility/hpx/util/decay.hpp>

#include <compatibility/hpx/util/always_void.hpp>

#include <compatibility/hpx/util/static.hpp>

#include <compatibility/hpx/util/void_guard.hpp>
```

```
#include <hpx/type_support/lazy_enable_if.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
template<typename T>
struct lazy_enable_if<true, T>
```

Public Types

```
typedef T::type type
```

```
#include <hpx/type_support/unwrap_ref.hpp>
```

```
template<typename T>
struct unwrap_reference<boost::reference_wrapper<T>>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<boost::reference_wrapper<T> const>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<std::reference_wrapper<T>>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<std::reference_wrapper<T> const>
```

Public Types

```
typedef T type
```

```
namespace hpx
```

```
namespace util
```

Functions

```
template<typename T>
unwrap_reference<T>::type &unwrap_ref (T &t)
```

```
template<typename T>
struct unwrap_reference
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<boost::reference_wrapper<T>>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<boost::reference_wrapper<T> const>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<std::reference_wrapper<T>>
```

Public Types

```
typedef T type
```

```
template<typename T>
struct unwrap_reference<std::reference_wrapper<T> const>
```

Public Types

```
typedef T type
```

```
#include <hpx/type_support/identity.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
template<typename T>
struct identity
```

Public Types

```
typedef T type

#include <hpx/type_support/pack.hpp>

namespace hpx

namespace util
```

Typedefs

```
using index_pack = pack_c<std::size_t, Is...>

template<typename ...Ts>
struct pack
```

Public Types

```
typedef pack type
```

Public Static Attributes

```
const std::size_t size = sizeof...(Ts)

template<typename T, T... Vs>
struct pack_c
    Subclassed by hpx::util::detail::make_index_pack_join< index_pack< Left... >, index_pack<
    Right... >>, hpx::util::make_index_pack< 1 >
```

Public Types

```
typedef pack_c type
```

Public Static Attributes

```
const std::size_t size = sizeof...(Vs)
```

```
#include <hpx/type_support/lazy_conditional.hpp>
```

```
#include <hpx/type_support/unused.hpp>
```

Defines

```
HPX_UNUSED(x)
```

```
namespace hpx
```

```
namespace util
```

Variables

```
constexpr unused_type unused = unused_type()
```

```
struct unused_type
```

Public Functions

```
constexpr HPX_HOST_DEVICE unused_type()
```

```
constexpr HPX_HOST_DEVICE unused_type (unused_type const&)
```

```
constexpr HPX_HOST_DEVICE unused_type (unused_type&&)
```

```
template<typename T>
```

```
constexpr HPX_HOST_DEVICE unused_type (T const&)
```

```
template<typename T>constexpr HPX_HOST_DEVICE unused_type const& hpx::util::unus
```

```
template<typename T>HPX_HOST_DEVICE unused_type& hpx::util::unused_type::operator
```

```
constexpr HPX_HOST_DEVICE unused_type const& hpx::util::unused_type::operator=(u
```

```
HPX_HOST_DEVICE unused_type& hpx::util::unused_type::operator=(unused_type const
```

```
constexpr HPX_HOST_DEVICE unused_type const& hpx::util::unused_type::operator=(u
```

```
HPX_HOST_DEVICE unused_type& hpx::util::unused_type::operator=(unused_type &&)
```

```
#include <hpx/type_support/detected.hpp>
```

```
namespace hpx
```

```
namespace util
```

Typedefs

```
using is_detected = typename detail::detector<nonesuch, void, Op, Args...>::value_t
```

```
using detected_t = typename detail::detector<nonesuch, void, Op, Args...>::type
```

```
using detected_or = detail::detector<Default, void, Op, Args...>
```

```
using detected_or_t = typename detected_or<Default, Op, Args...>::type
```

```
using is_detected_exact = std::is_same<Expected, detected_t<Op, Args...>>
```

```
using is_detected_convertible = std::is_convertible<detected_t<Op, Args...>, To>
```

```
struct nonesuch
```

Public Functions

```
nonesuch()  
~nonesuch()  
nonesuch(nonesuch const&)  
void operator=(nonesuch const&)
```

```
#include <hpx/type_support/decay.hpp>  
  
#include <hpx/type_support/always_void.hpp>  
  
namespace hpx
```

```
    namespace util
```

```
template<typename ...T>  
struct always_void
```

Public Types

```
typedef void type
```

```
#include <hpx/type_support/static.hpp>
```

Defines

```
HPX_EXPORT_STATIC_  
namespace hpx
```

```
    namespace util
```

```
template<typename T, typename Tag = T>  
struct static_
```

Public Types

```
typedef T value_type  
typedef T &reference  
typedef T const &const_reference
```

Public Functions

```
HPX_NON_COPYABLE (static_)

static_()

operator reference()

operator const_reference() const

reference get()

const_reference get() const
```

Private Types

```
typedef std::add_pointer<value_type>::type pointer

typedef std::aligned_storage<sizeof(value_type), std::alignment_of<value_type>::value>::type storage_type
```

Private Static Functions

```
static pointer get_address()
```

Private Static Attributes

```
static_<T, Tag>::storage_type data_

std::once_flag constructed_

struct default_constructor
```

Public Static Functions

```
template<>
static void construct()

struct destructor
```

Public Functions

```
template<>
~destructor()
```

```
#include <hpx/type_support/void_guard.hpp>
```

```
namespace hpx
```

```
namespace util
```

```
template<>
struct void_guard<void>
```

Public Functions

```
template<typename T>HPX_HOST_DEVICE void hpx::util::void_guard::operator,(T cons
```

```
resource_partitioner
```

```
#include <compatibility/hpx/runtime/resource/partitioner.hpp>
```

```
#include <compatibility/hpx/runtime/resource/partitioner_fwd.hpp>
```

```
#include <hpx/resource_partitioner/partitioner.hpp>
```

```
namespace hpx
```

```
namespace resource
```

```
class core
```

Public Functions

```
core (std::size_t id = invalid_core_id, numa_domain *domain = nullptr)
```

```
std::vector<pu> const &pus () const
```

```
std::size_t id () const
```

Private Functions

```
std::vector<core> cores_sharing_numa_domain()
```

Private Members

```
std::size_t id_
numa_domain *domain_
std::vector<pu> pus_
```

Private Static Attributes

```
constexpr const std::size_t invalid_core_id = std::size_t(-1)
```

Friends

```
friend hpx::resource::pu
friend hpx::resource::numa_domain

class numa_domain
```

Public Functions

```
numa_domain(std::size_t id = invalid_numa_domain_id)
std::vector<core> const &cores() const
std::size_t id() const
```

Private Members

```
std::size_t id_
std::vector<core> cores_
```

Private Static Attributes

```
constexpr const std::size_t invalid_numa_domain_id = std::size_t(-1)
```

Friends

```
friend hpx::resource::pu
friend hpx::resource::core

class partitioner
```

Public Functions

```

partitioner(util::function_nonsr<int> hpx::program_options::variables_map &vm
    > const &f, hpx::program_options::options_description const &desc_cmdline, int argc,
    char **argv, std::vector<std::string> ini_config, resource::partitioner_mode rmode = re-
    source::mode_default, runtime_mode mode = runtime_mode_default)

partitioner(util::function_nonsr<int> int, char **
    > const &f, int argc, char **argv, resource::partitioner_mode rmode = re-
    source::mode_default, hpx::runtime_mode mode = hpx::runtime_mode_default)

partitioner(util::function_nonsr<int> int, char **
    > const &f, int argc, char **argv, std::vector<std::string> const &cfg, re-
    source::partitioner_mode rmode = resource::mode_default, hpx::runtime_mode mode =
    hpx::runtime_mode_default)

partitioner(int argc, char **argv, resource::partitioner_mode rmode = re-
    source::mode_default, runtime_mode mode = runtime_mode_default)

partitioner(int argc, char **argv, std::vector<std::string> ini_config, re-
    source::partitioner_mode rmode = resource::mode_default, runtime_mode
    mode = runtime_mode_default)

partitioner(hpx::program_options::options_description const &desc_cmdline, int argc,
    char **argv, resource::partitioner_mode rmode = resource::mode_default,
    runtime_mode mode = runtime_mode_default)

partitioner(hpx::program_options::options_description const &desc_cmdline, int argc,
    char **argv, std::vector<std::string> ini_config, resource::partitioner_mode
    rmode = resource::mode_default, runtime_mode mode = run-
    time_mode_default)

partitioner(std::nullptr_t f, int argc, char **argv, resource::partitioner_mode
    rmode = resource::mode_default, hpx::runtime_mode mode =
    hpx::runtime_mode_default)

partitioner(std::nullptr_t f, int argc, char **argv, std::vector<std::string> const
    &cfg, resource::partitioner_mode rmode = resource::mode_default,
    hpx::runtime_mode mode = hpx::runtime_mode_default)

partitioner(std::nullptr_t f, hpx::program_options::options_description const
    &desc_cmdline, int argc, char **argv, std::vector<std::string> ini_config,
    resource::partitioner_mode rmode = resource::mode_default, runtime_mode
    mode = runtime_mode_default)

void create_thread_pool(std::string const &name, scheduling_policy
    sched = scheduling_policy::unspecified,
    hpx::threads::policies::scheduler_mode =
    hpx::threads::policies::scheduler_mode::default_mode)

void create_thread_pool(std::string const &name, scheduler_function sched-
    uler_creation)

void set_default_pool_name(std::string const &name)

const std::string &get_default_pool_name() const

void add_resource(hpx::resource::pu const &p, std::string const &pool_name,
    std::size_t num_threads = 1)

```

```
void add_resource (hpx::resource::pu const &p, std::string const &pool_name, bool exclusive, std::size_t num_threads = 1)

void add_resource (std::vector<hpx::resource::pu> const &pv, std::string const &pool_name, bool exclusive = true)

void add_resource (hpx::resource::core const &c, std::string const &pool_name, bool exclusive = true)

void add_resource (std::vector<hpx::resource::core> &cv, std::string const &pool_name, bool exclusive = true)

void add_resource (hpx::resource::numa_domain const &nd, std::string const &pool_name, bool exclusive = true)

void add_resource (std::vector<hpx::resource::numa_domain> const &ndv, std::string const &pool_name, bool exclusive = true)

std::vector<numa_domain> const &numa_domains() const

std::size_t get_number_requested_threads()

hpx::threads::topology const &get_topology() const
```

Private Members

```
detail::partitioner &partitioner_

class pu
```

Public Functions

```
pu (std::size_t id = invalid_pu_id, core *core = nullptr, std::size_t thread_occupancy = 0)

std::size_t id() const
```

Private Functions

```
std::vector<pu> pus_sharing_core()

std::vector<pu> pus_sharing_numa_domain()
```

Private Members

```
std::size_t id_

core *core_

std::size_t thread_occupancy_

std::size_t thread_occupancy_count_
```

Private Static Attributes

```
constexpr const std::size_t invalid_pu_id = std::size_t(-1)
```

Friends

```
friend hpx::resource::core
friend hpx::resource::numa_domain
```

```
#include <hpx/resource_partitioner/partitioner_fwd.hpp>
```

```
namespace hpx
```

```
namespace resource
```

TypeDefs

```
using scheduler_function = util::function_nonser<std::unique_ptr<hpx::threads::thread_pool_base> (hpx::thread
hpx::thread
```

Enums

enum partitioner_mode

This enumeration describes the modes available when creating a resource partitioner.

Values:

mode_default = 0

Default mode.

mode_allow_oversubscription = 1

Allow processing units to be oversubscribed, i.e. multiple worker threads to share a single processing unit.

mode_allow_dynamic_pools = 2

Allow worker threads to be added and removed from thread pools.

enum scheduling_policy

This enumeration lists the available scheduling policies (or schedulers) when creating thread pools.

Values:

user_defined = -2

unspecified = -1

local = 0

local_priority_fifo = 1

local_priority_lifo = 2

static_ = 3

static_priority = 4

abp_priority_fifo = 5

```
abp_priority_lifo = 6
shared_priority = 7
```

Functions

```
detail::partitioner &get_partitioner()
```

May be used anywhere in code and returns a reference to the single, global resource partitioner.

```
bool is_partitioner_valid()
```

Returns true if the resource partitioner has been initialized. Returns false otherwise.

plugin

```
#include <compatibility/hpx/util/plugin.hpp>

#include <compatibility/hpx/util/plugin/plugin_factory.hpp>

#include <compatibility/hpx/util/plugin/concrete_factory.hpp>

#include <compatibility/hpx/util/plugin/abstract_factory.hpp>

#include <compatibility/hpx/util/plugin/config.hpp>

#include <compatibility/hpx/util/plugin/dll.hpp>

#include <compatibility/hpx/util/plugin/plugin_wrapper.hpp>

#include <compatibility/hpx/util/plugin/export_plugin.hpp>

#include <compatibility/hpx/util/plugin/virtual_constructor.hpp>

#include <compatibility/hpx/traits/plugin_config_data.hpp>

#include <hpx/plugin.hpp>

#include <hpx/plugin/plugin_factory.hpp>

namespace hpx

    namespace util

        namespace plugin

            template<class BasePlugin>
            struct plugin_factory : public hpx::util::plugin::detail::plugin_factory_item<BasePlugin, detail::plugin_fa
```

Public Functions

```
plugin_factory(dll &d, std::string const &basename)
```

Private Types

```
template<>
using base_type = detail::plugin_factory_item<BasePlugin, detail::plugin_factory_item_base, typename v

template<class BasePlugin>
struct static_plugin_factory : public hpx::util::plugin::detail::static_plugin_factory_item<BasePlugin,
```

Public Functions

```
static_plugin_factory(get_plugins_list_type const &f)
```

Private Types

```
template<>
using base_type = detail::static_plugin_factory_item<BasePlugin, detail::static_plugin_factory_item_base, t
```

```
#include <hpx/plugin/concrete_factory.hpp>
```

```
#include <hpx/plugin/abstract_factory.hpp>
```

```
#include <hpx/plugin/config.hpp>
```

Defines

```
HPX_PLUGIN_EXPORT_API
HPX_PLUGIN_API
HPX_PLUGIN_ARGUMENT_LIMIT
HPX_PLUGIN_SYMBOLS_PREFIX_DYNAMIC
HPX_PLUGIN_SYMBOLS_PREFIX
HPX_PLUGIN_SYMBOLS_PREFIX_DYNAMIC_STR
HPX_PLUGIN_SYMBOLS_PREFIX_STR
namespace hpx

namespace util

namespace plugin
```

TypeDefs

```
using shared_ptr = boost::shared_ptr<T>
```

```
#include <hpx/plugin/dll.hpp>
```

Defines

```
HPX_HAS_DLOPEN
```

```
#include <hpx/plugin/plugin_wrapper.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace plugin
```

```
template<typename Wrapped, typename ...Parameters>
struct plugin_wrapper : public hpx::util::plugin::detail::dll_handle_holder, public Wrapped
```

Public Functions

```
plugin_wrapper(dll_handle dll, Parameters... parameters)
```

```
#include <hpx/plugin/export_plugin.hpp>
```

Defines

```
HPX_PLUGIN_NAME_2(name1, name2)
```

```
HPX_PLUGIN_NAME_3(name, base, cname)
```

```
HPX_PLUGIN_LIST_NAME_(prefix, name, base)
```

```
HPX_PLUGIN_EXPORTER_NAME_(prefix, name, base, cname)
```

```
HPX_PLUGIN_EXPORTER_INSTANCE_NAME_(prefix, name, base, cname)
```

```
HPX_PLUGIN_FORCE_LOAD_NAME_(prefix, name, base)
```

```
HPX_PLUGIN_LIST_NAME_(name, base)
```

```
HPX_PLUGIN_EXPORTER_NAME_(name, base, cname)
```

```
HPX_PLUGIN_EXPORTER_INSTANCE_NAME_(name, base, cname)
```

```
HPX_PLUGIN_FORCE_LOAD_NAME_(name, base)
```

```
HPX_PLUGIN_LIST_NAME_DYNAMIC_(name, base)
```

```
HPX_PLUGIN_EXPORTER_NAME_DYNAMIC_(name, base, cname)
```

```
HPX_PLUGIN_EXPORTER_INSTANCE_NAME_DYNAMIC_(name, base, cname)
```

```
HPX_PLUGIN_FORCE_LOAD_NAME_DYNAMIC(name, base)
HPX_PLUGIN_EXPORT_(prefix, name, BaseType, ActualType, actualname, classname)
HPX_PLUGIN_EXPORT(name, BaseType, ActualType, actualname, classname)
HPX_PLUGIN_EXPORT_DYNAMIC(name, BaseType, ActualType, actualname, classname)
HPX_PLUGIN_EXPORT_LIST_(prefix, name, classname)
HPX_PLUGIN_EXPORT_LIST(name, classname)
HPX_PLUGIN_EXPORT_LIST_DYNAMIC(name, classname)
```

```
#include <hpx/plugin/virtual_constructor.hpp>
```

```
namespace hpx
```

```
    namespace util
```

```
        namespace plugin
```

Typedefs

```
using exported_plugins_type = std::map<std::string, hpx::util::any_nonser>
typedef exported_plugins_type* HPX_PLUGIN_API* hpx::util::plugin::get_plugins_list_np
typedef exported_plugins_type* HPX_PLUGIN_API hpx::util::plugin::get_plugins_list_np
using dll_handle = shared_ptr<get_plugins_list_np>

template<typename BasePlugin>
struct virtual_constructor
```

Public Types

```
template<>
using type = hpx::util::pack<>
```

```
#include <hpx/plugin/traits/plugin_config_data.hpp>
```

```
namespace hpx
```

```
    namespace traits
```

```
template<typename Plugin, typename Enable = void>
struct plugin_config_data
```

Public Static Functions

```
static char const *call()
```

checkpoint

```
#include <compatibility/hpx/util/checkpoint.hpp>
```

This header defines the save_checkpoint and restore_checkpoint functions. These functions are designed to help HPX application developer's checkpoint their applications. Save_checkpoint serializes one or more objects and saves them as a byte stream. Restore_checkpoint converts the byte stream back into instances of the objects.

```
#include <hpx/checkpoint/checkpoint.hpp>
```

```
namespace hpx
```

```
namespace util
```

Functions

```
std::ostream &operator<< (std::ostream &ost, checkpoint const &ckp)  
Operator<< Overload
```

This overload is the main way to write data from a checkpoint to an object such as a file. Inside the function, the size of the checkpoint will be written to the stream before the checkpoint's data. The operator>> overload uses this to read the correct number of bytes. Be mindful of this additional write and read when you use different facilities to write out or read in data to a checkpoint!

Parameters

- ost: Output stream to write to.
- ckp: Checkpoint to copy from.

Return Operator<< returns the ostream object.

```
std::istream &operator>> (std::istream &ist, checkpoint &ckp)  
Operator>> Overload
```

This overload is the main way to read in data from an object such as a file to a checkpoint. It is important to note that inside the function, the first variable to be read is the size of the checkpoint. This size variable is written to the stream before the checkpoint's data in the operator<< overload. Be mindful of this additional read and write when you use different facilities to read in or write out data from a checkpoint!

Parameters

- ist: Input stream to write from.
- ckp: Checkpoint to write to.

Return Operator>> returns the ostream object.

```
template<typename T, typename ...Ts, typename U = typename std::enable_if<!hpx::traits::is_launch_policy<T>::value &  
hpx::future<checkpoint> save_checkpoint (T &&t, Ts&&... ts)  
Save_checkpoint
```

Save_checkpoint takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a shared_ptr to the

component or by passing a component's client instance to save_checkpoint. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used save_checkpoint will simply return a checkpoint object.

Template Parameters

- T: Containers passed to save_checkpoint to be serialized and placed into a checkpoint object.
- Ts: More containers passed to save_checkpoint to be serialized and placed into a checkpoint object.
- U: This parameter is used to make sure that T is not a launch policy or a checkpoint. This forces the compiler to choose the correct overload.

Parameters

- t: A container to restore.
- ts: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return Save_checkpoint returns a future to a checkpoint with one exception: if you pass hpx::launch::sync as the first argument. In this case save_checkpoint will simply return a checkpoint.

template<typename T, typename ...Ts>

`hpx::future<checkpoint> save_checkpoint (checkpoint &&c, T &&t, Ts&&... ts)`

Save_checkpoint - Take a pre-initialized checkpoint

Save_checkpoint takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a shared_ptr to the component or by passing a component's client instance to save_checkpoint. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used save_checkpoint will simply return a checkpoint object.

Template Parameters

- T: Containers passed to save_checkpoint to be serialized and placed into a checkpoint object.
- Ts: More containers passed to save_checkpoint to be serialized and placed into a checkpoint object.

Parameters

- c: Takes a pre-initialized checkpoint to copy data into.
- t: A container to restore.
- ts: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return Save_checkpoint returns a future to a checkpoint with one exception: if you pass hpx::launch::sync as the first argument. In this case save_checkpoint will simply return a checkpoint.

template<typename T, typename ...Ts>

`hpx::future<checkpoint> save_checkpoint (hpx::launch p, T &&t, Ts&&... ts)`

Save_checkpoint - Policy overload

Save_checkpoint takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a shared_ptr to the component or by passing a component's client instance to save_checkpoint. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used save_checkpoint will simply return a checkpoint object.

Template Parameters

- T: Containers passed to save_checkpoint to be serialized and placed into a checkpoint object.
- Ts: More containers passed to save_checkpoint to be serialized and placed into a checkpoint object.

Parameters

- p: Takes an HPX launch policy. Allows the user to change the way the function is launched i.e. async, sync, etc.

- `t`: A container to restore.
- `ts`: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return `Save_checkpoint` returns a future to a checkpoint with one exception: if you pass `hpx::launch::sync` as the first argument. In this case `save_checkpoint` will simply return a checkpoint.

```
template<typename T, typename ...Ts>
hpx::future<checkpoint> save_checkpoint (hpx::launch p, checkpoint &&c, T &&t, Ts&&... ts)
```

`Save_checkpoint` - Policy overload & pre-initialized checkpoint

`Save_checkpoint` takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a `shared_ptr` to the component or by passing a component's client instance to `save_checkpoint`. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used `save_checkpoint` will simply return a checkpoint object.

Template Parameters

- `T`: Containers passed to `save_checkpoint` to be serialized and placed into a checkpoint object.
- `Ts`: More containers passed to `save_checkpoint` to be serialized and placed into a checkpoint object.

Parameters

- `p`: Takes an HPX launch policy. Allows the user to change the way the function is launched i.e. `async`, `sync`, etc.
- `c`: Takes a pre-initialized checkpoint to copy data into.
- `t`: A container to restore.
- `ts`: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return `Save_checkpoint` returns a future to a checkpoint with one exception: if you pass `hpx::launch::sync` as the first argument. In this case `save_checkpoint` will simply return a checkpoint.

```
template<typename T, typename ...Ts, typename U = typename std::enable_if<!std::is_same<typename std::decay<T> checkpoint save_checkpoint (hpx::launch::sync_policy sync_p, T &&t, Ts&&... ts)
```

`Save_checkpoint` - Sync_policy overload

`Save_checkpoint` takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a `shared_ptr` to the component or by passing a component's client instance to `save_checkpoint`. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used `save_checkpoint` will simply return a checkpoint object.

Template Parameters

- `T`: Containers passed to `save_checkpoint` to be serialized and placed into a checkpoint object.
- `Ts`: More containers passed to `save_checkpoint` to be serialized and placed into a checkpoint object.
- `U`: This parameter is used to make sure that `T` is not a checkpoint. This forces the compiler to choose the correct overload.

Parameters

- `sync_p`: `hpx::launch::sync_policy`
- `t`: A container to restore.
- `ts`: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return `Save_checkpoint` which is passed `hpx::launch::sync_policy` will return a checkpoint which contains the serialized values checkpoint.

```
template<typename T, typename ...Ts>
checkpoint save_checkpoint (hpx::launch::sync_policy sync_p, checkpoint &&c, T &&t,
                           Ts&&... ts)
```

Save_checkpoint - Sync_policy overload & pre-init. checkpoint

Save_checkpoint takes any number of objects which a user may wish to store and returns a future to a checkpoint object. This function can also store a component either by passing a shared_ptr to the component or by passing a component's client instance to save_checkpoint. Additionally the function can take a policy as a first object which changes its behavior depending on the policy passed to it. Most notably, if a sync policy is used save_checkpoint will simply return a checkpoint object.

Template Parameters

- T: Containers passed to save_checkpoint to be serialized and placed into a checkpoint object.
- Ts: More containers passed to save_checkpoint to be serialized and placed into a checkpoint object.

Parameters

- sync_p: hpx::launch::sync_policy
- c: Takes a pre-initialized checkpoint to copy data into.
- t: A container to restore.
- ts: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

Return Save_checkpoint which is passed hpx::launch::sync_policy will return a checkpoint which contains the serialized values checkpoint.

```
template<typename T, typename ...Ts>
void restore_checkpoint (checkpoint const &c, T &t, Ts&... ts)
```

Restore_checkpoint

Restore_checkpoint takes a checkpoint object as a first argument and the containers which will be filled from the byte stream (in the same order as they were placed in save_checkpoint). Restore_checkpoint can resurrect a stored component in two ways: by passing in a instance of a component's shared_ptr or by passing in an instance of the component's client.

Return Restore_checkpoint returns void.

Template Parameters

- T: A container to restore.
- Ts: Other containers to restore. Containers must be in the same order that they were inserted into the checkpoint.

Parameters

- c: The checkpoint to restore.
- t: A container to restore.
- ts: Other containers to restore Containers must be in the same order that they were inserted into the checkpoint.

```
class checkpoint
#include <checkpoint.hpp> Checkpoint Object
```

Checkpoint is the container object which is produced by save_checkpoint and is consumed by a restore_checkpoint. A checkpoint may be moved into the save_checkpoint object to write the byte stream to the pre-created checkpoint object.

Checkpoints are able to store all containers which are able to be serialized including components.

Public Types

```
using const_iterator = std::vector::const_iterator
```

Public Functions

```
checkpoint()
checkpoint(checkpoint const &c)
checkpoint(checkpoint &&c)
~checkpoint()
checkpoint(std::vector<char> const &vec)
checkpoint(std::vector<char> &&vec)
checkpoint &operator=(checkpoint const &c)
checkpoint &operator=(checkpoint &&c)
const_iterator begin() const
const_iterator end() const
size_t size() const
```

Private Functions

```
template<typename Archive>
void serialize(Archive &arch, const unsigned int version)
```

Private Members

```
std::vector<char> data_
```

Friends

```
friend hpx::util::hpx::serialization::access
std::ostream &operator<< (std::ostream &ost, checkpoint const &ckp)
    Operator<< Overload
```

This overload is the main way to write data from a checkpoint to an object such as a file. Inside the function, the size of the checkpoint will be written to the stream before the checkpoint's data. The operator>> overload uses this to read the correct number of bytes. Be mindful of this additional write and read when you use different facilities to write out or read in data to a checkpoint!

Parameters

- ost: Output stream to write to.
- ckp: Checkpoint to copy from.

Return Operator<< returns the ostream object.

```
std::istream &operator>> (std::istream &ist, checkpoint &ckp)
    Operator>> Overload
```

This overload is the main way to read in data from an object such as a file to a checkpoint. It is important to note that inside the function, the first variable to be read is the size of the checkpoint. This size variable is written to the stream before the checkpoint's data in the operator<< overload. Be mindful of this additional read and write when you use different facilities to read in or write out data from a checkpoint!

Parameters

- `ist`: Input stream to write from.
- `ckp`: Checkpoint to write to.

Return Operator>> returns the ostream object.

```
template<typename T, typename ...Ts>
void restore_checkpoint (checkpoint const &c, T &t, Ts&... ts)
    Restore_checkpoint
```

`Restore_checkpoint` takes a checkpoint object as a first argument and the containers which will be filled from the byte stream (in the same order as they were placed in `save_checkpoint`). `Restore_checkpoint` can resurrect a stored component in two ways: by passing in a instance of a component's `shared_ptr` or by passing in an instance of the component's client.

Return `Restore_checkpoint` returns void.

Template Parameters

- `T`: A container to restore.
- `Ts`: Other containers to restore. Containers must be in the same order that they were inserted into the checkpoint.

Parameters

- `c`: The checkpoint to restore.
- `t`: A container to restore.
- `ts`: Other containers to restore. Containers must be in the same order that they were inserted into the checkpoint.

```
bool operator==(checkpoint const &lhs, checkpoint const &rhs)
```

```
bool operator!=(checkpoint const &lhs, checkpoint const &rhs)
```

2.10 Contributing to HPX

HPX development happens on Github. The following sections are a collection of useful information related to *HPX* development.

2.10.1 Release procedure for HPX

Below is a step by step procedure for making an *HPX* release. We aim to produce two releases per year: one in March-April, and one in September-October.

This is a living document and may not be totally current or accurate. It is an attempt to capture current practices in making an *HPX* release. Please update it as appropriate.

One way to use this procedure is to print a copy and check off the lines as they are completed to avoid confusion.

1. Notify developers that a release is imminent.

2. Make a list of examples and benchmarks that should not go into the release. Build all examples and benchmarks that will go in the release and make sure they build and run as expected.
 - Make sure all examples and benchmarks have example input files, and usage documentation, either in the form of comments or a readme.
3. Send the list of examples and benchmarks that will be included in the release to hpx-users@stellar.cct.lsu.edu and stellar@cct.lsu.edu, and ask for feedback. Update the list as necessary.
4. Write release notes in `docs/sphinx/releases/whats_new_$.VERSION.rst`. Keep adding merged PRs and closed issues to this until just before the release is made. Use `tools/generate_pr_issue_list.sh` to generate the lists. Add the new release notes to the table of contents in `docs/sphinx/releases.rst`.
5. Build the docs, and proof-read them. Update any documentation that may have changed, and correct any typos. Pay special attention to:
 - `$HPX_SOURCE/README.rst`
 - Update grant information
 - `docs/sphinx/releases/whats_new_$.VERSION.rst`
 - `docs/sphinx/about_hpx/people.rst`
 - Update collaborators
 - Update grant information
6. This step does not apply to patch releases. For both APEX and hpxMP:
 - Change the release branch to be the most current release tag available in the APEX/hpxMP `git_external` section in the main `CMakeLists.txt`. Please contact the maintainers of the respective packages to generate a new release to synchronize with the HPX release ([APEX²³⁰](#), [hpxMP²³¹](#)).
7. If there have been any commits to the release branch since the last release, create a tag from the old release branch before deleting the old release branch in the next step.
8. Unprotect the release branch in the github repository settings so that it can be deleted and recreated.
9. Delete the old release branch, and create a new one by branching a stable point from master. If you are creating a patch release, branch from the release tag for which you want to create a patch release.
 - `git push origin --delete release`
 - `git branch -D release`
 - `git checkout [stable point in master]`
 - `git branch release`
 - `git push origin release`
 - `git branch --set-upstream-to=origin/release release`
10. Protect the release branch again to disable deleting and force pushes.
11. Check out the release branch.
12. Make sure `HPX_VERSION_MAJOR/MINOR/SUBMINOR` in `CMakeLists.txt` contain the correct values. Change them if needed.

²³⁰ <http://github.com/khuck/xpress-apex>

²³¹ <https://github.com/STELLAR-GROUP/hpxMP>

13. This step does not apply to patch releases. Remove features which have been deprecated for at least 2 releases. This involves removing build options which enable those features from the main CMakeLists.txt and also deleting all related code and tests from the main source tree.

The general deprecation policy involves a three-step process we have to go through in order to introduce a breaking change:

- a. First release cycle: add a build option that allows for explicitly disabling any old (now deprecated) code.
- b. Second release cycle: turn this build option OFF by default.
- c. Third release cycle: completely remove the old code.

The main CMakeLists.txt contains a comment indicating for which version the breaking change was introduced first.

14. Switch Buildbot over to test the release branch.

- `https://github.com/STELLAR-GROUP/hermione-buildbot/blob/rostam/master/master.cfg`
- `branch field in c['change_source'] = GitPoller`

15. Repeat the following steps until satisfied with the release.

1. Change `HPX_VERSION_TAG` in CMakeLists.txt to `-rcN`, where N is the current iteration of this step. Start with `-rc1`.
2. Tag and create a pre-release on GitHub using the script `tools/roll_release.sh`. The script requires that you have the STELLAR Group signing key.
3. This step is not necessary for patch releases. Notify `hpx-users@stellar.cct.lsu.edu` and `stellar@cct.lsu.edu` of the availability of the release candidate. Ask users to test the candidate by checking out the release candidate tag.
4. Allow at least a week for testing of the release candidate.
 - Use `git merge` when possible, and fall back to `git cherry-pick` when needed. For patch releases `git cherry-pick` is most likely your only choice if there have been significant unrelated changes on master since the previous release.
 - Go back to the first step when enough patches have been added.
 - If there are no more patches, continue to make the final release.
16. Update any occurrences of the latest stable release to refer to the version about to be released. For example, `quickstart.rst` contains instructions to check out the latest stable tag. Make sure that refers to the new version.
17. Add a new entry to the RPM changelog (`cmake/packaging/rpm/Changelog.txt`) with the new version number and a link to the corresponding changelog.
18. Change `HPX_VERSION_TAG` in CMakeLists.txt to an empty string.
19. Add the release date to the caption of the current “What’s New” section in the docs, and change the value of `HPX_VERSION_DATE` in CMakeLists.txt.
20. Tag and create a release on GitHub using the script `tools/roll_release.sh`. The script requires that you have the STELLAR Group signing key.
21. Update the websites (stellar-group.org²³² and stellar.cct.lsu.edu²³³) with the following:
 - Download links on the downloads pages. Link to the release on GitHub.

²³² <https://stellar-group.org>

²³³ <https://stellar.cct.lsu.edu>

- Documentation links on the docs page (link to generated documentation on GitHub Pages). Follow the style of previous releases.
- A new blog post announcing the release, which links to downloads and the “What’s New” section in the documentation (see previous releases for examples).

22. Merge release branch into master.

23. Post-release cleanup. Create a new pull request against master with the following changes:

1. Modify the release procedure if necessary.
2. Change HPX_VERSION_TAG in CMakeLists.txt back to -trunk.

24. Switch Buildbot back to test the main branch

- `https://github.com/STELLAR-GROUP/hermione-buildbot/blob/rostam/master/master.cfg`
- branch field in `c['change_source'] = GitPoller`

25. Update Vcpkg (`https://github.com/Microsoft/vcpkg`) to pull from latest release.

- Update version number in CONTROL
- Update tag and SHA512 to that of the new release

26. Announce the release on `hpx-users@stellar.cct.lsu.edu`, `stellar@cct.lsu.edu`, `allcct@cct.lsu.edu`, `faculty@csc.lsu.edu`, `faculty@ece.lsu.edu`, `xpress@crest.iu.edu`, the *HPX* Slack channel, the IRC channel, Sonia Sachs, our list of external collaborators, `isocpp.org`, `reddit.com`, HPC Wire, Inside HPC, Heise Online, and a CCT press release.

27. Beer and pizza.

2.10.2 Testing *HPX*

To ensure correctness of *HPX*, we ship a large variety of unit and regression tests. The tests are driven by the CTest²³⁴ tool and are executed automatically by buildbot (see *HPX* Buildbot Website²³⁵) on each commit to the *HPX* Github²³⁶ repository. In addition, it is encouraged to run the test suite manually to ensure proper operation on your target system. If a test fails for your platform, we highly recommend submitting an issue on our *HPX* Issues²³⁷ tracker with detailed information about the target system.

Running tests manually

Running the tests manually is as easy as typing `make tests && make test`. This will build all tests and run them once the tests are built successfully. After the tests have been built, you can invoke separate tests with the help of the `ctest` command. You can list all available test targets using `make help | grep tests`. Please see the CTest Documentation²³⁸ for further details.

²³⁴ <https://gitlab.kitware.com/cmake/community/wikis/doc/ctest/Testing-With-CTest>

²³⁵ <http://rostam.cct.lsu.edu/>

²³⁶ <https://github.com/STELLAR-GROUP/hpx/>

²³⁷ <https://github.com/STELLAR-GROUP/hpx/issues>

²³⁸ <https://www.cmake.org/cmake/help/latest/manual/ctest.1.html>

Issue tracker

If you stumble over a bug or missing feature in *HPX*, please submit an issue to our [HPX Issues²³⁹](#) page. For more information on how to submit support requests or other means of getting in contact with the developers, please see the [Support Website²⁴⁰](#) page.

Continuous testing

In addition to manual testing, we run automated tests on various platforms. You can see the status of the current master head by visiting the [HPX Buildbot Website²⁴¹](#). We also run tests on all pull requests using both [CircleCI²⁴²](#) and a combination of [CDash²⁴³](#) and [pycicle²⁴⁴](#). You can see the dashboards here: [CircleCI HPX dashboard²⁴⁵](#) and [CDash HPX dashboard²⁴⁶](#).

2.10.3 Using docker for development

Although it can often be useful to set up a local development environment with system-provided or self-built dependencies, [Docker²⁴⁷](#) provides a convenient alternative to quickly get all the dependencies needed to start development of *HPX*. Our testing setup on [CircleCI²⁴⁸](#) uses a docker image to run all tests.

To get started you need to install [Docker²⁴⁹](#) using whatever means is most convenient on your system. Once you have [Docker²⁵⁰](#) installed, you can pull or directly run the docker image. The image is based on Debian and Clang, and can be found on [Docker Hub²⁵¹](#). To start a container using the *HPX* build environment, run:

```
docker run --interactive --tty stellargroup/build_env:ubuntu bash
```

You are now in an environment where all the *HPX* build and runtime dependencies are present. You can install additional packages according to your own needs. Please see the [Docker Documentation²⁵²](#) for more information on using [Docker²⁵³](#).

Warning: All changes made within the container are lost when the container is closed. If you want files to persist (e.g., the *HPX* source tree) after closing the container, you can bind directories from the host system into the container (see [Docker Documentation \(Bind mounts\)²⁵⁴](#)).

²³⁹ <https://github.com/STELLAR-GROUP/hpx/issues>

²⁴⁰ <https://stellar.cct.lsu.edu/support/>

²⁴¹ <http://rostam.cct.lsu.edu/>

²⁴² <https://circleci.com>

²⁴³ <https://www.kitware.com/cdash/project/about.html>

²⁴⁴ <https://github.com/biddisco/pycicle/>

²⁴⁵ <https://circleci.com/gh/STELLAR-GROUP/hpx>

²⁴⁶ <https://cdash.cscs.ch/index.php?project=HPX>

²⁴⁷ <https://www.docker.com>

²⁴⁸ <https://circleci.com>

²⁴⁹ <https://www.docker.com>

²⁵⁰ <https://www.docker.com>

²⁵¹ https://hub.docker.com/r/stellargroup/build_env/

²⁵² <https://docs.docker.com/>

²⁵³ <https://www.docker.com>

²⁵⁴ <https://docs.docker.com/storage/bind-mounts/>

2.10.4 Documentation

This documentation is built using Sphinx²⁵⁵, and an automatically generated API reference using Doxygen²⁵⁶ and Breathe²⁵⁷.

We always welcome suggestions on how to improve our documentation, as well as pull requests with corrections and additions.

Building documentation

Please see the *documentation prerequisites* section for details on what you need in order to build the HPX documentation. Enable building of the documentation by setting `HPX_WITH_DOCUMENTATION=ON` during CMake²⁵⁸ configuration. To build the documentation, build the `docs` target using your build tool. The default output format is HTML documentation. You can choose alternative output formats (single-page HTML, PDF, and man) with the `HPX_WITH_DOCUMENTATION_OUTPUT_FORMATS` CMake option.

Note: If you add new source files to the Sphinx documentation, you have to run CMake again to have the files included in the build.

Style guide

The documentation is written using reStructuredText. These are the conventions used for formatting the documentation:

- Use, at most, 80 characters per line.
- Top-level headings use over- and underlines with =.
- Sub-headings use only underlines with characters in decreasing level of importance: =, – and ..
- Use sentence case in headings.
- Refer to common terminology using :term:`Component`.
- Indent content of directives (.. directive::) by three spaces.
- For C++ code samples at the end of paragraphs, use :: and indent the code sample by 4 spaces.
 - For other languages (or if you don't want a colon at the end of the paragraph), use .. code-block:: language and indent by three spaces as with other directives.
- Use .. list-table:: to wrap tables with a lot of text in cells.

²⁵⁵ <http://www.sphinx-doc.org>

²⁵⁶ <https://www.doxygen.org>

²⁵⁷ <https://breathe.readthedocs.io/en/latest>

²⁵⁸ <https://www.cmake.org>

API documentation

The source code is documented using Doxygen. If you add new API documentation either to existing or new source files, make sure that you add the documented source files to the `doxygen_dependencies` variable in `docs/CMakeLists.txt`.

2.10.5 Module structure

This section explains the structure of an *HPX* module.

The tool `create_library_skeleton.py`²⁵⁹ can be used to generate a basic skeleton. To create a library skeleton, run the tool in the `libs` subdirectory with the module name as an argument:

```
./create_library_skeleton <lib_name>
```

This creates a skeleton with the necessary files for an *HPX* module. It will not create any actual source files. The structure of this skeleton is as follows:

- <lib_name>/
 - README.rst
 - CMakeLists.txt
 - cmake
 - docs/
 - * index.rst
 - examples/
 - * CMakeLists.txt
 - include/
 - * hpx/
 - <lib_name>
 - src/
 - * CMakeLists.txt
 - tests/
 - * CMakeLists.txt
 - * unit/
 - CMakeLists.txt
 - * regressions/
 - CMakeLists.txt
 - * performance/
 - CMakeLists.txt

²⁵⁹ https://github.com/STELLAR-GROUP/hpx/blob/master/libs/create_library_skeleton.py

A README.rst should be always included which explains the basic purpose of the library and a link to the generated documentation.

A main CMakeLists.txt is created in the root directory of the module. By default it contains a call to add_hpx_module which takes care of most of the boilerplate required for a module. You only need to fill in the source and header files in most cases.

add_hpx_module requires a module name. Optional flags are:

- DEPRECATION_WARNINGS: Enables deprecation warnings for the module.

Optional single-value arguments are:

- COMPATIBILITY_HEADERS: Can be ON, OFF, or left out. Enables compatibility headers. Creates a variable which can be turned on or off by the user when set to ON or OFF. If left out the option is completely disabled.
- INSTALL_BINARIES: Install the resulting library.

Optional multi-value arguments are:

- SOURCES: List of source files.
- HEADERS: List of header files.
- COMPAT_HEADERS: List of compatibility header files.
- DEPENDENCIES: Libraries that this module depends on, such as other modules.
- CMAKE_SUBDIRS: List of subdirectories to add to the module.

The include directory should contain only headers that other libraries need. For each of those headers, an automatic header test to check for self containment will be generated. Private headers should be placed under the src directory. This allows for clear separation. The cmake subdirectory may include additional CMake²⁶⁰ scripts needed to generate the respective build configurations.

Compatibility headers (forwarding headers for headers whose location is changed when creating a module, if moving them from the main library) should be placed in an include_compatibility directory. This directory is not created by default.

Documentation is placed in the docs folder. A empty skeleton for the index is created, which is picked up by the main build system and will be part of the generated documentation. Each header inside the include directory will automatically be processed by Doxygen and included into the documentation. If a header should be excluded from the API reference, a comment // sphinx:undocumented needs to be added.

Tests are placed in suitable subdirectories of tests.

When in doubt, consult existing modules for examples on how to structure the module.

2.11 Releases

2.11.1 HPX V1.4.0 (January 15, 2020)

General changes

- We have added the collectives all_to_all and all_reduce.
- We have added APIs for resiliency, which allows replication and replay for failed tasks. See the documentation for more details.
- Components can now be checkpointed.

²⁶⁰ <https://www.cmake.org>

- Performance improvements to schedulers and coroutines. A significant change is the addition of stackless coroutines. These are to be used for tasks that do not need to be suspended and can reduce overheads noticeably in applications with short tasks. A stackless coroutine can be created with the new stack size `thread_stacksize_nostack`.
- We have added an implementation of `unique_any`, which is a non-copyable version of `any`.
- The `shared_priority_queue_scheduler` has been improved. It now has lower overheads than the default scheduler in many situations. Unlike the default scheduler it fully supports NUMA scheduling hints. Enable it with the command line option `--hpx:queuing=shared-priority`. This scheduler should still be considered experimental, but its use is encouraged in real applications to help us make it production ready.
- We have added the performance counters `background-receive-duration` and `background-receive-overhead` for inspecting the time and overhead spent on receiving parcels in the background.
- Compilation time has been further improved when `HPX_WITH_NETWORKING=OFF`.
- We no longer require compiled Boost dependencies in certain configurations. This requires at least Boost 1.70, compiling on x86 with GCC 9, clang (libc++) 9, or VS2019 in C++17 mode. The dependency on Boost.Filesystem can explicitly be turned on with `HPX_FILESYSTEM_WITH_BOOST_FILESYSTEM_COMPATIBILITY=ON` (it is off by default if the standard library supports `std::filesystem`). Boost.ProgramOptions has been copied into the HPX repository. We have a compatibility layer for users who must explicitly use Boost.ProgramOptions instead of the ProgramOptions provided by HPX. To remove the dependency `HPX_PROGRAM_OPTIONS_WITH_BOOST_PROGRAM_OPTIONS_COMPATIBILITY` must be explicitly set to OFF. This option will be removed in a future release. We have also removed several other header-only dependencies on Boost.
- It is now possible to use the process affinity mask set by tools like `numactl` and various batch environments with the command line option `--hpx:use-process-mask`. Enabling this option implies `--hpx:ignore-batch-env`.
- It is now possible to create standalone thread pools without starting the runtime. See the `standalone_thread_pool_executor.cpp` test in the `execution` module for an example.
- Tasks annotated with `hpx::util::annotated_function` now have their correct name when using APEX to generate OTF2 files.
- Cloning of APEX was defective in previous releases (it required manual intervention to check out the correct tag or branch). This has been fixed.
- The option `HPX_WITH_MORE_THAN_64_THREADS` is now ignored and will be removed in a future release. The value is instead derived directly from `HPX_WITH_MAX_CPU_COUNT` option.
- We have deprecated compiling in C++11 mode. The next release will require a C++14 capable compiler.
- We have deprecated support for the Vc library. This option will be replaced with SIMD support from the standard library in a future release.
- We have significantly refactored our CMake setup. This is intended to be a non-breaking change and will allow for using HPX through CMake targets in the future.
- We have continued modularizing the HPX library. In the process we have rearranged many header files into module-specific directories. All moved headers have compatibility headers which forward from the old location to the new location, together with a deprecation warning. The compatibility headers will eventually be removed.
- We now enforce formatting with `clang-format` on the majority of our source files.
- We have added SPDX license tags to all files.
- Many bugfixes.

Breaking changes

- The `HPX_WITH_THREAD_COMPATIBILITY` option and the associated compatibility layer has been removed.
- The `HPX_WITH_INCLUSIVE_SCAN_COMPATIBILITY` option and the associated compatibility layer has been removed.
- The `HPX_WITH_UNWRAPPED_COMPATIBILITY` option and the associated compatibility layer has been removed.

Closed issues

- Issue #4282²⁶¹ - Build Issues with Release on Windows
- Issue #4278²⁶² - Build Issues with CMake 3.14.4
- Issue #4273²⁶³ - Clients of HPX 1.4.0-rc2 with APEX are not linked to `libhpx-apex`
- Issue #4269²⁶⁴ - Building HPX 1.4.0-rc2 with support for APEX fails
- Issue #4263²⁶⁵ - Compilation fail on latest master
- Issue #4232²⁶⁶ - Configure of HPX project using CMake FetchContent fails
- Issue #4223²⁶⁷ - “Re-using the `main()` function as the main HPX entry point” doesn’t work
- Issue #4220²⁶⁸ - HPX won’t compile - error building `resource_partitioner`
- Issue #4215²⁶⁹ - HPX 1.4.0rc1 does not link on s390x
- Issue #4204²⁷⁰ - Trouble compiling HPX with Intel compiler
- Issue #4199²⁷¹ - Refactor APEX to eliminate circular dependency
- Issue #4187²⁷² - HPX can’t build on OSX
- Issue #4185²⁷³ - Simple debug output for development
- Issue #4182²⁷⁴ - `@HPX_CONF_PREFIX@` is the empty string
- Issue #4169²⁷⁵ - HPX won’t build with APEX
- Issue #4163²⁷⁶ - Add back `HPX_LIBRARIES` and `HPX_INCLUDE_DIRS`
- Issue #4161²⁷⁷ - It should be possible to call `find_package(HPX)` multiple times
- Issue #4155²⁷⁸ - `get_self_id()` for stackless threads returns `invalid_thread_id`

²⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/4282>

²⁶² <https://github.com/STELLAR-GROUP/hpx/issues/4278>

²⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/4273>

²⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/4269>

²⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/4263>

²⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/4232>

²⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/4223>

²⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/4220>

²⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/4215>

²⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/4204>

²⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/4199>

²⁷² <https://github.com/STELLAR-GROUP/hpx/issues/4187>

²⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/4185>

²⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/4182>

²⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/4169>

²⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/4163>

²⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/4161>

²⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/4155>

- Issue #4151²⁷⁹ - build error with MPI code
- Issue #4150²⁸⁰ - hpx won't build on POWER9 with clang 8
- Issue #4148²⁸¹ - `cacheline_data` delivers poor performance with C++17 compared to C++14
- Issue #4144²⁸² - target general in `HPX_LIBRARIES` does not exist
- Issue #4134²⁸³ - CMake Error when `-DHPX_WITH_HPXMP=ON`
- Issue #4132²⁸⁴ - parallel fill leaves elements unfilled
- Issue #4123²⁸⁵ - PAPI performance counters are inaccessible
- Issue #4118²⁸⁶ - `static_chunk_size` is not obeyed in scan algorithms
- Issue #4115²⁸⁷ - dependency chaining error with APEX
- Issue #4107²⁸⁸ - Initializing runtime without entry point function and command line arguments
- Issue #4105²⁸⁹ - Bug in `hpx::bind=numa-balanced`
- Issue #4101²⁹⁰ - Bound tasks
- Issue #4100²⁹¹ - Add SPDX identifier to all files
- Issue #4085²⁹² - `hpx_topology` library should depend on `hwloc`
- Issue #4067²⁹³ - HPX fails to build on macOS
- Issue #4056²⁹⁴ - Building without thread manager idle backoff fails
- Issue #4052²⁹⁵ - Enforce `clang-format` style for modules
- Issue #4032²⁹⁶ - Simple hello world fails to launch correctly
- Issue #4030²⁹⁷ - Allow threads to skip context switching
- Issue #4029²⁹⁸ - Add support for `mimalloc`
- Issue #4005²⁹⁹ - Can't link HPX when APEX enabled
- Issue #4002³⁰⁰ - Missing header for algorithm module
- Issue #3989³⁰¹ - conversion from `long` to `unsigned int` requires a narrowing conversion on MSVC

²⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/4151>

²⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/4150>

²⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/4148>

²⁸² <https://github.com/STELLAR-GROUP/hpx/issues/4144>

²⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/4134>

²⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/4132>

²⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/4123>

²⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/4118>

²⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/4115>

²⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/4107>

²⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/4105>

²⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/4101>

²⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/4100>

²⁹² <https://github.com/STELLAR-GROUP/hpx/issues/4085>

²⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/4067>

²⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/4056>

²⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/4052>

²⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/4032>

²⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/4030>

²⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/4029>

²⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/4005>

³⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/4002>

³⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/3989>

- Issue #3958³⁰² - /statistics/average@ perf counter can't be created
- Issue #3953³⁰³ - CMake errors from HPX_AddPseudoDependencies
- Issue #3941³⁰⁴ - CMake error for APEX install target
- Issue #3940³⁰⁵ - Convert pseudo-doxygen function documentation into actual doxygen documentation
- Issue #3935³⁰⁶ - HPX compiler match too strict?
- Issue #3929³⁰⁷ - Buildbot failures on latest HPX stable
- Issue #3912³⁰⁸ - I recommend publishing a version that does not depend on the boost library
- Issue #3890³⁰⁹ - hpx.ini not working
- Issue #3883³¹⁰ - cuda compilation fails because of -faligned-new
- Issue #3879³¹¹ - HPX fails to configure with -DHPX_WITH_TESTS=OFF
- Issue #3871³¹² - dataflow does not support void allocators
- Issue #3867³¹³ - Latest HTML docs placed in wrong directory on GitHub pages
- Issue #3866³¹⁴ - Make sure all tests use HPX_TEST* macros and not HPX_ASSERT
- Issue #3857³¹⁵ - CMake all-keyword or all-plain for target_link_libraries
- Issue #3856³¹⁶ - hpx_setup_target adds rogue flags
- Issue #3850³¹⁷ - HPX fails to build on POWER8 with Clang7
- Issue #3848³¹⁸ - Remove lva member from thread_init_data
- Issue #3838³¹⁹ - hpx::parallel::count/count_if failing tests
- Issue #3651³²⁰ - hpx::parallel::transform_reduce with non const reference as lambda parameter
- Issue #3560³²¹ - Apex integration with HPX not working properly
- Issue #3322³²² - No warning when mixing debug/release builds

³⁰² <https://github.com/STELLAR-GROUP/hpx/issues/3958>

³⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/3953>

³⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3941>

³⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3940>

306 <https://github.com/STELLAR-GROUP/hpx/issues/3935>307 <https://github.com/STELLAR-GROUP/hpx/issues/3929>308 <https://github.com/STELLAR-GROUP/hpx/issues/3912>309 <https://github.com/STELLAR-GROUP/hpx/issues/3890>310 <https://github.com/STELLAR-GROUP/hpx/issues/3883>311 <https://github.com/STELLAR-GROUP/hpx/issues/3879>312 <https://github.com/STELLAR-GROUP/hpx/issues/3871>313 <https://github.com/STELLAR-GROUP/hpx/issues/3867>314 <https://github.com/STELLAR-GROUP/hpx/issues/3866>315 <https://github.com/STELLAR-GROUP/hpx/issues/3857>316 <https://github.com/STELLAR-GROUP/hpx/issues/3856>317 <https://github.com/STELLAR-GROUP/hpx/issues/3850>318 <https://github.com/STELLAR-GROUP/hpx/issues/3848>319 <https://github.com/STELLAR-GROUP/hpx/issues/3838>320 <https://github.com/STELLAR-GROUP/hpx/issues/3651>321 <https://github.com/STELLAR-GROUP/hpx/issues/3560>322 <https://github.com/STELLAR-GROUP/hpx/issues/3322>

Closed pull requests

- PR #4300³²³ - Checks for MPI_Init being called twice
- PR #4299³²⁴ - Small CMake fixes
- PR #4298³²⁵ - Remove extra call to annotate function that messes up traces
- PR #4296³²⁶ - Fixing collectives locking problem
- PR #4295³²⁷ - Do not check LICENSE_1_0.txt for inspect violations
- PR #4293³²⁸ - Applying two small changes fixing carious MSVC/Windows problems
- PR #4285³²⁹ - Delete apex.hpp
- PR #4276³³⁰ - Disable doxygen generation for hpx/debugging/print.hpp file
- PR #4275³³¹ - Make sure APEX is linked to even when not explicitly referenced
- PR #4272³³² - Fix pushing of documentation
- PR #4271³³³ - Updating APEX tag, don't create new task_wrapper on operator= of hpx_thread object
- PR #4268³³⁴ - Testing for noexcept function specializations in C++11/14 mode
- PR #4267³³⁵ - Fixing MSVC warning
- PR #4266³³⁶ - Make sure macOS Travis CI fails if build step fails
- PR #4264³³⁷ - Clean up compatibility header options
- PR #4262³³⁸ - Cleanup modules CMakeLists.txt
- PR #4261³³⁹ - Fixing HPX/APEX linking and dependencies for external projects like Phylanx
- PR #4260³⁴⁰ - Fix docs compilation problems
- PR #4258³⁴¹ - Couple of minor changes
- PR #4257³⁴² - Fix apex annotation for async dispatch
- PR #4256³⁴³ - Remove lambdas from assert expressions
- PR #4255³⁴⁴ - Ignoring lock in all_to_all and all_reduce

³²³ <https://github.com/STELLAR-GROUP/hpx/pull/4300>

³²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4299>

³²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4298>

³²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4296>

³²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4295>

³²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4293>

³²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4285>

³³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4276>

³³¹ <https://github.com/STELLAR-GROUP/hpx/pull/4275>

³³² <https://github.com/STELLAR-GROUP/hpx/pull/4272>

³³³ <https://github.com/STELLAR-GROUP/hpx/pull/4271>

³³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4268>

³³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4267>

³³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4266>

³³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4264>

³³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4262>

³³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4261>

³⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4260>

³⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/4258>

³⁴² <https://github.com/STELLAR-GROUP/hpx/pull/4257>

³⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/4256>

³⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4255>

- PR #4254³⁴⁵ - Adding action specializations for noexcept functions
- PR #4253³⁴⁶ - Move partlit.hpp to affinity module
- PR #4252³⁴⁷ - Make mismatching build types a hard error in CMake
- PR #4249³⁴⁸ - Scheduler improvement
- PR #4248³⁴⁹ - update hpxmp tag to v0.3.0
- PR #4245³⁵⁰ - Adding high performance channels
- PR #4244³⁵¹ - Ignore lock in ignore_while_locked_1485 test
- PR #4243³⁵² - Fix PAPI command line option documentation
- PR #4242³⁵³ - Ignore lock in target_distribution_policy
- PR #4241³⁵⁴ - Fix start_stop_callbacks test
- PR #4240³⁵⁵ - Mostly fix clang CUDA compilation
- PR #4238³⁵⁶ - Google Season of Docs updates to documentation; grammar edits.
- PR #4237³⁵⁷ - fixing annotated task to use the name, not the desc
- PR #4236³⁵⁸ - Move module print summary to modules
- PR #4235³⁵⁹ - Don't use alignas in cache_{aligned, line}_data
- PR #4234³⁶⁰ - Add basic overview sentence to all modules
- PR #4230³⁶¹ - Add OS X builds to Travis CI
- PR #4229³⁶² - Remove leftover queue compatibility checks
- PR #4226³⁶³ - Fixing APEX shutdown by explicitly shutting down throttling
- PR #4225³⁶⁴ - Allow CMAKE_INSTALL_PREFIX to be a relative path
- PR #4224³⁶⁵ - Deprecate verbs parcelport
- PR #4222³⁶⁶ - Update register_{thread, work} namespaces
- PR #4221³⁶⁷ - Changing HPX_GCC_VERSION check from 70000 to 70300

³⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4254>

³⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4253>

³⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4252>

³⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4249>

³⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4248>

³⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4245>

³⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/4244>

³⁵² <https://github.com/STELLAR-GROUP/hpx/pull/4243>

³⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/4242>

³⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4241>

³⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4240>

³⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4238>

³⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4237>

³⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4236>

³⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4235>

³⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4234>

³⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/4230>

³⁶² <https://github.com/STELLAR-GROUP/hpx/pull/4229>

³⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/4226>

³⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4225>

³⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4224>

³⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4222>

³⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4221>

- PR #4218³⁶⁸ - Google Season of Docs updates to documentation; grammar edits.
- PR #4217³⁶⁹ - Google Season of Docs updates to documentation; grammar edits.
- PR #4216³⁷⁰ - Fixing gcc warning on 32bit platforms (integer truncation)
- PR #4214³⁷¹ - Apex callback refactoring
- PR #4213³⁷² - Clean up allocator checks for dependent projects
- PR #4212³⁷³ - Google Season of Docs updates to documentation; grammar edits.
- PR #4211³⁷⁴ - Google Season of Docs updates to documentation; contributing to hpx
- PR #4210³⁷⁵ - Attempting to fix Intel compilation
- PR #4209³⁷⁶ - Fix CUDA 10 build
- PR #4205³⁷⁷ - Making sure that differences in CMAKE_BUILD_TYPE are not reported on multi-configuration cmake generators
- PR #4203³⁷⁸ - Deprecate Vc
- PR #4202³⁷⁹ - Fix CUDA configuration
- PR #4200³⁸⁰ - Making sure hpx_wrap is not passed on to linker on non-Linux systems
- PR #4198³⁸¹ - Fix execution_agent.cpp compilation with GCC 5
- PR #4197³⁸² - Remove deprecated options for 1.4.0 release
- PR #4196³⁸³ - minor fixes for building on OSX Darwin
- PR #4195³⁸⁴ - Use full clone on CircleCI for pushing stable tag
- PR #4193³⁸⁵ - Add scheduling hints to hello_world_distributed
- PR #4192³⁸⁶ - Set up CUDA in HPXConfig.cmake
- PR #4191³⁸⁷ - Export allocators root variables
- PR #4190³⁸⁸ - Don't use constexpr in thread_data with GCC <= 6
- PR #4189³⁸⁹ - Only use quick_exit if available
- PR #4188³⁹⁰ - Google Season of Docs updates to documentation; writing single node hpx applications

³⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4218>

³⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4217>

³⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4216>

³⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/4214>

³⁷² <https://github.com/STELLAR-GROUP/hpx/pull/4213>

³⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/4212>

³⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4211>

³⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4210>

³⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4209>

³⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4205>

³⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4203>

³⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4202>

³⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4200>

³⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/4198>

³⁸² <https://github.com/STELLAR-GROUP/hpx/pull/4197>

³⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/4196>

³⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4195>

³⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4193>

³⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4192>

³⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4191>

³⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4190>

³⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4189>

³⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4188>

- PR #4186³⁹¹ - correct vc to cuda in cuda cmake
- PR #4184³⁹² - Resetting some cached variables to make sure those are re-filled
- PR #4183³⁹³ - Fix hpxcxx configuration
- PR #4181³⁹⁴ - Rename base libraries var
- PR #4180³⁹⁵ - Move header left behind earlier to plugin module
- PR #4179³⁹⁶ - Moving zip_iterator and transform_iterator to iterator_support module
- PR #4178³⁹⁷ - Move checkpointing support to its own module
- PR #4177³⁹⁸ - Small const fix to basic_execution module
- PR #4176³⁹⁹ - Add back HPX_LIBRARIES and friends to HPXConfig.cmake
- PR #4175⁴⁰⁰ - Make Vc public and add it to HPXConfig.cmake
- PR #4173⁴⁰¹ - Wait for runtime to be running before returning from hpx::start
- PR #4172⁴⁰² - More protection against shutdown problems in error handling scenarios.
- PR #4171⁴⁰³ - Ignore lock in condition_variable::wait
- PR #4170⁴⁰⁴ - Adding APEX dependency to MPI parcelport
- PR #4168⁴⁰⁵ - Adding utility include
- PR #4167⁴⁰⁶ - Add a condition to setup the external libraries
- PR #4166⁴⁰⁷ - Add an INTERNAL_FLAGS option to link to hpx_internal_flags
- PR #4165⁴⁰⁸ - Forward HPX_* cmake cache variables to external projects
- PR #4164⁴⁰⁹ - Affinity and batch environment modules
- PR #4162⁴¹⁰ - Handle quick exit
- PR #4160⁴¹¹ - Using target_link_libraries for cmake versions >= 3.12
- PR #4159⁴¹² - Make sure HPX_WITH_NATIVE_TLS is forwarded to dependent projects
- PR #4158⁴¹³ - Adding allocator imported target as a dependency of allocator module

³⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/4186>

³⁹² <https://github.com/STELLAR-GROUP/hpx/pull/4184>

³⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/4183>

³⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4181>

³⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4180>

³⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4179>

³⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4178>

³⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4177>

³⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4176>

⁴⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4175>

⁴⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/4173>

⁴⁰² <https://github.com/STELLAR-GROUP/hpx/pull/4172>

⁴⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/4171>

⁴⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4170>

⁴⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4168>

⁴⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4167>

⁴⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4166>

⁴⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4165>

⁴⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4164>

⁴¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4162>

⁴¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/4160>

⁴¹² <https://github.com/STELLAR-GROUP/hpx/pull/4159>

⁴¹³ <https://github.com/STELLAR-GROUP/hpx/pull/4158>

- PR #4157⁴¹⁴ - Add `hpx_memory` as a dependency of parcelport plugins
- PR #4156⁴¹⁵ - Stackless coroutines now can refer to themselves (through `get_self()` and friends)
- PR #4154⁴¹⁶ - Added CMake policy CMP0060 for HPX applications.
- PR #4153⁴¹⁷ - add header `iomanip` to tests and tool
- PR #4152⁴¹⁸ - Casting MPI tag value
- PR #4149⁴¹⁹ - Add back private `m_desc` member variable in `program_options` module
- PR #4147⁴²⁰ - Resource partitioner and threadmanager modules
- PR #4146⁴²¹ - Google Season of Docs updates to documentation; creating hpx projects
- PR #4145⁴²² - Adding basic support for stackless threads
- PR #4143⁴²³ - Exclude `test_client_1950` from all target
- PR #4142⁴²⁴ - Add a new `thread_pool_executor`
- PR #4140⁴²⁵ - Google Season of Docs updates to documentation; why hpx
- PR #4139⁴²⁶ - Remove runtime includes from coroutines module
- PR #4138⁴²⁷ - Forking `boost::intrusive_ptr` and adding it as `hpx::intrusive_ptr`
- PR #4137⁴²⁸ - Fixing TSS destruction
- PR #4136⁴²⁹ - HPX.Compute modules
- PR #4133⁴³⁰ - Fix `block_executor`
- PR #4131⁴³¹ - Applying fixes based on reports from PVS Studio
- PR #4130⁴³² - Adding missing header to build system
- PR #4129⁴³³ - Fixing compilation if `HPX_WITH_DATAPAR_VC` is enabled
- PR #4128⁴³⁴ - Renaming `moveonly_any` to `unique_any`
- PR #4126⁴³⁵ - Attempt to fix `basic_any` constructor for gcc 7
- PR #4125⁴³⁶ - Changing `extra_archive_data` implementation

⁴¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4157>

⁴¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4156>

⁴¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4154>

⁴¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4153>

⁴¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4152>

⁴¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4149>

⁴²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4147>

⁴²¹ <https://github.com/STELLAR-GROUP/hpx/pull/4146>

⁴²² <https://github.com/STELLAR-GROUP/hpx/pull/4145>

⁴²³ <https://github.com/STELLAR-GROUP/hpx/pull/4143>

⁴²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4142>

⁴²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4140>

⁴²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4139>

⁴²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4138>

⁴²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4137>

⁴²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4136>

⁴³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4133>

⁴³¹ <https://github.com/STELLAR-GROUP/hpx/pull/4131>

⁴³² <https://github.com/STELLAR-GROUP/hpx/pull/4130>

⁴³³ <https://github.com/STELLAR-GROUP/hpx/pull/4129>

⁴³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4128>

⁴³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4126>

⁴³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4125>

- PR #4124⁴³⁷ - Don't link to Boost.System unless required
- PR #4122⁴³⁸ - Add kernel launch helper utility (+saxpy demo) and merge in octotiger changes
- PR #4121⁴³⁹ - Fixing migration test if networking is disabled.
- PR #4120⁴⁴⁰ - Google Season of Docs updates to documentation; hpx build system v1
- PR #4119⁴⁴¹ - Making sure chunk_size and max_chunk are actually applied to parallel algorithms if specified
- PR #4117⁴⁴² - Make CircleCI formatting check store diff
- PR #4116⁴⁴³ - Fix automatically setting C++ standard
- PR #4114⁴⁴⁴ - Module serialization
- PR #4113⁴⁴⁵ - Module datastructures
- PR #4111⁴⁴⁶ - Fixing performance regression introduced earlier
- PR #4110⁴⁴⁷ - Adding missing SPDX tags
- PR #4109⁴⁴⁸ - Overload for start without entry point/argv.
- PR #4108⁴⁴⁹ - Making sure C++ standard is properly detected and propagated
- PR #4106⁴⁵⁰ - use std::round for guaranteed rounding without errors
- PR #4104⁴⁵¹ - Extend scheduler_mode with new work_stealing and task assignment modes
- PR #4103⁴⁵² - Add this to lambda capture list
- PR #4102⁴⁵³ - Add spdx license and check
- PR #4099⁴⁵⁴ - Module coroutines
- PR #4098⁴⁵⁵ - Fix append module path in module CMakeLists template
- PR #4097⁴⁵⁶ - Function tests
- PR #4096⁴⁵⁷ - Removing return of thread_result_type from functions not needing them
- PR #4095⁴⁵⁸ - Stop-gap measure until cmake overhaul is in place
- PR #4094⁴⁵⁹ - Deprecate HPX_WITH_MORE_THAN_64_THREADS

⁴³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4124>

⁴³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4122>

⁴³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4121>

⁴⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4120>

⁴⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/4119>

⁴⁴² <https://github.com/STELLAR-GROUP/hpx/pull/4117>

⁴⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/4116>

⁴⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4114>

⁴⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4113>

⁴⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4111>

⁴⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4110>

⁴⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4109>

449 <https://github.com/STELLAR-GROUP/hpx/pull/4108>450 <https://github.com/STELLAR-GROUP/hpx/pull/4106>451 <https://github.com/STELLAR-GROUP/hpx/pull/4104>452 <https://github.com/STELLAR-GROUP/hpx/pull/4103>453 <https://github.com/STELLAR-GROUP/hpx/pull/4102>454 <https://github.com/STELLAR-GROUP/hpx/pull/4099>455 <https://github.com/STELLAR-GROUP/hpx/pull/4098>456 <https://github.com/STELLAR-GROUP/hpx/pull/4097>457 <https://github.com/STELLAR-GROUP/hpx/pull/4096>458 <https://github.com/STELLAR-GROUP/hpx/pull/4095>459 <https://github.com/STELLAR-GROUP/hpx/pull/4094>

- PR #4093⁴⁶⁰ - Fix initialization of `global_num_tasks` in `parallel_executor`
- PR #4092⁴⁶¹ - Add support for mi-malloc
- PR #4090⁴⁶² - Execution context
- PR #4089⁴⁶³ - Make counters in coroutines optional
- PR #4087⁴⁶⁴ - Making `hpx::util::any` compatible with C++17
- PR #4084⁴⁶⁵ - Making sure destination array for `std::transform` is properly resized
- PR #4083⁴⁶⁶ - Adapting `thread_queue_mc` to behave even if no 128bit atomics are available
- PR #4082⁴⁶⁷ - Fix compilation on GCC 5
- PR #4081⁴⁶⁸ - Adding option allowing to force using Boost.FileSystem
- PR #4080⁴⁶⁹ - Updating module dependencies
- PR #4079⁴⁷⁰ - Add missing tests for iterator_support module
- PR #4078⁴⁷¹ - Disable parcel-layer if networking is disabled
- PR #4077⁴⁷² - Add missing include that causes build fails
- PR #4076⁴⁷³ - Enable compatibility headers for functional module
- PR #4075⁴⁷⁴ - Coroutines module
- PR #4073⁴⁷⁵ - Use `configure_file` for generated files in modules
- PR #4071⁴⁷⁶ - Fixing MPI detection for PMIx
- PR #4070⁴⁷⁷ - Fix macOS builds
- PR #4069⁴⁷⁸ - Moving more facilities to the collectives module
- PR #4068⁴⁷⁹ - Adding main HPX `#include` directory to modules
- PR #4066⁴⁸⁰ - Switching the use of `message (STATUS " . . . ")` to `hpx_info`
- PR #4065⁴⁸¹ - Move Boost.Filesystem handling to filesystem module
- PR #4064⁴⁸² - Fix program_options test with older boost versions

⁴⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4093>⁴⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/4092>⁴⁶² <https://github.com/STELLAR-GROUP/hpx/pull/4090>⁴⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/4089>⁴⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4087>⁴⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4084>466 <https://github.com/STELLAR-GROUP/hpx/pull/4083>467 <https://github.com/STELLAR-GROUP/hpx/pull/4082>468 <https://github.com/STELLAR-GROUP/hpx/pull/4081>469 <https://github.com/STELLAR-GROUP/hpx/pull/4080>470 <https://github.com/STELLAR-GROUP/hpx/pull/4079>471 <https://github.com/STELLAR-GROUP/hpx/pull/4078>472 <https://github.com/STELLAR-GROUP/hpx/pull/4077>473 <https://github.com/STELLAR-GROUP/hpx/pull/4076>474 <https://github.com/STELLAR-GROUP/hpx/pull/4075>475 <https://github.com/STELLAR-GROUP/hpx/pull/4073>476 <https://github.com/STELLAR-GROUP/hpx/pull/4071>477 <https://github.com/STELLAR-GROUP/hpx/pull/4070>478 <https://github.com/STELLAR-GROUP/hpx/pull/4069>479 <https://github.com/STELLAR-GROUP/hpx/pull/4068>480 <https://github.com/STELLAR-GROUP/hpx/pull/4066>481 <https://github.com/STELLAR-GROUP/hpx/pull/4065>482 <https://github.com/STELLAR-GROUP/hpx/pull/4064>

- PR #4062⁴⁸³ - The `cpu_features` tool fails to compile on anything but x86 architectures
- PR #4061⁴⁸⁴ - Add `clang-format` checking step for modules
- PR #4060⁴⁸⁵ - Making sure `HPX_IDLE_BACKOFF_TIME_MAX` is always defined (even if its unused)
- PR #4059⁴⁸⁶ - Renaming module `hpx_parallel_executors` into `hpx_execution`
- PR #4058⁴⁸⁷ - Do not build networking tests when networking disabled
- PR #4057⁴⁸⁸ - Printing configuration summary for modules as well
- PR #4055⁴⁸⁹ - Google Season of Docs updates to documentation; hpx build systems
- PR #4054⁴⁹⁰ - Add troubleshooting section to manual
- PR #4051⁴⁹¹ - Add more variations to `future_overhead` test
- PR #4050⁴⁹² - Creating plugin module
- PR #4049⁴⁹³ - Move missing modules tests
- PR #4047⁴⁹⁴ - Add boost/filesystem headers to inspect deprecated headers
- PR #4045⁴⁹⁵ - Module functional
- PR #4043⁴⁹⁶ - Fix preconditions and error messages for suspension functions
- PR #4041⁴⁹⁷ - Pass `HPX_STANDARD` on to dependent projects via `HPXConfig.cmake`
- PR #4040⁴⁹⁸ - Program options module
- PR #4039⁴⁹⁹ - Moving non-serializable `any` (`any_nonser`) to datastructures module
- PR #4038⁵⁰⁰ - Adding MPark's variant (V1.4.0) to HPX
- PR #4037⁵⁰¹ - Adding resiliency module
- PR #4036⁵⁰² - Add C++17 filesystem compatibility header
- PR #4035⁵⁰³ - Fixing support for `mpirun`
- PR #4028⁵⁰⁴ - CMake to target based directives
- PR #4027⁵⁰⁵ - Remove GitLab CI configuration

⁴⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/4062>

⁴⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4061>

⁴⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4060>

⁴⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4059>

⁴⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4058>

⁴⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4057>

⁴⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4055>

⁴⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4054>

⁴⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/4051>

⁴⁹² <https://github.com/STELLAR-GROUP/hpx/pull/4050>

⁴⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/4049>

⁴⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4047>

⁴⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4045>

⁴⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4043>

⁴⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4041>

⁴⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4040>

⁴⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4039>

⁵⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4038>

⁵⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/4037>

⁵⁰² <https://github.com/STELLAR-GROUP/hpx/pull/4036>

⁵⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/4035>

⁵⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4028>

⁵⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4027>

- PR #4026⁵⁰⁶ - Threading refactoring
- PR #4025⁵⁰⁷ - Refactoring thread queue configuration options
- PR #4024⁵⁰⁸ - Fix padding calculation in `cache_aligned_data.hpp`
- PR #4023⁵⁰⁹ - Fixing Codacy issues
- PR #4022⁵¹⁰ - Make sure process mask option is passed to `affinity_data`
- PR #4021⁵¹¹ - Warn about compiling in C++11 mode
- PR #4020⁵¹² - Module concurrency
- PR #4019⁵¹³ - Module topology
- PR #4018⁵¹⁴ - Update deprecated header in `thread_queue_mc.hpp`
- PR #4015⁵¹⁵ - Avoid overwriting artifacts
- PR #4014⁵¹⁶ - Future overheads
- PR #4013⁵¹⁷ - Update URL to test output conversion script
- PR #4012⁵¹⁸ - Fix CUDA compilation
- PR #4011⁵¹⁹ - Fixing cyclic dependencies between modules
- PR #4010⁵²⁰ - Ignore stable tag on CircleCI
- PR #4009⁵²¹ - Check circular dependencies in a circle ci step
- PR #4008⁵²² - Extend cache aligned data to handle tuple-like data
- PR #4007⁵²³ - Fixing migration for components that have actions returning a client
- PR #4006⁵²⁴ - Move `is_value_proxy.hpp` to algorithms module
- PR #4004⁵²⁵ - Shorten CTest timeout on CircleCI
- PR #4003⁵²⁶ - Refactoring to remove (internal) dependencies
- PR #4001⁵²⁷ - Exclude tests from all target
- PR #4000⁵²⁸ - Module errors

⁵⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4026>

⁵⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4025>

⁵⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4024>

⁵⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4023>

⁵¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4022>

⁵¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/4021>

⁵¹² <https://github.com/STELLAR-GROUP/hpx/pull/4020>

⁵¹³ <https://github.com/STELLAR-GROUP/hpx/pull/4019>

⁵¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4018>

⁵¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4015>

⁵¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4014>

⁵¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4013>

⁵¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4012>

⁵¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/4011>

⁵²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/4010>

⁵²¹ <https://github.com/STELLAR-GROUP/hpx/pull/4009>

⁵²² <https://github.com/STELLAR-GROUP/hpx/pull/4008>

⁵²³ <https://github.com/STELLAR-GROUP/hpx/pull/4007>

⁵²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/4006>

⁵²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/4004>

⁵²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/4003>

⁵²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/4001>

⁵²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/4000>

- PR #3999⁵²⁹ - Enable support for compatibility headers for logging module
- PR #3998⁵³⁰ - Add process thread binding option
- PR #3997⁵³¹ - Export handle_assert function
- PR #3996⁵³² - Attempt to solve issue where `-latomic` does not support 128bit atomics
- PR #3993⁵³³ - Make sure `__LINE__` is an unsigned
- PR #3991⁵³⁴ - Fix dependencies and flags for header tests
- PR #3990⁵³⁵ - Documentation tags fixes
- PR #3988⁵³⁶ - Adding missing solution folder for format module test
- PR #3987⁵³⁷ - Move runtime-dependent functions out of command line handling
- PR #3986⁵³⁸ - Fix CMake configuration with PAPI on
- PR #3985⁵³⁹ - Module timing
- PR #3984⁵⁴⁰ - Fix default behaviour of paths in `add_hpx_component`
- PR #3982⁵⁴¹ - Parallel executors module
- PR #3981⁵⁴² - Segmented algorithms module
- PR #3980⁵⁴³ - Module logging
- PR #3979⁵⁴⁴ - Module util
- PR #3978⁵⁴⁵ - Fix clang-tidy step on CircleCI
- PR #3977⁵⁴⁶ - Fixing solution folders for moved components
- PR #3976⁵⁴⁷ - Module format
- PR #3975⁵⁴⁸ - Enable deprecation warnings on CircleCI
- PR #3974⁵⁴⁹ - Fix typos in documentation
- PR #3973⁵⁵⁰ - Fix compilation with GCC 9
- PR #3972⁵⁵¹ - Add condition to clone apex + use of new cmake var `APEX_ROOT`

⁵²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3999>

⁵³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3998>

⁵³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3997>

⁵³² <https://github.com/STELLAR-GROUP/hpx/pull/3996>

⁵³³ <https://github.com/STELLAR-GROUP/hpx/pull/3993>

⁵³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3991>

⁵³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3990>

⁵³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3988>

⁵³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3987>

⁵³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3986>

⁵³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3985>

⁵⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3984>

⁵⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3982>

⁵⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3981>

⁵⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3980>

⁵⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3979>

⁵⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3978>

⁵⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3977>

⁵⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3976>

⁵⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3975>

⁵⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3974>

⁵⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3973>

⁵⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3972>

- PR #3971⁵⁵² - Add testing module
- PR #3968⁵⁵³ - Remove unneeded file in hardware module
- PR #3967⁵⁵⁴ - Remove leftover PIC settings from main CMakeLists.txt
- PR #3966⁵⁵⁵ - Add missing export option in add_hpx_module
- PR #3965⁵⁵⁶ - Change current_function_helper back to non-constexpr
- PR #3964⁵⁵⁷ - Fixing merge problems
- PR #3962⁵⁵⁸ - Add a trait for std::array for unwrapping
- PR #3961⁵⁵⁹ - Making hpx::util::tuple<Ts...> and std::tuple<Ts...> convertible
- PR #3960⁵⁶⁰ - fix compilation with CUDA 10 and GCC 6
- PR #3959⁵⁶¹ - Fix C++11 incompatibility
- PR #3957⁵⁶² - Algorithms module
- PR #3956⁵⁶³ - [HPX_AddModule] Fix lower name var to upper
- PR #3955⁵⁶⁴ - Fix CMake configuration with examples off and tests on
- PR #3954⁵⁶⁵ - Move components to separate subdirectory in root of repository
- PR #3952⁵⁶⁶ - Update papi.cpp
- PR #3951⁵⁶⁷ - Exclude modules header tests from all target
- PR #3950⁵⁶⁸ - Adding all_reduce facility to collectives module
- PR #3949⁵⁶⁹ - This adds a configuration file that will cause for stale issues to be automatically closed
- PR #3948⁵⁷⁰ - Fixing ALPS environment
- PR #3947⁵⁷¹ - Add major compiler version check for building hpx as a binary package
- PR #3946⁵⁷² - [Modules] Move the location of the generated headers
- PR #3945⁵⁷³ - Simplify tests and examples cmake
- PR #3943⁵⁷⁴ - Remove example module

⁵⁵² <https://github.com/STELLAR-GROUP/hpx/pull/3971>

⁵⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3968>

⁵⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3967>

⁵⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3966>

⁵⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3965>

⁵⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3964>

⁵⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3962>

⁵⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3961>

⁵⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3960>

⁵⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3959>

⁵⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3957>

⁵⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3956>

⁵⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3955>

⁵⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3954>

⁵⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3952>

⁵⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3951>

⁵⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3950>

⁵⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3949>

⁵⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3948>

⁵⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/3947>

⁵⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3946>

⁵⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3945>

⁵⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3943>

- PR #3942⁵⁷⁵ - Add NOEXPORT option to `add_hpx_{component,library}`
- PR #3938⁵⁷⁶ - Use https for CDash submissions
- PR #3937⁵⁷⁷ - Add HPX_WITH_BUILD_BINARY_PACKAGE to the compiler check (refs #3935)
- PR #3936⁵⁷⁸ - Fixing installation of binaries on windows
- PR #3934⁵⁷⁹ - Add set function for `sliding_semaphore max_difference`
- PR #3933⁵⁸⁰ - Remove `cudadevrt` from compile/link flags as it breaks downstream projects
- PR #3932⁵⁸¹ - Fixing 3929
- PR #3931⁵⁸² - Adding `all_to_all`
- PR #3930⁵⁸³ - Add test demonstrating the use of broadcast with component actions
- PR #3928⁵⁸⁴ - fixed number of tasks and number of threads for heterogeneous slurm environments
- PR #3927⁵⁸⁵ - Moving Cache module's tests into separate solution folder
- PR #3926⁵⁸⁶ - Move unit tests to cache module
- PR #3925⁵⁸⁷ - Move version check to config module
- PR #3924⁵⁸⁸ - Add schedule hint executor parameters
- PR #3923⁵⁸⁹ - Allow aligning objects bigger than the cache line size
- PR #3922⁵⁹⁰ - Add Windows builds with Travis CI
- PR #3921⁵⁹¹ - Add ccls cache directory to gitignore
- PR #3920⁵⁹² - Fix `git_external` fetching of tags
- PR #3905⁵⁹³ - Correct rostambod url. Fix typo in doc
- PR #3904⁵⁹⁴ - Fix bug in `context_base.hpp`
- PR #3903⁵⁹⁵ - Adding new performance counters
- PR #3902⁵⁹⁶ - Add `add_hpx_module` function
- PR #3901⁵⁹⁷ - Factoring out container remapping into a separate trait

⁵⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3942>

⁵⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3938>

⁵⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3937>

⁵⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3936>

⁵⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3934>

⁵⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3933>

⁵⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3932>

⁵⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3931>

⁵⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3930>

⁵⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3928>

⁵⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3927>

⁵⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3926>

⁵⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3925>

⁵⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3924>

⁵⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3923>

⁵⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3922>

⁵⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3921>

⁵⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3920>

⁵⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3905>

⁵⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3904>

⁵⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3903>

⁵⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3902>

⁵⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3901>

- PR #3900⁵⁹⁸ - Making sure errors during command line processing are properly reported and will not cause assertions
- PR #3899⁵⁹⁹ - Remove old compatibility bases from make_action
- PR #3898⁶⁰⁰ - Make parameter size be of type size_t
- PR #3897⁶⁰¹ - Making sure all tests are disabled if HPX_WITH_TESTS=OFF
- PR #3895⁶⁰² - Add documentation for annotated_function
- PR #3894⁶⁰³ - Working around VS2019 problem with make_action
- PR #3892⁶⁰⁴ - Avoid MSVC compatibility warning in internal allocator
- PR #3891⁶⁰⁵ - Removal of the default intel config include
- PR #3888⁶⁰⁶ - Fix async_customization dataflow example and Clarify what's being tested
- PR #3887⁶⁰⁷ - Add Doxygen documentation
- PR #3882⁶⁰⁸ - Minor docs fixes
- PR #3880⁶⁰⁹ - Updating APEX version tag
- PR #3878⁶¹⁰ - Making sure symbols are properly exported from modules (needed for Windows/MacOS)
- PR #3877⁶¹¹ - Documentation
- PR #3876⁶¹² - Module hardware
- PR #3875⁶¹³ - Converted typedefs in actions submodule to using directives
- PR #3874⁶¹⁴ - Allow one to suppress target keywords in hpx_setup_target for backwards compatibility
- PR #3873⁶¹⁵ - Add scripts to create releases and generate lists of PRs and issues
- PR #3872⁶¹⁶ - Fix latest HTML docs location
- PR #3870⁶¹⁷ - Module cache
- PR #3869⁶¹⁸ - Post 1.3.0 version bumps
- PR #3868⁶¹⁹ - Replace the macro HPX_ASSERT by HPX_TEST in tests
- PR #3845⁶²⁰ - Assertion module

⁵⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3900>

⁵⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3899>

⁶⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3898>

⁶⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3897>

⁶⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3895>

⁶⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3894>

⁶⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3892>

⁶⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3891>

⁶⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3888>

⁶⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3887>

⁶⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3882>

⁶⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3880>

⁶¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3878>

⁶¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3877>

⁶¹² <https://github.com/STELLAR-GROUP/hpx/pull/3876>

⁶¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3875>

⁶¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3874>

⁶¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3873>

⁶¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3872>

⁶¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3870>

⁶¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3869>

⁶¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3868>

⁶²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3845>

- PR #3839⁶²¹ - Make tuple serialization non-intrusive
- PR #3832⁶²² - Config module
- PR #3799⁶²³ - Remove compat namespace and its contents
- PR #3701⁶²⁴ - MoodyCamel lockfree
- PR #3496⁶²⁵ - Disabling MPI's (deprecated) C++ interface
- PR #3192⁶²⁶ - Move type info into `hpx::debug` namespace and add print helper functions
- PR #3159⁶²⁷ - Support Checkpointing Components

2.11.2 HPX V1.3.0 (May 23, 2019)

General changes

- Performance improvements: the schedulers have significantly reduced overheads from removing false sharing and the parallel executor has been updated to create fewer futures.
- HPX now defaults to not turning on networking when running on one locality. This means that you can run multiple instances on the same system without adding command line options.
- Multiple issues reported by Clang sanitizers have been fixed.
- We have added (back) single-page HTML documentation and PDF documentation.
- We have started modularizing the HPX library. This is useful both for developers and users. In the long term users will be able to consume only parts of the HPX libraries if they do not require all the functionality that HPX currently provides.
- We have added an implementation of `function_ref`.
- The `barrier` and `latch` classes have gained a few additional member functions.

Breaking changes

- Executable and library targets are now created without the `_exe` and `_lib` suffix respectively. For example, the target `1d_stencil_1_exe` is now simply called `1d_stencil_1`.
- We have removed the following deprecated functionality: `queue`, `scoped_unlock`, and support for input iterators in algorithms.
- We have turned off the compatibility layer for `unwrapped` by default. The functionality will be removed in the next release. The option can still be turned on using the CMake⁶²⁸ option `HPX_WITH_UNWRAPPED_SUPPORT`. Likewise, `inclusive_scan` compatibility overloads have been turned off by default. They can still be turned on with `HPX_WITH_INCLUSIVE_SCAN_COMPATIBILITY`.
- The minimum compiler and dependency versions have been updated. We now support GCC from version 5 onwards, Clang from version 4 onwards, and Boost from version 1.61.0 onwards.

⁶²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3839>

⁶²² <https://github.com/STELLAR-GROUP/hpx/pull/3832>

⁶²³ <https://github.com/STELLAR-GROUP/hpx/pull/3799>

⁶²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3701>

⁶²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3496>

⁶²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3192>

⁶²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3159>

⁶²⁸ <https://www.cmake.org>

- The headers for preprocessor macros have moved as a result of the functionality being moved to a separate module. The old headers are deprecated and will be removed in a future version of HPX. You can turn off the warnings by setting `HPX_PREPROCESSOR_WITH_DEPRECATED_WARNINGS=OFF` or turn off the compatibility headers completely with `HPX_PREPROCESSOR_WITH_COMPATIBILITY_HEADERS=OFF`.

Closed issues

- Issue #3863⁶²⁹ - shouldn't “-faligned-new” be a usage requirement?
- Issue #3841⁶³⁰ - Build error with msvc 19 caused by SFINAE and C++17
- Issue #3836⁶³¹ - master branch does not build with idle rate counters enabled
- Issue #3819⁶³² - Add debug suffix to modules built in debug mode
- Issue #3817⁶³³ - `HPX_INCLUDE_DIRS` contains non-existent directory
- Issue #3810⁶³⁴ - Source groups are not created for files in modules
- Issue #3805⁶³⁵ - HPX won't compile with `-DHPX_WITH_APEX=TRUE`
- Issue #3792⁶³⁶ - Barrier Hangs When Locality Zero not included
- Issue #3778⁶³⁷ - Replace `throw()` with `noexcept`
- Issue #3763⁶³⁸ - configurable sort limit per task
- Issue #3758⁶³⁹ - dataflow doesn't convert `future<future<T>>` to `future<T>`
- Issue #3757⁶⁴⁰ - When compiling undefined reference to `hpx::hpx_check_version_1_2` HPX V1.2.1, Ubuntu 18.04.01 Server Edition
- Issue #3753⁶⁴¹ - `--hpx:list-counters=full` crashes
- Issue #3746⁶⁴² - Detection of MPI with pmix
- Issue #3744⁶⁴³ - Separate spinlock from same cacheline as internal data for all LCOs
- Issue #3743⁶⁴⁴ - `hpxcxx`'s shebang doesn't specify the python version
- Issue #3738⁶⁴⁵ - Unable to debug parcelport on a single node
- Issue #3735⁶⁴⁶ - Latest master: Can't compile in MSVC
- Issue #3731⁶⁴⁷ - `util::bound` seems broken on Clang with older libstdc++

⁶²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3863>

⁶³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3841>

⁶³¹ <https://github.com/STELLAR-GROUP/hpx/issues/3836>

⁶³² <https://github.com/STELLAR-GROUP/hpx/issues/3819>

⁶³³ <https://github.com/STELLAR-GROUP/hpx/issues/3817>

⁶³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3810>

⁶³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3805>

⁶³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3792>

⁶³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3778>

⁶³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3763>

⁶³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3758>

⁶⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3757>

⁶⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/3753>

⁶⁴² <https://github.com/STELLAR-GROUP/hpx/issues/3746>

⁶⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/3744>

⁶⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3743>

⁶⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3738>

⁶⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3735>

⁶⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3731>

- Issue #3724⁶⁴⁸ - Allow to pre-set command line options through environment
- Issue #3723⁶⁴⁹ - examples/resource_partitioner build issue on master branch / ubuntu 18
- Issue #3721⁶⁵⁰ - faced a building error
- Issue #3720⁶⁵¹ - Hello World example fails to link
- Issue #3719⁶⁵² - pkg-config produces invalid output: -l-pthread
- Issue #3718⁶⁵³ - Please make the python executable configurable through cmake
- Issue #3717⁶⁵⁴ - interested to contribute to the organisation
- Issue #3699⁶⁵⁵ - Remove 'HPX runtime' executable
- Issue #3698⁶⁵⁶ - Ignore all locks while handling asserts
- Issue #3689⁶⁵⁷ - Incorrect and inconsistent website structure <http://stellar.cct.lsu.edu/downloads/>.
- Issue #3681⁶⁵⁸ - Broken links on <http://stellar.cct.lsu.edu/2015/05/hpx-archives-now-on-gmane/>
- Issue #3676⁶⁵⁹ - HPX master built from source, cmake fails to link main.cpp example in docs
- Issue #3673⁶⁶⁰ - HPX build fails with std::atomic missing error
- Issue #3670⁶⁶¹ - Generate PDF again from documentation (with Sphinx)
- Issue #3643⁶⁶² - Warnings when compiling HPX 1.2.1 with gcc 9
- Issue #3641⁶⁶³ - Trouble with using ranges-v3 and hpx::parallel::reduce
- Issue #3639⁶⁶⁴ - util::unwrapping does not work well with member functions
- Issue #3634⁶⁶⁵ - The build fails if shared_future<>::then is called with a thread executor
- Issue #3622⁶⁶⁶ - VTune Amplifier 2019 not working with use_itt_notify=1
- Issue #3616⁶⁶⁷ - HPX Fails to Build with CUDA 10
- Issue #3612⁶⁶⁸ - False sharing of scheduling counters
- Issue #3609⁶⁶⁹ - executor_parameters timeout with gcc <= 7 and Debug mode
- Issue #3601⁶⁷⁰ - Misleading error message on power pc for rdtsc and rdtscp

⁶⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3724>

⁶⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3723>

⁶⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3721>

⁶⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/3720>

⁶⁵² <https://github.com/STELLAR-GROUP/hpx/issues/3719>

⁶⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/3718>

⁶⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3717>

⁶⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3699>

⁶⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3698>

⁶⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3689>

⁶⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3681>

⁶⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3676>

⁶⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3673>

⁶⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/3670>

⁶⁶² <https://github.com/STELLAR-GROUP/hpx/issues/3643>

⁶⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/3641>

⁶⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3639>

⁶⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3634>

⁶⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3622>

⁶⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3616>

⁶⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3612>

⁶⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3609>

⁶⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3601>

- Issue #3598⁶⁷¹ - Build of some examples fails when using Vc
- Issue #3594⁶⁷² - Error: The number of OS threads requested (20) does not match the number of threads to bind (12): HPX(bad_parameter)
- Issue #3592⁶⁷³ - Undefined Reference Error
- Issue #3589⁶⁷⁴ - include could not find load file: HPX_Utils.cmake
- Issue #3587⁶⁷⁵ - HPX won't compile on POWER8 with Clang 7
- Issue #3583⁶⁷⁶ - Fedora and openSUSE instructions missing on "Distribution Packages" page
- Issue #3578⁶⁷⁷ - Build error when configuring with HPX_HAVE_ALGORITHM_INPUT_ITERATOR_SUPPORT=ON
- Issue #3575⁶⁷⁸ - Merge openSUSE reproducible patch
- Issue #3570⁶⁷⁹ - Update HPX to work with the latest VC version
- Issue #3567⁶⁸⁰ - Build succeed and make failed for hpx:cout
- Issue #3565⁶⁸¹ - Polymorphic simple component destructor not getting called
- Issue #3559⁶⁸² - 1.2.0 is missing from download page
- Issue #3554⁶⁸³ - Clang 6.0 warning of hiding overloaded virtual function
- Issue #3510⁶⁸⁴ - Build on ppc64 fails
- Issue #3482⁶⁸⁵ - Improve error message when HPX_WITH_MAX_CPU_COUNT is too low for given system
- Issue #3453⁶⁸⁶ - Two HPX applications can't run at the same time.
- Issue #3452⁶⁸⁷ - Scaling issue on the change to 2 NUMA domains
- Issue #3442⁶⁸⁸ - HPX set_difference, set_intersection failure cases
- Issue #3437⁶⁸⁹ - Ensure parent_task pointer when child task is created and child/parent are on same locality
- Issue #3255⁶⁹⁰ - Suspension with lock for --hpx:list-component-types
- Issue #3034⁶⁹¹ - Use C++17 structured bindings for serialization
- Issue #2999⁶⁹² - Change thread scheduling use of size_t for thread indexing

⁶⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/3598>

⁶⁷² <https://github.com/STELLAR-GROUP/hpx/issues/3594>

⁶⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/3592>

⁶⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3589>

⁶⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3587>

⁶⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3583>

⁶⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3578>

⁶⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3575>

⁶⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3570>

⁶⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3567>

⁶⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/3565>

⁶⁸² <https://github.com/STELLAR-GROUP/hpx/issues/3559>

⁶⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/3554>

⁶⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3510>

⁶⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3482>

⁶⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3453>

⁶⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3452>

⁶⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3442>

⁶⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3437>

⁶⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3255>

⁶⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/3034>

⁶⁹² <https://github.com/STELLAR-GROUP/hpx/issues/2999>

Closed pull requests

- PR #3865⁶⁹³ - adds hpx_target_compile_option_if_available
- PR #3864⁶⁹⁴ - Helper functions that are useful in numa binding and testing of allocator
- PR #3862⁶⁹⁵ - Temporary fix to local_dataflow_boost_small_vector test
- PR #3860⁶⁹⁶ - Add cache line padding to intermediate results in for loop reduction
- PR #3859⁶⁹⁷ - Remove HPX_TLL_PUBLIC and HPX_TLL_PRIVATE from CMake files
- PR #3858⁶⁹⁸ - Add compile flags and definitions to modules
- PR #3851⁶⁹⁹ - update hpxmp release tag to v0.2.0
- PR #3849⁷⁰⁰ - Correct BOOST_ROOT variable name in quick start guide
- PR #3847⁷⁰¹ - Fix attach_debugger configuration option
- PR #3846⁷⁰² - Add tests for libs header tests
- PR #3844⁷⁰³ - Fixing source_groups in preprocessor module to properly handle compatibility headers
- PR #3843⁷⁰⁴ - This fixes the launch_process/launched_process pair of tests
- PR #3842⁷⁰⁵ - Fix macro call with ITTNNOTIFY enabled
- PR #3840⁷⁰⁶ - Fixing SLURM environment parsing
- PR #3837⁷⁰⁷ - Fixing misplaced #endif
- PR #3835⁷⁰⁸ - make all latch members protected for consistency
- PR #3834⁷⁰⁹ - Disable transpose_block_numa example on CircleCI
- PR #3833⁷¹⁰ - make latch **counter**_ protected for deriving latch in hpxmp
- PR #3831⁷¹¹ - Fix CircleCI config for modules
- PR #3830⁷¹² - minor fix: option HPX_WITH_TEST was not working correctly
- PR #3828⁷¹³ - Avoid for binaries that depend on HPX to directly link against internal modules
- PR #3827⁷¹⁴ - Adding shortcut for hpx::get_ptr<>(sync, id) for a local, non-migratable objects

⁶⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3865>

⁶⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3864>

⁶⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3862>

⁶⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3860>

⁶⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3859>

⁶⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3858>

⁶⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3851>

⁷⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3849>

⁷⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3847>

⁷⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3846>

⁷⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3844>

⁷⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3843>

⁷⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3842>

⁷⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3840>

⁷⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3837>

⁷⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3835>

⁷⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3834>

⁷¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3833>

⁷¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3831>

⁷¹² <https://github.com/STELLAR-GROUP/hpx/pull/3830>

⁷¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3828>

⁷¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3827>

- PR #3826⁷¹⁵ - Fix and update modules documentation
- PR #3825⁷¹⁶ - Updating default APEX version to 2.1.3 with HPX
- PR #3823⁷¹⁷ - Fix pkgconfig libs handling
- PR #3822⁷¹⁸ - Change includes in hpx_wrap.cpp to more specific includes
- PR #3821⁷¹⁹ - Disable barrier_3792 test when networking is disabled
- PR #3820⁷²⁰ - Assorted CMake fixes
- PR #3815⁷²¹ - Removing left-over debug output
- PR #3814⁷²² - Allow setting default scheduler mode via the configuration database
- PR #3813⁷²³ - Make the deprecation warnings issued by the old pp headers optional
- PR #3812⁷²⁴ - Windows requires to handle symlinks to directories differently from those linking files
- PR #3811⁷²⁵ - Clean up PP module and library skeleton
- PR #3806⁷²⁶ - Moving include path configuration to before APEX
- PR #3804⁷²⁷ - Fix latch
- PR #3803⁷²⁸ - Update hpcxx to look at lib64 and use python3
- PR #3802⁷²⁹ - Numa binding allocator
- PR #3801⁷³⁰ - Remove duplicated includes
- PR #3800⁷³¹ - Attempt to fix Posix context switching after lazy init changes
- PR #3798⁷³² - count and count_if accepts different iterator types
- PR #3797⁷³³ - Adding a couple of override keywords to overloaded virtual functions
- PR #3796⁷³⁴ - Re-enable testing all schedulers in shutdown_suspended_test
- PR #3795⁷³⁵ - Change std::terminate to std::abort in SIGSEGV handler
- PR #3794⁷³⁶ - Fixing #3792
- PR #3793⁷³⁷ - Extending migrate_polymorphic_component unit test

⁷¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3826>

⁷¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3825>

⁷¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3823>

⁷¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3822>

⁷¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3821>

⁷²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3820>

⁷²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3815>

⁷²² <https://github.com/STELLAR-GROUP/hpx/pull/3814>

⁷²³ <https://github.com/STELLAR-GROUP/hpx/pull/3813>

⁷²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3812>

⁷²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3811>

⁷²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3806>

⁷²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3804>

⁷²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3803>

⁷²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3802>

⁷³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3801>

⁷³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3800>

⁷³² <https://github.com/STELLAR-GROUP/hpx/pull/3798>

⁷³³ <https://github.com/STELLAR-GROUP/hpx/pull/3797>

⁷³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3796>

⁷³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3795>

⁷³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3794>

⁷³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3793>

- PR #3791⁷³⁸ - Change throw() to noexcept
- PR #3790⁷³⁹ - Remove deprecated options for 1.3.0 release
- PR #3789⁷⁴⁰ - Remove Boost filesystem compatibility header
- PR #3788⁷⁴¹ - Disabled even more spots that should not execute if networking is disabled
- PR #3787⁷⁴² - Bump minimal boost supported version to 1.61.0
- PR #3786⁷⁴³ - Bump minimum required versions for 1.3.0 release
- PR #3785⁷⁴⁴ - Explicitly set number of jobs for all ninja invocations on CircleCI
- PR #3784⁷⁴⁵ - Fix leak and address sanitizer problems
- PR #3783⁷⁴⁶ - Disabled even more spots that should not execute if networking is disabled
- PR #3782⁷⁴⁷ - Cherry-picked tuple and thread_init_data fixes from #3701
- PR #3781⁷⁴⁸ - Fix generic context coroutines after lazy stack allocation changes
- PR #3780⁷⁴⁹ - Rename hello world examples
- PR #3776⁷⁵⁰ - Sort algorithms now use the supplied chunker to determine the required minimal chunk size
- PR #3775⁷⁵¹ - Disable Boost auto-linking
- PR #3774⁷⁵² - Tag and push stable builds
- PR #3773⁷⁵³ - Enable migration of polymorphic components
- PR #3771⁷⁵⁴ - Fix link to stackoverflow in documentation
- PR #3770⁷⁵⁵ - Replacing constexpr if in brace-serialization code
- PR #3769⁷⁵⁶ - Fix SIGSEGV handler
- PR #3768⁷⁵⁷ - Adding flags to scheduler allowing to control thread stealing and idle back-off
- PR #3767⁷⁵⁸ - Fix help formatting in hpxrun.py
- PR #3765⁷⁵⁹ - Fix a couple of bugs in the thread test
- PR #3764⁷⁶⁰ - Workaround for SFINAE regression in msvc14.2

⁷³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3791>

⁷³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3790>

⁷⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3789>

⁷⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3788>

⁷⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3787>

⁷⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3786>

⁷⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3785>

⁷⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3784>

⁷⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3783>

⁷⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3782>

⁷⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3781>

⁷⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3780>

⁷⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3776>

⁷⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3775>

⁷⁵² <https://github.com/STELLAR-GROUP/hpx/pull/3774>

⁷⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3773>

⁷⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3771>

⁷⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3770>

⁷⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3769>

⁷⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3768>

⁷⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3767>

⁷⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3765>

⁷⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3764>

- PR #3762⁷⁶¹ - Prevent MSVC from prematurely instantiating things
- PR #3761⁷⁶² - Update python scripts to work with python 3
- PR #3760⁷⁶³ - Fix callable vtable for GCC4.9
- PR #3759⁷⁶⁴ - Rename PAGE_SIZE to PAGE_SIZE_ because AppleClang
- PR #3755⁷⁶⁵ - Making sure locks are not held during suspension
- PR #3754⁷⁶⁶ - Disable more code if networking is not available/not enabled
- PR #3752⁷⁶⁷ - Move util::format implementation to source file
- PR #3751⁷⁶⁸ - Fixing problems with `lcos::barrier` and `iostreams`
- PR #3750⁷⁶⁹ - Change error message to take into account `use_guard_page` setting
- PR #3749⁷⁷⁰ - Fix lifetime problem in `run_as_hpx_thread`
- PR #3748⁷⁷¹ - Fixed unusable behavior of the clang code analyzer.
- PR #3747⁷⁷² - Added `PMIX_RANK` to the defaults of `HPX_WITH_PARCELPORT_MPI_ENV`.
- PR #3745⁷⁷³ - Introduced `cache_aligned_data` and `cache_line_data` helper structure
- PR #3742⁷⁷⁴ - Remove more unused functionality from util/logging
- PR #3740⁷⁷⁵ - Fix includes in partitioned vector tests
- PR #3739⁷⁷⁶ - More fixes to make sure that `std::flush` really flushes all output
- PR #3737⁷⁷⁷ - Fix potential shutdown problems
- PR #3736⁷⁷⁸ - Fix `guided_pool_executor` after dataflow changes caused compilation fail
- PR #3734⁷⁷⁹ - Limiting executor
- PR #3732⁷⁸⁰ - More constrained bound constructors
- PR #3730⁷⁸¹ - Attempt to fix deadlocks during component loading
- PR #3729⁷⁸² - Add latch member function `count_up` and `reset`, requested by hpxMP
- PR #3728⁷⁸³ - Send even empty buffers on `hpx::endl` and `hpx::flush`

⁷⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3762>⁷⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3761>⁷⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3760>⁷⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3759>⁷⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3755>⁷⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3754>767 <https://github.com/STELLAR-GROUP/hpx/pull/3752>768 <https://github.com/STELLAR-GROUP/hpx/pull/3751>769 <https://github.com/STELLAR-GROUP/hpx/pull/3750>770 <https://github.com/STELLAR-GROUP/hpx/pull/3749>771 <https://github.com/STELLAR-GROUP/hpx/pull/3748>772 <https://github.com/STELLAR-GROUP/hpx/pull/3747>773 <https://github.com/STELLAR-GROUP/hpx/pull/3745>774 <https://github.com/STELLAR-GROUP/hpx/pull/3742>775 <https://github.com/STELLAR-GROUP/hpx/pull/3740>776 <https://github.com/STELLAR-GROUP/hpx/pull/3739>777 <https://github.com/STELLAR-GROUP/hpx/pull/3737>778 <https://github.com/STELLAR-GROUP/hpx/pull/3736>779 <https://github.com/STELLAR-GROUP/hpx/pull/3734>780 <https://github.com/STELLAR-GROUP/hpx/pull/3732>781 <https://github.com/STELLAR-GROUP/hpx/pull/3730>782 <https://github.com/STELLAR-GROUP/hpx/pull/3729>783 <https://github.com/STELLAR-GROUP/hpx/pull/3728>

- PR #3727⁷⁸⁴ - Adding example demonstrating how to customize the memory management for a component
- PR #3726⁷⁸⁵ - Adding support for passing command line options through the HPX_COMMANDLINE_OPTIONS environment variable
- PR #3722⁷⁸⁶ - Document known broken OpenMPI builds
- PR #3716⁷⁸⁷ - Add barrier reset function, requested by hpxMP for reusing barrier
- PR #3715⁷⁸⁸ - More work on functions and vtables
- PR #3714⁷⁸⁹ - Generate single-page HTML, PDF, manpage from documentation
- PR #3713⁷⁹⁰ - Updating default APEX version to 2.1.2
- PR #3712⁷⁹¹ - Update release procedure
- PR #3710⁷⁹² - Fix the C++11 build, after #3704
- PR #3709⁷⁹³ - Move some component_registry functionality to source file
- PR #3708⁷⁹⁴ - Ignore all locks while handling assertions
- PR #3707⁷⁹⁵ - Remove obsolete hpx runtime executable
- PR #3705⁷⁹⁶ - Fix and simplify make_ready_future overload sets
- PR #3704⁷⁹⁷ - Reduce use of binders
- PR #3703⁷⁹⁸ - Ini
- PR #3702⁷⁹⁹ - Fixing CUDA compiler errors
- PR #3700⁸⁰⁰ - Added barrier::increment function to increase total number of thread
- PR #3697⁸⁰¹ - One more attempt to fix migration...
- PR #3694⁸⁰² - Fixing component migration
- PR #3693⁸⁰³ - Print thread state when getting disallowed value in set_thread_state
- PR #3692⁸⁰⁴ - Only disable constexpr with clang-cuda, not nvcc+gcc
- PR #3691⁸⁰⁵ - Link with libsupc++ if needed for thread_local
- PR #3690⁸⁰⁶ - Remove thousands separators in set_operations_3442 to comply with C++11

⁷⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3727>

⁷⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3726>

⁷⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3722>

⁷⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3716>

⁷⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3715>

⁷⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3714>

⁷⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3713>

⁷⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3712>

⁷⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3710>

⁷⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3709>

⁷⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3708>

⁷⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3707>

⁷⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3705>

⁷⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3704>

⁷⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3703>

⁷⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3702>

⁸⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3700>

⁸⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3697>

⁸⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3694>

⁸⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3693>

⁸⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3692>

⁸⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3691>

⁸⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3690>

- PR #3688⁸⁰⁷ - Decouple serialization from function vtables
- PR #3687⁸⁰⁸ - Fix a couple of test failures
- PR #3686⁸⁰⁹ - Make sure tests.unit.build are run after install on CircleCI
- PR #3685⁸¹⁰ - Revise quickstart CMakeLists.txt explanation
- PR #3684⁸¹¹ - Provide concept emulation for Ranges-TS concepts
- PR #3683⁸¹² - Ignore uninitialized chunks
- PR #3682⁸¹³ - Ignore uninitialized chunks. Check proper indices.
- PR #3680⁸¹⁴ - Ignore uninitialized chunks. Check proper range indices
- PR #3679⁸¹⁵ - Simplify basic action implementations
- PR #3678⁸¹⁶ - Making sure HPX_HAVE_LIBATOMIC is unset before checking
- PR #3677⁸¹⁷ - Fix generated full version number to be usable in expressions
- PR #3674⁸¹⁸ - Reduce functional utilities call depth
- PR #3672⁸¹⁹ - Change new build system to use existing macros related to pseudo dependencies
- PR #3669⁸²⁰ - Remove indirection in function_ref when thread description is disabled
- PR #3668⁸²¹ - Unbreaking async_*cb* tests
- PR #3667⁸²² - Generate version.hpp
- PR #3665⁸²³ - Enabling MPI parcelport for gitlab runners
- PR #3664⁸²⁴ - making clang-tidy work properly again
- PR #3662⁸²⁵ - Attempt to fix exception handling
- PR #3661⁸²⁶ - Move lcos::latch to source file
- PR #3660⁸²⁷ - Fix accidentally explicit gid_type default constructor
- PR #3659⁸²⁸ - Parallel executor latch
- PR #3658⁸²⁹ - Fixing execution_parameters

⁸⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3688>

⁸⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3687>

⁸⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3686>

⁸¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3685>

⁸¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3684>

⁸¹² <https://github.com/STELLAR-GROUP/hpx/pull/3683>

⁸¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3682>

⁸¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3680>

⁸¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3679>

⁸¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3678>

⁸¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3677>

⁸¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3674>

⁸¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3672>

⁸²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3669>

⁸²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3668>

⁸²² <https://github.com/STELLAR-GROUP/hpx/pull/3667>

⁸²³ <https://github.com/STELLAR-GROUP/hpx/pull/3665>

⁸²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3664>

825 <https://github.com/STELLAR-GROUP/hpx/pull/3662>826 <https://github.com/STELLAR-GROUP/hpx/pull/3661>827 <https://github.com/STELLAR-GROUP/hpx/pull/3660>828 <https://github.com/STELLAR-GROUP/hpx/pull/3659>829 <https://github.com/STELLAR-GROUP/hpx/pull/3658>

- PR #3657⁸³⁰ - Avoid dangling references in wait_all
- PR #3656⁸³¹ - Avoiding lifetime problems with sync_put_parcel
- PR #3655⁸³² - Fixing nullptr dereference inside of function
- PR #3652⁸³³ - Attempt to fix thread_map_type definition with C++11
- PR #3650⁸³⁴ - Allowing for end iterator being different from begin iterator
- PR #3649⁸³⁵ - Added architecture identification to cmake to be able to detect timestamp support
- PR #3645⁸³⁶ - Enabling sanitizers on gitlab runner
- PR #3644⁸³⁷ - Attempt to tackle timeouts during startup
- PR #3642⁸³⁸ - Cleanup parallel partitioners
- PR #3640⁸³⁹ - Dataflow now works with functions that return a reference
- PR #3637⁸⁴⁰ - Merging the executor-enabled overloads of shared_future<>::then
- PR #3633⁸⁴¹ - Replace deprecated boost endian macros
- PR #3632⁸⁴² - Add instructions on getting HPX to documentation
- PR #3631⁸⁴³ - Simplify parcel creation
- PR #3630⁸⁴⁴ - Small additions and fixes to release procedure
- PR #3629⁸⁴⁵ - Modular pp
- PR #3627⁸⁴⁶ - Implement util::function_ref
- PR #3626⁸⁴⁷ - Fix cancelable_action_client example
- PR #3625⁸⁴⁸ - Added automatic serialization for simple structs (see #3034)
- PR #3624⁸⁴⁹ - Updating the default order of priority for thread_description
- PR #3621⁸⁵⁰ - Update copyright year and other small formatting fixes
- PR #3620⁸⁵¹ - Adding support for gitlab runner
- PR #3619⁸⁵² - Store debug logs and core dumps on CircleCI

⁸³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3657>

⁸³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3656>

⁸³² <https://github.com/STELLAR-GROUP/hpx/pull/3655>

⁸³³ <https://github.com/STELLAR-GROUP/hpx/pull/3652>

⁸³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3650>

⁸³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3649>

836 <https://github.com/STELLAR-GROUP/hpx/pull/3645>837 <https://github.com/STELLAR-GROUP/hpx/pull/3644>838 <https://github.com/STELLAR-GROUP/hpx/pull/3642>839 <https://github.com/STELLAR-GROUP/hpx/pull/3640>840 <https://github.com/STELLAR-GROUP/hpx/pull/3637>841 <https://github.com/STELLAR-GROUP/hpx/pull/3633>842 <https://github.com/STELLAR-GROUP/hpx/pull/3632>843 <https://github.com/STELLAR-GROUP/hpx/pull/3631>844 <https://github.com/STELLAR-GROUP/hpx/pull/3630>845 <https://github.com/STELLAR-GROUP/hpx/pull/3629>846 <https://github.com/STELLAR-GROUP/hpx/pull/3627>847 <https://github.com/STELLAR-GROUP/hpx/pull/3626>848 <https://github.com/STELLAR-GROUP/hpx/pull/3625>849 <https://github.com/STELLAR-GROUP/hpx/pull/3624>850 <https://github.com/STELLAR-GROUP/hpx/pull/3621>851 <https://github.com/STELLAR-GROUP/hpx/pull/3620>852 <https://github.com/STELLAR-GROUP/hpx/pull/3619>

- PR #3618⁸⁵³ - Various optimizations
- PR #3617⁸⁵⁴ - Fix link to the gpg key (#2)
- PR #3615⁸⁵⁵ - Fix unused variable warnings with networking off
- PR #3614⁸⁵⁶ - Restructuring counter data in scheduler to reduce false sharing
- PR #3613⁸⁵⁷ - Adding support for gitlab runners
- PR #3610⁸⁵⁸ - Don't wait for `stop_condition` in main thread
- PR #3608⁸⁵⁹ - Add inline keyword to `invalid_thread_id` definition for nvcc
- PR #3607⁸⁶⁰ - Adding configuration key that allows one to explicitly add a directory to the component search path
- PR #3606⁸⁶¹ - Add nvcc to exclude constexpress since is it not supported by nvcc
- PR #3605⁸⁶² - Add inline to definition of checkpoint stream operators to fix link error
- PR #3604⁸⁶³ - Use format for string formatting
- PR #3603⁸⁶⁴ - Improve the error message for using to less MAX_CPU_COUNT
- PR #3602⁸⁶⁵ - Improve the error message for to small values of MAX_CPU_COUNT
- PR #3600⁸⁶⁶ - Parallel executor aggregated
- PR #3599⁸⁶⁷ - Making sure networking is disabled for default one-locality-runs
- PR #3596⁸⁶⁸ - Store thread exit functions in `forward_list` instead of `deque` to avoid allocations
- PR #3590⁸⁶⁹ - Fix typo/mistake in thread queue `cleanup_terminated`
- PR #3588⁸⁷⁰ - Fix formatting errors in `launching_and_configuring_hpx_applications.rst`
- PR #3586⁸⁷¹ - Make bind propagate value category
- PR #3585⁸⁷² - Extend Cmake for building hpx as distribution packages (refs #3575)
- PR #3584⁸⁷³ - Untangle function storage from object pointer
- PR #3582⁸⁷⁴ - Towards Modularized HPX
- PR #3580⁸⁷⁵ - Remove extra || in merge.hpp

⁸⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3618>

⁸⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3617>

⁸⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3615>

⁸⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3614>

⁸⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3613>

⁸⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3610>

⁸⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3608>

⁸⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3607>

⁸⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3606>

⁸⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3605>

⁸⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3604>

⁸⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3603>

⁸⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3602>

⁸⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3600>

⁸⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3599>

⁸⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3596>

⁸⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3590>

⁸⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3588>

⁸⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/3586>

⁸⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3585>

⁸⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3584>

⁸⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3582>

⁸⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3580>

- PR #3577⁸⁷⁶ - Partially revert “Remove vtable empty flag”
- PR #3576⁸⁷⁷ - Make sure empty startup/shutdown functions are not being used
- PR #3574⁸⁷⁸ - Make sure DATAPAR settings are conveyed to depending projects
- PR #3573⁸⁷⁹ - Make sure HPX is usable with latest released version of Vc (V1.4.1)
- PR #3572⁸⁸⁰ - Adding test ensuring ticket 3565 is fixed
- PR #3571⁸⁸¹ - Make empty [unique_] function vtable non-dependent
- PR #3566⁸⁸² - Fix compilation with dynamic bitset for CPU masks
- PR #3563⁸⁸³ - Drop util::[unique_] function target_type
- PR #3562⁸⁸⁴ - Removing the target suffixes
- PR #3561⁸⁸⁵ - Replace executor traits return type deduction (keep non-SFINAE)
- PR #3557⁸⁸⁶ - Replace the last usages of boost::atomic
- PR #3556⁸⁸⁷ - Replace boost::scoped_array with std::unique_ptr
- PR #3552⁸⁸⁸ - (Re)move APEX readme
- PR #3548⁸⁸⁹ - Replace boost::scoped_ptr with std::unique_ptr
- PR #3547⁸⁹⁰ - Remove last use of Boost.Signals2
- PR #3544⁸⁹¹ - Post 1.2.0 version bumps
- PR #3543⁸⁹² - added Ubuntu dependency list to readme
- PR #3531⁸⁹³ - Warnings, warnings...
- PR #3527⁸⁹⁴ - Add CircleCI filter for building all tags
- PR #3525⁸⁹⁵ - Segmented algorithms
- PR #3517⁸⁹⁶ - Replace boost::regex with C++11 <regex>
- PR #3514⁸⁹⁷ - Cleaning up the build system
- PR #3505⁸⁹⁸ - Fixing type attribute warning for transfer_action

⁸⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3577>

⁸⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3576>

⁸⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3574>

⁸⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3573>

⁸⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3572>

⁸⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3571>

⁸⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3566>

⁸⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3563>

⁸⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3562>

⁸⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3561>

⁸⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3557>

⁸⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3556>

⁸⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3552>

⁸⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3548>

⁸⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3547>

⁸⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3544>

⁸⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3543>

⁸⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3531>

⁸⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3527>

⁸⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3525>

⁸⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3517>

⁸⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3514>

⁸⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3505>

- PR #3504⁸⁹⁹ - Add support for rpm packaging
- PR #3499⁹⁰⁰ - Improving spinlock pools
- PR #3498⁹⁰¹ - Remove thread specific ptr
- PR #3486⁹⁰² - Fix comparison for expect_connecting_localities config entry
- PR #3469⁹⁰³ - Enable (existing) code for extracting stack pointer on Power platform

2.11.3 HPX V1.2.1 (Feb 19, 2019)

General changes

This is a bugfix release. It contains the following changes:

- Fix compilation on ARM, s390x and 32-bit architectures.
- Fix a critical bug in the `future` implementation.
- Fix several problems in the CMake configuration which affects external projects.
- Add support for Boost 1.69.0.

Closed issues

- Issue #3638⁹⁰⁴ - Build HPX 1.2 with boost 1.69
- Issue #3635⁹⁰⁵ - Non-deterministic crashing on Stampede2
- Issue #3550⁹⁰⁶ - 1>e:000workhpxsrcthrow_exception.cpp(54): error C2440: ‘<function-style-cast>’: cannot convert from ‘boost::system::error_code’ to ‘hpx::exception’
- Issue #3549⁹⁰⁷ - HPX 1.2.0 does not build on i686, but release candidate did
- Issue #3511⁹⁰⁸ - Build on s390x fails
- Issue #3509⁹⁰⁹ - Build on armv7l fails

⁸⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3504>
⁹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3499>
⁹⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3498>
⁹⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3486>
⁹⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3469>
⁹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3638>
⁹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3635>
⁹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3550>
⁹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3549>
⁹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3511>
⁹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3509>

Closed pull requests

- PR #3695⁹¹⁰ - Don't install CMake templates and packaging files
- PR #3666⁹¹¹ - Fixing yet another race in future_data
- PR #3663⁹¹² - Fixing race between setting and getting the value inside future_data
- PR #3648⁹¹³ - Adding timestamp option for S390x platform
- PR #3647⁹¹⁴ - Blind attempt to fix warnings issued by gcc V9
- PR #3611⁹¹⁵ - Include GNUInstallDirs earlier to have it available for subdirectories
- PR #3595⁹¹⁶ - Use GNUInstallDirs lib path in pkgconfig config file
- PR #3593⁹¹⁷ - Add include(GNUInstallDirs) to HPXMacros.cmake
- PR #3591⁹¹⁸ - Fix compilation error on arm7 architecture. Compiles and runs on Fedora 29 on Pi 3.
- PR #3558⁹¹⁹ - Adding constructor *exception(boost::system::error_code const&)*
- PR #3555⁹²⁰ - cmake: make install locations configurable
- PR #3551⁹²¹ - Fix uint64_t causing compilation fail on i686

2.11.4 HPX V1.2.0 (Nov 12, 2018)

General changes

Here are some of the main highlights and changes for this release:

- Thanks to the work of our Google Summer of Code student, Nikunj Gupta, we now have a new implementation of `hpx_main.hpp` on supported platforms (Linux, BSD and MacOS). This is intended to be a less fragile drop-in replacement for the old implementation relying on preprocessor macros. The new implementation does not require changes if you are using the `CMake`⁹²² or `pkg-config`. The old behaviour can be restored by setting `HPX_WITH_DYNAMIC_HPX_MAIN=OFF` during `CMake`⁹²³ configuration. The implementation on Windows is unchanged.
- We have added functionality to allow passing scheduling hints to our schedulers. These will allow us to create executors that for example target a specific NUMA domain or allow for *HPX* threads to be pinned to a particular worker thread.
- We have significantly improved the performance of our futures implementation by making the shared state atomic.

⁹¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3695>

⁹¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3666>

⁹¹² <https://github.com/STELLAR-GROUP/hpx/pull/3663>

⁹¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3648>

⁹¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3647>

⁹¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3611>

⁹¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3595>

⁹¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3593>

⁹¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3591>

⁹¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3558>

⁹²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3555>

⁹²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3551>

⁹²² <https://www.cmake.org>

⁹²³ <https://www.cmake.org>

- We have replaced Boostbook by Sphinx for our documentation. This means the documentation is easier to navigate with built-in search and table of contents. We have also added a quick start section and restructured the documentation to be easier to follow for new users.
- We have added a new option to the `--hpx:threads` command line option. It is now possible to use `cores` to tell *HPX* to only use one worker thread per core, unlike the existing option `all` which uses one worker thread per processing unit (processing unit can be a hyperthread if hyperthreads are available). The default value of `--hpx:threads` has also been changed to `cores` as this leads to better performance in most cases.
- All command line options can now be passed alongside configuration options when initializing *HPX*. This means that some options that were previously only available on the command line can now be set as configuration options.
- HPXMP is a portable, scalable, and flexible application programming interface using the OpenMP specification that supports multi-platform shared memory multiprocessing programming in C and C++. HPXMP can be enabled within *HPX* by setting `DHPX_WITH_HPXMP=ON` during `CMake`⁹²⁴ configuration.
- Two new performance counters were added for measuring the time spent doing background work. `/threads/time/background-work-duration` returns the time spent doing background on a given thread or locality, while `/threads/time/background-overhead` returns the fraction of time spent doing background work with respect to the overall time spent running the scheduler. The new performance counters are disabled by default and can be turned on by setting `HPX_WITH_BACKGROUND_THREAD_COUNTERS=ON` during `CMake`⁹²⁵ configuration.
- The idling behaviour of *HPX* has been tweaked to allow for faster idling. This is useful in interactive applications where the *HPX* worker threads may not have work all the time. This behaviour can be tweaked and turned off as before with `HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF=OFF` during `CMake`⁹²⁶ configuration.
- It is now possible to register callback functions for *HPX* worker thread events. Callbacks can be registered for starting and stopping worker threads, and for when errors occur.

Breaking changes

- The implementation of `hpx_main.hpp` has changed. If you are using custom Makefiles you will need to make changes. Please see the documentation on [using Makefiles](#) for more details.
- The default value of `--hpx:threads` has changed from `all` to `cores`. The new option `cores` only starts one worker thread per core.
- We have dropped support for Boost 1.56 and 1.57. The minimal version of Boost we now test is 1.58.
- Our `boost::format`-based formatting implementation has been revised and replaced with a custom implementation. This changes the formatting syntax and requires changes if you are relying on `hpx::util::format` or `hpx::util::format_to`. The pull request for this change contains more information: [PR #3266](#)⁹²⁷.
- The following deprecated options have now been completely removed: `HPX_WITH_ASYNC_FUNCTION_COMPATIBILITY`, `HPX_WITH_LOCAL_DATAFLOW`, `HPX_WITH_GENERIC_EXECUTION_POLICY`, `HPX_WITH_BOOST_CHRONO_COMPATIBILITY`, `HPX_WITH_EXECUTOR_COMPATIBILITY`, `HPX_WITH_EXECUTION_POLICY_COMPATIBILITY`, and `HPX_WITH_TRANSFORM_REDUCE_COMPATIBILITY`.

⁹²⁴ <https://www.cmake.org>

⁹²⁵ <https://www.cmake.org>

⁹²⁶ <https://www.cmake.org>

⁹²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3266>

Closed issues

- Issue #3538⁹²⁸ - numa handling incorrect for hwloc 2
- Issue #3533⁹²⁹ - Cmake version 3.5.1does not work (git ff26b35 2018-11-06)
- Issue #3526⁹³⁰ - Failed building hpx-1.2.0-rc1 on Ubuntu16.04 x86-64 Virtualbox VM
- Issue #3512⁹³¹ - Build on aarch64 fails
- Issue #3475⁹³² - HPX fails to link if the MPI parcelport is enabled
- Issue #3462⁹³³ - CMake configuration shows a minor and inconsequential failure to create a symlink
- Issue #3461⁹³⁴ - Compilation Problems with the most recent Clang
- Issue #3460⁹³⁵ - Deadlock when create_partitioner fails (assertion fails) in debug mode
- Issue #3455⁹³⁶ - HPX build failing with HWLOC errors on POWER8 with hwloc 1.8
- Issue #3438⁹³⁷ - HPX no longer builds on IBM POWER8
- Issue #3426⁹³⁸ - hpx build failed on MacOS
- Issue #3424⁹³⁹ - CircleCI builds broken for forked repositories
- Issue #3422⁹⁴⁰ - Benchmarks in tests.performance.local are not run nightly
- Issue #3408⁹⁴¹ - CMake Targets for HPX
- Issue #3399⁹⁴² - processing unit out of bounds
- Issue #3395⁹⁴³ - Floating point bug in hpx/runtime/threads/policies/scheduler_base.hpp
- Issue #3378⁹⁴⁴ - compile error with lcos::communicator
- Issue #3376⁹⁴⁵ - Failed to build HPX with APEX using clang
- Issue #3366⁹⁴⁶ - Adapted Safe_Object example fails for –hpx:threads > 1
- Issue #3360⁹⁴⁷ - Segmentation fault when passing component id as parameter
- Issue #3358⁹⁴⁸ - HPX runtime hangs after multiple (~thousands) start-stop sequences
- Issue #3352⁹⁴⁹ - Support TCP provider in libfabric ParcelPort

928 <https://github.com/STELLAR-GROUP/hpx/issues/3538>

929 <https://github.com/STELLAR-GROUP/hpx/issues/3533>

930 <https://github.com/STELLAR-GROUP/hpx/issues/3526>

931 <https://github.com/STELLAR-GROUP/hpx/issues/3512>

932 <https://github.com/STELLAR-GROUP/hpx/issues/3475>

933 <https://github.com/STELLAR-GROUP/hpx/issues/3462>

934 <https://github.com/STELLAR-GROUP/hpx/issues/3461>

935 <https://github.com/STELLAR-GROUP/hpx/issues/3460>

936 <https://github.com/STELLAR-GROUP/hpx/issues/3455>

937 <https://github.com/STELLAR-GROUP/hpx/issues/3438>

938 <https://github.com/STELLAR-GROUP/hpx/issues/3426>

939 <https://github.com/STELLAR-GROUP/hpx/issues/3424>

940 <https://github.com/STELLAR-GROUP/hpx/issues/3422>

941 <https://github.com/STELLAR-GROUP/hpx/issues/3408>

942 <https://github.com/STELLAR-GROUP/hpx/issues/3399>

943 <https://github.com/STELLAR-GROUP/hpx/issues/3395>

944 <https://github.com/STELLAR-GROUP/hpx/issues/3378>

945 <https://github.com/STELLAR-GROUP/hpx/issues/3376>

946 <https://github.com/STELLAR-GROUP/hpx/issues/3366>

947 <https://github.com/STELLAR-GROUP/hpx/issues/3360>

948 <https://github.com/STELLAR-GROUP/hpx/issues/3358>

949 <https://github.com/STELLAR-GROUP/hpx/issues/3352>

- Issue #3342⁹⁵⁰ - undefined reference to __atomic_load_16
- Issue #3339⁹⁵¹ - setting command line options/flags from init cfg is not obvious
- Issue #3325⁹⁵² - AGAS migrates components prematurely
- Issue #3321⁹⁵³ - hpx bad_parameter handling is awful
- Issue #3318⁹⁵⁴ - Benchmarks fail to build with C++11
- Issue #3304⁹⁵⁵ - hpx::threads::run_as_hpx_thread does not properly handle exceptions
- Issue #3300⁹⁵⁶ - Setting pu step or offset results in no threads in default pool
- Issue #3297⁹⁵⁷ - Crash with APEX when running Phylanx lra_csv with > 1 thread
- Issue #3296⁹⁵⁸ - Building HPX with APEX configuration gives compiler warnings
- Issue #3290⁹⁵⁹ - make tests failing at hello_world_component
- Issue #3285⁹⁶⁰ - possible compilation error when “using namespace std;” is defined before including “hpx” headers files
- Issue #3280⁹⁶¹ - HPX fails on OSX
- Issue #3272⁹⁶² - CircleCI does not upload generated docker image any more
- Issue #3270⁹⁶³ - Error when compiling CUDA examples
- Issue #3267⁹⁶⁴ - tests.unit.host_.block_allocator fails occasionally
- Issue #3264⁹⁶⁵ - Possible move to Sphinx for documentation
- Issue #3263⁹⁶⁶ - Documentation improvements
- Issue #3259⁹⁶⁷ - set_parcel_write_handler test fails occasionally
- Issue #3258⁹⁶⁸ - Links to source code in documentation are broken
- Issue #3247⁹⁶⁹ - Rare tests.unit.host_.block_allocator test failure on 1.1.0-rc1
- Issue #3244⁹⁷⁰ - Slowing down and speeding up an interval_timer
- Issue #3215⁹⁷¹ - Cannot build both tests and examples on MSVC with pseudo-dependencies enabled
- Issue #3195⁹⁷² - Unnecessary customization point route causing performance penalty

⁹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3342>

⁹⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/3339>

⁹⁵² <https://github.com/STELLAR-GROUP/hpx/issues/3325>

⁹⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/3321>

⁹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3318>

⁹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3304>

⁹⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3300>

⁹⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3297>

⁹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3296>

⁹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3290>

⁹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3285>

⁹⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/3280>

⁹⁶² <https://github.com/STELLAR-GROUP/hpx/issues/3272>

⁹⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/3270>

⁹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3267>

⁹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3264>

⁹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3263>

⁹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3259>

⁹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3258>

⁹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3247>

⁹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3244>

⁹⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/3215>

⁹⁷² <https://github.com/STELLAR-GROUP/hpx/issues/3195>

- Issue #3088⁹⁷³ - A strange thing in parallel::sort.
- Issue #2650⁹⁷⁴ - libfabric support for passive endpoints
- Issue #1205⁹⁷⁵ - TSS is broken

Closed pull requests

- PR #3542⁹⁷⁶ - Fix numa lookup from pu when using hwloc 2.x
- PR #3541⁹⁷⁷ - Fixing the build system of the MPI parcelport
- PR #3540⁹⁷⁸ - Updating HPX people section
- PR #3539⁹⁷⁹ - Splitting test to avoid OOM on CircleCI
- PR #3537⁹⁸⁰ - Fix guided exec
- PR #3536⁹⁸¹ - Updating grants which support the LSU team
- PR #3535⁹⁸² - Fix hiding of docker credentials
- PR #3534⁹⁸³ - Fixing #3533
- PR #3532⁹⁸⁴ - fixing minor doc typo –hpx:print-counter-at arg
- PR #3530⁹⁸⁵ - Changing APEX default tag to v2.1.0
- PR #3529⁹⁸⁶ - Remove leftover security options and documentation
- PR #3528⁹⁸⁷ - Fix hwloc version check
- PR #3524⁹⁸⁸ - Do not build guided pool examples with older GCC compilers
- PR #3523⁹⁸⁹ - Fix logging regression
- PR #3522⁹⁹⁰ - Fix more warnings
- PR #3521⁹⁹¹ - Fixing argument handling in induction and reduction clauses for parallel::for_loop
- PR #3520⁹⁹² - Remove docs symlink and versioned docs folders
- PR #3519⁹⁹³ - hpxMP release
- PR #3518⁹⁹⁴ - Change all steps to use new docker image on CircleCI

⁹⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/3088>

⁹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2650>

⁹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1205>

⁹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3542>

⁹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3541>

⁹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3540>

⁹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3539>

⁹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3537>

⁹⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3536>

⁹⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3535>

⁹⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3534>

⁹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3532>

⁹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3530>

⁹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3529>

⁹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3528>

⁹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3524>

⁹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3523>

⁹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3522>

⁹⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3521>

⁹⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3520>

⁹⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3519>

⁹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3518>

- PR #3516⁹⁹⁵ - Drop usage of deprecated facilities removed in C++17
- PR #3515⁹⁹⁶ - Remove remaining uses of Boost.TypeTraits
- PR #3513⁹⁹⁷ - Fixing a CMake problem when trying to use libfabric
- PR #3508⁹⁹⁸ - Remove memory_block component
- PR #3507⁹⁹⁹ - Propagating the MPI compile definitions to all relevant targets
- PR #3503¹⁰⁰⁰ - Update documentation colors and logo
- PR #3502¹⁰⁰¹ - Fix bogus `throws` bindings in scheduled_thread_pool_Impl
- PR #3501¹⁰⁰² - Split parallel::remove_if tests to avoid OOM on CircleCI
- PR #3500¹⁰⁰³ - Support NONAMEPREFIX in add_hpx_library()
- PR #3497¹⁰⁰⁴ - Note that cuda support requires cmake 3.9
- PR #3495¹⁰⁰⁵ - Fixing dataflow
- PR #3493¹⁰⁰⁶ - Remove deprecated options for 1.2.0 part 2
- PR #3492¹⁰⁰⁷ - Add CUDA_LINK_LIBRARIES_KEYWORD to allow PRIVATE keyword in linkage t...
- PR #3491¹⁰⁰⁸ - Changing Base docker image
- PR #3490¹⁰⁰⁹ - Don't create tasks immediately with hpx::apply
- PR #3489¹⁰¹⁰ - Remove deprecated options for 1.2.0
- PR #3488¹⁰¹¹ - Revert "Use BUILD_INTERFACE generator expression to fix cmake flag exports"
- PR #3487¹⁰¹² - Revert "Fixing type attribute warning for transfer_action"
- PR #3485¹⁰¹³ - Use BUILD_INTERFACE generator expression to fix cmake flag exports
- PR #3483¹⁰¹⁴ - Fixing type attribute warning for transfer_action
- PR #3481¹⁰¹⁵ - Remove unused variables
- PR #3480¹⁰¹⁶ - Towards a more lightweight transfer action
- PR #3479¹⁰¹⁷ - Fix FLAGS - Use correct version of target_compile_options

⁹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3516>

⁹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3515>

⁹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3513>

⁹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3508>

⁹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3507>

¹⁰⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3503>

1001 <https://github.com/STELLAR-GROUP/hpx/pull/3502>1002 <https://github.com/STELLAR-GROUP/hpx/pull/3501>1003 <https://github.com/STELLAR-GROUP/hpx/pull/3500>1004 <https://github.com/STELLAR-GROUP/hpx/pull/3497>1005 <https://github.com/STELLAR-GROUP/hpx/pull/3495>1006 <https://github.com/STELLAR-GROUP/hpx/pull/3493>1007 <https://github.com/STELLAR-GROUP/hpx/pull/3492>1008 <https://github.com/STELLAR-GROUP/hpx/pull/3491>1009 <https://github.com/STELLAR-GROUP/hpx/pull/3490>1010 <https://github.com/STELLAR-GROUP/hpx/pull/3489>1011 <https://github.com/STELLAR-GROUP/hpx/pull/3488>1012 <https://github.com/STELLAR-GROUP/hpx/pull/3487>1013 <https://github.com/STELLAR-GROUP/hpx/pull/3485>1014 <https://github.com/STELLAR-GROUP/hpx/pull/3483>1015 <https://github.com/STELLAR-GROUP/hpx/pull/3481>1016 <https://github.com/STELLAR-GROUP/hpx/pull/3480>1017 <https://github.com/STELLAR-GROUP/hpx/pull/3479>

- PR #3478¹⁰¹⁸ - Making sure the application's exit code is properly propagated back to the OS
- PR #3476¹⁰¹⁹ - Don't print docker credentials as part of the environment.
- PR #3473¹⁰²⁰ - Fixing invalid cmake code if no jemalloc prefix was given
- PR #3472¹⁰²¹ - Attempting to work around recent clang test compilation failures
- PR #3471¹⁰²² - Enable jemalloc on windows
- PR #3470¹⁰²³ - Updates readme
- PR #3468¹⁰²⁴ - Avoid hang if there is an exception thrown during startup
- PR #3467¹⁰²⁵ - Add compiler specific fallthrough attributes if C++17 attribute is not available
- PR #3466¹⁰²⁶ - - bugfix : fix compilation with llvm-7.0
- PR #3465¹⁰²⁷ - This patch adds various optimizations extracted from the thread_local_allocator work
- PR #3464¹⁰²⁸ - Check for forked repos in CircleCI docker push step
- PR #3463¹⁰²⁹ - - cmake : create the parent directory before symlinking
- PR #3459¹⁰³⁰ - Remove unused/incomplete functionality from util/logging
- PR #3458¹⁰³¹ - Fix a problem with scope of CMAKE_CXX_FLAGS and hpx_add_compile_flag
- PR #3457¹⁰³² - Fixing more size_t -> int16_t (and similar) warnings
- PR #3456¹⁰³³ - Add #ifdefs to topology.cpp to support old hwloc versions again
- PR #3454¹⁰³⁴ - Fixing warnings related to silent conversion of size_t -> int16_t
- PR #3451¹⁰³⁵ - Add examples as unit tests
- PR #3450¹⁰³⁶ - Constexpr-fying bind and other functional facilities
- PR #3446¹⁰³⁷ - Fix some thread suspension timeouts
- PR #3445¹⁰³⁸ - Fix various warnings
- PR #3443¹⁰³⁹ - Only enable service pool config options if pools are enabled
- PR #3441¹⁰⁴⁰ - Fix missing closing brackets in documentation

¹⁰¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3478>

¹⁰¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3476>

¹⁰²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3473>

¹⁰²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3472>

¹⁰²² <https://github.com/STELLAR-GROUP/hpx/pull/3471>

¹⁰²³ <https://github.com/STELLAR-GROUP/hpx/pull/3470>

¹⁰²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3468>

¹⁰²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3467>

¹⁰²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3466>

¹⁰²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3465>

¹⁰²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3464>

¹⁰²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3463>

¹⁰³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3459>

¹⁰³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3458>

1032 <https://github.com/STELLAR-GROUP/hpx/pull/3457>1033 <https://github.com/STELLAR-GROUP/hpx/pull/3456>1034 <https://github.com/STELLAR-GROUP/hpx/pull/3454>1035 <https://github.com/STELLAR-GROUP/hpx/pull/3451>1036 <https://github.com/STELLAR-GROUP/hpx/pull/3450>1037 <https://github.com/STELLAR-GROUP/hpx/pull/3446>1038 <https://github.com/STELLAR-GROUP/hpx/pull/3445>1039 <https://github.com/STELLAR-GROUP/hpx/pull/3443>1040 <https://github.com/STELLAR-GROUP/hpx/pull/3441>

- PR #3439¹⁰⁴¹ - Use correct MPI CXX libraries for MPI parcelport
- PR #3436¹⁰⁴² - Add projection function to find_* (and fix very bad bug)
- PR #3435¹⁰⁴³ - Fixing 1205
- PR #3434¹⁰⁴⁴ - Fix threads cores
- PR #3433¹⁰⁴⁵ - Add Heise Online to release announcement list
- PR #3432¹⁰⁴⁶ - Don't track task dependencies for distributed runs
- PR #3431¹⁰⁴⁷ - Circle CI setting changes for hpxMP
- PR #3430¹⁰⁴⁸ - Fix unused params warning
- PR #3429¹⁰⁴⁹ - One thread per core
- PR #3428¹⁰⁵⁰ - This suppresses a deprecation warning that is being issued by MSVC 19.15.26726
- PR #3427¹⁰⁵¹ - Fixes #3426
- PR #3425¹⁰⁵² - Use source cache and workspace between job steps on CircleCI
- PR #3421¹⁰⁵³ - Add CDash timing output to future overhead test (for graphs)
- PR #3420¹⁰⁵⁴ - Add guided_pool_executor
- PR #3419¹⁰⁵⁵ - Fix typo in CircleCI config
- PR #3418¹⁰⁵⁶ - Add sphinx documentation
- PR #3415¹⁰⁵⁷ - Scheduler NUMA hint and shared priority scheduler
- PR #3414¹⁰⁵⁸ - Adding step to synchronize the APEX release
- PR #3413¹⁰⁵⁹ - Fixing multiple defines of APEX_HAVE_HPX
- PR #3412¹⁰⁶⁰ - Fixes linking with libhpx_wrap error with BSD and Windows based systems
- PR #3410¹⁰⁶¹ - Fix typo in CMakeLists.txt
- PR #3409¹⁰⁶² - Fix brackets and indentation in existing_performance_counters.qbk
- PR #3407¹⁰⁶³ - Fix unused param and extra ; warnings emitted by gcc 8.x

¹⁰⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3439>

¹⁰⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3436>

¹⁰⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3435>

¹⁰⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3434>

¹⁰⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3433>

¹⁰⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3432>

¹⁰⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3431>

¹⁰⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3430>

¹⁰⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3429>

¹⁰⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3428>

¹⁰⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3427>

¹⁰⁵² <https://github.com/STELLAR-GROUP/hpx/pull/3425>

¹⁰⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3421>

¹⁰⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3420>

¹⁰⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3419>

¹⁰⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3418>

¹⁰⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3415>

¹⁰⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3414>

¹⁰⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3413>

¹⁰⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3412>

¹⁰⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3410>

¹⁰⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3409>

¹⁰⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3407>

- PR #3406¹⁰⁶⁴ - Adding thread local allocator and use it for future shared states
- PR #3405¹⁰⁶⁵ - Adding DHPX_HAVE_THREAD_LOCAL_STORAGE=ON to builds
- PR #3404¹⁰⁶⁶ - fixing multiple definition of main() in linux
- PR #3402¹⁰⁶⁷ - Allow debug option to be enabled only for Linux systems with dynamic main on
- PR #3401¹⁰⁶⁸ - Fix cuda_future_helper.h when compiling with C++11
- PR #3400¹⁰⁶⁹ - Fix floating point exception scheduler_base idle backoff
- PR #3398¹⁰⁷⁰ - Atomic future state
- PR #3397¹⁰⁷¹ - Fixing code for older gcc versions
- PR #3396¹⁰⁷² - Allowing to register thread event functions (start/stop/error)
- PR #3394¹⁰⁷³ - Fix small mistake in primary_namespace_server.cpp
- PR #3393¹⁰⁷⁴ - Explicitly instantiate configured schedulers
- PR #3392¹⁰⁷⁵ - Add performance counters background overhead and background work duration
- PR #3391¹⁰⁷⁶ - Adapt integration of HPXMP to latest build system changes
- PR #3390¹⁰⁷⁷ - Make AGAS measurements optional
- PR #3389¹⁰⁷⁸ - Fix deadlock during shutdown
- PR #3388¹⁰⁷⁹ - Add several functionalities allowing to optimize synchronous action invocation
- PR #3387¹⁰⁸⁰ - Add cmake option to opt out of fail-compile tests
- PR #3386¹⁰⁸¹ - Adding support for boost::container::small_vector to dataflow
- PR #3385¹⁰⁸² - Adds Debug option for hpx initializing from main
- PR #3384¹⁰⁸³ - This hopefully fixes two tests that occasionally fail
- PR #3383¹⁰⁸⁴ - Making sure thread local storage is enable for hpxMP
- PR #3382¹⁰⁸⁵ - Fix usage of HPX_CAPTURE together with default value capture [=]
- PR #3381¹⁰⁸⁶ - Replace undefined instantiations of uniform_int_distribution

¹⁰⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3406>

¹⁰⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3405>

¹⁰⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3404>

¹⁰⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3402>

¹⁰⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3401>

¹⁰⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3400>

¹⁰⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3398>

¹⁰⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/3397>

¹⁰⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3396>

¹⁰⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3394>

¹⁰⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3393>

¹⁰⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3392>

¹⁰⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3391>

¹⁰⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3390>

¹⁰⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3389>

¹⁰⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3388>

¹⁰⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3387>

¹⁰⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3386>

¹⁰⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3385>

¹⁰⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3384>

¹⁰⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3383>

¹⁰⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3382>

¹⁰⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3381>

- PR #3380¹⁰⁸⁷ - Add missing semicolons to uses of HPX_COMPILER_FENCE
- PR #3379¹⁰⁸⁸ - Fixing #3378
- PR #3377¹⁰⁸⁹ - Adding build system support to integrate hpxmp into hpx at the user's machine
- PR #3375¹⁰⁹⁰ - Replacing wrapper for __libc_start_main with main
- PR #3374¹⁰⁹¹ - Adds hpx_wrap to HPX_LINK_LIBRARIES which links only when specified.
- PR #3373¹⁰⁹² - Forcing cache settings in HPXConfig.cmake to guarantee updated values
- PR #3372¹⁰⁹³ - Fix some more c++11 build problems
- PR #3371¹⁰⁹⁴ - Adds HPX_LINKER_FLAGS to HPX applications without editing their source codes
- PR #3370¹⁰⁹⁵ - util::format: add type_specifier<> specializations for %!s(MISSING) and %!l(MISSING)s
- PR #3369¹⁰⁹⁶ - Adding configuration option to allow explicit disable of the new hpx_main feature on Linux
- PR #3368¹⁰⁹⁷ - Updates doc with recent hpx_wrap implementation
- PR #3367¹⁰⁹⁸ - Adds Mac OS implementation to hpx_main.hpp
- PR #3365¹⁰⁹⁹ - Fix order of hpx libs in HPX_CONF_LIBRARIES.
- PR #3363¹¹⁰⁰ - Apex fixing null wrapper
- PR #3361¹¹⁰¹ - Making sure all parcels get destroyed on an HPX thread (TCP pp)
- PR #3359¹¹⁰² - Feature/improveerrorforcompiler
- PR #3357¹¹⁰³ - Static/dynamic executable implementation
- PR #3355¹¹⁰⁴ - Reverting changes introduced by #3283 as those make applications hang
- PR #3354¹¹⁰⁵ - Add external dependencies to HPX_LIBRARY_DIR
- PR #3353¹¹⁰⁶ - Fix libfabric tcp
- PR #3351¹¹⁰⁷ - Move obsolete header to tests directory.
- PR #3350¹¹⁰⁸ - Renaming two functions to avoid problem described in #3285
- PR #3349¹¹⁰⁹ - Make idle backoff exponential with maximum sleep time

¹⁰⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3380>

¹⁰⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3379>

¹⁰⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3377>

¹⁰⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3375>

¹⁰⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3374>

¹⁰⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3373>

¹⁰⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3372>

¹⁰⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3371>

¹⁰⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3370>

¹⁰⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3369>

¹⁰⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3368>

¹⁰⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3367>

¹⁰⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3365>

¹¹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3363>

¹¹⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3361>

¹¹⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3359>

¹¹⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3357>

¹¹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3355>

¹¹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3354>

¹¹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3353>

¹¹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3351>

¹¹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3350>

¹¹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3349>

- PR #3347¹¹¹⁰ - Replace *simple_component** with *component** in the Documentation
- PR #3346¹¹¹¹ - Fix CMakeLists.txt example in quick start
- PR #3345¹¹¹² - Fix automatic setting of HPX_MORE_THAN_64_THREADS
- PR #3344¹¹¹³ - Reduce amount of information printed for unknown command line options
- PR #3343¹¹¹⁴ - Safeguard HPX against destruction in global contexts
- PR #3341¹¹¹⁵ - Allowing for all command line options to be used as configuration settings
- PR #3340¹¹¹⁶ - Always convert inspect results to JUnit XML
- PR #3336¹¹¹⁷ - Only run docker push on master on CircleCI
- PR #3335¹¹¹⁸ - Update description of hpx.os_threads config parameter.
- PR #3334¹¹¹⁹ - Making sure early logging settings don't get mixed with others
- PR #3333¹¹²⁰ - Update CMake links and versions in documentation
- PR #3332¹¹²¹ - Add notes on target suffixes to CMake documentation
- PR #3331¹¹²² - Add quickstart section to documentation
- PR #3330¹¹²³ - Rename resource_partitioner test to avoid conflicts with pseudodependencies
- PR #3328¹¹²⁴ - Making sure object is pinned while executing actions, even if action returns a future
- PR #3327¹¹²⁵ - Add missing std::forward to tuple.hpp
- PR #3326¹¹²⁶ - Make sure logging is up and running while modules are being discovered.
- PR #3324¹¹²⁷ - Replace C++14 overload of std::equal with C++11 code.
- PR #3323¹¹²⁸ - Fix a missing apex thread data (wrapper) initialization
- PR #3320¹¹²⁹ - Adding support for -std=c++2a (define *HPX_WITH_CXX2A=On*)
- PR #3319¹¹³⁰ - Replacing C++14 feature with equivalent C++11 code
- PR #3317¹¹³¹ - Fix compilation with VS 15.7.1 and /std:c++latest
- PR #3316¹¹³² - Fix includes for 1d_stencil_*_omp examples

¹¹¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3347>

¹¹¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3346>

¹¹¹² <https://github.com/STELLAR-GROUP/hpx/pull/3345>

¹¹¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3344>

¹¹¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3343>

¹¹¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3341>

¹¹¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3340>

¹¹¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3336>

¹¹¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3335>

¹¹¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3334>

¹¹²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3333>

¹¹²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3332>

¹¹²² <https://github.com/STELLAR-GROUP/hpx/pull/3331>

¹¹²³ <https://github.com/STELLAR-GROUP/hpx/pull/3330>

¹¹²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3328>

¹¹²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3327>

¹¹²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3326>

¹¹²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3324>

¹¹²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3323>

¹¹²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3320>

¹¹³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3319>

¹¹³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3317>

¹¹³² <https://github.com/STELLAR-GROUP/hpx/pull/3316>

- PR #3314¹¹³³ - Remove some unused parameter warnings
- PR #3313¹¹³⁴ - Fix pu-step and pu-offset command line options
- PR #3312¹¹³⁵ - Add conversion of inspect reports to JUnit XML
- PR #3311¹¹³⁶ - Fix escaping of closing braces in format specification syntax
- PR #3310¹¹³⁷ - Don't overwrite user settings with defaults in registration database
- PR #3309¹¹³⁸ - Fixing potential stack overflow for dataflow
- PR #3308¹¹³⁹ - This updates the .clang-format configuration file to utilize newer features
- PR #3306¹¹⁴⁰ - Marking migratable objects in their gid to allow not handling migration in AGAS
- PR #3305¹¹⁴¹ - Add proper exception handling to run_as_hpx_thread
- PR #3303¹¹⁴² - Changed std::rand to a better inbuilt PRNG Generator
- PR #3302¹¹⁴³ - All non-migratable (simple) components now encode their lva and component type in their gid
- PR #3301¹¹⁴⁴ - Add nullptr_t overloads to resource partitioner
- PR #3298¹¹⁴⁵ - Apex task wrapper memory bug
- PR #3295¹¹⁴⁶ - Fix mistakes after merge of CircleCI config
- PR #3294¹¹⁴⁷ - Fix partitioned vector include in partitioned_vector_find tests
- PR #3293¹¹⁴⁸ - Adding emplace support to promise and make_ready_future
- PR #3292¹¹⁴⁹ - Add new cuda kernel synchronization with hpx::future demo
- PR #3291¹¹⁵⁰ - Fixes #3290
- PR #3289¹¹⁵¹ - Fixing Docker image creation
- PR #3288¹¹⁵² - Avoid allocating shared state for wait_all
- PR #3287¹¹⁵³ - Fixing /scheduler/utilization/instantaneous performance counter
- PR #3286¹¹⁵⁴ - dataflow() and future::then() use sync policy where possible
- PR #3284¹¹⁵⁵ - Background thread can use relaxed atomics to manipulate thread state

¹¹³³ <https://github.com/STELLAR-GROUP/hpx/pull/3314>

¹¹³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3313>

¹¹³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3312>

¹¹³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3311>

¹¹³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3310>

¹¹³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3309>

¹¹³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3308>

¹¹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3306>

¹¹⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3305>

¹¹⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3303>

¹¹⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3302>

¹¹⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3301>

¹¹⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3298>

¹¹⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3295>

¹¹⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3294>

¹¹⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3293>

¹¹⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3292>

¹¹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3291>

¹¹⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3289>

¹¹⁵² <https://github.com/STELLAR-GROUP/hpx/pull/3288>

¹¹⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3287>

¹¹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3286>

¹¹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3284>

- PR #3283¹¹⁵⁶ - Do not unwrap ready future
- PR #3282¹¹⁵⁷ - Fix virtual method override warnings in static schedulers
- PR #3281¹¹⁵⁸ - Disable set_area_membind_nodeset for OSX
- PR #3279¹¹⁵⁹ - Add two variations to the future_overhead benchmark
- PR #3278¹¹⁶⁰ - Fix circleci workspace
- PR #3277¹¹⁶¹ - Support external plugins
- PR #3276¹¹⁶² - Fix missing parenthesis in hello_compute.cu.
- PR #3274¹¹⁶³ - Reinit counters synchronously in reinit_counters test
- PR #3273¹¹⁶⁴ - Splitting tests to avoid compiler OOM
- PR #3271¹¹⁶⁵ - Remove leftover code from context_generic_context.hpp
- PR #3269¹¹⁶⁶ - Fix bulk_construct with count = 0
- PR #3268¹¹⁶⁷ - Replace constexpr with HPX_CXX14_CONSTEXPR and HPX_CONSTEXPR
- PR #3266¹¹⁶⁸ - Replace boost::format with custom sprintf-based implementation
- PR #3265¹¹⁶⁹ - Split parallel tests on CircleCI
- PR #3262¹¹⁷⁰ - Making sure documentation correctly links to source files
- PR #3261¹¹⁷¹ - Apex refactoring fix rebind
- PR #3260¹¹⁷² - Isolate performance counter parser into a separate TU
- PR #3256¹¹⁷³ - Post 1.1.0 version bumps
- PR #3254¹¹⁷⁴ - Adding trait for actions allowing to make runtime decision on whether to execute it directly
- PR #3253¹¹⁷⁵ - Bump minimal supported Boost to 1.58.0
- PR #3251¹¹⁷⁶ - Adds new feature: changing interval used in interval_timer (issue 3244)
- PR #3239¹¹⁷⁷ - Changing std::rand() to a better inbuilt PRNG generator.
- PR #3234¹¹⁷⁸ - Disable background thread when networking is off

¹¹⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3283>

¹¹⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3282>

¹¹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3281>

¹¹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3279>

¹¹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3278>

¹¹⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3277>

¹¹⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3276>

¹¹⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3274>

¹¹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3273>

¹¹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3271>

¹¹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3269>

¹¹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3268>

¹¹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3266>

¹¹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3265>

¹¹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3262>

¹¹⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/3261>

¹¹⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3260>

¹¹⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3256>

¹¹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3254>

¹¹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3253>

¹¹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3251>

¹¹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3239>

¹¹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3234>

- PR #3232¹¹⁷⁹ - Clean up suspension tests
- PR #3230¹¹⁸⁰ - Add optional scheduler mode parameter to create_thread_pool function
- PR #3228¹¹⁸¹ - Allow suspension also on static schedulers
- PR #3163¹¹⁸² - libfabric parcelport w/o HPX_PARCELPORT_LIBFABRIC_ENDPOINT_RDM
- PR #3036¹¹⁸³ - Switching to CircleCI 2.0

2.11.5 HPX V1.1.0 (Mar 24, 2018)

General changes

Here are some of the main highlights and changes for this release (in no particular order):

- We have changed the way *HPX* manages the processing units on a node. We do not longer implicitly bind all available cores to a single thread pool. The user has now full control over what processing units are bound to what thread pool, each with a separate scheduler. It is now also possible to create your own scheduler implementation and control what processing units this scheduler should use. We added the `hpx::resource::partitioner` that manages all available processing units and assigns resources to the used thread pools. Thread pools can be now be suspended/resumed independently. This functionality helps in running *HPX* concurrently to code that is directly relying on OpenMP¹¹⁸⁴ and/or MPI¹¹⁸⁵.
- We have continued to implement various parallel algorithms. *HPX* now almost completely implements all of the parallel algorithms as specified by the C++17 Standard¹¹⁸⁶. We have also continued to implement these algorithms for the distributed use case (for segmented data structures, such as `hpx::partitioned_vector`).
- Added a compatibility layer for `std::thread`, `std::mutex`, and `std::condition_variable` allowing for the code to use those facilities where available and to fall back to the corresponding Boost facilities otherwise. The CMake¹¹⁸⁷ configuration option `-DHPX_WITH_THREAD_COMPATIBILITY=On` can be used to force using the Boost equivalents.
- The parameter sequence for the `hpx::parallel::transform_inclusive_scan` overload taking one iterator range has changed (again) to match the changes this algorithm has undergone while being moved to C++17. The old overloads can be still enabled at configure time by passing `-DHPX_WITH_TRANSFORM_REDUCE_COMPATIBILITY=On` to CMake¹¹⁸⁸.
- The parameter sequence for the `hpx::parallel::inclusive_scan` overload taking one iterator range has changed to match the changes this algorithm has undergone while being moved to C++17. The old overloads can be still enabled at configure time by passing `-DHPX_WITH_INCLUSIVE_SCAN_COMPATIBILITY=On` to CMake.
- Added a helper facility `hpx::local_new` which is equivalent to `hpx::new_` except that it creates components locally only. As a consequence, the used component constructor may accept non-serializable argument types and/or non-const references or pointers.
- Removed the (broken) component type `hpx::lcos::queue<T>`. The old type is still available at configure time by passing `-DHPX_WITH_QUEUE_COMPATIBILITY=On` to CMake.

¹¹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3232>

¹¹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3230>

¹¹⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3228>

¹¹⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3163>

¹¹⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3036>

¹¹⁸⁴ <https://openmp.org/wp/>

¹¹⁸⁵ https://en.wikipedia.org/wiki/Message_Passing_Interface

¹¹⁸⁶ <http://www.open-std.org/jtc1/sc22/wg21>

¹¹⁸⁷ <https://www.cmake.org>

¹¹⁸⁸ <https://www.cmake.org>

- The parallel algorithms adopted for C++17 restrict the iterator categories usable with those to at least forward iterators. Our implementation of the parallel algorithms was supporting input iterators (and output iterators) as well by simply falling back to sequential execution. We have now made our implementations conforming by requiring at least forward iterators. In order to enable the old behavior use the compatibility option `-DPHX_WITH_ALGORITHM_INPUT_ITERATOR_SUPPORT=On` on the [CMake](#)¹¹⁸⁹ command line.
- We have added the functionalities allowing for LCOs being implemented using (simple) components. Before LCOs had to always be implemented using managed components.
- User defined components don't have to be default-constructible anymore. Return types from actions don't have to be default-constructible anymore either. Our serialization layer now in general supports non-default-constructible types.
- We have added a new launch policy `hpx::launch::lazy` that allows one to defer the decision on what launch policy to use to the point of execution. This policy is initialized with a function (object) that – when invoked – is expected to produce the desired launch policy.

Breaking changes

- We have dropped support for the gcc compiler version V4.8. The minimal gcc version we now test on is gcc V4.9. The minimally required version of [CMake](#)¹¹⁹⁰ is now V3.3.2.
- We have dropped support for the Visual Studio 2013 compiler version. The minimal Visual Studio version we now test on is Visual Studio 2015.5.
- We have dropped support for the Boost V1.51-V1.54. The minimal version of Boost we now test is Boost V1.55.
- We have dropped support for the `hpx::util::unwrapped` API. `hpx::util::unwrapped` will stay functional to some degree, until it finally gets removed in a later version of HPX. The functional usage of `hpx::util::unwrapped` should be changed to the new `hpx::util::unwrapping` function whereas the immediate usage should be replaced to `hpx::util::unwrap`.
- The performance counter names referring to properties as exposed by the threading subsystem have changes as those now additionally have to specify the thread-pool. See the corresponding documentation for more details.
- The overloads of `hpx::async` that invoke an action do not perform implicit unwrapping of the returned future anymore in case the invoked function does return a future in the first place. In this case `hpx::async` now returns a `hpx::future<future<T>>` making its behavior conforming to its local counterpart.
- We have replaced the use of `boost::exception_ptr` in our APIs with the equivalent `std::exception_ptr`. Please change your codes accordingly. No compatibility settings are provided.
- We have removed the compatibility settings for `HPX_WITH_COLOCATED_BACKWARDS_COMPATIBILITY` and `HPX_WITH_COMPONENT_GET_GID_COMPATIBILITY` as their life-cycle has reached its end.
- We have removed the experimental thread schedulers `hierarchy_scheduler`, `periodic_priority_scheduler` and `throttling_scheduler` in an effort to clean up and consolidate our thread schedulers.

¹¹⁸⁹ <https://www.cmake.org>

¹¹⁹⁰ <https://www.cmake.org>

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- PR #3250¹¹⁹¹ - Apex refactoring with guids
- PR #3249¹¹⁹² - Updating People.qbk
- PR #3246¹¹⁹³ - Assorted fixes for CUDA
- PR #3245¹¹⁹⁴ - Apex refactoring with guids
- PR #3242¹¹⁹⁵ - Modify task counting in thread_queue.hpp
- PR #3240¹¹⁹⁶ - Fixed typos
- PR #3238¹¹⁹⁷ - Readding accidentally removed std::abort
- PR #3237¹¹⁹⁸ - Adding Pipeline example
- PR #3236¹¹⁹⁹ - Fixing memory_block
- PR #3233¹²⁰⁰ - Make schedule_thread take suspended threads into account
- Issue #3226¹²⁰¹ - memory_block is breaking, signaling SIGSEGV on a thread on creation and freeing
- PR #3225¹²⁰² - Applying quick fix for hwloc-2.0
- Issue #3224¹²⁰³ - HPX counters crashing the application
- PR #3223¹²⁰⁴ - Fix returns when setting config entries
- Issue #3222¹²⁰⁵ - Errors linking libhpx.so
- Issue #3221¹²⁰⁶ - HPX on Mac OS X with HWLoc 2.0.0 fails to run
- PR #3216¹²⁰⁷ - Reorder a variadic array to satisfy VS 2017 15.6
- PR #3214¹²⁰⁸ - Changed prerequisites.qbk to avoid confusion while building boost
- PR #3213¹²⁰⁹ - Relax locks for thread suspension to avoid holding locks when yielding
- PR #3212¹²¹⁰ - Fix check in sequenced_executor test
- PR #3211¹²¹¹ - Use preinit_array to set argc/argv in init_globally example

¹¹⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3250>

¹¹⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3249>

¹¹⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3246>

¹¹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3245>

¹¹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3242>

¹¹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3240>

¹¹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3238>

¹¹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3237>

¹¹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3236>

¹²⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3233>

¹²⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/3226>

¹²⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3225>

¹²⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/3224>

¹²⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3223>

¹²⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3222>

¹²⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3221>

¹²⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3216>

¹²⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3214>

¹²⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3213>

¹²¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3212>

¹²¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3211>

- PR #3210¹²¹² - Adapted parallel::{search | search_n} for Ranges TS (see #1668)
- PR #3209¹²¹³ - Fix locking problems during shutdown
- Issue #3208¹²¹⁴ - init_globally throwing a run-time error
- PR #3206¹²¹⁵ - Addition of new arithmetic performance counter “Count”
- PR #3205¹²¹⁶ - Fixing return type calculation for bulk_then_execute
- PR #3204¹²¹⁷ - Changing std::rand() to a better inbuilt PRNG generator
- PR #3203¹²¹⁸ - Resolving problems during shutdown for VS2015
- PR #3202¹²¹⁹ - Making sure resource partitioner is not accessed if its not valid
- PR #3201¹²²⁰ - Fixing optional::swap
- Issue #3200¹²²¹ - hpx::util::optional fails
- PR #3199¹²²² - Fix sliding_semaphore test
- PR #3198¹²²³ - Set pre_main status before launching run_helper
- PR #3197¹²²⁴ - Update README.rst
- PR #3194¹²²⁵ - parallel::{fillfill_n} updated for Ranges TS
- PR #3193¹²²⁶ - Updating Runtime.cpp by adding correct description of Performance counters during register
- PR #3191¹²²⁷ - Fix sliding_semaphore_2338 test
- PR #3190¹²²⁸ - Topology improvements
- PR #3189¹²²⁹ - Deleting one include of median from BOOST library to arithmetics_counter file
- PR #3188¹²³⁰ - Optionally disable printing of diagnostics during terminate
- PR #3187¹²³¹ - Suppressing cmake warning issued by cmake > V3.11
- PR #3185¹²³² - Remove unused scoped_unlock, unlock_guard_try
- PR #3184¹²³³ - Fix nqueen example
- PR #3183¹²³⁴ - Add runtime start/stop, resume/suspend and OpenMP benchmarks

¹²¹² <https://github.com/STELLAR-GROUP/hpx/pull/3210>

¹²¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3209>

¹²¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3208>

¹²¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3206>

¹²¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3205>

¹²¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3204>

¹²¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3203>

¹²¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3202>

¹²²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3201>

¹²²¹ <https://github.com/STELLAR-GROUP/hpx/issues/3200>

¹²²² <https://github.com/STELLAR-GROUP/hpx/pull/3199>

¹²²³ <https://github.com/STELLAR-GROUP/hpx/pull/3198>

¹²²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3197>

¹²²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3194>

¹²²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3193>

¹²²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3191>

¹²²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3190>

¹²²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3189>

¹²³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3188>

¹²³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3187>

¹²³² <https://github.com/STELLAR-GROUP/hpx/pull/3185>

¹²³³ <https://github.com/STELLAR-GROUP/hpx/pull/3184>

¹²³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3183>

- Issue #3182¹²³⁵ - bulk_then_execute has unexpected return type/does not compile
- Issue #3181¹²³⁶ - hwloc 2.0 breaks topo class and cannot be used
- Issue #3180¹²³⁷ - Schedulers that don't support suspend/resume are unusable
- PR #3179¹²³⁸ - Various minor changes to support FLeCSI
- PR #3178¹²³⁹ - Fix #3124
- PR #3177¹²⁴⁰ - Removed allgather
- PR #3176¹²⁴¹ - Fixed Documentation for “using_hpx_pkgconfig”
- PR #3174¹²⁴² - Add hpx::iostreams::ostream overload to format_to
- PR #3172¹²⁴³ - Fix lifo queue backend
- PR #3171¹²⁴⁴ - adding the missing unset() function to cpu_mask() for case of more than 64 threads
- PR #3170¹²⁴⁵ - Add cmake flag -DHPX_WITHFAULT_TOLERANCE=ON (OFF by default)
- PR #3169¹²⁴⁶ - Adapted parallel::{count|count_if} for Ranges TS (see #1668)
- PR #3168¹²⁴⁷ - Changing used namespace for seq execution policy
- Issue #3167¹²⁴⁸ - Update GSoC projects
- Issue #3166¹²⁴⁹ - Application (Octotiger) gets stuck on hpx::finalize when only using one thread
- Issue #3165¹²⁵⁰ - Compilation of parallel algorithms with HPX_WITH_DATAPAR is broken
- PR #3164¹²⁵¹ - Fixing component migration
- PR #3162¹²⁵² - regex_from_pattern: escape regex special characters to avoid misinterpretation
- Issue #3161¹²⁵³ - Building HPX with hwloc 2.0.0 fails
- PR #3160¹²⁵⁴ - Fixing the handling of quoted command line arguments.
- PR #3158¹²⁵⁵ - Fixing a race with timed suspension (second attempt)
- PR #3157¹²⁵⁶ - Revert “Fixing a race with timed suspension”
- PR #3156¹²⁵⁷ - Fixing serialization of classes with incompatible serialize signature

¹²³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3182>

¹²³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3181>

¹²³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/3180>

¹²³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3179>

¹²³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3178>

¹²⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3177>

¹²⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3176>

¹²⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3174>

¹²⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3172>

¹²⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3171>

¹²⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3170>

¹²⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3169>

¹²⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3168>

¹²⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3167>

¹²⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3166>

¹²⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3165>

¹²⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3164>

¹²⁵² <https://github.com/STELLAR-GROUP/hpx/pull/3162>

¹²⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/3161>

¹²⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3160>

¹²⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3158>

¹²⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3157>

¹²⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3156>

- PR #3154¹²⁵⁸ - More refactorings based on clang-tidy reports
- PR #3153¹²⁵⁹ - Fixing a race with timed suspension
- PR #3152¹²⁶⁰ - Documentation for runtime suspension
- PR #3151¹²⁶¹ - Use small_vector only from boost version 1.59 onwards
- PR #3150¹²⁶² - Avoiding more stack overflows
- PR #3148¹²⁶³ - Refactoring component_base and base_action/transfer_base_action
- PR #3147¹²⁶⁴ - Move yield_while out of detail namespace and into own file
- PR #3145¹²⁶⁵ - Remove a leftover of the cxx11 std array cleanup
- PR #3144¹²⁶⁶ - Minor changes to how actions are executed
- PR #3143¹²⁶⁷ - Fix stack overhead
- PR #3142¹²⁶⁸ - Fix typo in config.hpp
- PR #3141¹²⁶⁹ - Fixing small_vector compatibility with older boost version
- PR #3140¹²⁷⁰ - is_heap_text fix
- Issue #3139¹²⁷¹ - Error in is_heap_tests.hpp
- PR #3138¹²⁷² - Partially reverting #3126
- PR #3137¹²⁷³ - Suspend speedup
- PR #3136¹²⁷⁴ - Revert “Fixing #2325”
- PR #3135¹²⁷⁵ - Improving destruction of threads
- Issue #3134¹²⁷⁶ - HPX_SERIALIZATION_SPLIT_FREE does not stop compiler from looking for serialize() method
- PR #3133¹²⁷⁷ - Make hwloc compulsory
- PR #3132¹²⁷⁸ - Update CXX14 constexpr feature test
- PR #3131¹²⁷⁹ - Fixing #2325
- PR #3130¹²⁸⁰ - Avoid completion handler allocation

¹²⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3154>

¹²⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3153>

¹²⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3152>

¹²⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3151>

¹²⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3150>

¹²⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3148>

¹²⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3147>

¹²⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3145>

¹²⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3144>

¹²⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3143>

¹²⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3142>

¹²⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3141>

¹²⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3140>

¹²⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/3139>

¹²⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3138>

¹²⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3137>

¹²⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3136>

¹²⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3135>

¹²⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3134>

¹²⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3133>

¹²⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3132>

¹²⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3131>

¹²⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3130>

- PR #3129¹²⁸¹ - Suspend runtime
- PR #3128¹²⁸² - Make docbook dtd and xsl path names consistent
- PR #3127¹²⁸³ - Add hpx::start nullptr overloads
- PR #3126¹²⁸⁴ - Cleaning up coroutine implementation
- PR #3125¹²⁸⁵ - Replacing nullptr with hpx::threads::invalid_thread_id
- Issue #3124¹²⁸⁶ - Add hello_world_component to CI builds
- PR #3123¹²⁸⁷ - Add new constructor.
- PR #3122¹²⁸⁸ - Fixing #3121
- Issue #3121¹²⁸⁹ - HPX_SMT_PAUSE is broken on non-x86 platforms when __GNUC__ is defined
- PR #3120¹²⁹⁰ - Don't use boost::intrusive_ptr for thread_id_type
- PR #3119¹²⁹¹ - Disable default executor compatibility with V1 executors
- PR #3118¹²⁹² - Adding performance_counter::reinit to allow for dynamically changing counter sets
- PR #3117¹²⁹³ - Replace uses of boost/experimental::optional with util::optional
- PR #3116¹²⁹⁴ - Moving background thread APEX timer #2980
- PR #3115¹²⁹⁵ - Fixing race condition in channel test
- PR #3114¹²⁹⁶ - Avoid using util::function for thread function wrappers
- PR #3113¹²⁹⁷ - cmake V3.10.2 has changed the variable names used for MPI
- PR #3112¹²⁹⁸ - Minor fixes to exclusive_scan algorithm
- PR #3111¹²⁹⁹ - Revert "fix detection of cxx11_std_atomic"
- PR #3110¹³⁰⁰ - Suspend thread pool
- PR #3109¹³⁰¹ - Fixing thread scheduling when yielding a thread id
- PR #3108¹³⁰² - Revert "Suspend thread pool"
- PR #3107¹³⁰³ - Remove UB from thread::id relational operators

¹²⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3129>

¹²⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3128>

¹²⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/3127>

¹²⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3126>

¹²⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3125>

¹²⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3124>

¹²⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3123>

¹²⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3122>

¹²⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3121>

¹²⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3120>

¹²⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3119>

¹²⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3118>

¹²⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/3117>

¹²⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3116>

¹²⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3115>

¹²⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3114>

¹²⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3113>

¹²⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3112>

¹²⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3111>

¹³⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3110>

¹³⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/3109>

¹³⁰² <https://github.com/STELLAR-GROUP/hpx/pull/3108>

¹³⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/3107>

- PR #3106¹³⁰⁴ - Add cmake test for std::decay_t to fix cuda build
- PR #3105¹³⁰⁵ - Fixing refcount for async traversal frame
- PR #3104¹³⁰⁶ - Local execution of direct actions is now actually performed directly
- PR #3103¹³⁰⁷ - Adding support for generic counter_raw_values performance counter type
- Issue #3102¹³⁰⁸ - Introduce generic performance counter type returning an array of values
- PR #3101¹³⁰⁹ - Revert “Adapting stack overhead limit for gcc 4.9”
- PR #3100¹³¹⁰ - Fix #3068 (condition_variable deadlock)
- PR #3099¹³¹¹ - Fixing lock held during suspension in papi counter component
- PR #3098¹³¹² - Unbreak broadcast_wait_for_2822 test
- PR #3097¹³¹³ - Adapting stack overhead limit for gcc 4.9
- PR #3096¹³¹⁴ - fix detection of cxx11_std_atomic
- PR #3095¹³¹⁵ - Add ciso646 header to get _LIBCPP_VERSION for testing inplace merge
- PR #3094¹³¹⁶ - Relax atomic operations on performance counter values
- PR #3093¹³¹⁷ - Short-circuit all_of/any_of/none_of instantiations
- PR #3092¹³¹⁸ - Take advantage of C++14 lambda capture initialization syntax, where possible
- PR #3091¹³¹⁹ - Remove more references to Boost from logging code
- PR #3090¹³²⁰ - Unify use of yield/yield_k
- PR #3089¹³²¹ - Fix a strange thing in parallel::detail::handle_exception. (Fix #2834.)
- Issue #3088¹³²² - A strange thing in parallel::sort.
- PR #3087¹³²³ - Fixing assertion in default_distribution_policy
- PR #3086¹³²⁴ - Implement parallel::remove and parallel::remove_if
- PR #3085¹³²⁵ - Addressing breaking changes in Boost V1.66
- PR #3084¹³²⁶ - Ignore build warnings round 2

¹³⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3106>

¹³⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3105>

¹³⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3104>

¹³⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3103>

¹³⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3102>

¹³⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3101>

¹³¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3100>

¹³¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/3099>

¹³¹² <https://github.com/STELLAR-GROUP/hpx/pull/3098>

¹³¹³ <https://github.com/STELLAR-GROUP/hpx/pull/3097>

¹³¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3096>

¹³¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3095>

¹³¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3094>

¹³¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3093>

¹³¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3092>

¹³¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3091>

¹³²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3090>

¹³²¹ <https://github.com/STELLAR-GROUP/hpx/pull/3089>

¹³²² <https://github.com/STELLAR-GROUP/hpx/issues/3088>

¹³²³ <https://github.com/STELLAR-GROUP/hpx/pull/3087>

¹³²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3086>

¹³²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3085>

¹³²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3084>

- PR #3083¹³²⁷ - Fix typo HPX_WITH_MM_PREFECTH
- PR #3081¹³²⁸ - Pre-decay template arguments early
- PR #3080¹³²⁹ - Suspend thread pool
- PR #3079¹³³⁰ - Ignore build warnings
- PR #3078¹³³¹ - Don't test inplace_merge with libc++
- PR #3076¹³³² - Fixing 3075: Part 1
- PR #3074¹³³³ - Fix more build warnings
- PR #3073¹³³⁴ - Suspend thread cleanup
- PR #3072¹³³⁵ - Change existing symbol_namespace::iterate to return all data instead of invoking a callback
- PR #3071¹³³⁶ - Fixing pack_traversal_async test
- PR #3070¹³³⁷ - Fix dynamic_counters_loaded_1508 test by adding dependency to memory_component
- PR #3069¹³³⁸ - Fix scheduling loop exit
- Issue #3068¹³³⁹ - hpx::lcos::condition_variable could be suspect to deadlocks
- PR #3067¹³⁴⁰ - #ifdef out random_shuffle deprecated in later c++
- PR #3066¹³⁴¹ - Make coalescing test depend on coalescing library to ensure it gets built
- PR #3065¹³⁴² - Workaround for minimal_timed_async_executor_test compilation failures, attempts to copy a deferred call (in unevaluated context)
- PR #3064¹³⁴³ - Fixing wrong condition in wrapper_heap
- PR #3062¹³⁴⁴ - Fix exception handling for execution::seq
- PR #3061¹³⁴⁵ - Adapt MSVC C++ mode handling to VS15.5
- PR #3060¹³⁴⁶ - Fix compiler problem in MSVC release mode
- PR #3059¹³⁴⁷ - Fixing #2931
- Issue #3058¹³⁴⁸ - minimal_timed_async_executor_test_exe fails to compile on master (d6f505c)
- PR #3057¹³⁴⁹ - Fix stable_merge_2964 compilation problems

¹³²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3083>

¹³²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3081>

¹³²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3080>

¹³³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3079>

¹³³¹ <https://github.com/STELLAR-GROUP/hpx/pull/3078>

¹³³² <https://github.com/STELLAR-GROUP/hpx/pull/3076>

¹³³³ <https://github.com/STELLAR-GROUP/hpx/pull/3074>

¹³³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3073>

¹³³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3072>

¹³³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3071>

¹³³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3070>

¹³³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3069>

¹³³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/3068>

¹³⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3067>

¹³⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/3066>

¹³⁴² <https://github.com/STELLAR-GROUP/hpx/pull/3065>

¹³⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/3064>

¹³⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3062>

¹³⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3061>

¹³⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3060>

¹³⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3059>

¹³⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/3058>

¹³⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3057>

- PR #3056¹³⁵⁰ - Fix some build warnings caused by unused variables/unnecessary tests
- PR #3055¹³⁵¹ - Update documentation for running tests
- Issue #3054¹³⁵² - Assertion failure when using bulk hpx::new_ in asynchronous mode
- PR #3052¹³⁵³ - Do not bind test running to cmake test build rule
- PR #3051¹³⁵⁴ - Fix HPX-Qt interaction in Qt example.
- Issue #3048¹³⁵⁵ - nqueen example fails occasionally
- PR #3047¹³⁵⁶ - Fixing #3044
- PR #3046¹³⁵⁷ - Add OS thread suspension
- PR #3042¹³⁵⁸ - PyCicle - first attempt at a build tool for checking PR's
- PR #3041¹³⁵⁹ - Fix a problem about asynchronous execution of parallel::merge and parallel::partition.
- PR #3040¹³⁶⁰ - Fix a mistake about exception handling in asynchronous execution of scan_partitioner.
- PR #3039¹³⁶¹ - Consistently use executors to schedule work
- PR #3038¹³⁶² - Fixing local direct function execution and lambda actions perfect forwarding
- PR #3035¹³⁶³ - Make parallel unit test names match build target/folder names
- PR #3033¹³⁶⁴ - Fix setting of default build type
- Issue #3032¹³⁶⁵ - Fix partitioner arg copy found in #2982
- Issue #3031¹³⁶⁶ - Errors linking libhpx.so due to missing references (master branch, commit 6679a8882)
- PR #3030¹³⁶⁷ - Revert "implement executor then interface with && forwarding reference"
- PR #3029¹³⁶⁸ - Run CI inspect checks before building
- PR #3028¹³⁶⁹ - Added range version of parallel::move
- Issue #3027¹³⁷⁰ - Implement all scheduling APIs in terms of executors
- PR #3026¹³⁷¹ - implement executor then interface with && forwarding reference
- PR #3025¹³⁷² - Fix typo uninitialized to uninitialized

¹³⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3056>

¹³⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/3055>

¹³⁵² <https://github.com/STELLAR-GROUP/hpx/issues/3054>

¹³⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/3052>

¹³⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3051>

¹³⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3048>

¹³⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3047>

¹³⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3046>

¹³⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3042>

¹³⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3041>

¹³⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3040>

¹³⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/3039>

¹³⁶² <https://github.com/STELLAR-GROUP/hpx/pull/3038>

¹³⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/3035>

¹³⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3033>

¹³⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/3032>

¹³⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/3031>

¹³⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3030>

¹³⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3029>

¹³⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3028>

¹³⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/3027>

¹³⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/3026>

¹³⁷² <https://github.com/STELLAR-GROUP/hpx/pull/3025>

- PR #3024¹³⁷³ - Inspect fixes
- PR #3023¹³⁷⁴ - P0356 Simplified partial function application
- PR #3022¹³⁷⁵ - Master fixes
- PR #3021¹³⁷⁶ - Segfault fix
- PR #3020¹³⁷⁷ - Disable command-line aliasing for applications that use user_main
- PR #3019¹³⁷⁸ - Adding enable_elasticity option to pool configuration
- PR #3018¹³⁷⁹ - Fix stack overflow detection configuration in header files
- PR #3017¹³⁸⁰ - Speed up local action execution
- PR #3016¹³⁸¹ - Unify stack-overflow detection options, remove reference to libsigsegv
- PR #3015¹³⁸² - Speeding up accessing the resource partitioner and the topology info
- Issue #3014¹³⁸³ - HPX does not compile on POWER8 with gcc 5.4
- Issue #3013¹³⁸⁴ - hello_world occasionally prints multiple lines from a single OS-thread
- PR #3012¹³⁸⁵ - Silence warning about casting away qualifiers in itt_notify.hpp
- PR #3011¹³⁸⁶ - Fix cpuset leak in hwloc_topology_info.cpp
- PR #3010¹³⁸⁷ - Remove useless decay_copy
- PR #3009¹³⁸⁸ - Fixing 2996
- PR #3008¹³⁸⁹ - Remove unused internal function
- PR #3007¹³⁹⁰ - Fixing wrapper_heap alignment problems
- Issue #3006¹³⁹¹ - hwloc memory leak
- PR #3004¹³⁹² - Silence C4251 (needs to have dll-interface) for future_data_void
- Issue #3003¹³⁹³ - Suspension of runtime
- PR #3001¹³⁹⁴ - Attempting to avoid data races in async_traversal while evaluating dataflow()
- PR #3000¹³⁹⁵ - Adding hpx::util::optional as a first step to replace experimental::optional

¹³⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/3024>

¹³⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3023>

¹³⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3022>

¹³⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3021>

¹³⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3020>

¹³⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3019>

¹³⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3018>

¹³⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3017>

¹³⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/3016>

¹³⁸² <https://github.com/STELLAR-GROUP/hpx/pull/3015>

¹³⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/3014>

¹³⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/3013>

¹³⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3012>

¹³⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/3011>

¹³⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/3010>

¹³⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/3009>

¹³⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/3008>

¹³⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/3007>

¹³⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/3006>

¹³⁹² <https://github.com/STELLAR-GROUP/hpx/pull/3004>

¹³⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/3003>

¹³⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/3001>

¹³⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/3000>

- PR #2998¹³⁹⁶ - Cleanup up and Fixing component creation and deletion
- Issue #2996¹³⁹⁷ - Build fails with HPX_WITH_HWLOC=OFF
- PR #2995¹³⁹⁸ - Push more future_data functionality to source file
- PR #2994¹³⁹⁹ - WIP: Fix throttle test
- PR #2993¹⁴⁰⁰ - Making sure -hpx:help does not throw for required (but missing) arguments
- PR #2992¹⁴⁰¹ - Adding non-blocking (on destruction) service executors
- Issue #2991¹⁴⁰² - run_as_os_thread locks up
- Issue #2990¹⁴⁰³ - --help will not work until all required options are provided
- PR #2989¹⁴⁰⁴ - Improve error messages caused by misuse of dataflow
- PR #2988¹⁴⁰⁵ - Improve error messages caused by misuse of .then
- Issue #2987¹⁴⁰⁶ - stack overflow detection producing false positives
- PR #2986¹⁴⁰⁷ - Deduplicate non-dependent thread_info logging types
- PR #2985¹⁴⁰⁸ - Adapted parallel::{all_oflany_oflnone_of} for Ranges TS (see #1668)
- PR #2984¹⁴⁰⁹ - Refactor one_size_heap code to simplify code
- PR #2983¹⁴¹⁰ - Fixing local_new_component
- PR #2982¹⁴¹¹ - Clang tidy
- PR #2981¹⁴¹² - Simplify allocator rebinding in pack traversal
- PR #2979¹⁴¹³ - Fixing integer overflows
- PR #2978¹⁴¹⁴ - Implement parallel:::inplace_merge
- Issue #2977¹⁴¹⁵ - Make hwloc compulsory instead of optional
- PR #2976¹⁴¹⁶ - Making sure client_base instance that registered the component does not unregister it when being destructed
- PR #2975¹⁴¹⁷ - Change version of pulled APEX to master
- PR #2974¹⁴¹⁸ - Fix domain not being freed at the end of scheduling loop

1396 <https://github.com/STELLAR-GROUP/hpx/pull/2998>

1397 <https://github.com/STELLAR-GROUP/hpx/issues/2996>

1398 <https://github.com/STELLAR-GROUP/hpx/pull/2995>

1399 <https://github.com/STELLAR-GROUP/hpx/pull/2994>

1400 <https://github.com/STELLAR-GROUP/hpx/pull/2993>

1401 <https://github.com/STELLAR-GROUP/hpx/pull/2992>

1402 <https://github.com/STELLAR-GROUP/hpx/issues/2991>

1403 <https://github.com/STELLAR-GROUP/hpx/issues/2990>

1404 <https://github.com/STELLAR-GROUP/hpx/pull/2989>

1405 <https://github.com/STELLAR-GROUP/hpx/pull/2988>

1406 <https://github.com/STELLAR-GROUP/hpx/issues/2987>

1407 <https://github.com/STELLAR-GROUP/hpx/pull/2986>

1408 <https://github.com/STELLAR-GROUP/hpx/pull/2985>

1409 <https://github.com/STELLAR-GROUP/hpx/pull/2984>

1410 <https://github.com/STELLAR-GROUP/hpx/pull/2983>

1411 <https://github.com/STELLAR-GROUP/hpx/pull/2982>

1412 <https://github.com/STELLAR-GROUP/hpx/pull/2981>

1413 <https://github.com/STELLAR-GROUP/hpx/pull/2979>

1414 <https://github.com/STELLAR-GROUP/hpx/pull/2978>

1415 <https://github.com/STELLAR-GROUP/hpx/issues/2977>

1416 <https://github.com/STELLAR-GROUP/hpx/pull/2976>

1417 <https://github.com/STELLAR-GROUP/hpx/pull/2975>

1418 <https://github.com/STELLAR-GROUP/hpx/pull/2974>

- PR #2973¹⁴¹⁹ - Fix small typos
- PR #2972¹⁴²⁰ - Adding uintstd.h header
- PR #2971¹⁴²¹ - Fall back to creating local components using local_new
- PR #2970¹⁴²² - Improve is_tuple_like trait
- PR #2969¹⁴²³ - Fix HPX_WITH_MORE_THAN_64_THREADS default value
- PR #2968¹⁴²⁴ - Cleaning up dataflow overload set
- PR #2967¹⁴²⁵ - Make parallel::merge is stable. (Fix #2964.)
- PR #2966¹⁴²⁶ - Fixing a couple of held locks during exception handling
- PR #2965¹⁴²⁷ - Adding missing #include
- Issue #2964¹⁴²⁸ - parallel merge is not stable
- PR #2963¹⁴²⁹ - Making sure any function object passed to dataflow is released after being invoked
- PR #2962¹⁴³⁰ - Partially reverting #2891
- PR #2961¹⁴³¹ - Attempt to fix the gcc 4.9 problem with the async pack traversal
- Issue #2959¹⁴³² - Program terminates during error handling
- Issue #2958¹⁴³³ - HPX_PLAIN_ACTION breaks due to missing include
- PR #2957¹⁴³⁴ - Fixing errors generated by mixing different attribute syntaxes
- Issue #2956¹⁴³⁵ - Mixing attribute syntaxes leads to compiler errors
- Issue #2955¹⁴³⁶ - Fix OS-Thread throttling
- PR #2953¹⁴³⁷ - Making sure any hpx.os_threads=N supplied through a -hpx::config file is taken into account
- PR #2952¹⁴³⁸ - Removing wrong call to cleanup_terminated_locked
- PR #2951¹⁴³⁹ - Revert “Make sure the function vtables are initialized before use”
- PR #2950¹⁴⁴⁰ - Fix a namespace compilation error when some schedulers are disabled
- Issue #2949¹⁴⁴¹ - master branch giving lockups on shutdown

¹⁴¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2973>

¹⁴²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2972>

¹⁴²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2971>

¹⁴²² <https://github.com/STELLAR-GROUP/hpx/pull/2970>

¹⁴²³ <https://github.com/STELLAR-GROUP/hpx/pull/2969>

¹⁴²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2968>

¹⁴²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2967>

¹⁴²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2966>

¹⁴²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2965>

¹⁴²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2964>

¹⁴²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2963>

¹⁴³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2962>

¹⁴³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2961>

¹⁴³² <https://github.com/STELLAR-GROUP/hpx/issues/2959>

¹⁴³³ <https://github.com/STELLAR-GROUP/hpx/issues/2958>

¹⁴³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2957>

¹⁴³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2956>

¹⁴³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2955>

¹⁴³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2953>

¹⁴³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2952>

¹⁴³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2951>

¹⁴⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2950>

¹⁴⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/2949>

- Issue #2947¹⁴⁴² - hpx.ini is not used correctly at initialization
- PR #2946¹⁴⁴³ - Adding explicit feature test for thread_local
- PR #2945¹⁴⁴⁴ - Make sure the function vtables are initialized before use
- PR #2944¹⁴⁴⁵ - Attempting to solve affinity problems on CircleCI
- PR #2943¹⁴⁴⁶ - Changing channel actions to be direct
- PR #2942¹⁴⁴⁷ - Adding split_future for std::vector
- PR #2941¹⁴⁴⁸ - Add a feature test to test for CXX11 override
- Issue #2940¹⁴⁴⁹ - Add split_future for future<vector<T>>
- PR #2939¹⁴⁵⁰ - Making error reporting during problems with setting affinity masks more verbose
- PR #2938¹⁴⁵¹ - Fix this various executors
- PR #2937¹⁴⁵² - Fix some typos in documentation
- PR #2934¹⁴⁵³ - Remove the need for “complete” SFINAE checks
- PR #2933¹⁴⁵⁴ - Making sure parallel::for_loop is executed in parallel if requested
- PR #2932¹⁴⁵⁵ - Classify chunk_size_iterator to input iterator tag. (Fix #2866)
- Issue #2931¹⁴⁵⁶ - –hpx:help triggers unusual error with clang build
- PR #2930¹⁴⁵⁷ - Add #include files needed to set _POSIX_VERSION for debug check
- PR #2929¹⁴⁵⁸ - Fix a couple of deprecated c++ features
- PR #2928¹⁴⁵⁹ - Fixing execution parameters
- Issue #2927¹⁴⁶⁰ - CMake warning: ... cycle in constraint graph
- PR #2926¹⁴⁶¹ - Default pool rename
- Issue #2925¹⁴⁶² - Default pool cannot be renamed
- Issue #2924¹⁴⁶³ - hpx:attach-debugger=startup does not work any more
- PR #2923¹⁴⁶⁴ - Alloc membind

1442 <https://github.com/STELLAR-GROUP/hpx/issues/2947>

1443 <https://github.com/STELLAR-GROUP/hpx/pull/2946>

1444 <https://github.com/STELLAR-GROUP/hpx/pull/2945>

1445 <https://github.com/STELLAR-GROUP/hpx/pull/2944>

1446 <https://github.com/STELLAR-GROUP/hpx/pull/2943>

1447 <https://github.com/STELLAR-GROUP/hpx/pull/2942>

1448 <https://github.com/STELLAR-GROUP/hpx/pull/2941>

1449 <https://github.com/STELLAR-GROUP/hpx/issues/2940>

1450 <https://github.com/STELLAR-GROUP/hpx/pull/2939>

1451 <https://github.com/STELLAR-GROUP/hpx/pull/2938>

1452 <https://github.com/STELLAR-GROUP/hpx/pull/2937>

1453 <https://github.com/STELLAR-GROUP/hpx/pull/2934>

1454 <https://github.com/STELLAR-GROUP/hpx/pull/2933>

1455 <https://github.com/STELLAR-GROUP/hpx/pull/2932>

1456 <https://github.com/STELLAR-GROUP/hpx/issues/2931>

1457 <https://github.com/STELLAR-GROUP/hpx/pull/2930>

1458 <https://github.com/STELLAR-GROUP/hpx/pull/2929>

1459 <https://github.com/STELLAR-GROUP/hpx/pull/2928>

1460 <https://github.com/STELLAR-GROUP/hpx/issues/2927>

1461 <https://github.com/STELLAR-GROUP/hpx/pull/2926>

1462 <https://github.com/STELLAR-GROUP/hpx/issues/2925>

1463 <https://github.com/STELLAR-GROUP/hpx/issues/2924>

1464 <https://github.com/STELLAR-GROUP/hpx/pull/2923>

- PR #2922¹⁴⁶⁵ - This fixes CircleCI errors when running with `-hpx:bind=none`
- PR #2921¹⁴⁶⁶ - Custom pool executor was missing priority and stacksize options
- PR #2920¹⁴⁶⁷ - Adding test to trigger problem reported in #2916
- PR #2919¹⁴⁶⁸ - Make sure the resource_partitioner is properly destructed on `hpx::finalize`
- Issue #2918¹⁴⁶⁹ - `hpx::init` calls wrong (first) callback when called multiple times
- PR #2917¹⁴⁷⁰ - Adding util::checkpoint
- Issue #2916¹⁴⁷¹ - Weird runtime failures when using a channel and chained continuations
- PR #2915¹⁴⁷² - Introduce executor parameters customization points
- Issue #2914¹⁴⁷³ - Task assignment to current Pool has unintended consequences
- PR #2913¹⁴⁷⁴ - Fix rp hang
- PR #2912¹⁴⁷⁵ - Update contributors
- PR #2911¹⁴⁷⁶ - Fixing CUDA problems
- PR #2910¹⁴⁷⁷ - Improve error reporting for process component on POSIX systems
- PR #2909¹⁴⁷⁸ - Fix typo in include path
- PR #2908¹⁴⁷⁹ - Use proper container according to iterator tag in benchmarks of parallel algorithms
- PR #2907¹⁴⁸⁰ - Optionally force-delete remaining channel items on close
- PR #2906¹⁴⁸¹ - Making sure generated performance counter names are correct
- Issue #2905¹⁴⁸² - collecting idle-rate performance counters on multiple localities produces an error
- Issue #2904¹⁴⁸³ - build broken for Intel 17 compilers
- PR #2903¹⁴⁸⁴ - Documentation Updates—Adding New People
- PR #2902¹⁴⁸⁵ - Fixing service_executor
- PR #2901¹⁴⁸⁶ - Fixing partitioned_vector creation
- PR #2900¹⁴⁸⁷ - Add numa-balanced mode to `hpx::bind`, spread cores over numa domains

¹⁴⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2922>

¹⁴⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2921>

¹⁴⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2920>

¹⁴⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2919>

¹⁴⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2918>

¹⁴⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2917>

¹⁴⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/2916>

¹⁴⁷² <https://github.com/STELLAR-GROUP/hpx/pull/2915>

¹⁴⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2914>

¹⁴⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2913>

¹⁴⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2912>

¹⁴⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2911>

¹⁴⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2910>

¹⁴⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2909>

¹⁴⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2908>

¹⁴⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2907>

¹⁴⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2906>

¹⁴⁸² <https://github.com/STELLAR-GROUP/hpx/issues/2905>

¹⁴⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/2904>

¹⁴⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2903>

¹⁴⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2902>

¹⁴⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2901>

¹⁴⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2900>

- Issue #2899¹⁴⁸⁸ - hpx::bind does not have a mode that balances cores over numa domains
- PR #2898¹⁴⁸⁹ - Adding missing #include and missing guard for optional code section
- PR #2897¹⁴⁹⁰ - Removing dependency on Boost.ICL
- Issue #2896¹⁴⁹¹ - Debug build fails without -fpermissive with GCC 7.1 and Boost 1.65
- PR #2895¹⁴⁹² - Fixing SLURM environment parsing
- PR #2894¹⁴⁹³ - Fix incorrect handling of compile definition with value 0
- Issue #2893¹⁴⁹⁴ - Disabling schedulers causes build errors
- PR #2892¹⁴⁹⁵ - added list serializer
- PR #2891¹⁴⁹⁶ - Resource Partitioner Fixes
- Issue #2890¹⁴⁹⁷ - Destroying a non-empty channel causes an assertion failure
- PR #2889¹⁴⁹⁸ - Add check for libatomic
- PR #2888¹⁴⁹⁹ - Fix compilation problems if HPX_WITH_ITT_NOTIFY=ON
- PR #2887¹⁵⁰⁰ - Adapt broadcast() to non-unwrapping async<Action>
- PR #2886¹⁵⁰¹ - Replace Boost.Random with C++11 <random>
- Issue #2885¹⁵⁰² - regression in broadcast?
- Issue #2884¹⁵⁰³ - linking -latomic is not portable
- PR #2883¹⁵⁰⁴ - Explicitly set -pthread flag if available
- PR #2882¹⁵⁰⁵ - Wrap boost::format uses
- Issue #2881¹⁵⁰⁶ - hpx not compiling with HPX_WITH_ITTNOTIFY=On
- Issue #2880¹⁵⁰⁷ - hpx::bind scatter/balanced give wrong pu masks
- PR #2878¹⁵⁰⁸ - Fix incorrect pool usage masks setup in RP/thread manager
- PR #2877¹⁵⁰⁹ - Require std::array by default
- PR #2875¹⁵¹⁰ - Deprecate use of BOOST_ASSERT

¹⁴⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2899>

¹⁴⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2898>

¹⁴⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2897>

¹⁴⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/2896>

¹⁴⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2895>

¹⁴⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2894>

¹⁴⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2893>

¹⁴⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2892>

¹⁴⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2891>

¹⁴⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2890>

¹⁴⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2889>

¹⁴⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2888>

¹⁵⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2887>

¹⁵⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2886>

¹⁵⁰² <https://github.com/STELLAR-GROUP/hpx/issues/2885>

¹⁵⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/2884>

¹⁵⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2883>

¹⁵⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2882>

¹⁵⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2881>

¹⁵⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2880>

¹⁵⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2878>

¹⁵⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2877>

¹⁵¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2875>

- PR #2874¹⁵¹¹ - Changed serialization of boost.variant to use variadic templates
- Issue #2873¹⁵¹² - building with parcelport_mpi fails on cori
- PR #2871¹⁵¹³ - Adding missing support for throttling scheduler
- PR #2870¹⁵¹⁴ - Disambiguate use of base_lco_with_value macros with channel
- Issue #2869¹⁵¹⁵ - Difficulty compiling HPX_REGISTER_CHANNEL_DECLARATION(double)
- PR #2868¹⁵¹⁶ - Removing unneeded assert
- PR #2867¹⁵¹⁷ - Implement parallel::unique
- Issue #2866¹⁵¹⁸ - The chunk_size_iterator violates multipass guarantee
- PR #2865¹⁵¹⁹ - Only use sched_getcpu on linux machines
- PR #2864¹⁵²⁰ - Create redistribution archive for successful builds
- PR #2863¹⁵²¹ - Replace casts/assignments with hard-coded memcpy operations
- Issue #2862¹⁵²² - sched_getcpu not available on MacOS
- PR #2861¹⁵²³ - Fixing unmatched header defines and recursive inclusion of threadmanager
- Issue #2860¹⁵²⁴ - Master program fails with assertion ‘type == data_type_address’ failed: HPX(assertion_failure)
- Issue #2852¹⁵²⁵ - Support for ARM64
- PR #2858¹⁵²⁶ - Fix misplaced #if #endif's that cause build failure without THREAD_CUMULATIVE_COUNTS
- PR #2857¹⁵²⁷ - Fix some listing in documentation
- PR #2856¹⁵²⁸ - Fixing component handling for lcos
- PR #2855¹⁵²⁹ - Add documentation for coarrays
- PR #2854¹⁵³⁰ - Support ARM64 in timestamps
- PR #2853¹⁵³¹ - Update Table 17. Non-modifying Parallel Algorithms in Documentation
- PR #2851¹⁵³² - Allowing for non-default-constructible component types

¹⁵¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2874>

¹⁵¹² <https://github.com/STELLAR-GROUP/hpx/issues/2873>

¹⁵¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2871>

¹⁵¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2870>

¹⁵¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2869>

¹⁵¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2868>

¹⁵¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2867>

¹⁵¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2866>

¹⁵¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2865>

¹⁵²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2864>

¹⁵²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2863>

¹⁵²² <https://github.com/STELLAR-GROUP/hpx/issues/2862>

¹⁵²³ <https://github.com/STELLAR-GROUP/hpx/pull/2861>

¹⁵²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2860>

¹⁵²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2852>

¹⁵²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2858>

¹⁵²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2857>

¹⁵²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2856>

¹⁵²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2855>

¹⁵³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2854>

¹⁵³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2853>

¹⁵³² <https://github.com/STELLAR-GROUP/hpx/pull/2851>

- PR #2850¹⁵³³ - Enable returning future<R> from actions where R is not default-constructible
- PR #2849¹⁵³⁴ - Unify serialization of non-default-constructable types
- Issue #2848¹⁵³⁵ - Components have to be default constructible
- Issue #2847¹⁵³⁶ - Returning a future<R> where R is not default-constructable broken
- Issue #2846¹⁵³⁷ - Unify serialization of non-default-constructible types
- PR #2845¹⁵³⁸ - Add Visual Studio 2015 to the tested toolchains in Appveyor
- Issue #2844¹⁵³⁹ - Change the appveyor build to use the minimal required MSVC version
- Issue #2843¹⁵⁴⁰ - multi node hello_world hangs
- PR #2842¹⁵⁴¹ - Correcting Spelling mistake in docs
- PR #2841¹⁵⁴² - Fix usage of std::aligned_storage
- PR #2840¹⁵⁴³ - Remove constexpr from a void function
- Issue #2839¹⁵⁴⁴ - memcpy buffer overflow: load_construct_data() and std::complex members
- Issue #2835¹⁵⁴⁵ - constexpr functions with void return type break compilation with CUDA 8.0
- Issue #2834¹⁵⁴⁶ - One suspicion in parallel::detail::handle_exception
- PR #2833¹⁵⁴⁷ - Implement parallel::merge
- PR #2832¹⁵⁴⁸ - Fix a strange thing in parallel::util::detail::handle_local_exceptions. (Fix #2818)
- PR #2830¹⁵⁴⁹ - Break the debugger when a test failed
- Issue #2831¹⁵⁵⁰ - parallel/executors/execution_fwd.hpp causes compilation failure in C++11 mode.
- PR #2829¹⁵⁵¹ - Implement an API for asynchronous pack traversal
- PR #2828¹⁵⁵² - Split unit test builds on CircleCI to avoid timeouts
- Issue #2827¹⁵⁵³ - failure to compile hello_world example with -Werror
- PR #2824¹⁵⁵⁴ - Making sure promises are marked as started when used as continuations
- PR #2823¹⁵⁵⁵ - Add documentation for partitioned_vector_view

¹⁵³³ <https://github.com/STELLAR-GROUP/hpx/pull/2850>

¹⁵³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2849>

¹⁵³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2848>

¹⁵³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2847>

¹⁵³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2846>

¹⁵³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2845>

¹⁵³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2844>

¹⁵⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2843>

¹⁵⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2842>

¹⁵⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2841>

¹⁵⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/2840>

¹⁵⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2839>

¹⁵⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2835>

¹⁵⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2834>

¹⁵⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2833>

¹⁵⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2832>

¹⁵⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2830>

¹⁵⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2831>

¹⁵⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2829>

¹⁵⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2828>

¹⁵⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/2827>

¹⁵⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2824>

¹⁵⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2823>

- Issue #2822¹⁵⁵⁶ - Yet another issue with wait_for similar to #2796
- PR #2821¹⁵⁵⁷ - Fix bugs and improve that about HPX_HAVE_CXX11_AUTO_RETURN_VALUE of CMake
- PR #2820¹⁵⁵⁸ - Support C++11 in benchmark codes of parallel::partition and parallel::partition_copy
- PR #2819¹⁵⁵⁹ - Fix compile errors in unit test of container version of parallel::partition
- Issue #2818¹⁵⁶⁰ - A strange thing in parallel::util::detail::handle_local_exceptions
- Issue #2815¹⁵⁶¹ - HPX fails to compile with HPX_WITH_CUDA=ON and the new CUDA 9.0 RC
- Issue #2814¹⁵⁶² - Using ‘gmakeN’ after ‘cmake’ produces error in src/CMakeFiles/hpx.dir/runtime/agas/addressing_service.cpp.o
- PR #2813¹⁵⁶³ - Properly support [[noreturn]] attribute if available
- Issue #2812¹⁵⁶⁴ - Compilation fails with gcc 7.1.1
- PR #2811¹⁵⁶⁵ - Adding hpx::launch::lazy and support for async, dataflow, and future::then
- PR #2810¹⁵⁶⁶ - Add option allowing to disable deprecation warning
- PR #2809¹⁵⁶⁷ - Disable throttling scheduler if HWLOC is not found/used
- PR #2808¹⁵⁶⁸ - Fix compile errors on some environments of parallel::partition
- Issue #2807¹⁵⁶⁹ - Difficulty building with HPX_WITH_HWLOC=Off
- PR #2806¹⁵⁷⁰ - Partitioned vector
- PR #2805¹⁵⁷¹ - Serializing collections with non-default constructible data
- PR #2802¹⁵⁷² - Fix FreeBSD 11
- Issue #2801¹⁵⁷³ - Rate limiting techniques in io_service
- Issue #2800¹⁵⁷⁴ - New Launch Policy: async_if
- PR #2799¹⁵⁷⁵ - Fix a unit test failure on GCC in tuple_cat
- PR #2798¹⁵⁷⁶ - bump minimum required cmake to 3.0 in test
- PR #2797¹⁵⁷⁷ - Making sure future::wait_for et.al. work properly for action results
- Issue #2796¹⁵⁷⁸ - wait_for does always in “deferred” state for calls on remote localities

¹⁵⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2822>

¹⁵⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2821>

¹⁵⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2820>

¹⁵⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2819>

¹⁵⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2818>

¹⁵⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/2815>

¹⁵⁶² <https://github.com/STELLAR-GROUP/hpx/issues/2814>

¹⁵⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/2813>

¹⁵⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2812>

¹⁵⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2811>

¹⁵⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2810>

¹⁵⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2809>

¹⁵⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2808>

¹⁵⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2807>

¹⁵⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2806>

¹⁵⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2805>

¹⁵⁷² <https://github.com/STELLAR-GROUP/hpx/pull/2802>

¹⁵⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2801>

¹⁵⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2800>

¹⁵⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2799>

¹⁵⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2798>

¹⁵⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2797>

¹⁵⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2796>

- Issue #2795¹⁵⁷⁹ - Serialization of types without default constructor
- PR #2794¹⁵⁸⁰ - Fixing test for partitioned_vector iteration
- PR #2792¹⁵⁸¹ - Implemented segmented find and its variations for partitioned vector
- PR #2791¹⁵⁸² - Circumvent scary warning about placement new
- PR #2790¹⁵⁸³ - Fix OSX build
- PR #2789¹⁵⁸⁴ - Resource partitioner
- PR #2788¹⁵⁸⁵ - Adapt parallel::is_heap and parallel::is_heap_until to Ranges TS
- PR #2787¹⁵⁸⁶ - Unwrap hotfixes
- PR #2786¹⁵⁸⁷ - Update CMake Minimum Version to 3.3.2 (refs #2565)
- Issue #2785¹⁵⁸⁸ - Issues with masks and cpuset
- PR #2784¹⁵⁸⁹ - Error with reduce and transform reduce fixed
- PR #2783¹⁵⁹⁰ - StackOverflow integration with libsigsegv
- PR #2782¹⁵⁹¹ - Replace boost::atomic with std::atomic (where possible)
- PR #2781¹⁵⁹² - Check for and optionally use [[deprecated]] attribute
- PR #2780¹⁵⁹³ - Adding empty (but non-trivial) destructor to circumvent warnings
- PR #2779¹⁵⁹⁴ - Exception info tweaks
- PR #2778¹⁵⁹⁵ - Implement parallel::partition
- PR #2777¹⁵⁹⁶ - Improve error handling in gather_here/gather_there
- PR #2776¹⁵⁹⁷ - Fix a bug in compiler version check
- PR #2775¹⁵⁹⁸ - Fix compilation when HPX_WITH_LOGGING is OFF
- PR #2774¹⁵⁹⁹ - Removing dependency on Boost.Date_Time
- PR #2773¹⁶⁰⁰ - Add sync_images() method to spmd_block class
- PR #2772¹⁶⁰¹ - Adding documentation for PAPI counters

¹⁵⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2795>

¹⁵⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2794>

¹⁵⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2792>

¹⁵⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2791>

¹⁵⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2790>

¹⁵⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2789>

¹⁵⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2788>

¹⁵⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2787>

¹⁵⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2786>

¹⁵⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2785>

¹⁵⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2784>

¹⁵⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2783>

¹⁵⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2782>

¹⁵⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2781>

¹⁵⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2780>

¹⁵⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2779>

¹⁵⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2778>

¹⁵⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2777>

¹⁵⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2776>

¹⁵⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2775>

¹⁵⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2774>

¹⁶⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2773>

¹⁶⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2772>

- PR #2771¹⁶⁰² - Removing boost preprocessor dependency
- PR #2770¹⁶⁰³ - Adding test, fixing deadlock in config registry
- PR #2769¹⁶⁰⁴ - Remove some other warnings and errors detected by clang 5.0
- Issue #2768¹⁶⁰⁵ - Is there iterator tag for HPX?
- PR #2767¹⁶⁰⁶ - Improvements to continuation annotation
- PR #2765¹⁶⁰⁷ - gcc split stack support for HPX threads #620
- PR #2764¹⁶⁰⁸ - Fix some uses of begin/end, remove unnecessary includes
- PR #2763¹⁶⁰⁹ - Bump minimal Boost version to 1.55.0
- PR #2762¹⁶¹⁰ - hpx::partitioned_vector serializer
- PR #2761¹⁶¹¹ - Adding configuration summary to cmake output and –hpx:info
- PR #2760¹⁶¹² - Removing 1d_hydro example as it is broken
- PR #2758¹⁶¹³ - Remove various warnings detected by clang 5.0
- Issue #2757¹⁶¹⁴ - In case of a “raw thread” is needed per core for implementing parallel algorithm, what is good practice in HPX?
- PR #2756¹⁶¹⁵ - Allowing for LCOs to be simple components
- PR #2755¹⁶¹⁶ - Removing make_index_pack_unrolled
- PR #2754¹⁶¹⁷ - Implement parallel::unique_copy
- PR #2753¹⁶¹⁸ - Fixing detection of [[fallthrough]] attribute
- PR #2752¹⁶¹⁹ - New thread priority names
- PR #2751¹⁶²⁰ - Replace boost::exception with proposed exception_info
- PR #2750¹⁶²¹ - Replace boost::iterator_range
- PR #2749¹⁶²² - Fixing hdf5 examples
- Issue #2748¹⁶²³ - HPX fails to build with enabled hdf5 examples
- Issue #2747¹⁶²⁴ - Inherited task priorities break certain DAG optimizations

¹⁶⁰² <https://github.com/STELLAR-GROUP/hpx/pull/2771>

¹⁶⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2770>

¹⁶⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2769>

¹⁶⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2768>

¹⁶⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2767>

¹⁶⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2765>

¹⁶⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2764>

¹⁶⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2763>

¹⁶¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2762>

¹⁶¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2761>

¹⁶¹² <https://github.com/STELLAR-GROUP/hpx/pull/2760>

¹⁶¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2758>

¹⁶¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2757>

¹⁶¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2756>

¹⁶¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2755>

¹⁶¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2754>

¹⁶¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2753>

¹⁶¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2752>

¹⁶²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2751>

¹⁶²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2750>

¹⁶²² <https://github.com/STELLAR-GROUP/hpx/pull/2749>

¹⁶²³ <https://github.com/STELLAR-GROUP/hpx/issues/2748>

¹⁶²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2747>

- Issue #2746¹⁶²⁵ - HPX segfaulting with valgrind
- PR #2745¹⁶²⁶ - Adding extended arithmetic performance counters
- PR #2744¹⁶²⁷ - Adding ability to statistics counters to reset base counter
- Issue #2743¹⁶²⁸ - Statistics counter does not support resetting
- PR #2742¹⁶²⁹ - Making sure Vc V2 builds without additional HPX configuration flags
- PR #2741¹⁶³⁰ - Deprecate unwrapped and implement unwrap and unwrapping
- PR #2740¹⁶³¹ - Coroutine stackoverflow detection for linux posix; Issue #2408
- PR #2739¹⁶³² - Add files via upload
- PR #2738¹⁶³³ - Appveyor support
- PR #2737¹⁶³⁴ - Fixing 2735
- Issue #2736¹⁶³⁵ - 1d_hydro example doesn't work
- Issue #2735¹⁶³⁶ - partitioned_vector_subview test failing
- PR #2734¹⁶³⁷ - Add C++11 range utilities
- PR #2733¹⁶³⁸ - Adapting iterator requirements for parallel algorithms
- PR #2732¹⁶³⁹ - Integrate C++ Co-arrays
- PR #2731¹⁶⁴⁰ - Adding on_migrated event handler to migratable component instances
- Issue #2729¹⁶⁴¹ - Add on_migrated() event handler to migratable components
- Issue #2728¹⁶⁴² - Why Projection is needed in parallel algorithms?
- PR #2727¹⁶⁴³ - Cmake files for StackOverflow Detection
- PR #2726¹⁶⁴⁴ - CMake for Stack Overflow Detection
- PR #2725¹⁶⁴⁵ - Implemented segmented algorithms for partitioned vector
- PR #2724¹⁶⁴⁶ - Fix examples in Action documentation
- PR #2723¹⁶⁴⁷ - Enable lcos::channel<T>::register_as

¹⁶²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2746>

¹⁶²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2745>

¹⁶²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2744>

¹⁶²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2743>

¹⁶²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2742>

¹⁶³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2741>

¹⁶³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2740>

¹⁶³² <https://github.com/STELLAR-GROUP/hpx/pull/2739>

¹⁶³³ <https://github.com/STELLAR-GROUP/hpx/pull/2738>

¹⁶³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2737>

¹⁶³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2736>

¹⁶³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2735>

¹⁶³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2734>

¹⁶³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2733>

¹⁶³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2732>

¹⁶⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2731>

¹⁶⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/2729>

¹⁶⁴² <https://github.com/STELLAR-GROUP/hpx/issues/2728>

¹⁶⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/2727>

¹⁶⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2726>

¹⁶⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2725>

¹⁶⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2724>

¹⁶⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2723>

- Issue #2722¹⁶⁴⁸ - channel register_as() failing on compilation
- PR #2721¹⁶⁴⁹ - Mind map
- PR #2720¹⁶⁵⁰ - reorder forward declarations to get rid of C++14-only auto return types
- PR #2719¹⁶⁵¹ - Add documentation for partitioned_vector and add features in pack.hpp
- Issue #2718¹⁶⁵² - Some forward declarations in execution_fwd.hpp aren't C++11-compatible
- PR #2717¹⁶⁵³ - Config support for fallthrough attribute
- PR #2716¹⁶⁵⁴ - Implement parallel::partition_copy
- PR #2715¹⁶⁵⁵ - initial import of icu string serializer
- PR #2714¹⁶⁵⁶ - initial import of valarray serializer
- PR #2713¹⁶⁵⁷ - Remove slashes before CMAKE_FILES_DIRECTORY variables
- PR #2712¹⁶⁵⁸ - Fixing wait for 1751
- PR #2711¹⁶⁵⁹ - Adjust code for minimal supported GCC having been bumped to 4.9
- PR #2710¹⁶⁶⁰ - Adding code of conduct
- PR #2709¹⁶⁶¹ - Fixing UB in destroy tests
- PR #2708¹⁶⁶² - Add inline to prevent multiple definition issue
- Issue #2707¹⁶⁶³ - Multiple defined symbols for task_block.hpp in VS2015
- PR #2706¹⁶⁶⁴ - Adding .clang-format file
- PR #2704¹⁶⁶⁵ - Add a synchronous mapping API
- Issue #2703¹⁶⁶⁶ - Request: Add the .clang-format file to the repository
- Issue #2702¹⁶⁶⁷ - STELLAR-GROUP/Vc slower than VCv1 possibly due to wrong instructions generated
- Issue #2701¹⁶⁶⁸ - Datapar with STELLAR-GROUP/Vc requires obscure flag
- Issue #2700¹⁶⁶⁹ - Naming inconsistency in parallel algorithms
- Issue #2699¹⁶⁷⁰ - Iterator requirements are different from standard in parallel copy_if.

¹⁶⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2722>

¹⁶⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2721>

¹⁶⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2720>

¹⁶⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2719>

¹⁶⁵² <https://github.com/STELLAR-GROUP/hpx/issues/2718>

¹⁶⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2717>

¹⁶⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2716>

¹⁶⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2715>

¹⁶⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2714>

¹⁶⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2713>

¹⁶⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2712>

¹⁶⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2711>

¹⁶⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2710>

¹⁶⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/2709>

¹⁶⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2708>

¹⁶⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/2707>

¹⁶⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2706>

¹⁶⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2704>

¹⁶⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2703>

¹⁶⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2702>

¹⁶⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2701>

¹⁶⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2700>

¹⁶⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2699>

- PR #2698¹⁶⁷¹ - Properly releasing parcelport write handlers
- Issue #2697¹⁶⁷² - Compile error in addressing_service.cpp
- Issue #2696¹⁶⁷³ - Building and using HPX statically: undefined references from runtime_support_server.cpp
- Issue #2695¹⁶⁷⁴ - Executor changes cause compilation failures
- PR #2694¹⁶⁷⁵ - Refining C++ language mode detection for MSVC
- PR #2693¹⁶⁷⁶ - P0443 r2
- PR #2692¹⁶⁷⁷ - Partially reverting changes to parcel_await
- Issue #2689¹⁶⁷⁸ - HPX build fails when HPX_WITH_CUDA is enabled
- PR #2688¹⁶⁷⁹ - Make Cuda Clang builds pass
- PR #2687¹⁶⁸⁰ - Add an is_tuple_like trait for sequenceable type detection
- PR #2686¹⁶⁸¹ - Allowing throttling scheduler to be used without idle backoff
- PR #2685¹⁶⁸² - Add support of std::array to hpx::util::tuple_size and tuple_element
- PR #2684¹⁶⁸³ - Adding new statistics performance counters
- PR #2683¹⁶⁸⁴ - Replace boost::exception_ptr with std::exception_ptr
- Issue #2682¹⁶⁸⁵ - HPX does not compile with HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF=OFF
- PR #2681¹⁶⁸⁶ - Attempt to fix problem in managed_component_base
- PR #2680¹⁶⁸⁷ - Fix bad size during archive creation
- Issue #2679¹⁶⁸⁸ - Mismatch between size of archive and container
- Issue #2678¹⁶⁸⁹ - In parallel algorithm, other tasks are executed to the end even if an exception occurs in any task.
- PR #2677¹⁶⁹⁰ - Adding include check for std::addressof
- PR #2676¹⁶⁹¹ - Adding parallel::destroy and destroy_n
- PR #2675¹⁶⁹² - Making sure statistics counters work as expected
- PR #2674¹⁶⁹³ - Turning assertions into exceptions

¹⁶⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2698>

¹⁶⁷² <https://github.com/STELLAR-GROUP/hpx/issues/2697>

¹⁶⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2696>

¹⁶⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2695>

¹⁶⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2694>

¹⁶⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2693>

¹⁶⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2692>

¹⁶⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2689>

¹⁶⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2688>

¹⁶⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2687>

¹⁶⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2686>

¹⁶⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2685>

¹⁶⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2684>

¹⁶⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2683>

¹⁶⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2682>

¹⁶⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2681>

¹⁶⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2680>

¹⁶⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2679>

¹⁶⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2678>

¹⁶⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2677>

¹⁶⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2676>

¹⁶⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2675>

¹⁶⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2674>

- PR #2673¹⁶⁹⁴ - Inhibit direct conversion from future<future<T>> -> future<void>
- PR #2672¹⁶⁹⁵ - C++17 invoke forms
- PR #2671¹⁶⁹⁶ - Adding uninitialized_value_construct and uninitialized_value_construct_n
- PR #2670¹⁶⁹⁷ - Integrate spmd multidimensional views for partitioned_vectors
- PR #2669¹⁶⁹⁸ - Adding uninitialized_default_construct and uninitialized_default_construct_n
- PR #2668¹⁶⁹⁹ - Fixing documentation index
- Issue #2667¹⁷⁰⁰ - Ambiguity of nested hpx::future<void>'s
- Issue #2666¹⁷⁰¹ - Statistics Performance counter is not working
- PR #2664¹⁷⁰² - Adding uninitialized_move and uninitialized_move_n
- Issue #2663¹⁷⁰³ - Seg fault in managed_component::get_base_gid, possibly cause by util::reinitializable_static
- Issue #2662¹⁷⁰⁴ - Crash in managed_component::get_base_gid due to problem with util::reinitializable_static
- PR #2665¹⁷⁰⁵ - Hide the detail namespace in doxygen per default
- PR #2660¹⁷⁰⁶ - Add documentation to hpx::util::unwrapped and hpx::util::unwrapped2
- PR #2659¹⁷⁰⁷ - Improve integration with vcpkg
- PR #2658¹⁷⁰⁸ - Unify access_data trait for use in both, serialization and de-serialization
- PR #2657¹⁷⁰⁹ - Removing hpx::lcos::queue<T>
- PR #2656¹⁷¹⁰ - Reduce MAX_TERMINATED_THREADS default, improve memory use on manycore cpus
- PR #2655¹⁷¹¹ - Maintenance for emulate-deleted macros
- PR #2654¹⁷¹² - Implement parallel is_heap and is_heap_until
- PR #2653¹⁷¹³ - Drop support for VS2013
- PR #2652¹⁷¹⁴ - This patch makes sure that all parcels in a batch are properly handled
- PR #2649¹⁷¹⁵ - Update docs (Table 18) - move transform to end
- Issue #2647¹⁷¹⁶ - hpx::parcelset::detail::parcel_data::**has_continuation** is uninitialized

¹⁶⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2673>

¹⁶⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2672>

¹⁶⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2671>

¹⁶⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2670>

¹⁶⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2669>

¹⁶⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2668>

¹⁷⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2667>

¹⁷⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/2666>

¹⁷⁰² <https://github.com/STELLAR-GROUP/hpx/pull/2664>

¹⁷⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/2663>

¹⁷⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2662>

¹⁷⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2665>

¹⁷⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2660>

¹⁷⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2659>

¹⁷⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2658>

¹⁷⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2657>

¹⁷¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2656>

¹⁷¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2655>

¹⁷¹² <https://github.com/STELLAR-GROUP/hpx/pull/2654>

¹⁷¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2653>

¹⁷¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2652>

¹⁷¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2649>

¹⁷¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2647>

- Issue #2644¹⁷¹⁷ - Some .vcxproj in the HPX.sln fail to build
- Issue #2641¹⁷¹⁸ - hpx::lcos::queue should be deprecated
- PR #2640¹⁷¹⁹ - A new throttling policy with public APIs to suspend/resume
- PR #2639¹⁷²⁰ - Fix a tiny typo in tutorial.
- Issue #2638¹⁷²¹ - Invalid return type ‘void’ of constexpr function
- PR #2636¹⁷²² - Add and use HPX_MSVC_WARNING_PRAGMA for #pragma warning
- PR #2633¹⁷²³ - Distributed define_spmd_block
- PR #2632¹⁷²⁴ - Making sure container serialization uses size-compatible types
- PR #2631¹⁷²⁵ - Add lcos::local::one_element_channel
- PR #2629¹⁷²⁶ - Move unordered_map out of parcelport into hpx/concurrent
- PR #2628¹⁷²⁷ - Making sure that shutdown does not hang
- PR #2627¹⁷²⁸ - Fix serialization
- PR #2626¹⁷²⁹ - Generate cmake_variables.qbk and cmake_toolchains.qbk outside of the source tree
- PR #2625¹⁷³⁰ - Supporting -std=c++17 flag
- PR #2624¹⁷³¹ - Fixing a small cmake typo
- PR #2622¹⁷³² - Update CMake minimum required version to 3.0.2 (closes #2621)
- Issue #2621¹⁷³³ - Compiling hpx master fails with /usr/bin/ld: final link failed: Bad value
- PR #2620¹⁷³⁴ - Remove warnings due to some captured variables
- PR #2619¹⁷³⁵ - LF multiple parcels
- PR #2618¹⁷³⁶ - Some fixes to libfabric that didn’t get caught before the merge
- PR #2617¹⁷³⁷ - Adding hpx::local_new
- PR #2616¹⁷³⁸ - Documentation: Extract all entities in order to autolink functions correctly
- Issue #2615¹⁷³⁹ - Documentation: Linking functions is broken

¹⁷¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2644>

¹⁷¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2641>

¹⁷¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2640>

¹⁷²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2639>

¹⁷²¹ <https://github.com/STELLAR-GROUP/hpx/issues/2638>

¹⁷²² <https://github.com/STELLAR-GROUP/hpx/pull/2636>

¹⁷²³ <https://github.com/STELLAR-GROUP/hpx/pull/2633>

¹⁷²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2632>

¹⁷²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2631>

¹⁷²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2629>

¹⁷²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2628>

¹⁷²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2627>

¹⁷²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2626>

¹⁷³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2625>

¹⁷³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2624>

¹⁷³² <https://github.com/STELLAR-GROUP/hpx/pull/2622>

¹⁷³³ <https://github.com/STELLAR-GROUP/hpx/issues/2621>

¹⁷³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2620>

¹⁷³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2619>

¹⁷³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2618>

¹⁷³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2617>

¹⁷³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2616>

¹⁷³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2615>

- PR #2614¹⁷⁴⁰ - Adding serialization for std::deque
- PR #2613¹⁷⁴¹ - We need to link with boost.thread and boost.chrono if we use boost.context
- PR #2612¹⁷⁴² - Making sure for_loop_n(par, ...) is actually executed in parallel
- PR #2611¹⁷⁴³ - Add documentation to invoke_fused and friends NFC
- PR #2610¹⁷⁴⁴ - Added reduction templates using an identity value
- PR #2608¹⁷⁴⁵ - Fixing some unused vars in inspect
- PR #2607¹⁷⁴⁶ - Fixed build for mingw
- PR #2606¹⁷⁴⁷ - Supporting generic context for boost >= 1.61
- PR #2605¹⁷⁴⁸ - Parcelport libfabric3
- PR #2604¹⁷⁴⁹ - Adding allocator support to promise and friends
- PR #2603¹⁷⁵⁰ - Barrier hang
- PR #2602¹⁷⁵¹ - Changes to scheduler to steal from one high-priority queue
- Issue #2601¹⁷⁵² - High priority tasks are not executed first
- PR #2600¹⁷⁵³ - Compat fixes
- PR #2599¹⁷⁵⁴ - Compatibility layer for threading support
- PR #2598¹⁷⁵⁵ - V1.1
- PR #2597¹⁷⁵⁶ - Release V1.0
- PR #2592¹⁷⁵⁷ - First attempt to introduce spmd_block in hpx
- PR #2586¹⁷⁵⁸ - local_segment in segmented_iterator_traits
- Issue #2584¹⁷⁵⁹ - Add allocator support to promise, packaged_task and friends
- PR #2576¹⁷⁶⁰ - Add missing dependencies of cuda based tests
- PR #2575¹⁷⁶¹ - Remove warnings due to some captured variables
- Issue #2574¹⁷⁶² - MSVC 2015 Compiler crash when building HPX

¹⁷⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2614>

¹⁷⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2613>

¹⁷⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2612>

¹⁷⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/2611>

¹⁷⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2610>

¹⁷⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2608>

¹⁷⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2607>

¹⁷⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2606>

¹⁷⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2605>

¹⁷⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2604>

¹⁷⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2603>

¹⁷⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2602>

¹⁷⁵² <https://github.com/STELLAR-GROUP/hpx/issues/2601>

¹⁷⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2600>

1754 <https://github.com/STELLAR-GROUP/hpx/pull/2599>1755 <https://github.com/STELLAR-GROUP/hpx/pull/2598>1756 <https://github.com/STELLAR-GROUP/hpx/pull/2597>1757 <https://github.com/STELLAR-GROUP/hpx/pull/2592>1758 <https://github.com/STELLAR-GROUP/hpx/pull/2586>1759 <https://github.com/STELLAR-GROUP/hpx/issues/2584>1760 <https://github.com/STELLAR-GROUP/hpx/pull/2576>1761 <https://github.com/STELLAR-GROUP/hpx/pull/2575>1762 <https://github.com/STELLAR-GROUP/hpx/issues/2574>

- Issue #2568¹⁷⁶³ - Remove throttle_scheduler as it has been abandoned
- Issue #2566¹⁷⁶⁴ - Add an inline versioning namespace before 1.0 release
- Issue #2565¹⁷⁶⁵ - Raise minimal cmake version requirement
- PR #2556¹⁷⁶⁶ - Fixing scan partitioner
- PR #2546¹⁷⁶⁷ - Broadcast async
- Issue #2543¹⁷⁶⁸ - make install fails due to a non-existing .so file
- PR #2495¹⁷⁶⁹ - wait_or_add_new returning thread_id_type
- Issue #2480¹⁷⁷⁰ - Unable to register new performance counter
- Issue #2471¹⁷⁷¹ - no type named ‘fcontext_t’ in namespace
- Issue #2456¹⁷⁷² - Re-implement hpx::util::unwrapped
- Issue #2455¹⁷⁷³ - Add more arithmetic performance counters
- PR #2454¹⁷⁷⁴ - Fix a couple of warnings and compiler errors
- PR #2453¹⁷⁷⁵ - Timed executor support
- PR #2447¹⁷⁷⁶ - Implementing new executor API (P0443)
- Issue #2439¹⁷⁷⁷ - Implement executor proposal
- Issue #2408¹⁷⁷⁸ - Stackoverflow detection for linux, e.g. based on libsigsegv
- PR #2377¹⁷⁷⁹ - Add a customization point for put_parcel so we can override actions
- Issue #2368¹⁷⁸⁰ - HPX_ASSERT problem
- Issue #2324¹⁷⁸¹ - Change default number of threads used to the maximum of the system
- Issue #2266¹⁷⁸² - hpx_0.9.99 make tests fail
- PR #2195¹⁷⁸³ - Support for code completion in VIM
- Issue #2137¹⁷⁸⁴ - Hpx does not compile over osx
- Issue #2092¹⁷⁸⁵ - make tests should just build the tests

¹⁷⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/2568>

¹⁷⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2566>

¹⁷⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2565>

¹⁷⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2556>

¹⁷⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2546>

¹⁷⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2543>

¹⁷⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2495>

¹⁷⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2480>

¹⁷⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/2471>

¹⁷⁷² <https://github.com/STELLAR-GROUP/hpx/issues/2456>

¹⁷⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2455>

¹⁷⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2454>

¹⁷⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2453>

¹⁷⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2447>

¹⁷⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2439>

¹⁷⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2408>

¹⁷⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2377>

¹⁷⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2368>

¹⁷⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/2324>

¹⁷⁸² <https://github.com/STELLAR-GROUP/hpx/issues/2266>

¹⁷⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2195>

¹⁷⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2137>

¹⁷⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2092>

- Issue #2026¹⁷⁸⁶ - Build HPX with Apple's clang
- Issue #1932¹⁷⁸⁷ - hpx with PBS fails on multiple localities
- PR #1914¹⁷⁸⁸ - Parallel heap algorithm implementations WIP
- Issue #1598¹⁷⁸⁹ - Disconnecting a locality results in segfault using heartbeat example
- Issue #1404¹⁷⁹⁰ - unwrapped doesn't work with movable only types
- Issue #1400¹⁷⁹¹ - hpx::util::unwrapped doesn't work with non-future types
- Issue #1205¹⁷⁹² - TSS is broken
- Issue #1126¹⁷⁹³ - vector<future<T>> does not work gracefully with dataflow, when_all and unwrapped
- Issue #1056¹⁷⁹⁴ - Thread manager cleanup
- Issue #863¹⁷⁹⁵ - Futures should not require a default constructor
- Issue #856¹⁷⁹⁶ - Allow runtimemode_connect to be used with security enabled
- Issue #726¹⁷⁹⁷ - Valgrind
- Issue #701¹⁷⁹⁸ - Add RCR performance counter component
- Issue #528¹⁷⁹⁹ - Add support for known failures and warning count/comparisons to hpx_run_tests.py

2.11.6 HPX V1.0.0 (Apr 24, 2017)

General changes

Here are some of the main highlights and changes for this release (in no particular order):

- Added the facility `hpx::split_future` which allows one to convert a `future<tuple<Ts...>>` into a `tuple<future<Ts>...>`. This functionality is not available when compiling *HPX* with VS2012.
- Added a new type of performance counter which allows one to return a list of values for each invocation. We also added a first counter of this type which collects a histogram of the times between parcels being created.
- Added new LCOs: `hpx::lcos::channel` and `hpx::lcos::local::channel` which are very similar to the well known channel constructs used in the Go language.
- Added new performance counters reporting the amount of data handled by the networking layer on a action-by-action basis (please see PR #2289¹⁸⁰⁰ for more details).
- Added a new facility `hpx::lcos::barrier`, replacing the equally named older one. The new facility has a slightly changed API and is much more efficient. Most notable, the new facility exposes a (global) function `hpx::lcos::barrier::synchronize()` which represents a global barrier across all localities.

¹⁷⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2026>

¹⁷⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1932>

¹⁷⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1914>

¹⁷⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1598>

¹⁷⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1404>

¹⁷⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1400>

¹⁷⁹² <https://github.com/STELLAR-GROUP/hpx/issues/1205>

¹⁷⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1126>

¹⁷⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1056>

¹⁷⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/863>

¹⁷⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/856>

¹⁷⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/726>

¹⁷⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/701>

¹⁷⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/528>

¹⁸⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2289>

- We have started to add support for vectorization to our parallel algorithm implementations. This support depends on using an external library, currently either Vc Library or [boost_simd](#). Please see [Issue #2333¹⁸⁰¹](#) for a list of currently supported algorithms. This is an experimental feature and its implementation and/or API might change in the future. Please see this [blog-post¹⁸⁰²](#) for more information.
- The parameter sequence for the `hpx::parallel::transform_reduce` overload taking one iterator range has changed to match the changes this algorithm has undergone while being moved to C++17. The old overload can be still enabled at configure time by specifying `-DHPX_WITH_TRANSFORM_REDUCE_COMPATIBILITY=On` to CMake.
- The algorithm `hpx::parallel::inner_product` has been renamed to `hpx::parallel::transform_reduce` to match the changes this algorithm has undergone while being moved to C++17. The old `inner_product` names can be still enabled at configure time by specifying `-DHPX_WITH_TRANSFORM_REDUCE_COMPATIBILITY=On` to CMake.
- Added versions of `hpx::get_ptr` taking client side representations for component instances as their parameter (instead of a global id).
- Added the helper utility `hpx::performance_counters::performance_counter_set` helping to encapsulate a set of performance counters to be managed concurrently.
- All execution policies and related classes have been renamed to be consistent with the naming changes applied for C++17. All policies now live in the namespace `hpx::parallel::execution`. The old names can be still enabled at configure time by specifying `-DHPX_WITH_EXECUTION_POLICY_COMPATIBILITY=On` to CMake.
- The thread scheduling subsystem has undergone a major refactoring which results in significant performance improvements. We have also improved the performance of creating `hpx::future` and of various facilities handling those.
- We have consolidated all of the code in HPX.Compute related to the integration of CUDA. `hpx::partitioned_vector` has been enabled to be usable with `hpx::compute::vector` which allows one to place the partitions on one or more GPU devices.
- Added new performance counters exposing various internals of the thread scheduling subsystem, such as the current idle- and busy-loop counters and instantaneous scheduler utilization.
- Extended and improved the use of the ITTNotify hooks allowing to collect performance counter data and function annotation information from within the Intel Amplifier tool.

Breaking changes

- We have dropped support for the gcc compiler versions V4.6 and 4.7. The minimal gcc version we now test on is gcc V4.8.
- We have removed (default) support for `boost::chrono` in interfaces, uses of it have been replaced with `std::chrono`. This facility can be still enabled at configure time by specifying `-DHPX_WITH_BOOST_CHRONO_COMPATIBILITY=On` to CMake.
- The parameter sequence for the `hpx::parallel::transform_reduce` overload taking one iterator range has changed to match the changes this algorithm has undergone while being moved to C++17.
- The algorithm `hpx::parallel::inner_product` has been renamed to `hpx::parallel::transform_reduce` to match the changes this algorithm has undergone while being moved to C++17.

¹⁸⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/2333>

¹⁸⁰² <http://stellar-group.org/2016/09/vectorized-cpp-parallel-algorithms-with-hpx/>

- the build options `HPX_WITH_COLOCATED_BACKWARDS_COMPATIBILITY` and `HPX_WITH_COMPONENT_GET_GID_COMPATIBILITY` are now disabled by default. Please change your code still depending on the deprecated interfaces.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- PR #2596¹⁸⁰³ - Adding apex data
- PR #2595¹⁸⁰⁴ - Remove obsolete file
- Issue #2594¹⁸⁰⁵ - FindOpenCL.cmake mismatch with the official cmake module
- PR #2592¹⁸⁰⁶ - First attempt to introduce spmd_block in hpx
- Issue #2591¹⁸⁰⁷ - Feature request: continuation (then) which does not require the callable object to take a future<R> as parameter
- PR #2588¹⁸⁰⁸ - Daint fixes
- PR #2587¹⁸⁰⁹ - Fixing transfer_(continuation)_action::schedule
- PR #2585¹⁸¹⁰ - Work around MSVC having an ICE when compiling with -Ob2
- PR #2583¹⁸¹¹ - changing 7zip command to 7za in roll_release.sh
- PR #2582¹⁸¹² - First attempt to introduce spmd_block in hpx
- PR #2581¹⁸¹³ - Enable annotated function for parallel algorithms
- PR #2580¹⁸¹⁴ - First attempt to introduce spmd_block in hpx
- PR #2579¹⁸¹⁵ - Make thread NICE level setting an option
- PR #2578¹⁸¹⁶ - Implementing enqueue instead of busy wait when no sender is available
- PR #2577¹⁸¹⁷ - Retrieve -std=c++11 consistent nvcc flag
- PR #2576¹⁸¹⁸ - Add missing dependencies of cuda based tests
- PR #2575¹⁸¹⁹ - Remove warnings due to some captured variables
- PR #2573¹⁸²⁰ - Attempt to resolve resolve_locality
- PR #2572¹⁸²¹ - Adding APEX hooks to background thread

¹⁸⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2596>

¹⁸⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2595>

¹⁸⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2594>

¹⁸⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2592>

¹⁸⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2591>

¹⁸⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2588>

¹⁸⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2587>

¹⁸¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2585>

¹⁸¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2583>

¹⁸¹² <https://github.com/STELLAR-GROUP/hpx/pull/2582>

¹⁸¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2581>

¹⁸¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2580>

¹⁸¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2579>

¹⁸¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2578>

¹⁸¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2577>

¹⁸¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2576>

¹⁸¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2575>

¹⁸²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2573>

¹⁸²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2572>

- PR #2571¹⁸²² - Pick up hpx.ignore_batch_env from config map
- PR #2570¹⁸²³ - Add commandline options –hpx:print-counters-locally
- PR #2569¹⁸²⁴ - Fix computeapi unit tests
- PR #2567¹⁸²⁵ - This adds another barrier::synchronize before registering performance counters
- PR #2564¹⁸²⁶ - Cray static toolchain support
- PR #2563¹⁸²⁷ - Fixed unhandled exception during startup
- PR #2562¹⁸²⁸ - Remove partitioned_vector.cu from build tree when nvcc is used
- Issue #2561¹⁸²⁹ - octo-tiger crash with commit 6e921495ff6c26f125d62629cbaad0525f14f7ab
- PR #2560¹⁸³⁰ - Prevent -Wundef warnings on Vc version checks
- PR #2559¹⁸³¹ - Allowing CUDA callback to set the future directly from an OS thread
- PR #2558¹⁸³² - Remove warnings due to float precisions
- PR #2557¹⁸³³ - Removing bogus handling of compile flags for CUDA
- PR #2556¹⁸³⁴ - Fixing scan partitioner
- PR #2554¹⁸³⁵ - Add more diagnostics to error thrown from find_appropriate_destination
- Issue #2555¹⁸³⁶ - No valid parcelport configured
- PR #2553¹⁸³⁷ - Add cmake cuda_arch option
- PR #2552¹⁸³⁸ - Remove incomplete datapar bindings to libflatarray
- PR #2551¹⁸³⁹ - Rename hwloc_topology to hwloc_topology_info
- PR #2550¹⁸⁴⁰ - Apex api updates
- PR #2549¹⁸⁴¹ - Pre-include defines.hpp to get the macro HPX_HAVE_CUDA value
- PR #2548¹⁸⁴² - Fixing issue with disconnect
- PR #2546¹⁸⁴³ - Some fixes around cuda clang partitioned_vector example
- PR #2545¹⁸⁴⁴ - Fix uses of the Vc2 datapar flags; the value, not the type, should be passed to functions

¹⁸²² <https://github.com/STELLAR-GROUP/hpx/pull/2571>

¹⁸²³ <https://github.com/STELLAR-GROUP/hpx/pull/2570>

¹⁸²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2569>

¹⁸²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2567>

¹⁸²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2564>

¹⁸²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2563>

¹⁸²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2562>

¹⁸²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2561>

¹⁸³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2560>

¹⁸³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2559>

1832 <https://github.com/STELLAR-GROUP/hpx/pull/2558>1833 <https://github.com/STELLAR-GROUP/hpx/pull/2557>1834 <https://github.com/STELLAR-GROUP/hpx/pull/2556>1835 <https://github.com/STELLAR-GROUP/hpx/pull/2554>1836 <https://github.com/STELLAR-GROUP/hpx/issues/2555>1837 <https://github.com/STELLAR-GROUP/hpx/pull/2553>1838 <https://github.com/STELLAR-GROUP/hpx/pull/2552>1839 <https://github.com/STELLAR-GROUP/hpx/pull/2551>1840 <https://github.com/STELLAR-GROUP/hpx/pull/2550>1841 <https://github.com/STELLAR-GROUP/hpx/pull/2549>1842 <https://github.com/STELLAR-GROUP/hpx/pull/2548>1843 <https://github.com/STELLAR-GROUP/hpx/pull/2546>1844 <https://github.com/STELLAR-GROUP/hpx/pull/2545>

- PR #2542¹⁸⁴⁵ - Make HPX_WITH_MALLOC easier to use
- PR #2541¹⁸⁴⁶ - avoid recompiles when enabling/disabling examples
- PR #2540¹⁸⁴⁷ - Fixing usage of target_link_libraries()
- PR #2539¹⁸⁴⁸ - fix RPATH behaviour
- Issue #2538¹⁸⁴⁹ - HPX_WITH_CUDA corrupts compilation flags
- PR #2537¹⁸⁵⁰ - Add output of a Bazel Skylark extension for paths and compile options
- PR #2536¹⁸⁵¹ - Add counter exposing total available memory to Windows as well
- PR #2535¹⁸⁵² - Remove obsolete support for security
- Issue #2534¹⁸⁵³ - Remove command line option --hpx:run-as-server
- PR #2533¹⁸⁵⁴ - Pre-cache locality endpoints during bootstrap
- PR #2532¹⁸⁵⁵ - Fixing handling of GIDs during serialization preprocessing
- PR #2531¹⁸⁵⁶ - Amend uses of the term “functor”
- PR #2529¹⁸⁵⁷ - added counter for reading available memory
- PR #2527¹⁸⁵⁸ - Facilities to create actions from lambdas
- PR #2526¹⁸⁵⁹ - Updated docs: HPX_WITH_EXAMPLES
- PR #2525¹⁸⁶⁰ - Remove warnings related to unused captured variables
- Issue #2524¹⁸⁶¹ - CMAKE failed because it is missing: TCMALLOC_LIBRARY TCMALLOC_INCLUDE_DIR
- PR #2523¹⁸⁶² - Fixing compose_cb stack overflow
- PR #2522¹⁸⁶³ - Instead of unlocking, ignore the lock while creating the message handler
- PR #2521¹⁸⁶⁴ - Create LPROGRESS_ logging macro to simplify progress tracking and timings
- PR #2520¹⁸⁶⁵ - Intel 17 support
- PR #2519¹⁸⁶⁶ - Fix components example
- PR #2518¹⁸⁶⁷ - Fixing parcel scheduling

¹⁸⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2542>
¹⁸⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2541>
¹⁸⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2540>
¹⁸⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2539>
¹⁸⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2538>
¹⁸⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2537>
¹⁸⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2536>
¹⁸⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2535>
¹⁸⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/2534>
¹⁸⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2533>
¹⁸⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2532>
¹⁸⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2531>
¹⁸⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2529>
¹⁸⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2527>
¹⁸⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2526>
¹⁸⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2525>
¹⁸⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/2524>
¹⁸⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2523>
¹⁸⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/2522>
¹⁸⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2521>
¹⁸⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2520>
¹⁸⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2519>
¹⁸⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2518>

- Issue #2517¹⁸⁶⁸ - Race condition during Parcel Coalescing Handler creation
- Issue #2516¹⁸⁶⁹ - HPX locks up when using at least 256 localities
- Issue #2515¹⁸⁷⁰ - error: Install cannot find “/lib/hpx/libparcel_coalescing.so.0.9.99” but I can see that file
- PR #2514¹⁸⁷¹ - Making sure that all continuations of a shared_future are invoked in order
- PR #2513¹⁸⁷² - Fixing locks held during suspension
- PR #2512¹⁸⁷³ - MPI Parcelpoint improvements and fixes related to the background work changes
- PR #2511¹⁸⁷⁴ - Fixing bit-wise (zero-copy) serialization
- Issue #2509¹⁸⁷⁵ - Linking errors in hwloc_topology
- PR #2508¹⁸⁷⁶ - Added documentation for debugging with core files
- PR #2506¹⁸⁷⁷ - Fixing background work invocations
- PR #2505¹⁸⁷⁸ - Fix tuple serialization
- Issue #2504¹⁸⁷⁹ - Ensure continuations are called in the order they have been attached
- PR #2503¹⁸⁸⁰ - Adding serialization support for Vc v2 (datapar)
- PR #2502¹⁸⁸¹ - Resolve various, minor compiler warnings
- PR #2501¹⁸⁸² - Some other fixes around cuda examples
- Issue #2500¹⁸⁸³ - nvcc / cuda clang issue due to a missing -DHPX_WITH_CUDA flag
- PR #2499¹⁸⁸⁴ - Adding support for std::array to wait_all and friends
- PR #2498¹⁸⁸⁵ - Execute background work as HPX thread
- PR #2497¹⁸⁸⁶ - Fixing configuration options for spinlock-deadlock detection
- PR #2496¹⁸⁸⁷ - Accounting for different compilers in CrayKNL toolchain file
- PR #2494¹⁸⁸⁸ - Adding component base class which ties a component instance to a given executor
- PR #2493¹⁸⁸⁹ - Enable controlling amount of pending threads which must be available to allow thread stealing
- PR #2492¹⁸⁹⁰ - Adding new command line option –hpx:print-counter-reset

¹⁸⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2517>

¹⁸⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2516>

¹⁸⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2515>

¹⁸⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2514>

¹⁸⁷² <https://github.com/STELLAR-GROUP/hpx/pull/2513>

¹⁸⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/2512>

¹⁸⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2511>

¹⁸⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2509>

¹⁸⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2508>

¹⁸⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2506>

¹⁸⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2505>

¹⁸⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2504>

¹⁸⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2503>

¹⁸⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2502>

¹⁸⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2501>

¹⁸⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/2500>

¹⁸⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2499>

¹⁸⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2498>

¹⁸⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2497>

¹⁸⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2496>

¹⁸⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2494>

¹⁸⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2493>

¹⁸⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2492>

- PR #2491¹⁸⁹¹ - Resolve ambiguities when compiling with APEX
- PR #2490¹⁸⁹² - Resuming threads waiting on future with higher priority
- Issue #2489¹⁸⁹³ - nvcc issue because -std=c++11 appears twice
- PR #2488¹⁸⁹⁴ - Adding performance counters exposing the internal idle and busy-loop counters
- PR #2487¹⁸⁹⁵ - Allowing for plain suspend to reschedule thread right away
- PR #2486¹⁸⁹⁶ - Only flag HPX code for CUDA if HPX_WITH_CUDA is set
- PR #2485¹⁸⁹⁷ - Making thread-queue parameters runtime-configurable
- PR #2484¹⁸⁹⁸ - Added atomic counter for parcel-destinations
- PR #2483¹⁸⁹⁹ - Added priority-queue lifo scheduler
- PR #2482¹⁹⁰⁰ - Changing scheduler to steal only if more than a minimal number of tasks are available
- PR #2481¹⁹⁰¹ - Extending command line option –hpx:print-counter-destination to support value ‘none’
- PR #2479¹⁹⁰² - Added option to disable signal handler
- PR #2478¹⁹⁰³ - Making sure the sine performance counter module gets loaded only for the corresponding example
- Issue #2477¹⁹⁰⁴ - Breaking at a throw statement
- PR #2476¹⁹⁰⁵ - Annotated function
- PR #2475¹⁹⁰⁶ - Ensure that using %osthread% during logging will not throw for non-hpx threads
- PR #2474¹⁹⁰⁷ - Remove now superficial non_direct actions from base_lco and friends
- PR #2473¹⁹⁰⁸ - Refining support for ITTNotify
- PR #2472¹⁹⁰⁹ - Some fixes around hpx compute
- Issue #2470¹⁹¹⁰ - redefinition of boost::detail::spinlock
- Issue #2469¹⁹¹¹ - Dataflow performance issue
- PR #2468¹⁹¹² - Perf docs update
- PR #2466¹⁹¹³ - Guarantee to execute remote direct actions on HPX-thread

¹⁸⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2491>

¹⁸⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2490>

¹⁸⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/2489>

¹⁸⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2488>

¹⁸⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2487>

¹⁸⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2486>

¹⁸⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2485>

¹⁸⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2484>

¹⁸⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2483>

¹⁹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2482>

¹⁹⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2481>

¹⁹⁰² <https://github.com/STELLAR-GROUP/hpx/pull/2479>

¹⁹⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2478>

¹⁹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2477>

¹⁹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2476>

¹⁹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2475>

¹⁹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2474>

¹⁹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2473>

¹⁹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2472>

¹⁹¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2470>

¹⁹¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/2469>

¹⁹¹² <https://github.com/STELLAR-GROUP/hpx/pull/2468>

¹⁹¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2466>

- PR #2465¹⁹¹⁴ - Improve demo : Async copy and fixed device handling
- PR #2464¹⁹¹⁵ - Adding performance counter exposing instantaneous scheduler utilization
- PR #2463¹⁹¹⁶ - Downcast to future<void>
- PR #2462¹⁹¹⁷ - Fixed usage of ITT-Notify API with Intel Amplifier
- PR #2461¹⁹¹⁸ - Cublas demo
- PR #2460¹⁹¹⁹ - Fixing thread bindings
- PR #2459¹⁹²⁰ - Make -std=c++11 nvcc flag consistent for in-build and installed versions
- Issue #2457¹⁹²¹ - Segmentation fault when registering a partitioned vector
- PR #2452¹⁹²² - Properly releasing global barrier for unhandled exceptions
- PR #2451¹⁹²³ - Fixing long shutdown times
- PR #2450¹⁹²⁴ - Attempting to fix initialization errors on newer platforms (Boost V1.63)
- PR #2449¹⁹²⁵ - Replace BOOST_COMPILER_FENCE with an HPX version
- PR #2448¹⁹²⁶ - This fixes a possible race in the migration code
- PR #2445¹⁹²⁷ - **Fixing dataflow et.al. for futures or future-ranges wrapped into ref()**
- PR #2444¹⁹²⁸ - Fix segfaults
- PR #2443¹⁹²⁹ - Issue 2442
- Issue #2442¹⁹³⁰ - Mismatch between #if/#endif and namespace scope brackets in this_thread_executers.hpp
- Issue #2441¹⁹³¹ - undeclared identifier BOOST_COMPILER_FENCE
- PR #2440¹⁹³² - Knl build
- PR #2438¹⁹³³ - Datapar backend
- PR #2437¹⁹³⁴ - Adapt algorithm parameter sequence changes from C++17
- PR #2436¹⁹³⁵ - Adapt execution policy name changes from C++17
- Issue #2435¹⁹³⁶ - Trunk broken, undefined reference to hpx::thread::interrupt(hpx::thread::id, bool)

¹⁹¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2465>

¹⁹¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2464>

¹⁹¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2463>

¹⁹¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2462>

¹⁹¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2461>

¹⁹¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2460>

¹⁹²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2459>

¹⁹²¹ <https://github.com/STELLAR-GROUP/hpx/issues/2457>

¹⁹²² <https://github.com/STELLAR-GROUP/hpx/pull/2452>

¹⁹²³ <https://github.com/STELLAR-GROUP/hpx/pull/2451>

¹⁹²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2450>

¹⁹²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2449>

¹⁹²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2448>

¹⁹²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2445>

¹⁹²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2444>

¹⁹²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2443>

¹⁹³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2442>

¹⁹³¹ <https://github.com/STELLAR-GROUP/hpx/issues/2441>

¹⁹³² <https://github.com/STELLAR-GROUP/hpx/pull/2440>

¹⁹³³ <https://github.com/STELLAR-GROUP/hpx/pull/2438>

¹⁹³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2437>

¹⁹³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2436>

¹⁹³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2435>

- PR #2434¹⁹³⁷ - More fixes to resource manager
- PR #2433¹⁹³⁸ - Added versions of `hpx::get_ptr` taking client side representations
- PR #2432¹⁹³⁹ - Warning fixes
- PR #2431¹⁹⁴⁰ - Adding facility representing set of performance counters
- PR #2430¹⁹⁴¹ - Fix parallel_executor thread spawning
- PR #2429¹⁹⁴² - Fix attribute warning for gcc
- Issue #2427¹⁹⁴³ - Seg fault running octo-tiger with latest HPX commit
- Issue #2426¹⁹⁴⁴ - Bug in 9592f5c0bc29806fce0dbe73f35b6ca7e027edcb causes immediate crash in Octo-tiger
- PR #2425¹⁹⁴⁵ - Fix nvcc errors due to constexpr specifier
- Issue #2424¹⁹⁴⁶ - Async action on component present on `hpx::find_here` is executing synchronously
- PR #2423¹⁹⁴⁷ - Fix nvcc errors due to constexpr specifier
- PR #2422¹⁹⁴⁸ - Implementing `hpx::this_thread` thread data functions
- PR #2421¹⁹⁴⁹ - Adding benchmark for `wait_all`
- Issue #2420¹⁹⁵⁰ - Returning object of a component client from another component action fails
- PR #2419¹⁹⁵¹ - Infiniband parcelport
- Issue #2418¹⁹⁵² - gcc + nvcc fails to compile code that uses `partitioned_vector`
- PR #2417¹⁹⁵³ - Fixing context switching
- PR #2416¹⁹⁵⁴ - Adding fixes and workarounds to allow compilation with nvcc/msvc (VS2015up3)
- PR #2415¹⁹⁵⁵ - Fix errors coming from hpx compute examples
- PR #2414¹⁹⁵⁶ - Fixing msvc12
- PR #2413¹⁹⁵⁷ - Enable cuda/nvcc or cuda/clang when using `add_hpx_executable()`
- PR #2412¹⁹⁵⁸ - Fix issue in `HPX_SetupTarget.cmake` when cuda is used
- PR #2411¹⁹⁵⁹ - This fixes the core compilation issues with MSVC12

¹⁹³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2434>

¹⁹³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2433>

¹⁹³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2432>

¹⁹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2431>

¹⁹⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2430>

¹⁹⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2429>

¹⁹⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/2427>

¹⁹⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2426>

¹⁹⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2425>

¹⁹⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2424>

¹⁹⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2423>

¹⁹⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2422>

¹⁹⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2421>

¹⁹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2420>

¹⁹⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2419>

¹⁹⁵² <https://github.com/STELLAR-GROUP/hpx/issues/2418>

¹⁹⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2417>

¹⁹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2416>

¹⁹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2415>

¹⁹⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2414>

¹⁹⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2413>

¹⁹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2412>

¹⁹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2411>

- Issue #2410¹⁹⁶⁰ - undefined reference to opal_hwloc191_hwloc_.....
- PR #2409¹⁹⁶¹ - Fixing locking for channel and receive_buffer
- PR #2407¹⁹⁶² - Solving #2402 and #2403
- PR #2406¹⁹⁶³ - Improve guards
- PR #2405¹⁹⁶⁴ - Enable parallel::for_each for iterators returning proxy types
- PR #2404¹⁹⁶⁵ - Forward the explicitly given result_type in the hpx invoke
- Issue #2403¹⁹⁶⁶ - datapar_execution + zip iterator: lambda arguments aren't references
- Issue #2402¹⁹⁶⁷ - datapar algorithm instantiated with wrong type #2402
- PR #2401¹⁹⁶⁸ - Added support for imported libraries to HPX_Libraries.cmake
- PR #2400¹⁹⁶⁹ - Use CMake policy CMP0060
- Issue #2399¹⁹⁷⁰ - Error trying to push back vector of futures to vector
- PR #2398¹⁹⁷¹ - Allow config #defines to be written out to custom config/defines.hpp
- Issue #2397¹⁹⁷² - CMake generated config defines can cause tedious rebuilds category
- Issue #2396¹⁹⁷³ - BOOST_ROOT paths are not used at link time
- PR #2395¹⁹⁷⁴ - Fix target_link_libraries() issue when HPX Cuda is enabled
- Issue #2394¹⁹⁷⁵ - Template compilation error using HPX_WITH_DATAPAR_LIBFLATARRAY
- PR #2393¹⁹⁷⁶ - Fixing lock registration for recursive mutex
- PR #2392¹⁹⁷⁷ - Add keywords in target_link_libraries in hpx_setup_target
- PR #2391¹⁹⁷⁸ - Clang goroutines
- Issue #2390¹⁹⁷⁹ - Adapt execution policy name changes from C++17
- PR #2389¹⁹⁸⁰ - Chunk allocator and pool are not used and are obsolete
- PR #2388¹⁹⁸¹ - Adding functionalities to datapar needed by octotiger
- PR #2387¹⁹⁸² - Fixing race condition for early parcels

¹⁹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2410>

¹⁹⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/2409>

¹⁹⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2407>

¹⁹⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/2406>

¹⁹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2405>

¹⁹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2404>

¹⁹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2403>

¹⁹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2402>

¹⁹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2401>

¹⁹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2400>

¹⁹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2399>

¹⁹⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2398>

¹⁹⁷² <https://github.com/STELLAR-GROUP/hpx/issues/2397>

¹⁹⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2396>

¹⁹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2395>

¹⁹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2394>

¹⁹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2393>

¹⁹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2392>

¹⁹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2391>

¹⁹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2390>

¹⁹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2389>

¹⁹⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2388>

¹⁹⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2387>

- Issue #2386¹⁹⁸³ - Lock registration broken for recursive_mutex
- PR #2385¹⁹⁸⁴ - Datapar zip iterator
- PR #2384¹⁹⁸⁵ - Fixing race condition in for_loop_reduction
- PR #2383¹⁹⁸⁶ - Continuations
- PR #2382¹⁹⁸⁷ - add LibFlatArray-based backend for datapar
- PR #2381¹⁹⁸⁸ - remove unused typedef to get rid of compiler warnings
- PR #2380¹⁹⁸⁹ - Tau cleanup
- PR #2379¹⁹⁹⁰ - Can send immediate
- PR #2378¹⁹⁹¹ - Renaming copy_helper/copy_n_helper/move_helper/move_n_helper
- Issue #2376¹⁹⁹² - Boost trunk's spinlock initializer fails to compile
- PR #2375¹⁹⁹³ - Add support for minimal thread local data
- PR #2374¹⁹⁹⁴ - Adding API functions set_config_entry_callback
- PR #2373¹⁹⁹⁵ - Add a simple utility for debugging that gives suspended task backtraces
- PR #2372¹⁹⁹⁶ - Barrier Fixes
- Issue #2370¹⁹⁹⁷ - Can't wait on a wrapped future
- PR #2369¹⁹⁹⁸ - Fixing stable_partition
- PR #2367¹⁹⁹⁹ - Fixing find_prefixes for Windows platforms
- PR #2366²⁰⁰⁰ - Testing for experimental/optional only in C++14 mode
- PR #2364²⁰⁰¹ - Adding set_config_entry
- PR #2363²⁰⁰² - Fix papi
- PR #2362²⁰⁰³ - Adding missing macros for new non-direct actions
- PR #2361²⁰⁰⁴ - Improve cmake output to help debug compiler incompatibility check
- PR #2360²⁰⁰⁵ - Fixing race condition in condition_variable

¹⁹⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/2386>

¹⁹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2385>

¹⁹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2384>

¹⁹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2383>

¹⁹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2382>

¹⁹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2381>

¹⁹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2380>

¹⁹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2379>

¹⁹⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2378>

¹⁹⁹² <https://github.com/STELLAR-GROUP/hpx/issues/2376>

¹⁹⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2375>

¹⁹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2374>

¹⁹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2373>

¹⁹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2372>

¹⁹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2370>

¹⁹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2369>

¹⁹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2367>

²⁰⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2366>

²⁰⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2364>

²⁰⁰² <https://github.com/STELLAR-GROUP/hpx/pull/2363>

²⁰⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2362>

²⁰⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2361>

²⁰⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2360>

- PR #2359²⁰⁰⁶ - Fixing shutdown when parcels are still in flight
- Issue #2357²⁰⁰⁷ - failed to insert console_print_action into typename_to_id_t registry
- PR #2356²⁰⁰⁸ - Fixing return type of get_iterator_tuple
- PR #2355²⁰⁰⁹ - Fixing compilation against Boost 1.62
- PR #2354²⁰¹⁰ - Adding serialization for mask_type if CPU_COUNT > 64
- PR #2353²⁰¹¹ - Adding hooks to tie in APEX into the parcel layer
- Issue #2352²⁰¹² - Compile errors when using intel 17 beta (for KNL) on edison
- PR #2351²⁰¹³ - Fix function vtable get_function_address implementation
- Issue #2350²⁰¹⁴ - Build failure - master branch (4de09f5) with Intel Compiler v17
- PR #2349²⁰¹⁵ - Enabling zero-copy serialization support for std::vector<>
- PR #2348²⁰¹⁶ - Adding test to verify #2334 is fixed
- PR #2347²⁰¹⁷ - Bug fixes for hpx::compute and hpx::lcos::channel
- PR #2346²⁰¹⁸ - Removing cmake "find" files that are in the APEX cmake Modules
- PR #2345²⁰¹⁹ - Implemented parallel::stable_partition
- PR #2344²⁰²⁰ - Making hpx::lcos::channel usable with basename registration
- PR #2343²⁰²¹ - Fix a couple of examples that failed to compile after recent api changes
- Issue #2342²⁰²² - Enabling APEX causes link errors
- PR #2341²⁰²³ - Removing cmake "find" files that are in the APEX cmake Modules
- PR #2340²⁰²⁴ - Implemented all existing datapar algorithms using Boost.SIMD
- PR #2339²⁰²⁵ - Fixing 2338
- PR #2338²⁰²⁶ - Possible race in sliding semaphore
- PR #2337²⁰²⁷ - Adjust osu_latency test to measure window_size parcels in flight at once
- PR #2336²⁰²⁸ - Allowing remote direct actions to be executed without spawning a task

²⁰⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2359>

²⁰⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2357>

²⁰⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2356>

²⁰⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2355>

²⁰¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2354>

²⁰¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2353>

²⁰¹² <https://github.com/STELLAR-GROUP/hpx/issues/2352>

²⁰¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2351>

²⁰¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2350>

2015 <https://github.com/STELLAR-GROUP/hpx/pull/2349>2016 <https://github.com/STELLAR-GROUP/hpx/pull/2348>2017 <https://github.com/STELLAR-GROUP/hpx/pull/2347>2018 <https://github.com/STELLAR-GROUP/hpx/pull/2346>2019 <https://github.com/STELLAR-GROUP/hpx/pull/2345>2020 <https://github.com/STELLAR-GROUP/hpx/pull/2344>2021 <https://github.com/STELLAR-GROUP/hpx/pull/2343>2022 <https://github.com/STELLAR-GROUP/hpx/issues/2342>2023 <https://github.com/STELLAR-GROUP/hpx/pull/2341>2024 <https://github.com/STELLAR-GROUP/hpx/pull/2340>2025 <https://github.com/STELLAR-GROUP/hpx/pull/2339>2026 <https://github.com/STELLAR-GROUP/hpx/pull/2338>2027 <https://github.com/STELLAR-GROUP/hpx/pull/2337>2028 <https://github.com/STELLAR-GROUP/hpx/pull/2336>

- PR #2335²⁰²⁹ - Making sure multiple components are properly initialized from arguments
- Issue #2334²⁰³⁰ - Cannot construct component with large vector on a remote locality
- PR #2332²⁰³¹ - Fixing hpx::lcos::local::barrier
- PR #2331²⁰³² - Updating APEX support to include OTF2
- PR #2330²⁰³³ - Support for data-parallelism for parallel algorithms
- Issue #2329²⁰³⁴ - Coordinate settings in cmake
- PR #2328²⁰³⁵ - fix LibGeoDecomp builds with HPX + GCC 5.3.0 + CUDA 8RC
- PR #2326²⁰³⁶ - Making scan_partitioner work (for now)
- Issue #2323²⁰³⁷ - Constructing a vector of components only correctly initializes the first component
- PR #2322²⁰³⁸ - Fix problems that bubbled up after merging #2278
- PR #2321²⁰³⁹ - Scalable barrier
- PR #2320²⁰⁴⁰ - Std flag fixes
- Issue #2319²⁰⁴¹ - -std=c++14 and -std=c++1y with Intel can't build recent Boost builds due to insufficient C++14 support; don't enable these flags by default for Intel
- PR #2318²⁰⁴² - Improve handling of -hpx:bind=<bind-spec>
- PR #2317²⁰⁴³ - Making sure command line warnings are printed once only
- PR #2316²⁰⁴⁴ - Fixing command line handling for default bind mode
- PR #2315²⁰⁴⁵ - Set id_retrieved if set_id is present
- Issue #2314²⁰⁴⁶ - Warning for requested/allocated thread discrepancy is printed twice
- Issue #2313²⁰⁴⁷ - --hpx:print-bind doesn't work with -hpx:pu-step
- Issue #2312²⁰⁴⁸ - --hpx:bind range specifier restrictions are overly restrictive
- Issue #2311²⁰⁴⁹ - hpx_0.9.99 out of project build fails
- PR #2310²⁰⁵⁰ - Simplify function registration
- PR #2309²⁰⁵¹ - Spelling and grammar revisions in documentation (and some code)

²⁰²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2335>

²⁰³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2334>

²⁰³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2332>

²⁰³² <https://github.com/STELLAR-GROUP/hpx/pull/2331>

²⁰³³ <https://github.com/STELLAR-GROUP/hpx/pull/2330>

²⁰³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2329>

²⁰³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2328>

²⁰³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2326>

²⁰³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2323>

²⁰³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2322>

²⁰³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2321>

²⁰⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2320>

²⁰⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/2319>

²⁰⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2318>

²⁰⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/2317>

²⁰⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2316>

²⁰⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2315>

²⁰⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2314>

²⁰⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2313>

²⁰⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2312>

²⁰⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2311>

²⁰⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2310>

²⁰⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2309>

- PR #2306²⁰⁵² - Correct minor typo in the documentation
- PR #2305²⁰⁵³ - Cleaning up and fixing parcel coalescing
- PR #2304²⁰⁵⁴ - Inspect checks for stream related includes
- PR #2303²⁰⁵⁵ - Add functionality allowing to enumerate threads of given state
- PR #2301²⁰⁵⁶ - Algorithm overloads fix for VS2013
- PR #2300²⁰⁵⁷ - Use <cstdint>, add inspect checks
- PR #2299²⁰⁵⁸ - Replace boost::[c]ref with std::[c]ref, add inspect checks
- PR #2297²⁰⁵⁹ - Fixing compilation with no hw_loc
- PR #2296²⁰⁶⁰ - Hpx compute
- PR #2295²⁰⁶¹ - Making sure for_loop(execution::par, 0, N, ...) is actually executed in parallel
- PR #2294²⁰⁶² - Throwing exceptions if the runtime is not up and running
- PR #2293²⁰⁶³ - Removing unused parcel port code
- PR #2292²⁰⁶⁴ - Refactor function vtables
- PR #2291²⁰⁶⁵ - Fixing 2286
- PR #2290²⁰⁶⁶ - Simplify algorithm overloads
- PR #2289²⁰⁶⁷ - Adding performance counters reporting parcel related data on a per-action basis
- Issue #2288²⁰⁶⁸ - Remove dormant parcelports
- Issue #2286²⁰⁶⁹ - adjustments to parcel handling to support parcelports that do not need a connection cache
- PR #2285²⁰⁷⁰ - add CMake option to disable package export
- PR #2283²⁰⁷¹ - Add more inspect checks for use of deprecated components
- Issue #2282²⁰⁷² - Arithmetic exception in executor static chunker
- Issue #2281²⁰⁷³ - For loop doesn't parallelize
- PR #2280²⁰⁷⁴ - Fixing 2277: build failure with PAPI

²⁰⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2306>

²⁰⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2305>

²⁰⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2304>

²⁰⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2303>

²⁰⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2301>

²⁰⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2300>

²⁰⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2299>

²⁰⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2297>

²⁰⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2296>

²⁰⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/2295>

²⁰⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2294>

²⁰⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/2293>

²⁰⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2292>

²⁰⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2291>

²⁰⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2290>

²⁰⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2289>

²⁰⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2288>

²⁰⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2286>

²⁰⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2285>

²⁰⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2283>

²⁰⁷² <https://github.com/STELLAR-GROUP/hpx/issues/2282>

²⁰⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/2281>

²⁰⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2280>

- PR #2279²⁰⁷⁵ - Child vs parent stealing
- Issue #2277²⁰⁷⁶ - master branch build failure (53c5b4f) with papi
- PR #2276²⁰⁷⁷ - Compile time launch policies
- PR #2275²⁰⁷⁸ - Replace boost::chrono with std::chrono in interfaces
- PR #2274²⁰⁷⁹ - Replace most uses of Boost.Assign with initializer list
- PR #2273²⁰⁸⁰ - Fixed typos
- PR #2272²⁰⁸¹ - Inspect checks
- PR #2270²⁰⁸² - Adding test verifying -Ihpx.os_threads=all
- PR #2269²⁰⁸³ - Added inspect check for now obsolete boost type traits
- PR #2268²⁰⁸⁴ - Moving more code into source files
- Issue #2267²⁰⁸⁵ - Add inspect support to deprecate Boost.TypeTraits
- PR #2265²⁰⁸⁶ - Adding channel LCO
- PR #2264²⁰⁸⁷ - Make support for std::ref mandatory
- PR #2263²⁰⁸⁸ - Constrain tuple_member forwarding constructor
- Issue #2262²⁰⁸⁹ - Test hpx.os_threads=all
- Issue #2261²⁰⁹⁰ - OS X: Error: no matching constructor for initialization of 'hpx::lcos::local::condition_variable_any'
- Issue #2260²⁰⁹¹ - Make support for std::ref mandatory
- PR #2259²⁰⁹² - Remove most of Boost.MPL, Boost.EnableIf and Boost.TypeTraits
- PR #2258²⁰⁹³ - Fixing #2256
- PR #2257²⁰⁹⁴ - Fixing launch process
- Issue #2256²⁰⁹⁵ - Actions are not registered if not invoked
- PR #2255²⁰⁹⁶ - Coalescing histogram
- PR #2254²⁰⁹⁷ - Silence explicit initialization in copy-constructor warnings

²⁰⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2279>

²⁰⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2277>

²⁰⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2276>

²⁰⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2275>

²⁰⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2274>

²⁰⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2273>

²⁰⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/2272>

²⁰⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2270>

²⁰⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2269>

²⁰⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2268>

²⁰⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2267>

²⁰⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2265>

²⁰⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2264>

²⁰⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2263>

²⁰⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2262>

²⁰⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2261>

²⁰⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/2260>

²⁰⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2259>

²⁰⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2258>

²⁰⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2257>

²⁰⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2256>

²⁰⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2255>

²⁰⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2254>

- PR #2253²⁰⁹⁸ - Drop support for GCC 4.6 and 4.7
- PR #2252²⁰⁹⁹ - Prepare V1.0
- PR #2251²¹⁰⁰ - Convert to 0.9.99
- PR #2249²¹⁰¹ - Adding iterator_facade and iterator_adaptor
- Issue #2248²¹⁰² - Need a feature to yield to a new task immediately
- PR #2246²¹⁰³ - Adding split_future
- PR #2245²¹⁰⁴ - Add an example for handing over a component instance to a dynamically launched locality
- Issue #2243²¹⁰⁵ - Add example demonstrating AGAS symbolic name registration
- Issue #2242²¹⁰⁶ - pkgconfig test broken on CentOS 7 / Boost 1.61
- Issue #2241²¹⁰⁷ - Compilation error for partitioned vector in hpx_compute branch
- PR #2240²¹⁰⁸ - Fixing termination detection on one locality
- Issue #2239²¹⁰⁹ - Create a new facility lcos::split_all
- Issue #2236²¹¹⁰ - hpx::cout vs. std::cout
- PR #2232²¹¹¹ - Implement local-only primary namespace service
- Issue #2147²¹¹² - would like to know how much data is being routed by particular actions
- Issue #2109²¹¹³ - Warning while compiling hpx
- Issue #1973²¹¹⁴ - Setting INTERFACE_COMPILE_OPTIONS for hpx_init in CMake taints Fortran_FLAGS
- Issue #1864²¹¹⁵ - run_guarded using bound function ignores reference
- Issue #1754²¹¹⁶ - Running with TCP parcelport causes immediate crash or freeze
- Issue #1655²¹¹⁷ - Enable zip_iterator to be used with Boost traversal iterator categories
- Issue #1591²¹¹⁸ - Optimize AGAS for shared memory only operation
- Issue #1401²¹¹⁹ - Need an efficient infiniband parcelport
- Issue #1125²¹²⁰ - Fix the IPC parcelport

²⁰⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2253>

²⁰⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2252>

²¹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2251>

²¹⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2249>

²¹⁰² <https://github.com/STELLAR-GROUP/hpx/issues/2248>

²¹⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2246>

²¹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2245>

²¹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2243>

²¹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2242>

²¹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2241>

²¹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2240>

²¹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2239>

²¹¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2236>

²¹¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2232>

²¹¹² <https://github.com/STELLAR-GROUP/hpx/issues/2147>

²¹¹³ <https://github.com/STELLAR-GROUP/hpx/issues/2109>

²¹¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1973>

²¹¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1864>

²¹¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1754>

²¹¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1655>

²¹¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1591>

²¹¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1401>

²¹²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1125>

- Issue #839²¹²¹ - Refactor ibverbs and shmem parcelport
- Issue #702²¹²² - Add instrumentation of parcel layer
- Issue #668²¹²³ - Implement ispc task interface
- Issue #533²¹²⁴ - Thread queue/deque internal parameters should be runtime configurable
- Issue #475²¹²⁵ - Create a means of combining performance counters into querysets

2.11.7 HPX V0.9.99 (Jul 15, 2016)

General changes

As the version number of this release hints, we consider this release to be a preview for the upcoming *HPX* V1.0. All of the functionalities we set out to implement for V1.0 are in place; all of the features we wanted to have exposed are ready. We are very happy with the stability and performance of *HPX* and we would like to present this release to the community in order for us to gather broad feedback before releasing V1.0. We still expect for some minor details to change, but on the whole this release represents what we would like to have in a V1.0.

Overall, since the last release we have had almost 1600 commits while closing almost 400 tickets. These numbers reflect the incredible development activity we have seen over the last couple of months. We would like to express a big ‘Thank you!’ to all contributors and those who helped to make this release happen.

The most notable addition in terms of new functionality available with this release is the full implementation of object migration (i.e. the ability to transparently move *HPX* components to a different compute node). Additionally, this release of *HPX* cleans up many minor issues and some API inconsistencies.

Here are some of the main highlights and changes for this release (in no particular order):

- We have fixed a couple of issues in AGAS and the parcel layer which have caused hangs, segmentation faults at exit, and a slowdown of applications over time. Fixing those has significantly increased the overall stability and performance of distributed runs.
- We have started to add parallel algorithm overloads based on the C++ Extensions for Ranges (N4560²¹²⁶) proposal. This also includes the addition of projections to the existing algorithms. Please see Issue #1668²¹²⁷ for a list of algorithms which have been adapted to N4560²¹²⁸.
- We have implemented index-based parallel for-loops based on a corresponding standardization proposal (P0075R1²¹²⁹). Please see Issue #2016²¹³⁰ for a list of available algorithms.
- We have added implementations for more parallel algorithms as proposed for the upcoming C++ 17 Standard. See Issue #1141²¹³¹ for an overview of which algorithms are available by now.
- We have started to implement a new prototypical functionality with *HPX.Compute* which uniformly exposes some of the higher level APIs to heterogeneous architectures (currently CUDA). This functionality is an early preview and should not be considered stable. It may change considerably in the future.
- We have pervasively added (optional) executor arguments to all API functions which schedule new work. Executors are now used throughout the code base as the main means of executing tasks.

²¹²¹ <https://github.com/STELLAR-GROUP/hpx/issues/839>

²¹²² <https://github.com/STELLAR-GROUP/hpx/issues/702>

²¹²³ <https://github.com/STELLAR-GROUP/hpx/issues/668>

²¹²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/533>

²¹²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/475>

²¹²⁶ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4560.pdf>

²¹²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1668>

²¹²⁸ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4560.pdf>

²¹²⁹ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2016/p0075r1.pdf>

²¹³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2016>

²¹³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1141>

- Added `hpx::make_future<R>(future<T> &&)` allowing to convert a future of any type `T` into a future of any other type `R`, either based on default conversion rules of the embedded types or using a given explicit conversion function.
- We finally finished the implementation of transparent migration of components to another locality. It is now possible to trigger a migration operation without ‘stopping the world’ for the object to migrate. *HPX* will make sure that no work is being performed on an object before it is migrated and that all subsequently scheduled work for the migrated object will be transparently forwarded to the new locality. Please note that the global id of the migrated object does not change, thus the application will not have to be changed in any way to support this new functionality. Please note that this feature is currently considered experimental. See [Issue #559²¹³²](#) and [PR #1966²¹³³](#) for more details.
- The `hpx::dataflow` facility is now usable with actions. Similarly to `hpx::async`, actions can be specified as an explicit template argument (`hpx::dataflow<Action>(target, ...)`) or as the first argument (`hpx::dataflow(Action(), target, ...)`). We have also enabled the use of distribution policies as the target for dataflow invocations. Please see [Issue #1265²¹³⁴](#) and [PR #1912²¹³⁵](#) for more information.
- Adding overloads of `gather_here` and `gather_there` to accept the plain values of the data to gather (in addition to the existing overloads expecting futures).
- We have cleaned up and refactored large parts of the code base. This helped reducing compile and link times of *HPX* itself and also of applications depending on it. We have further decreased the dependency of *HPX* on the Boost libraries by replacing part of those with facilities available from the standard libraries.
- Wherever possible we have removed dependencies of our API on Boost by replacing those with the equivalent facility from the C++11 standard library.
- We have added new performance counters for parcel coalescing, file-IO, the AGAS cache, and overall scheduler time. Resetting performance counters has been overhauled and fixed.
- We have introduced a generic client type `hpx::components::client<>` and added support for using it with `hpx::async`. This removes the necessity to implement specific client types for every component type without losing type safety. This deemphasizes the need for using the low level `hpx::id_type` for referencing (possibly remote) component instances. The plan is to deprecate the direct use of `hpx::id_type` in user code in the future.
- We have added a special iterator which supports automatic prefetching of one or more arrays for speeding up loop-like code (see `hpx::parallel::util::make_prefetcher_context()`).
- We have extended the interfaces exposed from executors (as proposed by [N4406²¹³⁶](#)) to accept an arbitrary number of arguments.

Breaking changes

- In order to move the dataflow facility to namespace `hpx` we added a definition of `hpx::dataflow` which might create ambiguities in existing codes. The previous definition of this facility (`hpx::lcos::local::dataflow`) has been deprecated and is available only if the constant `-DHGX_WITH_LOCAL_DATAFLOW_COMPATIBILITY=On` to [CMake²¹³⁷](#) is defined at configuration time. Please explicitly qualify all uses of the dataflow facility if you enable this compatibility setting and encounter ambiguities.

²¹³² <https://github.com/STELLAR-GROUP/hpx/issues/559>

²¹³³ <https://github.com/STELLAR-GROUP/hpx/pull/1966>

²¹³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1265>

²¹³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1912>

²¹³⁶ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4406.pdf>

²¹³⁷ <https://www.cmake.org>

- The adaptation of the C++ Extensions for Ranges (N4560²¹³⁸) proposal imposes some breaking changes related to the return types of some of the parallel algorithms. Please see Issue #1668²¹³⁹ for a list of algorithms which have already been adapted.
- The facility `hpx::lcos::make_future_void()` has been replaced by `hpx::make_future<void>()`.
- We have removed support for Intel V13 and gcc 4.4.x.
- We have removed (default) support for the generic `hpx::parallel::execution_policy` because it was removed from the Parallelism TS (`_cpp11_n4104_`) while it was being added to the upcoming C++17 Standard. This facility can be still enabled at configure time by specifying `-DHPX_WITH_GENERIC_EXECUTION_POLICY=On` to CMake.
- Uses of `boost::shared_ptr` and related facilities have been replaced with `std::shared_ptr` and friends. Uses of `boost::unique_lock`, `boost::lock_guard` etc. have also been replaced by the equivalent (and equally named) tools available from the C++11 standard library.
- Facilities that used to expect an explicit `boost::unique_lock` now take an `std::unique_lock`. Additionally, `condition_variable` no longer aliases `condition_variable_any`; its interface now only works with `std::unique_lock<local::mutex>`.
- Uses of `boost::function`, `boost::bind`, `boost::tuple` have been replaced by the corresponding facilities in *HPX* (`hpx::util::function`, `hpx::util::bind`, and `hpx::util::tuple`, respectively).

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- PR #2250²¹⁴⁰ - change default chunker of parallel executor to static one
- PR #2247²¹⁴¹ - HPX on ppc64le
- PR #2244²¹⁴² - Fixing MSVC problems
- PR #2238²¹⁴³ - Fixing small typos
- PR #2237²¹⁴⁴ - Fixing small typos
- PR #2234²¹⁴⁵ - Fix broken add test macro when extra args are passed in
- PR #2231²¹⁴⁶ - Fixing possible race during future awaiting in serialization
- PR #2230²¹⁴⁷ - Fix stream nvcc
- PR #2229²¹⁴⁸ - Fixed run_as_hpx_thread
- PR #2228²¹⁴⁹ - On prefetching_test branch : adding prefetching_iterator and related tests used for prefetching containers within lambda functions

²¹³⁸ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4560.pdf>

²¹³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1668>

²¹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2250>

²¹⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2247>

²¹⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2244>

²¹⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/2238>

²¹⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2237>

²¹⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2234>

²¹⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2231>

²¹⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2230>

²¹⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2229>

²¹⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2228>

- PR #2227²¹⁵⁰ - Support for HPXCL's opencl::event
- PR #2226²¹⁵¹ - Preparing for release of V0.9.99
- PR #2225²¹⁵² - fix issue when compiling components with hpxcxx
- PR #2224²¹⁵³ - Compute alloc fix
- PR #2223²¹⁵⁴ - Simplify promise
- PR #2222²¹⁵⁵ - Replace last uses of boost::function by util::function_nonsr
- PR #2221²¹⁵⁶ - Fix config tests
- PR #2220²¹⁵⁷ - Fixing gcc 4.6 compilation issues
- PR #2219²¹⁵⁸ - nullptr support for [unique_]function
- PR #2218²¹⁵⁹ - Introducing clang tidy
- PR #2216²¹⁶⁰ - Replace NULL with nullptr
- Issue #2214²¹⁶¹ - Let inspect flag use of NULL, suggest nullptr instead
- PR #2213²¹⁶² - Require support for nullptr
- PR #2212²¹⁶³ - Properly find jemalloc through pkg-config
- PR #2211²¹⁶⁴ - Disable a couple of warnings reported by Intel on Windows
- PR #2210²¹⁶⁵ - Fixed host::block_allocator::bulk_construct
- PR #2209²¹⁶⁶ - Started to clean up new sort algorithms, made things compile for sort_by_key
- PR #2208²¹⁶⁷ - A couple of fixes that were exposed by a new sort algorithm
- PR #2207²¹⁶⁸ - Adding missing includes in /hpx/include/serialization.hpp
- PR #2206²¹⁶⁹ - Call package_action::get_future before package_action::apply
- PR #2205²¹⁷⁰ - The indirect_packaged_task::operator() needs to be run on a HPX thread
- PR #2204²¹⁷¹ - Variadic executor parameters
- PR #2203²¹⁷² - Delay-initialize members of partitioned iterator

²¹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2227>

²¹⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2226>

²¹⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2225>

²¹⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2224>

²¹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2223>

²¹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2222>

2156 <https://github.com/STELLAR-GROUP/hpx/pull/2221>2157 <https://github.com/STELLAR-GROUP/hpx/pull/2220>2158 <https://github.com/STELLAR-GROUP/hpx/pull/2219>2159 <https://github.com/STELLAR-GROUP/hpx/pull/2218>2160 <https://github.com/STELLAR-GROUP/hpx/pull/2216>2161 <https://github.com/STELLAR-GROUP/hpx/issues/2214>2162 <https://github.com/STELLAR-GROUP/hpx/pull/2213>2163 <https://github.com/STELLAR-GROUP/hpx/pull/2212>2164 <https://github.com/STELLAR-GROUP/hpx/pull/2211>2165 <https://github.com/STELLAR-GROUP/hpx/pull/2210>2166 <https://github.com/STELLAR-GROUP/hpx/pull/2209>2167 <https://github.com/STELLAR-GROUP/hpx/pull/2208>2168 <https://github.com/STELLAR-GROUP/hpx/pull/2207>2169 <https://github.com/STELLAR-GROUP/hpx/pull/2206>2170 <https://github.com/STELLAR-GROUP/hpx/pull/2205>2171 <https://github.com/STELLAR-GROUP/hpx/pull/2204>2172 <https://github.com/STELLAR-GROUP/hpx/pull/2203>

- PR #2202²¹⁷³ - Added segmented fill for hpx::vector
- Issue #2201²¹⁷⁴ - Null Thread id encountered on partitioned_vector
- PR #2200²¹⁷⁵ - Fix hangs
- PR #2199²¹⁷⁶ - Deprecating hpx/traits.hpp
- PR #2198²¹⁷⁷ - Making explicit inclusion of external libraries into build
- PR #2197²¹⁷⁸ - Fix typo in QT CMakeLists
- PR #2196²¹⁷⁹ - Fixing a gcc warning about attributes being ignored
- PR #2194²¹⁸⁰ - Fixing partitioned_vector_spmd_foreach example
- Issue #2193²¹⁸¹ - partitioned_vector_spmd_foreach seg faults
- PR #2192²¹⁸² - Support Boost.Thread v4
- PR #2191²¹⁸³ - HPX.Compute prototype
- PR #2190²¹⁸⁴ - Spawning operation on new thread if remaining stack space becomes too small
- PR #2189²¹⁸⁵ - Adding callback taking index and future to when_each
- PR #2188²¹⁸⁶ - Adding new example demonstrating receive_buffer
- PR #2187²¹⁸⁷ - Mask 128-bit ints if CUDA is being used
- PR #2186²¹⁸⁸ - Make startup & shutdown functions unique_function
- PR #2185²¹⁸⁹ - Fixing logging output not to cause hang on shutdown
- PR #2184²¹⁹⁰ - Allowing component clients as action return types
- Issue #2183²¹⁹¹ - Enabling logging output causes hang on shutdown
- Issue #2182²¹⁹² - 1d_stencil seg fault
- Issue #2181²¹⁹³ - Setting small stack size does not change default
- PR #2180²¹⁹⁴ - Changing default bind mode to balanced
- PR #2179²¹⁹⁵ - adding prefetching_iterator and related tests used for prefetching containers within lambda functions

²¹⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/2202>

²¹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2201>

²¹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2200>

²¹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2199>

²¹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2198>

²¹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2197>

²¹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2196>

²¹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2194>

²¹⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/2193>

²¹⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2192>

²¹⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2191>

²¹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2190>

²¹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2189>

²¹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2188>

²¹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2187>

²¹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2186>

²¹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2185>

²¹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2184>

²¹⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/2183>

²¹⁹² <https://github.com/STELLAR-GROUP/hpx/issues/2182>

²¹⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/2181>

²¹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2180>

²¹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2179>

- PR #2177²¹⁹⁶ - Fixing 2176
- Issue #2176²¹⁹⁷ - Launch process test fails on OSX
- PR #2175²¹⁹⁸ - Fix unbalanced config/warnings includes, add some new ones
- PR #2174²¹⁹⁹ - Fix test categorization : regression not unit
- Issue #2172²²⁰⁰ - Different performance results
- Issue #2171²²⁰¹ - “negative entry in reference count table” running octotiger on 32 nodes on queenbee
- Issue #2170²²⁰² - Error while compiling on Mac + boost 1.60
- PR #2168²²⁰³ - Fixing problems with is_bitwise_serializable
- Issue #2167²²⁰⁴ - startup & shutdown function should accept unique_function
- Issue #2166²²⁰⁵ - Simple receive_buffer example
- PR #2165²²⁰⁶ - Fix wait all
- PR #2164²²⁰⁷ - Fix wait all
- PR #2163²²⁰⁸ - Fix some typos in config tests
- PR #2162²²⁰⁹ - Improve #includes
- PR #2160²²¹⁰ - Add inspect check for missing #include <list>
- PR #2159²²¹¹ - Add missing finalize call to stop test hanging
- PR #2158²²¹² - Algo fixes
- PR #2157²²¹³ - Stack check
- Issue #2156²²¹⁴ - OSX reports stack space incorrectly (generic context coroutines)
- Issue #2155²²¹⁵ - Race condition suspected in runtime
- PR #2154²²¹⁶ - Replace boost::detail::atomic_count with the new util::atomic_count
- PR #2153²²¹⁷ - Fix stack overflow on OSX
- PR #2152²²¹⁸ - Define is_bitwise_serializable as is_trivially_copyable when available

²¹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2177>

²¹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2176>

²¹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2175>

²¹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2174>

²²⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2172>

²²⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/2171>

²²⁰² <https://github.com/STELLAR-GROUP/hpx/issues/2170>

²²⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2168>

²²⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2167>

²²⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2166>

²²⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2165>

²²⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2164>

²²⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2163>

²²⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2162>

²²¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2160>

²²¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2159>

²²¹² <https://github.com/STELLAR-GROUP/hpx/pull/2158>

²²¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2157>

²²¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2156>

²²¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2155>

²²¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2154>

²²¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2153>

²²¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2152>

- PR #2151²²¹⁹ - Adding missing <cstring> for std::mem* functions
- Issue #2150²²²⁰ - Unable to use component clients as action return types
- PR #2149²²²¹ - std::memmove copies bytes, use bytes*sizeof(type) when copying larger types
- PR #2146²²²² - Adding customization point for parallel copy/move
- PR #2145²²²³ - Applying changes to address warnings issued by latest version of PVS Studio
- Issue #2148²²²⁴ - hpx::parallel::copy is broken after trivially copyable changes
- PR #2144²²²⁵ - Some minor tweaks to compute prototype
- PR #2143²²²⁶ - Added Boost version support information over OSX platform
- PR #2142²²²⁷ - Fixing memory leak in example
- PR #2141²²²⁸ - Add missing specializations in execution policies
- PR #2139²²²⁹ - This PR fixes a few problems reported by Clang's Undefined Behavior sanitizer
- PR #2138²²³⁰ - Revert "Adding fedora docs"
- PR #2136²²³¹ - Removed double semicolon
- PR #2135²²³² - Add deprecated #include check for hpx_fwd.hpp
- PR #2134²²³³ - Resolved memory leak in stencil_8
- PR #2133²²³⁴ - Replace uses of boost pointer containers
- PR #2132²²³⁵ - Removing unused typedef
- PR #2131²²³⁶ - Add several include checks for std facilities
- PR #2130²²³⁷ - Fixing parcel compression, adding test
- PR #2129²²³⁸ - Fix invalid attribute warnings
- Issue #2128²²³⁹ - hpx::init seems to segfault
- PR #2127²²⁴⁰ - Making executor_traits N-nary
- PR #2126²²⁴¹ - GCC 4.6 fails to deduce the correct type in lambda

²²¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2151>

²²²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2150>

²²²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2149>

²²²² <https://github.com/STELLAR-GROUP/hpx/pull/2146>

²²²³ <https://github.com/STELLAR-GROUP/hpx/pull/2145>

²²²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2148>

²²²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2144>

²²²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2143>

²²²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2142>

²²²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2141>

²²²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2139>

²²³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2138>

²²³¹ <https://github.com/STELLAR-GROUP/hpx/pull/2136>

²²³² <https://github.com/STELLAR-GROUP/hpx/pull/2135>

²²³³ <https://github.com/STELLAR-GROUP/hpx/pull/2134>

²²³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2133>

²²³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2132>

²²³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2131>

²²³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2130>

²²³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2129>

²²³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2128>

²²⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2127>

²²⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2126>

- PR #2125²²⁴² - Making parcel coalescing test actually test something
- Issue #2124²²⁴³ - Make a testcase for parcel compression
- Issue #2123²²⁴⁴ - hpx/hpx/runtime/applier_fwd.hpp - Multiple defined types
- Issue #2122²²⁴⁵ - Exception in primary_namespace::resolve_free_list
- Issue #2121²²⁴⁶ - Possible memory leak in 1d_stencil_8
- PR #2120²²⁴⁷ - Fixing 2119
- Issue #2119²²⁴⁸ - reduce_by_key compilation problems
- Issue #2118²²⁴⁹ - Premature unwrapping of boost::ref'ed arguments
- PR #2117²²⁵⁰ - Added missing initializer on last constructor for thread_description
- PR #2116²²⁵¹ - Use a lightweight bind implementation when no placeholders are given
- PR #2115²²⁵² - Replace boost::shared_ptr with std::shared_ptr
- PR #2114²²⁵³ - Adding hook functions for executor_parameter_traits supporting timers
- Issue #2113²²⁵⁴ - Compilation error with gcc version 4.9.3 (MacPorts gcc49 4.9.3_0)
- PR #2112²²⁵⁵ - Replace uses of safe_bool with explicit operator bool
- Issue #2111²²⁵⁶ - Compilation error on QT example
- Issue #2110²²⁵⁷ - Compilation error when passing non-future argument to unwrapped continuation in dataflow
- Issue #2109²²⁵⁸ - Warning while compiling hpx
- Issue #2109²²⁵⁹ - Stack trace of last bug causing issues with octotiger
- Issue #2108²²⁶⁰ - Stack trace of last bug causing issues with octotiger
- PR #2107²²⁶¹ - Making sure that a missing parcel_coalescing module does not cause startup exceptions
- PR #2106²²⁶² - Stop using hpx_fwd.hpp
- Issue #2105²²⁶³ - coalescing plugin handler is not optional any more
- Issue #2104²²⁶⁴ - Make executor_traits N-nary

²²⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2125>

²²⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/2124>

²²⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2123>

²²⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2122>

²²⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2121>

²²⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2120>

²²⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2119>

²²⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2118>

²²⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2117>

²²⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2116>

²²⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2115>

²²⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2114>

²²⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2113>

²²⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2112>

²²⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2111>

²²⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2110>

²²⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2109>

²²⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2109>

²²⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2108>

²²⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/2107>

²²⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2106>

²²⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/2105>

²²⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2104>

- Issue #2103²²⁶⁵ - Build error with octotiger and hpx commit e657426d
- PR #2102²²⁶⁶ - Combining thread data storage
- PR #2101²²⁶⁷ - Added repartition version of 1d stencil that uses any performance counter
- PR #2100²²⁶⁸ - Drop obsolete TR1 result_of protocol
- PR #2099²²⁶⁹ - Replace uses of boost::bind with util::bind
- PR #2098²²⁷⁰ - Deprecated inspect checks
- PR #2097²²⁷¹ - Reduce by key, extends #1141
- PR #2096²²⁷² - Moving local cache from external to hpx/util
- PR #2095²²⁷³ - Bump minimum required Boost to 1.50.0
- PR #2094²²⁷⁴ - Add include checks for several Boost utilities
- Issue #2093²²⁷⁵ - /.../local_cache.hpp(89): error #303: explicit type is missing (“int” assumed)
- PR #2091²²⁷⁶ - Fix for Raspberry pi build
- PR #2090²²⁷⁷ - Fix storage size for util::function<>
- PR #2089²²⁷⁸ - Fix #2088
- Issue #2088²²⁷⁹ - More verbose output from cmake configuration
- PR #2087²²⁸⁰ - Making sure init_globally always executes hpx_main
- Issue #2086²²⁸¹ - Race condition with recent HPX
- PR #2085²²⁸² - Adding #include checker
- PR #2084²²⁸³ - Replace boost lock types with standard library ones
- PR #2083²²⁸⁴ - Simplify packaged task
- PR #2082²²⁸⁵ - Updating APEX version for testing
- PR #2081²²⁸⁶ - Cleanup exception headers
- PR #2080²²⁸⁷ - Make call_once variadic

²²⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2103>

²²⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2102>

²²⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2101>

²²⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2100>

²²⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2099>

²²⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2098>

²²⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/2097>

²²⁷² <https://github.com/STELLAR-GROUP/hpx/pull/2096>

²²⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/2095>

²²⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2094>

²²⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2093>

²²⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2091>

²²⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2090>

²²⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2089>

²²⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2088>

²²⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2087>

²²⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/2086>

²²⁸² <https://github.com/STELLAR-GROUP/hpx/pull/2085>

²²⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/2084>

²²⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2083>

²²⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2082>

²²⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2081>

²²⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2080>

- Issue #2079²²⁸⁸ - With GNU C++, line 85 of hpx/config/version.hpp causes link failure when linking application
- Issue #2078²²⁸⁹ - Simple test fails with _GLIBCXX_DEBUG defined
- PR #2077²²⁹⁰ - Instantiate board in nqueen client
- PR #2076²²⁹¹ - Moving coalescing registration to TUs
- PR #2075²²⁹² - Fixed some documentation typos
- PR #2074²²⁹³ - Adding flush-mode to message handler flush
- PR #2073²²⁹⁴ - Fixing performance regression introduced lately
- PR #2072²²⁹⁵ - Refactor local::condition_variable
- PR #2071²²⁹⁶ - Timer based on boost::asio::deadline_timer
- PR #2070²²⁹⁷ - Refactor tuple based functionality
- PR #2069²²⁹⁸ - Fixed typos
- Issue #2068²²⁹⁹ - Seg fault with octotiger
- PR #2067²³⁰⁰ - Algorithm cleanup
- PR #2066²³⁰¹ - Split credit fixes
- PR #2065²³⁰² - Rename HPX_MOVABLE_BUT_NOT_COPYABLE to HPX_MOVABLE_ONLY
- PR #2064²³⁰³ - Fixed some typos in docs
- PR #2063²³⁰⁴ - Adding example demonstrating template components
- Issue #2062²³⁰⁵ - Support component templates
- PR #2061²³⁰⁶ - Replace some uses of lexical_cast<string> with C++11 std::to_string
- PR #2060²³⁰⁷ - Replace uses of boost::noncopyable with HPX_NON_COPYABLE
- PR #2059²³⁰⁸ - Adding missing for_loop algorithms
- PR #2058²³⁰⁹ - Move several definitions to more appropriate headers
- PR #2057²³¹⁰ - Simplify assert_owns_lock and ignore_while_checking

²²⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/2079>

²²⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2078>

²²⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2077>

²²⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2076>

²²⁹² <https://github.com/STELLAR-GROUP/hpx/pull/2075>

²²⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/2074>

²²⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2073>

²²⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2072>

²²⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2071>

²²⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2070>

²²⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2069>

²²⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2068>

²³⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2067>

²³⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/2066>

²³⁰² <https://github.com/STELLAR-GROUP/hpx/pull/2065>

²³⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/2064>

²³⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2063>

²³⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2062>

²³⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2061>

²³⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2060>

²³⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2059>

²³⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2058>

²³¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2057>

- PR #2056²³¹¹ - Replacing std::result_of with util::result_of
- PR #2055²³¹² - Fix process launching/connecting back
- PR #2054²³¹³ - Add a forwarding coroutine header
- PR #2053²³¹⁴ - Replace uses of boost::unordered_map with std::unordered_map
- PR #2052²³¹⁵ - Rewrite tuple unwrap
- PR #2050²³¹⁶ - Replace uses of BOOST_SCOPED_ENUM with C++11 scoped enums
- PR #2049²³¹⁷ - Attempt to narrow down split_credit problem
- PR #2048²³¹⁸ - Fixing gcc startup hangs
- PR #2047²³¹⁹ - Fixing when_xxx and wait_xxx for MSVC12
- PR #2046²³²⁰ - adding persistent_auto_chunk_size and related tests for for_each
- PR #2045²³²¹ - Fixing HPX_HAVE_THREAD_BACKTRACE_DEPTH build time configuration
- PR #2044²³²² - Adding missing service executor types
- PR #2043²³²³ - Removing ambiguous definitions for is_future_range and future_range_traits
- PR #2042²³²⁴ - Clarify that HPX builds can use (much) more than 2GB per process
- PR #2041²³²⁵ - Changing future_iterator_traits to support pointers
- Issue #2040²³²⁶ - Improve documentation memory usage warning?
- PR #2039²³²⁷ - Coroutine cleanup
- PR #2038²³²⁸ - Fix cmake policy CMP0042 warning MACOSX_RPATH
- PR #2037²³²⁹ - Avoid redundant specialization of [**unique**]function_nonser
- PR #2036²³³⁰ - nvcc dies with an internal error upon pushing/popping warnings inside templates
- Issue #2035²³³¹ - Use a less restrictive iterator definition in hpx::lcos::detail::future_iterator_traits
- PR #2034²³³² - Fixing compilation error with thread queue wait time performance counter
- Issue #2033²³³³ - Compilation error when compiling with thread queue waittime performance counter

²³¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/2056>

²³¹² <https://github.com/STELLAR-GROUP/hpx/pull/2055>

²³¹³ <https://github.com/STELLAR-GROUP/hpx/pull/2054>

²³¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2053>

²³¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2052>

²³¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2050>

²³¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2049>

²³¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2048>

²³¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2047>

²³²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2046>

²³²¹ <https://github.com/STELLAR-GROUP/hpx/pull/2045>

²³²² <https://github.com/STELLAR-GROUP/hpx/pull/2044>

²³²³ <https://github.com/STELLAR-GROUP/hpx/pull/2043>

²³²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2042>

²³²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2041>

²³²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2040>

²³²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2039>

²³²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2038>

²³²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2037>

²³³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2036>

²³³¹ <https://github.com/STELLAR-GROUP/hpx/issues/2035>

²³³² <https://github.com/STELLAR-GROUP/hpx/pull/2034>

²³³³ <https://github.com/STELLAR-GROUP/hpx/issues/2033>

- Issue #2032²³³⁴ - Ambiguous template instantiation for is_future_range and future_range_traits.
- PR #2031²³³⁵ - Don't restart timer on every incoming parcel
- PR #2030²³³⁶ - Unify handling of execution policies in parallel algorithms
- PR #2029²³³⁷ - Make pkg-config .pc files use .dylib on OSX
- PR #2028²³³⁸ - Adding process component
- PR #2027²³³⁹ - Making check for compiler compatibility independent on compiler path
- PR #2025²³⁴⁰ - Fixing inspect tool
- PR #2024²³⁴¹ - Intel13 removal
- PR #2023²³⁴² - Fix errors related to older boost versions and parameter pack expansions in lambdas
- Issue #2022²³⁴³ - gmake fail: "No rule to make target /usr/lib46/libboost_context-mt.so"
- PR #2021²³⁴⁴ - Added Sudoku example
- Issue #2020²³⁴⁵ - Make errors related to init_globally.cpp example while building HPX out of the box
- PR #2019²³⁴⁶ - Fixed some compilation and cmake errors encountered in nqueen example
- PR #2018²³⁴⁷ - For loop algorithms
- PR #2017²³⁴⁸ - Non-recursive at_index implementation
- Issue #2016²³⁴⁹ - Add index-based for-loops
- Issue #2015²³⁵⁰ - Change default bind-mode to balanced
- PR #2014²³⁵¹ - Fixed dataflow if invoked action returns a future
- PR #2013²³⁵² - Fixing compilation issues with external example
- PR #2012²³⁵³ - Added Sierpinski Triangle example
- Issue #2011²³⁵⁴ - Compilation error while running sample hello_world_component code
- PR #2010²³⁵⁵ - Segmented move implemented for hpx::vector
- Issue #2009²³⁵⁶ - pkg-config order incorrect on 14.04 / GCC 4.8

²³³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2032>

²³³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2031>

²³³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2030>

²³³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2029>

²³³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2028>

²³³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2027>

²³⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2025>

²³⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/2024>

²³⁴² <https://github.com/STELLAR-GROUP/hpx/pull/2023>

²³⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/2022>

²³⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/2021>

²³⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2020>

²³⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/2019>

²³⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/2018>

²³⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2017>

²³⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/2016>

²³⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/2015>

²³⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/2014>

²³⁵² <https://github.com/STELLAR-GROUP/hpx/pull/2013>

²³⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/2012>

²³⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2011>

²³⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/2010>

²³⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/2009>

- Issue #2008²³⁵⁷ - Compilation error in dataflow of action returning a future
- PR #2007²³⁵⁸ - Adding new performance counter exposing overall scheduler time
- PR #2006²³⁵⁹ - Function includes
- PR #2005²³⁶⁰ - Adding an example demonstrating how to initialize HPX from a global object
- PR #2004²³⁶¹ - Fixing 2000
- PR #2003²³⁶² - Adding generation parameter to gather to enable using it more than once
- PR #2002²³⁶³ - Turn on position independent code to solve link problem with hpx_init
- Issue #2001²³⁶⁴ - Gathering more than once segfaults
- Issue #2000²³⁶⁵ - Undefined reference to hpx::assertion_failed
- Issue #1999²³⁶⁶ - Seg fault in hpx::lcos::base_lco_with_value<*>::set_value_nonvirt() when running octo-tiger
- PR #1998²³⁶⁷ - Detect unknown command line options
- PR #1997²³⁶⁸ - Extending thread description
- PR #1996²³⁶⁹ - Adding natvis files to solution (MSVC only)
- Issue #1995²³⁷⁰ - Command line handling does not produce error
- PR #1994²³⁷¹ - Possible missing include in test_utils.hpp
- PR #1993²³⁷² - Add missing LANGUAGES tag to a hpx_add_compile_flag_if_available() call in CMake-Lists.txt
- PR #1992²³⁷³ - Fixing shared_executor_test
- PR #1991²³⁷⁴ - Making sure the winsock library is properly initialized
- PR #1990²³⁷⁵ - Fixing bind_test placeholder ambiguity coming from boost-1.60
- PR #1989²³⁷⁶ - Performance tuning
- PR #1987²³⁷⁷ - Make configurable size of internal storage in util::function
- PR #1986²³⁷⁸ - AGAS Refactoring+1753 Cache mods
- PR #1985²³⁷⁹ - Adding missing task_block::run() overload taking an executor

²³⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/2008>

²³⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/2007>

²³⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/2006>

²³⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/2005>

²³⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/2004>

²³⁶² <https://github.com/STELLAR-GROUP/hpx/pull/2003>

²³⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/2002>

²³⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/2001>

²³⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/2000>

²³⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1999>

²³⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1998>

²³⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1997>

²³⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1996>

²³⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1995>

²³⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1994>

²³⁷² <https://github.com/STELLAR-GROUP/hpx/pull/1993>

²³⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1992>

²³⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1991>

²³⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1990>

²³⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1989>

²³⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1987>

²³⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1986>

²³⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1985>

- PR #1984²³⁸⁰ - Adding an optimized LRU Cache implementation (for AGAS)
- PR #1983²³⁸¹ - Avoid invoking migration table look up for all objects
- PR #1981²³⁸² - Replacing uintptr_t (which is not defined everywhere) with std::size_t
- PR #1980²³⁸³ - Optimizing LCO continuations
- PR #1979²³⁸⁴ - Fixing Cori
- PR #1978²³⁸⁵ - Fix test check that got broken in hasty fix to memory overflow
- PR #1977²³⁸⁶ - Refactor action traits
- PR #1976²³⁸⁷ - Fixes typo in README.rst
- PR #1975²³⁸⁸ - Reduce size of benchmark timing arrays to fix test failures
- PR #1974²³⁸⁹ - Add action to update data owned by the partitioned_vector component
- PR #1972²³⁹⁰ - Adding partitioned_vector SPMD example
- PR #1971²³⁹¹ - Fixing 1965
- PR #1970²³⁹² - Papi fixes
- PR #1969²³⁹³ - Fixing continuation recursions to not depend on fixed amount of recursions
- PR #1968²³⁹⁴ - More segmented algorithms
- Issue #1967²³⁹⁵ - Simplify component implementations
- PR #1966²³⁹⁶ - Migrate components
- Issue #1964²³⁹⁷ - fatal error: ‘boost/lockfree/detail/branch_hints.hpp’ file not found
- Issue #1962²³⁹⁸ - parallel:copy_if has race condition when used on in place arrays
- PR #1963²³⁹⁹ - Fixing Static Parcelport initialization
- PR #1961²⁴⁰⁰ - Fix function target
- Issue #1960²⁴⁰¹ - Papi counters don’t reset
- PR #1959²⁴⁰² - Fixing 1958

²³⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1984>

²³⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/1983>

²³⁸² <https://github.com/STELLAR-GROUP/hpx/pull/1981>

²³⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/1980>

²³⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1979>

²³⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1978>

²³⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1977>

²³⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1976>

²³⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1975>

²³⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1974>

²³⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1972>

²³⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1971>

²³⁹² <https://github.com/STELLAR-GROUP/hpx/pull/1970>

²³⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/1969>

²³⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1968>

²³⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1967>

²³⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1966>

²³⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1964>

²³⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1962>

²³⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1963>

²⁴⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1961>

²⁴⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/1960>

²⁴⁰² <https://github.com/STELLAR-GROUP/hpx/pull/1959>

- Issue #1958²⁴⁰³ - inclusive_scan gives incorrect results with non-commutative operator
- PR #1957²⁴⁰⁴ - Fixing #1950
- PR #1956²⁴⁰⁵ - Sort by key example
- PR #1955²⁴⁰⁶ - Adding regression test for #1946: Hang in wait_all() in distributed run
- Issue #1954²⁴⁰⁷ - HPX releases should not use -Werror
- PR #1953²⁴⁰⁸ - Adding performance analysis for AGAS cache
- PR #1952²⁴⁰⁹ - Adapting test for explicit variadics to fail for gcc 4.6
- PR #1951²⁴¹⁰ - Fixing memory leak
- Issue #1950²⁴¹¹ - Simplify external builds
- PR #1949²⁴¹² - Fixing yet another lock that is being held during suspension
- PR #1948²⁴¹³ - Fixed container algorithms for Intel
- PR #1947²⁴¹⁴ - Adding workaround for tagged_tuple
- Issue #1946²⁴¹⁵ - Hang in wait_all() in distributed run
- PR #1945²⁴¹⁶ - Fixed container algorithm tests
- Issue #1944²⁴¹⁷ - assertion ‘p.destination_locality() == hpx::get_locality()’ failed
- PR #1943²⁴¹⁸ - Fix a couple of compile errors with clang
- PR #1942²⁴¹⁹ - Making parcel coalescing functional
- Issue #1941²⁴²⁰ - Re-enable parcel coalescing
- PR #1940²⁴²¹ - Touching up make_future
- PR #1939²⁴²² - Fixing problems in over-subscription management in the resource manager
- PR #1938²⁴²³ - Removing use of unified Boost.Thread header
- PR #1937²⁴²⁴ - Cleaning up the use of Boost.Accumulator headers
- PR #1936²⁴²⁵ - Making sure interval timer is started for aggregating performance counters

²⁴⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1958>

²⁴⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1957>

²⁴⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1956>

²⁴⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1955>

²⁴⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1954>

²⁴⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1953>

²⁴⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1952>

²⁴¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1951>

²⁴¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1950>

²⁴¹² <https://github.com/STELLAR-GROUP/hpx/pull/1949>

²⁴¹³ <https://github.com/STELLAR-GROUP/hpx/pull/1948>

²⁴¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1947>

²⁴¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1946>

²⁴¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1945>

²⁴¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1944>

²⁴¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1943>

²⁴¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1942>

²⁴²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1941>

²⁴²¹ <https://github.com/STELLAR-GROUP/hpx/pull/1940>

²⁴²² <https://github.com/STELLAR-GROUP/hpx/pull/1939>

²⁴²³ <https://github.com/STELLAR-GROUP/hpx/pull/1938>

²⁴²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1937>

²⁴²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1936>

- PR #1935²⁴²⁶ - Tagged results
- PR #1934²⁴²⁷ - Fix remote async with deferred launch policy
- Issue #1933²⁴²⁸ - Floating point exception in `statistics_counter<boost::accumulators::tag::mean>::get_mean()`
- PR #1932²⁴²⁹ - Removing superfluous includes of boost/lockfree/detail/branch_hints.hpp
- PR #1931²⁴³⁰ - fix compilation with clang 3.8.0
- Issue #1930²⁴³¹ - Missing online documentation for HPX 0.9.11
- PR #1929²⁴³² - LWG2485: get() should be overloaded for const tuple&&
- PR #1928²⁴³³ - Revert “Using ninja for circle-ci builds”
- PR #1927²⁴³⁴ - Using ninja for circle-ci builds
- PR #1926²⁴³⁵ - Fixing serialization of std::array
- Issue #1925²⁴³⁶ - Issues with static HPX libraries
- Issue #1924²⁴³⁷ - Performance degrading over time
- Issue #1923²⁴³⁸ - serialization of std::array appears broken in latest commit
- PR #1922²⁴³⁹ - Container algorithms
- PR #1921²⁴⁴⁰ - Tons of smaller quality improvements
- Issue #1920²⁴⁴¹ - Seg fault in hpx::serialization::output_archive::add_gid when running octotiger
- Issue #1919²⁴⁴² - Intel 15 compiler bug preventing HPX build
- PR #1918²⁴⁴³ - Address sanitizer fixes
- PR #1917²⁴⁴⁴ - Fixing compilation problems of parallel::sort with Intel compilers
- PR #1916²⁴⁴⁵ - Making sure code compiles if HPX_WITH_HWLOC=Off
- Issue #1915²⁴⁴⁶ - max_cores undefined if HPX_WITH_HWLOC=Off
- PR #1913²⁴⁴⁷ - Add utility member functions for partitioned_vector
- PR #1912²⁴⁴⁸ - Adding support for invoking actions to dataflow

²⁴²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1935>

²⁴²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1934>

²⁴²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1933>

²⁴²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1932>

²⁴³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1931>

²⁴³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1930>

²⁴³² <https://github.com/STELLAR-GROUP/hpx/pull/1929>

²⁴³³ <https://github.com/STELLAR-GROUP/hpx/pull/1928>

²⁴³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1927>

²⁴³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1926>

²⁴³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1925>

²⁴³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1924>

²⁴³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1923>

²⁴³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1922>

²⁴⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1921>

²⁴⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/1920>

²⁴⁴² <https://github.com/STELLAR-GROUP/hpx/issues/1919>

²⁴⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/1918>

²⁴⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1917>

²⁴⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1916>

²⁴⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1915>

²⁴⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1913>

²⁴⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1912>

- PR #1911²⁴⁴⁹ - Adding first batch of container algorithms
- PR #1910²⁴⁵⁰ - Keep cmake_module_path
- PR #1909²⁴⁵¹ - Fix mpirun with pbs
- PR #1908²⁴⁵² - Changing parallel::sort to return the last iterator as proposed by N4560
- PR #1907²⁴⁵³ - Adding a minimum version for Open MPI
- PR #1906²⁴⁵⁴ - Updates to the Release Procedure
- PR #1905²⁴⁵⁵ - Fixing #1903
- PR #1904²⁴⁵⁶ - Making sure std containers are cleared before serialization loads data
- Issue #1903²⁴⁵⁷ - When running octotiger, I get: assertion '(*new_gids_)[gid].size() == 1' failed: HPX(assertion_failure)
- Issue #1902²⁴⁵⁸ - Immediate crash when running hpx/octotiger with _GLIBCXX_DEBUG defined.
- PR #1901²⁴⁵⁹ - Making non-serializable classes non-serializable
- Issue #1900²⁴⁶⁰ - Two possible issues with std::list serialization
- PR #1899²⁴⁶¹ - Fixing a problem with credit splitting as revealed by #1898
- Issue #1898²⁴⁶² - Accessing component from locality where it was not created segfaults
- PR #1897²⁴⁶³ - Changing parallel::sort to return the last iterator as proposed by N4560
- Issue #1896²⁴⁶⁴ - version 1.0?
- Issue #1895²⁴⁶⁵ - Warning comment on numa_allocator is not very clear
- PR #1894²⁴⁶⁶ - Add support for compilers that have thread_local
- PR #1893²⁴⁶⁷ - Fixing 1890
- PR #1892²⁴⁶⁸ - Adds typed future_type for executor_traits
- PR #1891²⁴⁶⁹ - Fix wording in certain parallel algorithm docs
- Issue #1890²⁴⁷⁰ - Invoking papi counters give segfault
- PR #1889²⁴⁷¹ - Fixing problems as reported by clang-check

²⁴⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1911>

²⁴⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1910>

²⁴⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/1909>

²⁴⁵² <https://github.com/STELLAR-GROUP/hpx/pull/1908>

²⁴⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/1907>

²⁴⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1906>

²⁴⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1905>

²⁴⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1904>

²⁴⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1903>

²⁴⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1902>

²⁴⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1901>

²⁴⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1900>

²⁴⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/1899>

²⁴⁶² <https://github.com/STELLAR-GROUP/hpx/issues/1898>

²⁴⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/1897>

²⁴⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1896>

²⁴⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1895>

²⁴⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1894>

²⁴⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1893>

²⁴⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1892>

²⁴⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1891>

²⁴⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1890>

²⁴⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1889>

- PR #1888²⁴⁷² - WIP parallel is_heap
- PR #1887²⁴⁷³ - Fixed resetting performance counters related to idle-rate, etc
- Issue #1886²⁴⁷⁴ - Run hpx with qsub does not work
- PR #1885²⁴⁷⁵ - Warning cleaning pass
- PR #1884²⁴⁷⁶ - Add missing parallel algorithm header
- PR #1883²⁴⁷⁷ - Add feature test for thread_local on Clang for TLS
- PR #1882²⁴⁷⁸ - Fix some redundant qualifiers
- Issue #1881²⁴⁷⁹ - Unable to compile Octotiger using HPX and Intel MPI on SuperMIC
- Issue #1880²⁴⁸⁰ - clang with libc++ on Linux needs TLS case
- PR #1879²⁴⁸¹ - Doc fixes for #1868
- PR #1878²⁴⁸² - Simplify functions
- PR #1877²⁴⁸³ - Removing most usage of Boost.Config
- PR #1876²⁴⁸⁴ - Add missing parallel algorithms to algorithm.hpp
- PR #1875²⁴⁸⁵ - Simplify callables
- PR #1874²⁴⁸⁶ - Address long standing FIXME on using std::unique_ptr with incomplete types
- PR #1873²⁴⁸⁷ - Fixing 1871
- PR #1872²⁴⁸⁸ - Making sure PBS environment uses specified node list even if no PBS_NODEFILE env is available
- Issue #1871²⁴⁸⁹ - Fortran checks should be optional
- PR #1870²⁴⁹⁰ - Touch local::mutex
- PR #1869²⁴⁹¹ - Documentation refactoring based off #1868
- PR #1867²⁴⁹² - Embrace static_assert
- PR #1866²⁴⁹³ - Fix #1803 with documentation refactoring
- PR #1865²⁴⁹⁴ - Setting OUTPUT_NAME as target properties

²⁴⁷² <https://github.com/STELLAR-GROUP/hpx/pull/1888>

²⁴⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1887>

²⁴⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1886>

²⁴⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1885>

²⁴⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1884>

²⁴⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1883>

²⁴⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1882>

²⁴⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1881>

²⁴⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1880>

²⁴⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/1879>

²⁴⁸² <https://github.com/STELLAR-GROUP/hpx/pull/1878>

²⁴⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/1877>

²⁴⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1876>

²⁴⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1875>

²⁴⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1874>

²⁴⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1873>

²⁴⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1872>

²⁴⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1871>

²⁴⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1870>

²⁴⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1869>

²⁴⁹² <https://github.com/STELLAR-GROUP/hpx/pull/1867>

²⁴⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/1866>

²⁴⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1865>

- PR #1863²⁴⁹⁵ - Use SYSTEM for boost includes
- PR #1862²⁴⁹⁶ - Minor cleanups
- PR #1861²⁴⁹⁷ - Minor Corrections for Release
- PR #1860²⁴⁹⁸ - Fixing hpx gdb script
- Issue #1859²⁴⁹⁹ - reset_active_counters resets times and thread counts before some of the counters are evaluated
- PR #1858²⁵⁰⁰ - Release V0.9.11
- PR #1857²⁵⁰¹ - removing diskperf example from 9.11 release
- PR #1856²⁵⁰² - fix return in packaged_task_base::reset()
- Issue #1842²⁵⁰³ - Install error: file INSTALL cannot find libhpx_parcel_coalescing.so.0.9.11
- PR #1839²⁵⁰⁴ - Adding fedora docs
- PR #1824²⁵⁰⁵ - Changing version on master to V0.9.12
- PR #1818²⁵⁰⁶ - Fixing #1748
- Issue #1815²⁵⁰⁷ - seg fault in AGAS
- Issue #1803²⁵⁰⁸ - wait_all documentation
- Issue #1796²⁵⁰⁹ - Outdated documentation to be revised
- Issue #1759²⁵¹⁰ - glibc munmap_chunk or free(): invalid pointer on SuperMIC
- Issue #1753²⁵¹¹ - HPX performance degrades with time since execution begins
- Issue #1748²⁵¹² - All public HPX headers need to be self contained
- PR #1719²⁵¹³ - How to build HPX with Visual Studio
- Issue #1684²⁵¹⁴ - Race condition when using –hpx:connect?
- PR #1658²⁵¹⁵ - Add serialization for std::set (as there is for std::vector and std::map)
- PR #1641²⁵¹⁶ - Generic client
- Issue #1632²⁵¹⁷ - heartbeat example fails on separate nodes

²⁴⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1863>

²⁴⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1862>

²⁴⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1861>

²⁴⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1860>

²⁴⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1859>

²⁵⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1858>

²⁵⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/1857>

²⁵⁰² <https://github.com/STELLAR-GROUP/hpx/pull/1856>

²⁵⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1842>

²⁵⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1839>

²⁵⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1824>

²⁵⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1818>

²⁵⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1815>

²⁵⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1803>

²⁵⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1796>

²⁵¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1759>

²⁵¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1753>

²⁵¹² <https://github.com/STELLAR-GROUP/hpx/issues/1748>

²⁵¹³ <https://github.com/STELLAR-GROUP/hpx/pull/1719>

²⁵¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1684>

²⁵¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1658>

²⁵¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1641>

²⁵¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1632>

- PR #1603²⁵¹⁸ - Adds preferred namespace check to inspect tool
- Issue #1559²⁵¹⁹ - Extend inspect tool
- Issue #1523²⁵²⁰ - Remote async with deferred launch policy never executes
- Issue #1472²⁵²¹ - Serialization issues
- Issue #1457²⁵²² - Implement N4392: C++ Latches and Barriers
- PR #1444²⁵²³ - Enabling usage of moveonly types for component construction
- Issue #1407²⁵²⁴ - The Intel 13 compiler has failing unit tests
- Issue #1405²⁵²⁵ - Allow component constructors to take movable only types
- Issue #1265²⁵²⁶ - Enable dataflow() to be usable with actions
- Issue #1236²⁵²⁷ - NUMA aware allocators
- Issue #802²⁵²⁸ - Fix Broken Examples
- Issue #559²⁵²⁹ - Add hpx::migrate facility
- Issue #449²⁵³⁰ - Make actions with template arguments usable and add documentation
- Issue #279²⁵³¹ - Refactor addressing_service into a base class and two derived classes
- Issue #224²⁵³² - Changing thread state metadata is not thread safe
- Issue #55²⁵³³ - Uniform syntax for enums should be implemented

2.11.8 HPX V0.9.11 (Nov 11, 2015)

Our main focus for this release was the design and development of a coherent set of higher-level APIs exposing various types of parallelism to the application programmer. We introduced the concepts of an `executor`, which can be used to customize the `where` and `when` of execution of tasks in the context of parallelizing codes. We extended all APIs related to managing parallel tasks to support executors which gives the user the choice of either using one of the predefined executor types or to provide its own, possibly application specific, executor. We paid very close attention to align all of these changes with the existing C++ Standards documents or with the ongoing proposals for standardization.

This release is the first after our change to a new development policy. We switched all development to be strictly performed on branches only, all direct commits to our main branch (`master`) are prohibited. Any change has to go through a peer review before it will be merged to `master`. As a result the overall stability of our code base has significantly increased, the development process itself has been simplified. This change manifests itself in a large number of pull-requests which have been merged (please see below for a full list of closed issues and pull-requests).

²⁵¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1603>

²⁵¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1559>

²⁵²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1523>

²⁵²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1472>

²⁵²² <https://github.com/STELLAR-GROUP/hpx/issues/1457>

²⁵²³ <https://github.com/STELLAR-GROUP/hpx/pull/1444>

²⁵²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1407>

²⁵²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1405>

²⁵²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1265>

²⁵²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1236>

²⁵²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/802>

²⁵²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/559>

²⁵³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/449>

²⁵³¹ <https://github.com/STELLAR-GROUP/hpx/issues/279>

²⁵³² <https://github.com/STELLAR-GROUP/hpx/issues/224>

²⁵³³ <https://github.com/STELLAR-GROUP/hpx/issues/55>

All in all for this release, we closed almost 100 issues and merged over 290 pull-requests. There have been over 1600 commits to the master branch since the last release.

General changes

- We are moving into the direction of unifying managed and simple components. As such, the classes `hpx::components::component` and `hpx::components::component_base` have been added which currently just forward to the currently existing simple component facilities. The examples have been converted to only use those two classes.
- Added integration with the [CircleCI²⁵³⁴](https://circleci.com/gh/STELLAR-GROUP/hpx) hosted continuous integration service. This gives us constant and immediate feedback on the health of our master branch.
- The compiler configuration subsystem in the build system has been reimplemented. Instead of using Boost.Config we now use our own lightweight set of cmake scripts to determine the available language and library features supported by the used compiler.
- The API for creating instances of components has been consolidated. All component instances should be created using the `hpx::new_` only. It allows one to instantiate both, single component instances and multiple component instances. The placement of the created components can be controlled by special distribution policies. Please see the corresponding documentation outlining the use of `hpx::new_`.
- Introduced four new distribution policies which can be used with many API functions which traditionally expected to be used with a locality id. The new distribution policies are:
 - `hpx::components::default_distribution_policy` which tries to place multiple component instances as evenly as possible.
 - `hpx::components::colocating_distribution_policy` which will refer to the locality where a given component instance is currently placed.
 - `hpx::components::binpacking_distribution_policy` which will place multiple component instances as evenly as possible based on any performance counter.
 - `hpx::components::target_distribution_policy` which allows one to represent a given locality in the context of a distribution policy.
- The new distribution policies can now be also used with `hpx::async`. This change also deprecates `hpx::async_colocated(id, ...)` which now is replaced by a distribution policy: `hpx::async(hpx::colocated(id), ...)`.
- The `hpx::vector` and `hpx::unordered_map` data structures can now be used with the new distribution policies as well.
- The parallel facility `hpx::parallel::task_region` has been renamed to `hpx::parallel::task_block` based on the changes in the corresponding standardization proposal [N4411²⁵³⁵](#).
- Added extensions to the parallel facility `hpx::parallel::task_block` allowing to combine a `task_block` with an execution policy. This implies a minor breaking change as the `hpx::parallel::task_block` is now a template.
- Added new LCOs: `hpx::lcos::latch` and `hpx::lcos::local::latch` which semantically conform to the proposed `std::latch` (see [N4399²⁵³⁶](#)).
- Added performance counters exposing data related to data transferred by input/output (filesystem) operations (thanks to Maciej Brodowicz).

²⁵³⁴ <https://circleci.com/gh/STELLAR-GROUP/hpx>

²⁵³⁵ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4411.pdf>

²⁵³⁶ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4399.html>

- Added performance counters allowing to track the number of action invocations (local and remote invocations).
- Added new command line options `-hpx:print-counter-at` and `-hpx:reset-counters`.
- The `hpx::vector` component has been renamed to `hpx::partitioned_vector` to make it explicit that the underlying memory is not contiguous.
- Introduced a completely new and uniform higher-level parallelism API which is based on executors. All existing parallelism APIs have been adapted to this. We have added a large number of different executor types, such as a numa-aware executor, a this-thread executor, etc.
- Added support for the MingW toolchain on Windows (thanks to Eric Lemanissier).
- HPX now includes support for APEX, (Autonomic Performance Environment for eXascale). APEX is an instrumentation and software adaptation library that provides an interface to TAU profiling / tracing as well as runtime adaptation of HPX applications through policy definitions. For more information and documentation, please see <https://github.com/khuck/xpress-apex>. To enable APEX at configuration time, specify `-DHPX_WITH_APEX=On`. To also include support for TAU profiling, specify `-DHPX_WITH_TAU=On` and specify the `-DTAU_ROOT`, `-DTAU_ARCH` and `-DTAU_OPTIONS` cmake parameters.
- We have implemented many more of the [Using parallel algorithms](#). Please see [Issue #1141](#)²⁵³⁷ for the list of all available parallel algorithms (thanks to Daniel Bourgeois and John Biddiscombe for contributing their work).

Breaking changes

- We are moving into the direction of unifying managed and simple components. In order to stop exposing the old facilities, all examples have been converted to use the new classes. The breaking change in this release is that performance counters are now a `hpx::components::component_base` instead of `hpx::components::managed_component_base`.
- We removed the support for stackless threads. It turned out that there was no performance benefit when using stackless threads. As such, we decided to clean up our codebase. This feature was not documented.
- The CMake project name has changed from ‘hpx’ to ‘HPX’ for consistency and compatibility with naming conventions and other CMake projects. Generated config files go into `<prefix>/lib/cmake/HPX` and not `<prefix>/lib/cmake/hpx`.
- The macro `HPX_REGISTER_MINIMAL_COMPONENT_FACTORY` has been deprecated. Please use [`HPX_REGISTER_COMPONENT`](#) instead. The old macro will be removed in the next release.
- The obsolete `distributing_factory` and `binpacking_factory` components have been removed. The corresponding functionality is now provided by the `hpx::new_` API function in conjunction with the `hpx::default_layout` and `hpx::binpacking` distribution policies ([`hpx::components::default_distribution_policy`](#) and [`hpx::components::binpacking_distribution_policy`](#))
- The API function `hpx::new_colocated` has been deprecated. Please use the consolidated API `hpx::new_` in conjunction with the new `hpx::colocated` distribution policy ([`hpx::components::colocating_distribution_policy`](#)) instead. The old API function will still be available for at least one release of `HPX` if the configuration variable `HPX_WITH_COLOCATED_BACKWARDS_COMPATIBILITY` is enabled.
- The API function `hpx::async_colocated` has been deprecated. Please use the consolidated API `hpx::async` in conjunction with the new `hpx::colocated` distribution policy ([`hpx::components::colocating_distribution_policy`](#)) instead. The old API function will still be available for at least one release of `HPX` if the configuration variable `HPX_WITH_COLOCATED_BACKWARDS_COMPATIBILITY` is enabled.

²⁵³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1141>

- The obsolete `remote_object` component has been removed.
- Replaced the use of `Boost.Serialization` with our own solution. While the new version is mostly compatible with `Boost.Serialization`, this change requires some minor code modifications in user code. For more information, please see the corresponding announcement²⁵³⁸ on the hpx-users@stellar.cct.lsu.edu mailing list.
- The names used by `cmake` to influence various configuration options have been unified. The new naming scheme relies on all configuration constants to start with `HPX_WITH_...`, while the preprocessor constant which is used at build time starts with `HPX_HAVE_...`. For instance, the former `cmake` command line `-DHPX_MALLOC=...` now has to be specified a `-DHPX_WITH_MALLOC=...` and will cause the preprocessor constant `HPX_HAVE_MALLOC` to be defined. The actual name of the constant (i.e. `MALLOC`) has not changed. Please see the corresponding documentation for more details (*CMake variables used to configure HPX*).
- The `get_gid()` functions exposed by the component base classes `hpx::components::server::simple_component_base`, `hpx::components::server::managed_component` and `hpx::components::server::fixed_component_base` have been replaced by two new functions: `get_unmanaged_id()` and `get_id()`. To enable the old function name for backwards compatibility, use the `cmake` configuration option `HPX_WITH_COMPONENT_GET_GID_COMPATIBILITY=On`.
- All functions which were named `get_gid()` but were returning `hpx::id_type` have been renamed to `get_id()`. To enable the old function names for backwards compatibility, use the `cmake` configuration option `HPX_WITH_COMPONENT_GET_GID_COMPATIBILITY=On`.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- PR #1855²⁵³⁹ - Completely removing external/endian
- PR #1854²⁵⁴⁰ - Don't pollute CMAKE_CXX_FLAGS through `find_package()`
- PR #1853²⁵⁴¹ - Updating CMake configuration to get correct version of TAU library
- PR #1852²⁵⁴² - Fixing Performance Problems with MPI Parcelport
- PR #1851²⁵⁴³ - Fixing `hpx_add_link_flag()` and `hpx_remove_link_flag()`
- PR #1850²⁵⁴⁴ - Fixing 1836, adding `parallel::sort`
- PR #1849²⁵⁴⁵ - Fixing configuration for use of more than 64 cores
- PR #1848²⁵⁴⁶ - Change default APEX version for release
- PR #1847²⁵⁴⁷ - Fix `client_base::then` on release
- PR #1846²⁵⁴⁸ - Removing broken `lcos::local::channel` from release
- PR #1845²⁵⁴⁹ - Adding example demonstrating a possible safe-object implementation to release

²⁵³⁸ <http://thread.gmane.org/gmane.comp.lib.hpx.devel/196>

²⁵³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1855>

²⁵⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1854>

²⁵⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/1853>

²⁵⁴² <https://github.com/STELLAR-GROUP/hpx/pull/1852>

²⁵⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/1851>

²⁵⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1850>

²⁵⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1849>

²⁵⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1848>

²⁵⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1847>

²⁵⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1846>

²⁵⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1845>

- PR #1844²⁵⁵⁰ - Removing stubs from accumulator examples
- PR #1843²⁵⁵¹ - Don't pollute CMAKE_CXX_FLAGS through find_package()
- PR #1841²⁵⁵² - Fixing client_base<>::then
- PR #1840²⁵⁵³ - Adding example demonstrating a possible safe-object implementation
- PR #1838²⁵⁵⁴ - Update version rc1
- PR #1837²⁵⁵⁵ - Removing broken lcos::local::channel
- PR #1835²⁵⁵⁶ - Adding explicit move constructor and assignment operator to hpx::lcos::promise
- PR #1834²⁵⁵⁷ - Making hpx::lcos::promise move-only
- PR #1833²⁵⁵⁸ - Adding fedora docs
- Issue #1832²⁵⁵⁹ - hpx::lcos::promise<> must be move-only
- PR #1831²⁵⁶⁰ - Fixing resource manager gcc5.2
- PR #1830²⁵⁶¹ - Fix intel13
- PR #1829²⁵⁶² - Unbreaking thread test
- PR #1828²⁵⁶³ - Fixing #1620
- PR #1827²⁵⁶⁴ - Fixing a memory management issue for the Parquet application
- Issue #1826²⁵⁶⁵ - Memory management issue in hpx::lcos::promise
- PR #1825²⁵⁶⁶ - Adding hpx::components::component and hpx::components::component_base
- PR #1823²⁵⁶⁷ - Adding git commit id to circleci build
- PR #1822²⁵⁶⁸ - applying fixes suggested by clang 3.7
- PR #1821²⁵⁶⁹ - Hyperlink fixes
- PR #1820²⁵⁷⁰ - added parallel multi-locality sanity test
- PR #1819²⁵⁷¹ - Fixing #1667
- Issue #1817²⁵⁷² - Hyperlinks generated by inspect tool are wrong

²⁵⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1844>

²⁵⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/1843>

²⁵⁵² <https://github.com/STELLAR-GROUP/hpx/pull/1841>

²⁵⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/1840>

²⁵⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1838>

²⁵⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1837>

²⁵⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1835>

²⁵⁵⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1834>

²⁵⁵⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1833>

²⁵⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1832>

²⁵⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1831>

²⁵⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/1830>

²⁵⁶² <https://github.com/STELLAR-GROUP/hpx/pull/1829>

²⁵⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/1828>

²⁵⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1827>

²⁵⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1826>

²⁵⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1825>

²⁵⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1823>

²⁵⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1822>

²⁵⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1821>

²⁵⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1820>

²⁵⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1819>

²⁵⁷² <https://github.com/STELLAR-GROUP/hpx/issues/1817>

- PR #1816²⁵⁷³ - Support hpxrx
- PR #1814²⁵⁷⁴ - Fix async to dispatch to the correct locality in all cases
- Issue #1813²⁵⁷⁵ - async(launch::..., action(), ...) always invokes locally
- PR #1812²⁵⁷⁶ - fixed syntax error in CMakeLists.txt
- PR #1811²⁵⁷⁷ - Agas optimizations
- PR #1810²⁵⁷⁸ - drop superfluous typedefs
- PR #1809²⁵⁷⁹ - Allow HPX to be used as an optional package in 3rd party code
- PR #1808²⁵⁸⁰ - Fixing #1723
- PR #1807²⁵⁸¹ - Making sure resolve_localities does not hang during normal operation
- Issue #1806²⁵⁸² - Spinlock no longer movable and deletes operator '=' , breaks MiniGhost
- Issue #1804²⁵⁸³ - register_with_basename causes hangs
- PR #1801²⁵⁸⁴ - Enhanced the inspect tool to take user directly to the problem with hyperlinks
- Issue #1800²⁵⁸⁵ - Problems compiling application on smic
- PR #1799²⁵⁸⁶ - Fixing cv exceptions
- PR #1798²⁵⁸⁷ - Documentation refactoring & updating
- PR #1797²⁵⁸⁸ - Updating the activeharmony CMake module
- PR #1795²⁵⁸⁹ - Fixing cv
- PR #1794²⁵⁹⁰ - Fix connect with hpx::runtime_mode_connect
- PR #1793²⁵⁹¹ - fix a wrong use of HPX_MAX_CPU_COUNT instead of HPX_HAVE_MAX_CPU_COUNT
- PR #1792²⁵⁹² - Allow for default constructed parcel instances to be moved
- PR #1791²⁵⁹³ - Fix connect with hpx::runtime_mode_connect
- Issue #1790²⁵⁹⁴ - assertion action_.get () failed: HPX(assertion_failure) when running Octotiger with pull request 1786
- PR #1789²⁵⁹⁵ - Fixing discover_counter_types API function

²⁵⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1816>

²⁵⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1814>

²⁵⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1813>

²⁵⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1812>

²⁵⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1811>

²⁵⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1810>

²⁵⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1809>

²⁵⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1808>

²⁵⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/1807>

²⁵⁸² <https://github.com/STELLAR-GROUP/hpx/issues/1806>

²⁵⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/1804>

²⁵⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1801>

²⁵⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1800>

²⁵⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1799>

²⁵⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1798>

²⁵⁸⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1797>

²⁵⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1795>

²⁵⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1794>

²⁵⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1793>

²⁵⁹² <https://github.com/STELLAR-GROUP/hpx/pull/1792>

²⁵⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/1791>

²⁵⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1790>

²⁵⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1789>

- Issue #1788²⁵⁹⁶ - connect with hpx::runtime_mode_connect
- Issue #1787²⁵⁹⁷ - discover_counter_types not working
- PR #1786²⁵⁹⁸ - Changing addressing_service to use std::unordered_map instead of std::map
- PR #1785²⁵⁹⁹ - Fix is_iterator for container algorithms
- PR #1784²⁶⁰⁰ - Adding new command line options:
- PR #1783²⁶⁰¹ - Minor changes for APEX support
- PR #1782²⁶⁰² - Drop legacy forwarding action traits
- PR #1781²⁶⁰³ - Attempt to resolve the race between cv::wait_xxx and cv::notify_all
- PR #1780²⁶⁰⁴ - Removing serialize_sequence
- PR #1779²⁶⁰⁵ - Fixed #1501: hwloc configuration options are wrong for MIC
- PR #1778²⁶⁰⁶ - Removing ability to enable/disable parcel handling
- PR #1777²⁶⁰⁷ - Completely removing stackless threads
- PR #1776²⁶⁰⁸ - Cleaning up util/plugin
- PR #1775²⁶⁰⁹ - Agas fixes
- PR #1774²⁶¹⁰ - Action invocation count
- PR #1773²⁶¹¹ - replaced MSVC variable with WIN32
- PR #1772²⁶¹² - Fixing Problems in MPI parcelport and future serialization.
- PR #1771²⁶¹³ - Fixing intel 13 compiler errors related to variadic template template parameters for `lcos::when_` tests
- PR #1770²⁶¹⁴ - Forwarding decay to `std::`:
- PR #1769²⁶¹⁵ - Add more characters with special regex meaning to the existing patch
- PR #1768²⁶¹⁶ - Adding test for receive_buffer
- PR #1767²⁶¹⁷ - Making sure that uptime counter throws exception on any attempt to be reset
- PR #1766²⁶¹⁸ - Cleaning up code related to throttling scheduler

²⁵⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1788>

²⁵⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1787>

²⁵⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1786>

²⁵⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1785>

²⁶⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1784>

²⁶⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/1783>

²⁶⁰² <https://github.com/STELLAR-GROUP/hpx/pull/1782>

²⁶⁰³ <https://github.com/STELLAR-GROUP/hpx/pull/1781>

²⁶⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1780>

²⁶⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1779>

²⁶⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1778>

²⁶⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1777>

²⁶⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1776>

²⁶⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1775>

²⁶¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1774>

²⁶¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1773>

²⁶¹² <https://github.com/STELLAR-GROUP/hpx/pull/1772>

²⁶¹³ <https://github.com/STELLAR-GROUP/hpx/pull/1771>

²⁶¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1770>

²⁶¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1769>

²⁶¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1768>

²⁶¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1767>

²⁶¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1766>

- PR #1765²⁶¹⁹ - Restricting thread_data to creating only with intrusive_pointers
- PR #1764²⁶²⁰ - Fixing 1763
- Issue #1763²⁶²¹ - UB in thread_data::operator delete
- PR #1762²⁶²² - Making sure all serialization registries/factories are unique
- PR #1761²⁶²³ - Fixed #1751: hpx::future::wait_for fails a simple test
- PR #1758²⁶²⁴ - Fixing #1757
- Issue #1757²⁶²⁵ - pinning not correct using –hpx:bind
- Issue #1756²⁶²⁶ - compilation error with MinGW
- PR #1755²⁶²⁷ - Making output serialization const-correct
- Issue #1753²⁶²⁸ - HPX performance degrades with time since execution begins
- Issue #1752²⁶²⁹ - Error in AGAS
- Issue #1751²⁶³⁰ - hpx::future::wait_for fails a simple test
- PR #1750²⁶³¹ - Removing hpx_fwd.hpp includes
- PR #1749²⁶³² - Simplify result_of and friends
- PR #1747²⁶³³ - Removed superfluous code from message_buffer.hpp
- PR #1746²⁶³⁴ - Tuple dependencies
- Issue #1745²⁶³⁵ - Broken when_some which takes iterators
- PR #1744²⁶³⁶ - Refining archive interface
- PR #1743²⁶³⁷ - Fixing when_all when only a single future is passed
- PR #1742²⁶³⁸ - Config includes
- PR #1741²⁶³⁹ - Os executors
- Issue #1740²⁶⁴⁰ - hpx::promise has some problems
- PR #1739²⁶⁴¹ - Parallel composition with generic containers

²⁶¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1765>

²⁶²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1764>

²⁶²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1763>

²⁶²² <https://github.com/STELLAR-GROUP/hpx/pull/1762>

²⁶²³ <https://github.com/STELLAR-GROUP/hpx/pull/1761>

²⁶²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1758>

²⁶²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1757>

²⁶²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1756>

²⁶²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1755>

²⁶²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1753>

²⁶²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1752>

²⁶³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1751>

²⁶³¹ <https://github.com/STELLAR-GROUP/hpx/pull/1750>

²⁶³² <https://github.com/STELLAR-GROUP/hpx/pull/1749>

²⁶³³ <https://github.com/STELLAR-GROUP/hpx/pull/1747>

²⁶³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1746>

²⁶³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1745>

²⁶³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1744>

²⁶³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1743>

²⁶³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1742>

²⁶³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1741>

²⁶⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1740>

²⁶⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/1739>

- Issue #1738²⁶⁴² - After building program and successfully linking to a version of hpx DHPX_DIR seems to be ignored
- Issue #1737²⁶⁴³ - Uptime problems
- PR #1736²⁶⁴⁴ - added convenience c-tor and begin()/end() to serialize_buffer
- PR #1735²⁶⁴⁵ - Config includes
- PR #1734²⁶⁴⁶ - Fixed #1688: Add timer counters for tfunc_total and exec_total
- Issue #1733²⁶⁴⁷ - Add unit test for hpx/lcos/local/receive_buffer.hpp
- PR #1732²⁶⁴⁸ - Renaming get_os_thread_count
- PR #1731²⁶⁴⁹ - Basename registration
- Issue #1730²⁶⁵⁰ - Use after move of thread_init_data
- PR #1729²⁶⁵¹ - Rewriting channel based on new gate component
- PR #1728²⁶⁵² - Fixing #1722
- PR #1727²⁶⁵³ - Fixing compile problems with apply_colocated
- PR #1726²⁶⁵⁴ - Apex integration
- PR #1725²⁶⁵⁵ - fixed test timeouts
- PR #1724²⁶⁵⁶ - Renaming vector
- Issue #1723²⁶⁵⁷ - Drop support for intel compilers and gcc 4.4. based standard libs
- Issue #1722²⁶⁵⁸ - Add support for detecting non-ready futures before serialization
- PR #1721²⁶⁵⁹ - Unifying parallel executors, initializing from launch policy
- PR #1720²⁶⁶⁰ - dropped superfluous typedef
- Issue #1718²⁶⁶¹ - Windows 10 x64, VS 2015 - Unknown CMake command “add_hpx_pseudo_target”.
- PR #1717²⁶⁶² - Timed executor traits for thread-executors
- PR #1716²⁶⁶³ - serialization of arrays didn't work with non-pod types. fixed
- PR #1715²⁶⁶⁴ - List serialization

²⁶⁴² <https://github.com/STELLAR-GROUP/hpx/issues/1738>

²⁶⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1737>

²⁶⁴⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1736>

²⁶⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1735>

²⁶⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1734>

²⁶⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1733>

²⁶⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1732>

²⁶⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1731>

²⁶⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1730>

²⁶⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/1729>

²⁶⁵² <https://github.com/STELLAR-GROUP/hpx/pull/1728>

²⁶⁵³ <https://github.com/STELLAR-GROUP/hpx/pull/1727>

²⁶⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1726>

²⁶⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1725>

²⁶⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1724>

²⁶⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1723>

²⁶⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1722>

²⁶⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1721>

²⁶⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1720>

²⁶⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/1718>

²⁶⁶² <https://github.com/STELLAR-GROUP/hpx/pull/1717>

²⁶⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/1716>

²⁶⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1715>

- PR #1714²⁶⁶⁵ - changing misspellings
- PR #1713²⁶⁶⁶ - Fixed distribution policy executors
- PR #1712²⁶⁶⁷ - Moving library detection to be executed after feature tests
- PR #1711²⁶⁶⁸ - Simplify parcel
- PR #1710²⁶⁶⁹ - Compile only tests
- PR #1709²⁶⁷⁰ - Implemented timed executors
- PR #1708²⁶⁷¹ - Implement parallel::executor_traits for thread-executors
- PR #1707²⁶⁷² - Various fixes to threads::executors to make custom schedulers work
- PR #1706²⁶⁷³ - Command line option –hpx:cores does not work as expected
- Issue #1705²⁶⁷⁴ - command line option –hpx:cores does not work as expected
- PR #1704²⁶⁷⁵ - vector deserialization is speeded up a little
- PR #1703²⁶⁷⁶ - Fixing shared_mutexes
- Issue #1702²⁶⁷⁷ - Shared_mutex does not compile with no_mutex cond_var
- PR #1701²⁶⁷⁸ - Add distribution_policy_executor
- PR #1700²⁶⁷⁹ - Executor parameters
- PR #1699²⁶⁸⁰ - Readers writer lock
- PR #1698²⁶⁸¹ - Remove leftovers
- PR #1697²⁶⁸² - Fixing held locks
- PR #1696²⁶⁸³ - Modified Scan Partitioner for Algorithms
- PR #1695²⁶⁸⁴ - This thread executors
- PR #1694²⁶⁸⁵ - Fixed #1688: Add timer counters for tfunc_total and exec_total
- PR #1693²⁶⁸⁶ - Fix #1691: is_executor template specification fails for inherited executors
- PR #1692²⁶⁸⁷ - Fixed #1662: Possible exception source in coalescing_message_handler

²⁶⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1714>

²⁶⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1713>

²⁶⁶⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1712>

²⁶⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1711>

²⁶⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1710>

²⁶⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1709>

²⁶⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1708>

²⁶⁷² <https://github.com/STELLAR-GROUP/hpx/pull/1707>

²⁶⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1706>

²⁶⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1705>

²⁶⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1704>

²⁶⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1703>

²⁶⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1702>

²⁶⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1701>

²⁶⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1700>

²⁶⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1699>

²⁶⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/1698>

²⁶⁸² <https://github.com/STELLAR-GROUP/hpx/pull/1697>

²⁶⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/1696>

²⁶⁸⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1695>

²⁶⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1694>

²⁶⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1693>

²⁶⁸⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1692>

- Issue #1691²⁶⁸⁸ - is_executor template specification fails for inherited executors
- PR #1690²⁶⁸⁹ - added macro for non-intrusive serialization of classes without a default c-tor
- PR #1689²⁶⁹⁰ - Replace value_or_error with custom storage, unify future_data state
- Issue #1688²⁶⁹¹ - Add timer counters for tfunc_total and exec_total
- PR #1687²⁶⁹² - Fixed interval timer
- PR #1686²⁶⁹³ - Fixing cmake warnings about not existing pseudo target dependencies
- PR #1685²⁶⁹⁴ - Converting partitioners to use bulk async execute
- PR #1683²⁶⁹⁵ - Adds a tool for inspect that checks for character limits
- PR #1682²⁶⁹⁶ - Change project name to (uppercase) HPX
- PR #1681²⁶⁹⁷ - Counter shortnames
- PR #1680²⁶⁹⁸ - Extended Non-intrusive Serialization to Ease Usage for Library Developers
- PR #1679²⁶⁹⁹ - Working on 1544: More executor changes
- PR #1678²⁷⁰⁰ - Transpose fixes
- PR #1677²⁷⁰¹ - Improve Boost compatibility check
- PR #1676²⁷⁰² - 1d stencil fix
- Issue #1675²⁷⁰³ - hpx project name is not HPX
- PR #1674²⁷⁰⁴ - Fixing the MPI parcelport
- PR #1673²⁷⁰⁵ - added move semantics to map/vector deserialization
- PR #1672²⁷⁰⁶ - Vs2015 await
- PR #1671²⁷⁰⁷ - Adapt transform for #1668
- PR #1670²⁷⁰⁸ - Started to work on #1668
- PR #1669²⁷⁰⁹ - Add this_thread_executors
- Issue #1667²⁷¹⁰ - Apple build instructions in docs are out of date

²⁶⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1691>

²⁶⁸⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1690>

²⁶⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1689>

²⁶⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1688>

²⁶⁹² <https://github.com/STELLAR-GROUP/hpx/pull/1687>

²⁶⁹³ <https://github.com/STELLAR-GROUP/hpx/pull/1686>

²⁶⁹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1685>

²⁶⁹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1683>

²⁶⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1682>

²⁶⁹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1681>

²⁶⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1680>

²⁶⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1679>

²⁷⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1678>

²⁷⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/1677>

²⁷⁰² <https://github.com/STELLAR-GROUP/hpx/pull/1676>

²⁷⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1675>

²⁷⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1674>

²⁷⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1673>

²⁷⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1672>

²⁷⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1671>

²⁷⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1670>

²⁷⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1669>

²⁷¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1667>

- PR #1666²⁷¹¹ - Apex integration
- PR #1665²⁷¹² - Fixes an error with the whitespace check that showed the incorrect location of the error
- Issue #1664²⁷¹³ - Inspect tool found incorrect endline whitespace
- PR #1663²⁷¹⁴ - Improve use of locks
- Issue #1662²⁷¹⁵ - Possible exception source in coalescing_message_handler
- PR #1661²⁷¹⁶ - Added support for 128bit number serialization
- PR #1660²⁷¹⁷ - Serialization 128bits
- PR #1659²⁷¹⁸ - Implemented inner_product and adjacent_diff algos
- PR #1658²⁷¹⁹ - Add serialization for std::set (as there is for std::vector and std::map)
- PR #1657²⁷²⁰ - Use of shared_ptr in io_service_pool changed to unique_ptr
- Issue #1656²⁷²¹ - 1d_stencil codes all have wrong factor
- PR #1654²⁷²² - When using runtime_mode_connect, find the correct localhost public ip address
- PR #1653²⁷²³ - Fixing 1617
- PR #1652²⁷²⁴ - Remove traits::action_may_require_id_splitting
- PR #1651²⁷²⁵ - Fixed performance counters related to AGAS cache timings
- PR #1650²⁷²⁶ - Remove leftovers of traits::type_size
- PR #1649²⁷²⁷ - Shorten target names on Windows to shorten used path names
- PR #1648²⁷²⁸ - Fixing problems introduced by merging #1623 for older compilers
- PR #1647²⁷²⁹ - Simplify running automatic builds on Windows
- Issue #1646²⁷³⁰ - Cache insert and update performance counters are broken
- Issue #1644²⁷³¹ - Remove leftovers of traits::type_size
- Issue #1643²⁷³² - Remove traits::action_may_require_id_splitting
- PR #1642²⁷³³ - Adds spell checker to the inspect tool for qbk and doxygen comments

²⁷¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1666>

²⁷¹² <https://github.com/STELLAR-GROUP/hpx/pull/1665>

²⁷¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1664>

²⁷¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1663>

²⁷¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1662>

²⁷¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1661>

²⁷¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1660>

²⁷¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1659>

²⁷¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1658>

²⁷²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1657>

²⁷²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1656>

²⁷²² <https://github.com/STELLAR-GROUP/hpx/pull/1654>

²⁷²³ <https://github.com/STELLAR-GROUP/hpx/pull/1653>

²⁷²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1652>

²⁷²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1651>

²⁷²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1650>

²⁷²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1649>

²⁷²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1648>

²⁷²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1647>

²⁷³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1646>

²⁷³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1644>

²⁷³² <https://github.com/STELLAR-GROUP/hpx/issues/1643>

²⁷³³ <https://github.com/STELLAR-GROUP/hpx/pull/1642>

- PR #1640²⁷³⁴ - First step towards fixing 688
- PR #1639²⁷³⁵ - Re-apply remaining changes from limit_dataflow_recursion branch
- PR #1638²⁷³⁶ - This fixes possible deadlock in the test ignore_while_locked_1485
- PR #1637²⁷³⁷ - Fixing hpx::wait_all() invoked with two vector<future<T>>
- PR #1636²⁷³⁸ - Partially re-apply changes from limit_dataflow_recursion branch
- PR #1635²⁷³⁹ - Adding missing test for #1572
- PR #1634²⁷⁴⁰ - Revert “Limit recursion-depth in dataflow to a configurable constant”
- PR #1633²⁷⁴¹ - Add command line option to ignore batch environment
- PR #1631²⁷⁴² - hpx::lcos::queue exhibits strange behavior
- PR #1630²⁷⁴³ - Fixed endline_whitespace_check.cpp to detect lines with only whitespace
- Issue #1629²⁷⁴⁴ - Inspect trailing whitespace checker problem
- PR #1628²⁷⁴⁵ - Removed meaningless const qualifiers. Minor icpc fix.
- PR #1627²⁷⁴⁶ - Fixing the queue LCO and add example demonstrating its use
- PR #1626²⁷⁴⁷ - Deprecating get_gid(), add get_id() and get_unmanaged_id()
- PR #1625²⁷⁴⁸ - Allowing to specify whether to send credits along with message
- Issue #1624²⁷⁴⁹ - Lifetime issue
- Issue #1623²⁷⁵⁰ - hpx::wait_all() invoked with two vector<future<T>> fails
- PR #1622²⁷⁵¹ - Executor partitioners
- PR #1621²⁷⁵² - Clean up coroutines implementation
- Issue #1620²⁷⁵³ - Revert #1535
- PR #1619²⁷⁵⁴ - Fix result type calculation for hpx::make_continuation
- PR #1618²⁷⁵⁵ - Fixing RDTSC on Xeon/Phi
- Issue #1617²⁷⁵⁶ - hpx cmake not working when run as a subproject

²⁷³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1640>

²⁷³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1639>

²⁷³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1638>

²⁷³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1637>

²⁷³⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1636>

²⁷³⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1635>

²⁷⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1634>

²⁷⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/1633>

²⁷⁴² <https://github.com/STELLAR-GROUP/hpx/pull/1631>

²⁷⁴³ <https://github.com/STELLAR-GROUP/hpx/pull/1630>

²⁷⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1629>

²⁷⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1628>

²⁷⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1627>

²⁷⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1626>

²⁷⁴⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1625>

²⁷⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1624>

²⁷⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1623>

²⁷⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/1622>

²⁷⁵² <https://github.com/STELLAR-GROUP/hpx/pull/1621>

²⁷⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1620>

²⁷⁵⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1619>

²⁷⁵⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1618>

²⁷⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1617>

- Issue #1616²⁷⁵⁷ - cmake problem resulting in RDTSC not working correctly for Xeon Phi creates very strange results for duration counters
- Issue #1615²⁷⁵⁸ - hpx::make_continuation requires input and output to be the same
- PR #1614²⁷⁵⁹ - Fixed remove copy test
- Issue #1613²⁷⁶⁰ - Dataflow causes stack overflow
- PR #1612²⁷⁶¹ - Modified foreach partitioner to use bulk execute
- PR #1611²⁷⁶² - Limit recursion-depth in dataflow to a configurable constant
- PR #1610²⁷⁶³ - Increase timeout for CircleCI
- PR #1609²⁷⁶⁴ - Refactoring thread manager, mainly extracting thread pool
- PR #1608²⁷⁶⁵ - Fixed running multiple localities without localities parameter
- PR #1607²⁷⁶⁶ - More algorithm fixes to adjacentfind
- Issue #1606²⁷⁶⁷ - Running without localities parameter binds to bogus port range
- Issue #1605²⁷⁶⁸ - Too many serializations
- PR #1604²⁷⁶⁹ - Changes the HPX image into a hyperlink
- PR #1601²⁷⁷⁰ - Fixing problems with remove_copy algorithm tests
- PR #1600²⁷⁷¹ - Actions with ids cleanup
- PR #1599²⁷⁷² - Duplicate binding of global ids should fail
- PR #1598²⁷⁷³ - Fixing array access
- PR #1597²⁷⁷⁴ - Improved the reliability of connecting/disconnecting localities
- Issue #1596²⁷⁷⁵ - Duplicate id binding should fail
- PR #1595²⁷⁷⁶ - Fixing more cmake config constants
- PR #1594²⁷⁷⁷ - Fixing preprocessor constant used to enable C++11 chrono
- PR #1593²⁷⁷⁸ - Adding operator|() for hpx::launch
- Issue #1592²⁷⁷⁹ - Error (typo) in the docs

²⁷⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1616>

²⁷⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1615>

²⁷⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1614>

²⁷⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1613>

²⁷⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/1612>

²⁷⁶² <https://github.com/STELLAR-GROUP/hpx/pull/1611>

²⁷⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/1610>

²⁷⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1609>

²⁷⁶⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1608>

²⁷⁶⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1607>

²⁷⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1606>

²⁷⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1605>

²⁷⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1604>

²⁷⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1601>

²⁷⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1600>

²⁷⁷² <https://github.com/STELLAR-GROUP/hpx/pull/1599>

²⁷⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1598>

²⁷⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1597>

²⁷⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1596>

²⁷⁷⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1595>

²⁷⁷⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1594>

²⁷⁷⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1593>

²⁷⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1592>

- Issue #1590²⁷⁸⁰ - CMake fails when CMAKE_BINARY_DIR contains ‘+’.
- Issue #1589²⁷⁸¹ - Disconnecting a locality results in segfault using heartbeat example
- PR #1588²⁷⁸² - Fix doc string for config option HPX_WITH_EXAMPLES
- PR #1586²⁷⁸³ - Fixing 1493
- PR #1585²⁷⁸⁴ - Additional Check for Inspect Tool to detect Endline Whitespace
- Issue #1584²⁷⁸⁵ - Clean up coroutines implementation
- PR #1583²⁷⁸⁶ - Adding a check for end line whitespace
- PR #1582²⁷⁸⁷ - Attempt to fix assert firing after scheduling loop was exited
- PR #1581²⁷⁸⁸ - Fixed adjacentfind_binary test
- PR #1580²⁷⁸⁹ - Prevent some of the internal cmake lists from growing indefinitely
- PR #1579²⁷⁹⁰ - Removing type_size trait, replacing it with special archive type
- Issue #1578²⁷⁹¹ - Remove demangle_helper
- PR #1577²⁷⁹² - Get ptr problems
- Issue #1576²⁷⁹³ - Refactor async, dataflow, and future::then
- PR #1575²⁷⁹⁴ - Fixing tests for parallel rotate
- PR #1574²⁷⁹⁵ - Cleaning up schedulers
- PR #1573²⁷⁹⁶ - Fixing thread pool executor
- PR #1572²⁷⁹⁷ - Fixing number of configured localities
- PR #1571²⁷⁹⁸ - Reimplement decay
- PR #1570²⁷⁹⁹ - Refactoring async, apply, and dataflow APIs
- PR #1569²⁸⁰⁰ - Changed range for mach-o library lookup
- PR #1568²⁸⁰¹ - Mark decltype support as required
- PR #1567²⁸⁰² - Removed const from algorithms

2780 <https://github.com/STELLAR-GROUP/hpx/issues/1590>

2781 <https://github.com/STELLAR-GROUP/hpx/issues/1589>

2782 <https://github.com/STELLAR-GROUP/hpx/pull/1588>

2783 <https://github.com/STELLAR-GROUP/hpx/pull/1586>

2784 <https://github.com/STELLAR-GROUP/hpx/pull/1585>

2785 <https://github.com/STELLAR-GROUP/hpx/issues/1584>

2786 <https://github.com/STELLAR-GROUP/hpx/pull/1583>

2787 <https://github.com/STELLAR-GROUP/hpx/pull/1582>

2788 <https://github.com/STELLAR-GROUP/hpx/pull/1581>

2789 <https://github.com/STELLAR-GROUP/hpx/pull/1580>

2790 <https://github.com/STELLAR-GROUP/hpx/pull/1579>

2791 <https://github.com/STELLAR-GROUP/hpx/issues/1578>

2792 <https://github.com/STELLAR-GROUP/hpx/pull/1577>

2793 <https://github.com/STELLAR-GROUP/hpx/issues/1576>2794 <https://github.com/STELLAR-GROUP/hpx/pull/1575>2795 <https://github.com/STELLAR-GROUP/hpx/pull/1574>2796 <https://github.com/STELLAR-GROUP/hpx/pull/1573>2797 <https://github.com/STELLAR-GROUP/hpx/pull/1572>2798 <https://github.com/STELLAR-GROUP/hpx/pull/1571>2799 <https://github.com/STELLAR-GROUP/hpx/pull/1570>2800 <https://github.com/STELLAR-GROUP/hpx/pull/1569>2801 <https://github.com/STELLAR-GROUP/hpx/pull/1568>2802 <https://github.com/STELLAR-GROUP/hpx/pull/1567>

- Issue #1566²⁸⁰³ - CMAKE Configuration Test Failures for clang 3.5 on debian
- PR #1565²⁸⁰⁴ - Dylib support
- PR #1564²⁸⁰⁵ - Converted partitioners and some algorithms to use executors
- PR #1563²⁸⁰⁶ - Fix several #includes for Boost.Preprocessor
- PR #1562²⁸⁰⁷ - Adding configuration option disabling/enabling all message handlers
- PR #1561²⁸⁰⁸ - Removed all occurrences of boost::move replacing it with std::move
- Issue #1560²⁸⁰⁹ - Leftover HPX_REGISTER_ACTION_DECLARATION_2
- PR #1558²⁸¹⁰ - Revisit async/apply SFINAE conditions
- PR #1557²⁸¹¹ - Removing type_size trait, replacing it with special archive type
- PR #1556²⁸¹² - Executor algorithms
- PR #1555²⁸¹³ - Remove the necessity to specify archive flags on the receiving end
- PR #1554²⁸¹⁴ - Removing obsolete Boost.Serialization macros
- PR #1553²⁸¹⁵ - Properly fix HPX_DEFINE_*_ACTION macros
- PR #1552²⁸¹⁶ - Fixed algorithms relying on copy_if implementation
- PR #1551²⁸¹⁷ - Pxfs - Modifying FindOrangeFS.cmake based on OrangeFS 2.9.X
- Issue #1550²⁸¹⁸ - Passing plain identifier inside HPX_DEFINE_PLAIN_ACTION_1
- PR #1549²⁸¹⁹ - Fixing intel14/libstdc++4.4
- PR #1548²⁸²⁰ - Moving raw_ptr to detail namespace
- PR #1547²⁸²¹ - Adding support for executors to future.then
- PR #1546²⁸²² - Executor traits result types
- PR #1545²⁸²³ - Integrate executors with dataflow
- PR #1543²⁸²⁴ - Fix potential zero-copy for primarynamespace::bulk_service_async et.al.
- PR #1542²⁸²⁵ - Merging HPX0.9.10 into pxfs branch

²⁸⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1566>

²⁸⁰⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1565>

²⁸⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1564>

²⁸⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1563>

²⁸⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1562>

²⁸⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1561>

²⁸⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1560>

²⁸¹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1558>

²⁸¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1557>

²⁸¹² <https://github.com/STELLAR-GROUP/hpx/pull/1556>

²⁸¹³ <https://github.com/STELLAR-GROUP/hpx/pull/1555>

²⁸¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1554>

²⁸¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1553>

²⁸¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1552>

²⁸¹⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1551>

²⁸¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1550>

²⁸¹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1549>

²⁸²⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1548>

²⁸²¹ <https://github.com/STELLAR-GROUP/hpx/pull/1547>

²⁸²² <https://github.com/STELLAR-GROUP/hpx/pull/1546>

²⁸²³ <https://github.com/STELLAR-GROUP/hpx/pull/1545>

²⁸²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1543>

²⁸²⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1542>

- PR #1541²⁸²⁶ - Removed stale cmake tests, unused since the great cmake refactoring
- PR #1540²⁸²⁷ - Fix idle-rate on platforms without TSC
- PR #1539²⁸²⁸ - Reporting situation if zero-copy-serialization was performed by a parcel generated from a plain apply/async
- PR #1538²⁸²⁹ - Changed return type of bulk executors and added test
- Issue #1537²⁸³⁰ - Incorrect cpuid config tests
- PR #1536²⁸³¹ - Changed return type of bulk executors and added test
- PR #1535²⁸³² - Make sure promise::get_gid() can be called more than once
- PR #1534²⁸³³ - Fixed async_callback with bound callback
- PR #1533²⁸³⁴ - Updated the link in the documentation to a publically- accessible URL
- PR #1532²⁸³⁵ - Make sure sync primitives are not copyable nor movable
- PR #1531²⁸³⁶ - Fix unwrapped issue with future ranges of void type
- PR #1530²⁸³⁷ - Serialization complex
- Issue #1528²⁸³⁸ - Unwrapped issue with future<void>
- Issue #1527²⁸³⁹ - HPX does not build with Boost 1.58.0
- PR #1526²⁸⁴⁰ - Added support for boost.multi_array serialization
- PR #1525²⁸⁴¹ - Properly handle deferred futures, fixes #1506
- PR #1524²⁸⁴² - Making sure invalid action argument types generate clear error message
- Issue #1522²⁸⁴³ - Need serialization support for boost multi array
- Issue #1521²⁸⁴⁴ - Remote async and zero-copy serialization optimizations don't play well together
- PR #1520²⁸⁴⁵ - Fixing UB while registering polymorphic classes for serialization
- PR #1519²⁸⁴⁶ - Making detail::condition_variable safe to use
- PR #1518²⁸⁴⁷ - Fix when_some bug missing indices in its result
- Issue #1517²⁸⁴⁸ - Typo may affect CMake build system tests

²⁸²⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1541>

²⁸²⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1540>

²⁸²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1539>

²⁸²⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1538>

²⁸³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1537>

²⁸³¹ <https://github.com/STELLAR-GROUP/hpx/pull/1536>

²⁸³² <https://github.com/STELLAR-GROUP/hpx/pull/1535>

²⁸³³ <https://github.com/STELLAR-GROUP/hpx/pull/1534>

²⁸³⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1533>

²⁸³⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1532>

²⁸³⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1531>

²⁸³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1530>

²⁸³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1528>

²⁸³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1527>

²⁸⁴⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1526>

²⁸⁴¹ <https://github.com/STELLAR-GROUP/hpx/pull/1525>

²⁸⁴² <https://github.com/STELLAR-GROUP/hpx/pull/1524>

²⁸⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1522>

²⁸⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1521>

²⁸⁴⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1520>

²⁸⁴⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1519>

²⁸⁴⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1518>

²⁸⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1517>

- PR #1516²⁸⁴⁹ - Fixing Posix context
- PR #1515²⁸⁵⁰ - Fixing Posix context
- PR #1514²⁸⁵¹ - Correct problems with loading dynamic components
- PR #1513²⁸⁵² - Fixing intel glibc4 4
- Issue #1508²⁸⁵³ - memory and papi counters do not work
- Issue #1507²⁸⁵⁴ - Unrecognized Command Line Option Error causing exit status 0
- Issue #1506²⁸⁵⁵ - Properly handle deferred futures
- PR #1505²⁸⁵⁶ - Adding #include - would not compile without this
- Issue #1502²⁸⁵⁷ - boost::filesystem::exists throws unexpected exception
- Issue #1501²⁸⁵⁸ - hwloc configuration options are wrong for MIC
- PR #1504²⁸⁵⁹ - Making sure boost::filesystem::exists() does not throw
- PR #1500²⁸⁶⁰ - Exit application on --hpx:version/-v and --hpx:info
- PR #1498²⁸⁶¹ - Extended task block
- PR #1497²⁸⁶² - Unique ptr serialization
- PR #1496²⁸⁶³ - Unique ptr serialization (closed)
- PR #1495²⁸⁶⁴ - Switching circleci build type to debug
- Issue #1494²⁸⁶⁵ - --hpx:version/-v does not exit after printing version information
- Issue #1493²⁸⁶⁶ - add an hpx_ prefix to libraries and components to avoid name conflicts
- Issue #1492²⁸⁶⁷ - Define and ensure limitations for arguments to async/apply
- PR #1489²⁸⁶⁸ - Enable idle rate counter on demand
- PR #1488²⁸⁶⁹ - Made sure detail::condition_variable can be safely destroyed
- PR #1487²⁸⁷⁰ - Introduced default (main) template implementation for ignore_while_checking
- PR #1486²⁸⁷¹ - Add HPX inspect tool

²⁸⁴⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1516>

²⁸⁵⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1515>

²⁸⁵¹ <https://github.com/STELLAR-GROUP/hpx/pull/1514>

²⁸⁵² <https://github.com/STELLAR-GROUP/hpx/pull/1513>

²⁸⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1508>

²⁸⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1507>

²⁸⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1506>

²⁸⁵⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1505>

²⁸⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1502>

²⁸⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1501>

²⁸⁵⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1504>

²⁸⁶⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1500>

²⁸⁶¹ <https://github.com/STELLAR-GROUP/hpx/pull/1498>

²⁸⁶² <https://github.com/STELLAR-GROUP/hpx/pull/1497>

²⁸⁶³ <https://github.com/STELLAR-GROUP/hpx/pull/1496>

²⁸⁶⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1495>

²⁸⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1494>

²⁸⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1493>

²⁸⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1492>

²⁸⁶⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1489>

²⁸⁶⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1488>

²⁸⁷⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1487>

²⁸⁷¹ <https://github.com/STELLAR-GROUP/hpx/pull/1486>

- Issue #1485²⁸⁷² - ignore_while_locked doesn't support all Lockable types
- PR #1484²⁸⁷³ - Docker image generation
- PR #1483²⁸⁷⁴ - Move external endian library into HPX
- PR #1482²⁸⁷⁵ - Actions with integer type ids
- Issue #1481²⁸⁷⁶ - Sync primitives safe destruction
- Issue #1480²⁸⁷⁷ - Move external/boost/endian into hpx/util
- Issue #1478²⁸⁷⁸ - Boost inspect violations
- PR #1479²⁸⁷⁹ - Adds serialization for arrays; some further/minor fixes
- PR #1477²⁸⁸⁰ - Fixing problems with the Intel compiler using a GCC 4.4 std library
- PR #1476²⁸⁸¹ - Adding hpx::lcos::latch and hpx::lcos::local::latch
- Issue #1475²⁸⁸² - Boost inspect violations
- PR #1473²⁸⁸³ - Fixing action move tests
- Issue #1471²⁸⁸⁴ - Sync primitives should not be movable
- PR #1470²⁸⁸⁵ - Removing hpx::util::polymorphic_factory
- PR #1468²⁸⁸⁶ - Fixed container creation
- Issue #1467²⁸⁸⁷ - HPX application fail during finalization
- Issue #1466²⁸⁸⁸ - HPX doesn't pick up Torque's nodefile on SuperMIC
- Issue #1464²⁸⁸⁹ - HPX option for pre and post bootstrap performance counters
- PR #1463²⁸⁹⁰ - Replacing async_colocated(id, ...) with async(colocated(id), ...)
- PR #1462²⁸⁹¹ - Consolidated task_region with N4411
- PR #1461²⁸⁹² - Consolidate inconsistent CMake option names
- Issue #1460²⁸⁹³ - Which malloc is actually used? or at least which one is HPX built with
- Issue #1459²⁸⁹⁴ - Make cmake configure step fail explicitly if compiler version is not supported

²⁸⁷² <https://github.com/STELLAR-GROUP/hpx/issues/1485>

²⁸⁷³ <https://github.com/STELLAR-GROUP/hpx/pull/1484>

²⁸⁷⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1483>

²⁸⁷⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1482>

²⁸⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1481>

²⁸⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1480>

²⁸⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1478>

²⁸⁷⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1479>

²⁸⁸⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1477>

²⁸⁸¹ <https://github.com/STELLAR-GROUP/hpx/pull/1476>

²⁸⁸² <https://github.com/STELLAR-GROUP/hpx/issues/1475>

²⁸⁸³ <https://github.com/STELLAR-GROUP/hpx/pull/1473>

²⁸⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1471>

²⁸⁸⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1470>

²⁸⁸⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1468>

²⁸⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1467>

²⁸⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1466>

²⁸⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1464>

²⁸⁹⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1463>

²⁸⁹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1462>

²⁸⁹² <https://github.com/STELLAR-GROUP/hpx/pull/1461>

²⁸⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1460>

²⁸⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1459>

- Issue #1458²⁸⁹⁵ - Update parallel::task_region with N4411
- PR #1456²⁸⁹⁶ - Consolidating new_<>()
- Issue #1455²⁸⁹⁷ - Replace async_colocated(id, ...) with async(colocated(id), ...)
- PR #1454²⁸⁹⁸ - Removed harmful std::moves from return statements
- PR #1453²⁸⁹⁹ - Use range-based for-loop instead of Boost.Foreach
- PR #1452²⁹⁰⁰ - C++ feature tests
- PR #1451²⁹⁰¹ - When serializing, pass archive flags to traits::get_type_size
- Issue #1450²⁹⁰² - traits::get_type_size needs archive flags to enable zero_copy optimizations
- Issue #1449²⁹⁰³ - “couldn’t create performance counter” - AGAS
- Issue #1448²⁹⁰⁴ - Replace distributing factories with new_<T[]>(...)
- PR #1447²⁹⁰⁵ - Removing obsolete remote_object component
- PR #1446²⁹⁰⁶ - Hpx serialization
- PR #1445²⁹⁰⁷ - Replacing travis with circleci
- PR #1443²⁹⁰⁸ - Always stripping HPX command line arguments before executing start function
- PR #1442²⁹⁰⁹ - Adding –hpx:bind=none to disable thread affinities
- Issue #1439²⁹¹⁰ - Libraries get linked in multiple times, RPATH is not properly set
- PR #1438²⁹¹¹ - Removed superfluous typedefs
- Issue #1437²⁹¹² - hpx::init() should strip HPX-related flags from argv
- Issue #1436²⁹¹³ - Add strong scaling option to htts
- PR #1435²⁹¹⁴ - Adding async_cb, async_continue_cb, and async_colocated_cb
- PR #1434²⁹¹⁵ - Added missing install rule, removed some dead CMake code
- PR #1433²⁹¹⁶ - Add GitExternal and SubProject cmake scripts from eyescale/cmake repo
- Issue #1432²⁹¹⁷ - Add command line flag to disable thread pinning

²⁸⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1458>

²⁸⁹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1456>

²⁸⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1455>

²⁸⁹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1454>

²⁸⁹⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1453>

²⁹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1452>

²⁹⁰¹ <https://github.com/STELLAR-GROUP/hpx/pull/1451>

²⁹⁰² <https://github.com/STELLAR-GROUP/hpx/issues/1450>

²⁹⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1449>

²⁹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1448>

²⁹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1447>

²⁹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1446>

²⁹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1445>

²⁹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1443>

²⁹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/pull/1442>

²⁹¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1439>

²⁹¹¹ <https://github.com/STELLAR-GROUP/hpx/pull/1438>

²⁹¹² <https://github.com/STELLAR-GROUP/hpx/issues/1437>

²⁹¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1436>

²⁹¹⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1435>

²⁹¹⁵ <https://github.com/STELLAR-GROUP/hpx/pull/1434>

²⁹¹⁶ <https://github.com/STELLAR-GROUP/hpx/pull/1433>

²⁹¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1432>

- PR #1431²⁹¹⁸ - Fix #1423
- Issue #1430²⁹¹⁹ - Inconsistent CMake option names
- Issue #1429²⁹²⁰ - Configure setting HPX_HAVE_PARCELPORT_MPI is ignored
- PR #1428²⁹²¹ - Fixes #1419 (closed)
- PR #1427²⁹²² - Adding stencil_iterator and transform_iterator
- PR #1426²⁹²³ - Fixes #1419
- PR #1425²⁹²⁴ - During serialization memory allocation should honour allocator chunk size
- Issue #1424²⁹²⁵ - chunk allocation during serialization does not use memory pool/allocator chunk size
- Issue #1423²⁹²⁶ - Remove HPX_STD_UNIQUE_PTR
- Issue #1422²⁹²⁷ - hpx:threads=all allocates too many os threads
- PR #1420²⁹²⁸ - added .travis.yml
- Issue #1419²⁹²⁹ - Unify enums: hpx::runtime::state and hpx::state
- PR #1416²⁹³⁰ - Adding travis builder
- Issue #1414²⁹³¹ - Correct directory for dispatch_gcc46.hpp iteration
- Issue #1410²⁹³² - Set operation algorithms
- Issue #1389²⁹³³ - Parallel algorithms relying on scan partitioner break for small number of elements
- Issue #1325²⁹³⁴ - Exceptions thrown during parcel handling are not handled correctly
- Issue #1315²⁹³⁵ - Errors while running performance tests
- Issue #1309²⁹³⁶ - hpx::vector partitions are not easily extendable by applications
- PR #1300²⁹³⁷ - Added serialization/de-serialization to examples.tuplespace
- Issue #1251²⁹³⁸ - hpx::threads::get_thread_count doesn't consider pending threads
- Issue #1008²⁹³⁹ - Decrease in application performance overtime; occasional spikes of major slowdown
- Issue #1001²⁹⁴⁰ - Zero copy serialization raises assert

²⁹¹⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1431>

²⁹¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1430>

²⁹²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1429>

²⁹²¹ <https://github.com/STELLAR-GROUP/hpx/pull/1428>

²⁹²² <https://github.com/STELLAR-GROUP/hpx/pull/1427>

²⁹²³ <https://github.com/STELLAR-GROUP/hpx/pull/1426>

²⁹²⁴ <https://github.com/STELLAR-GROUP/hpx/pull/1425>

²⁹²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1424>

²⁹²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1423>

²⁹²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1422>

²⁹²⁸ <https://github.com/STELLAR-GROUP/hpx/pull/1420>

²⁹²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1419>

²⁹³⁰ <https://github.com/STELLAR-GROUP/hpx/pull/1416>

²⁹³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1414>

²⁹³² <https://github.com/STELLAR-GROUP/hpx/issues/1410>

²⁹³³ <https://github.com/STELLAR-GROUP/hpx/issues/1389>

²⁹³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1325>

²⁹³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1315>

²⁹³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1309>

²⁹³⁷ <https://github.com/STELLAR-GROUP/hpx/pull/1300>

²⁹³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1251>

²⁹³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1008>

²⁹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1001>

- Issue #721²⁹⁴¹ - Make HPX usable for Xeon Phi
- Issue #524²⁹⁴² - Extend scheduler to support threads which can't be stolen

2.11.9 HPX V0.9.10 (Mar 24, 2015)

General changes

This is the 12th official release of *HPX*. It coincides with the 7th anniversary of the first commit to our source code repository. Since then, we have seen over 12300 commits amounting to more than 220000 lines of C++ code.

The major focus of this release was to improve the reliability of large scale runs. We believe to have achieved this goal as we now can reliably run *HPX* applications on up to ~24k cores. We have also shown that *HPX* can be used with success for symmetric runs (applications using both, host cores and Intel Xeon/Phi coprocessors). This is a huge step forward in terms of the usability of *HPX*. The main focus of this work involved isolating the causes of the segmentation faults at start up and shut down. Many of these issues were discovered to be the result of the suspension of threads which hold locks.

A very important improvement introduced with this release is the refactoring of the code representing our parcel-port implementation. Parcel- ports can now be implemented by 3rd parties as independent plugins which are dynamically loaded at runtime (static linking of parcel-ports is also supported). This refactoring also includes a massive improvement of the performance of our existing parcel-ports. We were able to significantly reduce the networking latencies and to improve the available networking bandwidth. Please note that in this release we disabled the ibverbs and ipc parcel ports as those have not been ported to the new plugin system yet (see Issue #839²⁹⁴³).

Another corner stone of this release is our work towards a complete implementation of `__cpp11_n4104` (Working Draft, Technical Specification for C++ Extensions for Parallelism). This document defines a set of parallel algorithms to be added to the C++ standard library. We now have implemented about 75% of all specified parallel algorithms (see [link hpx.manual.parallel.parallel_algorithms Parallel Algorithms] for more details). We also implemented some extensions to `__cpp11_n4104` allowing to invoke all of the algorithms asynchronously.

This release adds a first implementation of `hpx::vector` which is a distributed data structure closely aligned to the functionality of `std::vector`. The difference is that `hpx::vector` stores the data in partitions where the partitions can be distributed over different localities. We started to work on allowing to use the parallel algorithms with `hpx::vector`. At this point we have implemented only a few of the parallel algorithms to support distributed data structures (like `hpx::vector`) for testing purposes (see Issue #1338²⁹⁴⁴ for a documentation of our progress).

Breaking changes

With this release we put a lot of effort into changing the code base to be more compatible to C++11. These changes have caused the following issues for backward compatibility:

- Move to Variadics- All of the API now uses variadic templates. However, this change required to modify the argument sequence for some of the exiting API functions (`hpx::async_continue`, `hpx::apply_continue`, `hpx::when_each`, `hpx::wait_each`, synchronous invocation of actions).
- Changes to Macros- We also removed the macros `HPX_STD_FUNCTION` and `HPX_STD_TUPLE`. This shouldn't affect any user code as we replaced `HPX_STD_FUNCTION` with `hpx::util::function_nonser` which was the default expansion used for this macro. All *HPX* API functions which expect a `hpx::util::function_nonser` (or a `hpx::util::unique_function_nonser`) can now be transparently called with a compatible

²⁹⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/721>

²⁹⁴² <https://github.com/STELLAR-GROUP/hpx/issues/524>

²⁹⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/839>

²⁹⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1338>

`std::function` instead. Similarly, `HPX_STD_TUPLE` was replaced by its default expansion as well: `hpx::util::tuple`.

- Changes to `hpx::unique_future`- `hpx::unique_future`, which was deprecated in the previous release for `hpx::future` is now completely removed from *HPX*. This completes the transition to a completely standards conforming implementation of `hpx::future`.
- Changes to Supported Compilers. Finally, in order to utilize more C++11 semantics, we have officially dropped support for GCC 4.4 and MSVC 2012. Please see our [Prerequisites](#) page for more details.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- Issue #1402²⁹⁴⁵ - Internal shared_future serialization copies
- Issue #1399²⁹⁴⁶ - Build takes unusually long time...
- Issue #1398²⁹⁴⁷ - Tests using the scan partitioner are broken on at least gcc 4.7 and intel compiler
- Issue #1397²⁹⁴⁸ - Completely remove `hpx::unique_future`
- Issue #1396²⁹⁴⁹ - Parallel scan algorithms with different initial values
- Issue #1395²⁹⁵⁰ - Race Condition - 1d_stencil_8 - SuperMIC
- Issue #1394²⁹⁵¹ - “suspending thread while at least one lock is being held” - 1d_stencil_8 - SuperMIC
- Issue #1393²⁹⁵² - SEGFAULT in 1d_stencil_8 on SuperMIC
- Issue #1392²⁹⁵³ - Fixing #1168
- Issue #1391²⁹⁵⁴ - Parallel Algorithms for scan partitioner for small number of elements
- Issue #1387²⁹⁵⁵ - Failure with more than 4 localities
- Issue #1386²⁹⁵⁶ - Dispatching unhandled exceptions to outer user code
- Issue #1385²⁹⁵⁷ - Adding Copy algorithms, fixing `parallel::copy_if`
- Issue #1384²⁹⁵⁸ - Fixing 1325
- Issue #1383²⁹⁵⁹ - Fixed #504: Refactor Dataflow LCO to work with futures, this removes the dataflow component as it is obsolete
- Issue #1382²⁹⁶⁰ - `is_sorted`, `is_sorted_until` and `is_partitioned` algorithms
- Issue #1381²⁹⁶¹ - fix for CMake versions prior to 3.1

²⁹⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1402>

²⁹⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1399>

²⁹⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1398>

²⁹⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1397>

²⁹⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1396>

²⁹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1395>

²⁹⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/1394>

²⁹⁵² <https://github.com/STELLAR-GROUP/hpx/issues/1393>

²⁹⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1392>

²⁹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1391>

²⁹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1387>

²⁹⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1386>

²⁹⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1385>

²⁹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1384>

²⁹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1383>

²⁹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1382>

²⁹⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/1381>

- Issue #1380²⁹⁶² - resolved warning in CMake 3.1 and newer
- Issue #1379²⁹⁶³ - Compilation error with papi
- Issue #1378²⁹⁶⁴ - Towards safer migration
- Issue #1377²⁹⁶⁵ - HPXConfig.cmake should include TCMALLOC_LIBRARY and TCMALLOC_INCLUDE_DIR
- Issue #1376²⁹⁶⁶ - Warning on uninitialized member
- Issue #1375²⁹⁶⁷ - Fixing 1163
- Issue #1374²⁹⁶⁸ - Fixing the MSVC 12 release builder
- Issue #1373²⁹⁶⁹ - Modifying parallel search algorithm for zero length searches
- Issue #1372²⁹⁷⁰ - Modifying parallel search algorithm for zero length searches
- Issue #1371²⁹⁷¹ - Avoid holding a lock during agas::inref while doing a credit split
- Issue #1370²⁹⁷² - --hpx:bind throws unexpected error
- Issue #1369²⁹⁷³ - Getting rid of (void) in loops
- Issue #1368²⁹⁷⁴ - Variadic templates support for tuple
- Issue #1367²⁹⁷⁵ - One last batch of variadic templates support
- Issue #1366²⁹⁷⁶ - Fixing symbolic namespace hang
- Issue #1365²⁹⁷⁷ - More held locks
- Issue #1364²⁹⁷⁸ - Add counters 1363
- Issue #1363²⁹⁷⁹ - Add thread overhead counters
- Issue #1362²⁹⁸⁰ - Std config removal
- Issue #1361²⁹⁸¹ - Parcelport plugins
- Issue #1360²⁹⁸² - Detuplify transfer_action
- Issue #1359²⁹⁸³ - Removed obsolete checks
- Issue #1358²⁹⁸⁴ - Fixing 1352

²⁹⁶² <https://github.com/STELLAR-GROUP/hpx/issues/1380>

²⁹⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/1379>

²⁹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1378>

²⁹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1377>

²⁹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1376>

²⁹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1375>

²⁹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1374>

²⁹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1373>

²⁹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1372>

²⁹⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/1371>

²⁹⁷² <https://github.com/STELLAR-GROUP/hpx/issues/1370>

²⁹⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/1369>

²⁹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1368>

²⁹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1367>

²⁹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1366>

²⁹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1365>

²⁹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1364>

²⁹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1363>

²⁹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1362>

²⁹⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/1361>

²⁹⁸² <https://github.com/STELLAR-GROUP/hpx/issues/1360>

²⁹⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/1359>

²⁹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1358>

- Issue #1357²⁹⁸⁵ - Variadic templates support for runtime_support and components
- Issue #1356²⁹⁸⁶ - fixed coordinate test for intel13
- Issue #1355²⁹⁸⁷ - fixed coordinate.hpp
- Issue #1354²⁹⁸⁸ - Lexicographical Compare completed
- Issue #1353²⁹⁸⁹ - HPX should set Boost_ADDITIONAL_VERSIONS flags
- Issue #1352²⁹⁹⁰ - Error: Cannot find action “ in type registry: HPX(bad_action_code)
- Issue #1351²⁹⁹¹ - Variadic templates support for applicers
- Issue #1350²⁹⁹² - Actions simplification
- Issue #1349²⁹⁹³ - Variadic when and wait functions
- Issue #1348²⁹⁹⁴ - Added hpx_init header to test files
- Issue #1347²⁹⁹⁵ - Another batch of variadic templates support
- Issue #1346²⁹⁹⁶ - Segmented copy
- Issue #1345²⁹⁹⁷ - Attempting to fix hangs during shutdown
- Issue #1344²⁹⁹⁸ - Std config removal
- Issue #1343²⁹⁹⁹ - Removing various distribution policies for hpx::vector
- Issue #1342³⁰⁰⁰ - Inclusive scan
- Issue #1341³⁰⁰¹ - Exclusive scan
- Issue #1340³⁰⁰² - Adding parallel::count for distributed data structures, adding tests
- Issue #1339³⁰⁰³ - Update argument order for transform_reduce
- Issue #1337³⁰⁰⁴ - Fix dataflow to handle properly ranges of futures
- Issue #1336³⁰⁰⁵ - dataflow needs to hold onto futures passed to it
- Issue #1335³⁰⁰⁶ - Fails to compile with msvc14
- Issue #1334³⁰⁰⁷ - Examples build problem

²⁹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1357>

²⁹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1356>

²⁹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1355>

²⁹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1354>

²⁹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1353>

²⁹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1352>

²⁹⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1351>

²⁹⁹² <https://github.com/STELLAR-GROUP/hpx/issues/1350>

²⁹⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1349>

²⁹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1348>

²⁹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1347>

²⁹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1346>

²⁹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1345>

²⁹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1344>

²⁹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1343>

³⁰⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1342>

³⁰⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/1341>

³⁰⁰² <https://github.com/STELLAR-GROUP/hpx/issues/1340>

³⁰⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1339>

³⁰⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1337>

³⁰⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1336>

³⁰⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1335>

³⁰⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1334>

- Issue #1333³⁰⁰⁸ - Distributed transform reduce
- Issue #1332³⁰⁰⁹ - Variadic templates support for actions
- Issue #1331³⁰¹⁰ - Some ambiguous calls of map::erase have been prevented by adding additional check in locality constructor.
- Issue #1330³⁰¹¹ - Defining Plain Actions does not work as described in the documentation
- Issue #1329³⁰¹² - Distributed vector cleanup
- Issue #1328³⁰¹³ - Sync docs and comments with code in hello_world example
- Issue #1327³⁰¹⁴ - Typos in docs
- Issue #1326³⁰¹⁵ - Documentation and code diverged in Fibonacci tutorial
- Issue #1325³⁰¹⁶ - Exceptions thrown during parcel handling are not handled correctly
- Issue #1324³⁰¹⁷ - fixed bandwidth calculation
- Issue #1323³⁰¹⁸ - mmap() failed to allocate thread stack due to insufficient resources
- Issue #1322³⁰¹⁹ - HPX fails to build aa182cf
- Issue #1321³⁰²⁰ - Limiting size of outgoing messages while coalescing parcels
- Issue #1320³⁰²¹ - passing a future with launch::deferred in remote function call causes hang
- Issue #1319³⁰²² - An exception when tries to specify number high priority threads with abp-priority
- Issue #1318³⁰²³ - Unable to run program with abp-priority and numa-sensitivity enabled
- Issue #1317³⁰²⁴ - N4071 Search/Search_n finished, minor changes
- Issue #1316³⁰²⁵ - Add config option to make -Ihpix.run_hpx_main!=1 the default
- Issue #1314³⁰²⁶ - Variadic support for async and apply
- Issue #1313³⁰²⁷ - Adjust when_any/some to the latest proposed interfaces
- Issue #1312³⁰²⁸ - Fixing #857: hpx::naming::locality leaks parcelport specific information into the public interface
- Issue #1311³⁰²⁹ - Distributed get'er/set'er_values for distributed vector

³⁰⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1333>

³⁰⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1332>

³⁰¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1331>

³⁰¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1330>

³⁰¹² <https://github.com/STELLAR-GROUP/hpx/issues/1329>

³⁰¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1328>

³⁰¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1327>

³⁰¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1326>

³⁰¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1325>

³⁰¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1324>

³⁰¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1323>

³⁰¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1322>

³⁰²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1321>

³⁰²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1320>

³⁰²² <https://github.com/STELLAR-GROUP/hpx/issues/1319>

³⁰²³ <https://github.com/STELLAR-GROUP/hpx/issues/1318>

³⁰²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1317>

³⁰²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1316>

³⁰²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1314>

³⁰²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1313>

³⁰²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1312>

³⁰²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1311>

- Issue #1310³⁰³⁰ - Crashing in hpx::parcelset::policies::mpi::connection_handler::handle_messages() on Super-MIC
- Issue #1308³⁰³¹ - Unable to execute an application with –hpx:threads
- Issue #1307³⁰³² - merge_graph linking issue
- Issue #1306³⁰³³ - First batch of variadic templates support
- Issue #1305³⁰³⁴ - Create a compiler wrapper
- Issue #1304³⁰³⁵ - Provide a compiler wrapper for hpx
- Issue #1303³⁰³⁶ - Drop support for GCC44
- Issue #1302³⁰³⁷ - Fixing #1297
- Issue #1301³⁰³⁸ - Compilation error when tried to use boost range iterators with wait_all
- Issue #1298³⁰³⁹ - Distributed vector
- Issue #1297³⁰⁴⁰ - Unable to invoke component actions recursively
- Issue #1294³⁰⁴¹ - HDF5 build error
- Issue #1275³⁰⁴² - The parcelport implementation is non-optimal
- Issue #1267³⁰⁴³ - Added classes and unit tests for local_file, orangefs_file and pxfs_file
- Issue #1264³⁰⁴⁴ - Error “assertion ‘!m_fun’ failed” randomly occurs when using TCP
- Issue #1254³⁰⁴⁵ - thread binding seems to not work properly
- Issue #1220³⁰⁴⁶ - parallel::copy_if is broken
- Issue #1217³⁰⁴⁷ - Find a better way of fixing the issue patched by #1216
- Issue #1168³⁰⁴⁸ - Starting HPX on Cray machines using aprun isn’t working correctly
- Issue #1085³⁰⁴⁹ - Replace startup and shutdown barriers with broadcasts
- Issue #981³⁰⁵⁰ - With SLURM, –hpx:threads=8 should not be necessary
- Issue #857³⁰⁵¹ - hpx::naming::locality leaks parcelport specific information into the public interface
- Issue #850³⁰⁵² - “flush” not documented

³⁰³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1310>

³⁰³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1308>

³⁰³² <https://github.com/STELLAR-GROUP/hpx/issues/1307>

³⁰³³ <https://github.com/STELLAR-GROUP/hpx/issues/1306>

³⁰³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1305>

³⁰³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1304>

³⁰³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1303>

³⁰³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1302>

³⁰³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1301>

³⁰³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1298>

³⁰⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1297>

³⁰⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/1294>

³⁰⁴² <https://github.com/STELLAR-GROUP/hpx/issues/1275>

³⁰⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1267>

³⁰⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1264>

³⁰⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1254>

³⁰⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1220>

³⁰⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1217>

³⁰⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1168>

³⁰⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1085>

³⁰⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/981>

³⁰⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/857>

³⁰⁵² <https://github.com/STELLAR-GROUP/hpx/issues/850>

- Issue #763³⁰⁵³ - Create buildbot instance that uses std::bind as HPX_STD_BIND
- Issue #680³⁰⁵⁴ - Convert parcel ports into a plugin system
- Issue #582³⁰⁵⁵ - Make exception thrown from HPX threads available from hpx::init
- Issue #504³⁰⁵⁶ - Refactor Dataflow LCO to work with futures
- Issue #196³⁰⁵⁷ - Don't store copies of the locality network metadata in the gva table

2.11.10 HPX V0.9.9 (Oct 31, 2014, codename Spooky)

General changes

We have had over 1500 commits since the last release and we have closed over 200 tickets (bugs, feature requests, pull requests, etc.). These are by far the largest numbers of commits and resolved issues for any of the *HPX* releases so far. We are especially happy about the large number of people who contributed for the first time to *HPX*.

- We completed the transition from the older (non-conforming) implementation of hpx::future to the new and fully conforming version by removing the old code and by renaming the type hpx::unique_future to hpx::future. In order to maintain backwards compatibility with existing code which uses the type hpx::unique_future we support the configuration variable `HPX_UNIQUE_FUTURE_ALIAS`. If this variable is set to ON while running cmake it will additionally define a template alias for this type.
- We rewrote and significantly changed our build system. Please have a look at the new (now generated) documentation here: [HPX build system](#). Please revisit your build scripts to adapt to the changes. The most notable changes are:
 - `HPX_NO_INSTALL` is no longer necessary.
 - For external builds, you need to set `HPX_DIR` instead of `HPX_ROOT` as described here: [Using HPX with CMake-based projects](#).
 - IDEs that support multiple configurations (Visual Studio and XCode) can now be used as intended. that means no build dir.
 - Building HPX statically (without dynamic libraries) is now supported (`-DHPX_STATIC_LINKING=On`).
 - Please note that many variables used to configure the build process have been renamed to unify the naming conventions (see the section [CMake variables used to configure HPX](#) for more information).
 - This also fixes a long list of issues, for more information see [Issue #1204³⁰⁵⁸](#).
- We started to implement various proposals to the C++ Standardization committee related to parallelism and concurrency, most notably [N4409³⁰⁵⁹](#) (Working Draft, Technical Specification for C++ Extensions for Parallelism), [N4411³⁰⁶⁰](#) (Task Region Rev. 3), and [N4313³⁰⁶¹](#) (Working Draft, Technical Specification for C++ Extensions for Concurrency).

³⁰⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/763>

³⁰⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/680>

³⁰⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/582>

³⁰⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/504>

³⁰⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/196>

³⁰⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1204>

³⁰⁵⁹ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4409.pdf>

³⁰⁶⁰ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/n4411.pdf>

³⁰⁶¹ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

- We completely remodeled our automatic build system to run builds and unit tests on various systems and compilers. This allows us to find most bugs right as they were introduced and helps to maintain a high level of quality and compatibility. The newest build logs can be found at [HPX Buildbot Website³⁰⁶²](#).

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- Issue #1296³⁰⁶³ - Rename make_error_future to make_exceptional_future, adjust to N4123
- Issue #1295³⁰⁶⁴ - building issue
- Issue #1293³⁰⁶⁵ - Transpose example
- Issue #1292³⁰⁶⁶ - Wrong abs() function used in example
- Issue #1291³⁰⁶⁷ - non-synchronized shift operators have been removed
- Issue #1290³⁰⁶⁸ - RDTSCP is defined as true for Xeon Phi build
- Issue #1289³⁰⁶⁹ - Fixing 1288
- Issue #1288³⁰⁷⁰ - Add new performance counters
- Issue #1287³⁰⁷¹ - Hierarchy scheduler broken performance counters
- Issue #1286³⁰⁷² - Algorithm cleanup
- Issue #1285³⁰⁷³ - Broken Links in Documentation
- Issue #1284³⁰⁷⁴ - Uninitialized copy
- Issue #1283³⁰⁷⁵ - missing boost::scoped_ptr includes
- Issue #1282³⁰⁷⁶ - Update documentation of build options for schedulers
- Issue #1281³⁰⁷⁷ - reset idle rate counter
- Issue #1280³⁰⁷⁸ - Bug when executing on Intel MIC
- Issue #1279³⁰⁷⁹ - Add improved when_all/wait_all
- Issue #1278³⁰⁸⁰ - Implement improved when_all/wait_all
- Issue #1277³⁰⁸¹ - feature request: get access to argc argv and variables_map

³⁰⁶² <http://rostam.cct.lsu.edu/>

³⁰⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/1296>

³⁰⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1295>

³⁰⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1293>

³⁰⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1292>

³⁰⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1291>

³⁰⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1290>

³⁰⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1289>

³⁰⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1288>

³⁰⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/1287>

³⁰⁷² <https://github.com/STELLAR-GROUP/hpx/issues/1286>

³⁰⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/1285>

³⁰⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1284>

³⁰⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1283>

³⁰⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1282>

³⁰⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1281>

³⁰⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1280>

³⁰⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1279>

³⁰⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1278>

³⁰⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/1277>

- Issue #1276³⁰⁸² - Remove merging map
- Issue #1274³⁰⁸³ - Weird (wrong) string code in papi.cpp
- Issue #1273³⁰⁸⁴ - Sequential task execution policy
- Issue #1272³⁰⁸⁵ - Avoid CMake name clash for Boost.Thread library
- Issue #1271³⁰⁸⁶ - Updates on HPX Test Units
- Issue #1270³⁰⁸⁷ - hpx/util/safe_lexical_cast.hpp is added
- Issue #1269³⁰⁸⁸ - Added default value for “LIB” cmake variable
- Issue #1268³⁰⁸⁹ - Memory Counters not working
- Issue #1266³⁰⁹⁰ - FindHPX.cmake is not installed
- Issue #1263³⁰⁹¹ - apply_remote test takes too long
- Issue #1262³⁰⁹² - Chrono cleanup
- Issue #1261³⁰⁹³ - Need make install for papi counters and this builds all the examples
- Issue #1260³⁰⁹⁴ - Documentation of Stencil example claims
- Issue #1259³⁰⁹⁵ - Avoid double-linking Boost on Windows
- Issue #1257³⁰⁹⁶ - Adding additional parameter to create_thread
- Issue #1256³⁰⁹⁷ - added buildbot changes to release notes
- Issue #1255³⁰⁹⁸ - Cannot build MiniGhost
- Issue #1253³⁰⁹⁹ - hpx::thread defects
- Issue #1252³¹⁰⁰ - HPX_PREFIX is too fragile
- Issue #1250³¹⁰¹ - switch_to_fiber_emulation does not work properly
- Issue #1249³¹⁰² - Documentation is generated under Release folder
- Issue #1248³¹⁰³ - Fix usage of hpx_generic_coroutine_context and get tests passing on powerpc
- Issue #1247³¹⁰⁴ - Dynamic linking error

³⁰⁸² <https://github.com/STELLAR-GROUP/hpx/issues/1276>

³⁰⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/1274>

³⁰⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1273>

³⁰⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1272>

³⁰⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1271>

³⁰⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1270>

³⁰⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1269>

³⁰⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1268>

³⁰⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1266>

³⁰⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1263>

³⁰⁹² <https://github.com/STELLAR-GROUP/hpx/issues/1262>

³⁰⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1261>

³⁰⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1260>

³⁰⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1259>

³⁰⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1257>

³⁰⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1256>

³⁰⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1255>

³⁰⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1253>

³¹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1252>

³¹⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/1250>

³¹⁰² <https://github.com/STELLAR-GROUP/hpx/issues/1249>

³¹⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1248>

³¹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1247>

- Issue #1246³¹⁰⁵ - Make cpuid.cpp C++11 compliant
- Issue #1245³¹⁰⁶ - HPX fails on startup (setting thread affinity mask)
- Issue #1244³¹⁰⁷ - HPX_WITH_RDTSC configure test fails, but should succeed
- Issue #1243³¹⁰⁸ - CTest dashboard info for CSCS CDash drop location
- Issue #1242³¹⁰⁹ - Mac fixes
- Issue #1241³¹¹⁰ - Failure in Distributed with Boost 1.56
- Issue #1240³¹¹¹ - fix a race condition in examples.diskperf
- Issue #1239³¹¹² - fix wait_each in examples.diskperf
- Issue #1238³¹¹³ - Fixed #1237: hpx::util::portable_binary_iarchive failed
- Issue #1237³¹¹⁴ - hpx::util::portable_binary_iarchive faileds
- Issue #1235³¹¹⁵ - Fixing clang warnings and errors
- Issue #1234³¹¹⁶ - TCP runs fail: Transport endpoint is not connected
- Issue #1233³¹¹⁷ - Making sure the correct number of threads is registered with AGAS
- Issue #1232³¹¹⁸ - Fixing race in wait_xxx
- Issue #1231³¹¹⁹ - Parallel minmax
- Issue #1230³¹²⁰ - Distributed run of 1d_stencil_8 uses less threads than spec. & sometimes gives errors
- Issue #1229³¹²¹ - Unstable number of threads
- Issue #1228³¹²² - HPX link error (cmake / MPI)
- Issue #1226³¹²³ - Warning about struct/class thread_counters
- Issue #1225³¹²⁴ - Adding parallel::replace etc
- Issue #1224³¹²⁵ - Extending dataflow to pass through non-future arguments
- Issue #1223³¹²⁶ - Remaining find algorithms implemented, N4071
- Issue #1222³¹²⁷ - Merging all the changes

³¹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1246>

³¹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1245>

³¹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1244>

³¹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1243>

³¹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1242>

³¹¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1241>

³¹¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1240>

³¹¹² <https://github.com/STELLAR-GROUP/hpx/issues/1239>

³¹¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1238>

³¹¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1237>

³¹¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1235>

³¹¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1234>

³¹¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1233>

³¹¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1232>

³¹¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1231>

³¹²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1230>

³¹²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1229>

³¹²² <https://github.com/STELLAR-GROUP/hpx/issues/1228>

³¹²³ <https://github.com/STELLAR-GROUP/hpx/issues/1226>

³¹²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1225>

³¹²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1224>

³¹²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1223>

³¹²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1222>

- Issue #1221³¹²⁸ - No error output when using mpirun with hpx
- Issue #1219³¹²⁹ - Adding new AGAS cache performance counters
- Issue #1216³¹³⁰ - Fixing using futures (clients) as arguments to actions
- Issue #1215³¹³¹ - Error compiling simple component
- Issue #1214³¹³² - Stencil docs
- Issue #1213³¹³³ - Using more than a few dozen MPI processes on SuperMike results in a seg fault before getting to hpx_main
- Issue #1212³¹³⁴ - Parallel rotate
- Issue #1211³¹³⁵ - Direct actions cause the future's shared_state to be leaked
- Issue #1210³¹³⁶ - Refactored local::promise to be standard conformant
- Issue #1209³¹³⁷ - Improve command line handling
- Issue #1208³¹³⁸ - Adding parallel::reverse and parallel::reverse_copy
- Issue #1207³¹³⁹ - Add copy_backward and move_backward
- Issue #1206³¹⁴⁰ - N4071 additional algorithms implemented
- Issue #1204³¹⁴¹ - Cmake simplification and various other minor changes
- Issue #1203³¹⁴² - Implementing new launch policy for (local) async: hpx::launch::fork.
- Issue #1202³¹⁴³ - Failed assertion in connection_cache.hpp
- Issue #1201³¹⁴⁴ - pkg-config doesn't add mpi link directories
- Issue #1200³¹⁴⁵ - Error when querying time performance counters
- Issue #1199³¹⁴⁶ - library path is now configurable (again)
- Issue #1198³¹⁴⁷ - Error when querying performance counters
- Issue #1197³¹⁴⁸ - tests fail with intel compiler
- Issue #1196³¹⁴⁹ - Silence several warnings
- Issue #1195³¹⁵⁰ - Rephrase initializers to work with VC++ 2012

³¹²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1221>

³¹²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1219>

³¹³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1216>

³¹³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1215>

³¹³² <https://github.com/STELLAR-GROUP/hpx/issues/1214>

³¹³³ <https://github.com/STELLAR-GROUP/hpx/issues/1213>

³¹³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1212>

³¹³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1211>

³¹³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1210>

³¹³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1209>

³¹³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1208>

³¹³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1207>

³¹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1206>

³¹⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/1204>

³¹⁴² <https://github.com/STELLAR-GROUP/hpx/issues/1203>

³¹⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1202>

³¹⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1201>

³¹⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1200>

³¹⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1199>

³¹⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1198>

³¹⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1197>

³¹⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1196>

³¹⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1195>

- Issue #1194³¹⁵¹ - Simplify parallel algorithms
- Issue #1193³¹⁵² - Adding parallel::equal
- Issue #1192³¹⁵³ - HPX(out_of_memory) on including <hpx/hpx.hpp>
- Issue #1191³¹⁵⁴ - Fixing #1189
- Issue #1190³¹⁵⁵ - Chrono cleanup
- Issue #1189³¹⁵⁶ - Deadlock .. somewhere? (probably serialization)
- Issue #1188³¹⁵⁷ - Removed future::get_status()
- Issue #1186³¹⁵⁸ - Fixed FindOpenCL to find current AMD APP SDK
- Issue #1184³¹⁵⁹ - Tweaking future unwrapping
- Issue #1183³¹⁶⁰ - Extended parallel::reduce
- Issue #1182³¹⁶¹ - future::unwrap hangs for launch::deferred
- Issue #1181³¹⁶² - Adding all_of, any_of, and none_of and corresponding documentation
- Issue #1180³¹⁶³ - hpx::cout defect
- Issue #1179³¹⁶⁴ - hpx::async does not work for member function pointers when called on types with self-defined unary operator*
- Issue #1178³¹⁶⁵ - Implemented variadic hpx::util::zip_iterator
- Issue #1177³¹⁶⁶ - MPI parcelport defect
- Issue #1176³¹⁶⁷ - HPX_DEFINE_COMPONENT_CONST_ACTION_TPL does not have a 2-argument version
- Issue #1175³¹⁶⁸ - Create util::zip_iterator working with util::tuple<>
- Issue #1174³¹⁶⁹ - Error Building HPX on linux, root_certificate_authority.cpp
- Issue #1173³¹⁷⁰ - hpx::cout output lost
- Issue #1172³¹⁷¹ - HPX build error with Clang 3.4.2
- Issue #1171³¹⁷² - CMAKE_INSTALL_PREFIX ignored
- Issue #1170³¹⁷³ - Close hpx_benchmarks repository on Github

³¹⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/1194>

³¹⁵² <https://github.com/STELLAR-GROUP/hpx/issues/1193>

³¹⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1192>

³¹⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1191>

³¹⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1190>

³¹⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1189>

³¹⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1188>

³¹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1186>

³¹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1184>

³¹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1183>

³¹⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/1182>

³¹⁶² <https://github.com/STELLAR-GROUP/hpx/issues/1181>

³¹⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/1180>

³¹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1179>

³¹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1178>

³¹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1177>

³¹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1176>

³¹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1175>

³¹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1174>

³¹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1173>

³¹⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/1172>

³¹⁷² <https://github.com/STELLAR-GROUP/hpx/issues/1171>

³¹⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/1170>

- Issue #1169³¹⁷⁴ - Buildbot emails have syntax error in url
- Issue #1167³¹⁷⁵ - Merge partial implementation of standards proposal N3960
- Issue #1166³¹⁷⁶ - Fixed several compiler warnings
- Issue #1165³¹⁷⁷ - cmake warns: “tests.regressions.actions” does not exist
- Issue #1164³¹⁷⁸ - Want my own serialization of hpx::future
- Issue #1162³¹⁷⁹ - Segfault in hello_world example
- Issue #1161³¹⁸⁰ - Use HPX_ASSERT to aid the compiler
- Issue #1160³¹⁸¹ - Do not put -DNDEBUG into hpx_application.pc
- Issue #1159³¹⁸² - Support Clang 3.4.2
- Issue #1158³¹⁸³ - Fixed #1157: Rename when_n/wait_n, add when_xxx_n/wait_xxx_n
- Issue #1157³¹⁸⁴ - Rename when_n/wait_n, add when_xxx_n/wait_xxx_n
- Issue #1156³¹⁸⁵ - Force inlining fails
- Issue #1155³¹⁸⁶ - changed header of printout to be compatible with python csv module
- Issue #1154³¹⁸⁷ - Fixing iostreams
- Issue #1153³¹⁸⁸ - Standard manipulators (like std::endl) do not work with hpx::ostream
- Issue #1152³¹⁸⁹ - Functions revamp
- Issue #1151³¹⁹⁰ - Suppressing cmake 3.0 policy warning for CMP0026
- Issue #1150³¹⁹¹ - Client Serialization error
- Issue #1149³¹⁹² - Segfault on Stampede
- Issue #1148³¹⁹³ - Refactoring mini-ghost
- Issue #1147³¹⁹⁴ - N3960 copy_if and copy_n implemented and tested
- Issue #1146³¹⁹⁵ - Stencil print
- Issue #1145³¹⁹⁶ - N3960 hpx::parallel::copy implemented and tested

³¹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1169>

³¹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1167>

³¹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1166>

³¹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1165>

³¹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1164>

³¹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1162>

³¹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1161>

³¹⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/1160>

³¹⁸² <https://github.com/STELLAR-GROUP/hpx/issues/1159>

³¹⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/1158>

³¹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1157>

³¹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1156>

³¹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1155>

³¹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1154>

³¹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1153>

³¹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1152>

³¹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1151>

³¹⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1150>

³¹⁹² <https://github.com/STELLAR-GROUP/hpx/issues/1149>

³¹⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1148>

³¹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1147>

³¹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1146>

³¹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1145>

- Issue #1144³¹⁹⁷ - OpenMP examples 1d_stencil do not build
- Issue #1143³¹⁹⁸ - 1d_stencil OpenMP examples do not build
- Issue #1142³¹⁹⁹ - Cannot build HPX with gcc 4.6 on OS X
- Issue #1140³²⁰⁰ - Fix OpenMP lookup, enable usage of config tests in external CMake projects.
- Issue #1139³²⁰¹ - hpx/hpx/config/compiler_specific.hpp
- Issue #1138³²⁰² - clean up pkg-config files
- Issue #1137³²⁰³ - Improvements to create binary packages
- Issue #1136³²⁰⁴ - HPX_GCC_VERSION not defined on all compilers
- Issue #1135³²⁰⁵ - Avoiding collision between winsock2.h and windows.h
- Issue #1134³²⁰⁶ - Making sure, that hpx::finalize can be called from any locality
- Issue #1133³²⁰⁷ - 1d stencil examples
- Issue #1131³²⁰⁸ - Refactor unique_function implementation
- Issue #1130³²⁰⁹ - Unique function
- Issue #1129³²¹⁰ - Some fixes to the Build system on OS X
- Issue #1128³²¹¹ - Action future args
- Issue #1127³²¹² - Executor causes segmentation fault
- Issue #1124³²¹³ - Adding new API functions: register_id_with_basename, unregister_id_with_basename, find_ids_from_basename; adding test
- Issue #1123³²¹⁴ - Reduce nesting of try-catch construct in encode_parcels?
- Issue #1122³²¹⁵ - Client base fixes
- Issue #1121³²¹⁶ - Update hpxrun.py.in
- Issue #1120³²¹⁷ - HTTS2 tests compile errors on v110 (VS2012)
- Issue #1119³²¹⁸ - Remove references to boost::atomic in accumulator example
- Issue #1118³²¹⁹ - Only build test thread_pool_executor_1114_test if HPX_LOCAL_SCHEDULER is set

³¹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1144>

³¹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1143>

³¹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1142>

³²⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1140>

³²⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/1139>

³²⁰² <https://github.com/STELLAR-GROUP/hpx/issues/1138>

³²⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1137>

³²⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1136>

³²⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1135>

³²⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1134>

³²⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1133>

³²⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1131>

³²⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1130>

³²¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1129>

³²¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1128>

³²¹² <https://github.com/STELLAR-GROUP/hpx/issues/1127>

³²¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1124>

³²¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1123>

³²¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1122>

³²¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1121>

³²¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1120>

³²¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1119>

³²¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1118>

- Issue #1117³²²⁰ - local_queue_executor linker error on vc110
- Issue #1116³²²¹ - Disabled performance counter should give runtime errors, not invalid data
- Issue #1115³²²² - Compile error with Intel C++ 13.1
- Issue #1114³²²³ - Default constructed executor is not usable
- Issue #1113³²²⁴ - Fast compilation of logging causes ABI incompatibilities between different NDEBUG values
- Issue #1112³²²⁵ - Using thread_pool_executors causes segfault
- Issue #1111³²²⁶ - hpx::threads::get_thread_data always returns zero
- Issue #1110³²²⁷ - Remove unnecessary null pointer checks
- Issue #1109³²²⁸ - More tests adjustments
- Issue #1108³²²⁹ - Clarify build rules for “libboost_atomic-mt.so”?
- Issue #1107³²³⁰ - Remove unnecessary null pointer checks
- Issue #1106³²³¹ - network_storage benchmark improvements, adding legends to plots and tidying layout
- Issue #1105³²³² - Add more plot outputs and improve instructions doc
- Issue #1104³²³³ - Complete quoting for parameters of some CMake commands
- Issue #1103³²³⁴ - Work on test/scripts
- Issue #1102³²³⁵ - Changed minimum requirement of window install to 2012
- Issue #1101³²³⁶ - Changed minimum requirement of window install to 2012
- Issue #1100³²³⁷ - Changed readme to no longer specify using MSVC 2010 compiler
- Issue #1099³²³⁸ - Error returning futures from component actions
- Issue #1098³²³⁹ - Improve storage test
- Issue #1097³²⁴⁰ - data_actions quickstart example calls missing function decorate_action of data_get_action
- Issue #1096³²⁴¹ - MPI parcelport broken with new zero copy optimization
- Issue #1095³²⁴² - Warning C4005: _WIN32_WINNT: Macro redefinition

³²²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1117>

³²²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1116>

³²²² <https://github.com/STELLAR-GROUP/hpx/issues/1115>

³²²³ <https://github.com/STELLAR-GROUP/hpx/issues/1114>

³²²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1113>

³²²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1112>

3226 <https://github.com/STELLAR-GROUP/hpx/issues/1111>3227 <https://github.com/STELLAR-GROUP/hpx/issues/1110>3228 <https://github.com/STELLAR-GROUP/hpx/issues/1109>

3229 <https://github.com/STELLAR-GROUP/hpx/issues/1108>

3230 <https://github.com/STELLAR-GROUP/hpx/issues/1107>

3231 <https://github.com/STELLAR-GROUP/hpx/issues/1106>

3232 <https://github.com/STELLAR-GROUP/hpx/issues/1105>

3233 <https://github.com/STELLAR-GROUP/hpx/issues/1104>

3234 <https://github.com/STELLAR-GROUP/hpx/issues/1103>

3235 <https://github.com/STELLAR-GROUP/hpx/issues/1102>

3236 <https://github.com/STELLAR-GROUP/hpx/issues/1101>

3237 <https://github.com/STELLAR-GROUP/hpx/issues/1100>

3238 <https://github.com/STELLAR-GROUP/hpx/issues/1099>

3239 <https://github.com/STELLAR-GROUP/hpx/issues/1098>

3240 <https://github.com/STELLAR-GROUP/hpx/issues/1097>

3241 <https://github.com/STELLAR-GROUP/hpx/issues/1096>

3242 <https://github.com/STELLAR-GROUP/hpx/issues/1095>

2.11. Releases

1075

- Issue #1094³²⁴³ - Syntax error for -DHPX_UNIQUE_FUTURE_ALIAS in master
- Issue #1093³²⁴⁴ - Syntax error for -DHPX_UNIQUE_FUTURE_ALIAS
- Issue #1092³²⁴⁵ - Rename unique_future<> back to future<>
- Issue #1091³²⁴⁶ - Inconsistent error message
- Issue #1090³²⁴⁷ - On windows 8.1 the examples crashed if using more than one os thread
- Issue #1089³²⁴⁸ - Components should be allowed to have their own executor
- Issue #1088³²⁴⁹ - Add possibility to select a network interface for the ibverbs parcelport
- Issue #1087³²⁵⁰ - ibverbs and ipc parcelport uses zero copy optimization
- Issue #1083³²⁵¹ - Make shell examples copyable in docs
- Issue #1082³²⁵² - Implement proper termination detection during shutdown
- Issue #1081³²⁵³ - Implement thread_specific_ptr for hpx::threads
- Issue #1072³²⁵⁴ - make install not working properly
- Issue #1070³²⁵⁵ - Complete quoting for parameters of some CMake commands
- Issue #1059³²⁵⁶ - Fix more unused variable warnings
- Issue #1051³²⁵⁷ - Implement when_each
- Issue #973³²⁵⁸ - Would like option to report hwloc bindings
- Issue #970³²⁵⁹ - Bad flags for Fortran compiler
- Issue #941³²⁶⁰ - Create a proper user level context switching class for BG/Q
- Issue #935³²⁶¹ - Build error with gcc 4.6 and Boost 1.54.0 on hpx trunk and 0.9.6
- Issue #934³²⁶² - Want to build HPX without dynamic libraries
- Issue #927³²⁶³ - Make hpx/lcos/reduce.hpp accept futures of id_type
- Issue #926³²⁶⁴ - All unit tests that are run with more than one thread with CTest/hpx_run_test should configure hpx.os_threads
- Issue #925³²⁶⁵ - regression_dataflow_791 needs to be brought in line with HPX standards

³²⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1094>

³²⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1093>

³²⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1092>

³²⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1091>

³²⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1090>

³²⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1089>

³²⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1088>

³²⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1087>

³²⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/1083>

³²⁵² <https://github.com/STELLAR-GROUP/hpx/issues/1082>

³²⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1081>

³²⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1072>

³²⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1070>

³²⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1059>

³²⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1051>

³²⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/973>

³²⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/970>

³²⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/941>

³²⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/935>

³²⁶² <https://github.com/STELLAR-GROUP/hpx/issues/934>

³²⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/927>

³²⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/926>

³²⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/925>

- Issue #899³²⁶⁶ - Fix race conditions in regression tests
- Issue #879³²⁶⁷ - Hung test leads to cascading test failure; make tests should support the MPI parcelport
- Issue #865³²⁶⁸ - future<T> and friends shall work for movable only Ts
- Issue #847³²⁶⁹ - Dynamic libraries are not installed on OS X
- Issue #816³²⁷⁰ - First Program tutorial pull request
- Issue #799³²⁷¹ - Wrap lexical_cast to avoid exceptions
- Issue #720³²⁷² - broken configuration when using cmake on Ubuntu
- Issue #622³²⁷³ - --hpx:hpx and --hpx:debug-hpx-log is nonsensical
- Issue #525³²⁷⁴ - Extend barrier LCO test to run in distributed
- Issue #515³²⁷⁵ - Multi-destination version of hpx::apply is broken
- Issue #509³²⁷⁶ - Push Boost.Atomic changes upstream
- Issue #503³²⁷⁷ - Running HPX applications on Windows should not require setting %PATH%
- Issue #461³²⁷⁸ - Add a compilation sanity test
- Issue #456³²⁷⁹ - hpx_run_tests.py should log output from tests that timeout
- Issue #454³²⁸⁰ - Investigate threadmanager performance
- Issue #345³²⁸¹ - Add more versatile environmental/cmake variable support to hpx_find_* CMake macros
- Issue #209³²⁸² - Support multiple configurations in generated build files
- Issue #190³²⁸³ - hpx::cout should be a std::ostream
- Issue #189³²⁸⁴ - iostreams component should use startup/shutdown functions
- Issue #183³²⁸⁵ - Use Boost.ICL for correctness in AGAS
- Issue #44³²⁸⁶ - Implement real futures

³²⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/899>

³²⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/879>

³²⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/865>

³²⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/847>

³²⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/816>

³²⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/799>

³²⁷² <https://github.com/STELLAR-GROUP/hpx/issues/720>

³²⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/622>

³²⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/525>

³²⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/515>

³²⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/509>

³²⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/503>

³²⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/461>

³²⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/456>

³²⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/454>

³²⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/345>

³²⁸² <https://github.com/STELLAR-GROUP/hpx/issues/209>

³²⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/190>

³²⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/189>

³²⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/183>

³²⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/44>

2.11.11 HPX V0.9.8 (Mar 24, 2014)

We have had over 800 commits since the last release and we have closed over 65 tickets (bugs, feature requests, etc.).

With the changes below, *HPX* is once again leading the charge of a whole new era of computation. By intrinsically breaking down and synchronizing the work to be done, *HPX* insures that application developers will no longer have to fret about where a segment of code executes. That allows coders to focus their time and energy to understanding the data dependencies of their algorithms and thereby the core obstacles to an efficient code. Here are some of the advantages of using *HPX*:

- *HPX* is solidly rooted in a sophisticated theoretical execution model – ParalleX
- *HPX* exposes an API fully conforming to the C++11 and the draft C++14 standards, extended and applied to distributed computing. Everything programmers know about the concurrency primitives of the standard C++ library is still valid in the context of *HPX*.
- It provides a competitive, high performance implementation of modern, future-proof ideas which gives an smooth migration path from today's mainstream techniques
- There is no need for the programmer to worry about lower level parallelization paradigms like threads or message passing; no need to understand pthreads, MPI, OpenMP, or Windows threads, etc.
- There is no need to think about different types of parallelism such as tasks, pipelines, or fork-join, task or data parallelism.
- The same source of your program compiles and runs on Linux, BlueGene/Q, Mac OS X, Windows, and Android.
- The same code runs on shared memory multi-core systems and supercomputers, on handheld devices and Intel® Xeon Phi™ accelerators, or a heterogeneous mix of those.

General changes

- A major API breaking change for this release was introduced by implementing `hpx::future` and `hpx::shared_future` fully in conformance with the C++11 Standard³²⁸⁷. While `hpx::shared_future` is new and will not create any compatibility problems, we revised the interface and implementation of the existing `hpx::future`. For more details please see the mailing list archive³²⁸⁸. To avoid any incompatibilities for existing code we named the type which implements the `std::future` interface as `hpx::unique_future`. For the next release this will be renamed to `hpx::future`, making it full conforming to C++11 Standard³²⁸⁹.
- A large part of the code base of *HPX* has been refactored and partially re-implemented. The main changes were related to
 - The threading subsystem: these changes significantly reduce the amount of overheads caused by the schedulers, improve the modularity of the code base, and extend the variety of available scheduling algorithms.
 - The parcel subsystem: these changes improve the performance of the *HPX* networking layer, modularize the structure of the parcelports, and simplify the creation of new parcelports for other underlying networking libraries.
 - The API subsystem: these changes improved the conformance of the API to C++11 Standard, extend and unify the available API functionality, and decrease the overheads created by various elements of the API.
 - The robustness of the component loading subsystem has been improved significantly, allowing to more portably and more reliably register the components needed by an application as startup. This additionally speeds up general application initialization.

³²⁸⁷ <http://www.open-std.org/jtc1/sc22/wg21>

³²⁸⁸ <http://mail.cct.lsu.edu/pipermail/hpx-users/2014-January/000141.html>

³²⁸⁹ <http://www.open-std.org/jtc1/sc22/wg21>

- We added new API functionality like `hpx::migrate` and `hpx::copy_component` which are the basic building blocks necessary for implementing higher level abstractions for system-wide load balancing, runtime-adaptive resource management, and object-oriented checkpointing and state-management.
- We removed the use of C++11 move emulation (using Boost.Move), replacing it with C++11 rvalue references. This is the first step towards using more and more native C++11 facilities which we plan to introduce in the future.
- We improved the reference counting scheme used by *HPX* which helps managing distributed objects and memory. This improves the overall stability of *HPX* and further simplifies writing real world applications.
- The minimal Boost version required to use *HPX* is now V1.49.0.
- This release coincides with the first release of HPXPI (V0.1.0), the first implementation of the XPI specification³²⁹⁰.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- Issue #1086³²⁹¹ - Expose internal boost::shared_array to allow user management of array lifetime
- Issue #1083³²⁹² - Make shell examples copyable in docs
- Issue #1080³²⁹³ - /threads{locality#*/total}/count/cumulative broken
- Issue #1079³²⁹⁴ - Build problems on OS X
- Issue #1078³²⁹⁵ - Improve robustness of component loading
- Issue #1077³²⁹⁶ - Fix a missing enum definition for ‘take’ mode
- Issue #1076³²⁹⁷ - Merge Jb master
- Issue #1075³²⁹⁸ - Unknown CMake command “add_hpx_pseudo_target”
- Issue #1074³²⁹⁹ - Implement apply_continue_callback and apply_colocated_callback
- Issue #1073³³⁰⁰ - The new apply_colocated and async_colocated functions lead to automatic registered functions
- Issue #1071³³⁰¹ - Remove deferred_packaged_task
- Issue #1069³³⁰² - serialize_buffer with allocator fails at destruction
- Issue #1068³³⁰³ - Coroutine include and forward declarations missing
- Issue #1067³³⁰⁴ - Add allocator support to util::serialize_buffer

³²⁹⁰ <https://github.com/STELLAR-GROUP/hpxpi/blob/master/spec.pdf?raw=true>

³²⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1086>

³²⁹² <https://github.com/STELLAR-GROUP/hpx/issues/1083>

³²⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/1080>

³²⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1079>

³²⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1078>

³²⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1077>

³²⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1076>

³²⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1075>

³²⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1074>

³³⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1073>

³³⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/1071>

³³⁰² <https://github.com/STELLAR-GROUP/hpx/issues/1069>

³³⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/1068>

³³⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1067>

- Issue #1066³³⁰⁵ - Allow for MPI_Init being called before HPX launches
- Issue #1065³³⁰⁶ - AGAS cache isn't used/populated on worker localities
- Issue #1064³³⁰⁷ - Reorder includes to ensure ws2 includes early
- Issue #1063³³⁰⁸ - Add hpx::runtime::suspend and hpx::runtime::resume
- Issue #1062³³⁰⁹ - Fix async_continue to properly handle return types
- Issue #1061³³¹⁰ - Implement async_colocated and apply_colocated
- Issue #1060³³¹¹ - Implement minimal component migration
- Issue #1058³³¹² - Remove HPX_UTIL_TUPLE from code base
- Issue #1057³³¹³ - Add performance counters for threading subsystem
- Issue #1055³³¹⁴ - Thread allocation uses two memory pools
- Issue #1053³³¹⁵ - Work stealing flawed
- Issue #1052³³¹⁶ - Fix a number of warnings
- Issue #1049³³¹⁷ - Fixes for TLS on OSX and more reliable test running
- Issue #1048³³¹⁸ - Fixing after 588 hang
- Issue #1047³³¹⁹ - Use port '0' for networking when using one locality
- Issue #1046³³²⁰ - composable_guard test is broken when having more than one thread
- Issue #1045³³²¹ - Security missing headers
- Issue #1044³³²² - Native TLS on FreeBSD via __thread
- Issue #1043³³²³ - async et.al. compute the wrong result type
- Issue #1042³³²⁴ - async et.al. implicitly unwrap reference_wrappers
- Issue #1041³³²⁵ - Remove redundant costly Kleene stars from regex searches
- Issue #1040³³²⁶ - CMake script regex match patterns has unnecessary kleenes
- Issue #1039³³²⁷ - Remove use of Boost.Move and replace with std::move and real rvalue refs

³³⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1066>

³³⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1065>

³³⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1064>

³³⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1063>

³³⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1062>

³³¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1061>

³³¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/1060>

³³¹² <https://github.com/STELLAR-GROUP/hpx/issues/1058>

³³¹³ <https://github.com/STELLAR-GROUP/hpx/issues/1057>

³³¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1055>

³³¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1053>

³³¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1052>

³³¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1049>

³³¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1048>

³³¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1047>

³³²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1046>

³³²¹ <https://github.com/STELLAR-GROUP/hpx/issues/1045>

³³²² <https://github.com/STELLAR-GROUP/hpx/issues/1044>

³³²³ <https://github.com/STELLAR-GROUP/hpx/issues/1043>

³³²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1042>

³³²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1041>

³³²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1040>

³³²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1039>

- Issue #1038³³²⁸ - Bump minimal required Boost to 1.49.0
- Issue #1037³³²⁹ - Implicit unwrapping of futures in async broken
- Issue #1036³³³⁰ - Scheduler hangs when user code attempts to “block” OS-threads
- Issue #1035³³³¹ - Idle-rate counter always reports 100% idle rate
- Issue #1034³³³² - Symbolic name registration causes application hangs
- Issue #1033³³³³ - Application options read in from an options file generate an error message
- Issue #1032³³³⁴ - `hpx::id_type` local reference counting is wrong
- Issue #1031³³³⁵ - Negative entry in reference count table
- Issue #1030³³³⁶ - Implement condition_variable
- Issue #1029³³³⁷ - Deadlock in thread scheduling subsystem
- Issue #1028³³³⁸ - HPX-thread cumulative count performance counters report incorrect value
- Issue #1027³³³⁹ - Expose `hpx::thread_interrupted` error code as a separate exception type
- Issue #1026³³⁴⁰ - Exceptions thrown in asynchronous calls can be lost if the value of the future is never queried
- Issue #1025³³⁴¹ - `future::wait_for/wait_until` do not remove callback
- Issue #1024³³⁴² - Remove dependence to boost assert and create hpx assert
- Issue #1023³³⁴³ - Segfaults with tcmalloc
- Issue #1022³³⁴⁴ - prerequisites link in readme is broken
- Issue #1020³³⁴⁵ - HPX Deadlock on external synchronization
- Issue #1019³³⁴⁶ - Convert using BOOST_ASSERT to HPX_ASSERT
- Issue #1018³³⁴⁷ - compiling bug with gcc 4.8.1
- Issue #1017³³⁴⁸ - Possible crash in io_pool executor
- Issue #1016³³⁴⁹ - Crash at startup
- Issue #1014³³⁵⁰ - Implement Increment/Decrement Merging

³³²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1038>

³³²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1037>

³³³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1036>

³³³¹ <https://github.com/STELLAR-GROUP/hpx/issues/1035>

³³³² <https://github.com/STELLAR-GROUP/hpx/issues/1034>

³³³³ <https://github.com/STELLAR-GROUP/hpx/issues/1033>

³³³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1032>

³³³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1031>

³³³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1030>

³³³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1029>

³³³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1028>

³³³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1027>

³³⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1026>

³³⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/1025>

³³⁴² <https://github.com/STELLAR-GROUP/hpx/issues/1024>

³³⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/1023>

³³⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1022>

³³⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1020>

³³⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1019>

³³⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1018>

³³⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1017>

³³⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1016>

³³⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/1014>

- Issue #1013³³⁵¹ - Add more logging channels to enable greater control over logging granularity
- Issue #1012³³⁵² - --hpx:debug-hpx-log and --hpx:debug-agas-log lead to non-thread safe writes
- Issue #1011³³⁵³ - After installation, running applications from the build/staging directory no longer works
- Issue #1010³³⁵⁴ - Mergable decrement requests are not being merged
- Issue #1009³³⁵⁵ - --hpx:list-symbolic-names crashes
- Issue #1007³³⁵⁶ - Components are not properly destroyed
- Issue #1006³³⁵⁷ - Segfault/hang in set_data
- Issue #1003³³⁵⁸ - Performance counter naming issue
- Issue #982³³⁵⁹ - Race condition during startup
- Issue #912³³⁶⁰ - OS X: component type not found in map
- Issue #663³³⁶¹ - Create a buildbot slave based on Clang 3.2/OSX
- Issue #636³³⁶² - Expose this_locality::apply<act>(p1, p2); for local execution
- Issue #197³³⁶³ - Add --console=address option for PBS runs
- Issue #175³³⁶⁴ - Asynchronous AGAS API

2.11.12 HPX V0.9.7 (Nov 13, 2013)

We have had over 1000 commits since the last release and we have closed over 180 tickets (bugs, feature requests, etc.).

General changes

- Ported HPX to BlueGene/Q
- Improved HPX support for Xeon/Phi accelerators
- Reimplemented hpx::bind, hpx::tuple, and hpx::function for better performance and better compliance with the C++11 Standard. Added hpx::mem_fn.
- Reworked hpx::when_all and hpx::when_any for better compliance with the ongoing C++ standardization effort, added heterogeneous version for those functions. Added hpx::when_any_swapped.
- Added hpx::copy as a precursor for a migrate functionality
- Added hpx::get_ptr allowing to directly access the memory underlying a given component

³³⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/1013>

³³⁵² <https://github.com/STELLAR-GROUP/hpx/issues/1012>

³³⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/1011>

³³⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/1010>

³³⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1009>

3356 <https://github.com/STELLAR-GROUP/hpx/issues/1007>3357 <https://github.com/STELLAR-GROUP/hpx/issues/1006>3358 <https://github.com/STELLAR-GROUP/hpx/issues/1003>3359 <https://github.com/STELLAR-GROUP/hpx/issues/982>3360 <https://github.com/STELLAR-GROUP/hpx/issues/912>3361 <https://github.com/STELLAR-GROUP/hpx/issues/663>3362 <https://github.com/STELLAR-GROUP/hpx/issues/636>3363 <https://github.com/STELLAR-GROUP/hpx/issues/197>3364 <https://github.com/STELLAR-GROUP/hpx/issues/175>

- Added the `hpx::lcos::broadcast`, `hpx::lcos::reduce`, and `hpx::lcos::fold` collective operations
- Added `hpx::get_locality_name` allowing to retrieve the name of any of the localities for the application.
- Added support for more flexible thread affinity control from the HPX command line, such as new modes for `--hpx:bind` (`balanced`, `scattered`, `compact`), improved default settings when running multiple localities on the same node.
- Added experimental executors for simpler thread pooling and scheduling. This API may change in the future as it will stay aligned with the ongoing C++ standardization efforts.
- Massively improved the performance of the HPX serialization code. Added partial support for zero copy serialization of array and bitwise-copyable types.
- General performance improvements of the code related to threads and futures.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release.

- Issue #1005^{[3365](https://github.com/STELLAR-GROUP/hpx/issues/1005)} - Allow one to disable array optimizations and zero copy optimizations for each parcelport
- Issue #1004^{[3366](https://github.com/STELLAR-GROUP/hpx/issues/1004)} - Generate new HPX logo image for the docs
- Issue #1002^{[3367](https://github.com/STELLAR-GROUP/hpx/issues/1002)} - If MPI parcelport is not available, running HPX under mpirun should fail
- Issue #1001^{[3368](https://github.com/STELLAR-GROUP/hpx/issues/1001)} - Zero copy serialization raises assert
- Issue #1000^{[3369](https://github.com/STELLAR-GROUP/hpx/issues/1000)} - Can't connect to a HPX application running with the MPI parcelport from a non MPI parcelport locality
- Issue #999^{[3370](https://github.com/STELLAR-GROUP/hpx/issues/999)} - Optimize `hpx::when_n`
- Issue #998^{[3371](https://github.com/STELLAR-GROUP/hpx/issues/998)} - Fixed const-correctness
- Issue #997^{[3372](https://github.com/STELLAR-GROUP/hpx/issues/997)} - Making `serialize_buffer::data()` type save
- Issue #996^{[3373](https://github.com/STELLAR-GROUP/hpx/issues/996)} - Memory leak in `hpx::lcos::promise`
- Issue #995^{[3374](https://github.com/STELLAR-GROUP/hpx/issues/995)} - Race while registering pre-shutdown functions
- Issue #994^{[3375](https://github.com/STELLAR-GROUP/hpx/issues/994)} - `thread_rescheduling` regression test does not compile
- Issue #992^{[3376](https://github.com/STELLAR-GROUP/hpx/issues/992)} - Correct comments and messages
- Issue #991^{[3377](https://github.com/STELLAR-GROUP/hpx/issues/991)} - `setcap cap_sys_rawio=ep` for power profiling causes an HPX application to abort
- Issue #989^{[3378](https://github.com/STELLAR-GROUP/hpx/issues/989)} - Jacobi hangs during execution

³³⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/1005>

³³⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/1004>

³³⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/1002>

³³⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/1001>

³³⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/1000>

³³⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/999>

³³⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/998>

³³⁷² <https://github.com/STELLAR-GROUP/hpx/issues/997>

³³⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/996>

³³⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/995>

³³⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/994>

³³⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/992>

³³⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/991>

³³⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/989>

- Issue #988³³⁷⁹ - multiple_init test is failing
- Issue #986³³⁸⁰ - Can't call a function called "init" from "main" when using <hpx/hpx_main.hpp>
- Issue #984³³⁸¹ - Reference counting tests are failing
- Issue #983³³⁸² - thread_suspension_executor test fails
- Issue #980³³⁸³ - Terminating HPX threads don't leave stack in virgin state
- Issue #979³³⁸⁴ - Static scheduler not in documents
- Issue #978³³⁸⁵ - Preprocessing limits are broken
- Issue #977³³⁸⁶ - Make tests.regressions.lcos.future_hang_on_get shorter
- Issue #976³³⁸⁷ - Wrong library order in pkgconfig
- Issue #975³³⁸⁸ - Please reopen #963
- Issue #974³³⁸⁹ - Option pu-offset ignored in fixing_588 branch
- Issue #972³³⁹⁰ - Cannot use MKL with HPX
- Issue #969³³⁹¹ - Non-existent INI files requested on the command line via --hpx:config do not cause warnings or errors.
- Issue #968³³⁹² - Cannot build examples in fixing_588 branch
- Issue #967³³⁹³ - Command line description of --hpx:queuing seems wrong
- Issue #966³³⁹⁴ - --hpx:print-bind physical core numbers are wrong
- Issue #965³³⁹⁵ - Deadlock when building in Release mode
- Issue #963³³⁹⁶ - Not all worker threads are working
- Issue #962³³⁹⁷ - Problem with SLURM integration
- Issue #961³³⁹⁸ - --hpx:print-bind outputs incorrect information
- Issue #960³³⁹⁹ - Fix cut and paste error in documentation of get_thread_priority
- Issue #959³⁴⁰⁰ - Change link to boost.atomic in documentation to point to boost.org
- Issue #958³⁴⁰¹ - Undefined reference to intrusive_ptr_release

³³⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/988>

³³⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/986>

³³⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/984>

³³⁸² <https://github.com/STELLAR-GROUP/hpx/issues/983>

³³⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/980>

³³⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/979>

³³⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/978>

³³⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/977>

³³⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/976>

³³⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/975>

³³⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/974>

³³⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/972>

³³⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/969>

³³⁹² <https://github.com/STELLAR-GROUP/hpx/issues/968>

³³⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/967>

³³⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/966>

³³⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/965>

³³⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/963>

³³⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/962>

³³⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/961>

³³⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/960>

³⁴⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/959>

³⁴⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/958>

- Issue #957³⁴⁰² - Make tuple standard compliant
- Issue #956³⁴⁰³ - Segfault with a3382fb
- Issue #955³⁴⁰⁴ --hpx:nodes and --hpx:nodenames do not work with foreign nodes
- Issue #954³⁴⁰⁵ - Make order of arguments for hpx::async and hpx::broadcast consistent
- Issue #953³⁴⁰⁶ - Cannot use MKL with HPX
- Issue #952³⁴⁰⁷ - register_[pre_]shutdown_function never throw
- Issue #951³⁴⁰⁸ - Assert when number of threads is greater than hardware concurrency
- Issue #948³⁴⁰⁹ - HPX_HAVE_GENERIC_CONTEXT_COROUTINES conflicts with HPX_HAVE_FIBER_BASED_COROUTINES
- Issue #947³⁴¹⁰ - Need MPI_THREAD_MULTIPLE for backward compatibility
- Issue #946³⁴¹¹ - HPX does not call MPI_Finalize
- Issue #945³⁴¹² - Segfault with hpx::lcos::broadcast
- Issue #944³⁴¹³ - OS X: assertion pu_offset_ < hardware_concurrency failed
- Issue #943³⁴¹⁴ - #include <hpx/hpx_main.hpp> does not work
- Issue #942³⁴¹⁵ - Make the BG/Q work with -O3
- Issue #940³⁴¹⁶ - Use separator when concatenating locality name
- Issue #939³⁴¹⁷ - Refactor MPI parcelport to use MPI_Wait instead of multiple MPI_Test calls
- Issue #938³⁴¹⁸ - Want to officially access client_base::gid_
- Issue #937³⁴¹⁹ - client_base::gid_ should be private``
- Issue #936³⁴²⁰ - Want doxygen-like source code index
- Issue #935³⁴²¹ - Build error with gcc 4.6 and Boost 1.54.0 on hpx trunk and 0.9.6
- Issue #933³⁴²² - Cannot build HPX with Boost 1.54.0
- Issue #932³⁴²³ - Components are destructed too early
- Issue #931³⁴²⁴ - Make HPX work on BG/Q

³⁴⁰² <https://github.com/STELLAR-GROUP/hpx/issues/957>

³⁴⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/956>

³⁴⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/955>

³⁴⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/954>

³⁴⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/953>

³⁴⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/952>

³⁴⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/951>

³⁴⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/948>

³⁴¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/947>

³⁴¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/946>

³⁴¹² <https://github.com/STELLAR-GROUP/hpx/issues/945>

³⁴¹³ <https://github.com/STELLAR-GROUP/hpx/issues/944>

³⁴¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/943>

³⁴¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/942>

³⁴¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/940>

³⁴¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/939>

³⁴¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/938>

³⁴¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/937>

³⁴²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/936>

³⁴²¹ <https://github.com/STELLAR-GROUP/hpx/issues/935>

³⁴²² <https://github.com/STELLAR-GROUP/hpx/issues/933>

³⁴²³ <https://github.com/STELLAR-GROUP/hpx/issues/932>

³⁴²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/931>

- Issue #930³⁴²⁵ - make git-docs is broken
- Issue #929³⁴²⁶ - Generating index in docs broken
- Issue #928³⁴²⁷ - Optimize `hpx::util::static_` for C++11 compilers supporting magic statics
- Issue #924³⁴²⁸ - Make `kill_process_tree` (in `process.py`) more robust on Mac OSX
- Issue #923³⁴²⁹ - Correct BLAS and RNPL cmake tests
- Issue #922³⁴³⁰ - Cannot link against BLAS
- Issue #921³⁴³¹ - Implement `hpx::mem_fn`
- Issue #920³⁴³² - Output locality with `--hpx:print-bind`
- Issue #919³⁴³³ - Correct grammar; simplify boolean expressions
- Issue #918³⁴³⁴ - Link to `hello_world.cpp` is broken
- Issue #917³⁴³⁵ - adapt cmake file to new boostbook version
- Issue #916³⁴³⁶ - fix problem building documentation with `xsltproc >= 1.1.27`
- Issue #915³⁴³⁷ - Add another TBBMalloc library search path
- Issue #914³⁴³⁸ - Build problem with Intel compiler on Stampede (TACC)
- Issue #913³⁴³⁹ - fix error messages in fibonacci examples
- Issue #911³⁴⁴⁰ - Update OS X build instructions
- Issue #910³⁴⁴¹ - Want like to specify `MPI_ROOT` instead of compiler wrapper script
- Issue #909³⁴⁴² - Warning about `void*` arithmetic
- Issue #908³⁴⁴³ - Buildbot for MIC is broken
- Issue #906³⁴⁴⁴ - Can't use `--hpx:bind=balanced` with multiple MPI processes
- Issue #905³⁴⁴⁵ - `--hpx:bind` documentation should describe full grammar
- Issue #904³⁴⁴⁶ - Add `hpx::lcos::fold` and `hpx::lcos::inverse_fold` collective operation
- Issue #903³⁴⁴⁷ - Add `hpx::when_any_swapped()`

³⁴²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/930>

³⁴²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/929>

³⁴²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/928>

³⁴²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/924>

³⁴²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/923>

³⁴³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/922>

³⁴³¹ <https://github.com/STELLAR-GROUP/hpx/issues/921>

³⁴³² <https://github.com/STELLAR-GROUP/hpx/issues/920>

³⁴³³ <https://github.com/STELLAR-GROUP/hpx/issues/919>

³⁴³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/918>

³⁴³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/917>

³⁴³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/916>

³⁴³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/915>

³⁴³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/914>

³⁴³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/913>

³⁴⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/911>

³⁴⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/910>

³⁴⁴² <https://github.com/STELLAR-GROUP/hpx/issues/909>

³⁴⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/908>

³⁴⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/906>

³⁴⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/905>

³⁴⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/904>

³⁴⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/903>

- Issue #902³⁴⁴⁸ - Add `hpx::lcos::reduce` collective operation
- Issue #901³⁴⁴⁹ - Web documentation is not searchable
- Issue #900³⁴⁵⁰ - Web documentation for trunk has no index
- Issue #898³⁴⁵¹ - Some tests fail with GCC 4.8.1 and MPI parcel port
- Issue #897³⁴⁵² - HWLOC causes failures on Mac
- Issue #896³⁴⁵³ - pu-offset leads to startup error
- Issue #895³⁴⁵⁴ - `hpx::get_locality_name` not defined
- Issue #894³⁴⁵⁵ - Race condition at shutdown
- Issue #893³⁴⁵⁶ - --`hpx:print-bind` switches `std::cout` to hexadecimal mode
- Issue #892³⁴⁵⁷ - `hwloc_topology_load` can be expensive – don't call multiple times
- Issue #891³⁴⁵⁸ - The documentation for `get_locality_name` is wrong
- Issue #890³⁴⁵⁹ - --`hpx:print-bind` should not exit
- Issue #889³⁴⁶⁰ --`hpx:debug-hpx-log=FILE` does not work
- Issue #888³⁴⁶¹ - MPI parcelport does not exit cleanly for --`hpx:print-bind`
- Issue #887³⁴⁶² - Choose thread affinities more cleverly
- Issue #886³⁴⁶³ - Logging documentation is confusing
- Issue #885³⁴⁶⁴ - Two threads are slower than one
- Issue #884³⁴⁶⁵ - `is_callable` failing with member pointers in C++11
- Issue #883³⁴⁶⁶ - Need help with `is_callable`_test
- Issue #882³⁴⁶⁷ - `tests/regressions.lcos.future_hang_on_get` does not terminate
- Issue #881³⁴⁶⁸ - `tests/regressions/block_matrix/matrix.hh` won't compile with GCC 4.8.1
- Issue #880³⁴⁶⁹ - HPX does not work on OS X
- Issue #878³⁴⁷⁰ - `future::unwrap` triggers assertion

³⁴⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/902>

³⁴⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/901>

³⁴⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/900>

³⁴⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/898>

³⁴⁵² <https://github.com/STELLAR-GROUP/hpx/issues/897>

³⁴⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/896>

³⁴⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/895>

³⁴⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/894>

³⁴⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/893>

³⁴⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/892>

³⁴⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/891>

³⁴⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/890>

³⁴⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/889>

³⁴⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/888>

³⁴⁶² <https://github.com/STELLAR-GROUP/hpx/issues/887>

³⁴⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/886>

³⁴⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/885>

³⁴⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/884>

³⁴⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/883>

³⁴⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/882>

³⁴⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/881>

³⁴⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/880>

³⁴⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/878>

- Issue #877³⁴⁷¹ - “make tests” has build errors on Ubuntu 12.10
- Issue #876³⁴⁷² - tcmalloc is used by default, even if it is not present
- Issue #875³⁴⁷³ - global_fixture is defined in a header file
- Issue #874³⁴⁷⁴ - Some tests take very long
- Issue #873³⁴⁷⁵ - Add block-matrix code as regression test
- Issue #872³⁴⁷⁶ - HPX documentation does not say how to run tests with detailed output
- Issue #871³⁴⁷⁷ - All tests fail with “make test”
- Issue #870³⁴⁷⁸ - Please explicitly disable serialization in classes that don’t support it
- Issue #868³⁴⁷⁹ - boost_any test failing
- Issue #867³⁴⁸⁰ - Reduce the number of copies of hpx::function arguments
- Issue #863³⁴⁸¹ - Futures should not require a default constructor
- Issue #862³⁴⁸² - value_or_error shall not default construct its result
- Issue #861³⁴⁸³ - HPX_UNUSED macro
- Issue #860³⁴⁸⁴ - Add functionality to copy construct a component
- Issue #859³⁴⁸⁵ - hpx::endl should flush
- Issue #858³⁴⁸⁶ - Create hpx::get_ptr<> allowing to access component implementation
- Issue #855³⁴⁸⁷ - Implement hpx::INVOKER
- Issue #854³⁴⁸⁸ - hpx/hpx.hpp does not include hpx/include/iostreams.hpp
- Issue #853³⁴⁸⁹ - Feature request: null future
- Issue #852³⁴⁹⁰ - Feature request: Locality names
- Issue #851³⁴⁹¹ - hpx::cout output does not appear on screen
- Issue #849³⁴⁹² - All tests fail on OS X after installing
- Issue #848³⁴⁹³ - Update OS X build instructions

³⁴⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/877>

³⁴⁷² <https://github.com/STELLAR-GROUP/hpx/issues/876>

³⁴⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/875>

³⁴⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/874>

³⁴⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/873>

³⁴⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/872>

³⁴⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/871>

³⁴⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/870>

³⁴⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/868>

³⁴⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/867>

³⁴⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/863>

³⁴⁸² <https://github.com/STELLAR-GROUP/hpx/issues/862>

³⁴⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/861>

³⁴⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/860>

³⁴⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/859>

³⁴⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/858>

³⁴⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/855>

³⁴⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/854>

³⁴⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/853>

³⁴⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/852>

³⁴⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/851>

³⁴⁹² <https://github.com/STELLAR-GROUP/hpx/issues/849>

³⁴⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/848>

- Issue #846³⁴⁹⁴ - Update hpx_external_example
- Issue #845³⁴⁹⁵ - Issues with having both debug and release modules in the same directory
- Issue #844³⁴⁹⁶ - Create configuration header
- Issue #843³⁴⁹⁷ - Tests should use CTest
- Issue #842³⁴⁹⁸ - Remove buffer_pool from MPI parcelport
- Issue #841³⁴⁹⁹ - Add possibility to broadcast an index with hpx::lcos::broadcast
- Issue #838³⁵⁰⁰ - Simplify util::tuple
- Issue #837³⁵⁰¹ - Adopt boost::tuple tests for util::tuple
- Issue #836³⁵⁰² - Adopt boost::function tests for util::function
- Issue #835³⁵⁰³ - Tuple interface missing pieces
- Issue #833³⁵⁰⁴ - Partially preprocessing files not working
- Issue #832³⁵⁰⁵ - Native papi counters do not work with wild cards
- Issue #831³⁵⁰⁶ - Arithmetics counter fails if only one parameter is given
- Issue #830³⁵⁰⁷ - Convert hpx::util::function to use new scheme for serializing its base pointer
- Issue #829³⁵⁰⁸ - Consistently use decay<T> instead of remove_const< remove_reference<T>>
- Issue #828³⁵⁰⁹ - Update future implementation to N3721 and N3722
- Issue #827³⁵¹⁰ - Enable MPI parcelport for bootstrapping whenever application was started using mpirun
- Issue #826³⁵¹¹ - Support command line option --hpx:print-bind even if --hpx::bind was not used
- Issue #825³⁵¹² - Memory counters give segfault when attempting to use thread wild cards or numbers only total works
- Issue #824³⁵¹³ - Enable lambda functions to be used with hpx::async/hpx::apply
- Issue #823³⁵¹⁴ - Using a hashing filter
- Issue #822³⁵¹⁵ - Silence unused variable warning
- Issue #821³⁵¹⁶ - Detect if a function object is callable with given arguments

³⁴⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/846>

³⁴⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/845>

³⁴⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/844>

³⁴⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/843>

³⁴⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/842>

³⁴⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/841>

³⁵⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/838>

³⁵⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/837>

³⁵⁰² <https://github.com/STELLAR-GROUP/hpx/issues/836>

³⁵⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/835>

³⁵⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/833>

³⁵⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/832>

³⁵⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/831>

³⁵⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/830>

³⁵⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/829>

³⁵⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/828>

³⁵¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/827>

³⁵¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/826>

³⁵¹² <https://github.com/STELLAR-GROUP/hpx/issues/825>

³⁵¹³ <https://github.com/STELLAR-GROUP/hpx/issues/824>

³⁵¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/823>

³⁵¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/822>

³⁵¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/821>

- Issue #820³⁵¹⁷ - Allow wildcards to be used for performance counter names
- Issue #819³⁵¹⁸ - Make the AGAS symbolic name registry distributed
- Issue #818³⁵¹⁹ - Add future::then() overload taking an executor
- Issue #817³⁵²⁰ - Fixed typo
- Issue #815³⁵²¹ - Create an lco that is performing an efficient broadcast of actions
- Issue #814³⁵²² - Papi counters cannot specify thread#* to get the counts for all threads
- Issue #813³⁵²³ - Scoped unlock
- Issue #811³⁵²⁴ - simple_central_tuplespace_client run error
- Issue #810³⁵²⁵ - ostream error when << any objects
- Issue #809³⁵²⁶ - Optimize parcel serialization
- Issue #808³⁵²⁷ - HPX applications throw exception when executed from the build directory
- Issue #807³⁵²⁸ - Create performance counters exposing overall AGAS statistics
- Issue #795³⁵²⁹ - Create timed make_ready_future
- Issue #794³⁵³⁰ - Create heterogeneous when_all/when_any/etc.
- Issue #721³⁵³¹ - Make HPX usable for Xeon Phi
- Issue #694³⁵³² - CMake should complain if you attempt to build an example without its dependencies
- Issue #692³⁵³³ - SLURM support broken
- Issue #683³⁵³⁴ - python/hpx/process.py imports epoll on all platforms
- Issue #619³⁵³⁵ - Automate the doc building process
- Issue #600³⁵³⁶ - GTC performance broken
- Issue #577³⁵³⁷ - Allow for zero copy serialization/networking
- Issue #551³⁵³⁸ - Change executable names to have debug postfix in Debug builds
- Issue #544³⁵³⁹ - Write a custom .lib file on Windows pulling in hpx_init and hpx.dll, phase out hpx_init

³⁵¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/820>

³⁵¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/819>

³⁵¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/818>

³⁵²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/817>

³⁵²¹ <https://github.com/STELLAR-GROUP/hpx/issues/815>

³⁵²² <https://github.com/STELLAR-GROUP/hpx/issues/814>

³⁵²³ <https://github.com/STELLAR-GROUP/hpx/issues/813>

³⁵²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/811>

³⁵²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/810>

³⁵²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/809>

³⁵²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/808>

³⁵²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/807>

³⁵²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/795>

³⁵³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/794>

³⁵³¹ <https://github.com/STELLAR-GROUP/hpx/issues/721>

³⁵³² <https://github.com/STELLAR-GROUP/hpx/issues/694>

³⁵³³ <https://github.com/STELLAR-GROUP/hpx/issues/692>

³⁵³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/683>

³⁵³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/619>

³⁵³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/600>

³⁵³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/577>

³⁵³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/551>

³⁵³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/544>

- Issue #534³⁵⁴⁰ - hpx::init should take functions by `std::function` and should accept all forms of `hpx_main`
- Issue #508³⁵⁴¹ - FindPackage fails to set `FOO_LIBRARY_DIR`
- Issue #506³⁵⁴² - Add cmake support to generate ini files for external applications
- Issue #470³⁵⁴³ - Changing build-type after configure does not update boost library names
- Issue #453³⁵⁴⁴ - Document `hpx_run_tests.py`
- Issue #445³⁵⁴⁵ - Significant performance mismatch between MPI and HPX in SMP for allgather example
- Issue #443³⁵⁴⁶ - Make docs viewable from build directory
- Issue #421³⁵⁴⁷ - Support multiple HPX instances per node in a batch environment like PBS or SLURM
- Issue #316³⁵⁴⁸ - Add message size limitation
- Issue #249³⁵⁴⁹ - Clean up locking code in big boot barrier
- Issue #136³⁵⁵⁰ - Persistent CMake variables need to be marked as cache variables

2.11.13 HPX V0.9.6 (Jul 30, 2013)

We have had over 1200 commits since the last release and we have closed roughly 140 tickets (bugs, feature requests, etc.).

General changes

The major new features in this release are:

- We further consolidated the API exposed by *HPX*. We aligned our APIs as much as possible with the existing C++11 Standard³⁵⁵¹ and related proposals to the C++ standardization committee (such as N3632³⁵⁵² and N3857³⁵⁵³).
- We implemented a first version of a distributed AGAS service which essentially eliminates all explicit AGAS network traffic.
- We created a native ibverbs parcelport allowing to take advantage of the superior latency and bandwidth characteristics of Infiniband networks.
- We successfully ported *HPX* to the Xeon Phi platform.
- Support for the SLURM scheduling system was implemented.
- Major efforts have been dedicated to improving the performance counter framework, numerous new counters were implemented and new APIs were added.

³⁵⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/534>

³⁵⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/508>

³⁵⁴² <https://github.com/STELLAR-GROUP/hpx/issues/506>

³⁵⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/470>

³⁵⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/453>

3545 <https://github.com/STELLAR-GROUP/hpx/issues/445>3546 <https://github.com/STELLAR-GROUP/hpx/issues/443>3547 <https://github.com/STELLAR-GROUP/hpx/issues/421>3548 <https://github.com/STELLAR-GROUP/hpx/issues/316>

3549 <https://github.com/STELLAR-GROUP/hpx/issues/249>

3550 <https://github.com/STELLAR-GROUP/hpx/issues/136>

3551 <http://www.open-std.org/jtc1/sc22/wg21>

3552 <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2013/n3632.html>

3553 <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n3857.pdf>

2.11. Releases

1091

- We added a modular parcel compression system allowing to improve bandwidth utilization (by reducing the overall size of the transferred data).
- We added a modular parcel coalescing system allowing to combine several parcels into larger messages. This reduces latencies introduced by the communication layer.
- Added an experimental executors API allowing to use different scheduling policies for different parts of the code. This API has been modelled after the Standards proposal N3562³⁵⁵⁴. This API is bound to change in the future, though.
- Added minimal security support for localities which is enforced on the parcelport level. This support is preliminary and experimental and might change in the future.
- We created a parcelport using low level MPI functions. This is in support of legacy applications which are to be gradually ported and to support platforms where MPI is the only available portable networking layer.
- We added a preliminary and experimental implementation of a tuple-space object which exposes an interface similar to such systems described in the literature (see for instance The Linda Coordination Language³⁵⁵⁵).

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release. This is again a very long list of newly implemented features and fixed issues.

- Issue #806³⁵⁵⁶ - make (all) in examples folder does nothing
- Issue #805³⁵⁵⁷ - Adding the introduction and fixing DOCBOOK dependencies for Windows use
- Issue #804³⁵⁵⁸ - Add stackless (non-suspendable) thread type
- Issue #803³⁵⁵⁹ - Create proper serialization support functions for util::tuple
- Issue #800³⁵⁶⁰ - Add possibility to disable array optimizations during serialization
- Issue #798³⁵⁶¹ - HPX_LIMIT does not work for local dataflow
- Issue #797³⁵⁶² - Create a parcelport which uses MPI
- Issue #796³⁵⁶³ - Problem with Large Numbers of Threads
- Issue #793³⁵⁶⁴ - Changing dataflow test case to hang consistently
- Issue #792³⁵⁶⁵ - CMake Error
- Issue #791³⁵⁶⁶ - Problems with local::dataflow
- Issue #790³⁵⁶⁷ - wait_for() doesn't compile
- Issue #789³⁵⁶⁸ - HPX with Intel compiler segfaults

³⁵⁵⁴ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2013/n3562.pdf>

³⁵⁵⁵ [https://en.wikipedia.org/wiki/Linda_\(coordination_language\)](https://en.wikipedia.org/wiki/Linda_(coordination_language))

³⁵⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/806>

³⁵⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/805>

³⁵⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/804>

³⁵⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/803>

³⁵⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/800>

³⁵⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/798>

³⁵⁶² <https://github.com/STELLAR-GROUP/hpx/issues/797>

³⁵⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/796>

³⁵⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/793>

³⁵⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/792>

³⁵⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/791>

³⁵⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/790>

³⁵⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/789>

- Issue #788³⁵⁶⁹ - Intel compiler support
- Issue #787³⁵⁷⁰ - Fixed SFINAEd specializations
- Issue #786³⁵⁷¹ - Memory issues during benchmarking.
- Issue #785³⁵⁷² - Create an API allowing to register external threads with HPX
- Issue #784³⁵⁷³ - util::plugin is throwing an error when a symbol is not found
- Issue #783³⁵⁷⁴ - How does hpx:bind work?
- Issue #782³⁵⁷⁵ - Added quotes around STRING REPLACE potentially empty arguments
- Issue #781³⁵⁷⁶ - Make sure no exceptions propagate into the thread manager
- Issue #780³⁵⁷⁷ - Allow arithmetics performance counters to expand its parameters
- Issue #779³⁵⁷⁸ - Test case for 778
- Issue #778³⁵⁷⁹ - Swapping futures segfaults
- Issue #777³⁵⁸⁰ - hpx::lcos::details::when_xxx don't restore completion handlers
- Issue #776³⁵⁸¹ - Compiler chokes on dataflow overload with launch policy
- Issue #775³⁵⁸² - Runtime error with local dataflow (copying futures?)
- Issue #774³⁵⁸³ - Using local dataflow without explicit namespace
- Issue #773³⁵⁸⁴ - Local dataflow with unwrap: functor operators need to be const
- Issue #772³⁵⁸⁵ - Allow (remote) actions to return a future
- Issue #771³⁵⁸⁶ - Setting HPX_LIMIT gives huge boost MPL errors
- Issue #770³⁵⁸⁷ - Add launch policy to (local) dataflow
- Issue #769³⁵⁸⁸ - Make compile time configuration information available
- Issue #768³⁵⁸⁹ - Const correctness problem in local dataflow
- Issue #767³⁵⁹⁰ - Add launch policies to async
- Issue #766³⁵⁹¹ - Mark data structures for optimized (array based) serialization

³⁵⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/788>

³⁵⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/787>

³⁵⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/786>

³⁵⁷² <https://github.com/STELLAR-GROUP/hpx/issues/785>

³⁵⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/784>

³⁵⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/783>

3575 <https://github.com/STELLAR-GROUP/hpx/issues/782>3576 <https://github.com/STELLAR-GROUP/hpx/issues/781>3577 <https://github.com/STELLAR-GROUP/hpx/issues/780>3578 <https://github.com/STELLAR-GROUP/hpx/issues/779>

3579 <https://github.com/STELLAR-GROUP/hpx/issues/778>

3580 <https://github.com/STELLAR-GROUP/hpx/issues/777>

3581 <https://github.com/STELLAR-GROUP/hpx/issues/776>

3582 <https://github.com/STELLAR-GROUP/hpx/issues/775>

3583 <https://github.com/STELLAR-GROUP/hpx/issues/774>

3584 <https://github.com/STELLAR-GROUP/hpx/issues/773>

3585 <https://github.com/STELLAR-GROUP/hpx/issues/772>

3586 <https://github.com/STELLAR-GROUP/hpx/issues/771>

3587 <https://github.com/STELLAR-GROUP/hpx/issues/770>

3588 <https://github.com/STELLAR-GROUP/hpx/issues/769>

3589 <https://github.com/STELLAR-GROUP/hpx/issues/768>

3590 <https://github.com/STELLAR-GROUP/hpx/issues/767>

3591 <https://github.com/STELLAR-GROUP/hpx/issues/766>

- Issue #765³⁵⁹² - Align hpx::any with N3508: Any Library Proposal (Revision 2)
- Issue #764³⁵⁹³ - Align hpx::future with newest N3558: A Standardized Representation of Asynchronous Operations
- Issue #762³⁵⁹⁴ - added a human readable output for the ping pong example
- Issue #761³⁵⁹⁵ - Ambiguous typename when constructing derived component
- Issue #760³⁵⁹⁶ - Simple components can not be derived
- Issue #759³⁵⁹⁷ - make install doesn't give a complete install
- Issue #758³⁵⁹⁸ - Stack overflow when using locking_hook<>
- Issue #757³⁵⁹⁹ - copy paste error; unsupported function overloading
- Issue #756³⁶⁰⁰ - GTCX runtime issue in Gordon
- Issue #755³⁶⁰¹ - Papi counters don't work with reset and evaluate API's
- Issue #753³⁶⁰² - cmake bugfix and improved component action docs
- Issue #752³⁶⁰³ - hpx simple component docs
- Issue #750³⁶⁰⁴ - Add hpx::util::any
- Issue #749³⁶⁰⁵ - Thread phase counter is not reset
- Issue #748³⁶⁰⁶ - Memory performance counter are not registered
- Issue #747³⁶⁰⁷ - Create performance counters exposing arithmetic operations
- Issue #745³⁶⁰⁸ - apply_callback needs to invoke callback when applied locally
- Issue #744³⁶⁰⁹ - CMake fixes
- Issue #743³⁶¹⁰ - Problem Building github version of HPX
- Issue #742³⁶¹¹ - Remove HPX_STD_BIND
- Issue #741³⁶¹² - assertion 'px != 0' failed: HPX(assertion_failure) for low numbers of OS threads
- Issue #739³⁶¹³ - Performance counters do not count to the end of the program or evalution
- Issue #738³⁶¹⁴ - Dedicated AGAS server runs don't work; console ignores -a option.

³⁵⁹² <https://github.com/STELLAR-GROUP/hpx/issues/765>

³⁵⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/764>

³⁵⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/762>

³⁵⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/761>

³⁵⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/760>

³⁵⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/759>

³⁵⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/758>

³⁵⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/757>

3600 <https://github.com/STELLAR-GROUP/hpx/issues/756>3601 <https://github.com/STELLAR-GROUP/hpx/issues/755>3602 <https://github.com/STELLAR-GROUP/hpx/issues/753>3603 <https://github.com/STELLAR-GROUP/hpx/issues/752>3604 <https://github.com/STELLAR-GROUP/hpx/issues/750>3605 <https://github.com/STELLAR-GROUP/hpx/issues/749>3606 <https://github.com/STELLAR-GROUP/hpx/issues/748>3607 <https://github.com/STELLAR-GROUP/hpx/issues/747>3608 <https://github.com/STELLAR-GROUP/hpx/issues/745>3609 <https://github.com/STELLAR-GROUP/hpx/issues/744>3610 <https://github.com/STELLAR-GROUP/hpx/issues/743>3611 <https://github.com/STELLAR-GROUP/hpx/issues/742>3612 <https://github.com/STELLAR-GROUP/hpx/issues/741>3613 <https://github.com/STELLAR-GROUP/hpx/issues/739>3614 <https://github.com/STELLAR-GROUP/hpx/issues/738>

- Issue #737³⁶¹⁵ - Missing bind overloads
- Issue #736³⁶¹⁶ - Performance counter wildcards do not always work
- Issue #735³⁶¹⁷ - Create native ibverbs parcelport based on rdma operations
- Issue #734³⁶¹⁸ - Threads stolen performance counter total is incorrect
- Issue #733³⁶¹⁹ - Test benchmarks need to be checked and fixed
- Issue #732³⁶²⁰ - Build fails with Mac, using mac ports clang-3.3 on latest git branch
- Issue #731³⁶²¹ - Add global start/stop API for performance counters
- Issue #730³⁶²² - Performance counter values are apparently incorrect
- Issue #729³⁶²³ - Unhandled switch
- Issue #728³⁶²⁴ - Serialization of hpx::util::function between two localities causes seg faults
- Issue #727³⁶²⁵ - Memory counters on Mac OS X
- Issue #725³⁶²⁶ - Restore original thread priority on resume
- Issue #724³⁶²⁷ - Performance benchmarks do not depend on main HPX libraries
- Issue #723³⁶²⁸ - [teletype]-hpx:nodes=``cat \$PBS_NODEFILE`` works; -hpx:nodefile=\$PBS_NODEFILE does not.[c++]
- Issue #722³⁶²⁹ - Fix binding const member functions as actions
- Issue #719³⁶³⁰ - Create performance counter exposing compression ratio
- Issue #718³⁶³¹ - Add possibility to compress parcel data
- Issue #717³⁶³² - strip_credit_from_gid has misleading semantics
- Issue #716³⁶³³ - Non-option arguments to programs run using pbsdsh must be before --hpx:nodes, contrary to directions
- Issue #715³⁶³⁴ - Re-thrown exceptions should retain the original call site
- Issue #714³⁶³⁵ - failed assertion in debug mode
- Issue #713³⁶³⁶ - Add performance counters monitoring connection caches

³⁶¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/737>

³⁶¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/736>

³⁶¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/735>

³⁶¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/734>

³⁶¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/733>

³⁶²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/732>

³⁶²¹ <https://github.com/STELLAR-GROUP/hpx/issues/731>

³⁶²² <https://github.com/STELLAR-GROUP/hpx/issues/730>

³⁶²³ <https://github.com/STELLAR-GROUP/hpx/issues/729>

³⁶²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/728>

³⁶²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/727>

³⁶²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/725>

³⁶²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/724>

³⁶²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/723>

³⁶²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/722>

³⁶³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/719>

³⁶³¹ <https://github.com/STELLAR-GROUP/hpx/issues/718>

³⁶³² <https://github.com/STELLAR-GROUP/hpx/issues/717>

³⁶³³ <https://github.com/STELLAR-GROUP/hpx/issues/716>

³⁶³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/715>

³⁶³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/714>

³⁶³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/713>

- Issue #712³⁶³⁷ - Adjust parcel related performance counters to be connection type specific
- Issue #711³⁶³⁸ - configuration failure
- Issue #710³⁶³⁹ - Error “timed out while trying to find room in the connection cache” when trying to start multiple localities on a single computer
- Issue #709³⁶⁴⁰ - Add new thread state ‘staged’ referring to task descriptions
- Issue #708³⁶⁴¹ - Detect/mitigate bad non-system installs of GCC on Redhat systems
- Issue #707³⁶⁴² - Many examples do not link with Git HEAD version
- Issue #706³⁶⁴³ - hpx::init removes portions of non-option command line arguments before last = sign
- Issue #705³⁶⁴⁴ - Create rolling average and median aggregating performance counters
- Issue #704³⁶⁴⁵ - Create performance counter to expose thread queue waiting time
- Issue #703³⁶⁴⁶ - Add support to HPX build system to find librctool.a and related headers
- Issue #699³⁶⁴⁷ - Generalize instrumentation support
- Issue #698³⁶⁴⁸ - compilation failure with hwloc absent
- Issue #697³⁶⁴⁹ - Performance counter counts should be zero indexed
- Issue #696³⁶⁵⁰ - Distributed problem
- Issue #695³⁶⁵¹ - Bad perf counter time printed
- Issue #693³⁶⁵² - --help doesn't print component specific command line options
- Issue #692³⁶⁵³ - SLURM support broken
- Issue #691³⁶⁵⁴ - exception while executing any application linked with hwloc
- Issue #690³⁶⁵⁵ - thread_id_test and thread_launcher_test failing
- Issue #689³⁶⁵⁶ - Make the buildbots use hwloc
- Issue #687³⁶⁵⁷ - compilation error fix (hwloc_topology)
- Issue #686³⁶⁵⁸ - Linker Error for Applications
- Issue #684³⁶⁵⁹ - Pinning of service thread fails when number of worker threads equals the number of cores

³⁶³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/712>

³⁶³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/711>

³⁶³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/710>

³⁶⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/709>

³⁶⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/708>

³⁶⁴² <https://github.com/STELLAR-GROUP/hpx/issues/707>

³⁶⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/706>

³⁶⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/705>

³⁶⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/704>

³⁶⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/703>

³⁶⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/699>

³⁶⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/698>

³⁶⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/697>

³⁶⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/696>

³⁶⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/695>

³⁶⁵² <https://github.com/STELLAR-GROUP/hpx/issues/693>

³⁶⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/692>

³⁶⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/691>

³⁶⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/690>

³⁶⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/689>

³⁶⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/687>

³⁶⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/686>

³⁶⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/684>

- Issue #682³⁶⁶⁰ - Add performance counters exposing number of stolen threads
- Issue #681³⁶⁶¹ - Add apply_continue for asynchronous chaining of actions
- Issue #679³⁶⁶² - Remove obsolete async_callback API functions
- Issue #678³⁶⁶³ - Add new API for setting/triggering LCOs
- Issue #677³⁶⁶⁴ - Add async_continue for true continuation style actions
- Issue #676³⁶⁶⁵ - Buildbot for gcc 4.4 broken
- Issue #675³⁶⁶⁶ - Partial preprocessing broken
- Issue #674³⁶⁶⁷ - HPX segfaults when built with gcc 4.7
- Issue #673³⁶⁶⁸ - use_guard_pages has inconsistent preprocessor guards
- Issue #672³⁶⁶⁹ - External build breaks if library path has spaces
- Issue #671³⁶⁷⁰ - release tarballs are tarbombs
- Issue #670³⁶⁷¹ - CMake won't find Boost headers in layout=versioned install
- Issue #669³⁶⁷² - Links in docs to source files broken if not installed
- Issue #667³⁶⁷³ - Not reading ini file properly
- Issue #664³⁶⁷⁴ - Adapt new meanings of ‘const’ and ‘mutable’
- Issue #661³⁶⁷⁵ - Implement BTL Parcel port
- Issue #655³⁶⁷⁶ - Make HPX work with the “decltype” result_of
- Issue #647³⁶⁷⁷ - documentation for specifying the number of high priority threads --hpx:high-priority-threads
- Issue #643³⁶⁷⁸ - Error parsing host file
- Issue #642³⁶⁷⁹ - HWLoc issue with TAU
- Issue #639³⁶⁸⁰ - Logging potentially suspends a running thread
- Issue #634³⁶⁸¹ - Improve error reporting from parcel layer
- Issue #627³⁶⁸² - Add tests for async and apply overloads that accept regular C++ functions

³⁶⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/682>

³⁶⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/681>

³⁶⁶² <https://github.com/STELLAR-GROUP/hpx/issues/679>

³⁶⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/678>

³⁶⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/677>

³⁶⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/676>

³⁶⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/675>

³⁶⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/674>

³⁶⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/673>

³⁶⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/672>

³⁶⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/671>

³⁶⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/670>

³⁶⁷² <https://github.com/STELLAR-GROUP/hpx/issues/669>

³⁶⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/667>

³⁶⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/664>

³⁶⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/661>

³⁶⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/655>

³⁶⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/647>

³⁶⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/643>

³⁶⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/642>

³⁶⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/639>

³⁶⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/634>

³⁶⁸² <https://github.com/STELLAR-GROUP/hpx/issues/627>

- Issue #626³⁶⁸³ - hpx/future.hpp header
- Issue #601³⁶⁸⁴ - Intel support
- Issue #557³⁶⁸⁵ - Remove action codes
- Issue #531³⁶⁸⁶ - AGAS request and response classes should use switch statements
- Issue #529³⁶⁸⁷ - Investigate the state of hwloc support
- Issue #526³⁶⁸⁸ - Make HPX aware of hyper-threading
- Issue #518³⁶⁸⁹ - Create facilities allowing to use plain arrays as action arguments
- Issue #473³⁶⁹⁰ - hwloc thread binding is broken on CPUs with hyperthreading
- Issue #383³⁶⁹¹ - Change result type detection for hpx::util::bind to use result_of protocol
- Issue #341³⁶⁹² - Consolidate route code
- Issue #219³⁶⁹³ - Only copy arguments into actions once
- Issue #177³⁶⁹⁴ - Implement distributed AGAS
- Issue #43³⁶⁹⁵ - Support for Darwin (Xcode + Clang)

2.11.14 HPX V0.9.5 (Jan 16, 2013)

We have had over 1000 commits since the last release and we have closed roughly 150 tickets (bugs, feature requests, etc.).

General changes

This release is continuing along the lines of code and API consolidation, and overall usability improvements. We dedicated much attention to performance and we were able to significantly improve the threading and networking subsystems.

We successfully ported *HPX* to the Android platform. *HPX* applications now not only can run on mobile devices, but we support heterogeneous applications running across architecture boundaries. At the Supercomputing Conference 2012 we demonstrated connecting Android tablets to simulations running on a Linux cluster. The Android tablet was used to query performance counters from the Linux simulation and to steer its parameters.

We successfully ported *HPX* to Mac OSX (using the Clang compiler). Thanks to Pyry Jähkälä for contributing the corresponding patches. Please see the section *How to install HPX on OS X (Mac)* for more details.

We made a special effort to make *HPX* usable in highly concurrent use cases. Many of the *HPX* API functions which possibly take longer than 100 microseconds to execute now can be invoked asynchronously. We added uniform

³⁶⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/626>

³⁶⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/601>

³⁶⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/557>

³⁶⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/531>

3687 <https://github.com/STELLAR-GROUP/hpx/issues/529>3688 <https://github.com/STELLAR-GROUP/hpx/issues/526>3689 <https://github.com/STELLAR-GROUP/hpx/issues/518>3690 <https://github.com/STELLAR-GROUP/hpx/issues/473>3691 <https://github.com/STELLAR-GROUP/hpx/issues/383>3692 <https://github.com/STELLAR-GROUP/hpx/issues/341>3693 <https://github.com/STELLAR-GROUP/hpx/issues/219>3694 <https://github.com/STELLAR-GROUP/hpx/issues/177>3695 <https://github.com/STELLAR-GROUP/hpx/issues/43>

support for composing futures which simplifies to write asynchronous code. HPX actions (function objects encapsulating possibly concurrent remote function invocations) are now well integrated with all other API facilities such like `hpx::bind`.

All of the API has been aligned as much as possible with established paradigms. HPX now mirrors many of the facilities as defined in the C++11 Standard, such as `hpx::thread`, `hpx::function`, `hpx::future`, etc.

A lot of work has been put into improving the documentation. Many of the API functions are documented now, concepts are explained in detail, and examples are better described than before. The new documentation index enables finding information with lesser effort.

This is the first release of HPX we perform after the move to [Github³⁶⁹⁶](#). This step has enabled a wider participation from the community and further encourages us in our decision to release HPX as a true open source library (HPX is licensed under the very liberal [Boost Software License³⁶⁹⁷](#)).

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release. This is by far the longest list of newly implemented features and fixed issues for any of HPX' releases so far.

- Issue #666³⁶⁹⁸ - Segfault on calling `hpx::finalize` twice
- Issue #665³⁶⁹⁹ - Adding declaration `num_of_cores`
- Issue #662³⁷⁰⁰ - `pkgconfig` is building wrong
- Issue #660³⁷⁰¹ - Need `uninterrupt` function
- Issue #659³⁷⁰² - Move our logging library into a different namespace
- Issue #658³⁷⁰³ - Dynamic performance counter types are broken
- Issue #657³⁷⁰⁴ - HPX v0.9.5 (RC1) `hello_world` example segfaulting
- Issue #656³⁷⁰⁵ - Define the affinity of parcel-pool, io-pool, and timer-pool threads
- Issue #654³⁷⁰⁶ - Integrate the Boost `auto_index` tool with documentation
- Issue #653³⁷⁰⁷ - Make HPX build on OS X + Clang + `libc++`
- Issue #651³⁷⁰⁸ - Add fine-grained control for thread pinning
- Issue #650³⁷⁰⁹ - Command line no error message when using `-hpx:(anything)`
- Issue #645³⁷¹⁰ - Command line aliases don't work in `[teletype]``@file``[c++]`
- Issue #644³⁷¹¹ - Terminated threads are not always properly cleaned up

³⁶⁹⁶ <https://github.com/STELLAR-GROUP/hpx/>

³⁶⁹⁷ https://www.boost.org/LICENSE_1_0.txt

³⁶⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/666>

³⁶⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/665>

³⁷⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/662>

³⁷⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/660>

³⁷⁰² <https://github.com/STELLAR-GROUP/hpx/issues/659>

³⁷⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/658>

³⁷⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/657>

³⁷⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/656>

³⁷⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/654>

³⁷⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/653>

³⁷⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/651>

³⁷⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/650>

³⁷¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/645>

³⁷¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/644>

- Issue #640³⁷¹² - `future_data<T>::set_on_completed_` used without locks
- Issue #638³⁷¹³ - hpx build with intel compilers fails on linux
- Issue #637³⁷¹⁴ - `--copy-dt-needed-entries` breaks with gold
- Issue #635³⁷¹⁵ - Boost V1.53 will add Boost.Lockfree and Boost.Atomic
- Issue #633³⁷¹⁶ - Re-add examples to final 0.9.5 release
- Issue #632³⁷¹⁷ - Example `thread_aware_timer` is broken
- Issue #631³⁷¹⁸ - FFT application throws error in parcellayer
- Issue #630³⁷¹⁹ - Event synchronization example is broken
- Issue #629³⁷²⁰ - Waiting on futures hangs
- Issue #628³⁷²¹ - Add an `HPX_ALWAYS_ASSERT` macro
- Issue #625³⁷²² - Port coroutines context switch benchmark
- Issue #621³⁷²³ - New INI section for stack sizes
- Issue #618³⁷²⁴ - `pkg_config` support does not work with a HPX debug build
- Issue #617³⁷²⁵ - `hpx/external/logging/boost/logging/detail/cache_before_init.hpp:139:67:` error: ‘`get_thread_id`’ was not declared in this scope
- Issue #616³⁷²⁶ - Change `wait_xxx` not to use locking
- Issue #615³⁷²⁷ - Revert visibility ‘fix’ (`fb0b6b8245dad1127b0c25ebaf9386b3945cca9`)
- Issue #614³⁷²⁸ - Fix Dataflow linker error
- Issue #613³⁷²⁹ - `find_here` should throw an exception on failure
- Issue #612³⁷³⁰ - Thread phase doesn’t show up in debug mode
- Issue #611³⁷³¹ - Make stack guard pages configurable at runtime (initialization time)
- Issue #610³⁷³² - Co-Locate Components
- Issue #609³⁷³³ - `future_overhead`
- Issue #608³⁷³⁴ - `--hpx:list-counter-infos` problem

³⁷¹² <https://github.com/STELLAR-GROUP/hpx/issues/640>

³⁷¹³ <https://github.com/STELLAR-GROUP/hpx/issues/638>

³⁷¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/637>

³⁷¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/635>

³⁷¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/633>

³⁷¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/632>

³⁷¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/631>

³⁷¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/630>

³⁷²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/629>

³⁷²¹ <https://github.com/STELLAR-GROUP/hpx/issues/628>

³⁷²² <https://github.com/STELLAR-GROUP/hpx/issues/625>

³⁷²³ <https://github.com/STELLAR-GROUP/hpx/issues/621>

³⁷²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/618>

³⁷²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/617>

³⁷²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/616>

³⁷²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/615>

³⁷²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/614>

³⁷²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/613>

³⁷³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/612>

³⁷³¹ <https://github.com/STELLAR-GROUP/hpx/issues/611>

³⁷³² <https://github.com/STELLAR-GROUP/hpx/issues/610>

³⁷³³ <https://github.com/STELLAR-GROUP/hpx/issues/609>

³⁷³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/608>

- Issue #607³⁷³⁵ - Update Boost.Context based backend for coroutines
- Issue #606³⁷³⁶ - 1d_wave_equation is not working
- Issue #605³⁷³⁷ - Any C++ function that has serializable arguments and a serializable return type should be remotable
- Issue #604³⁷³⁸ - Connecting localities isn't working anymore
- Issue #603³⁷³⁹ - Do not verify any ini entries read from a file
- Issue #602³⁷⁴⁰ - Rename argument_size to type_size/ added implementation to get parcel size
- Issue #599³⁷⁴¹ - Enable locality specific command line options
- Issue #598³⁷⁴² - Need an API that accesses the performance counter reporting the system uptime
- Issue #597³⁷⁴³ - compiling on ranger
- Issue #595³⁷⁴⁴ - I need a place to store data in a thread self pointer
- Issue #594³⁷⁴⁵ - 32/64 interoperability
- Issue #593³⁷⁴⁶ - Warn if logging is disabled at compile time but requested at runtime
- Issue #592³⁷⁴⁷ - Add optional argument value to --hpx:list-counters and --hpx:list-counter-infos
- Issue #591³⁷⁴⁸ - Allow for wildcards in performance counter names specified with --hpx:print-counter
- Issue #590³⁷⁴⁹ - Local promise semantic differences
- Issue #589³⁷⁵⁰ - Create API to query performance counter names
- Issue #587³⁷⁵¹ - Add get_num_localities and get_num_threads to AGAS API
- Issue #586³⁷⁵² - Adjust local AGAS cache size based on number of localities
- Issue #585³⁷⁵³ - Error while using counters in HPX
- Issue #584³⁷⁵⁴ - counting argument size of actions, initial pass.
- Issue #581³⁷⁵⁵ - Remove RemoteResult template parameter for future<>
- Issue #580³⁷⁵⁶ - Add possibility to hook into actions

³⁷³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/607>

³⁷³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/606>

³⁷³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/605>

³⁷³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/604>

³⁷³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/603>

³⁷⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/602>

³⁷⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/599>

³⁷⁴² <https://github.com/STELLAR-GROUP/hpx/issues/598>

³⁷⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/597>

³⁷⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/595>

³⁷⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/594>

³⁷⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/593>

³⁷⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/592>

³⁷⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/591>

³⁷⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/590>

³⁷⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/589>

³⁷⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/587>

³⁷⁵² <https://github.com/STELLAR-GROUP/hpx/issues/586>

³⁷⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/585>

³⁷⁵⁴ <https://github.com/STELLAR-GROUP/hpx/issues/584>

³⁷⁵⁵ <https://github.com/STELLAR-GROUP/hpx/issues/581>

³⁷⁵⁶ <https://github.com/STELLAR-GROUP/hpx/issues/580>

- Issue #578³⁷⁵⁷ - Use angle brackets in HPX error dumps
- Issue #576³⁷⁵⁸ - Exception incorrectly thrown when --help is used
- Issue #575³⁷⁵⁹ - HPX(bad_component_type) with gcc 4.7.2 and boost 1.51
- Issue #574³⁷⁶⁰ - --hpx:connect command line parameter not working correctly
- Issue #571³⁷⁶¹ - hpx::wait () (callback version) should pass the future to the callback function
- Issue #570³⁷⁶² - hpx::wait should operate on boost::arrays and std::lists
- Issue #569³⁷⁶³ - Add a logging sink for Android
- Issue #568³⁷⁶⁴ - 2-argument version of HPX_DEFINE_COMPONENT_ACTION
- Issue #567³⁷⁶⁵ - Connecting to a running HPX application works only once
- Issue #565³⁷⁶⁶ - HPX doesn't shutdown properly
- Issue #564³⁷⁶⁷ - Partial preprocessing of new component creation interface
- Issue #563³⁷⁶⁸ - Add hpx::start/hpx::stop to avoid blocking main thread
- Issue #562³⁷⁶⁹ - All command line arguments swallowed by hpx
- Issue #561³⁷⁷⁰ - Boost.Tuple is not move aware
- Issue #558³⁷⁷¹ - boost::shared_ptr<> style semantics/syntax for client classes
- Issue #556³⁷⁷² - Creation of partially preprocessed headers should be enabled for Boost newer than V1.50
- Issue #555³⁷⁷³ - BOOST_FORCEINLINE does not name a type
- Issue #554³⁷⁷⁴ - Possible race condition in thread get_id()
- Issue #552³⁷⁷⁵ - Move enable client_base
- Issue #550³⁷⁷⁶ - Add stack size category 'huge'
- Issue #549³⁷⁷⁷ - ShenEOS run seg-faults on single or distributed runs
- Issue #545³⁷⁷⁸ - AUTOGLOB broken for add_hpx_component
- Issue #542³⁷⁷⁹ - FindHPX_HDF5 still searches multiple times

³⁷⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/578>

³⁷⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/576>

³⁷⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/575>

³⁷⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/574>

³⁷⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/571>

³⁷⁶² <https://github.com/STELLAR-GROUP/hpx/issues/570>

3763 <https://github.com/STELLAR-GROUP/hpx/issues/569>3764 <https://github.com/STELLAR-GROUP/hpx/issues/568>3765 <https://github.com/STELLAR-GROUP/hpx/issues/567>3766 <https://github.com/STELLAR-GROUP/hpx/issues/565>3767 <https://github.com/STELLAR-GROUP/hpx/issues/564>3768 <https://github.com/STELLAR-GROUP/hpx/issues/563>3769 <https://github.com/STELLAR-GROUP/hpx/issues/562>3770 <https://github.com/STELLAR-GROUP/hpx/issues/561>3771 <https://github.com/STELLAR-GROUP/hpx/issues/558>3772 <https://github.com/STELLAR-GROUP/hpx/issues/556>3773 <https://github.com/STELLAR-GROUP/hpx/issues/555>3774 <https://github.com/STELLAR-GROUP/hpx/issues/554>3775 <https://github.com/STELLAR-GROUP/hpx/issues/552>3776 <https://github.com/STELLAR-GROUP/hpx/issues/550>3777 <https://github.com/STELLAR-GROUP/hpx/issues/549>3778 <https://github.com/STELLAR-GROUP/hpx/issues/545>3779 <https://github.com/STELLAR-GROUP/hpx/issues/542>

- Issue #541³⁷⁸⁰ - Quotes around application name in hpx::init
- Issue #539³⁷⁸¹ - Race condition occurring with new lightweight threads
- Issue #535³⁷⁸² - hpx_run_tests.py exits with no error code when tests are missing
- Issue #530³⁷⁸³ - Thread description(<unknown>) in logs
- Issue #523³⁷⁸⁴ - Make thread objects more lightweight
- Issue #521³⁷⁸⁵ - hpx::error_code is not usable for lightweight error handling
- Issue #520³⁷⁸⁶ - Add full user environment to HPX logs
- Issue #519³⁷⁸⁷ - Build succeeds, running fails
- Issue #517³⁷⁸⁸ - Add a guard page to linux coroutine stacks
- Issue #516³⁷⁸⁹ - hpx::thread::detach suspends while holding locks, leads to hang in debug
- Issue #514³⁷⁹⁰ - Preprocessed headers for <hpx/apply.hpp> don't compile
- Issue #513³⁷⁹¹ - Buildbot configuration problem
- Issue #512³⁷⁹² - Implement action based stack size customization
- Issue #511³⁷⁹³ - Move action priority into a separate type trait
- Issue #510³⁷⁹⁴ - trunk broken
- Issue #507³⁷⁹⁵ - no matching function for call to boost::scoped_ptr<hpx::threads::topology>::scoped_ptr(hpx::threads::topology)
- Issue #505³⁷⁹⁶ - undefined_symbol regression test currently failing
- Issue #502³⁷⁹⁷ - Adding OpenCL and OCLM support to HPX for Windows and Linux
- Issue #501³⁷⁹⁸ - find_package(HPX) sets cmake output variables
- Issue #500³⁷⁹⁹ - wait_any/wait_all are badly named
- Issue #499³⁸⁰⁰ - Add support for disabling pbs support in pbs runs
- Issue #498³⁸⁰¹ - Error during no-cache runs
- Issue #496³⁸⁰² - Add partial preprocessing support to cmake

³⁷⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/541>

³⁷⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/539>

³⁷⁸² <https://github.com/STELLAR-GROUP/hpx/issues/535>

³⁷⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/530>

³⁷⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/523>

³⁷⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/521>

³⁷⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/520>

³⁷⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/519>

³⁷⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/517>

³⁷⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/516>

³⁷⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/514>

³⁷⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/513>

³⁷⁹² <https://github.com/STELLAR-GROUP/hpx/issues/512>

³⁷⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/511>

³⁷⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/510>

³⁷⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/507>

³⁷⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/505>

³⁷⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/502>

³⁷⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/501>

³⁷⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/500>

³⁸⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/499>

³⁸⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/498>

³⁸⁰² <https://github.com/STELLAR-GROUP/hpx/issues/496>

- Issue #495³⁸⁰³ - Support HPX modules exporting startup/shutdown functions only
- Issue #494³⁸⁰⁴ - Allow modules to specify when to run startup/shutdown functions
- Issue #493³⁸⁰⁵ - Avoid constructing a string in make_success_code
- Issue #492³⁸⁰⁶ - Performance counter creation is no longer synchronized at startup
- Issue #491³⁸⁰⁷ - Performance counter creation is no longer synchronized at startup
- Issue #490³⁸⁰⁸ - Sheneos on_completed_bulk seg fault in distributed
- Issue #489³⁸⁰⁹ - compiling issue with g++44
- Issue #488³⁸¹⁰ - Adding OpenCL and OCLM support to HPX for the MSVC platform
- Issue #487³⁸¹¹ - FindHPX.cmake problems
- Issue #485³⁸¹² - Change distributing_factory and binpacking_factory to use bulk creation
- Issue #484³⁸¹³ - Change HPX_DONT_USE_PREPROCESSED_FILES to HPX_USE_PREPROCESSED_FILES
- Issue #483³⁸¹⁴ - Memory counter for Windows
- Issue #479³⁸¹⁵ - strange errors appear when requesting performance counters on multiple nodes
- Issue #477³⁸¹⁶ - Create (global) timer for multi-threaded measurements
- Issue #472³⁸¹⁷ - Add partial preprocessing using Wave
- Issue #471³⁸¹⁸ - Segfault stack traces don't show up in release
- Issue #468³⁸¹⁹ - External projects need to link with internal components
- Issue #462³⁸²⁰ - Startup/shutdown functions are called more than once
- Issue #458³⁸²¹ - Consolidate hpx::util::high_resolution_timer and hpx::util::high_resolution_clock
- Issue #457³⁸²² - index out of bounds in allgather_and_gate on 4 cores or more
- Issue #448³⁸²³ - Make HPX compile with clang
- Issue #447³⁸²⁴ - 'make tests' should execute tests on local installation
- Issue #446³⁸²⁵ - Remove SVN-related code from the codebase

³⁸⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/495>

³⁸⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/494>

³⁸⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/493>

³⁸⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/492>

³⁸⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/491>

³⁸⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/490>

³⁸⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/489>

³⁸¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/488>

³⁸¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/487>

³⁸¹² <https://github.com/STELLAR-GROUP/hpx/issues/485>

³⁸¹³ <https://github.com/STELLAR-GROUP/hpx/issues/484>

³⁸¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/483>

³⁸¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/479>

³⁸¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/477>

³⁸¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/472>

³⁸¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/471>

³⁸¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/468>

³⁸²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/462>

³⁸²¹ <https://github.com/STELLAR-GROUP/hpx/issues/458>

³⁸²² <https://github.com/STELLAR-GROUP/hpx/issues/457>

³⁸²³ <https://github.com/STELLAR-GROUP/hpx/issues/448>

³⁸²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/447>

³⁸²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/446>

- Issue #444³⁸²⁶ - race condition in smp
- Issue #441³⁸²⁷ - Patched Boost.Serialization headers should only be installed if needed
- Issue #439³⁸²⁸ - Components using HPX_REGISTER_STARTUP_MODULE fail to compile with MSVC
- Issue #436³⁸²⁹ - Verify that no locks are being held while threads are suspended
- Issue #435³⁸³⁰ - Installing HPX should not clobber existing Boost installation
- Issue #434³⁸³¹ - Logging external component failed (Boost 1.50)
- Issue #433³⁸³² - Runtime crash when building all examples
- Issue #432³⁸³³ - Dataflow hangs on 512 cores/64 nodes
- Issue #430³⁸³⁴ - Problem with distributing factory
- Issue #424³⁸³⁵ - File paths referring to XSL-files need to be properly escaped
- Issue #417³⁸³⁶ - Make dataflow LCOs work out of the box by using partial preprocessing
- Issue #413³⁸³⁷ - hpx_svnversion.py fails on Windows
- Issue #412³⁸³⁸ - Make hpx::error_code equivalent to hpx::exception
- Issue #398³⁸³⁹ - HPX clobbers out-of-tree application specific CMake variables (specifically CMAKE_BUILD_TYPE)
- Issue #394³⁸⁴⁰ - Remove code generating random port numbers for network
- Issue #378³⁸⁴¹ - ShenEOS scaling issues
- Issue #354³⁸⁴² - Create a coroutines wrapper for Boost.Context
- Issue #349³⁸⁴³ - Commandline option --localities=N/-1N should be necessary only on AGAS locality
- Issue #334³⁸⁴⁴ - Add auto_index support to cmake based documentation toolchain
- Issue #318³⁸⁴⁵ - Network benchmarks
- Issue #317³⁸⁴⁶ - Implement network performance counters
- Issue #310³⁸⁴⁷ - Duplicate logging entries
- Issue #230³⁸⁴⁸ - Add compile time option to disable thread debugging info

³⁸²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/444>

³⁸²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/441>

³⁸²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/439>

³⁸²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/436>

³⁸³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/435>

³⁸³¹ <https://github.com/STELLAR-GROUP/hpx/issues/434>

³⁸³² <https://github.com/STELLAR-GROUP/hpx/issues/433>

³⁸³³ <https://github.com/STELLAR-GROUP/hpx/issues/432>

³⁸³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/430>

³⁸³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/424>

³⁸³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/417>

³⁸³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/413>

³⁸³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/412>

³⁸³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/398>

³⁸⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/394>

³⁸⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/378>

³⁸⁴² <https://github.com/STELLAR-GROUP/hpx/issues/354>

³⁸⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/349>

³⁸⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/334>

³⁸⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/318>

³⁸⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/317>

³⁸⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/310>

³⁸⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/230>

- Issue #171³⁸⁴⁹ - Add an INI option to turn off deadlock detection independently of logging
- Issue #170³⁸⁵⁰ - OSHL internal counters are incorrect
- Issue #103³⁸⁵¹ - Better diagnostics for multiple component/action registrations under the same name
- Issue #48³⁸⁵² - Support for Darwin (Xcode + Clang)
- Issue #21³⁸⁵³ - Build fails with GCC 4.6

2.11.15 HPX V0.9.0 (Jul 5, 2012)

We have had roughly 800 commits since the last release and we have closed approximately 80 tickets (bugs, feature requests, etc.).

General changes

- Significant improvements made to the usability of *HPX* in large-scale, distributed environments.
- Renamed `hpx::lcos::packaged_task` to `hpx::lcos::packaged_action` to reflect the semantic differences to a `packaged_task` as defined by the C++11 Standard³⁸⁵⁴.
- *HPX* now exposes `hpx::thread` which is compliant to the C++11 `std::thread` type except that it (purely locally) represents an *HPX* thread. This new type does not expose any of the remote capabilities of the underlying *HPX*-thread implementation.
- The type `hpx::lcos::future` is now compliant to the C++11 `std::future<>` type. This type can be used to synchronize both, local and remote operations. In both cases the control flow will ‘return’ to the future in order to trigger any continuation.
- The types `hpx::lcos::local::promise` and `hpx::lcos::local::packaged_task` are now compliant to the C++11 `std::promise<>` and `std::packaged_task<>` types. These can be used to create a future representing local work only. Use the types `hpx::lcos::promise` and `hpx::lcos::packaged_action` to wrap any (possibly remote) action into a future.
- `hpx::thread` and `hpx::lcos::future` are now cancelable.
- Added support for sequential and logic composition of `hpx::lcos::futures`. The member function `hpx::lcos::future::when` permits futures to be sequentially composed. The helper functions `hpx::wait_all`, `hpx::wait_any`, and `hpx::wait_n` can be used to wait for more than one future at a time.
- *HPX* now exposes `hpx::apply` and `hpx::async` as the preferred way of creating (or invoking) any deferred work. These functions are usable with various types of functions, function objects, and actions and provide a uniform way to spawn deferred tasks.
- *HPX* now utilizes `hpx::util::bind` to (partially) bind local functions and function objects, and also actions. Remote bound actions can have placeholders as well.
- *HPX* continuations are now fully polymorphic. The class `hpx::actions::forwarding_continuation` is an example of how the user can write their own types of continuations. It can be used to execute any function as an continuation of a particular action.

³⁸⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/171>

³⁸⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/170>

³⁸⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/103>

³⁸⁵² <https://github.com/STELLAR-GROUP/hpx/issues/48>

³⁸⁵³ <https://github.com/STELLAR-GROUP/hpx/issues/21>

³⁸⁵⁴ <http://www.open-std.org/jtc1/sc22/wg21>

- Reworked the action invocation API to be fully conformant to normal functions. Actions can now be invoked using `hpx::apply`, `hpx::async`, or using the `operator()` implemented on actions. Actions themselves can now be cheaply instantiated as they do not have any members anymore.
- Reworked the lazy action invocation API. Actions can now be directly bound using `hpx::util::bind` by passing an action instance as the first argument.
- A minimal *HPX* program now looks like this:

```
#include <hpx/hpx_init.hpp>

int hpx_main()
{
    return hpx::finalize();
}

int main()
{
    return hpx::init();
}
```

This removes the immediate dependency on the `Boost.Program Options`³⁸⁵⁵ library.

Note: This minimal version of an *HPX* program does not support any of the default command line arguments (such as `-help`, or command line options related to PBS). It is suggested to always pass `argc` and `argv` to *HPX* as shown in the example below.

- In order to support those, but still not to depend on `Boost.Program Options`³⁸⁵⁶, the minimal program can be written as:

```
#include <hpx/hpx_init.hpp>

// The arguments for hpx_main can be left off, which very similar to the
// behavior of ``main()`` as defined by C++.
int hpx_main(int argc, char* argv[])
{
    return hpx::finalize();
}

int main(int argc, char* argv[])
{
    return hpx::init(argc, argv);
}
```

- Added performance counters exposing the number of component instances which are alive on a given locality.
- Added performance counters exposing the number of messages sent and received, the number of parcels sent and received, the number of bytes sent and received, the overall time required to send and receive data, and the overall time required to serialize and deserialize the data.
- Added a new component: `hpx::components::binpacking_factory` which is equivalent to the existing `hpx::components::distributing_factory` component, except that it equalizes the overall population of the components to create. It exposes two factory methods, one based on the number of existing instances of the component type to create, and one based on an arbitrary performance counter which will be queried for all relevant localities.

³⁸⁵⁵ https://www.boost.org/doc/html/program_options.html

³⁸⁵⁶ https://www.boost.org/doc/html/program_options.html

- Added API functions allowing to access elements of the diagnostic information embedded in the given exception: `hpx::get_locality_id`, `hpx::get_host_name`, `hpx::get_process_id`, `hpx::get_function_name`, `hpx::get_file_name`, `hpx::get_line_number`, `hpx::get_os_thread`, `hpx::get_thread_id`, and `hpx::get_thread_description`.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release:

- Issue #71³⁸⁵⁷ - GIDs that are not serialized via `handle_gid` should raise an exception
- Issue #105³⁸⁵⁸ - Allow for `hpx::util::functions` to be registered in the AGAS symbolic namespace
- Issue #107³⁸⁵⁹ - Nasty threadmanger race condition (reproducible in `sheneos_test`)
- Issue #108³⁸⁶⁰ - Add millisecond resolution to HPX logs on Linux
- Issue #110³⁸⁶¹ - Shutdown hang in distributed with release build
- Issue #116³⁸⁶² - Don't use TSS for the applier and runtime pointers
- Issue #162³⁸⁶³ - Move local synchronous execution shortcut from `hpx::function` to the applier
- Issue #172³⁸⁶⁴ - Cache sources in CMake and check if they change manually
- Issue #178³⁸⁶⁵ - Add an INI option to turn off ranged-based AGAS caching
- Issue #187³⁸⁶⁶ - Support for disabling performance counter deployment
- Issue #202³⁸⁶⁷ - Support for sending performance counter data to a specific file
- Issue #218³⁸⁶⁸ - boost.coroutines allows different stack sizes, but stack pool is unaware of this
- Issue #231³⁸⁶⁹ - Implement movable `boost::bind`
- Issue #232³⁸⁷⁰ - Implement movable `boost::function`
- Issue #236³⁸⁷¹ - Allow binding `hpx::util::function` to actions
- Issue #239³⁸⁷² - Replace `hpx::function` with `hpx::util::function`
- Issue #240³⁸⁷³ - Can't specify `RemoteResult` with `Icos::async`
- Issue #242³⁸⁷⁴ - `REGISTER_TEMPLATE` support for plain actions
- Issue #243³⁸⁷⁵ - `handle_gid` support for `hpx::util::function`

³⁸⁵⁷ <https://github.com/STELLAR-GROUP/hpx/issues/71>

³⁸⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/105>

³⁸⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/107>

³⁸⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/108>

³⁸⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/110>

³⁸⁶² <https://github.com/STELLAR-GROUP/hpx/issues/116>

³⁸⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/162>

³⁸⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/172>

3865 <https://github.com/STELLAR-GROUP/hpx/issues/178>3866 <https://github.com/STELLAR-GROUP/hpx/issues/187>3867 <https://github.com/STELLAR-GROUP/hpx/issues/202>3868 <https://github.com/STELLAR-GROUP/hpx/issues/218>3869 <https://github.com/STELLAR-GROUP/hpx/issues/231>3870 <https://github.com/STELLAR-GROUP/hpx/issues/232>3871 <https://github.com/STELLAR-GROUP/hpx/issues/236>3872 <https://github.com/STELLAR-GROUP/hpx/issues/239>3873 <https://github.com/STELLAR-GROUP/hpx/issues/240>3874 <https://github.com/STELLAR-GROUP/hpx/issues/242>3875 <https://github.com/STELLAR-GROUP/hpx/issues/243>

- Issue #245³⁸⁷⁶ - `*_c_cache` code throws an exception if the queried GID is not in the local cache
- Issue #246³⁸⁷⁷ - Undefined references in dataflow/adaptive1d example
- Issue #252³⁸⁷⁸ - Problems configuring sheneos with CMake
- Issue #254³⁸⁷⁹ - Lifetime of components doesn't end when client goes out of scope
- Issue #259³⁸⁸⁰ - CMake does not detect that MSVC10 has lambdas
- Issue #260³⁸⁸¹ - `io_service_pool` segfault
- Issue #261³⁸⁸² - Late parcel executed outside of pxthread
- Issue #263³⁸⁸³ - Cannot select allocator with CMake
- Issue #264³⁸⁸⁴ - Fix allocator select
- Issue #267³⁸⁸⁵ - Runtime error for `hello_world`
- Issue #269³⁸⁸⁶ - `pthread_affinity_np` test fails to compile
- Issue #270³⁸⁸⁷ - Compiler noise due to `-Wcast-qual`
- Issue #275³⁸⁸⁸ - Problem with configuration tests/include paths on Gentoo
- Issue #325³⁸⁸⁹ - Sheneos is 200-400 times slower than the fortran equivalent
- Issue #331³⁸⁹⁰ - `hpx::init` and `hpx_main()` should not depend on `program_options`
- Issue #333³⁸⁹¹ - Add doxygen support to CMake for doc toolchain
- Issue #340³⁸⁹² - Performance counters for parcels
- Issue #346³⁸⁹³ - Component loading error when running `hello_world` in distributed on MSVC2010
- Issue #362³⁸⁹⁴ - Missing initializer error
- Issue #363³⁸⁹⁵ - Parcel port serialization error
- Issue #366³⁸⁹⁶ - Parcel buffering leads to types incompatible exception
- Issue #368³⁸⁹⁷ - Scalable alternative to `rand()` needed for *HPX*
- Issue #369³⁸⁹⁸ - IB over IP is substantially slower than just using standard TCP/IP

³⁸⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/245>

³⁸⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/246>

³⁸⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/252>

³⁸⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/254>

³⁸⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/259>

³⁸⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/260>

³⁸⁸² <https://github.com/STELLAR-GROUP/hpx/issues/261>

³⁸⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/263>

³⁸⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/264>

³⁸⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/267>

³⁸⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/269>

³⁸⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/270>

³⁸⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/275>

³⁸⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/325>

³⁸⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/331>

³⁸⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/333>

³⁸⁹² <https://github.com/STELLAR-GROUP/hpx/issues/340>

³⁸⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/346>

³⁸⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/362>

³⁸⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/363>

³⁸⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/366>

³⁸⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/368>

³⁸⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/369>

- Issue #374³⁸⁹⁹ - `hpx::lcos::wait` should work with dataflows and arbitrary classes meeting the future interface
- Issue #375³⁹⁰⁰ - Conflicting/ambiguous overloads of `hpx::lcos::wait`
- Issue #376³⁹⁰¹ - `Find_HPX.cmake` should set CMake variable `HPX_FOUND` for out of tree builds
- Issue #377³⁹⁰² - `ShenEOS` interpolate bulk and `interpolate_one_bulk` are broken
- Issue #379³⁹⁰³ - Add support for distributed runs under SLURM
- Issue #382³⁹⁰⁴ - `_Unwind_Word` not declared in `boost.backtrace`
- Issue #387³⁹⁰⁵ - Doxygen should look only at list of specified files
- Issue #388³⁹⁰⁶ - Running `make install` on an out-of-tree application is broken
- Issue #391³⁹⁰⁷ - Out-of-tree application segfaults when running in `qsub`
- Issue #392³⁹⁰⁸ - Remove `HPX_NO_INSTALL` option from `cmake` build system
- Issue #396³⁹⁰⁹ - Pragma related warnings when compiling with older gcc versions
- Issue #399³⁹¹⁰ - Out of tree component build problems
- Issue #400³⁹¹¹ - Out of source builds on Windows: linker should not receive compiler flags
- Issue #401³⁹¹² - Out of source builds on Windows: components need to be linked with `hpx_serialization`
- Issue #404³⁹¹³ - gfortran fails to link automatically when fortran files are present
- Issue #405³⁹¹⁴ - Inability to specify linking order for external libraries
- Issue #406³⁹¹⁵ - Adapt action limits such that dataflow applications work without additional defines
- Issue #415³⁹¹⁶ - `locality_results` is not a member of `hpx::components::server`
- Issue #425³⁹¹⁷ - Breaking changes to `traits::*result` wrt `std::vector<id_type>`
- Issue #426³⁹¹⁸ - AUTOGLOB needs to be updated to support fortran

³⁸⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/374>

³⁹⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/375>

³⁹⁰¹ <https://github.com/STELLAR-GROUP/hpx/issues/376>

³⁹⁰² <https://github.com/STELLAR-GROUP/hpx/issues/377>

³⁹⁰³ <https://github.com/STELLAR-GROUP/hpx/issues/379>

³⁹⁰⁴ <https://github.com/STELLAR-GROUP/hpx/issues/382>

³⁹⁰⁵ <https://github.com/STELLAR-GROUP/hpx/issues/387>

³⁹⁰⁶ <https://github.com/STELLAR-GROUP/hpx/issues/388>

³⁹⁰⁷ <https://github.com/STELLAR-GROUP/hpx/issues/391>

³⁹⁰⁸ <https://github.com/STELLAR-GROUP/hpx/issues/392>

³⁹⁰⁹ <https://github.com/STELLAR-GROUP/hpx/issues/396>

³⁹¹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/399>

³⁹¹¹ <https://github.com/STELLAR-GROUP/hpx/issues/400>

³⁹¹² <https://github.com/STELLAR-GROUP/hpx/issues/401>

³⁹¹³ <https://github.com/STELLAR-GROUP/hpx/issues/404>

³⁹¹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/405>

³⁹¹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/406>

³⁹¹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/415>

³⁹¹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/425>

³⁹¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/426>

2.11.16 HPX V0.8.1 (Apr 21, 2012)

This is a point release including important bug fixes for *HPX V0.8.0 (Mar 23, 2012)*.

General changes

- *HPX* does not need to be installed anymore to be functional.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this point release:

- Issue #295³⁹¹⁹ - Don't require install path to be known at compile time.
- Issue #371³⁹²⁰ - Add hpx iostreams to standard build.
- Issue #384³⁹²¹ - Fix compilation with GCC 4.7.
- Issue #390³⁹²² - Remove keep_factory_alive startup call from ShenEOS; add shutdown call to H5close.
- Issue #393³⁹²³ - Thread affinity control is broken.

Bug fixes (commits)

Here is a list of the important commits included in this point release:

- r7642 - External: Fix backtrace memory violation.
- **r7775 - Components: Fix symbol visibility bug with component startup** providers. This prevents one components providers from overriding another components.
- r7778 - Components: Fix startup/shutdown provider shadowing issues.

2.11.17 HPX V0.8.0 (Mar 23, 2012)

We have had roughly 1000 commits since the last release and we have closed approximately 70 tickets (bugs, feature requests, etc.).

General changes

- Improved PBS support, allowing for arbitrary naming schemes of node-hostnames.
- Finished verification of the reference counting framework.
- Implemented decrement merging logic to optimize the distributed reference counting system.
- Restructured the LCO framework. Renamed `hpx::lcos::eager_future<>` and `hpx::lcos::lazy_future<>` into `hpx::lcos::packaged_task` and `hpx::lcos::deferred_packaged_task`. Split `hpx::lcos::promise` into `hpx::lcos::packaged_task` and `hpx::lcos::future`. Added 'local' futures (in namespace `hpx::lcos::local`).

³⁹¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/295>

³⁹²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/371>

³⁹²¹ <https://github.com/STELLAR-GROUP/hpx/issues/384>

³⁹²² <https://github.com/STELLAR-GROUP/hpx/issues/390>

³⁹²³ <https://github.com/STELLAR-GROUP/hpx/issues/393>

- Improved the general performance of local and remote action invocations. This (under certain circumstances) drastically reduces the number of copies created for each of the parameters and return values.
- Reworked the performance counter framework. Performance counters are now created only when needed, which reduces the overall resource requirements. The new framework allows for much more flexible creation and management of performance counters. The new sine example application demonstrates some of the capabilities of the new infrastructure.
- Added a buildbot-based continuous build system which gives instant, automated feedback on each commit to SVN.
- Added more automated tests to verify proper functioning of *HPX*.
- Started to create documentation for *HPX* and its API.
- Added documentation toolchain to the build system.
- Added dataflow LCO.
- Changed default *HPX* command line options to have `hpx:` prefix. For instance, the former option `--threads` is now `--hpx:threads`. This has been done to make ambiguities with possible application specific command line options as unlikely as possible. See the section [HPX Command Line Options](#) for a full list of available options.
- Added the possibility to define command line aliases. The former short (one-letter) command line options have been predefined as aliases for backwards compatibility. See the section [HPX Command Line Options](#) for a detailed description of command line option aliasing.
- Network connections are now cached based on the connected host. The number of simultaneous connections to a particular host is now limited. Parcels are buffered and bundled if all connections are in use.
- Added more refined thread affinity control. This is based on the external library Portable Hardware Locality (HWLOC).
- Improved support for Windows builds with CMake.
- Added support for components to register their own command line options.
- Added the possibility to register custom startup/shutdown functions for any component. These functions are guaranteed to be executed by an *HPX* thread.
- Added two new experimental thread schedulers: `hierarchy_scheduler` and `periodic_priority_scheduler`. These can be activated by using the command line options `--hpx:queuing=hierarchy` or `--hpx:queuing=periodic`.

Example applications

- Graph500 performance benchmark³⁹²⁴ (thanks to Matthew Anderson for contributing this application).
- GTC (Gyrokinetic Toroidal Code)³⁹²⁵: a skeleton for particle in cell type codes.
- Random Memory Access: an example demonstrating random memory accesses in a large array
- ShenEOS example³⁹²⁶, demonstrating partitioning of large read-only data structures and exposing an interpolation API.
- Sine performance counter demo.
- Accumulator examples demonstrating how to write and use *HPX* components.

³⁹²⁴ <http://www.graph500.org/>

³⁹²⁵ <http://www.nersc.gov/research-and-development/benchmarking-and-workload-characterization/nersc-6-benchmarks/gtc/>

³⁹²⁶ <http://stellarcollapse.org/equationofstate>

- Quickstart examples (like hello_world, fibonacci, quicksort, factorial, etc.) demonstrating simple *HPX* concepts which introduce some of the concepts in *HPX*.
- Load balancing and work stealing demos.

API changes

- Moved all local LCOs into a separate namespace `hpx::lcos::local` (for instance, `hpx::lcos::local_mutex` is now `hpx::lcos::local::mutex`).
- Replaced `hpx::actions::function` with `hpx::util::function`. Cleaned up related code.
- Removed `hpx::traits::handle_gid` and moved handling of global reference counts into the corresponding serialization code.
- Changed terminology: `prefix` is now called `locality_id`, renamed the corresponding API functions (such as `hpx::get_prefix`, which is now called `hpx::get_locality_id`).
- Adding `hpx::find_remote_localities`, and `hpx::get_num_localities`.
- Changed performance counter naming scheme to make it more bash friendly. The new performance counter naming scheme is now

```
/object{parentname#parentindex/instance#index}/counter#parameters
```

- Added `hpx::get_worker_thread_num` replacing `hpx::threadmanager_base::get_thread_num`.
- Renamed `hpx::get_num_os_threads` to `hpx::get_os_threads_count`.
- Added `hpx::threads::get_thread_count`.
- Restructured the Futures sub-system, renaming types in accordance with the terminology used by the C++11 ISO standard.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release:

- Issue #31³⁹²⁷ - Specialize `handle_gid<>` for examples and tests
- Issue #72³⁹²⁸ - Fix AGAS reference counting
- Issue #104³⁹²⁹ - heartbeat throws an exception when decrefing the performance counter it's watching
- Issue #111³⁹³⁰ - throttle causes an exception on the target application
- Issue #142³⁹³¹ - One failed component loading causes an unrelated component to fail
- Issue #165³⁹³² - Remote exception propagation bug in AGAS reference counting test
- Issue #186³⁹³³ - Test credit exhaustion/splitting (e.g. `prepare_gid` and symbol NS)
- Issue #188³⁹³⁴ - Implement remaining AGAS reference counting test cases

³⁹²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/31>

³⁹²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/72>

³⁹²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/104>

³⁹³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/111>

³⁹³¹ <https://github.com/STELLAR-GROUP/hpx/issues/142>

³⁹³² <https://github.com/STELLAR-GROUP/hpx/issues/165>

³⁹³³ <https://github.com/STELLAR-GROUP/hpx/issues/186>

³⁹³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/188>

- Issue #258³⁹³⁵ - No type checking of GIDs in stubs classes
- Issue #271³⁹³⁶ - Seg fault/shared pointer assertion in distributed code
- Issue #281³⁹³⁷ - CMake options need descriptive text
- Issue #283³⁹³⁸ - AGAS caching broken (gva_cache needs to be rewritten with ICL)
- Issue #285³⁹³⁹ - HPX_INSTALL root directory not the same as CMAKE_INSTALL_PREFIX
- Issue #286³⁹⁴⁰ - New segfault in dataflow applications
- Issue #289³⁹⁴¹ - Exceptions should only be logged if not handled
- Issue #290³⁹⁴² - c++11 tests failure
- Issue #293³⁹⁴³ - Build target for component libraries
- Issue #296³⁹⁴⁴ - Compilation error with Boost V1.49rc1
- Issue #298³⁹⁴⁵ - Illegal instructions on termination
- Issue #299³⁹⁴⁶ - gravity aborts with multiple threads
- Issue #301³⁹⁴⁷ - Build error with Boost trunk
- Issue #303³⁹⁴⁸ - Logging assertion failure in distributed runs
- Issue #304³⁹⁴⁹ - Exception ‘what’ strings are lost when exceptions from decode_parcel are reported
- Issue #306³⁹⁵⁰ - Performance counter user interface issues
- Issue #307³⁹⁵¹ - Logging exception in distributed runs
- Issue #308³⁹⁵² - Logging deadlocks in distributed
- Issue #309³⁹⁵³ - Reference counting test failures and exceptions
- Issue #311³⁹⁵⁴ - Merge AGAS remote_interface with the runtime_support object
- Issue #314³⁹⁵⁵ - Object tracking for id_types
- Issue #315³⁹⁵⁶ - Remove handle_gid and handle credit splitting in id_type serialization
- Issue #320³⁹⁵⁷ - applier::get_locality_id() should return an error value (or throw an exception)

³⁹³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/258>

³⁹³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/271>

³⁹³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/281>

³⁹³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/283>

³⁹³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/285>

³⁹⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/286>

³⁹⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/289>

³⁹⁴² <https://github.com/STELLAR-GROUP/hpx/issues/290>

³⁹⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/293>

3944 <https://github.com/STELLAR-GROUP/hpx/issues/296>3945 <https://github.com/STELLAR-GROUP/hpx/issues/298>3946 <https://github.com/STELLAR-GROUP/hpx/issues/299>3947 <https://github.com/STELLAR-GROUP/hpx/issues/301>3948 <https://github.com/STELLAR-GROUP/hpx/issues/303>3949 <https://github.com/STELLAR-GROUP/hpx/issues/304>3950 <https://github.com/STELLAR-GROUP/hpx/issues/306>3951 <https://github.com/STELLAR-GROUP/hpx/issues/307>3952 <https://github.com/STELLAR-GROUP/hpx/issues/308>3953 <https://github.com/STELLAR-GROUP/hpx/issues/309>3954 <https://github.com/STELLAR-GROUP/hpx/issues/311>3955 <https://github.com/STELLAR-GROUP/hpx/issues/314>3956 <https://github.com/STELLAR-GROUP/hpx/issues/315>3957 <https://github.com/STELLAR-GROUP/hpx/issues/320>

- Issue #321³⁹⁵⁸ - Optimization for id_types which are never split should be restored
- Issue #322³⁹⁵⁹ - Command line processing ignored with Boost 1.47.0
- Issue #323³⁹⁶⁰ - Credit exhaustion causes object to stay alive
- Issue #324³⁹⁶¹ - Duplicate exception messages
- Issue #326³⁹⁶² - Integrate Quickbook with CMake
- Issue #329³⁹⁶³ - --help and –version should still work
- Issue #330³⁹⁶⁴ - Create pkg-config files
- Issue #337³⁹⁶⁵ - Improve usability of performance counter timestamps
- Issue #338³⁹⁶⁶ - Non-std exceptions deriving from std::exceptions in tfunc may be sliced
- Issue #339³⁹⁶⁷ - Decrease the number of send_pending_parcels threads
- Issue #343³⁹⁶⁸ - Dynamically setting the stack size doesn't work
- Issue #351³⁹⁶⁹ - ‘make install’ does not update documents
- Issue #353³⁹⁷⁰ - Disable FIXMEs in the docs by default; add a doc developer CMake option to enable FIXMEs
- Issue #355³⁹⁷¹ - ‘make’ doesn’t do anything after correct configuration
- Issue #356³⁹⁷² - Don’t use hpx::util::static_ in topology code
- Issue #359³⁹⁷³ - Infinite recursion in hpx::tuple serialization
- Issue #361³⁹⁷⁴ - Add compile time option to disable logging completely
- Issue #364³⁹⁷⁵ - Installation seriously broken in r7443

2.11.18 HPX V0.7.0 (Dec 12, 2011)

We have had roughly 1000 commits since the last release and we have closed approximately 120 tickets (bugs, feature requests, etc.).

³⁹⁵⁸ <https://github.com/STELLAR-GROUP/hpx/issues/321>
³⁹⁵⁹ <https://github.com/STELLAR-GROUP/hpx/issues/322>
³⁹⁶⁰ <https://github.com/STELLAR-GROUP/hpx/issues/323>
³⁹⁶¹ <https://github.com/STELLAR-GROUP/hpx/issues/324>
³⁹⁶² <https://github.com/STELLAR-GROUP/hpx/issues/326>
³⁹⁶³ <https://github.com/STELLAR-GROUP/hpx/issues/329>
³⁹⁶⁴ <https://github.com/STELLAR-GROUP/hpx/issues/330>
³⁹⁶⁵ <https://github.com/STELLAR-GROUP/hpx/issues/337>
³⁹⁶⁶ <https://github.com/STELLAR-GROUP/hpx/issues/338>
³⁹⁶⁷ <https://github.com/STELLAR-GROUP/hpx/issues/339>
³⁹⁶⁸ <https://github.com/STELLAR-GROUP/hpx/issues/343>
³⁹⁶⁹ <https://github.com/STELLAR-GROUP/hpx/issues/351>
³⁹⁷⁰ <https://github.com/STELLAR-GROUP/hpx/issues/353>
³⁹⁷¹ <https://github.com/STELLAR-GROUP/hpx/issues/355>
³⁹⁷² <https://github.com/STELLAR-GROUP/hpx/issues/356>
³⁹⁷³ <https://github.com/STELLAR-GROUP/hpx/issues/359>
³⁹⁷⁴ <https://github.com/STELLAR-GROUP/hpx/issues/361>
³⁹⁷⁵ <https://github.com/STELLAR-GROUP/hpx/issues/364>

General changes

- Completely removed code related to deprecated AGAS V1, started to work on AGAS V2.1.
- Started to clean up and streamline the exposed APIs (see ‘API changes’ below for more details).
- Revamped and unified performance counter framework, added a lot of new performance counter instances for monitoring of a diverse set of internal *HPX* parameters (queue lengths, access statistics, etc.).
- Improved general error handling and logging support.
- Fixed several race conditions, improved overall stability, decreased memory footprint, improved overall performance (major optimizations include native TLS support and ranged-based AGAS caching).
- Added support for running *HPX* applications with PBS.
- Many updates to the build system, added support for gcc 4.5.x and 4.6.x, added C++11 support.
- Many updates to default command line options.
- Added many tests, set up buildbot for continuous integration testing.
- Better shutdown handling of distributed applications.

Example applications

- quickstart/factorial and quickstart/fibonacci, future-recursive parallel algorithms.
- quickstart/hello_world, distributed hello world example.
- quickstart/rma, simple remote memory access example
- quickstart/quicksort, parallel quicksort implementation.
- gtc, gyrokinetic torodial code.
- bfs, breadth-first-search, example code for a graph application.
- sheneos, partitioning of large data sets.
- accumulator, simple component example.
- balancing/os_thread_num, balancing/px_thread_phase, examples demonstrating load balancing and work stealing.

API changes

- Added `hpx::find_all_localities`.
- Added `hpx::terminate` for non-graceful termination of applications.
- Added `hpx::lcos::async` functions for simpler asynchronous programming.
- Added new AGAS interface for handling of symbolic namespace (`hpx::agas::*`).
- Renamed `hpx::components::wait` to `hpx::lcos::wait`.
- Renamed `hpx::lcos::future_value` to `hpx::lcos::promise`.
- Renamed `hpx::lcos::recursive_mutex` to `hpx::lcos::local_recursive_mutex`, `hpx::lcos::mutex` to `hpx::lcos::local_mutex`
- Removed support for Boost versions older than V1.38, recommended Boost version is now V1.47 and newer.
- Removed `hpx::process` (this will be replaced by a real process implementation in the future).

- Removed non-functional LCO code (`hpx::lcos::dataflow`, `hpx::lcos::thunk`, `hpx::lcos::dataflow_variable`).
- Removed deprecated `hpx::naming::full_address`.

Bug fixes (closed tickets)

Here is a list of the important tickets we closed for this release:

- Issue #28³⁹⁷⁶ - Integrate Windows/Linux CMake code for *HPX* core
- Issue #32³⁹⁷⁷ - `hpx::cout()` should be `hpx::cout`
- Issue #33³⁹⁷⁸ - AGAS V2 legacy client does not properly handle `error_code`
- Issue #60³⁹⁷⁹ - AGAS: allow for registerid to optionally take ownership of the gid
- Issue #62³⁹⁸⁰ - adaptive1d compilation failure in Fusion
- Issue #64³⁹⁸¹ - Parcel subsystem doesn't resolve domain names
- Issue #83³⁹⁸² - No error handling if no console is available
- Issue #84³⁹⁸³ - No error handling if a hosted locality is treated as the bootstrap server
- Issue #90³⁹⁸⁴ - Add general commandline option `-N`
- Issue #91³⁹⁸⁵ - Add possibility to read command line arguments from file
- Issue #92³⁹⁸⁶ - Always log exceptions/errors to the log file
- Issue #93³⁹⁸⁷ - Log the command line/program name
- Issue #95³⁹⁸⁸ - Support for distributed launches
- Issue #97³⁹⁸⁹ - Attempt to create a bad component type in AMR examples
- Issue #100³⁹⁹⁰ - factorial and factorial_get examples trigger AGAS component type assertions
- Issue #101³⁹⁹¹ - Segfault when `hpx::process::here()` is called in fibonacci2
- Issue #102³⁹⁹² - unknown_component_address in int_object_semaphore_client
- Issue #114³⁹⁹³ - marduk raises assertion with default parameters
- Issue #115³⁹⁹⁴ - Logging messages for SMP runs (on the console) shouldn't be buffered

³⁹⁷⁶ <https://github.com/STELLAR-GROUP/hpx/issues/28>

³⁹⁷⁷ <https://github.com/STELLAR-GROUP/hpx/issues/32>

³⁹⁷⁸ <https://github.com/STELLAR-GROUP/hpx/issues/33>

³⁹⁷⁹ <https://github.com/STELLAR-GROUP/hpx/issues/60>

³⁹⁸⁰ <https://github.com/STELLAR-GROUP/hpx/issues/62>

³⁹⁸¹ <https://github.com/STELLAR-GROUP/hpx/issues/64>

³⁹⁸² <https://github.com/STELLAR-GROUP/hpx/issues/83>

³⁹⁸³ <https://github.com/STELLAR-GROUP/hpx/issues/84>

³⁹⁸⁴ <https://github.com/STELLAR-GROUP/hpx/issues/90>

³⁹⁸⁵ <https://github.com/STELLAR-GROUP/hpx/issues/91>

³⁹⁸⁶ <https://github.com/STELLAR-GROUP/hpx/issues/92>

³⁹⁸⁷ <https://github.com/STELLAR-GROUP/hpx/issues/93>

³⁹⁸⁸ <https://github.com/STELLAR-GROUP/hpx/issues/95>

³⁹⁸⁹ <https://github.com/STELLAR-GROUP/hpx/issues/97>

³⁹⁹⁰ <https://github.com/STELLAR-GROUP/hpx/issues/100>

³⁹⁹¹ <https://github.com/STELLAR-GROUP/hpx/issues/101>

³⁹⁹² <https://github.com/STELLAR-GROUP/hpx/issues/102>

³⁹⁹³ <https://github.com/STELLAR-GROUP/hpx/issues/114>

³⁹⁹⁴ <https://github.com/STELLAR-GROUP/hpx/issues/115>

- Issue #119³⁹⁹⁵ - marduk linking strategy breaks other applications
- Issue #121³⁹⁹⁶ - pbsdsh problem
- Issue #123³⁹⁹⁷ - marduk, dataflow and adaptive1d fail to build
- Issue #124³⁹⁹⁸ - Lower default preprocessing arity
- Issue #125³⁹⁹⁹ - Move hpx::detail::diagnostic_information out of the detail namespace
- Issue #126⁴⁰⁰⁰ - Test definitions for AGAS reference counting
- Issue #128⁴⁰⁰¹ - Add averaging performance counter
- Issue #129⁴⁰⁰² - Error with endian.hpp while building adaptive1d
- Issue #130⁴⁰⁰³ - Bad initialization of performance counters
- Issue #131⁴⁰⁰⁴ - Add global startup/shutdown functions to component modules
- Issue #132⁴⁰⁰⁵ - Avoid using auto_ptr
- Issue #133⁴⁰⁰⁶ - On Windows hpx.dll doesn't get installed
- Issue #134⁴⁰⁰⁷ - HPX_LIBRARY does not reflect real library name (on Windows)
- Issue #135⁴⁰⁰⁸ - Add detection of unique_ptr to build system
- Issue #137⁴⁰⁰⁹ - Add command line option allowing to repeatedly evaluate performance counters
- Issue #139⁴⁰¹⁰ - Logging is broken
- Issue #140⁴⁰¹¹ - CMake problem on windows
- Issue #141⁴⁰¹² - Move all non-component libraries into \$PREFIX/lib/hpx
- Issue #143⁴⁰¹³ - adaptive1d throws an exception with the default command line options
- Issue #146⁴⁰¹⁴ - Early exception handling is broken
- Issue #147⁴⁰¹⁵ - Sheneos doesn't link on Linux
- Issue #149⁴⁰¹⁶ - sheneos_test hangs
- Issue #154⁴⁰¹⁷ - Compilation fails for r5661

³⁹⁹⁵ <https://github.com/STELLAR-GROUP/hpx/issues/119>

³⁹⁹⁶ <https://github.com/STELLAR-GROUP/hpx/issues/121>

³⁹⁹⁷ <https://github.com/STELLAR-GROUP/hpx/issues/123>

³⁹⁹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/124>

³⁹⁹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/125>

⁴⁰⁰⁰ <https://github.com/STELLAR-GROUP/hpx/issues/126>

4001 <https://github.com/STELLAR-GROUP/hpx/issues/128>4002 <https://github.com/STELLAR-GROUP/hpx/issues/129>4003 <https://github.com/STELLAR-GROUP/hpx/issues/130>4004 <https://github.com/STELLAR-GROUP/hpx/issues/131>4005 <https://github.com/STELLAR-GROUP/hpx/issues/132>4006 <https://github.com/STELLAR-GROUP/hpx/issues/133>4007 <https://github.com/STELLAR-GROUP/hpx/issues/134>4008 <https://github.com/STELLAR-GROUP/hpx/issues/135>4009 <https://github.com/STELLAR-GROUP/hpx/issues/137>4010 <https://github.com/STELLAR-GROUP/hpx/issues/139>4011 <https://github.com/STELLAR-GROUP/hpx/issues/140>4012 <https://github.com/STELLAR-GROUP/hpx/issues/141>4013 <https://github.com/STELLAR-GROUP/hpx/issues/143>4014 <https://github.com/STELLAR-GROUP/hpx/issues/146>4015 <https://github.com/STELLAR-GROUP/hpx/issues/147>4016 <https://github.com/STELLAR-GROUP/hpx/issues/149>4017 <https://github.com/STELLAR-GROUP/hpx/issues/154>

- Issue #155⁴⁰¹⁸ - Sine performance counters example chokes on chrono headers
- Issue #156⁴⁰¹⁹ - Add build type to –version
- Issue #157⁴⁰²⁰ - Extend AGAS caching to store gid ranges
- Issue #158⁴⁰²¹ - r5691 doesn't compile
- Issue #160⁴⁰²² - Re-add AGAS function for resolving a locality to its prefix
- Issue #168⁴⁰²³ - Managed components should be able to access their own GID
- Issue #169⁴⁰²⁴ - Rewrite AGAS future pool
- Issue #179⁴⁰²⁵ - Complete switch to request class for AGAS server interface
- Issue #182⁴⁰²⁶ - Sine performance counter is loaded by other examples
- Issue #185⁴⁰²⁷ - Write tests for symbol namespace reference counting
- Issue #191⁴⁰²⁸ - Assignment of read-only variable in point_geometry
- Issue #200⁴⁰²⁹ - Seg faults when querying performance counters
- Issue #204⁴⁰³⁰ - –ifnames and suffix stripping needs to be more generic
- Issue #205⁴⁰³¹ - –list-* and –print-counter-* options do not work together and produce no warning
- Issue #207⁴⁰³² - Implement decrement entry merging
- Issue #208⁴⁰³³ - Replace the spinlocks in AGAS with hpx::lcos::local_mutexes
- Issue #210⁴⁰³⁴ - Add an –ifprefix option
- Issue #214⁴⁰³⁵ - Performance test for PX-thread creation
- Issue #216⁴⁰³⁶ - VS2010 compilation
- Issue #222⁴⁰³⁷ - r6045 context_linux_x86.hpp
- Issue #223⁴⁰³⁸ - fibonacci hangs when changing the state of an active thread
- Issue #225⁴⁰³⁹ - Active threads end up in the FEB wait queue
- Issue #226⁴⁰⁴⁰ - VS Build Error for Accumulator Client

⁴⁰¹⁸ <https://github.com/STELLAR-GROUP/hpx/issues/155>

⁴⁰¹⁹ <https://github.com/STELLAR-GROUP/hpx/issues/156>

⁴⁰²⁰ <https://github.com/STELLAR-GROUP/hpx/issues/157>

⁴⁰²¹ <https://github.com/STELLAR-GROUP/hpx/issues/158>

⁴⁰²² <https://github.com/STELLAR-GROUP/hpx/issues/160>

⁴⁰²³ <https://github.com/STELLAR-GROUP/hpx/issues/168>

⁴⁰²⁴ <https://github.com/STELLAR-GROUP/hpx/issues/169>

⁴⁰²⁵ <https://github.com/STELLAR-GROUP/hpx/issues/179>

⁴⁰²⁶ <https://github.com/STELLAR-GROUP/hpx/issues/182>

⁴⁰²⁷ <https://github.com/STELLAR-GROUP/hpx/issues/185>

⁴⁰²⁸ <https://github.com/STELLAR-GROUP/hpx/issues/191>

⁴⁰²⁹ <https://github.com/STELLAR-GROUP/hpx/issues/200>

⁴⁰³⁰ <https://github.com/STELLAR-GROUP/hpx/issues/204>

⁴⁰³¹ <https://github.com/STELLAR-GROUP/hpx/issues/205>

⁴⁰³² <https://github.com/STELLAR-GROUP/hpx/issues/207>

⁴⁰³³ <https://github.com/STELLAR-GROUP/hpx/issues/208>

⁴⁰³⁴ <https://github.com/STELLAR-GROUP/hpx/issues/210>

⁴⁰³⁵ <https://github.com/STELLAR-GROUP/hpx/issues/214>

⁴⁰³⁶ <https://github.com/STELLAR-GROUP/hpx/issues/216>

⁴⁰³⁷ <https://github.com/STELLAR-GROUP/hpx/issues/222>

⁴⁰³⁸ <https://github.com/STELLAR-GROUP/hpx/issues/223>

⁴⁰³⁹ <https://github.com/STELLAR-GROUP/hpx/issues/225>

⁴⁰⁴⁰ <https://github.com/STELLAR-GROUP/hpx/issues/226>

- Issue #228⁴⁰⁴¹ - Move all traits into namespace hpx::traits
- Issue #229⁴⁰⁴² - Invalid initialization of reference in thread_init_data
- Issue #235⁴⁰⁴³ - Invalid GID in iostreams
- Issue #238⁴⁰⁴⁴ - Demangle type names for the default implementation of get_action_name
- Issue #241⁴⁰⁴⁵ - C++11 support breaks GCC 4.5
- Issue #247⁴⁰⁴⁶ - Reference to temporary with GCC 4.4
- Issue #248⁴⁰⁴⁷ - Seg fault at shutdown with GCC 4.4
- Issue #253⁴⁰⁴⁸ - Default component action registration kills compiler
- Issue #272⁴⁰⁴⁹ - G++ unrecognized command line option
- Issue #273⁴⁰⁵⁰ - quicksort example doesn't compile
- Issue #277⁴⁰⁵¹ - Invalid CMake logic for Windows

2.12 About HPX

2.12.1 History

The development of High Performance ParalleX (*HPX*) began in 2007. At that time, Hartmut Kaiser became interested in the work done by the ParalleX group at the Center for Computation and Technology (CCT)⁴⁰⁵², a multi-disciplinary research institute at Louisiana State University (LSU)⁴⁰⁵³. The ParalleX group was working to develop a new and experimental execution model for future high performance computing architectures. This model was christened ParalleX. The first implementations of ParalleX were crude, and many of those designs had to be discarded entirely. However, over time the team learned quite a bit about how to design a parallel, distributed runtime system which implements the concepts of ParalleX.

From the very beginning, this endeavour has been a group effort. In addition to a handful of interested researchers, there have always been graduate and undergraduate students participating in the discussions, design, and implementation of *HPX*. In 2011 we decided to formalize our collective research efforts by creating the STE||AR⁴⁰⁵⁴ group (Systems Technology, Emergent Parallelism, and Algorithm Research). Over time, the team grew to include researchers around the country and the world. In 2014, the STE||AR⁴⁰⁵⁵ Group was reorganized to become the international community it is today. This consortium of researchers aims to develop stable, sustainable, and scalable tools which will enable application developers to exploit the parallelism latent in the machines of today and tomorrow. Our goal of the *HPX* project is to create a high quality, freely available, open source implementation of ParalleX concepts for conventional and future systems by building a modular and standards conforming runtime system for SMP and distributed application environments. The API exposed by *HPX* is conformant to the interfaces defined by the

⁴⁰⁴¹ <https://github.com/STELLAR-GROUP/hpx/issues/228>

⁴⁰⁴² <https://github.com/STELLAR-GROUP/hpx/issues/229>

⁴⁰⁴³ <https://github.com/STELLAR-GROUP/hpx/issues/235>

⁴⁰⁴⁴ <https://github.com/STELLAR-GROUP/hpx/issues/238>

⁴⁰⁴⁵ <https://github.com/STELLAR-GROUP/hpx/issues/241>

⁴⁰⁴⁶ <https://github.com/STELLAR-GROUP/hpx/issues/247>

⁴⁰⁴⁷ <https://github.com/STELLAR-GROUP/hpx/issues/248>

⁴⁰⁴⁸ <https://github.com/STELLAR-GROUP/hpx/issues/253>

⁴⁰⁴⁹ <https://github.com/STELLAR-GROUP/hpx/issues/272>

⁴⁰⁵⁰ <https://github.com/STELLAR-GROUP/hpx/issues/273>

⁴⁰⁵¹ <https://github.com/STELLAR-GROUP/hpx/issues/277>

⁴⁰⁵² <https://www.cct.lsu.edu>

⁴⁰⁵³ <https://www.lsu.edu>

⁴⁰⁵⁴ <https://stellar-group.org>

⁴⁰⁵⁵ <https://stellar-group.org>

C++11/14 ISO standard and adheres to the programming guidelines used by the [Boost⁴⁰⁵⁶](#) collection of C++ libraries. We steer the development of *HPX* with real world applications and aim to provide a smooth migration path for domain scientists.

To learn more about [STE||AR⁴⁰⁵⁷](#) and ParalleX, see [People](#) and [Why HPX?](#).

2.12.2 People

The [STE||AR⁴⁰⁵⁸](#) Group (pronounced as stellar) stands for “Systems Technology, Emergent Parallelism, and Algorithm Research”. We are an international group of faculty, researchers, and students working at various institutions around the world. The goal of the [STE||AR⁴⁰⁵⁹](#) Group is to promote the development of scalable parallel applications by providing a community for ideas, a framework for collaboration, and a platform for communicating these concepts to the broader community.

Our work is focused on building technologies for scalable parallel applications. *HPX*, our general purpose C++ runtime system for parallel and distributed applications, is no exception. We use *HPX* for a broad range of scientific applications, helping scientists and developers to write code which scales better and shows better performance compared to more conventional programming models such as MPI.

HPX is based on *ParalleX* which is a new (and still experimental) parallel execution model aiming to overcome the limitations imposed by the current hardware and the techniques we use to write applications today. Our group focuses on two types of applications - those requiring excellent strong scaling, allowing for a dramatic reduction of execution time for fixed workloads and those needing highest level of sustained performance through massive parallelism. These applications are presently unable (through conventional practices) to effectively exploit a relatively small number of cores in a multi-core system. By extension, these application will not be able to exploit high-end exascale computing systems which are likely to employ hundreds of millions of such cores by the end of this decade.

Critical bottlenecks to the effective use of new generation high performance computing (HPC) systems include:

- *Starvation*: due to lack of usable application parallelism and means of managing it,
- *Overhead*: reduction to permit strong scalability, improve efficiency, and enable dynamic resource management,
- *Latency*: from remote access across system or to local memories,
- *Contention*: due to multicore chip I/O pins, memory banks, and system interconnects.

The *ParalleX* model has been devised to address these challenges by enabling a new computing dynamic through the application of message-driven computation in a global address space context with lightweight synchronization. The work on *HPX* is centered around implementing the concepts as defined by the *ParalleX* model. *HPX* is currently targeted at conventional machines, such as classical Linux based Beowulf clusters and SMP nodes.

We fully understand that the success of *HPX* (and *ParalleX*) is very much the result of the work of many people. To see a list of who is contributing see our tables below.

⁴⁰⁵⁶ <https://www.boost.org/>

⁴⁰⁵⁷ <https://stellar-group.org>

⁴⁰⁵⁸ <https://stellar-group.org>

⁴⁰⁵⁹ <https://stellar-group.org>

HPX contributors

Table 2.39: Contributors

Name	Institution	Email
Hartmut Kaiser	Center for Computation and Technology (CCT) ⁴⁰⁶⁰ , Louisiana State University (LSU) ⁴⁰⁶¹	hkaiser@cct.lsu.edu
Thomas Heller	Department of Computer Science 3 - Computer Architecture ⁴⁰⁶² , Friedrich-Alexander University Erlangen-Nuremberg (FAU) ⁴⁰⁶³	thom.heller@gmail.com
Agustin Berge	Center for Computation and Technology (CCT) ⁴⁰⁶⁴ , Louisiana State University (LSU) ⁴⁰⁶⁵	kaballo86@hotmail.com
Mikael Simberg	Swiss National Supercomputing Centre ⁴⁰⁶⁶	simbergm@scs.ch
John Biddiscombe	Swiss National Supercomputing Centre ⁴⁰⁶⁷	biddisco@scs.ch
Anton Bikineev	Center for Computation and Technology (CCT) ⁴⁰⁶⁸ , Louisiana State University (LSU) ⁴⁰⁶⁹	ant.bikineev@gmail.com
Martin Stumpf	Department of Computer Science 3 - Computer Architecture ⁴⁰⁷⁰ , Friedrich-Alexander University Erlangen-Nuremberg (FAU) ⁴⁰⁷¹	martin.h.stumpf@gmail.com
Bryce Adelstein Lelbach	NVIDIA ⁴⁰⁷²	brycelelbach@gmail.com
Shuangyang Yang	Center for Computation and Technology (CCT) ⁴⁰⁷³ , Louisiana State University (LSU) ⁴⁰⁷⁴	syang16@cct.lsu.edu
Jeroen Habraken	Technische Universiteit Eindhoven ⁴⁰⁷⁵	jhabraken@cct.lsu.edu
Steven Brandt	Center for Computation and Technology (CCT) ⁴⁰⁷⁶ , Louisiana State University (LSU) ⁴⁰⁷⁷	sbrandt@cct.lsu.edu
Antoine Tran Tan	Center for Computation and Technology (CCT) ⁴⁰⁷⁸ , Louisiana State University (LSU) ⁴⁰⁷⁹	antoine.trantan@iri.fr
Adrian Serio	Center for Computation and Technology (CCT) ⁴⁰⁸⁰ , Louisiana State University (LSU) ⁴⁰⁸¹	aserio@cct.lsu.edu
Maciej Brodowicz	Center for Research in Extreme Scale Technologies (CREST) ⁴⁰⁸² , Indiana University (IU) ⁴⁰⁸³	mbrodowi@indiana.edu

Contributors to this document

Table 2.40: Documentation authors

Name	Institution	Email
Hartmut Kaiser	Center for Computation and Technology (CCT) ⁴⁰⁸⁴ , Louisiana State University (LSU) ⁴⁰⁸⁵	hkaiser@cct.lsu.edu
Thomas Heller	Department of Computer Science 3 - Computer Architecture ⁴⁰⁸⁶ , Friedrich-Alexander University Erlangen-Nuremberg (FAU) ⁴⁰⁸⁷	thom.heller@gmail.com
Bryce Adelstein Lelbach	NVIDIA ⁴⁰⁸⁸	brycelelbach@gmail.com
Vinay C Amatya	Center for Computation and Technology (CCT) ⁴⁰⁸⁹ , Louisiana State University (LSU) ⁴⁰⁹⁰	vamatya@cct.lsu.edu
Steven Brandt	Center for Computation and Technology (CCT) ⁴⁰⁹¹ , Louisiana State University (LSU) ⁴⁰⁹²	sbrandt@cct.lsu.edu
Maciej Brodowicz	Center for Research in Extreme Scale Technologies (CREST) ⁴⁰⁹³ , Indiana University (IU) ⁴⁰⁹⁴	mbrodowi@indiana.edu
Adrian Serio	Center for Computation and Technology (CCT) ⁴⁰⁹⁵ , Louisiana State University (LSU) ⁴⁰⁹⁶	aserio@cct.lsu.edu

4060 <https://www.cct.lsu.edu>4061 <https://www.lsu.edu>4062 <https://www3.cs.fau.de>4063 <https://www.fau.de>4064 <https://www.cct.lsu.edu>4065 <https://www.lsu.edu>4066 <https://www.csес.ch>4067 <https://www.csес.ch>4068 <https://www.cct.lsu.edu>4069 <https://www.lsu.edu>4070 <https://www3.cs.fau.de>4071 <https://www.fau.de>4072 <https://nvidia.com/>4073 <https://www.cct.lsu.edu>4074 <https://www.lsu.edu>4075 <https://www.tui.nl>4076 <https://www.cct.lsu.edu>4077 <https://www.lsu.edu>4078 <https://www.cct.lsu.edu>4079 <https://www.lsu.edu>4080 <https://www.cct.lsu.edu>4081 <https://www.lsu.edu>4082 <https://pti.iu.edu>4083 <https://www.iu.edu>4084 <https://www.cct.lsu.edu>4085 <https://www.lsu.edu>4086 <https://www3.cs.fau.de>4087 <https://www.fau.de>4088 <https://nvidia.com/>4089 <https://www.cct.lsu.edu>4090 <https://www.lsu.edu>4091 <https://www.cct.lsu.edu>4092 <https://www.lsu.edu>4093 <https://pti.iu.edu>4094 <https://www.iu.edu>4095 <https://www.cct.lsu.edu>4096 <https://www.lsu.edu>

Acknowledgements

Thanks also to the following people who contributed directly or indirectly to the project through discussions, pull requests, documentation patches, etc.

- Auriane Reverdell ([Swiss National Supercomputing Centre⁴⁰⁹⁷](#)), for her tireless work on refactoring our CMake setup and modularizing *HPX*.
- Christopher Hinz, for his work on refactoring our CMake setup.
- Weile Wei, for fixing *HPX* builds with CUDA on Summit.
- Severin Strobl, for fixing our CMake setup related to linking and adding new entry points to the *HPX* runtime.
- Rebecca Stobaugh, for her major documentation review and contributions during and after the 2019 Google Season of Documentation.
- Jan Melech, for adding automatic serialization of simple structs.
- Austin McCartney, for adding concept emulation of the Ranges TS bidirectional and random access iterator concepts.
- Marco Diers, reporting and fixing issues related PMIx.
- Maximilian Bremer, for reporting multiple issues and extending the component migration tests.
- Piotr Mikolajczyk, for his improvements and fixes to the set and count algorithms.
- Grant Rostig, for reporting several deficiencies on our web pages.
- Jakub Golinowski, for implementing an *HPX* backend for OpenCV and in the process improving documentation and reporting issues.
- Mikael Simberg ([Swiss National Supercomputing Centre⁴⁰⁹⁸](#)), for his tireless help cleaning up and maintaining *HPX*.
- Tianyi Zhang, for his work on HPXMP.
- Shahrzad Shirzad, for her contributions related to Phylanx.
- Christopher Ogle, for his contributions to the parallel algorithms.
- Surya Priy, for his work with statistic performance counters.
- Anushi Maheshwari, for her work on random number generation.
- Bruno Pitrus, for his work with parallel algorithms.
- Nikunj Gupta, for rewriting the implementation of `hpx_main.hpp` and for his fixes for tests.
- Christopher Taylor, for his interest in *HPX* and the fixes he provided.
- Shoshana Jakobovits, for her work on the resource partitioner.
- Denis Blank, who re-wrote our unwrapped function to accept plain values arbitrary containers, and properly deal with nested futures.
- Ajai V. George, who implemented several of the parallel algorithms.
- Taeguk Kwon, who worked on implementing parallel algorithms as well as adapting the parallel algorithms to the Ranges TS.
- Zach Byerly ([Louisiana State University \(LSU\)⁴⁰⁹⁹](#)), who in his work developing applications on top of *HPX* opened tickets and contributed to the *HPX* examples.

⁴⁰⁹⁷ <https://www.cscs.ch>

⁴⁰⁹⁸ <https://www.cscs.ch>

⁴⁰⁹⁹ <https://www.lsu.edu>

- Daniel Estermann, for his work porting *HPX* to the Raspberry Pi.
- Alireza Kheirkhahan ([Louisiana State University \(LSU\)](#)⁴¹⁰⁰), who built and administered our local cluster as well as his work in distributed IO.
- Abhimanyu Rawat, who worked on stack overflow detection.
- David Pfander, who improved signal handling in *HPX*, provided his optimization expertise, and worked on incorporating the Vc vectorization into *HPX*.
- Denis Demidov, who contributed his insights with VexCL.
- Khalid Hasanov, who contributed changes which allowed to run *HPX* on 64Bit power-pc architectures.
- Zahra Khatami ([Louisiana State University \(LSU\)](#)⁴¹⁰¹), who contributed the prefetching iterators and the persistent auto chunking executor parameters implementation.
- Marcin Copik, who worked on implementing GPU support for C++AMP and HCC. He also worked on implementing a HCC backend for *HPX.Compute*.
- Minh-Khanh Do, who contributed the implementation of several segmented algorithms.
- Bibek Wagle ([Louisiana State University \(LSU\)](#)⁴¹⁰²), who worked on fixing and analyzing the performance of the *parcel* coalescing plugin in *HPX*.
- Lukas Troska, who reported several problems and contributed various test cases allowing to reproduce the corresponding issues.
- Andreas Schaefer, who worked on integrating his library ([LibGeoDecomp](#)⁴¹⁰³) with *HPX*. He reported various problems and submitted several patches to fix issues allowing for a better integration with [LibGeoDecomp](#)⁴¹⁰⁴.
- Satyaki Upadhyay, who contributed several examples to *HPX*.
- Brandon Cordes, who contributed several improvements to the inspect tool.
- Harris Brakmic, who contributed an extensive build system description for building *HPX* with Visual Studio.
- Parsa Amini ([Louisiana State University \(LSU\)](#)⁴¹⁰⁵), who refactored and simplified the implementation of *AGAS* in *HPX* and who works on its implementation and optimization.
- Luis Martinez de Bartolome who implemented a build system extension for *HPX* integrating it with the [Conan](#)⁴¹⁰⁶ C/C++ package manager.
- Vinay C Amatyia ([Louisiana State University \(LSU\)](#)⁴¹⁰⁷), who contributed to the documentation and provided some of the *HPX* examples.
- Kevin Huck and Nick Chaimov ([University of Oregon](#)⁴¹⁰⁸), who contributed the integration of APEX (Autonomic Performance Environment for eXascale) with *HPX*.
- Francisco Jose Tapia, who helped with implementing the parallel sort algorithm for *HPX*.
- Patrick Diehl, who worked on implementing CUDA support for our companion library targeting GPGPUs ([HPXCL](#)⁴¹⁰⁹).
- Eric Lemanissier contributed fixes to allow compilation using the MingW toolchain.

⁴¹⁰⁰ <https://www.lsu.edu>⁴¹⁰¹ <https://www.lsu.edu>⁴¹⁰² <https://www.lsu.edu>⁴¹⁰³ <https://www.libgeodecomp.org/>⁴¹⁰⁴ <https://www.libgeodecomp.org/>⁴¹⁰⁵ <https://www.lsu.edu>⁴¹⁰⁶ <https://www.conan.io/>⁴¹⁰⁷ <https://www.lsu.edu>⁴¹⁰⁸ <https://uoregon.edu/>⁴¹⁰⁹ <https://github.com/STELLAR-GROUP/hpxcl/>

- Nidhi Makhijani who helped cleaning up some enum consistencies in *HPX* and contributed to the resource manager used in the thread scheduling subsystem. She also worked on *HPX* in the context of the Google Summer of Code 2015.
- Larry Xiao, Devang Bacharwar, Marcin Copik, and Konstantin Kronfeldner who worked on *HPX* in the context of the Google Summer of Code program 2015.
- Daniel Bourgeois (Center for Computation and Technology (CCT)⁴¹¹⁰) who contributed to *HPX* the implementation of several parallel algorithms (as proposed by N4313⁴¹¹¹).
- Anuj Sharma and Christopher Bross (Department of Computer Science 3 - Computer Architecture⁴¹¹²), who worked on *HPX* in the context of the Google Summer of Code⁴¹¹³ program 2014.
- Martin Stumpf (Department of Computer Science 3 - Computer Architecture⁴¹¹⁴), who rebuilt our contiguous testing infrastructure (see the *HPX* Buildbot Website⁴¹¹⁵). Martin is also working on *HPXCL*⁴¹¹⁶ (mainly all work related to OpenCL⁴¹¹⁷) and implementing an *HPX* backend for POCL⁴¹¹⁸, a portable computing language solution based on OpenCL⁴¹¹⁹.
- Grant Mercer (University of Nevada, Las Vegas⁴¹²⁰), who helped creating many of the parallel algorithms (as proposed by N4313⁴¹²¹).
- Damond Howard (Louisiana State University (LSU)⁴¹²²), who works on *HPXCL*⁴¹²³ (mainly all work related to CUDA⁴¹²⁴).
- Christoph Junghans (Los Alamos National Lab), who helped making our buildsystem more portable.
- Antoine Tran Tan (Laboratoire de Recherche en Informatique, Paris), who worked on integrating *HPX* as a backend for NT2⁴¹²⁵. He also contributed an implementation of an API similar to Fortran co-arrays on top of *HPX*.
- John Biddiscombe (Swiss National Supercomputing Centre⁴¹²⁶), who helped with the BlueGene/Q port of *HPX*, implemented the parallel sort algorithm, and made several other contributions.
- Erik Schnetter (Perimeter Institute for Theoretical Physics), who greatly helped to make *HPX* more robust by submitting a large amount of problem reports, feature requests, and made several direct contributions.
- Mathias Gaunard (Metascale), who contributed several patches to reduce compile time warnings generated while compiling *HPX*.
- Andreas Buhr, who helped with improving our documentation, especially by suggesting some fixes for inconsistencies.
- Patricia Grubel (New Mexico State University⁴¹²⁷), who contributed the description of the different *HPX* thread scheduler policies and is working on the performance analysis of our thread scheduling subsystem.

⁴¹¹⁰ <https://www.cct.lsu.edu>

⁴¹¹¹ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

⁴¹¹² <https://www3.cs.fau.de>

⁴¹¹³ <https://developers.google.com/open-source/soc/>

⁴¹¹⁴ <https://www3.cs.fau.de>

⁴¹¹⁵ <http://rostam.cct.lsu.edu/>

⁴¹¹⁶ <https://github.com/STELLAR-GROUP/hpxcl/>

⁴¹¹⁷ <https://www.khronos.org/opencl/>

⁴¹¹⁸ <https://portablecl.org/>

⁴¹¹⁹ <https://www.khronos.org/opencl/>

⁴¹²⁰ <https://www.unlv.edu>

⁴¹²¹ <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2014/n4313.html>

⁴¹²² <https://www.lsu.edu>

⁴¹²³ <https://github.com/STELLAR-GROUP/hpxcl/>

⁴¹²⁴ https://www.nvidia.com/object/cuda_home_new.html

⁴¹²⁵ <https://www.numscale.com/nt2/>

⁴¹²⁶ <https://www.cscs.ch>

⁴¹²⁷ <https://www.nmsu.edu>

- Lars Viklund, whose wit, passion for testing, and love of odd architectures has been an amazing contribution to our team. He has also contributed platform specific patches for FreeBSD and MSVC12.
- Agustin Berge, who contributed patches fixing some very nasty hidden template meta-programming issues. He rewrote large parts of the API elements ensuring strict conformance with C++11/14.
- Anton Bikineev for contributing changes to make using `boost::lexical_cast` safer, he also contributed a thread safety fix to the iostreams module. He also contributed a complete rewrite of the serialization infrastructure replacing Boost.Serialization inside *HPX*.
- Pyry Jakkola, who contributed the Mac OS build system and build documentation on how to build *HPX* using Clang and libc++.
- Mario Mulansky, who created an *HPX* backend for his Boost.Odeint library, and who submitted several test cases allowing us to reproduce and fix problems in *HPX*.
- Rekha Raj, who contributed changes to the description of the Windows build instructions.
- Jeremy Kemp who worked on an *HPX* OpenMP backend and added regression tests.
- Alex Nagelberg for his work on implementing a C wrapper API for *HPX*.
- Chen Guo, helvihartmann, Nicholas Pezolano, and John West who added and improved examples in *HPX*.
- Joseph Kleinhenz, Markus Elfring, Kirill Kropivnyansky, Alexander Neundorf, Bryant Lam, and Alex Hirsch who improved our CMake.
- Tapasweni Pathak, Praveen Velliengiri, Jean-Loup Tastet, Michael Levine, Aalekh Nigam, HadrienG2, Prayag Verma, Islada, Alex Myczko, and Avyav Kumar who improved the documentation.
- Jayesh Badwaik, J. F. Bastien, Christoph Garth, Christopher Hinz, Brandon Kohn, Mario Lang, Maikel Nadolski, pierrele, hendrx, Dekken, woodmeister123, xaguilar, Andrew Kemp, Dylan Stark, Matthew Anderson, Jeremy Wilke, Jiazheng Yuan, CyberDrudge, david8dixon, Maxwell Reeser, Raffaele Solca, Marco Ippolito, Jules Penuchot, Weile Wei, Severin Strobl, Kor de Jong, albestro, Jeff Trull, and Yuri Victorovich who contributed to the general improvement of *HPX*.

In addition to the people who worked directly with *HPX* development we would like to acknowledge the NSF, DoE, DARPA, Center for Computation and Technology (CCT)⁴¹²⁸, Department of Computer Science 3 - Computer Architecture⁴¹²⁹, and Swiss National Supercomputing Centre⁴¹³⁰ who fund and support our work. We would also like to thank the following organizations for granting us allocations of their compute resources: LSU HPC, LONI, XSEDE, NERSC, and the Gauss Center for Supercomputing.

HPX is currently funded by the following grants:

- The National Science Foundation through awards 1240655 (STAR), 1339782 (STORM), and 1737785 (Phylanx). Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation.
- The Department of Energy (DoE) through the awards DE-AC52-06NA25396 (FLeCSI) and DE-NA0003525 (Resilience). Neither the United States Government nor any agency thereof, nor any of their employees, makes any warranty, express or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government or any agency thereof. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States Government or any agency thereof.

⁴¹²⁸ <https://www.cct.lsu.edu>

⁴¹²⁹ <https://www3.cs.fau.de>

⁴¹³⁰ <https://www.cscs.ch>

- The Defense Technical Information Center (DTIC) under contract FA8075-14-D-0002/0007. Neither the United States Government nor any agency thereof, nor any of their employees makes any warranty, express or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights.
- The Bavarian Research Foundation (Bayerische Forschungsstiftung) through the grant AZ-987-11.
- The European Commission's Horizon 2020 programme through the grant H2020-EU.1.2.2. 671603 (AllScale).

**CHAPTER
THREE**

INDEX

- genindex

INDEX

Symbols

--itt_heap_function (*C++ type*), 778
--hpx:affinity arg
 command line option, 131
--hpx:agas arg
 command line option, 130
--hpx:app-config arg
 command line option, 132
--hpx:attach-debugger arg
 command line option, 132
--hpx:bind arg
 command line option, 131
--hpx:config arg
 command line option, 132
--hpx:connect
 command line option, 130
--hpx:console
 command line option, 130
--hpx:cores arg
 command line option, 131
--hpx:debug-agas-log [arg]
 command line option, 132
--hpx:debug-app-log [arg]
 command line option, 132
--hpx:debug-clp
 command line option, 132
--hpx:debug-hpx-log [arg]
 command line option, 132
--hpx:debug-parcel-log [arg]
 command line option, 132
--hpx:debug-timing-log [arg]
 command line option, 132
--hpx:dump-config
 command line option, 132
--hpx:dump-config-initial
 command line option, 132
--hpx:endnodes
 command line option, 130
--hpx:exit
 command line option, 132
--hpx:expect-connecting-localities
 command line option, 131
--hpx:help
 command line option, 130
--hpx:high-priority-threads arg
 command line option, 131
--hpx:hpx arg
 command line option, 130
--hpx:ifprefix arg
 command line option, 130
--hpx:ifsuffix arg
 command line option, 130
--hpx:iftransform arg
 command line option, 131
--hpx:ignore-batch-env
 command line option, 131
--hpx:info
 command line option, 130
--hpx:ini arg
 command line option, 132
--hpx:list-component-types
 command line option, 132
--hpx:list-counter-infos
 command line option, 133
--hpx:list-counters
 command line option, 133
--hpx:list-symbolic-names
 command line option, 132
--hpx:localities arg
 command line option, 131
--hpx:no-csv-header
 command line option, 133
--hpx:node arg
 command line option, 131
--hpx:nodofile arg
 command line option, 130
--hpx:nodes arg
 command line option, 130
--hpx:numa-sensitive
 command line option, 131
--hpx:options-file arg
 command line option, 130
--hpx:print-bind
 command line option, 131

```
--hpx:print-counter
    command line option, 133
--hpx:print-counter-at arg
    command line option, 133
--hpx:print-counter-destination
    command line option, 133
--hpx:print-counter-format
    command line option, 133
--hpx:print-counter-interval
    command line option, 133
--hpx:print-counter-reset
    command line option, 133
--hpx:print-counters-locally
    command line option, 133
--hpx:pu-offset
    command line option, 131
--hpx:pu-step
    command line option, 131
--hpx:queuing arg
    command line option, 131
--hpx:reset-counters
    command line option, 133
--hpx:run-agas-server
    command line option, 130
--hpx:run-agas-server-only
    command line option, 130
--hpx:run-hpx-main
    command line option, 130
--hpx:threads arg
    command line option, 131
--hpx:use-process-mask
    command line option, 131
--hpx:version
    command line option, 130
--hpx:worker
    command line option, 130
```

A

Action, [18](#)
Active Global Address Space, [17](#)
AGAS, [17](#)
AMPLIFIER_ROOT:PATH
 command line option, 90
applier (*C++ type*), 271

B

boost (*C++ type*), [591](#), [592](#), [785](#)
boost::accumulators (*C++ type*), [591](#), [592](#)
boost::accumulators::extract (*C++ type*),
 [591](#), [592](#)
boost::accumulators::extract::histogram
 (*C++ member*), [592](#)
boost::accumulators::extract::rolling_max
 (*C++ member*), [591](#)

```
boost::accumulators::extract::rolling_min
    (C++ member), 592
boost::accumulators::tag (C++ type), 591,  
    592
boost::accumulators::tag::histogram
    (C++ class), 592
boost::accumulators::tag::histogram::impl
    (C++ class), 592
boost::accumulators::tag::histogram::impl::apply
    (C++ class), 592
boost::accumulators::tag::histogram::impl::apply::t
    (C++ type), 593
boost::accumulators::tag::rolling_max
    (C++ class), 591
boost::accumulators::tag::rolling_max::impl
    (C++ class), 591
boost::accumulators::tag::rolling_max::impl::apply
    (C++ class), 591
boost::accumulators::tag::rolling_max::impl::apply::t
    (C++ type), 591
boost::accumulators::tag::rolling_min
    (C++ class), 592
boost::accumulators::tag::rolling_min::impl
    (C++ class), 592
boost::accumulators::tag::rolling_min::impl::apply
    (C++ class), 592
boost::accumulators::tag::rolling_min::impl::apply::t
    (C++ type), 592
boost::lockfree (C++ type), 785
boost::lockfree::deque (C++ class), 785
boost::lockfree::deque::~deque (C++ func-
    tion), 785
boost::lockfree::deque::alloc_node (C++ 
    function), 786
boost::lockfree::deque::anchor (C++ type),
    785
boost::lockfree::deque::anchor_
    (C++ member), 786
boost::lockfree::deque::anchor_pair
    (C++ type), 785
boost::lockfree::deque::atomic_anchor_pair
    (C++ type), 785
boost::lockfree::deque::atomic_node_pointer
    (C++ type), 785
boost::lockfree::deque::dealloc_node
    (C++ function), 786
boost::lockfree::deque::deque (C++ func-
    tion), 785
boost::lockfree::deque::empty (C++ func-
    tion), 785
boost::lockfree::deque::HPX_NON_COPYABLE
    (C++ function), 785
boost::lockfree::deque::is_lock_free
    (C++ function), 785
```

boost::lockfree::deque::node (*C++ type*), 785
boost::lockfree::deque::node_allocator (*C++ type*), 785
boost::lockfree::deque::node_pointer (*C++ type*), 785
boost::lockfree::deque::pool (*C++ type*), 785
boost::lockfree::deque::pool_ (*C++ member*), 786
boost::lockfree::deque::pop_left (*C++ function*), 785
boost::lockfree::deque::pop_right (*C++ function*), 786
boost::lockfree::deque::push_left (*C++ function*), 785
boost::lockfree::deque::push_right (*C++ function*), 785
boost::lockfree::deque::stabilize (*C++ function*), 786
boost::lockfree::deque::stabilize_left (*C++ function*), 786
boost::lockfree::deque::stabilize_right (*C++ function*), 786
boost::lockfree::deque::tag_t (*C++ type*), 785
boost::lockfree::deque_anchor (*C++ class*), 786
boost::lockfree::deque_anchor::atomic_nodeopbinlockfree::lpush (*C++ enumerator*), 785
boost::lockfree::deque_anchor::atomic_pabost::lockfree::stable (*C++ enumerator*), 785
boost::lockfree::deque_anchor::cas (*C++ function*), 787
boost::lockfree::deque_anchor::dequeue_anchorbootstrap_logging (*C++ class*), 834
boost::lockfree::deque_anchor::is_lock_free (*C++ function*), 834
boost::lockfree::deque_anchor::left (*C++ function*), 786
boost::lockfree::deque_anchor::lrs (*C++ function*), 786
boost::lockfree::deque_anchor::node (*C++ type*), 786
boost::lockfree::deque_anchor::node_pointer (*C++ type*), 786
boost::lockfree::deque_anchor::operator!= (*C++ function*), 787
boost::lockfree::deque_anchor::operator== (*C++ function*), 787
boost::lockfree::deque_anchor::pair (*C++ type*), 786
boost::lockfree::deque_anchor::pair_ (*C++ member*), 787
boost::lockfree::deque_anchor::right (*C++ function*), 787
boost::lockfree::deque_anchor::status (*C++ function*), 787
boost::lockfree::deque_anchor::tag (*C++ function*), 787
boost::lockfree::deque_anchor::tag_t (*C++ type*), 786
boost::lockfree::deque_node (*C++ class*), 787
boost::lockfree::deque_node::atomic_pointer (*C++ type*), 787
boost::lockfree::deque_node::data (*C++ member*), 788
boost::lockfree::deque_node::deque_node (*C++ function*), 787
boost::lockfree::deque_node::left (*C++ member*), 788
boost::lockfree::deque_node::pointer (*C++ type*), 787
boost::lockfree::deque_node::right (*C++ member*), 788
boost::lockfree::deque_node::tag_t (*C++ type*), 787
boost::lockfree::deque_status_type (*C++ enum*), 785
boost::lockfree::deque<T, freelist_t, Alloc>::padding (*C++ member*), 786
boost::lockfree::lpush (*C++ enumerator*), 785
boost::lockfree::rpush (*C++ enumerator*), 785
boost::lockfree::stable (*C++ enumerator*), 785
BOOST_ROOT:PATH command line option, 90
BREATHE_APIDOC_ROOT:PATH command line option, 46
C
call_once (*C++ function*), 602
command line option
--hpx:affinity arg, 131
--hpx:agas arg, 130
--hpx:app-config arg, 132
--hpx:attach-debugger arg, 132
--hpx:bind arg, 131
--hpx:config arg, 132
--hpx:connect, 130
--hpx:console, 130
--hpx:cores arg, 131
--hpx:debug-agas-log [arg], 132
--hpx:debug-app-log [arg], 132

C

```
call_once (C++ function), 602
command line option
    --hpx:affinity arg, 131
    --hpx:agas arg, 130
    --hpx:app-config arg, 132
    --hpx:attach-debugger arg, 132
    --hpx:bind arg, 131
    --hpx:config arg, 132
    --hpx:connect, 130
    --hpx:console, 130
    --hpx:cores arg, 131
    --hpx:debug-agas-log [arg], 132
    --hpx:debug-app-log [arg], 132
```

```
--hpx:debug-clp, 132
--hpx:debug-hpx-log [arg], 132
--hpx:debug-parcel-log [arg], 132
--hpx:debug-timing-log [arg], 132
--hpx:dump-config, 132
--hpx:dump-config-initial, 132
--hpx:endnodes, 130
--hpx:exit, 132
--hpx:expect-connecting-localities,
    131
--hpx:help, 130
--hpx:high-priority-threads arg, 131
--hpx:hpx arg, 130
--hpx:ifprefix arg, 130
--hpx:ifsuffix arg, 130
--hpx:iftransform arg, 131
--hpx:ignore-batch-env, 131
--hpx:info, 130
--hpx:ini arg, 132
--hpx:list-component-types, 132
--hpx:list-counter-infos, 133
--hpx:list-counters, 133
--hpx:list-symbolic-names, 132
--hpx:localities arg, 131
--hpx:no-csv-header, 133
--hpx:node arg, 131
--hpx:nodefile arg, 130
--hpx:nodes arg, 130
--hpx:numa-sensitive, 131
--hpx:options-file arg, 130
--hpx:print-bind, 131
--hpx:print-counter, 133
--hpx:print-counter-at arg, 133
--hpx:print-counter-destination, 133
--hpx:print-counter-format, 133
--hpx:print-counter-interval, 133
--hpx:print-counter-reset, 133
--hpx:print-counters-locally, 133
--hpx:pu-offset, 131
--hpx:pu-step, 131
--hpx:queuing arg, 131
--hpx:reset-counters, 133
--hpx:run-agas-server, 130
--hpx:run-agas-server-only, 130
--hpx:run-hpx-main, 130
--hpx:threads arg, 131
--hpx:use-process-mask, 131
--hpx:version, 130
--hpx:worker, 130
AMPLIFIER_ROOT:PATH, 90
BOOST_ROOT:PATH, 90
BREATHE_APIDOC_ROOT:PATH, 46
DOXYGEN_ROOT:PATH, 45
HDF5_ROOT:PATH, 90

HPX_AFFINITY_WITH_COMPATIBILITY_HEADERS:BOOL,
    84
HPX_AFFINITY_WITH_DEPRECATED_WARNINGS:BOOL,
    84
HPX_AFFINITY_WITH_TESTS:BOOL, 84
HPX_ALGORITHMS_WITH_DEPRECATED_WARNINGS:BOOL,
    84
HPX_ALGORITHMS_WITH_TESTS:BOOL, 84
HPX_ALLOCATOR_SUPPORT_WITH_COMPATIBILITY_HEADER
    84
HPX_ALLOCATOR_SUPPORT_WITH_DEPRECATED_WARNINGS
    84
HPX_ALLOCATOR_SUPPORT_WITH_TESTS:BOOL,
    84
HPX_ASSERTION_WITH_COMPATIBILITY_HEADERS:BOOL,
    84
HPX_ASSERTION_WITH_DEPRECATED_WARNINGS:BOOL,
    84
HPX_ASSERTION_WITH_TESTS:BOOL, 84
HPX_BASIC_EXECUTION_WITH_COMPATIBILITY_HEADERS:
    84
HPX_BASIC_EXECUTION_WITH_DEPRECATED_WARNINGS:B
    84
HPX_BASIC_EXECUTION_WITH_TESTS:BOOL,
    84
HPX_BATCH_ENVIRONMENTS_WITH_COMPATIBILITY_HEADE
    84
HPX_BATCH_ENVIRONMENTS_WITH_DEPRECATED_WARNING
    84
HPX_BATCH_ENVIRONMENTS_WITH_TESTS:BOOL,
    84
HPX_CACHE_WITH_COMPATIBILITY_HEADERS:BOOL,
    84
HPX_CACHE_WITH_DEPRECATED_WARNINGS:BOOL,
    84
HPX_CACHE_WITH_TESTS:BOOL, 84
HPX_CHECKPOINT_WITH_COMPATIBILITY_HEADERS:BOOL,
    84
HPX_CHECKPOINT_WITH_DEPRECATED_WARNINGS:BOOL,
    84
HPX_CHECKPOINT_WITH_TESTS:BOOL, 85
HPX_COLLECTIVES_WITH_COMPATIBILITY_HEADERS:BOOL
    85
HPX_COLLECTIVES_WITH_DEPRECATED_WARNINGS:BOOL,
    85
HPX_COLLECTIVES_WITH_TESTS:BOOL, 85
HPX_COMPUTE_CUDA_WITH_DEPRECATED_WARNINGS:BOOL
    85
HPX_COMPUTE_CUDA_WITH_TESTS:BOOL, 85
HPX_COMPUTE_WITH_DEPRECATED_WARNINGS:BOOL,
    85
HPX_COMPUTE_WITH_TESTS:BOOL, 85
HPX_CONCEPTS_WITH_COMPATIBILITY_HEADERS:BOOL,
    85
```

HPX_CONCEPTS_WITH_DEPRECATED_WARNINGS:BOOL, 86
85 HPX_FUNCTIONAL_WITH_TESTS:BOOL, 86
HPX_CONCEPTS_WITH_TESTS:BOOL, 85 HPX_HARDWARE_WITH_COMPATIBILITY_HEADERS:BOOL,
HPX_CONCURRENCY_WITH_COMPATIBILITY_HEADERS:BOOL,
85 HPX_HARDWARE_WITH_DEPRECATED_WARNINGS:BOOL,
HPX_CONCURRENCY_WITH_DEPRECATED_WARNINGS:BOOL, 86
85 HPX_HARDWARE_WITH_TESTS:BOOL, 87
HPX_CONCURRENCY_WITH_TESTS:BOOL, 85 HPX_HASHING_WITH_COMPATIBILITY_HEADERS:BOOL,
HPX_CONFIG_WITH_COMPATIBILITY_HEADERS:BOOL, 87
85 HPX_HASHING_WITH_DEPRECATED_WARNINGS:BOOL,
HPX_CONFIG_WITH_DEPRECATED_WARNINGS:BOOL, 87
85 HPX_HASHING_WITH_TESTS:BOOL, 87
HPX_CONFIG_WITH_TESTS:BOOL, 85 HPX_ITERATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS
HPX_COROUTINES_WITH_COMPATIBILITY_HEADERS:BOOL,
85 HPX_ITERATOR_SUPPORT_WITH_DEPRECATED_WARNINGS:
HPX_COROUTINES_WITH_DEPRECATED_WARNINGS:BOOL,
85 HPX_ITERATOR_SUPPORT_WITH_TESTS:BOOL,
HPX_COROUTINES_WITH_TESTS:BOOL, 85 87
HPX_DATASTRUCTURES_WITH_ADAPTER_STD_TUPLE:BOOL LOCAL_LCOS_WITH_COMPATIBILITY_HEADERS:BOOL,
85 87
HPX_DATASTRUCTURES_WITH_COMPATIBILITY_HEADERS:BOOL LOCAL_LCOS_WITH_DEPRECATED_WARNINGS:BOOL,
86 87
HPX_DATASTRUCTURES_WITH_DEPRECATED_WARNINGS:BOOL LOCAL_LCOS_WITH_TESTS:BOOL, 87
86 HPX_LOGGING_WITH_COMPATIBILITY_HEADERS:BOOL,
HPX_DATASTRUCTURES_WITH_TESTS:BOOL, 86 87
HPX_LOGGING_WITH_DEPRECATED_WARNINGS:BOOL,
HPX_DEBUGGING_WITH_COMPATIBILITY_HEADERS:BOOL,
86 HPX_LOGGING_WITH_TESTS:BOOL, 87
HPX_DEBUGGING_WITH_DEPRECATED_WARNINGS:BOOL LOCAL_MEMORY_WITH_DEPRECATED_WARNINGS:BOOL,
86 87
HPX_DEBUGGING_WITH_TESTS:BOOL, 86 HPX_MEMORY_WITH_TESTS:BOOL, 87
HPX_ERRORS_WITH_COMPATIBILITY_HEADERS:BOOL HPX_NAMING_BASE_WITH_COMPATIBILITY_HEADERS:BOOL
86 87
HPX_ERRORS_WITH_DEPRECATED_WARNINGS:BOOL HPX_NAMING_BASE_WITH_DEPRECATED_WARNINGS:BOOL,
86 87
HPX_ERRORS_WITH_TESTS:BOOL, 86 HPX_NAMING_BASE_WITH_TESTS:BOOL, 87
HPX_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL PLUGIN_WITH_COMPATIBILITY_HEADERS:BOOL,
86 87
HPX_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL PLUGIN_WITH_DEPRECATED_WARNINGS:BOOL,
86 88
HPX_EXECUTION_WITH_TESTS:BOOL, 86 HPX_PLUGIN_WITH_TESTS:BOOL, 88
HPX_FILESYSTEM_WITH_BOOST_FILESYSTEM_COMPILER_PREPROCESSOR_WITH_COMPATIBILITY_HEADERS:BOOL
86 88
HPX_FILESYSTEM_WITH_DEPRECATED_WARNINGS:BOOL PLUGIN_WITH_DEPRECATED_WARNINGS:BOOL
86 88
HPX_FILESYSTEM_WITH_TESTS:BOOL, 86 HPX_PREPROCESSOR_WITH_TESTS:BOOL, 88
HPX_FORMAT_WITH_COMPATIBILITY_HEADERS:BOOL HPX_PROGRAM_OPTIONS_WITH_BOOST_PROGRAM_OPTIONS_
86 88
HPX_FORMAT_WITH_DEPRECATED_WARNINGS:BOOL HPX_PROGRAM_OPTIONS_WITH_DEPRECATED_WARNINGS:BOOL
86 88
HPX_FORMAT_WITH_TESTS:BOOL, 86 HPX_PROGRAM_OPTIONS_WITH_TESTS:BOOL,
HPX_FUNCTIONAL_WITH_COMPATIBILITY_HEADERS:BOOL,
86 HPX_RESILIENCY_WITH_DEPRECATED_WARNINGS:BOOL,
HPX_FUNCTIONAL_WITH_DEPRECATED_WARNINGS:BOOL,

HPX_RESILIENCY_WITH_TESTS:BOOL, 88
HPX_RESOURCE_PARTITIONER_WITH_COMPATIBILITY_HEADERS:BOOL,
88
HPX_RESOURCE_PARTITIONER_WITH_DEPRECATED_WARNINGS:BOOL,
88
HPX_RESOURCE_PARTITIONER_WITH_TESTS:BOOL, 89
HPX_TIMING_WITH_TESTS:BOOL, 89
HPX_TOPOLOGY_WITH_COMPATIBILITY_HEADERS:BOOL,
HPX_SCHEDULER_MAX_TERMINATED_THREADS:STRING, 89
76
HPX_SEGMENTED_ALGORITHMS_WITH_DEPRECATED_WARNINGS:BOOL,
88
HPX_SEGMENTED_ALGORITHMS_WITH_TESTS:BOOL, 89
HPX_TYPE_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL,
88
HPX_SERIALIZATION_WITH_BOOST_TYPES:BOOL, HPX_TYPE_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL
88
HPX_SERIALIZATION_WITH_COMPATIBILITY_HEADERS:BOOL,
88
HPX_UTIL_WITH_COMPATIBILITY_HEADERS:BOOL,
88
HPX_SERIALIZATION_WITH_DEPRECATED_WARNINGS:BOOL,
88
HPX_UTIL_WITH_DEPRECATED_WARNINGS:BOOL,
88
HPX_UTIL_WITH_TESTS:BOOL, 90
HPX_STATIC_REINIT_WITH_COMPATIBILITY_HEADERS:BOOL,
88
ACTION_BASE_COMPATIBILITY:BOOL,
88
HPX_STATIC_REINIT_WITH_DEPRECATED_WARNINGS:BOOL,
88
AGAS_DUMP_REFCNT_ENTRIES:BOOL,
78
HPX_STATIC_REINIT_WITH_TESTS:BOOL, 79
89
HPX_ATTACH_DEBUGGER_ON_TEST_FAILURE:BOOL,
HPX_STATISTICS_WITH_COMPATIBILITY_HEADERS:BOOL,
89
HPX_WITH_AUTOMATIC_SERIALIZATION_REGISTRATION:BOOL,
89
HPX_STATISTICS_WITH_DEPRECATED_WARNINGS:BOOL,
89
HPX_WITH_BENCHMARK_SCRIPTS_PATH:PATH,
89
HPX_STATISTICS_WITH_TESTS:BOOL, 73
89
HPX_SYNCHRONIZATION_WITH_COMPATIBILITY_HEADERS:BOOL,
89
BUILD_BINARY_PACKAGE:BOOL,
73
HPX_SYNCHRONIZATION_WITH_DEPRECATED_WARNINGS:BOOL,
89
HPX_COMPILE_ONLY_TESTS:BOOL, 75
89
HPX_COMPILER_WARNINGS:BOOL, 73
HPX_SYNCHRONIZATION_WITH_TESTS:BOOL, 73
89
HPX_COMPILER_WARNINGS_AS_ERRORS:BOOL,
73
HPX_TESTING_WITH_COMPATIBILITY_HEADERS:BOOL,
89
HPX_WITH_COMPRESSION_BZIP2:BOOL, 73
89
HPX_WITH_COMPRESSION_SNAPPY:BOOL, 73
HPX_TESTING_WITH_DEPRECATED_WARNINGS:BOOL,
89
HPX_WITH_COMPRESSION_ZLIB:BOOL, 73
89
HPX_WITH_COROUTINE_COUNTERS:BOOL, 76
HPX_TESTING_WITH_TESTS:BOOL, 73
89
HPX_WITH_CUDA:BOOL, 73
HPX_THREAD_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL,
89
HPX_WITH_DATAPAR_VC:BOOL, 73
HPX_THREAD_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL,
89
HPX_WITH_DEFAULT_TARGETS:BOOL, 75
89
HPX_WITH_DEPRECATED_WARNINGS:BOOL,
HPX_THREAD_SUPPORT_WITH_TESTS:BOOL, 73
89
HPX_WITH_DISABLED_SIGNAL_EXCEPTION_HANDLERS:BOOL,
HPX_THREADMANAGER_WITH_COMPATIBILITY_HEADERS:BOOL,
89
HPX_WITH_DOCUMENTATION:BOOL, 75
HPX_THREADMANAGER_WITH_DEPRECATED_WARNINGS:BOOL,
89
HPX_DOCUMENTATION_OUTPUT_FORMATS:STRING,
75
HPX_THREADMANAGER_WITH_TESTS:BOOL, 73
HPX_DYNAMIC_HPX_MAIN:BOOL, 73

HPX_WITH_EXAMPLES:BOOL, 75
 HPX_WITH_EXAMPLES_HDF5:BOOL, 75
 HPX_WITH_EXAMPLES_OPENMP:BOOL, 75
 HPX_WITH_EXAMPLES_QT4:BOOL, 75
 HPX_WITH_EXAMPLES_QTHREADS:BOOL, 75
 HPX_WITH_EXAMPLES_TBB:BOOL, 75
 HPX_WITH_EXECUTABLE_PREFIX:STRING,
 75
 HPX_WITH_FAIL_COMPILE_TESTS:BOOL, 75
 HPX_WITHFAULT_TOLERANCE:BOOL, 73
 HPX_WITH_FORTRAN:BOOL, 73
 HPX_WITH_FULL_RPATH:BOOL, 73
 HPX_WITH_GCC_VERSION_CHECK:BOOL, 73
 HPX_WITH_GENERIC_CONTEXT_COROUTINES:BOOLHPX_WITH_TESTS_DEBUG_LOG_DESTINATION:STRING,
 73
 HPX_WITH_GOOGLE_PERFTOOLS:BOOL, 79
 HPX_WITH_HIDDEN_VISIBILITY:BOOL, 73
 HPX_WITH_IO_COUNTERS:BOOL, 75
 HPX_WITH_IO_POOL:BOOL, 77
 HPX_WITH_ITTNNOTIFY:BOOL, 79
 HPX_WITH_LOGGING:BOOL, 73
 HPX_WITH_MALLOC:STRING, 74
 HPX_WITH_MAX_CPU_COUNT:STRING, 77
 HPX_WITH_MAX_NUMA_DOMAIN_COUNT:STRING,
 77
 HPX_WITH_MORE_THAN_64_THREADS:BOOL,
 77
 HPX_WITH_NATIVE_TLS:BOOL, 74
 HPX_WITH_NETWORKING:BOOL, 78
 HPX_WITH_NICE_THREADLEVEL:BOOL, 74
 HPX_WITH_PAPI:BOOL, 79
 HPX_WITH_PARCEL_COALESCING:BOOL, 74
 HPX_WITH_PARCEL_PROFILING:BOOL, 78
 HPX_WITH_PARCELPORT_ACTION_COUNTERS:BOOLHPX_WITH_THREAD_IDLE_RATES:BOOL, 77
 78
 HPX_WITH_PARCELPORT_LIBFABRIC:BOOL,
 78
 HPX_WITH_PARCELPORT_MPI:BOOL, 78
 HPX_WITH_PARCELPORT_MPI_ENV:STRING,
 78
 HPX_WITH_PARCELPORT_MPI_MULTITHREADED:BOOLHPX_WITH_THREAD_SCHEDULERS:STRING,
 78
 HPX_WITH_PARCELPORT_TCP:BOOL, 78
 HPX_WITH_PARCELPORT_VERBS:BOOL, 78
 HPX_WITH_PSEUDO_DEPENDENCIES:BOOL,
 75
 HPX_WITH_REGISTER_THREAD_COMPATIBILITY:BOOL, 74
 HPX_WITH_RUN_MAIN_EVERYWHERE:BOOL,
 74
 HPX_WITH_SANITIZERS:BOOL, 79
 HPX_WITH_SCHEDULER_LOCAL_STORAGE:BOOL,
 77
 HPX_WITH_SPINLOCK_DEADLOCK_DETECTION:BOOL,
 77
 HPX_WITH_SPINLOCK_POOL_NUM:STRING,
 77
 HPX_WITH_STACKOVERFLOW_DETECTION:BOOL,
 74
 HPX_WITH_STACKTRACES:BOOL, 77
 HPX_WITH_STATIC_LINKING:BOOL, 74
 HPX_WITH_SWAP_CONTEXT_EMULATION:BOOL,
 77
 HPX_WITH_TESTS:BOOL, 75
 HPX_WITH_TESTS_BENCHMARKS:BOOL, 75
 HPX_WITH_TESTS_DEBUG_LOG:BOOL, 79
 HPX_WITH_TESTS_DEBUG_LOG_DESTINATION:STRING,
 79
 HPX_WITH_TESTS_EXAMPLES:BOOL, 75
 HPX_WITH_TESTS_EXTERNAL_BUILD:BOOL,
 75
 HPX_WITH_TESTS_HEADERS:BOOL, 76
 HPX_WITH_TESTS_REGRESSIONS:BOOL, 76
 HPX_WITH_TESTS_UNIT:BOOL, 76
 HPX_WITH_THREAD_BACKTRACE_DEPTH:STRING,
 77
 HPX_WITH_THREAD_BACKTRACE_ON_SUSPENSION:BOOL,
 77
 HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATES:BOOL
 77
 HPX_WITH_THREAD_CUMULATIVE_COUNTS:BOOL,
 77
 HPX_WITH_THREAD_DEBUG_INFO:BOOL, 79
 HPX_WITH_THREAD_DESCRIPTION_FULL:BOOL,
 79
 HPX_WITH_THREAD_GUARD_PAGE:BOOL, 79
 HPX_WITH_THREAD_IDLE_RATES:BOOL, 77
 HPX_WITH_THREAD_LOCAL_STORAGE:BOOL,
 77
 HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF:BOOL,
 77
 HPX_WITH_THREAD_QUEUE_WAITTIME:BOOL,
 77
 HPX_WITH_THREAD_TARGET_ADDRESS:BOOL,
 77
 HPX_WITH_THREAD_STACK_MMAP:BOOL, 77
 HPX_WITH_THREAD_STEALING_COUNTS:BOOL,
 77
 HPX_WITH_TIMER_POOL:BOOL, 78
 HPX_WITH_TOOLS:BOOL, 76
 HPX_WITH_VALGRIND:BOOL, 80
 HPX_WITH_VERIFY_LOCKS:BOOL, 80
 HPX_WITH_VERIFY_LOCKS_BACKTRACE:BOOL,
 80
 HPX_WITH_VERIFY_LOCKS_GLOBALLY:BOOL,

80
HPX_WITH_VIM_YCM:BOOL, 74
HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD:STRING, 74
HWLOC_ROOT:PATH, 90
PAPI_ROOT:PATH, 90
SPHINX_ROOT:PATH, 46
Component, 18

D

DECL_TYPE_SPECIFIER (*C macro*), 573
DOXYGEN_ROOT:PATH
 command line option, 45

G

get_memory_page_size (*C++ function*), 445

H

HDF5_ROOT:PATH
 command line option, 90

hpx (*C++ type*), 271, 374, 375, 380, 385, 390, 394–396, 401, 416, 417, 419–423, 425–434, 436, 438, 439, 445–450, 453, 460, 464, 467–470, 472, 473, 478–482, 485, 490, 492–499, 505, 508–510, 512–514, 516, 518–521, 523, 525–533, 535, 536, 538–540, 542, 544–547, 549, 555, 556, 560, 562–564, 567, 568, 570, 573–575, 577, 578, 580, 582, 583, 585, 587, 589, 591, 593–597, 599–603, 605, 607–609, 611–613, 615–620, 623, 625, 627, 628, 630, 632, 633, 636, 639, 642, 647, 649–652, 654, 656, 657, 661, 662, 664, 666, 669, 677, 681, 682, 684, 685, 687, 688, 690, 692, 695, 696, 699, 700, 702, 704, 705, 707, 708, 710, 713, 715, 716, 719, 722, 724, 725, 727–729, 731, 733, 736, 737, 741, 743, 745–747, 750, 752, 754, 756, 758, 761, 763, 766, 768, 769, 771, 773–776, 780, 783, 788, 807–809, 811, 812, 814–817, 820, 825, 826, 828, 830, 834, 838, 840, 843–846, 848, 849, 851, 853, 854, 856, 857, 859, 861, 862, 864, 865, 867, 869, 870, 873, 875, 877, 879, 881–884, 887, 890, 891, 893–909, 911, 915–920

hpx::actions (*C++ type*), 333
hpx::applier (*C++ type*), 333
hpx::applier::get_applier (*C++ function*), 334
hpx::applier::get_applier_ptr (*C++ function*), 334
hpx::assertion (*C++ type*), 773, 774
hpx::assertion::assertion_handler (*C++ type*), 774
hpx::assertion::force_linking (*C++ function*), 774

hpx::assertion::force_linking_helper
 (*C++ class*), 774
hpx::assertion::force_linking_helper::handle_assertion
 (*C++ member*), 774
hpx::assertion::operator<< (*C++ function*), 775
hpx::assertion::set_assertion_handler
 (*C++ function*), 774
hpx::assertion::source_location (*C++ class*), 775
hpx::assertion::source_location::file_name
 (*C++ member*), 775
hpx::assertion::source_location::function_name
 (*C++ member*), 775
hpx::assertion::source_location::line_number
 (*C++ member*), 775
hpx::assertion_failure (*C++ enumerator*), 829
hpx::bad_action_code (*C++ enumerator*), 828
hpx::bad_component_type (*C++ enumerator*), 828

hpx::bad_function_call (*C++ enumerator*), 830
hpx::bad_parameter (*C++ enumerator*), 828
hpx::bad_plugin_type (*C++ enumerator*), 830
hpx::bad_request (*C++ enumerator*), 828
hpx::bad_response_type (*C++ enumerator*), 829
hpx::basic_execution (*C++ type*), 478–481
hpx::basic_execution::agent_base (*C++ class*), 481
hpx::basic_execution::agent_base::~agent_base
 (*C++ function*), 481
hpx::basic_execution::agent_base::abort
 (*C++ function*), 482
hpx::basic_execution::agent_base::context
 (*C++ function*), 481
hpx::basic_execution::agent_base::description
 (*C++ function*), 481
hpx::basic_execution::agent_base::resume
 (*C++ function*), 482
hpx::basic_execution::agent_base::sleep_for
 (*C++ function*), 482
hpx::basic_execution::agent_base::sleep_until
 (*C++ function*), 482
hpx::basic_execution::agent_base::suspend
 (*C++ function*), 481
hpx::basic_execution::agent_base::yield
 (*C++ function*), 481
hpx::basic_execution::agent_base::yield_k
 (*C++ function*), 481
hpx::basic_execution::agent_ref (*C++ class*), 479
hpx::basic_execution::agent_ref::abort
 (*C++ function*), 479
hpx::basic_execution::agent_ref::agent_ref
 (*C++ function*), 479

```

hpx::basic_execution::agent_ref::impl_ hpx::basic_execution::this_thread::yield_k
    (C++ member), 480           (C++ function), 478
hpx::basic_execution::agent_ref::operator< hpx::broken_promise (C++ enumerator), 829
    bool (C++ function), 479      hpx::broken_task (C++ enumerator), 829
hpx::basic_execution::agent_ref::operator= hpx::commandline_option_error (C++ enu-
    (C++ function), 479        merator), 829
hpx::basic_execution::agent_ref::ref     hpx::components (C++ type), 334
    (C++ function), 479        hpx::components::binpacked (C++ member),
hpx::basic_execution::agent_ref::reset   337
    (C++ function), 479        hpx::components::binpacking_distribution_policy
hpx::basic_execution::agent_ref::resume  (C++ class), 260
    (C++ function), 479        hpx::components::binpacking_distribution_policy::b
hpx::basic_execution::agent_ref::sleep_for (C++ function), 261
    (C++ function), 479        hpx::components::binpacking_distribution_policy::b
hpx::basic_execution::agent_ref::sleep_until (C++ function), 261
    (C++ function), 479        hpx::components::binpacking_distribution_policy::c
hpx::basic_execution::agent_ref::suspend  (C++ function), 261
    (C++ function), 479        hpx::components::binpacking_distribution_policy::g
hpx::basic_execution::agent_ref::yield   (C++ function), 262
    (C++ function), 479        hpx::components::binpacking_distribution_policy::g
hpx::basic_execution::agent_ref::yield_k  (C++ function), 262
    (C++ function), 479        hpx::components::binpacking_distribution_policy::o
hpx::basic_execution::context_base (C++
    class), 481           hpx::components::colocated (C++ member),
hpx::basic_execution::context_base::~context_base 337
    (C++ function), 481           hpx::components::colocating_distribution_policy
hpx::basic_execution::context_base::resource  (C++ class), 262
    (C++ function), 481           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::resource_base  (C++ class), 480
    (C++ function), 480           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::resource_base::~resource_base 6Casefunction, 263
    (C++ function), 480           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::this_thread (C++
    type), 478           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::this_thread::agent  (C++ function), 263
    (C++ function), 478           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::this_thread::reset_agent (C++ class), 260
    (C++ class), 478           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::this_thread::reset_agent::(C++ type), 260
    (C++ function), 478           hpx::components::colocating_distribution_policy::a
hpx::basic_execution::this_thread::reset_agent::(C++ function), 262
    (C++ member), 478           hpx::components::colocating_distribution_policy::c
hpx::basic_execution::this_thread::reset_agent::(C++ function), 262
    (C++ function), 478           hpx::components::colocating_distribution_policy::c
hpx::basic_execution::this_thread::reset_agent::(C++ function), 262
    (C++ member), 478           hpx::components::colocating_distribution_policy::g
hpx::basic_execution::this_thread::sleep_for  (C++ function), 263
    (C++ function), 478           hpx::components::colocating_distribution_policy::g
hpx::basic_execution::this_thread::sleep_until (C++ function), 263
    (C++ function), 478           hpx::components::colocating_distribution_policy::o
hpx::basic_execution::this_thread::suspend  (C++ function), 262
    (C++ function), 478           hpx::components::copy (C++ function), 335
hpx::basic_execution::this_thread::yieldhpx::components::default_binpacking_counter_name
    (C++ function), 478           (C++ member), 337

```

```
hpx::components::default_distribution::pohpx:compute::host::block_allocator::construct  
    (C++ class), 263  
                                         (C++ function), 880  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::deallocate  
    (C++ function), 265  
                                         (C++ function), 880  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::destroy  
    (C++ function), 265  
                                         (C++ function), 880  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::difference_type  
    (C++ function), 264  
                                         (C++ type), 879  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::executor_  
    (C++ function), 264  
                                         (C++ member), 881  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::executor_type  
    (C++ class), 260  
                                         (C++ type), 879  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::is_always_equal  
    (C++ type), 260  
                                         (C++ type), 879  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::max_size  
    (C++ function), 264  
                                         (C++ function), 880  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::operator=  
    (C++ function), 264  
                                         (C++ function), 880  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::pointer  
    (C++ function), 264  
                                         (C++ type), 879  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::propagate_on_c  
    (C++ function), 265  
                                         (C++ type), 879  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::rebind  
    (C++ function), 265  
                                         (C++ class), 881  
hpx::components::default_distribution::pohpx:compute::host::block_allocator::reference  
    (C++ function), 264  
                                         (C++ type), 879  
hpx::components::default_layout     (C++ member), 337  
hpx::compute::migrate (C++ function), 336, 337  
hpx::components::migrate_from_storage (C++ function), 334  
hpx::components::migrate_to_storage (C++ function), 334, 335  
hpx::compute (C++ type), 870, 873, 875, 877, 879, 881, 882  
hpx::compute::host (C++ type), 873, 875, 877, 879  
hpx::compute::host::block_allocator  
    (C++ class), 879  
hpx::compute::host::block_allocator::address (C++ function), 880  
hpx::compute::host::block_allocator::allocate (C++ function), 880  
hpx::compute::host::block_allocator::block_alll (C++ function), 880  
hpx::compute::host::block_allocator::bulk_constl (C++ function), 880  
hpx::compute::host::block_allocator::bulk_destl (C++ member), 880  
hpx::compute::host::block_allocator::const_pointl (C++ type), 879  
hpx::compute::host::block_allocator::const_refl (C++ member), 879  
hpx::compute::host::block_executor  
    (C++ class), 873  
hpx::compute::host::block_executor::async_execute  
    (C++ function), 874  
hpx::compute::host::block_executor::block_executor  
    (C++ class), 874  
hpx::compute::host::block_executor::bulk_async_execl (C++ function), 874  
hpx::compute::host::block_executor::bulk_sync_execl (C++ function), 874  
hpx::compute::host::block_executor::current_<T>l (C++ function), 874  
hpx::compute::host::block_executor::executor_paraml (C++ member), 875  
hpx::compute::host::block_executor::executors_<T>l (C++ member), 875  
hpx::compute::host::block_executor::init_executors_<T>l (C++ member), 875
```

hpx::compute::host::block_executor::operator= (*C++ function*, 878)
hpx::compute::host::block_executor::post (*C++ function*, 878)
hpx::compute::host::block_executor::sync_execut(*C++ function*, 878)
hpx::compute::host::block_executor::targets (*C++ member*, 877)
hpx::compute::host::get_local_targets (*C++ function*, 875)
hpx::compute::host::get_targets (*C++ function*, 875)
hpx::compute::host::numa_domains (*C++ function*, 879)
hpx::compute::host::target (*C++ class*, 875)
hpx::compute::host::target::get_future (*C++ function*, 876)
hpx::compute::host::target::get_local_ta*gments* (*C++ function*, 876)
hpx::compute::host::target::get_locality (*C++ function*, 876)
hpx::compute::host::target::get_targets (*C++ function*, 876)
hpx::compute::host::target::handle_ (*C++ member*, 876)
hpx::compute::host::target::locality_ (*C++ member*, 876)
hpx::compute::host::target::native_handl*e* (*C++ function*, 876)
hpx::compute::host::target::native_handl*e* (*C++ class*, 876)
hpx::compute::host::target::native_handle_type (*C++ function*, 877)
hpx::compute::host::target::native_handle_type (*C++ member*, 877)
hpx::compute::host::target::native_handle_type (*C++ function*, 877)
hpx::compute::host::target::native_handle_type (*C++ member*, 877)
hpx::compute::host::target::native_handle_type (*C++ function*, 877)
hpx::compute::host::target::num_pus (*C++ function*, 876)
hpx::compute::host::target::serialize (*C++ function*, 876)
hpx::compute::host::target::synchronize (*C++ function*, 876)
hpx::compute::host::target::target (*C++ function*, 876)
hpx::compute::host::target_distribution_policy (*member*, 873)
hpx::compute::host::target_distribution_policy (*C++ class*, 877)
hpx::compute::host::target_distribution_policy (*C++ type*, 877)
hpx::compute::host::target_distribution_policy::bulk_create (*C++ function*, 878)
hpx::compute::host::target_distribution_policy::clear (*C++ function*, 878)
hpx::compute::host::target_distribution_policy::cre*ate* (*C++ function*, 878)
hpx::compute::host::target_distribution_policy::operator= (*C++ function*, 878)
hpx::compute::host::target_layout (*C++ member*, 878)
hpx::compute::swap (*C++ function*, 870)
hpx::compute::traits (*C++ type*, 881, 882)
hpx::compute::traits::access_target::read (*C++ function*, 881, 882)
hpx::compute::traits::access_target::target_type (*C++ type*, 881, 882)
hpx::compute::traits::access_target::write (*C++ function*, 881, 882)
hpx::compute::traits::access_target<host::target> (*C++ class*, 881, 882)
hpx::compute::traits::access_target<std::vector<host::target>> (*C++ class*, 881, 882)
hpx::compute::traits::allocator_traits (*C++ class*, 882)
hpx::compute::traits::allocator_traits::access_target (*C++ type*, 883)
hpx::compute::traits::allocator_traits::base_type (*C++ type*, 883)
hpx::compute::traits::allocator_traits::const_reference (*C++ type*, 883)
hpx::compute::traits::allocator_traits::reference (*C++ type*, 883)
hpx::compute::traits::allocator_traits::target_type (*C++ type*, 883)
hpx::compute::vector (*C++ class*, 870)
hpx::compute::vector::~vector (*C++ func*t**, 870)
hpx::compute::vector::device (*ta*int**, 871)
hpx::compute::vector::access_target (*ta*int**, 871)
hpx::compute::vector::alloc_ (*C++ mem*ber**, 871)
hpx::compute::vector::alloc_type (*C++ type*, 873)
hpx::compute::vector::begin (*C++ function*, 872)
hpx::compute::vector::capacity (*C++ function*, 872)
hpx::compute::vector::capacity_ (*C++ type*, 872)
hpx::compute::vector::cbegin (*C++ func*t**, 872)
hpx::compute::vector::cend (*C++ function*, 872)
hpx::compute::vector::clear (*C++ function*, 872)

872
hpx::compute::vector::const_iterator
(C++ type), 871
hpx::compute::vector::const_pointer
(C++ type), 871
hpx::compute::vector::const_reference
(C++ type), 871
hpx::compute::vector::const_reverse_iterator
(C++ type), 871
hpx::compute::vector::data (C++ function),
871, 872
hpx::compute::vector::data_ (C++ member),
873
hpx::compute::vector::device_data (C++
function), 872
hpx::compute::vector::difference_type
(C++ type), 871
hpx::compute::vector::empty (C++ function),
872
hpx::compute::vector::end (C++ function),
872
hpx::compute::vector::get_allocator
(C++ function), 871
hpx::compute::vector::iterator (C++ type),
871
hpx::compute::vector::operator= (C++
function), 871
hpx::compute::vector::pointer (C++ type),
871
hpx::compute::vector::reference (C++
type), 871
hpx::compute::vector::resize (C++ func-
tion), 872
hpx::compute::vector::reverse_iterator
(C++ type), 871
hpx::compute::vector::size (C++ function),
872
hpx::compute::vector::size_ (C++ member),
873
hpx::compute::vector::size_type (C++
type), 871
hpx::compute::vector::swap (C++ function),
872
hpx::compute::vector::value_type (C++
type), 871
hpx::compute::vector::vector (C++ func-
tion), 871
hpx::config (C++ type), 490
hpx::config::force_linking (C++ function),
491
hpx::config::force_linking_helper (C++
class), 491
hpx::config::force_linking_helper::boost
(C++ member), 491
hpx::config::force_linking_helper::hpx
hpx::config::force_linking_helper::hpx_version
(C++ member), 491
hpx::create_binary_filter (C++ function),
297
hpx::custom_exception_info_handler_type
(C++ type), 821
hpx::deadlock (C++ enumerator), 829
hpx::disconnect (C++ function), 292, 293
hpx::duplicate_component_address (C++
enumerator), 828
hpx::duplicate_component_id (C++ enumera-
tor), 830
hpx::duplicate_console (C++ enumerator), 828
hpx::dynamic_link_failure (C++ enumerator),
829
hpx::error (C++ enum), 828
hpx::error_code (C++ class), 818
hpx::error_code::clear (C++ function), 820
hpx::error_code::error_code (C++ function),
818–820
hpx::error_code::exception_ (C++ member),
820
hpx::error_code::get_message (C++ func-
tion), 820
hpx::error_code::operator= (C++ function),
820
hpx::error_info (C++ class), 831
hpx::error_info::_value (C++ member), 831
hpx::error_info::error_info (C++ function),
831
hpx::error_info<Tag, Type>::tag (C++
type), 831
hpx::error_info<Tag, Type>::type (C++
type), 831
hpx::evaluate_active_counters (C++ func-
tion), 296
hpx::exception (C++ class), 823
hpx::exception::~exception (C++ function),
824
hpx::exception::exception (C++ function),
823, 824
hpx::exception::get_error (C++ function),
824
hpx::exception::get_error_code (C++ func-
tion), 824
hpx::exception_info (C++ class), 831
hpx::exception_info::_data (C++ member),
832
hpx::exception_info::~exception_info
(C++ function), 831
hpx::exception_info::exception_info
(C++ function), 831
hpx::exception_info::get (C++ function), 831
hpx::exception_info::node_ptr (C++ type),

832
`hpx::exception_info::operator= (C++ function)`, 831
`hpx::exception_info::set (C++ function)`, 831
`hpx::exception_list (C++ class)`, 826
`hpx::exception_list::begin (C++ function)`, 826
`hpx::exception_list::end (C++ function)`, 826
`hpx::exception_list::exception_list_type (C++ type)`, 827
`hpx::exception_list::exceptions_ (C++ member)`, 827
`hpx::exception_list::iterator (C++ type)`, 826
`hpx::exception_list::mtx_ (C++ member)`, 827
`hpx::exception_list::mutex_type (C++ type)`, 827
`hpx::exception_list::size (C++ function)`, 826
`hpx::filesystem (C++ type)`, 775
`hpx::filesystem::basename (C++ function)`, 775
`hpx::filesystem::canonical (C++ function)`, 775
`hpx::filesystem::initial_path (C++ function)`, 775
`hpx::filesystem_error (C++ enumerator)`, 830
`hpx::finalize (C++ function)`, 291, 292
`hpx::find_all_from_basename (C++ function)`, 297
`hpx::find_all_localities (C++ function)`, 301
`hpx::find_from_basename (C++ function)`, 297, 298
`hpx::find_here (C++ function)`, 300
`hpx::find_locality (C++ function)`, 303
`hpx::find_remote_localities (C++ function)`, 302
`hpx::find_root_locality (C++ function)`, 301
`hpx::future_already_retrieved (C++ enumerator)`, 829
`hpx::future_can_not_be_cancelled (C++ enumerator)`, 829
`hpx::future_cancelled (C++ enumerator)`, 829
`hpx::future_does_not_support_cancellation (C++ enumerator)`, 829
`hpx::get_colocation_id (C++ function)`, 303, 304
`hpx::get_error (C++ function)`, 821, 822
`hpx::get_error_file_name (C++ function)`, 822
`hpx::get_error_function_name (C++ function)`, 822
`hpx::get_error_line_number (C++ function)`, 822
`hpx::get_error.what (C++ function)`, 821
`hpx::get_exception_info (C++ function)`, 830
`hpx::get_hpx_category (C++ function)`, 818
`hpx::get_hpx_rethrow_category (C++ function)`, 818
`hpx::get_initial_num_localities (C++ function)`, 306
`hpx::get_locality (C++ function)`, 294
`hpx::get_locality_id (C++ function)`, 306
`hpx::get_locality_name (C++ function)`, 306
`hpx::get_num_localities (C++ function)`, 307
`hpx::get_num_worker_threads (C++ function)`, 295
`hpx::get_os_thread_count (C++ function)`, 308
`hpx::get_ptr (C++ function)`, 304, 305
`hpx::get_runtime_instance_number (C++ function)`, 294
`hpx::get_runtime_mode_from_name (C++ function)`, 309
`hpx::get_runtime_mode_name (C++ function)`, 309
`hpx::get_system_uptime (C++ function)`, 295
`hpx::get_thread_name (C++ function)`, 308
`hpx::get_worker_thread_num (C++ function)`, 308
`hpx::init (C++ function)`, 272–280
`hpx::internal_server_error (C++ enumerator)`, 828
`hpx::invalid_data (C++ enumerator)`, 829
`hpx::invalid_status (C++ enumerator)`, 828
`hpx::invoke_with_exception_info (C++ function)`, 830, 831
`hpx::is_running (C++ function)`, 295
`hpx::is_starting (C++ function)`, 294
`hpx::is_stopped (C++ function)`, 295
`hpx::is_stopped_or_shutting_down (C++ function)`, 295
`hpx::kernel_error (C++ enumerator)`, 829
`hpx::launch (C++ class)`, 265
`hpx::launch::apply (C++ member)`, 265
`hpx::launch::deferred (C++ member)`, 265
`hpx::launch::fork (C++ member)`, 265
`hpx::launch::launch (C++ function)`, 265
`hpx::launch::select (C++ member)`, 266
`hpx::launch::sync (C++ member)`, 265
`hpx::lcos (C++ type)`, 549, 555, 556, 560, 563, 564, 567, 568, 575, 577, 578, 580, 582, 583, 585, 587, 589, 594–597, 599–603, 605, 607–609, 611
`hpx::lcos::all_reduce (C++ function)`, 582, 583
`hpx::lcos::all_to_all (C++ function)`, 589, 590
`hpx::lcos::barrier (C++ class)`, 578
`hpx::lcos::barrier::barrier (C++ function)`, 578, 579

hpx::lcos::barrier::synchronize (*C++ function*), 579
hpx::lcos::barrier::wait (*C++ function*), 579
hpx::lcos::define_spmd_block (*C++ function*), 587
hpx::lcos::gather_here (*C++ function*), 580, 581
hpx::lcos::gather_there (*C++ function*), 580, 581
hpx::lcos::latch (*C++ class*), 585
hpx::lcos::latch::base_type (*C++ type*), 586
hpx::lcos::latch::count_down (*C++ function*), 586
hpx::lcos::latch::count_down_and_wait (*C++ function*), 586
hpx::lcos::latch::is_ready (*C++ function*), 586
hpx::lcos::latch::latch (*C++ function*), 585, 586
hpx::lcos::latch::wait (*C++ function*), 586
hpx::lcos::local (*C++ type*), 549, 555, 556, 560, 563, 564, 567, 568, 594–597, 599–603, 605, 607–609, 611
hpx::lcos::local::and_gate (*C++ class*), 564
hpx::lcos::local::and_gate::and_gate (*C++ function*), 565
hpx::lcos::local::and_gate::base_type (*C++ type*), 565
hpx::lcos::local::and_gate::get_future (*C++ function*), 565
hpx::lcos::local::and_gate::get_shared_future (*C++ function*), 565
hpx::lcos::local::and_gate::operator= (*C++ function*), 565
hpx::lcos::local::and_gate::set (*C++ function*), 565
hpx::lcos::local::and_gate::synchronize (*C++ function*), 565
hpx::lcos::local::barrier (*C++ class*), 600
hpx::lcos::local::barrier::~barrier (*C++ function*), 600
hpx::lcos::local::barrier::barrier (*C++ function*), 600
hpx::lcos::local::barrier::cond_ (*C++ member*), 601
hpx::lcos::local::barrier::count_up (*C++ function*), 600
hpx::lcos::local::barrier::mtx_ (*C++ member*), 601
hpx::lcos::local::barrier::mutex_type (*C++ type*), 601
hpx::lcos::local::barrier::number_of_threads (*C++ member*), 601
hpx::lcos::local::barrier::reset (*C++ function*), 600
hpx::lcos::local::barrier::total_ (*C++ member*), 601
hpx::lcos::local::barrier::wait (*C++ function*), 600
hpx::lcos::local::base_and_gate (*C++ class*), 565
hpx::lcos::local::base_and_gate::base_and_gate (*C++ function*), 565
hpx::lcos::local::base_and_gate::condition_list_type (*C++ type*), 566
hpx::lcos::local::base_and_gate::conditions_ (*C++ member*), 566
hpx::lcos::local::base_and_gate::generation (*C++ function*), 565
hpx::lcos::local::base_and_gate::generation_ (*C++ member*), 566
hpx::lcos::local::base_and_gate::get_future (*C++ function*), 565, 566
hpx::lcos::local::base_and_gate::get_shared_future (*C++ function*), 565, 566
hpx::lcos::local::base_and_gate::init_locked (*C++ function*), 566
hpx::lcos::local::base_and_gate::manage_condition (*C++ class*), 566
hpx::lcos::local::base_and_gate::mtx_ (*C++ member*), 566
hpx::lcos::local::base_and_gate::mutex_type (*C++ type*), 566
hpx::lcos::local::base_and_gate::next_generation (*C++ function*), 565
hpx::lcos::local::base_and_gate::operator= (*C++ function*), 565
hpx::lcos::local::base_and_gate::promise_ (*C++ member*), 566
hpx::lcos::local::base_and_gate::received_segments_ (*C++ member*), 566
hpx::lcos::local::base_and_gate::set (*C++ function*), 565, 566
hpx::lcos::local::base_and_gate::synchronize (*C++ function*), 565, 566
hpx::lcos::local::base_and_gate::test_condition (*C++ function*), 566
hpx::lcos::local::base_and_gate::trigger_conditions (*C++ function*), 566
hpx::lcos::local::base_and_gate<Mutex>::manage_condition (*C++ function*), 567
hpx::lcos::local::base_and_gate<Mutex>::manage_condition_ (*C++ function*), 567
hpx::lcos::local::base_and_gate<Mutex>::manage_condition_ (*C++ member*), 567
hpx::lcos::local::base_and_gate<Mutex>::manage_condition_ (*C++ function*), 567
hpx::lcos::local::base_and_gate<Mutex>::manage_condition_ (*C++ member*), 567
hpx::lcos::local::base_and_gate<Mutex>::manage_condition_ (*C++ function*), 567

(C++ member), 567
`hpx::lcos::local::base_channel_mpsc` (C++ class), 602
`hpx::lcos::local::base_channel_mpsc::~base_channel_mpsc` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::base_channel_mpsc` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::buf` (C++ type), 602
`hpx::lcos::local::base_channel_mpsc::buf::buf` (C++ member), 603
`hpx::lcos::local::base_channel_mpsc::cap` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::close` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::get` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::head` (C++ member), 603
`hpx::lcos::local::base_channel_mpsc::is_empty` (C++ function), 603
`hpx::lcos::local::base_channel_mpsc::is_full` (C++ function), 603
`hpx::lcos::local::base_channel_mpsc::open` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::set` (C++ function), 602
`hpx::lcos::local::base_channel_mpsc::size` (C++ member), 603
`hpx::lcos::local::base_channel_mpsc::tail` (C++ member), 603
`hpx::lcos::local::base_channel_mpsc::tail_data` (C++ class), 603
`hpx::lcos::local::base_channel_mpsc<T, Mutex>::mutex_type` (C++ type), 602
`hpx::lcos::local::base_channel_mpsc<T, Mutex>::tail_data::mtx_` (C++ member), 603
`hpx::lcos::local::base_channel_mpsc<T, Mutex>::tail_data::tail_` (C++ member), 603
`hpx::lcos::local::base_trigger` (C++ class), 560
`hpx::lcos::local::base_trigger::base_trigger` (C++ function), 561
`hpx::lcos::local::base_trigger::condition` (C++ type), 561
`hpx::lcos::local::base_trigger::condition::hpx::lcos::local::bounded_channel::head_` (C++ member), 561
`hpx::lcos::local::base_trigger::generator` (C++ function), 561
`hpx::lcos::local::base_trigger::generator::hpx::lcos::local::bounded_channel::is_empty` (C++ function), 561
`hpx::lcos::local::base_trigger::generator::hpx::lcos::local::bounded_channel::is_full` (C++ function), 561
`hpx::lcos::local::base_trigger::get_futu` (C++ type), 561
`hpx::lcos::local::base_trigger::get_futu::hpx::lcos::local::bounded_channel::mtx_` (C++ function), 561
`hpx::lcos::local::base_trigger::manage_condition` (C++ class), 561
`hpx::lcos::local::base_trigger::operator=` (C++ function), 561
`hpx::lcos::local::base_trigger::promise_` (C++ member), 561
`hpx::lcos::local::base_trigger::set` (C++ function), 561
`hpx::lcos::local::base_trigger::synchronize` (C++ function), 561
`hpx::lcos::local::base_trigger::test_condition` (C++ function), 561
`hpx::lcos::local::base_trigger::trigger_conditions` (C++ function), 561
`hpx::lcos::local::base_trigger<Mutex>::manage_cond` (C++ function), 562
`hpx::lcos::local::base_trigger<Mutex>::open` (C++ function), 562
`hpx::lcos::local::base_trigger<Mutex>::set` (C++ function), 562
`hpx::lcos::local::base_trigger<Mutex>::size` (C++ member), 562
`hpx::lcos::local::base_trigger<Mutex>::tail` (C++ member), 562
`hpx::lcos::local::bounded_channel` (C++ class), 611
`hpx::lcos::local::bounded_channel::~bounded_channel` (C++ function), 611
`hpx::lcos::local::bounded_channel::bounded_channel` (C++ function), 611
`hpx::lcos::local::bounded_channel::buffer_` (C++ member), 612
`hpx::lcos::local::bounded_channel::capacity` (C++ function), 611
`hpx::lcos::local::bounded_channel::close` (C++ function), 611, 612
`hpx::lcos::local::bounded_channel::closed_` (C++ member), 612
`hpx::lcos::local::bounded_channel::get` (C++ function), 611
`hpx::lcos::local::bounded_channel::head_` (C++ member), 612
`hpx::lcos::local::bounded_channel::is_empty` (C++ function), 612
`hpx::lcos::local::bounded_channel::is_full` (C++ function), 612

(*C++ member*), 612
hpx::lcos::local::bounded_channel::operator= (*C++ function*), 552
(*C++ function*), 611
hpx::lcos::local::bounded_channel::set (*C++ function*), 611
hpx::lcos::local::bounded_channel::size_ (*C++ member*), 612
hpx::lcos::local::bounded_channel::tail_ (*C++ member*), 612
hpx::lcos::local::bounded_channel<T, Mutex>::mutex_type (*C++ type*), 612
hpx::lcos::local::call_once (*C++ function*), 601
hpx::lcos::local::channel (*C++ class*), 549
hpx::lcos::local::channel::base_type (*C++ type*), 550
hpx::lcos::local::channel::channel (*C++ function*), 550
hpx::lcos::local::channel::value_type (*C++ type*), 550
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::closed_ (*C++ class*), 550
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::get (*C++ type*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::head_ (*C++ member*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::is_empty (*C++ function*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::is_full (*C++ member*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::operator= (*C++ function*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::set (*C++ function*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::size_ (*C++ function*), 551
hpx::lcos::local::channel_async_iterator hpx::lcos::local::channel_spsc::tail_ (*C++ function*), 551
hpx::lcos::local::channel_iterator (*C++ class*), 551
hpx::lcos::local::channel_iterator::base hpx::lcos::local::condition_variable (*C++ type*), 551, 552
hpx::lcos::local::channel_iterator::chan hpx::lcos::local::condition_variable::cond_ (*C++ member*), 552, 553
hpx::lcos::local::channel_iterator::chan hpx::lcos::local::condition_variable::mtx_ (*C++ function*), 551, 552
hpx::lcos::local::channel_iterator::data hpx::lcos::local::condition_variable::mutex_type (*C++ member*), 552, 553
hpx::lcos::local::channel_iterator::dereference hpx::lcos::local::condition_variable::notify_all (*C++ function*), 552
hpx::lcos::local::channel_iterator::equal hpx::lcos::local::condition_variable::notify_one (*C++ function*), 552
hpx::lcos::local::channel_iterator::get_hpx::lcos::local::condition_variable::wait (*C++ function*), 552
hpx::lcos::local::channel_iterator::increment hpx::lcos::local::channel_iterator::increment (*C++ function*), 552
hpx::lcos::local::channel_iterator::operator= hpx::lcos::local::channel_iterator<void> (*C++ class*), 552
hpx::lcos::local::channel_mpmc (*C++ type*), 611
hpx::lcos::local::channel_mpsc (*C++ type*), 602
hpx::lcos::local::channel_spsc (*C++ class*), 594
hpx::lcos::local::channel_spsc::~channel_spsc (*C++ function*), 594
hpx::lcos::local::channel_spsc::buffer_ (*C++ member*), 595
hpx::lcos::local::channel_spsc::capacity (*C++ function*), 594
hpx::lcos::local::channel_spsc::channel_spdc (*C++ function*), 594
hpx::lcos::local::channel_spsc::close (*C++ function*), 594
hpx::lcos::local::channel_spdc::closed_ (*C++ member*), 595
hpx::lcos::local::channel_spdc::get (*C++ function*), 594
hpx::lcos::local::channel_spdc::head_ (*C++ member*), 595
hpx::lcos::local::channel_spdc::is_empty (*C++ function*), 595
hpx::lcos::local::channel_spdc::is_full (*C++ function*), 595
hpx::lcos::local::channel_spdc::operator= (*C++ function*), 594
hpx::lcos::local::channel_spdc::set (*C++ function*), 594
hpx::lcos::local::channel_spdc::size_ (*C++ member*), 595
hpx::lcos::local::channel_spdc::tail_ (*C++ member*), 595
hpx::lcos::local::channel_spdc::increment hpx::lcos::local::channel_spdc::increment (*C++ function*), 595
hpx::lcos::local::channel_spdc::operator= hpx::lcos::local::channel_spdc::operator= (*C++ function*), 594
hpx::lcos::local::channel_spdc::wait hpx::lcos::local::channel_spdc::wait (*C++ function*), 595

```

hpx::lcos::local::condition_variable::waithpx::lcos::local::counting_semaphore_var::try_wait  

    (C++ function), 606  

    (C++ function), 596
hpx::lcos::local::condition_variable::waithpx::lcos::local::counting_semaphore_var::wait  

    (C++ function), 606  

    (C++ function), 596
hpx::lcos::local::condition_variable_anyhpx::lcos::local::cv_status (C++ enum)  

    (C++ class), 606  

    605
hpx::lcos::local::condition_variable_anyhpx::lcos::local::error (C++ enumerator)  

    (C++ member), 607  

    605
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event (C++ class), 607  

    (C++ member), 607  

    hpx::lcos::local::event::cond_ (C++ member)
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event::event (C++ function), 608  

    (C++ type), 607  

    hpx::lcos::local::event::event_ (C++ function), 607
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event::event (C++ function), 608  

    (C++ function), 606  

    hpx::lcos::local::event::event_ (C++ function), 608
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event::mtx_ (C++ member), 608  

    (C++ function), 606  

    hpx::lcos::local::event::mutex_type
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event::wait (C++ type), 608  

    (C++ function), 606, 607  

    hpx::lcos::local::event::occurred (C++ function), 607
hpx::lcos::local::condition_variable_anyhpx::lcos::local::event::wait (function), 607  

    (C++ function), 606  

    hpx::lcos::local::event::reset (C++ function), 607
hpx::lcos::local::conditional_trigger  

    (C++ class), 567  

    hpx::lcos::local::event::set (C++ function), 608
hpx::lcos::local::conditional_trigger::cond_hpx::lcos::local::event::set_locked  

    (C++ member), 568  

    hpx::lcos::local::event::set (C++ function), 608
hpx::lcos::local::conditional_trigger::condhpx::lcos::local::event::wait (C++ function), 607  

    (C++ function), 567  

    hpx::lcos::local::event::wait (C++ function), 607
hpx::lcos::local::conditional_trigger::get_futurehpx::lcos::local::event::wait_locked  

    (C++ function), 567  

    hpx::lcos::local::event::wait_locked
hpx::lcos::local::conditional_trigger::operator<hpx::lcos::local::guard (C++ class), 608  

    (C++ function), 567  

    hpx::lcos::local::guard (C++ class), 563
hpx::lcos::local::conditional_trigger::phoenixhpx::lcos::local::guard::~guard (C++ function)  

    (C++ member), 568  

    hpx::lcos::local::guard (C++ class), 563
hpx::lcos::local::conditional_trigger::repeathpx::lcos::local::guard::guard (C++ function), 567  

    (C++ function), 563
hpx::lcos::local::conditional_trigger::skiphpx::lcos::local::guard::task (C++ member)  

    (C++ function), 567  

    hpx::lcos::local::guard::task (C++ member), 564
hpx::lcos::local::counting_semaphore  

    (C++ type), 596  

    hpx::lcos::local::guard_set (C++ class), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::~guard_set  

    (C++ class), 596  

    (C++ function), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::semaphore_guard_set::add  

    (C++ function), 596  

    (C++ function), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::get  

    (C++ member), 597  

    (C++ function), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::guard_set  

    (C++ type), 597  

    (C++ function), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::guards  

    (C++ member), 597  

    (C++ member), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::size  

    (C++ function), 597  

    (C++ function), 564
hpx::lcos::local::counting_semaphore_varhpx::lcos::local::guard_set::sort  

    (C++ function), 597  

    (C++ function), 564

```

hpx::lcos::local::guard_set::sorted
 (C++ member), 564

hpx::lcos::local::latch (C++ class), 603

hpx::lcos::local::latch::~latch (C++
 function), 604

hpx::lcos::local::latch::abort_all (C++
 function), 604

hpx::lcos::local::latch::cond_ (C++ mem-
 ber), 605

hpx::lcos::local::latch::count_down
 (C++ function), 604

hpx::lcos::local::latch::count_down_and_hpx::lcos::local::once_flag::event_
 (C++ function), 604

hpx::lcos::local::latch::count_up (C++
 function), 604

hpx::lcos::local::latch::counter_ (C++
 member), 605

hpx::lcos::local::latch::HPX_NON_COPYABLE
 (C++ function), 604

hpx::lcos::local::latch::is_ready (C++
 function), 604

hpx::lcos::local::latch (C++ func-
 tion), 604

hpx::lcos::local::latch::mtx_ (C++ mem-
 ber), 605

hpx::lcos::local::latch::mutex_type
 (C++ type), 605

hpx::lcos::local::latch::notified_ (C++
 member), 605

hpx::lcos::local::latch::reset (C++ func-
 tion), 605

hpx::lcos::local::latch::wait (C++ func-
 tion), 604

hpx::lcos::local::mutex (C++ class), 609

hpx::lcos::local::mutex::~mutex (C++
 function), 610

hpx::lcos::local::mutex::cond_ (C++ mem-
 ber), 610

hpx::lcos::local::mutex::HPX_NON_COPYABLE
 (C++ function), 610

hpx::lcos::local::mutex::lock (C++ func-
 tion), 610

hpx::lcos::local::mutex::mtx_ (C++ mem-
 ber), 610

hpx::lcos::local::mutex::mutex (C++ func-
 tion), 610

hpx::lcos::local::mutex::mutex_type
 (C++ type), 610

hpx::lcos::local::mutex::owner_id_ (C++
 member), 610

hpx::lcos::local::mutex::try_lock (C++
 function), 610

hpx::lcos::local::mutex::unlock (C++
 function), 610

hpx::lcos::local::no_mutex (C++ class), 594

hpx::lcos::local::no_mutex::lock (C++
 function), 594

hpx::lcos::local::no_mutex::try_lock
 (C++ function), 594

hpx::lcos::local::no_mutex::unlock (C++
 function), 594

hpx::lcos::local::no_timeout (C++ enumera-
 tor), 605

hpx::lcos::local::once_flag (C++ class),
 601

hpx::lcos::local::once_flag::event_
 (C++ member), 601

hpx::lcos::local::once_flag::HPX_NON_COPYABLE
 (C++ function), 601

hpx::lcos::local::once_flag::once_flag
 (C++ function), 601

hpx::lcos::local::once_flag::status_
 (C++ member), 601

hpx::lcos::local::one_element_channel
 (C++ class), 553

hpx::lcos::local::one_element_channel::base_type
 (C++ type), 553, 554

hpx::lcos::local::one_element_channel::one_element_
 (C++ function), 553

hpx::lcos::local::one_element_channel::value_type
 (C++ type), 553

hpx::lcos::local::one_element_channel<void>
 (C++ class), 553

hpx::lcos::local::packaged_task::function_
 (C++ member), 556

hpx::lcos::local::packaged_task::function_type
 (C++ type), 556

hpx::lcos::local::packaged_task::get_future
 (C++ function), 556

hpx::lcos::local::packaged_task::invoke_impl
 (C++ function), 556

hpx::lcos::local::packaged_task::operator()
 (C++ function), 556

hpx::lcos::local::packaged_task::operator=
 (C++ function), 556

hpx::lcos::local::packaged_task::packaged_task
 (C++ function), 556

hpx::lcos::local::packaged_task::promise_
 (C++ member), 556

hpx::lcos::local::packaged_task::reset
 (C++ function), 556

hpx::lcos::local::packaged_task::set_exception
 (C++ function), 556

hpx::lcos::local::packaged_task::swap
 (C++ function), 556

hpx::lcos::local::packaged_task::valid
 (C++ function), 556

hpx::lcos::local::packaged_task<R(Ts...)>

```

(C++ class), 555
hpx::lcos::local::promise (C++ class), 568
hpx::lcos::local::promise::~promise
    (C++ function), 568, 569
hpx::lcos::local::promise::base_type
    (C++ type), 569, 570
hpx::lcos::local::promise::operator=
    (C++ function), 568, 569
hpx::lcos::local::promise::promise (C++
    function), 568, 569
hpx::lcos::local::promise::set_exceptionhpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::can_be_deleted_
    (C++ member), 558
hpx::lcos::local::promise::set_value
    (C++ function), 568, 569
hpx::lcos::local::promise::swap   (C++
    function), 568, 569
hpx::lcos::local::promise::valid  (C++
    function), 568, 569
hpx::lcos::local::promise<R&> (C++ class),
    569
hpx::lcos::local::promise<void>   (C++
    class), 569
hpx::lcos::local::receive_buffer  (C++ class), 556
hpx::lcos::local::receive_buffer::~receive_buffer(C++ function), 558
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::buffer_map_ Mutex>::entry_data::value_set_
    (C++ member), 557, 559
hpx::lcos::local::receive_buffer::bufferhpax::lcos::local::receive_buffer<T,
    (C++ type), 557, 559
hpx::lcos::local::receive_buffer::buffer_promis(C++ function), 558
    (C++ type), 557, 559
hpx::lcos::local::receive_buffer::cancel_waiting Mutex>::erase_on_exit::buffer_map_
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::empty hpx::lcos::local::receive_buffer<T,
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::entry_data   (C++ function), 558
    (C++ class), 557
hpx::lcos::local::receive_buffer::erase_on_exitMutex>::erase_on_exit::erase_on_exit
    (C++ class), 558
hpx::lcos::local::receive_buffer::get_bufhpax::lcos::local::receive_buffer<void,
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::iterat(hpx::lcos::local::receive_buffer<void,
    (C++ type), 557, 559
hpx::lcos::local::receive_buffer::mtx_ hpx::lcos::local::receive_buffer<void,
    (C++ member), 557, 559
hpx::lcos::local::receive_buffer::mutex_type   hpx::lcos::local::receive_buffer<void,
    (C++ type), 557, 559
hpx::lcos::local::receive_buffer::operator=     Mutex>::entry_data::cancel (C++
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::receivhpx::lcos::local::receive_buffer<void,
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::receive_buffe(C++ function), 559
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::store_received
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer::try_receive
    (C++ function), 557, 559
hpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::entry_data
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::get_future
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::HPX_NON_COPYABLE
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::promise_
    (C++ member), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::entry_data::set_value
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::erase_on_exit::~erase_on_exit
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::erase_on_exit::buffer_map_
    (C++ member), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::erase_on_exit::erase_on_exit
    (C++ function), 558
hpx::lcos::local::receive_buffer<T,
    Mutex>::erase_on_exit::it_
    (C++ member), 558
hpx::lcos::local::receive_buffer<void,
    Mutex> (C++ class), 558
hpx::lcos::local::receive_buffer<void,
    Mutex>::entry_data (C++ class), 559
hpx::lcos::local::receive_buffer<void,
    Mutex>::entry_data::can_be_deleted_
    (C++ member), 560
hpx::lcos::local::receive_buffer<void,
    Mutex>::entry_data::cancel (C++
    function), 560
hpx::lcos::local::receive_buffer<void,
    Mutex>::entry_data::entry_data
    (C++ function), 559
hpx::lcos::local::receive_buffer<void,
    Mutex>::entry_data::store_received
    (C++ function), 559

```

```
    Mutex>::entry_data::get_future  
    (C++ function), 559  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::entry_data::HPX_NON_COPYABLE  
    (C++ function), 559  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::entry_data::promise_  
    (C++ member), 560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::entry_data::set_value  
    (C++ function), 559  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::entry_data::value_set_  
    (C++ member), 560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::erase_on_exit (C++ class),  
    560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::erase_on_exit::~erase_on_exit  
    (C++ function), 560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::erase_on_exit::buffer_map_  
    (C++ member), 560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::erase_on_exit::erase_on_exit  
    (C++ function), 560  
hpx::lcos::local::receive_buffer<void,  
    Mutex>::erase_on_exit::it_ (C++  
    member), 560  
hpx::lcos::local::receive_channel (C++  
    class), 554  
hpx::lcos::local::receive_channel::base_type  
    (C++ type), 554  
hpx::lcos::local::receive_channel::receive_channel  
    (C++ function), 554  
hpx::lcos::local::receive_channel<void>  
    (C++ class), 554  
hpx::lcos::local::recursive_mutex (C++  
    type), 608  
hpx::lcos::local::run_guarded (C++ func-  
    tion), 563  
hpx::lcos::local::send_channel (C++  
    class), 555  
hpx::lcos::local::send_channel::base_type  
    (C++ type), 555  
hpx::lcos::local::send_channel::send_channel  
    (C++ function), 555  
hpx::lcos::local::send_channel<void>  
    (C++ class), 555  
hpx::lcos::local::shared_mutex (C++ type),  
    599  
hpx::lcos::local::sliding_semaphore  
    (C++ type), 597  
hpx::lcos::local::sliding_semaphore_var  
    (C++ class), 597  
    hpx::lcos::local::sliding_semaphore_var::mtx_  
        (C++ member), 599  
    hpx::lcos::local::sliding_semaphore_var::mutex_type  
        (C++ type), 598  
    hpx::lcos::local::sliding_semaphore_var::sem_  
        (C++ member), 599  
    hpx::lcos::local::sliding_semaphore_var::set_max_d  
        (C++ function), 598  
    hpx::lcos::local::sliding_semaphore_var::signal  
        (C++ function), 598  
    hpx::lcos::local::sliding_semaphore_var::signal_all  
        (C++ function), 598  
    hpx::lcos::local::sliding_semaphore_var::sliding_se  
        (C++ function), 598  
    hpx::lcos::local::sliding_semaphore_var::try_wait  
        (C++ function), 598  
    hpx::lcos::local::sliding_semaphore_var::wait  
        (C++ function), 598  
        hpx::lcos::local::spinlock (C++ class), 595  
        hpx::lcos::local::spinlock::~spinlock  
    hpx::lcos::local::spinlock::acquire_lock  
        (C++ function), 595  
    hpx::lcos::local::spinlock::lock (C++  
        function), 595  
    hpx::lcos::local::spinlock::relinquish_lock  
        (C++ function), 595  
    hpx::lcos::local::spinlock::spinlock  
        (C++ function), 595  
    hpx::lcos::local::spinlock::try_lock  
        (C++ function), 595  
    hpx::lcos::local::spinlock::unlock (C++  
        function), 595  
    hpx::lcos::local::spinlock::v_ (C++ mem-  
        ber), 596  
    hpx::lcos::local::spinlock_no_backoff  
        (C++ class), 599  
    hpx::lcos::local::spinlock_no_backoff::~spinlock_no  
        (C++ function), 599  
    hpx::lcos::local::spinlock_no_backoff::spinlock_no  
        (C++ function), 599  
    hpx::lcos::local::spinlock_no_backoff::try_lock  
        (C++ function), 599  
    hpx::lcos::local::spinlock_no_backoff::unlock  
        (C++ function), 599  
    hpx::lcos::local::spinlock_no_backoff::v_  
        (C++ member), 600
```

```

hpx::lcos::local::spinlock_pool      (C++ type), 588
    class), 608
hpx::lcos::local::spinlock_pool::pool_ (C++ member), 609
hpx::lcos::local::spinlock_pool::scoped_lock (C++ class), 609
hpx::lcos::local::spinlock_pool::spinlock_for (C++ function), 587
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::~scoped_lock (C++ function), 609
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::HPX_NON_COPYABLE (C++ function), 609
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::lock (C++ function), 609
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::scoped_lock (C++ function), 609
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::sp_ (C++ member), 609
hpx::lcos::local::spinlock_pool<Tag, N>::scoped_lock::unlock (C++ function), 609
hpx::lcos::local::swap (C++ function), 568
hpx::lcos::local::timed_mutex (C++ class), 610
hpx::lcos::local::timed_mutex::~timed_muhpx::length_error (C++ enumerator), 830
    (C++ function), 610
hpx::lcos::local::timed_mutex::HPX_NON_COPYABLE (C++ type), 562
    (C++ function), 610
hpx::lcos::local::timed_mutex::timed_mutex (C++ function), 610
hpx::lcos::local::timed_mutex::try_lock_for (C++ function), 610
hpx::lcos::local::timed_mutex::try_lock_until (C++ member), 563
hpx::lcos::local::timeout (C++ enumerator), 605
hpx::lcos::local::trigger (C++ class), 562
hpx::lcos::local::trigger::base_type (C++ type), 562
hpx::lcos::local::trigger::operator= (C++ function), 562
hpx::lcos::local::trigger::synchronize (C++ function), 562
hpx::lcos::local::trigger::trigger (C++ function), 562
hpx::lcos::spmd_block (C++ class), 587
hpx::lcos::spmd_block::barrier_ (C++ member), 588
hpx::lcos::spmd_block::barrier_type
hpx::lcos::spmd_block::barriers_ (C++ member), 588
hpx::lcos::spmd_block::get_images_per_locality
hpx::lcos::spmd_block::get_num_images
hpx::lcos::spmd_block::hash_ (C++ member), 588
hpx::lcos::spmd_block::image_id_ (C++ member), 588
hpx::lcos::spmd_block::images_per_locality_ (C++ member), 588
hpx::lcos::spmd_block::name_ (C++ member), 588
hpx::lcos::spmd_block::num_images_ (C++ member), 588
hpx::lcos::spmd_block::serialize (C++ function), 588
hpx::lcos::spmd_block::spmd_block (C++ function), 587
hpx::lcos::spmd_block::sync_all (C++ function), 587
hpx::lcos::spmd_block::sync_images (C++ function), 587, 588
hpx::lcos::spmd_block::table_type (C++ type), 588
hpx::lcos::spmd_block::this_image (C++ function), 587
hpx::length_error (C++ enumerator), 830
hpx::lightweight (C++ enumerator), 825
hpx::local_lcos::force_linking (C++ function), 563
hpx::local_lcos::force_linking_helper (C++ class), 563
hpx::local_lcos::force_linking_helper::free
hpx::lock_error (C++ enumerator), 828
hpx::make_error_code (C++ function), 817, 818
hpx::make_success_code (C++ function), 818
hpx::memory (C++ type), 482
hpx::memory::const_pointer_cast (C++ function), 483
hpx::memory::dynamic_pointer_cast (C++ function), 483
hpx::memory::get_pointer (C++ function), 483
hpx::memory::intrusive_ptr (C++ class), 483
hpx::memory::intrusive_ptr::~intrusive_ptr
hpx::memory::intrusive_ptr::detach (C++ function), 484
hpx::memory::intrusive_ptr::get (C++ function), 484

```

hpx::memory::intrusive_ptr::intrusive_ptr
 (*C++ function*), 484
hpx::memory::intrusive_ptr::operator
 bool (*C++ function*), 484
hpx::memory::intrusive_ptr::operator*
 (*C++ function*), 484
hpx::memory::intrusive_ptr::operator=
 (*C++ function*), 484
hpx::memory::intrusive_ptr::operator->
 (*C++ function*), 484
hpx::memory::intrusive_ptr::px (*C++ mem-
ber*), 485
hpx::memory::intrusive_ptr::reset (*C++
function*), 484
hpx::memory::intrusive_ptr::swap (*C++
function*), 484
hpx::memory::intrusive_ptr<T>::element_type
 (*C++ type*), 483
hpx::memory::intrusive_ptr<T>::this_type
 (*C++ type*), 484
hpx::memory::operator!= (*C++ function*), 482,
 483
hpx::memory::operator== (*C++ function*), 482,
 483
hpx::memory::operator< (*C++ function*), 483
hpx::memory::operator<< (*C++ function*), 483
hpx::memory::static_pointer_cast (*C++
function*), 483
hpx::memory::swap (*C++ function*), 483
hpx::migration_needs_retry (*C++ enumera-
tor*), 830
hpx::naming (*C++ type*), 338, 809
hpx::naming::address_type (*C++ type*), 809
hpx::naming::component_type (*C++ type*), 809
hpx::naming::invalid_locality_id (*C++
member*), 809
hpx::naming::unmanaged (*C++ function*), 338
hpx::network_error (*C++ enumerator*), 828
hpx::no_registered_console (*C++ enumera-
tor*), 829
hpx::no_state (*C++ enumerator*), 829
hpx::no_success (*C++ enumerator*), 828
hpx::not_implemented (*C++ enumerator*), 828
hpx::null_thread_id (*C++ enumerator*), 829
hpx::out_of_memory (*C++ enumerator*), 828
hpx::out_of_range (*C++ enumerator*), 830
hpx::parallel (*C++ type*), 338, 401, 416, 417, 419–
 422, 425–433, 435, 438, 439, 613, 615–620,
 623, 625, 627, 628, 630, 632, 633, 636, 639,
 642, 647, 649–652, 654, 656, 657, 661, 662,
 664, 666, 669, 677, 681, 682, 684, 685, 687,
 688, 690, 692, 695, 696, 699, 700, 702, 704,
 705, 707, 708, 710, 713, 715, 716, 719, 722,
 724, 725, 727–729, 731, 733, 736, 737, 741,
 743, 745–747, 750, 752, 754, 756, 758, 761,
 763, 766, 768, 769, 771, 875
hpx::parallel::execution (*C++ type*), 401,
 416, 417, 419–422, 425–433, 435, 438, 875
hpx::parallel::execution::auto_chunk_size
 (*C++ class*), 438
hpx::parallel::execution::auto_chunk_size::auto_chu
 (*C++ function*), 438
hpx::parallel::execution::default_executor
 (*C++ type*), 422
hpx::parallel::execution::distribution_policy_execu
 (*C++ class*), 427
hpx::parallel::execution::distribution_policy_execu
 (*C++ function*), 428
hpx::parallel::execution::distribution_policy_execu
 (*C++ member*), 428
hpx::parallel::execution::dynamic_chunk_size
 (*C++ class*), 432
hpx::parallel::execution::executor_context
 (*C++ class*), 417
hpx::parallel::execution::executor_context<Executor
 (*C++ type*), 418
hpx::parallel::execution::executor_execution_catego
 (*C++ class*), 418
hpx::parallel::execution::executor_execution_catego
 (*C++ type*), 873, 875
hpx::parallel::execution::executor_execution_catego
 (*C++ class*), 873, 875
hpx::parallel::execution::executor_execution_catego
 (*C++ type*), 418
hpx::parallel::execution::executor_execution_catego
 (*C++ type*), 418
hpx::parallel::execution::executor_index
 (*C++ class*), 418
hpx::parallel::execution::executor_index<Executor>
 (*C++ type*), 418
hpx::parallel::execution::executor_index<Executor>
 (*C++ type*), 418
hpx::parallel::execution::executor_parameters_join
 (*C++ class*), 434
hpx::parallel::execution::executor_parameters_join
 (*C++ class*), 434
hpx::parallel::execution::executor_parameters_join
 (*C++ type*), 434
hpx::parallel::execution::executor_parameters_join
 (*C++ type*), 434
hpx::parallel::execution::executor_parameters_type
 (*C++ class*), 418
hpx::parallel::execution::executor_parameters_type
 (*C++ type*), 418
hpx::parallel::execution::executor_parameters_type
 (*C++ type*), 418

hpx::parallel::execution::executor_shape^{hp}x::parallel::execution::is_timed_executor_t
 (C++ class), 418
 (C++ type), 417

hpx::parallel::execution::executor_shape^{hp}x::parallel::shape_type^{hp}x::is_two_way_executor_t
 (C++ type), 419
 (C++ type), 419

hpx::parallel::execution::executor_shape^{hp}x::parallel::type^{hp}x::join_executor_parameters
 (C++ type), 419
 (C++ function), 434

hpx::parallel::execution::extract_execut^{hp}x::parameters::execution::local_priority_queue_attachment
 (C++ class), 420
 (C++ type), 431

hpx::parallel::execution::extract_execut^{hp}x::parameters::local_priority_queue_execution_enable
 (C++ type), 421
 (C++ type), 438

hpx::parallel::execution::extract_execut^{hp}x::parameters::local_priority_queue_os_ctype
 typename hpx::util::always_void<typename(C++ type), 426
 Executor::executor_parameters_type^{hp}x::type^{hp}x::parallel::execution::main_pool_executor
 (C++ class), 420, 421
 (C++ type), 430

hpx::parallel::execution::extract_execut^{hp}x::parameters::make_distribution_policy
 typename hpx::util::always_void<typename(C++ function), 427
 Executor::executor_parameters_type^{hp}x::type^{hp}x::parallel::execution::parallel_execution_tag
 (C++ type), 420, 421
 (C++ class), 432

hpx::parallel::execution::extract_has_val^{hp}biparal^{hp}hks::size^{hp}x::parallel_executor
 (C++ class), 421
 (C++ type), 429

hpx::parallel::execution::extract_has_val^{hp}biparal^{hp}hks::size^{hp}x::parallel_executor_aggregation
 Enable>::type (C++ type), 421
 (C++ type), 435

hpx::parallel::execution::extract_has_val^{hp}biparal^{hp}hks::size^{hp}x::parallel_parallel_policy
 typename hpx::util::always_void<typename(C++ class), 401
 Parameters::has_variable_chunk_size^{hp}x::type^{hp}x::parallel::execution::parallel_policy::exec_<
 (C++ class), 420, 421
 (C++ member), 403

hpx::parallel::execution::extract_has_val^{hp}biparal^{hp}hks::size^{hp}x::parallel_parallel_policy::execution
 typename hpx::util::always_void<typename(C++ type), 401
 Parameters::has_variable_chunk_size^{hp}x::type^{hp}x::parallel::execution::parallel_policy::executor_<
 (C++ type), 420, 421
 (C++ function), 402

hpx::parallel::execution::guided_chunk_sh^{hp}x::parallel::execution::parallel_policy::executor_<
 (C++ class), 433
 (C++ type), 401

hpx::parallel::execution::guided_chunk_sh^{hp}x::parallel::guiding_chunk_size^{hp}x::parallel::execution::parallel_policy::executor_<
 (C++ function), 433
 (C++ type), 401

hpx::parallel::execution::io_pool_execut^{hp}x::parallel::execution::parallel_policy::on
 (C++ type), 430
 (C++ function), 402

hpx::parallel::execution::is_async_execu^{hp}x::parallel::execution::parallel_policy::operator<
 (C++ class), 416
 (C++ function), 402

hpx::parallel::execution::is_bulk_one_wa^{hp}x::parallel::execution::parallel_policy::parameters
 (C++ type), 419
 (C++ function), 402

hpx::parallel::execution::is_bulk_two_wa^{hp}x::parallel::execution::parallel_policy::params_<
 (C++ type), 419
 (C++ member), 403

hpx::parallel::execution::is_execution_ph^{hp}x::parallel::execution::parallel_policy::rebind
 (C++ class), 416
 (C++ class), 403

hpx::parallel::execution::is_executor_pa^{hp}x::parallel::execution::parallel_policy::rebind:
 (C++ type), 420
 (C++ type), 403

hpx::parallel::execution::is_never_block^{hp}x::operator_value^{hp}x::execution::parallel_policy::serializ_<
 (C++ type), 419
 (C++ function), 402

hpx::parallel::execution::is_one_way_exeh^{hp}x::parallel::execution::parallel_policy::with
 (C++ type), 419
 (C++ function), 402

hpx::parallel::execution::is_parallel_ex^{hp}x::parallel::execution::parallel_policy_executor
 (C++ class), 416
 (C++ class), 429

hpx::parallel::execution::is_sequenced_ex^{hp}x::parallel::policy^{hp}x::execution::parallel_policy_executor
 (C++ class), 417
 (C++ type), 429

```
hpx::parallel::execution::parallel_policy_executor(C++ type), 405 or_parameters_type
    (C++ type), 429                                     hpx::parallel::execution::parallel_task_policy::on
hpx::parallel::execution::parallel_policy_executor(C++ function), 405 policy_executor
    (C++ function), 430                                     hpx::parallel::execution::parallel_task_policy::op
hpx::parallel::execution::parallel_policy_executor(C++ function), 405
    (C++ class), 435                                     hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ function), 405 bulk_async_execute
    (C++ function), 434–436                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ member), 406
    (C++ function), 434–436                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ class), 406
    (C++ class), 434                                     hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ function), 406 bulk_async_execute
    (C++ function), 434–436                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ member), 406
    (C++ function), 434–436                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ class), 406
    (C++ class), 434, 435                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ type), 406 negotiated<hpx::launch>
    (C++ type), 434, 435                                hpx::parallel::execution::parallel_task_policy::par
hpx::parallel::execution::parallel_policy_executor(C++ function), 406
    (C++ type), 434, 435                                hpx::parallel::execution::parallel_task_policy::se
hpx::parallel::execution::parallel_policy_executor(C++ type), 406
    (C++ type), 434, 435                                hpx::parallel::execution::parallel_task_policy::wi
hpx::parallel::execution::parallel_policy_executor(C++ function), 406
    (C++ type), 435                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_executor(C++ class), 406
    (C++ type), 435                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ type), 406
    (C++ class), 403                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 406 category
    (C++ type), 403                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ type), 406
    (C++ function), 404                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ type), 406 parameters_type
    (C++ type), 403                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 406
    (C++ type), 403                                     hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 407
    (C++ function), 403                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 407
    (C++ function), 403                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 407
    (C++ function), 403                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 407
    (C++ function), 403                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ class), 407
    (C++ function), 404                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ class), 407
    (C++ function), 404                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ function), 407
    (C++ class), 404                                hpx::parallel::execution::parallel_task_policy::shin
hpx::parallel::execution::parallel_policy_shim(C++ type), 408 rebind::type (C++
    (C++ function), 404                                type), 408
hpx::parallel::execution::parallel_policy_shim(C++ type), 408 Parameters>::rebind::type (C++
    type), 404                                         (C++ type), 422
hpx::parallel::execution::parallel_unsequenced_policy(C++ type), 404
    (C++ class), 404                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ class), 408
    (C++ class), 404                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ member), 409
    (C++ member), 406                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ type), 408 category
    (C++ type), 405                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ function), 408
    (C++ function), 405                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ type), 408 parameters_type
    (C++ type), 405                                     hpx::parallel::execution::parallel_unsequenced_policy
hpx::parallel::execution::parallel_task_policy(C++ type), 408 type
```

```

hpx::parallel::execution::parallel_unseqhphpx::parallel::execution::sequenced_policy_shim::ex
    (C++ function), 408                               (C++ type), 411
hpx::parallel::execution::parallel_unseqhphpx::parallel::execution::parameters::sequenced_policy_shim::ex
    (C++ function), 408                               (C++ function), 412
hpx::parallel::execution::parallel_unseqhphpx::parallel::execution::parameters::sequenced_policy_shim::ex
    (C++ member), 409                               (C++ type), 411
hpx::parallel::execution::parallel_unseqhphpx::parallel::execution::parameters::sequenced_policy_shim::ex
    (C++ function), 408                               (C++ type), 411
hpx::parallel::execution::parcel_pool_exhphpx::parallel::execution::sequenced_policy_shim::on
    (C++ type), 430                               (C++ function), 411
hpx::parallel::execution::persistent_authhpx::parallel::execution::sequenced_policy_shim::op
    (C++ class), 431                               (C++ function), 411
hpx::parallel::execution::persistent_authhpx::parallel::execution::parameters::sequenced_policy_shim::pa
    (C++ function), 431                               (C++ function), 412
hpx::parallel::execution::pool_executor hpx::parallel::execution::sequenced_policy_shim::re
    (C++ type), 426                               (C++ class), 412
hpx::parallel::execution::rebind_executorhpx::parallel::execution::sequenced_policy_shim::w
    (C++ class), 425                               (C++ function), 411
hpx::parallel::execution::rebind_executorhpx::parallel::execution::sequenced_policy_shim<Ex
    (C++ type), 426                               Parameters>::rebind::type (C++ f
hpx::parallel::execution::sequenced_executionhpx::parallel::execution::type, 412
    (C++ class), 432                               hpx::parallel::execution::sequenced_task_policy
hpx::parallel::execution::sequenced_executor (C++ class), 412
    (C++ class), 425                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy (C++ member), 414
    (C++ class), 409                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::try
    (C++ type), 412                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::try
    (C++ function), 412                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::try
    (C++ type), 413                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::try
    (C++ function), 413                               hpx::parallel::execution::sequenced_task_policy::e
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::parameters_type
    (C++ type), 409                               hpx::parallel::execution::sequenced_task_policy::o
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::type, 413
    (C++ type), 409                               hpx::parallel::execution::sequenced_task_policy::o
hpx::parallel::execution::sequenced_policy::exechpx::parallel::execution::type
    (C++ function), 409                               hpx::parallel::execution::sequenced_task_policy::o
hpx::parallel::execution::sequenced_policy::on (C++ function), 413
    (C++ function), 409                               hpx::parallel::execution::sequenced_task_policy::p
hpx::parallel::execution::sequenced_policy::openhpx::parallel::execution::type, 413
    (C++ type), 409                               hpx::parallel::execution::sequenced_task_policy::o
hpx::parallel::execution::sequenced_policy::openhpx::parallel::execution::type
    (C++ function), 409                               hpx::parallel::execution::sequenced_task_policy::o
hpx::parallel::execution::sequenced_policy::parhpx::parallel::execution::member, 414
    (C++ function), 410                               hpx::parallel::execution::sequenced_task_policy::r
hpx::parallel::execution::sequenced_policy::parhpx::parallel::execution::type, 414
    (C++ member), 410                               hpx::parallel::execution::sequenced_task_policy::r
hpx::parallel::execution::sequenced_policy::parhpx::parallel::execution::type
    (C++ class), 410                               hpx::parallel::execution::sequenced_task_policy::r
hpx::parallel::execution::sequenced_policy::relhpx::parallel::execution::type, 414
    (C++ type), 411                               hpx::parallel::execution::sequenced_task_policy::w
hpx::parallel::execution::sequenced_policy::serhpx::parallel::execution::type, 413
    (C++ function), 410                               hpx::parallel::execution::sequenced_task_policy::sh
hpx::parallel::execution::sequenced_policy::withpx::parallel::execution::type, 414
    (C++ function), 409                               hpx::parallel::execution::sequenced_task_policy::sh
hpx::parallel::execution::sequenced_policy_shim(C++ type), 414
    (C++ class), 411                               hpx::parallel::execution::sequenced_task_policy_sh

```

(C++ function), 415
 hpx::parallel::execution::sequenced_task_policy(C++ function), 412
 (C++ type), 414
 hpx::parallel::execution::sequenced_task_policy(C++ function), 412
 (C++ type), 414
 hpx::parallel::execution::sequenced_task_policy(C++ type), 430
 (C++ function), 415
 hpx::parallel::execution::sequenced_task_policy(C++ class), 432
 (C++ function), 415
 hpx::parallel::execution::sequenced_task:parameter_projected (C++ class), 620
 hpx::parallel::execution::sequenced_task:rebands::projected::iterator_type (C++ class), 415
 (C++ type), 620
 hpx::parallel::execution::sequenced_task:withits::projected::projector_type (C++ function), 415
 (C++ type), 620
 hpx::parallel::execution::sequenced_task:projected::traits::projected_range::iterator_type Parameters>::rebind::type (C++ type), 619
 (C++ type), 416
 hpx::parallel::execution::sequenced_timed_executor(C++ type), 619
 (C++ type), 422
 hpx::parallel::execution::service_executor Rng, typename std::enable_if<hpx::traits::is_range<Rng>::value>
 hpx::parallel::execution::static_chunk_size (C++ class), 619
 (C++ class), 426
 hpx::parallel::execution::static_chunk_shape::parallel_chunk_size::accumulate (C++ function), 614
 (C++ function), 427
 hpx::parallel::execution::task(C++ member), 401
 hp^x::parallel::util::accumulate_n (C++ function), 614
 hpx::parallel::execution::thread_pool_ex:parallel::util::cancellation_token (C++ class), 428
 (C++ class), 615
 hpx::parallel::execution::thread_pool_ex:parallel::util::cancelation_token::cancel (C++ type), 428
 (C++ function), 616
 hpx::parallel::execution::thread_pool_ex:parallel::util::cancellation_token::cancelation_token::cancel (C++ type), 428
 (C++ function), 616
 hpx::parallel::execution::thread_pool_ex:parallel::util::flag_type (C++ member), 429
 (C++ type), 616
 hpx::parallel::execution::thread_pool_ex:parallel::util::field::pbol::execution_token::get_data (C++ function), 428
 (C++ function), 616
 hpx::parallel::execution::timed_executor:parallel::util::cancellation_token::was_cancelled (C++ class), 422
 (C++ function), 616
 hpx::parallel::execution::timed_executor:parallel::util::cancellation_token::was_cancelled (C++ function), 422
 (C++ member), 616
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::comp_ (C++ type), 422
 (C++ member), 618
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::comp_ (C++ member), 423
 (C++ function), 617, 618
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::operator() (C++ member), 423
 (C++ function), 617, 618
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::proj1_ (C++ type), 422
 (C++ member), 618
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::proj2_ (C++ type), 422
 (C++ member), 618
 hpx::parallel::execution::timed_executor:parallel::util::compare_projected::proj_ (C++ function), 422
 (C++ member), 618

hpx::parallel::util::compare_projected<Compare,697,752
 Proj1, Proj2> (*C++ class*), 618 hpx::parallel::v1::copy (*C++ function*), 692,
 hpx::parallel::util::compare_projected<Compare,750
 Proj> (*C++ class*), 617 hpx::parallel::v1::copy_if (*C++ function*),
 hpx::parallel::util::copy (*C++ function*),
 615 693, 751
 hpx::parallel::util::copy_synchronize
 (*C++ function*), 615 hpx::parallel::v1::copy_n (*C++ function*),
 693
 hpx::parallel::util::invoke_projected
 (*C++ class*), 617 720,
 hpx::parallel::util::invoke_projected:::invoke_projected
 (*C++ function*), 617 766
 (*C++ function*), 617 hpx::parallel::v1::count_if (*C++ function*),
 hpx::parallel::util::invoke_projected:::operator<>
 (*C++ function*), 617 720, 767
 (*C++ member*), 617 hpx::parallel::v1::destroy (*C++ function*),
 hpx::parallel::util::invoke_projected:::pred_<>
 (*C++ type*), 617 661
 (*C++ type*), 617 hpx::parallel::v1::destroy_n (*C++ func-
 hpx::parallel::util::invoke_projected:::pred_type<>
 (*C++ type*), 617 61
 (*C++ type*), 617 hpx::parallel::v1::equal (*C++ function*), 637,
 638
 hpx::parallel::util::invoke_projected:::proj_<>
 (*C++ member*), 617 664, 665
 (*C++ member*), 617 hpx::parallel::v1::fill (*C++ function*), 684,
 hpx::parallel::util::invoke_projected:::proj_type<>
 (*C++ type*), 617 745
 (*C++ type*), 617 hpx::parallel::v1::fill_n (*C++ function*),
 684, 745
 hpx::parallel::util::loop_idx_n (*C++
 function*), 614 684, 745
 hpx::parallel::util::loop_with_cleanup (*C++
 function*), 614 642
 (*C++ function*), 614 hpx::parallel::v1::find_end (*C++ function*),
 644, 734
 hpx::parallel::util::loop_with_cleanup_n
 (*C++
 function*), 646, 735
 hpx::parallel::util::loop_with_cleanup_n
 (*C++
 function*), 643
 hpx::parallel::util::make_prefetcher_con
 (*C++
 function*), 618 643
 (*C++ function*), 618 hpx::parallel::v1::for_each (*C++ function*),
 723, 768
 hpx::parallel::util::move_n (*C++ function*),
 615 hpx::parallel::v1::for_each_n (*C++ function*),
 722
 hpx::parallel::util::projection_identity
 (*C++ class*), 615 687, 746
 (*C++ type*), 621, 623, 625, 627, hpx::parallel::v1::generate_n (*C++ func-
 628, 630, 632, 633, 636, 639, 642, 647, 649,
 651, 652, 654, 656, 657, 661, 662, 664, 666,
 677, 681, 682, 684, 685, 687, 688, 690, 692,
 695, 696, 699, 700, 702, 704, 705, 707, 708,
 710, 713, 715, 716, 719, 722, 724, 725, 727,
 729, 731, 733, 736, 737, 741, 743, 745–747,
 750, 752, 754, 756, 758, 761, 763, 766, 768,
 769, 771 687
 hpx::parallel::v1::adjacent_difference
 (*C++ function*), 623 hpx::parallel::v1::generate_n (*C++ func-
 (*C++ function*), 715 687
 (*C++ function*), 715 hpx::parallel::v1::includes (*C++ function*),
 727
 hpx::parallel::v1::all_of (*C++ function*),
 698, 753 727
 hpx::parallel::v1::any_of (*C++ function*),
 625***

hpx::parallel::v1::is_sorted_until (*C++ function*), 626
hpx::parallel::v1::lexicographical_comparision (*C++ function*), 656
hpx::parallel::v1::max_element (*C++ function*), 711, 759
hpx::parallel::v1::merge (*C++ function*), 725, 771
hpx::parallel::v1::min_element (*C++ function*), 710, 758
hpx::parallel::v1::minmax_element (*C++ function*), 712, 760
hpx::parallel::v1::mismatch (*C++ function*), 705, 706
hpx::parallel::v1::move (*C++ function*), 650, 736
hpx::parallel::v1::none_of (*C++ function*), 697, 752
hpx::parallel::v1::partition (*C++ function*), 667, 741
hpx::parallel::v1::partition_copy (*C++ function*), 668, 742
hpx::parallel::v1::reduce (*C++ function*), 639–641
hpx::parallel::v1::reduce_by_key (*C++ function*), 648
hpx::parallel::v1::remove (*C++ function*), 709, 756
hpx::parallel::v1::remove_copy (*C++ function*), 621, 730
hpx::parallel::v1::remove_copy_if (*C++ function*), 621, 730
hpx::parallel::v1::remove_if (*C++ function*), 708, 757
hpx::parallel::v1::replace (*C++ function*), 657, 737
hpx::parallel::v1::replace_copy (*C++ function*), 659, 739
hpx::parallel::v1::replace_copy_if (*C++ function*), 660, 740
hpx::parallel::v1::replace_if (*C++ function*), 658, 738
hpx::parallel::v1::reverse (*C++ function*), 699, 754
hpx::parallel::v1::reverse_copy (*C++ function*), 700, 755
hpx::parallel::v1::rotate (*C++ function*), 681, 743
hpx::parallel::v1::rotate_copy (*C++ function*), 681, 744
hpx::parallel::v1::search (*C++ function*), 690, 747
hpx::parallel::v1::search_n (*C++ function*), 691, 748
hpx::parallel::v1::set_difference (*C++ function*), 702
hpx::parallel::v1::set_intersection (*C++ function*), 662
hpx::parallel::v1::set_symmetric_difference (*C++ function*), 630
hpx::parallel::v1::set_union (*C++ function*), 685
hpx::parallel::v1::sort (*C++ function*), 724, 769
hpx::parallel::v1::sort_by_key (*C++ function*), 707
hpx::parallel::v1::stable_partition (*C++ function*), 666
hpx::parallel::v1::swap_ranges (*C++ function*), 651
hpx::parallel::v1::transform (*C++ function*), 716–718, 763–765
hpx::parallel::v1::transform_exclusive_scan (*C++ function*), 695
hpx::parallel::v1::transform_inclusive_scan (*C++ function*), 634, 635
hpx::parallel::v1::transform_reduce (*C++ function*), 652, 653, 682
hpx::parallel::v1::uninitialized_copy (*C++ function*), 654
hpx::parallel::v1::uninitialized_copy_n (*C++ function*), 655
hpx::parallel::v1::uninitialized_default_construct (*C++ function*), 688
hpx::parallel::v1::uninitialized_default_construct_n (*C++ function*), 689
hpx::parallel::v1::uninitialized_fill (*C++ function*), 701
hpx::parallel::v1::uninitialized_fill_n (*C++ function*), 701
hpx::parallel::v1::uninitialized_move (*C++ function*), 629
hpx::parallel::v1::uninitialized_move_n (*C++ function*), 629
hpx::parallel::v1::uninitialized_value_construct (*C++ function*), 627
hpx::parallel::v1::uninitialized_value_construct_n (*C++ function*), 627
hpx::parallel::v1::unique (*C++ function*), 713, 761
hpx::parallel::v1::unique_copy (*C++ function*), 714, 762
hpx::parallel::v2 (*C++ type*), 338, 650, 669, 728
hpx::parallel::v2::define_task_block (*C++ function*), 338
hpx::parallel::v2::define_task_block_restore_thread (*C++ function*), 339
hpx::parallel::v2::for_loop (*C++ function*),


```

hpx::program_options::error_with_option_hpx::pgetan option'san option'sopenenoneonemaking_helper::parse_<  

    (C++ function), 455                               (C++ member), 469
hpx::program_options::error_with_option_hpx::pgetan option'san option'sopenenoneoneprefixingixing_helper::split_u<  

    (C++ function), 455                               (C++ member), 469
hpx::program_options::error_with_option_hpx::pgetoptioptionnamename:full_match   (C++  

    (C++ function), 455                           enumerator), 460
hpx::program_options::error_with_option_hpx::pmogramrogramopemphemphateinclude_positional  

    (C++ member), 455                           (C++ enumerator), 464
hpx::program_options::error_with_option_hpx::pmearmearmagationsations::invalid_bool_value  

    (C++ member), 456                           (C++ class), 456
hpx::program_options::error_with_option_hpx::pmoptioptionstst:invalid_bool_value  

    (C++ member), 456                           (C++ enumerator), 459
hpx::program_options::error_with_option_hpx::pmognatognatoptionosos:defaultalid_bool_value::invalid_b<  

    (C++ member), 456                           (C++ function), 456
hpx::program_options::error_with_option_hpx::pmognatognatoptionosos:invalid_command_line_style  

    (C++ member), 456                           (C++ class), 456
hpx::program_options::error_with_option_hpx::pogramogramoptionensens::invalid_command_line_style::<  

    (C++ function), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::psetramramoptionnamename:invalid_command_line_syntax  

    (C++ function), 455                           (C++ class), 456
hpx::program_options::error_with_option_hpx::psetramramoptionnamename:invalid_config_file_syntax:  

    (C++ function), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::psetramramoptionosos:defaultalid_config_file_syntax::<  

    (C++ function), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::psetramramoptionosos:invalid_config_file_syntax::<  

    (C++ type), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::pogramogramoptionnsns:invalid_config_file_syntax::<  

    (C++ function), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::pogramogramoptionnsns:invalid_config_file_syntax::<  

    (C++ function), 455                           (C++ function), 456
hpx::program_options::error_with_option_hpx::pogramogramoptionnsns:invalid_option  

    (C++ function), 455                           (C++ enumerator), 459
hpx::program_options::exclude_positionalhpx::program_options::invalid_option_value  

    (C++ enumerator), 464                           (C++ class), 456
hpx::program_options::ext_parser   (C++ hpx::program_options::invalid_option_value  

    type), 464                           (C++ enumerator), 459
hpx::program_options::extra_parameter      hpx::program_options::invalid_option_value::invalid_o<  

    (C++ enumerator), 457                           (C++ function), 457
hpx::program_options::force_linking       hpx::program_options::invalid_syntax  

    (C++ function), 469                           (C++ class), 457
hpx::program_options::force_linkinghpx::program_options::invalid_syntax::~invalid_synt<  

    (C++ class), 469                           (C++ function), 457
hpx::program_options::force_linkinghpx::program_options::programrogramoptiongg:challll_syntax::get_template  

    (C++ member), 469                           (C++ function), 457
hpx::program_options::force_linkinghpx::program_options::programrogramoptiongg:challll_syntax::invalid_synt<  

    (C++ member), 469                           (C++ function), 457
hpx::program_options::force_linkinghpx::program_options::programrogramoptiongg:challll_syntax::kind  

    (C++ member), 469                           (C++ function), 457
hpx::program_options::force_linkinghpx::program_options::programrogramoptiongg:invalid_syntax::m_kind  

    (C++ member), 469                           (C++ member), 458
hpx::program_options::force_linkinghpx::program_options::programrogramoptiongg:invalid_syntax::tokens  

    (C++ member), 469                           (C++ function), 457

```

hpx::program_options::kind_t (*C++ enum*), hpx::program_options::option_description::m_value_size
457, 459
hpx::program_options::long_adjacent_not_hipflow::program_options::option_description::match
(*C++ enumerator*), 457
hpx::program_options::long_not_allowed hpx::program_options::option_description::option_descriptions
(*C++ enumerator*), 457
hpx::program_options::match_result (*C++ enum*) hpx::program_options::option_description::semantic
enum), 460
hpx::program_options::missing_parameter hpx::program_options::option_description::set_names
(*C++ enumerator*), 457
hpx::program_options::multiple_occurrence hpx::program_options::optional (*C++ type*),
(*C++ class*), 458
hpx::program_options::multiple_occurrence hpx::program_options::options_description
(*C++ function*), 458
hpx::program_options::multiple_occurrence hpx::program_options::option_exceptions::add
(*C++ function*), 458
hpx::program_options::multiple_values hpx::program_options::options_description::add_option
(*C++ class*), 458
hpx::program_options::multiple_values::~hpx::program_options::options_description::approximate
(*C++ function*), 458
hpx::program_options::multiple_values::mbkippeogramme::options_description::belong_to
(*C++ function*), 458
hpx::program_options::multiple_values_nohipflow::program_options::options_description::find
(*C++ enumerator*), 459
hpx::program_options::no_match (*C++ enumerator*), 460
hpx::program_options::notify (*C++ function*), 450
hpx::program_options::option (*C++ type*), 472
hpx::program_options::option_description hpx::program_options::options_description::groups
(*C++ class*), 460
hpx::program_options::option_description hpx::program_options::options_description::m_caption
(*C++ function*), 461
hpx::program_options::option_description hpx::program_options::options_description::m_default
(*C++ function*), 461
hpx::program_options::option_description hpx::program_options::options_description::m_line_index
(*C++ function*), 461
hpx::program_options::option_description hpx::program_options::options_description::m_min_depth
(*C++ function*), 461
hpx::program_options::option_description hpx::program_options::options_description::m_option
(*C++ function*), 461
hpx::program_options::option_description hpx::program_options::options_description::name2index
(*C++ function*), 461
hpx::program_options::option_description hpx::keyprogram_options::options_description::options
(*C++ function*), 461
hpx::program_options::option_description hpx::longprogram_options::options_description::options
(*C++ function*), 461
hpx::program_options::option_description hpx::longprogram_options::options_description::print
(*C++ function*), 461
hpx::program_options::option_description hpx::option::options::options_description::easy_init
(*C++ member*), 462
hpx::program_options::option_description hpx::longprogram_options::options_description::easy_init
(*C++ member*), 462
hpx::program_options::option_description hpx::shortprogram_options::options_description::easy_init
(*C++ member*), 462

```

hpx::program_options::options_description hpx::program_options::strip_prefixes
    (C++ member), 464                               (C++ function), 453
hpx::program_options::parse_command_line hpx::program_options::too_many_positional_options_error
    (C++ function), 464                               (C++ class), 458
hpx::program_options::parse_config_file hpx::program_options::too_many_positional_options_error
    (C++ function), 464                               (C++ function), 459
hpx::program_options::parse_config_file_hp::pari::program_options::typed_value (C++ type)
    (C++ type), 468                               (C++ class), 473
hpx::program_options::parse_config_file_hp::pari::program_options::typed_value::apply_default
    (C++ type), 468                               (C++ function), 474
hpx::program_options::parse_config_file_hp::pari::program_options::typed_value::composing
    (C++ type), 468                               (C++ function), 474
hpx::program_options::parse_environment hpx::program_options::typed_value::default_value
    (C++ function), 465                               (C++ function), 473
hpx::program_options::parse_environment1 hpx::program_options::typed_value::implicit_value
    (C++ type), 468                               (C++ function), 474
hpx::program_options::parse_environment2 hpx::program_options::typed_value::is_composing
    (C++ type), 468                               (C++ function), 474
hpx::program_options::parse_environment3 hpx::program_options::typed_value::is_required
    (C++ type), 468                               (C++ function), 474
hpx::program_options::parsed_options      hpx::program_options::typed_value::m_composing
    (C++ type), 464                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_default_value
    (C++ class), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_default_value
    (C++ function), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_implicit
    (C++ member), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_implicit
    (C++ member), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_implicit_value
    (C++ function), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_multitoken
    (C++ function), 449                               (C++ member), 475
hpx::program_options::positional_options hpx::program_options::typed_value::m_no_pifor
    (C++ function), 449                               (C++ member), 475
hpx::program_options::reading_file (C++ class) hpx::program_options::typed_value::m_required
    (C++ member), 458                               (C++ member), 475
hpx::program_options::reading_file::read hpx::program_options::typed_value::m_store_to
    (C++ function), 458                               (C++ member), 475
hpx::program_options::required_option      hpx::program_options::typed_value::m_value_name
    (C++ class), 458                               (C++ member), 475
hpx::program_options::required_option::~hpx::program_options::typed_value::m_zero_tokens
    (C++ function), 458                               (C++ member), 475
hpx::program_options::required_option::rhs hpx::program_options::typed_value::max_tokens
    (C++ function), 458                               (C++ function), 474
hpx::program_options::short_adjacent_noth hpx::program_options::typed_value::min_tokens
    (C++ enumerator), 457                               (C++ function), 474
hpx::program_options::split_unix     (C++ function) hpx::program_options::typed_value::multitoken
    (C++ function), 465                               (C++ function), 474
hpx::program_options::split_unix_type hpx::program_options::typed_value::name
    (C++ type), 468                               (C++ function), 474
hpx::program_options::store (C++ function), 450 hpx::program_options::typed_value::notifier
                                            (C++ function), 474

```

```
hpx::program_options::typed_value::notify hpx::program_options::validation_error::get_template
    (C++ function), 474                                (C++ function), 459
hpx::program_options::typed_value::require hpx::program_options::validation_error::kind
    (C++ function), 474                                (C++ function), 459
hpx::program_options::typed_value::typed hpx::program_options::validation_error::m_kind
    (C++ function), 473                                (C++ member), 460
hpx::program_options::typed_value::value hpx::program_options::validation_error::validation_
    (C++ function), 474                                (C++ function), 459
hpx::program_options::typed_value::value hpx::program_options::value (C++ function),
    (C++ function), 474                                473
hpx::program_options::typed_value::xparse hpx::program_options::value_semantic
    (C++ function), 474                                (C++ class), 476
hpx::program_options::typed_value::zero hpx::program_options::value_semantic::~value_semantic
    (C++ function), 474                                (C++ function), 476
hpx::program_options::typed_value_base hpx::program_options::value_semantic::apply_default
    (C++ class), 475                                (C++ function), 476
hpx::program_options::typed_value_base::hpx::program_options::value_semantic::is_composing
    (C++ function), 475                                (C++ function), 476
hpx::program_options::typed_value_base::hpx::program_options::value_semantic::is_required
    (C++ function), 475                                (C++ function), 476
hpx::program_options::unknown_option hpx::program_options::value_semantic::max_tokens
    (C++ class), 459                                (C++ function), 476
hpx::program_options::unknown_option::uh hpx::program_options::value_semantic::min_tokens
    (C++ function), 459                                (C++ function), 476
hpx::program_options::unknown_option::unkn hpx::program_options::value_semantic::name
    (C++ function), 459                                (C++ function), 476
hpx::program_options::unrecognized_line hpx::program_options::value_semantic::notify
    (C++ enumerator), 457                            (C++ function), 476
hpx::program_options::untypesd_value hpx::program_options::value_semantic::parse
    (C++ class), 475                                (C++ function), 476
hpx::program_options::untypesd_value::app hpx::program_options::value_semantic_codecvt_helper
    (C++ function), 476                            (C++ class), 477
hpx::program_options::untypesd_value::is_hpx::program_options::value_semantic_codecvt_helper
    (C++ function), 475                            (C++ function), 477
hpx::program_options::untypesd_value::is_hpx::program_options::value_semantic_codecvt_helper
    (C++ function), 475                            (C++ function), 477
hpx::program_options::untypesd_value::m_z hpx::program_options::value_semantic_codecvt_helper
    (C++ member), 476                            (C++ class), 477
hpx::program_options::untypesd_value::max hpx::program_options::value_semantic_codecvt_helper
    (C++ function), 475                            (C++ class), 477
hpx::program_options::untypesd_value::min hpx::program_options::variable_value
    (C++ function), 475                            (C++ class), 451
hpx::program_options::untypesd_value::nam hpx::program_options::variable_value::as
    (C++ function), 475                            (C++ function), 451
hpx::program_options::untypesd_value::not hpx::program_options::variable_value::defaulted
    (C++ function), 476                            (C++ function), 451
hpx::program_options::untypesd_value::un hpx::program_options::variable_value::empty
    (C++ function), 475                            (C++ function), 451
hpx::program_options::untypesd_value::xpah hpx::program_options::variable_value::m_defaulted
    (C++ function), 476                            (C++ member), 452
hpx::program_options::validation_error hpx::program_options::variable_value::m_value_sema
    (C++ class), 459                            (C++ member), 452
hpx::program_options::validation_error::hpx::lidatbameoptions::variable_value::v
    (C++ function), 459                            (C++ member), 452
```

hpx::program_options::variable_value::value *function*, 813
(C++ function), 451 hpx::resiliency::async_replicate_validate
hpx::program_options::variable_value::variable (C++ function), 813
(C++ function), 451 hpx::resiliency::async_replicate_vote
(C++ function), 812
hpx::program_options::variables_map (C++ class), 452 hpx::resiliency::async_replicate_vote_validate
(C++ function), 452 hpx::resiliency::dataflow_replay (C++
(C++ function), 452 hpx::resiliency::dataflow_replay_validate
hpx::program_options::variables_map::clear (C++ function), 812
(C++ function), 452 hpx::resiliency::dataflow_replicate
hpx::program_options::variables_map::get (C++ function), 817
(C++ function), 452 hpx::resiliency::dataflow_replicate_validate
hpx::program_options::variables_map::m_final (C++ function), 816
(C++ member), 452 hpx::resiliency::force_linking (C++ func-
hpx::program_options::variables_map::m_required (C++ function), 816
(C++ member), 452 hpx::resiliency::force_linking_helper
(C++ function), 452 hpx::resiliency::force_linking_helper::full_version
hpx::program_options::variables_map::operator[] (C++ function), 815
(C++ function), 452 hpx::resiliency::force_linking_helper::full_version
hpx::program_options::variables_map::variables (C++ function), 815
(C++ function), 452 hpx::resiliency::force_linking_helper::major_version
hpx::program_options::wcommand_line_parser (C++ type), 464 hpx::resiliency::force_linking_helper::minor_version
(C++ type), 464 hpx::resiliency::force_linking_helper::subminor_version
hpx::program_options::woption (C++ type), 472 hpx::resiliency::full_version (C++ func-
hpx::program_options::wparsed_options (C++ type), 464 tion), 814
hpx::program_options::wvalue (C++ func-
tion), 473 hpx::resiliency::full_version_str (C++
hpx::promise_already_satisfied (C++ enum-
erator), 829 function), 814
hpx::register_on_exit (C++ function), 294 hpx::resiliency::functional (C++ type),
hpx::register_pre_shutdown_function (C++ function), 309 812, 813
hpx::register_pre_startup_function (C++ function), 309 hpx::resiliency::functional::async_replay
hpx::register_shutdown_function (C++ function), 309 (C++ class), 812
hpx::register_startup_function (C++ function), 310 hpx::resiliency::functional::async_replay::operator
(C++ function), 812
hpx::register_thread (C++ function), 294 hpx::resiliency::functional::async_replay_validate
hpx::register_with_basename (C++ function), 299 (C++ class), 812
hpx::reinit_active_counters (C++ function), 296 hpx::resiliency::functional::async_replay_validate
(C++ function), 812
hpx::repeated_request (C++ enumerator), 828 hpx::resiliency::functional::async_replicate
hpx::report_error (C++ function), 309 (C++ class), 813
hpx::reset_active_counters (C++ function), 295 hpx::resiliency::functional::async_replicate::operator
(C++ function), 813
hpx::resiliency (C++ type), 811, 812, 814–816 hpx::resiliency::functional::async_replicate_validate
(C++ function), 811 (C++ class), 813
hpx::resiliency::async_replay (C++ func-
tion), 811 hpx::resiliency::functional::async_replicate_validate
(C++ function), 813
hpx::resiliency::async_replay_validate (C++ function), 811 hpx::resiliency::functional::async_replicate_validate
(C++ class), 813

(C++ function), 814
hpx::resiliency::functional::async_replicate_vtox (C++ class), 814
hpx::resiliency::functional::async_replicate_vtox_member (C++ function), 814
hpx::resiliency::functional::async_replicate_vtox_member (C++ class), 814
hpx::resiliency::functional::async_replicate_vtox_member (C++ function), 814
hpx::resiliency::functional::major_version (C++ function), 814
hpx::resiliency::minor_version (C++ function), 814
hpx::resiliency::subminor_version (C++ function), 814
hpx::resource (C++ type), 911, 915
hpx::resource::abp_priority_fifo (C++ enumerator), 915
hpx::resource::abp_priority_lifo (C++ enumerator), 915
hpx::resource::core (C++ class), 911
hpx::resource::core::core (C++ function), 911
hpx::resource::core::cores_sharing_numa (C++ function), 911
hpx::resource::core::domain_ (C++ member), 912
hpx::resource::core::id (C++ function), 911
hpx::resource::core::id_ (C++ member), 912
hpx::resource::core::invalid_core_id (C++ member), 912
hpx::resource::core::pus (C++ function), 911
hpx::resource::core::pus_ (C++ member), 912
hpx::resource::get_partitioner (C++ function), 916
hpx::resource::is_partitioner_valid (C++ function), 916
hpx::resource::local (C++ enumerator), 915
hpx::resource::local_priority_fifo (C++ enumerator), 915
hpx::resource::local_priority_lifo (C++ enumerator), 915
hpx::resource::mode_allow_dynamic_pools (C++ enumerator), 915
hpx::resource::mode_allow_oversubscription (C++ enumerator), 915
hpx::resource::mode_default (C++ enumerator), 915
hpx::resource::numa_domain (C++ class), 912
hpx::resource::numa_domain::cores (C++ function), 912
hpx::resource::numa_domain::cores_ (C++ member), 912
hpx::resource::numa_domain::id (C++ function), 912
hpx::resource::numa_domain::id_ (C++ operator()), 912
hpx::resource::numa_domain::invalid_numa_domain_id (C++ operator()), 912
hpx::resource::partitioner (C++ class), 912
hpx::resource::partitioner::add_resource (C++ function), 913, 914
hpx::resource::partitioner::create_thread_pool (C++ function), 913
hpx::resource::partitioner::get_default_pool_name (C++ function), 913
hpx::resource::partitioner::get_number_requested_threads (C++ function), 914
hpx::resource::partitioner::get_topology (C++ function), 914
hpx::resource::partitioner::numa_domains (C++ function), 914
hpx::resource::partitioner::partitioner (C++ function), 913
hpx::resource::partitioner::partitioner_ (C++ member), 914
hpx::resource::partitioner::set_default_pool_name (C++ function), 913
hpx::resource::partitioner_mode (C++ enum), 915
hpx::resource::pu (C++ class), 914
hpx::resource::pu::core_ (C++ member), 914
hpx::resource::pu::id (C++ function), 914
hpx::resource::pu::id_ (C++ member), 914
hpx::resource::pu::invalid_pu_id (C++ member), 915
hpx::resource::pu::pu (C++ function), 914
hpx::resource::pu::pus_sharing_core (C++ function), 914
hpx::resource::pu::pus_sharing_numa_domain (C++ function), 914
hpx::resource::pu::thread_occupancy_ (C++ member), 914
hpx::resource::pu::thread_occupancy_count_ (C++ member), 914
hpx::resource::scheduler_function (C++ type), 915
hpx::resource::scheduling_policy (C++ enum), 915
hpx::resource::shared_priority (C++ enumerator), 916
hpx::resource::static_ (C++ enumerator), 915
hpx::resource::static_priority (C++ enumerator), 915
hpx::resource::unspecified (C++ enumerator), 915

tor), 915

`hpx::resource::user_defined (C++ enumerator), 915`

`hpx::resume (C++ function), 294`

`hpx::rethrow (C++ enumerator), 825`

`hpx::runtime_mode (C++ enum), 271`

`hpx::runtime_mode_connect (C++ enumerator), 271`

`hpx::runtime_mode_console (C++ enumerator), 271`

`hpx::runtime_mode_default (C++ enumerator), 271`

`hpx::runtime_mode_invalid (C++ enumerator), 271`

`hpx::runtime_mode_last (C++ enumerator), 271`

`hpx::runtime_mode_worker (C++ enumerator), 271`

`hpx::serialization (C++ type), 485, 523, 525–533, 535, 536, 538–540, 542, 544–546, 883, 894–896, 899`

`hpx::serialization::access (C++ class), 536`

`hpx::serialization::access::get_name (C++ function), 536`

`hpx::serialization::access::has_serialize (C++ class), 537`

`hpx::serialization::access::has_serialize::test (C++ class), 535`

`hpx::serialization::access::has_serialize::value (C++ function), 535, 536`

`hpx::serialization::access::load_base_object (C++ function), 536`

`hpx::serialization::access::save_base_object (C++ function), 536`

`hpx::serialization::access::serialize (C++ function), 536`

`hpx::serialization::access::serialize_dispatch (C++ class), 537`

`hpx::serialization::all_archive_flags (C++ enumerator), 524`

`hpx::serialization::archive_flags (C++ enum), 524`

`hpx::serialization::array (C++ class), 541`

`hpx::serialization::array::address (C++ function), 541`

`hpx::serialization::array::array (C++ function), 541`

`hpx::serialization::array::count (C++ function), 541`

`hpx::serialization::array::m_element_count (C++ member), 542`

`hpx::serialization::array::m_t (C++ member), 542`

`hpx::serialization::array::serialize (C++ function), 541`

`hpx::serialization::array::serialize_optimized (C++ function), 541`

`hpx::serialization::array<T>::value_type (C++ type), 541`

`hpx::serialization::base_object (C++ function), 535`

`hpx::serialization::base_object_type`

`hpx::serialization::base_object::base_object_type (C++ function), 535`

`hpx::serialization::base_object::base_object_type::d_`

`hpx::serialization::base_object::base_object_type::HPX_SERIALIZATION_ID (C++ function), 535, 536`

`hpx::serialization::base_object::load (C++ function), 535, 536`

`hpx::serialization::base_object::save (C++ function), 535, 536`

`hpx::serialization::base_object_type::load (C++ function), 535, 536`

`hpx::serialization::base_object_type::save (C++ function), 535, 536`

`hpx::serialization::base_object_type<Derived, Base>::Base::type (C++ class), 534, 536`

`hpx::serialization::basic_archive (C++ class), 524`

`hpx::serialization::basic_archive::~basic_archive (C++ function), 524`

`hpx::serialization::basic_archive::basic_archive (C++ function), 525`

`hpx::serialization::basic_archive::basic_archive::basic_archive (C++ function), 525`

`hpx::serialization::basic_archive::polymorphic::current_pos (C++ function), 524`

`hpx::serialization::basic_archive::polymorphic::enable_array_operations (C++ function), 524`

`hpx::serialization::basic_archive::polymorphic::enable_archive::enable_data_change (C++ function), 524`

`hpx::serialization::basic_archive::enable_compression (C++ function), 524`

(C++ function), 524
hpx::serialization::basic_archive::endian~~hpx~~serialization::chunk_data::pos_ (C++ member), 530
(C++ function), 524
hpx::serialization::basic_archive::endian~~hpx~~serialization::chunk_type (C++ enum), 530
(C++ function), 524
hpx::serialization::basic_archive::extra~~hpx~~serialization::chunk_type_index (C++ enumerator), 530
(C++ function), 525
hpx::serialization::basic_archive::flags~~hpx~~serialization::chunk_type_pointer (C++ enumerator), 530
(C++ function), 524
hpx::serialization::basic_archive::flags~~hpx~~serialization::copy (C++ enumerator), 542
(C++ member), 525
hpx::serialization::basic_archive::get~~hpx~~serialization::create_index_chunk (C++ function), 530
(C++ function), 525
hpx::serialization::basic_archive::invok~~hpx~~serialization::create_pointer_chunk (C++ function), 530
(C++ function), 524
hpx::serialization::basic_archive::is~~hpx~~serialization::current_pos (C++ function), 524
(C++ function), 524
hpx::serialization::basic_archive::load~~hpx~~serialization::disable_array_optimization (C++ enumerator), 524
(C++ function), 524
hpx::serialization::basic_archive::npos hpx::serialization::disable_data_chunking (C++ enumerator), 524
(C++ member), 525
hpx::serialization::basic_archive::opera~~hpx~~serialization::enable_compression (C++ enumerator), 524
(C++ function), 525
hpx::serialization::basic_archive::resethpx::serialization::endian_big (C++ enumerator), 524
(C++ function), 525
hpx::serialization::basic_archive::save~~hpx~~serialization::endian_little (C++ enumerator), 524
(C++ function), 524
hpx::serialization::basic_archive::size~~hpx~~serialization::erased_input_container (C++ class), 528
(C++ member), 525
hpx::serialization::basic_archive::try~~hpx~~serialization::erased_input_container::~erased_ (C++ function), 528
(C++ function), 525
hpx::serialization::binary_filter (C++ hpx::serialization::erased_input_container::is_pre_ class), 533
(C++ function), 528
hpx::serialization::binary_filter::~binary~~hpx~~filter::erased_input_container::load_b_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::flush~~hpx~~serialization::erased_input_container::load_b_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::HPX_SERIALISATI~~ON_ZEROMORETHAN_ABSTRACT~~container::set_fi_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::init~~hpx~~serialization::erased_output_container (C++ class), 528
(C++ function), 533
hpx::serialization::binary_filter::load hpx::serialization::erased_output_container::~erased_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::save hpx::serialization::erased_output_container::flush_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::serializ~~hpx~~serialization::erased_output_container::get_nu_ (C++ function), 528
(C++ function), 533
hpx::serialization::binary_filter::set~~hpx~~length::alization::erased_output_container::is_p_ (C++ function), 528
(C++ function), 533
hpx::serialization::chunk_data (C++ hpx::serialization::erased_output_container::reset_ union), 530
(C++ function), 528
hpx::serialization::chunk_data::cpos_ hpx::serialization::erased_output_container::save_k_ (C++ function), 528
(C++ member), 530
hpx::serialization::chunk_data::index_ hpx::serialization::erased_output_container::save_k_ (C++ function), 528
(C++ function), 530


```
hpx::serialization::operator>> (C++ function), 525, 535, 540, 541           hpx::serialization::output_container::get_num_chunks (C++ function), 527
hpx::serialization::operator<< (C++ function), 525, 535, 540, 541           hpx::serialization::output_container::is_preprocessed (C++ function), 527
hpx::serialization::output_archive (C++ class), 531                         hpx::serialization::output_container::output_container (C++ function), 527
hpx::serialization::output_archive::basehpx serialization::output_container::reset (C++ type), 531          (C++ function), 527
hpx::serialization::output_archive::buffhpx serialization::output_container::save_binary (C++ member), 532          (C++ function), 527
hpx::serialization::output_archive::bytehpx serialization::output_container::save_binary (C++ function), 531          (C++ function), 527
hpx::serialization::output_archive::currhpx serialization::output_container::set_filter (C++ function), 531          (C++ function), 527
hpx::serialization::output_archive::flushhpx serialization::output_container<Container, Chunker>::access_traits (C++ type), 531
hpx::serialization::output_archive::get_num_chunks (C++ function), 531           hpx::serialization::reference (C++ enum)
hpx::serialization::output_archive::invoke_implementor, 542
hpx::serialization::output_archive::is_preprocessed (C++ function), 531           hpx::serialization::reverse_bytes (C++ function)
hpx::serialization::output_archive::make_flags (C++ function), 532           hpx::serialization::save (C++ function), 485
hpx::serialization::output_archive::output_archive (C++ function), 524
hpx::serialization::output_archive::reset (C++ function), 531           hpx::serialization::save_binary (C++ function), 529, 530, 540
hpx::serialization::output_archive::save (C++ class), 530           hpx::serialization::save_construct_data
hpx::serialization::output_archive::save_bitwise (C++ function), 530           hpx::serialization::serialization_chunk
hpx::serialization::output_archive::save_binary (C++ member), 530           hpx::serialization::serialization_chunk::data_
hpx::serialization::output_archive::save_binary (C++ member), 530           hpx::serialization::serialization_chunk::rkey_
hpx::serialization::output_archive::save_binary (C++ member), 530           hpx::serialization::serialization_chunk::size_
hpx::serialization::output_archive::save_bitwise (C++ function), 531           hpx::serialization::serialization_chunk::type_
hpx::serialization::output_archive::save_integral (C++ member), 530           hpx::serialization::serialize (C++ function)
hpx::serialization::output_archive::save_integral (C++ function), 532           hpx::serialization::serialize (function), 521, 527, 529, 539, 540, 544, 546, 883, 894–896, 899
hpx::serialization::output_archive::savehpx::nonserializable_type::polymorphize_buffer (C++ class), 542
hpx::serialization::output_container (C++ class), 526           hpx::serialization::serialize_buffer::alloc_ (C++ member), 543
hpx::serialization::output_container::~hpx::putstream::deinitialization::serialize_buffer::begin (C++ function), 543
hpx::serialization::output_container::chhpx serialization::serialize_buffer::data (C++ function), 543
hpx::serialization::output_container::cohhpx serialization::serialize_buffer::data_ (C++ member), 543
hpx::serialization::output_container::cuhpx serialization::serialize_buffer::data_array (C++ member), 543
hpx::serialization::output_container::flhpx serialization::serialize_buffer::deleter (C++ function), 544
```

```

hpx::serialization::serialize_buffer::end      829
    (C++ function), 543                      hpx::task_block_not_active (C++ enumerator),
hpx::serialization::serialize_buffer::load     830
    (C++ function), 543                      hpx::task_canceled_exception (C++ enumerator),
hpx::serialization::serialize_buffer::no_deleter, 830
    (C++ function), 544                      hpx::task_moved (C++ enumerator), 829
hpx::serialization::serialize_buffer::ophpx::task::hpx::this_thread (C++ type), 343
    (C++ function), 543                      hpx::this_thread::get_executor (C++ function),
hpx::serialization::serialize_buffer::save      345
    (C++ function), 543                      hpx::this_thread::get_pool (C++ function),
hpx::serialization::serialize_buffer::size      343–345
    (C++ function), 543                      hpx::this_thread::suspend (C++ function),
hpx::serialization::serialize_buffer::size      343–345
    (C++ function), 543                      hpx::thread_cancelled (C++ enumerator), 830
hpx::serialization::serialize_buffer::size      343–345
    (C++ member), 543                      hpx::thread_interrupted (C++ class), 824
hpx::serialization::serialize_buffer<T,           hpx::thread_not_interruptable (C++ enumerator), 830
    Allocator>::allocator_type (C++ type), 543  hpx::thread_resource_error (C++ enumerator), 829
hpx::serialization::serialize_buffer<T,           hpx::threads (C++ type), 346, 423, 425, 436, 439,
    Allocator>::value_type (C++ type), 542       570, 610, 783, 884, 887, 890, 891, 904
hpx::serialization::serialize_non_intrusive,      hpx::threads::active (C++ enumerator), 887
    (C++ function), 536, 538                  436
hpx::serialization::serialize_non_intrusive,      hpx::threads::async_execute (C++ function),
    (C++ class), 536, 538                  424
hpx::serialization::take (C++ enumerator),        hpx::threads::async_execute_after (C++ type)
    542                                         typename std::enable_if<has_serialize_additionfunction>::type>
hpx::serialization_error (C++ enumerator),        hpx::threads::async_execute_at (C++ function),
    829                                         424
hpx::service_unavailable (C++ enumerator),       hpx::threads::bulk_async_execute (C++ function),
    828                                         437
hpx::set_custom_exception_info_handler          hpx::threads::bulk_sync_execute (C++ function),
    (C++ function), 821                         437
hpx::set_lco_error (C++ function), 313–315
hpx::set_lco_value (C++ function), 311, 312
hpx::set_lco_value_unmanaged (C++ function),   hpx::threads::bulk_then_execute (C++ function),
    312, 313                                     437
hpx::set_pre_exception_handler (C++ function), hpx::threads::coroutines (C++ type), 885,
    821                                         891
hpx::shutdown_function_type (C++ type), 271
hpx::split_future (C++ function), 318, 319
hpx::start (C++ function), 281–291
hpx::start_active_counters (C++ function),   hpx::threads::coroutines::context_state
    295                                         (C++ enum), 886
hpx::startup_function_type (C++ type), 271
hpx::startup_timed_out (C++ enumerator), 829
hpx::stop (C++ function), 293
hpx::stop_active_counters (C++ function),   hpx::threads::coroutines::coroutine
    296                                         (C++ class), 891
hpx::success (C++ enumerator), 828
hpx::suspend (C++ function), 293
hpx::task_already_started (C++ enumerator),

```


(C++ function), 347
hpx::threads::get_self_id (*C++ function*), 346, 611
hpx::threads::get_self_id_data (*C++ function*), 346
hpx::threads::get_self_ptr (*C++ function*), 346, 611
hpx::threads::get_self_ptr_checked (*C++ function*), 346
hpx::threads::get_self_stacksize (*C++ function*), 346
hpx::threads::get_stack_size (*C++ function*), 352
hpx::threads::get_stack_size_name (*C++ function*), 889
hpx::threads::get_thread_count (*C++ function*), 347
hpx::threads::get_thread_description (*C++ function*), 349
hpx::threads::get_thread_interruption_enhanced (*C++ function*), 350
hpx::threads::get_thread_interruption_repeated (*C++ function*), 351
hpx::threads::get_thread_lco_description (*C++ function*), 349
hpx::threads::get_thread_phase (*C++ function*), 350
hpx::threads::get_thread_priority (*C++ function*), 352
hpx::threads::get_thread_priority_name (*C++ function*), 889
hpx::threads::get_thread_state (*C++ function*), 349
hpx::threads::get_thread_state_ex_name (*C++ function*), 889
hpx::threads::get_thread_state_name (*C++ function*), 889
hpx::threads::hardware_concurrency (*C++ function*), 440
hpx::threads::has_pending_closures (*C++ function*), 423
hpx::threads::hpx_hwloc_bitmap_wrapper (*C++ class*), 440
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper (*C++ function*), 440
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper::default_mode (*C++ enumerator*), 354
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper::delay_exit (*C++ enumerator*), 353
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper::do_background_work (*C++ enumerator*), 353
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper::enable_elasticity (*C++ enumerator*), 353
hpx::threads::hpx_hwloc_bitmap_wrapper::HPX_NORESUME_Policies::enable_idle_backoff (*C++ enumerator*), 353
hpx::threads::hpx_hwloc_bitmap_wrapper::hpx_hwloc_bitmap_wrapper::policies::enable_stealing

(C++ enumerator), 353
hpx::threads::policies::enable_stealing_numa
 (C++ enumerator), 353
hpx::threads::policies::fast_idle_mode
 (C++ enumerator), 353
hpx::threads::policies::nothing_special
 (C++ enumerator), 353
hpx::threads::policies::reduce_thread_priority
 (C++ enumerator), 353
hpx::threads::policies::scheduler_mode
 (C++ enum), 353
hpx::threads::policies::steal_after_local
 (C++ enumerator), 353
hpx::threads::policies::steal_high_priority_flag
 (C++ enumerator), 353
hpx::threads::post (C++ function), 436, 437
hpx::threads::post_after (C++ function), 424
hpx::threads::post_at (C++ function), 424
hpx::threads::processing_units_count
 (C++ function), 423
hpx::threads::set_scheduler_mode (C++
 function), 423
hpx::threads::set_thread_description
 (C++ function), 349
hpx::threads::set_thread_interruption_en-
 hpx::threads::set_thread_lco_description
 (C++ function), 349
hpx::threads::set_thread_state (C++ func-
 tion), 347, 348
hpx::threads::staged (C++ enumerator), 887
hpx::threads::suspended (C++ enumerator),
 887
hpx::threads::sync_execute (C++ function),
 436
hpx::threads::sync_execute_after (C++
 function), 424
hpx::threads::sync_execute_at (C++ func-
 tion), 424
hpx::threads::terminated (C++ enumerator),
 887
hpx::threads::then_execute (C++ function),
 436
hpx::threads::thread_id (C++ class), 890
hpx::threads::thread_id::get (C++ func-
 tion), 890
hpx::threads::thread_id::operator bool
 (C++ function), 890
hpx::threads::thread_id::operator=(C++
 function), 890
hpx::threads::thread_id::reset (C++ func-
 tion), 890
hpx::threads::thread_id::thrd_ (C++ mem-
 ber), 891
 hpx::threads::thread_id (C++
 function), 890
 hpx::threads::thread_id::repr
 (C++ type), 891
 hpx::threads::thread_id_type (C++ type),
 611
 hpx::threads::thread_pool_base (C++
 function), 269
 hpx::threads::thread_pool_base::resume_direct
 (C++ function), 269
 hpx::threads::thread_pool_base::resume_processing_u-
 (C++ function), 269
 hpx::threads::thread_pool_base::suspend_direct
 (C++ function), 269
 hpx::threads::thread_pool_base::suspend_processing_u-
 (C++ function), 269
 hpx::threads::thread_pool_init_parameters
 (C++ class), 269
 hpx::threads::thread_pool_init_parameters::affinity
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::index_
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::max_back-
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::max_busy
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::max_idle
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::mode_
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::name_
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::network
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::notifier
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::num_thre-
 (C++ member), 270
 hpx::threads::thread_pool_init_parameters::thread_o
 (C++ function), 270
 hpx::threads::thread_priority (C++ enum),
 887
 hpx::threads::thread_priority_boost
 (C++ enumerator), 888
 hpx::threads::thread_priority_bound
 (C++ enumerator), 888
 hpx::threads::thread_priority_default
 (C++ enumerator), 887
 hpx::threads::thread_priority_high (C++
 enumerator), 888
 hpx::threads::thread_priority_high_recursive
 (C++ enumerator), 888

hpx::threads::thread_priority_low (*C++ enumerator*), 887
 hpx::threads::thread_priority_normal (*C++ enumerator*), 887
 hpx::threads::thread_priority_unknown (*C++ enumerator*), 887
 hpx::threads::thread_schedule_hint (*C++ class*), 889
 hpx::threads::thread_schedule_hint::hint (*C++ member*), 890
 hpx::threads::thread_schedule_hint::mode (*C++ member*), 890
 hpx::threads::thread_schedule_hint::thread_sche*C++ function*, 570
 hpx::threads::thread_schedule_hint_mode (*C++ enum*), 889
 hpx::threads::thread_schedule_hint_mode_none (*C++ enumerator*), 889
 hpx::threads::thread_schedule_hint_mode_numa (*C++ enumerator*), 889
 hpx::threads::thread_schedule_hint_mode_thread (*C++ function*), 572
 hpx::threads::thread_self (*C++ type*), 611
 hpx::threads::thread_stacksize (*C++ enum*), 888
 hpx::threads::thread_stacksize_current (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_default (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_huge (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_large (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_maximal (*C++ enumerator*), 889
 hpx::threads::thread_stacksize_medium (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_minimal (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_nostack (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_small (*C++ enumerator*), 888
 hpx::threads::thread_stacksize_unknown (*C++ enumerator*), 888
 hpx::threads::thread_state_enum (*C++ enum*), 887
 hpx::threads::thread_state_ex_enum (*C++ enum*), 888
 hpx::threads::threadmanager (*C++ class*), 570
 hpx::threads::threadmanager::~threadmanager
 hpx::threads::threadmanager::abort_all_shpenddeadl*deadl*
 hpx::threads::threadmanager::create_pools (*C++ function*), 570
 hpx::threads::threadmanager::default_pool (*C++ function*), 570
 hpx::threads::threadmanager::default_scheduler (*C++ function*), 570
 hpx::threads::threadmanager::deinit_tss (*C++ function*), 572
 hpx::threads::threadmanager::enumerate_threads (*C++ function*), 572
 hpx::threads::threadmanager::expand_pool (*C++ function*), 572
 hpx::threads::threadmanager::get_background_thread (*C++ function*), 572
 hpx::threads::threadmanager::get_cumulative_duration (*C++ function*), 572
 hpx::threads::threadmanager::get_os_thread_count (*C++ function*), 572
 hpx::threads::threadmanager::get_os_thread_handle (*C++ function*), 572
 hpx::threads::threadmanager::get_pool (*C++ function*), 570
 hpx::threads::threadmanager::get_pool_numa_bitmap (*C++ function*), 572
 hpx::threads::threadmanager::get_queue_length (*C++ function*), 572
 hpx::threads::threadmanager::get_thread_count (*C++ function*), 572
 hpx::threads::threadmanager::get_thread_count_active (*C++ function*), 573
 hpx::threads::threadmanager::get_thread_count_pending (*C++ function*), 573
 hpx::threads::threadmanager::get_thread_count_staged (*C++ function*), 573
 hpx::threads::threadmanager::get_thread_count_suspended (*C++ function*), 573
 hpx::threads::threadmanager::get_thread_count_terminated (*C++ function*), 573
 hpx::threads::threadmanager::get_thread_count_unknown (*C++ function*), 573
 hpx::threads::threadmanager::get_used_processing_units (*C++ function*), 572
 hpx::threads::threadmanager::init (*C++ function*), 570
 hpx::threads::threadmanager::init_tss (*C++ function*), 572
 hpx::threads::threadmanager::mtx_ (*C++ function*)

member), 573
hpx::threads::threadmanager::mutex_type (C++ type), 573
hpx::threads::threadmanager::network_background (C++ member), 573
hpx::threads::threadmanager::notification_polid (C++ type), 570
hpx::threads::threadmanager::notifier_ (C++ member), 573
hpx::threads::threadmanager::num_threads_ (C++ member), 573
hpx::threads::threadmanager::pool_type (C++ type), 570
hpx::threads::threadmanager::pool_vector (C++ type), 570
hpx::threads::threadmanager::pools_ (C++ member), 573
hpx::threads::threadmanager::print_pools (C++ function), 570
hpx::threads::threadmanager::register_thread (C++ function), 571
hpx::threads::threadmanager::register_work (C++ function), 570
hpx::threads::threadmanager::remove_scheduler_m (C++ function), 572
hpx::threads::threadmanager::report_error (C++ function), 572
hpx::threads::threadmanager::reset_thread_distr (C++ function), 572
hpx::threads::threadmanager::resume (C++ function), 571
hpx::threads::threadmanager::run (C++ function), 571
hpx::threads::threadmanager::scheduler_type (C++ type), 570
hpx::threads::threadmanager::set_scheduler_mode (C++ function), 572
hpx::threads::threadmanager::shrink_pool (C++ function), 572
hpx::threads::threadmanager::status (C++ function), 571
hpx::threads::threadmanager::stop (C++ function), 571
hpx::threads::threadmanager::suspend (C++ function), 571
hpx::threads::threadmanager::threadmanager (C++ function), 570
hpx::threads::threadmanager::threads_lookup_ (C++ member), 573
hpx::threads::topology (C++ class), 441
hpx::threads::topology::~topology (C++ function), 441
hpx::threads::topology::allocate (C++ function), 443
hpx::threads::topology::allocate_membind (C++ function), 443
hpx::threads::topology::bitmap_to_mask (C++ function), 444
hpx::threads::topology::core_affinity_masks_ (C++ type), 445
hpx::threads::topology::core_numbers_ (C++ member), 445
hpx::threads::topology::core_offset (C++ member), 445
hpx::threads::topology::cpuset_to_nodeset (C++ function), 443
hpx::threads::topology::deallocate (C++ function), 444
hpx::threads::topology::empty_mask (C++ member), 445
hpx::threads::topology::extract_node_count (C++ function), 444
hpx::threads::topology::extract_node_mask (C++ function), 444
hpx::threads::topology::get_area_membind_nodeset (C++ function), 443
hpx::threads::topology::get_core_affinity_mask (C++ function), 442
hpx::threads::topology::get_core_number (C++ function), 443
hpx::threads::topology::get_cpubind_mask (C++ function), 443
hpx::threads::topology::get_machine_affinity_mask (C++ function), 441
hpx::threads::topology::get numa_domain (C++ function), 444
hpx::threads::topology::get numa_node_affinity_mask (C++ function), 441
hpx::threads::topology::get numa_node_number (C++ function), 441
hpx::threads::topology::get number_of_core_pus (C++ function), 443
hpx::threads::topology::get number_of_cores (C++ function), 443
hpx::threads::topology::get number_of_numa_node_cores (C++ function), 443
hpx::threads::topology::get number_of_numa_node_pus (C++ function), 443
hpx::threads::topology::get number_of_numa_nodes (C++ function), 443
hpx::threads::topology::get number_of_pus (C++ function), 443
hpx::threads::topology::get number_of_socket_cores (C++ function), 443
hpx::threads::topology::get number_of_socket_pus (C++ function), 443

```

hpx::threads::topology::get_number_of_sockets  

    (C++ function), 443
hpx::threads::topology::print_affinity_mask  

    (C++ function), 442
hpx::threads::topology::get_pu_number      hpx::threads::topology::print_hwloc  

    (C++ function), 443                      (C++ function), 444
hpx::threads::topology::get_service_affinity_masks  

    (C++ function), 441
hpx::threads::topology::get_socket_affinity_masks  

    (C++ function), 441
hpx::threads::topology::get_socket_number hpx::threads::topology::pu_offset (C++  

    (C++ function), 441                      member), 445
hpx::threads::topology::get_thread_affinity_masks  

    (C++ function), 442
hpx::threads::topology::get_thread_affinity_masks  

    (C++ function), 442
hpx::threads::topology::init_core_affinity_masks  

    (C++ function), 444
hpx::threads::topology::init_core_affinity_masks  

    (C++ function), 444
hpx::threads::topology::init_core_number hpx::threads::topology::socket_numbers_<br>  

    (C++ function), 444                      (C++ member), 445
hpx::threads::topology::init_machine_affinity_masks  

    (C++ function), 444
hpx::threads::topology::init_node_number hpx::threads::topology::topo (C++ mem-  

    (C++ function), 444                      ber), 445
hpx::threads::topology::init_num_of_pus hpx::threads::topology::topo_mtx (C++  

    (C++ function), 444                      member), 445
hpx::threads::topology::init numa_node affinity_masks::topology::topology (C++<br>  

    (C++ function), 444                      function), 441
hpx::threads::topology::init numa_node affinity_masks::from_topology::update_to_log  

    (C++ function), 444
hpx::threads::topology::init numa_node affinity_masks::topology (C++ enum-  

    (C++ function), 444                      erator), 887
hpx::threads::topology::init numa_node affinity_masks::topology (C++ enum-  

    (C++ function), 444                      erator), 887
hpx::threads::topology::init_socket_affinity_mask  

    (C++ function), 444
hpx::threads::topology::init_socket_affinity_mask  

    (C++ function), 444
hpx::threads::topology::init_socket_number hpx::threads::wait_abort (C++ enum-  

    (C++ function), 444                      erator), 888
hpx::threads::topology::init_socket_number hpx::threads::wait_sighaled (C++ enum-  

    (C++ function), 444                      erator), 888
hpx::threads::topology::init_thread_affinity_mask  

    (C++ function), 444
hpx::threads::topology::machine_affinity_mask<br>  

    (C++ member), 445
hpx::threads::topology::mask_to_bitmap hpx::threads::wait_timeout (C++ enum-  

    (C++ function), 444                      erator), 888
hpx::threads::topology::memory_page_size_ hpx::threads::wait_unknown (C++ enum-  

    (C++ member), 445                      erator), 888
hpx::threads::topology::mutex_type (C++<br>  

    type), 444
hpx::threads::topology::num_of_pus_ hpx::traits (C++ type), 354, 394, 419, 421, 547,  

    (C++ member), 445                      612, 620, 867, 870, 902, 903, 919
hpx::threads::topology::traits (C++ class), 547
hpx::threads::topology::default_serialization_access_data<br>
    (C++ class), 547
hpx::threads::topology::tolerate_node_faults (C++ function), 825
hpx::threads::topology::traits (C++ enum), 825
hpx::threads::topology::throws (C++ member), 826
hpx::threads::topology::default_serialization_access_data (C++ class), 547
hpx::threads::topology::traits (C++ type), 354, 394, 419, 421, 547, 612, 620, 867, 870, 902, 903, 919
hpx::threads::topology::default_serialization_access_data::flush (C++ function), 547
hpx::threads::topology::default_serialization_access_data::flush (C++ function), 547
hpx::threads::topology::default_serialization_access_data::is_<br>
    (C++ function), 547

```

(*C++ function*), 547
hpx::traits::default_serialization_access_data (*C++ type*), 420
 (*C++ function*), 547
 hpx::traits::is_one_way_executor_t (*C++ type*), 420
hpx::traits::default_serialization_access_data (*types*), 420
 (*C++ function*), 547
 hpx::traits::is_tuple_like (*C++ class*), 394
hpx::traits::default_serialization_access_data (*types*), 420
 (*C++ function*), 547
 hpx::traits::is_two_way_executor_t (*C++ type*), 420
hpx::traits::default_serialization_access_data (*types*), 420
 (*C++ type*), 547
 hpx::traits::is_inplace_type_of (C++ class), 919
hpx::traits::executor_context_t (*C++ type*), 419
 (C++ hpx::traits::plugin_config_data::call
 type), 419
 (*C++ function*), 920
hpx::traits::executor_execution_category (*C++ type*), 419
 (C++ hpx::traits::projected_iterator
 class), 620
hpx::traits::executor_index_t (*C++ type*), 419
 (C++ hpx::traits::projected_iterator::local_iterator
 type), 619, 620
hpx::traits::executor_parameters_type_t (*C++ type*), 419
 (C++ hpx::traits::projected_iterator::type
 type), 619, 620
hpx::traits::executor_shape_t (*C++ type*), 419
 (C++ hpx::traits::projected_iterator<Iterator,
 typename hpx::util::always_void<typename
 type>::type::proxy_type
 type), 619, 620
hpx::traits::functional (*C++ type*), 867
hpx::traits::functional::lift_zipped_iterators (*C++ class*), 619, 620
 (*C++ function*), 867
 (C++ hpx::traits::projected_iterator<Iterator,
 typename hpx::traits::enable_if<is_segmented_iterator
 type>::type), 619, 620
hpx::traits::functional::lift_zipped_iterator (*C++ type*), 867
 (C++ hpx::traits::range_traits::iterator_type
 type), 870
hpx::traits::functional::lift_zipped_iterator (*C++ class*), 866, 867
 (*C++ function*), 870
 (C++ hpx::traits::range_traits<R, true> (*C++*
 class), 870
hpx::traits::get_function_address (*C++ class*), 902
 (*C++ function*), 902, 903
hpx::traits::get_function_address::call
 (*C++ function*), 902, 903
hpx::traits::get_function_address<R
 (* (Ts...)> (*C++ class*), 903
hpx::traits::get_function_address<R
 (Obj::*) (Ts...) const > (*C++ class*), 902, 903
hpx::traits::get_function_address<R
 (Obj::*) (Ts...)> (*C++ class*), 902, 903
hpx::traits::get_function_annotation
 (*C++ class*), 903
hpx::traits::get_function_annotation::call
 (*C++ function*), 903
 (C++ hpx::traits::segmented_local_iterator_traits::local
 type), 613
hpx::traits::is_bind_expression (*C++ class*), 902
 (*C++ function*), 902
hpx::traits::is_bulk_one_way_executor_t (*C++ type*), 420
 (C++ hpx::traits::segmented_local_iterator_traits::remote
 type), 613
hpx::traits::is_bulk_two_way_executor_t (*C++ type*), 420
 (C++ hpx::traits::serialization_access_data
 class), 547
hpx::traits::is_executor_any_t (*C++ type*), 420
 (C++ hpx::traits::serialization_access_data::flush
 type), 548
hpx::traits::is_executor_parameters_t (*C++ type*), 421
 (C++ hpx::traits::serialization_access_data::init_data
 type), 548

hpx::traits::serialization_access_data::read *function*, 396
(C++ function), 548
hpx::traits::serialization_access_data::resize (C++ function), 396
(C++ function), 548
hpx::traits::serialization_access_data::size (C++ function), 396
(C++ function), 548
hpx::traits::serialization_access_data::write *member*, 397
(C++ function), 548
hpx::trigger_lco_event (C++ function), 310, 311
hpx::unhandled_exception (C++ enumerator), 829
hpx::uninitialized_value (C++ enumerator), 829
hpx::unknown_component_address (C++ enumerator), 828
hpx::unknown_error (C++ enumerator), 830
hpx:: unregister_thread (C++ function), 294
hpx:: unregister_with_basename (C++ function), 300
hpx::util (C++ type), 354, 374, 375, 380, 385, 390, 394–396, 445–447, 480, 492–499, 505, 508–510, 512–514, 516, 518–521, 573, 574, 591, 593, 776, 780, 784, 788, 807–809, 811, 834, 838, 840, 843–846, 848, 849, 851, 853, 854, 856, 857, 859, 861, 862, 864, 865, 867, 869, 883, 893–901, 905–909, 911, 916–920
hpx::util:: _optional_swap (C++ type), 389
hpx::util:: _optional_swap::check_swap (C++ function), 389
hpx::util:: adl_barrier (C++ type), 393
hpx::util:: allocator_deleter (C++ class), 446
hpx::util:: allocator_deleter:: alloc_ (C++ member), 446
hpx::util:: allocator_deleter::operator() (C++ function), 446
hpx::util:: always_void (C++ class), 909
hpx::util:: always_void:: type (C++ type), 909
hpx::util:: any_cast (C++ function), 380, 381
hpx::util:: any_nonser (C++ type), 380
hpx::util:: atomic_count (C++ class), 396
hpx::util:: atomic_count::atomic_count (C++ function), 396
hpx::util:: atomic_count::HPX_NON_COPYABLE (C++ function), 396
hpx::util:: atomic_count::operator long (C++ function), 396
hpx::util:: atomic_count::operator++ (C++ function), 396
hpx::util:: atomic_count::operator+= (C++ function), 396
hpx::util:: atomic_count::operator=(C++ function), 396
hpx::util:: atomic_count::operator-= (C++ function), 396
hpx::util:: atomic_count::operator-- (C++ function), 396
hpx::util:: atomic_count::value_ (C++ function), 396
hpx::util:: attach_debugger (C++ function), 354
hpx::util:: bad_any_cast (C++ class), 381
hpx::util:: bad_any_cast::bad_any_cast (C++ function), 381
hpx::util:: bad_any_cast::from (C++ member), 381
hpx::util:: bad_any_cast::to (C++ member), 381
hpx::util:: bad_any_cast::what (C++ function), 381
hpx::util:: bad_lexical_cast (C++ class), 574
hpx::util:: bad_lexical_cast::~bad_lexical_cast (C++ function), 575
hpx::util:: bad_lexical_cast::bad_lexical_cast (C++ function), 575
hpx::util:: bad_lexical_cast::source (C++ member), 575
hpx::util:: bad_lexical_cast::source_type (C++ function), 575
hpx::util:: bad_lexical_cast::target (C++ member), 575
hpx::util:: bad_lexical_cast::target_type (C++ function), 575
hpx::util:: bad_lexical_cast::what (C++ function), 575
hpx::util:: bad_optional_access (C++ class), 387
hpx::util:: bad_optional_access::bad_optional_access (C++ function), 387
hpx::util:: barrier (C++ class), 776
hpx::util:: barrier:: ~barrier (C++ function), 777
hpx::util:: barrier::cond_ (C++ member), 777
hpx::util:: barrier::mtx_ (C++ member), 777
hpx::util:: barrier::mutex_type (C++ type), 777
hpx::util:: barrier::number_of_threads_ (C++ member), 777
hpx::util:: barrier::total_ (C++ member), 777
hpx::util:: barrier::wait (C++ function), 777
hpx::util:: basic_any:: ~basic_any (C++ function), 381

function), 507
hpx::util::cache::lru_cache::evict (C++ function), 508
hpx::util::cache::lru_cache::get_entry (C++ function), 506
hpx::util::cache::lru_cache::get_statistics (C++ function), 508
hpx::util::cache::lru_cache::holds_key (C++ function), 506
hpx::util::cache::lru_cache::insert (C++ function), 506
hpx::util::cache::lru_cache::insert_nonexist (C++ function), 507
hpx::util::cache::lru_cache::key_type (C++ type), 505
hpx::util::cache::lru_cache::lru_cache (C++ function), 505
hpx::util::cache::lru_cache::map_ (C++ member), 508
hpx::util::cache::lru_cache::map_type (C++ type), 505
hpx::util::cache::lru_cache::max_size_ (C++ member), 508
hpx::util::cache::lru_cache::reserve (C++ function), 506
hpx::util::cache::lru_cache::size (C++ function), 505
hpx::util::cache::lru_cache::size_type (C++ type), 505
hpx::util::cache::lru_cache::statistics_ (C++ member), 508
hpx::util::cache::lru_cache::statistics_type (C++ type), 505
hpx::util::cache::lru_cache::storage_ (C++ member), 508
hpx::util::cache::lru_cache::storage_type (C++ type), 505
hpx::util::cache::lru_cache::touch (C++ function), 508
hpx::util::cache::lru_cache::update (C++ function), 507
hpx::util::cache::lru_cache::update_if (C++ function), 507
hpx::util::cache::lru_cache::update_on_exit (C++ type), 508
hpx::util::cache::policies (C++ type), 519
hpx::util::cache::policies::always (C++ class), 519
hpx::util::cache::policies::always::operator (C++ function), 520
hpx::util::cache::statistics (C++ type), 514, 516, 518
hpx::util::cache::statistics::local_full_hpkmstfilescache::statistics::local_statistics::ev

```
hpx::util::cache::statistics::local_statistics::statistics::no_statistics::get_up-  
    (C++ member), 517  
                                              (C++ function), 519  
hpx::util::cache::statistics::local_statistics::get_and_update::statistics::no_statistics::got_en-  
    (C++ function), 517  
                                              (C++ function), 518  
hpx::util::cache::statistics::local_statistics::get_and_update::statistics::no_statistics::got_hi-  
    (C++ function), 517  
                                              (C++ function), 518  
hpx::util::cache::statistics::local_statistics::get_and_update::statistics::no_statistics::got_in-  
    (C++ function), 517  
                                              (C++ function), 518  
hpx::util::cache::statistics::local_statistics::get_and_update::statistics::no_statistics::got_ma-  
    (C++ function), 517  
                                              (C++ function), 518  
hpx::util::cache::statistics::local_statistics::get_and_update::statistics::no_statistics::update::  
    (C++ function), 517  
                                              (C++ class), 519  
hpx::util::cache::statistics::local_statistics::hit::cache::statistics::no_statistics::update::  
    (C++ function), 517  
                                              (C++ function), 519  
hpx::util::cache::statistics::local_statistics::hit::aligned_data (C++ class),  
    (C++ member), 517  
                                              784  
hpx::util::cache::statistics::local_statistics::hit::aligned_data::cache_aligned_data  
    (C++ function), 517  
                                              (C++ function), 783, 784  
hpx::util::cache::statistics::local_statistics::hit::aligned_data::data_::  
    (C++ member), 517  
                                              (C++ member), 783, 784  
hpx::util::cache::statistics::local_statistics::hit::aligned_data<Data,  
    (C++ function), 517  
                                              NeedsPadding>::cacheline_pad  
hpx::util::cache::statistics::local_statistics::method (C++ member), 784  
    (C++ function), 517  
                                              hpx::util::cache_aligned_data<Data,  
hpx::util::cache::statistics::method::false_type> (C++ class), 783, 784  
    (C++ member), 517  
                                              hpx::util::cache_line_data (C++ type), 784  
hpx::util::cache::statistics::method::method::calculate_fanout (C++ function),  
    (C++ enum), 518  
                                              374  
hpx::util::cache::statistics::method::method_erase::entry::checkpoint (C++ class), 923  
    (C++ enumerator), 518  
                                              hpx::util::checkpoint::~checkpoint (C++  
hpx::util::cache::statistics::method_get_entry::function), 924  
    (C++ enumerator), 518  
                                              hpx::util::checkpoint::begin (C++ func-  
hpx::util::cache::statistics::method_insert_entry::function), 924  
    (C++ enumerator), 518  
                                              hpx::util::checkpoint::checkpoint (C++  
hpx::util::cache::statistics::method_update_entry::function), 924  
    (C++ enumerator), 518  
                                              hpx::util::checkpoint::const_iterator  
hpx::util::cache::statistics::no_statistics (C++ type), 924  
    (C++ class), 518  
                                              hpx::util::checkpoint::data_ (C++ mem-  
hpx::util::cache::statistics::no_statistics::clear::function), 924  
    (C++ function), 518  
                                              hpx::util::checkpoint::end (C++ function),  
hpx::util::cache::statistics::no_statistics::erase_entry_count  
    (C++ function), 519  
                                              hpx::util::checkpoint::operator= (C++  
hpx::util::cache::statistics::no_statistics::get_entry_time::function), 924  
    (C++ function), 519  
                                              hpx::util::checkpoint::serialize (C++  
hpx::util::cache::statistics::no_statistics::get_entry_time::function), 924  
    (C++ function), 518  
                                              hpx::util::checkpoint::size (C++ function),  
hpx::util::cache::statistics::no_statistics::get_entry_time  
    (C++ function), 519  
                                              hpx::util::counter_sanity (C++ enumerator),  
hpx::util::cache::statistics::no_statistics::get_insert_entry_count  
    (C++ function), 518  
                                              hpx::util::counter_test (C++ enumerator),  
hpx::util::cache::statistics::no_statistics::get_insert_entry_time  
    (C++ function), 519  
                                              hpx::util::counter_type (C++ enum), 448  
hpx::util::cache::statistics::no_statistics::get_update::function (C++ type), 883  
    (C++ function), 518  
                                              hpx::util::debug::cxx_type_id (C++ type),
```

884
`hpx::util::debug::cxxabi_demangle_helper`
`hpx::util::demangle_helper` (*C++ class*), 884
`hpx::util::demangle_helper::type` `hpx::util::function<R(Ts...)>::VTable`
`(C++ type)`, 884
`hpx::util::demangle_helper::type` (*C++ function*), 884
`hpx::util::print_type` (*C++ function*), 884
`hpx::util::type_id` (*C++ class*), 884
`hpx::util::type_id::typeid` (*C++ member*), 884
`hpx::util::deferred_call` (*C++ function*), 896
`hpx::util::detected_or` (*C++ type*), 908
`hpx::util::detected_or_t` (*C++ type*), 908
`hpx::util::detected_t` (*C++ type*), 908
`hpx::util::enable_lock_detection` (*C++ function*), 480
`hpx::util::fibhash` (*C++ function*), 521
`hpx::util::force_error_on_lock` (*C++ function*), 480
`hpx::util::format` (*C++ function*), 574
`hpx::util::format_to` (*C++ function*), 574
`hpx::util::from_string` (*C++ function*), 574
`hpx::util::function::function` (*C++ function*), 897
`hpx::util::function::operator=` (*C++ function*), 897
`hpx::util::function::result_type` (*C++ type*), 897
`hpx::util::function_nonser` (*C++ type*), 897
`hpx::util::function_ref::assign` (*C++ function*), 898
`hpx::util::function_ref::function_ref` (*C++ function*), 898
`hpx::util::function_ref::get_function_address` (*C++ function*), 898
`hpx::util::function_ref::get_function_annotation` (*C++ function*), 898
`hpx::util::function_ref::get_function_annotation` (*C++ function*), 898
`hpx::util::function_ref::get_vtable` (*C++ function*), 899
`hpx::util::function_ref::object` (*C++ member*), 899
`hpx::util::function_ref::operator()` (*C++ function*), 898
`hpx::util::function_ref::operator=` (*C++ function*), 898
`hpx::util::function_ref::swap` (*C++ function*), 898
`hpx::util::function_ref<R(Ts...)>` (*C++ class*), 898
`hpx::util::function_ref<R(Ts...)>::vptr` `hpx::util::hardware::unbounded_shifter<T>` (*C++ class*), 807
`(C++ member)`, 899
`(C++ type)`, 899
`Serializable` (*C++ class*), 897
`Serializable::base_type` (*C++ type*), 898
`hpx::util::functional` (*C++ type*), 357, 901
`hpx::util::functional::invoke` (*C++ class*), 901
`hpx::util::functional::invoke_r` (*C++ class*), 901
`hpx::util::functional::unwrap` (*C++ class*), 270
`hpx::util::functional::unwrap_all` (*C++ class*), 270
`hpx::util::functional::unwrap_n` (*C++ class*), 270
`hpx::util::get_and_reset_value` (*C++ function*), 374
`hpx::util::hardware` (*C++ type*), 807, 808
`hpx::util::hardware::cpuid` (*C++ function*), 808
`hpx::util::hardware::cpuid_register` (*C++ class*), 808
`hpx::util::hardware::cpuidex` (*C++ function*), 808
`hpx::util::hardware::eax` (*C++ enumerator*), 808
`hpx::util::hardware::ebx` (*C++ enumerator*), 808
`hpx::util::hardware::ecx` (*C++ enumerator*), 808
`hpx::util::hardware::edx` (*C++ enumerator*), 808
`hpx::util::hardware::get_bit_range` (*C++ function*), 807
`hpx::util::hardware::has_bit_set` (*C++ function*), 807
`hpx::util::hardware::info` (*C++ enum*), 808
`hpx::util::hardware::pack_bits` (*C++ function*), 807
`hpx::util::hardware::timestamp` (*C++ function*), 809
`hpx::util::hardware::unbounded_shifter` (*C++ class*), 807
`hpx::util::hardware::unbounded_shifter::shl` (*C++ function*), 807
`hpx::util::hardware::unbounded_shifter::shr` (*C++ function*), 807
`hpx::util::hardware::unbounded_shifter<0, T>` (*C++ class*), 807
`hpx::util::hardware::unbounded_shl` (*C++*

tion), 781
 hpx::util::itt::event_tick (*C++ function*), 780
 hpx::util::itt::frame_context (*C++ class*), 781
 hpx::util::itt::frame_context::~frame_context (*C++ function*), 781
 hpx::util::itt::frame_context::frame_context (*C++ function*), 781
 hpx::util::itt::heap_allocate (*C++ class*), 781
 hpx::util::itt::heap_allocate::~heap_allocate (*C++ function*), 781
 hpx::util::itt::heap_allocate::heap_allocate (*C++ function*), 781
 hpx::util::itt::heap_free (*C++ class*), 781
 hpx::util::itt::heap_free::~heap_free (*C++ function*), 781
 hpx::util::itt::heap_free::heap_free (*C++ function*), 781
 hpx::util::itt::heap_function (*C++ class*), 781
 hpx::util::itt::heap_function::~heap_function (*C++ function*), 781
 hpx::util::itt::heap_function::heap_function (*C++ function*), 781
 hpx::util::itt::heap_internal_access (*C++ class*), 781
 hpx::util::itt::heap_internal_access::~heap_internal_access (*C++ function*), 781
 hpx::util::itt::heap_internal_access::heap_internal_access (*C++ function*), 781
 hpx::util::itt::id (*C++ class*), 781
 hpx::util::itt::id::~id (*C++ function*), 782
 hpx::util::itt::id::id (*C++ function*), 782
 hpx::util::itt::mark_context (*C++ class*), 782
 hpx::util::itt::mark_context::~mark_context (*C++ function*), 782
 hpx::util::itt::mark_context::mark_context (*C++ function*), 782
 hpx::util::itt::mark_event (*C++ class*), 782
 hpx::util::itt::mark_event::~mark_event (*C++ function*), 782
 hpx::util::itt::mark_event::mark_event (*C++ function*), 782
 hpx::util::itt::stack_context (*C++ class*), 782
 hpx::util::itt::stack_context::~stack_context (*C++ function*), 782
 hpx::util::itt::stack_context::stack_context (*C++ function*), 782
 hpx::util::string_handle (*C++ class*), 782
 hpx::util::itt::string_handle::string_handle (*C++ function*), 782
 hpx::util::itt::task (*C++ class*), 782
 hpx::util::itt::task::~task (*C++ function*), 782
 hpx::util::itt::thread_domain (*C++ class*), 782
 hpx::util::itt::undo_frame_context (*C++ class*), 783
 hpx::util::itt::undo_mark_context (*C++ class*), 783
 hpx::util::itt::undo_mark_context::~undo_mark_context (*C++ function*), 783
 hpx::util::itt::undo_mark_context::undo_mark_context (*C++ function*), 783
 hpx::util::jenkins_hash (*C++ class*), 520
 hpx::util::jenkins_hash::~jenkins_hash (*C++ function*), 520
 hpx::util::jenkins_hash::hash (*C++ function*), 521
 hpx::util::jenkins_hash::jenkins_hash (*C++ function*), 520
 hpx::util::jenkins_hash::operator () (*C++ function*), 520
 hpx::util::jenkins_hash::reset (*C++ function*), 520
 hpx::util::jenkins_hash::seed_ (*C++ member*), 521
 hpx::util::jenkins_hash::set_seed (*C++ function*), 520
 hpx::util::jenkins_hash::size_type (*C++ type*), 520
 hpx::util::jenkins_hash::swap (*C++ function*), 521
 hpx::util::lazy_enable_if::type (*C++ type*), 905
 hpx::util::lazy_enable_if<true, T> (*C++ class*), 905
 hpx::util::logging (*C++ type*), 834, 838, 840, 843–846, 848, 849, 851, 853, 854, 856, 857, 859, 861
 hpx::util::logging::array (*C++ type*), 844
 hpx::util::logging::array::ptr_holder (*C++ class*), 844

```
hpx::util::logging::array::ptr_holder::appendutil::logging::destination::file::operator==  
    (C++ function), 844  
    (C++ function), 858  
hpx::util::logging::array::ptr_holder::ahpxy_typel::logging::destination::file_settings  
    (C++ type), 844  
    (C++ class), 858  
hpx::util::logging::array::ptr_holder::dhp::util::logging::destination::file_settings::do_...  
    (C++ function), 844  
    (C++ member), 858  
hpx::util::logging::array::ptr_holder::gahpxptml::logging::destination::file_settings::ext_...  
    (C++ function), 844  
    (C++ member), 858  
hpx::util::logging::array::ptr_holder::mhpxraytil::logging::destination::file_settings::fil_...  
    (C++ member), 845  
    (C++ function), 858  
hpx::util::logging::array::ptr_holder::phpxtypeil::logging::destination::file_settings::fla_...  
    (C++ type), 844  
    (C++ type), 858  
hpx::util::logging::array::ptr_holder::vhpxme_typel::logging::destination::file_settings::flu_...  
    (C++ type), 844  
    (C++ member), 858  
hpx::util::logging::destination    (C++ hpx::util::logging::destination::file_settings::in_...  
    type), 857, 859, 861  
    (C++ member), 858  
hpx::util::logging::destination::cerr    hpx::util::logging::destination::named  
    (C++ class), 861  
    (C++ class), 859  
hpx::util::logging::destination::cerr::oppxatotil::logging::destination::named::add  
    (C++ function), 861  
    (C++ function), 860  
hpx::util::logging::destination::cerr::oppxatotil::logging::destination::named::configure_...  
    (C++ function), 861  
    (C++ function), 860  
hpx::util::logging::destination::converthpx::util::logging::destination::named::del  
    (C++ type), 859  
    (C++ function), 860  
hpx::util::logging::destination::converthpdo_utnverflogging::destination::named::named  
    (C++ function), 859  
    (C++ function), 860  
hpx::util::logging::destination::converthpwriteil::logging::destination::named::non_const_...  
    (C++ function), 859  
    (C++ type), 860  
hpx::util::logging::destination::cout    hpx::util::logging::destination::named::operator()  
    (C++ class), 861  
    (C++ function), 860  
hpx::util::logging::destination::cout::oppxatotil::logging::destination::named::operator==  
    (C++ function), 861  
    (C++ function), 860  
hpx::util::logging::destination::cout::oppxatotil::logging::destination::named::string  
    (C++ function), 861  
    (C++ function), 860  
hpx::util::logging::destination::dbg_winhpx::util::logging::destination::stream  
    (C++ class), 861  
    (C++ class), 861  
hpx::util::logging::destination::dbg_winhpx::oppxatotil::logging::destination::stream::clear  
    (C++ function), 861  
    (C++ function), 862  
hpx::util::logging::destination::dbg_winhpx::oppxatotil::logging::destination::stream::non_const_...  
    (C++ function), 861  
    (C++ type), 862  
hpx::util::logging::destination::file    hpx::util::logging::destination::stream::operator()  
    (C++ class), 857  
    (C++ function), 862  
hpx::util::logging::destination::file::chpxfigutel::logging::destination::stream::operator==  
    (C++ function), 858  
    (C++ function), 862  
hpx::util::logging::destination::file::fhp::util::logging::destination::stream::set_stream  
    (C++ function), 858  
    (C++ function), 862  
hpx::util::logging::destination::file::mhpx::util::logging::destination::stream::stream  
    (C++ member), 858  
    (C++ function), 862  
hpx::util::logging::destination::file::mhpxtypeil::logging::destination::stream::stream_type  
    (C++ type), 858  
    (C++ type), 862  
hpx::util::logging::destination::file::nbpxconstlcombgxngasformat_and_write  
    (C++ type), 858  
    (C++ type), 834  
hpx::util::logging::destination::file::oppxatotil::logging::format_and_write::simple  
    (C++ function), 858  
    (C++ class), 835
```


hpx::util::logging::formatter::spacer_t (C++ type), 847
(C++ class), 851
hpx::util::logging::formatter::spacer_t::m_pref(C++ function), 847
(C++ member), 852, 853
hpx::util::logging::formatter::spacer_t::m_suff(C++ function), 847
(C++ member), 852, 853
hpx::util::logging::formatter::spacer_t::operator==(C++ function), 847
(C++ member), 852, 853
hpx::util::logging::formatter::spacer_t::operator=(C++ function), 847
(C++ member), 852
hpx::util::logging::formatter::spacer_t::param (C++ function), 847
(C++ type), 852
hpx::util::logging::formatter::spacer_t::parse_format
(C++ function), 852, 853
hpx::util::logging::formatter::spacer_t::msg_route (C++ type),
hpx::util::logging::formatter::spacer_t::parse_format
(C++ function), 852, 853
hpx::util::logging::formatter::spacer_t::msg_route::formatter_and_destination
(C++ type), 852
hpx::util::logging::formatter::spacer_t::msg_route::formatter_and_destination
(C++ function), 852, 853
hpx::util::logging::formatter::spacer_t<convert(C++ function), 836
original_formatter, true> (C++ class), 852
hpx::util::logging::formatter::spacer_t::msg_route::formatter_and_destination
(C++ function), 836
hpx::util::logging::formatter::thread_id hpx::util::logging::msg_route::formatter_and_destination
(C++ type), 853
hpx::util::logging::formatter::thread_id hpx::util::logging::msg_route::formatter_and_destination
(C++ class), 853
hpx::util::logging::formatter::thread_id hpx::convrtlogging::msg_route::simple
(C++ type), 854
hpx::util::logging::formatter::thread_id hpx::operatorlogging::msg_route::simple::append_destination
(C++ function), 854
hpx::util::logging::formatter::thread_id hpx::operatorlogging::msg_route::simple::append_formatter
(C++ function), 854
hpx::util::logging::formatter::time hpx::util::logging::msg_route::simple::d_array
(C++ type), 847
hpx::util::logging::formatter::time_strf hpx::util::logging::msg_route::simple::del_destination
(C++ type), 856
hpx::util::logging::formatter::time_strf hpx::util::logging::msg_route::simple::del_formatter
(C++ class), 856
hpx::util::logging::formatter::time_strf hpx::convrtlogging::msg_route::simple::destination
(C++ type), 857
hpx::util::logging::formatter::time_strf hpx::operatorlogging::msg_route::simple::f_array
(C++ member), 857
hpx::util::logging::formatter::time_strf hpx::operatorlogging::msg_route::simple::formatter_pt
(C++ member), 857
hpx::util::logging::formatter::time_strf hpx::operatorlogging::msg_route::simple::m_to_write
(C++ function), 857
hpx::util::logging::formatter::time_strf hpx::operatorlogging::msg_route::simple::simple
(C++ function), 857
hpx::util::logging::formatter::time_strf hpx::operatorlogging::msg_route::simple::write
(C++ function), 857
hpx::util::logging::formatter::time_t hpx::util::logging::msg_route::simple::write_info
(C++ class), 847
hpx::util::logging::formatter::time_t::cbpfigurel::logging::msg_route::simple::write_info:
(C++ function), 847
hpx::util::logging::formatter::time_t::cbpxertifypelogging::msg_route::simple::write_info:
(C++ type), 847
hpx::util::logging::formatter::time_t::nbpxconstlcombeginbasep_equal (C++ type),

843 hpx::util::logging::op_equal::operator== (C++ function), 843
hpx::util::logging::op_equal::same_type_op_equal(C++ function), 841
(C++ class), 843 hpx::util::logging::optimize::cache_string_one_string
hpx::util::logging::op_equal::same_type_op_equal(C++ function), 841
(C++ function), 843 hpx::util::logging::optimize::cache_string_one_string
hpx::util::logging::op_equal::same_type_op_equal(C++ function), 842
(C++ class), 843 hpx::util::logging::optimize::operator<<
hpx::util::logging::op_equal::same_type_op_equal(C++ function), 841
(C++ class), 843 hpx::util::logging::writer(C++ type), 838
hpx::util::logging::op_equal::same_type_bpxequal_lt_to_poggisage_typeeop: named_type
(C++ function), 844 (C++ class), 838
hpx::util::logging::op_equal::same_type_bpxequal_lt_to_poggisnglswriter::named_write::add_destination
(C++ function), 843 (C++ function), 839
hpx::util::logging::op_equal::same_type_bpxequal_lt_to_poggisngt:typeeop: eqamedtowrite::add_format
(C++ function), 844 (C++ function), 839
hpx::util::logging::optimize (C++ type), hpx::util::logging::writer::named_write::destination
840 (C++ function), 838, 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logging::writer::named_write::format
(C++ class), 841 (C++ function), 838, 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logghagstwingeonnamed_write::m_destination
(C++ function), 842 (C++ function), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logghagstwingeonnamed_write::m_destination
(C++ function), 841 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggdingstrwngter::named_write::m_destination
(C++ function), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggwingwriter::named_write::m_format_size
(C++ function), 841, 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logggingwriter::named_write::m_format_size
(C++ function), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggging:msgiter::named_write::m_format_k
(C++ member), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggging:mwgitempushnamed_write::m_format_k
(C++ member), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggging:swzeter::named_write::m_format_s
(C++ member), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logggesgrverappendnamed_write::m_writer
(C++ member), 842 (C++ member), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:logggesgrverprependnamed_write::named_write
(C++ member), 842 (C++ function), 838
hpx::util::logging::optimize::cache_strihgxoneatstr:loggging::writer::named_write::operator()
(C++ member), 842 (C++ function), 839
hpx::util::logging::optimize::cache_strihgxoneatstr:loggratorwriter::named_write::parse_destination
const std::string& (C++ function), (C++ class), 839
842 hpx::util::logging::writer::named_write::parse_destination
hpx::util::logging::optimize::cache_string_one_(C++ function), 840 string
(C++ function), 842 hpx::util::logging::writer::named_write::parse_destination
hpx::util::logging::optimize::cache_string_one_(C++ function), 840 append
(C++ function), 841 hpx::util::logging::writer::named_write::parse_destination
hpx::util::logging::optimize::cache_string_one_(C++ function), 840 prepend
(C++ function), 841 hpx::util::logging::writer::named_write::parse_destination
hpx::util::logging::optimize::cache_string_one_(C++ function), 840 string
(C++ function), 842 hpx::util::logging::writer::named_write::parse_destination

(C++ member), 840
hpx::util::logging::writer::named_write::hpx::parse_format&opt (C++ member), 387
(C++ class), 840 hpx::util::nullopt_t (C++ class), 387
hpx::util::logging::writer::named_write::hpx::parse_format&opt::add::nullopt_t (C++
(C++ function), 840 function), 387
hpx::util::logging::writer::named_write::hpx::parse_format&shot (C++ function), 896
(C++ function), 840 hpx::util::operator!= (C++ function), 385, 386
hpx::util::logging::writer::named_write::hpx::parse_format&at::get (C++ function), 385, 386
(C++ function), 840 hpx::util::operator> (C++ function), 385–387
hpx::util::logging::writer::named_write::hpx::parse_format&at::has (C++ function), 385, 386
(C++ function), 840 hpx::util::operator>> (C++ function), 380, 920
hpx::util::logging::writer::named_write::hpx::parse_format&atom (C++ function), 385, 386
(C++ member), 840 hpx::util::operator<= (C++ function), 385–387
hpx::util::logging::writer::named_write::hpx::piat&despenat&< (C++ function), 380, 920
(C++ function), 839 hpx::util::optional (C++ class), 387
hpx::util::logging::writer::named_write::hpx::piat&elfomat&::~optional (C++ func-
(C++ function), 839 tion), 388
hpx::util::logging::writer::named_write::hpx::itat&ndlconfig&::emplace (C++ func-
(C++ function), 839 tion), 388
hpx::util::logging::writer::named_write::hpx::itat::optional::empty_ (C++ member),
(C++ function), 839 389
hpx::util::make_any_nonsr (C++ function), hpx::util::optional::has_value (C++ func-
381 tion), 388
hpx::util::make_iterator_range (C++ func- hpx::util::optional::operator bool (C++
tion), 866 function), 388
hpx::util::make_optional (C++ function), 387 hpx::util::optional::operator* (C++ func-
hpx::util::make_streamable_any_nonsr
(C++ function), 381 tion), 388
hpx::util::make_streamable_unique_any_nonsr
(C++ function), 381 hpx::util::optional::operator= (C++ func-
tion), 388
hpx::util::make_transform_iterator (C++
function), 864 hpx::util::optional::operator-> (C++
function), 388
hpx::util::make_unique_any_nonsr (C++
function), 381 hpx::util::optional::optional (C++ func-
tion), 388
hpx::util::manage_config (C++ class), 375 hpx::util::optional::reset (C++ function),
hpx::util::manage_config::add (C++ func- 389
tion), 376 hpx::util::optional::storage_ (C++ mem-
ber), 389
hpx::util::manage_config::config_ (C++
member), 376 hpx::util::optional::swap (C++ function),
388 hpx::util::optional::value (C++ function),
388
hpx::util::manage_config::get_value
(C++ function), 376 hpx::util::optional::value_or (C++ func-
tion), 388
hpx::util::manage_config::manage_config
(C++ function), 376 hpx::util::optional<T>::value_type (C++
type), 388
hpx::util::manage_config::map_type (C++
type), 375 hpx::util::pack (C++ class), 907
hpx::util::may_attach_debugger (C++ func- hpx::util::pack::size (C++ member), 907
tion), 354 hpx::util::pack::type (C++ type), 907
hpx::util::mem_fn (C++ function), 901 hpx::util::pack_c (C++ class), 907
hpx::util::nonesuch (C++ class), 908 hpx::util::pack_c::size (C++ member), 907
hpx::util::nonesuch::~nonesuch (C++ func- hpx::util::pack_c::type (C++ type), 907
tion), 909 hpx::util::placeholders (C++ type), 899
hpx::util::nonesuch::nonesuch (C++ func- hpx::util::plugin (C++ type), 916–919
tion), 909 hpx::util::plugin::dll_handle (C++ type),
hpx::util::nonesuch::operator= (C++ func-

919
 hpx::util::plugin::exported_plugins_typehpx::util::reinitializable_static::HPX_NON_COPYABLE
(C++ type), 919
 hpx::util::plugin::plugin_factory (*C++ hpx::util::reinitializable_static::operator class*), 916
 hpx::util::plugin::plugin_factory::pluginfactory::reinitializable_static::operator reference (*C++ function*), 810
 hpx::util::plugin::plugin_factory<BasePlugin>::reinitializable_static::pointer
(C++ type), 917
 hpx::util::plugin::plugin_wrapper (*C++ hpx::util::reinitializable_static::reference class*), 918
 hpx::util::plugin::plugin_wrapper::pluginwrapper::reinitializable_static::reinitializable_<*C++ function*>, 918
 hpx::util::plugin::shared_ptr (*C++ type*), hpx::util::reinitializable_static::storage_type
 918
 hpx::util::plugin::static_plugin_factoryhpx::util::reinitializable_static::value_construct
(C++ class), 917
 hpx::util::plugin::static_plugin_factoryhpx::util::reinitializable_static::value_construct
(C++ function), 917
 hpx::util::plugin::static_plugin_factoryhpx::util::reinitializable_static::value_type
(C++ type), 917
 hpx::util::plugin::virtual_constructor hpx::util::report_errors (*C++ function*), 448
(C++ class), 919
 hpx::util::plugin::virtual_constructor<BasePlugin>::reset_ignored_all (*C++ function*), 480
 hpx::util::print_cdash_timing (*C++ function*), 448
 hpx::util::range_adl (*C++ type*), 865
 hpx::util::register_lock (*C++ function*), 480
 hpx::util::reinit_construct (*C++ function*),
 811
 hpx::util::reinit_destruct (*C++ function*),
 811
 hpx::util::reinit_register (*C++ function*),
 811
 hpx::util::reinitializable_static (*C++ class*), 810
 hpx::util::reinitializable_static (*C++ member*), 810
 hpx::util::reinitializable_static::const_refere
(C++ type), 810
 hpx::util::reinitializable_static::constructed_494
(C++ member), 811
 hpx::util::reinitializable_static::data_hpx::util::seedenumerator, 520
(C++ member), 811
 hpx::util::reinitializable_static::default_cons
(C++ function), 810
 hpx::util::reinitializable_static::default_cons
(C++ function), 810
 hpx::util::reinitializable_static::destrhpx::util::spinlock (~spinlock (*C++ func-
tion*)), 776
 hpx::util::reinitializable_static::get hpx::util::spinlock::HPX_NON_COPYABLE
(C++ function), 810
 hpx::util::reinitializable_static::get_a
(C++ function), 776
 hpx::util::reinitializable_static::get_a
(C++ function), 776

776
hpx::util::spinlock::m (*C++ member*), 776
hpx::util::spinlock::native_handle (*C++ function*), 776
hpx::util::spinlock::native_handle_type (*C++ type*), 776
hpx::util::spinlock::spinlock (*C++ function*), 776
hpx::util::spinlock::try_lock (*C++ function*), 776
hpx::util::spinlock::unlock (*C++ function*), 776
hpx::util::spinlock_pool (*C++ class*), 788
hpx::util::spinlock_pool::pool_ (*C++ member*), 788
hpx::util::spinlock_pool::scoped_lock (*C++ class*), 788
hpx::util::spinlock_pool::spinlock_for (*C++ function*), 788
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::~scoped_lock
 (*C++ function*), 788
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::HPX_NON_COPYABLE
 (*C++ function*), 788
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::lock (*C++ function*), 788
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::scoped_lock
 (*C++ function*), 788
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::sp_ (*C++ member*), 789
hpx::util::spinlock_pool<Tag,
 N>::scoped_lock::unlock (*C++ function*), 788
hpx::util::static_ (*C++ class*), 909
hpx::util::static_::const_reference
 (*C++ type*), 909
hpx::util::static_::constructed_ (*C++ member*), 910
hpx::util::static_::data_ (*C++ member*), 910
hpx::util::static_::default_constructor
 (*C++ class*), 910
hpx::util::static_::destructor
 (*C++ class*), 910
hpx::util::static_::get (*C++ function*), 910
hpx::util::static_::get_address
 (*C++ function*), 910
hpx::util::static_::HPX_NON_COPYABLE
 (*C++ function*), 910
hpx::util::static_::operator
 const_reference (*C++ function*), 910
hpx::util::static_::operator reference
 (*C++ function*), 910
hpx::util::static_::pointer (*C++ type*), 910
hpx::util::static_::reference (*C++ type*),
 909
hpx::util::static_::static_ (*C++ function*),
 910
hpx::util::static_::storage_type (*C++ type*), 910
hpx::util::static_::value_type (*C++ type*),
 909
hpx::util::static_<T,
 Tag>::default_constructor::construct
 (*C++ function*), 910
hpx::util::static_<T,
 Tag>::destructor::~destructor
 (*C++ function*), 910
hpx::util::steady_duration (*C++ class*), 492
hpx::util::steady_duration::_rel_time
 (*C++ member*), 493
hpx::util::steady_duration::from_now
 (*C++ function*), 493
hpx::util::steady_duration::steady_duration
 (*C++ function*), 493
hpx::util::steady_duration::value
 (*C++ function*), 493
hpx::util::steady_duration::value_type
 (*C++ type*), 493
hpx::util::steady_time_point (*C++ class*),
 493
hpx::util::steady_time_point::_abs_time
 (*C++ member*), 493
hpx::util::steady_time_point::steady_time_point
 (*C++ function*), 493
hpx::util::steady_time_point::value
 (*C++ function*), 493
hpx::util::steady_time_point::value_type
 (*C++ type*), 493
hpx::util::streamable_any_nonser
 (*C++ type*), 380
hpx::util::streamable_unique_any_nonser
 (*C++ type*), 380
hpx::util::streamable_unique_wany_nonser
 (*C++ type*), 380
hpx::util::streamable_wany_nonser
 (*C++ type*), 380
hpx::util::swap (*C++ function*), 380, 387
hpx::util::test_failure_handler_type
 (*C++ type*), 448
hpx::util::thread_specific_ptr
 (*C++ class*), 394
hpx::util::thread_specific_ptr::element_type
 (*C++ type*), 395

```

hpx::util::thread_specific_ptr::get           (C++ type), 390, 393
    (C++ function), 395
hpx::util::thread_specific_ptr::operator*     tuple<Ts...>> (C++ class), 393
    (C++ function), 395
hpx::util::thread_specific_ptr::operator->   tuple<Ts...>>::type (C++ type),
    (C++ function), 395
hpx::util::thread_specific_ptr::ptr_          hpx::util::tuple_element<I,
    (C++ member), 395
hpx::util::thread_specific_ptr::reset         hpx::util::tuple_size (C++ class), 393
    (C++ function), 395
hpx::util::to_string (C++ function), 574
hpx::util::transform_iterator (C++ class), 864
hpx::util::transform_iterator::base_type       hpx::util::unique_any_nonser (C++ type),
    (C++ type), 864
hpx::util::transform_iterator::dereference     380
    (C++ function), 865
hpx::util::transform_iterator::transformer     hpx::util::unique_function::operator=
    (C++ function), 864
    (C++ type), 894
hpx::util::transform_iterator::transformer_     hpx::util::unique_function::result_type
    (C++ member), 865
hpx::util::traverse_pack_async (C++ function), 354
hpx::util::traverse_pack_async_allocator        hpx::util::unique_function::unique_function
    (C++ function), 355
hpx::util::tuple (C++ class), 391
hpx::util::tuple::impl (C++ member), 392
hpx::util::tuple::tuple (C++ function), 391, 392
hpx::util::tuple_element<0, std::pair<T0, T1>> (C++ class), 389, 392
hpx::util::tuple_element<0, std::pair<T0, T1>>::type (C++ type), 389, 392
hpx::util::tuple_element<1, std::pair<T0, T1>> (C++ class), 389, 392
hpx::util::tuple_element<1, std::pair<T0, T1>>::type (C++ type), 390, 392
hpx::util::tuple_element<I, boost::array<Type, Size>> (C++ class), 390, 392
hpx::util::tuple_element<I, boost::array<Type, Size>>::type (C++ type), 390, 393
hpx::util::tuple_element<I, std::array<Type, Size>> (C++ class), 390, 393
hpx::util::tuple_element<I, std::array<Type, Size>>::type (C++ type), 390, 393
hpx::util::unlock_guard (C++ class), 395
hpx::util::unlock_guard::~unlock_guard          (C++ function), 395
hpx::util::unlock_guard::HPX_NON_COPYABLE      (C++ function), 395
hpx::util::unlock_guard::m_ (C++ member), 396
hpx::util::unlock_guard::unlock_guard          (C++ function), 395
hpx::util::unlock_guard<Mutex>::mutex_type    (C++ type), 395
hpx::util::unregister_lock (C++ function), 480
hpx::util::unused (C++ member), 908
hpx::util::unused_type (C++ class), 908
hpx::util::unused_type::unused_type            (C++ function), 908
hpx::util::unwrap (C++ function), 356
hpx::util::unwrap_all (C++ function), 357
hpx::util::unwrap_n (C++ function), 356
hpx::util::unwrap_ref (C++ function), 906
hpx::util::unwrap_reference (C++ class), 906
hpx::util::unwrap_reference::type              (C++ type), 905, 906
hpx::util::unwrap_reference<boost::reference_wrapper<T>>::operator= (C++ class), 905, 906
hpx::util::unwrap_reference<boost::reference_wrapper<T>>::operator<=> (C++ class), 905, 906
hpx::util::unwrap_reference<std::reference_wrapper<T>>::operator= (C++ class), 905, 906

```

hpx::util::unwrap_reference<std::reference_wrapper<T>> (C++ class), 905, 906

hpx::util::unwrapping (C++ function), 357

hpx::util::unwrapping_all (C++ function), 357

hpx::util::unwrapping_n (C++ function), 357

hpx::util::verify_no_locks (C++ function), 480

hpx::util::void_guard<void> (C++ class), 911

hpx::util::zip_iterator (C++ class), 867

hpx::util::zip_iterator::base_type (C++ type), 868

hpx::util::zip_iterator::zip_iterator (C++ function), 868

hpx::util::zip_iterator<tuple<Ts...>> (C++ class), 868

hpx::version_too_new (C++ enumerator), 828

hpx::version_too_old (C++ enumerator), 828

hpx::version_unknown (C++ enumerator), 828

hpx::wait_all (C++ function), 319, 320

hpx::wait_all_n (C++ function), 320

hpx::wait_any (C++ function), 322, 323

hpx::wait_any_n (C++ function), 323

hpx::wait_each (C++ function), 330, 331

hpx::wait_each_n (C++ function), 331

hpx::wait_some (C++ function), 325–327

hpx::wait_some_n (C++ function), 327

hpx::when_all (C++ function), 320, 321

hpx::when_all_n (C++ function), 321

hpx::when_any (C++ function), 324

hpx::when_any_n (C++ function), 325

hpx::when_any_result (C++ class), 270

hpx::when_any_result::futures (C++ member), 271

hpx::when_any_result::index (C++ member), 271

hpx::when_each (C++ function), 332, 333

hpx::when_each_n (C++ function), 333

hpx::when_some (C++ function), 328, 329

hpx::when_some_n (C++ function), 330

hpx::when_some_result (C++ class), 271

hpx::when_some_result::futures (C++ member), 271

hpx::when_some_result::indices (C++ member), 271

hpx::yield_aborted (C++ enumerator), 829

HPX_AFFINITY_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

command line option, 84

HPX_AFFINITY_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_AFFINITY_WITH_TESTS:BOOL (C macro), 488

command line option, 84

HPX_AGAS_BOOTSTRAP_PREFIX (C macro), 488

HPX_ALGORTITHMS_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_ALGORITHMS_WITH_TESTS:BOOL (C macro), 488

command line option, 84

HPX_ALLOCATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

command line option, 84

HPX_ALLOCATOR_SUPPORT_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_APPLICATION_STRING (C macro), 487

HPX_ASSERT (C macro), 773

HPX_ASSERT_CURRENT_FUNCTION (C macro), 774

HPX_ASSERT_MSG (C macro), 773

HPX_ASSERT_OWNS_LOCK (C macro), 396

HPX_ASSERTION_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

command line option, 84

HPX_ASSERTION_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_ASSERTION_WITH_TESTS:BOOL (C macro), 488

command line option, 84

HPX_BASIC_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 484

command line option, 84

HPX_BASIC_EXECUTION_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 484

command line option, 84

HPX_BATCH_ENVIRONMENTS_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

command line option, 84

HPX_BATCH_ENVIRONMENTS_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_BATCH_ENVIRONMENTS_WITH_TESTS:BOOL (C macro), 488

command line option, 84

HPX_BUILD_TYPE (C macro), 489

HPX_BUSY_LOOP_COUNT_MAX (C macro), 488

HPX_CACHE_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

command line option, 84

HPX_CACHE_WITH_DEPRECATED_WARNINGS:BOOL (C macro), 488

command line option, 84

HPX_CACHE_WITH_TESTS:BOOL (C macro), 488

command line option, 84

HPX_CAPTURE_FORWARD (C macro), 490

HPX_CAPTURE_MOVE (C macro), 490

HPX_CHECKPOINT_WITH_COMPATIBILITY_HEADERS:BOOL (C macro), 488

```

    command line option, 84
HPX_CHECKPOINT_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 84
HPX_CHECKPOINT_WITH_TESTS:BOOL
    command line option, 85
HPX_CLANG_VERSION (C macro), 488
HPX_COLD (C macro), 492
HPX_COLLECTIVES_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 85
HPX_COLLECTIVES_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_COLLECTIVES_WITH_TESTS:BOOL
    command line option, 85
HPX_COMPILER_FENCE (C macro), 492
HPX_COMPONENT_NAME (C macro), 487
HPX_COMPONENT_STRING (C macro), 487
HPX_COMPUTE_CUDA_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_COMPUTE_CUDA_WITH_TESTS:BOOL
    command line option, 85
HPX_COMPUTE_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_COMPUTE_TESTS:BOOL
    command line option, 85
HPX_CONCEPT_ASSERT (C macro), 548
HPX_CONCEPTQUIRES (C macro), 548
HPX_CONCEPTQUIRES_ (C macro), 548
HPX_CONCEPTS_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 85
HPX_CONCEPTS_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_CONCEPTS_WITH_TESTS:BOOL
    command line option, 85
HPX_CONCURRENCY_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 85
HPX_CONCURRENCY_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_CONCURRENCY_WITH_TESTS:BOOL
    command line option, 85
HPX_CONFIG_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 85
HPX_CONFIG_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_CONFIG_TESTS:BOOL
    command line option, 85
HPX_CONNECTING_IP_PORT (C macro), 485
HPX_CONSTEXPR (C macro), 489
HPX_CONSTEXPR_OR_CONST (C macro), 489
HPX_CONTINUATION_MAX_RECURSION_DEPTH (C macro), 488
HPX_COROUTINE_NUM_HEAPS (C macro), 487
HPX_COROUTINES_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 85
HPX_COROUTINES_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 85
HPX_COROUTINES_WITH_TESTS:BOOL
    command line option, 85
HPX_DEBUG (C macro), 489
HPX_DEBUGGING_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 86
HPX_DEBUGGING_DEPRECATED_WARNINGS:BOOL
    command line option, 86
HPX_DEBUGGING_TESTS:BOOL
    command line option, 86
HPX_DECLARE_PLAIN_ACTION (C macro), 365
HPX_DEFINE_COMPONENT_ACTION (C macro), 364
HPX_DEFINE_ERROR_INFO (C macro), 830
HPX_DEFINE_PLAIN_ACTION (C macro), 365
HPX_DEPRECATED (C macro), 491
HPX_ERRORS_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 86
HPX_ERRORS_WITH_DEPRECATED_WARNINGS:BOOL
    command line option, 86
HPX_ERRORS_TESTS:BOOL
    command line option, 86
HPX_EXECUTABLE_EXTENSION (C macro), 487
HPX_EXECUTION_WITH_COMPATIBILITY_HEADERS:BOOL
    command line option, 86
HPX_EXECUTION_DEPRECATED_WARNINGS:BOOL
    command line option, 86
HPX_EXECUTION_TESTS:BOOL
    command line option, 86
HPX_EXPORT (C macro), 491
HPX_EXPORT_REINITIALIZABLE_STATIC (C macro), 809
HPX_EXPORT_STATIC_ (C macro), 909
HPX_EXPORT_THREAD_SPECIFIC_PTR (C macro), 394
HPX_FALLTHROUGH (C macro), 491
HPX_FILESYSTEM_WITH_BOOST_FILESYSTEM_COMPATIBILITY
    command line option, 86
HPX_FILESYSTEM_DEPRECATED_WARNINGS:BOOL
    command line option, 86
HPX_FILESYSTEM_TESTS:BOOL
    command line option, 86
HPX_FORCEINLINE (C macro), 491
HPX_FORMAT_EXPORT (C macro), 573
HPX_FORMAT_DEPRECATED_HEADERS:BOOL
    command line option, 86
HPX_FORMAT_TESTS:BOOL
    command line option, 86

```

command line option, 86
HPX_FORMAT_WITH_TESTS:BOOL
 command line option, 86
HPX_FUNCTIONAL_WITH_COMPATIBILITY_HEADERS
 command line option, 86
HPX_FUNCTIONAL_WITH_DEPRECATED_WARNINGS
 command line option, 86
HPX_FUNCTIONAL_WITH_TESTS:BOOL
 command line option, 86
HPX_GCC_VERSION (*C macro*), 488
HPX_GLOBALCREDIT_INITIAL (*C macro*), 487
HPX_HARDWARE_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 87
HPX_HARDWARE_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
HPX_HARDWARE_WITH_TESTS:BOOL
 command line option, 87
HPX_HAS_DLOPEN (*C macro*), 918
HPX_HAS_MEMBER_XXX_TRAIT_DEF (*C macro*), 549
HPX_HAS_XXX_TRAIT_DEF (*C macro*), 548
HPX_HASHING_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 87
HPX_HASHING_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
HPX_HASHING_WITH_TESTS:BOOL
 command line option, 87
HPX_HAVE_THREAD_BACKTRACE_DEPTH
 (*C macro*), 487
HPX_HOT (*C macro*), 492
HPX_HUGE_STACK_SIZE (*C macro*), 489
HPX_IDLE_BACKOFF_TIME_MAX (*C macro*), 488
HPX_IDLE_LOOP_COUNT_MAX (*C macro*), 487
HPX_INI_PATH_DELIMITER (*C macro*), 487
HPX_INITIAL_AGAS_MAX_PENDING_REFCNT
 (*C macro*), 487
HPX_INITIAL_GID_RANGE (*C macro*), 488
HPX_INITIAL_IP_ADDRESS (*C macro*), 485
HPX_INITIAL_IP_PORT (*C macro*), 485
HPX_INTEL_VERSION (*C macro*), 488
HPX_INVOKE (*C macro*), 900
HPX_INVOKE_R (*C macro*), 900
HPX_IS_BITWISE_SERIALIZABLE (*C macro*), 547
HPX_ITERATOR_SUPPORT_WITH_COMPATIBILITY_HEADERS
 command line option, 87
HPX_ITERATOR_SUPPORT_WITH_DEPRECATED_WARNINGS
 command line option, 87
HPX_ITERATOR_SUPPORT_WITH_TESTS:BOOL
 command line option, 87
HPX_ITT_COUNTER_CREATE (*C macro*), 778
HPX_ITT_COUNTER_CREATE_TYPED (*C macro*), 778
HPX_ITT_COUNTER_DESTROY (*C macro*), 778
HPX_ITT_COUNTER_SET_VALUE (*C macro*), 778
HPX_ITT_DOMAIN_CREATE (*C macro*), 778
HPX_ITT_FRAME_BEGIN (*C macro*), 777
HPX_ITT_FRAME_END (*C macro*), 777
HPX_ITT_HEAP_ALLOCATE_BEGIN (*C macro*), 778
HPX_ITT_HEAP_ALLOCATE_END (*C macro*), 778
HPX_ITT_HEAP_FREE_BEGIN (*C macro*), 778
HPX_ITT_HEAP_FREE_END (*C macro*), 778
HPX_ITT_HEAP_FUNCTION_CREATE (*C macro*), 778
HPX_ITT_HEAP_INTERNAL_ACCESS_BEGIN
 (*C macro*), 778
HPX_ITT_HEAP_INTERNAL_ACCESS_END
 (*C macro*), 778
HPX_ITT_HEAP_REALLOCATE_BEGIN (*C macro*),
 778
HPX_ITT_HEAP_REALLOCATE_END (*C macro*), 778
HPX_ITT_ID_CREATE (*C macro*), 778
HPX_ITT_ID_DESTROY (*C macro*), 778
HPX_ITT_MAKE_ID (*C macro*), 778
HPX_ITT_MARK (*C macro*), 777
HPX_ITT_MARK_CREATE (*C macro*), 777
HPX_ITT_MARK_OFF (*C macro*), 777
HPX_ITT_METADATA_ADD (*C macro*), 778
HPX_ITT_STACK_CALLEE_ENTER (*C macro*), 777
HPX_ITT_STACK_CALLEE_LEAVE (*C macro*), 777
HPX_ITT_STACK_CREATE (*C macro*), 777
HPX_ITT_STACK_DESTROY (*C macro*), 777
HPX_ITT_STRING_HANDLE_CREATE (*C macro*), 778
HPX_ITT_SYNC_ACQUIRED (*C macro*), 777
HPX_ITT_SYNC_CANCEL (*C macro*), 777
HPX_ITT_SYNC_CREATE (*C macro*), 777
HPX_ITT_SYNC_DESTROY (*C macro*), 777
HPX_ITT_SYNC_PREPARE (*C macro*), 777
HPX_ITT_SYNC_RELEASED (*C macro*), 777
HPX_ITT_SYNC_RELEASENING (*C macro*), 777
HPX_ITT_SYNC_RENAME (*C macro*), 777
HPX_ITT_TASK_BEGIN (*C macro*), 778
HPX_ITT_TASK_BEGIN_ID (*C macro*), 778
HPX_ITT_TASK_END (*C macro*), 778
HPX_ITT_THREAD_IGNORE (*C macro*), 778
HPX_ITT_THREAD_SET_NAME (*C macro*), 777
HPX_LARGE_STACK_SIZE (*C macro*), 489
HPX_LIKELY (*C macro*), 490
HPX_LOCAL_LCOS_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 87
HPX_NGSCADOLCOS_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
HPX_NGSCADOLCOS_WITH_TESTS:BOOL
 command line option, 87
HPX_LOGGING_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 87
HPX_LOGGING_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
HPX_LOGGING_WITH_TESTS:BOOL
 command line option, 87
HPX_MAKE_DLL_STRING (*C macro*), 487
HPX_MANGLE_NAME (*C macro*), 487

HPX_MANGLE_STRING (*C macro*), 487
 HPX_MAX_NETWORK_RETRIES (*C macro*), 487
 HPX_MEDIUM_STACK_SIZE (*C macro*), 489
 HPX_MEMORY_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
 HPX_MEMORY_WITH_TESTS:BOOL
 command line option, 87
 HPX_MINGW (*C macro*), 488
 HPX_MSVC (*C macro*), 488
 HPX_NAMING_BASE_WITH_COMPATIBILITY_HEADERBX~~B~~PLUGIN_FORCE_LOAD_NAME (*C macro*), 918
 command line option, 87
 HPX_NAMING_BASE_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 87
 HPX_NAMING_BASE_WITH_TESTS:BOOL
 command line option, 87
 HPX_NATIVE_MIC (*C macro*), 489
 HPX_NATIVE_TLS (*C macro*), 492
 HPX_NETWORK_RETRIES_SLEEP (*C macro*), 487
 HPX_NOINLINE (*C macro*), 491
 HPX_NON_COPYABLE (*C macro*), 490
 HPX_NORETURN (*C macro*), 491
 HPX_NUM_IO_POOL_SIZE (*C macro*), 487
 HPX_NUM_PARCEL_POOL_SIZE (*C macro*), 487
 HPX_NUM_TIMER_POOL_SIZE (*C macro*), 487
 HPX_ONCE_INIT (*C macro*), 601
 HPX_PARCEL_BOOTSTRAP (*C macro*), 486
 HPX_PARCEL_IPC_DATA_BUFFER_CACHE_SIZE
 (*C macro*), 486
 HPX_PARCEL_MAX_CONNECTIONS (*C macro*), 486
 HPX_PARCEL_MAX_CONNECTIONS_PER_LOCALITY
 (*C macro*), 486
 HPX_PARCEL_MAX_MESSAGE_SIZE (*C macro*), 486
 HPX_PARCEL_MAX_OUTBOUND_MESSAGE_SIZE (*C
 macro*), 486
 HPX_PARCEL_MPI_MAX_REQUESTS (*C macro*), 486
 HPX_PARCEL_SERIALIZATION_OVERHEAD
 (*C
 macro*), 486
 HPX_PATH_DELIMITERS (*C macro*), 487
 HPX_PLAIN_ACTION (*C macro*), 365
 HPX_PLAIN_ACTION_ID (*C macro*), 366
 HPX_PLUGIN_API (*C macro*), 917
 HPX_PLUGIN_ARGUMENT_LIMIT (*C macro*), 917
 HPX_PLUGIN_COMPONENT_PREFIX (*C macro*), 487
 HPX_PLUGIN_EXPORT (*C macro*), 919
 HPX_PLUGIN_EXPORT_ (*C macro*), 919
 HPX_PLUGIN_EXPORT_API (*C macro*), 917
 HPX_PLUGIN_EXPORT_DYNAMIC (*C macro*), 919
 HPX_PLUGIN_EXPORT_LIST (*C macro*), 919
 HPX_PLUGIN_EXPORT_LIST_ (*C macro*), 919
 HPX_PLUGIN_EXPORT_LIST_DYNAMIC (*C macro*),
 919
 HPX_PLUGIN_EXPORTER_INSTANCE_NAME
 (*C
 macro*), 918
 HPX_PLUGIN_EXPORTER_INSTANCE_NAME_ (*C
 macro*), 918
 HPX_PLUGIN_EXPORTER_INSTANCE_NAME_DYNAMIC
 (*C
 macro*), 918
 HPX_PLUGIN_EXPORTER_NAME (*C macro*), 918
 HPX_PLUGIN_EXPORTER_NAME_ (*C macro*), 918
 HPX_PLUGIN_EXPORTER_NAME_DYNAMIC (*C
 macro*), 918
 HPX_PLUGIN_FORCE_LOAD_NAME (*C macro*), 918
 HPX_PLUGIN_FORCE_LOAD_NAME_ (*C macro*), 918
 HPX_PLUGIN_FORCE_LOAD_NAME_DYNAMIC (*C
 macro*), 918
 HPX_PLUGIN_LIST_NAME (*C macro*), 918
 HPX_PLUGIN_LIST_NAME_ (*C macro*), 918
 HPX_PLUGIN_LIST_NAME_DYNAMIC (*C macro*), 918
 HPX_PLUGIN_NAME (*C macro*), 487
 HPX_PLUGIN_NAME_2 (*C macro*), 918
 HPX_PLUGIN_NAME_3 (*C macro*), 918
 HPX_PLUGIN_PLUGIN_PREFIX (*C macro*), 487
 HPX_PLUGIN_STRING (*C macro*), 487
 HPX_PLUGIN_SYMBOLS_PREFIX (*C macro*), 917
 HPX_PLUGIN_SYMBOLS_PREFIX_DYNAMIC (*C
 macro*), 917
 HPX_PLUGIN_SYMBOLS_PREFIX_DYNAMIC_STR
 (*C macro*), 917
 HPX_PLUGIN_SYMBOLS_PREFIX_STR (*C macro*),
 917
 HPX_PLUGIN_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 87
 HPX_PLUGIN_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 88
 HPX_PLUGIN_WITH_TESTS:BOOL
 command line option, 88
 HPX_PP_CAT (*C macro*), 398
 HPX_PP_EXPAND (*C macro*), 397
 HPX_PP_NARGS (*C macro*), 398
 HPX_PP_STRINGIZE (*C macro*), 398
 HPX_PP_STRIP_PARENS (*C macro*), 397
 HPX_PREPROCESSOR_WITH_COMPATIBILITY_HEADERS:BOOL
 command line option, 88
 HPX_PREPROCESSOR_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 88
 HPX_PREPROCESSOR_WITH_TESTS:BOOL
 command line option, 88
 HPX_PROGRAM_OPTIONS_IMPLICIT_VALUE_NEXT_TOKEN
 (*C macro*), 467
 HPX_PROGRAM_OPTIONS_VERSION (*C macro*), 467
 HPX_PROGRAM_OPTIONS_WITH_BOOST_PROGRAM_OPTIONS_COM
 command line option, 88
 HPX_PROGRAM_OPTIONS_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 88
 HPX_PROGRAM_OPTIONS_WITH_TESTS:BOOL
 command line option, 88
 HPX_PURE (*C macro*), 492

HPX_REGISTER_ACTION (<i>C macro</i>), 363	
HPX_REGISTER_ACTION_DECLARATION macro), 362	(C
HPX_REGISTER_ACTION_DECLARATION_ macro), 363	(C
HPX_REGISTER_ACTION_DECLARATION_1 macro), 363	(C
HPX_REGISTER_ACTION_ID (<i>C macro</i>), 363	
HPX_REGISTER_ALLREDUCE (<i>C macro</i>), 582	
HPX_REGISTER_ALLREDUCE_DECLARATION macro), 582	(C
HPX_REGISTER_ALTOALL (<i>C macro</i>), 589	
HPX_REGISTER_ALTOALL_DECLARATION macro), 589	(C
HPX_REGISTER_COMPONENT (<i>C macro</i>), 368	
HPX_REGISTER_GATHER (<i>C macro</i>), 579	
HPX_REGISTER_GATHER_DECLARATION macro), 579	(C
HPX_RESILIENCY_VERSION_DATE (<i>C macro</i>), 814	
HPX_RESILIENCY_VERSION_FULL (<i>C macro</i>), 814	
HPX_RESILIENCY_VERSION_MAJOR (<i>C macro</i>), 814	
HPX_RESILIENCY_VERSION_MINOR (<i>C macro</i>), 814	
HPX_RESILIENCY_VERSION_SUBMINOR macro), 814	(C
HPX_RESILIENCY_WITH_DEPRECATED_WARNINGS	HB BO STATISTICS_WITH_TESTS:BOOL
command line option, 88	command line option, 89
HPX_RESILIENCY_WITH_TESTS:BOOL	
command line option, 88	HPX_SUPER PURE (<i>C macro</i>), 492
HPX_RESOURCE_PARTITIONER_WITH_COMPATIBILITY_HEADERS	HB BO OL option, 89
command line option, 88	HPX_SYNCHRONIZATION_WITH_COMPATIBILITY_HEADERS:BOOL
HPX_RESOURCE_PARTITIONER_WITH_DEPRECATED_WARNINGS	HB BO LE option, 89
command line option, 88	HPX_SYNCHRONIZATION_WITH_TESTS:BOOL
HPX_RESOURCE_PARTITIONER_WITH_TESTS:BOOL	command line option, 89
command line option, 88	HPX_TEST (<i>C macro</i>), 447
HPX_RUNTIME_INSTANCE_LIMIT (<i>C macro</i>), 486	HPX_TEST_EQ (<i>C macro</i>), 447
HPX_SANITY (<i>C macro</i>), 447	HPX_TEST_EQ_MSG (<i>C macro</i>), 447
HPX_SANITY_EQ (<i>C macro</i>), 447	HPX_TEST_LT (<i>C macro</i>), 447
HPX_SANITY_EQ_MSG (<i>C macro</i>), 447	HPX_TEST_LTE (<i>C macro</i>), 447
HPX_SANITY_LT (<i>C macro</i>), 447	HPX_TEST_MSG (<i>C macro</i>), 447
HPX_SANITY_LTE (<i>C macro</i>), 447	HPX_TEST_NEQ (<i>C macro</i>), 447
HPX_SANITY_MSG (<i>C macro</i>), 447	HPX_TEST_NEQ_MSG (<i>C macro</i>), 447
HPX_SANITY_NEQ (<i>C macro</i>), 447	HPX_TEST_RANGE (<i>C macro</i>), 447
HPX_SANITY_RANGE (<i>C macro</i>), 447	HPX_TEST_THROW (<i>C macro</i>), 447
HPX_SCHEDULER_MAX_TERMINATED_THREADS:STR EN TESTING_WITH_COMPATIBILITY_HEADERS:BOOL	command line option, 89
command line option, 76	command line option, 89
HPX_SEGMENTED_ALGORITHMS_WITH_DEPRECATED_WARNINGS	NP WARNINGS GB BO TH_DEPRECATED_WARNINGS:BOOL
command line option, 88	command line option, 89
HPX_SEGMENTED_ALGORITHMS_WITH_TESTS:BOOL	HPX TESTING WITH_TESTS:BOOL
command line option, 88	command line option, 89
HPX_SERIALIZATION_SPLIT_FREE (<i>C macro</i>), 532	HPX_THREAD_QUEUE_MAX_ADD_NEW_COUNT (<i>C</i>
HPX_SERIALIZATION_SPLIT_FREE_TEMPLATE (<i>C macro</i>), 532	macro), 488
HPX_SERIALIZATION_SPLIT_MEMBER (<i>C macro</i>),	HPX_THREAD_QUEUE_MAX_DELETE_COUNT (<i>C</i>
532	macro), 488

HPX_WITH_DEPRECATED_WARNINGS:BOOL
 command line option, 73

HPX_WITH_DISABLED_SIGNAL_EXCEPTION_HANDLER:BOOL
 command line option, 73

HPX_WITH_DOCUMENTATION:BOOL
 command line option, 75

HPX_WITH_DOCUMENTATION_OUTPUT_FORMATS:STRING
 command line option, 75

HPX_WITH_DYNAMIC_HPX_MAIN:BOOL
 command line option, 73

HPX_WITH_EXAMPLES:BOOL
 command line option, 75

HPX_WITH_EXAMPLES_HDF5:BOOL
 command line option, 75

HPX_WITH_EXAMPLES_OPENMP:BOOL
 command line option, 75

HPX_WITH_EXAMPLES_QT4:BOOL
 command line option, 75

HPX_WITH_EXAMPLES_QTHREADS:BOOL
 command line option, 75

HPX_WITH_EXAMPLES_TBB:BOOL
 command line option, 75

HPX_WITH_EXECUTABLE_PREFIX:STRING
 command line option, 75

HPX_WITH_FAIL_COMPILE_TESTS:BOOL
 command line option, 75

HPX_WITH_FAULT_TOLERANCE:BOOL
 command line option, 73

HPX_WITH_FORTRAN:BOOL
 command line option, 73

HPX_WITH_FULL_RPATH:BOOL
 command line option, 73

HPX_WITH_GCC_VERSION_CHECK:BOOL
 command line option, 73

HPX_WITH_GENERIC_CONTEXT_COROUTINES:BOOL
 command line option, 73

HPX_WITH_GOOGLE_PERFTOOLS:BOOL
 command line option, 79

HPX_WITH_HIDDEN_VISIBILITY:BOOL
 command line option, 73

HPX_WITH_IO_COUNTERS:BOOL
 command line option, 75

HPX_WITH_IO_POOL:BOOL
 command line option, 77

HPX_WITH_ITTNOTIFY:BOOL
 command line option, 79

HPX_WITH_LOGGING:BOOL
 command line option, 73

HPX_WITH_MALLOC:STRING
 command line option, 74

HPX_WITH_MAX_CPU_COUNT:STRING
 command line option, 77

HPX_WITH_MAX_NUMA_DOMAIN_COUNT:STRING
 command line option, 77

HPX_WITH_MORE_THAN_64_THREADS:BOOL
 command line option, 77

HPX_WITH_NATIVE_TLS:BOOL
 command line option, 74

HPX_WITH_NETWORKING:BOOL
 command line option, 78

HPX_WITH_NICE_THREADLEVEL:BOOL
 command line option, 74

HPX_WITH_PAPI:BOOL
 command line option, 79

HPX_WITH_PARCEL_COALESCING:BOOL
 command line option, 74

HPX_WITH_PARCEL_PROFILING:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_ACTION_COUNTERS:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_LIBFABRIC:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_MPI:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_MPI_ENV:STRING
 command line option, 78

HPX_WITH_PARCELPORT_MPI_MULTITHREADED:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_TCP:BOOL
 command line option, 78

HPX_WITH_PARCELPORT_VERBS:BOOL
 command line option, 78

HPX_WITH_PSEUDO_DEPENDENCIES:BOOL
 command line option, 75

HPX_WITH_REGISTER_THREAD_COMPATIBILITY:BOOL
 command line option, 74

HPX_WITH_RUN_MAIN_EVERYWHERE:BOOL
 command line option, 74

HPX_WITH_SANITIZERS:BOOL
 command line option, 79

HPX_WITH_SCHEDULER_LOCAL_STORAGE:BOOL
 command line option, 77

HPX_WITH_SPINLOCK_DEADLOCK_DETECTION:BOOL
 command line option, 77

HPX_WITH_SPINLOCK_POOL_NUM:STRING
 command line option, 77

HPX_WITH_STACKOVERFLOW_DETECTION:BOOL
 command line option, 74

HPX_WITH_STACKTRACES:BOOL
 command line option, 77

HPX_WITH_STATIC_LINKING:BOOL
 command line option, 74

HPX_WITH_SWAP_CONTEXT_EMULATION:BOOL
 command line option, 77

HPX_WITH_TESTS:BOOL
 command line option, 75

HPX_WITH_TESTS_BENCHMARKS:BOOL
 command line option, 75

HPX_WITH_TESTS_DEBUG_LOG:BOOL	HPX_WITH_VERIFY_LOCKS_GLOBALLY:BOOL
command line option, 79	command line option, 80
HPX_WITH_TESTS_DEBUG_LOG_DESTINATION:STRING	HPX_WITH_VIM_YCM:BOOL
command line option, 79	command line option, 74
HPX_WITH_TESTS_EXAMPLES:BOOL	HPX_WITH_ZERO_COPY_SERIALIZATION_THRESHOLD:STRING
command line option, 75	command line option, 74
HPX_WITH_TESTS_EXTERNAL_BUILD:BOOL	HPX_WRAPPER_HEAP_STEP (<i>C macro</i>), 488
command line option, 75	
HPX_WITH_TESTS_HEADERS:BOOL	HWLOC_ROOT:PATH
command line option, 76	command line option, 90
HPX_WITH_TESTS_REGRESSIONS:BOOL	
command line option, 76	
HPX_WITH_TESTS_UNIT:BOOL	itt_counter_create (<i>C++ function</i>), 779
command line option, 76	itt_counter_create_typed (<i>C++ function</i>), 779
HPX_WITH_THREAD_BACKTRACE_DEPTH:STRING	itt_counter_destroy (<i>C++ function</i>), 779
command line option, 77	itt_counter_set_value (<i>C++ function</i>), 779
HPX_WITH_THREAD_BACKTRACE_ON_SUSPENSION:BOOL	itt_domain_create (<i>C++ function</i>), 779
command line option, 77	itt_event_create (<i>C++ function</i>), 779
HPX_WITH_THREAD_CREATION_AND_CLEANUP_RATE:FLOAT	itt_event_end (<i>C++ function</i>), 779
command line option, 77	itt_event_start (<i>C++ function</i>), 779
HPX_WITH_THREAD_CUMULATIVE_COUNTS:BOOL	itt_frame_begin (<i>C++ function</i>), 779
command line option, 77	itt_frame_end (<i>C++ function</i>), 779
HPX_WITH_THREAD_DEBUG_INFO:BOOL	itt_heap_allocate_begin (<i>C++ function</i>), 779
command line option, 79	itt_heap_allocate_end (<i>C++ function</i>), 779
HPX_WITH_THREAD_DESCRIPTION_FULL:BOOL	itt_heap_free_begin (<i>C++ function</i>), 779
command line option, 79	itt_heap_free_end (<i>C++ function</i>), 779
HPX_WITH_THREAD_GUARD_PAGE:BOOL	itt_heap_function_create (<i>C++ function</i>), 779
command line option, 79	itt_heap_internal_access_begin (<i>C++ func-</i>
HPX_WITH_THREAD_IDLE_RATES:BOOL	<i>tion</i>), 779
command line option, 77	itt_heap_internal_access_end (<i>C++ func-</i>
HPX_WITH_THREAD_LOCAL_STORAGE:BOOL	<i>tion</i>), 779
command line option, 77	itt_heap_reallocate_begin (<i>C++ function</i>),
HPX_WITH_THREAD_MANAGER_IDLE_BACKOFF:BOOL	779
command line option, 77	itt_heap_reallocate_end (<i>C++ function</i>), 779
HPX_WITH_THREAD_QUEUE_WAITTIME:BOOL	itt_id_create (<i>C++ function</i>), 779
command line option, 77	itt_id_destroy (<i>C++ function</i>), 779
HPX_WITH_THREAD_SCHEDULERS:STRING	itt_make_id (<i>C++ function</i>), 779
command line option, 77	itt_mark (<i>C++ function</i>), 779
HPX_WITH_THREAD_STACK_MMAP:BOOL	itt_mark_create (<i>C++ function</i>), 779
command line option, 77	itt_mark_off (<i>C++ function</i>), 779
HPX_WITH_THREAD_STEALING_COUNTS:BOOL	itt_metadata_add (<i>C++ function</i>), 779, 780
command line option, 77	itt_stack_create (<i>C++ function</i>), 779
HPX_WITH_THREAD_TARGET_ADDRESS:BOOL	itt_stack_destroy (<i>C++ function</i>), 779
command line option, 77	itt_stack_enter (<i>C++ function</i>), 779
HPX_WITH_TIMER_POOL:BOOL	itt_stack_leave (<i>C++ function</i>), 779
command line option, 78	itt_string_handle_create (<i>C++ function</i>), 779
HPX_WITH_TOOLS:BOOL	itt_sync_acquired (<i>C++ function</i>), 778
command line option, 76	itt_sync_cancel (<i>C++ function</i>), 778
HPX_WITH_VALGRIND:BOOL	itt_sync_create (<i>C++ function</i>), 778
command line option, 80	itt_sync_destroy (<i>C++ function</i>), 778
HPX_WITH_VERIFY_LOCKS:BOOL	itt_sync_prepare (<i>C++ function</i>), 778
command line option, 80	itt_sync_released (<i>C++ function</i>), 778
HPX_WITH_VERIFY_LOCKS_BACKTRACE:BOOL	itt_sync_releasing (<i>C++ function</i>), 778
command line option, 80	itt_sync_rename (<i>C++ function</i>), 778
	itt_task_begin (<i>C++ function</i>), 779

itt_task_end (*C++ function*), 779
itt_thread_ignore (*C++ function*), 779
itt_thread_set_name (*C++ function*), 779

L

LAGAS_ (*C macro*), 833
LAGAS_CONSOLE_ (*C macro*), 833
LAGAS_ENABLED (*C macro*), 833
LAPP_ (*C macro*), 833
LAPP_CONSOLE_ (*C macro*), 833
LAPP_ENABLED (*C macro*), 834
LAS_ (*C macro*), 833
LBT_ (*C macro*), 833
lbt_ (*C++ member*), 834
LCO, 18
LDEB_ (*C macro*), 833
LDEB_CONSOLE_ (*C macro*), 833
LDEB_ENABLED (*C macro*), 834
LERR_ (*C macro*), 833
LFATAL_ (*C macro*), 833
LHPX_ (*C macro*), 833
LHPX_CONSOLE_ (*C macro*), 833
LHPX_ENABLED (*C macro*), 833
Lightweight Control Object, 18
LLCO_ (*C macro*), 833
Local Control Object, 18
Locality, 17
LOSH_ (*C macro*), 833
LPICS_ (*C macro*), 833
LPROGRESS_ (*C macro*), 833
LPT_ (*C macro*), 833
LPT_CONSOLE_ (*C macro*), 833
LPT_ENABLED (*C macro*), 833
LRT_ (*C macro*), 833
LTIM_ (*C macro*), 833
LTIM_CONSOLE_ (*C macro*), 833
LTIM_ENABLED (*C macro*), 833
LTM_ (*C macro*), 833

M

make_error_code (*C++ function*), 820
moodycamel (*C++ type*), 789
moodycamel::AllocationMode (*C++ enum*), 791
moodycamel::CanAlloc (*C++ enumerator*), 791
moodycamel::CannotAlloc (*C++ enumerator*), 791
moodycamel::ConcurrentQueue (*C++ class*), 789
moodycamel::ConcurrentQueue::~ConcurrentQueue (*C++ function*), 790
moodycamel::ConcurrentQueue::add_block_t_{0_free_list} (*C++ function*), 792
moodycamel::ConcurrentQueue::add_blocks_t_{0_free_list} (*C++ function*), 792

moodycamel::ConcurrentQueue::add_producer (*C++ function*), 792
moodycamel::ConcurrentQueue::Block (*C++ class*), 793
moodycamel::ConcurrentQueue::BLOCK_SIZE (*C++ member*), 791
moodycamel::ConcurrentQueue::ConcurrentQueue (*C++ function*), 790
moodycamel::ConcurrentQueue::consumer_token_t (*C++ type*), 789
moodycamel::ConcurrentQueue::create (*C++ function*), 793
moodycamel::ConcurrentQueue::create_array (*C++ function*), 793
moodycamel::ConcurrentQueue::destroy (*C++ function*), 793
moodycamel::ConcurrentQueue::destroy_array (*C++ function*), 793
moodycamel::ConcurrentQueue::enqueue (*C++ function*), 790
moodycamel::ConcurrentQueue::enqueue_bulk (*C++ function*), 790
moodycamel::ConcurrentQueue::EXPLICIT_BLOCK_EMPTY_C (*C++ member*), 791
moodycamel::ConcurrentQueue::EXPLICIT_CONSUMER_C (*C++ member*), 791
moodycamel::ConcurrentQueue::EXPLICIT_INITIAL_INDE (*C++ member*), 791
moodycamel::ConcurrentQueue::ExplicitProducer (*C++ class*), 794
moodycamel::ConcurrentQueue::ExplicitProducer::Blo (*C++ class*), 795
moodycamel::ConcurrentQueue::ExplicitProducer::Blo (*C++ class*), 795
moodycamel::ConcurrentQueue::FreeList (*C++ class*), 795
moodycamel::ConcurrentQueue::freeList (*C++ member*), 792
moodycamel::ConcurrentQueue::FreeListNode (*C++ class*), 796
moodycamel::ConcurrentQueue::get_or_add_implicit_p (*C++ function*), 792
moodycamel::ConcurrentQueue::globalExplicitConsumer (*C++ member*), 792
moodycamel::ConcurrentQueue::IMPLICIT_INITIAL_INDE (*C++ member*), 791
moodycamel::ConcurrentQueue::ImplicitProducer (*C++ class*), 797
moodycamel::ConcurrentQueue::ImplicitProducer::Blo (*C++ class*), 798
moodycamel::ConcurrentQueue::ImplicitProducer::Blo (*C++ class*), 798
moodycamel::ConcurrentQueue::ImplicitProducerHash (*C++ class*), 798

moodycamel::ConcurrentQueue::implicitPromoddyCamel::ConcurrentQueue::swap (C++
function), 790
moodycamel::ConcurrentQueue::implicitPromoddyCamel::ConcurrentQueue::swap_implicit_producer
(C++ member), 792
moodycamel::ConcurrentQueue::implicitPromoddyCamel::ConcurrentQueue::swap_internal
(C++ function), 791
moodycamel::ConcurrentQueue::ImplicitPromoddyCamel::ConcurrentQueue::try_dequeue
(C++ class), 798
moodycamel::ConcurrentQueue::index_t moodycamel::ConcurrentQueue::try_dequeue_bulk
(C++ type), 789
moodycamel::ConcurrentQueue::INITIAL_IMPMODdyCamel::ConcurrentQueue::try_dequeue_bulk_from
(C++ member), 791
moodycamel::ConcurrentQueue::initialBlockBody moodycamel::ConcurrentQueue::try_dequeue_from_produ
(C++ member), 792
moodycamel::ConcurrentQueue::initialBlockBody moodycamel::ConcurrentQueue::try_dequeue_non_inter
(C++ member), 792
moodycamel::ConcurrentQueue::initialImplBody moodyCamel::ConcurrentQueue::try_enqueue
(C++ member), 792
moodycamel::ConcurrentQueue::initialImplBody moodyCamel::ConcurrentQueue::try_get_block_from_f
(C++ member), 792
moodycamel::ConcurrentQueue::inner_enqueue moodycamel::ConcurrentQueue::try_get_block_from_in
(C++ function), 791
moodycamel::ConcurrentQueue::inner_enqueue moodyCamel::ConcurrentQueue::update_current_produc
(C++ function), 791, 792
moodycamel::ConcurrentQueue::is_lock_free moodycamel::ConcurrentQueue<T,
(C++ function), 791 Traits>::Block::Block (C++ func-
moodycamel::ConcurrentQueue::nextExplicitConsumer, 793
(C++ member), 792 moodycamel::ConcurrentQueue<T,
moodycamel::ConcurrentQueue::operator= Traits>::Block::dummy (C++ mem-
(C++ function), 790 ber), 794
moodycamel::ConcurrentQueue::populate_immoddyCamel::ConcurrentQueue<T,
(C++ function), 792 Traits>::Block::dynamicallyAllocated
moodycamel::ConcurrentQueue::populate_initial_immoddyCamel::ConcurrentQueue<T,
(C++ member), 794 moodycamel::ConcurrentQueue<T,
moodycamel::ConcurrentQueue::producer_token_t Traits>::Block::elements (C++
(C++ type), 789 member), 794
moodycamel::ConcurrentQueue::ProducerBasmoodycamel::ConcurrentQueue<T,
(C++ class), 799 Traits>::Block::elementsCompletelyDequeued
moodycamel::ConcurrentQueue::producerCount (C++ member), 794
(C++ member), 792 moodycamel::ConcurrentQueue<T,
moodycamel::ConcurrentQueue::producerListTail Traits>::Block::freeListNext
(C++ member), 792 (C++ member), 794
moodycamel::ConcurrentQueue::recycle_or_moodycamel::ConcurrentQueue<T,
(C++ function), 792 Traits>::Block::freeListRefs
moodycamel::ConcurrentQueue::reown_producers (C++ member), 794
(C++ function), 792 moodycamel::ConcurrentQueue<T,
moodycamel::ConcurrentQueue::requisition_block Traits>::Block::is_empty (C++
(C++ function), 792 function), 793
moodycamel::ConcurrentQueue::size_approxmoodycamel::ConcurrentQueue<T,
(C++ function), 791 Traits>::Block::next (C++ mem-
moodycamel::ConcurrentQueue::size_t ber), 794
(C++ type), 789 moodycamel::ConcurrentQueue<T,
moodycamel::ConcurrentQueue::size_t

`Traits>::Block::operator[] (C++
function), 793`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::reset_empty
(C++ function), 793`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::set_all_empty
(C++ function), 793`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::set_empty (C++
function), 793`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::set_many_empty
(C++ function), 793`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::shouldBeOnFreeList
(C++ member), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::Block::[anonymous]
(C++ member), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::~ExplicitProducer
(C++ function), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::blockIndex
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexEntyyaibase::FreeList::add_knowingRefCount_is_z
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexEntyyaib$ockFreeList::FreeList
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexHeadeitentFreeList::freeListHead
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexHeadeit$onFreeList::head_unsafe
(C++ function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexHeadeit$prevFreeList::operator=
(C++ function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::BlockIndexHeadeitsizeFreeList::REFS_MASK
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::dequeue
(C++ function), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::dequeue_bulk
(C++ function), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::enqueue
(C++ function), 794`

`Traits>::ExplicitProducer::enqueue_bulk
(C++ function), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::ExplicitProducer
(C++ function), 794`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::new_block_index
(C++ function), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::pr_blockIndexEntr
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::pr_blockIndexFront
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::pr_blockIndexRaw
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::pr_blockIndexSize
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::ExplicitProducer::pr_blockIndexSlot
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::add (C++ func-
tion), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::add_knowingRefCount_is_z
(C++ function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::head_unsafe
(C++ function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::operator= (C++
function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::REFS_MASK
(C++ member), 795`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::SHOULD_BE_ON_FREELIST
(C++ member), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::swap (C++
function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::try_get (C++
function), 796`

`moodycamel::ConcurrentQueue<T,
Traits>::FreeList::try_set (C++
function), 796`

Traits>::FreeListNode::freeListNext (C++ member), 797	Traits>::ImplicitProducer::ImplicitProducer (C++ function), 797
moodycamel::ConcurrentQueue<T, Traits>::FreeListNode::FreeListNode (C++ function), 797	moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::insert_block_index (C++ function), 797
moodycamel::ConcurrentQueue<T, Traits>::FreeListNode::freeListRefs (C++ member), 797	moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::INVALID_BLOCK_BASE (C++ member), 798
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::~ImplicitProducer (C++ function), 797	moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::new_block_index (C++ function), 797
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::blockIndex (C++ member), 798	moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::nextBlockIndexCap (C++ member), 798
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexEntryKey (C++ member), 798	moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::rewind_block_index (C++ function), 797
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexEntryValueImplicitProducerHash::capacity (C++ member), 798	moodycamel::ConcurrentQueue<T, ImplicitProducerHash::capacity (C++ member), 798
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexHeadEntriesImplicitProducerHash::entries (C++ member), 798	moodycamel::ConcurrentQueue<T, ImplicitProducerHash::entries (C++ member), 798
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexHeadEntriesImplicitProducerHash::prev (C++ member), 798	moodycamel::ConcurrentQueue<T, ImplicitProducerHash::prev (C++ member), 798
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexHeadEntriesImplicitProducerKVP::ImplicitProducer (C++ member), 798	moodycamel::ConcurrentQueue<T, ImplicitProducerKVP::ImplicitProducer (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexHeadEntriesImplicitProducerKVP::key (C++ member), 798	moodycamel::ConcurrentQueue<T, ImplicitProducerKVP::key (C++ member), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::BlockIndexHeadEntriesImplicitProducerKVP::operator=	moodycamel::ConcurrentQueue<T, ImplicitProducerKVP::operator= (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::dequeue (C++ function), 797	moodycamel::ConcurrentQueue<T, ImplicitProducerKVP::swap (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::dequeue_bulk (C++ function), 797	moodycamel::ConcurrentQueue<T, ImplicitProducerKVP::value (C++ member), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::enqueue (C++ function), 797	moodycamel::ConcurrentQueue<T, ProducerBase::~ProducerBase (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::enqueue_bulk (C++ function), 797	moodycamel::ConcurrentQueue<T, ProducerBase::dequeue (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::get_block_indexExhaustiveIndex::ProducerIndexBase::dequeue_bulk (C++ function), 797	moodycamel::ConcurrentQueue<T, ProducerIndexBase::dequeue_bulk (C++ function), 799
moodycamel::ConcurrentQueue<T, Traits>::ImplicitProducer::get_block_indexExhaustiveIndex::ProducerIndexBase::dequeueOptimisticCount (C++ function), 797	moodycamel::ConcurrentQueue<T, ProducerIndexBase::dequeueOptimisticCount (C++ member), 800
moodycamel::ConcurrentQueue<T,	moodycamel::ConcurrentQueue<T,

```
Traits>::ProducerBase::dequeueOverrun  
    (C++ member), 800  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::getTail  
    (C++ function), 799  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::headIndex  
    (C++ member), 800  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::isExplicit  
    (C++ member), 800  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::next_prod  
    (C++ function), 799  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::parent  
    (C++ member), 800  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::ProducerBase  
    (C++ function), 799  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::size_approx  
    (C++ function), 799  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::tailBlock  
    (C++ member), 800  
moodycamel::ConcurrentQueue<T,  
    Traits>::ProducerBase::tailIndex  
    (C++ member), 800  
moodycamel::ConcurrentQueueDefaultTraits  
    (C++ class), 800  
moodycamel::ConcurrentQueueDefaultTraits  
    (C++ member), 800  
moodycamel::ConsumerToken (C++ class), 800  
moodycamel::ConsumerToken::ConsumerToken  
    (C++ function), 801  
moodycamel::ConsumerToken::currentProducer  
    (C++ member), 801  
moodycamel::ConsumerToken::desiredProducer  
    (C++ member), 801  
moodycamel::ConsumerToken::initialOffset  
    (C++ member), 801  
moodycamel::ConsumerToken::itemsConsumedFromCurrent  
    (C++ member), 801  
moodycamel::ConsumerToken::lastKnownGlobalOffset  
    (C++ member), 801  
moodycamel::ConsumerToken::operator=  
    (C++ function), 801  
moodycamel::ConsumerToken::swap (C++  
    function), 801  
moodycamel::details (C++ type), 802  
moodycamel::details::_hash_32_or_64  
    (C++ class), 803  
moodycamel::details::_hash_32_or_64::hash  
    (C++ function), 803  
moodycamel::details::_hash_32_or_64<1>  
    (C++ class), 803  
moodycamel::details::align_for (C++ func-  
    tion), 802  
moodycamel::details::ceil_to_pow_2 (C++  
    function), 802  
moodycamel::details::circular_less_than  
    (C++ function), 802  
moodycamel::details::ConcurrentQueueProducerType  
    (C++ class), 803  
moodycamel::details::ConcurrentQueueProducerType  
    (C++ function), 803  
moodycamel::details::EmptyCounterThreshold  
    (C++ member), 803  
moodycamel::details::Consumer:consumeNonQoSOrderBeforeReplace  
    (C++ member), 803  
moodycamel::details::FinalIndexProducer  
    (C++ class), 803  
moodycamel::details::const_numeric_max  
    (C++ class), 803  
moodycamel::details::IndexIndexExcept  
    (C++ function), 802  
moodycamel::details::hash_thread_id  
    (C++ function), 802  
moodycamel::details::ImbalanceFaitesProduquerHashHeadZed  
    (C++ member), 803  
moodycamel::details::invalid_thread_id2  
    (C++ member), 803  
moodycamel::details::max_align_t (C++  
    union), 803  
moodycamel::details::max_align_t::x  
    (C++ member), 804  
moodycamel::details::max_align_t::y
```

(*C++ member*), 804
moodycamel::details::max_align_t :: *z*
 (*C++ member*), 804
moodycamel::details::nomove (*C++ function*),
 802
moodycamel::details::nomove_if (*C++ class*), 804
moodycamel::details::nomove_if::eval
 (*C++ function*), 804
moodycamel::details::nomove_if<false>
 (*C++ class*), 804
moodycamel::details::static_is_lock_free
 (*C++ class*), 804
moodycamel::details::static_is_lock_free
 (*C++ class*), 805
moodycamel::details::std_max_align_t
 (*C++ type*), 802
moodycamel::details::swap_relaxed (*C++ function*), 802
moodycamel::details::thread_id (*C++ function*), 802
moodycamel::details::thread_id_converter
 (*C++ class*), 806
moodycamel::details::thread_id_converter
 (*C++ function*), 806
moodycamel::details::thread_id_converter
 (*C++ type*), 806
moodycamel::details::thread_id_converter
 (*C++ type*), 806
moodycamel::details::thread_id_t (*C++ type*), 802
moodycamel::details::value (*C++ enumerator*), 804–806
moodycamel::details::[anonymous] (*C++ enum*), 804–806
moodycamel::explicit_context (*C++ enumerator*), 791
moodycamel::implicit_context (*C++ enumerator*), 791
moodycamel::InnerQueueContext (*C++ enum*),
 791
moodycamel::ProducerToken (*C++ class*), 801
moodycamel::ProducerToken :: ~*ProducerToken*
 (*C++ function*), 801
moodycamel::ProducerToken :: operator=
 (*C++ function*), 801
moodycamel::ProducerToken :: producer
 (*C++ member*), 802
moodycamel::ProducerToken :: ProducerToken
 (*C++ function*), 801
moodycamel::ProducerToken :: swap (*C++ function*), 801
moodycamel::ProducerToken :: valid (*C++ function*), 801
moodycamel::swap (*C++ function*), 789, 793
MOODYCAMEL_CATCH (*C macro*), 789
MOODYCAMEL_DELETE_FUNCTION (*C macro*), 789
MOODYCAMEL_EXCEPTIONS_ENABLED (*C macro*),
 789
MOODYCAMEL_NOEXCEPT (*C macro*), 789
MOODYCAMEL_NOEXCEPT_ASSIGN (*C macro*), 789
MOODYCAMEL_NOEXCEPTCTOR (*C macro*), 789
MOODYCAMEL_RETHROW (*C macro*), 789
MOODYCAMEL_THREADLOCAL (*C macro*), 789
MOODYCAMEL_THROW (*C macro*), 789
MOODYCAMEL_TRY (*C macro*), 789
O
operator!= (*C++ function*), 480, 891, 925
operator== (*C++ function*), 480, 544, 876, 891, 925
operator> (*C++ function*), 891
operator>= (*C++ function*), 891
operator>> (*C++ function*), 924
operator< (*C++ function*), 509, 510, 512–514, 891
operator<= (*C++ function*), 891
operator<< (*C++ function*), 441, 463, 480, 834, 891,
 924
P
PathRootHash_t
 command line option, 90
PathRootHash_t :: *path_id_numeric_size_t*
 Process, 17
R
restore_checkpoint (*C++ function*), 925
run_guarded (*C++ function*), 564
S
SPHINX_ROOT:PATH
 command line option, 46
std (*C++ type*), 357, 389, 485
std::hash::operator() (*C++ function*), 265,
 385, 389, 482, 485
std::hash<:hpx::threads::thread_id>
 (*C++ class*), 265

```
std::hash<hpx::memory::intrusive_ptr<T>>
    (C++ class), 482, 485
std::hash<hpx::memory::intrusive_ptr<T>>::result_type
    (C++ type), 482, 485
std::hash<hpx::util::optional<T>> (C++ class), 385, 389
store (C++ function), 452, 453
```